

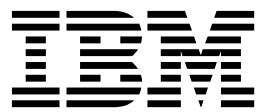
IBM Spectrum Scale
Version 4 Release 2.3

*Command and Programming
Reference*



IBM Spectrum Scale
Version 4 Release 2.3

*Command and Programming
Reference*



Note

Before using this information and the product it supports, read the information in "Notices" on page 1125.

This edition applies to version 4 release 2 modification 3 of the following products, and to all subsequent releases and modifications until otherwise indicated in new editions:

- IBM Spectrum Scale ordered through Passport Advantage® (product number 5725-Q01)
- IBM Spectrum Scale ordered through AAS/eConfig (product number 5641-GPF)
- IBM Spectrum Scale for Linux on Z (product number 5725-S28)
- IBM Spectrum Scale for IBM ESS (product number 5765-ESS)

Significant changes or additions to the text and illustrations are indicated by a vertical line (!) to the left of the change.

IBM welcomes your comments; see the topic "How to send your comments" on page xxi. When you send information to IBM, you grant IBM a nonexclusive right to use or distribute the information in any way it believes appropriate without incurring any obligation to you.

© Copyright IBM Corporation 2014, 2018.

US Government Users Restricted Rights – Use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

Contents

Tables	vii
About this information	ix
Prerequisite and related information	xx
Conventions used in this information.	xx
How to send your comments	xxi
Summary of changes.	xxiii
Chapter 1. Command reference.	1
gpfs.snap command	6
mmaddcallback command	10
mmadddisk command.	23
mmaddnode command	29
mmadquery command.	32
mmafmconfig command	37
mmafmctl command	40
mmafmlocal command	54
mmapplypolicy command	56
mmauth command	67
mmbbackup command	72
mmbbackupconfig command	81
mmblock command	83
mmbuildgpl command	88
mmcallehome command	90
mmces command	103
mmcesdr command	113
mmchattr command	122
mmchcluster command	128
mmchconfig command	132
mmchdisk command	161
mmcheckquota command	169
mmchfileset command	173
mmchfs command.	179
mmchlicense command	185
mmchmgr command	188
mmchnode command	190
mmchnodeclass command	195
mmchnsd command	198
mmchpolicy command	201
mmchpool command.	204
mmchqos command	206
mmclidecode command	213
mmclone command	215
mmcloudgateway command	218
mmcrcluster command	239
mmcrfileset command	244
mmcrfs command	250
mmcrnodeclass command	260
mmcrnsd command	262
mmcrsnapshot command	267
mmdefedquota command	272
mmdefquotaoff command	275
mmdefquotao command	278
mmdefragfs command	281
mmdelacl command	284
mmdelcallback command	286
mmddeldisk command	287
mmdefileset command	292
mmdeflfs command	295
mmdelnode command	297
mmdelnodeclass command	300
mmdelnsd command	302
mmdelesnapshot command	304
mmdf command	308
mmdiag command	312
mmdsh command	319
mmeditacl command	322
mmmedquota command	325
mmexportfs command	329
mmfsck command.	331
mmfsctl command.	343
mmgetacl command	347
mmgetstate command	350
mmhadoopctl command	353
mmhealth command	355
mmimgbackup command	367
mmimgrestore command	371
mmimportfs command	374
mmkeyserv command	378
mmlinkfileset command	390
mmlsattr command	392
mmllcallback command	395
mmllcluster command	397
mmllconfig command	400
mmlldisk command	402
mmllfileset command	406
mmllfs command	411
mmlllicense command	415
mmllsmgr command	417
mmllsmount command	419
mmllsnodeclass command	421
mmllnsd command	423
mmllspolicy command	426
mmllspool command	428
mmllsqos command	430
mmllsquota command	435
mmllssnapshot command	439
mmmigratefs command	442
mmmount command	444
mmnetverify command	446
mmnfs command	454
mmnsdiscover command	464
mmobj command	466
mmperfmon command	481
mmppmon command	492
mmprotocoltrace command.	497
mmpsnap command	501
mmputacl command	504
mmquotaoff command	507
mmquotaon command	509

mmremotecluster command	511
mmremotefs command	514
mmrepquota command	517
mmrestoreconfig command.	521
mmrestorefs command	525
mmrestripefile command	529
mmrestripefs command	532
mmrpldisk command.	539
mmsdrestore command.	546
mmsetquota command	548
mmshutdown command.	552
mmsmb command.	554
mmsnapdir command	565
mmstartup command.	569
mmtracectl command.	571
mmumount command	575
mmunlinkfileset command	578
mmuserauth command	581
mmwinstervctl command	602
spectrumscale command.	604

Chapter 2. IBM Spectrum Scale Data Management API for GPFS information 621

Overview of IBM Spectrum Scale Data Management API for GPFS	621
Management API for GPFS.	621
GPFS-specific DMAPI events	621
DMAPI functions	622
DMAPI configuration attributes	626
DMAPI restrictions for GPFS	627
Concepts of IBM Spectrum Scale Data Management API for GPFS	628
Sessions	628
Data management events	628
Mount and unmount	630
Tokens and access rights	631
Parallelism in Data Management applications	632
Data Management attributes	633
Support for NFS	633
Quota	633
Memory mapped files	633
Administration of IBM Spectrum Scale Data Management API for GPFS.	634
Required files for implementation of Data Management applications	634
GPFS configuration attributes for DAPI	635
Enabling DAPI for a file system	636
Initializing the Data Management application	637
Specifications of enhancements for IBM Spectrum Scale Data Management API for GPFS	637
Enhancements to data structures	638
Usage restrictions on DAPI functions.	639
Definitions for GPFS-specific DAPI functions	641
Semantic changes to DAPI functions	654
GPFS-specific DAPI events	655
Additional error codes returned by DAPI functions	656
Failure and recovery of IBM Spectrum Scale Data Management API for GPFS.	657
Single-node failure	658

Session failure and recovery	659
Event recovery	659
Loss of access rights	660
DODeferred deletions	660
DM application failure	660

Chapter 3. GPFS programming

interfaces 663	
gpfs_acl_t structure	666
gpfs_clone_copy() subroutine	667
gpfs_clone_snap() subroutine	669
gpfs_clone_split() subroutine	671
gpfs_clone_unsnap() subroutine	673
gpfs_close_inodescan() subroutine	675
gpfs_cmp_fssnapid() subroutine	676
gpfs_declone() subroutine	678
gpfs_direntx_t structure	680
gpfs direntx64_t structure	682
gpfs_fcntl() subroutine	684
gpfs_fgetattrs() subroutine	687
gpfs_fputatrrs() subroutine	689
gpfs_fputatrrswithpathname() subroutine	691
gpfs_free_fssnaphandle() subroutine	693
gpfs_fssnap_handle_t structure	694
gpfs_fssnap_id_t structure	695
gpfs_fstat() subroutine	696
gpfs_fstat_x() subroutine	698
gpfs_get_fnname_from_fssnaphandle() subroutine	700
gpfs_get_fssnaphandle_by_fssnapid() subroutine	701
gpfs_get_fssnaphandle_by_name() subroutine	703
gpfs_get_fssnaphandle_by_path() subroutine	705
gpfs_get_fssnapid_from_fssnaphandle() subroutine	707
gpfs_get.pathname_from_fssnaphandle() subroutine	709
gpfs_get_snapdirname() subroutine	711
gpfs_get_snapname_from_fssnaphandle() subroutine	713
gpfs_getacl() subroutine	715
gpfs_iattr_t structure	717
gpfs_iattr64_t structure	720
gpfs_iclose() subroutine	724
gpfs_ifile_t structure	726
gpfs_igetattrs() subroutine	727
gpfs_igetattrsx() subroutine	729
gpfs_igetfilesetname() subroutine	731
gpfs_igetstoragepool() subroutine.	733
gpfs_ipopen() subroutine	735
gpfs_ipopen64() subroutine	737
gpfs_iputatrrs() subroutine	739
gpfs_iread() subroutine	742
gpfs_ireaddir() subroutine	744
gpfs_ireaddir64() subroutine	746
gpfs_ireadlink() subroutine	748
gpfs_ireadlink64() subroutine	750
gpfs_ireadx() subroutine	752
gpfs_iscan_t structure	755
gpfs_lib_init() subroutine	756
gpfs_lib_term() subroutine	757
gpfs_next_inode() subroutine	758
gpfs_next_inode64() subroutine	760
gpfs_next_inode_with_xatrrs() subroutine	762

gpfs_next_inode_with_xattr64() subroutine	764	Snapshots/{snapshotName}: DELETE	907
gpfs_next_xattr() subroutine	766	Version 2	908
gpfs_opaque_acl_t structure	768	CES/addresses: GET	909
gpfs_open_inodescan() subroutine	769	CES/addresses/{cesAddress}: GET	913
gpfs_open_inodescan64() subroutine.	772	CES/services: GET	916
gpfs_open_inodescan_with_xattr() subroutine	775	CES/services/{service}: GET	919
gpfs_open_inodescan_with_xattr64() subroutine	778	Cluster: GET	922
gpfs_prealloc() subroutine	781	Config: GET	926
gpfs_putacl() subroutine.	784	Filesystems: GET	931
gpfs_quotactl() subroutine	786	Filesystems/{filesystemName}: GET	939
gpfs_quotaInfo_t structure	789	Filesystems/{filesystemName}/acl/{path}: GET	946
gpfs_seek_inode() subroutine	791	Filesystems/{filesystemName}/acl/{path}: PUT	949
gpfs_seek_inode64() subroutine	793	Filesystems/{filesystemName}/disks: GET.	954
gpfs_stat() subroutine	795	Filesystems/{filesystemName}/disks/	
gpfs_stat_inode() subroutine	797	{diskName}: GET	959
gpfs_stat_inode64() subroutine	799	Filesystems/{filesystemName}/filesets: GET	963
gpfs_stat_inode_with_xattrs() subroutine	801	Filesystems/{filesystemName}/filesets: POST	970
gpfs_stat_inode_with_xattr64() subroutine	803	Filesystems/{filesystemName}/filesets/	
gpfs_stat_x() subroutine	805	{filesetName}: DELETE	975
gpfsFcctlHeader_t structure	807	Filesystems/{filesystemName}/filesets/	
gpfsGetDataBlkDiskIdx_t structure	808	{filesetName}: GET	978
gpfsGetFileName_t structure	811	Filesystems/{filesystemName}/filesets/	
gpfsGetReplication_t structure.	812	{filesetName}: PUT	984
gpfsGetSetXAttr_t structure	814	Filesystems/{filesystemName}/filesets/	
gpfsGetSnapshotName_t structure	816	{filesetName}/link: DELETE	990
gpfsGetStoragePool_t structure	817	Filesystems/{filesystemName}/filesets/	
gpfsListXAttr_t structure	818	{filesetName}/link: POST	993
gpfsRestripeData_t structure	819	Filesystems/{filesystemName}/filesets/	
gpfsSetReplication_t structure	821	{filesetName}/quotas: GET	996
gpfsSetStoragePool_t structure.	823	Filesystems/{filesystemName}/filesets/	
Chapter 4. GPFS user exits	825	{filesetName}/quotas: POST	1000
mmsdrbackup user exit	826	Filesystems/{filesystemName}/filesets/	
nsddevices user exit	827	{filesetName}/snapshots: GET	1004
syncfsconfig user exit.	828	Filesystems/{filesystemName}/filesets/	
Chapter 5. IBM Spectrum Scale		{filesetName}/snapshots: POST	1008
management API commands	829	Filesystems/{filesystemName}/filesets/	
Version 1 - deprecated	829	{filesetName}/snapshots/{snapshotName}:	
CES addresses: GET	830	GET	1014
CES addresses/{cesAddress}: GET	832	Filesystems/{filesystemName}/owner/{path}:	
CES services: GET.	834	GET	1017
CES services/{service}: GET	837	Filesystems/{filesystemName}/owner/{path}:	
Config: GET.	840	PUT	1020
Cluster: GET	841	Filesystems/{filesystemName}/quotas: GET	1023
Filesets: GET	847	Filesystems/{filesystemName}/quotas: POST	1027
Filesets/{filename}: GET	853	Filesystems/{filesystemName}/snapshots: GET	1031
Filesets: POST	859	Filesystems/{filesystemName}/snapshots:	
Filesets/{filename}: PUT	864	POST.	1034
Filesets/{filename}: DELETE	868	Filesystems/{filesystemName}/snapshots/	
Filesystems: GET	870	{snapshotName}: DELETE.	1037
Filesystems/{filesystemName}: GET	876	Filesystems/{filesystemName}/snapshots/	
Info: GET.	882	{snapshotName}: GET	1040
Nodes: GET	885	Info: GET	1043
Nodes/{name}: GET	889	Jobs: GET	1045
Quotas: GET	892	Jobs/{jobID}: GET	1049
Quotas: POST	896	NFS/exports: GET	1053
S snapshots: GET.	899	NFS/exports: POST	1057
S snapshots/{snapshotName}: GET.	902	NFS/exports/{exportPath}: GET.	1060
S snapshots: POST	905	NFS/exports/{exportPath}: PUT.	1064

NFS/exports/{exportPath}: DELETE	1068	IBM and accessibility	1123
Nodes: GET	1071	Notices	1125
Nodes/{name}: GET	1078	Trademarks	1126
Nodes/{name}/health/events: GET	1083	Terms and conditions for product documentation	1127
Nodes/{name}/health/states: GET	1087	IBM Online Privacy Statement	1127
NSDs: GET	1091		
NSDs/{nsdName}: GET	1097		
SMB/shares: GET	1100	Glossary	1129
SMB/shares/{shareName}: GET	1105		
SMB/shares: POST	1110	Index	1135
SMB/shares/{shareName}: PUT	1115		
SMB/shares/{shareName}: DELETE	1120		

Accessibility features for IBM

Spectrum Scale	1123
Accessibility features	1123
Keyboard navigation	1123

Tables

1.	IBM Spectrum Scale library information units	x	34.	List of parameters	959
2.	Conventions	xxi	35.	List of parameters	963
3.	List of changes in documentation	xxx	36.	List of request parameters	970
4.	GPFS commands	1	37.	List of request parameters	975
5.	Global events and supported parameters	15	38.	List of parameters	978
6.	Local events and supported parameters	16	39.	List of request parameters	984
7.	Query details by type	33	40.	List of request parameters	990
8.	key-value	94	41.	List of request parameters	993
9.	key-value	95	42.	List of request parameters	996
10.	Settings for debugDataControl	141	43.	List of request parameters	1004
11.	Allocation of IOPS	208	44.	List of request parameters	1008
12.	GPFS commands that support QoS	209	45.	List of request parameters	1011
13.	Contents of columns <code>input1</code> and <code>input2</code> depending on the value in column <code>Buf type</code>	313	46.	List of request parameters	1014
14.	<code>mmkeyserv server show</code>	382	47.	List of request parameters	1017
15.	<code>mmkeyserv tenant show</code>	383	48.	List of request parameters	1020
16.	Information and error messages	446	49.	List of request parameters	1023
17.	Shortcut terms for network checks	449	50.	List of request parameters	1027
18.	Network checks	449	51.	List of request parameters	1031
19.	Restoring a global snapshot	526	52.	List of request parameters	1034
20.	Restoring a fileset snapshot	526	53.	List of request parameters	1037
21.	DMAPI configuration attributes	627	54.	List of request parameters	1040
22.	Specific DMAPI functions and associated error codes	657	55.	List of request parameters	1045
23.	GPFS programming interfaces	663	56.	List of request parameters	1049
24.	GPFS user exits	825	57.	List of request parameters	1053
25.	List of parameters	909	58.	List of request parameters	1057
26.	List of parameters	913	59.	List of request parameters	1060
27.	List of parameters	919	60.	List of request parameters	1064
28.	List of parameters	926	61.	List of request parameters	1068
29.	List of parameters	931	62.	List of request parameters	1078
30.	List of parameters	939	63.	List of request parameters	1083
31.	List of parameters	946	64.	List of request parameters	1087
32.	List of parameters	949	65.	List of request parameters	1091
33.	List of parameters	954	66.	List of request parameters	1097
			67.	List of request parameters	1100
			68.	List of request parameters	1105

About this information

This edition applies to IBM Spectrum Scale™ version 4.2.3 for AIX®, Linux, and Windows.

IBM Spectrum Scale is a file management infrastructure, based on IBM® General Parallel File System (GPFS™) technology, which provides unmatched performance and reliability with scalable access to critical file data.

To find out which version of IBM Spectrum Scale is running on a particular AIX node, enter:

```
lslpp -l gpfs\*
```

To find out which version of IBM Spectrum Scale is running on a particular Linux node, enter:

```
rpm -qa | grep gpfs      (for SLES and Red Hat Enterprise Linux)  
dpkg -l | grep gpfs     (for Ubuntu Linux)
```

To find out which version of IBM Spectrum Scale is running on a particular Windows node, open **Programs and Features** in the control panel. The IBM Spectrum Scale installed program name includes the version number.

Which IBM Spectrum Scale information unit provides the information you need?

The IBM Spectrum Scale library consists of the information units listed in Table 1 on page x.

To use these information units effectively, you must be familiar with IBM Spectrum Scale and the AIX, Linux, or Windows operating system, or all of them, depending on which operating systems are in use at your installation. Where necessary, these information units provide some background information relating to AIX, Linux, or Windows. However, more commonly they refer to the appropriate operating system documentation.

Note: Throughout this documentation, the term “Linux” refers to all supported distributions of Linux, unless otherwise specified.

Table 1. IBM Spectrum Scale library information units

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Concepts, Planning, and Installation Guide</i>	<p>This guide provides the following information:</p> <p>Product overview</p> <ul style="list-style-type: none"> • Overview of IBM Spectrum Scale • GPFS architecture • Protocols support overview: Integration of protocol access methods with GPFS • Active File Management • AFM-based Asynchronous Disaster Recovery (AFM DR) • Data protection and disaster recovery in IBM Spectrum Scale • Introduction to IBM Spectrum Scale GUI • IBM Spectrum Scale management API • Introduction to Cloud services • IBM Spectrum Scale in an OpenStack cloud deployment • IBM Spectrum Scale product editions • IBM Spectrum Scale license designation • Capacity based licensing • IBM Spectrum Storage™ Suite <p>Planning</p> <ul style="list-style-type: none"> • Planning for GPFS • Planning for protocols • Considerations for GPFS applications • Firewall recommendations • Planning for cloud services 	System administrators, analysts, installers, planners, and programmers of IBM Spectrum Scale clusters who are very experienced with the operating systems on which each IBM Spectrum Scale cluster is based

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Concepts, Planning, and Installation Guide</i>	<p>Installing and upgrading</p> <ul style="list-style-type: none"> • Steps for establishing and starting your IBM Spectrum Scale cluster • Installing IBM Spectrum Scale on Linux nodes and deploying protocols • Installing IBM Spectrum Scale on AIX nodes • Installing IBM Spectrum Scale on Windows nodes • Installing cloud services on IBM Spectrum Scale nodes • Installing and configuring IBM Spectrum Scale management API • Installing Active File Management • Installing and upgrading AFM-based Disaster Recovery • Installing call home • Migration, coexistence and compatibility • Steps to permanently uninstall GPFS and/or Protocols 	System administrators, analysts, installers, planners, and programmers of IBM Spectrum Scale clusters who are very experienced with the operating systems on which each IBM Spectrum Scale cluster is based

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Administration Guide</i>	<p>This guide provides the following information:</p> <p>Configuring</p> <ul style="list-style-type: none"> • Configuring the GPFS cluster • Configuring the CES and protocol configuration • Configuring and tuning your system for GPFS • Parameters for performance tuning and optimization • Configuring and tuning your system for Cloud services • Configuring Active File Management • Configuring AFM-based DR • Tuning for Kernel NFS backend on AFM and AFM DR <p>Administering</p> <ul style="list-style-type: none"> • Performing GPFS administration tasks • Verifying network operation with the mmnetverify command • Managing file systems • File system format changes between versions of IBM Spectrum Scale • Managing disks • Managing protocol services • Managing protocol user authentication • Managing protocol data exports • Managing object storage • Managing GPFS quotas • Managing GUI users • Managing GPFS access control lists • Considerations for GPFS applications • Accessing a remote GPFS file system 	System administrators or programmers of IBM Spectrum Scale systems

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Administration Guide</i>	<ul style="list-style-type: none"> • Information lifecycle management for IBM Spectrum Scale • Creating and maintaining snapshots of file systems • Creating and managing file clones • Scale Out Backup and Restore (SOBAR) • Data Mirroring and Replication • Implementing a clustered NFS environment on Linux • Implementing Cluster Export Services • Identity management on Windows • Protocols cluster disaster recovery • File Placement Optimizer • Encryption • Managing certificates to secure communications between GUI web server and web browsers • Securing protocol data • Cloud services: Transparent cloud tiering and Cloud data sharing • Highly-available write cache (HAWC) • Local read-only cache • Miscellaneous advanced administration • GUI limitations 	System administrators or programmers of IBM Spectrum Scale systems

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Problem Determination Guide</i>	<p>This guide provides the following information:</p> <p>Monitoring</p> <ul style="list-style-type: none"> • Performance monitoring • Monitoring system health through the IBM Spectrum Scale GUI • Monitoring system health by using the mmhealth command • Monitoring events through callbacks • Monitoring capacity through GUI • Monitoring AFM and AFM DR • GPFS SNMP support • Monitoring the IBM Spectrum Scale system by using call home • Monitoring the health of cloud services <p>Troubleshooting</p> <ul style="list-style-type: none"> • Best practices for troubleshooting • Understanding the system limitations • Collecting details of the issues • Managing deadlocks • Installation and configuration issues • Upgrade issues • Network issues • File system issues • Disk issues • Security issues • Protocol issues • Disaster recovery issues • Performance issues • GUI issues • AFM issues • AFM DR issues • Transparent cloud tiering issues • Recovery procedures • Support for troubleshooting • References 	<p>System administrators of GPFS systems who are experienced with the subsystems used to manage disks and who are familiar with the concepts presented in the <i>IBM Spectrum Scale: Concepts, Planning, and Installation Guide</i></p>

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Command and Programming Reference</i>	<p>This guide provides the following information:</p> <p>Command reference</p> <ul style="list-style-type: none"> • gpfs.snap command • mmaddcallback command • mmadddisk command • mmaddnode command • mmadquery command • mmafmconfig command • mmafmctl command • mmafmlocal command • mmapplypolicy command • mmauth command • mmbackup command • mmbackupconfig command • mmblock command • mmbuildgpl command • mmcallhome command • mmces command • mmcesdr command • mmchattr command • mmchcluster command • mmchconfig command • mmchdisk command • mmcheckquota command • mmchfileset command • mmchfs command • mmchlicense command • mmchmgr command • mmchnode command • mmchnodeclass command • mmchnsd command • mmchpolicy command • mmchpool command • mmchqos command • mmclidecode command • mmclone command • mmcloudgateway command • mmcrcluster command • mmcrfileset command • mmcrfs command • mmcrnodeclass command • mmcrnsd command • mmcrsnapshot command 	<ul style="list-style-type: none"> • System administrators of IBM Spectrum Scale systems • Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Command and Programming Reference</i>	<ul style="list-style-type: none"> • mmdefquota command • mmdefquotoff command • mmdefquoton command • mmdefragfs command • mmdelacl command • mmdelcallback command • mmdeldisk command • mmdelfileset command • mmdelfs command • mmdelnodename command • mmdelnodeclass command • mmdelnsd command • mmdelsnapshot command • mmdf command • mmdiag command • mmdsh command • mmeditacl command • mmedquota command • mmexportfs command • mmfsck command • mmfsctl command • mmgetacl command • mmgetstate command • mmhadoopctl command • mmhealth command • mmimgbackup command • mmimgrestore command • mmimportfs command • mmkeyserv command • mmlinkfilesset command • mmlsattr command • mmlscallback command • mmlscluster command • mmlsconfig command • mmlsdisk command 	<ul style="list-style-type: none"> • System administrators of IBM Spectrum Scale systems • Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Command and Programming Reference</i>	<ul style="list-style-type: none"> • mmlsfileset command • mmlsfs command • mmlslicense command • mmlsmgr command • mmlsmount command • mmlsnodeclass command • mmlsnsd command • mmlspolicy command • mmlspool command • mmlsqos command • mmlsquota command • mmlsnapshot command • mmmigratefs command • mmmount command • mmnetverify command • mmnfs command • mmnsdiscover command • mmobj command • mmperfmon command • mmpmon command • mmprotocoltrace command • mmptsnap command • mmputacl command • mmquotaoff command • mmquotaon command • mmremotecluster command • mmremotefs command • mmrepquota command • mmrestoreconfig command • mmrestoref command • mmrestripefile command • mmrestripefs command • mmrpldisk command • mmsdrrestore command • mmsetquota command • mmshutdown command • mmsmb command • mmsnapdir command • mmstartup command • mmtracectl command • mmumount command • mmunlinkfileset command • mmuserauth command • mmwinservctl command • spectrumscale command 	<ul style="list-style-type: none"> • System administrators of IBM Spectrum Scale systems • Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Command and Programming Reference</i>	Programming reference <ul style="list-style-type: none">• IBM Spectrum Scale Data Management API for GPFS information• GPFS programming interfaces• GPFS user exits• IBM Spectrum Scale management API commands	<ul style="list-style-type: none">• System administrators of IBM Spectrum Scale systems• Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Big Data and Analytics Guide</i>	<p>This guide provides the following information:</p> <p>IBM Spectrum Scale support for Hadoop</p> <ul style="list-style-type: none"> • HDFS transparency • Supported IBM Spectrum Scale storage modes • Hadoop cluster planning • Installation and configuration of HDFS transparency • Application interaction with HDFS transparency • Upgrading the HDFS Transparency cluster • Rolling upgrade for HDFS Transparency • Security • Advanced features • Hadoop distribution support • Limitations and differences from native HDFS • Problem determination <p>BigInsights® 4.2.5 and Hortonworks Data Platform 2.6</p> <ul style="list-style-type: none"> • Planning <ul style="list-style-type: none"> – Hardware requirements – Preparing the environment – Preparing a stanza file • Installation <ul style="list-style-type: none"> – Set up – Installation of software stack – BigInsights value-add services on IBM Spectrum Scale • Upgrading software stack <ul style="list-style-type: none"> – Migrating from BI IOP to HDP – Upgrading IBM Spectrum Scale service MPack – Upgrading HDFS Transparency – Upgrading IBM Spectrum Scale file system – Upgrading to BI IOP 4.2.5 	<ul style="list-style-type: none"> • System administrators of IBM Spectrum Scale systems • Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Table 1. IBM Spectrum Scale library information units (continued)

Information unit	Type of information	Intended users
<i>IBM Spectrum Scale: Big Data and Analytics Guide</i>	<ul style="list-style-type: none"> • Configuration <ul style="list-style-type: none"> – Setting up High Availability [HA] – IBM Spectrum Scale configuration parameter checklist – Dual-network deployment – Manually starting services in Ambari – Setting up local repository – Configuring LogSearch – Setting IBM Spectrum Scale configuration for BigSQL • Administration <ul style="list-style-type: none"> – IBM Spectrum Scale-FPO deployment – Ranger – Kerberos – Short-circuit read (SSR) – Disabling short circuit write – IBM Spectrum Scale service management – Ambari node management – Restricting root access – IBM Spectrum Scale management GUI – IBM Spectrum Scale versus Native HDFS • Troubleshooting <ul style="list-style-type: none"> – Snap data collection • Limitations <ul style="list-style-type: none"> – Limitations and information • FAQ <ul style="list-style-type: none"> – General – Service fails to start – Service check failures 	<ul style="list-style-type: none"> • System administrators of IBM Spectrum Scale systems • Application programmers who are experienced with IBM Spectrum Scale systems and familiar with the terminology and concepts in the XDSM standard

Prerequisite and related information

For updates to this information, see IBM Spectrum Scale in IBM Knowledge Center(www.ibm.com/support/knowledgecenter/STXKQY/ibmspectrumscale_welcome.html).

For the latest support information, see the IBM Spectrum Scale FAQ in IBM Knowledge Center(www.ibm.com/support/knowledgecenter/STXKQY/gpfsclustersfaq.html).

Conventions used in this information

Table 2 on page xxi describes the typographic conventions used in this information. UNIX file name conventions are used throughout this information.

Note: Users of IBM Spectrum Scale for Windows must be aware that on Windows, UNIX-style file names need to be converted appropriately. For example, the GPFS cluster configuration data is stored in the /var/mmfs/gen/mmsdrfs file. On Windows, the UNIX namespace starts under the %SystemDrive%\cygwin64 directory, so the GPFS cluster configuration data is stored in the C:\cygwin64\var\mmfs\gen\mmsdrfs file.

Table 2. Conventions

Convention	Usage
bold	Bold words or characters represent system elements that you must use literally, such as commands, flags, values, and selected menu options. Depending on the context, bold typeface sometimes represents path names, directories, or file names.
<u>bold underlined</u>	<u>bold underlined</u> keywords are defaults. These take effect if you do not specify a different keyword.
constant width	Examples and information that the system displays appear in <code>constant-width</code> typeface. Depending on the context, <code>constant-width</code> typeface sometimes represents path names, directories, or file names.
<i>italic</i>	<i>Italic</i> words or characters represent variable values that you must supply. <i>Italics</i> are also used for information unit titles, for the first use of a glossary term, and for general emphasis in text.
<key>	Angle brackets (less-than and greater-than) enclose the name of a key on the keyboard. For example, <Enter> refers to the key on your terminal or workstation that is labeled with the word <i>Enter</i> .
\	In command examples, a backslash indicates that the command or coding example continues on the next line. For example: <code>mkcondition -r IBM.FileSystem -e "PercentTotUsed > 90" \ -E "PercentTotUsed < 85" -m p "FileSystem space used"</code>
{item}	Braces enclose a list from which you must choose an item in format and syntax descriptions.
[item]	Brackets enclose optional items in format and syntax descriptions.
<Ctrl-x>	The notation <Ctrl-x> indicates a control character sequence. For example, <Ctrl-c> means that you hold down the control key while pressing <c>.
item...	Ellipses indicate that you can repeat the preceding item one or more times.
	In <i>synopsis</i> statements, vertical lines separate a list of choices. In other words, a vertical line means <i>Or</i> . In the left margin of the document, vertical lines indicate technical changes to the information.

Note: CLI options that accept a list of option values delimit with a comma and no space between values. As an example, to display the state on three nodes use **mmgetstate -N NodeA,NodeB,NodeC**. Exceptions to this syntax are listed specifically within the command.

How to send your comments

Your feedback is important in helping us to produce accurate, high-quality information. If you have any comments about this information or any other IBM Spectrum Scale documentation, send your comments to the following e-mail address:

mhvrcfs@us.ibm.com

Include the publication title and order number, and, if applicable, the specific location of the information about which you have comments (for example, a page number or a table number).

To contact the IBM Spectrum Scale development organization, send your comments to the following e-mail address:

gpfs@us.ibm.com

Summary of changes

This topic summarizes changes to the IBM Spectrum Scale licensed program and the IBM Spectrum Scale library. Within each information unit in the library, a vertical line (|) to the left of text and illustrations indicates technical changes or additions that are made to the previous edition of the information.

| **Summary of changes**
| **for IBM Spectrum Scale version 4 release 2.3**
| **as updated, June 2018**

| This release of the IBM Spectrum Scale licensed program and the IBM Spectrum Scale library includes the following improvements:

| **Added support for OpenStack Mitaka packages**

| Support for OpenStack Mitaka packages has been added for the object protocol. For more information, see *Protocols support overview: Integration of protocol access methods with GPF in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

| **Authentication considerations changes**

| The following changes are done:

- | • Authentication support matrix has been divided to separate out the File and object protocols and accordingly, the corresponding explanation is modified.
- | • The matrix is further divided based on the authentication service that is used.
- | • A diagram is added to explain the high-level flow of authentication for File protocols.
- | • "Authentication for file access" topic is renamed to "Authentication and ID mapping for file access".

| For more information, see the *Authentication considerations* topic in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

| **Big data and analytics changes**

| For information on changes in IBM Spectrum Scale Big Data and Analytics support, see *Big Data and Analytics - summary of changes*.

| **Directory preallocation**

| In environments in which many files are added to and removed from a directory in a short time, you can improve performance by setting the minimum compaction size of the directory. The minimum compaction size is the number of directory slots, including both full and empty slots, that a directory is allowed to retain when it is compacted. For more information, see *gpfs_prealloc() subroutine, mmchattr command, and mmclsattr command in IBM Spectrum Scale: Command and Programming Reference*.

| **Express Edition no longer available**

| IBM Spectrum Scale Express Edition is no longer available. For information on migrating from IBM Spectrum Scale Express Edition 4.2.2.x or earlier to IBM Spectrum Scale Standard Edition 4.2.3.x, see *Migrating from Express Edition to Standard Edition* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

| **FPO enhancements**

| FPO performs the following functions:

- | • Provides QoS support for auto recovery
- | • Supports locality-aware data copy
- | • Uses the **mmrestripefile** command to check whether the replicas of data blocks are matched for one file

Installation toolkit support for `gpfs.adv` and `gpfs.crypto` packages

The installation toolkit now supports installation, deployment, and upgrade of `gpfs.adv` and `gpfs.crypto` packages.

Installation toolkit support for populating cluster definition file

The installation toolkit now supports populating the cluster definition file with the current cluster state. For more information, see *Populating cluster definition file with current cluster state using the installation toolkit* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

Installation toolkit support for Red Hat Enterprise Linux 7.4 and 7.5

The installation toolkit now also supports Red Hat Enterprise Linux 7.4 and 7.5 on x86_64, PPC64, and PPC64LE architectures. For more information, see *Installation prerequisites* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

IBM Spectrum Scale GUI changes

The following main changes are made in the IBM Spectrum Scale GUI:

- Supports mounting and unmounting of file systems on selected nodes or group of nodes using GUI. For more information, see *Mounting a file system through GUI* and *Unmounting a file system through GUI* topics in *IBM Spectrum Scale: Administration Guide*.
- Added new **Storage > Pools** page. The Pools page provides details about configuration, health, capacity, and performance aspects of storage pools.
- Added new **Files > Active File Management** page. This new GUI page helps to view the configuration, health status, and performance of AFM, AFM DR, and gateway nodes.
- Added new **Monitoring > Tips** page. The tip events give recommendations to the user to avoid certain issues that might occur in the future. A tip disappears from the GUI when the problem behind the tip event is resolved.
- Added option to select events of type “tip” in the **Settings > Event Notifications > Email Recipients** page. You can configure whether to send email to the recipients if a tip event is reported in the system.
- Added detailed view in the **Files > Filesets** page. You can access the detailed view of individual filesets either by double-clicking the individual filesets or by selecting **View Details** option.
- Modified the **Storage > NSDs** page to list the rack, position, and node of the NSD in an FPO-enabled environment. This helps to sort the NSDs based on these parameters. The failure group definition is also modified to accommodate these new parameters.
- Added the **Customize the number of replicas** option in the **Files > Information Lifecycle** page to specify the number of replicas in a file placement rule.
- Modified the **Settings > Event Notifications** page to accept both IP address and host name for the email server.
- Added **Nodes** and **File Systems** tabs in the detailed view that is available in the **Files > Transparent Cloud Tiering** page.
- Added a separate **Properties** tab in the detailed view that is available in the **Monitoring > Nodes , Files > File Systems , and Storage > NSDs** pages.

IBM Spectrum Scale functionality to support GDPR requirements

To understand the requirements of EU General Data Protection Regulation (GDPR) compliance that are applicable to unstructured data storage and how IBM Spectrum Scale helps to address them, see the IBM Spectrum Scale functionality to support GDPR requirements technote.

Introduction of IBM Spectrum Scale management API Version 2

The architecture and syntax of IBM Spectrum Scale management API is changed. The new implementation is based on the GUI stack. The GUI server is managing and processing the API requests and commands. Version 2 has the following features:

- Reuses the GUI deployment's backend infrastructure, which makes introduction of new API commands easier.

- No separate configuration is required as the GUI installation takes care of the basic deployment.
- Fixes scalability issues and introduces new features such as filter parameter, field parameter, and paging.
- Supports large clusters with thousands of nodes.
- All POST, PUT, and DELETE requests are completed asynchronously. A "jobs" object is created immediately when such a request is submitted.
- The APIs are driven by the same WebSphere® server and object cache that is used by the IBM Spectrum Scale GUI.
- The **mmrest** command is no longer required for configuring the management API. The IBM Spectrum Scale GUI installation and configuration takes care of the API infrastructure configuration. For more information on how to configure IBM Spectrum Scale management API Version 2, see *Configuring IBM Spectrum Scale management API* in IBM Spectrum Scale:Administration Guide.

As the syntax and architecture are changed for the API, modified the entire set of commands, which were available in the Version 1. New API commands are also added for improved flexibility. For more information about the available commands, see *IBM Spectrum Scale management API commands* in *IBM Spectrum Scale: Command and Programming Reference*. You can also access the documentation corresponding to each API command from the GUI itself. The API documentation is available in the GUI at: <https://<IP address or host name of API server>:<port>/ibm/api/explorer/>. For example: <https://sclegui.ibm.com:443/ibm/api/explorer>.

Linux on Z enhancements

The following changes are made:

- IBM Spectrum Scale for Linux on Z now supports Remote Cluster Mount (Multi-cluster).
- SLES 12.2 and RHEL 7.3 are now supported by IBM Spectrum Scale for Linux on Z.

mmcallhome command: Addition of --long option to mmcallhome group list command

The **--long** option displays the long admin node names. For more information, see *mmcallhome command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmchconfig command: Setting an InfiniBand partition key

The **--verbsRdmaPkey** attribute specifies an InfiniBand partition key for a connection between a node and an InfiniBand server that is included in an InfiniBand partition. For more information, see *mmchconfig command* in the *IBM Spectrum Scale: Command and Programming Reference*.

mmdiag command: Status and queue statistics for NSD queues

The **--nsd** parameter displays the status and queue statistics for NSD queues.

For more information, see *mmdiag command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmfsck command: Severity of errors

The command displays a summary of the errors that were found that includes the severity of each error: **CRITICAL**, **NONCRITICAL**, or **HARMLESS**. You must specify the verbose or semi-verbose parameter to get this output. For more information, see *mmfsck command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmhealth command: Addition of new options to command

Addition of **AFM** and **THRESHOLD** options to the **mmhealth node show** and **mmhealth cluster show** commands. The **AFM** option displays the health status of a gateway node or cluster. The **THRESHOLD** option monitors whether the node-related thresholds rules evaluation is running as expected, and if the health status has changed as a result of the threshold limits being crossed.

Addition of **--clear** option to the **mmhealth node eventlog** command. This option clears the event log's database.

Addition of **threshold add** and **threshold delete** option to the **mmhealth** command. This option allows users to create and delete threshold rule.

Addition of **event hide**, **event unhide**, and **list hidden** options to the **mmhealth** command. The **event hide** option hides the specified TIP events, while the **event unhide** option reveals all TIP events that were previously hidden. The **list hidden** option shows all the TIP events that are added to the list of hidden events.

Addition of **config interval** option to the **mmhealth** command. The **config interval** option allows you to set the interval for monitoring the whole cluster.

For more information, see *mmhealth command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmkeyserv command: Updating a certificate or a connection

You can now get a fresh certificate from an Remote Key Management (RKM) server without rebuilding the connection. You can also temporarily update a connection by adding backup servers, reordering the list of backup servers, or changing the timeout, number of retries, or retry interval. For more information, see *mmkeyserv command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmlslicense command: Displaying disk and cluster size information

You can now get information about disk and cluster size with the **mmlslicense** command. For more information, see *mmlslicense command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmnetverify command: Enhancements

Several enhancements increase the capabilities of the **mmnetverify** command. Network checks are added to measure the total bandwidth, to check connectivity with the CTDB port, and to check connectivity with servers that are used with the Object protocol. If there are multiple local nodes, the command is run on all the local nodes in parallel. The lists of local nodes and target nodes accept node classes. The **--ces-override** parameter causes the command to consider all the nodes in the configuration to be CES-enabled. For more information, see *mmnetverify command* in *IBM Spectrum Scale: Command and Programming Reference*.

mmrestripfile command: Fix inconsistencies between file data and replicas

The **-c** option compares the data of individual files with their replicas and attempts to fix any inconsistencies. For more information, see *mmrestripfile command* in *IBM Spectrum Scale: Command and Programming Reference*.

Monitoring of AFM and AFM DR

Using commands:

- Functionality added to **mmhealth**, **mmdiag**, and **mmperfmon**.

Using IBM Spectrum Scale GUI:

- Added new **Files > Active File Management** page. This new GUI page helps to view the configuration, health status, and performance of AFM, AFM DR, and gateway nodes.

Mount options specific to IBM Spectrum Scale: syncnfs is now the default on Linux nodes

In the mount options specific to IBM Spectrum Scale, **syncnfs** is now the default on Linux nodes.

On AIX nodes, **nosyncnfs** is the default. For more information, see *Mount options specific to IBM Spectrum Scale* in *IBM Spectrum Scale: Command and Programming Reference*.

Protocol support on remotely mounted file systems

You can create an NFS/SMB export on a file system that is mounted from a remote cluster. For more information, see the *Using NFS/SMB protocol over remote cluster mounts* topic in the *IBM Spectrum Scale: Administration Guide*.

Tip added to event status to inform users when a configuration is not optimal

A new event type **TIP** is added to system health monitoring. A **Tip** is similar to a state-changing event, but can be hidden by the user. Like state-changing events, a tip is removed automatically if

the problem is resolved. For more information on **Tip**, see *Event type and monitoring status for system health* in the *IBM Spectrum Scale: Problem Determination Guide*

Quality of Service for I/O operations (QoS): Detailed statistics

You can now display more detailed statistics about IOPS rates for the QoS programs that are running on each node. The statistics are intended to be used as input for programs that analyze and display data. For more information, see *mmchqos command* and *mmlsqos command* in *IBM Spectrum Scale: Command and Programming Reference*.

Support for Samba 4.5

Transparent cloud tiering enhancements.

The following changes are done:

- Support for configuring and deploying WORM solutions. Your files will remain WORM-compliant, both in the file system and on the cloud object storage. For more information, see the *Deploying WORM solutions* topic in the *IBM Spectrum Scale: Administration Guide*.
- Support for configuring Transparent cloud tiering with a proxy server.
- Support for configuring cloud retention time, which overrides the default value.
- Support for restoring only the file stubs from the cloud storage tier in situations where files are deleted from the local file system.
- Support for Power8 Little Endian platform.

Note: This feature is available from 4.2.3.1 onwards.

- Substantial improvement in the performance when files are transparently recalled from the storage tier.
- Support for manually deleting orphaned cloud objects before retention time expires. For more information, see the *Manually deleting cloud objects before retention time* topic in the *IBM Spectrum Scale: Administration Guide*.
- Support for migrating files in the co-resident state, by which applications can directly access data without performing any recall operation. For more information, see the *Pre-migrating files to the cloud storage tier* topic in the *IBM Spectrum Scale: Administration Guide*

-Y option

Added the -Y option to the following commands:

• mmblock	• mmhealth	• mmlsfileset	• mmlsnodeclass	• mmnetverify
• mmccloudgateway	• mmkeyserv	• mmlsfs	• mmlsnsd	• mmnfs
• mmdf	• mmlscluster	• mmlslicense	• mmlspolicy	• mmrepquota
• mmdiag	• mmlsconfig	• mmlsmgr	• mmlsquota	• mmsmb
• mmgetstate	• mmlsdisk	• mmlsmount	• mmlssnapshot	• mmuserauth

Documented commands, structures, and subroutines

The following lists the modifications to the documented commands, structures, and subroutines:

New commands

The following commands are new:

- **mmclidecode**

New structures

There are no new structures.

New subroutines

There are no new subroutines.

Changed commands

The following commands were changed:

- **mmadquery**
- **mmbackup**
- **mmblock**
- **mmcallhome**
- **mmces**
- **mmcesdr**
- **mmchattr**
- **mmchconfig**
- **mmchqos**
- **mmcloudgateway**
- **mmcrnsd**
- **mmdf**
- **mmdiag**
- **mmfsck**
- **mmgetstate**
- **mmhadoopctl**
- **mmhealth**
- **mmimgbackup**
- **mmimgrestore**
- **mmkeyserv**
- **mmlsattr**
- **mmlscluster**
- **mmlsconfig**
- **mmldisk**
- **mmlsfileset**
- **mmlsfs**
- **mmlslicense**
- **mmlsmgr**
- **mmlsmount**
- **mmlsnodeclass**
- **mmlsnsd**
- **mmlspolicy**
- **mmlsqos**
- **mmlsquota**
- **mmlssnapshot**
- **mmnetverify**
- **mmnfs**
- **mmprotocoltrace**
- **mmrepquota**
- **mmssmb**
- **mmuserauth**
- **spectrumscale**

| **Changed structures**
| The following structures were changed:
| • **gpfs_iattr64_t**

| **Changed subroutines**
| The following subroutines were changed:
| • **gpfs_prealloc**

| **Deleted commands**
| **mmrest**

| **Deleted structures**
| There are no deleted structures.

| **Deleted subroutines**
| There are no deleted subroutines.

| **Messages**
| The following are the new, changed, and deleted messages:

| **New messages**
| 6027-1525, 6027-1756, 6027-2392, 6027-2393, 6027-2503, 6027-2504, and 6027-3258

| **Changed messages**
| 6027-1023, 6027-1725

| **Deleted messages**
| None.

| **Changes in documentation**

| **Big data and analytics support**
| Moved the entire big data and analytics support information to a new section. See the topic *Big data and analytics support* in *IBM Spectrum Scale: Big Data and Analytics Guide*.

| **Restructured events page**
| The events page was split up into 19 different pages, with a separate page for each component. See the topic *Events* in the *IBM Spectrum Scale: Problem Determination Guide*.

| **Renamed “REST API” to “IBM Spectrum Scale management API” in the documentation.**

| **List of documentation changes in product guides and respective Knowledge Center sections**
| The following is a list of documentation changes including changes in topic titles, changes in placement of topics, and deleted topics:

| **Table 3. List of changes in documentation**

Guide	Knowledge center section	List of changes
Concepts, Planning, and Installation Guide	Product overview	<p>Under <i>IBM Spectrum Scale management API</i></p> <ul style="list-style-type: none"> Moved the IBM Spectrum Scale management API topics from the <i>Administering</i> section to the <i>IBM Spectrum Scale management API</i> section.
	Planning	<p>Under <i>Planning for protocols</i> → <i>Authentication considerations</i></p> <ul style="list-style-type: none"> Changed the title <i>Authentication for file access</i> to <i>Authentication and ID mapping for file access</i> <p>Under <i>Planning for protocols</i> → <i>Planning for SMB</i></p> <ul style="list-style-type: none"> The <i>SMB share limitations</i> topic under <i>Administering</i> → <i>Managing protocol data exports</i> → <i>Managing SMB shares</i> has been removed. Limitations from the <i>SMB share limitations</i> topic have been added in the <i>SMB limitations</i> topic.
	Installing and upgrading	<ul style="list-style-type: none"> Removed <i>Installing the Scale Management server (REST API)</i> section. Moved <i>Manually upgrading pmswift</i> and <i>Manually upgrading the Performance Monitoring tool</i> from <i>Manually installing the Performance Monitoring tool</i> to the <i>Migration, coexistence and compatibility</i> section. Moved <i>Upgrading IBM Spectrum Scale components with the installation toolkit</i> from <i>Using the spectrumscale installation toolkit to perform installation tasks: Explanations and examples</i> to the <i>Migration, coexistence and compatibility</i> section. Created the <i>Upgrading Object packages</i> section: <ul style="list-style-type: none"> Moved the <i>Upgrading Object packages to version 4.2.2.x from 4.2.2.x</i> topic from <i>Migrating to IBM Spectrum Scale 4.2.2.x from IBM Spectrum Scale 4.2.0.x or later</i> to the <i>Upgrading Object packages</i> section. Added the <i>Upgrading Object packages to version 4.2.3.x from 4.2.2.x</i> topic.
Administration Guide	Configuring	<ul style="list-style-type: none"> Removed the <i>Configuring and starting the Scale Management server (REST API)</i> section. Removed the <i>Enabling Cloud services performance monitoring metrics on the GUI</i> topic under the <i>Configuring</i> → <i>Configuring and tuning your system for Cloud services</i> section.
	Administering	<p>Under <i>File Placement Optimizer</i></p> <ul style="list-style-type: none"> Added <i>Data locality based copy</i> in <i>IBM Spectrum Scale: Administration Guide</i> section. Added <i>mmgetlocation</i> in <i>IBM Spectrum Scale: Administration Guide</i> Data locality restore section is renamed to <i>Data locality</i> in <i>IBM Spectrum Scale: Administration Guide</i>. Added the maintenance steps for IBM Spectrum Scale FPO. Added the performance tuning steps for IBM Spectrum Scale Sharing Nothing Cluster.

| *Table 3. List of changes in documentation (continued)*

Guide	Knowledge center section	List of changes
Problem Determination Guide	Monitoring	<p>Under <i>Monitoring AFM and AFM DR</i>, renamed the following topics:</p> <ul style="list-style-type: none"> • <i>Fileset states for AFM</i> to <i>Monitoring fileset states for AFM</i> • <i>Fileset states for AFM DR</i> to <i>Monitoring fileset states for AFM DR</i> • <i>Callback events for AFM and AFM DR</i> to <i>Monitoring callback events for AFM and AFM DR</i> • <i>Prefetch</i> to <i>Monitoring prefetch</i> • Moved the <i>Monitoring callback events for AFM and AFM DR</i> topic under <i>Monitoring health and events</i> • Moved the <i>Monitoring with mmpmon</i> topic under <i>Monitoring performance</i> <p>Restructured the following topics:</p> <ul style="list-style-type: none"> • <i>Monitoring system health by using the mmhealth command</i> This topic has been split into four topics: <ul style="list-style-type: none"> – Monitoring the health of a node – Event type and monitoring status for system health – Threshold monitoring for system health – Use cases • <i>Monitoring the IBM Spectrum Scale system by using call home</i> This topic has been split into four topics: <ul style="list-style-type: none"> – Understanding call home – Configuring call home to enable manual and automated data upload – Monitoring, uploading, and sharing collected data with IBM Support – Use cases • <i>List of performance metrics</i> This topic has been split into three topics: <ul style="list-style-type: none"> – Linux metrics – GPFS metrics – Protocol metrics
	Troubleshooting	<ul style="list-style-type: none"> • Under <i>SMB issues</i>, added a new topic <i>Slow access to SMB caused by contended access to files or directories</i>.
Command and Programming Reference	Command reference	<ul style="list-style-type: none"> • Removed the mmrest command man page.
	Programming reference	<ul style="list-style-type: none"> • Added documentation for each IBM Spectrum Scale management API Version 2 command.

Changes in the Library and related publications section

- Under *Library and related publications*, the following topics were updated:
 - *Redbooks®, Redpapers™, and Blueprints*: Six new links added.
 - *ISV links*
 - *Applying IBM Spectrum Scale → Using AFM with object*
- Under *AFM-based Disaster Recovery* section, the *Fallback of multiple filesets* use case was added.

Chapter 1. Command reference

A list of all the GPFS commands and a short description of each is presented in this topic.

Table 4 summarizes the GPFS-specific commands.

Table 4. GPFS commands

Command	Purpose
"gpfs.snap command" on page 6	Creates an informational system snapshot at a single point in time. This system snapshot consists of information such as cluster configuration, disk configuration, network configuration, network status, GPFS logs, dumps, and traces.
"mmaddcallback command" on page 10	Registers a user-defined command that GPFS will execute when certain events occur.
"mmadddisk command" on page 23	Adds disks to a GPFS file system.
"mmaddnode command" on page 29	Adds nodes to a GPFS cluster.
"mmadquery command" on page 32	
"mmafmconfig command" on page 37	Can be used to manage home caching behavior and mapping of gateways and home NFS exported servers.
"mmafmctl command" on page 40	This command is for various operations and reporting information on all filesets. It is recommended to read the <i>IBM Spectrum Scale: Administration Guide AFM</i> and <i>AFM Disaster Recovery</i> chapters in conjunction with this manual for detailed description of the functions.
"mmafmlocal command" on page 54	Provides a list of cached files and file statistics such as inode number, allocated blocks, and so on.
"mmapplypolicy command" on page 56	Deletes files, migrates files between storage pools, or does file compression or decompression in a file system as directed by policy rules.
"mmauth command" on page 67	Manages secure access to GPFS file systems.
"mmbbackup command" on page 72	Performs a backup of a GPFS file system or independent fileset to an IBM Spectrum Protect™ server.
"mmbbackupconfig command" on page 81	Collects GPFS file system configuration information.
"mmblock command" on page 83	Manages the iSCSI block service.
"mmbuildgpl command" on page 88	Manages prerequisite packages for Linux and builds the GPFS portability layer.
"mmccallhome command" on page 90	Manages the call home operations.
"mmces command" on page 103	Manages CES configuration.
"mmcesdr command" on page 113	Manages protocol cluster disaster recovery.
"mmchattr command" on page 122	Changes attributes of one or more GPFS files.
"mmchcluster command" on page 128	Changes GPFS cluster configuration data.
"mmchconfig command" on page 132	Changes GPFS configuration parameters.
"mmchdisk command" on page 161	Changes state or parameters of one or more disks in a GPFS file system.
"mmcheckquota command" on page 169	Checks file system user, group and fileset quotas.
"mmchfileset command" on page 173	Changes the attributes of a GPFS fileset.

Table 4. GPFS commands (continued)

Command	Purpose
"mmchfs command" on page 179	Changes the attributes of a GPFS file system.
"mmchlicense command" on page 185	Controls the type of GPFS license associated with the nodes in the cluster.
"mmchmgr command" on page 188	Assigns a new file system manager node or cluster manager node.
"mmchnode command" on page 190	Changes node attributes.
"mmchnodeclass command" on page 195	Changes user-defined node classes.
"mmchnsd command" on page 198	Changes Network Shared Disk (NSD) configuration attributes.
"mmchpolicy command" on page 201	Establishes policy rules for a GPFS file system.
"mmchpool command" on page 204	Modifies storage pool properties.
"mmchqos command" on page 206	Changes the Quality of Service for I/O operations (QoS) settings for a file system.
"mmclidecode command" on page 213	Decodes the parseable command output field.
"mmclone command" on page 215	Creates and manages file clones.
"mmcloudgateway command" on page 218	Creates and manages the cloud storage tier.
"mmcrcluster command" on page 239	Creates a GPFS cluster from a set of nodes.
"mmcrfileset command" on page 244	Creates a GPFS fileset.
"mmcrfs command" on page 250	Creates a GPFS file system.
"mmcrnodeclass command" on page 260	Creates user-defined node classes.
"mmcrnsd command" on page 262	Creates Network Shared Disks (NSDs) used by GPFS.
"mmcrsnapshot command" on page 267	Creates a snapshot of a file system or fileset at a single point in time.
"mmdefedquota command" on page 272	Sets default quota limits.
"mmdefquotaoff command" on page 275	Deactivates default quota limit usage.
"mmdefquotao command" on page 278	Activates default quota limit usage.
"mmdefragfs command" on page 281	Reduces disk fragmentation by increasing the number of full free blocks available to the file system.
"mmdelacl command" on page 284	Deletes a GPFS access control list.
"mmdelcallback command" on page 286	Deletes one or more user-defined callbacks from the GPFS system.
"mmdeldisk command" on page 287	Deletes disks from a GPFS file system.
"mmdelfileset command" on page 292	Deletes a GPFS fileset.
"mmdefefs command" on page 295	Removes a GPFS file system.
"mmdelnodenode command" on page 297	Removes one or more nodes from a GPFS cluster.
"mmdelnodenodeclass command" on page 300	Deletes user-defined node classes.
"mmdelnsd command" on page 302	Deletes Network Shared Disks (NSDs) from the GPFS cluster.
"mmdeletesnapshot command" on page 304	Deletes a GPFS snapshot.
"mmdf command" on page 308	Queries available file space on a GPFS file system.
"mmdiag command" on page 312	Displays diagnostic information about the internal GPFS state on the current node.
"mmdsh command" on page 319	Runs commands on multiple nodes or network connected hosts at the same time.

Table 4. GPFS commands (continued)

Command	Purpose
“mmeditacl command” on page 322	Creates or changes a GPFS access control list.
“mmedquota command” on page 325	Sets quota limits.
“mmexportfs command” on page 329	Retrieves the information needed to move a file system to a different cluster.
“mmfsck command” on page 331	Checks and repairs a GPFS file system.
“mmfsctl command” on page 343	Issues a file system control request.
“mmgetacl command” on page 347	Displays the GPFS access control list of a file or directory.
“mmgetstate command” on page 350	Displays the state of the GPFS daemon on one or more nodes.
“mmhadoopctl command” on page 353	Installs and sets up the GPFS connector for a Hadoop distribution; starts or stops the GPFS connector daemon on a node.
“mmimgbackup command” on page 367	Performs a backup of a single GPFS file system metadata image.
“mmimgrestore command” on page 371	Restores a single GPFS file system from a metadata image.
“mmimportfs command” on page 374	Imports into the cluster one or more file systems that were created in another GPFS cluster.
“mmkeyserv command” on page 378	Manages encryption key servers and clients.
“mmlinkfileset command” on page 390	Creates a junction that references the root directory of a GPFS fileset.
“mmlsattr command” on page 392	Queries file attributes.
“mmlscallback command” on page 395	Lists callbacks that are currently registered in the GPFS system.
“mmlscluster command” on page 397	Displays the current configuration information for a GPFS cluster.
“mmlsconfig command” on page 400	Displays the current configuration data for a GPFS cluster.
“mmlsdisk command” on page 402	Displays the current configuration and state of the disks in a file system.
“mmlsfileset command” on page 406	Displays attributes and status for GPFS filesets.
“mmlsfs command” on page 411	Displays file system attributes.
“mmlslicense command” on page 415	Displays information about the GPFS node licensing designation or about disk and cluster capacity.
“mmlsmgr command” on page 417	Displays which node is the file system manager for the specified file systems or which node is the cluster manager.
“mmlsmount command” on page 419	Lists the nodes that have a given GPFS file system mounted.
“mmlsnodeclass command” on page 421	Displays node classes defined in the system.
“mmlsnsd command” on page 423	Displays Network Shared Disk (NSD) information for the GPFS cluster.
“mmlspolicy command” on page 426	Displays policy information.
“mmlspool command” on page 428	Displays information about the known storage pools.
“mmlsquota command” on page 435	Displays quota information for a user, group, or fileset.
“mmlsqos command” on page 430	Displays the I/O performance values of a file system, when you enable Quality of Service for I/O operations (QoS) with the mmchqos command.
“mmlssnapshot command” on page 439	Displays GPFS snapshot information.

Table 4. GPFS commands (continued)

Command	Purpose
" mmmigratefs command " on page 442	Performs needed conversions to support new file system features.
" mmmount command " on page 444	Mounts GPFS file systems on one or more nodes in the cluster.
" mmnetverify command " on page 446	Verifies network configuration and operation in a cluster.
" mmnfs command " on page 454	Manages NFS exports and configuration.
" mmnsdiscover command " on page 464	RedisCOVERS paths to the specified network shared disks.
" mmobj command " on page 466	Manages configuration of Object protocol service, and administers storage policies for object storage, unified file and object access, and multi-region object deployment.
" mmperfmon command " on page 481	Configures the Performance Monitoring tool and lists the performance metrics.
" mmpmon command " on page 492	Manages performance monitoring and displays performance information.
" mmprotocoltrace command " on page 497	Starts, stops, and monitors tracing for the CES protocols.
" mmssnap command " on page 501	Creates or deletes identical snapshots on the cache and home clusters, or shows the status of snapshots that have been queued up on the gateway nodes.
" mmputacl command " on page 504	Sets the GPFS access control list for the specified file or directory.
" mmquotaoff command " on page 507	Deactivates quota limit checking.
" mmquotao command " on page 509	Activates quota limit checking.
" mmremotecluster command " on page 511	Manages information about remote GPFS clusters.
" mmremotefs command " on page 514	Manages information needed for mounting remote GPFS file systems.
" mmrepquota command " on page 517	Displays file system user, group, and fileset quotas.
" mmrestoreconfig command " on page 521	Restores file system configuration information.
" mmrestorefs command " on page 525	Restores a file system or an independent fileset from a snapshot.
" mmrestripefile command " on page 529	Rebalances or restores the replication factor of the specified files, or performs any incomplete or deferred file compression or decompression.
" mmrestripefs command " on page 532	Rebalances or restores the replication factor of all the files in a file system. Alternatively, this command performs any incomplete or deferred file compression or decompression of all the files in a file system.
" mmrpldisk command " on page 539	Replaces the specified disk.
" mmsdrrestore command " on page 546	Restores the latest GPFS system files on the specified nodes.
" mmsetquota command " on page 548	Sets quota limits.
" mmshutdown command " on page 552	Unmounts all GPFS file systems and stops GPFS on one or more nodes.
" mmsmb command " on page 554	Administers SMB shares, export ACLs, and global configuration.
" mmsnapdir command " on page 565	Controls how the special directories that connect to snapshots appear.
" mmstartup command " on page 569	Starts the GPFS subsystem on one or more nodes.

Table 4. GPFS commands (continued)

Command	Purpose
"mmtracectl command" on page 571	Sets up and enables GPFS tracing.
"mmumount command" on page 575	Unmounts GPFS file systems on one or more nodes in the cluster.
"mmunlinkfileset command" on page 578	Removes the junction to a GPFS fileset.
"mmuserauth command" on page 581	Manages the authentication configuration of file and object access protocols. The configuration allows protocol access methods to authenticate users who need to access data that is stored on the system over these protocols.
"mmwinservcctl command" on page 602	Manages the mmwinserv Windows service.
"spectrumscale command" on page 604	Installs and configures GPFS; adds nodes to a cluster; deploys and configures protocols, performance monitoring tools, and authentication services; and upgrades GPFS and protocols.

The following commands are specific to IBM Spectrum Scale RAID and are documented in *IBM Spectrum Scale RAID: Administration*:

- **mmaddcomp**
- **mmaddcompspec**
- **mmaddpdisk**
- **mmchcarrier**
- **mmchcomp**
- **mmchcomploc**
- **mmchenclosure**
- **mmchfirmware**
- **mmchpdisk**
- **mmchrecoverygroup**
- **mmcrrecoverygroup**
- **mmcrvdisk**
- **mmdelcomp**
- **mmdelcomploc**
- **mmdelcompspec**
- **mmdelpdisk**
- **mmdelrecoverygroup**
- **mmdelvdisk**
- **mmdiscovercomp**
- **mmgetdisktopology**
- **mmlscomp**
- **mmlscomploc**
- **mmlscompspec**
- **mmlsenclosure**
- **mmlsfirmware**
- **mmlspdisk**
- **mmlsrecoverygroup**
- **mmlsrecoverygroupevents**
- **mmlsvdisk**
- **mmsyncdisplayid**

gpfs.snap

gpfs.snap command

Creates an informational system snapshot at a single point in time. This system snapshot consists of information such as cluster configuration, disk configuration, network configuration, network status, GPFS logs, dumps, and traces.

Synopsis

```
gpfs.snap [-d OutputDirectory] [-m | -z]
           [-a | -N {Node[,Node...] | NodeFile | NodeClass}]
           [--check-space | --no-check-space | --check-space-only]
           [--cloud-gateway {BASIC | FULL} ] [--full-collection] [--deadlock [--quick] |
             --limit-large-files {YYYY:MM:DD:HH:MM | NumberOfDaysBack | latest}]
           [--exclude-aix-disk-attr] [--exclude-aix-lvm] [--exclude-merge-logs]
           [--exclude-net] [--gather-logs] [--mmdf] [--performance] [--prefix]
           [--protocol ProtocolType[,ProtocolType,...]] [--timeout Seconds]
           [--purge-files KeepNumberOfDaysBack] [--hadoop]
```

Availability

- Available on all IBM Spectrum Scale editions.

Description

Use the **gpfs.snap** command as the main tools to gather data when a GPFS problem is encountered, such as a hung file system, a hung GPFS command, or a daemon assert.

The **gpfs.snap** command gathers information (for example, GPFS internal dumps, traces, and kernel thread dumps) to solve a GPFS problem.

Note: By default, large debug files are now a delta collection, which means that they are only collected when there are new files since the previous run of **gpfs.snap**. To override this default behavior, use either the **--limit-large-files** or **--full-collection** options.

Note: This is a service tool and options might change dynamically. The tool impacts performance and occupies disk space when it runs.

Parameters

-d *OutputDirectory*

Specifies the output directory. The default is /tmp/gpfs.snapOut.

-m Specifying this option is equivalent to specifying **--exclude-merge-logs** with **-N**.

-z

Collects **gpfs.snap** data only from the node on which the command is invoked. No master data is collected.

-a

Directs **gpfs.snap** to collect data from all nodes in the cluster. This is the default.

-N *{Node[,Node ...] | NodeFile | NodeClass}*

Specifies the nodes from which to collect **gpfs.snap** data. This option supports all defined node classes. For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in *IBM Spectrum Scale: Administration Guide*.

--check-space

Specifies that space checking is performed before collecting data.

--no-check-space

Specifies that no space checking is performed. This is the default.

--check-space-only

Specifies that only space checking is performed. No data is collected.

--cloud-gateway {BASIC | FULL}

With the BASIC option, when the Transparent cloud tiering service is enabled, the snap will collect information such as logs, traces, Java™ cores, along with minimal system and IBM Spectrum Scale cluster information specific to transparent cloud tiering. No customer sensitive information is collected.

Note: The default behavior of the **gpfs.snap** command includes basic information of Transparent cloud tiering, in addition to the GPFS information.

With the FULL option, extra details such as Java Heap dump are collected, along with the information captured with the BASIC option.

--full-collection

Specifies that all large debug files are collected instead of the default behavior that only collects new files since the previous run of **gpfs.snap**.

--deadlock

Collects only the minimum amount of data necessary to debug a deadlock problem. Part of the data collected is the output of the **mmfsadm dump all** command. This option ignores all other options except for **-a**, **-N**, **-d**, and **--prefix**.

--quick

Collects less data when specified along with the **--deadlock** option. The output includes **mmfsadm dump most**, **mmfsadm dump kthreads**, and 10 seconds of trace in addition to the usual **gpfs.snap** output.

--limit-large-files {YYYY:MM:DD:HH:MM | NumberOfDaysBack | latest}]

Specifies a time limit to reduce the number of large files collected.

--exclude-aix-disk-attr

Specifies that data about AIX disk attributes will not be collected. Collecting data about AIX disk attributes on an AIX node that has a large number of disks could be very time-consuming, so using this option could help improve performance.

--exclude-aix-lvm

Specifies that data about the AIX Logical Volume Manager (LVM) will not be collected.

--exclude-merge-logs

Specifies that merge logs and waiters will not be collected.

--exclude-net

Specifies that network-related information will not be collected.

--gather-logs

Gathers, merges, and chronologically sorts all of the **mmfs.log** files. The results are stored in the directory specified with **-d** option.

--mmdf

Specifies that **mmdf** output will be collected.

--performance

Specifies that performance data is to be gathered.

Note: The performance script can take up to 30 minutes to run; therefore, it is not included when all other types of protocol information are gathered by default. Specifying this option is the only way to turn on the gathering of performance data.

--prefix

Specifies that the prefix name **gpfs.snap** will be added to the tar file.

gpfs.snap

--protocol *ProtocolType[,ProtocolType,...]*

Specifies the type (or types) of protocol information to be gathered. By default, whenever any protocol is enabled on a file system, information is gathered for all types of protocol information (except for performance data; see the **--performance** option). However, when the **--protocol** option is specified, the automatic gathering of all protocol information is turned off, and only the specified type of protocol information will be gathered. The following values for *ProtocolType* are accepted:

- smb**
- nfs**
- object**
- authentication**
- ces**
- core**
- none**

--timeout *Seconds*

Specifies the timeout value, in seconds, for all commands.

--purge-files *KeepNumberOfDaysBack*

Specifies that large debug files will be deleted from the cluster nodes based on the *KeepNumberOfDaysBack* value. If **0** is specified, all of the large debug files will be deleted. If a value greater than **0** is specified, large debug files that are older than the number of days specified will be deleted. For example, if the value **2** is specified, the previous two days of large debug files are retained.

This option is not compatible with many of the **gpfs.snap** options because it only removes files and does not collect any **gpfs.snap** data.

--hadoop

Specifies that Hadoop data is to be gathered.

Use the **-z** option to generate a non-master snapshot. This is useful if there are many nodes on which to take a snapshot, and only one master snapshot is needed. For a GPFS problem within a large cluster (hundreds or thousands of nodes), one strategy might call for a single master snapshot (one invocation of **gpfs.snap** with no options), and multiple non-master snapshots (multiple invocations of **gpfs.snap** with the **-z** option).

Use the **-N** option to obtain **gpfs.snap** data from multiple nodes in the cluster. When the **-N** option is used, the **gpfs.snap** command takes non-master snapshots of all the nodes specified with this option and a master snapshot of the node on which it was invoked.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **gpfs.snap** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Examples

- To collect **gpfs.snap** on all nodes with the default data, issue the command:

```
(09:25:47) c34f2n03:~ # gpfs.snap
gpfs.snap started at Mon Feb  8 09:25:54 EST 2016.
Gathering common data.
Gathering Linux specific data...
Gathering trace reports and internal dumps...
gpfs.snap: Spawning remote gpfs.snap calls. Master is c34f2n03.
This may take a while.

Copying file
/tmp/gpfs.snapOut/18720/gpfs.snap.c13c1apv7_0208092648.out.tar.gz from c13c1apv7.gpfs.net ...
gpfs.snap.c13c1apv7_0208092648.out.tar.gz 100% 592KB 592.2KB/s 00:00
Successfully copied file
/tmp/gpfs.snapOut/18720/gpfs.snap.c13c1apv7_0208092648.out.tar.gz from c13c1apv7.gpfs.net.
```

```
Copying file
/tmp/gpfs.snapOut/18720/gpfs.snap.c6f2bc4n8_0208092705.out.tar.gz from c6f2bc4n8.gpfs.net ...
gpfs.snap.c6f2bc4n8_0208092705.out.tar.gz 100% 928KB 927.9KB/s 00:00
Successfully copied file
/tmp/gpfs.snapOut/18720/gpfs.snap.c6f2bc4n8_0208092705.out.tar.gz from c6f2bc4n8.gpfs.net.
Gathering cluster wide protocol data
Packaging master node data.
Writing * to file
/tmp/gpfs.snapOut/18720/collect/gpfs.snap.c34f2n03_master_0208092554.out.tar.gz
Packaging all data.
Writing . to file /tmp/gpfs.snapOut/18720/all.0208092554.tar
gpfs.snap completed at Mon Feb  8 09:26:45 EST 2016
#####
Send file /tmp/gpfs.snapOut/18720/all.0208092554.tar to IBM Service
Examine previous messages to determine additional required data.
#####
```

After this command customer would send the tar file (highlighted) to IBM service as per the message

- To collect **gpfs.snap** on specific nodes, issue the command:

```
(09:32:38) c34f2n03:~ # gpfs.snap -N c34f2n03,c13c1apv7
gpfs.snap started at Mon Feb  8 09:32:48 EST 2016.
Gathering common data.
Gathering Linux specific data...
Gathering trace reports and internal dumps...
gpfs.snap: Spawning remote gpfs.snap calls. Master is c34f2n03.
This may take a while.

Copying file
/tmp/gpfs.snapOut/23453/gpfs.snap.c13c1apv7_0208093340.out.tar.gz from c13c1apv7.gpfs.net ...
gpfs.snap.c13c1apv7_0208093340.out.tar.gz 100% 583KB 583.1KB/s 00:00
Successfully copied file
/tmp/gpfs.snapOut/23453/gpfs.snap.c13c1apv7_0208093340.out.tar.gz from c13c1apv7.gpfs.net.
Gathering cluster wide protocol data
Packaging master node data.
Writing * to file /tmp/gpfs.snapOut/23453/collect/gpfs.snap.c34f2n03_master_0208093248.out.tar.gz
Packaging all data.
Writing . to file /tmp/gpfs.snapOut/23453/all.0208093248.tar
gpfs.snap completed at Mon Feb  8 09:33:34 EST 2016
#####
Send file /tmp/gpfs.snapOut/23453/all.0208093248.tar to IBM Service
Examine previous messages to determine additional required data.
#####
```

Location

/usr/lpp/mmfs/bin

mmaddcallback

mmaddcallback command

Registers a user-defined command that GPFS will execute when certain events occur.

Synopsis

```
mmaddcallback CallbackIdentifier --command CommandPathname
    --event Event [,Event...] [--priority Value]
    [--async | --sync [--timeout Seconds] [--onerror Action]]
    [-N {Node [,Node...] | NodeFile | NodeClass}]
    [--parms ParameterString ...]
```

or

```
mmaddcallback {-S Filename | --spec-file Filename}
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmaddcallback** command to register a user-defined command that GPFS executes when certain events occur.

The callback mechanism is intended to provide notifications when node and cluster events occur. Invoking complex or long-running commands, or commands that involve GPFS files, may cause unexpected and undesired results, including loss of file system availability. This is particularly true when the **--sync** option is specified.

Note: For documentation about local events (callbacks) and variables for IBM Spectrum Scale RAID, see the separate publication *IBM Spectrum Scale RAID: Administration*.

Parameters

CallbackIdentifier

Specifies a user-defined unique name that identifies the callback. It can be up to 255 characters long. It cannot contain special characters (for example, a colon, semicolon, blank, tab, or comma) and it cannot start with the letters gpfs or mm (which are reserved for GPFS internally defined callbacks).

--command *CommandPathname*

Specifies the full path name of the executable to run when the event occurs. On Windows, *CommandPathname* must be a Korn shell script because it will be invoked in the Cygwin **ksh** environment.

The executable called by the callback facility must be installed on all nodes on which the callback can be triggered. Place the executable in a local file system (not in a GPFS file system) so that it is accessible even when the GPFS file system is unavailable.

--event *Event [,Event...]*

Specifies a list of events that trigger the callback. The value defines when the callback is invoked. There are two kinds of events: global events and local events. A global event triggers a callback on all nodes in the cluster, such as a **nodeLeave** event, which informs all nodes in the cluster that a node has failed. A local event triggers a callback only on the node on which the event occurred, such as mounting a file system on one of the nodes.

Table 5 on page 15 lists the supported global events and their parameters.

Table 6 on page 16 lists the supported local events and their parameters.

Local events for IBM Spectrum Scale RAID are documented in *IBM Spectrum Scale RAID: Administration*.

--priority *Value*

Specifies a floating point number that controls the order in which callbacks for a given event are run. Callbacks with a smaller numerical value are run before callbacks with a larger numerical value. Callbacks that do not have an assigned priority are run last. If two callbacks have the same priority, the order in which they are run is undetermined.

--async | --sync [--timeout Seconds] [--onerror Action]

Specifies whether GPFS will wait for the user program to complete and for how long it will wait. The default is **--async** (GPFS invokes the command asynchronously). **--onerror Action** specifies one of the following actions that GPFS is to take if the callback command returns a nonzero error code:

continue

GPFS ignores the result from executing the user-provided command. This is the default.

quorumLoss

The node executing the user-provided command will voluntarily resign as, or refrain from taking over as, cluster manager. This action is valid only in conjunction with the **tiebreakerCheck** event.

shutdown

GPFS will be shut down on the node executing the user-provided command.

-N {Node[,Node...] | NodeFile | NodeClass}

Defines the set of nodes on which the callback is invoked. For global events, the callback is invoked only on the specified set of nodes. For local events, the callback is invoked only if the node on which the event occurred is one of the nodes specified by the **-N** option. The default is **-N all**. For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

--parms *ParameterString* ...

Specifies parameters to be passed to the executable specified with the **--command** parameter. The **--parms** parameter can be specified multiple times.

When the callback is invoked, the combined parameter string is tokenized on white-space boundaries. Constructs of the form `%name` and `%name.qualifier` are assumed to be GPFS variables and are replaced with their appropriate values at the time of the event. If a variable does not have a value in the context of a particular event, the string **UNDEFINED** is returned instead.

GPFS recognizes the following variables:

%blockLimit

Specifies the current hard quota limit in KB.

%blockQuota

Specifies the current soft quota limit in KB.

%blockUsage

Specifies the current usage in KB for quota-related events.

%ccrObjectName

Specifies the name of the modified object.

%ccrObjectValue

Specifies the value of the modified object.

%ccrObjectVersion

Specifies the version of the modified object.

%clusterManager[.qualifier]

Specifies the current cluster manager node.

%clusterName

Specifies the name of the cluster where this callback was triggered.

mmaddcallback

%ckDataLen

Specifies the length of data involved in a checksum mismatch.

%ckErrorCountClient

Specifies the cumulative number of errors for the client side in a checksum mismatch.

%ckErrorCountNSD

Specifies the cumulative number of errors for the NSD side in a checksum mismatch.

%ckErrorCountServer

Specifies the cumulative number of errors for the server side in a checksum mismatch.

%ckNSD

Specifies the NSD involved.

%ckOtherNode

Specifies the IP address of the other node in an NSD checksum event.

%ckReason

Specifies the reason string indicating why a checksum mismatch callback was invoked.

%ckReportingInterval

Specifies the error-reporting interval in effect at the time of a checksum mismatch.

%ckRole

Specifies the role (client or server) of a GPFS node.

%ckStartSector

Specifies the starting sector of a checksum mismatch.

%daName

Specifies the name of the declustered array involved.

%daRemainingRedundancy

Specifies the remaining fault tolerance in a declustered array.

%diskName

Specifies a disk or a comma-separated list of disk names for which this callback is triggered.

%downNodes[.qualifier]

Specifies a comma-separated list of nodes that are currently down. Only nodes local to the given cluster are listed. Nodes which are in a remote cluster but have temporarily joined the cluster are not included.

%eventName

Specifies the name of the event that triggered this callback.

%eventNode[.qualifier]

Specifies a node or comma-separated list of nodes on which this callback is triggered. Note that the list may include nodes which are not local to the given cluster, but have temporarily joined the cluster to mount a file system provided by the local cluster. Those remote nodes could leave the cluster if there is a node failure or if the file systems are unmounted.

%filesLimit

Specifies the current hard quota limit for the number of files.

%filesQuota

Specifies the current soft quota limit for the number of files.

%filesUsage

Specifies the current number of files for quota-related events.

%filesetName

Specifies the name of a fileset for which the callback is being executed.

%filesetSize

Specifies the size of the fileset.

%fsErr

Specifies the file system structure error code.

%fsName

Specifies the file system name for file system events.

%hardLimit

Specifies the hard limit for the block.

%homeServer

Specifies the name of the home server.

%inodeLimit

Specifies the hard limit of the inode.

%inodeQuota

Specifies the soft limit of the inode.

%inodeUsage

Specifies the total number of files in the fileset.

%myNode[.qualifier]

Specifies the node where callback script is invoked.

%nodeName

Specifies the node name to which the request is sent.

%nodeNames

Specifies a space-separated list of node names to which the request is sent.

%pcacheEvent

Specifies the pcache related events.

%pdFru

Specifies the FRU (field replaceable unit) number of the pdisk.

%pdLocation

The physical location code of a pdisk.

%pdName

The name of the pdisk involved.

%pdPath

The block device path of the pdisk.

%pdPriority

The replacement priority of the pdisk.

%pdState

The state of the pdisk involved.

%pdWwn

The worldwide name of the pdisk.

%prep0AlreadyCachedFiles

Specifies the number of files that are cached. These number of files are not read into cache because data is same between cache and home.

%prep0CompletedReads

Specifies the number of reads executed during a prefetch operation.

%prep0Data

Specifies the total data read from the home as part of a prefetch operation.

mmaddcallback

%prepopFailedReads

Specifies the number of files for which prefetch failed. Messages are logged to indicate the failure. However, there is no indication about the file names that failed to read.

%quorumNodes[.qualifier]

Specifies a comma-separated list of quorum nodes.

%quotaEventType

Specifies either the **blockQuotaExceeded** event or the **inodeQuotaExceeded** event. These events are related to soft quota limit being exceeded,

%quotaID

Specifies the numerical ID of the quota owner (UID, GID, or fileset ID).

%quotaOwnerName

Specifies the name of the quota owner (user name, group name, or fileset name).

%quotaType

Specifies the type of quota for quota-related events. Possible values are **USR**, **GRP**, or **FILESET**.

%reason

Specifies the reason for triggering the event. For the **preUnmount** and **unmount** events, the possible values are **normal** and **forced**. For the **preShutdown** and **shutdown** events, the possible values are **normal** and **abnormal**. For all other events, the value is **UNDEFINED**.

%requestType

Specifies the type of request to send to the target nodes.

%rgCount

The number of recovery groups involved.

%rgErr

A code from a recovery group, where 0 indicates no error.

%rgName

The name of the recovery group involved.

%rgReason

The reason string indicating why a recovery group callback was invoked.

%senseDataFormatted

Sense data for the specific fileset structure error in a formatted string output.

%senseDataHex

Sense data for the specific fileset structure error in Big endian hex output.

%snapshotID

Specifies the identifier of the new snapshot.

%snapshotName

Specifies the name of the new snapshot.

%softLimit

Specifies the soft limit of the block.

%storagePool

Specifies the storage pool name for space-related events.

%upNodes[.qualifier]

Specifies a comma-separated list of nodes that are currently up. Only nodes local to the given cluster are listed. Nodes which are in a remote cluster but have temporarily joined the cluster are not included.

%userName

Specifies the user name.

%waiterLength

Specifies the length of the waiter in seconds.

Variables recognized by IBM Spectrum Scale RAID are documented in *IBM Spectrum Scale RAID: Administration*.

Variables that represent node identifiers accept an optional qualifier that can be used to specify how the nodes are to be identified. When specifying one of these optional qualifiers, separate it from the variable with a period, as shown here:

variable.qualifier

The value for *qualifier* can be one of the following:

ip Specifies that GPFS should use the nodes' IP addresses.

name

Specifies that GPFS should use fully-qualified node names. This is the default.

shortName

Specifies that GPFS should strip the domain part of the node names.

Events and supported parameters

Table 5. Global events and supported parameters

Global event	Supported parameters
afmFilesetExpired Triggered when the contents of a fileset expire either as a result of the fileset being disconnected for the expiration timeout value or when the fileset is marked as expired using the AFM administration commands.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmFilesetUnexpired Triggered when the contents of a fileset become unexpired either as a result of the reconnection to home or when the fileset is marked as unexpired using the AFM administration commands.	%fsName %filesetName %pcacheEvent %homeServer %reason
nodeJoin Triggered when one or more nodes join the cluster.	%eventNode
nodeLeave Triggered when one or more nodes leave the cluster.	%eventNode
quorumReached Triggered when a quorum has been established in the GPFS cluster. This event is triggered only on the cluster manager, not on all the nodes in the cluster.	%quorumNodes
quorumLoss Triggered when quorum has been lost in the GPFS cluster.	N/A
quorumNodeJoin Triggered when one or more quorum nodes join the cluster.	%eventNode
quorumNodeLeave Triggered when one or more quorum nodes leave the cluster.	%eventNode

mmaddcallback

Table 5. Global events and supported parameters (continued)

Global event	Supported parameters
clusterManagerTakeOver Triggered when a new cluster manager node is elected. This happens when a cluster first starts up or when the current cluster manager fails or resigns and a new node takes over as cluster manager.	N/A

Table 6. Local events and supported parameters

Local event	Supported parameters
afmCmdRequeued Triggered during replication when messages are queued up again because of errors. These messages are retried after 15 minutes.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmFilesetUnmounted Triggered when the fileset is moved to an Unmounted state because NFS server is not reachable or remote cluster mount is not available for GPFS Native protocol.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmHomeConnected Triggered when a gateway node connects to the afmTarget of the fileset that it is serving. This event is local on gateway nodes.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmHomeDisconnected Triggered when a gateway node gets disconnected from the afmTarget of the fileset that it is serving. This event is local on gateway nodes.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmManualResyncComplete Triggered when a manual resync is completed.	<code>%fsName %filesetName %reason</code>
afmPrepopEnd Triggered when all the files specified by a prefetch operation have been completed. This event is local on a gateway node.	<code>%fsName %filesetName %prepopCompletedReads %prepopFailedReads %prepopAlreadyCachedFiles %prepData</code>
afmQueueDropped Triggered when replication encounters an issue that cannot be corrected. After the queue is dropped, next recovery action attempts to fix the error and continue to replicate.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmRecoveryFail Triggered when recovery fails. The recovery action is retried after 300 seconds. If recovery keeps failing, fileset is moved to a resync state if the fileset mode allows it.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>
afmRecoveryStart Triggered when AFM recovery starts. This event is local on gateway nodes.	<code>%fsName %filesetName %pcacheEvent %homeServer %reason</code>

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
afmRecoveryEnd Triggered when AFM recovery ends. This event is local on gateway nodes.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmRPOMiss Triggered when Recovery Point Objective (RPO) is missed on DR primary filesets, RPO Manager keeps retrying the snapshots. This event occurs when there is lot of data to replicate for the RPO snapshot to be taken or there is an error such as, deadlock and recovery keeps failing.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmFilesetCreate Triggered when an AFM fileset is created.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmFilesetLink Triggered when an AFM fileset is linked.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmFilesetChange Triggered when an AFM fileset is changed. If a fileset is renamed the new name is part of %reason.	%fsName %filesetName %pcacheEvent %homeServer %reason
afmFilesetUnlink Triggered when a AFM fileset is linked.	%fsName %filesetName %pcacheEvent %homeServer %reason
ccrFileChange Triggered when CCR fput operation takes place.	%ccrObjectName %ccrObjectVersion
ccrVarChange Triggered when CCR vput operation takes place.	%ccrObjectName %ccrObjectValue %ccrObjectVersion
daRebuildFailed The daRebuildFailed callback is generated when the spare space in a declustered array has been exhausted, and vdisk tracks involving damaged pdisks can no longer be rebuilt. The occurrence of this event indicates that fault tolerance in the declustered array has become degraded and that disk maintenance should be performed immediately. The daRemainingRedundancy parameter indicates how much fault tolerance remains in the declustered array.	%myNode %rgName %daName %daRemainingRedundancy
deadlockDetected Triggered when a node detects a potential deadlock. If the exit code of the registered callback for this event is 1, debug data will not be collected. See the <i>/usr/lpp/mmfs/samples/deadlockdetected.sample</i> file for an example of using the deadlockDetected event.	%eventName %myNode %waiterLength
deadlockOverload Triggered when an overload event occurs. The event is local to the node detecting the overload condition.	%eventName %nodeName

mmaddcallback

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
diskFailure Triggered on the file system manager when the status of a disk in a file system changes to down .	%eventName %diskName %fsName
filesetLimitExceeded Triggered when the file system manager detects that a fileset quota has been exceeded. This is a variation of softQuotaExceeded that applies only to fileset quotas. It exists only for compatibility (and may be deleted in a future version); therefore, using softQuotaExceeded is recommended instead.	%filesetName %fsName %filesetSize %softLimit %hardLimit %inodeUsage %inodeQuota %inodeLimit %quotaEventType
fsstruct Triggered when the file system manager detects a file system structure (FS Struct) error. For more information about FS Struct errors, see “mmfsck command” on page 331 and the following topics in <i>IBM Spectrum Scale: Problem Determination Guide</i> : <ul style="list-style-type: none">• MMFS_FSSTRUCT• Reliability, Availability, and Serviceability (RAS) events• Information to collect before contacting the IBM Support Center	%fsName %fsErr %senseDataFormatted %senseDataHex
lowDiskSpace Triggered when the file system manager detects that disk space usage has reached the high occupancy threshold that is specified in the current policy rule. The event is generated every two minutes until the condition no longer exists. For more information, see the topic <i>Using thresholds with external pools</i> in the <i>IBM Spectrum Scale: Administration Guide</i> .	%storagePool %fsName
noDiskSpace Triggered when the file system encounters a disk, or storage pool that has run out of space or an inodespace has run out of inodes. An inode space can be an entire file system or an independent fileset. Use the noSpaceEventInterval configuration attribute of the mmchconfig command to control the time interval between two noDiskSpace events. The default value is 120 seconds. When a storage pool runs out of disk space, %reason is “diskspace”, %storagePool is the name of the pool that ran out of disk space, and %filesetName is “UNDEFINED”. When a fileset runs out of inode space, %reason is “inodespace”, %filesetName is the name of the independent fileset that owns the affected inode space, and %storagePool is “UNDEFINED”.	%storagePool %fsName %reason %filesetName

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
nsdCksumMismatch The nsdCksumMismatch callback is generated whenever transmission of vdisk data by the NSD network layer fails to verify the data checksum. This can indicate problems in the network between the GPFS client node and a recovery group server. The first error between a given client and server generates the callback; subsequent callbacks are generated for each ckReportingInterval occurrence.	%myNode %ckRole %ckOtherNode %ckNSD %ckReason %ckStartSector %ckDataLen %ckErrorCountClient %ckErrorCountServer %ckErrorCountNSD %ckReportingInterval
pdFailed The pdFailed callback is generated whenever a pdisk in a recovery group is marked as dead, missing, failed, or readonly.	%myNode %rgName %daName %pdName %pdLocation %pdFru %pdWwn %pdState
pdPathDown The pdPathDown callback is generated whenever one of the block device paths to a pdisk disappears or becomes inoperative. The occurrence of this event can indicate connectivity problems with the JBOD array in which the pdisk resides.	%myNode %rgName %daName %pdName %pdLocation %pdFru %pdWwn %pdPath
pdReplacePdisk The pdReplacePdisk callback is generated whenever a pdisk is marked for replacement according to the replace threshold setting of the declustered array in which it resides.	%myNode %rgName %daName %pdName %pdLocation %pdFru %pdWwn %pdState %pdPriority
pdRecovered The pdRecovered callback is generated whenever a missing pdisk is rediscovered. The following parameters are available to this callback: %myNode, %rgName, %daName, %pdName, %pdLocation, %pdFru, and %pdWwn.	%myNode %rgName %daName %pdName %pdLocation %pdFru %pdWwn
preMount, preUnmount, mount, unmount These events are triggered when a file system is about to be mounted or unmounted or has been mounted or unmounted successfully. These events are generated for explicit mount and unmount commands, a remount after GPFS recovery and a forced unmount when GPFS panics and shuts down.	%fsName %reason
preRGRelinquish The preRGRelinquish callback is invoked on a recovery group server prior to relinquishing service of recovery groups. The rgName parameter may be passed into the callback as the keyword value _ALL_ , indicating that the recovery group server is about to relinquish service for all recovery groups it is serving; the rgCount parameter will be equal to the number of recovery groups being relinquished. Additionally, the callback will be invoked with the rgName of each individual recovery group and an rgCount of 1 whenever the server relinquishes serving recovery group rgName .	%myNode %rgName %rgErr %rgCount %rgReason

mmaddcallback

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
preRGTakeover The preRGTakeover callback is invoked on a recovery group server prior to attempting to open and serve recovery groups. The rgName parameter may be passed into the callback as the keyword value _ALL_ , indicating that the recovery group server is about to open multiple recovery groups; this is typically at server startup, and the parameter rgCount will be equal to the number of recovery groups being processed. Additionally, the callback will be invoked with the rgName of each individual recovery group and an rgCount of 1 whenever the server checks to determine whether it should open and serve recovery group rgName .	%myNode %rgName %rgErr %rgCount %rgReason
preShutdown Triggered when GPFS detects a failure and is about to shut down.	%reason
preStartup Triggered after the GPFS daemon completes its internal initialization and joins the cluster, but before the node runs recovery for any file systems that were already mounted, and before the node starts accepting user initiated sessions.	N/A
postRGRelinquish The postRGRelinquish callback is invoked on a recovery group server after it has relinquished serving recovery groups. If multiple recovery groups have been relinquished, the callback will be invoked with rgName keyword _ALL_ and an rgCount equal to the total number of involved recovery groups. The callback will also be triggered for each individual recovery group.	%myNode %rgName %rgErr %rgCount %rgReason
postRGTakeover The postRGTakeover callback is invoked on a recovery group server after it has checked, attempted, or begun to serve a recovery group. If multiple recovery groups have been taken over, the callback will be invoked with rgName keyword _ALL_ and an rgCount equal to the total number of involved recovery groups. The callback will also be triggered for each individual recovery group.	%myNode %rgName %rgErr %rgCount %rgReason
rgOpenFailed The rgOpenFailed callback will be invoked on a recovery group server when it fails to open a recovery group that it is attempting to serve. This may be due to loss of connectivity to some or all of the disks in the recovery group; the rgReason string will indicate why the recovery group could not be opened.	%myNode %rgName %rgErr %rgReason

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
rgPanic The rgPanic callback will be invoked on a recovery group server when it is no longer able to continue serving a recovery group. This may be due to loss of connectivity to some or all of the disks in the recovery group; the rgReason string will indicate why the recovery group can no longer be served.	%myNode %rgName %rgErr %rgReason
sendRequestToNodes Triggered when a node sends a request for collecting expel-related debug data. For this event, the %requestType is requestExpelData .	%eventName %requestType %nodeNames
shutdown Triggered when GPFS completes the shutdown.	%reason
snapshotCreated Triggered after a snapshot is created, and run before the file system is resumed. This event helps correlate the timing of DMAPI events with the creation of a snapshot. GPFS must wait for snapshotCreated to exit before it resumes the file system, so the ordering of DMAPI events and snapshot creation is known. The %filesetName is the name of the fileset whose snapshot was created. For file system level snapshots that affect all filesets, %filesetName is set to global .	%snapshotID %snapshotName %fsName %filesetName
softQuotaExceeded Triggered when the file system manager detects that a soft quota limit (for either files or blocks) has been exceeded. This event is triggered only on the file system manager. Therefore, this event must be handled on all manager nodes.	%fsName %filesetName %quotaId %quotaType %quotaOwnerName %blockUsage %blockQuota %blockLimit %filesUsage %filesQuota %filesLimit
startup Triggered after a successful GPFS startup before the node is ready for user initiated sessions. After this event is triggered GPFS proceeds to finish starting including mounting all file systems defined to mount on startup.	N/A
tiebreakerCheck Triggered when the cluster manager detects a lease timeout on a quorum node before GPFS runs the algorithm that decides if the node will remain in the cluster. This event is generated only in configurations that use tiebreaker disks. Note: Before you add or delete the tiebreakerCheck event, you must stop the GPFS daemon on all the nodes in the cluster.	N/A
traceConfigChanged Triggered when GPFS tracing configuration is changed.	N/A

mmaddcallback

Table 6. Local events and supported parameters (continued)

Local event	Supported parameters
usageUnderSoftQuota Triggered when the file system manager detects that quota usage has dropped below soft limits and grace time is reset.	%fsName %filesetName %fsName %quotaId %quotaType %quotaOwnerName %blockUsage %blockQuota %blockLimit %filesUsage %filesQuota %filesLimit

Options

-S *Filename* | **--spec-file** *Filename*

Specifies a file with multiple callback definitions, one per line. The first token on each line must be the callback identifier.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmaddcallback** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

1. To register command **/tmp/myScript** to run after GPFS startup, issue this command:

```
mmaddcallback test1 --command=/tmp/myScript --event startup
```

The system displays information similar to:

```
mmaddcallback: Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.
```

2. To register a callback on the NFS servers to export or to unexport a particular file system after it has been mounted or before it has been unmounted, issue this command:

```
mmaddcallback NFSExport --command /usr/local/bin/NFSExport --event mount,preUnmount -N nfsserver1, nfsserver2 --parms "%eventName %fsName" --parms "%eventName %fsName"
```

The system displays information similar to:

```
mmaddcallback: 6027-1371 Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.
```

See also

- “**mmdelcallback** command” on page 286
- “**mmlscallback** command” on page 395

Location

`/usr/lpp/mmfs/bin`

mmadddisk command

Adds disks to a GPFS file system.

Synopsis

```
| mmadddisk Device {"DiskDesc[;DiskDesc...]" | -F StanzaFile} [-a] [-r]
| [-v {yes | no}] [-N {Node[,Node...] | NodeFile | NodeClass}]
| [--qos QOSClass]
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmadddisk** command to add disks to a GPFS file system. When the **-r** flag is specified, the command rebalances an existing file system after it adds the disks. The command does not require the file system to be mounted. The file system can be in use.

The actual number of disks in your file system might be constrained by products other than GPFS that you installed. See to the individual product documentation.

To add disks to a GPFS file system, first decide which of the following two tasks you want to perform:

1. Create new disks with the **mmcrnsd** command.

In this case, you must also decide whether to create a new set of NSD and pools stanzas or use the rewritten NSD and pool stanzas that the **mmcrnsd** command produces. In a rewritten file, the disk usage, failure group, and storage pool values are the same as the values that are specified in the **mmcrnsd** command.

2. Select disks no longer in use in any file system. To display the disks that are not in use, run the following command:

```
mm1snsd -F
```

Note: If **mmadddisk** fails with a NO_SPACE error, try one of the following actions:

- Rebalance the file system.
- Run the command **mmfsck -y** to deallocate unreferenced subblocks.
- Create a pool with larger disks and move data from the old pool to the new one.

Parameters

Device

The device name of the file system to which the disks are added. File system names need not be fully qualified. **fs0** is as acceptable as **/dev/fs0**.

This parameter must be first.

DiskDesc

A descriptor for each disk to be added. Each descriptor is delimited by a semicolon (;) and the entire list must be enclosed in quotation marks (' or "). The use of disk descriptors is discouraged.

-F StanzaFile

Specifies a file that contains the NSD stanzas and pool stanzas for the disks to be added to the file system.

NSD stanzas have this format:

```
%nsd:
  nsd=NsName
  usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}
```

mmadddisk

failureGroup=FailureGroup
pool=StoragePool
servers=ServerList
device=DiskName

where:

nsd=NsdName

The name of an NSD previously created by the **mmcrnsd** command. For a list of available disks, run the **mmlsnsd -F** command. This clause is mandatory for the **mmadddisk** command.

usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}

Specifies the type of data to be stored on the disk:

dataAndMetadata

Indicates that the disk contains both data and metadata. This value is the default for disks in the system pool.

dataOnly

Indicates that the disk contains data and does not contain metadata. This value is the default for disks in storage pools other than the system pool.

metadataOnly

Indicates that the disk contains metadata and does not contain data.

descOnly

Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster recovery configurations. For more information, see the *IBM Spectrum Scale: Administration Guide* and search for “Synchronous mirroring utilizing GPFS replication”

failureGroup=FailureGroup

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

Disk that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

pool=StoragePool

Specifies the storage pool to which the disk is to be assigned. If this name is not provided, the default is **system**.

Only the system storage pool can contain **metadataOnly**, **dataAndMetadata**, or **descOnly** disks. Disks in other storage pools must be **dataOnly**.

servers=ServerList

A comma-separated list of NSD server nodes. This clause is ignored by the **mmadddisk** command.

device=DiskName

The block device name of the underlying disk device. This clause is ignored by the **mmadddisk** command.

Note: An NSD belonging to a tiebreaker disk is not allowed to be added to a file system if NSD format conversion is required.

Pool stanzas have this format:

```
%pool:  
  pool=StoragePoolName  
  blockSize=BlockSize  
  usage={dataOnly | metadataOnly | dataAndMetadata}  
  layoutMap={scatter | cluster}  
  allowWriteAffinity={yes | no}  
  writeAffinityDepth={0 | 1 | 2}  
  blockGroupFactor=BlockGroupFactor
```

where:

pool=StoragePoolName

Is the name of a storage pool.

blockSize=BlockSize

Specifies the block size of the disks in the storage pool.

usage={dataOnly | metadataOnly | dataAndMetadata}

Specifies the type of data to be stored in the storage pool:

dataAndMetadata

Indicates that the disks in the storage pool contain both data and metadata. This is the default for disks in the system pool.

dataOnly

Indicates that the disks contain data and do not contain metadata. This is the default for disks in storage pools other than the system pool.

metadataOnly

Indicates that the disks contain metadata and do not contain data.

layoutMap={scatter | cluster}

Specifies the block allocation map type. When allocating blocks for a given file, GPFS first uses a round-robin algorithm to spread the data across all disks in the storage pool. After a disk is selected, the location of the data block on the disk is determined by the block allocation map type. If **cluster** is specified, GPFS attempts to allocate blocks in clusters. Blocks that belong to a particular file are kept adjacent to each other within each cluster. If **scatter** is specified, the location of the block is chosen randomly.

The **cluster** allocation method may provide better disk performance for some disk subsystems in relatively small installations. The benefits of clustered block allocation diminish when the number of nodes in the cluster or the number of disks in a file system increases, or when the file system's free space becomes fragmented. The **cluster** allocation method is the default for GPFS clusters with eight or fewer nodes and for file systems with eight or fewer disks.

The **scatter** allocation method provides more consistent file system performance by averaging out performance variations due to block location (for many disk subsystems, the location of the data relative to the disk edge has a substantial effect on performance). This allocation method is appropriate in most cases and is the default for GPFS clusters with more than eight nodes or file systems with more than eight disks.

The block allocation map type cannot be changed after the storage pool has been created.

allowWriteAffinity={yes | no}

Indicates whether the File Placement Optimizer (FPO) feature is to be enabled for the storage pool. For more information on FPO, see the *File Placement Optimizer* section in the *IBM Spectrum Scale: Administration Guide*.

writeAffinityDepth={0 | 1 | 2}

Specifies the allocation policy to be used by the node writing the data.

mmadddisk

A write affinity depth of 0 indicates that each replica is to be striped across the disks in a cyclical fashion with the restriction that no two disks are in the same failure group. By default, the unit of striping is a block; however, if the block group factor is specified in order to exploit chunks, the unit of striping is a chunk.

A write affinity depth of 1 indicates that the first copy is written to the writer node. The second copy is written to a different rack. The third copy is written to the same rack as the second copy, but on a different half (which can be composed of several nodes).

A write affinity depth of 2 indicates that the first copy is written to the writer node. The second copy is written to the same rack as the first copy, but on a different half (which can be composed of several nodes). The target node is determined by a hash value on the fileset ID of the file, or it is chosen randomly if the file does not belong to any fileset. The third copy is striped across the disks in a cyclical fashion with the restriction that no two disks are in the same failure group. The following conditions must be met while using a write affinity depth of 2 to get evenly allocated space in all disks:

1. The configuration in disk number, disk size, and node number for each rack must be similar.
2. The number of nodes must be the same in the bottom half and the top half of each rack.

This behavior can be altered on an individual file basis by using the **--write-affinity-failure-group** option of the **mmchattr** command.

This parameter is ignored if write affinity is disabled for the storage pool.

blockGroupFactor=BlockGroupFactor

Specifies how many file system blocks are laid out sequentially on disk to behave like a single large block. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

See *File Placement Optimizer in IBM Spectrum Scale: Administration Guide*.

-a Specifies asynchronous processing. If this flag is specified, the **mmadddisk** command returns after the file system descriptor is updated and the rebalancing scan is started; it does not wait for rebalancing to finish. If no rebalancing is requested (the **-r** flag not specified), this option has no effect.

| **-r** Rebalance all existing files in the file system to use the new disks.

| **Note:** Rebalancing of files is an I/O intensive and time-consuming operation, and is important only for file systems with large files that are mostly invariant. In many cases, normal file update and creation will rebalance the file system over time.

-v {yes | no}

Verify that specified disks do not belong to an existing file system. The default is **-v yes**. Specify **-v no** only when you want to reuse disks that are no longer needed for an existing file system. If the command is interrupted for any reason, use the **-v no** option on the next invocation of the command.

Important: Using **-v no** on a disk that already belongs to a file system corrupts that file system. This problem is not detected until the next time that file system is mounted.

-N {Node[,Node...] | NodeFile | NodeClass}

Specifies the nodes that are to participate in the restriping of the file system after the specified disks are available for use by GPFS. This parameter must be used with the **-r** option. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands in the IBM Spectrum Scale: Administration Guide*.

--qos QoSClass

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default

to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “[mmchqos command](#)” on page 206. Specify one of the following QoS classes:

maintenance

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

- other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

Exit status

- 0** Successful completion.

nonzero

- A failure occurred.

Security

You must have root authority to run the **mmaddirdisk** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

1. Assume that the file `./newNSDstanza` contains the following NSD stanza:

```
%nsd: nsd=gpfsl0nsd
      servers=k148n07,k148n06
      usage=dataOnly
      failureGroup=5
      pool=pool2
```

To add the disk that is defined in this stanza, run the following command:

```
mmaddirdisk fs1 -F ./newNSDstanza -r
```

The command displays information like the following example:

```
GPFS: 6027-531 The following disks of fs1 will be formatted on node
k148n07.kgn.ibm.com:
  gpfsl0nsd: size 2202 MB
Extending Allocation Map
Creating Allocation Map for storage pool 'pool2'
  75 % complete on Thu Feb 16 13:57:52 2006
  100 % complete on Thu Feb 16 13:57:54 2006
Flushing Allocation Map for storage pool 'pool2'
GPFS: 6027-535 Disks up to size 24 GB can be added to storage
pool pool2.
Checking allocation map for storage pool system
  62 % complete on Thu Feb 16 13:58:03 2006
  100 % complete on Thu Feb 16 13:58:06 2006
Checking allocation map for storage pool pool1
  62 % complete on Thu Feb 16 13:58:11 2006
  100 % complete on Thu Feb 16 13:58:14 2006
Checking allocation map for storage pool pool2
  63 % complete on Thu Feb 16 13:58:19 2006
```

mmadddisk

```
100 % complete on Thu Feb 16 13:58:22 2006
GPFS: 6027-1503 Completed adding disks to file system fs1.
mmadddisk: 6027-1371 Propagating the cluster configuration data to all
    affected nodes. This is an asynchronous process.
Restriping fs1 ...
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
68 % complete on Thu Feb 16 13:59:06 2006
100 % complete on Thu Feb 16 13:59:07 2006
GPFS: 6027-552 Scan completed successfully.
Done
```

See also

- “[mmchdisk command](#)” on page 161
- “[mmcrnsd command](#)” on page 262
- “[mmdeldisk command](#)” on page 287
- “[mmlsdisk command](#)” on page 402
- “[mmlsnsd command](#)” on page 423
- “[mmlspool command](#)” on page 428

Location

/usr/lpp/mmfs/bin

mmaddnode command

Adds nodes to a GPFS cluster.

Synopsis

mmaddnode -N {NodeDesc[,NodeDesc...]} | NodeFile}

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmaddnode** command to add nodes to an existing GPFS cluster. On each new node, a mount point directory and character mode device is created for each GPFS file system.

Follow these rules when adding nodes to a GPFS cluster:

- You may issue the command only from a node that already belongs to the GPFS cluster.
- While a node may mount file systems from multiple clusters, the node itself may only be added to a single cluster using the **mmcrcluster** or **mmaddnode** command.
- The nodes must be available for the command to be successful. If any of the nodes listed are not available when the command is issued, a message listing those nodes is displayed. You must correct the problem on each node and reissue the command to add those nodes.
- After the nodes are added to the cluster, use the **mmchlicense** command to designate appropriate GPFS licenses to the new nodes.

Parameters

-N NodeDesc[,NodeDesc...] | **NodeFile**

Specifies node descriptors, which provide information about nodes to be added to the cluster.

NodeFile

Specifies a file containing a list of node descriptors, one per line, to be added to the cluster.

NodeDesc[,NodeDesc...]

Specifies the list of nodes and node designations to be added to the GPFS cluster. Node descriptors are defined as:

NodeName:NodeDesignations:AdminNodeName

where:

NodeName

Specifies the host name or IP address of the node for GPFS daemon-to-daemon communication.

The host name or IP address must refer to the communication adapter over which the GPFS daemons communicate. Aliased interfaces are not allowed. Use the original address or a name that is resolved by the **host** command to that original address. You can specify a node using any of these forms:

- Short host name (for example, h135n01)
- Long, fully-qualified, host name (for example, h135n01.ibm.com)
- IP address (for example, 7.111.12.102). IPv6 addresses must be enclosed in brackets (for example, [2001:192::192:168:115:124]).

Regardless of which form you use, GPFS will resolve the input to a host name and an IP address and will store these in its configuration files. It is expected that those values will not change while the node belongs to the cluster.

mmaddnode

NodeDesignations

An optional, "-" separated list of node roles:

- **manager** | **client** – Indicates whether a node is part of the node pool from which file system managers and token managers can be selected. The default is **client**.
- **quorum** | **nonquorum** – Indicates whether a node is counted as a quorum node. The default is **nonquorum**.

Note: If you are designating a new node as a quorum node, and **adminMode central** is in effect for the cluster, GPFS must be down on all nodes in the cluster. Alternatively, you may choose to add the new nodes as **nonquorum** and once GPFS has been successfully started on the new nodes, you can change their designation to **quorum** using the **mmchnode** command.

AdminNodeName

Specifies an optional field that consists of a node interface name to be used by the administration commands to communicate between nodes. If *AdminNodeName* is not specified, the *NodeName* value is used.

Note: *AdminNodeName* must be a resolvable network host name. For more information, see the topic *GPFS node adapter interface names* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

You must provide a *NodeDesc* for each node to be added to the GPFS cluster.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmaddnode** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Examples

To add nodes **k164n06** and **k164n07** as **quorum** nodes, designating **k164n06** to be available as a **manager** node, issue this command:

```
mmaddnode -N k164n06:quorum-manager,k164n07:quorum
```

To confirm the addition, issue this command:

```
mmlscluster
```

The system displays information similar to:

```
GPFS cluster information
=====
GPFS cluster name:      cluster1.kgn.ibm.com
GPFS cluster id:        680681562214606028
GPFS UID domain:       cluster1.kgn.ibm.com
Remote shell command:   /usr/bin/ssh
Remote file copy command: /usr/bin/scp
Repository type:        server-based
```

GPFS cluster configuration servers:

Primary server: k164n07.kgn.ibm.com
Secondary server: k164n04.kgn.ibm.com

Node	Daemon node name	IP address	Admin node name	Designation
1	k164n04.kgn.ibm.com	198.117.68.68	k164n04.kgn.ibm.com	quorum
2	k164n07.kgn.ibm.com	198.117.68.71	k164n07.kgn.ibm.com	quorum
3	k164n06.kgn.ibm.com	198.117.68.70	k164n06.kgn.ibm.com	quorum-manager

See also

- “mmchconfig command” on page 132
- “mmcrcluster command” on page 239
- “mmchcluster command” on page 128
- “mmdelnnode command” on page 297
- “mmlscluster command” on page 397

Location

/usr/lpp/mmfs/bin

mmadquery

mmadquery command

Queries and validates Active Directory (AD) server settings.

Synopsis

```
mmadquery list {user | uids | gids | groups | dc | trusts | idrange} [Options]
```

or

```
mmadquery check {uids | gids | idrange} [Options]
```

or

```
mmadquery stats {user | uids}
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmadquery** command to query an AD Server for users, groups, user IDs, group IDs, known domain controller and trusts, and to run consistency checks.

Parameters

user

Queries and lists the defined users.

uids

Queries and lists the defined users with user IDs and group IDs.

gids

Queries and lists the defined groups with group IDs.

groups

Queries and lists the defined groups.

dc

Queries and lists the defined domain controllers.

trusts

Queries and lists the defined trusts.

idrange

Queries and lists the ID range used by a given AD server.

Options

--server SERVER

Specifies the IP address of the AD server you want to query. If you do not specify a server, **mmadquery** attempts to get the AD server from the /etc/resolv.conf file (nameserver).

Note: This option should be used along with the domain option, which is provided in the following section.

--domain DOMAIN

Specifies the Windows domain. If you do not specify a domain, **mmadquery** uses nslookup to determine the domain based on the server.

Note: This option should be used along with the server option.

```

--user USER
    Specifies the AD user used to run the LDAP query against the AD server. The default is Administrator.

--pwd-file File
    Specifies the file that contains a password to use for authentication.

--filter FILTER
|   Specifies a search phrase to limit the number of LDAP objects, thus is applied only to first
|   column of output. Every LDAP object beginning with the search phrase is queried.

--CSV
|   Shows output in machine parseable (CSV) format.

--debug or -d
|   Shows debugging information

--basedn or -b
|   Includes basedn for LDAP objects queried in query output. This option is not supported when
|   querying idrange or running a 'stats' query.

--traverse
|   Traverses all known domains and provide query output for all domains that are detected.

--long or -L
|   Indicates that you want to see more details. For more information, see Level of query detail
|   below. This option is not supported for the "stats" queries.

```

Level of query detail

Table 7. Query details by type

Query	Additional content
User	Group membership
DC	Operating system
UIDs	GID, Primary Group ID
Trusts	DC

Exit status

- 0** No errors found.
- 1** No arguments specified.
- 10** Failed a check.
- 11** Unable to determine the AD server to check.
- 12** Unable to determine the domain.
- 13** Failed to construct a basedn for an LDAP query.
- 99** Access to the AD server failed, can be incorrect password, user, or domain.

Security

You must have root authority to run the mmadquery command. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

mmadquery

Examples

1. To show a list of users for the AD server, run this command:

```
mmadquery list user --pwd-file /tmp/mmadquery.cfg
```

The system displays information similar to:

```
USER from server 9.155.106.234 (domain subdom1.mzdom.com)
User
-----
Administrator
Guest
krbtgt
MZDOM$
aduser1
aduser2
Taduser3
```

2. To show a list of groups for the AD server, run this command:

```
mmadquery list groups --pwd-file /tmp/mmadquery.cfg
```

The system displays information similar to:

```
GROUPS from server 9.155.106.234 (domain subdom1.mzdom.com)
Group
-----
Domain Computers
Cert Publishers
Domain Users
Domain Guests
RAS and IAS Servers
Domain Admins
Schema Admins
Enterprise Admins
Group Policy Creator Owners
Allowed RODC Password Replication Group
Denied RODC Password Replication Group
Enterprise Read-only Domain Controllers
Domain Controllers
Read-only Domain Controllers
DnsAdmins
DnsUpdateProxy
UNIXGRP
unmapped group
bla
```

3. To check user IDs against locally defined ID mapping range, issue the following command:

```
mmadquery check uids --pwd-file /tmp/mmadquery.cfg -L
```

The system displays information similar to:

UIDS from server 9.155.106.234 (domain subdom1.mzdom.com)					
User	SID				ID
Guest	S-1-5-21-2808815044-4164012579-2832416960-501	-	-	-	514
SUBDOM1\$	S-1-5-21-2808815044-4164012579-2832416960-1103	-	-	-	513
Administrator	S-1-5-21-2808815044-4164012579-2832416960-500	-	-	-	513
krbtgt	S-1-5-21-2808815044-4164012579-2832416960-502	-	-	-	513
User 1	S-1-5-21-2808815044-4164012579-2832416960-1107	-	-	-	513
aduser1	S-1-5-21-2808815044-4164012579-2832416960-1601	aduser1	20000007	20000008	513
User 2	S-1-5-21-2808815044-4164012579-2832416960-1110	aduser	10001	20000009	513

WARNING: UID of user User 2 outside id mapping range 'mzdom'.

4. To show a list of users with group membership by domain, run this command:

```
mmadquery list user -L --pwd-file /tmp/mmadquery.cfg --traverse
```

The system displays information similar to:

USER from server 9.155.106.232 (domain mzdom.com)		Groups
User		Groups
Guest		Guests
SUBDOM1\$		
Administrator	Group Policy Creator Owners,Enterprise Admins,Schema Admins,Domain Admins,Administrators	
krbtgt		Denied RODC Password Replication Group

```

aduser1                               Administrators
aduser2                               bla,unmapped group
aduser3
aduser4

USER from server 9.155.106.234 (domain subdom1.mzdom.com)
User                                Groups
-----
Administrator Group Policy Creator Owners,Domain Admins,Administrators
Guest                               Guests
krbtgt      Denied RODC Password Replication Group Administrators
MZDOM$ 
aduser1
aduser2
aduser3
aduser4

```

5. To show the number of users by group and domain, run this command:

```
mmadquery stats user -L --pwd-file /tmp/mmadquery.cfg --traverse
```

The system displays information similar to:

```
USER from server 9.155.106.232 (domain mzdom.com)
```

	Group Count
TOTAL	7
Guests	1
Group Policy Creator Owners	1
Enterprise Admins	1
Schema Admins	1
Domain Admins	1
Administrators	2
Denied RODC Password Replication Group	1
bla	1
unmapped group	1

```
USER from server 198.51.100.13 (domain subdom1.mzdom.com)
```

	Group Count
TOTAL	7
Group Policy Creator Owners	1
Domain Admins	1
Administrators	2
Guests	1
Denied RODC Password Replication Group	1

6. To show a list of the number of unmapped users, run this command:

```
mmadquery stats uids --pwd-file /tmp/mmadquery.cfg
```

The system displays information similar to:

```
UIDS from server 9.155.106.232 (domain mzdom.com)
```

Group	Count
TOTAL	7
MAPPED	2
UN-MAPPED	5

7. To check group IDs against locally defined ID map, run this command:

```
mmadquery check gids -L --pwd-file /tmp/mmadquery.cfg
```

The system displays information similar to:

```
GIDS from server 9.155.106.232 (domain w2k8r2-dom02.mzdom.com)
```

```

GIDS from server 9.155.106.232 (domain w2k8r2-dom02.mzdom.com)
Group                                SID   UID   UIDNumber   GIDNumber
-----
Domain Computers S-1-5-21-2808815044-4164012579-2832416960-515   -   -   -
Cert Publishers S-1-5-21-2808815044-4164012579-2832416960-517   -   -   -
Domain Users S-1-5-21-2808815044-4164012579-2832416960-513   -   -   20000008 -
Domain Guests S-1-5-21-2808815044-4164012579-2832416960-514   -   -   -
RAS and IAS Servers S-1-5-21-2808815044-4164012579-2832416960-553 -   -   -
Domain Admins S-1-5-21-2808815044-4164012579-2832416960-512 -   -   -
Schema Admins S-1-5-21-2808815044-4164012579-2832416960-518 -   -   -
Enterprise Admins S-1-5-21-2808815044-4164012579-2832416960-519 -   -   -
Group Policy Creator Owners S-1-5-21-2808815044-4164012579-2832416960-520 -   -   -
Allowed RODC Password Replication Group S-1-5-21-2808815044-4164012579-2832416960-571 -   -   -
Denied RODC Password Replication Group S-1-5-21-2808815044-4164012579-2832416960-57 -   -   -
Enterprise Read-only Domain Controllers S-1-5-21-2808815044-4164012579-2832416960-498 -   -   -
Domain Controllers S-1-5-21-2808815044-4164012579-2832416960-516 -   -   -
Read-only Domain Controllers S-1-5-21-2808815044-4164012579-2832416960-521 -   -   -
DnsAdmins S-1-5-21-2808815044-4164012579-2832416960-1101 -   -   -
DnsUpdateProxy S-1-5-21-2808815044-4164012579-2832416960-1102 -   -   -

```

mmapquery

```
UNIXGRP S-1-5-21-2808815044-4164012579-2832416960-1104 - - 200002222 -
unmapped group S-1-5-21-2808815044-4164012579-2832416960-1603 - - -
bia S-1-5-21-2808815044-4164012579-2832416960-1604 - - -
```

-WARNING: GID of group 'UNIXGRP' outside id mapping range 'mzdom'.

8. To show a list of domain controllers, run the following command:

```
mmapquery list dc L --pwd-file /tmp/mmapquery.cfg
```

The system displays information similar to:

DC	Hostname	Operating System
WW2K8R2-DOM03	w2k8r2-dom03.mzdom.com	Windows Server 2008 R2 Standard
WW2K8R2-DOM02	w2k8r2-dom02.mzdom.com	Windows Server 2008 R2 Standard

9. To show a list of trusts, run the following command:

```
mmapquery list trusts --pwd-file /tmp/mmapquery.cfg
```

The system displays information similar to:

DC	Trust Type
subdom1.mzdom.com	Within Forest bi-directional
w2k12dom.com	Forest Transitive outbound

10. To show a list of ID ranges and to check whether any IDs on the Ad server are outside of the locally defined ID range, run this command:

```
mmapquery check idrange --pwd-file /tmp/mmapquery.cfg
```

The system displays information similar to:

IDRange
msdom.com 10001-200000000 20000000-25999999

WARNING: IDs from domain 'mzdom.com' are outside locally defined id mapping range 'mzdom'.

11. To show a list of ID ranges by domain, run this command:

```
mmapquery list idrange --pwd-file /tmp/mmapquery.cfg -L --traverse
```

The system displays information similar to:

IDRange
mzdom.com 10001--260000009 10000000-29999999

IDRange
subdom1.mzdom.com 200000001-26000010 10000000-29999999

Location

/usr/lpp/mmfs/bin

mmafmconfig command

Can be used to manage home caching behavior and mapping of gateways and home NFS exported servers.

Synopsis

You can use the **mmafmconfig** command to -

- set up or update mapping for parallel data transfers by using add, update, or delete options.
- enable or disable extended attributes/sparse file support from the AFM cache.

mmafmconfig {add | update} MapName --export-map ExportServerMap

or

mmafmconfig delete {MapName | all}

or

mmafmconfig show [MapName | all]

or

mmafmconfig {enable | disable} ExportPath

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

You can use this command to configure a home cluster for enabling support of extended attributes /sparse files on AFM cache filesets pointing to this home. You must run the **mmafmconfig enable** command on the home path. Running this command creates the .afm directory which contains the control-enabled, directio .afmctl file. The **mmafmconfig disable** command removes the .afm directory from the home path and subsequently, the cache does not support sparse files and files with extended attributes.

You can also use the **mmafmconfig** command with add, update, delete, or show options on the cache site to manage mapping of gateway node with home NFS servers for parallel data transfers.

Ensure that you run **mmafmconfig enable** at home fileset before the cache fileset is linked. As an administrator, if you run this command at home after linking fileset at cache, unlink and relink fileset at cache.

Parameters

MapName

Specifies the name that uniquely identifies the mapping of the gateway nodes with the home NFS exported servers.

--export-map *ExportServerMap*

Specifies a comma-separated list of pairs of home NFS exported server nodes (*ExportServer*) and gateway nodes (*GatewayNode*), in the following format:

[*ExportServer/GatewayNode*] [, *ExportServer/GatewayNode*] [, ...]

where:

ExportServer

Is the IP address or host name of a member node in the home cluster *MapName*.

mmafmcconfig

GatewayNode

Specifies a gateway node in the cache cluster (the cluster where the command is issued).

enable

Enables extended attributes or sparse files functions on the AFM cache. Run at the home cluster only.

disable

Disables extended attributes or sparse files functions on the AFM cache. Run at the home cluster only.

ExportPath

Specifies the root of the home exported directory for enabling or disabling the AFM features.

add

Sets up maps for parallel data transfers. Run at cache only.

delete

Deletes maps for parallel data transfers. Run at cache only.

update

Updates maps for parallel data transfers. Run at cache only.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmafmcconfig** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Example

The following is an example of a mapping for NFS targets, adding all gateway nodes and NFS servers, and using these mapping for creating AFM filesets. The four cache gateway nodes assumed for this example are -hs22n18, hs22n19, hs22n20, and hs22n21. The gateway nodes are mapped to two home NFS servers js22n01 with an IP address 192.168.200.11 and js22n02 with an IP address 192.168.200.12.

1. Issue the following command:

```
#mmafmcconfig add mapping1 --export-map js22n01/hs22n18,js22n02/hs22n19,  
js22n01/hs22n20,js22n02/hs22n21
```

The system displays output similar to:

```
mmafmcconfig: Command successfully completed.
```

2. Issue the following command:

```
# mmafmcconfig show
```

The system displays output similar to:

```
Map name: mapping1
```

```
Export server map: 192.168.200.11/hs22n18.gpfs.net,192.168.200.12/  
hs22n19.gpfs.net,92.168.200.11/hs22n20.gpfs.net,192.168.200.12/hs22n21.gpfs.net
```

3. Issue the following commands to create filesets:

```
#Create filesets using these mappings:
```

```
mmcrfileset gpfs1 sw1 --inode-space new -p afmmode=sw,afmtarget=nfs://mapping1/gpfs/gpfs2/
swhome
mmcrfileset gpfs1 ro1 --inode-space new -p afmmode=ro,afmtarget=nfs://mapping1/gpfs/gpfs2/
swhome
```

See also

- “mmafmctl command” on page 40
- “mmafmlocal command” on page 54
- “mmchconfig command” on page 132
- “mmchfileset command” on page 173
- “mmchfs command” on page 179
- “mmcrfileset command” on page 244
- “mmcrfs command” on page 250
- “mmlsconfig command” on page 400
- “mmlsfileset command” on page 406
- “mmlsfs command” on page 411

Location

/usr/lpp/mmfs/bin

mmafmctl

mmafmctl command

This command is for various operations and reporting information on all filesets. It is recommended to read the *IBM Spectrum Scale: Administration Guide* AFM and AFM Disaster Recovery chapters in conjunction with this manual for detailed description of the functions.

Synopsis

To use the AFM DR functions correctly, use all commands enlisted in this chapter in accordance with the steps described in the AFM-based DR chapter in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

AFM read only mode is referred as RO, single writer mode is referred as SW, independent writer mode is referred as IW and local update mode is referred as LU in this manual.

```
mmafmctl Device {resync | expire | unexpire} -j FilesetName
```

or

```
mmafmctl Device {getstate | resumeRequeued} [-j FilesetName]
```

or

```
mmafmctl Device flushPending [-j FilesetName [--list-file ListFile]]  
[-s LocalWorkDirectory]
```

or

```
mmafmctl Device failover -j FilesetName  
--new-target NewAfmTarget [--target-only] [-s LocalWorkDirectory]
```

or

```
mmafmctl Device prefetch -j FilesetName [--metadata-only]  
[--list-file ListFile] |  
{--home-list-file HomeListFile} |  
{--home-inode-file PolicyListFile}  
[--home-fs-path HomeFileSystemPath]  
[-s LocalWorkDirectory]
```

or

```
mmafmctl Device evict -j FilesetName  
[--safe-limit SafeLimit] [--order {LRU | SIZE}]  
[--log-file LogFile] [--filter Attribute=Value ...]  
[--list-file ListFile] [--file FilePath]
```

or

```
mmafmctl Device fallback -j FilesetName {--start --failover-time Time} | --stop  
[-sLocalWorkDirectory]
```

or

```
mmafmctl Device failoverToSecondary -j FilesetName [--norestore | --restore ]
```

or

```
mmafmctl Device convertToPrimary -j FilesetName  
[ --afmtarget Target { --inband | --outband | --secondary-snapname SnapshotName }]  
[ --check-metadata | --nocheck-metadata ] [--rpo RPO] [-s LocalWorkDirectory]
```

or

```
mmafmctl Device convertToSecondary -j FilesetName --primaryid PrimaryId [ --force ]
```

or

```
mmafmctl Device changeSecondary -j FilesetName
--new-target NewAfmTarget [ --target-only | --inband | --outband ]
[-s LocalWorkDirectory]
```

or

```
mmafmctl Device replacePrimary -j FilesetName
```

or

```
mmafmctl Device fallbackToPrimary -j FilesetName {--start | --stop} [ --force ]
```

or

```
mmafmctl Device {applyUpdates |getPrimaryId } -j FilesetName
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

The usage of options of this command for different operations on both AFM (RO/SW/IW/LU) filesets and AFM primary/secondary filesets are explained with examples.

File system should be mounted on all gateway nodes for **mmafmctl** functions to work.

Parameters

Device

Specifies the device name of the file system.

-j FilesetName

Specifies the fileset name.

-s LocalWorkDirectory

Specifies the temporary working directory.

1. This section describes:

```
mmafmctl Device {resync | expire | unexpire} -j FilesetName
```

resync

This option is available only for SW cache. In case of inadvertent changes made at home of an SW fileset, such as delete of a file or change of data in a file etc., the administrator can correct the home by sending all contents from cache to home using this option. The limitation of this option that renamed files at home may not be fixed by **resync**. Using **resync** requires the cache to be either in NeedsResync or Active state.

expire | unexpire

This option is available only for RO cache. When an RO cache is disconnected, the cached contents are still accessible for the user. However the administrator can define a time from home beyond which access to the cached contents becomes stale. Such an event would occur automatically after disconnection (when cached contents are no longer accessible) and is called *expiration*; the cache is said to be expired. This state can also be forced manually using the **expire** parameter.

When the home comes back or reconnects, the cache contents become automatically accessible again and the cache is said to un-expire. This can be forced manually using the **unexpire** parameter.

The manual expiration and un-expiration can be forced on a cache even when the home is in a connected state. For expiring a fileset manually the **afmExpirationTimeout** needs to have been set on the fileset. If a cache is expired using this manual method, it will also have to be manually unexpired.

mmafmctl

2. This section describes:

```
mmafmctl Device {getstate | resumeRequeued} [-j FilesetName]
```

getstate

This option is applicable for all AFM (RO/SW/IW/LU) and AFM primary filesets. It displays the status of the fileset in the following fields:

Fileset Name

The name of the fileset.

Fileset Target

The host server and the exported path on it.

Gateway Node

Primary gateway of the fileset. This gateway node is handling requests for this fileset.

Queue Length

Current length of the queue on the primary gateway.

Queue numExec

Number of operations played at home since the fileset is last Active.

Cache State

- Cache states applicable for all AFM RO/SW/IW/LU filesets:
Active, Inactive, Dirty, Disconnected, Unmounted
- Cache states applicable for RO filesets:
Expired
- Cache states applicable for SW and IW filesets:
Recovery, FlushOnly, QueueOnly, Dropped, NeedsResync, FailoverInProgress
- Cache states applicable for IW filesets:
FallbackInProgress, FallbackCompleted, NeedsFallback
- Cache states applicable for AFM primary filesets:
PrimInitInProgress, PrimInitFail, Active, Inactive, Dirty, Disconnected, Unmounted, FallbackInProgress, Recovery, FlushOnly, QueueOnly, Dropped, NeedsResync

For more information on all cache states, see the AFM and AFM-based DR chapters in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

resumeRequeued

This option is applicable for SW/IW and primary filesets. If there are operations in the queue that were re-queued due to errors at home, the Administrator should correct those errors and can run this option to retry the re-queued operations.

3. This section describes:

```
mmafmctl Device flushPending [-j FilesetName [--list-file ListFile]]  
[-s LocalWorkDirectory]
```

flushPending

Flushes all point-in-time pending messages in the normal queue on the fileset to home. Requeued messages and messages in the priority queue for the fileset are not flushed by this command.

When **--list-file** *ListFile* is specified, the messages pending on the files listed in the list file are flushed to home. *ListFile* contains a list of files that you want to flush, one file per line. All files must have absolute path names, specified from the fileset linked path. If the list of files has filenames with special characters, use a policy to generate the listfile. Edit to remove all entries other than the filenames. FlushPending is applicable for SW/IW and primary filesets.

4. This section describes:

```
mmafmctl Device failover -j FilesetName
--new-target NewAfmTarget [--target-only] [-s LocalWorkDirectory]
```

This option is applicable only for SW/IW filesets. This option pushes all the data from cache to home. It should be used only in case home is completely lost due to a disaster and a new home is being set up. Failover often takes a long time to complete; status can be checked using the **afmManualResyncComplete** callback or via **mmafmctl getstate** command.

--new-target *NewAfmTarget*

Specifies a new home server and path, replacing the home server and path originally set by the **afmTarget** parameter of the **mmcrfileset** command. Specified in either of the following formats:

Protocol://[Host|Map]/Path

or

{Host|Map}:Path

where:

Protocol://

Specifies the transport protocol. Valid values are **nfs://** or **gpfs://**.

Host|Map

Host

Specifies the server domain name system (DNS) name or IP address.

Map

Specifies the export map name.

Notes:

1. When specifying **nfs://** as the value for *Protocol://*, you must provide a value for *Host* or *Map*.
2. When specifying **gpfs://** as the value for *Protocol://*, do not provide a value for *Host*. However, provide a value for *Map* if it refers to an export map entry.

Path

Specifies the export path.

It is possible to change the protocol along with the target using failover. For example, a cache using an NFS target `bear110:/gpfs/gpfsA/home` can be switched to a GPFS target whose remote file system is mounted at `/gpfs/fs1`, and vice-versa, as follows:

```
mmafmctl fs0 failover -j afm-mc1 --new-target gpfs:///gpfs/fs1
mmafmctl fs0 failover -j afm-mc1 --new-target nfs://bear110/gpfs/gpfsA/home
```

Note that in the first command, `///` is needed because *Host* is not provided.

--target-only

This is used if the user wants to change the mount path/IP address in the target path. The new NFS server should be in the same home cluster and should be of the same architecture as the existing NFS server in the target path. This option should not be used to change the target location or protocol.

5. This section describes:

```
mmafmctl Device prefetch -j FilesetName [--metadata-only]
[--list-file ListFile] |
[--home-fs-path HomeFileSystemPath]
[-s LocalWorkDirectory]
```

This option is used for fetching file contents from home before the application requests for the contents. This reduces the network delay when the application is in progress. You can also use this option to move files over the WAN when the WAN usage is low. These files might be the files that are accessed during high WAN usage. Thus, you can use this option for better WAN management.

mmafmctl

You can use the prefetch option to -

- populate data
- populate metadata
- view prefetch statistics

If you run **prefetch** without providing any options, it displays statistics of the last **prefetch** command run on the fileset.

Prefetch completion can be monitored using the **afmPrepopEnd** event.

--metadata-only

Prefetches only the metadata and not the actual data. This is useful in migration scenarios. This option requires the list of files whose metadata you want. Hence it must be combined with a list file option.

--list-file *ListFile*

The specified file is a file containing a list of files, and needs to be prefetched, one file per line. All files must have fully qualified path names.

If the list of files to be prefetched have filenames with special characters then a policy must be used to generate the listfile. Remove entries from the file other than the filenames.

An indicative list of files:

- files with fully qualified names from cache
- files with fully qualified names from home
- list of files from home generated using policy. Do not edit.

--home-list-file *HomeListFile*

Contains a list of files from the home cluster that needs to be prefetched, one file per line. All files must have fully qualified path names. If the list of files has filenames with special characters, use a policy to generate the listfile. Edit to remove all entries other than the filenames.

This command is deprecated. Use **-list-file** instead.

--home-inode-file *PolicyListFile*

Contains a list of files from the home cluster that needs to be prefetched in the cache. Do not edit the file. The file is generated using policy.

This command is deprecated. Use **-list-file** instead.

--home-fs-path *HomeFileSystemPath*

Specifies the full path to the fileset at the home cluster and can be used in conjunction with *ListFile*.

You must use this option, when in the NSD protocol the mount point on the gateway nodes of the **afmTarget** filesets does not match the mount point on the Home cluster.

For example, **mmafmctl gpfs1 prefetch -j cache1 --home-fs-path /gpfs/remotefs1**

In this example, the file system is mounted on the :

- home cluster at /gpfs/homefs1
- gateway nodes at /gpfs/remotefs1

Prefetch is an asynchronous process and you can use the fileset when prefetch is in progress. You can monitor Prefetch using the **afmPrepopEnd** event. AFM can prefetch the data using the **mmafmctl prefetch** command (which specifies a list of files to prefetch). Prefetch always pulls the complete file contents from home and AFM automatically sets a file as cached when it is completely prefetched.

6. This section describes:

```
mmafmctl Device evict -j FilesetName
  [--safe-limit SafeLimit] [--order {LRU | SIZE}]
  [--log-file LogFile] [--filter Attribute=Value ...]
  [--list-file ListFile] [--file FilePath]
```

This option is applicable for RO/SW/IW/LU filesets. When cache space exceeds the allocated quota, data blocks from non-dirty are automatically de-allocated with the eviction process. This option can be used for a file that is specifically to be de-allocated based on some criteria. All options can be combined with each other.

--safe-limit SafeLimit

This is a compulsory parameter for the manual evict option, for order and filter attributes. Specifies target quota limit (which is used as the low water mark) for eviction in bytes; must be less than the soft limit. This parameter can be used alone or can be combined with one of the following parameters (order or filter attributes). Specify the parameter in bytes.

--order LRU | SIZE

Specifies the order in which files are to be chosen for eviction:

LRU

Least recently used files are to be evicted first.

SIZE

Larger-sized files are to be evicted first.

--log-file LogFile

Specifies the file where the eviction log is to be stored. The default is that no logs are generated.

--filter Attribute=Value

Specifies attributes that enable you to control how data is evicted from the cache. Valid attributes are:

FILENAME=FileName

Specifies the name of a file to be evicted from the cache. This uses an SQL-type search query. If the same file name exists in more than one directory, it will evict all the files with that name. The complete path to the file should not be given here.

MINFILESIZE=Size

Sets the minimum size of a file to evict from the cache. This value is compared to the number of blocks allocated to a file (**KB_ALLOCATED**), which may differ slightly from the file size.

MAXFILESIZE=Size

Sets the maximum size of a file to evict from the cache. This value is compared to the number of blocks allocated to a file (**KB_ALLOCATED**), which may differ slightly from the file size.

--list-file ListFile

Contains a list of files that you want to evict, one file per line. All files must have fully qualified path names. Filesystem quotas need not be specified. If the list of files has filenames with special characters, use a policy to generate the listfile. Edit to remove all entries other than the filenames.

--file FilePath

The fully qualified name of the file that needs to be evicted. Filesystem quotas need not be specified.

Possible combinations of *safelimit*, *order*, and *filter* are:

only Safe limit

Safe limit + LRU

Safe limit + SIZE

Safe limit + FILENAME

Safe limit + MINFILESIZE

Safe limit + MAXFILESIZE

Safe limit + LRU + FILENAME

Safe limit + LRU + MINFILESIZE

Safe limit + LRU + MAXFILESIZE

mmafmctl

Safe limit + SIZE + FILENAME
Safe limit + SIZE + MINFILESIZE
Safe limit + SIZE + MAXFILESIZE

7. This section describes:

```
mmafmctl Device failback -j FilesetName {--start --failover-time Time} | --stop  
[-sLocalWorkDirectory]
```

failback is applicable only for IW filesets.

failback --start --failover-time Time

Specifies the point in time at the home cluster, from which the independent-writer cache taking over as writer should sync up. *Time* can be specified in date command format with time zones. It will use the cluster's time-zone and year by default.

failback --stop

An option to be run after the failback process is complete and the fileset moves to **FallbackCompleted** state. This will move the fileset to **Active** state.

8. This section describes:

```
mmafmctl Device failoverToSecondary -j FilesetName [--norestore | --restore ]
```

This is to be run on a secondary fileset.

When primary experiences a disaster, all applications will need to be moved to the secondary to ensure business continuity. The secondary has to be first converted to an acting primary using this option.

There is a choice of restoring the latest snapshot data on the secondary during the failover process or leave the data as is using the **--norestore** option. Once this is complete, the secondary becomes ready to host applications.

--norestore

Specifies that restoring from the latest RPO snapshot is not required. This is the default setting.

--restore

Specifies that data must be restored from the latest RPO snapshot.

9. This section describes:

```
mmafmctl Device convertToPrimary -j FilesetName  
[ --afmtarget Target { --inband | --outband | --secondary-snapname SnapshotName }]  
[ --check-metadata | --nocheck-metadata ] [ --rpo RPO ] [-s LocalWorkDirectory]
```

This is to be run on a GPFS fileset or SW/IW fileset which is intended to be converted to primary.

--afmtargetTarget

Specifies the secondary that needs to be configured for this primary. Need not be used for AFM filesets as target would already have been defined.

--inband

Used for inband trucking. *Inband trucking* is copying the data while setting up a primary/secondary relationship from GPFS filesets, where primary site has contents and secondary site is empty.

--outband

Used for outband trucking. *Outband trucking* is copying data manually using other ways such as **ftp**, **scp**, **rsync** etc. This should be completed before the relationship is established.

--check-metadata

This checks if the disallowed types (like immutable/append-only files) are present in the GPFS fileset on the primary site before the conversion. Conversion with this option fails if such files exist.

For SW/IW filesets, presence of orphans and incomplete directories are also checked. SW/IW filesets should have established contact with at least once home for this option to succeed.

--nocheck-metadata

Used if one needs to proceed with conversion without checking for appendonly/imutable files.

--secondary-snapname SnapshotName

Used while establishing a new primary for an existing secondary or acting primary during failback.

--rpo RPO

Specifies the RPO interval in minutes for this primary fileset. Disabled by default.

10. This section describes:

```
mmafmctl Device convertToSecondary -j FilesetName --primaryid PrimaryId [ --force ]
```

This is to be run on a GPFS fileset on the secondary site. This converts a GPFS independent fileset to a secondary and sets the primary ID.

--primaryid PrimaryId

Specifies the ID of the primary with which the secondary will be associated.

--force

If **convertToSecondary** failed or got interrupted, it will not create afmctl file at the secondary. The command should be rerun with the **--force** option.

11. This section describes:

```
mmafmctl Device changeSecondary -j FilesetName  
--new-target NewAfmTarget [ --target-only | --inband | --outband ]  
[-s LocalWorkDirectory]
```

This is to be run on a primary fileset only.

A disaster at the secondary can take place due to which secondary is not available.

Run this command on the primary when a secondary fails and this primary needs to be connected with a new secondary. On the new secondary site a new GPFS independent fileset has to be created. Data on the primary can be copied to the new GPFS fileset that was created with this command using other means such as **ftp**, **scp** etc. Alternatively it can be decided that data will be trucked using the relationship.

--new-target NewAfmTarget

Used to mention the new secondary.

--inband | --outband

Used based on the method used to truck data.

--target-only

Used when you want to change the IP address or NFS server name for the same target path. The new NFS server must be in the same home cluster and must be of the same architecture(power or x86) as the existing NFS server in the target path. This option can be used to move from NFS to a mapping target.

12. This section describes:

```
mmafmctl Device replacePrimary -j FilesetName
```

This is used on an acting primary only. This creates a latest snapshot of the acting primary. This command deletes any old RPO snapshots on the acting primary and creates a new initial RPO snapshot psnap0.

This RPO snapshot is used in the setup of the new primary.

mmafmctl

13. This section describes:

```
mmafmctl Device fallbackToPrimary -j FilesetName {--start | --stop} [--force]
```

This is to be run on an old primary that came back after the disaster, or on a new primary that is to be configured after an old primary went down with a disaster. The new primary should have been converted from GPFS to primary using convertToPrimary option.

--start

Restores the primary to the contents from the last RPO on the primary before the disaster. This option will put the primary in read-only mode, to avoid accidental corruption until the fallback process is completed. In case of new primary that is setup using convertToPrimary, the fallback --start does no change.

--stop

Used to complete the Fallback process. This will put the fileset in read-write mode. The primary is now ready for starting applications.

--force

Used if --stop or --start does not complete successfully due to any errors, and not allow fallbackToPrimary to stop or start again.

14. This section describes:

```
mmafmctl Device {applyUpdates |getPrimaryId } -j FilesetName
```

Both options are intended for the primary fileset.

applyUpdates

Run this on the primary after running the **fallback --start** command. All the differences can be brought over in one go or through multiple iterations. For minimizing application downtime, this command can be run multiple times to bring the primary's contents in sync with the acting primary. When the contents are as close as possible or minimal, applications should take a downtime and then this command should be run one last time.

It is possible that **applyUpdates** fails with an error during instances when the acting primary is overloaded. In such cases the command has to be run again.

getPrimaryID

Used to get primary Id of a primary fileset.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmafmctl** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Examples

1. Running resync on SW:

```
# mmafmctl fs1 resync -j sw1  
mmafmctl: Performing resync of fileset: sw1
```

```
# mmafmctl fs1 getstate -j sw1
Fileset Name Fileset Target          Cache State Gateway Node Queue Length Queue numExec
----- -----
sw1      nfs://c26c3apv2/gpfs/homefs1/newdir1 Dirty      c26c2apv1   4067      10844
```

2. Expiring a RO fileset:

```
# mmafmctl fs1 expire -j ro1
```

```
# mmafmctl fs1 getstate -j ro1
Fileset Name Fileset Target          Cache State Gateway Node Queue Length Queue numExec
----- -----
ro1      gpfs:///gpfs/remotefs1/dir1 Expired    c26c4apv1   0           4
```

3. Unexpiring a RO fileset:

```
# mmafmctl fs1 unexpire -j ro1
```

```
# mmafmctl fs1 getstate -j ro1
Fileset Name Fileset Target          Cache State Gateway Node Queue Length Queue numExec
----- -----
ro1      gpfs:///gpfs/remotefs1/dir1 Active     c26c4apv1   0           4
```

4. Run flushPending on SW fileset:

```
// Populate the fileset with data
```

```
# mmafmctl fs1 getstate -j sw1
Fileset Name Fileset Target          Cache State Gateway Node Queue Length Queue numExec
----- -----
sw1      gpfs:///gpfs/remotefs1/dir1 Dirty      c26c2apv1   5671      293
```

Get the list of files newly created using policy:

```
RULE EXTERNAL LIST 'L' RULE 'List' LIST 'L' WHERE PATH_NAME LIKE '%'
```

```
# mmapplypolicy /gpfs/fs1/sw1/migrateDir.popFSDir.22655 -P p1 -f p1.res -L 1 -N mount -I defer
```

Policy created this file, this should be hand-edited to retain only the names:

```
11012030 65537 0 -- /gpfs/fs1/sw1/migrateDir.popFSDir.22655/file_with_posix_acl1
11012032 65537 0 -- /gpfs/fs1/sw1/migrateDir.popFSDir.22655/populateFS.log
```

```
11012033 65537 0 --
```

```
/gpfs/fs1/sw1/migrateDir.popFSDir.22655/sparse_file_0_with_0_levels_indirection
```

```
# cat p1.res.list | awk '{print $5}' > lfile
```

```
# mmafmctl fs1 flushPending -j sw1 --list-file=lfile
```

5. Failover of SW to a new home:

```
# mmafmctl fs1 getstate -j sw1
```

```
Fileset Name Fileset Target          Cache State Gateway Node Queue Length Queue numExec
----- -----
sw1      gpfs:///gpfs/remotefs1/dir1 Dirty      c26c2apv1   785       5179
```

```
# mmcrfileset homefs1 newdir1 --inode-space=new
Fileset newdir1 created with id 219 root inode 52953091.
```

```
# mmmlinkfileset homefs1 newdir1 -J /gpfs/homefs1/newdir1
Fileset newdir1 linked at /gpfs/homefs1/newdir1
```

```
# mmafmconfig enable /gpfs/homefs1/newdir1
```

```
# mmafmctl fs1 failover -j sw1 --new-target=c26c3apv1:/gpfs/homefs1/newdir1
mmafmctl: Performing failover to nfs://c26c3apv1/gpfs/homefs1/newdir1
Fileset sw1 changed.
```

```
mmafmctl: Failover in progress. This may take while...
```

```
Check fileset state or register for callback to know the completion status.
```

Callback registered, logged into mmfs.log:

Thu May 21 03:06:18.303 2015: [I] Calling User Exit Script callback7: event

mmafmctl

```
afmManualResyncComplete, Async command recovery.sh
```

```
# mmafmctl fs1 getstate -j sw1
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
----- -----
sw1      nfs://c26c3apv1/gpfs/homefs1/newdir1 Active      c26c2apv1  0        5250
```

6. Changing target of SW fileset:

Changing to another NFS server in the same home cluster using --target-only option:

```
# mmafmctl fs1 failover -j sw1 --new-target=c26c3apv2:/gpfs/homefs1/newdir1 --target-only
mmafmctl: Performing failover to nfs://c26c3apv2/gpfs/homefs1/newdir1
Fileset sw1 changed.
```

```
# mmafmctl fs1 getstate -j sw1
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
----- -----
sw1      nfs://c26c3apv2/gpfs/homefs1/newdir1 Active      c26c2apv1  0        5005
```

7. Metadata population using prefetch:

```
# mmafmctl fs1 getstate -j ro
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
----- -----
ro       nfs://c26c3apv1/gpfs/homefs1/dir3 Active      c26c2apv2  0        7
```

List Policy:

```
RULE EXTERNAL LIST 'List' RULE 'List' LIST 'List' WHERE PATH_NAME LIKE '%'
```

Run the policy at home:

```
mmaplypolicy /gpfs/homefs1/dir3 -P px -f px.res -L 1 -N mount -I defer
```

Policy created this file, this should be hand-edited to retain only file names.

This file can be used at the cache to populate metadata.

```
# mmafmctl fs1 prefetch -j ro --metadata-only -list-file=px.res.list.List
mmafmctl: Performing prefetching of fileset: ro
```

Prefetch end can be monitored using this event:

```
Thu May 21 06:49:34.748 2015: [I] Calling User Exit Script prep0: event afmPrepopEnd,
Async command prep0.sh.
```

The statistics of the last prefetch command is viewed by:

```
mmafmctl fs1 prefetch -j ro
Fileset Name Async Read (Pending) Async Read (Failed) Async Read (Already Cached) Async Read (Total) Async Read (Data in Bytes)
----- -----
ro       0           1           0           7           0
```

8. Prefetch of data using --home-list-file option:

```
# cat /lfile1
/gpfs/homefs1/dir3/file1
/gpfs/homefs1/dir3/dir1/file1
# mmafmctl fs1 prefetch -j ro --home-list-file=/lfile1
mmafmctl: Performing prefetching of fileset: ro
```

```
# mmafmctl fs1 prefetch -j ro
Fileset Name Async Read (Pending) Async Read (Failed) Async Read (Already Cached) Async Read (Total) Async Read (Data in Bytes)
----- -----
ro       0           0           0           2           122880
```

9. Prefetch of data using --home-inode-file option:

Inode file is created using the above policy at home, and should be used as such without hand-editing.

List Policy:

```
RULE EXTERNAL LIST 'List' RULE 'List' LIST 'List' WHERE PATH_NAME LIKE '%'
```

Run the policy at home:

```
# mmapplypolicy /gpfs/homefs1/dir3 -P px -f px.res -L 1 -N mount -I defer
```

```
# cat /lfile2
113289 65538 0 -- /gpfs/homefs1/dir3/file2
113292 65538 0 -- /gpfs/homefs1/dir3/dir1/file2
# mmafmctl fs1 prefetch -j ro --home-inode-file=/lfile2
mmafmctl: Performing prefetching of fileset: ro

mmafmctl fs1 prefetch -j ro
Fileset Name Async Read (Pending) Async Read (Failed) Async Read (Already Cached) Async Read (Total) Async Read (Data in Bytes)
-----
```

Fileset	Name	Async Read (Pending)	Async Read (Failed)	Async Read (Already Cached)	Async Read (Total)	Async Read (Data in Bytes)
ro		0	0	2	2	0

10. Using --home-fs-path option for a target with NSD protocol:

```
# mmafmctl fs1 getstate -j ro2
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
-----
```

Fileset	Name	Target	Cache State	Gateway	Node	Queue Length	Queue numExec
ro2		gpfs:///gpfs/remotefs1/dir3	Active	c26c4apv1	0	7	

```
# cat /lfile2
113289 65538 0 -- /gpfs/homefs1/dir3/file2
113292 65538 0 -- /gpfs/homefs1/dir3/dir1/file2
# mmafmctl fs1 prefetch -j ro2 --home-fs-path=/gpfs/homefs1/dir3
mmafmctl: Performing prefetching of fileset: ro2
```

```
# mmafmctl fs1 prefetch -j ro2
Fileset Name Async Read (Pending) Async Read (Failed) Async Read (Already Cached) Async Read (Total) Async Read (Data in Bytes)
-----
```

Fileset	Name	Async Read (Pending)	Async Read (Failed)	Async Read (Already Cached)	Async Read (Total)	Async Read (Data in Bytes)
ro2		0	0	0	2	122880

11. Manually evicting using safe-limit and filename parameters:

```
# ls -lis /gpfs/fs1/ro2/file10M_1
12605961 10240 -rw-r--r-- 1 root root 10485760 May 21 07:44 /gpfs/fs1/ro2/file10M_1
```

```
# mmafmctl fs1 evict -j ro2 --safe-limit=1 --filter FILENAME=file10M_1
```

```
# ls -lis /gpfs/fs1/ro2/file10M_1
12605961 0 -rw-r--r-- 1 root root 10485760 May 21 07:44 /gpfs/fs1/ro2/file10M_1
```

12. IW Fallback:

```
# mmafmctl fs1 getstate -j iw1
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
-----
```

Fileset	Name	Target	Cache State	Gateway	Node	Queue Length	Queue numExec
iw1		nfs://c26c3apv1/gpfs/homefs1/dir3	Active	c25m4n03	0	8	

```
# touch file3 file4
```

```
# mmafmctl fs1 getstate -j iw1
Fileset Name Fileset Target Cache State Gateway Node Queue Length Queue numExec
-----
```

Fileset	Name	Target	Cache State	Gateway	Node	Queue Length	Queue numExec
iw1		nfs://c26c3apv1/gpfs/homefs1/dir3	Dirty	c25m4n03	2	11	

Unlink IW fileset feigning failure:

```
# munlinkfileset fs1 iw1 -f
Fileset iw1 unlinked.
```

Write from IW home, assuming applications failed over to home:

```
Thu May 21 08:20:41 4]dir3# touch file5 file6
```

Relink IW back on the cache cluster, assuming it came back up:

```
# mmlinkfileset fs1 iw1 -J /gpfs/fs1/iw1
Fileset iw1 linked at /gpfs/fs1/iw1
```

Run fallback on IW:

```
# mmafmctl fs1 fallback -j iw1 --start --failover-time='May 21 08:20:41'
```

```
# mmafmctl fs1 getstate -j iw1
```

mmafctl

```
Fileset Name Fileset Target          Cache State   Gateway Node Queue Length Queue numExec
----- ----- -----
iw1      nfs://c26c3apv1/gpfs/homefs1/dir3 FallbackInProg c25m4n03    0           0
```

```
# mmafctl fs1 fallback -j iw1 -stop
```

```
# mmafctl fs1 getstate -j iw1
```

```
Fileset Name Fileset Target          Cache State   Gateway Node Queue Length Queue numExec
----- ----- -----
iw1      nfs://c26c3apv1/gpfs/homefs1/dir3 Active       c25m4n03    0           3
```

13. Manual evict using the --list-file option:

```
[root@c21f2n08 ~]# ls -lshi /gpfs/fs1/evictCache
total 6.0M
```

```
27858308 1.0M -rw-r--r-- 1 root root 1.0M Feb  5 02:07 file1M
27858307 2.0M -rw-r--r-- 1 root root 2.0M Feb  5 02:07 file2M
27858306 3.0M -rw-r--r-- 1 root root 3.0M Feb  5 02:07 file3M
```

```
[root@c21f2n08 ~]# echo "RULE EXTERNAL LIST 'HomePREPDAEMON' RULE 'ListLargeFiles'
LIST 'HomePREPDAEMON' WHERE PATH_NAME LIKE '%'" > /tmp/evictionPolicy.pol
```

```
[root@c21f2n08 ~]# mmapplypolicy /gpfs/fs1/evictCache -I defer -P /tmp/evictionPolicy.pol -f /tmp/evictionList
```

```
#Edited list of files to be evicted
```

```
[root@c21f2n08 ~]# cat /tmp/evictionList.list.HomePREPDAEMON
27858306 605742886 0 -- /gpfs/fs1/evictCache/file3M
```

```
[root@c21f2n08 ~]# mmafctl fs1 evict -j evictCache --list-file /tmp/evictionList.list.HomePREPDAEMON
```

```
[root@c21f2n08 ~]# ls -lshi /gpfs/fs1/evictCache
total 3.0M
```

```
27858308 1.0M -rw-r--r-- 1 root root 1.0M Feb  5 02:07 file1M
27858307 2.0M -rw-r--r-- 1 root root 2.0M Feb  5 02:07 file2M
27858306 0 -rw-r--r-- 1 root root 3.0M Feb  5 02:07 file3M
```

14. Manual evict using the --file option:

```
[root@c21f2n08 ~]# ls -lshi /gpfs/fs1/evictCache
total 3.0M
```

```
27858308 1.0M -rw-r--r-- 1 root root 1.0M Feb  5 02:07 file1M
27858307 2.0M -rw-r--r-- 1 root root 2.0M Feb  5 02:07 file2M
27858306 0 -rw-r--r-- 1 root root 3.0M Feb  5 02:07 file3M
```

```
[root@c21f2n08 ~]# mmafctl fs1 evict -j evictCache --file /gpfs/fs1/evictCache/file1M
```

```
[root@c21f2n08 ~]# ls -lshi /gpfs/fs1/evictCache
total 0
```

```
27858308 0 -rw-r--r-- 1 root root 1.0M Feb  5 02:07 file1M
27858307 0 -rw-r--r-- 1 root root 2.0M Feb  5 02:07 file2M
27858306 0 -rw-r--r-- 1 root root 3.0M Feb  5 02:07 file3M
```

See also

- “mmafmconfig command” on page 37
- “mmafmlocal command” on page 54
- “mmchattr command” on page 122
- “mmchconfig command” on page 132
- “mmchfileset command” on page 173
- “mmchfs command” on page 179
- “mmcrfileset command” on page 244
- “mmcrfs command” on page 250
- “mmlsconfig command” on page 400
- “mmlsfileset command” on page 406

- “`mmlsfs` command” on page 411
- “`mmpsnap` command” on page 501

See the AFM and AFM-based DR chapters in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide* for details.

Location

`/usr/lpp/mmfs/bin`

mmafmlocal

mmafmlocal command

Provides a list of cached files and file statistics such as inode number, allocated blocks, and so on.

Synopsis

mmafmlocal ls [FileName ...]

or

mmafmlocal stat FileName ...

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

The **mmafmlocal** command provides information about files that exist in the cache.

Parameters

ls Lists files with data that is in the cache already. This parameter is valid for fully-cached files only.

FileName

Specifies the name of a file to be listed.

stat

Displays statistics for files. If the file is not cached already, the number of allocated blocks is zero. This parameter is valid for partially-cached and fully-cached files.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmafmlocal** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

1. To list information of all the cached files in a fileset:

```
mmafmlocal ls
```

The system displays information similar to:

```
total 10240
-rwxrwxrwx 1 root root 10485760 May 24 09:20 file1
```

2. To list information of a specific cached file in a fileset:

```
mmafmlocal ls file2
```

The system displays information similar to:

```
-rwxrwxrwx 1 root root 10485760 May 24 09:20 file2
```

3. To list the file statistics:

```
mmafmlocal stat file2
```

The system displays information similar to:

```
File: file2
  Inode number: 1582477
  Device ID: 0x2C (44)
  Size: 10485760
  Blocks: 20480
  Block size: 262144
  Links: 1
  Uid: 0 (root)
  Gid: 0 (root)
  Mode: 0100777
  Access time: 1464281093 (Thu May 26 16:44:53 2016 UTC)
  Modify time: 1464096038 (Tue May 24 13:20:38 2016 UTC)
  Change time: 1464096038 (Tue May 24 13:20:38 2016 UTC)
```

See also

- “mmafmconfig command” on page 37
- “mmafmctl command” on page 40
- “mmchattr command” on page 122
- “mmchconfig command” on page 132
- “mmchfileset command” on page 173
- “mmchfs command” on page 179
- “mmcrfileset command” on page 244
- “mmcrfs command” on page 250
- “mmlsconfig command” on page 400
- “mmlsfileset command” on page 406
- “mmlsfs command” on page 411
- “mmpssnap command” on page 501

Location

/usr/lpp/mmfs/bin

mmapplypolicy

mmapplypolicy command

Deletes files, migrates files between storage pools, or does file compression or decompression in a file system as directed by policy rules.

Synopsis

```
mmapplypolicy {Device|Directory}
  [-A IscanBuckets] [-a IscanThreads] [-B MaxFiles]
  [-D yyyy-mm-dd[@hh:mm[:ss]]] [-e] [-f FileListPrefix]
  [-g GlobalWorkDirectory] [-I {yes|defer|test|prepare}]
  [-i InputFileList] [-L n] [-M name=value...] [-m ThreadLevel]
  [-N {all | mount | Node[,Node...] | NodeFile | NodeClass}]
  [-n DirThreadLevel] [-P PolicyFile] [-q] [-r FileListPathname...]
  [-S SnapshotName] [-s LocalWorkDirectory]
  [--choice-algorithm {best | exact | fast}]
  [--maxdepth MaxDirectoryDepth]
  [--max-merge-files MaxFiles] [--max-sort-bytes MaxBytes]
  [--other-sort-options SortOptions] [--qos QosClass]
  [--scope {filesystem | fileset | inodespace}]
  [--single-instance] [--sort-buffer-size Size]
  [--sort-command SortCommand] [--split-filelists-by-weight]
  [--split-margin n.n]
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

You can use the **mmapplypolicy** command to apply rules that manage the following types of tasks:

- Migration and replication of file data to and from storage pools.
- Deleting files.
- File compression or decompression. For more information, see the topic *File compression* in the *IBM Spectrum Scale: Administration Guide*.

For more information about policy rules, see the topic *Policies for automating file management* in the *IBM Spectrum Scale: Administration Guide*.

You can run the **mmapplypolicy** command from any node in the cluster that has mounted the file system.

The **mmapplypolicy** command does not affect placement rules (for example, the **SET POOL** and **RESTORE** rule) that are installed for a file system by the **mmchpolicy** command. To display the currently installed rules, issue the **mmlspolicy** command.

A given file can match more than one list rule, but will be included in a given list only once. *ListName* provides the binding to an **EXTERNAL LIST** rule that specifies the executable program to use when processing the generated list.

The **EXTERNAL POOL** rule defines an external storage pool. This rule does not match files, but serves to define the binding between the policy language and the external storage manager that implements the external storage.

Any given file is a potential candidate for at most one **MIGRATE** or **DELETE** operation during one invocation of the **mmapplypolicy** command. That same file may also match the first applicable **LIST** rule.

A file that matches an **EXCLUDE** rule is not subject to any subsequent **MIGRATE**, **DELETE**, or **LIST** rules. You should carefully consider the order of rules within a policy to avoid unintended consequences.

For detailed information on GPFS policies, see the *IBM Spectrum Scale: Administration Guide*.

This command cannot be run from a Windows node. The GPFS API, documented functions in **gpfs.h** are not implemented on Windows, however the policy language does support the Windows file attributes, so you can manage your GPFS Windows files using the **mmapplypolicy** command running on an AIX or Linux node.

Note: To terminate **mmapplypolicy**, use the **kill** command to send a **SIGTERM** signal to the process group running **mmapplypolicy**.

For example, on Linux if you wanted to terminate **mmapplypolicy** on a process group whose ID is 3813, you would enter the following:

```
kill -s SIGTERM -- -3813
```

If you need to determine which process group is running **mmapplypolicy**, you can use the following command (which also tells you which process groups are running **tsapolicy** and **mmhelp-apolicy**):

```
mmdsh -N all ps auxw | grep policy
```

The system displays output similar to the following:

```
root      31666  0.0  0.0  84604  2328 ?          S1    07:29   0:00 /usr/lpp/mmfs/bin/mmhelp-apolicy na -X
          10.222.4.12 -s /tmp -Y -x 36845 -m 24 -n 24 -a 2 -L 1 -d 00 -z 15
root      3813   0.3  0.1  68144  4792 pts/1    S     07:29   0:00 /bin/ksh /usr/lpp/mmfs/bin/mmapplypolicy
          /mak/millions -P /ghome/makaplan/policies/1p.policy -N all
root      3847   127  0.1  455228  5808 pts/1    S1    07:29   0:38 /usr/lpp/mmfs/bin/tsapolicy /mak/millions
          -P /ghome/makaplan/policies/1p.policy -I yes -L 1 -X 10.222.4.12 -N all
root      3850   0.0  0.0  84832  1620 pts/1    S1    07:29   0:00 /usr/lpp/mmfs/bin/tsapolicy /mak/millions
          -P /ghome/makaplan/policies/1p.policy -I yes -L 1 -X 10.222.4.12 -N all
root      3891   0.0  0.0  61156   768 pts/2    S+   07:29   0:00 grep policy
```

Parameters

Device

Specifies the device name of the file system from which files will have the policy rules applied. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**. If specified, this must be the first parameter.

Directory

Specifies the fully-qualified path name of a GPFS file system subtree from which files will have the policy rules applied. If specified, this must be the first parameter.

-A *IscanBuckets*

Specifies the number of buckets of inode numbers (number of inode/filelists) to be created by the parallel directory scan and processed by the parallel inode scan. Affects the execution of the high-performance protocol that is used when both **-g** and **-N** are specified.

Tip: Set this parameter to the expected number of files to be scanned divided by one million. Then each bucket will have about one million files.

-a *IscanThreads*

Specifies the number of threads and sort pipelines each node will run during the parallel inode scan and policy evaluation. It affects the execution of the high-performance protocol that is used when both **-g** and **-N** are specified. The default is 2. Using a moderately larger number can significantly improve performance, but might "strain" the resources of the node. In some environments a large value for this parameter can lead to a command failure.

Tip: Set this parameter to the number of CPU "cores" implemented on a typical node in your GPFS cluster.

mmapplypolicy

-B MaxFiles

Specifies how many files are passed for each invocation of the EXEC script. The default value is 100.

If the number of files exceeds the value specified for *MaxFiles*, **mmapplypolicy** invokes the external program multiple times.

For more information about file list records, refer to the *IBM Spectrum Scale: Administration Guide*.

-D yyyy-mm-dd[hh:mm[:ss]]

Specifies a date and optionally a (UTC) time as *year-month-day* at *hour:minute:second*.

The **mmapplypolicy** command evaluates policy rules as if it were running on the date and time specified by the **-D** flag. This can be useful for planning or testing policies, to see how the **mmapplypolicy** command would act in the future. If this flag is omitted, the **mmapplypolicy** command uses the current date and (UTC) time. If a date is specified but not a time, the time is assumed to be 00:00:00.

-e

Causes **mmapplypolicy** to re-evaluate and revalidate the following conditions immediately before executing the policy action for each chosen file:

- That the PATH_NAME still leads to the chosen file, and that the INODE and GENERATION values are the same.
- That the rule (iRule) still applies to, and is a first matching rule for, the chosen file.

Note: The **-e** option is particularly useful with **-r**, but can be used apart from it. It is useful because in the time that elapses after the policy evaluation and up to the policy execution, it is possible that the chosen pathname no longer refers to the same inode (for example the original file was removed or renamed), or that some of the attributes of the chosen file have changed in some way so that the chosen file no longer satisfies the conditions of the rule. In general, the longer the elapsed time, the more likely it is that conditions have changed (depending on how the file system is being used). For example, if files are only written once and never renamed or erased, except by policy rules that call for deletion after an expiration interval, then it is probably not necessary to re-evaluate with the **-e** option.

For more information about **-r**, see *IBM Spectrum Scale: Administration Guide*.

-f FileListPrefix

Specifies the location (a path name or file name prefix or directory) in which the file lists for external pool and list operations are stored when either the **-I defer** or **-I prepare** option is chosen. The default location is *LocalWorkDirectory/mmapplypolicy.processid*.

-g GlobalWorkDirectory

Specifies a *global* directory to be used for temporary storage during **mmapplypolicy** command processing. The specified directory must exist within a shared file system. It must also be mounted and available for writing and reading from each of the nodes specified by the **-N** option. When both **-N** and **-g** are specified, **mmapplypolicy** uses high performance and fault-tolerant protocols during execution.

Note: The **-g** option should specify a directory (for temporary or work files) within a GPFS file system that is accessible from each node specified with the **-N** option. The directory can either be in the file system being operated upon by **mmapplypolicy** or in another file system.

There is no default value for **-g**.

If the **-g** option is specified, but not the **-s** option, the directory specified by **-g** is used for all temporary files required by **mmapplypolicy**. If both the **-g** and **-s** options are specified, temporary files may be stored in each. In general, temporary files that are only written and read by a single node are stored in the local work directory specified by the **-s** option, while temporary files that must be accessed by more than one node are stored in the global work directory specified by the **-g** option.

-I {yes | defer | test | prepare}

Specifies what actions the **mmapplypolicy** command performs on files:

yes

Indicates that all applicable policy rules are run, and the data movement between pools is done during the processing of the **mmapplypolicy** command. All defined external lists will be executed. This is the default action.

defer

Indicates that all applicable policy rules are run, but actual data movement between pools is deferred until the next **mmrestripefs** or **mmrestripefile** command. See also “-f FileListPrefix” on page 58.

test

Indicates that all policy rules are evaluated, but the **mmapplypolicy** command only displays the actions that would be performed had **-I defer** or **-I yes** been specified. There is no actual deletion of files or data movement between pools. This option is intended for testing the effects of particular policy rules.

prepare

Indicates that all policy execution is deferred and that **mmapplypolicy** only prepares file lists that are suitable for execution with the **-r** option. Records are written for each of the chosen files and are stored in one or more file lists, under a path name that is specified by the **-f** option or in the default local work directory. The actual data movement occurs when the command is rerun with the **-r** option.

-i InputFileList

Specifies the path name for a user-provided input file list. This file list enables you to specify multiple starter directories or files. It can be in either of the following formats:

simple format file list

A list of records with the following format:

PATH_NAME

Each record represents either a single file or a directory. When a directory is specified, the command processes the entire subtree that is rooted at the specified path name

File names can contain spaces and special characters; however, the special characters '\\' and '\\n' must be escaped with the '\' character similarly to the way **mmapplypolicy** writes path names in file lists for external pool and list operations.

The end-of-record character must be '\n'.

Example:

```
/mak/ea  
/mak/old news  
/mak/special\\stuff
```

/usr/lpp/mmfssamples/ilm/mmglobexpf.sample is an example of a script that can be used to generate simple format file lists.

expert format file list

A list of records with the following format:

```
INODE:GENERATION:path-length!PATH_NAME end-of-record-character
```

Each record represents exactly one file.

The INODE and GENERATION values must be specified in hexadecimal format (%11x). If you do not know the generation number or inode number, specify 0 and GPFS will look it up for you.

mmapplypolicy

The *path-length* value must be specified in decimal format (%d). The *path-length* value is followed by the delimiter !.

The end-of-record character must be \n or \0.

Example (the end-of-record characters are invisible):

```
00009a00:0:8!d14/f681  
00009a01:1002:8!d14/f682
```

When you use an expert format file list, the directory scan phase is skipped and only the files that are specified with the *InputFileList* parameter are tested against the policy rules.

For more information, see the *IBM Spectrum Scale: Administration Guide*.

With either format, if a path name is not fully qualified, it is assumed to be relative to one of the following:

- the *Directory* parameter on the **mmapplypolicy** command invocation

Or

- the mount point of the GPFS file system, if *Device* is specified as the first argument

-L n

Controls the level of information displayed by the **mmapplypolicy** command. Larger values indicate the display of more detailed information. These terms are used:

candidate file

A file that matches a **MIGRATE**, **DELETE**, or **LIST** policy rule.

chosen file

A candidate file that has been scheduled for action.

These are the valid values for *n*:

- 0 Displays only serious errors.
- 1 Displays some information as the command runs, but not for each file. This is the default.
- 2 Displays each chosen file and the scheduled migration or deletion action.
- 3 Displays the same information as 2, plus each candidate file and the applicable rule.
- 4 Displays the same information as 3, plus each explicitly **EXCLUDEed** or **LISTed** file, and the applicable rule.
- 5 Displays the same information as 4, plus the attributes of candidate and **EXCLUDEed** or **LISTed** files.
- 6 Displays the same information as 5, plus non-candidate files and their attributes.

For examples and more information on this flag, see the section: *The mmapplypolicy -L command* in the *IBM Spectrum Scale: Problem Determination Guide*.

-M name=value...

Indicates a string substitution that will be made in the text of the policy rules before the rules are interpreted. This allows the administrator to reuse a single policy rule file for incremental backups without editing the file for each backup.

-m ThreadLevel

The number of threads that are created and dispatched within each **mmapplypolicy** process during the policy execution phase. The default value is 24.

-N {all | mount | Node[,Node...] | NodeFile | NodeClass}

Specifies the list of nodes that will run parallel instances of policy code in the GPFS home cluster. This command supports all defined node classes. The default is to run on the node where the **mmapplypolicy** command is running or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

-n DirThreadLevel...

The number of threads that will be created and dispatched within each **mmapplypolicy** process during the directory scan phase. The default is 24.

-P PolicyFile

Specifies the name of the file containing the policy rules to be applied. If not specified, the policy rules currently in effect for the file system are used. Use the **mmlspolicy** command to display the current policy rules.

- q** When specified, **mmapplypolicy** dispatches bunches of files from the file lists specified by the **-r** option in a round-robin fashion, so that the multithreaded (**-m**) and node parallel (**-N**) policy execution works on all the file lists "at the same time." When **-q** is not specified, policy execution works on the file lists sequentially. In either case bunches of files are dispatched for parallel execution to multiple threads (**-m**) on each of the possibly multiple nodes (**-N**).

-r FileListPathname...

Specifies one or more file lists of files for policy execution. The file lists that were used as input for **-r** were created by issuing **mmapplypolicy** with the **-I prepare** flag. You can specify several file lists by doing one of the following:

- Provide the path name of a directory of file lists, **or**
- Specify the **-r** option several times, each time with the path name of a different file list.

You can use this parameter to logically continue where **mmapplypolicy** left off when you specified the **-I prepare** option. To do this, invoke **mmapplypolicy** with all the same parameters and options (except the **-I prepare** option), and now substitute the **-r** option for **-f**. In between the invocations, you can process, reorder, filter, or edit the file lists that were created when you invoked **-I prepare**. You can specify any or all of the resulting file lists with the **-r** option.

The format of the records in each file list file can be expressed as:

```
iAggregate:WEIGHT:INODE:GENERATION:SIZE:iRule:resourceId:attr_flags:  
path-length!PATH_NAME:pool-length!POOL_NAME  
[;show-length>!SHOW]end-of-record-character
```

For more information about file list records, refer to the *IBM Spectrum Scale: Administration Guide*.

-S SnapshotName

Specifies the name of a global snapshot for file system backup operations or for migrating snapshot data. The name appears as a subdirectory of the **.snapshots** directory in the file system root and can be found with the **mmlssnapshot** command.

Note: GPFS snapshots are read-only. Do not use deletion rules with **-S SnapshotName**.

| **Note:** Do not use the **-S** option to scan files in a fileset snapshot. Instead, specify the full path of a directory in the snapshot as the first parameter of the command, as in the following example:

```
| mmapplypolicy <directory_path> -P policyfile ...
```

-s LocalWorkDirectory

Specifies the directory to be used for temporary storage during **mmapplypolicy** command processing. The default directory is **/tmp**. The **mmapplypolicy** command stores lists of candidate and chosen files in temporary files within this directory.

When you execute **mmapplypolicy**, it creates several temporary files and file lists. If the specified file system or directories contain many files, this can require a significant amount of temporary storage. The required storage is proportional to the number of files (NF) being acted on and the average length of the path name to each file (AVPL). To make a rough estimate of the space required, estimate NF and assume an AVPL of 80 bytes. With an AVPL of 80, the space required is roughly (300 X NF) bytes of temporary space.

mmapplypolicy

--choice-algorithm {best | exact | fast}

Specifies one of the following types of algorithms that the policy engine is to use when selecting candidate files:

best

Chooses the optimal method based on the rest of the input parameters.

exact

Sorts all of the candidate files completely by weight, then serially considers each file from highest weight to lowest weight, choosing feasible candidates for migration, deletion, or listing according to any applicable rule LIMITs and current storage-pool occupancy. This is the default.

fast

Works together with the parallelized -g /shared-tmp -N node-list selection method. The **fast** choice method does not completely sort the candidates by weight. It uses a combination of statistical, heuristic, and parallel computing methods to favor higher weight candidate files over those of lower weight, but the set of chosen candidates may be somewhat different than those of the **exact** method, and the order in which the candidates are migrated, deleted, or listed is somewhat more random. The **fast** method uses statistics gathered during the policy evaluation phase. The **fast** choice method is especially fast when the collected statistics indicate that either all or none of the candidates are feasible.

--maxdepth MaxDirectoryDepth

Specifies how deeply in the hierarchy of the starting directory to apply policies to files. If this parameter is omitted, the command applies the policies to all the subdirectories of the starting directory. A value of 0 causes the policies to be applied only to the files in the starting directory.

--max-merge-files MaxFiles

Specifies the maximum number of files to be passed as input to the **sort** command for sorting and merging.

The **mmapplypolicy** command must do multiple, potentially large sorts and merges of temporary files as part of its processing. The **--max-merge-files** parameter specifies the maximum number of files that the **mmapplypolicy** command passes as input to a single **sort** command for sorting and merging.

If you set this value too high, the result can be excessive memory usage. The operating system can respond by terminating some of the **sort** processes, which causes the **mmapplypolicy** command to run for a longer time or to return with an error.

The default value of this parameter is 12. In general, it is a good practice to accept the default value of this parameter or to test carefully if you specify an overriding value.

See the related parameters **--max-sort-bytes** and **--sort-buffer-size**.

--max-sort-bytes MaxBytes

Specifies the maximum number of bytes to be passed as input files to the **sort** command. This parameter does not apply to merges in which each of the input files has already been sorted.

The **mmapplypolicy** command must do multiple, potentially large sorts and merges of temporary files as part of its processing. The **--max-sort-bytes** parameter specifies the maximum number of bytes that the **mmapplypolicy** command can pass to a single instance of the **sort** command in one or more files.

If you set this value too high, the result can be excessive memory usage. The operating system can respond by terminating some of the **sort** processes, which causes the **mmapplypolicy** command to run for a longer time or to return with an error.

The default value of this parameter is 411 MB. In general, it is a good practice to accept the default value of this parameter or to test carefully if you specify an overriding value.

See the related parameters **--max-merge-files** and **--sort-buffer-size**.

--scope {filesystem | inodespace | fileset}

If *Device* is specified, the directory traversal starts at the root of the file system. **--scope** indicates one of the following levels of scope to be applied to the policy scan:

filesystem

The scan will involve the objects in the entire file system subtree pointed to by the *Directory* parameter. This is the default.

fileset

The scope of the scan is limited to the objects in the same fileset as the directory pointed to by the *Directory* parameter.

inodespace

The scope is limited to objects in the same single inode space from which the directory pointed to by the *Directory* parameter is allocated. The scan may span more than one fileset, if those filesets share the same inode space.

--qos QoSClass

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “[mmchqos command](#)” on page 206. Specify one of the following QoS classes:

maintenance

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

other This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

--other-sort-options SortOptions

Passes options to the sort command (either the default sort command provided with the operating system, or an alternative sort command specified by the **--sort-command** parameter).

--single-instance

Ensures that, for the specified file system, only one instance of **mmapplypolicy** invoked with the **--single-instance** option can execute at one time. If another instance of **mmapplypolicy** invoked with the **--single-instance** option is currently executing, this invocation will do nothing but terminate.

--sort-buffer-size Size

Sets the sort-buffer size that is passed to the **sort** command. This parameter limits memory usage by the **sort** commands that the **mmapplypolicy** command calls to do sorts and merges.

The **mmapplypolicy** command must do multiple, potentially large sorts and merges of temporary files as part of its processing. It calls the operating-system **sort** command each time that it must do a sort. It can have multiple instances of the **sort** command running at the same time. If the numbers of items to be sorted are very large, the result can be excessive memory usage. The operating system can respond by terminating some of the sort processes, which causes the **mmapplypolicy** command to run for a longer time or to return with an error.

To prevent excessive memory consumption, you can set the **--sort-buffer-size** parameter to a lower value than its default. The **--sort-buffer-size** parameter is the value that the **mmapplypolicy** command passes to the **sort** command in the **buffer-size** parameter. The default value is 8%. If you need a lower value, you might set it to 5%.

You can specify the **--sort-buffer-size** parameter in any format that the **sort** program's **buffer-size** parameter accepts, such as "5%" or "1M".

mmapplypolicy

In general, accept the default value of this parameter unless the system has excessive memory consumption that is attributable to large sort operations by the **mmapplypolicy** command.

See the related parameters **--max-merge-files** and **max-sort-bytes**.

--sort-command SortCommand

Specifies the fully-qualified path name for a Posix-compliant sort command to be used instead of the default, standard command provided by the operating system.

Before specifying an alternative sort command (and for information about a suggested sort command), see *Improving performance with the --sort-command parameter in IBM Spectrum Scale: Administration Guide*.

--split-filelists-by-weight

Specifies that each of the generated file lists contain elements with the same **WEIGHT** value. This can be useful in conjunction with the **LIST** rule and the **WEIGHT (DIRECTORY_HASH)** clause. In this case, each generated list will contain files from the same directory.

Note: If you want all of the files from a given directory to appear in just one list, you might have to specify a sufficiently large **-B** value.

--split-margin n.n

A floating-point number that specifies the percentage within which the **fast**-choice algorithm is allowed to deviate from the **LIMIT** and **THRESHOLD** targets specified by the policy rules. For example if you specified a **THRESHOLD** number of 80% and a split-margin value of 0.2, the **fast**-choice algorithm could finish choosing files when it reached 80.2%, or it might choose files that bring the occupancy down to 79.8%. A nonzero value for split-margin can greatly accelerate the execution of the **fast**-choice algorithm when there are many small files. The default is 0.2.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmapplypolicy** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

Examples

1. This command displays the actions that would occur if a policy were applied, but does not apply the policy at this time:

```
mmapplypolicy fs1 -P policyfile -I test
```

The system displays output similar to:

```
[I] GPFS current data pool utilization in KB and %
sp1    9728    19531264    0.049807%
sp2    4608    19531264    0.023593%
system 105216  19531264   0.538706%
[I] Loaded policy rules from fs1.pol.
Evaluating MIGRATE/DELETE/EXCLUDE rules with CURRENT_TIMESTAMP = 2009-02-27@20:00:22 UTC
parsed 2 Placement Rules, 0 Restore Rules, 3 Migrate/Delete/Exclude Rules,
      0 List Rules, 0 External Pool/List Rules
RULE 'sp1' SET POOL 'sp1' WHERE name like '%.sp1' or name like '%.tmp'
```

```

RULE 'default' SET POOL 'system'

RULE 'exclude *.save files' EXCLUDE WHERE NAME LIKE '%.save'

/* Deletion rule */
RULE 'delete' DELETE FROM POOL 'sp1' WHERE NAME LIKE '%.tmp'

/* Migration rule */
RULE 'migration to system pool' MIGRATE FROM POOL 'sp1' TO POOL 'system' WHERE NAME LIKE '%sp1%'
[I] Directories scan: 11 files, 1 directories, 0 other objects, 0 'skipped' files and/or errors.
[I] Inodes scan: 11 files, 1 directories, 0 other objects, 0 'skipped' files and/or errors.
[I] Summary of Rule Applicability and File Choices:
  Rule# Hit_Cnt KB_Hit Chosen KB_Chosen KB_Inv Rule
  0      3       1536   0       0       RULE 'exclude *.save files' EXCLUDE WHERE(.)
  1      3       1536   3       1536   0       RULE 'delete' DELETE FROM POOL 'sp1' WHERE(.)
  2      2       1024   2       1024   0       RULE 'migration to system pool' MIGRATE FROM POOL \
                                             'sp1' TO POOL 'system' WHERE(.)

[I] Files with no applicable rules: 4.

```

```

[I] GPFS Policy Decisions and File Choice Totals:
  Choose to migrate 1024KB: 2 of 2 candidates;
  Choose to premigrate 0KB: 0 candidates;
  Already co-managed 0KB: 0 candidates;
  Choose to delete 1536KB: 3 of 3 candidates;
  Choose to list 0KB: 0 of 0 candidates;
  0KB of chosen data is illplaced or illreplicated;
Predicted Data Pool Utilization in KB and %:
sp1    7168    19531264    0.036700%
sp2    4608    19531264    0.023593%
system 106240   19531264   0.543948%

```

2. This command applies a policy immediately:

```
mmapplypolicy fs1 -P policyfile
```

The system displays output similar to:

```

[I] GPFS current data pool utilization in KB and %
sp1    9728    19531264    0.049807%
sp2    4608    19531264    0.023593%
system 105216   19531264   0.538706%
[I] Loaded policy rules from fs1.pol.
Evaluating MIGRATE/DELETE/EXCLUDE rules with CURRENT_TIMESTAMP = 2009-02-27@20:2
5:34 UTC
parsed 2 Placement Rules, 0 Restore Rules, 3 Migrate/Delete/Exclude Rules,
  0 List Rules, 0 External Pool/List Rules
RULE 'sp1' SET POOL 'sp1' WHERE name like '%.sp1' or name like '%.tmp'
RULE 'default' SET POOL 'system'

RULE 'exclude *.save files' EXCLUDE WHERE NAME LIKE '%.save'

/* Deletion rule */
RULE 'delete' DELETE FROM POOL 'sp1' WHERE NAME LIKE '%.tmp'

/* Migration rule */
RULE 'migration to system pool' MIGRATE FROM POOL 'sp1' TO POOL 'system' WHERE NAME LIKE '%sp1%'
[I] Directories scan: 11 files, 1 directories, 0 other objects, 0 'skipped' files and/or errors.
[I] Inodes scan: 11 files, 1 directories, 0 other objects, 0 'skipped' files and/or errors.
[I] Summary of Rule Applicability and File Choices:
  Rule# Hit_Cnt KB_Hit Chosen KB_Chosen KB_Inv Rule
  0      3       3072   0       0       RULE 'exclude *.save files' EXCLUDE WHERE(.)
  1      3       3072   3       3072   0       RULE 'delete' DELETE FROM POOL 'sp1' WHERE(.)
  2      2       2048   2       2048   0       RULE 'migration to system pool' MIGRATE FROM POOL \
                                             'sp1' TO POOL 'system' WHERE(.)

[I] Files with no applicable rules: 4.

```

mmapplypolicy

```
[I] GPFS Policy Decisions and File Choice Totals:  
Chose to migrate 2048KB: 2 of 2 candidates;  
Chose to premigrate 0KB: 0 candidates;  
Already co-managed 0KB: 0 candidates;  
Chose to delete 3072KB: 3 of 3 candidates;  
Chose to list 0KB: 0 of 0 candidates;  
0KB of chosen data is illplaced or illreplicated;  
Predicted Data Pool Utilization in KB and %:  
sp1    4608    19531264    0.023593%  
sp2    4608    19531264    0.023593%  
system 107264  19531264    0.549191%  
[I] A total of 5 files have been migrated, deleted or processed by an EXTERNAL E  
XEC/script;  
0 'skipped' files and/or errors.
```

Additional examples of GPFS policies and using the **mmapplypolicy** command are in the *IBM Spectrum Scale: Administration Guide*.

See also

- “**mmchpolicy** command” on page 201
- “**mmcrsnapshot** command” on page 267
- “**mmlspolicy** command” on page 426
- “**mmlssnapshot** command” on page 439
- “**mmsnapdir** command” on page 565

Location

/usr/lpp/mmfs/bin

mmauth command

Manages secure access to GPFS file systems.

Synopsis

```
mmauth genkey {new | commit | propagate [-N {Node[,Node...] | NodeFile | NodeClass}]}
```

or

```
mmauth add RemoteClusterName -k KeyFile [-l CipherList]
```

or

```
mmauth update RemoteClusterName {[ -C NewClusterName] [-k KeyFile] [-l CipherList]}
```

or

```
mmauth delete {RemoteClusterName | all}
```

or

```
mmauth grant {RemoteClusterName | all} -f {Device | all} [-a {rw | ro}] [-r {uid:gid | no}]
```

or

```
mmauth deny {RemoteClusterName | all} -f {Device | all}
```

or

```
mmauth show [RemoteClusterName | all | ciphers]
```

or

```
mmauth gencert --cname CanonicalName --cert ServerCertificateFile --out OutputKeystoreFile
--label ClientCertificateLabel [--pwd-file KeystorePasswordFile]
```

Availability

Available on all IBM Spectrum Scale editions.

Description

The **mmauth** command prepares a cluster to grant secure access to file systems owned locally. The **mmauth** command also prepares a cluster to receive secure access to file systems owned by another cluster. Use the **mmauth** command to generate a public/private key pair for the local cluster. A public/private key pair must be generated on both the cluster owning the file system and the cluster desiring access to the file system. The administrators of the clusters are responsible for exchanging the public portion of the public/private key pair. Use the **mmauth** command to add or delete permission for a cluster to mount file systems owned by the local cluster.

When a cluster generates a new public/private key pair, administrators of clusters participating in remote file system mounts are responsible for exchanging their respective public key file `/var/mmfsssl/id_rsa.pub` generated by this command.

The administrator of a cluster desiring to mount a file system from another cluster must provide the received key file as input to the **mmremotecluster** command. The administrator of a cluster allowing another cluster to mount a file system must provide the received key file to the **mmauth** command.

The keyword appearing after **mmauth** determines which action is performed:

mmauth

add

Adds a cluster and its associated public key to the list of clusters authorized to connect to this cluster for the purpose of mounting file systems owned by this cluster.

delete

Deletes a cluster and its associated public key from the list of clusters authorized to mount file systems owned by this cluster.

deny

Denies a cluster the authority to mount a specific file system owned by this cluster.

gencert

Creates a client keystore with the keys and certificates required to communicate with the ISKLM key server.

genkey

Controls the generation and propagation of the OpenSSL key files:

new

Generates a new public/private key pair for this cluster. The key pair is placed in **/var/mmfs/ssl**. This must be done at least once before **cipherList**, the GPFS configuration parameter that enables GPFS with OpenSSL, is set.

The new key is in addition to the currently in effect committed key. Both keys are accepted until the administrator runs **mmauth genkey commit**.

commit

Commits the new public/private key pair for this cluster. Once **mmauth genkey commit** is run, the old key pair will no longer be accepted, and remote clusters that have not updated their keys (by running **mmauth update** or **mmremotecluster update**) will be disconnected.

propagate

Ensures that the currently in effect key files are placed in **/var/mmfs/ssl** on the nodes specified with the **-N** parameter. This may be necessary if the key files are lost and **adminMode central** is in effect for the cluster.

grant

Allows a cluster to mount a specific file system owned by this cluster.

show

Shows the list of clusters authorized to mount file system owned by this cluster.

update

Updates the public key and other information associated with a cluster authorized to mount file systems owned by this cluster.

When the local cluster name (or '.') is specified, **mmauth update -l** can be used to set the *cipherList* value for the local cluster. Note that you cannot use this command to change the name of the local cluster. Use the **mmchcluster** command for this purpose.

Parameters

-N {Node[,Node...] | NodeFile | NodeClass}

Specifies the nodes on which the key files should be restored. The default is **-N all**.

For general information on how to specify node names, see *Specifying nodes as input to GPFS(tm) commands in IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

RemoteClusterName

Specifies the remote cluster name requesting access to local GPFS file systems.

a11

Indicates all remote clusters defined to the local cluster.

ciphers

Shows the supported ciphers.

Options**-a {rw | ro}**

Specifies the type of access allowed:

ro Specifies read-only access.

rw Specifies read/write access. This is the default.

-c NewClusterName

Specifies a new, fully-qualified cluster name for the already-defined cluster *RemoteClusterName*.

-f {Device | a11}

Specifies the device name for a file system owned by this cluster. The *Device* argument is required. If **all** is specified, the command applies to all file systems owned by this cluster at the time that the command is issued.

-k KeyFile

Specifies the public key file generated by the **mmauth** command in the cluster requesting to remotely mount the local GPFS file system.

-l CipherList

Sets the security mode for communications between the current cluster and the remote cluster that is specified in the *RemoteClusterName* parameter. There are three security modes:

EMPTY

The sending node and the receiving node do not authenticate each other, do not encrypt transmitted data, and do not check data integrity.

AUTHONLY

The sending and receiving nodes authenticate each other, but they do not encrypt transmitted data and do not check data integrity. This mode is the default in IBM Spectrum Scale V4.2 or later.

Cipher

The sending and receiving nodes authenticate each other, encrypt transmitted data, and check data integrity. To set this mode, you must specify the name of a supported cipher, such as AES128-GCM-SHA256.

For more information about the security mode and supported ciphers, see the topic *Security mode* in the *IBM Spectrum Scale: Administration Guide*.

-r {uid:gid | no}

Specifies a root credentials remapping (*root squash*) option. The UID and GID of all processes with root credentials from the remote cluster will be remapped to the specified values. The default is not to remap the root UID and GID. The *uid* and *gid* must be specified as unsigned integers or as symbolic names that can be resolved by the operating system to a valid UID and GID. Specifying **no**, **off**, or **DEFAULT** turns off the remapping.

For more information, see the *IBM Spectrum Scale: Administration Guide* and search on *root squash*.

--cname CanonicalName

Specifies the canonical name of the client used in the certificate.

--cert ServerCertificateFile

Specifies the path name to a file containing an ISKLM certificate.

mmauth

--out OutputKeystoreFile

Specifies the path name for the file that will contain the keystore.

--pwd-file KeystorePasswordFile

Specifies the keystore password file. If omitted, you will be prompted to enter the keystore password. A maximum of 20 characters are allowed. The **--pwd** *KeystorePassword* option is considered deprecated and may be removed in a future release.

--label ClientCertificateLabel

Specifies the label of the client certificate within the keystore. A maximum of 20 characters are allowed.

Exit status

- 0 Successful completion. After a successful completion of the **mmauth** command, the configuration change request will have been propagated to all nodes in the cluster.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmauth** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

Examples

1. This is an example of an **mmauth genkey new** command:

```
mmauth genkey new
```

The output is similar to this:

```
Generating RSA private key, 512 bit long modulus
.....+++++++
e is 65537 (0x10001)
mmauth: Command successfully completed
mmauth: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

2. This is an example of an **mmauth genkey commit** command:

```
mmauth genkey commit
```

The output is similar to this:

```
mmauth: Command successfully completed
mmauth: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

3. This is an example of an **mmauth add** command:

```
mmauth add clustA.kgn.ibm.com -k /u/admin/keys/clustA.pub
```

The output is similar to this:

```
mmauth: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

4. This is an example of an **mmauth update** command:

```
mmauth update clustA.kgn.ibm.com -k /u/admin/keys/clustA_new.pub
```

The output is similar to this:

```
mmauth: Propagating the cluster configuration data to all  
affected nodes. This is an asynchronous process.
```

5. This is an example of an **mmauth grant** command:

```
mmauth grant clustA.kgn.ibm.com -f /dev/gpfs1 -a ro
```

The output is similar to this:

```
mmauth: Propagating the cluster configuration data to all  
affected nodes. This is an asynchronous process.
```

6. This is an example of an **mmauth show** command:

```
mmauth show all
```

The output is similar to this:

```
Cluster name: clustA.kgn.ibm.com  
Cipher list: AES128-SHA  
SHA digest: a3917c8282fca7a27d951566940768dcd241902b  
File system access: gpfs1 (ro)  
Cluster name: clustB.kgn.ibm.com (this cluster)  
Cipher list: AES128-SHA  
SHA digest: 6ba5e3c1038246fe30f3fc8c1181fbb2130d7a8a  
SHA digest (new): 3c1038246fe30f3fc8c1181fbb2130d7a8a9ab4d  
File system access: (all rw)
```

For **clustB.kgn.ibm.com**, the **mmauth genkey new** command has been issued, but the **mmauth genkey commit** command has not yet been issued.

For more information on the SHA digest, see *The SHA digest in IBM Spectrum Scale: Problem Determination Guide*.

See also

- “**mmremoteefs** command” on page 514
- “**mmremotecluster** command” on page 511

See also the topic about accessing GPFS file systems from other GPFS clusters in the *IBM Spectrum Scale: Administration Guide*.

Location

/usr/lpp/mmf/bin

mmbackup

mmbackup command

Performs a backup of a GPFS file system or independent fileset to an IBM Spectrum Protect server.

Synopsis

```
mmbackup {Device | Directory} [-t {full | incremental}]
  [-N {Node[,Node...]} | NodeFile | NodeClass}]
  [-g GlobalWorkDirectory] [-s LocalWorkDirectory]
  [-S SnapshotName] [-f] [-q] [-v] [-d]
  [-a IscanfThreads] [-n DirThreadLevel]
  [-m ExecThreads | [[--expire-threads ExpireThreads] [--backup-threads BackupThreads]]]
  [-B MaxFiles | [[--max-backup-count MaxBackupCount] [--max-expire-count MaxExpireCount]]]
  [[--max-backup-size MaxBackupSize] [--qos QoSClass] [--quote | --noquote]
  [--rebuild] [--scope {filesystem | inodespace}]
  [--backup-migrated | --skip-migrated] [--tsm-servers TSMServer[,TSMServer...]]
  [--tsm-errorlog TSMServerLogFile] [-L n] [-P PolicyFile]
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

Use the **mmbackup** command to back up the user data from a GPFS file system or independent fileset to a TSM server or servers. The **mmbackup** command can be used only to back up file systems that are owned by the local cluster.

Attention: In GPFS V4.1 and later, a full backup (-t full) with **mmbackup** is required if a full backup was never performed with GPFS 3.3 or later. For more information, see the topic *File systems backed up using GPFS 3.2 or earlier versions of mmbackup* in the *IBM Spectrum Scale: Administration Guide*.

The IBM Spectrum Protect Backup-Archive client must be installed and at the same version on all the nodes that are executing the **mmbackup** command or are named in a node specification with -N. For more information about TSM requirements for the **mmbackup** command, see the topic *GPFS port usage* in the *IBM Spectrum Scale: Administration Guide*.

- | You can run multiple simultaneous instances of **mmbackup** if they are on different file systems or filesets.

See the topic *Backup considerations for using IBM Spectrum Protect* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide* and the topic *Configuration reference for using IBM Spectrum Protect with IBM Spectrum Scale* in the *IBM Spectrum Scale: Administration Guide*.

Parameters

Device

The device name for the file system to be backed up. File system names do not need to be fully qualified. **fs0** is as acceptable as **/dev/fs0**.

Directory

Specifies either the mount point of a GPFS file system or the path to an independent fileset root to back up. **/gpfs/fs0** can be used to specify the GPFS file system called **fs0**.

Note: A snapshot directory path is not permitted. To back up a snapshot, use **-S SnapshotName** instead. Do not use a subdirectory path as this leads to inconsistent backups.

-t {full | incremental}

Specifies whether to perform a full backup of all of the files in the file system, or an incremental backup of only those files that changed since the last backup was performed. The default is an **incremental** backup.

A full backup expires all GPFS 3.2 format TSM inventory if all previous backups were incremental and the GPFS 3.2 backup format was in use previously. If **mmbackup** on GPFS 3.2 was first used, and then only incremental backups were done on GPFS 3.4 or 3.5, then TSM will still contain 3.2 format backup inventory. This inventory will automatically be marked for expiration by **mmbackup** after a successful full or incremental backup.

Note: Do not use **-t full** with an unlinked fileset. Use an **EXCLUDE** statement to exclude directories or files from the backup operation.

-N {Node[,Node...] | NodeFile | NodeClass}

Specifies the list of nodes that run parallel instances of the backup process. The IBM Spectrum Protect Backup-Archive client must be installed on all nodes that are specified with this parameter. This command supports all defined node classes. The default is to run only on the node where the **mmbackup** command is running or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

-g GlobalWorkDirectory

Specifies the directory to be used for temporary files that need to be shared between the **mmbackup** worker nodes. Defaults to the value specified with the **-s** option or **/tmp**.

-s LocalWorkDirectory

Specifies the local directory to be used for temporary storage during **mmbackup** command processing. The default directory is **/tmp**. A *LocalWorkDirectory* must exist on each node that is used to run the **mmbackup** command or specified in a node specification with **-N**.

-S SnapshotName

Specifies the name of a global snapshot for any backup operations or a fileset-level snapshot if **--scope inodespace** is also specified for a fileset backup. The snapshot must be created before the **mmbackup** command is used. Snapshot names can be found with the **mm1ssnapshot** command. If a fileset is not present in the named snapshot and fileset-level backup is invoked with **--scope inodespace**, an error is returned. If a fileset-level snapshot name is used with a file system backup, an error is returned.

The use of **-S SnapshotName** is recommended because it provides **mmbackup** and TSM a consistent view of GPFS from which to perform backup. Deletion of the snapshot used for backup can be performed using the **mmdeletSnapshot** command after **mmbackup** completes.

-f Specifies that processing should continue when unlinked filesets are detected. All files that belong to unlinked filesets are ignored.

Note: Because **-f** has a large impact on performance, avoid using it unless performing a backup operation with unlinked filesets is necessary.

-q Performs a query operation before beginning **mmbackup**. The TSM server may have data stored already that is not recognized as having been backed up by **mmbackup** and its own shadow database. To properly compute the set of files that currently need to be backed up, **mmbackup** can perform a TSM query and process the results to update its shadow database. Use the **-q** switch to perform this query and then immediately commence the requested backup operation.

Note: Do not use **-q** with the **--rebuild** parameter.

-v Specifies verbose message output. Use this flag to cause **mmbackup** to issue more verbose messages about its processing. See also “Environment” on page 76.

-d Gathers debugging information that is useful to the IBM Support Center when diagnosing problems.

mmbackup

-a *IscanThreads*

Specifies the number of threads and sort pipelines each node runs during parallel inode scan and policy evaluation. The default value is 2.

-n *DirThreadLevel*

Specifies the number of threads that are created and dispatched within each **mmapplypolicy** process during the directory scan phase. The default value is 24.

-m *ExecThreads*

Specifies the number of threads created and dispatched within each **mmapplypolicy** process during the policy execution phase. The default value for **mmapplypolicy** is 24; however, the default value for **mmbackup** is 1. This option cannot be used with the **--expire-threads** and **--backup-threads** options.

--expire-threads *ExpireThreads*

Specifies the number of worker threads permitted on each node to perform **dsmc expire** tasks in parallel. This option cannot be used with the **-m** option. Valid values are 1 - 32. The default value is 4.

--backup-threads *BackupThreads*

Specifies the number of worker threads permitted on each node to perform **dsmc selective** or **dsmc incremental** tasks in parallel. This option cannot be used with the **-m** option. Valid values are 1 - 32. The default value is 1.

-B *MaxFiles*

Specifies the maximum number of objects in a bunch for each invocation of the **mmapplypolicy** EXEC script. If this option is not specified, the ideal bunch count is automatically computed. Valid values are 100 - 2147483647.

--max-backup-count *MaxBackupCount*

Specifies the maximum number of objects in a bunch for each **dsmc selective** or **dsmc incremental** command. This option cannot be used with the **-B** option. Valid values are 100 - 2147483647.

--max-expire-count *MaxExpireCount*

Specifies the maximum number of objects in a bunch for each **dsmc expire** command. This option cannot be used with the **-B** option. Valid values are 100 - 2147483647.

--max-backup-size *MaxBackupSize*

Specifies a policy limit on the content size in kilobytes for each **dsmc selective** or **dsmc incremental** command.

--qos *QoSClass*

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos** command” on page 206. Specify one of the following QoS classes:

maintenance

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

other This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

--quote | --noquote

Specifies whether to decorate (or not to decorate) file-list entries with quotation marks. The **mmbackup** command uses file lists to convey the lists of files to the IBM Spectrum Protect Backup-Archive client program. The file lists may or may not require each file name to be surrounded by quotation marks

depending on TSM client configuration options. If certain TSM client configuration options are in use, the quotation marks should not be added to file-list entries. Use the **--noquote** option in these instances.

--rebuild

Specifies whether to rebuild the **mmbackup** shadow database from the inventory of the TSM server. This option is similar to the **-q** option; however, no backup operation proceeds after the shadow database is rebuilt. This option should be used if the shadow database of **mmbackup** is known to be out of date, but the rebuilding operation must be done at a time when the TSM server is less loaded than during the normal time **mmbackup** is run.

If there are backup files with the old snapshot name */Device/.snapshots/.mmbuSnapshot* in the inventory of the TSM server, those files will be expired from the TSM server after the shadow database is rebuilt and any successful incremental or full backup completes.

Note: Do not use **--rebuild** with the **-q** parameter.

--scope {filesystem | inodespace}

Specifies that one of the following traversal scopes be applied to the policy scan and backup candidate selection:

filesystem

Scans all the objects in the file system specified by *Device* or mounted at the *Directory* specified. This is the default behavior.

inodespace

Specifies that the scan is limited in scope to objects in the same single inode space from which the *Directory* is allocated. The scan might span more than one fileset if those filesets share inode space; for example, dependent filesets.

| **--backup-migrated | --skip-migrated**

| HSM migrated objects are normally skipped even if they need to be backed up.

| **--backup-migrated**

| When specified, HSM migrated objects that need to be backed up are sent for backup. This can cause an HSM recall and other associated costs. When the same IBM Spectrum Protect server instance is used for HSM and backup, this option protects migrated, changed objects more efficiently. These objects are copied internally on the IBM Spectrum Protect server from the HSM storage pool directly to the BACKUP storage pool without incurring object recalls and their local storage costs. After backup, these objects will remain migrated. This occurs only if the object has no ACL or XAttrs attached.

| **--skip-migrated**

| When specified, HSM migrated objects that need to be backed up are not sent for backup. The pathnames of these objects are listed in a file, *mmbackup.hsmMigFiles.TSMServer*, located in the *Directory/Device* being backed up. The pathnames that are listed can then be recalled by the administrator in a staged manner. This is the default action that is taken for HSM migrated objects.

--tsm-servers TSMServer[,TSMServer...]

Specifies the name of the TSM server or servers to be used for this backup. The TSM servers specified are each used for the backup task specified.

If this option is not specified, the **mmbackup** command will backup to the servers that are specified in the **dsm.sys** file.

--tsm-errorlog TSMErrorLogFile

Specifies the pathname of the log file to pass to IBM Spectrum Protect Backup-Archive client commands.

mmbackup

-L n

Controls the level of information displayed by the **mmbackup** command. Larger values indicate the display of more detailed information. *n* should be one of the following values:

- 0** Displays only serious errors. This is the default.
- 1** Displays some information as the command runs, but not for each file.
- 2** Displays each chosen file and the scheduled migration or deletion action.
- 3** Displays the same information as 2, plus each candidate file and the applicable rule.
- 4** Displays the same information as 3, plus each explicitly EXCLUDEed or LISTed file, and the applicable rule.
- 5** Displays the same information as 4, plus the attributes of candidate and EXCLUDEed or LISTed files.
- 6** Displays the same information as 5, plus non-candidate files and their attributes.

-P PolicyFile

Specifies a customized policy rules file for the backup.

Environment

The behavior of **mmbackup** can be influenced by several environment variables when set.

Variables that apply to IBM Spectrum Protect Backup-Archive client program dsmc

MMBACKUP_DSMC_MISC

The value of this variable is passed as arguments to **dsmc restore** and **dsmc query {backup,incl,excl,session}** commands.

MMBACKUP_DSMC_BACKUP

The value of this variable is passed as arguments to **dsmc**, **dsmc selective**, and **dsmc incremental** commands.

MMBACKUP_DSMC_EXPIRE

The value of this variable is passed as arguments to **dsmc expire** commands.

Variables that change mmbackup output progress reporting

MMBACKUP_PROGRESS_CONTENT

Controls what progress information is displayed to the user as **mmbackup** runs. It is a bit field with the following bit meanings:

0x01

Specifies that basic text progress for each server is to be displayed.

0x02

Specifies that additional text progress for phases within each server is to be displayed.

0x04

Specifies that numerical information about files being considered is to be displayed.

MMBACKUP_PROGRESS_INTERVAL

Controls how frequently status callouts are made. The value is the minimum number of seconds between calls to the **MMBACKUP_PROGRESS_CALLOUT** script or program. It does not affect how frequently messages are displayed, except for the messages of **MMBACKUP_PROGRESS_CONTENT** category 0x04.

MMBACKUP_PROGRESS_CALLOUT

Specifies the path to a program or script to be called with a formatted argument, as described in the topic **MMBACKUP_PROGRESS_CALLOUT environment variable** in the *IBM Spectrum Scale: Administration Guide*.

Variables that change mmbackup debugging facilities

In case of a failure, certain debugging and data collection can be enabled by setting the specified environment variable value.

DEBUGmmbackup

This variable controls what debugging features are enabled. It is interpreted as a bitmask with the following bit meanings:

0x001

Specifies that basic debug messages are printed to STDOUT. There are multiple components that comprise **mmbackup**, so the debug message prefixes can vary. Some examples include:

```
mmbackup:mbbackup.sh
DEBUGtsbackup33:
```

0x002

Specifies that temporary files are to be preserved for later analysis.

0x004

Specifies that all **dsmc** command output is to be mirrored to STDOUT.

DEBUGmmcmi

This variable controls debugging facilities in the **mmbackup** helper program **mmcmi**, which is used when the cluster **minReleaseLevel** is less than 3.5.0.11.

DEBUGtsbuhelper

This variable controls debugging facilities in the **mmbackup** helper program **tsbuhelper**, which is used when the cluster **minReleaseLevel** is greater than or equal to 3.5.0.11.

Variables that change mmbackup record locations

MMBACKUP_RECORD_ROOT

Specifies an alternative directory name for storing all temporary and permanent records for the backup. The directory name specified must be an existing directory and it cannot contain special characters (for example, a colon, semicolon, blank, tab, or comma).

The directory specified for **MMBACKUP_RECORD_ROOT** must be accessible from each node specified with the **-N** option.

Exit status

- 0** Successful completion. All of the eligible files were backed up.
- 1** Partially successful completion. Some files, but not all eligible files, were backed up. The shadow database or databases reflect the correct inventory of the TSM server. Invoke **mmbackup** again to complete the backup of eligible files.
- 2** A failure occurred that prevented backing up some or all files or recording any progress in the shadow database or databases. Correct any known problems and invoke **mmbackup** again to complete the backup of eligible files. If some files were backed up, using the **-q** or **--rebuild** option can help avoid backing up some files additional times.

Security

You must have root authority to run the **mmbackup** command.

The node on which the command is issued, and all other IBM Spectrum Protect Backup-Archive client nodes, must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

mmbbackup

Examples

1. To perform an incremental backup of the file system **gpfs0**, issue this command:

```
mmbbackup gpfs0
```

The system displays information similar to:

```
-----  
mmbbackup: Backup of /gpfs/gpfs0 begins at Mon Apr 7 15:37:50 EDT 2014.  
-----  
Mon Apr 7 15:38:04 2014 mmbbackup:Scanning file system gpfs0  
Mon Apr 7 15:38:14 2014 mmbbackup:Determining file system changes for gpfs0 [balok1].  
Mon Apr 7 15:38:14 2014 mmbbackup:changed=364, expired=0, unsupported=0 for server [balok1]  
Mon Apr 7 15:38:14 2014 mmbbackup:Sending files to the TSM server [364 changed, 0 expired].  
mmbbackup: TSM Summary Information:  
    Total number of objects inspected:      364  
    Total number of objects backed up:      364  
    Total number of objects updated:        0  
    Total number of objects rebound:        0  
    Total number of objects deleted:       0  
    Total number of objects expired:       0  
    Total number of objects failed:        0  
    Total number of bytes transferred:     2179695902  
-----  
mmbbackup: Backup of /gpfs/gpfs0 completed successfully at Mon Apr 7 15:41:09 EDT 2014.  
-----
```

2. To recreate a shadow database for **gpfs0**, issue this command:

```
mmbbackup gpfs0 --rebuild
```

The system displays information similar to:

```
-----  
mmbbackup: Shadow database rebuild of /gpfs/gpfs0 begins at Tue Apr 8 09:44:59 EDT 2014.  
-----  
Tue Apr 8 09:45:11 2014 mmbbackup:Querying files currently backed up in TSM server:balok1.  
Tue Apr 8 09:45:14 2014 mmbbackup:Built query data file from TSM server: balok1 rc = 0  
Tue Apr 8 09:45:17 2014 mmbbackup:Scanning file system gpfs0  
Tue Apr 8 09:45:26 2014 mmbbackup:Reconstructing previous shadow file /gpfs/gpfs0/.mmbbackupShadow.1.balok1 from  
query data for balok1  
Tue Apr 8 09:45:26 2014 mmbbackup:Done with shadow file database rebuilds  
-----  
mmbbackup: Shadow database rebuild of /gpfs/gpfs0 completed successfully at Tue Apr 8 09:45:26 EDT 2014.  
-----
```

3. To perform an incremental backup of the file system **gpfs0** with more progress information displayed, first issue this command:

```
export MMBACKUP_PROGRESS_CONTENT=3
```

Next, issue the **mmbbackup** command:

```
mmbbackup gpfs0
```

The system displays information similar to:

```
-----  
mmbbackup: Backup of /gpfs/gpfs0 begins at Mon Apr 7 16:02:28 EDT 2014.  
-----  
Mon Apr 7 16:02:33 2014 mmbbackup:Begin checking server and shadow file for balok1  
Mon Apr 7 16:02:37 2014 mmbbackup:Querying TSM server balok1 for options  
Mon Apr 7 16:02:40 2014 mmbbackup:Found old shadow DB for balok1 present in /gpfs/gpfs0/.mmbbackupShadow.1.balok1  
Mon Apr 7 16:02:40 2014 mmbbackup:Checking format version of old shadow DB  
Mon Apr 7 16:02:40 2014 mmbbackup:Found old shadow DB version: 1400  
Mon Apr 7 16:02:40 2014 mmbbackup:Previous shadow /gpfs/gpfs0/.mmbbackupShadow.1.balok1 state: present  
Mon Apr 7 16:02:40 2014 mmbbackup:Generating policy rules file:/var/mmfs/mmbbackup/.mmbbackupRules.gpfs0 to  
use /gpfs/gpfs0/.mmbbackupCfg/BAexecScript.gpfs0  
Mon Apr 7 16:02:42 2014 mmbbackup:Completed policy rule generation. 2 Exclude Dir directives, 1 Exclude File  
directives, 1 Include directives, 0 Warnings.  
Mon Apr 7 16:02:42 2014 mmbbackup:Scanning file system gpfs0  
Mon Apr 7 16:02:51 2014 mmbbackup:File system scan of gpfs0 is complete.  
Mon Apr 7 16:02:51 2014 mmbbackup:Calculating backup and expire lists for server balok1  
Mon Apr 7 16:02:51 2014 mmbbackup:Determining file system changes for gpfs0 [balok1].  
Mon Apr 7 16:02:51 2014 mmbbackup:changed=364, expired=0, unsupported=0 for server [balok1]  
Mon Apr 7 16:02:51 2014 mmbbackup:Finished calculating lists [364 changed, 0 expired] for server balok1.  
-----
```

```

Mon Apr  7 16:02:51 2014 mmbackup: Sending files to the TSM server [364 changed, 0 expired].
Mon Apr  7 16:02:51 2014 mmbackup: Performing backup operations
Mon Apr  7 16:05:40 2014 mmbackup: Completed policy backup run with 0 policy errors, 0 files failed, 0 severe
errors, returning rc=0.
Mon Apr  7 16:05:40 2014 mmbackup: Policy for backup returned 0 Highest TSM error 0
mmbackup: TSM Summary Information:
      Total number of objects inspected:      364
      Total number of objects backed up:      364
      Total number of objects updated:        0
      Total number of objects rebound:        0
      Total number of objects deleted:        0
      Total number of objects expired:        0
      Total number of objects failed:        0
      Total number of bytes transferred:    2179695902
Mon Apr  7 16:05:40 2014 mmbackup: analyzing: results from balok1.
Mon Apr  7 16:05:40 2014 mmbackup: Copying updated shadow file to the TSM server
Mon Apr  7 16:05:44 2014 mmbackup: Done working with files for TSM Server: balok1.
Mon Apr  7 16:05:44 2014 mmbackup: Completed backup and expire jobs.
Mon Apr  7 16:05:44 2014 mmbackup: TSM server balok1
      had 0 failures or excluded paths and returned 0.
      Its shadow database has been updated. Shadow DB state:updated
Mon Apr  7 16:05:44 2014 mmbackup: Completed successfully. exit 0
-----  

mmbackup: Backup of /gpfs/gpfs0 completed successfully at Mon Apr  7 16:05:44 EDT 2014.
-----
```

4. To perform an incremental backup of the objects in the inode space of the gpfs/testfs/infs2 directory to the **balok1** server, issue this command:

```
mmbackup /gpfs/testfs/infs2 -t incremental --scope inodespace --tsm-servers balok1
```

The system displays information similar to:

```

mmbackup: Backup of /gpfs/testfs/infs2 begins at Wed May 27 12:58:39 EDT 2015.
-----  

Wed May 27 12:58:48 2015 mmbackup: Scanning fileset testfs.indfs2
Wed May 27 12:58:53 2015 mmbackup: Determining fileset changes for testfs.indfs2 [balok1].
Wed May 27 12:58:53 2015 mmbackup: changed=2, expired=2, unsupported=0 for server [balok1]
Wed May 27 12:58:53 2015 mmbackup: Sending files to the TSM server [2 changed, 2 expired].
mmbackup: TSM Summary Information:
      Total number of objects inspected:      4
      Total number of objects backed up:      2
      Total number of objects updated:        0
      Total number of objects rebound:        0
      Total number of objects deleted:        0
      Total number of objects expired:        2
      Total number of objects failed:        0
      Total number of objects encrypted:     0
      Total number of bytes inspected:      53934
      Total number of bytes transferred:    53995
-----  

mmbackup: Backup of /gpfs/testfs/infs2 completed successfully at Wed May 27 12:59:31 EDT 2015.
-----
```

5. To perform an incremental backup of a global snapshot that is called **backupsnap6** to the **balok1** server, issue this command:

```
mmbackup testfs -t incremental -S backupsnap6 --scope filesystem --tsm-servers balok1
```

The system displays information similar to:

```

mmbackup: Backup of /gpfs/testfs begins at Wed May 27 13:08:45 EDT 2015.
-----  

Wed May 27 13:08:50 2015 mmbackup: Scanning file system testfs
Wed May 27 13:08:53 2015 mmbackup: Determining file system changes for testfs [balok1].
Wed May 27 13:08:53 2015 mmbackup: changed=130, expired=100, unsupported=0 for server [balok1]
Wed May 27 13:08:53 2015 mmbackup: Sending files to the TSM server [130 changed, 100 expired].
-----  

Wed May 27 13:08:59 2015 mmbackup: Policy for expiry returned 9 Highest TSM error 0
```

mmbackup

```
mmbackup: TSM Summary Information:  
    Total number of objects inspected:      230  
    Total number of objects backed up:       130  
    Total number of objects updated:         0  
    Total number of objects rebound:         0  
    Total number of objects deleted:        0  
    Total number of objects expired:        100  
    Total number of objects failed:         0  
    Total number of objects encrypted:      0  
    Total number of bytes inspected:       151552  
    Total number of bytes transferred:     135290  
-----  
mmbackup: Backup of /gpfs/testfs completed successfully at Wed May 27 13:09:05 EDT 2015.  
-----
```

See also

- “mmapplypolicy command” on page 56

Location

/usr/lpp/mmfs/bin

mmbackupconfig command

Collects GPFS file system configuration information.

Synopsis

mmbackupconfig *Device* **-o** *OutputFile*

Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

Description

The **mmbackupconfig** command, in conjunction with the **mmrestoreconfig** command, can be used to collect basic file system configuration information that can later be used to restore the file system. The configuration information backed up by this command includes block size, replication factors, number and size of disks, storage pool layout, filessets and junction points, policy rules, quota information, and a number of other file system attributes.

This command does not back up user data or individual file attributes.

For more information about the **mmimgbackup** and **mmimgrestore** commands, see the topic *Scale out Backup and Restore (SOBAR)* in the *IBM Spectrum Scale: Administration Guide*.

Parameters

Device

The device name of the file system to be backed up. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

-o *OutputFile*

The path name of a file to which the file system information is to be written. This file must be provided as input to the subsequent **mmrestoreconfig** command.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmbackupconfig** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

To backup file system **fsiam2** to output file **backup.config.fsiam2** issue:

```
mmbackupconfig fsiam2 -o backup.config.fsiam2
```

The system displays information similar to:

mmbackupconfig

```
mmbackupconfig: Processing file system fsiam2 ...
mmbackupconfig: Command successfully completed
```

See also

- “[mmimgbackup command](#)” on page 367
- “[mmimgrstore command](#)” on page 371
- “[mmrestoreconfig command](#)” on page 521

Location

/usr/lpp/mmfs/bin

mmblock command

Manages the iSCSI block service.

Synopsis

```
| mmblock host add -H | --host-alias HostAlias -i | --iscsi-name iscsiInitiator
  or
| mmblock host list [-Y]
  or
| mmblock host remove -H | --host-alias HostAlias
  or
| mmblock volume add -v|--volume-name VolumeName -f|--file FilePath
  or
| mmblock volume remove -v | --volume-name VolumeName
  or
| mmblock volume list [-Y]
  or
| mmblock map-volume add -H|--host-alias HostAlias -v|--volume-name VolumeName[-l|--lun LUNID]
  or
| mmblock map-volume remove -H|--host-alias HostAlias -v|--volume-name VolumeName
  or
| mmblock map-volume list {-H | --host-alias HostAlias | -a | -all} [-f | --file] [-Y]
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher.

Description

Important: The block service is intended to be used only for the remote booting of diskless nodes. This service is not intended for performance-critical block storage access, and is not supported for such usage. Contact IBM if you have questions about your use case and to get instructions for enabling the **mmblock** command. To contact IBM, perform one of the following steps:

- Contact an IBM Client Technical Specialist.
- Send an email to scale@us.ibm.com.

The environment in which you plan to use the block service must meet the following requirements:

- The environment is a valid CES environment.
- The block service is used only for diskless boot of an OS image.

Provide information that supports these two requirements when you contact IBM.

The **mmces** command supports the iSCSI protocol by the implementation of the BLOCK service. See the first set of examples at the end of this topic.

The **mmblock** command manages the BLOCK service and provides the following features:

mmblock

- Administering block hosts.
- Administering block volumes.
- Administering block volume mappings.

Parameters

host

Administers block hosts.

add

Defines a host.

-H | --host-alias HostAlias

Specifies the name of the host.

-i | --iscsi-name iscsiInitiator

Specifies the name of the iSCSI in the IQN or the EUI format.

list

Lists all the created hosts.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

remove

Removes a defined host.

-H | --host-alias HostAlias

Specifies the name of the host that must be deleted.

volume

Manages storage volume objects.

add

Creates a volume.

-v | --volume-name VolumeName

Specifies the name of the volume.

-f | --file FilePath

Specifies the filepath. The filepath must be absolute. For example, /ibm/gpfs/xyz can be provided as the absolute filepath. Before you run the command, preallocate the file. For example, use dd if=/dev/zero of=/ibm/gpfs/xyz bs=1M count=1024 to allocate a file of 1GiB in size.

remove

-v | --volume-name VolumeName

Specifies the name of the volume that must be removed.

list

Lists all the defined volumes.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

map-volume

Manages storage volume mappings.

add

Maps a volume to a host.

- H | --host-alias HostAlias**
Specifies the name of the host.
- v | --volume-name VolumeName**
Specifies the name of the volume.
- l | --lun LUN**
Specifies the LUN (Logical Unit Number) of the mapping. The value is an integer starting from 0. If LUN is not provided, the system selects an unused LUN starting from 0.

remove

Removes a mapping.

- H | --host-alias HostAlias**
Specifies the name of the host.
- v | --volume-name VolumeName**
Specifies the name of the volume.

list

Lists the mapping of a host.

- H | --host-alias HostAlias**
Specifies the name of the host that must be listed.
- a | -all**
Lists the map volume of all the hosts.
- f | --file**
Lists the file path of the volume.
- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

Exit status

0 Successful completion.

Nonzero

A failure occurred.

Security

You must have root authority to run the **mmblock** command.

The node on which the command is issued must be able to execute remote shell commands on any other CES node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

mmblock

Examples

The following examples show how to use the **mmblock** command to manage the block service:

1. To add a host:

```
mmblock host add --host-alias host0 --iscsi-name iqn.1994-05.com.example:7c9c9c3c93c1
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmblock host add: The Host 'host0' was created successfully.
```

2. To list all the created hosts:

```
mmblock host list
```

The system displays the following output:

Host Name	Initiator
host0	iqn.1994-05.com.example:7c9c9c3c93c1
host1	iqn.2001-04.com.example:diskarrays-sn-a8675309
host2	iqn.2001-04.com.example.storage:tape.sys1.xyz
host4	iqn.1995-11.com.example.ssp:customers.4567.disks.107
host5	eui.02004567a425678d

3. To remove a defined host:

```
mmblock host remove --host-alias host0
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmblock: The Host 'host0' was removed successfully.
```

4. To manage the storage volume objects:

```
mmblock volume add --volume-name vol1 --file /ibm/gpfs/VOL1
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmblock: The Volume 'vol1' was created successfully.
```

5. To remove a volume:

```
mmblock volume remove --volume-name vol1
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmblock: The Volume 'vol1' was removed successfully.
```

6. To list all the created volumes:

```
mmblock volume list
```

The system displays the following output:

Volume Name	Backend File
vol1	/ibm/gpfs/VOL1

7. To map a volume to a host:

```
mmblock map-volume add --host-alias host1 --volume-name vol1
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmcesblockcmd: The Volume 'vol1' was mapped at LUN 1 of Host 'host1' successfully.
```

8. To remove a mapping:

```
mmblock map-volume remove --host-alias host1 --volume-name vol1
```

The system displays the following output:

```
Configuration successfully updated. Now apply settings on the Protocol nodes.  
192.168.40.68: Done  
mmblock: The Volume 'vol1' was removed from LUN 1 of Host 'host1' successfully.
```

9. To list mapping of a host:

```
mmblock map-volume list --host-alias host1
```

The system displays the following output:

LUN	Volume Name
1	vol1

Location

/usr/lpp/mmfs/bin

mmbuildgpl

mmbuildgpl command

Manages prerequisite packages for Linux and builds the GPFS portability layer.

Synopsis

```
mmbuildgpl [--quiet] [--build-package] [-v]
```

Availability

Available on all IBM Spectrum Scale editions. Available and needed only on Linux.

Description

Use the **mmbuildgpl** command to manage and verify prerequisite packages for Linux and build the GPFS portability layer. If all packages are installed correctly, **mmbuildgpl** builds the GPFS portability layer. If any packages are missing, the package names are displayed. The missing packages can be installed manually.

Parameters

--quiet

Specifies that when there are any missing packages, the **mmbuildgpl** command installs the prerequisite packages automatically by using the default package manager.

--build-package

Builds an installable package (**gpfs.gplbin**) for the portability layer binaries after compilation is successful. This option builds an RPM package on SLES and RHEL Linux and a Debian package on Debian and Ubuntu Linux.

When the command finishes, it displays the location of the generated package similar to:

```
Wrote: /root/rpmbuild/RPMS/x86_64/gpfs.gplbin-3.10.0-229.e17.x86_64-4.2.3-0.x86_64.rpm
```

or

```
Wrote: /tmp/deb/gpfs.gplbin_4.2.3-0_amd64.deb
```

You can then copy the generated package to other machines for deployment. The generated package can *only* be deployed to machines with identical architecture, distribution level, Linux kernel, and IBM Spectrum Scale maintenance level. If the **gpfs.gplbin** package is installed, installing the **gpfs.gpl** package is not required.

Note: During the package generation, temporary files are written to the **/tmp/rpm** or **/tmp/deb** directory, so be sure there is sufficient space available. By default, the generated package goes to **/usr/src/packages/RPMS/<arch>** for SUSE Linux Enterprise Server, **/usr/src/redhat/RPMS/<arch>** for Red Hat Enterprise Linux, and **/tmp/deb** for Debian and Ubuntu Linux.

-v Specifies that the output is verbose and contains information for debugging purposes.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmbuildgpl** command.

Examples

To build the GPFS portability layer, issue this command:

```
mmbuildgpl
```

The system displays information similar to:

```
-----  
mmbuildgpl: Building GPL module begins at Fri Jun 13 16:37:50 EDT 2014.  
-----  
  
Verifying Kernel Header...  
  kernel version = 3001300 (3.0.13-0.27)  
  kernel module dir = /lib/modules/3.0.13-0.27-default/build/include  
  kernel source dir = /usr/src/linux-3.0.13-0.27/include  
  Found a valid kernel include directory: /lib/modules/3.0.13-0.27-default/build/include  
Verifying Compiler...  
  make is present at /usr/bin/make  
  cpp is present at /usr/bin/cpp  
  gcc is present at /usr/bin/gcc  
  g++ is present at /usr/bin/g++  
  ld is present at /usr/bin/ld  
make World ...  
make Install ...  
-----  
mmbuildgpl: Building GPL module completed successfully at Fri Jun 13 16:39:08 EDT 2014.  
-----
```

See also

- *Building the GPFS portability layer on Linux nodes in IBM Spectrum Scale: Concepts, Planning, and Installation Guide.*

Location

/usr/lpp/mmfs/bin

mmcallhome

mmcallhome command

Manages the call home operations.

Synopsis

```
mmcallhome group add groupName server [--node {all | childNode [,childNode...]}]
```

or

```
| mmcallhome group list [--long] [-Y]
```

or

```
mmcallhome group delete GroupName
```

or

```
mmcallhome group auto [--server {all | {ServerName[,{ServerName...]} } | --force |  
  |--enable [LICENSE | ACCEPT] | --disable ]]
```

or

```
| mmcallhome capability list [-Y]
```

or

```
mmcallhome capability enable
```

or

```
mmcallhome capability disable
```

or

```
| mmcallhome info list [-Y]
```

or

```
mmcallhome info change [ --customer-name CustomerName | --customer-id CustomerId  
  | --email Email | --country-code CountryCode ]
```

or

```
mmcallhome proxy enable [--with-proxy-auth]
```

or

```
mmcallhome proxy disable
```

or

```
| mmcallhome proxy list [-Y]
```

or

```
mmcallhome proxy change [ --proxy-location ProxyLocation | --proxy-port ProxyPort  
  | --proxy-username ProxyUsername | --proxy-password ProxyPassword ]
```

or

```
| mmcallhome schedule list [-Y]
```

or

```
mmcallhome schedule add --task { DAILY | WEEKLY }
```

or

```
mmcallhome schedule delete --task { DAILY | WEEKLY }
```

or

```
mmcallhome run GatherSend --task { DAILY | WEEKLY }
```

or

```
mmcallhome run SendFile --file
```

or

```
mmcallhome status list [ --task{ DAILY | WEEKLY | sendfile} ]  
| [-n num] [--verbose] [-Y]
```

or

```
mmcallhome status delete {--task{ DAILY | WEEKLY | sendfile} | --startTime time |  
--startTImeBefore time | --all }
```

or

```
mmcallhome test connection
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher.

Description

Use the **mmcallhome** command to configure, enable, run, schedule, and monitor call home related tasks in the GPFS cluster. The call home feature organizes the cluster in groups. This feature is configured for local and global configurations. The local configuration of a node is about creating the membership of the node in a group configuration. A node cannot be a member of more than one group. The call home command is executed on the node only if the node is a member of a call home group. If a global configuration is predefined, and a new call home group is created, then the newly created group inherits the predefined global configuration. The **proxy change**, **info change**, and **capability enable/disable** commands support only global configuration, whereas the **schedule** command can change only local configuration.

Note: After the installation of call home feature, no further configuration changes can be made. Even the **list** command can not be used for any configuration changes. Only if no configuration has been done previously for the call home feature, a default configuration can be done after the call home installation. The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms.

By using this command, predefined data from each node can be collected on a regular basis and uploaded to IBM. IBM support and development teams can use this data to understand how the customers are using IBM Spectrum Scale. In case of issues, the data can be referenced for problem analysis. The data can also possibly be used to provide advice to customers regarding failure prevention.

Since IBM Spectrum Scale consists of multiple nodes, the call home feature introduces the concept of the call home group to manage them. A *call home group* consists of one gateway node (which is defined as a call home node) and one or more client nodes (which are defined as call home child nodes). The call home node initiates the data collection from the call home child nodes and uploads data to IBM using the HTTPS protocol. Since the call home group can be configured independently, the group concept can be used for special conditions, such as for split clusters, that requires all the group members to be on the same side to avoid unnecessary data transfer over large distance. Also, a call home group can be mapped to a node group or other cluster specific attributes that needs to keep all Linux nodes in one group and

mmcallhome

all AIX nodes in the other group. The call home node needs to have access to the external network via port 443. The maximum number of nodes per group should not exceed 32. Multiple call home groups can be defined within an IBM Spectrum Scale cluster.

For more information about the call home feature, see the *Monitoring the IBM Spectrum Scale(tm) system remotely by using call home* section in *IBM Spectrum Scale: Problem Determination Guide*.

Parameters

group

Manages topology with one of the following actions:

add

Creates a call home group, which is a group of nodes consisting of one call home node and multiple call home child nodes. Multiple call home groups can be configured within a GPFS cluster.

The call home node initiates data collection within the call home group and uploads the data package to the IBM server.

group

Specifies the name of the call home group.

Note: The group name can consist of any alphanumeric characters and these non-alphanumeric characters: '-', '_', and '..'.

Important: The group name cannot be `global`. Call home uses `global` as the default name for the group that contains the global values which are applied to all groups.

server

Specifies the name of the call home server belonging to the call home group.

Note: The server name can consist of any alphanumeric characters and these non-alphanumeric characters: '-', '_', and '.'.

--node childNode

Specifies the call home child nodes.

Note: The child node name can consist of any alphanumeric characters and these non-alphanumeric characters: '-', '_', and '.'.

--node all

Selects all Linux nodes in the GPFS cluster. If the number of nodes exceeds 32, the command will fail. When this parameter is omitted, only the call home node will be added to the child node. Additionally, call home node will be always added to the child node group.

list

Displays the configured call home groups.

--long

Displays the node names as long admin node names. Without `--long` the node names are listed as short admin node names.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of `mmclidecode`. Use the `mmclidecode` command to decode the field.

delete

Deletes the specified call home group.

GroupName

Specifies the name of the call home group that needs to be deleted.

auto

Enables automatic creation of a call home group.

--server ServerName

Specifies one or more call home servers. The server must be able to access the call home functionality through internet. If no server is specified, the system detects call home node automatically. In this scenario, the system checks if the detected node can access internet. If a server is specified, the defined nodes are used as call home nodes without any further check.

If a proxy is needed, specify the proxy by using the **mmcallhome proxy** command.

Note: If no --server (option) is specified, the system automatically identifies the possible call home node. If a proxy is defined and enabled in the global configuration, then the connection through proxy is used in identification of the possible call home node. If the proxy is disabled in the global configuration, then a direct connection from the nodes to the IBM EDGE server in the internet is used.

all

Specifies that each node is a call home server. This configuration avoids heavy work load for scheduled call home tasks on call home servers serving large groups. While using **all** option, it is recommended to use a proxy to access the internet.

Note: Multiple servers can be specified by repeating this option or by specifying a string, containing either a comma or a blank to separate the list of servers. If a group exists then the specified server must not be a member of that group. Each call home server must be able to access the internet either directly or via a proxy.

--force

Creates new groups after deleting the existing groups.

Note: If this option is defined and no server node has been specified, potential server nodes get detected automatically. If no server gets detected, an error is returned and the existing groups does not get deleted.

--enable

Enables the cluster for call home functionality. If no other option is defined, the **enable** parameter shows the license and asks for acceptance by default.

license

Shows license and terminate.

accept

Does not show license and assumes that the license is accepted.

Note: It is not possible to use this option multiple times or together with the --disable option.

--disable

Disable call home.

Note: All groups are disabled.

capability

Manages the overall call home activities with one of the following actions:

mmcallhome

list

Displays the configured customer information such as the current enable or disable status, call home node, and call home child nodes.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmc1idecode**. Use the **mmc1idecode** command to decode the field.

enable

Enables the call home service.

disable

Disables the currently running call home service.

info

Manages customer data with one of the following actions:

list

Displays the configured parameter values.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmc1idecode**. Use the **mmc1idecode** command to decode the field.

change

Sets parameter values.

[--key value]

Indicates a placeholder pointing to the table below.

Table 8. key-value

Key	Value
customer-name	Business/company name This name can consist of any alphanumeric characters and these non-alphanumeric characters: '-', '_', '.', ',', '''
customer-id	Customer ID This can consist of any alphanumeric characters and these non-alphanumeric characters: '-', '_', ''
email Customer	E-mail ID of the customer. All alphanumeric and non-alphanumeric characters are supported.
country-code 2 alphabet ISO	Country code (Example, US)

proxy

Configures proxy-related parameters with one of the following actions:

enable

Enables proxy access.

[--with-proxy-auth]

Enables user ID and password authentication to the proxy server.

disable

Disables proxy access.

list

Displays the currently configured proxy-related parameter values.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

change

Modifies the proxy configuration.

[--key value]

Indicates a placeholder pointing to the table below.

Table 9. key-value

Key	Value
proxy-location	Proxy server address (IP address/fully qualified domain name)
proxy-port	Proxy server port number
proxy-username	Proxy server user name This name can consist of any alphanumeric and non-alphanumeric characters.
proxy-password	Proxy server password This can consist of any alphanumeric and non-alphanumeric characters.

schedule

Configures scheduling of call home tasks with one of the following actions:

list

Displays the registered gather-send tasks. A gather-send task is a process that runs on the call home node to collect data from the child nodes and upload the data to the configured server. The gather-end configuration file will include information about what needs to be collected from the child nodes.

-Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

add

Registers the specified task to cron.

--task daily

Specifies the configuration file for the daily task.

--task weekly

Specifies the configuration file for the weekly task.

delete

Removes a daily or weekly task from cron with one of the following options:

--task daily

Specifies the daily task that needs to be removed from cron.

- **Started time:** Timestamp when the gather-send task is invoked.
- **Updated time:** Timestamp when the status is updated.
- **Status:** Success/minor error/failed/running.
- **RC or Step:** When the status is failed/success, the return code of the task is shown. See below for the return code description of the gather-send task. When the status is running, the step is shown. See below for the step description.
 - **Package file name:** Name of the created data package to upload.
 - **additional info:** Any additional info for the task.

Gather-send task return codes:

- 0 - Success
- 1 - Successfully uploaded after a few send retries
- 2 - Some gather commands failed but logs collected from all child nodes and successfully uploaded
- 3 - Could not collect logs from some nodes but the call home node created the data package and successfully uploaded
- 4 - Error in command parameter
- 5 - Call home is disabled
- 6 - Another gather-send task for the same configuration file is already running
- 7 - Error in gather-send configuration file
- 8 - Data package created but sender failed
- 9 - Internal error
- 10 - Critical error
- 99 - Unknown

Gather-send task steps:

- step 1 - Initializing
- step 2 - Each call home child nodes gathering logs
- step 3 - Pulling log collection from child nodes
- step 4 - Creating data package
- step 5 - Calling send task
- step 6 - Final status

delete

Deletes the status log of the specified configuration file with the following options:

--task [daily | weekly]

Specifies the log of the daily or weekly task.

--task sendfile

Specifies the log of the tasks initiated by the "run sendfile" command.

[-n num]

Specifies the number of entry per gather-send task to show.

--startTime starttime

Specifies the start time of the log to delete.

--startTimeBefore starttime

All logs older than the time specified by this option will be deleted..

--all

All logs will be deleted.

test

Executes detailed system checks:

mmcallhome

connection

Checks the connectivity between the IBM e-support infrastructure and the call home node of the current group. A proxy is used if the call home proxy settings are enabled. If the proxy setting is disabled, direct connections are attempted.

Note: This command can be executed only from the nodes which are a part of a currently existing call home group.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmcallhome** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

1. To configure a call home group, issue this command:

```
mmcallhome group add group1 themisto0 --node themisto0,themisto1,themisto2
```

The system displays output similar to this:

```
Call home group group1 has been created
```

2. To view the configured call home groups, issue this command:

```
mmcallhome group list
```

The system displays output similar to this:

Call Home Group	Call Home Node	Call Home Child Nodes
group1	test-11	test-11,test-12,test-13

You can also give the same command with the **--long** option to view the configured call home groups with their long names:

```
mmcallhome group list --long
```

The system displays output similar to this:

Call Home Group	Call Home Node	Call Home Child Nodes
group1	test-11.localnet.com	test-11.localnet.com,test-12.localnet.com, test-13.localnet.com

3. To change customer information such as customer name, customer ID, and the country code, issue this command:

```
mmcallhome info change --customer-name "SpectrumScaleTest" --customerid  
"1234" --country-code "JP"
```

The system displays output similar to this:

```
Success
```

4. To view the customer information, issue this command:

```
mmcallhome info list
```

The system displays output similar to this:

Parameter	Value
customer-name	SpectrumScaleTest
customer-id	1234
callhome-method	ethernet
country-code	JP

- To create a call home group automatically, issue this command:

```
mmcallhome group auto
```

The system displays output similar to this:

```
mmcallhome group auto: [I] In progress: Create 1 new call home groups.
mmcallhome group auto: [I] In progress: Nodes without call home: 1
See /var/adm/ras/mmfs.log.latest for details.
group: autoGroup_1 successfully added
```

- To create a call home group automatically and enable the cluster for call home functionality by displaying options for acceptance, issue this command:

```
mmcallhome group auto --enable
```

The system displays output similar to this:

```
By accepting this request, you agree to allow IBM
and its subsidiaries to store and use your contact information and your support
information anywhere they do business worldwide. For more information, please
refer to the Program license agreement and documentation. If you agree, please
respond with "accept" for acceptance, else with "not accepted" to decline.
(accept / not accepted)
accept
mmcallhome group auto: [I] In progress: Create 1 new call home groups.
mmcallhome group auto: [I] In progress: Nodes without call home: 1
See /var/adm/ras/mmfs.log.latest for details.
group: autoGroup_1 successfully added
```

Note: To accept the call home functionality, type **accept** manually. Type **mmcallhome group auto --enable accept** to avoid the explicit acceptance from the user.

- To use create new group after deleting the existing group, issue this command:

```
mmcallhome group auto --force
```

The system displays output similar to this:

```
mmcallhome group auto: [I] In progress: Create 1 new call home groups.
mmcallhome group auto: [I] In progress: Nodes without call home: 1
See /var/adm/ras/mmfs.log.latest for details.
mmcallhome group auto: [I] In progress: Delete existing groups.
Call home group autoGroup_1 has been deleted
group: autoGroup_2 successfully added
```

- To enable the call home service, issue this command:

```
mmcallhome capability enable
```

The system displays output similar to this:

```
Call home node: themisto0
Call home child nodes to collect data: themisto0 themisto1 themisto2 (total 3 nodes)
Excluded nodes:
SSH Access Check: OK
Data package directory: /tmp/mmfs/callhome
Success
```

- To register a daily task with cron, issue this command:

```
mmcallhome schedule add --task daily
```

The system displays output similar to this:

mmcallhome

```
/etc/cron.d/gpfscallhome_gatherSend_daily.conf registered  
41 command entries are defined for this task
```

10. To register a weekly task with cron, issue this command:

```
mmcallhome schedule add --task weekly
```

The system displays output similar to this:

```
/etc/cron.d/gpfscallhome_gatherSend_weekly.conf registered  
14 command entries are defined for this task
```

11. To list the registered tasks for gather-send, issue this command:

```
mmcallhome schedule list
```

The system displays output similar to this:

```
Registered Tasks for gatherSend:  
Conffile          CronParameters  
daily.conf        3 2 * * *  
weekly.conf       54 3 * * sun
```

Note: The CronParameter indicates the date and time settings for the execution of the command. It displays the values for minutes (0-59), hours (0-23), day of month (1-31), month (1-12 or jan-dec) and day of week (0-6, where sun=0 or sun-sat). For example CronParameter 54 3 * * sun indicates that the command executes on every Sunday at 3.54 am.

12. To monitor the call home tasks, issue this command:

```
mmcallhome status list
```

The system displays output similar to this:

Task	Start time	Status	Package file name
daily	20150930132656.582	success	...aultDaily.g_daily.20150930132656582.c10.DC
daily	20150930133134.802	success	...aultDaily.g_daily.20150930133134802.c10.DC
daily	20150930133537.509	success	...aultDaily.g_daily.20150930133537509.c10.DC
daily	20150930133923.063	success	...aultDaily.g_daily.20150930133923063.c10.DC
RunSendFile	20150930133422.843	success	...group2.MyTestData.s_file.20150930133422843.c10.DC

13. To set the parameters for the proxy server, issue this command:

```
| mmcallhome proxy change --proxy-location okapi --proxy-port 80  
| --proxy-username root --proxy-password <password>
```

The system displays output similar to this:

```
| Call home proxy-port has been set to 80  
| Call home proxy-username has been set to root  
| Call home proxy-password has been set to <password>  
| Call home proxy-location has been set to okapi
```

14. To view the proxy server parameters, issue this command:

```
mmcallhome proxy list
```

The system displays output similar to this:

Status
proxy-enabled NO
proxy-auth-enabled false
Parameter
proxy-location okapi
proxy-port 80
proxy-username root
proxy-password xxxxx

15. To invoke a one-time gather-send task, issue this command:

```
mmcallhome run GatherSend --task daily
```

The system displays output similar to this:

```
Starting one time run using daily.conf
```

- To run one-time send command to upload a file, issue this command:

```
mmcallhome run SendFile --file /ibm/gpfs0/testDir/testFile --desc MyTestData
```

The system displays output similar to this:

```
Running.sendFile... (In case of network errors, it may take over 20 minutes for
retries.)
StartTime=20150930193046.693
Successfully uploaded the given file
Run mmcallhome status ls -v to see the package name
```

- To view the status of the currently running and the already completed call home tasks, issue this command:

```
mmcallhome status list --verbose
```

The system displays output similar to this:

Task	Start time	Updated time	Status	RC or Step
	Package file name			
	[additional info: value]			
	31790849425327.4_2_1_0.x.abc.autoGroup_1.gat_weekly.g_weekly.20160412160447854.c10.DC			
	31790849425327.4_2_1_0.x.abc.autoGroup_1.gat_weekly.g_weekly.20160412173941161.c10.DC			
weekly	20160412174030.803	20160412174034	failed	RC=6 (lock err)
	NA			
	31790849425327.4_2_1_0.x.abc.autoGroup_1.gat_weekly.g_weekly.20160412175159390.c10.DC			

Note: Sometimes the output of `mmcallhome status list --verbose` displays single line without detailed information about RC indicating successful completion of call home tasks. The failed status indicates that there was an issue with the call home task and the RC numeral indicates the respective issue. If the value of RC is zero, it indicates that the upload procedure is successful, but some automatically resolvable issue occurred while uploading the data. The value, `RC != 0`, indicates that the upload procedure is not successful. The detailed information about the upload procedure is available in the logs.

- To test the connection, issue this command:

```
mmcallhome test connection
```

The system displays output similar to this:

```
Starting connectivity test between the call home node and IBM
Call home node: themisto0
Starting time: Wed Sep 30 14:37:51 JST 2015
Testing connection via proxy server (no authentication required)
User: NA Pass: NA Host: okapi Port: 80
Testing prefix Edge_SP_Config:
Edge_SP_Config_1: 129.42.56.189 OK
Testing prefix Edge_Profile:
Edge_Profile_1: 129.42.56.189 OK
Testing prefix Edge_Status_Report:
Edge_Status_Report_1: 129.42.56.189 OK
```

- To check the version and the subversion of the mmcallhome program, issue this command:

```
mmcallhome version
```

The system displays output similar to this:

```
Version: 4.2.0-005
```

mmcallhome

See also

- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfss command” on page 454
- “mmobj command” on page 466
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

Location

/usr/lpp/mmfs/bin

mmces command

Manage CES (Cluster Export Services) configuration.

Synopsis

```
mmces address add [--ces-node Node] [--attribute Attribute]
                  [--ces-group Group] {--ces-ip {IP[,IP...]} }
```

or

```
mmces address remove --ces-ip {IP[,IP...]} }
```

or

```
mmces address move --ces-ip {IP[,IP...]} --ces-node Node
```

or

```
mmces address move --rebalance
```

or

```
mmces address change
      {--ces-ip IP | --remove-attribute}
      --attribute Attribute[,Attribute...]
```

```
mmces address change
      [--ces-ip IP[,IP...]]
      [--attribute Attribute[,Attribute...]]
      [--ces-group Group]
      [--remove-attribute]
      [--remove-group]
```

or

```
mmces address list [-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
mmces address list [--ces-ip IP[,IP...]] [--ces-group Group[,Group...]]
                  [--attribute Attribute[,Attribute...]] [--by-node]
```

or

```
mmces address policy [none | balanced-load | node-affinity | even-coverage]
```

or

```
mmces node {suspend | resume} [-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
```

or

```
mmces node list [--ces-group Group[,Group...]] [--verbose]
```

or

```
mmces service {enable | disable} {NFS | SMB | OBJ | BLOCK}
```

or

```
mmces service {start | stop} {NFS | SMB | OBJ | BLOCK}
      [-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
```

or

```
mmces service list [-N {Node[,Node...]} | NodeFile | NodeClass} | -a] [--verbose]
```

or

```
mmces log level [new-level]
```

mmces

or

```
mmces events active [NFS | OBJ | SMB | BLOCK | AUTH | AUTH_OBJ | NETWORK]
[-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
```

or

```
mmces events list [NFS | OBJ | SMB | BLOCK | AUTH | AUTH_OBJ | NETWORK]
[--time {hour | day | week | month}]
[--severity {INFO | WARNING | ERROR | SEVERE}]
[-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
```

or

```
mmces state show [NFS | OBJ | SMB | BLOCK | AUTH | AUTH_OBJ | NETWORK | CES]
[-N {Node[,Node...]} | NodeFile | NodeClass} | -a]
```

or

```
mmces state cluster [NFS | OBJ | SMB | BLOCK | AUTH | AUTH_OBJ | NETWORK | CES]
```

Availability

Available with IBM Spectrum Scale Standard Edition or higher.

Description

Use the **mmces** command to manage protocol addresses, services, node state, logging level and load balancing. The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms. Before using the **mmces** command, CES has to be installed and configured in the system. For more information on how to install CES, see the *Installing Cluster Export Services as part of installing IBM Spectrum Scale on Red Hat Enterprise Linux 7.x systems* section in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

CES currently supports the NFS, SMB, BLOCK, and Object services. Each service can be enabled or disabled with the **mmces service** command. Enabling a service is a CES cluster-wide operation. In addition, enabled services can be started and stopped on individual nodes.

Clients access the CES services using one or more IP addresses in the CES address pool. Addresses can be added to and removed from the pool using the **mmces address add** and **mmces address remove** commands. Existing addresses can be reassigned to another node with the **mmces address move** command.

Addresses can have one or more attributes associated with them. An address attribute is a tag that the services can identify a specific address as having a special meaning, which is defined by the service protocol. Addresses can have multiple attributes, but an attribute can only be associated with a single address. The supported attributes are **object_singleton_node** and **object_database_node**. These attributes are used to manage object protocol-related services. For more information, see the *Object services and object protocol nodes* table in the *Understanding and managing Object services* section in the *IBM Spectrum Scale: Administration Guide*.

Addresses can have a policy associated with them, and that policy determines how addresses are automatically distributed. The allowed policies are **none**, **balanced-load**, **node-affinity**, and **even-coverage**. A policy of **none** means addresses are not distributed automatically.

A CES node can be placed in a suspended state. When a node is in suspended state, all of the CES addresses for that node are reassigned to other nodes, and the node will not accept new address assignments. Any services that are started when the node is suspended continue to run. The suspended state is persistent, which means nodes remain suspended following a reboot. Use the **mmces node** command to suspend and resume a node.

Parameters

address

Manages CES addresses with one of the following actions:

add

Adds the addresses specified by the **--ces-ip** parameter to the CES address pool and assigns them to a node. The node to which an address is assigned will configure its network interfaces to accept network communication destined for the address. CES addresses must be different from IP addresses used for GPFS or CNFS communication.

If **--ces-node** is specified with **add**, all addresses specified with the **--ces-ip** parameter will be assigned to this node. If **--ces-node** is not specified, the addresses will be distributed among the CES nodes.

If an attribute is specified with **--attribute**, there can only be one address specified with the **--ces-ip** parameter.

If **--ces-group** is specified with **add**, all new addresses will be associated with the specified group. The result can be viewed with the "mmces address list" command.

Note: The provided addresses or hostnames must be resolvable by forward and reverse name resolution (DNS or /etc/hosts on all CES nodes). Otherwise you get the following error message:
Cannot resolve <ip address>; Name or service not known

You can also perform a manual check by running the following command: **mmcni host <ip address>**.

Ensure that the netmask (PREFIX) setting in the ifcfg-<interface> files is correct.

remove

Removes the addresses specified by the **--ces-ip** parameter from the CES address pool. The node to which the address is assigned reconfigures its network interfaces to no longer accept communication for that address.

move

Moves addresses.

If the **--ces-ip** parameter is specified, the addresses specified by *IP* are moved from one CES node to another. The addresses are reassigned to the node specified by the **--ces-node** parameter.

If the **--rebalance** parameter is specified, the addresses are distributed within 60 seconds based on the currently configured distribution policy. If the policy is currently undefined or **none**, the **even-coverage** policy is applied.

Use this command with caution because IP movement will trigger CES protocol recovery.

change

Changes or removes address attributes.

If the **--ces-ip** parameter is specified:

- The command associates the attributes that are specified by the **--attribute** parameter with the address that is specified by the **--ces-ip** parameter. If an attribute is already associated with another address, that association is ended.
- If the **--remove-attribute** parameter is specified, the command removes the attributes that are specified by the **--attribute** parameter from the addresses that are specified by the **--ces-ip** parameter.
- The command associates the groups that are specified by the **--ces-group** parameter with the address that is specified by the **--ces-ip** parameter

mmces

- If the **--remove-group** parameter is specified, the command removes the groups that are specified by the **--ces-group** parameter from the addresses that are specified by the **--ces-ip** parameter.

If the **--ces-ip** parameter is not specified:

- If the **--remove-attribute** parameter is specified, the command removes the attributes that are specified by the **--attribute** parameter from their current associations.
- If the **--remove-group** parameter is specified, the command removes the groups that are specified by the **--group** parameter from their current associations.
- Specifying **--remove-group** with the groups specified by the **--group** parameter removes the groups from their current associations.

list

Lists the CES addresses along with group, attribute and node assignments.

Options:

- ces-ip** List only the addresses provided.
- ces-group** List only addresses whose group assignment matches one of the groups provided.
- attribute** List only addresses whose attributes match one of the attributes provided.
- by-node** List addresses by node, using the output format from IBM Spectrum Scale V4.1.1 and later.

| **Note:** The [-N {Node[,Node...] | NodeFile | NodeClass} | -a] option is deprecated for IBM
| Spectrum Scale version 4.2.3 and might be removed in a later release.

policy

Sets the CES address distribution policy.

node

Manages CES node state with one of the following actions:

suspend

Suspends the specified nodes. If neither the **-N** or **-a** parameters are specified, only the local node is suspended.

When a node is suspended, all addresses assigned to the node are reassigned to other nodes and the node will not accept any subsequent address assignments. Suspending a node will trigger CES protocol recovery if the node has CES addresses assigned. This does not stop any services, and some services like CTDB are still a part of the cluster infrastructure.

resume

Resumes the specified suspended nodes. If neither the **-N** or **-a** parameters are specified, only the local node is resumed.

When a suspended node is resumed (no longer suspended), the node accepts subsequent address assignments.

list

Lists the specified nodes along with their current node state. If the **-N** parameter is not specified, all nodes are listed.

verbose

Lists the addresses assigned to the nodes.

--ces-group

Lists the nodes belonging to the specified groups.

service

Manages protocol services with one of the following actions:

enable

Enables and starts the specified service on all CES nodes.

disable

Disables and stops the specified service on all CES nodes.

Note: Disabling a service will discard any configuration data from the CES cluster and needs to be used with caution. If applicable, backup any relevant configuration data. Subsequent service enablement will start with a clean configuration.

start

Starts the specified service on the nodes specified. If neither the **-N** or **-a** parameters are specified, the service is started on the local node.

stop

Stops the specified service on the nodes specified. If neither the **-N** or **-a** parameters are specified, the service is stopped on the local node.

Note: If a service is stopped on a node that has CES addresses assigned, clients will not be able to access the service using any of the addresses assigned to that node. Access to the data from clients is not possible any more for services that are stopped. This state is not persistent, so after a reboot all the services become active again.

list

Lists the state of the enabled services.

log level

Sets and checks the CES log level. The CES log level determines how much information related to the CES nodes is entered into the GPFS log file. Values can be from **0** (less logging) to **3** (more logging).

events

Shows one of the following CES events that occurred on a node or nodes:

active

Lists all events that are currently contributing to making the state of a component unhealthy. If no component is specified, active events for all components are listed. If neither the **-N** or **-a** parameters are specified, the active events for the local node are listed. If there are multiple events shown by the command they will be listed in the order we recommend they be fixed, with the most important event to fix at the top.

list

Lists the events that occurred on a node or nodes, whether or not they are currently contributing to the state of a component. If no component is specified, events for all components are listed. If **--time** is specified, only events from the previous chosen interval are listed, otherwise all events are listed. If **--severity** is specified, only events of the chosen severity are listed, otherwise all events are listed. If neither the **-N** or **-a** parameters are specified, the events for the local node are listed.

Events older than 180 days are removed from the list. A maximum of 10,000 events are saved in the list.

state

Shows the state of one or more nodes in the cluster.

show

Shows the state of the specified service on the nodes specified. If no service is specified, all services will be displayed. If neither the **-N** or **-a** parameters are specified, the state of the local node is shown.

cluster

Shows the combined state for the services across the whole CES cluster. If no service is specified an aggregated state will be displayed for each service, where healthy means the service is healthy

mmces

on all nodes, degraded means the service is not healthy on one or all nodes, and failed means that the service is not available on any node. If a service is specified the state of that service will be listed for each node, along with the name of any event that is contributing to an unhealthy state.

--ces-node

Indicates that the command applies only to the specified CES node name.

--attribute

Specifies either a single attribute or a comma-separated list of attributes as indicated in the command syntax.

--ces-ip

Specifies either a single or comma-separated list of DNS qualified host names or IP addresses as indicated in the command syntax.

--rebalance

Distributes addresses immediately based on the currently configured distribution policy. If the policy is currently undefined or **none**, the **even-coverage** policy is applied.

none

Specifies that addresses are not distributed automatically.

balanced-load

Distributes addresses dynamically in order to approach an optimized load distribution throughout the cluster. The network and CPU load on all the nodes is monitored and addresses are moved based on given policies.

Addresses that were recently moved or addresses with attributes are not moved.

node-affinity

Attempts to keep addresses associated with the node to which they were assigned. Address node associations are created with the **--ces-node** parameter of the **mmces address add** command or the **mmces address move** command. Automatic movements of addresses do not change the association. Addresses that were enabled without a node specification do not have a node association. Addresses that are associated with a node but assigned to a different node are moved back to the associated node.

Addresses that were recently moved or addresses with attributes are not moved.

even-coverage

Attempts to evenly distribute all of the addresses among the available nodes.

Addresses that were recently moved or addresses with attributes are not moved.

--remove-attribute

Indicates that the specified attributes should be removed.

-N {Node[,Node...] | NodeFile | NodeClass}

Indicates that the command applies only to the specified node names.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

-a Specifies that the command applies to all CES nodes.

NFS

Specifies that the command applies to the NFS service.

OBJ

Specifies that the command applies to the Object service.

SMB

Specifies that the command applies to the SMB service.

BLOCK

Specifies that the command applies to the BLOCK service.

AUTH

Specifies that the command applies to the AUTH service.

NETWORK

Specifies that the command applies to the NETWORK service.

CES

Specifies that the command applies to the CES service.

--verbose

Specifies that the output is verbose.

new-level

Sets the log level to a new value. If the *new-level* parameter is not specified, the current log level is displayed.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

Note: Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

--time

Lists the previous events from one of the following intervals:

hour

Lists the events from the past hour.

day

Lists the events from the past day.

week

Lists the events from the past week.

month

Lists the events from the past month.

The events are listed whether or not they are currently contributing to the state of a component.

--severity

Specifies that only events for one of the following severities are listed:

INFO

Lists only informational events.

WARNING

Lists only warning events.

ERROR

Lists only error events.

SEVERE

Lists only severe events.

Exit status

- 0 Successful completion.

nonzero

A failure has occurred.

mmces

Security

You must have root authority to run the **mmces** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

Examples

1. To add an address to a specified node, issue this command:

```
mmces address add --ces-node node1 --ces-ip 10.1.2.3
```

When this command is successful, the system does not display output.

2. To add several addresses to a specified node, issue this command:

```
mmces address add --ces-node node1 --ces-ip 10.1.2.3,10.1.2.4
```

When this command is successful, the system does not display output.

3. To add an address with the attribute **object_singleton_node** to a specified node, issue this command:

```
mmces address add --ces-node node1 --ces-ip 10.1.2.3 --attribute object_singleton_node
```

When this command is successful, the system does not display output.

4. To add addresses which are distributed among the CES nodes, issue this command:

```
mmces address add --ces-ip 10.1.2.3,10.1.2.4,10.1.2.5,10.1.2.6
```

When this command is successful, the system does not display output.

5. To remove several addresses, issue this command:

```
mmces address remove --ces-ip 10.1.2.3,10.1.2.4
```

When this command is successful, the system does not display output.

6. To associate the attribute **object_singleton_node** to the address **10.1.2.3**, issue this command:

```
mmces address change --ces-ip 10.1.2.3 --attribute object_singleton_node
```

When this command is successful, the system does not display output.

7. To remove the attribute **object_singleton_node**, issue this command:

```
mmces address change --remove-attribute --attribute object_singleton_node
```

When this command is successful, the system does not display output.

8. To move an address to another node, issue this command:

```
mmces address move --ces-ip 10.0.100.231 --ces-node node2
```

When this command is successful, the system does not display output.

9. To suspend a group of nodes, issue this command:

```
mmces node suspend -N node1,node2,node3
```

The system displays output similar to this:

```
Node now in suspended state.
```

10. To resume the current node, issue this command:

```
mmces node resume
```

The system displays output similar to this:

Node no longer in suspended state.

- To enable the Object service in the CES cluster, issue this command:

```
mmces service enable obj
```

When this command is successful, the system does not display output.

- To disable the NFS service in the CES cluster, issue this command:

```
mmces service disable nfs
```

When this command is successful, the system does not display output.

- To stop the SMB service on a few nodes, issue this command:

```
mmces service stop smb -N node1,node2,node3
```

When this command is successful, the system does not display output.

- To start the SMB service on all CES nodes, issue this command:

```
mmces service start smb -a
```

When this command is successful, the system does not display output.

- To show which services are enabled and which are running all CES nodes, issue this command:

```
mmces service list -a
```

The system displays output similar to this:

```
Enabled services: NFS OBJ
node1: NFS is running, OBJ is running
node2: NFS is running, OBJ is running
node3: NFS is running, OBJ is running
```

- To display the current CES log level, issue this command:

```
mmces log level
```

The system displays output similar to this:

```
CES log level is currently set to 1
```

- To set the CES log level to 2, issue this command:

```
mmces log level 2
```

The system displays output similar to this:

```
mmchconfig: Command successfully completed
mmchconfig: Propagating the cluster configuration data to all affected nodes.
This is an asynchronous process.
```

- To display the state of all CES components on the local node, issue this command:

```
mmces state show
```

The system displays output similar to this:

NODE	AUTH	BLOCK	NETWORK	AUTH_OBJ	NFS	OBJ	SMB	CES
node1	DISABLED	DISABLED	HEALTHY	DISABLED	DISABLED	DISABLED	HEALTHY	HEALTHY

- To display the state of the NFS component on all nodes, issue this command:

```
mmces state cluster NFS
```

The system displays output similar to this:

NODE	COMPONENT	STATE	EVENTS
node1	NFS	HEALTHY	
node2	NFS	FAILED	nfsd_down
node3	NFS	HEALTHY	

- To display a list of active events of all CES components on the local node, issue this command:

```
mmces events active
```

mmces

The system displays output similar to this:

Node	Component	Event Name	Severity	Details
node1	NFS	nfsd_down	ERROR	NFSD process not running

21. To display a list of all NFS events from the last hour on the local node, issue this command:

```
mmces events list
```

The system displays output similar to this:

Node	Timestamp	Event Name	Severity	Details
node1	2015-05-13 10:57:52.124369+00:00UTC	nfsd_down	ERROR	NFSD process not running
node1	2015-05-13 10:58:06.809071+00:00UTC	cesnodestatechange_info	INFO	A CES node state changed
node1	2015-05-13 10:58:07.137343+00:00U	ganeshagrace_info	INFO	Ganesha NFS is set to grace

22. To enable the block service, issue this command:

```
mmces service enable BLOCK
```

The system displays the following prompt:

```
Block device support in Spectrum Scale is intended for use only in diskless node  
remote boot (non-performance-critical), and is not suited for high-bandwidth  
block device access needs. Confirm that this matches your use case before enabling  
the block service. If you have any questions contact scale@us.ibm.com  
Do you want to continue to enable BLOCK service? (yes/no)
```

After you press Y, the system displays the following output:

```
c40bbc1xn12.gpfs.net: Loading and configuring SCST  
mmchconfig: Command successfully completed  
mmchconfig: Propagating the cluster configuration data to all  
affected nodes. This is an asynchronous process.
```

See also

- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmobj command” on page 466
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

Location

/usr/lpp/mmfs/bin

mmcesdr command

Manages protocol cluster disaster recovery.

Synopsis

```
mmcesdr primary config --output-file-path FilePath --ip-list IPAddress[,IPAddress,...]
  [--allowed-nfs-clients {--all | --gateway-nodes | IPAddress[,IPAddress,...]}]
  [--rpo RPOValue] [--inband] [-v]

or

mmcesdr primary backup [-v]

or

mmcesdr primary restore [--new-primary] [--input-file-path FilePath]
  [--file-config {--recreate | --restore}] [-v]

or

mmcesdr primary update {--obj | --nfs | --smb | --ces} [-v]

or

mmcesdr primary fallback --prep-outband-transfer --input-file-path FilePath [-v]

or

mmcesdr primary fallback --convert-new --output-file-path FilePath --input-file-path FilePath [-v]

or

mmcesdr primary fallback {--start | --apply-updates | --stop [--force]} [--input-file-path FilePath] [-v]

or

mmcesdr secondary config --input-file-path FilePath [--prep-outband-transfer] [--inband] [-v]

or

mmcesdr secondary failover [--input-file-path FilePath] [--file-config {--recreate | --restore}]
  [--data {--restore | --norestore}] [-v]

or

mmcesdr secondary fallback --generate-recovery-snapshots --output-file-path FilePath
  [--input-file-path FilePath] [-v]

or

mmcesdr secondary fallback --post-failback-complete [--input-file-path FilePath]
  [--file-config {--recreate | --restore}] [-v]

or

mmcesdr secondary fallback --post-failback-complete --new-primary --input-file-path FilePath
  [--file-config {--recreate | --restore}] [-v]
```

Availability

Available with IBM Spectrum Scale Advanced Edition or IBM Spectrum Scale Data Management Edition.

Description

Use the **mmcesdr** command to manage protocol cluster disaster recovery.

mmcesdr

You can use the **mmcesdr primary config** command to perform initial configuration for protocols disaster recovery on the primary cluster and to generate a configuration file that is used on the secondary cluster. The protocol configuration data can be backed up using the **mmcesdr primary backup** command and the backed-up data can be restored using the **mmcesdr primary restore** command. The backed-up configuration information for the primary cluster can be updated by using the **mmcesdr primary update** command. You can use the **mmcesdr primary fallback** command to fail back the client operations to the primary cluster.

You can use the **mmcesdr secondary config** command to perform initial configuration for protocols disaster recovery on the secondary cluster by using the configuration file generated from the primary cluster. The secondary read-only filesets can be converted into read-write primary filesets using the **mmcesdr secondary failover** command. You can use the **mmcesdr secondary fallback** command to either generate a snapshot for each acting primary fileset or complete the fallback process, and convert the acting primary filesets on the secondary cluster back into secondary filesets.

For information on detailed steps for protocols disaster recovery, see *Protocols cluster disaster recovery* in *IBM Spectrum Scale: Administration Guide*.

The **mmcesdr** log file is at `/var/adm/ras/mmcesdr.log`. This log file is included with the CES information generated by the **gpfs.snap** command. The **gpfs.snap** command generates the CES information by default, if a protocol is enabled.

Parameters

primary

This command is run on the primary cluster.

config

Perform initial configuration of protocol cluster disaster recovery.

--output-file-path *FilePath*

Specifies the path to store output of the generated configuration file, which is always named `DR_Config`.

--ip-list *IPAddress[, IPAddress, ...]*

Comma-separated list of public IP addresses on the secondary cluster to be used for active file management (AFM) DR-related NFS exports.

--allowed-nfs-clients {--all | --gateway-nodes | *IPAddress[, IPAddress, ...]*}

Optional. Specifies the entities that can connect to the AFM DR-related NFS shares, where:

--all

Specifies that all clients must be allowed to connect to the AFM DR-related NFS shares. If omitted, the default value of `--all` is used.

--gateway-nodes

Specifies the gateway nodes currently defined on the primary that must be allowed to connect to the AFM DR-related NFS shares.

IPAddress[, IPAddress, ...]

Specifies the comma-separated list of IP addresses that must be allowed to connect to the AFM DR-related NFS shares.

--rpo *RPOValue*

Optional. Specifies the integer value of the recovery point objective (RPO) to be used for the AFM DR filesets. By default, this parameter is disabled. The valid range is: $720 \leq RPO \leq 2147483647$. The minimum value that can be set for the RPO is 720.

Note: Setting the value of RPO to less than 720 generates an error.

--inband

Optional. Specifies to use the inband (across the WAN) method of initial data transfer from primary to secondary cluster. If omitted, the default value of outband is used.

backup

Backs up all protocol configuration and CES configuration into a dedicated, independent fileset with each protocol in its own subdirectory.

restore

Restores object, NFS, and SMB protocol configuration and CES configuration from the configuration data backed up.

--new-primary

Optional. Performs restore operation to a newly failed back primary cluster.

--input-file-path *FilePath*

Optional. Specifies the original configuration file that was used to set up the secondary cluster. If not specified, the file that is saved in the configuration independent fileset is used as default.

--file-config {--recreate | --restore}

Optional. Specifies whether SMB and NFS exports are re-created, or if the entire protocol configuration is restored. If not specified, the SMB and NFS exports are re-created by default.

update

Updates the backed-up copy of the protocol configuration or CES configuration.

--obj

Specifies the backed up-copy of the object protocol configuration to be updated with the current object configuration.

--nfs

Specifies the backed-up copy of the IBM NFSv4 stack protocol configuration to be updated with the current IBM NFSv4 stack configuration.

--smb

Specifies the backed-up copy of the SMB protocol configuration to be updated with the current SMB configuration.

--ces

Specifies the backed-up copy of the CES configuration to be updated with the current CES configuration.

failback

Used for several options to failback client operations to a primary cluster.

Failback involves transfer of data from the acting primary (secondary) cluster to the old primary cluster as well as restoring protocol and possibly CES configuration information and transformation of protected filesets to primary filesets.

--prep-outband-transfer

Creates independent filesets that out of band data is transferred to.

--input-file-path *FilePath*

Specifies the configuration file that is the output from the **mmcesdr secondary failback --generate-recovery-snapshots** command.

failback

Used for several options to failback client operations to a primary cluster.

mmcesdr

Fallback involves transfer of data from the acting primary (secondary) cluster to the old primary cluster as well as restoring protocol and possibly CES configuration information and transformation of protected filesets to primary filesets.

--convert-new

Specifies that the fallback is not going to the old primary but instead a new primary. This step specifically converts the newly created independent filesets to primary AFM DR filesets.

--output-file-path *FilePath*

Specifies the path to store output of generated configuration file, **DR_Config**, with the new AFM primary IDs.

--input-file-path *FilePath*

Specifies the configuration file that is the output from the **mmcesdr secondary fallback --generate-recovery-snapshots** command.

fallback

Used for several options to fallback client operations to a primary cluster.

Fallback involves transfer of data from the acting primary (secondary) cluster to the old primary cluster as well as restoring protocol and possibly CES configuration information and transformation of protected filesets to primary filesets.

--start

Begins the fallback process and restores the data to the last RPO snapshot.

--apply-updates

Transfers data that was written to the secondary cluster while failover was in-place.

Note: The use of this option might need to be done more than once depending on the system load.

--stop [--force]

Completes the transfer of data process and puts the filesets in the read-write mode.
Optionally, if this fails you can use the **--force** option.

Note: In addition to using these options, after stopping the data transfer, you need to use the **mmcesdr primary restore** command to restore the protocol and the CES configuration.

--input-file-path *FilePath*

Optional. Specifies the original configuration file that was used to set up the secondary cluster. If not provided, the default is to use the one saved in the configuration independent fileset.

secondary

This command is run on the secondary cluster.

config

Perform initial configuration of protocol cluster disaster recovery.

--prep-outband-transfer

Creates independent filesets that out of band data is transferred to as part of the initial configuration. If out of band data transfer is used for DR configuration, this option must be used before data is transferred from the primary to the secondary using out of band transfer. If out of band transfer is used, this command is run once with this option and then again after the data is transferred without the option.

--input-file-path *FilePath*

Specifies the path of the configuration file generated from the configuration step of the primary cluster.

--inband

Optional. Specifies to use the inband (across the WAN) method of initial data transfer from primary to secondary cluster. If omitted, the default value of outband is used.

Note: If **--inband** is used for the primary configuration, it must also be used for the secondary configuration.

failover

Converts secondary filesets from read-only to read-write primary filesets and converts the secondary protocol configurations to those of the failed primary.

--input-file-path *FilePath*

Optional. Specifies the original configuration file that was used to set up the secondary cluster. If not specified, the file that is saved in the configuration independent fileset is used as default.

--file-config {--recreate | --restore}

Optional. Specifies whether SMB and NFS exports are re-created, or if the entire protocol configuration is restored. If not specified, the SMB and NFS exports are re-created by default.

--data {--restore | --norestore}

Optional. Specifies that data must be restored from the latest RPO snapshot or that restoring from the latest RPO snapshot is not required. Default is that restoring from the latest RPO snapshot is not required.

failback

Runs one of the two failback options: either generates a snapshot for each acting primary fileset or completes the failback process and convert the acting primary filesets on the secondary cluster back into secondary filesets

--generate-recovery-snapshots

Generates the psnap0 snapshot for each acting primary fileset and stores in the default snapshot location for use in creation of a new primary cluster with new primary filesets to fail back to. The files within the snapshot need to be manually transported to the new primary.

--output-file-path *FilePath*

Specifies the path to store output of generated snapshot recovery configuration file.

--input-file-path *FilePath*

Optional. Specifies the path of the original configuration file that was used to set up the secondary cluster. If not provided, the default is to use the one saved in the configuration independent fileset.

failback

Runs one of the two failback options: either generates a snapshot for each acting primary fileset or completes the failback process and convert the acting primary filesets on the secondary cluster back into secondary filesets

--post-failback-complete

Completes the failback process by converting the acting primary filesets back into secondary, read-only filesets and ensures that the proper NFS exports for AFM DR exist.

--new-primary

Performs the failback operation to a newly, failed back primary cluster.

--input-file-path *FilePath*

Specifies the path of the updated configuration file that is created from the **mmcesdr primary failback --convert-new** command, which includes updated AFM primary IDs.

--file-config {--recreate | --restore}

mmcesdr

Optional. Specifies whether SMB and NFS exports are re-created, or if the entire protocol configuration is restored. If not specified, the SMB and NFS exports are re-created by default.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmcesdr** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Examples

1. Issue the following command on the primary cluster to configure independent fileset exports as AFM DR filesets and backup configuration information:

```
mmcesdr primary config --output-file-path /root/ --ip-list "9.11.102.211,9.11.102.210" --rpo 720 --inband
```

The system displays output similar to this:

```
Performing step 1/5, configuration fileset creation/verification.  
Successfully completed step 1/5, configuration fileset creation/verification.  
Performing step 2/5, protocol and export services configuration backup.  
Successfully completed step 2/5, protocol and export services configuration backup.  
Performing step 3/5, determination of protocol exports to protect with AFM DR.  
WARNING: Export /gpfs/fs0/nfs-ganesha-dep of type nfs will NOT be protected through AFM DR because it is a dependent fileset.  
Not all exports of type NFS-ganesha will be protected through AFM DR, rc: 2  
WARNING: Export /gpfs/fs0/smb-dep of type smb will NOT be protected through AFM DR because it is a dependent fileset.  
Not all exports of type SMB will be protected through AFM DR, rc: 2  
Completed with errors step 3/5, determination of protocol exports to protect with AFM DR.  
Performing step 4/5, conversion of protected filesets into AFM DR primary filesets.  
Successfully completed step 4/5, conversion of protected filesets into AFM DR primary filesets.  
Performing step 5/5, creation of output DR configuration file.  
Successfully completed step 5/5, creation of output DR configuration file.
```

File to be used with secondary cluster in next step of cluster DR setup: /root//DR_Config

2. Issue the following command on the secondary cluster to create the independent filesets that are a part of the pair of AFM DR filesets associated with those on the primary cluster:

```
mmcesdr secondary config --input-file-path /root/ --inband
```

In addition to fileset creation, this command also creates the necessary NFS exports and converts the independent filesets to AFM DR secondary filesets. The system displays output similar to this:

```
Performing step 1/3, creation of independent filesets to be used for AFM DR.  
Successfully completed step 1/3, creation of independent filesets to be used for AFM DR.  
Performing step 2/3, creation of NFS exports to be used for AFM DR.  
Successfully completed step 2/3, creation of NFS exports to be used for AFM DR.  
Performing step 3/3, conversion of independent filesets to AFM DR secondary filesets.  
Successfully completed step 3/3, conversion of independent filesets to AFM DR secondary filesets.
```

3. Issue the following command on the primary cluster to configure independent fileset exports as AFM DR filesets, back up configuration information, and facilitate outband data transfer.

Note: The outband data transfer is the default method of data transfer from the primary cluster to the secondary cluster when AFM DR fileset relationships are first set up.

:

```
mmcesdr primary config --output-file-path /root/ --ip-list "9.11.102.211,9.11.102.210" --rpo 720
```

The system displays output similar to this:

```
Performing step 1/5, configuration fileset creation/verification.  
Successfully completed step 1/5, configuration fileset creation/verification.  
Performing step 2/5, protocol and export services configuration backup.
```

Successfully completed step 2/5, protocol and export services configuration backup.
 Performing step 3/5, determination of protocol exports to protect with AFM DR.
 Successfully completed step 3/5, determination of protocol exports to protect with AFM DR.
 Performing step 4/5, conversion of protected filesets into AFM DR primary filesets.
 Successfully completed step 4/5, conversion of protected filesets into AFM DR primary filesets.
 Performing step 5/5, creation of output DR configuration file.
 Successfully completed step 5/5, creation of output DR configuration file.

File to be used with secondary cluster in next step of cluster DR setup: /root//DR_Config

4. Issue the following command on the secondary cluster to create the independent filesets that will later be paired with those on the primary cluster to form AFM DR pairs as part of failing back to a new primary cluster:

```
mmcesdr secondary config --input-file-path /root --prep-outband-transfer
```

The system displays output similar to this:

```
Creating independent filesets to be used as recipients of AFM DR outband transfer of data.  

Transfer all data on primary cluster for fileset fs0:combo1 to fileset fs0:combo1 on secondary cluster.  

Transfer all data on primary cluster for fileset fs0:combo2 to fileset fs0:combo2 on secondary cluster.  

Transfer all data on primary cluster for fileset fs0:nfs-ganeshal to fileset fs0:nfs-ganeshal on secondary cluster.  

Transfer all data on primary cluster for fileset fs0:nfs-ganeshal2 to fileset fs0:nfs-ganeshal2 on secondary cluster.  

Transfer all data on primary cluster for fileset fs0:smb1 to fileset fs0:smb1 on secondary cluster.  

Transfer all data on primary cluster for fileset fs0:smb2 to fileset fs0:smb2 on secondary cluster.  

Transfer all data on primary cluster for fileset fs1:async_dr to fileset fs1:async_dr on secondary cluster.  

Transfer all data on primary cluster for fileset fs1:obj_softpolicy1 to fileset fs1:obj_softpolicy1 on secondary cluster.  

mmcesdr: CES Object protocol is not enabled but there is an object related export present.  

Skipping clearing out the object related files and directories from export.  

Transfer all data on primary cluster for fileset fs1:obj_softpolicy2 to fileset fs1:obj_softpolicy2 on secondary cluster.  

mmcesdr: CES Object protocol is not enabled but there is an object related export present.  

Skipping clearing out the object related files and directories from export.  

Transfer all data on primary cluster for fileset fs1:object_fileset to fileset fs1:object_fileset on secondary cluster.  

mmcesdr: CES Object protocol is not enabled but there is an object related export present.  

Skipping clearing out the object related files and directories from export.  

Successfully completed creating independent filesets to be used as recipients of AFM DR outband transfer of data.  

Transfer data from primary cluster through outbound trucking to the newly created independent filesets before proceeding to the next step.
```

5. After all the data has been transferred to the secondary, issue the following command to complete the setup on the secondary:

```
mmcesdr secondary config --input-file-path /root
```

The system displays output similar to this:

```
Performing step 1/3, verification of independent filesets to be used for AFM DR.  

Successfully completed step 1/3, creation of independent filesets to be used for AFM DR.  

Successfully completed 1/3, verification of independent filesets to be used for AFM DR.  

Performing step 2/3, creation of NFS exports to be used for AFM DR.  

Successfully completed step 2/3, creation of NFS exports to be used for AFM DR.  

Performing step 3/3, conversion of independent filesets to AFM DR secondary filesets.  

Successfully completed step 3/3, conversion of independent filesets to AFM DR secondary filesets.
```

6. Issue the following command on the secondary cluster after the primary cluster has failed:

```
mmcesdr secondary failover
```

The system displays output similar to this:

```
Performing step 1/4, saving current NFS configuration to restore after failback.  

Successfully completed step 1/4, saving current NFS configuration to restore after failback.  

Performing step 2/4, failover of secondary filesets to primary filesets.  

Successfully completed step 2/4, failover of secondary filesets to primary filesets.  

Performing step 3/4, protocol configuration(exports) restore.  

Successfully completed step 3/4, protocol configuration(exports) restore.  

Performing step 4/4, create/verify NFS AFM DR transport exports.  

Successfully completed step 4/4, create/verify NFS AFM DR transport exports.
```

7. Issue the following command on the secondary cluster to prepare recovery snapshots that contain data that is transferred to the new primary cluster:

```
mmcesdr secondary failback --generate-recovery-snapshots  
--output-file-path "/root/" --input-file-path "/root/"
```

The system displays output similar to this:

```
Performing step 1/2, generating recovery snapshots for all AFM DR acting primary filesets.  

Transfer all data under snapshot located on acting primary cluster at:  

/gpfs/fs0/combo1/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-1 to  

fileset link point of fileset fs0:combo1 on new primary cluster.  

Transfer all data under snapshot located on acting primary cluster at:  

/gpfs/fs0/combo2/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-2 to  

fileset link point of fileset fs0:combo2 on new primary cluster.
```

mmcesdr

```
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs0/nfs-ganeshal/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-3 to  
fileset link point of fileset fs0:nfs-ganeshal on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs0/nfs-ganeshal2/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-4 to  
fileset link point of fileset fs0:nfs-ganeshal2 on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs0/smb1/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-5 to  
fileset link point of fileset fs0:smb1 on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs0/smb2/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEBF-6 to  
fileset link point of fileset fs0:smb2 on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs1/.async_dr/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEC-2 to  
fileset link point of fileset fs1:async_dr on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs1/obj_softpolicy1/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEC-3 to  
fileset link point of fileset fs1:obj_softpolicy1 on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs1/obj_softpolicy2/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEC-4 to  
fileset link point of fileset fs1:obj_softpolicy2 on new primary cluster.  
Transfer all data under snapshot located on acting primary cluster at:  
/gpfs/fs1/object_fileset/.snapshots/psnap0-newprimary-base-rpo-090B66F65623DEC-1 to  
fileset link point of fileset fs1:object_fileset on new primary cluster.  
Successfully completed step 1/2, generating recovery snapshots for all AFM DR acting primary filesets.  
Performing step 2/2, creation of recovery output file for failback to new primary.  
Successfully completed step 2/2, creation of recovery output file for failback to new primary.
```

File to be used with new primary cluster in next step of failback to new primary cluster: /root//DR_Config

8. Issue the following command on the primary cluster to restore the protocol and export services configuration information:

```
mmcesdr primary restore --new-primary
```

The system displays output similar to this:

```
Restoring cluster and enabled protocol configurations/exports.  
Successfully completed restoring cluster and enabled protocol configurations/exports.
```

9. Issue the following command on the secondary cluster to restore the protocol and export services configuration information:

```
mmcesdr secondary fallback --post-failback-complete --new-primary --input-file-path "/root"
```

The system displays output similar to this:

```
Performing step 1/2, converting protected filesets back into AFM DR secondary filesets.  
Successfully completed step 1/2, converting protected filesets back into AFM DR secondary filesets.  
Performing step 2/2, restoring AFM DR-based NFS share configuration.  
Successfully completed step 2/2, restoring AFM DR-based NFS share configuration.
```

10. Issue the following command on the primary cluster to back up configuration:

```
mmcesdr primary backup
```

The system displays output similar to this:

```
Performing step 1/2, configuration fileset creation/verification.  
Successfully completed step 1/2, configuration fileset creation/verification.  
Performing step 2/2, protocol and export services configuration backup.  
Successfully completed step 2/2, protocol and export services configuration backup.
```

11. Issue the following command on the primary cluster to restore configuration when the primary cluster is not in a protocols DR relationship with another cluster:

```
mmcesdr primary restore --file-config --restore
```

The system displays output similar to this:

```
Restoring cluster and enabled protocol configurations/exports.  
Successfully completed restoring cluster and enabled protocol configurations/exports.
```

```
=====  
= If all steps completed successfully, remove and then re-create file  
= authentication on the Primary Cluster.  
= Once this is complete, Protocol Cluster Configuration Restore will be complete.  
=====
```

See also

- “mmafmctl command” on page 40
- “mmces command” on page 103
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmobj command” on page 466
- “mmpsnap command” on page 501
- “mmrestorefs command” on page 525
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

Location

/usr/lpp/mmfs/bin

mmchattr command

Changes attributes of one or more GPFS files.

Synopsis

```
mmchattr [-m MetadataReplicas] [-M MaxMetadataReplicas]
          [-r DataReplicas] [-R MaxDataReplicas] [-P DataPoolName]
          [-D {yes | no}] [-I {yes | defer}] [-i {yes | no}]
          [-a {yes | no}] [-1]
          [{--set-attr AttributeName[=Value] [--pure-attr-create | --pure-attr-replace]} |
           {--delete-attr AttributeName [--pure-attr-delete]}]
          [--hex-attr] [--hex-attr-name] [--no-attr-ctime]
          [--compact[=NumDirectoryEntries]]
          [--compression {yes | no}]
          [--block-group-factor BlockGroupFactor]
          [--write-affinity-depth WriteAffinityDepth]
          [--write-affinity-failure-group "WadfgValueString"]
          [--indefinite-retention {yes | no}]
          [--expiration-time yyyy-mm-dd[@hh:mm:ss]]
          Filename [Filename...]
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmchattr** command to change the replication attributes, storage pool assignments, retention and immutability attributes, I/O caching policy, and file compression or decompression of files in the file system.

The replication factor must be less than or equal to the maximum replication factor for the file. If insufficient space is available in the file system to increase the number of replicas to the value requested, the **mmchattr** command ends. However, some blocks of the file might have their replication factor increased after the **mmchattr** command ends. If more free space becomes available in the file system later (when, for example, you add another disk to the file system), you can then issue the **mmrestripefs** command with the **-r** or **-b** option to complete the replication of the file. The **mmrestripefile** command can be used in a similar manner. You can use the **mmlsattr** command to display the replication values.

Data of a file is stored in a specific storage pool. A storage pool is a collection of disks or RAIDs with similar properties. Because these storage devices have similar properties, you can manage them as a group. You can use storage pools to do the following tasks:

- Partition storage for the file system
- Assign file storage locations
- Improve system performance
- Improve system reliability

The Direct I/O caching policy bypasses file cache and transfers data directly from disk into the user space buffer, as opposed to using the normal cache policy of placing pages in kernel memory. Applications with poor cache hit rates or a large amount of I/O might benefit from the use of Direct I/O.

The **mmchattr** command can be run against a file in use.

You must have write permission for the files whose attributes you are changing.

Parameters

-m MetadataReplicas

Specifies how many copies of the file system's metadata to create. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of the *MaxMetadataReplicas* attribute of the file.

-M MaxMetadataReplicas

Specifies the maximum number of copies of indirect blocks for a file. Space is reserved in the inode for all possible copies of pointers to indirect blocks. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be less than the value of the *DefaultMetadataReplicas* attribute of the file.

-r DataReplicas

Specifies how many copies of the file data to create. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of the *MaxDataReplicas* attribute of the file.

-R MaxDataReplicas

Specifies the maximum number of copies of data blocks for a file. Space is reserved in the inode and indirect blocks for all possible copies of pointers to data blocks. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be less than the value of the *DefaultDataReplicas* attribute of the file.

-P DataPoolName

Changes the assigned storage pool of the file to the specified *DataPoolName*. The caller must have superuser or root privileges to change the assigned storage pool.

-D {yes | no}

Enable or disable the Direct I/O caching policy for files.

-I {yes | defer}

Specifies whether replication and migration between pools, or file compression or decompression, is to be performed immediately (**-I yes**), or deferred until a later call to **mmrestripefs** or **mmrestripefile** (**-I defer**). By deferring the operation, you can complete it when the system is not loaded with processes or I/O. Also, if multiple files are affected, the data movement can be done in parallel. The default is **-I yes**. For more information about file compression and decompression, see the **--compression** option in this topic.

-i {yes | no}

Specifies whether the file is immutable (**-i yes**) or not immutable (**-i no**).

Note: The immutability attribute is specific to the current instance of the file. Restoring an image of the file to another location does not retain the immutability option. You must set it yourself.

-a {yes | no}

Specifies whether the file is in **appendOnly** mode (**-a yes**) or not (**-a no**).

Notes:

1. The **appendOnly** setting is specific to the current instance of the file. Restoring an image of the file to another location does not retain the **appendOnly** mode. You must set it yourself.
2. **appendOnly** mode is not supported for AFM filesets.

-l

Specifies that this command works only with regular files and directories and does not follow symlinks. The default is to follow symlinks.

--set-attr AttributeName [=Value]

Sets the specified extended attribute name to the specified *Value* for each file. If no *Value* is specified, **--set-attr AttributeName** sets the extended attribute name to a zero-length value.

--pure-attr-create

When this option is used, the command fails if the specified extended attribute exists.

mmchattr

--pure-attr-replace

When this option is used, the command fails if the specified extended attribute does not exist.

--delete-attr *AttributeName*

Removes the extended attribute.

For example, To remove wad, wadfg, and bgf, enter the following command:

```
mmchattr --delete-attr gpfs.WAD,gpfs.WADFG,gpfs.BGF
```

--pure-attr-delete

When this option is used, the command fails if the specified extended attribute does not exist.

--hex-attr

Inputs the attribute value in hex.

--hex-attr-name

Inputs the attribute name in hex.

--no-attr-ctime

Changes the attribute without setting the ctime of the file. This is restricted to root only.

--compact [=NumDirectoryEntries]

Sets the minimum compaction size of the directories that are specified in the *Filename* parameter. The minimum compaction size is the number of directory slots, including both full and empty slots, that a directory is allowed to retain when it is automatically compacted. In IBM Spectrum Scale v4.1 or later, by default, a directory is automatically compacted as far as possible. However, in systems in which many files are added to and removed from a directory in a short time, file system performance might be improved by setting the minimum compaction size of the directory.

The **compact** parameter sets the minimum compaction size of a directory to the specified number of slots. For example, if a directory contains 5,000 files and you set the minimum compaction size to 50,000, then the file system adds 45,000 directory slots. The directory can grow beyond 50,000 entries, but the file system does not allow the directory to be compacted below 50,000 slots.

Set *NumDirectoryEntries* to the total number of directory slots that you want to keep, including files that the directory already contains. You can specify the number of directory slots either as an integer or as an integer followed by the letter k (1000 slots) or m (1,000,000 slots). If you expect the average length of file names to be greater than 19 bytes, calculate the number of slots by the following formula:

```
NumDirectoryEntries = n * ( 1 + ceiling((namelen - 19)/32)
```

where:

n Specifies the number of entries (file names) in the directory.

ceiling()

Is a function that rounds a fractional number up to the next highest integer. For example, ceiling(1.03125) returns 2.

namelen

Specifies the expected average length of file names.

For example, if you want 50,000 entries with an average file name length of 48, then
NumDirectoryEntries = 50000 * (1 + 1) = 100000.

To restore the default behavior of the file system, specify the compact=0. The directory is compacted as far as possible.

To see the current value of this parameter, run the **mmlsattr** command with the **-L** option. For more information, see the topic “**mmlsattr command**” on page 392. To set or read the value of this parameter in a program, see the topics “**gpfs_malloc()** subroutine” on page 781 and “**gpfs_fstat_x()** subroutine” on page 698.

| **Note:** If the file system was created earlier than IBM Spectrum Scale v4.1, the *NumDirectoryEntries* value is not supported. The **compact** parameter specified by itself converts the directory to v4.1 format and compacts the directory as far as possible.

--compression {yes | no}

Compresses or decompresses the specified files. You can use the **-I defer** option to defer the operation until a later call to **mmrestripefs** or **mmrestripefile**. For more information, see the topic *File compression* in the *IBM Spectrum Scale: Administration Guide*.

--block-group-factor BlockGroupFactor

Specifies how many file system blocks are laid out sequentially on disk to behave like a single large block. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

--write-affinity-depth WriteAffinityDepth

Specifies the allocation policy to be used. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

--write-affinity-failure-group "WadfgValueString"

Indicates the range of nodes (in a shared nothing architecture) where replicas of blocks in the file are to be written. You use this parameter to determine the layout of a file in the cluster so as to optimize the typical access patterns of your applications. This applies only to a new data block layout; it does not migrate previously existing data blocks.

"*WadfgValueString*" is a semicolon-separated string identifying one or more failure groups in the following format:

FailureGroup1[;FailureGroup2[;FailureGroup3]]

where each *FailureGroupx* is a comma-separated string identifying the rack (or range of racks), location (or range of locations), and node (or range of nodes) of the failure group in the following format:

Rack1{::Rack2{...{:Rackx}}},Location1{::Location2{...{:Locationx}}},ExtLg1{::ExtLg2{...{:ExtLgx}}}

For example, the following value

1,1,1:2;2,1,1:2;2,0,3:4

means that the first failure group is on rack 1, location 1, extLg 1 or 2; the second failure group is on rack 2, location 1, extLg 1 or 2; and the third failure group is on rack 2, location 0, extLg 3 or 4.

If the end part of a failure group string is missing, it is interpreted as 0. For example, the following are interpreted the same way:

2
2,0
2,0,0

Notes:

- Only the end part of a failure group string can be left off. The missing end part may be the third field only, or it may be both the second and third fields; however, if the third field is provided, the second field must also be provided. The first field must *always* be provided. In other words, every comma must both follow and precede a number; therefore, *none* of the following are valid:

2,0,
2,
,0,0
0,,0
,,0

- Wildcard characters (*) are supported in these fields.

mmchattr

--indefinite-retention {yes | no}

Turns indefinite retention on or off. An alternative form of this parameter is **-e {yes | no}**. See **--expiration-time**.

--expiration-time yyyy-mm-dd[@hh:mm:ss]

Specifies the expiration time. An alternative form of this parameter is **-E yyyy-mm-dd[@hh:mm:ss]**. Expiration time and indefinite retention are independent attributes. You can change the value of either one without affecting the value of the other.

Filename

The name of the file to be changed. You must enter at least one file name; if you specify more than one, delimit each file name by a space. Wildcard characters are supported in file names; for example, **project*.sched**.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have write access to the file to run the **mmchattr** command.

You can issue the **mmchattr** command only from a node in the GPFS cluster where the file system is mounted.

Examples

1. To change the metadata replication factor to 2 and the data replication factor to 2 for the **project7.resource** file in file system **fs1**, issue this command:

```
mmchattr -m 2 -r 2 /fs1/project7.resource
```

To confirm the change, issue this command:

```
mmlsattr project7.resource
```

The system displays information similar to:

```
replication factors
metadata(max) data(max) file      [flags]
-----
2 ( 2 ) 2 ( 2 ) /fs1/project7.resource
```

2. Migrating data from one storage pool to another using the **mmchattr** command with the **-I defer** option, or the **mmapplypolicy** command with the **-I defer** option causes the data to be ill-placed. This means that the storage pool assignment for the file has changed, but the file data has not yet been migrated to the assigned storage pool.

The **mmlsattr -L** command causes show ill-placed flags on the files that are ill-placed. The **mmrestripefs**, or **mmrestripefile** command can be used to migrate data to the correct storage pool, and the ill-placed flag is cleared. This is an example of an ill-placed file:

```
mmlsattr -L 16Kfile6.tmp
```

The system displays output similar to this:

```
file name:          16Kfile6.tmp
metadata replication: 1 max 2
data replication:   1 max 2
immutable:          no
appendOnly:         no
flags:              directio
storage pool name: system
```

```
fileset name:      root
snapshot name:
creation time:    Thu Mar 28 14:49:23 2013
Misc attributes:  ARCHIVE
```

3. The following example shows the result of using the **--set-attr** parameter.

```
mmchattr --set-attr user.pfs001=testuser 16Kfile7.tmp
mmlsattr -L -d 16Kfile7.tmp
```

The system displays output similar to this:

```
file name:          16Kfile7.tmp
metadata replication: 1 max 2
data replication:   1 max 2
immutable:          no
appendOnly:         no
flags:
storage pool name: system
fileset name:       root
snapshot name:
creation Time:     Fri Feb 24 12:00:13 2012
Misc attributes:   ARCHIVE
user.pfs001:        "testuser"
```

4. To set the write affinity failure group for a file and to see the results, issue these commands:

```
mmchattr --write-affinity-failure-group="64,0,0;128,0,1;128,0,2" /gpfs1/testfile
mmlsattr -L /gpfs1/testfile
```

The system displays output similar to this:

```
file name:          /gpfs1/testfile
metadata replication: 3 max 3
data replication:   3 max 3
immutable:          no
appendOnly:         no
flags:
storage pool name: system
fileset name:       root
snapshot name:
Write Affinity Depth Failure Group(FG) Map for copy:1 64,0,0
Write Affinity Depth Failure Group(FG) Map for copy:2 128,0,1
Write Affinity Depth Failure Group(FG) Map for copy:3 128,0,2
creation time:      Wed Sep 12 02:53:18 2012
Misc attributes:   ARCHIVE
```

See also

- “mmcrfs command” on page 250
- “mmlsattr command” on page 392
- “mmlsfs command” on page 411

Location

/usr/lpp/mmfs/bin

mmchcluster

mmchcluster command

Changes GPFS cluster configuration data.

Synopsis

```
mmchcluster --ccr-enable
```

or

```
mmchcluster {[--ccr-disable] [-p PrimaryServer] [-s SecondaryServer]}
```

or

```
mmchcluster -p LATEST
```

or

```
mmchcluster {[[-r RemoteShellCommand] [-R RemoteFileCopyCommand] [--no-use-sudo-wrapper]]} |  
--use-sudo-wrapper
```

or

```
mmchcluster -C ClusterName
```

Availability

Available on all IBM Spectrum Scale editions.

Description

The **mmchcluster** command serves several purposes. You can use it to do the following:

1. Change the remote shell and remote file copy programs to be used by the nodes in the cluster.
2. Change the cluster name.
3. Enable or disable the cluster configuration repository (CCR).

When using the traditional server-based (non-CCR) configuration repository, you can also do the following:

1. Change the primary or secondary GPFS cluster configuration server.
2. Synchronize the primary GPFS cluster configuration server.

To display current system information for the cluster, issue the **mmlscluster** command.

For information on how to specify node names, see the topic *Specifying nodes as inputs to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

When issuing the **mmchcluster** command with the **-p** or **-s** options, the specified nodes must be available in order for the command to succeed. If any of the nodes listed are not available when the command is issued, a message listing those nodes is displayed. You must correct the problem on each node and reissue the command.

Attention: The **mmchcluster** command, when issued with either the **-p** or **-s** option, is designed to operate in an environment where the current primary and secondary cluster configuration servers are **not** available. As a result, the command can run without obtaining its regular serialization locks. To assure smooth transition to a new cluster configuration server, no other GPFS commands (**mm** commands) should be running when the command is issued, nor should any other command be issued until the **mmchcluster** command has successfully completed.

Parameters

--ccr-enable

Enables the configuration server repository (CCR), which stores redundant copies of configuration data files on all quorum nodes. The advantage of CCR over the traditional primary or backup configuration server semantics is that when using CCR, all GPFS administration commands as well as file system mounts and daemon startups work normally as long as a majority of quorum nodes are accessible.

For more information, see the topic *Cluster configuration data files* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

The CCR operation requires the use of the GSKit toolkit for authenticating network connections. As such, the **gpfs.gskit** package, which is available on all Editions, should be installed.

--ccr-disable

Reverts to the traditional primary or backup configuration server semantics and destroys the CCR environment. GPFS must be shut down on all the nodes before disabling CCR.

-p PrimaryServer

Change the primary server node for the GPFS cluster data. This may be specified as a short or long node name, an IP address, or a node number.

LATEST – Synchronize all of the nodes in the GPFS cluster ensuring they are using the most recently specified primary GPFS cluster configuration server. If an invocation of the **mmchcluster** command fails, you are prompted to reissue the command and specify **LATEST** on the **-p** option to synchronize all of the nodes in the GPFS cluster. Synchronization provides for all nodes in the GPFS cluster to use the most recently specified primary GPFS cluster configuration server.

This option only applies when the traditional server-based configuration (non-CCR) repository is used.

-s SecondaryServer

Change the secondary server node for the GPFS cluster data. To remove the secondary GPFS server and continue operating without it, specify a null string, "", as the parameter. This may be specified as a short or long nodename, an IP address, or a node number.

This option only applies when the traditional server-based configuration (non-CCR) repository is used.

-r RemoteShellCommand

Specifies the fully-qualified path name for the remote shell program to be used by GPFS.

The remote shell command must adhere to the same syntax format as the **ssh** command, but may implement an alternate authentication mechanism.

-R RemoteFileCopy

Specifies the fully-qualified path name for the remote file copy program to be used by GPFS.

The remote copy command must adhere to the same syntax format as the **scp** command, but may implement an alternate authentication mechanism.

--no-use-sudo-wrapper

Specifies that the cluster reverts to using the default remote shell program and remote copy program. For more information, see the topic *Running IBM Spectrum Scale without remote root login* in the *IBM Spectrum Scale: Administration Guide*.

--use-sudo-wrapper

Specifies that the nodes in the cluster call the ssh sudo wrapper script and the scp sudo wrapper script as the remote shell program and the remote copy program. For more information, see the topic *Running IBM Spectrum Scale without remote root login* in the *IBM Spectrum Scale: Administration Guide*.

-c ClusterName

Specifies a new name for the cluster. If the user-provided name contains dots, it is assumed to be a

mmchcluster

fully qualified domain name. Otherwise, to make the cluster name unique, the domain of the first quorum node or, if specified, the primary configuration server will be appended to the user-provided name.

Since each cluster is managed independently, there is no automatic coordination and propagation of changes between clusters like there is between the nodes within a cluster. This means that if you change the name of the cluster, you should notify the administrators of all other GPFS clusters that can mount your file systems so that they can update their own environments.

Before running this option, ensure that all GPFS daemons on all nodes have been stopped.

See the **mmauth**, **mmremotecluster**, and **mmremoteefs** commands.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmchcluster** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

Examples

To change the primary GPFS server for the cluster, issue this command:

```
mmchcluster -p k164n06
```

The system displays output similar to:

```
mmchcluster: Command successfully completed
```

To confirm the change, issue this command:

```
mmlscluster
```

The system displays information similar to:

```
GPFS cluster information
```

```
=====
```

```
GPFS cluster name:      cluster1.kgn.ibm.com
GPFS cluster id:       680681562214606028
GPFS UID domain:      cluster1.kgn.ibm.com
Remote shell command:   /usr/bin/rsh
Remote file copy command: /usr/bin/rcp
```

```
GPFS cluster configuration servers:
```

```
-----
```

```
Primary server:    k164n06.kgn.ibm.com
Secondary server:  k164n05.kgn.ibm.com
```

```
Node Daemon node name     IP address     Admin node name     Designation
```

```
-----
```

1	k164n04.kgn.ibm.com	198.117.68.68	k164n04.kgn.ibm.com	quorum
2	k164n05.kgn.ibm.com	198.117.68.71	k164n05.kgn.ibm.com	quorum
3	k164n06.kgn.ibm.com	198.117.68.70	k164n06.kgn.ibm.com	

See also

- “[mmaddnode command](#)” on page 29
- “[mmchnode command](#)” on page 190
- “[mmcrcluster command](#)” on page 239
- “[mmdelnod command](#)” on page 297
- “[mmlscluster command](#)” on page 397
- “[mmremotecluster command](#)” on page 511

Location

/usr/lpp/mmfs/bin

mmchconfig

mmchconfig command

Changes GPFS configuration parameters.

Synopsis

```
mmchconfig Attribute=value[,Attribute=value...] [-i | -I]  
[-N {Node[,Node...]} | NodeFile | NodeClass]
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmchconfig** command to change the GPFS configuration attributes on a single node, a set of nodes, or globally for the entire cluster.

Results

The configuration is updated on the specified nodes.

Parameters

-I Specifies that the changes take effect immediately, but do not persist when GPFS is restarted. This option is valid only for the following attributes:

- **deadlockBreakupDelay**
- **deadlockDataCollectionDailyLimit**
- **deadlockDataCollectionMinInterval**
- **deadlockDetectionThreshold**
- **deadlockDetectionThresholdForShortWaiters**
- **deadlockOverloadThreshold**
- **dmapiMountTimeout**
- **dmapiSessionFailureTimeout**
- **expelDataCollectionDailyLimit**
- **expelDataCollectionMinInterval**
- **fastestPolicyCmpThreshold**
- **fastestPolicyMaxValidPeriod**
- **fastestPolicyMinDiffPercent**
- **fastestPolicyNumReadSamples**
- **fileHeatLossPercent**
- **fileHeatPeriodMinutes**
- **ignorePrefetchLUNCount**
- **IrocData**
- **IrocDataMaxFileSize**
- **IrocDataStubFileSize**
- **IrocDirectories**
- **IrocInodes**
- **maxMBpS**
- **nfsPrefetchStrategy**
- **nsdBufSpace**

- nsdInlineWriteMax
- nsdMultiQueue
- pagepool
- pitWorkerThreadsPerNode
- readReplicaPolicy
- seqDiscardThreshold
- syncbuffsperiteration
- systemLogLevel
- unmountOnDiskFail
- verbsRdmaRoCEToS
- verbsRdmasesPerConnection
- verbsRdmasesPerNode
- verbsSendBufferMemoryMB
- worker1Threads (only when adjusting value down)
- writebehindThreshold

-i Specifies that the changes take effect immediately and are permanent. This option is valid only for the following attributes:

- cesSharedRoot
- cnfsGrace
- cnfsMountdPort
- cnfsNFSDprocs
- cnfsReboot
- cnfsSharedRoot
- cnfsVersions
- commandAudit
- dataDiskWaitTimeForRecovery
- dataStructureDump
- deadlockBreakupDelay
- deadlockDataCollectionDailyLimit
- deadlockDataCollectionMinInterval
- deadlockDetectionThreshold
- deadlockDetectionThresholdForShortWaiters
- deadlockOverloadThreshold
- debugDataControl
- disableInodeUpdateOnFDatasync
- dmapiMountTimeout
- dmapiSessionFailureTimeout
- expelDataCollectionDailyLimit
- expelDataCollectionMinInterval
- fastestPolicyCmpThreshold
- fastestPolicyMaxValidPeriod
- fastestPolicyMinDiffPercent
- fastestPolicyNumReadSamples
- fileHeatLossPercent
- fileHeatPeriodMinutes

mmchconfig

- **forceLogWriteOnFdatasync**
- **ignorePrefetchLUNCount**
- **IrocData**
- **IrocDataMaxFileSize**
- **IrocDataStubFileSize**
- **IrocDirectories**
- **IrocInodes**
- **maxDownDisksForRecovery**
- **maxFailedNodesForRecovery**
- **maxMBpS**
- **metadataDiskWaitTimeForRecovery**
- **minDiskWaitTimeForRecovery**
- **mmfsLogTimeStampISO8601**
- **nfsPrefetchStrategy**
- **nsdBufSpace**
- **nsdInlineWriteMax**
- **nsdMultiQueue**
- **pagepool**
- **pitWorkerThreadsPerNode**
- **readReplicaPolicy**
- **restripeOnDiskFailure**
- **seqDiscardThreshold**
- **syncbuffsperiteration**
- **systemLogLevel**
- **unmountOnDiskFail**
- **verbsRdmaRoCEToS**
- **verbsRdmasesPerConnection**
- **verbsRdmasesPerNode**
- **verbsSendBufferMemoryMB**
- **worker1Threads** (only when adjusting value down)
- **writebehindThreshold**

-N {Node[,Node...] | NodeFile | NodeClass}

Specifies the set of nodes to which the configuration changes apply. The default is **-N all**.

For information on how to specify node names, see the topic *Specifying nodes as inputs to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

To see a complete list of the attributes for which the **-N** flag is valid, see the see the table "Configuration attributes on the **mmchconfig** command" in the topic *Changing the GPFS cluster configuration data* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

Attribute=value

Specifies the name of the attribute to be changed and its associated *value*. More than one attribute and value pair can be specified. To restore the GPFS default setting for an attribute, specify **DEFAULT** as its *value*.

This command accepts the following attributes:

adminMode

Specifies whether all nodes in the cluster are used for issuing GPFS administration commands or just a subset of the nodes. Valid values are:

allToAll

Indicates that all nodes in the cluster are used for running GPFS administration commands and that all nodes are able to execute remote commands on any other node in the cluster without the need of a password.

central

Indicates that only a subset of the nodes is used for running GPFS commands and that only those nodes are able to execute remote commands on the rest of the nodes in the cluster without the need of a password.

For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

afmAsyncDelay

Specifies (in seconds) the amount of time by which write operations are delayed (because write operations are asynchronous with respect to remote clusters). For write-intensive applications that keep writing to the same set of files, this delay is helpful because it replaces multiple writes to the home cluster with a single write containing the latest data. However, setting a very high value weakens the consistency of data on the remote cluster.

This configuration parameter is applicable only for writer caches (SW, IW, and primary), where data from cache is pushed to home.

Valid values are between 1 and 2147483647. The default is 15.

afmDirLookupRefreshInterval

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a directory, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of that directory has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 60. In situations where home cluster data changes frequently, a value of 0 is recommended.

afmDirOpenRefreshInterval

Controls the frequency of data revalidations that are triggered by such I/O operations as **read** or **write** (specified in seconds). After a directory has been cached, **open** requests resulting from I/O operations on that object are directed to the cached directory until the specified amount of time has passed. Once the specified amount of time has passed, the **open** request gets directed to a gateway node rather than to the cached directory.

Valid values are between 0 and 2147483647. The default is 60. Setting a lower value guarantees a higher level of consistency.

afmDisconnectTimeout

The Waiting period in seconds to detect the status of the home cluster. If the home cluster is inaccessible, the metadata server (MDS) changes the state to 'disconnected'.

afmExpirationTimeout

Is used with **afmDisconnectTimeout** (which can be set only through **mmchconfig**) to control how long a network outage between the cache and home clusters can continue before the data in the cache is considered out of sync with home. After **afmDisconnectTimeout** expires, cached data remains available until **afmExpirationTimeout** expires, at which point the cached data is considered expired and cannot be read until a reconnect occurs.

Valid values are 0 through 2147483647. The default is disable.

mmchconfig

afmFileLookupRefreshInterval

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a file, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of the file has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 30. In situations where home cluster data changes frequently, a value of 0 is recommended.

afmFileOpenRefreshInterval

Controls the frequency of data revalidations that are triggered by such I/O operations as **read** or **write** (specified in seconds). After a file has been cached, **open** requests resulting from I/O operations on that object are directed to the cached file until the specified amount of time has passed. Once the specified amount of time has passed, the **open** request gets directed to a gateway node rather than to the cached file.

Valid values are 0 through 2147483647. The default is 30. Setting a lower value guarantees a higher level of consistency.

afmHardMemThreshold

Sets a limit to the maximum amount of memory that AFM can use on each gateway node to record changes to the file system. After this limit is reached, the filesset goes into a 'dropped' state.

Exceeding the limit and the filesset going into a 'dropped' state due to accumulated pending requests might occur if -

- the cache cluster is disconnected for an extended period of time.
- the connection with the home cluster is on a low bandwidth.

afmHashVersion

Specifies the version of hashing algorithm to be used to assign AFM and AFM ADR filessets across gateway nodes, thus running as few recoveries as possible. This minimizes impact of gateway nodes joining or leaving the active cluster.

Valid values are 1,2 or 4. Default value is 2.

afmMaxParallelRecoveries

Specifies the number of filessets per gateway node on which event recovery is run. The default value is 0. When the value is 0, event recovery is run on all filessets of the gateway node.

afmNumReadThreads

Defines the number of threads that can be used on each participating gateway node during parallel read. The default value of this parameter is 1; that is, one reader thread will be active on every gateway node for each big read operation qualifying for splitting per the parallel read threshold value. The valid range of values is 1 to 64.

afmNumWriteThreads

Defines the number of threads that can be used on each participating gateway node during parallel write. The default value of this parameter is 1; that is, one writer thread will be active on every gateway node for each big write operation qualifying for splitting per the parallel write threshold value. Valid values can range from 1 to 64.

afmParallelReadChunkSize

Defines the minimum chunk size of the read that needs to be distributed among the gateway nodes during parallel reads. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at filesset level using **mmcrfilesset** or **mmchfilesset** commands.

afmParallelReadThreshold

Defines the threshold beyond which parallel reads become effective. Reads are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value is 1024 MiB. The valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

afmParallelWriteChunkSize

Defines the minimum chunk size of the write that needs to be distributed among the gateway nodes during parallel writes. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

afmParallelWriteThreshold

Defines the threshold beyond which parallel writes become effective. Writes are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value of this parameter is 1024 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

afmReadSparseThreshold

Specifies the size in MB for files in cache beyond which sparseness is maintained. For all files below the specified threshold, sparseness is not maintained.

afmSecondaryRW

Specifies if the secondary is read-write or not.

yes

Specifies that the secondary is read-write.

no

Specifies that the secondary is not read-write.

afmShowHomeSnapshot

Controls the visibility of the home snapshot directory in cache. For this to be visible in cache, this variable has to be set to **yes**, and the snapshot directory name in cache and home should not be the same.

yes

Specifies that the home snapshot link directory is visible.

no

Specifies that the home snapshot link directory is not visible.

See *Peer snapshot -psnap* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

atimeDeferredSeconds

Controls the update behavior of **atime** when the **relatime** option is enabled. The default value is 86400 seconds (24 hours). A value of 0 effectively disables **relatime** and causes the behavior to be the same as the **atime** setting.

For more information, see the topic *GPFS-specific mount options* in the *IBM Spectrum Scale: Administration Guide*.

autoload

Starts GPFS automatically whenever the nodes are rebooted. Valid values are **yes** or **no**.

The **-N** flag is valid for this attribute.

automountDir

Specifies the directory to be used by the Linux automounter for GPFS file systems that are being mounted automatically. The default directory is `/gpfs/automountdir`. This parameter does not apply to AIX and Windows environments.

mmchconfig

cesSharedRoot

Specifies a directory in a GPFS file system to be used by the Cluster Export Services (CES) subsystem. For the CES shared root, the recommended value is a dedicated file system, but it is not enforced. The CES shared root can also be a part of an existing GPFS file system. In any case, **cesSharedRoot** must reside on GPFS and must be available when it is configured through **mmchconfig**.

GPFS must be down on all CES nodes in the cluster when changing the **cesSharedRoot** attribute.

cifsBypassTraversalChecking

Controls the GPFS behavior while performing access checks for directories

GPFS grants the SEARCH access when the following conditions are met:

- The object is a directory
- The parameter value is **yes**
- The calling process is a Samba process

GPFS grants the SEARCH access regardless of the mode or ACL.

cipherList

Sets the security mode for the cluster. The security mode determines the level of the security that the cluster provides for communications between nodes in the cluster and also for communications with other clusters. There are three security modes:

EMPTY

The sending node and the receiving node do not authenticate each other, do not encrypt transmitted data, and do not check data integrity.

AUTHONLY

The sending and receiving nodes authenticate each other, but they do not encrypt transmitted data and do not check data integrity. This mode is the default in IBM Spectrum Scale V4.2 or later.

Cipher

The sending and receiving nodes authenticate each other, encrypt transmitted data, and check data integrity. To set this mode, you must specify the name of a supported cipher, such as AES128-GCM-SHA256.

Note: Although after **mmchconfig** is issued, the mmfsd daemon accepts the new **cipherList** immediately, it uses the cipher only for new TCP/TLS connections to other nodes. Existing mmfsd daemon connections remain with the prior **cipherList** settings. For the new **cipherList** to take complete effect immediately, GPFS needs to be restarted on all nodes in a rolling fashion, one node at a time, to prevent cluster outage. When a cipher other than **AUTHONLY** or **EMPTY** is in effect, it can lead to significant performance degradation, as this results in encryption and data integrity verification of the transmitted data.

For more information about the security mode and supported ciphers, see the topic *Security mode* in the *IBM Spectrum Scale: Administration Guide*.

cnfsGrace

Specifies the number of seconds a CNFS node will deny new client requests after a node failover or fallback, to allow clients with existing locks to reclaim them without the possibility of some other client that is being granted a conflicting access. For v3, only new lock requests are denied. For v4, new lock, read, and write requests are rejected. The **cnfsGrace** value also determines the time period for the server lease.

Valid values are 10 - 600. The default is 90 seconds. While a short grace period is good for fast server failover, it comes at the cost of increased load on server to effect lease renewal.

GPFS must be down on all CNFS nodes in the cluster when changing the **cnfsGrace** attribute.

cnfsMountdPort

Specifies the port number to be used for **rpc.mountd**. See the *IBM Spectrum Scale: Administration Guide* for restrictions and additional information.

cnfsNFSDprocs

Specifies the number of **nfsd** kernel threads. The default is 32.

cnfsReboot

Specifies whether the node reboots when CNFS monitoring detects an unrecoverable problem that can be handled only by node failover.

Valid values are **yes** or **no**. The default is **yes** and recommended. If node reboot is not desired for other reasons, it should be noted that clients that were communicating with the failing node are likely to get errors or hang. CNFS failover is only guaranteed with **cnfsReboot** enabled.

The **-N** flag is valid for this attribute.

cnfsSharedRoot

Specifies a directory in a GPFS file system to be used by the clustered NFS subsystem.

GPFS must be down on all CNFS nodes in the cluster when changing the **cnfsSharedRoot** attribute.

See the *IBM Spectrum Scale: Administration Guide* for restrictions and additional information.

cnfsVersions

Specifies a comma-separated list of protocol versions that CNFS should start and monitor.

The default is 3,4.

GPFS must be down on all CNFS nodes in the cluster when changing the **cnfsVersions** attribute.

See the *IBM Spectrum Scale: Administration Guide* for additional information.

commandAudit

Controls the logging of audit messages for GPFS commands that change the configuration of the cluster. This attribute is not supported on Windows operating systems. For more information, see the topic *Audit messages for cluster configuration changes* in the *IBM Spectrum Scale: Problem Determination Guide*.

on Starts audit messages. Messages go to syslog and the GPFS log.

syslogOnly

Starts audit messages. Messages go to syslog only. This value is the default.

off Stops audit messages.

The **-N** flag is valid for this attribute.

dataDiskCacheProtectionMethod

The **dataDiskCacheProtectionMethod** parameter defines the cache protection method for disks that are used for the GPFS file system. The valid values for this parameter are 0, 1, and 2.

The default value is 0. The default value indicates that the disks are Power-Protected and, when the down disk is started, only the standard GPFS log recovery is required. If the value of this parameter is 1, the disks are Power-Protected with no disk cache. GPFS works the same as before. If the value of this parameter is 2, when a node stops functioning, files that have data in disk cache must be recovered to a consistent state when the disk is started.

This parameter impacts only disks in the FPO storage pool. If the physical disk-write cache is enabled, the value of this parameter must be set to 2. Otherwise, maintain the default.

dataDiskWaitTimeForRecovery

Specifies a period, in seconds, during which the recovery of **dataOnly** disks is suspended to give the disk subsystem a chance to correct itself. This parameter is taken into account when the

mmchconfig

affected disks belong to a single failure group. If more than one failure group is affected, the delay is based on the value of **minDiskWaitTimeForRecovery**.

Valid values are 0 - 3600 seconds. The default is 3600. If **restripeOnDiskFailure** is **no**, **dataDiskWaitTimeForRecovery** has no effect.

dataStructureDump

Specifies a path for storing dumps. You can specify a directory or a symbolic link. The default is to store dumps in /tmp/mmfs. This attribute takes effect immediately whether or not -i is specified.

It is a good idea to create a directory or a symbolic link for problem determination information. Do not put it in a GPFS file system, because it might not be available if GPFS fails. When a problem occurs, GPFS can write 200 MiB or more of problem determination data into the directory. Copy and delete the files promptly so that you do not get a **NOSPACE** error if another failure occurs.

Important: Before you change the value of **dataStructureDump**, stop the GPFS trace. Otherwise you will lose GPFS trace data. Restart the GPFS trace afterward. For more information, see the topic *Generating GPFS trace reports* in the *IBM Spectrum Scale: Problem Determination Guide*.

The -N flag is valid for this attribute.

deadlockBreakupDelay

Specifies how long to wait after a deadlock is detected before attempting to break up the deadlock. Enough time must be provided to allow the debug data collection to complete.

The default is 0, which means that the automated deadlock breakup is disabled. A positive value enables the automated deadlock breakup. If automated deadlock breakup is to be enabled, a delay of 300 seconds or longer is recommended.

deadlockDataCollectionDailyLimit

Specifies the maximum number of times that debug data can be collected each day.

The default is 3. If the value is 0, then no debug data is collected when a potential deadlock is detected.

deadlockDataCollectionMinInterval

Specifies the minimum interval between two consecutive collections of debug data.

The default is 3600 seconds.

deadlockDetectionThreshold

Specifies the initial deadlock detection threshold. The effective deadlock detection threshold adjusts itself over time. A suspected deadlock is detected when a waiter waits longer than the effective deadlock detection threshold.

The default is 300 seconds. If the value is 0, then automated deadlock detection is disabled.

deadlockDetectionThresholdForShortWaiters

Specifies the deadlock detection threshold for short waiters that should never be long.

The default is 60 seconds.

deadlockOverloadThreshold

Specifies the threshold for detecting a cluster overload condition. If the overload index on a node exceeds the **deadlockOverloadThreshold**, then the effective **deadlockDetectionThreshold** is raised. The overload index is calculated heuristically and is based mainly on the I/O completion times.

The default is 1. If the value is 0, then overload detection is disabled.

debugDataControl

Controls the amount of debug data that is collected. This attribute takes effect immediately whether or not -i is specified. The -N flag is valid for this attribute.

none No debug data is collected.

light The minimum amount of debug data that is most important for debugging issues is collected. This is the default value.

medium

More debug data is collected.

heavy The maximum amount of debug data is collected, targeting internal test systems.

verbose

Needed only for troubleshooting special cases and can result in large dumps.

The following table provides more information about these settings:

Table 10. Settings for debugDataControl

Setting	Collect dump data	Collect trace information (if tracing is already running)	Collect a short sample of trace information ¹
none	No	No	No
light (default)	Yes	Yes	No
medium	Yes, more dump data	Yes	No
heavy (for internal test teams)	Yes, even more dump data	Yes	Yes
verbose (for developers)	Yes, all dump data	Yes	Yes

¹If trace is not running, turn tracing on, let it run for 20 seconds, and then turn trace off.

defaultHelperNodes

For commands that distribute work among a set of nodes, the **defaultHelperNodes** parameter specifies the nodes to be used. When you specify values, follow the rules that are described for the **-N** parameter.

To override this setting when you use such commands, explicitly specify the helper nodes with **-N**.

The commands that use **-N** for this purpose are the following: **mmaddirdisk**, **mmapplypolicy**, **mmbbackup**, **mmchdisk**, **mmcheckquota**, **mmdefragfs**, **mmdeldisk**, **mmdelsnapshot**, **mmfileid**, **mmfsck**, **mmimgbackup**, **mmimgrestore**, **mmrestorefs**, **mmrestripefs**, and **mmrpldisk**.

NodeClass values are listed.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

defaultMountDir

Specifies the default parent directory for GPFS file systems. The default value is **/gpfs**. If an explicit mount directory is not provided with the **mmcrfs**, **mmchfs**, or **mmremotesfs** command, the default mount point is set to *DefaultMountDir/DeviceName*.

disableInodeUpdateOnFdatasync

Controls the inode update on **fdatasync** for mtime and atime updates. Valid values are **yes** or **no**.

When **disableInodeUpdateOnFdatasync** is set to **yes**, the inode object is not updated on disk for mtime and atime updates on **fdatasync()** calls. File size updates are always synced to the disk.

When **disableInodeUpdateOnFdatasync** is set to **no**, the inode object is updated with the current mtime on **fdatasync()** calls. This is the default.

dmapiDataEventRetry

Controls how GPFS handles data events that are enabled again immediately after the event is handled by the DMAPI application. Valid values are as follows:

mmchconfig

- 1 Specifies that GPFS always regenerates the event as long as it is enabled. This value should be used only when the DMAPI application recalls and migrates the same file in parallel by many processes at the same time.
- 0 Specifies to never regenerate the event. This value should not be used if a file could be migrated and recalled at the same time.

RetryCount

Specifies the number of times the data event should be retried. The default is 2.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

dmapiEventTimeout

Controls the blocking of file operation threads of NFS, while in the kernel waiting for the handling of a DMAPI synchronous event. The parameter value is the maximum time, in milliseconds, the thread blocks. When this time expires, the file operation returns ENOTREADY, and the event continues asynchronously. The NFS server is expected to repeatedly retry the operation, which eventually finds the response of the original event and continue. This mechanism applies only to read, write, and truncate event types, and only when such events come from NFS server threads. The timeout value is given in milliseconds. The value 0 indicates immediate timeout (fully asynchronous event). A value greater than or equal to 86400000 (which is 24 hours) is considered *infinity* (no timeout, fully synchronous event). The default value is 86400000.

For the parameter change to take effect, restart the GPFS daemon on the nodes that are specified in the -N option. If the -N option is not used, restart the GPFS daemon on all nodes.

For further information about DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

The -N flag is valid for this attribute.

dmapiMountEvent

Controls the generation of the **mount**, **preunmount**, and **unmount** events. Valid values are:

all

mount, **preunmount**, and **unmount** events are generated on each node. This is the default behavior.

SessionNode

mount, **preunmount**, and **unmount** events are generated on each node and are delivered to the session node, but the session node does not deliver the event to the DMAPI application unless the event is originated from the **SessionNode** itself.

LocalNode

mount, **preunmount**, and **unmount** events are generated only if the node is a session node.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

dmapiMountTimeout

Controls the blocking of **mount** operations, waiting for a disposition for the mount event to be set. This timeout is activated, at most once on each node, by the first external mount of a file system that has DMAPI enabled, and only if there has never before been a mount disposition. Any **mount** operation on this node that starts while the timeout period is active waits for the mount disposition. The parameter value is the maximum time, in seconds, that the **mount** operation waits for a disposition. When this time expires and there is still no disposition for the mount event, the **mount** operation fails, returning the **EIO** error. The timeout value is given in full seconds. The value 0 indicates immediate timeout (immediate failure of the mount operation).

A value greater than or equal to 86400 (which is 24 hours) is considered *infinity* (no timeout, indefinite blocking until there is a disposition). The default value is 60.

The **-N** flag is valid for this attribute.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

dmapiSessionFailureTimeout

Controls the blocking of file operation threads, while in the kernel, waiting for the handling of a DMAPI synchronous event that is enqueued on a session that has experienced a failure. The parameter value is the maximum time, in seconds, the thread waits for the recovery of the failed session. When this time expires and the session has not yet recovered, the event is canceled and the file operation fails, returning the **EIO** error. The timeout value is given in full seconds. The value 0 indicates immediate timeout (immediate failure of the file operation). A value greater than or equal to 86400 (which is 24 hours) is considered *infinity* (no timeout, indefinite blocking until the session recovers). The default value is 0.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

The **-N** flag is valid for this attribute.

enableIPv6

Controls whether the GPFS daemon communicates through the IPv6 network. The following values are valid:

no Specifies that the GPFS daemon does not communicate through the IPv6 network. This is the default.

yes

Specifies that the GPFS daemon communicates through the IPv6 network. **yes** requires that the daemon be down on all nodes.

prepare

After the command completes, the daemons can be recycled on all nodes at a time chosen by the user (before proceeding to run the command with **commit** specified).

commit

Verifies that all currently active daemons have received the new value, allowing the user to add IPv6 nodes to the cluster.

Note:

Before changing the value of **enableIPv6**, the GPFS daemon on the primary configuration server must be inactive. After changing the parameter, the GPFS daemon on the rest of nodes in the cluster should be recycled. This can be done one node at a time.

To use IPv6 addresses for GPFS, the operating system must be properly configured as IPv6 enabled, and IPv6 addresses must be configured on all the nodes within the cluster.

enforceFilesetQuotaOnRoot

Controls whether fileset quotas should be enforced for the root user the same way as for any other users. Valid values are **yes** or **no**. The default is **no**.

expelDataCollectionDailyLimit

Specifies the maximum number of times that debug data associated with expelling nodes can be collected in a 24-hour period. Sometimes exceptions are made to help capture the most relevant debug data.

The default is 3. If the value is 0, then no expel-related debug data is collected.

mmchconfig

expelDataCollectionMinInterval

Specifies the minimum interval, in seconds, between two consecutive expel-related data collection attempts on the same node.

The default is 3600 seconds.

failureDetectionTime

Indicates to GPFS the amount of time it takes to detect that a node has failed.

GPFS must be down on all the nodes when changing the **failureDetectionTime** attribute.

fastestPolicyCmpThreshold

Indicates the disk comparison count threshold, above which GPFS forces selection of this disk as the preferred disk to read and update its current speed.

Valid values are ≥ 3 . The default is 50. In a system with SSD and regular disks, the value of the **fastestPolicyCmpThreshold** parameter can be set to a greater number to let GPFS refresh the speed statistics for slower disks less frequently.

fastestPolicyMaxValidPeriod

Indicates the time period after which the disk's current evaluation is considered invalid (even if its comparison count has exceeded the threshold) and GPFS prefers to read this disk in the next selection to update its latest speed evaluation.

Valid values are ≥ 1 in seconds. The default is 600 (10 minutes).

fastestPolicyMinDiffPercent

A percentage value indicating how GPFS selects the fastest between two disks. For example, if you use the default **fastestPolicyMinDiffPercent** value of 50, GPFS selects a disk as faster only if it is 50% faster than the other. Otherwise, the disks remain in the existing read order.

Valid values are 0 - 100 in percentage points. The default is 50.

fastestPolicyNumReadSamples

Controls how many read samples are taken to evaluate the disk's recent speed.

Valid values are 3 - 100. The default is 5.

fileHeatLossPercent

Specifies the reduction rate of **FILE_HEAT** value for every **fileHeatPeriodMinutes** of file inactivity. The default value is 10.

fileHeatPeriodMinutes

Specifies the inactivity time before a file starts to lose **FILE_HEAT** value. The default value is 0, which means that **FILE_HEAT** is not tracked.

FIPS1402mode

Controls whether GPFS operates in FIPS 140-2 mode, which requires using a FIPS-compliant encryption module for all encryption and decryption activity. Valid values are **yes** or **no**. The default value is **no**.

For FIPS 140-2 considerations, see *Encryption in IBM Spectrum Scale: Administration Guide*.

forceLogWriteOnFdatasync

Controls forcing log writes to disk. Valid values are **yes** or **no**.

When **forceLogWriteOnFdatasync** is set to **yes**, the GPFS log record is flushed to disk every time **fdatasync()** is invoked. This is the default.

When **forceLogWriteOnFdatasync** is set to **no**, the GPFS log record is flushed only when a new block is written to the file.

ignorePrefetchLUNCount

The GPFS client node calculates the number of sequential access prefetch and write-behind threads to run concurrently for each file system by using the count of the number of LUNs in the

file system and the value of **maxMBpS**. However, if the LUNs being used are composed of multiple physical disks, this calculation can underestimate the amount of IO that can be done concurrently.

Setting the value of the **ignorePrefetchLUNCount** parameter to **yes** does not include the LUN count and uses the **maxMBpS** value to dynamically determine the number of threads to schedule the **prefetchThreads** value.

This parameter impacts only the GPFS client node. The GPFS NSD server does not include this parameter in the calculation.

The valid values for this parameter are **yes** and **no**. The default value is **no** and can be used in traditional LUNs where one LUN maps to a single disk or an n+mP array. Set the value of this parameter to **yes** when the LUNs presented to GPFS are made up of a large numbers of physical disks.

The **-N** flag is valid for this attribute.

IrocData

Controls whether user data is populated into the local read-only cache. Other configuration options can be used to select the data that is eligible for the local read-only cache. When using more than one such configuration option, data that matches any of the specified criteria is eligible to be saved.

Valid values are **yes** or **no**. The default value is **yes**.

If **IrocData** is set to **yes**, by default the data that was not already in the cache when accessed by a user is subsequently saved to the local read-only cache. The default behavior can be overridden using the **IrocDataMaxFileSize** and **IrocDataStubFileSize** configuration options to save all data from small files or all data from the initial portion of large files.

IrocDataMaxFileSize

Limits the data that can be saved in the local read-only cache to only the data from small files.

A value of -1 indicates that all data is eligible to be saved. A value of 0 indicates that small files are not to be saved. A positive value indicates the maximum size of a file to be considered for the local read-only cache. For example, a value of 32768 indicates that files with 32 KB of data or less are eligible to be saved in the local read-only cache. The default value is 0.

IrocDataStubFileSize

Limits the data that can be saved in the local read-only cache to only the data from the first portion of all files.

A value of -1 indicates that all file data is eligible to be saved. A value of 0 indicates that stub data is not eligible to be saved. A positive value indicates that the initial portion of each file that is eligible is to be saved. For example, a value of 32768 indicates that the first 32 KB of data from each file is eligible to be saved in the local read-only cache. The default value is 0.

IrocDirectories

Controls whether directory blocks is populated into the local read-only cache. The option also controls other file system metadata such as indirect blocks, symbolic links, and extended attribute overflow blocks.

Valid values are **yes** or **no**. The default value is **yes**.

IrocInodes

Controls whether inodes from open files is populated into the local read-only cache; the cache contains the full inode, including all disk pointers, extended attributes, and data.

Valid values are **yes** or **no**. The default value is **yes**.

maxblocksize

Changes the maximum file system block size. Valid block sizes are 64 KiB, 128 KiB, 256 KiB, 512

mmchconfig

- | KiB, 1 MiB, 2 MiB, 4 MiB, 8 MiB, and 16 MiB. The default maximum block size is 1 MiB. Specify this value with the character **K** or **M**; for example, use **2M** to specify a block size of 2 MiB.
- | File systems with block sizes larger than the specified value cannot be created or mounted unless the block size is increased.
- | GPFS must be down on all the nodes in the cluster when changing the **maxblocksize** attribute.
- | The **-N** flag is valid for this attribute.

maxBufferDescs

Valid values are from 512 to 10,000,000.

Without explicit setting, it is set to a value of $10 * \text{maxFilesToCache}$ up to **pagepool** size/16 KB. Each buffer descriptor caches maximum block size data for a file. When caching small files, it does not need to be more than a small multiple of **maxFilesToCache** since only OpenFile objects can cache data blocks. When an application needs to cache large files, **maxBufferDescs** can be tuned to ensure that there are enough to cache large files.

For example, if you have 10,000 buffer descriptors that are configured and a 1MiB file system block size, you do not have enough buffer descriptors to cache a 20 GiB file. To cache a 20 GiB file, increase **maxBufferDescs** to at least 20,480 (20 GiB/1MiB=20,480).

The **-N** flag is valid for this attribute.

maxDownDisksForRecovery

Specifies the maximum number of disks that might experience a failure and still be subject to an automatic recovery attempt. If this value is exceeded, no automatic recovery actions take place.

Valid values are in the range 0 - 300. The default is 16. If **restripeOnDiskFailure** is **no**, **maxDownDisksForRecovery** has no effect.

maxFailedNodesForRecovery

Specifies the maximum number of nodes that might be unavailable before automatic disk recovery actions are canceled.

Valid values are in the range 0 - 300. The default is 3. If **restripeOnDiskFailure** is **no**, **maxFailedNodesForRecovery** has no effect.

maxFcntlRangesPerFile

Specifies the number of **fctl** locks that are allowed per file. The default is 200. The minimum value is 10 and the maximum value is 200000.

maxFilesToCache

Specifies the number of inodes to cache for recently used files that have been closed.

Storing the inode of a file in cache permits faster re-access to the file. The default is 4000, but increasing this number might improve throughput for workloads with high file reuse. However, increasing this number excessively might cause paging at the file system manager node. The value should be large enough to handle the number of concurrently open files plus allow caching of recently used files.

The **-N** flag is valid for this attribute.

maxMBps

Specifies an estimate of how many megabytes of data can be transferred per second into or out of a single node. The default is 2048 MiB per second. The value is used in calculating the amount of I/O that can be done to effectively prefetch data for readers and write-behind data from writers. By lowering this value, you can artificially limit how much I/O one node can put on all of the disk servers.

The **-N** flag is valid for this attribute.

maxMissedPingTimeout

See the **minMissedPingTimeout** parameter.

maxStatCache

Specifies the number of inodes to keep in the stat cache. The stat cache maintains only enough inode information to perform a query on the file system. The valid range for **maxStatCache** is 0 - 100,000,000. If you do not specify values for **maxFilesToCache** and **maxStatCache**, the default value of **maxFilesToCache** is 4000 and the default value of **maxStatCache** is 1000. If you specify a value for **maxFilesToCache** but not for **maxStatCache**, the default value of **maxStatCache** is $4 * \text{maxFilesToCache}$.

The -N flag is valid for this attribute.

Note: The stat cache is not effective on the Linux platform unless the Local Read-Only Cache (LROC) is configured. Therefore on Linux, the setting for **maxStatCache** depends on whether the Local Read-Only Cache (LROC) is enabled on the node:

- If LROC is not enabled on the node, set **maxStatCache** to 0.
- If LROC is enabled on the node, set **maxStatCache** to an appropriate nonzero value based on the number of nodes in the cluster, the number of token managers in the cluster, the size of the LROC device, and any other relevant factors. For more information about LROC, see the topic *Local read-only cache* in the *IBM Spectrum Scale: Administration Guide*.

This restriction applies to all the versions and distributions of Linux that IBM Spectrum Scale supports.

metadataDiskWaitTimeForRecovery

Specifies a period, in seconds, during which the recovery of metadata disks is suspended to give the disk subsystem a chance to correct itself. This parameter is taken into account when the affected disks belong to a single failure group. If more than one failure group is affected, the delay is based on the value of **minDiskWaitTimeForRecovery**.

Valid values are 0 - 3600 seconds. The default is 2400. If **restripeOnDiskFailure** is **no**, **metadataDiskWaitTimeForRecovery** has no effect.

minDiskWaitTimeForRecovery

Specifies a period, in seconds, during which the recovery of disks is suspended to give the disk subsystem a chance to correct itself. This parameter is taken into account when more than one failure group is affected. If the affected disks belong to a single failure group, the delay is based on the values of **dataDiskWaitTimeForRecovery** and **metadataDiskWaitTimeForRecovery**.

Valid values are 0 - 3600 seconds. The default is 1800. If **restripeOnDiskFailure** is **no**, **minDiskWaitTimeForRecovery** has no effect.

minMissedPingTimeout

The **minMissedPingTimeout** and **maxMissedPingTimeout** parameters set limits on the calculation of **missedPingTimeout** (MPT). The MPT is the allowable time for pings sent from the Cluster Manager (CM) to a node that has not renewed its lease to fail. The default MPT value is 5 seconds less than **leaseRecoveryWait**. The CM will wait the MPT seconds after the lease has expired before declaring a node out of the cluster. The values of the **minMissedPingTimeout** and **maxMissedPingTimeout** are in seconds; the default values are 3 and 60 respectively. If these values are changed, only GPFS on the quorum nodes that elect the CM must be recycled to take effect.

This parameter can be used to cover over a central network switch failure timeout or other network glitches that might be longer than **leaseRecoveryWait**. This might prevent false node down conditions, but it extends the time for node recovery to finish and might block other nodes from progressing if the failing node holds the tokens for many shared files.

As is the case with **leaseRecoveryWait**, a node is usually expelled from the cluster if there is a problem with the network or the node runs out of resources like paging. For example, if there is

mmchconfig

an application that is running on a node that is paging the machine too much or overrunning network capacity, GPFS might not have the chance to contact the Cluster Manager node to renew the lease within the timeout period.

The default value of this parameter is 3. A valid value is any number in the range 1 - 300.

mmapRangeLock

Specifies POSIX or non-POSIX **mmap** byte-range semantics. Valid values are **yes** or **no** (**yes** is the default). A value of **yes** indicates POSIX byte-range semantics apply to **mmap** operations. A value of **no** indicates non-POSIX **mmap** byte-range semantics apply to **mmap** operations.

If using InterProcedural Analysis (IPA), turn off this option:

```
mmchconfig mmapRangeLock=no -i
```

This allows more lenient intranode locking, but imposes internode whole file range tokens on files using **mmap** while writing.

mmfsLogTimeStampISO8601

Controls the time stamp format for GPFS log entries. Specify **yes** to use the ISO 8601 time stamp format for log entries or **no** to use the earlier time stamp format. The default value is **yes**. You can specify the log time stamp format for the entire cluster or for individual nodes. You can have different log time stamp formats on different nodes of the cluster. For more information, see the *Time stamp in GPFS log entries* topic in *IBM Spectrum Scale: Problem Determination Guide*.

The **-N** flag is valid for this attribute. This attribute takes effect immediately, whether or not **-i** is specified.

nfsPrefetchStrategy

With the **nfsPrefetchStrategy** parameter, GPFS optimizes prefetching for NFS file-style access patterns. This parameter defines a window of the number of blocks around the current position that are treated as fuzzy-sequential access. The value of this parameter can improve the performance while reading large files sequentially. However, because of kernel scheduling, some read requests that come to GPFS are not sequential. If the file system block size is smaller than the read request sizes, increasing the value of this parameter provides a bigger window of blocks. The default value is 0. A valid value is any number in the range 0 - 10.

Setting the value of **nfsPrefetchStrategy** to 1 or greater can improve the sequential read performance when large files are accessed by using NFS and the filesystem block size is smaller than the NFS transfer block size.

nistCompliance

Controls whether GPFS operates in the NIST 800-131A mode. (This applies to security transport only, not to encryption, as encryption always uses NIST-compliant mechanisms.)

Valid values are:

off

Specifies that there is no compliance to NIST standards. For clusters that are operating below the GPFS 4.1 level, this is the default.

SP800-131A

Specifies that security transport is to follow the NIST SP800-131A recommendations. For clusters at the GPFS 4.1 level or higher, this is the default.

Note: In a remote cluster setup, all clusters must have the same **nistCompliance** value.

noSpaceEventInterval

Specifies the time interval between calling a callback script of two **noDiskSpace** events of a file system. The default value is 120 seconds. If this value is set to zero, the **noDiskSpace** event is generated every time the file system encounters the **noDiskSpace** event. The **noDiskSpace** event is generated when a callback script is registered for this event with the **mmaddcallback** command.

nsdBufSpace

This option specifies the percentage of the page pool that is reserved for the network transfer of NSD requests. Valid values are within the range of 10 to 70. The default value is 30. On IBM Spectrum Scale RAID recovery group NSD servers, this value should be decreased to its minimum of 10, since vdisk-based NSDs are served directly from the RAID buffer pool (as governed by **nsdRAIDBufferPoolSizePct**). On all other NSD servers, increasing either this value or the amount of page pool, or both, could improve NSD server performance. On NSD client-only nodes, this parameter is ignored. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

The -N flag is valid for this attribute.

nsdInlineWriteMax

The **nsdInlineWriteMax** parameter specifies the maximum transaction size that can be sent as embedded data in an NSD-write RPC.

In most cases, the NSD-write RPC exchange performs the following steps:

1. An RPC is sent from the client to the server to request a write.
2. A GetData RPC is sent back from the server to the client to request the data.

Note: For data smaller than **nsdInlineWriteMax**, GPFS sends that amount of write data directly without the GetData RPC from the server to the client.

The default value of this parameter is 1024. A valid value is any number in the range 0 - 8M.

nsdMaxWorkerThreads

The **nsdMaxWorkerThreads** parameter sets the maximum number of NSD threads that can be involved in NSD I/O operations on an NSD server to the storage system to which the server is connected. On 64-bit architectures, the maximum value of the **workerThreads**, **prefetchThreads**, and **nsdMaxWorkerThreads** is less than 8192. The minimum value of **nsdMaxWorkerThreads** is 8 and the default value is 512. The default value works for most of the use cases.

This value should be scaled with the number of NSDs the server is connected with and not NSD clients in the cluster. When the NSDs to which the server is doing I/O are constructed from high performance storage device, such as IBM FlashSystem® or ESS, increased value of **nsdMaxWorkerThreads** can obtain better performance.

nsdMinWorkerThreads

The **nsdMinWorkerThreads** parameter sets a lower bound on number of active NSD I/O threads on an NSD server node that executes I/O operations against NSDs. The value represents the minimum number of NSD I/O threads required to execute the I/O operations.

The default value of this parameter is 16 and the minimum value is 1. For limits on setting the number of NSD threads, see the description of **nsdMaxWorkerThreads**.

nsdMultiQueue

The **nsdMultiQueue** parameter sets the number of queues. The default value of this parameter is 256. A valid value is any number in the range 2 - 512.

nsdRAIDTracks

This option specifies the number of tracks in the IBM Spectrum Scale RAID buffer pool, or 0 if this node does not have a IBM Spectrum Scale RAID vdisk buffer pool. This controls whether IBM Spectrum Scale RAID services are configured. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

Valid values are: 0; 256 or greater.

The -N flag is valid for this attribute.

nsdRAIDBufferPoolSizePct

This option specifies the percentage of the page pool that is used for the IBM Spectrum Scale RAID vdisk buffer pool. Valid values are within the range of 10 to 90. The default is 50 when

mmchconfig

IBM Spectrum Scale RAID is configured on the node in question; 0 when it is not. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

The -N flag is valid for this attribute.

nsdServerWaitTimeForMount

When mounting a file system whose disks depend on NSD servers, this option specifies the number of seconds to wait for those servers to come up. The decision to wait is controlled by the criteria managed by the **nsdServerWaitTimeWindowOnMount** option.

Valid values are 0 - 1200 seconds. The default is 300. A value of zero indicates that no waiting is done. The interval for checking is 10 seconds. If **nsdServerWaitTimeForMount** is 0, **nsdServerWaitTimeWindowOnMount** has no effect.

The mount thread waits when the daemon delays for safe recovery. The mount wait for NSD servers to come up, which is covered by this option, occurs after expiration of the recovery wait allows the mount thread to proceed.

The -N flag is valid for this attribute.

nsdServerWaitTimeWindowOnMount

Specifies a window of time (in seconds) during which a mount can wait for NSD servers as described for the **nsdServerWaitTimeForMount** option. The window begins when quorum is established (at cluster startup or subsequently), or at the last known failure times of the NSD servers required to perform the mount.

Valid values are 1 - 1200 seconds. The default is 600. If **nsdServerWaitTimeForMount** is 0, **nsdServerWaitTimeWindowOnMount** has no effect.

The -N flag is valid for this attribute.

When a node rejoins the cluster after being removed for any reason, the node resets all the failure time values that it knows about. Therefore, when a node rejoins the cluster it believes that the NSD servers have not failed. From the perspective of a node, old failures are no longer relevant.

GPFS checks the cluster formation criteria first. If that check falls outside the window, GPFS then checks for NSD server fail times being within the window.

numaMemoryInterleave

In a Linux NUMA environment, the default memory policy is to allocate memory from the local NUMA node of the CPU from which the allocation request was made. This parameter is used to change to an interleave memory policy for GPFS by starting GPFS with **numactl --interleave=all**.

Valid values are **yes** and **no**. The default is **no**.

If you run IBM Spectrum Scale on a multi-processor system, you should set this variable to **yes** and restart GPFS.

Before using this parameter, ensure that the Linux **numactl** package has been installed.

pagepool

Changes the size of the cache on each node. The default value is either one-third of the physical memory on the node or 1 GiB, whichever is smaller. This applies to new installations only; on upgrades the existing default value is kept.

The maximum GPFS page pool size depends on the value of the **pagepoolMaxPhysMemPct** parameter and the amount of physical memory on the node. You can specify this value with the suffix **K**, **M**, or **G**, for example, **128M**.

The -N flag is valid for this attribute.

pagepoolMaxPhysMemPct

Percentage of physical memory that can be assigned to the page pool. Valid values are 10 - 90 percent. The default is 75 percent (with the exception of Windows, where the default is 50 percent).

The **-N** flag is valid for this attribute.

pitWorkerThreadsPerNode

Controls the maximum number of threads to be involved in parallel processing on each node that is serving as a Parallel Inode Traversal (PIT) worker.

By default, when a command that uses the PIT engine is run, the file system manager asks all nodes in the local cluster to serve as PIT workers; however, you can specify an exact set of nodes to serve as PIT workers by using the **-N** option of a PIT command. Note that the current file system manager node is a mandatory participant, even if it is not in the list of nodes you specify. On each participating node, up to **pitWorkerThreadsPerNode** can be involved in parallel processing. The range of accepted values is 0 to 8192. The default value varies within the 2-16 range, depending on the file system configuration.

prefetchPct

GPFS uses the **prefetchPct** parameter as a guideline to limit the page pool space that is to be used for prefetch and write-behind buffers for active sequential streams. The default value of the **prefetchPct** parameter is 20% of the **pagepool** value. If the workload is sequential with very little caching of small files or random IO, increase the value of this parameter to 60% of the **pagepool** value, so that each stream can have more buffers that are cached for prefetch and write-behind operations.

The default value of this parameter is 20. The valid value can be any numbers in the range 0 - 60.

prefetchThreads

Controls the maximum number of threads that are dedicated to prefetching data for files that are read sequentially, or to handle sequential write-behind.

Functions in the GPFS daemon dynamically determine the actual degree of parallelism for prefetching data. The default value is 72. The minimum value is 2. The maximum value of **prefetchThreads** plus **worker1Threads** plus **nsdMaxWorkerThreads** is 8192 on all 64-bit platforms.

The **-N** flag is valid for this attribute.

profile

Specifies a predefined profile of attributes to be applied. System-defined profiles are located in `/usr/lpp/mmfss/profiles/`. All the configuration attributes listed under a cluster stanza are changed as a result of this command. The following system-defined profile names are accepted:

- **gpfsProtocolDefaults**
- **gpfsProtocolRandomIO**

A user's profiles must be installed in `/var/mmfs/etc/`. The profile file specifies GPFS configuration parameters with values different than the documented defaults. A user-defined profile must not begin with the string 'gpfs' and must have the .profile suffix.

User-defined profiles consist of the following stanzas:

```
%cluster:  
[CommaSeparatedNodesOrNodeClasses:]ClusterConfigurationAttribute=Value  
...
```

File system attributes and values are ignored.

A sample file can be found in `/usr/lpp/mmfss/samples/sample.profile`. See the **mmchconfig** command for a detailed description of the different configuration parameters. User-defined profiles should be used only by experienced administrators. When in doubt, use the **mmchconfig** command instead.

readReplicaPolicy

Specifies the location from which the disk is to read replicas. By default, GPFS reads the first replica whether there is a replica on the local disk or not. When **readReplicaPolicy=local** is specified, the policy reads replicas from the local disk if the local disk has data; for performance

mmchconfig

considerations, this is the recommended setting for FPO environments. When **readReplicaPolicy=fastest** is specified, the policy reads replicas from the disk considered the fastest based on the read I/O statistics of the disk. You can tune the way the system determines the fastest policy using the following parameters:

- **fastestPolicyNumReadSamples**
- **fastestPolicyCmpThreshold**
- **fastestPolicyMaxValidPeriod**
- **fastestPolicyMinDiffPercent**

In a system with SSD and regular disks, the value of **fastestPolicyCmpThreshold** can be set to a greater number to let GPFS refresh the speed statistics for the slower disks less frequently. The default value is maintained for all other configurations. The default value of this parameter is default. The valid values are default, local, and fastest.

To return this attribute to the default setting, specify **readReplicaPolicy=DEFAULT -i**.

release=LATEST

Changes the IBM Spectrum Scale configuration information to the latest format that is supported by the currently installed level of the product. Perform this operation after you have migrated all the nodes in the cluster to the latest level of the product. For more information, see the topic *Completing the migration to a new level of IBM Spectrum Scale* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

The command tries to access each node in the cluster to verify the level of the installed code. If the command cannot reach one or more nodes, you must rerun the command until it verifies the information for all the nodes.

The command fails with an error message if the **cipherList** configuration attribute of the cluster is not set to **AUTHONLY** or higher. For more information, see the topic *Completing the migration to a new level of IBM Spectrum Scale* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

restripeOnDiskFailure

Specifies whether GPFS will attempt to automatically recover from certain common disk failure situations.

When a disk experiences a failure and becomes unavailable, the recovery procedure first attempts to restart the disk and if this fails, the disk is suspended and its data that is moved to other disks. Similarly, when a node joins the cluster, all disks for which the node is responsible are checked and an attempt is made to restart any that are in a down state.

Whether a file system is a subject of a recovery attempt is determined by the max replication values for the file system. If the **mmlsfs -M** or **-R** value is greater than one, then the recovery code is executed. The recovery actions are asynchronous and GPFS continues its processing while the recovery attempts take place. The results from the recovery actions and any errors that are encountered are recorded in the `/var/adm/ras/autorecovery.log.<timestamp>` log.

For more information on GPFS disk fail auto recovery, see Big Data best practices in the IBM Spectrum Scale wiki in developerWorks®.

rpcPerfNumberDayIntervals

Controls the number of days that aggregated RPC data is saved. Every day the previous 24 hours of one-hour RPC data is aggregated into a one-day interval.

The default value for **rpcPerfNumberDayIntervals** is 30, which allows the previous 30 days of one-day intervals to be displayed. To conserve memory, fewer intervals can be configured to reduce the number of recent one-day intervals that can be displayed. The values that are allowed for **rpcPerfNumberDayIntervals** are in the range 4 - 60.

rpcPerfNumberHourIntervals

Controls the number of hours that aggregated RPC data is saved. Every hour the previous 60 minutes of 1-minute RPC data is aggregated into a one-hour interval.

The default value for **rpcPerfNumberHourIntervals** is 24, which allows the previous day's worth of one-hour intervals to be displayed. To conserve memory, fewer intervals can be configured to reduce the number of recent one-hour intervals that can be displayed. The values that are allowed for **rpcPerfNumberHourIntervals** are 4, 6, 8, 12, or 24.

rpcPerfNumberMinuteIntervals

Controls the number of minutes that aggregated RPC data is saved. Every minute the previous 60 seconds of 1-second RPC data is aggregated into a 1-minute interval.

The default value for **rpcPerfNumberMinuteIntervals** is 60, which allows the previous hour's worth of 1-minute intervals to be displayed. To conserve memory, fewer intervals can be configured to reduce the number of recent 1-minute intervals that can be displayed. The values that are allowed for **rpcPerfNumberMinuteIntervals** are 4, 5, 6, 10, 12, 15, 20, 30, or 60.

rpcPerfNumberSecondIntervals

Controls the number of seconds that aggregated RPC data is saved. Every second RPC data is aggregated into a 1-second interval.

The default value for **rpcPerfNumberSecondIntervals** is 60, which allows the previous minute's worth of 1-second intervals to be displayed. To conserve memory, fewer intervals can be configured to reduce the number of recent 1-second intervals that can be displayed. The values that are allowed for **rpcPerfNumberSecondIntervals** are 4, 5, 6, 10, 12, 15, 20, 30, or 60.

rpcPerfRawExecBufferSize

Specifies the number of bytes to allocate for the buffer that is used to store raw RPC execution statistics. For each RPC received by a node, 16 bytes of associated data is saved in this buffer when the RPC completes. This circular buffer must be large enough to hold 1 second's worth of raw execution statistics.

The default value for **rpcPerfRawExecBufferSize** is 2 MiB, which produces 131072 entries. The data in this buffer is processed every second. It is a good idea to set the buffer size 10% to 20% larger than what is needed to hold 1 second's worth of data.

rpcPerfRawStatBufferSize

Specifies the number of bytes to allocate for the buffer that is used to store raw RPC performance statistics. For each RPC sent to another node, 56 bytes of associated data is saved in this buffer when the reply is received. This circular buffer must be large enough to hold 1 second's worth of raw performance statistics.

The default value for **rpcPerfRawStatBufferSize** is 6 MiB, which produces 112347 entries. The data in this buffer is processed every second. It is a good idea to set the buffer size 10% to 20% larger than what is needed to hold one second's worth of data.

seqDiscardThreshold

With the **seqDiscardThreshold** parameter, GPFS detects a sequential read or write access pattern and specifies what has to be done with the page pool buffer after it is consumed or flushed by write-behind threads. This is the highest performing option in a case where a very large file is read or written sequentially. The default for this value is 1 MiB, which means that if a file is sequentially read and is greater than 1 MiB, GPFS does not keep the data in cache after consumption. There are some instances where large files are reread by multiple processes such as data analytics. In some cases, you can improve the performance of these applications by increasing the value of the **seqDiscardThreshold** parameter so that it is larger than the sets of files that have to be cached. If the value of the **seqDiscardthreshold** parameter is increased, GPFS attempts to keep as much data in cache as possible for the files that are below the threshold.

The value of **seqDiscardThreshold** is file size in bytes. The default is 1/ MB. Increase this value if you want to cache files that are sequentially read or written and are larger than 1 MiB in size. Ensure that there are enough buffer descriptors to cache the file data. For more information about buffer descriptors, see the **maxBufferDescs** parameter.

mmchconfig

sidAutoMapRangeLength

Controls the length of the reserved range for Windows SID to UNIX ID mapping. See *Identity management on Windows* in the *IBM Spectrum Scale: Administration Guide* for additional information.

sidAutoMapRangeStart

Specifies the start of the reserved range for Windows SID to UNIX ID mapping. See *Identity management on Windows* in the *IBM Spectrum Scale: Administration Guide* for additional information.

subnets

Specifies subnets that are used to communicate between nodes in a GPFS cluster or a remote GPFS cluster.

The subnets option must use the following format:

`subnets="Subnet[/ClusterName[;ClusterName...]][Subnet[/ClusterName[;ClusterName...]]...]"`

where:

Subnet

Is a subnet specification such as 192.168.2.0.

ClusterName

Can be either a cluster name or a shell-style regular expression, which is used to match cluster names, such as:

CL[23].kgn.ibm.com

Matches CL2.kgn.ibm.com and CL3.kgn.ibm.com.

CL[0-7].kgn.ibm.com

Matches CL0.kgn.ibm.com, CL1.kgn.ibm.com, ... CL7.kgn.ibm.com.

CL*.ibm.com

Matches any cluster name that starts with CL and ends with .ibm.com.

CL?.kgn.ibm.com

Matches any cluster name that starts with CL, is followed by any one character, and then ends with .kgn.ibm.com.

The order in which you specify the subnets determines the order in which GPFS uses these subnets to establish connections to the nodes within the cluster. GPFS follows the network settings of the operating system for a specified subnet address, including the network mask. For example, if you specify `subnets="192.168.2.0"` and a 23-bit mask is configured, then the subnet spans IP addresses 192.168.2.0 - 192.168.3.255. In contrast, with a 25-bit mask, the subnet spans IP addresses 192.168.2.0 - 192.168.2.127.

GPFS does not impose limits on the number of bits in the subnet mask.

This feature cannot be used to establish fault tolerance or automatic failover. If the interface corresponding to an IP address in the list is down, GPFS does not use the next one on the list.

When you use subnets, both the interface corresponding to the daemon address and the interface that matches the subnet settings must be operational.

For more information about subnets, see *Using remote access with public and private IP addresses* in the *IBM Spectrum Scale: Administration Guide*.

Specifying a cluster name or a cluster name pattern for each subnet is needed only when a private network is shared across clusters. If the use of a private network is confined within the local cluster, then you do not need to specify the cluster name in the subnet specification.

Limitation and fix: Although there is no upper limit to the number of subnets that can be specified in the **subnets** option, a limit does exist as to the number of subnets that are listed in

the **subnets** option that a given node can be a part of. That limit is seven for nodes that do not have a fix that increases the limit and 64 for nodes that do have the fix. For example, the 7-subnet limit precludes the effective use of more than seven network interfaces on a node if each interface belongs to a distinct subnet that is listed in the **subnets** option.

The fix that increases the limit to 64 is included as part of the following APARs: IJ06771 for Version 4.1.1, IJ06770 for 4.2.3, and IJ06762 for 5.0.1.

If a node exceeds the limit, then some or all of its network interfaces that belong to the subnets in the **subnets** option might not be used in communicating with other nodes, with the primary GPFS daemon interface being used instead.

syncBufsPerIteration

This parameter is used to expedite buffer flush and the rename operations that are done by MapReduce jobs.

The default value is 100. It should be set to 1 for the GPFS FPO cluster for Big Data applications. Keep it as the default value for all other cases.

syncSambaMetadataOps

Is used to enable and disable the syncing of metadata operations that are issued by the SMB server.

If set to **yes**, **fsync()** is used after each metadata operation to provide reasonable failover behavior on node failure. This ensures that the node taking over can see the metadata changes. Enabling **syncSambaMetadataOps** can affect performance due to more sync operations.

If set to **no**, the additional sync overhead is avoided at the potential risk of losing metadata updates after a failure.

systemLogLevel

Specifies the minimum severity level for messages that are sent to the system log. The severity levels from highest to lowest priority are: **alert**, **critical**, **error**, **warning**, **notice**, **configuration**, **informational**, **detail**, and **debug**. The value that is specified for this attribute can be any severity level, or the value **none** can be specified so no messages are sent to the system log. The default value is **notice**.

GPFS generates some critical log messages that are always sent to the system logging service. This attribute only affects messages originating in the GPFS daemon (**mmfsd**). Log messages originating in some administrative commands are only stored in the GPFS log file.

This attribute is only valid for Linux nodes.

tiebreakerDisks

Controls whether GPFS uses the node-quorum-with-tiebreaker algorithm in place of the regular node-based quorum algorithm. See the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*. To enable this feature, specify the names of 1 - 3 disks. Separate the NSD names with a semicolon (;) and enclose the list in quotes. The disks do not have to belong to any particular file system, but must be directly accessible from the quorum nodes. For example:

```
tiebreakerDisks="gpfs1nsd;gpfs2nsd;gpfs3nsd"
```

To disable this feature, use:

```
tiebreakerDisks=no
```

When you change the tiebreaker disks, be aware of the following requirements:

- In a CCR-based cluster, if the disks that are specified in the **tiebreakerDisks** parameter belong to any file system, the GPFS daemon does not need to be up on all of the nodes in the cluster. However, any file system that contains the disks must be available to run the **mmlsfs** command.

mmchconfig

- In a traditional server-based (non-CCR) configuration repository cluster, the GPFS daemon must be shut down on all the nodes of the cluster.

Note: When you add or delete a **tiebreakerCheck** event, IBM Spectrum Scale must be down on all the nodes of the cluster. For more information, see “**mmaddcallback command**” on page 10 and “**mmdelcallback command**” on page 286.

tscCmdPortRange=Min-Max

Specifies the range of port numbers to be used for extra TCP/IP ports that some administration commands need for their processing. Defining a port range makes it easier for you set firewall rules that allow incoming traffic on only those ports. For more information, see the topic *IBM Spectrum Scale port usage* in the *IBM Spectrum Scale: Administration Guide*.

If you used the **spectrumscale** installation toolkit to install a version of IBM Spectrum Scale that is earlier than version 5.0.0, then this attribute is initialized to 60000-61000. Otherwise, this attribute is initially undefined and the port numbers are dynamically assigned from the range of ephemeral ports that are provided by the operating system.

uidDomain

Specifies the UID domain name for the cluster.

GPFS must be down on all the nodes when changing the **uidDomain** attribute.

See the IBM white paper entitled *UID Mapping for GPFS in a Multi-cluster Environment* in IBM Knowledge Center([www.ibm.com/support/knowledgecenter/SSFKCN/com.ibm.cluster.gpfs.doc/gpfs_uid/uid_gpfs.html](http://www.ibm.com/support/knowledgecenter/SSFKCN/com.ibm.cluster.gpfs/doc/gpfs_uid/uid_gpfs.html)).

unmountOnDiskFail

Controls how the daemon responds when it detects a disk failure:

yes The local node force-unmounts the file system that contains the failed disk. Other file systems on the local node and all nodes in the cluster continue to function normally, if they can. The local node can remount the file system when the disk problem is resolved. Use this setting in the following cases:

- You are using SAN-attached disks in large multinode configurations and you are not using replication.
- You have a node that hosts **descOnly** disks. See *Establishing disaster recovery for your GPFS cluster* in the *IBM Spectrum Scale: Administration Guide*.

no The **no** option is the default value. The daemon marks the disk as failed, notifies all nodes that use this disk that it has failed, and continues as long as it can without using the disk. You can make the disk active again with the **mmchdisk** command. This setting is appropriate when the node is using metadata-and-data replication, because the cluster can work from the replica until the failed disk is active again. The default setting for **unmountOnDiskFail** causes file system panic/unmount if the number of failure groups with down disk is the same or more than the metadata replication.

meta This option is like **no** except that the file system remains mounted unless it cannot access any replica of the metadata.

The **-N** flag is valid for this attribute.

usePersistentReserve

Specifies whether to enable or disable Persistent Reserve (PR) on the disks. Valid values are **yes** or **no** (**no** is the default). GPFS must be stopped on all nodes when setting this attribute.

To enable PR and to obtain recovery performance improvements, your cluster requires a specific environment:

- All disks must be PR-capable.
- On AIX, all disks must be hdisk; on Linux, they must be generic (**/dev/sd***) or DM-MP (**/dev/dm-***) disks.

- If the disks have defined NSD servers, all NSD server nodes must be running the same operating system (AIX or Linux).
- If the disks are SAN-attached to all nodes, all nodes in the cluster must be running the same operating system (AIX or Linux).

For more information, see *Reduced recovery time using Persistent Reserve* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

verbsPorts

Specifies the InfiniBand device names and port numbers that are used for RDMA transfers between an NSD client and server. You must enable **verbsRdma** to enable **verbsPorts**.

The format for verbsPorts is:

```
verbsPorts="Device/Port/Fabric[ Device/Port/Fabric ...]"
```

In this format, *Device* is the HCA device name (such as `mthca0` or `m1x4_0`); *Port* is the one-based port number (such as 1 or 2); and *Fabric* is a value to identify different InfiniBand (IB) fabrics (IB subnets on different switches).

If you do not specify a port number, GPFS uses port 1 as the default. If the fabric number is not specified, the fabric number is 0.

For example:

```
verbsPorts="m1x4_0/1/7 m1x4_0/2/8"
```

creates two RDMA connections between the NSD client and server by using both ports of a dual ported adapter with a fabric identifier 7 on port 1 and fabric identifier 8 on port 2.

Another example, without the fabric number:

```
verbsPorts="mthca0/1 mthca0/2"
```

creates two RDMA connections between the NSD client and server using both ports of a dual ported adapter, with the fabric identifier defaulting to 0.

A third example, without port or fabric number:

```
verbsPorts="m1x4_0 m1x4_1"
```

uses port 1 on each HCA device for the RDMA connections.

The **-N** flag is valid for this attribute.

verbsRdma

Enables or disables InfiniBand RDMA using the Verbs API for data transfers between an NSD client and NSD server. Valid values are **enable** or **disable**.

The **-N** flag is valid for this attribute.

verbsRdmaCm

Enables or disables the RDMA Connection Manager (RDMA CM or RDMA_CM) using the RDMA_CM API for establishing connections between an NSD client and NSD server. Valid values are **enable** or **disable**. You must enable **verbsRdma** to enable **verbsRdmaCm**.

If RDMA CM is enabled for a node, the node is only able to establish RDMA connections using RDMA CM to other nodes with **verbsRdmaCm** enabled. RDMA CM enablement requires IPOIB (IP over InfiniBand) with an active IP address for each port. Although IPv6 must be enabled, the GPFS implementation of RDMA CM does not currently support IPv6 addresses, so an IPv4 address must be used.

If **verbsRdmaCm** is not enabled when **verbsRdma** is enabled, the older method of RDMA connection prevails.

The **-N** flag is valid for this attribute.

mmchconfig

verbsRdmaPkey

Specifies an InfiniBand partition key for a connection between the specified node and an Infiniband server that is included in an InfiniBand partition. This parameter is valid only if **verbsRdmaCm** is set to **disable**.

Only one partition key is supported per IBM Spectrum Scale cluster.

The **-N** flag is valid for this attribute.

verbsRdmaRoCEToS

Specifies the Type of Service (ToS) value for clusters using RDMA over Converged Ethernet (RoCE). Acceptable values for this parameter are 0, 8, 16, and 24. The default value is -1.

If the user-specified value is neither the default nor an acceptable value, the script exits with an error message to indicate that no change has been made. However, a RoCE cluster continues to operate with an internally set ToS value of 0 even if the **mmchconfig** command failed. Different ToS values can be set for different nodes or groups of nodes.

The **-N** flag is valid for this attribute.

The **verbsPorts** parameter can use IP netmask/subnet to specify network interfaces to use for RDMA CM. However, this format is allowed only when **verbsRdmaCm=yes**. Otherwise these entries are ignored. This allows the use of VLANs and multiple IP interfaces per IB device in general.

verbsRdmaSend

Enables or disables the use of InfiniBand RDMA rather than TCP for most GPFS daemon-to-daemon communication. When disabled, only data transfers between an NSD client and NSD server are eligible for RDMA. Valid values are **enable** or **disable**. The default value is **disable**. The **verbsRdma** option must be enabled for **verbsRdmaSend** to have any effect.

The **-N** flag is valid for this attribute.

verbsRdmasPerConnection

Sets the maximum number of simultaneous RDMA data transfer requests allowed per connection. The default value for **verbsRdmasPerConnection** is 16.

verbsRdmasPerNode

Sets the maximum number of simultaneous RDMA data transfer requests allowed per node. The default value for **verbsRdmasPerNode** is 1000. The value for **verbsRdmasPerNode** is limited by the value of the **nsdMaxWorkerThreads** setting used.

verbsSendBufferMemoryMB

Sets the amount of page pool memory (in MiB) to reserve as dedicated buffer space for use by the **verbsRdmaSend** feature. If the value is unreasonably small or large (for example, larger than **pagepool**), the actual memory that is used is adjusted to a more appropriate value. If the value is zero (the default), a value is calculated based on the maximum number of RDMA allowed per node (**verbsRdmasPerNode**). This option has no effect unless **verbsRdmaSend** is enabled.

workerThreads

Controls an integrated group of variables that tune file system performance. Use this variable to tune file systems in environments that are capable of high sequential or random read/write workloads or small-file activity. For new installations of the product, this variable is preferred over **worker1Threads** and **prefetchThreads**.

The default value is 48. If protocols are installed, then the default value is 512. The valid range is 1-8192. However, the maximum value of **workerThreads** plus **prefetchThreads** plus **nsdMaxWorkerThreads** is 8192. The **-N** flag is valid with this variable.

This variable controls both internal and external variables. The internal variables include maximum settings for concurrent file operations, for concurrent threads that flush dirty data and metadata, and for concurrent threads that prefetch data and metadata. You can further adjust the external variables with the **mmchconfig** command:

```
logBufferCount
prefetchThreads
worker3Threads
```

The **prefetchThreads** parameter is described in this help topic. See the Tuning Parameters article in the IBM Spectrum Scale wiki in developerWorks for descriptions of the **logBufferCount** and **worker3Threads** parameters.

Important: After you set **workerThreads** to a non-default value, avoid setting **worker1Threads**. If you do, at first only **worker1Threads** is changed. But when IBM Spectrum Scale is restarted, all corresponding variables are automatically tuned according to the value of **worker1Threads**, instead of **workerThreads**.

worker1Threads

For some categories of file I/O, this variable controls the maximum number of concurrent file I/O operations. You can increase this value to increase the I/O performance of the file system. However, increasing this variable beyond some point might begin to degrade file system performance.

Important: After you set **workerThreads** to a non-default value, avoid setting **worker1Threads**. If you do, at first only **worker1Threads** is changed. But when IBM Spectrum Scale is restarted, all corresponding variables are automatically tuned according to the value of **worker1Threads**, instead of **workerThreads**.

This attribute is primarily used for random read or write requests that cannot be pre-fetched, random I/O requests, or small file activity. The default value is 48. The minimum value is 1. The maximum value of **prefetchThreads** plus **worker1Threads** plus **nsdMaxWorkerThreads** is 8192 on all 64-bit platforms.

The **-N** flag is valid for this attribute.

writebehindThreshold

The **writebehindThreshold** parameter specifies the point at which GPFS starts flushing new data out of the page pool for a file that is being written sequentially. Until the file size reaches this threshold, no write-behind is started because the full blocks are filled.

Increasing this value defers write-behind for new larger files, which can be useful. The workload folder contains temporary files that are smaller than the value of **writebehindThreshold** and are deleted before they are flushed from cache. The default value of this parameter is 512 KiB. If the value is too large, there might be too many dirty buffers that the sync thread has to flush at the next sync interval, causing a surge in disk IO. Keeping the value small ensures a smooth flow of dirty data to disk.

Note: If you set new values for **afmParallelReadChunkSize**, **afmParallelReadThreshold**, **afmParallelWriteChunkSize**, and **afmParallelWriteThreshold**; you need not relink filesets for the new values to take effect.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmchconfig** command.

mmchconfig

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

Examples

To change the maximum file system block size that is allowed to 8 MiB, issue this command:

```
| mmchconfig maxblocksize=4M
```

The system displays information similar to the following:

```
Verifying GPFS is stopped on all nodes ...
mmchconfig: Command successfully completed
mmchconfig: Propagating the cluster configuration data to all
    affected nodes. This is an asynchronous process.
```

To confirm the change, issue the following command:

```
mmlsconfig
```

The system displays information similar to the following:

```
Configuration data for cluster ib.cluster:
-----
```

```
clusterName ib.cluster
clusterId 13882433899463047326
autoload no
| minReleaseLevel 4.1.1.0
| dmapiFileHandleSize 32
| maxblocksize 1M
pagepool 2g
[c21f1n18]
pagepool 5g
[common]
verbsPorts mthca0/1
verbsRdma enable
subnets 10.168.80.0
adminMode central
```

```
File systems in cluster ib.cluster:
-----
```

```
/dev/fs1
```

See also

- “[mmaddnode command](#)” on page 29
- “[mmchnode command](#)” on page 190
- “[mmcrcluster command](#)” on page 239
- “[mmdelnnode command](#)” on page 297
- “[mmlsconfig command](#)” on page 400
- “[mmlscluster command](#)” on page 397

Location

/usr/lpp/mmfs/bin

mmchdisk command

Changes state or parameters of one or more disks in a GPFS file system.

Synopsis

```
mmchdisk Device {resume | start} -a
  [-N {Node[,Node...] | NodeFile | NodeClass}]
  [--inode-criteria CriteriaFile]
  [-o InodeResultFile]
  [--qos QOSClass]
```

or

```
mmchdisk Device {suspend | empty | resume | stop | start | change}
  {-d "DiskDesc[;DiskDesc...]" | -F StanzaFile}
  [-N {Node[,Node...] | NodeFile | NodeClass}]
  [--inode-criteria CriteriaFile]
  [-o InodeResultFile]
  [--qos QOSClass]
```

Availability

Available on all IBM Spectrum Scale editions.

Description

Use the **mmchdisk** command to change the state or the parameters of one or more disks in a GPFS file system.

- | The state of a disk is a combination of its status and availability, displayed with the **mmlsdisk** command.
- | Disk status is normally either **ready**, **emptied**, **suspended**, or **to be emptied**. A transitional status such as **replacing**, **replacement**, or **being emptied** might appear if a disk is being deleted or replaced. An **emptied** disk indicates that there is not any file system data or metadata on the disk and that new data will not be put on the disk. A disk status of **emptied** means that the disk can be removed without needing to migrate data. A suspended or being emptied disk is one that the user has decided not to place any new data on. Existing data on a suspended or being emptied disk may still be read or updated.
- | Typically, a disk is suspended before it is removed from a file system. When a disk is suspended, the **mmrestripefs** command migrates the data off that disk. Disk availability is either **up** or **down**.

Be sure to use **stop** before you take a disk offline for maintenance. You should also use **stop** when a disk has become temporarily inaccessible due to a disk failure that is repairable without loss of data on that disk (for example, an adapter failure or a failure of the disk electronics).

The *Disk Usage* (**dataAndMetadata**, **dataOnly**, **metadataOnly**, or **descOnly**) and *Failure Group* parameters of a disk are adjusted with the **change** option. See the *Recoverability considerations* topic in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

- | The **mmchdisk change** command does not move data or metadata that resides on the disk. After you issue the **mmchdisk change** command for *DiskUsage* or *State*, you might have to issue the **mmrestripefs** command with the **-r** option to relocate data so that it conforms to the new disk parameters.

The **mmchdisk** command can be issued for a mounted or unmounted file system. When maintenance is complete or the failure has been repaired, use the **mmchdisk** command with the **start** option. If the failure cannot be repaired without loss of data, you can use the **mmdeldisk** command.

Note:

1. The **mmchdisk** command cannot be used to change the NSD servers associated with the disk. Use the **mmchnsd** command for this purpose.

mmchdisk

2. Similarly, the **mmchdisk** command cannot be used to change the storage pool for the disk. Use the **mmdeldisk** and **mmadddisk** commands to move a disk from one storage pool to another.

Prior to GPFS 3.5, the disk information for the **mmchdisk** change option was specified in the form of disk descriptors defined as follows (with the second, third, sixth and seventh fields reserved):

DiskName:::DiskUsage:FailureGroup:::

- | For backward compatibility, the **mmchdisk** command still accepts the traditional disk descriptor format, but its use is deprecated.

Parameters

Device

The device name of the file system to which the disks belong. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

suspend

or

empty

- | Instructs GPFS to stop allocating space on the specified disk. Put a disk in this state when you are preparing to remove the file system data from the disk or if you want to prevent new data from being put on the disk. This is a user-initiated state that GPFS never enters without an explicit command to change the disk state. Existing data on a suspended disk may still be read or updated.
- | A disk remains in a suspended or to be emptied state until it is explicitly resumed. Restarting GPFS or rebooting nodes does not restore normal access to a suspended disk.

resume

Informs GPFS that a disk previously suspended is now available for allocating new space. If the disk is currently in a stopped state, it remains stopped until you specify the **start** option. Otherwise, normal read and write access to the disk resumes.

stop

Instructs GPFS to stop any attempts to access the specified disks. Use this option to tell the file system manager that a disk has failed or is currently inaccessible because of maintenance.

A disk remains stopped until it is explicitly started by the **mmchdisk** command with the **start** option.

- | You cannot run **mmchdisk stop** on a file system with a default replication setting of 1.

start

Informs GPFS that disks previously stopped are now accessible. This is accomplished by first changing the disk availability from **down** to **recovering**. The file system metadata is then scanned and any missing updates (replicated data that was changed while the disk was **down**) are repaired. If this operation is successful, the availability is changed to **up**. If the metadata scan fails, availability is set to **unrecovered**. This situation can occur if too many disks in the file system are **down**. The metadata scan can be re-initiated at a later time by issuing the **mmchdisk start** command again.

If more than one disk in the file system is down, they must all be started at the same time by issuing the **mmchdisk Device start -a** command. If you start them separately and metadata is stored on any disk that remains down, the **mmchdisk start** command fails.

change

Instructs GPFS to change the disk usage parameter, the failure group parameter, or both, according to the values specified in the NSD stanzas.

-d "DiskDesc[;DiskDesc...]"

A descriptor for each disk to be changed.

Specify only disk names when using the **suspend**, **resume**, **stop**, or **start** options. Delimit multiple disk names with semicolons and enclose the list in quotation marks. For example, "gpfs1nsd;gpfs2nsd"

When using the **change** option, include the disk name and any new *Disk Usage* and *Failure Group* positional parameter values in the descriptor. Delimit descriptors with semicolons and enclose the list in quotation marks; for example, "gpfs1nsd::dataOnly;gpfs2nsd::metadataOnly:12".

The use of disk descriptors is discouraged.

-F StanzaFile

Specifies a file containing the NSD stanzas for the disks to be changed. NSD stanzas have this format:

```
%nsd:
  nsd=NsdName
  usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}
  failureGroup=FailureGroup
  pool=StoragePool
  servers=ServerList
  device=DiskName
```

where:

nsd=NsdName

The name of the NSD to change. For a list of disks that belong to a particular file system, issue the **mmlsnsd -f Device**, **mmlsfs Device -d**, or **mmlsdisk Device** command. The **mmlsdisk Device** command will also show the current disk usage and failure group values for each of the disks. This clause is mandatory for the **mmchdisk** command.

usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}

Specifies the type of data to be stored on the disk:

dataAndMetadata

Indicates that the disk contains both data and metadata. This is the default for disks in the system pool.

dataOnly

Indicates that the disk contains data and does not contain metadata. This is the default for disks in storage pools other than the system pool.

metadataOnly

Indicates that the disk contains metadata and does not contain data.

descOnly

Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster-recovery configurations. For more information, see the *IBM Spectrum Scale: Administration Guide* and search on "Synchronous mirroring utilizing GPFS replication"

This clause is meaningful only for the **mmchdisk change** option.

failureGroup=FailureGroup

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

mmchdisk

Disks that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

pool=StoragePool

Specifies the storage pool to which the disk is to be assigned. If this name is not provided, the default is **system**.

Only the system storage pool can contain **metadataOnly**, **dataAndMetadata**, or **descOnly** disks. Disks in other storage pools must be **dataOnly**.

servers=ServerList

A comma-separated list of NSD server nodes. This clause is ignored by the **mmadddisk** command.

device=DiskName

The block device name of the underlying disk device. This clause is ignored by the **mmadddisk** command.

-a Specifies to change the state of all of the disks belonging to the file system, *Device*. This operand is valid only on the **resume** and **start** options.

-N {Node[,Node...] | NodeFile | NodeClass }

Specifies a list of nodes that should be used for making the requested disk changes. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

--inode-criteria CriteriaFile

Specifies the interesting inode criteria flag, where *CriteriaFile* is one of the following:

BROKEN

Indicates that a file has a data block with all of its replicas on disks that have been removed.

Note: **BROKEN** is always included in the list of flags even if it is not specified.

dataUpdateMiss

Indicates that at least one data block was not updated successfully on all replicas.

exposed

Indicates an inode with an exposed risk; that is, the file has data where all replicas are on suspended disks. This could cause data to be lost if the suspended disks have failed or been removed.

i11Compressed

Indicates an inode in which file compression or decompression is deferred, or in which a compressed file is partly decompressed to allow the file to be written into or memory-mapped.

i11Placed

Indicates an inode with some data blocks that might be stored in an incorrect storage pool.

i11Replicated

Indicates that the file has a data block that does not meet the setting for the replica.

metaUpdateMiss

Indicates that there is at least one metadata block that has not been successfully updated to all replicas.

unbalanced

Indicates that the file has a data block that is not well balanced across all the disks in all failure groups.

Note: If a file matches *any* of the specified interesting flags, all of its interesting flags (even those not specified) will be displayed.

-o InodeResultFile

Contains a list of the inodes that met the interesting inode flags that were specified on the **--inode-criteria** parameter. The output file contains the following:

INODE_NUMBER

This is the inode number.

DISKADDR

Specifies a dummy address for later **tsfindinode** use.

SNAPSHOT_ID

This is the snapshot ID.

ISGLOBAL_SNAPSHOT

Indicates whether or not the inode is in a global snapshot. Files in the live file system are considered to be in a global snapshot.

INDEPENDENT_FSETID

Indicates the independent fileset to which the inode belongs.

MEMO (INODE_FLAGS FILE_TYPE [ERROR])

Indicates the inode flag and file type that will be printed:

Inode flags:

BROKEN
exposed
dataUpdateMiss
i11Compressed
i11Placed
i11Replicated
metaUpdateMiss
unbalanced

File types:

BLK_DEV
CHAR_DEV
DIRECTORY
FIFO
LINK
LOGFILE
REGULAR_FILE
RESERVED
SOCK
UNLINKED
DELETED

Notes:

1. An error message will be printed in the output file if an error is encountered when repairing the inode.
2. **DISKADDR**, **ISGLOBAL_SNAPSHOT**, and **FSET_ID** work with the **tsfindinode** tool (`/usr/lpp/mmf/bin/tsfindinode`) to find the file name for each inode. **tsfindinode** uses the output file to retrieve the file name for each interesting inode.

--qos QoSClass

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **other** QoS class. (Unlike other commands that have the **--qos** option, the **mmchdisk** command runs in the **other** class by default.) This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

mmchdisk

maintenance

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

other

This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

Exit status

0 Successful completion.

nonzero

A failure has occurred.

Security

You must have root authority to run the **mmchdisk** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

Examples

1. To suspend active disk **gpfs2nsd**, issue this command:

```
mmchdisk fs0 suspend -d gpfs2nsd
```

To confirm the change, issue this command:

```
mmlsdisk fs0
```

In IBM Spectrum Scale versions earlier than V4.1.1, the product displays information similar to the following example:

disk name	driver type	sector size	failure group	holds metadata	status data	storage availability	pool
gpfs2nsd	nsd	512	2	yes	suspended	up	system
hd3vsdn01	nsd	512	2	yes	ready	up	system
hd27n01	nsd	512	8	yes	ready	up	system
hd28n01	nsd	512	8	yes	ready	up	system
hd29n01	nsd	512	8	yes	ready	up	system
hd10vsdn09	nsd	512	4003	no	ready	up	sp1
hd11vsdn10	nsd	512	4003	no	ready	up	sp1

Note: In product versions earlier than V4.1.1, the **mmlsdisk** command lists the disk status as suspended. In product versions V4.1.1 and later, the **mmlsdisk** command lists the disk status as to be emptied with both **mmchdisk suspend** or **mmchdisk empty** commands.

2. To empty active disk **gpfs1nsd**, issue this command:

```
mmchdisk fs0 empty -d gpfs1nsd
```

To confirm the change, issue this command:

```
mmlsdisk fs0 -L
```

In product version V4.1.1 and later, the system displays information similar to the following example:

```

disk      driver sector failure holds   holds           storage
name      type    size   group metadata data  status     availability disk id pool
-----  -----  -----  -----  -----  -----  -----  -----
gpfs1nsd nsd    512    -1 Yes     Yes  to be emptied up          1 system
gpfs2nsd nsd    512    -1 Yes     Yes  to be emptied up          2 system
gpfs3nsd nsd    512    -1 Yes     Yes  ready      up          3 system
gpfs4nsd nsd    512    -1 Yes     Yes  ready      up          4 system
Number of quorum disks: 3
Read quorum value: 2
Write quorum value: 2
Attention: Due to an earlier configuration change the file system
may contain data that is at risk of being lost.

```

3. To specify that metadata should no longer be stored on disk **gpfs1nsd**, issue this command:
- ```
mmchdisk fs0 change -d "gpfs1nsd:::dataOnly"
```

To confirm the change, issue this command:

```
mmlsdisk fs0
```

The system displays information similar to:

```

disk driver sector failure holds holds status storage
name type size group metadata data status availability pool
----- ----- ----- ----- ----- ----- ----- -----
hd2vsdn01 nsd 512 2 yes yes ready up system
hd3vsdn01 nsd 512 2 yes yes ready up system
hd27n01 nsd 512 8 yes yes ready up system
gpfs1nsd nsd 512 8 no yes ready up system
hd29n01 nsd 512 8 yes yes ready up system
hd10vsdn09 nsd 512 4003 no yes ready up sp1
hd11vsdn10 nsd 512 4003 no yes ready up sp1

```

4. To start a disk and check for files matching the interesting inode criteria located on the disk, issue this command:

```
mmchdisk fs1 start -d vmip2_nsd3 /tmp/crit --inode-criteria
```

The system displays information similar to:

```

mmnsddiscover: Attempting to rediscover the disks. This may take a while ...
mmnsddiscover: Finished.
vmip2.gpfs.net: GPFS: 6027-1805 [N] Rediscovered nsd server access to vmip2_nsd3.
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
Scanning file system metadata for data storage pool
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
100.00 % complete on Wed Apr 15 10:20:37 2015 65792 inodes with total 398 MB data processed)
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-3312 No inode was found matching the criteria.

```

The disk was started successfully. No files matching the requested criteria were found.

## See also

- *Displaying GPFS disk states* in the *IBM Spectrum Scale: Administration Guide*.
- “[mmadddisk command](#)” on page 23
- “[mmchnsd command](#)” on page 198
- “[mmdeldisk command](#)” on page 287
- “[mmlsdisk command](#)” on page 402
- “[mmlsnsd command](#)” on page 423

## **mmchdisk**

- “mmrpldisk command” on page 539

## **Location**

/usr/lpp/mmfs/bin

---

## mmcheckquota command

Checks file system user, group and fileset quotas.

### Synopsis

```
mmcheckquota [-v] [-N {Node[,Node...] | NodeFile | NodeClass}]
[--qos QosClass] {-a | Device [Device ...]}
```

or

```
mmcheckquota {-u UserQuotaFile | -g GroupQuotaFile | -j FilesetQuotaFile}
[--qos QosClass] Device
```

or

```
mmcheckquota --backup backupDir Device
```

### Availability

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

### Description

The **mmcheckquota** command serves two purposes:

1. Count inode and space usage in a file system by user, group and fileset, and write the collected data into quota files.

**Note:** In cases where small files do not have an additional block allocated for them, quota usage may show less space usage than expected.

2. Replace either the user, group, or fileset quota files, for the file system designated by *Device*, thereby restoring the quota files for the file system. These files must be contained in the root directory of *Device*. If a backup copy does not exist, an empty file is created when the **mmcheckquota** command is issued.

The **mmcheckquota** command counts inode and space usage for a file system and writes the collected data into quota files. Indications leading you to the conclusion you should run the **mmcheckquota** command include:

- MMFS\_QUOTA error log entries. This error log entry is created when the quota manager has a problem reading or writing the quota file.
- Quota information is lost due to a node failure. A node failure could leave users unable to open files or deny them disk space that their quotas should allow.
- The in-doubt value is approaching the quota limit.

The sum of the in-doubt value and the current usage may not exceed the hard limit. Consequently, the actual block space and number of files available to the user or the group may be constrained by the in-doubt value. If the in-doubt value approaches a significant percentage of the quota, use the **mmcheckquota** command to account for the lost space and files.

- User, group, or fileset quota files are corrupted.

The **mmcheckquota** command is I/O-intensive and should be run when the system load is light. When issuing the **mmcheckquota** command on a mounted file system, negative in-doubt values may be reported if the quota server processes a combination of up-to-date and back-level information. This is a transient situation and can be ignored.

If a file system is ill-replicated, the **mmcheckquota** command will not be able to determine exactly how many valid replicas actually exist for some of the blocks. If this happens, the used block count results from **mmcheckquota** will not be accurate. It is recommended that you run **mmcheckquota** to restore

## **mmcheckquota**

accurate usage count after the file system is no longer ill-replicated.

### **Parameters**

**-a** Checks all GPFS file systems in the cluster from which the command is issued.

**--backup *BackupDirectory***

Specifies a backup directory, which must be in the same GPFS file system as the root directory of *Device*.

In IBM Spectrum Scale V4.1.1 and later, you can use this parameter to copy quota files. The command copies three quota files to the specified directory.

*Device*

Specifies the device name of the file system. File system names do not need to be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

**-g *GroupQuotaFileName***

Replaces the current group quota file with the file indicated.

When replacing quota files with the **-g** option, the quota file must be in the root directory of the GPFS file system.

**-j *FilesetQuotaFilename***

Replaces the current fileset quota file with the file indicated.

When replacing quota files with the **-j** option, the quota file must be in the root directory of the GPFS file system.

**-N {*Node*[,*Node...*] | *NodeFile* | *NodeClass*}**

Specifies the nodes that will participate in a parallel quota check of the system. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

**-u *UserQuotaFilename***

Replaces the current user quota file with the file indicated.

When replacing quota files with the **-u** option, the quota file must be in the root directory of the GPFS file system.

**--qos *QoSClass***

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **Options**

**-v** Reports discrepancies between calculated and recorded disk quotas.

## Exit status

**0** Successful completion.

### nonzero

A failure has occurred.

## Security

You must have root authority to run the **mmcheckquota** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmcheckquota** command is issued.

## Examples

1. To check quotas for file system **fs0**, issue this command:

```
mmcheckquota fs0
```

The system displays information only if a problem is found.

2. To check quotas for all file systems, issue this command:

```
mmcheckquota -a
```

The system displays information only if a problem is found or if quota management is not enabled for a file system:

```
fs2: no quota management installed
fs3: no quota management installed
```

3. To report discrepancies between calculated and recorded disk quotas, issue this command:

```
mmcheckquota -v fs1
```

The system displays information similar to:

```
fs1: Start quota check
1 % complete on Fri Apr 17 13:07:47 2009
6 % complete on Fri Apr 17 13:07:48 2009
11 % complete on Fri Apr 17 13:07:49 2009
17 % complete on Fri Apr 17 13:07:50 2009
22 % complete on Fri Apr 17 13:07:51 2009
28 % complete on Fri Apr 17 13:07:52 2009
33 % complete on Fri Apr 17 13:07:53 2009
38 % complete on Fri Apr 17 13:07:54 2009
44 % complete on Fri Apr 17 13:07:55 2009
49 % complete on Fri Apr 17 13:07:56 2009
55 % complete on Fri Apr 17 13:07:57 2009
61 % complete on Fri Apr 17 13:07:58 2009
66 % complete on Fri Apr 17 13:07:59 2009
72 % complete on Fri Apr 17 13:08:00 2009
78 % complete on Fri Apr 17 13:08:01 2009
83 % complete on Fri Apr 17 13:08:02 2009
89 % complete on Fri Apr 17 13:08:03 2009
94 % complete on Fri Apr 17 13:08:04 2009
Finished scanning the inodes for fs1.
Merging results from scan.
fs1: quota check found the following differences:
USR 0: 288400 subblocks counted (was 288466); 24 inodes counted (was 81)
USR 60011: 50 subblocks counted (was 33); 2 inodes counted (was 20)
USR 60012: 225 subblocks counted (was 223); 9 inodes counted (was 4)
USR 60013: 175 subblocks counted (was 146); 7 inodes counted (was 26)
```

## **mmcheckquota**

```
USR 60014: 200 subblocks counted (was 178); 8 inodes counted (was 22)
USR 60015: 275 subblocks counted (was 269); 11 inodes counted (was 0)
USR 60019: 0 subblocks counted (was 9); 0 inodes counted (was 5)
USR 60020: 0 subblocks counted (was 1); 0 inodes counted (was 3)
GRP 0: 28845098 subblocks counted (was 28844639); 14 inodes counted (was 91)
FILESET 0: 28849125 subblocks counted (was 28848717); 105 inodes counted (was 24)
```

### **See also**

- “[mmedquota command](#)” on page 325
- “[mmfsck command](#)” on page 331
- “[mmlsquota command](#)” on page 435
- “[mmquotaon command](#)” on page 509
- “[mmquotaooff command](#)” on page 507
- “[mmrepquota command](#)” on page 517

### **Location**

/usr/lpp/mmfs/bin

---

## mmchfileset command

Changes the attributes of a GPFS fileset.

### Synopsis

```
mmchfileset Device {FilesetName | -J JunctionPath}
[-j NewFilesetName] [-t NewComment] [-p afmAttribute=Value...]
[--allow-permission-change PermissionChangeMode]
[--inode-limit MaxNumInodes[:NumInodesToPreallocate]]
[--iam-mode Mode]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher.

### Description

The **mmchfileset** command changes attributes for an existing GPFS fileset.

For information on GPFS filesets, see the *IBM Spectrum Scale: Administration Guide*.

### Parameters

#### *Device*

The device name of the file system that contains the fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

#### *FilesetName*

Specifies the name of the fileset.

#### -J *JunctionPath*

Specifies the junction path name for the fileset.

A junction is a special directory entry that connects a name in a directory of one fileset to the root directory of another fileset.

#### -j *NewFilesetName*

Specifies the new name that is to be given to the fileset. This name must be less than 256 characters in length. The root fileset cannot be renamed.

#### -t *NewComment*

Specifies an optional comment that appears in the output of the **mmlsfileset** command. This comment must be less than 256 characters in length. This option cannot be used on the root fileset.

#### -p *afmAttribute=Value*

Specifies an AFM configuration attribute and its value. More than one **-p** option can be specified.

The following AFM configuration attributes are valid:

##### **afmAsyncDelay**

Specifies (in seconds) the amount of time by which write operations are delayed (because write operations are asynchronous with respect to remote clusters). For write-intensive applications that keep writing to the same set of files, this delay is helpful because it replaces multiple writes to the home cluster with a single write containing the latest data. However, setting a very high value weakens the consistency of data on the remote cluster.

This configuration parameter is applicable only for writer caches (Single writer, Independent writer, and Primary ) in which data from cache is pushed to home.

Valid values are between 1 and 2147483647. The default is 15.

## **mmchfileset**

### **afmDirLookupRefreshInterval**

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a directory, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of that directory has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 60. In situations where home cluster data changes frequently, a value of 0 is recommended.

### **afmDirOpenRefreshInterval**

Controls the frequency of data revalidations that are triggered by such I/O operations as **read** or **write** (specified in seconds). After a directory has been cached, **open** requests resulting from I/O operations on that object are directed to the cached directory until the specified amount of time has passed. Once the specified amount of time has passed, the **open** request gets directed to a gateway node rather than to the cached directory.

Valid values are between 0 and 2147483647. The default is 60. Setting a lower value guarantees a higher level of consistency.

### **afmEnableAutoEviction**

Enables eviction on a given fileset. A **yes** value specifies that eviction is allowed on the fileset. A **no** value specifies that eviction is not allowed on the fileset.

See also the topic about cache eviction in the *IBM Spectrum Scale: Administration Guide*.

### **afmExpirationTimeout**

Is used with **afmDisconnectTimeout** (which can be set only through **mmchconfig**) to control how long a network outage between the cache and home clusters can continue before the data in the cache is considered out of sync with home. After **afmDisconnectTimeout** expires, cached data remains available until **afmExpirationTimeout** expires, at which point the cached data is considered expired and cannot be read until a reconnect occurs.

Valid values are 0 through 2147483647. The default is disable.

### **afmFileLookupRefreshInterval**

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a file, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of the file has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 30. In situations where home cluster data changes frequently, a value of 0 is recommended.

### **afmFileOpenRefreshInterval**

Controls the frequency of data revalidations that are triggered by such I/O operations as **read** or **write** (specified in seconds). After a file has been cached, **open** requests resulting from I/O operations on that object are directed to the cached file until the specified amount of time has passed. Once the specified amount of time has passed, the **open** request gets directed to a gateway node rather than to the cached file.

Valid values are 0 through 2147483647. The default is 30. Setting a lower value guarantees a higher level of consistency.

### **afmMode**

Specifies the mode in which the cache operates. Valid values are the following:

#### **single-writer | sw**

Specifies single-writer mode.

#### **read-only | ro**

Specifies read-only mode. (For **mmcrfileset**, this is the default value.)

**local-updates | lu**

Specifies local-updates mode.

**independent-writer | iw**

Specifies independent-writer mode.

**Primary | drp**

Specifies the primary mode for AFM asynchronous data replication.

**Secondary | drs**

Specifies the secondary mode for AFM asynchronous data replication.

Changing from single-writer/read-only modes to read-only/local-updates/single-writer is supported. When changing from read-only to single-writer, the read-only cache is up-to-date. When changing from single-writer to read-only, all requests from cache should have been played at home. Changing from local-updates to read-only/local-updates/single-writer is restricted. A typical dataset is set up to include a single cache cluster in single-writer mode (which generates the data) and one or more cache clusters in local-updates or read-only mode. AFM single-writer/independent-writer filesets can be converted to primary. Primary/secondary filesets cannot be converted to AFM filesets.

In case of AFM asynchronous data replication, the **mmchfileset** command cannot be used to convert to primary from secondary. For detailed information, see *AFM-based Asynchronous Disaster Recovery (AFM DR)* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

For more information, see the topic about caching modes in the *IBM Spectrum Scale: Administration Guide* chapter about active file management.

**afmNumFlushThreads**

Defines the number of threads used on each gateway to synchronize updates to the home cluster. The default value is 4, which is sufficient for most installations. The current maximum value is 1024, which is too high for most installations; setting this parameter to such an extreme value should be avoided.

**afmNumReadThreads**

Defines the number of threads that can be used on each participating gateway node during parallel read. The default value of this parameter is 1; that is, one reader thread will be active on every gateway node for each big read operation qualifying for splitting per the parallel read threshold value. The valid range of values is 1 to 64.

**afmNumWriteThreads**

Defines the number of threads that can be used on each participating gateway node during parallel write. The default value of this parameter is 1; that is, one writer thread will be active on every gateway node for each big write operation qualifying for splitting per the parallel write threshold value. Valid values can range from 1 to 64.

**afmParallelReadChunkSize**

Defines the minimum chunk size of the read that needs to be distributed among the gateway nodes during parallel reads. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

**afmParallelReadThreshold**

Defines the threshold beyond which parallel reads become effective. Reads are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value is 1024 MiB. The valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

## **mmchfileset**

### **afmParallelWriteChunkSize**

Defines the minimum chunk size of the write that needs to be distributed among the gateway nodes during parallel writes. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

### **afmParallelWriteThreshold**

Defines the threshold beyond which parallel writes become effective. Writes are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value of this parameter is 1024 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

### **afmPrefetchThreshold**

Controls partial file caching and prefetching. Valid values are the following:

- 0** Enables full file prefetching. This is useful for sequentially accessed files that are read in their entirety, such as image files, home directories, and development environments. The file will be prefetched after three blocks have been read into the cache.

#### **1-99**

Specifies the percentage of file size that must be cached before the entire file is prefetched. A large value is suitable for a file accessed either randomly or sequentially but partially, for which it might be useful to ingest the rest of the file when most of it has been accessed.

#### **100**

Disables full file prefetching. This value only fetches and caches data that is read by the application. This is useful for large random-access files, such as databases, that are either too big to fit in the cache or are never expected to be read in their entirety. When all data blocks are accessed in the cache, the file is marked as cached.

0 is the default value.

For local-updates mode, the whole file is prefetched when the first update is made.

### **afmPrimaryId**

Specifies the unique primary ID of the primary fileset for asynchronous data replication. This is used for connecting a secondary to a primary.

### **afmReadSparseThreshold**

Specifies the size in MB for files in cache beyond which sparseness is maintained. For all files below the specified threshold, sparseness is not maintained.

### **afmRPO**

Specifies the recovery point objective (RPO) interval for an AFM DR fileset. This attribute is disabled by default. You can specify a value with the suffix M for minutes, H for hours, or W for weeks. For example, for 12 hours specify 12H. If you do not add a suffix, the value is assumed to be in minutes. The range of valid values is 720 minutes - 2147483647 minutes.

### **afmShowHomeSnapshot**

Controls the visibility of the home snapshot directory in cache. For this to be visible in cache, this variable has to be set to yes, and the snapshot directory name in cache and home should not be the same.

#### **yes**

Specifies that the home snapshot link directory is visible.

**no** Specifies that the home snapshot link directory is not visible.

See *Peer snapshot -psnap* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

**afmTarget**

The only allowed value is **disable**. It is used to convert AFM filesets to regular independent filesets; for example:

```
mmchfileset fs1 ro -p afmTarget=disable
```

After an AFM fileset is converted to a regular fileset, the fileset cannot be changed back to an AFM fileset.

**--allow-permission-change** *PermissionChangeMode*

Specifies the new permission change mode. This mode controls how **chmod** and ACL operations are handled on objects in the fileset. Valid modes are as follows:

**chmodOnly**

Specifies that only the UNIX change mode operation (**chmod**) is allowed to change access permissions (ACL commands and API will not be accepted).

**setAclOnly**

Specifies that permissions can be changed using ACL commands and API only (**chmod** will not be accepted).

**chmodAndSetAcl**

Specifies that **chmod** and ACL operations are permitted. If the **chmod** command (or **setattr** file operation) is issued, the result depends on the type of ACL that was previously controlling access to the object:

- If the object had a Posix ACL, it will be modified accordingly.
- If the object had an NFSv4 ACL, it will be replaced by the given UNIX mode bits.

**Note:** This is the default setting when a fileset is created.

**chmodAndUpdateAcl**

Specifies that **chmod** and ACL operations are permitted. If **chmod** is issued, the ACL will be updated by privileges derived from UNIX mode bits.

**--inode-limit** *MaxNumInodes* [:*NumInodesToPreallocate*]

Specifies the new inode limit for the inode space owned by the specified fileset. The *FilesetName* or *JunctionPath* must refer to an independent fileset. The *NumInodesToPreallocate* specifies an optional number of additional inodes to pre-allocate for the inode space. Use the **mmchfs** command to change inode limits for the root fileset.

The *MaxNumInodes* and *NumInodesToPreallocate* values can be specified with a suffix, for example 100K or 2M.

**--iam-mode** *Mode*

Specifies the integrated archive manager (IAM) mode for the fileset. IAM modes can be used to modify some of the file-operation restrictions that normally apply to immutable files. The following values (listed in order of strictness) are accepted:

**ad** | **advisory**

**nc** | **noncompliant**

**co** | **compliant**

For more information about IAM modes, see the topic about immutability and `appendOnly` restrictions in *Information lifecycle management for IBM Spectrum Scale of IBM Spectrum Scale: Administration Guide*.

**Note:** If you set new values for **afmParallelReadChunkSize**, **afmParallelReadThreshold**, **afmParallelWriteChunkSize**, and **afmParallelWriteThreshold**; you need not relink filesets for the new values to take effect.

## **mmchfileset**

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority or be a fileset owner to run the **mmchfileset** command with the **-t** option. All other options require root authority.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. This command renames fileset **fset1** to **fset2** and gives it the comment "first fileset":

```
mmchfileset gpfs1 fset1 -j fset2 -t 'first fileset'
```

The system displays a message similar to:

```
Fileset 'fset1' changed.
```

2. To confirm the change, issue this command:

```
mmlsfileset gpfs1 -L
```

The system displays information similar to:

```
Filesets in file system 'gpfs1':
Name Id RootInode ParentId Created InodeSpace MaxInodes AllocInodes Comment
root 0 3 -- Mon Apr 12 16:31:05 2010 0 8001536 8001536 root fileset
fset2 2 13569 0 Mon Apr 12 16:32:28 2010 0 0 0 first fileset
```

### **See also**

- “[mmchfs command](#)” on page 179
- “[mmcrfileset command](#)” on page 244
- “[mmdelfileset command](#)” on page 292
- “[mmlinkfileset command](#)” on page 390
- “[mmlsfileset command](#)” on page 406
- “[mmunlinkfileset command](#)” on page 578

### **Location**

/usr/lpp/mmfs/bin

---

## mmchfs command

Changes the attributes of a GPFS file system.

### Synopsis

```
mmchfs Device [-A {yes | no | automount}] [-D {posix | nfs4}] [-E {yes | no}]
[-k {posix | nfs4 | all}] [-K {no | whenpossible | always}]
[-L LogFileSize] [-m DefaultMetadataReplicas] [-n NumNodes]
[-o MountOptions] [-r DefaultDataReplicas] [-S {yes | no | relatime}]
[-T Mountpoint] [-t DriveLetter] [-V {full | compat}] [-z {yes | no}]
[--filesedf | --nofilesedf]
[--inode-limit MaxNumInodes[:NumInodesToPreallocate]]
[--log-replicas LogReplicas] [--mount-priority Priority]
[--perfileset-quota | --noperfileset-quota]
[--rapid-repair | --norapid-repair]
[--write-cache-threshold HAWCThreshold]
```

or

```
mmchfs Device -Q {yes | no}
```

or

```
mmchfs Device -W NewDeviceName
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmchfs** command to change the attributes of a GPFS file system.

### Parameters

#### *Device*

The device name of the file system to be changed.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**. However, file system names must be unique across GPFS clusters.

This must be the first parameter.

#### -A {yes | no | automount}

Indicates when the file system is to be mounted:

##### **yes**

When the GPFS daemon starts.

##### **no**

Manual mount.

##### **automount**

On non-Windows nodes, when the file system is first accessed. On Windows nodes, when the GPFS daemon starts.

#### Note:

- The file system must be unmounted prior to changing the automount settings.
- IBM Spectrum Protect for Space Management does not support file systems with the **-A** option set to automount.

#### -D {nfs4 | posix}

Specifies whether a deny-write open lock will block writes, which is expected and required by NFS V4, Samba, and Windows. File systems supporting NFS V4 must have **-D nfs4** set. The option **-D**

## **mmchfs**

**posix** allows NFS writes even in the presence of a deny-write open lock. If you intend to export the file system using NFS V4 or Samba, or mount your file system on Windows, you must use **-D nfs4**. For NFS V3 (or if the file system is not NFS exported at all) use **-D posix**.

### **-E {yes | no}**

Specifies whether to report exact **mtime** values. If **-E no** is specified, the **mtime** value is periodically updated. If you desire to always display exact modification times, specify **-E yes**.

### **-k {posix | nfs4 | all}**

Specifies the type of authorization supported by the file system:

#### **posix**

Traditional GPFS ACLs only (NFS V4 and Windows ACLs are not allowed). Authorization controls are unchanged from earlier releases.

#### **nfs4**

Support for NFS V4 and Windows ACLs only. Users are not allowed to assign traditional GPFS ACLs to any file system objects (directories and individual files).

#### **all**

Any supported ACL type is permitted. This includes traditional GPFS (**posix**) and NFS V4 and Windows ACLs (**nfs4**).

The administrator is allowing a mixture of ACL types. For example, **fileA** may have a **posix** ACL, while **fileB** in the same file system may have an NFS V4 ACL, implying different access characteristics for each file depending on the ACL type that is currently assigned.

Avoid specifying **nfs4** or **all** unless files will be exported to NFS V4 or Samba clients, or the file system will be mounted on Windows. NFS V4 and Windows ACLs affect file attributes (mode) and have access and authorization characteristics that are different from traditional GPFS ACLs.

### **-K {no | whenpossible | always}**

Specifies whether strict replication is to be enforced:

**no** Strict replication is not enforced. GPFS will try to create the needed number of replicas, but will still return EOK as long as it can allocate at least one replica.

#### **whenpossible**

Strict replication is enforced provided the disk configuration allows it. If there is only one failure group, strict replication will not be enforced.

#### **always**

Strict replication is enforced.

For more information, see the topic *Strict replication* in the *IBM Spectrum Scale: Problem Determination Guide*.

### **-L LogFileSize**

Specifies the new size of the internal log files. The **LogFileSize** specified must be a multiple of the metadata block size. The minimum size is 256 KB and the maximum size is 1024 MB. Specify this value with the K or M character, for example: 8M.

You must restart the GPFS daemons before the new log file size takes effect. The GPFS daemons can be restarted one node at a time. When the GPFS daemon is restarted on the last node in the cluster, the new log size becomes effective.

### **-m DefaultMetaReplicas**

Changes the default number of metadata replicas. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of **MaxMetaReplicas** set when the file system was created.

Changing the default replication settings using the **mmchfs** command does not change the replication setting of existing files. After running the **mmchfs** command, the **mmrestripefs** command with the **-R** option can be used to change *all* existing files or you can use the **mmchattr** command to change a small number of existing files.

**-n NumNodes**

Changes the number of nodes for a file system but does not change the existing system metadata structures. This setting is just an estimate and can be used only to affect the layout of the system metadata for storage pools created after the setting is changed.

**-o MountOptions**

Specifies the mount options to pass to the mount command when mounting the file system. For a detailed description of the available mount options, see *GPFS-specific mount options* in the *IBM Spectrum Scale: Administration Guide*.

**-Q {yes | no}**

If **-Q yes** is specified, quotas are activated automatically when the file system is mounted. If **-Q no** is specified, the quota files remain in the file system, but are not used.

For more information, see the topic *Enabling and disabling GPFS quota management* in the *IBM Spectrum Scale: Administration Guide*.

**-r DefaultDataReplicas**

Changes the default number of data replicas. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of *MaxDataReplicas* set when the file system was created.

Changing the default replication settings using the **mmchfs** command does not change the replication setting of existing files. After running the **mmchfs** command, the **mmrestripefs** command with the **-R** option can be used to change *all* existing files or you can use the **mmchattr** command to change a small number of existing files.

**-S {yes | no | relatime}**

Suppress the periodic updating of the value of **atime** as reported by the **gpfs\_stat()**, **gpfs\_fstat()**, **stat()**, and **fstat()** calls. If **yes** is specified, these calls report the last time the file was accessed when the file system was mounted with **-S no**.

If **relatime** is specified, the file access time is updated only if the existing access time is older than the value of the **atimeDeferredSeconds** configuration attribute or the existing file modification time is greater than the existing access time.

**-T Mountpoint**

Change the mount point of the file system starting at the next mount of the file system.

The file system must be unmounted on all nodes prior to issuing the command.

**-t DriveLetter**

Changes the Windows drive letter for the file system.

The file system must be unmounted on all nodes prior to issuing the command.

**-V {full | compat}**

Changes the file system format to the latest format supported by the currently installed level of GPFS. This *may* cause the file system to become permanently incompatible with earlier releases of GPFS.

**Note:** The **-V** option cannot be used to make file systems created prior to GPFS 3.2.1.5 available to Windows nodes. Windows nodes can mount only file systems that are created with GPFS 3.2.1.5 or later.

Before issuing **-V**, see *Migration, coexistence and compatibility* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*. Ensure that all nodes in the cluster have been migrated to the latest level of GPFS code and that you have successfully run the **mmchconfig release=LATEST** command.

## **mmchfs**

For information about specific file system format and function changes when you upgrade to the current release, see the topic *File system format changes between versions of GPFS* in the *IBM Spectrum Scale: Administration Guide*.

### **full**

Enables all new functionality that requires different on-disk data structures. Nodes in remote clusters running an older GPFS version will no longer be able to mount the file system. If there are any nodes running an older GPFS version that have the file system mounted at the time the command is issued, the **mmchfs** command will fail.

### **compat**

Enables only backward-compatible format changes. Nodes in remote clusters that are running GPFS 3.5 will still be able to mount the file system. Nodes running GPFS 3.4 or earlier will no longer be able to mount the file system.

#### **-W NewDeviceName**

Assign *NewDeviceName* to be the device name for the file system.

#### **-z {yes | no}**

Enable or disable DMAPI on the file system. Turning this option on will require an external data management application such as IBM Spectrum Protect hierarchical storage management (HSM) before the file system can be mounted.

| **Note:** IBM Spectrum Protect for Space Management does not support file systems with the -A option set to automount.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events* in the *IBM Spectrum Scale: Command and Programming Reference*.

#### **--filesetdf**

Specifies that when quotas are enforced for a fileset, the numbers reported by the **df** command are based on the quotas for the fileset (rather than the entire file system). This option affects the **df** command behavior only on Linux nodes.

#### **--nofilesetdf**

Specifies that when quotas are enforced for a fileset, the numbers reported by the **df** command are based on the quotas for the entire file system (rather than individual filesets).

#### **--inode-limit MaxNumInodes[:NumInodesToPreallocate]**

*MaxNumInodes* specifies the maximum number of files that can be created. Allowable values range from the current number of created inodes (determined by issuing the **mmdf** command with -F), through the maximum number of files possibly supported as constrained by the formula:

**maximum number of files = (total file system space) / (inode size + subblock size)**

**Note:** This formula works only for simpler configurations. For complex configurations, such as separation of data and metadata, this formula might not provide an accurate result.

If your file system has additional disks added or the number of inodes was insufficiently sized at file system creation, you can change the number of inodes and hence the maximum number of files that can be created.

For file systems that will be doing parallel file creates, if the total number of free inodes is not greater than 5% of the total number of inodes, there is the potential for slowdown in file system access. Take this into consideration when changing your file system.

*NumInodesToPreallocate* specifies the number of inodes that will be pre-allocated by the system right away. If this number is not specified, GPFS allocates inodes dynamically as needed.

The *MaxNumInodes* and *NumInodesToPreallocate* values can be specified with a suffix, for example 100K or 2M. Note that in order to optimize file system operations, the number of inodes that are actually created may be greater than the specified value.

This option applies only to the root fileset. When there are multiple inode spaces, use the **--inode-space** option of the **mmchfileset** command to alter the inode limits of independent filesets. The **mmchfileset** command can also be used to modify the root inode space. The **--inode-space** option of the **mmlsfs** command shows the sum of all inode spaces.

#### **--log-replicas** *LogReplicas*

Specifies the number of recovery log replicas. Valid values are **1**, **2**, **3**, or **DEFAULT**. If **DEFAULT** is specified, the number of log replicas is the same as the number of metadata replicas currently in effect for the file system and will change when this number is changed.

Changing the default replication settings using the **mmchfs** command does not change the replication setting of existing files. After running the **mmchfs** command, the **mmrestripefs** command with the **-R** option can be used to change existing log files.

This option is only applicable if the recovery log is stored in the **system.log** storage pool.

#### **--mount-priority** *Priority*

Controls the order in which the individual file systems are mounted at daemon startup or when one of the **all** keywords is specified on the **mmmount** command.

File systems with higher *Priority* numbers are mounted after file systems with lower numbers. File systems that do not have mount priorities are mounted last. A value of zero indicates no priority.

#### **--profileset-quota**

Sets the scope of user and group quota limit checks to the individual fileset level, rather than to the entire file system.

#### **--noperfileset-quota**

Sets the scope of user and group quota limit checks to the entire file system, rather than to individual filesets.

#### **--rapid-repair**

Keeps track of incomplete replication on an individual file block basis (as opposed to the entire file). This may result in a faster repair time when very large files are only partially ill-replicated.

#### **--norapid-repair**

Specifies that replication status is kept on a whole file basis (rather than on individual block basis).

#### **--write-cache-threshold** *HAWCThreshold*

Specifies the maximum length (in bytes) of write requests that will be initially buffered in the highly-available write cache before being written back to primary storage. Only synchronous write requests are guaranteed to be buffered in this fashion.

A value of 0 disables this feature. 64K is the maximum supported value. Specify in multiples of 4K.

This feature can be enabled or disabled at any time (the file system does not need to be unmounted). For more information about this feature, see the topic *Highly-available write cache (HAWC)* in the *IBM Spectrum Scale: Administration Guide*.

## Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmchfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For

## **mmchfs**

more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

To change the default replicas for metadata to 2 and the default replicas for data to 2 for new files created in the **fs0** file system, issue this command:

```
mmchfs fs0 -m 2 -r 2
```

To confirm the change, issue this command:

```
mmlsfs fs0 -m -r
```

The system displays information similar to:

| flag | value | description                         |
|------|-------|-------------------------------------|
| -m   | 2     | Default number of metadata replicas |
| -r   | 2     | Default number of data replicas     |

### **See also**

- “mmchfileset command” on page 173
- “mmcrfs command” on page 250
- “mmdelfs command” on page 295
- “mmdf command” on page 308
- “mmfsck command” on page 331
- “mmlsfs command” on page 411
- “mmrestripefs command” on page 532

### **Location**

/usr/lpp/mmfs/bin

---

## mmchlicense command

Controls the type of GPFS license associated with the nodes in the cluster.

### Synopsis

```
mmchlicense {client|fpo|server} [--accept] -N {Node[,Node...] | NodeFile | NodeClass}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmchlicense** command to change the type of GPFS license associated with the nodes in the cluster.

For information on IBM Spectrum Scale license designation, see *IBM Spectrum Scale license designation* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### Parameters

#### **client | fpo | server**

The type of GPFS license to be assigned to the nodes specified with the **-N** parameter.

##### **client**

The IBM Spectrum Scale Client license permits exchange of data between nodes that locally mount the same GPFS file system. No other export of the data is permitted. The GPFS client may not be used for nodes to share GPFS data directly through any application, service, protocol or method, such as Network File System (NFS), Common Internet File System (CIFS), File Transfer Protocol (FTP), or Hypertext Transfer Protocol (HTTP). For these functions, an IBM Spectrum Scale Server license would be required. The use of any of the following components or functions of IBM Spectrum Scale Client is not authorized:

- Configuring a virtual server in the following IBM Spectrum Scale roles: Configuration Manager, Quorum node, Manager node, Network Shared Disk (NSD) Server node, Cluster Export Services node (also known as Protocol node), Advanced File Management (AFM) Gateway node, Transparent Cloud Tiering Gateway node.
- Exporting IBM Spectrum Scale data to virtual servers that do not have a valid IBM Spectrum Scale license through any application, protocol or method, including Network File System (NFS), Server Message Block (SMB), File Transfer Protocol (FTP), Hypertext Transfer Protocol (HTTP), Object Protocol (OpenStack Swift, Amazon S3 API).

##### **server**

The IBM Spectrum Scale Server license permits the licensed node to perform GPFS management functions such as cluster configuration manager, quorum node, manager node, and Network Shared Disk (NSD) server. In addition, the IBM Spectrum Scale Server license permits the licensed node to share GPFS data directly through any application, service protocol or method such as NFS, CIFS, FTP, or HTTP. Therefore, protocol nodes also require an IBM Spectrum Scale Server license.

##### **fpo**

The IBM Spectrum Scale FPO license permits the licensed node to perform NSD server functions for sharing GPFS data with other nodes that have an IBM Spectrum Scale FPO or IBM Spectrum Scale Server license. This license cannot be used to share data with nodes that have an IBM Spectrum Scale Client license or non-GPFS nodes. The use of any of the following components or functions of IBM Spectrum Scale FPO is not authorized:

## **mmchlicense**

- Configuring a virtual server in the following IBM Spectrum Scale roles: Configuration Manager, Quorum node, Manager node, Cluster Export Services node (also known as Protocol node), Advanced File Management (AFM) Gateway node, Transparent Cloud Tiering Gateway node.
- Configuring a virtual server as an IBM Spectrum Scale Network Shared Disk (NSD) Server node for providing IBM Spectrum Scale data access to virtual servers that do not have a valid IBM Spectrum Scale Server or IBM Spectrum Scale FPO license entitlement.
- Exporting IBM Spectrum Scale data to virtual servers that do not have a valid IBM Spectrum Scale license through any application, protocol or method, including Network File System (NFS), Server Message Block (SMB), File Transfer Protocol (FTP), Hypertext Transfer Protocol (HTTP), Object Protocol (OpenStack Swift, Amazon S3 API).

The full text of the Licensing Agreement is provided with the installation media and can be found at the IBM Software license agreements website ([www.ibm.com/software/sla/sladb.nsf](http://www.ibm.com/software/sla/sladb.nsf)).

### **--accept**

Indicates that you accept the applicable licensing terms. The license acceptance prompt will be suppressed.

### **-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes that are to be assigned the specified license type.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands in the IBM Spectrum Scale: Administration Guide*.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmchlicense** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

To designate nodes k145n04 and k145n05 as possessing a GPFS server license, issue this command:

```
mmchlicense server --accept -N k145n04,k145n05
```

The system displays information similar to:

```
The following nodes will be designated as possessing GPFS server licenses:
```

```
 k145n04.kgn.ibm.com
```

```
 k145n05.kgn.ibm.com
```

```
mmchlicense: Command successfully completed
```

```
mmchlicense: Propagating the cluster configuration data to all
```

```
 affected nodes. This is an asynchronous process.
```

## **See also**

- “[mmlslicense command](#)” on page 415

**Location**

/usr/lpp/mmfs/bin

## **mmchmgr**

---

### **mmchmgr command**

Assigns a new file system manager node or cluster manager node.

#### **Synopsis**

```
mmchmgr {Device | -c} [Node]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmchmgr** command assigns a new file system manager node or cluster manager node.

#### **Parameters**

##### *Device*

The device name of the file system for which the file system manager node is to be changed. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

**-c** Changes the cluster manager node.

##### *Node*

The target node to be appointed as either the new cluster manager node or the new file system manager node. Target nodes for manager functions are selected according to the following criteria:

- Target nodes for the cluster manager function must be specified from the list of quorum nodes.
- Target nodes for the file system manager function should be specified from the list of manager nodes, although this is not strictly required.

If *Node* is not specified, the new manager is selected automatically.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmchmgr** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

#### **Examples**

1. Assume the file system manager for the file system **gpfsl1** is currently **k164n05**. To migrate the file system manager responsibilities to **k164n06**, issue this command:

```
mmchmgr gpfsl1 k164n06
```

The system displays information similar to:

```
GPFS: 6027-628 Sending migrate request to current manager node 89.116.68.69 (k164n05).
GPFS: 6027-629 [N] Node 89.116.68.69 (k164n05) resigned as manager for gpfs1.
GPFS: 6027-630 [N] Node 89.116.68.70 (k164n06) appointed as manager for gpfs1.
```

To verify the change, issue the command:

```
mmlsmgr gpfs1
```

The system displays information similar to:

```
file system manager node [from 89.116.68.69 (k164n05)]

gpfs1 89.116.68.69 (k164n06)
```

2. To change the cluster manager node, issue the command:

```
mmchmgr -c c5n107
```

The system displays information similar to:

```
Appointing node 9.114.132.107 (c5n107) as cluster manager
Node 9.114.132.107 (c5n107) has taken over as cluster manager
```

To verify the change, issue the command:

```
mmlsmgr -c
```

The system displays information similar to:

```
Cluster manager node: 9.114.132.107 (c5n107)
```

## See also

- “mmlsmgr command” on page 417

## Location

/usr/lpp/mmfs/bin

## **mmchnode**

---

### **mmchnode command**

Changes node attributes.

#### **Synopsis**

```
mmchnode change-options -N {Node[,Node...] | NodeFile | NodeClass}
[--cloud-gateway-nodeclass CloudGatewayNodeClass]
```

or

```
mmchnode {-S Filename | --spec-file=Filename}
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmchnode** command to change one or more attributes on a single node or on a set of nodes. If conflicting node designation attributes are specified for a given node, the last value is used. If any of the attributes represent a node-unique value, the **-N** option must resolve to a single node.

Do not use the **mmchnode** command to change the gateway node role while IO is happening on the fileset. Run the **flushpending** command to flush any pending messages from queues before running the **mmchnode** command for the gateway node role changes.

#### **Parameters**

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes whose states are to be changed.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

**[--cloud-gateway-nodeclass CloudGatewayNodeClass]**

Use this option to specify a node class you will use for Transparent cloud tiering management along with the **-N** option where you will specify individual node names. Both **-N** with a node list and **--cloud-gateway-node** with a node class will be required.

**-S Filename | --spec-file=Filename**

Specifies a file with a detailed description of the changes to be made. Each line represents the changes to an individual node and has the following format:

*node-identifier change-options*

*change-options*

A blank-separated list of attribute[=value] pairs. The following attributes can be specified:

**--admin-interface={hostname | ip\_address}**

Specifies the name of the node to be used by GPFS administration commands when communicating between nodes. The admin node name must be specified as an IP address or a hostname that is resolved by the host command to the desired IP address. If the keyword DEFAULT is specified, the admin interface for the node is set to be equal to the daemon interface for the node.

**--client**

Specifies that the node should not be part of the pool of nodes from which cluster managers, file system managers, and token managers are selected.

**--cloud-gateway-enable**

Enables one or more nodes as Transparent cloud tiering nodes on the cluster based on the **-N** option parameters.

**--cloud-gateway-disable**

Disables one or more Transparent cloud tiering nodes from the cluster based on the -N option parameters. Only disable a Transparent cloud tiering node if you no longer need it to migrate or recall data from the configured cloud.

**--ces-enable**

Enables Cluster Export Services (CES) on the node.

**--ces-disable**

Disables CES on the node.

**--ces-group=Group[,Group...]**

Adds one or more groups to the specified nodes.

**--no ces-group=Group[,Group...]**

Removes one or more groups from the specified nodes.

**--cnfs-disable**

Disables the CNFS functionality of a CNFS member node.

**--cnfs-enable**

Enables a previously-disabled CNFS member node.

**--cnfs-groupid=groupid**

Specifies a failover recovery group for the node. If the keyword DEFAULT is specified, the CNFS recovery group for the node is set to zero.

For more information, see *Implementing a clustered NFS environment on Linux in IBM Spectrum Scale: Administration Guide*.

**--cnfs-interface=ip\_address\_list**

A comma-separated list of host names or IP addresses to be used for GPFS cluster NFS serving.

The specified IP addresses can be real or virtual (aliased). These addresses must be configured to be static (not DHCP) and to not start at boot time.

The GPFS daemon interface for the node cannot be a part of the list of CNFS IP addresses.

If the keyword DEFAULT is specified, the CNFS IP address list is removed and the node is no longer considered a member of CNFS.

If **adminMode** central is in effect for the cluster, all CNFS member nodes must be able to execute remote commands without the need for a password.

For more information, see *Implementing a clustered NFS environment on Linux in IBM Spectrum Scale: Administration Guide*.

**--daemon-interface={hostname | ip\_address}**

Specifies the host name or IP address to be used by the GPFS daemons for node-to-node communication. The host name or IP address must refer to the communication adapter over which the GPFS daemons communicate. Alias interfaces are not allowed. Use the original address or a name that is resolved by the host command to the original address.

Before you specify this option, you must stop GPFS on all the nodes in the cluster. You cannot use the keyword DEFAULT with this option.

You cannot specify the **--daemon-interface** option for a quorum node if CCR is enabled.

Temporarily change the node to a nonquorum node. Then run the **mmchnode** command with the **--daemon-interface** option against the nonquorum node. Finally, change the node back into a quorum node.

**--gateway | --nogateway**

Specifies whether the node is to be designated as a gateway node or not.

## **mmchnode**

### **--manager | --nomanager**

Designates the node as part of the pool of nodes from which file system managers and token managers are selected.

### **--nonquorum**

Designates the node as a non-quorum node. If two or more quorum nodes are downgraded at the same time, GPFS must be stopped on all nodes in the cluster. GPFS does not have to be stopped if the nodes are downgraded one at a time.

### **--perfmon | --noperfmon**

Specifies whether the node is to be designated as a perfmon node or not.

### **--nosnmp-agent**

Stops the SNMP subagent and specifies that the node should no longer serve as an SNMP collector node. For more information, see *GPFS SNMP support in IBM Spectrum Scale: Problem Determination Guide*.

### **--quorum**

Designates the node as a quorum node.

**Note:** If you are designating a node as a quorum node, and **adminMode central** is in effect for the cluster, you must ensure that GPFS is up and running on that node (**mmgetstate** reports the state of the node as **active**).

### **--snmp-agent**

Designates the node as an SNMP collector node. If the GPFS daemon is active on this node, the SNMP subagent will be started as well. For more information, *GPFS SNMP support in IBM Spectrum Scale: Problem Determination Guide*.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmchnode** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To change nodes k145n04 and k145n05 to be both quorum and manager nodes, issue this command:

```
mmchnode --quorum --manager -N k145n04,k145n05
```

The system displays information similar to:

```
Wed May 16 04:50:24 EDT 2007: mmchnode: Processing node k145n04.kgn.ibm.com
Wed May 16 04:50:24 EDT 2007: mmchnode: Processing node k145n05.kgn.ibm.com
mmchnode: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

After completion, **mmlscluster** displays information similar to:

```
GPFS cluster information
=====
GPFS cluster name: mynodes.kgn.ibm.com
GPFS cluster id: 680681553700098206
```

```

GPFS UID domain: mynodes.kgn.ibm.com
Remote shell command: /usr/bin/ssh
Remote file copy command: /usr/bin/scp

GPFS cluster configuration servers:

Primary server: k145n04.kgn.ibm.com
Secondary server: k145n06.kgn.ibm.com

Node Daemon node name IP address Admin node name Designation

1 k145n04.kgn.ibm.com 9.114.68.68 k145n04.kgn.ibm.com quorum-manager
2 k145n05.kgn.ibm.com 9.114.68.69 k145n05.kgn.ibm.com quorum-manager
3 k145n06.kgn.ibm.com 9.114.68.70 k145n06.kgn.ibm.com

```

2. To change nodes k145n04 and k145n05 to be both quorum and manager nodes, and node k45n06 to be a non-quorum node, issue this command:

```
mmchnode -S /tmp/specFile
```

Where the contents of /tmp/specFile are:

```
k145n04 --quorum --manager
k145n05 --quorum --manager
k145n06 --nonquorum
```

The system displays information similar to:

```

Wed May 16 05:23:31 EDT 2007: mmchnode: Processing node k145n04
Wed May 16 05:23:32 EDT 2007: mmchnode: Processing node k145n05
Wed May 16 05:23:32 EDT 2007: mmchnode: Processing node k145n06
Verifying GPFS is stopped on all nodes ...
mmchnode: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.

```

And **mmlscluster** displays information similar to:

```

GPFS cluster information
=====
GPFS cluster name: mynodes.kgn.ibm.com
GPFS cluster id: 680681553700098206
GPFS UID domain: mynodes.kgn.ibm.com
Remote shell command: /usr/bin/rsh
Remote file copy command: /usr/bin/rcp

GPFS cluster configuration servers:

Primary server: k145n04.kgn.ibm.com
Secondary server: k145n06.kgn.ibm.com

Node Daemon node name IP address Admin node name Designation

1 k145n04.kgn.ibm.com 9.114.68.68 k145n04.kgn.ibm.com quorum-manager
2 k145n05.kgn.ibm.com 9.114.68.69 k145n05.kgn.ibm.com quorum-manager
3 k145n06.kgn.ibm.com 9.114.68.70 k145n06.kgn.ibm.com

```

3. To enable all the nodes specified in the node class TCTNodeClass1 as Transparent cloud tiering nodes, issue this command:

```
mmchnode --cloud-gateway-enable -N TCTNodeClass1
```

The system displays output similar to this:

```

Wed May 11 12:51:37 EDT 2016: mmchnode: Processing node c350f2u18
mmchnode: Verifying media for Transparent Cloud Tiering nodes...
mmchnode: node c350f2u18 media checks passed.

```

```

Wed May 11 12:51:38 EDT 2016: mmchnode: Processing node c350f2u22.pk.labs.ibm.com
mmchnode: node c350f2u22.pok.stglabs.ibm.com media checks passed.

```

```
Wed May 11 12:51:41 EDT 2016: mmchnode: Processing node c350f2u26.pk.labs.ibm.com
```

## **mmchnode**

```
mmchnode: node c350f2u26.pok.stglabs.ibm.com media checks passed.
```

```
mmchnode: Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.
```

You can verify the Transparent cloud tiering nodes by issuing this command:

```
mmcloudgateway node list
```

4. To designate only a few nodes (node1 and node2) in the node class, *TCTNodeClass1*, as Transparent cloud tiering server nodes, issue this command:

```
mmchnode --cloud-gateway-enable -N node1,node2 --cloud-gateway-nodeclass TCTNodeClass1
```

**Note:** It only designates node1 and node2 as Transparent cloud tiering server nodes from the node class, *TCTNodeClass1*. Administrators can continue to use the node class for other purposes.

5. To disable all Transparent cloud tiering nodes from the node class, *TCTNodeClass1*, issue this command:

```
mmchnode --cloud-gateway-disable -N TCTNodeClass1
```

The system displays output similar to this:

```
Thu May 12 16:10:11 EDT 2016: mmchnode: Processing node c350f2u18
mmchnode: Verifying Transparent Cloud Tiering node c350f2u18 can be disabled...
mmchnode: Node c350f2u18 passed disable checks.
```

```
Thu May 12 16:10:11 EDT 2016: mmchnode: Processing node c350f2u22.pk.labs.ibm.com
mmchnode: Verifying Transparent Cloud Tiering node c350f2u22.pk.labs.ibm.com can be disabled...
mmchnode: Node c350f2u22.pok.stglabs.ibm.com passed disable checks.
```

```
Thu May 12 16:10:14 EDT 2016: mmchnode: Processing node c350f2u26.pk.labs.ibm.com
mmchnode: Verifying Transparent Cloud Tiering node c350f2u26.pk.labs.ibm.com can be disabled...
mmchnode: Node c350f2u26.pk.labs.ibm.com passed disable checks.
```

mmchnode: Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.

6. To disable only a few nodes (node1 and node2) from the node class, *TCTNodeClass1*, as Transparent cloud tiering server nodes, issue this command:

```
mmchnode --cloud-gateway-disable -N node1,node2 --cloud-gateway-nodeclass TCTNodeClass1
```

**Note:** It only disables node1 and node2 as Transparent cloud tiering server nodes from the node class, *TCTNodeClass1*.

## **See also**

- “mmchconfig command” on page 132
- “mmlscluster command” on page 397

## **Location**

/usr/lpp/mmfs/bin

---

## mmchnodeclass command

Changes user-defined node classes.

### Synopsis

```
mmchnodeclass ClassName {add | delete | replace}
-N {Node[,Node...] | NodeFile | NodeClass}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmchnodeclass** command to make changes to existing user-defined node classes.

### Parameters

#### *ClassName*

Specifies the name of an existing user-defined node class to modify.

#### **add**

Adds the nodes specified with the -N option to *ClassName*.

#### **delete**

Deletes the nodes specified with the -N option from *ClassName*.

#### **replace**

Replaces all *ClassName* members with a new list of nodes specified with the -N option.

#### -N {*Node*[,*Node*...] | *NodeFile* | *NodeClass*}

Specifies the member names of nodes and node classes that will be used for the **add**, **delete**, or **replace** action.

*NodeClass* cannot be used to add members that already contain other node classes. For example, two user-defined node classes called **siteA** and **siteB** were used to create a new node class called **siteAandB**, as follows:

```
mmcrnodeclass siteAandB -N siteA,siteB
```

The **siteAandB** node class cannot later be specified for *NodeClass* when adding to existing node classes.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### Security

You must have root authority to run the **mmchnodeclass** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **mmchnodeclass**

### **Examples**

To display the current members of a user-defined node class called **siteA**, issue this command:

```
mm1snodeclass siteA
```

The system displays information similar to:

| Node Class Name | Members             |
|-----------------|---------------------|
| siteA           | c8f2c4vp1,c8f2c4vp2 |

To add node **c8f2c1vp4** to the member list of the user-defined node class **siteA**, issue this command:

```
mmchnodeclass siteA add -N c8f2c1vp4
```

The system displays information similar to:

```
mmchnodeclass: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

To display the updated member list of **siteA**, issue this command:

```
mm1snodeclass siteA
```

The system displays information similar to:

| Node Class Name | Members                       |
|-----------------|-------------------------------|
| siteA           | c8f2c1vp4,c8f2c4vp1,c8f2c4vp2 |

To delete node **c8f2c4vp2** from the member list of **siteA**, issue this command:

```
mmchnodeclass siteA delete -N c8f2c4vp2
```

The system displays information similar to:

```
mmchnodeclass: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

To display the updated member list of **siteA**, issue this command:

```
mm1snodeclass siteA
```

The system displays information similar to:

| Node Class Name | Members             |
|-----------------|---------------------|
| siteA           | c8f2c1vp4,c8f2c4vp1 |

To replace all the current members of **siteA** with the members of node class **linuxNodes**, issue this command:

```
mmchnodeclass siteA replace -N linuxNodes
```

The system displays information similar to:

```
mmchnodeclass: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

To display the updated member list of **siteA**, issue this command:

```
mm1snodeclass siteA
```

The system displays information similar to:

| Node Class Name | Members    |
|-----------------|------------|
| siteA           | linuxNodes |

**See also**

- “[mmcrnodeclass command](#)” on page 260
- “[mmdelnodeclass command](#)” on page 300
- “[mmlsnodeclass command](#)” on page 421

**Location**

/usr/lpp/mmfs/bin

## **mmchnsd**

---

### **mmchnsd command**

Changes Network Shared Disk (NSD) configuration attributes.

#### **Synopsis**

```
mmchnsd {"DiskDesc[;DiskDesc...]" | -F StanzaFile}
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmchnsd** command serves several purposes. You can use it to:

- Specify a server list for an NSD that does not have one.
- Change the NSD server nodes specified in the server list.
- Delete the server list. The disk must now be SAN-attached to all nodes in the cluster on which the file system will be mounted.

You must follow these rules when changing NSDs:

- Identify the disks by the NSD names that were given to them by the **mmcrlnsd** command.
- Explicitly specify values for all NSD servers on the list even if you are only changing one of the values.
- Unmount the file system that contains the NSD being changed prior to issuing the **mmchnsd** command.
- Connect the NSD to the new nodes prior to issuing the **mmchnsd** command.
- **mmchnsd** cannot change the disk usage or failure group for an NSD. Use the **mmchdisk** command to change these attributes.
- To move a disk from one storage pool to another, use the **mmdeldisk** and **mmadddisk** commands.
- You cannot change the name of the NSD.

Prior to GPFS 3.5, the disk information was specified in the form of disk descriptors defined as:

DiskName:ServerList:

For backward compatibility, the **mmchnsd** command will still accept the traditional disk descriptors but their use is discouraged.

#### **Parameters**

##### *DiskDesc*

A descriptor for each NSD to be changed. Each descriptor is separated by a semicolon (;). The entire list must be enclosed in single or double quotation marks. The use of disk descriptors is discouraged.

##### **-F StanzaFile**

Specifies a file containing the NSD stanzas for the disks to be changed. NSD stanzas have this format:

```
%nsd:
 nsd=NsdName
 servers=ServerList
 usage=DiskUsage
 failureGroup=FailureGroup
 pool=StoragePool
 device=DiskName
```

where:

**nsd=NsName**

Is the NSD name that was given to the disk by the **mmcrnsd** command. This clause is mandatory for the **mmchnsd** command.

**servers=ServerList**

Is a comma-separated list of NSD server nodes. You can specify up to eight NSD servers in this list. The defined NSD will preferentially use the first server on the list. If the first server is not available, the NSD will use the next available server on the list.

When specifying server nodes for your NSDs, the output of the **mmlscluster** command lists the host name and IP address combinations recognized by GPFS. The utilization of aliased host names not listed in the **mmlscluster** command output may produce undesired results.

If you do not define a *ServerList*, GPFS assumes that the disk is SAN-attached to all nodes in the cluster. If all nodes in the cluster do not have access to the disk, or if the file system to which the disk belongs is to be accessed by other GPFS clusters, you must specify a value for *ServerList*.

To remove the NSD server list, do not specify a value for *ServerList* (remove or comment out the **servers=ServerList** clause of the NSD stanza).

**usage=DiskUsage**

Specifies the type of data to be stored on the disk. If this clause is specified, the value must match the type of usage already in effect for the disk; **mmchnsd** cannot be used to change this value.

**failureGroup=FailureGroup**

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

Disk that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

If this clause is specified, the value must match the failure group already in effect for the disk; **mmchnsd** cannot be used to change this value.

**pool=StoragePool**

Specifies the storage pool to which the disk is to be assigned. If this clause is specified, the value must match the storage pool already in effect for the disk; **mmchnsd** cannot be used to change this value.

**device=DiskName**

Is the block device name of the underlying disk device. This clause is ignored by the **mmchnsd** command.

## Exit status

0      Successful completion.

**nonzero**

A failure has occurred.

## **mmchnsd**

### **Security**

You must have root authority to run the **mmchnsd** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

If the disk **gpfs1nsd** is currently defined with **k145n05** as the first server and **k145n07** as the second server, and you want to replace **k145n05** with **k145n09**, create a file **./newNSDstanza** that contains:

```
%nsd: nsd=gpfs1nsd
 servers=k145n09,k148n07
```

Issue this command:

```
mmchnsd -F ./newNSDstanza
```

To confirm the changes, issue this command:

```
mmlsnsd -d gpfs1nsd
```

The system displays information similar to:

| File system | Disk name | NSD servers                                     |
|-------------|-----------|-------------------------------------------------|
| fs2         | gpfs1nsd  | k145n09.ppd.pok.ibm.com,k145n07.ppd.pok.ibm.com |

### **See also**

- “[mmchdisk command](#)” on page 161
- “[mmcrcluster command](#)” on page 239
- “[mmcrnsd command](#)” on page 262
- “[mmlsnsd command](#)” on page 423

### **Location**

/usr/lpp/mmfs/bin

## mmchpolicy command

Establishes policy rules for a GPFS file system.

### Synopsis

```
mmchpolicy Device PolicyFilename [-t DescriptiveName] [-I {yes | test}]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher.

### Description

Use the **mmchpolicy** command to establish the rules for policy-based lifecycle management of the files in a given GPFS file system. Some of the things that can be controlled with the help of policy rules are:

- File placement at creation time
- Snapshot data placement during file writes and deletes
- Replication factors
- Movement of data between storage pools
- File deletion

The **mmapplypolicy** command must be run to move data between storage pools or delete files.

Policy changes take effect immediately on all nodes that have the affected file system mounted. For nodes that do not have the file system mounted, policy changes take effect upon the next mount of the file system.

**For file systems that are created at or upgraded to product version V4.1.1 or later:** If there are no **SET POOL** policy rules installed to a file system by **mmchpolicy**, the system acts as if the single rule **SET POOL 'first-data-pool'** is in effect, where *first-data-pool* is the firstmost non-system pool that is available for file data storage, if such a non-system pool is available. ("Firstmost" is the first according to an internal index of all pools.) However, if there are no policy rules installed and there is no non-system pool, the system acts as if **SET POOL 'system'** is in effect.

This change applies only to file systems that were created at or upgraded to V4.1.1 or later. Until a file system is upgraded, if no **SET POOL** rules are present (set by **mmchpolicy**) for the file system, all data will be stored in the '**system**' pool.

For information on GPFS policies, see the *IBM Spectrum Scale: Administration Guide*.

### Parameters

#### *Device*

Specifies the device name of the file system for which policy information is to be established or changed. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**. This must be the first parameter.

#### *PolicyFilename*

Specifies the name of the file that contains the policy rules. If you specify **DEFAULT**, GPFS replaces the current policy file with a single policy rule that assigns all newly-created files to the **system** storage pool.

### Options

#### **-I {yes | test}**

Specifies whether to activate the rules in the policy file *PolicyFileName*.

## **mmchpolicy**

### **yes**

The policy rules are validated and immediately activated. This is the default.

### **test**

The policy rules are validated, but not installed.

### **-t DescriptiveName**

Specifies an optional descriptive name to be associated with the policy rules. The string must be less than 256 characters in length. If not specified, the descriptive name defaults to the base name portion of the *PolicyFileName* parameter.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmchpolicy** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. This command validates a policy before it is installed:

```
mmchpolicy fs2 fs2.pol -I test
```

The system displays output similar to:

```
Validated policy `fs2.pol': parsed 3 Placement Rules, 0 Restore Rules,
 3 Migrate/Delete/Exclude Rules, 0 List Rules, 0 External Pool/List Rules
```

2. This command installs a policy:

```
mmchpolicy fs2 fs2.pol
```

The system displays output similar to:

```
Validated policy `fs2.pol': parsed 1 Placement Rules, 0 Restore Rules,
 0 Migrate/Delete/Exclude Rules, 1 List Rules, 1 External Pool/List Rules
Policy `fs2.pol' installed and broadcast to all nodes.
```

To confirm the change, issue this command:

```
mmlspolicy fs2
```

The system displays output similar to:

```
Policy file for file system '/dev/fs2':
 Installed by root@k155n11.kgn.ibm.com on Mon Dec 12
 16:56:31 2005.
 First line from original file 'fs2.pol' was:
 /* This is the policy for the fs2 GPFS file system. */
```

## **See also**

- “[mmapplypolicy command](#)” on page 56
- “[mmlspolicy command](#)” on page 426

**Location**

/usr/lpp/mmfs/bin

## **mmchpool**

---

### **mmchpool command**

Modifies storage pool properties.

#### **Synopsis**

```
mmchpool Device {PoolName[,PoolName...] | all}
 [--block-group-factor BlockGroupFactor]
 [--write-affinity-depth WriteAffinityDepth]
```

or

```
mmchpool Device -F PoolDescriptorFile
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

When running the **mmchpool** command, the file system must be unmounted on all nodes.

#### **Description**

Use the **mmchpool** command to change storage pool properties.

#### **Parameters**

##### *Device*

Specifies the device name of the file system for which storage pool information is to be changed. File system names do not need to be fully qualified; for example, fs0 is as acceptable as /dev/fs0.

##### *PoolName[,PoolName...]*

Specifies one or more storage pools for which attributes will be changed.

##### *all*

Changes the attributes for all the storage pools in the specified file system.

##### **--block-group-factor** *BlockGroupFactor*

Specifies how many file system blocks are laid out sequentially on disk to behave like a single large block. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

##### **--write-affinity-depth** *WriteAffinityDepth*

Specifies the allocation policy to be used. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

##### **-F** *PoolDescriptorFile*

Specifies a file used to describe the storage pool attributes. The file contains one line per storage pool, in the following format:

```
%pool:name:blockSize:diskUsage:reserved:maxDiskSize:allocationType:allowWriteAffinity:writeAffinityDepth:blockGroupFactor:
```

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmchpool** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Example

For example, to change the **writeAffinityDepth** to 2 for FPO pool pool1 of file system fs1, issue this command:

```
mmchpool fs1 pool1 --write-affinity-depth 2
```

To confirm the change, issue this command:

```
mmlspool fs1 pool1 -L
```

The system displays information similar to the following:

```
mmlspool fs_1 p1 -L
Pool:
 name = pool1
 poolID = 65537
 blockSize = 4 MB
 usage = dataOnly
 maxDiskSize = 11 TB
 layoutMap = cluster
 allowWriteAffinity = yes
 writeAffinityDepth = 2
 blockGroupFactor = 128
```

## See also

- “mmlspool command” on page 428

## Location

/usr/lpp/mmfs/bin

## **mmchqos**

---

### **mmchqos command**

Changes the Quality of Service for I/O operations (QoS) settings for a file system.

#### **Synopsis**

```
mmchqos Device --enable [--reset] [--force] [--fine-stats Seconds] [--pid-stats {yes|no}]
[[--N {Node[,Node...]} | NodeFile | NodeClass}]
[-C {all | all_remote | ClusterName[,ClusterName...]}]
pool=PoolName[,QOSClass={nnnIOPS | unlimited}][,QOSClass=...]
...
or
mmchqos Device --enable [--reset] [--force] -F StanzaFile [--fine-stats Seconds] [--pid-stats {yes|no}]
or
mmchqos Device --disable
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

With the **mmchqos** command, you can regulate I/O access to a specified storage pool by allocating shares of I/O operations to two QoS classes:

- A **maintenance** class for I/O-intensive, potentially long-running GPFS commands. Typically you assign fewer IOPS to this class to prevent the I/O-intensive commands from dominating file system performance and significantly delaying other tasks.
- An **other** class for all other processes. Typically you assign more IOPS or **unlimited** to this class so that normal processes have greater access to I/O resources and finish more quickly.

A third class, **misc**, is used to count the IOPS that some critical file system processes consume. You cannot assign IOPS to this class, but its count of IOPS is displayed in the output of the **mmlsqos** command.

When QoS is enabled, it restricts the active processes in a QoS class from collectively consuming more than the number of IOPS that you allocate to the class. It queues further I/O attempts until more I/O operations become available.

Remember the following points:

- You can allocate shares of IOPS separately for each storage pool.
- QoS divides each IOPS allocation equally among the specified nodes that have mounted the file system. See Table 11 on page 208.
- Allocations persist across unmounting and remounting the file system.
- QoS stops applying allocations when you unmount the file system and resumes when you remount it.
- When you change allocations or mount the file system, a brief delay due to reconfiguration occurs before QoS starts applying allocations.

For more information about this command, see *Setting the Quality of Service for I/O operations (QoS)* in *IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system to which the command applies.

**--enable**

Causes QoS to start or to continue to apply IOPS allocations.

- If you are specifying the **--enable** option for the first time, then QoS sets the QoS classes of pools that you do not specify in the command to **unlimited** IOPS.
- On subsequent enables, QoS sets only the IOPS allocations that you specify in the command. It does not disturb other IOPS allocations. Compare the **--reset** parameter.

**--disable**

Causes QoS to stop applying IOPS allocations. Lets the file system run without any participation by QoS.

**--reset**

Causes QoS to set any QoS classes that you do not specify in the same command to **unlimited** IOPS.

When you enter multiple **mmchqos** commands for different storage pools, QoS typically records the settings for each pool and regulates the I/O consumption of each pool accordingly. However, with the **--reset** parameter, QoS discards the settings for all pools that are not specified in the same command and sets their QoS classes to **unlimited**. You can use this feature to reset the settings for any pools that you no longer want QoS to regulate and monitor.

**--force**

Causes QoS to accept an IOPS value lower than 100 IOPS.

Assigning less than 100 IOPS to a class is typically ineffective, because processes in that class run for an indefinitely long time. Therefore, the **mmchqos** command rejects IOPS values less than 100IOPS with an error message, unless you specify the **--force** option.

**--fine-stats Seconds**

Specifies how many seconds of fine-grained statistics to save in memory for the **mmlsqos** command to display. The default value is 0, which indicates that no fine-grained statistics are saved. The valid range is 0 - 3840 seconds.

**Note:** The value that you specify for *Seconds* is mapped to a larger actual value. The **mmlsqos** command reports the actual value.

Fine-grained statistics are taken at one-second intervals and contain more information than regular statistics. For more information, see the topic “**mmlsqos** command” on page 430.

**--pid-stats {yes|no}**

Enables or disables the keeping of fine-grained statistics for each QoS program that is active on each node of the cluster. The default value is **no**. This parameter is effective only if **--fine-stats** is not zero.

When this parameter is disabled, statistics reflect the combined data of all the QoS processes running on each node.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies nodes in the local cluster to which the IOPS are to be allocated. QoS divides the allocated IOPS equally among the specified nodes that have mounted the file system.

**a11**

All the nodes in the cluster where the command is entered.

**Node[,Node...]**

The specified nodes.

**NodeFile**

A file that contains a list of nodes.

**NodeClass**

The nodes in the specified class.

## **mmchqos**

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

**-C {all | all\_remote | ClusterName[,ClusterName...]}**

Specifies clusters to which the IOPS are to be allocated. For each cluster, the IOPS are divided equally among the nodes that have mounted the file system.

**all**

All clusters, including both the local cluster (the cluster where the command is entered) and all remote clusters.

**all\_remote**

All clusters other than the one where the command was entered.

*ClusterName[,ClusterName...]*

The specified clusters.

If neither the **-N** nor the **-C** parameter is specified, QoS divides the IOPS as described in the last row of Table 11.

*Table 11. Allocation of IOPS*

| Option                                                          | Allocation of IOPS                                                                                                                                                                                                                                                                                                                                                                                                                                |
|-----------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>-N {all   Node[Node...]   NodeFile   NodeClass}</b>          | The IOPS are divided equally among the specified nodes that have mounted the file system.                                                                                                                                                                                                                                                                                                                                                         |
| <b>-C {all   all_remote   ClusterName[,ClusterName...]}</b>     | For each cluster, the IOPS are divided equally among the nodes that have mounted the file system.                                                                                                                                                                                                                                                                                                                                                 |
| Neither the <b>-N</b> nor the <b>-C</b> parameter is specified. | <ul style="list-style-type: none"><li>For the <b>maintenance</b> class: The IOPS are divided equally among all the nodes in the local cluster that have mounted the file system.</li><li><b>other</b> class: The IOPS are divided among all the nodes in the local cluster that have mounted the file system. Also, for each remote cluster, the IOPS are divided among all the nodes in the cluster that have mounted the file system.</li></ul> |

**pool=PoolName**

Specifies a storage pool to whose **maintenance** or **other** class the IOPS are to be allocated. If you specify an asterisk (\*) as the pool name, then the IOPS are allocated to the QoS classes of the unspecified pools. The *unspecified pools* are storage pools that you have not specified by name in any previous **mmchqos** command.

**QOSClass={nnnIOPS | unlimited}**

Identifies a QoS class and allocates IOPS to it.

**QOSClass**

The QoS class to which IOPS are assigned. You can specify one of the following classes:

**maintenance**

Most I/O-intensive, potentially slow-running GPFS administration commands run in this class by default. See the list of commands that support QoS in Table 12 on page 209.

Typically, you allocate fewer IOPS to this QoS class so that the commands that belong to it do not reduce overall file system performance.

When you start one of these commands, you can explicitly assign it to either QoS class. The assignment is effective only for the instance of the command that you are starting. In certain situations, you might assign one of these commands to the **other** class so that it runs faster and completes sooner.

**other** All other processes that use I/O run in this class by default. Typically you assign more IOPS or **unlimited** to this class so that normal processes have greater access to I/O resources and finish more quickly.

Some I/O-intensive, potentially slow-running GPFS administration commands run in this class by default. (Currently just one: **mmchdisk**.) See the list of commands that support QoS in Table 12. When you start one of these commands, you can explicitly assign it to either QoS class. The assignment is effective only for the instance of the command that you are starting. In certain situations, you might want to assign one of these commands to the **maintenance** class so that normal processes can finish more quickly.

The following table lists the GPFS commands that support QoS and the QoS class that the command runs in by default:

Table 12. GPFS commands that support QoS

| Commands that support QoS | Default QoS class |
|---------------------------|-------------------|
| <b>mmadddisk</b>          | maintenance       |
| <b>mmapplypolicy</b>      | maintenance       |
| <b>mmbackup</b>           | maintenance       |
| <b>mmchdisk</b>           | other             |
| <b>mmcheckquota</b>       | maintenance       |
| <b>mmdefragfs</b>         | maintenance       |
| <b>mmdeldisk</b>          | maintenance       |
| <b>mmdelfileset</b>       | maintenance       |
| <b>mmdelsnapshot</b>      | maintenance       |
| <b>mmdf</b>               | maintenance       |
| <b>mmfileid</b>           | maintenance       |
| <b>mmfsck</b>             | maintenance       |
| <b>mmimgbackup</b>        | maintenance       |
| <b>mmimgrestore</b>       | maintenance       |
| <b>mmlssnapshot</b>       | maintenance       |
| <b>mmrestripefs</b>       | maintenance       |
| <b>mmrpldisk</b>          | maintenance       |

### nnnIOPS

You can use the following values for IOPS:

- A value in the range **0IOPS - 1999999999IOPS**. For an IOPS value less than 100, you must specify the **-force** option. Otherwise, QoS displays an error message like the following one: `maintenance=99iops is not reasonable. To insist, try --force.`

QoS divides the IOPS allocation equally among the specified nodes that have mounted the file system.

You can type nnnIOPS either with or without the trailing characters IOPS. So either of the following two examples is valid:

- (1) `maintenance=400`
- (2) `maintenance=400IOPS`

- **unlimited**: QoS does not restrict access to I/O operations.

### -F StanzaFile

Specifies a file that contains QoS stanzas that describe allocations of IOPS. QoS stanzas have the following format:

## **mmchqos**

```
%qos:
 pool=PoolName
 class=QOSClass
 iops=Value
 nodes={Node[,Node...] | NodeFile | NodeClass}
 cluster={all | all_remote | ClusterName[,ClusterName...]}
```

where:

### **%qos**

Identifies the stanza as a QoS stanza.

### **pool=PoolName**

Specifies a storage pool to whose **maintenance** or **other** class the IOPS are allocated. If you specify an asterisk (\*) as the pool name, then the IOPS are allocated to the QoS classes of unspecified pools. *Unspecified pools* are storage pools that you have not specified by name in any previous **mmchqos** command.

### **class=QOSClass**

Specifies a QoS class. It must be **maintenance** or **other**.

### **iops=Value**

Specifies the IOPS that are to be allocated to the QoS class. QoS divides the allocated IOPS among the specified nodes that have mounted the file system.

### **nodes={Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes to which the IOPS are allocated.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

### **cluster={all | all\_remote | ClusterName[,ClusterName...]}**

Specifies clusters to which the IOPS are to be allocated. For each cluster, the IOPS are divided equally among the nodes that have mounted the file system.

#### **all**

All clusters, including both the local cluster (the cluster where the command is entered) and all remote clusters.

#### **all\_remote**

All clusters other than the one where the command was entered.

#### **ClusterName[,ClusterName...]**

The specified clusters.

## **Exit status**

**0**      Successful completion.

### **Nonzero**

A failure occurred.

## **Security**

You must have root authority to run the **mmchqos** command.

The node on which you enter the command must be able to execute remote shell commands on any other administration node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

- The following command enables QoS. If you are using the **--enable** option for the first time, then QoS sets the value of all the QoS classes to **unlimited**. If not, then the QoS classes retain their current settings:

```
mmchqos fs0 --enable
```

- The following command disables QoS but does not change the current allocations of IOPS:

```
mmchqos fs0 --disable
```

- The following command enables QoS and allocates 123 IOPS to the **maintenance** class of each unspecified pool:

```
mmchqos fs0 --enable pool=*,maintenance=123IOPS
```

You could also type the command like this:

```
mmchqos fs0 --enable pool=*,maintenance=123
```

- The following command enables QoS and allocates 222 IOPS to the **maintenance** class of each unspecified pool. It also allocates 576 IOPS to the **maintenance** class of the pool **mySSDs**. You might make an allocation like this one to favor a pool of high-speed storage (**mySSDs**) that you expect to be accessed frequently. The command does not change the allocations of the QoS **other** classes:

| Storage pool          | maintenance class | other class |
|-----------------------|-------------------|-------------|
| Each unspecified pool | 222IOPS           | Unchanged   |
| mySSDs                | 567IOPS           | Unchanged   |

```
mmchqos fs0 --enable pool=*,maintenance=222IOPS pool=mySSDs,maintenance=567IOPS
```

- The following command enables QoS and allocates IOPS to the classes of the unspecified pools and three named pools:

| Storage pool          | maintenance class | other class      |
|-----------------------|-------------------|------------------|
| Each unspecified pool | 444IOPS           | 555IOPS          |
| system                | 111IOPS           | <b>unlimited</b> |
| second                | 222IOPS           | <b>unlimited</b> |
| third                 | 333IOPS           | <b>unlimited</b> |

The command is all on one line:

```
mmchqos fs0 --enable pool=*,maintenance=444IOPS,other=555iops
pool=system,maintenance=111IOPS,other=unlimited
pool=second,maintenance=222IOPS,other=unlimited
pool=third,other=unlimited,maintenance=333IOPS
```

- The following command enables QoS and allocates IOPS to the classes of three named pools:

| Storage pool          | maintenance class | other class      |
|-----------------------|-------------------|------------------|
| Each unspecified pool | Unchanged         | Unchanged        |
| system                | 111IOPS           | <b>unlimited</b> |
| second                | 222IOPS           | <b>unlimited</b> |
| third                 | 333IOPS           | <b>unlimited</b> |

The command is all on one line:

```
mmchqos fs0 --enable pool=system,maintenance=111IOPS,other=unlimited
pool=second,maintenance=222IOPS,other=unlimited
pool=third,other=unlimited,maintenance=333IOPS
```

- The following command enables QoS and allocates IOPS to both classes of the system pool. Also, because the command contains the **--reset** parameter, it sets both classes of all the other storage pools

## mmchqos

in the file system to **unlimited**. The reset affects not only any unspecified pools, but also any named pools that are not explicitly mentioned in this command.

| Storage pool                               | maintenance class | other class      |
|--------------------------------------------|-------------------|------------------|
| system                                     | 111IOPS           | <b>unlimited</b> |
| Each unspecified pool, if any              | <b>unlimited</b>  | <b>unlimited</b> |
| Each named pool, if any, other than system | <b>unlimited</b>  | <b>unlimited</b> |

The command is all on one line:

```
mmchqos fs0 --enable --reset pool=system,maintenance=111IOPS,other=unlimited
```

8. The first part of the following command assigns IOPS to the QoS classes of the unspecified pools. It assigns 222 IOPS to the **maintenance** class of each unspecified pool.

The second part of the command allocates 456 IOPS to the **other** class of the storage pool mySAN, rather than allocating to it the typical value **unlimited**. You might make an allocation like this one to a SAN controller that serves both GPFS and other systems.

| Storage pool          | maintenance class | other class      |
|-----------------------|-------------------|------------------|
| Each unspecified pool | 222IOPS           | <b>unlimited</b> |
| mySAN                 | 123IOPS           | 456IOPS          |

The command is all on one line:

```
mmchqos fs0 --enable pool=*,maintenance=222IOPS
pool=mySAN,other=456IOPS,maintenance=123IOPS
```

9. The following command allocates IOPS as they are specified in the stanza file `qos.stanza`:

```
mmchqos fs0 --enable -F /tmp/qos.stanza
```

The stanza file contains only one stanza:

```
%qos:
pool=sp1
class=maintenance
iops=800
nodes=node1,aixNodes
```

The command divides the allocated IOPS equally among the nodes node1 and the class aixNodes. Assuming that the node class contains three nodes, then 200 IOPS are assigned to each of the four nodes:

| Storage pool or node | maintenance class | other class |
|----------------------|-------------------|-------------|
| sp1 and node1        | 200 IOPS          | Unchanged   |

## See also

- “[mmlsqos command](#)” on page 430
- Setting the Quality of Service for I/O operations (QoS) in the IBM Spectrum Scale: Administration Guide*.

## Location

/usr/lpp/mmfs/bin

## mmclidecode command

- Decodes the parseable command output field.

### Synopsis

**mmclidecode** *EncodedString*

### Availability

- Available with IBM Spectrum Scale Standard Edition or higher.

### Description

- The parseable output of a command might contain a colon (:) which interferes with the field delimiter character colon (:) of the command output. Therefore, the filed output that might contain colon needs to be encoded. Since, Percent-encoding, also known as URL encoding, is used to encode the output fields, the characters in the following table are also encoded if present in the output field.

| Character | Encoded text |
|-----------|--------------|
| !         | %21          |
| #         | %23          |
| \$        | %24          |
| &         | %26          |
| '         | %27          |
| (         | %28          |
| )         | %29          |
| *         | %2A          |
| +         | %2B          |
| ,         | %2C          |
| /         | %2F          |
| :         | %3A          |
| ;         | %3B          |
| =         | %3D          |
| ?         | %3F          |
| @         | %40          |
| [         | %5B          |
| ]         | %5D          |

- The output fields can be decoded using the **mmclidecode** command.

- Note:** If you run the **mmclidecode** command on non-encoded strings, the command leaves them intact.

### Parameters

#### EncodedString

- The encoded string to decode.

### Exit status

- 0 Successful completion.

## **mmclidecode**

- | **Nonzero**
  - | A failure occurred.

### **Security**

- | The **mmclidecode** execution does not require root user.

### **Examples**

- | 1. The date is encoded due to the presence of colon (:). It can be decoded as:

```
| mmclidecode "Mon Jan 30 10%3A43%3A29 2017"
| Mon Jan 30 10:43:29 2017
```

- | 2. The encoded path is decoded as follows:

```
| mmclidecode "%2Fmnt%2Fgpfso"
| /mnt/gpfso
```

### **Location**

- | /usr/lpp/mmfs/bin

---

## mmclone command

Creates and manages file clones.

### Synopsis

**mmclone snap** *SourceFile* [*CloneParentFile*]

or

**mmclone copy** *CloneParentFile* *TargetFile*

or

**mmclone split** *Filenname* [*Filenname...*]

or

**mmclone redirect** *Filenname* [*Filenname...*]

or

**mmclone show** *Filenname* [*Filenname...*]

### Availability

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

### Description

Use the **mmclone** command to create and manage file clones. Clones are writable snapshots of individual files. Cloning a file is similar to creating a copy of a file, but the creation process is faster and more space efficient because no additional disk space is consumed until the clone or the original file is modified. The keyword specified after **mmclone** determines which action is performed:

#### **snap**

Creates a read-only snapshot of an existing file for the purpose of cloning. This read-only snapshot becomes known as the clone parent.

If only one file is specified with the **mmclone snap** command, it will convert that file to a clone parent without creating a separate clone parent file. When using this method to create a clone parent, the specified file cannot be open for writing or have hard links.

#### **copy**

Creates a file clone from a clone parent created with the **mmclone snap** command or from a file in a snapshot.

#### **split**

Splits a file clone from all clone parents.

#### **redirect**

Splits a file clone from the immediate clone parent only.

#### **show**

Displays the current status for one or more specified files. When a file is a clone, the report will show the parent inode number. When a file was cloned from a file in a snapshot, **mmclone show** displays the snapshot and fileset information.

The Depth field in the **mmclone show** output denotes the distance of the file from the root of the clone tree of which it is a member. The root of a clone tree has depth 0. This field is blank if the file in question is not a clone. This field is not updated when a clone's ancestor is redirected or split from the clone tree. However, even if a clone's ancestor has been split or redirected, the depth of the clone should always be greater than that of each of its ancestors.

## **mmclone**

The maximum depth for a clone tree is 1000.

**Note:** The **mmclone** command does not copy extended attributes.

If a snapshot has file clones, those file clones should be deleted or split from their clone parents prior to deleting the snapshot. Use the **mmclone split** or **mmclone redirect** command to split file clones. Use a regular delete (**rm**) command to delete a file clone. If a snapshot is deleted that contains a clone parent, any attempts to read a block that refers to the missing snapshot will return an error. A policy file can be created to help determine if a snapshot has file clones.

For more information about file clones and policy files, see the *IBM Spectrum Scale: Administration Guide*.

### **Parameters**

#### *SourceFile*

Specifies the name of a file to clone.

#### *CloneParentFile*

When *CloneParentFile* is specified with a **mmclone snap** command, it indicates the name of the read-only clone parent that will be created from *SourceFile*.

When *CloneParentFile* is specified with a **mmclone copy** command, it indicates the name of a read-only clone parent. The *CloneParentFile* can be a clone parent created with the **mmclone snap** command or a file in a snapshot.

#### *TargetFile*

Specifies the name of the writable file clone that will be created from *CloneParentFile*.

#### *Filename*

Specifies the name of one or more files to **split**, **redirect**, or **show**.

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

To run the **mmclone** command, you must have read access to the source file that will be cloned, and write access to the directory where the file clone will be created.

### **Examples**

1. To create a clone parent called **base.img** from a file called **test01.img**, issue this command:

```
mmclone snap test01.img base.img
```

To use this clone parent to create a file clone called **test02.img**, issue this command:

```
mmclone copy base.img test02.img
```

After the file clone is created, use the **mmclone show** command to show information about all **img** files in the current directory:

```
mmclone show *.img
```

The system displays output similar to the following:

| Parent | Depth | Parent inode | File name  |
|--------|-------|--------------|------------|
| yes    | 0     |              | base.img   |
| no     | 1     | 148488       | test01.img |
| no     | 1     | 148488       | test02.img |

2. To create a file clone called **file1.clone** from a file called **file1** in the **snap1** snapshot, issue this command:

```
mmclone copy /fs1/.snapshots/snap1/file1 file1.clone
```

## See also

- “**mmcsnapshot command**” on page 267
- “**mmdelsnapshot command**” on page 304

## Location

/usr/lpp/mmfs/bin

## mmcloudgateway

### mmcloudgateway command

Creates and manages the cloud storage tier.

#### Synopsis

```
mmcloudgateway account pre-test --cloud-type {S3 | SWIFT | SWIFT-K EYSTONE | SWIFT3
| CLEVERSAFE | CLEVERSAFE-NEW}
[--container-name ContainerName] {--username UserName
[--pwd-file PasswordFile] |
--src-keystore-path SourceKeystorePath
--src-keystore-alias-name SourceKeystoreAliasName
--src-keystore-type SourceKeystoreType
[--src-keystore-pwd-file SourceKeystorePasswordFile]}
[--cloud-url CloudURL] [--tenant-id TenantID]
[--location Location] [--object-size ObjectSize]
[--server-cert-path ServerCertPath]
```

or

```
mmcloudgateway account create --cloud-nodeclass CloudNodeClass --cloud-name CloudName
--cloud-type {S3 | SWIFT
| SWIFT-KEYSTONE | SWIFT3 | CLEVERSAFE | CLEVERSAFE-NEW}
{--username UserName [--pwd-file PasswordFile] |
--src-keystore-path SourceKeystorePath
--src-keystore-alias-name SourceKeystoreAliasName
--src-keystore-type SourceKeystoreType
[--src-keystore-pwd-file SourceKeystorePasswordFile]}
--enable {TRUE | FALSE} [--cloud-url CloudURL]
[--tenant-id TenantID] [--location Location]
[--meta-location MetaLocation]
[--mpu-parts-size MPUPartsSize]
[--server-cert-path ServerCertPath]
[--enc-enable {TRUE | FALSE }]
[--etag-enable {TRUE | FALSE }]
```

or

```
mmcloudgateway account test --cloud-nodeclass CloudNodeClass --cloud-name CloudName
```

or

```
mmcloudgateway account update --cloud-nodeclass CloudNodeClass --cloud-name CloudName
{[--pwd-file PasswordFile |
--src-keystore-path SourceKeystorePath
--src-keystore-alias-name SourceKeystoreAliasName
--src-keystore-type SourceKeystoreType
--src-keystore-pwd-file SourceKeystorePasswordFile]} |
[--mpu-parts-size MPUPartsSize]
```

or

```
mmcloudgateway account delete --cloud-nodeclass CloudNodeClass --cloud-name CloudName
```

or

```
| mmcloudgateway account list [--cloud-nodeclass CloudNodeClass] [--cloud-name CloudName] [-Y]
```

or

```
mmcloudgateway config set --cloud-nodeclass CloudNodeClass
[--port Port]
[--slice-size SliceSize]
[--migrate-threadpool-size MigrateThreadpoolSize]
[--recall-threadpool-size RecallThreadpoolSize]
[--tracing-enable {TRUE | FALSE }]
[--tracing-level {comp=level [,comp=level ...]}]
[--audit-enable { TRUE | FALSE }] [--rotate-key]
```

```

| [--proxy-ip ProxyIP --proxy-port ProxyPort]
| [--cloud-retention-period-days CloudRetentionPeriodDays]
| [--rkm-enable { TRUE --rkm-servername RKMServerName
| --rkm-port RKMPort --rkm-username RKMUserName
| [--pwd-file PasswordFile] | FALSE}]

or

| mmcloudgateway config unset --cloud-nodeclass CloudNodeClass
| {[--slice-size] | [--migrate-threadpool-size] |
| [--recall-threadpool-size] | [
| [--proxy-ip] | [--proxy-port] | [--cloud-retention-period-days]
| [--tracing-enable] | [
| [--tracing-level] | [--audit-enable]}]

or

| mmcloudgateway config list --cloud-nodeclass CloudNodeClass [-Y]

or

| mmcloudgateway filesystem create --cloud-nodeclass CloudNodeClass
| --file-system FileSystem --container-prefix
| ContainerPrefix {[--override-container-name]
| [--thumbnail-size ThumbnailSize]}

or

| mmcloudgateway filesystem delete --cloud-nodeclass CloudNodeClass --file-system FileSystem

or

| mmcloudgateway filesystem list --cloud-nodeclass CloudNodeClass [-Y]

or

| mmcloudgateway node list [-Y]

or

| mmcloudgateway service start [-N { Node [,Node ...] | NodeFile | NodeClass }]

or

| mmcloudgateway service stop [-N { Node [,Node ...] | NodeFile | NodeClass }]

or

| mmcloudgateway service status [-N { Node [,Node ...] | NodeFile | NodeClass }] [-Y]

or

| mmcloudgateway service version [-N { Node [,Node ...] | NodeFile | NodeClass }] [-Y]

or

| mmcloudgateway files migrate [-v] [--co-resident-state] [--] File [File ...]

or

| mmcloudgateway files recall [-v] [--] File [File ...]

or

| mmcloudgateway files restore [-v] [--overwrite] [--restore-stubs-only]
| { -F FileListFile | [--dry-run]
| [--restore-location RestoreLocation]
| [--id ID] [--] File}
```

## **mmcloudgateway**

or

```
mmcloudgateway files delete {-delete-local-file | -recall-cloud-file | --require-local-file}
[--keep-last-cloud-file] [--] File [File ...]
```

or

```
| mmcloudgateway files destroy
| [--cloud-retention-period-days CloudRetentionPeriodDays [--preview]]
| [--timeout Minutes] --filesystem-path FileSystemPath
```

or

```
mmcloudgateway files reconcile Device
```

or

```
mmcloudgateway files cloudList {--path Path [--recursive [--depth Depth]] [--file File] |
[--file-versions File] | --files-usage --path Path
--depth Depth] | --reconcile-status --path Path}
```

or

```
mmcloudgateway files rebuildDB Device
```

or

```
mmcloudgateway files list File [File ...]
```

or

```
mmcloudgateway files import
[--container Container] [--import-only-stub]
[--import-metadata]
{[--directory Directory] |
[--directory-root DirectoryRoot] |
[--target-name TargetName] } File [File]
```

or

```
mmcloudgateway files export [--tag Tag]
[--target-name TargetName]
[--container Container]
[--manifest-file ManifestFile]
[--export-metadata [--fail-if-metadata-too-big]]
[--strip-filesystem-root] File [File]
```

## **Availability**

Available with IBM Spectrum Scale Advanced Edition or IBM Spectrum Scale Data Management Edition.

## **Description**

Use the **mmcloudgateway** command to manage and administer the Transparent cloud tiering feature.

## **Parameters**

### **account**

Manages cloud storage accounts with one of the following actions:

### **pre-test**

Verifies that the cloud storage account settings are correct before the account is actually created. If no errors, you can go ahead and create a cloud storage.

**--cloud-type CloudType**

Specifies the name of the object storage provider. For a list of supported cloud providers, see the *Supported cloud providers* topic in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

**Note:** Incorrect cloud type affects migration of large files that are transferred by using multipart upload (MPU) technology.

**--container containerName**

When you specify this, pre-test will use the specified container name instead of creating a new container (or vault or a bucket). If you do not specify this, pre-test will attempt to create a container or a vault on the object storage system, and it will be a regular vault.

**Note:** The container names should only have lower case characters. It should not have a combination of upper and lower cases. It is also recommended that the names comply with DNS naming conventions.

**--username UserName**

Specifies the user name of the cloud object storage account. For Amazon S3 and IBM Cloud Object Storage, it represents the access key.

**Note:** Skip this parameter for locked vaults.

**--pwd-file PasswordFile**

Specifies a file that includes the password. For Amazon S3 and IBM Cloud Object Storage, password represents the secret access key.

**Note:** Skip this parameter for locked vaults.

**--cloud-url CloudURL**

Specifies the URL of the cloud storage provider.

**Note:** Optional for S3, but mandatory for Swift, Swift-Keystone, and IBM Cloud Object Storage. For object storage providers that require an end-point URL, the end-point URL need not have a container/vault name.

**--tenant-id Tenant ID**

Specifies the tenant ID for the cloud storage provider account.

**Note:** Optional for cloud type “S3” but mandatory for cloud types “Swift” and “Swift-Keystone”

**--location Location**

Specifies the preferred location of the object cloud storage provider. The valid values for S3 are the following:

- us-west-1
- us-west-2
- EU or eu-west-1
- eu-central-1
- apsoutheast-1
- ap-southeast-2
- ap-northeast-1
- ap-northeast-2
- sa-east-1

## **mmcloudgateway**

For IBM Cloud Object Storage, use the provisioning code for vault provisioning template. Do not use the name of the provisioning template. If a location is not used to provide provisioning code, then the default provisioning template is used.

### **--server-cert-path ServerCertPath**

Specifies the certificate path for the self-signed certificates that are presented by the private object storage servers. This is required only when the cloud URL uses https.

### **--object size ObjectSize**

Size of the objects to be transferred to the cloud for the pre-test.

### **--src-keystore-path SourceKeystorePath**

Specifies the keystore path for the X.509 certificate and the private key that is used to authenticate the vault.

### **--src-keystore-alias-name SourceKeystoreAliasName**

Specifies the alias name that needs to be imported to the Transparent cloud tiering keystore from the specified keystore.

### **--src-keystore-type SourceKeystoreType**

Specifies the type of source keystore. Allowed keystore types are JKS, JCEKS.

### **--src-keystore-pwd-file SourceKeystorePasswordField**

Specifies the password file of the source keystore.

**Note:** If you do not supply this option with the keystore options, you will be prompted for the password.

### **create**

Creates a cloud account definition to access the remote cloud.

#### **--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

#### **--cloud-name CloudName**

Specifies a name that uniquely identifies the cloud object storage account on the node.

**Note:** No special characters are allowed in the name except for “-” and “\_”.

#### **--cloud-type CloudType**

Specifies the name of the object storage provider.

#### **--username Name**

Specifies the user name of the cloud object storage account. For Amazon S3 and IBM Cloud Object Storage, it represents the access key.

**Note:** Skip this parameter for locked vaults.

#### **--pwd-file PasswordFile**

Specifies a file that includes the password.

**Note:** Skip this parameter for locked vaults.

#### **--enable {TRUE | FALSE}**

Enables or disables the added cloud object storage account for Transparent cloud tiering node.

**Note:** This property is always enabled in the 4.2.1 release.

#### **--cloud-url CloudURL**

Specifies the URL of the cloud storage provider. You can either specify host name or IP address of the cloud storage provider server.

**Note:** Optional for S3, but mandatory for Swift, Swift-Keystone, and IBM Cloud Object Storage. For object storage providers that require an end-point URL, the end-point URL need not have a container/vault name.

**--tenant-id Tenant ID**

Specifies the tenant ID for the cloud storage provider account.

**Note:** Optional for cloud type “S3” but mandatory for cloud types “Swift” and “Swift-Keystone.”

**--location Location**

Specifies the preferred location of the object cloud storage provider. The valid values for S3 are the following:

- us-west-1
- us-west-2
- EU or eu-west-1
- eu-central-1
- apsoutheast-1
- ap-southeast-2
- ap-northeast-1
- ap-northeast-2
- sa-east-1

For IBM Cloud Object Storage, use the provisioning code for vault provisioning template. Do not use the name of the provisioning template. If a location is not used to provide provisioning code, then the default provisioning template is used.

**--meta-location MetaLocation**

For IBM Cloud Object Storage, two containers are created on the cloud storage tier, one for data and one for metadata. You can use this attribute to specify different vault provisioning codes for metadata containers. Data container should have indexing disabled, whereas metadata container should have indexing enabled.

**--mpu-parts-size MPUPartsSize**

Specifies multi-part upload size in MB. Value that is allowed is 5 - 32, default being 32.

**--server-cert-path ServerCertPath**

Specifies the certificate path for the self-signed certificates that are presented by the private object storage servers. This is required only when the cloud URL uses https.

**--enc-enable [TRUE | FALSE]**

Specifies whether you want to enable encryption on the data that is transferred to the object storage.

**--etag-enable [TRUE | FALSE]**

Specifies whether you want to enable an integrity check on the data that is migrated or recalled to or from the cloud storage.

**--src-keystore-path SourceKeystorePath**

Specifies the keystore path for the X.509 certificate and the private key that is used to authenticate the vault.

**--src-keystore-alias-name SourceKeystoreAliasName**

Specifies the alias name that needs to be imported to the Transparent cloud tiering keystore from the specified keystore.

**--src-keystore-type SourceKeystoreType**

Specifies the type of source keystore. Allowed keystore types are JKS, JCEKS.

## **mmcloudgateway**

```
| --src-keystore-pwd-file SourceKeystorePasswordFile
| Specifies the password file of the source keystore.

test
 Verifies the validity of the cloud account.

--cloud-nodeclass CloudNodeClass
 Specifies the node class that was created by using the mmcrnodeclass command.

--cloud-name CloudName
 Specifies the unique name that was provided to the cloud object storage account on the
 Transparent cloud tiering node.

update
 Updates the cloud account attributes (except the cloud type) that are provided while you create a
 cloud account. To update the cloud type, you must delete the cloud account and create a new
 one.

--cloud-nodeclass CloudNodeClass
 Specifies the node class that was created by using the mmcrnodeclass command.

--cloud-name CloudName
 Specifies the unique name that was provided to the cloud object storage account on the node.

--pwd-file PasswordFile
 Specifies the file that includes the password.

Note: This only updates the password retained within the IBM Spectrum Scale system. You
must have already updated the password within the cloud account before running this step.

--mpu-parts-size MPUPartsSize
 Specifies multi-part upload size in MB. Value that is allowed is 5 - 32, default being 32.

--src-keystore-path SourceKeystorePath
 Specifies the keystore path for the X.509 certificate and the private key that is used to
 authenticate the vault.

--src-keystore-alias-name SourceKeystoreAliasName
 Specifies the alias name that needs to be imported to the Transparent cloud tiering keystore
 from the specified keystore.

--src-keystore-type SourceKeystoreType
 Specifies the type of source keystore. Allowed keystore types are JKS, JCEKS.

--src-keystore-pwd-file SourceKeystorePasswordFile
 Specifies the password file of the source keystore.

delete
 Deletes the cloud storage tier that is created by using the mmcloudgateway account create
 command.

--cloud-nodeclass CloudNodeClass
 Specifies the node class that was created by using the mmcrnodeclass command.

--cloud-name CloudName
 Specifies the cloud account that you want to delete.

list
 Lists the registered cloud accounts. Displays more information about the configured cloud
 account such as the cloud provider name, cloud provider tenant ID, cloud provider URL.

--cloud-nodeclass CloudNodeClass
 Specifies the node class that was created by using the mmcrnodeclass command.
```

**--cloud-name CloudName**

Specifies the cloud account that you want to delete.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**config**

Configures and tunes the Transparent cloud tiering node parameters with one of the following actions:

**set**

Sets the following system parameters, overriding the default values:

**--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

**--port Port**

Specifies the port the Transparent cloud tiering service listens on incoming configuration and data migration/recall commands.

**--slice-size SliceSize**

Specifies the internal unit of transferring data within Transparent cloud tiering modules. Higher slice size indicates better performance. Default value is 256 KB.

**--migrate-threadpool-size MigrateThreadpoolSize**

Enables administrators to increase the parallelism for data migration within the Transparent cloud tiering service. Higher threadpool-size indicates better performance. Its value can be within 1 to 32, with default being 16.

**--recall-threadpool-size RecallThreadpoolSize**

Enables administrators to increase the parallelism for data migration within the Transparent cloud tiering service. Higher threadpool-size indicates better performance. Its value can be within 1 to 32, with default being 16.

**--tracing-enable {TRUE | FALSE}**

Tracing level is to set non-default tracing levels for various Transparent cloud tiering internal components to generate more debug data if any problems occur.

**--audit-enable {TRUE | FALSE}**

Enables administrator to change the default audit behavior to be able to record important events within the Transparent cloud tiering service.

**--rotate-key**

Enables administrator to generate a new key as and when needed according to the security requirements.

**--rkm-enable**

Enables the IBM Security Key Lifecycle Manager.

**--rkm-servername RKMServerName**

Specifies the host name or IP address of the IBM Security Lifecycle Manager server.

**--rkm-port RKMPort**

Specifies the port number on which the IBM Security Key Lifecycle Manager server listens for requests. Default value is 9080.

**--rkm-username RKMUserName**

Specifies the user name of the IBM Security Key Lifecycle Manager server REST Global Admin. Default value is SKLMAdmin.

## **mmcloudgateway**

```
--pwd-file PasswordFile
 Specifies the password file of the IBM Security Key Lifecycle Manager server REST Global Admin.

--proxy-ip ProxyIP
 If you require a proxy server to access your cloud over the network, specify the proxy server IP address in dotted decimal format. You must specify the --proxy-port option when enabling a proxy server.

 Note: This supports only IPv4.

--proxy-port ProxyPort
 Specifies the port number on which the proxy server listens for requests.

--cloud-retention-period-days CloudRetentionPeriodDays
 Specifies the number of days for which the deleted files from the file system need to be retained on the cloud. The number of days specified here will override the default value (30 days) in the policy.

unset
 Unsets the configured attributes and retains the default values.

--cloud-nodeclass CloudNodeClass
 Specifies the node class that was created by using the mmcrnodeclass command.

--slice-size SliceSize
 Specifies the internal unit of transferring data within modules. Higher slice size indicates better performance. Default value is 256 KB.

--recall-threadpool-size
 Enables administrators to increase the parallelism for data recall within the service. Higher threadpool-size indicates better performance. Its value can be within 1 to 32, with default being 16.

--proxy-ip ProxyIP
 Specifies the proxy server IP address that you want to remove.

 Note: The proxy server settings will be in effect until the Transparent cloud tiering servers are restarted. You can specify one of the attributes (either IP address or port number) to unset the proxy settings completely.

--proxy-port ProxyPort
 Specifies the port number of the proxy server that you want to remove.

--cloud-retention-period-days
 Specifies to unset this parameter.

--tracing-enable
 Tracing level is to set non-default tracing levels for various internal components to generate more debug data if any problems occur.

--tracing-level
 Enables administrators to print trace messages of the internal components in a file.

--audit-enable
 Enables administrator to change the default audit behavior to be able to record important events within the service.

list
 Lists the current configurations such as IP address, port number, thread-pool size, tracing level, slice size.
```

**--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**filesystem**

Enables the mapping of a file system to a node class.

**create**

Specifies the node class and a container name.

**--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

**--file-system FileSystem**

Specifies the file system.

**--container-prefix ContainerPrefix**

Specifies the cloud container.

**--override-container-name**

Overrides the container name and make it unique by attaching the file system ID to the container prefix.

**Note:** This option should not be used for locked vaults. If you use this option, a regular container will be created.

**--thumbnail-size ThumbnailSize**

Specifies the number of bytes that Transparent cloud tiering must store on the local file system for displaying thumbnail of files that are migrated to a storage tier. The value that you specify is applicable to each file in the file system that is managed by Transparent cloud tiering. Valid range is 1 - 1048576 bytes (1 MB). If you specify a value that is lower than the file system block size, then the file system block size is used. For example, if you specify a value of 128 KB and the file system block size is 256 KB, then 256 KB data of each file is stored locally and used for thumbnail.

After the thumbnail-size parameter is enabled with a filesystem, you can verify the setting by using the **mmcloudgateway filesystem list** command, and additionally, the thumbnails are displayed when you browse the files on Windows Explorer or a similar tool.

**Note:** Thumbnail is disabled by default. If you do not specify a valid value, then thumbnail is not enabled for the file system. You should judiciously make a decision to enable this option, as once enabled, you cannot disable this.

**delete**

Deletes the cloud file system association to the node class.

**--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

**--file-system FileSystem**

Specifies the file system that is attached to the cloud storage tier.

**list**

Lists the table of file systems associated with the specified node class.

**--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

## **mmcloudgateway**

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### **node**

Enables administrators to manage registration of Transparent cloud tiering nodes within a cluster and also display the node class the nodes are part of.

#### **list**

Lists the Transparent cloud tiering nodes that are registered to the cluster.

##### **--cloud-nodeclass CloudNodeClass**

Specifies the node class that was created by using the **mmcrnodeclass** command.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### **service**

Manages the Transparent cloud tiering service with these options:

#### **start**

Starts the Transparent cloud tiering service for a node or set of nodes and make the service available for file movement.

**-N** Specifies the nodes.

##### **Node[,Node...]**

Specifies the list of nodes where the service needs to be started.

##### **NodeFile**

Specifies a file, containing the list of nodes where the service needs to be started.

##### **NodeClass**

Specifies the node class.

#### **stop**

Stops the Transparent cloud tiering service for a node or set of nodes.

**-N** Specifies the nodes.

##### **Node[,Node...]**

Specifies the list of nodes where the service needs to be stopped.

##### **NodeFile**

Specifies a file, containing the list of nodes where the service needs to be started.

##### **NodeClass**

Specifies the node class.

#### **status**

Displays detailed status of the Transparent cloud tiering service including running state of the daemon service, cloud account name, and its connectivity status. For more information on various statuses that are associated with the Transparent cloud tiering service, see the "Transparent cloud tiering service status description" topic in the *IBM Spectrum Scale: Problem Determination Guide*.

**-N** Specifies the nodes.

**Node[,Node...]**

Specifies the list of nodes where the status of the service needs to be checked.

**NodeFile**

Specifies a file, containing the list of nodes where the status of the service needs to be started.

**NodeClass**

Specifies the node class.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**version**

Displays the Transparent cloud tiering version number associated with each node in a node class.

| **Note:** This command can run on all nodes in the cluster and will display if a node is a cloud node or not. If it is, it will show more column-based details. This makes this command different from others since normally the commands act only on cloud nodes for report generation.

- N Use this option before specifying any node or node class.

**Node[,Node...]**

Specifies the list of nodes for which the versions need to be checked.

**NodeFile**

Specifies a file, containing the list of nodes for which the versions need to be started.

**NodeClass**

Specifies the node class whose version is displayed at a cluster level.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**files**

Manages the resident and co-resident files on the cloud tier, with the following options:

**migrate**

Migrates the specified files or file sets to the cloud storage tier.

- v Specifies the verbose message.

**--co-resident-state**

Indicates that the local files will remain local and cloud resident after migration to the cloud. You can open such files from the local file system without recalling them from the cloud storage tier.

**--File [File ...]**

Specifies multiple files that need to be migrated to the cloud storage tier. This parameter must be a complete file name. It cannot be a fragment of a file name and it cannot be a path.

**recall**

Recalls the specified files or file sets from the cloud storage tier.

- v Specifies the verbose message.

## **mmcloudgateway**

### **--File [File ...]**

Specifies multiple files that need to be recalled from the storage tier. This parameter must be a complete file name. It cannot be a fragment of a file name and it cannot be a path.

### **restore**

Restores a file or list of files from the cloud storage tier when the local files are lost. The files to be restored along with their options can be either specified at the command line, or in a separate file provided by the *-F* option. For more information, see the *Restoring files* topic in the *IBM Spectrum Scale: Administration Guide*.

#### **-v**

Specifies the verbose message.

#### **--overwrite**

Overwrite the files if needed. If this option is not set, files will not be overwritten, and the files that are retrieved from the cloud will remain in temporary locations.

#### **--restore-stubs-only**

Restores only the file stubs.

#### **-F** Loads file arguments from the given filename.

#### **--dry-run**

Queries the local database and prints what would have been sent to the server. Does not contact the server. This is intended for debugging.

#### **--restore-location RestoreLocation**

Specifies the target location of the files to be restored.

#### **--id Id**

Specifies the version ID of a file if the file has multiple versions.

### **File**

Specifies the files to be restored.

### **delete**

Deletes the specified files or file sets.

#### **--delete-local-file**

Deletes the local files and the corresponding cloud objects.

#### **--recall-cloud-file**

Recalls the files from the cloud before they are deleted on the cloud. The status of local files becomes resident after the operation.

#### **--require-local-file**

Removes the extended attributes from a co-resident file and makes it resident, without deleting the corresponding cloud objects. The option requires the file data to be present on the file system and will not work on a non-resident file.

#### **--keep-last-cloud-file**

This option deletes all the versions of the file except the last one from the cloud. For example, if a file has three versions on the cloud, then versions 1 and 2 are deleted and version 3 is retained.

#### **--File [File ...]**

Specifies multiple files. This parameter must be a complete file name. It cannot be a fragment of a file name and it cannot be a path.

### **destroy**

Manually cleans up the cloud objects of the deleted files before the retention period expires.

#### **--preview**

Displays how many objects will be cleaned up and how much space will be reclaimed.

```

| --cloud-retention-period-days
| Specifies the number of days for which the deleted files from the file system need to be
| retained on the cloud.
|
| --timeout
| Specifies the duration of time (in minutes) for which the command should run. If this value is
| not specified, the command will run until all the specified files are deleted.
|
| --filesystem-path
| The path to the file system where the files are migrated from.

reconcile
 Reconciles files between your file system and the cloud storage tier.

 Device
 Specifies the device name associated with the file system.

 cloudList
 Lists the files on the cloud.

 --path Path
 Lists files and directories under the specified path.

 --recursive
 List all files in all directories under the current directory.

 --depth Depth
 List directories up to the specified depth under the specified path. Default is to list up to the
 full depth. Specify 0 to list only the current directory.

 --file [File]
 Specifies the names of the files that need to be listed. This parameter must be a complete file
 name. It cannot be a fragment of a file name and it cannot be a path.

 --file-versions File
 Displays information about all versions of the files specified by the full path.

 --files-usage --path Path
 Displays cloud data and metadata space usage under the specified path.

 --reconcile-status --path Path
 Displays the progress of the reconcile operation.

rebuildDB
 Rebuilds the database manually.

 Device
 Specifies the device name associated with the file system whose database is corrupted and
 which is in need of manual recovery.

list
 Lists the files and the associated states.

 --File [File ...]
 Specifies the names of the files that need to be listed. This parameter must be a complete file
 name. It cannot be a fragment of a file name and it cannot be a path.

import
 Imports data from a storage server.

 --container Container
 Specifies the name of the cloud container to import from. If no container option is specified,
 the default configured container name is used.

```

## **mmcloudgateway**

### **--import-only-stub**

Creates only a stub file, data in the file is not imported.

### **--import-metadata**

Attempts to restore metadata of the file from the cloud object. The data is in IBM Spectrum Scale format. The cloud object must have been exported using the **--export-metadata** option. If metadata is not attached to the cloud object, this option has no effect.

### **--directory**

Imports files into the given directory using only the filename from the cloud. Mutually exclusive with the **--target-name** and **--directory-root** options.

### **--directory-root**

Imports files starting at the given directory, keeping the cloud naming hierarchy intact. Mutually exclusive with the **--directory** and **--target-name** options.

### **--target-name**

Imports a single file from the cloud to the specified target name. Mutually exclusive with the **--directory** and **--directory-root** options.

## **export**

Exports files to the cloud.

### **--tag Tag**

Specifies an optional identifier to associate with the files. This ID will be stored in the manifest file if one is specified.

### **--target-name TargetName**

Export a single file to the cloud to the specified target name.

### **--container Container**

Specifies the name of the cloud container to export to. If no container option is specified, the default configured container name will be used.

### **--manifest-file ManifestFile**

Specifies a manifest file that will contain an entry for each file exported to the cloud. Entries will be in the CSV format of : Tag, Container, TimeStamp of Blob on cloud, or file name.

### **--export-metadata**

Attempts to attach a file's metadata to the cloud object. This is IBM Spectrum Scale specific data and format, and it contains the user-defined attributes, ACLs, etc. A file exported with this option can be fully restored by the corresponding import command. This metadata is stored in the blob metadata, and as such there is limited space available, and the metadata might not be written if it is too large.

### **--fail-if-metadata-too-big**

If the metadata of the file is very large, it causes the entire export to fail. Valid only with the **--export-metadata** option.

### **--strip-filesystem-root**

Removes the root of the IBM Spectrum Scale file system from the name as stored on the cloud. This could be used to export /filesystem1/dir/file and then import that file into a differently named file system root directory.

## **Security**

You must have root authority to run the **mmcloudgateway** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

- To pre-validate the account setting for the Cleversafe®-new account, issue the following command:

```
mmcloudgateway account pre-test --cloud-type cleversafe-new --username
9r38AFy2FK1Fs0I3D0cg --cloud-url http://9.11.10.38
```

The system displays an output similar to this:

```
Please enter a password:
```

```
Confirm your password:
```

```
Validation under progress. This may take a while...
```

```
Cloud object storage provider validation is under progress. Please wait for results to publish..
```

```
Cloud Provider Validator Tool Results:
```

```

```

```
Summary:
```

```
Provider compatibility with Transparent Cloud Tiering : Compatible
Configured Network Adapter Speed : 10000Mb/s
Sample data size used to calculate IOPS: 1KB
Approx. IOPS number for Puts (Single Thread) : 31.0
Approx. IOPS number for Gets (Single Thread) : 122.0
```

```

```

```
Details:
```

```
Successful Vault operations : PUT(Vault), PUT(Object), GET(Object), DELETE(Object)
Successful Object operations: PUT, GET, HEAD, DELETE
Unsuccessful Vault operations : None
Unsuccessful Object operations : None
```

```

```

```
mmcloudgateway: Command completed.
```

- To pre-validate the account settings for a locked vault that is created for deploying WORM solution, issue the following command:

```
mmcloudgateway account pre-test --cloud-type CLEVERSAFE-NEW --cloud-url
https://198.51.100.0 --server-cert-path/root/cleversafe.pk.stg.ibm.com
--src-keystore-path /root/testalias.ssl/testalias.jks --src-keystore-alias-name
testalias --src-keystore-type JKS --src-keystore-pwd-file /root/pwd
--container-name newvaultmarchnew
```

The system displays output similar to this:

```
Validation under progress. This may take a while...
```

```
Cloud object storage provider validation is under progress. Please wait for results to publish..
```

```
Cloud Provider Validator Tool Results:
```

```

```

```
Summary:
```

```
Provider compatibility with Transparent Cloud Tiering : Compatible
Configured Network Adapter Speed : 10000Mb/s
Sample data size used to calculate IOPS: 1KB
Approx. IOPS number for Puts (Single Thread) : 4.0
Approx. IOPS number for Gets (Single Thread) : 6.0
```

```

```

```
Details:
```

```
Successful Vault operations : PUT(Object), GET(Object), DELETE(Object)
Successful Object operations: PUT, GET, HEAD, DELETE
Unsuccessful Vault operations : None
Unsuccessful Object operations : None
```

```

```

```
mmcloudgateway: You can now delete the password file '/root/pwd'
```

```
mmcloudgateway: Command completed.
```

- To view the registered nodes in a node class, issue this command:

```
mmcloudgateway node list
```

The system displays output similar to this:

## **mmcloudgateway**

| Node | Cloud node name     | Cloud Node Class |
|------|---------------------|------------------|
| 8    | c35f1m4n09.gpfs.net | CloudNodesClass1 |
| 9    | c34f2n06.gpfs.net   | CloudNodesClass2 |

4. To start the Transparent cloud tiering service on the node class *TCTNodeClass1*, issue this command:

```
mmcloudgateway service start -N TCTNodeClass1
```

The system displays output similar to this:

```
mmcloudgateway: Sending the command to node c350f1u1b4.
Starting the Transparent Cloud Tiering service...
mmcloudgateway: The command completed on node c350f1u1b4.
```

*mmcloudgateway: Command completed.*

5. To verify the status of the Transparent cloud tiering service, issue this command:

```
mmcloudgateway service status -N TCTNodeClass1
```

The system displays output similar to this:

| Node | Daemon node name    | TCT Server Status | TCT Filesystem Status | TCT Account Status |
|------|---------------------|-------------------|-----------------------|--------------------|
| 9    | c34f2n06.gpfs.net   | Not Configured    | Not Configured        | Not Configured     |
| 8    | c35f1m4n09.gpfs.net | Not Configured    | Not Configured        | Not Configured     |

6. To create a file system association, issue this command:

```
mmcloudgateway filesystem create --cloud-nodeclass TCTNodeClass1
--filesystem /dev/gpfs0 --container-prefix multinode
```

The system displays output similar to this:

```
mmcloudgateway: Sending the Transparent Cloud Tiering request to the first successful server.
mmcloudgateway: This may take a while...
mmcloudgateway: Command completed successfully on cf2u18.
mmcloudgateway: Command completed.
```

7. To create a cloud storage account with IBM Cloud Object Storage version 3.7.2 and above as cloud type, issue this command:

```
mmcloudgateway account create --cloud-nodeclass TCTNodeClass1 --cloud-name tctnew
--cloud-type cleversafe-new --username "XYZ" --pwd-file PFile --enable TRUE
--cloud-url http://192.0.2.0
```

The system displays output similar to this:

```
Configured account options from TCTNodeClass1:

Cloud Provider Name: tctnew
Cloud Provider Tenant Id: admin
Cloud Provider URL: http://192.0.2.0:5000/v2.0
Cloud Provider Type: swift-keystone
Cloud Provider User Name: admin
Cloud Provider Enabled: true
Filesystem: ga
Container: gajan
Metadata Container: gajan.meta
```

8. To create a cloud storage tier with Swift-Keystone as cloud type, issue this command:

```
mmcloudgateway account create --cloud-nodeclass TCTNodeClass1 --cloud-name tct
--cloud-type SWIFT-KEYSTONE --username admin --pwd-file PFile --enable TRUE
--cloud-url http://10.11.12.13:5000/v2.0 --tenant-id admin --enc-enable TRUE
--cloud-nodeclass cloud
```

The system displays output similar to this:

```
Configured account options from TCTNodeClass1:

Cloud Provider Name: mcstore
Cloud Provider Tenant Id: admin
```

```

Cloud Provider URL: http://192.168.2.11:5000/v2.0
Cloud Provider Type: swift-keystone
Cloud Provider User Name: admin
Cloud Provider Enabled: true
Filesystem: ga
Container: gajan
Metadata Container: gajan.meta

```

- To create a cloud account for the S3 cloud type, issue a command similar to this:

```
mmcloudgateway account create --cloud-nodeclass TCTNodeClass1 --cloud-name clouptest
--cloud-type s3 --username admin --pwd-file MyFile --enable true
```

The system displays output similar to this:

```

mmcloudgateway: Sending the Transparent Cloud Tiering request to the first successful server.
mmcloudgateway: This may take a while...
mmcloudgateway: Command completed successfully on c350f2u18.
mmcloudgateway: Command completed.

```

- To create a cloud account for deploying a WORM solution by using locked vaults, issue a command like the following:

```
mmcloudgateway account create --cloud-nodeclass NodeClass1 --cloud-name clouptest
--cloud-type CLEVERSAFE-NEW --cloud-url https://198.51.100.0 --enable true
--server-cert-path /root/cleversafe.pk.stg.ibm.com --src-keystore-path
/root/test/testalias.ssl/testalias.jks --src-keystore-alias-name testalias
--src-keystore-type JKS --src-keystore-pwd-file /root/pwd/file.txt
```

The system displays output similar to this:

```
mmcloudgateway: Sending the command to the first successful node starting with c350f3u30
mmcloudgateway: This may take a while...
```

Note: Please ensure to keep a backup of the Source Key Store used to import the private key and certificates. Transparent Cloud Tiering will remove the private key and certificate from the trust store if the account delete command is run.

mmcloudgateway: Command completed successfully on c350f3u30.

mmcloudgateway: You can now delete the password file '/root/pwd/file.txt'

mmcloudgateway: Command completed.

- To update a cloud account for a locked vault, issue the following command:

```
mmcloudgateway account update --cloud-nodeclass NodeClass1 --cloud-name clouptest
--src-keystore-path /root/test/testalias.ssl/testalias.jks --src-keystore-alias-name
testaliasnew --src-keystore-type JKS --src-keystore-pwd-file /root/pwd/file.txt
```

The system displays output similar to this:

```
mmcloudgateway: Sending the command to the first successful node starting with jupiter-vm641.pok.stglabs.ibm.com
mmcloudgateway: This may take a while...
```

Note: Please ensure to keep a backup of the Source Key Store used to import the private key and certificates. Transparent Cloud Tiering will remove the private key and certificate from the trust store if the account delete command is run.

mmcloudgateway: Command completed successfully on jupiter.pk.stglabs.ibm.com.

mmcloudgateway: You can now delete the password file '/root/pwd'

mmcloudgateway: Command completed.

- To list the cloud accounts for all node classes, issue this command:

```
mmcloudgateway account list
```

The system displays output similar to this:

```
Cloud names from CloudNodeClass1:

```

```
TCT1
```

## **mmcloudgateway**

```
| Cloud names from CloudNodeClass2:
```

```
|-----|
| TCT2
```

13. To list the cloud accounts with cloud name, *mcstore* and node classes *CloudNodeClass1* and *CloudNodeClass2*, issue this command:

```
| mmcloudgateway account list --cloud-name mcstore
```

The system displays output similar to this:

```
| Configured account options from CloudNodeClass1:
```

```
|-----|
| Cloud Provider Name: mcstore
| Cloud Provider Tenant Id: admin
| Cloud Provider URL: http://192.168.2.11:5000/v2.0
| Cloud Provider Type: swift-keystone
| Cloud Provider User Name: admin
| Cloud Provider Enabled: true
| Filesystem: ga
| Container: gajan
| Metadata Container: gajan.meta
```

```
| Configured account options from CloudNodeClass2:
```

```
|-----|
| Cloud Provider Name: mcstore
| Cloud Provider Tenant Id: admin
| Cloud Provider URL: http://192.168.2.11:5000/v2.0
| Cloud Provider Type: swift-keystone
| Cloud Provider User Name: admin
| Cloud Provider Enabled: true
| Filesystem: ga
| Container: gajan
| Metadata Container: gajan.meta
```

14. To update a cloud storage with an MPU part size of 32, issue this command:

```
| mmcloudgateway account update --cloud-nodeclass cloud --cloud-name mcstore
| --mpu-parts-size 32
```

The system displays output similar to this:

```
| mmcloudgateway: Sending the Transparent Cloud Tiering request to the first sucessful server.
| mmcloudgateway: This may take a while...
| mmcloudgateway: Command completed successfully on c80f4m5n01.gpfs.net.
| mmcloudgateway: Command completed.
```

15. To verify that the cloud storage tier is active, issue this command:

```
| mmcloudgateway account test --cloud-nodeclass cloud --cloud-name mcstore
```

The system displays output similar to this:

```
| mmcloudgateway: Sending the Transparent Cloud Tiering request to the first sucessful server.
| mmcloudgateway: This may take a while...
| Cloud Status : Configured cloud account is Active!
| mmcloudgateway: Command completed successfully on vmip206.gpfs.net.
| mmcloudgateway: Command completed.
```

16. To list the current cloud configuration, issue this command:

```
| mmcloudgateway config list
```

The system displays output similar to this:

```
| audit.enable=true
| ip-address=9.47.83.202
| port=8085
| sliceSize=262144
| threadpool.poolSize=16
| tracing.enable=true
| tracing.level=ALL=4
```

- To migrate a file (file1) to the configured cloud storage tier, issue this command:

```
mmcloudgateway files migrate file1
```

The system displays output similar to this:

```
mmcloudgateway: Command completed.
```

- To migrate multiple files (file1 and file2) to the configured cloud storage tier, issue this command:

```
mmcloudgateway files migrate file1 file2
```

The system displays output similar to this:

```
mmcloudgateway: Command completed.
```

- To verify that the file is migrated to the configured cloud storage tier, issue this command:

```
mmcloudgateway files list file1
```

The system displays output similar to this:

```
File name : /gpfs/girish/file1
On-line size : 45
Used blocks : 0
Data Version : 1
Meta Version : 1
State : Non-resident
Base Name : 7448805A60ED1970.17F2AFD45704E1E4.52E20457CA532F09.0000000000000000.57B76BDD.000000000000100B
```

**Note:** The **State** is displaying as **Non-resident**. This means that the file is successfully migrated to the cloud storage tier.

- To recall a file from the configured cloud storage tier, issue this command:

```
mmcloudgateway files recall file1
```

The system displays output similar to this:

```
mmcloudgateway: Command completed.
```

**Note:** If you run the **mmcloudgateway filesystem list file1** command, the value of the **State** attribute is displayed as **Co-resident**. This means that the file is successfully recalled.

- To recall multiple files (file1 and file2) from the configured cloud storage tier, issue this command:

```
mmcloudgateway files recall file1 file2
```

The system displays output similar to this:

```
mmcloudgateway: Command completed.
```

- To delete the association between the file system and the cloud storage tier, issue this command:

```
mmcloudgateway filesystem delete --cloud-nodeclass TCTNodeClass1 --filesystem /dev/gpfs0
```

The system displays output similar to this:

```
mmcloudgateway: Sending the Transparent Cloud Tiering request to the first successful server.
```

```
mmcloudgateway: This may take a while...
```

```
mmcloudgateway: Command completed successfully on c350f2u18.
```

```
mmcloudgateway: Command completed.
```

- To delete a cloud storage account, issue this command:

```
mmcloudgateway account delete --cloud-nodeclass TCTNodeClass1 --cloud-name mycloud
```

The system displays output similar to this:

```
mmcloudgateway: Sending the Transparent Cloud Tiering request to the first successful server.
```

```
mmcloudgateway: This may take a while...
```

```
mmcloudgateway: Command completed successfully on c350f2u18.
```

```
mmcloudgateway: Command completed.
```

- To export a local file named /dir1/dir2/file1 to the cloud and store it in a container named MyContainer, issue this command:

## **mmcloudgateway**

```
mmcloudgateway files export --container MyContainer --tag MRI_Images --export-metadata
--manifest-file /dir/ManifestFile /dir1/dir2/file1
```

**Note:** A manifest file will be created, and the object exported to the cloud will have an entry in that manifest file, tagged with MRI\_Images.

25. To import files from the cloud, issue the following command(s):

```
mmcloudgateway files import --directory /localdir /dir1/dir2/file1
mmcloudgateway files import --directory-root /localdir /dir1/dir2/file1
```

**Note:** This command creates a local directory structure as necessary when importing the file from the cloud. If the **--directory** option is specified, only the file name of the cloud object is used.

26. To check the Transparent cloud tiering service version of the node class, TCTNodeClass1, issue this command:

```
mmcloudgateway service version -N TCTNodeClass1
```

The system displays output similar to this:

```
Cluster minReleaseLevel: 4.2.2.0
```

| Node | Daemon node name           | TCT Type | TCT Version | Equivalent Product Version |
|------|----------------------------|----------|-------------|----------------------------|
| 1    | c350f2u22.pk.slabs.ibm.com | Server   | 1.1.2       | 4.2.2                      |
| 3    | c350f3u30                  | Server   | 1.1.2       | 4.2.2                      |
| 2    | c350f3u6.pok.slabs.ibm.com | Server   | 1.1.2       | 4.2.2                      |

## **See also**

- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmchnode command” on page 190
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmobj command” on page 466
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

## **Location**

/usr/lpp/mmfs/bin

## mmcrcluster command

Creates a GPFS cluster from a set of nodes.

### Synopsis

```
mmcrcluster -N {NodeDesc[,NodeDesc...] | NodeFile}
 [--ccr-enable | {--ccr-disable -p PrimaryServer [-s SecondaryServer]}]
 [-r RemoteShellCommand] [-R RemoteFileCopyCommand] | --use-sudo-wrapper]
 [-C ClusterName] [-U DomainName] [-A]
 [-c ConfigFile | --profile ProfileName]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmcrcluster** command to create a GPFS cluster.

Upon successful completion of the **mmcrcluster** command, the **/var/mmfss/gen/mmsdrfs** and the **/var/mmfss/gen/mmfsNodeData** files are created on each of the nodes in the cluster. Do not delete these files under any circumstances. For more information, see *Quorum in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

Follow these rules when creating your GPFS cluster:

- While a node may mount file systems from multiple clusters, the node itself may only be added to a single cluster using the **mmcrcluster** or **mmaddnode** command.
- The nodes must be available for the command to be successful. If any of the nodes listed are not available when the command is issued, a message listing those nodes is displayed. You must correct the problem on each node and issue the **mmaddnode** command to add those nodes.
- Designate at least one but not more than seven nodes as quorum nodes. How many quorum nodes altogether you will have depends on whether you intend to use the node quorum with tiebreaker algorithm or the regular node based quorum algorithm. For more information, see *Quorum in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.
- After the nodes are added to the cluster, use the **mmchlicense** command to designate appropriate GPFS licenses to the new nodes.
- Clusters that will include both UNIX and Windows nodes must use **ssh** and **scp** for the remote shell and copy commands. For more information, see *Installing and configuring OpenSSH on Windows nodes* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.
- Carefully consider the remote execution and remote copy tooling you want to use within your cluster. Once a cluster has been created, it is complicated to change, especially if additional nodes are added. The default tools as specified under **-r RemoteShellCommand** and **-R RemoteFileCopyCommand** by default use **/usr/bin/ssh** and **/usr/bin/scp** respectively. For more information, see *GPFS cluster creation considerations* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### Parameters

**-N NodeDesc[,NodeDesc...] | NodeFile**

Specifies node descriptors, which provide information about nodes to be added to the cluster.

*NodeFile*

Specifies a file containing a list of node descriptors, one per line, to be added to the cluster.

*NodeDesc[,NodeDesc...]*

Specifies the list of nodes and node designations to be added to the GPFS cluster. Node descriptors are defined as:

## **mmcrcluster**

NodeName:NodeDesignations:AdminNodeName

where:

### *NodeName*

Specifies the host name or IP address of the node for GPFS daemon-to-daemon communication. For hosts with multiple adapters, see the *IBM Spectrum Scale: Administration Guide* and search on *Using remote access with public and private IP addresses*.

The host name or IP address must refer to the communication adapter over which the GPFS daemons communicate. Aliased interfaces are not allowed. Use the original address or a name that is resolved by the **host** command to that original address. You can specify a node using any of these forms:

- Short host name (for example, h135n01)
- Long, fully-qualified, host name (for example, h135n01.ibm.com)
- IP address (for example, 7.111.12.102). IPv6 addresses must be enclosed in brackets (for example, [2001:192::192:168:115:124]).

Regardless of which form you use, GPFS will resolve the input to a host name and an IP address and will store these in its configuration files. It is expected that those values will not change while the node belongs to the cluster.

### *NodeDesignations*

An optional, "-" separated list of node roles:

- **manager** | **client** – Indicates whether a node is part of the node pool from which file system managers and token managers can be selected. The default is **client**.
- **quorum** | **nonquorum** – Indicates whether a node is counted as a quorum node. The default is **nonquorum**.

### *AdminNodeName*

Specifies an optional field that consists of a node name to be used by the administration commands to communicate between nodes. If *AdminNodeName* is not specified, the *NodeName* value is used.

**Note:** *AdminNodeName* must be a resolvable network host name. For more information, see the topic *GPFS node adapter interface names* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

You must provide a *NodeDesc* for each node to be added to the GPFS cluster.

### **--ccr-enable**

Enables the configuration server repository (CCR), which stores redundant copies of configuration data files on all quorum nodes. All GPFS administration commands, as well as file system mounts and daemon startups, work normally as long as a majority of quorum nodes are accessible. This is the default.

The CCR operation requires the use of the GSKit toolkit for authenticating network connections. As such, the **gpfs.gskit** package, which is available on all Editions, should be installed.

### **--ccr-disable**

Indicates that the traditional primary/backup server-based configuration repository (non-CCR, earlier than GPFS 4.1) is to be used.

When using this option you must also specify a primary configuration server (**-p** option). It is suggested that you also specify a secondary GPFS cluster configuration server (**-s** option) to prevent the loss of configuration data in the event your primary GPFS cluster configuration server goes down. When the GPFS daemon starts up, at least one of the two GPFS cluster configuration servers must be accessible.

If your primary GPFS cluster configuration server fails and you have not designated a secondary server, the GPFS cluster configuration files are inaccessible, and any GPFS administration commands that are issued fail. File system mounts or daemon startups also fail if no GPFS cluster configuration server is available.

You are strongly advised to designate the cluster configuration servers as quorum nodes.

**-p PrimaryServer**

Specifies the primary GPFS cluster configuration server node used to store the GPFS configuration data. This node must be a member of the GPFS cluster. This option is necessary only when **-ccr-disable** is specified.

**-s SecondaryServer**

Specifies the secondary GPFS cluster configuration server node used to store the GPFS cluster data. This node must be a member of the GPFS cluster. This option is necessary only when **-ccr-disable** is specified.

**-r RemoteShellCommand**

Specifies the fully-qualified path name for the remote shell program to be used by GPFS. The default value is **/usr/bin/ssh**.

The remote shell command must adhere to the same syntax format as the **ssh** command, but may implement an alternate authentication mechanism.

**-R RemoteFileCopy**

Specifies the fully-qualified path name for the remote file copy program to be used by GPFS. The default value is **/usr/bin/scp**.

The remote copy command must adhere to the same syntax format as the **scp** command, but may implement an alternate authentication mechanism.

**--use-sudo-wrapper**

Specifies that the nodes in the cluster call the ssh sudo wrapper script and the scp sudo wrapper script as the remote shell program and the remote copy program. For more information, see *Running IBM Spectrum Scale without remote root login* in *IBM Spectrum Scale: Administration Guide*.

**-c ClusterName**

Specifies a name for the cluster. If the user-provided name contains dots, it is assumed to be a fully qualified domain name. Otherwise, to make the cluster name unique, the domain of the primary configuration server will be appended to the user-provided name.

If the **-C** flag is omitted, the cluster name defaults to the name of the primary GPFS cluster configuration server.

**-u DomainName**

Specifies the UID domain name for the cluster.

**-A** Specifies that GPFS daemons are to be automatically started when nodes come up. The default is not to start daemons automatically.

**-c ConfigFile**

Specifies a file containing GPFS configuration parameters with values different than the documented defaults. A sample file can be found in **/usr/lpp/mmfs/samples/mmfs.cfg.sample**. See the **mmchconfig** command for a detailed description of the different configuration parameters.

The **-c ConfigFile** parameter should be used only by experienced administrators. Use this file to set up only those parameters that appear in the **mmfs.cfg.sample** file. Changes to any other values may be ignored by GPFS. When in doubt, use the **mmchconfig** command instead.

**--profile ProfileName**

Specifies a predefined profile of attributes to be applied. System-defined profiles are located in **/usr/lpp/mmfs/profiles/**. All the configuration attributes listed under a cluster stanza will be changed as a result of this command.

## **mmcrcluster**

The following system-defined profile names are accepted:

- **gpfsProtocolDefaults**
- **gpfsProtocolRandomIO**

A user's profiles must be installed in `/var/mmfs/etc/`. The profile file specifies GPFS configuration parameters with values different than the documented defaults. A user-defined profile must not begin with the string 'gpfs' and must have the .profile suffix.

User-defined profiles consist of the following stanzas:

```
%cluster:
[CommaSeparatedNodesOrNodeClasses:]ClusterConfigurationAttribute=Value
...
%filesystem:
FilesystemConfigurationAttribute=Value
```

See the **mmchconfig** command for a detailed description of the different configuration parameters. A sample file can be found in `/usr/lpp/mmfs/samples/sample.profile`.

**Note:** User-defined profiles should be used only by experienced administrators. When in doubt, use the **mmchconfig** command instead.

### **Exit status**

**0**      Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmcrcluster** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

To create a GPFS cluster made of all of the nodes listed in the file `/u/admin/nodelist`, using node **k164n05** as the primary server, and node **k164n04** as the secondary server, issue:

```
mmcrcluster -N /u/admin/nodelist -p k164n05 -s k164n04
```

where `/u/admin/nodelist` has these contents:

```
k164n04.kgn.ibm.com:quorum
k164n05.kgn.ibm.com:quorum
k164n06.kgn.ibm.com
```

The output of the command is similar to:

```
Mon May 10 10:59:09 EDT 2010: mmcrcluster:
 Processing node k164n04.kgn.ibm.com
Mon May 10 10:59:09 EDT 2010: mmcrcluster:
 Processing node k164n05.kgn.ibm.com
Mon May 10 10:59:09 EDT 2010: mmcrcluster:
 Processing node k164n06.kgn.ibm.com
mmcrcluster: Command successfully completed
```

```
mmcrcluster: Warning: Not all nodes have proper
GPFS license designations.
Use the mmchlicense command to designate
licenses as needed.
```

To confirm the creation, issue this command:

```
mmlscluster
```

The system displays information similar to:

```
GPFS cluster information
=====
GPFS cluster name: k164n05.kgn.ibm.com
GPFS cluster id: 680681562214606028
GPFS UID domain: k164n05.kgn.ibm.com
Remote shell command: /usr/bin/ssh
Remote file copy command: /usr/bin/scp

GPFS cluster configuration servers:

Primary server: k164n05.kgn.ibm.com
Secondary server: k164n04.kgn.ibm.com

Node Daemon node name IP address Admin node name Designation

1 k164n04.kgn.ibm.com 198.117.68.68 k164n04.kgn.ibm.com quorum
2 k164n05.kgn.ibm.com 198.117.68.71 k164n05.kgn.ibm.com quorum
3 k164n06.kgn.ibm.com 198.117.68.70 k164n06.kgn.ibm.com
```

## See also

- “[mmaddnode command](#)” on page 29
- “[mmchconfig command](#)” on page 132
- “[mmdelnod command](#)” on page 297
- “[mmlscluster command](#)” on page 397
- “[mmlsconfig command](#)” on page 400

## Location

/usr/lpp/mmfs/bin

## **mmcrfileset**

---

### **mmcrfileset command**

Creates a GPFS fileset.

#### **Synopsis**

```
mmcrfileset Device FilesetName [-p afmAttribute=Value...] [-t Comment]
[--inode-space {new [--inode-limit MaxNumInodes[:NumInodesToPreallocate]] | ExistingFileset}]
[--allow-permission-change PermissionChangeMode]
```

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

The **mmcrfileset** command constructs a new fileset with the specified name. The new fileset is empty except for a root directory, and does not appear in the directory namespace until the **mmlinkfileset** command is issued. The **mmcrfileset** command is separate from the **mmlinkfileset** command to allow the administrator to establish policies and quotas on the fileset before it is linked into the namespace.

For information on filesets, see the *Filesets* section in the *IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system to contain the new fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

##### *FilesetName*

Specifies the name of the fileset to be created.

Note the following restrictions on fileset names:

- The name must be unique within the file system.
- The length of the name must be in the range 1-255.
- The name **root** is reserved for the fileset of the root directory of the file system.
- The name cannot be the reserved word *new*. However, the character string *new* can appear within a fileset name.
- The name cannot begin with a hyphen (-).
- The name cannot contain the following characters: / ? \$ & \* ( ) ^ # | [ ] \
- The name cannot contain a white-space character such as blank space or tab.

##### **-p** *afmAttribute=Value*

Specifies an AFM configuration parameter and its value. More than one **-p** option can be specified.

The following AFM configuration parameter is required for the **mmcrfileset** command:

##### **afmTarget**

Identifies the home that is associated with the cache; specified in either of the following forms:

*Protocol://[Host|Map]/Path*

or

*{Host|Map}:Path*

where:

*Protocol://*

Specifies the transport protocol. Valid values are **nfs://** or **gpfs://**.

*Host | Map*

*Host*

Specifies the server domain name system (DNS) name or IP address.

*Map*

Specifies the export map name.

**Notes:**

- When specifying **nfs://** as the value for *Protocol://*, you must provide a value for *Host* or *Map*.
- When specifying **gpfs://** as the value for *Protocol://*, do not provide a value for *Host*. However, provide a value for *Map* if it refers to an export map entry.

*Path*

Specifies the export path.

For example:

- The following command creates a single-writer AFM fileset in a GPFS file system *fs3* with a remote file system mounted at */gpfs/thefs1* from node *c41bn3*, using protocol **nfs://**:

```
mmcrfileset fs3 singleWriter2 -p afmtarget=nfs://c41bn3/gpfs/thefs1/target2 -p afemode=sw --inode-space new
```

Fileset singleWriter2 created with id 23 root inode 3145731.

- The following command creates a single-writer AFM fileset in a GPFS file system *fs3* with a GPFS remote file system mounted at */gpfs/thefs1*, using protocol **gpfs://**:

```
mmcrfileset fs3 singleWriter1 -p afmtarget=gpfs:///gpfs/thefs1/target1 -p afemode=sw --inode-space new
```

Fileset singleWriter1 created with id 21 root inode 2883587.

Note that in this case **///** is needed because *Host* is not provided.

The following optional AFM configuration parameters are also valid:

**afmAsyncDelay**

Specifies (in seconds) the amount of time by which write operations are delayed (because write operations are asynchronous with respect to remote clusters). For write-intensive applications that keep writing to the same set of files, this delay is helpful because it replaces multiple writes to the home cluster with a single write containing the latest data. However, setting a very high value weakens the consistency of data on the remote cluster.

This configuration parameter is applicable only for writer caches (SW, IW, and primary), where data from cache is pushed to home.

Valid values are between 1 and 2147483647. The default is 15.

**afmDirLookupRefreshInterval**

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a directory, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of that directory has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 60. In situations where home cluster data changes frequently, a value of 0 is recommended.

**afmDirOpenRefreshInterval**

Controls the frequency of data revalidations that are triggered by such I/O operations as **read** or **write** (specified in seconds). After a directory has been cached, **open** requests resulting from I/O operations on that object are directed to the cached directory until the specified amount of time has passed. Once the specified amount of time has passed, the **open** request gets directed to a gateway node rather than to the cached directory.

## **mmcrfileset**

Valid values are between 0 and 2147483647. The default is 60. Setting a lower value guarantees a higher level of consistency.

### **afmEnableAutoEviction**

Enables eviction on a given fileset. A **yes** value specifies that eviction is allowed on the fileset. A **no** value specifies that eviction is not allowed on the fileset.

See also the topic about cache eviction in the *IBM Spectrum Scale: Administration Guide*.

### **afmExpirationTimeout**

Is used with **afmDisconnectTimeout** (which can be set only through **mmchconfig**) to control how long a network outage between the cache and home clusters can continue before the data in the cache is considered out of sync with home. After **afmDisconnectTimeout** expires, cached data remains available until **afmExpirationTimeout** expires, at which point the cached data is considered expired and cannot be read until a reconnect occurs.

Valid values are 0 through 2147483647. The default is disable.

### **afmFileLookupRefreshInterval**

Controls the frequency of data revalidations that are triggered by such lookup operations as **ls** or **stat** (specified in seconds). When a lookup operation is done on a file, if the specified amount of time has passed, AFM sends a message to the home cluster to find out whether the metadata of the file has been modified since the last time it was checked. If the time interval has not passed, AFM does not check the home cluster for updates to the metadata.

Valid values are 0 through 2147483647. The default is 30. In situations where home cluster data changes frequently, a value of 0 is recommended.

### **afmMode**

Specifies the mode in which the cache operates. Valid values are the following:

#### **single-writer | sw**

Specifies single-writer mode.

#### **read-only | ro**

Specifies read-only mode. (For **mmcrfileset**, this is the default value.)

#### **local-updates | lu**

Specifies local-updates mode.

#### **independent-writer | iw**

Specifies independent-writer mode.

#### **Primary | drp**

Specifies the primary mode for AFM asynchronous data replication.

#### **Secondary | drs**

Specifies the secondary mode for AFM asynchronous data replication.

Changing from single-writer/read-only modes to read-only/local-updates/single-writer is supported. When changing from read-only to single-writer, the read-only cache is up-to-date. When changing from single-writer to read-only, all requests from cache should have been played at home. Changing from local-updates to read-only/local-updates/single-writer is restricted. A typical dataset is set up to include a single cache cluster in single-writer mode (which generates the data) and one or more cache clusters in local-updates or read-only mode. AFM single-writer/independent-writer filesets can be converted to primary. Primary/secondary filesets cannot be converted to AFM filesets.

In case of AFM asynchronous data replication, the **mmchfileset** command cannot be used to convert to primary from secondary. For detailed information, see *AFM-based Asynchronous Disaster Recovery (AFM DR)* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

For more information, see the topic about caching modes in the *IBM Spectrum Scale: Administration Guide* chapter about active file management.

#### **afmNumFlushThreads**

Defines the number of threads used on each gateway to synchronize updates to the home cluster. The default value is 4, which is sufficient for most installations. The current maximum value is 1024, which is too high for most installations; setting this parameter to such an extreme value should be avoided.

#### **afmParallelReadChunkSize**

Defines the minimum chunk size of the read that needs to be distributed among the gateway nodes during parallel reads. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

#### **afmParallelReadThreshold**

Defines the threshold beyond which parallel reads become effective. Reads are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value is 1024 MiB. The valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

#### **afmParallelWriteChunkSize**

Defines the minimum chunk size of the write that needs to be distributed among the gateway nodes during parallel writes. Values are interpreted in terms of bytes. The default value of this parameter is 128 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

#### **afmParallelWriteThreshold**

Defines the threshold beyond which parallel writes become effective. Writes are split into chunks when file size exceeds this threshold value. Values are interpreted in terms of MiB. The default value of this parameter is 1024 MiB, and the valid range of values is 0 to 2147483647. It can be changed cluster wide with the **mmchconfig** command. It can be set at fileset level using **mmcrfileset** or **mmchfileset** commands.

#### **afmPrefetchThreshold**

Controls partial file caching and prefetching. Valid values are the following:

- 0** Enables full file prefetching. This is useful for sequentially accessed files that are read in their entirety, such as image files, home directories, and development environments. The file will be prefetched after three blocks have been read into the cache.

#### **1-99**

Specifies the percentage of file size that must be cached before the entire file is prefetched. A large value is suitable for a file accessed either randomly or sequentially but partially, for which it might be useful to ingest the rest of the file when most of it has been accessed.

#### **100**

Disables full file prefetching. This value only fetches and caches data that is read by the application. This is useful for large random-access files, such as databases, that are either too big to fit in the cache or are never expected to be read in their entirety. When all data blocks are accessed in the cache, the file is marked as cached.

0 is the default value.

For local-updates mode, the whole file is prefetched when the first update is made.

## **mmcrfileset**

### **afmPrimaryId**

Specifies the unique primary ID of the primary fileset for asynchronous data replication. This is used for connecting a secondary to a primary.

### **afmRPO**

Specifies the recovery point objective (RPO) interval for an AFM DR fileset. This attribute is disabled by default. You can specify a value with the suffix M for minutes, H for hours, or W for weeks. For example, for 12 hours specify 12H. If you do not add a suffix, the value is assumed to be in minutes. The range of valid values is 720 minutes - 2147483647 minutes.

### **afmShowHomeSnapshot**

Controls the visibility of the home snapshot directory in cache. For this to be visible in cache, this variable has to be set to yes, and the snapshot directory name in cache and home should not be the same.

#### **yes**

Specifies that the home snapshot link directory is visible.

**no** Specifies that the home snapshot link directory is not visible.

See *Peer snapshot -psnap* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### **-t Comment**

Specifies an optional comment that appears in the output of the **mmlsfileset** command. This comment must be less than 256 characters in length.

### **--inode-space {new | ExistingFileset}**

Specifies the type of fileset to create, which controls how inodes are allocated:

#### **new**

Creates an independent fileset and its own dedicated inode space.

#### *ExistingFileset*

Creates a dependent fileset that will share inode space with the specified *ExistingFileset*. The *ExistingFileset* can be **root** or any other independent fileset.

If **--inode-space** is not specified, a dependent fileset will be created in the root inode space.

### **--inode-limit MaxNumInodes [:NumInodesToPreallocate]**

Specifies the inode limit for the new inode space. The *NumInodesToPreallocate* specifies an optional number of inodes to preallocate when the fileset is created. This option is valid only when creating an independent fileset with the **--inode-space new** parameter.

### **--allow-permission-change PermissionChangeMode**

Specifies the new permission change mode. This mode controls how **chmod** and ACL operations are handled on objects in the fileset. Valid modes are as follows:

#### **chmodOnly**

Specifies that only the UNIX change mode operation (**chmod**) is allowed to change access permissions (ACL commands and API will not be accepted).

#### **setAclOnly**

Specifies that permissions can be changed using ACL commands and API only (**chmod** will not be accepted).

#### **chmodAndSetAcl**

Specifies that **chmod** and ACL operations are permitted. If the **chmod** command (or **setattr** file operation) is issued, the result depends on the type of ACL that was previously controlling access to the object:

- If the object had a Posix ACL, it will be modified accordingly.
- If the object had an NFSv4 ACL, it will be replaced by the given UNIX mode bits.

**Note:** This is the default setting when a fileset is created.

#### chmodAndUpdateAcl

Specifies that **chmod** and ACL operations are permitted. If **chmod** is issued, the ACL will be updated by privileges derived from UNIX mode bits.

### Exit status

**0** Successful completion.

#### nonzero

A failure has occurred.

### Security

You must have root authority to run the **mmcrfileset** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### Examples

1. This example creates a fileset in file system **gpfs1**:

```
mmcrfileset gpfs1 fset1
```

The system displays output similar to:

Fileset fset1 created with id 1.

2. This example adds **fset2** in file system **gpfs1** with the comment "another fileset":

```
mmcrfileset gpfs1 fset2 -t "another fileset"
```

The system displays output similar to:

Fileset fset2 created with id 2.

To confirm the change, issue this command:

```
mmlsfileset gpfs1 -L
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |    |           |          |                          |            |           |             |                 |
|----------------------------------|----|-----------|----------|--------------------------|------------|-----------|-------------|-----------------|
| Name                             | Id | RootInode | ParentId | Created                  | InodeSpace | MaxInodes | AllocInodes | Comment         |
| root                             | 0  | 3         | --       | Mon Apr 12 16:31:05 2010 | 0          | 8001536   | 8001536     | root fileset    |
| fset1                            | 1  | 13568     | 0        | Mon Apr 12 16:32:28 2010 | 0          | 0         | 0           |                 |
| fset2                            | 2  | 13569     | 0        | Mon Apr 12 16:32:28 2010 | 0          | 0         | 0           | another fileset |

### See also

- “[mmchfileset command](#)” on page 173
- “[mmdelfileset command](#)” on page 292
- “[mmlinkfileset command](#)” on page 390
- “[mmlsfileset command](#)” on page 406
- “[mmunlinkfileset command](#)” on page 578

### Location

/usr/lpp/mmfs/bin

---

## mmcrfs command

Creates a GPFS file system.

### Synopsis

```
mmcrfs Device {"DiskDesc[;DiskDesc...]" | -F StanzaFile}
 [-A {yes | no | automount}] [-B BlockSize] [-D {posix | nfs4}]
 [-E {yes | no}] [-i InodeSize] [-j {cluster | scatter}]
 [-k {posix | nfs4 | all}] [-K {no | whenpossible | always}]
 [-L LogFileSize] [-m DefaultMetadataReplicas]
 [-M MaxMetadataReplicas] [-n NumNodes] [-Q {yes | no}]
 [-r DefaultDataReplicas] [-R MaxDataReplicas]
 [-S {yes | no | relative}] [-T Mountpoint] [-t DriveLetter]
 [-v {yes | no}] [-z {yes | no}] [--filesedf | --nofilesedf]
 [--inode-limit MaxNumInodes[:NumInodesToPreallocate]]
 [--log-replicas LogReplicas] [--metadata-block-size MetadataBlockSize]
 [--perfileset-quota | --noperfileset-quota]
 [--mount-priority Priority] [--version VersionString]
 [--write-cache-threshold HAWCThreshold]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmcrfs** command to create a GPFS file system. The first two parameters *must* be *Device* and either *DiskDescList* or *StanzaFile* and they *must* be in that order. The block size and replication factors chosen affect file system performance. A maximum of 256 file systems can be mounted in a GPFS cluster at one time, including remote file systems.

When deciding on the maximum number of files (number of inodes) in a file system, consider that for file systems that will be doing parallel file creates, if the total number of free inodes is not greater than 5% of the total number of inodes, there is the potential for slowdown in file system access. The total number of inodes can be increased using the **mmchfs** command.

When deciding on a block size for a file system, consider these points:

1. Supported block sizes are 64 KiB, 128 KiB, 256 KiB, 512 KiB, 1 MiB, 2 MiB, 4 MiB, 8 MiB, and 16 MiB.
2. The GPFS block size determines:
  - The minimum disk space allocation unit. The minimum amount of space that file data can occupy is a sub-block. A sub-block is 1/32 of the block size.
  - The maximum size of a read or write request that GPFS sends to the underlying disk driver.
3. From a performance perspective, it is recommended that you set the GPFS block size to match the application buffer size, the RAID stripe size, or a multiple of the RAID stripe size. If the GPFS block size does not match the RAID stripe size, performance may be severely degraded, especially for write operations. If IBM Spectrum Scale RAID is in use, the block size must equal the vdisk track size. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.
4. In file systems with a high degree of variance in the size of files within the file system, using a small block size would have a large impact on performance when accessing large files. In this kind of system it is suggested that you use a block size of 256 KB (8 KB sub-block). Even if only 1% of the files are large, the amount of space taken by the large files usually dominates the amount of space used on disk, and the waste in the sub-block used for small files is usually insignificant. For further performance information, see the GPFS white papers in the Techdocs Library ([www.ibm.com/support/techdocs/atスマート.nsf/Web/WhitePapers](http://www.ibm.com/support/techdocs/atスマート.nsf/Web/WhitePapers)).
5. The effect of block size on file system performance largely depends on the application I/O pattern.
  - A larger block size is often beneficial for large sequential read and write workloads.

- A smaller block size is likely to offer better performance for small file, small random read and write, and metadata-intensive workloads.
- The efficiency of many algorithms that rely on caching file data in a GPFS page pool depends more on the number of blocks cached rather than the absolute amount of data. For a page pool of a given size, a larger file system block size would mean fewer blocks cached. Therefore, when you create file systems with a block size larger than the default of 256 KB, it is recommended that you increase the page pool size in proportion to the block size.
  - The file system block size must not exceed the value of the GPFS maximum file system block size. The default maximum block size is 1 MiB. If a larger block size is desired, use the **mmchconfig** command to increase the **maxblocksize** configuration parameter before starting GPFS.

## Results

Upon successful completion of the **mmcrfs** command, these tasks are completed on all GPFS nodes:

- Mount point directory is created.
- File system is formatted.

Prior to GPFS 3.5, the disk information was specified in the form of disk descriptors defined as follows (with the second, third, and sixth fields reserved):

**DiskName::DiskUsage::FailureGroup::StoragePool:**

For backward compatibility, the **mmcrfs** command will still accept the traditional disk descriptors, but their use is discouraged.

## Parameters

### *Device*

The device name of the file system to be created.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**. However, file system names must be unique within a GPFS cluster. Do not specify an existing entry in **/dev**.

This must be the first parameter.

### **"DiskDesc[;DiskDesc...]"**

A descriptor for each disk to be included. Each descriptor is separated by a semicolon (;). The entire list must be enclosed in quotation marks (' or "). The use of disk descriptors is discouraged.

### **-F StanzaFile**

Specifies a file containing the NSD stanzas and pool stanzas for the disks to be added to the file system.

NSD stanzas have this format:

```
%nsd:
 nsd=NsibName
 usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}
 failureGroup=FailureGroup
 pool=StoragePool
 servers=ServerList
 device=DiskName
```

where:

**nsd=NsibName**

The name of an NSD previously created by the **mmcrnsd** command. For a list of available disks, issue the **mmlsnsd -F** command. This clause is mandatory for the **mmcrfs** command.

**usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}**

Specifies the type of data to be stored on the disk:

## **mmcrfs**

### **dataAndMetadata**

Indicates that the disk contains both data and metadata. This is the default for disks in the system pool.

### **dataOnly**

Indicates that the disk contains data and does not contain metadata. This is the default for disks in storage pools other than the system pool.

### **metadataOnly**

Indicates that the disk contains metadata and does not contain data.

### **descOnly**

Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster-recovery configurations. For more information, see the *IBM Spectrum Scale: Administration Guide* and search for “Synchronous mirroring utilizing GPFS replication”

### **failureGroup=FailureGroup**

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

Disk that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

### **pool=StoragePool**

Specifies the storage pool to which the disk is to be assigned. If this name is not provided, the default is **system**.

Only the system storage pool can contain **metadataOnly**, **dataAndMetadata**, or **descOnly** disks. Disks in other storage pools must be **dataOnly**.

### **servers=ServerList**

A comma-separated list of NSD server nodes. This clause is ignored by the **mmcrfs** command.

### **device=DiskName**

The block device name of the underlying disk device. This clause is ignored by the **mmcrfs** command.

Pool stanzas have this format:

```
%pool:
 pool=StoragePoolName
 blockSize=BlockSize
 usage={dataOnly | metadataOnly | dataAndMetadata}
 layoutMap={scatter | cluster}
 allowWriteAffinity={yes | no}
 writeAffinityDepth={0 | 1 | 2}
 blockGroupFactor=BlockGroupFactor
```

where:

### **pool=StoragePoolName**

Is the name of a storage pool.

**blockSize=BlockSize**

Specifies the block size of the disks in the storage pool.

**usage={dataOnly | metadataOnly | dataAndMetadata}**

Specifies the type of data to be stored in the storage pool:

**dataAndMetadata**

Indicates that the disks in the storage pool contain both data and metadata. This is the default for disks in the system pool.

**dataOnly**

Indicates that the disks contain data and do not contain metadata. This is the default for disks in storage pools other than the system pool.

**metadataOnly**

Indicates that the disks contain metadata and do not contain data.

**layoutMap={scatter | cluster}**

Specifies the block allocation map type. When allocating blocks for a given file, GPFS first uses a round-robin algorithm to spread the data across all disks in the storage pool. After a disk is selected, the location of the data block on the disk is determined by the block allocation map type. If **cluster** is specified, GPFS attempts to allocate blocks in clusters. Blocks that belong to a particular file are kept adjacent to each other within each cluster. If **scatter** is specified, the location of the block is chosen randomly.

The **cluster** allocation method may provide better disk performance for some disk subsystems in relatively small installations. The benefits of clustered block allocation diminish when the number of nodes in the cluster or the number of disks in a file system increases, or when the file system's free space becomes fragmented. The **cluster** allocation method is the default for GPFS clusters with eight or fewer nodes and for file systems with eight or fewer disks.

The **scatter** allocation method provides more consistent file system performance by averaging out performance variations due to block location (for many disk subsystems, the location of the data relative to the disk edge has a substantial effect on performance). This allocation method is appropriate in most cases and is the default for GPFS clusters with more than eight nodes or file systems with more than eight disks.

The block allocation map type cannot be changed after the storage pool has been created.

**allowWriteAffinity={yes | no}**

Indicates whether the File Placement Optimizer (FPO) feature is to be enabled for the storage pool. For more information on FPO, see the *File Placement Optimizer* section in the *IBM Spectrum Scale: Administration Guide*.

**writeAffinityDepth={0 | 1 | 2}**

Specifies the allocation policy to be used by the node writing the data.

A write affinity depth of 0 indicates that each replica is to be striped across the disks in a cyclical fashion with the restriction that no two disks are in the same failure group. By default, the unit of striping is a block; however, if the block group factor is specified in order to exploit chunks, the unit of striping is a chunk.

A write affinity depth of 1 indicates that the first copy is written to the writer node. The second copy is written to a different rack. The third copy is written to the same rack as the second copy, but on a different half (which can be composed of several nodes).

A write affinity depth of 2 indicates that the first copy is written to the writer node. The second copy is written to the same rack as the first copy, but on a different half (which can be composed of several nodes). The target node is determined by a hash value on the fileset ID of the file, or it is chosen randomly if the file does not belong to any fileset. The third copy is striped across the

## **mmcrfs**

disks in a cyclical fashion with the restriction that no two disks are in the same failure group. The following conditions must be met while using a write affinity depth of 2 to get evenly allocated space in all disks:

1. The configuration in disk number, disk size, and node number for each rack must be similar.
2. The number of nodes must be the same in the bottom half and the top half of each rack.

This behavior can be altered on an individual file basis by using the **--write-affinity-failure-group** option of the **mmchattr** command.

This parameter is ignored if write affinity is disabled for the storage pool.

### **blockGroupFactor=BlockGroupFactor**

Specifies how many file system blocks are laid out sequentially on disk to behave like a single large block. This option only works if **--allow-write-affinity** is set for the data pool. This applies only to a new data block layout; it does not migrate previously existing data blocks.

See the section about File Placement Optimizer in the *IBM Spectrum Scale: Administration Guide*.

#### **-A {yes | no | automount}**

Indicates when the file system is to be mounted:

##### **yes**

When the GPFS daemon starts. This is the default.

##### **no**

Manual mount.

##### **automount**

On non-Windows nodes, when the file system is first accessed. On Windows nodes, when the GPFS daemon starts.

|     **Note:** IBM Spectrum Protect for Space Management does not support file systems with the **-A** option set to **automount**.

#### **-B BlockSize**

Specifies the size of data blocks. Must be 64 KiB, 128 KiB, 256 KiB (the default), 512 KiB, 1 MiB, 2 MiB, 4 MiB, 8 MiB, or 16 MiB. Specify this value with the character **K** or **M**, for example **512K**.

#### **-D {nfs4 | posix}**

Specifies whether a deny-write open lock will block writes, which is expected and required by NFS V4. File systems supporting NFS V4 must have **-D nfs4** set. The option **-D posix** allows NFS writes even in the presence of a deny-write open lock. If you intend to export the file system using NFS V4 or Samba, you must use **-D nfs4**. For NFS V3 (or if the file system is not NFS exported at all) use **-D posix**. The default is **-D nfs4**.

#### **-E {yes | no}**

Specifies whether to report *exact mtime* values (**-E yes**), or to periodically update the **mtime** value for a file system (**-E no**). If it is more desirable to display exact modification times for a file system, specify or use the default **-E yes**.

#### **-i InodeSize**

Specifies the byte size of inodes. Supported inode sizes are 512, 1024, and 4096 bytes. The default is 4096.

#### **-j {cluster | scatter}**

Specifies the default block allocation map type to be used if **layoutMap** is not specified for a given storage pool.

#### **-k {posix | nfs4 | all}**

Specifies the type of authorization supported by the file system:

**posix**

Traditional GPFS ACLs only (NFS V4 and Windows ACLs are not allowed). Authorization controls are unchanged from earlier releases.

**nfs4**

Support for NFS V4 and Windows ACLs only. Users are not allowed to assign traditional GPFS ACLs to any file system objects (directories and individual files).

**all**

Any supported ACL type is permitted. This includes traditional GPFS (**posix**) and NFS V4 and Windows ACLs (**nfs4**).

The administrator is allowing a mixture of ACL types. For example, **fileA** may have a **posix** ACL, while **fileB** in the same file system may have an NFS V4 ACL, implying different access characteristics for each file depending on the ACL type that is currently assigned. The default is **-k all**.

Avoid specifying **nfs4** or **all** unless files will be exported to NFS V4 or Samba clients, or the file system will be mounted on Windows. NFS V4 and Windows ACLs affect file attributes (mode) and have access and authorization characteristics that are different from traditional GPFS ACLs.

**-K {no | whenpossible | always}**

Specifies whether strict replication is to be enforced:

**no** Indicates that strict replication is not enforced. GPFS will try to create the needed number of replicas, but will still return EOK as long as it can allocate at least one replica.

**whenpossible**

Indicates that strict replication is enforced provided the disk configuration allows it. If the number of failure groups is insufficient, strict replication will not be enforced. This is the default value.

**always**

Indicates that strict replication is enforced.

For more information, see the topic "Strict replication" in the *IBM Spectrum Scale: Problem Determination Guide*.

**-L LogFileSize**

Specifies the size of the internal log files. The *LogFileSize* specified must be a multiple of the metadata block size. The default size is 4 MB or the metadata block size, whichever is larger. The minimum size is 256 KB and the maximum size is 1024 MB. Specify this value with the K or M character, for example: 8M.

In most cases, allowing the log file size to default works well. An increased log file size is useful for file systems that have a large amount of metadata activity, such as creating and deleting many small files or performing extensive block allocation and deallocation of large files.

**-m DefaultMetadataReplicas**

Specifies the default number of copies of inodes, directories, and indirect blocks for a file. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of *MaxMetadataReplicas*. The default is 1.

**-M MaxMetadataReplicas**

Specifies the default maximum number of copies of inodes, directories, and indirect blocks for a file. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be less than the value of *DefaultMetadataReplicas*. The default is 2.

**-n NumNodes**

The estimated number of nodes that will mount the file system in the local cluster and all remote clusters. This is used as a best guess for the initial size of some file system data structures. The

## **mmcrfs**

default is 32. This value can be changed after the file system has been created but it does not change the existing data structures. Only the newly created data structure is affected by the new value. For example, *new storage pool*.

When you create a GPFS file system, you might want to overestimate the number of nodes that will mount the file system. GPFS uses this information for creating data structures that are essential for achieving maximum parallelism in file system operations (For more information, see *GPFS architecture* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide* ). If you are sure there will never be more than 64 nodes, allow the default value to be applied. If you are planning to add nodes to your system, you should specify a number larger than the default.

### **-Q {yes | no}**

Activates quotas automatically when the file system is mounted. The default is **-Q no**. Issue the **mmdefedquota** command to establish default quota values. Issue the **mmedquota** command to establish explicit quota values.

To activate GPFS quota management after the file system has been created:

1. Mount the file system.
2. To establish default quotas:
  - a. Issue the **mmdefedquota** command to establish default quota values.
  - b. Issue the **mmdefquotaon** command to activate default quotas.
3. To activate explicit quotas:
  - a. Issue the **mmedquota** command to activate quota values.
  - b. Issue the **mmquotaon** command to activate quota enforcement.

### **-r DefaultDataReplicas**

Specifies the default number of copies of each data block for a file. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be greater than the value of *MaxDataReplicas*. The default is 1.

### **-R MaxDataReplicas**

Specifies the default maximum number of copies of data blocks for a file. Valid values are 1, 2, and (for GPFS V3.5.0.7 and later) 3. This value cannot be less than the value of *DefaultDataReplicas*. The default is 2.

### **-S {yes | no | relatime}**

Suppresses the periodic updating of the value of **atime** as reported by the **gpfs\_stat()**, **gpfs\_fstat()**, **stat()**, and **fstat()** calls. The default value is **-S no**. Specifying **-S yes** for a new file system results in reporting the time the file system was created.

If **relatime** is specified, the file access time is updated only if the existing access time is older than the value of the **atimeDeferredSeconds** configuration attribute or the existing file modification time is greater than the existing access time.

### **-T MountPoint**

Specifies the mount point directory of the GPFS file system. If it is not specified, the mount point will be set to *DefaultMountDir/Device*. The default value for *DefaultMountDir* is /gpfs but, it can be changed with the **mmchconfig** command.

### **-t DriveLetter**

Specifies the drive letter to use when the file system is mounted on Windows.

### **-v {yes | no}**

Verifies that specified disks do not belong to an existing file system. The default is **-v yes**. Specify **-v no** only when you want to reuse disks that are no longer needed for an existing file system. If the command is interrupted for any reason, use **-v no** on the next invocation of the command.

**Important:** Using **-v no** on a disk that already belongs to a file system will corrupt that file system. This will not be noticed until the next time that file system is mounted.

**-z {yes | no}**

Enable or disable DMAPI on the file system. Turning this option on will require an external data management application such as IBM Spectrum Protect hierarchical storage management (HSM) before the file system can be mounted. The default is **-z no**.

For further information regarding DMAPI for GPFS, see *GPFS-specific DMAPI events in the IBM Spectrum Scale: Command and Programming Reference*.

| **Note:** IBM Spectrum Protect for Space Management does not support file systems with the **-A** option set to **automount**.

**--filesetdf**

Specifies that when quotas are enforced for a fileset, the numbers reported by the **df** command are based on the quotas for the fileset (rather than the entire file system). This option affects the **df** command behavior only on Linux nodes.

**--nofilesetdf**

Specifies that when quotas are enforced for a fileset, the numbers reported by the **df** command are based on the quotas for the entire file system (rather than individual filesets). This is the default.

**--inode-limit MaxNumInodes[:NumInodesToPreallocate]**

Specifies the maximum number of files in the file system.

For file systems on which you intend to create files in parallel, if the total number of free inodes is not greater than 5% of the total number of inodes, file system access might slow down. Take this into consideration when creating your file system.

The parameter *NumInodesToPreallocate* specifies the number of inodes that the system will immediately preallocate. If you do not specify a value for *NumInodesToPreallocate*, GPFS will dynamically allocate inodes as needed.

You can specify the *NumInodes* and *NumInodesToPreallocate* values with a suffix, for example 100K or 2M. Note that in order to optimize file system operations, the number of inodes that are actually created may be greater than the specified value.

**--log-replicas LogReplicas**

Specifies the number of recovery log replicas. Valid values are **1**, **2**, **3**, or **DEFAULT**. If not specified, or if **DEFAULT** is specified, the number of log replicas is the same as the number of metadata replicas currently in effect for the file system.

This option is only applicable if the recovery log is stored in the **system.log** storage pool.

**--metadata-block-size MetadataBlockSize**

Specifies the block size for the system storage pool, provided its usage is set to **metadataOnly**. Valid values are the same as those listed for **-B BlockSize**.

**--profileset-quota**

Sets the scope of user and group quota limit checks to the individual fileset level (rather than the entire file system).

**--noperfileset-quota**

Sets the scope of user and group quota limit checks to the entire file system (rather than per individual fileset). This is the default.

**--mount-priority Priority**

Controls the order in which the individual file systems are mounted at daemon startup or when one of the **all** keywords is specified on the **mmmount** command.

File systems with higher *Priority* numbers are mounted after file systems with lower numbers. File systems that do not have mount priorities are mounted last. A value of zero indicates no priority. This is the default.

## **mmcrfs**

### **--version VersionString**

Enable only the file system features that are compatible with the specified release. The lowest allowed *VersionString* is 3.1.0.0.

The default value is the current product version, which enables all currently available features but prevents nodes that are running earlier GPFS releases from accessing the file system. Windows nodes can mount only file systems that are created with GPFS 3.2.1.5 or later.

### **--profile ProfileName**

Specifies a predefined profile of attributes to be applied. System-defined profiles are located in /usr/lpp/mmfs/profiles/. All the file system attributes listed under a file system stanza will be changed as a result of this command. The following system-defined profile names are accepted:

- gpfssProtocolDefaults
- gpfssProtocolRandomIO

The file system attributes will be applied at file system creation. If there is a current profile in place on the system (use mmlsconfig profile to check), then the file system will be created with those attributes and values listed in the profile's file system stanza. The default is to use whatever attributes and values associate with the current profile setting.

Furthermore, any and all file system attributes from an installed profile file can be by-passed with '--profile=userDefinedProfile', where the userDefinedProfile is a profile file has been installed by the user in /var/mmfs/etc/.

User-defined profiles consist of the following stanzas:

```
%cluster:
[CommaSeparatedNodesOrNodeClasses:]ClusterConfigurationAttribute=Value
...
%filesystem:
FilesystemConfigurationAttribute=Value
...
```

A sample file can be found in /usr/lpp/mmfs/samples/sample.profile. See the **mmchconfig** command for a detailed description of the different configuration parameters.

User-defined profiles should be used only by experienced administrators. When in doubt, use the **mmchconfig** command instead.

### **--write-cache-threshold HAWCThreshold**

Specifies the maximum length (in bytes) of write requests that will be initially buffered in the highly-available write cache before being written back to primary storage. Only synchronous write requests are guaranteed to be buffered in this fashion.

A value of 0 disables this feature. 64K is the maximum supported value. Specify in multiples of 4K.

This feature can be enabled or disabled at any time (the file system does not need to be unmounted). For more information about this feature, see the topic *Highly-available write cache (HAWC)* in the *IBM Spectrum Scale: Administration Guide*.

## **Exit status**

**0**      Successful completion.

**nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmcrfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For

more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

This example shows how to create a file system named **gpfs1** using three disks, each with a block size of 512 KB, allowing metadata and data replication to be 2, turning quotas on, and creating /gpfs1 as the mount point. The NSD stanzas describing the three disks are assumed to have been placed in file/tmp/freedisks. To complete this task, issue the command:

```
mmcrfs gpfs1 -F /tmp/freedisks -B 512K -m 2 -r 2 -Q yes -T /gpfs1
```

The system displays output similar to:

```
The following disks of gpfs1 will be formatted on node c21f1n13:
hd2n97: size 1951449088 KB
hd3n97: size 1951449088 KB
hd4n97: size 1951449088 KB
Formatting file system ...
Disks up to size 16 TB can be added to storage pool 'system'.
Creating Inode File
Creating Allocation Maps
Creating Log Files
Clearing Inode Allocation Map
Clearing Block Allocation Map
Formatting Allocation Map for storage pool 'system'
 19 % complete on Tue Feb 28 18:03:20 2012
 42 % complete on Tue Feb 28 18:03:25 2012
 62 % complete on Tue Feb 28 18:03:30 2012
 79 % complete on Tue Feb 28 18:03:35 2012
 96 % complete on Tue Feb 28 18:03:40 2012
100 % complete on Tue Feb 28 18:03:41 2012
Completed creation of file system /dev/gpfs1.
mmcrfs: Propagating the cluster configuration data to all
 affected nodes. This is an asynchronous process.
```

## See also

- “[mmchfs command](#)” on page 179
- “[mmdelfs command](#)” on page 295
- “[mmdf command](#)” on page 308
- “[mmedquota command](#)” on page 325
- “[mmfsck command](#)” on page 331
- “[mmlsfs command](#)” on page 411
- “[mmlspool command](#)” on page 428

## Location

/usr/lpp/mmfs/bin

## **mmcrnodeclass**

---

### **mmcrnodeclass command**

Creates user-defined node classes.

#### **Synopsis**

```
mmcrnodeclass ClassName -N {Node[,Node...] | NodeFile | NodeClass}
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmcrnodeclass** command to create user-defined node classes. After a node class is created, it can be specified as an argument on commands that accept the **-N NodeClass** option.

#### **Parameters**

##### *ClassName*

Specifies a name that uniquely identifies the user-defined node classes to create. An existing node name cannot be specified. Class names that end with **nodes** or system-defined node classes are reserved for use by GPFS.

##### **-N {*Node[ ,Node...]* | *NodeFile* | *NodeClass*}**

Specifies the nodes and node classes that will become members of the user-defined node class *ClassName*.

*NodeClass* cannot be a node class that already contains other node classes. For example, two user-defined node classes called **siteA** and **siteB** could be used to create a new node class called **siteAandB**, as follows:

```
mmcrnodeclass siteAandB -N siteA,siteB
```

The **siteAandB** node class cannot later be specified for *NodeClass* when creating new node classes.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmcrnodeclass** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

#### **Examples**

To create a user-defined node class called **siteA** that contains nodes **c8f2c4vp1** and **c8f2c4vp2**, issue this command:

```
mmcrnodeclass siteA -N c8f2c4vp1,c8f2c4vp2
```

The system displays information similar to:

**mmcrnodeclass:** Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.

To display the member list of **siteA**, issue this command:

```
mmlsnodeclass siteA
```

The system displays information similar to:

| Node Class Name | Members |
|-----------------|---------|
|-----------------|---------|

|       |                     |
|-------|---------------------|
| siteA | c8f2c4vp1,c8f2c4vp2 |
|-------|---------------------|

## See also

- “[mmchnodeclass command](#)” on page 195
- “[mmdelnodetree command](#)” on page 300
- “[mmlsnodetree command](#)” on page 421

## Location

```
/usr/lpp/mmfs/bin
```

## **mmcrnsd**

---

### **mmcrnsd command**

Creates Network Shared Disks (NSDs) used by GPFS.

#### **Synopsis**

```
mmcrnsd -F StanzaFile [-v {yes | no}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmcrnsd** command is used to create cluster-wide names for NSDs used by GPFS.

This is the first GPFS step in preparing disks for use by a GPFS file system. The input to this command consists of a file containing NSD stanzas describing the properties of the disks to be created. This file may be updated as necessary by the **mmcrnsd** command and can be supplied as input to the **mmcrfs**, **mmadddisk**, or **mmrpldisk** command.

The names created by the **mmcrnsd** command are necessary since disks connected to multiple nodes may have different disk device names on each node. The NSD names uniquely identify each disk. This command must be run for all disks that are to be used in GPFS file systems. The **mmcrnsd** command is also used to assign each disk an NSD server list that can be used for I/O operations on behalf of nodes that do not have direct access to the disk.

- | To identify that a disk has been processed by the **mmcrnsd** command, a unique NSD volume ID is written to the disk. All of the NSD commands (**mmcrnsd**, **mmlsnsd**, and **mmdelnsd**) use this unique NSD volume ID to identify and process NSDs.

After the NSDs are created, the GPFS cluster data is updated and they are available for use by GPFS.

**Note:** It is customary to use whole LUNs as NSDs. This generally provides the best performance and fault isolation. When SCSI-3 PR is in use, whole LUN use is required. In other deployment scenarios, it is possible to use disk partitions rather than whole LUNs, as long as care is taken to ensure that sharing of the same LUN through multiple partitions does not have an undesirable performance impact.

On Windows, GPFS will only create NSDs from empty disk drives. **mmcrnsd** accepts Windows *Basic* disks or *Unknown/Not Initialized* disks. It always re-initializes these disks so that they become *Basic GPT Disks* with a single *GPFS partition*. NSD data is stored in GPFS partitions. This allows other operating system components to recognize the disks are used. **mmdelnsd** deletes the partition tables created by **mmcrnsd**.

#### **Results**

Upon successful completion of the **mmcrnsd** command, these tasks are completed:

- NSDs are created.
- The *StanzaFile* contains NSD names to be used as input to the **mmcrfs**, **mmadddisk**, or the **mmrpldisk** commands.
- | • A unique NSD volume ID to identify each disk as an NSD has been written to the disk.
- An entry for each new disk is created in the GPFS cluster data.

## Parameters

### **-F StanzaFile**

Specifies the file containing the NSD stanzas for the disks to be created. NSD stanzas have this format:

```
%nsd: device=DiskName
 nsd=NsdName
 servers=ServerList
 usage={dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}
 failureGroup=FailureGroup
 pool=StoragePool
```

where:

#### **device=DiskName**

On UNIX, the block device name appearing in /dev for the disk you want to define as an NSD. Examples of disks that are accessible through a block device are SAN-attached disks. If server nodes are specified, *DiskName* must be the /dev name for the disk device of the first listed NSD server node.

On Windows, the disk number (for example, 3) of the disk you want to define as an NSD. Disk numbers appear in Windows Disk Management console and the DISKPART command line utility. If a server node is specified, *DiskName* must be the disk number from the first NSD server node defined in the server list.

For the latest supported disk types, see the IBM Spectrum Scale FAQ in IBM Knowledge Center([www.ibm.com/support/knowledgecenter/STXKQY/gpfscustersfaq.html](http://www.ibm.com/support/knowledgecenter/STXKQY/gpfscustersfaq.html)).

This clause is mandatory for the **mmcrnsd** command.

#### **nsd=NsdName**

Specify the name you desire for the NSD to be created. This name must not already be used as another GPFS disk name, and it must not begin with the reserved string 'gpfs'.

**Note:** This name can contain only the following characters: 'A' through 'Z', 'a' through 'z', '0' through '9', or '\_' (the underscore). All other characters are not valid.

If this clause is not specified, GPFS will generate a unique name for the disk and will add the appropriate **nsd=NsdName** clause to the stanza file. The NSD is assigned a name according to the convention:

#### **gpfsNNnsd**

where NN is a unique nonnegative integer not used in any prior NSD.

#### **servers=ServerList**

Is a comma-separated list of NSD server nodes. You may specify up to eight NSD servers in this list. The defined NSD will preferentially use the first server on the list. If the first server is not available, the NSD will use the next available server on the list.

When specifying server nodes for your NSDs, the output of the **mmlscluster** command lists the host name and IP address combinations recognized by GPFS. The utilization of aliased host names not listed in the **mmlscluster** command output may produce undesired results.

There are two cases where a server list either must be omitted or is optional:

- For IBM Spectrum Scale RAID, a server list is not allowed. The servers are determined from the underlying vdisk definition. For more information about IBM Spectrum Scale RAID, see the *IBM Spectrum Scale RAID: Administration and Programming Reference (SA23-1354)*.
- For SAN configurations where the disks are SAN-attached to all nodes in the cluster, a server list is optional. However, if all nodes in the cluster do not have access to the disk, or if the file system to which the disk belongs is to be accessed by other GPFS clusters, you must specify a server list.

## **mmcrnsd**

**usage={dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}**

Specifies the type of data to be stored on the disk:

### **dataAndMetadata**

Indicates that the disk contains both data and metadata. This is the default for disks in the system pool.

### **dataOnly**

Indicates that the disk contains data and does not contain metadata. This is the default for disks in storage pools other than the system pool.

### **metadataOnly**

Indicates that the disk contains metadata and does not contain data.

### **descOnly**

Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster recovery configurations. For more information, see the help topic "Synchronous mirroring utilizing GPFS replication" in *IBM Spectrum Scale: Administration Guide*.

### **localCache**

Indicates that the disk is to be used as a local read-only cache device.

This clause is ignored by the **mmcrnsd** command and is passed unchanged to the output file produced by the **mmcrnsd** command.

### **failureGroup=FailureGroup**

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

Disk that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

This clause is ignored by the **mmcrnsd** command, and is passed unchanged to the output file produced by the **mmcrnsd** command.

### **pool=StoragePool**

Specifies the name of the storage pool to which the NSD is assigned. This clause is ignored by the **mmcrnsd** command and is passed unchanged to the output file produced by the **mmcrnsd** command.

The default value for **pool** is **system**.

### **-v {yes | no}**

Verify the disks are not already formatted as an NSD.

- | A value of **-v yes** specifies that the NSDs are to be created only if each disk has not been formatted by a previous invocation of the **mmcrnsd** command, as indicated by the NSD volume ID on the disk.
- | A value of **-v no** specifies that the disks are to be created irrespective of their previous state. The default is **-v yes**.

| **Important:** Using **-v no** when a disk already belongs to a file system can corrupt that file system by  
| making that physical disk undiscoverable by that file system. This will not be noticed until the next  
| time that file system is mounted.

Upon successful completion of the **mmcrnsd** command, the *StanzaFile* file is rewritten to reflect changes made by the command, as follows:

- If an NSD stanza is found to be in error, the stanza is commented out.
- If an **nsd=NsdName** clause is not specified, and an NSD name is generated by GPFS, an **nsd=** clause will be inserted in the corresponding stanza.

You must have **write** access to the directory where the *StanzaFile* file is located in order to rewrite the created NSD information.

The disk usage, failure group, and storage pool specifications are preserved only if you use the rewritten file produced by the **mmcrnsd** command. If you do not use this file, you must either accept the default values or specify new values when creating NSD stanzas for other commands.

Prior to GPFS 3.5, the disk information was specified in the form of disk descriptors defined as:

**DiskName:ServerList::DiskUsage:FailureGroup:DesiredName:StoragePool**

For backward compatibility, the **mmcrnsd** command will still accept the traditional disk descriptors, but their use is discouraged. For additional information about GPFS stanzas, see the help topic "Stanza files" in the *IBM Spectrum Scale: Administration Guide*.

## Exit status

**0** Successful completion.

### nonzero

A failure has occurred.

## Security

You must have root authority to run the **mmcrnsd** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

To create two new NSDs from the stanza file `/tmp/newNSDstanza` containing:

```
%nsd: device=/dev/sdav1
servers=k145n05,k145n06
failureGroup=4
```

```
%nsd:
device=/dev/sdav2
nsd=sd2pA
servers=k145n06,k145n05
usage=dataonly
failureGroup=5
pool=poolA
```

Issue this command:

```
mmcrnsd -F /tmp/newNSDstanza
```

## **mmcrnsd**

The output is similar to this:

```
mmcrnsd: Processing disk sdav1
mmcrnsd: Processing disk sdav2
mmcrnsd: 6027-1371 Propagating the cluster configuration data to all
 affected nodes. This is an asynchronous process.
```

As a result, two NSDs will be created. The first one will be assigned a name by the system, for example, **gpfs1023nsd**. The second disk will be assigned the name **sd2pA** (as indicated by the **nsd=** clause in the stanza).

The newNSDstanza file will be rewritten and will look like this (note the addition of an **nsd=** clause in the first stanza):

```
%nsd: nsd=gpfs1023nsd device=/dev/sdav1
 servers=k145n05,k145n06
 failureGroup=4

%nsd:
 device=/dev/sdav2
 nsd=sd2pA
 servers=k145n06,k145n05
 usage=dataOnly
 failureGroup=5
 pool=poolA
```

## **See also**

- “[mmadddisk command](#)” on page 23
- “[mmchnsd command](#)” on page 198
- “[mmcrfs command](#)” on page 250
- “[mmdeldisk command](#)” on page 287
- “[mmdelnsd command](#)” on page 302
- “[mmlsnsd command](#)” on page 423
- “[mmrpldisk command](#)” on page 539

## **Location**

/usr/lpp/mmfs/bin

---

## mmcrsnapshot command

Creates a snapshot of a file system or fileset at a single point in time.

### Synopsis

```
mmcrsnapshot Device [[Fileset]:]Snapshot[[[Fileset]:]Snapshot[-j FilesetName[,FilesetName...]]]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmcrsnapshot** command to create global snapshots or fileset snapshots at a single point in time. System data and existing snapshots are not copied. The snapshot function allows a backup or mirror program to run concurrently with user updates and still obtain a consistent copy of the file system as of the time the copy was created. Snapshots also provide an online backup capability that allows easy recovery from common problems such as accidental deletion of a file, and comparison with older versions of a file.

In IBM Spectrum Scale Release 4.2.1 and later, snapshot commands support the specification of multiple snapshots. Users can easily create and delete multiple snapshots. Also, system performance is increased by batching operations and reducing overhead.

In this release, the following new usages of the **mmcrsnapshot** command have been introduced:

```
mmcrsnapshot device [[Fileset]:]Snapshot[[[Fileset]:]Snapshot...]
mmcrsnapshot device [[Fileset]:]Snapshot[[[Fileset]:]Snapshot...] -j Fileset
mmcrsnapshot device Snapshot -j Fileset1,Fileset2,...
```

A global snapshot is an exact copy of changed data in the active files and directories of a file system. Snapshots of a file system are read-only and they appear in a **.snapshots** directory located in the file system root directory. The files and attributes of the file system can be changed only in the active copy.

A fileset snapshot is an exact copy of changed data in the active files and directories of an independent fileset plus all dependent filesets. Fileset snapshots are read-only and they appear in a **.snapshots** directory located in the root directory of the fileset. The files and attributes of the fileset can be changed only in the active copy.

Snapshots may be deleted only by issuing the **mmdelsnapshot** command. The **.snapshots** directory cannot be deleted, though it can be renamed with the **mmsnapdir** command using the **-s** option.

Because global snapshots are not full, independent copies of the entire file system, they should not be used as protection against media failures. For information about protection against media failures, see the *Recoverability considerations* topic in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

For more information on global snapshots, see *Creating and maintaining snapshots of GPFSfile systems* in the *IBM Spectrum Scale: Administration Guide*.

For more information on fileset snapshots, see *Fileset-level snapshots* in the *IBM Spectrum Scale: Administration Guide*.

### Parameters

#### *Device*

The device name of the file system for which the snapshot is to be created. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

## **mmcrsnapshot**

This must be the first parameter.

### *Fileset*

The name of the fileset that contains the fileset snapshot to be created. If *Fileset* is not specified, the **mmcrsnapshot** command creates a global snapshot named Snapshot.

**Note:** Ensure that multiple snapshots and multiple filesets are not used together.

### *Snapshot*

Specifies the name given to the snapshot.

The snapshot names are separated by a comma.

The snapshot specifier describes global and fileset snapshots. For example, *Fileset1:Snapshot1* specifies a fileset snapshot named Snapshot1 for fileset Fileset1. If *Fileset1* is empty, *Snapshot1* is a global snapshot named Snapshot1.

Each global snapshot name must be unique from any other global snapshots. If you do not want to traverse the root of the file system to access the global snapshot, a more convenient mechanism that enables a connection in each directory of the active file system can be enabled with the **-a** option of the **mmsnapdir** command.

**Note:** Ensure that the snapshot name is using the "@GMT-yyyy.MM.dd-HH.mm.ss" format in order to be identifiable by the Windows VSS.

For a fileset snapshot, *Snapshot* appears as a subdirectory of the **.snapshots** directory in the root directory of the fileset. Fileset snapshot names can be duplicated across different filesets. A fileset snapshot can also have the same name as a global snapshot. The **mmsnapdir** command provides an option to make global snapshots also available through the **.snapshots** in the root directory of all independent filesets.

#### **Note:**

- Ensure that the snapshot name does not include a colon (:) and a comma (,).
- Ensure that multiple snapshots and multiple filesets are not used together.

### **-j** *FilesetName*

Creates a snapshot that includes the specified fileset and all the dependent filesets that share the same inode space. *FilesetName* refers to an independent fileset. If **-j** is not specified, the **mmcrsnapshot** command creates a global snapshot that includes all filesets.

If **-j** is specified, fileset names can be considered as a list separated by commas.

#### **Note:**

- When a list of snapshots separated by a comma (,) is used with the **-j** option, the fileset is applicable to each snapshot that does not use the colon (:) syntax. The fileset name must not consist of a white space.
- Only one global snapshot and one snapshot can be specified in each fileset. Therefore, multiple snapshots must not be created with the **-j** option. IBM recommends that at the most one global snapshot and one fileset snapshot must be included in each independent fileset. If multiple snapshots are specified with the **-j** option, a snapshot with the same name is created for each fileset that has been listed.

## **Exit status**

0      Successful completion.

### **nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmcrsnapshot** command when creating global snapshots.

Independent fileset owners can run the **mmcrsnapshot** command to create snapshots of the filesets they own.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To create a global snapshot **snap1**, for the file system **fs1**, issue this command:

```
mmcrsnapshot fs1 snap1
```

The output is similar to this:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
Snapshot snap1 created with id 1.
```

Before issuing the command, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3
```

After the command has been issued, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3
```

If a second snapshot were to be created at a later time, the first snapshot would remain as is. Snapshots are made only of active file systems, not existing snapshots. For example:

```
mmcrsnapshot fs1 snap2
```

The output is similar to this:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
Snapshot snap2 created with id 2.
```

After the command has been issued, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3

/fs1/.snapshots/snap2/file1
/fs1/.snapshots/snap2/userA/file2
/fs1/.snapshots/snap2/userA/file3
```

2. To create a snapshot **Snap3** of the fileset **FsetF5-V2** for the file system **fs1**, issue this command:

```
mmcrsnapshot fs1 Snap3 -j FsetF5-V2
```

## **mmcrsnapshot**

The system displays output similar to:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
Snapshot Snap3 created with id 69.
```

To display the snapshot that contains the **FsetF5-V2** fileset, issue this command:

```
mmlssnapshot device -j FsetF5-V2
```

The system displays output similar to:

```
Snapshots in file system fs1:
Directory SnapId Status Created Fileset
Snap3 69 Valid Wed Feb 1 12:55:51 2012 FsetF5-V2
```

3. To create a snapshot of the gpfs0 file system that can be viewed over SMB protocol with Windows VSS, issue this command:

```
mmcrsnapshot gpfs0 $(date --utc +@GMT-%Y.%m.%d-%H.%M.%S)
```

The system displays output similar to:

```
mmcrsnapshot gpfs0 $(date --utc +@GMT-%Y.%m.%d-%H.%M.%S)
Flushing dirty data for snapshot @GMT-2015.10.02-21.03.33...
Quiescing all file system operations.
Snapshot @GMT-2015.10.02-21.03.33 created with id 7.
```

4. To create a snapshot for the fset1, fset2, fset3 filesets for the file system **fs1**, run the following command:

```
mmcrsnapshot fs1 snap1 -j fset1,fset2,fset3
```

The system displays the following output:

```
Flushing dirty data for snapshot fset1:snap1 fset2:snap1 fset3:snap1 (1..3) of 3...
Quiescing all file system operations.
Snapshot fset1:snap1 created with id 1.
Snapshot fset2:snap1 created with id 2.
Snapshot fset3:snap1 created with id 3.
```

5. To specify different snapshot names for each fileset (fset1, fset2, and fset3) for the file system **fs1**, run the following command:

```
mmcrsnapshot fs1 fset1:snapA,fset2:snapB,fset3:snapC
```

The system displays the following output:

```
Flushing dirty data for snapshot fset1:snapA fset2:snapB fset3:snapC (1..3) of 3...
Quiescing all file system operations.
Snapshot fset1:snapA created with id 4.
Snapshot fset2:snapB created with id 5.
Snapshot fset3:snapC created with id 6.
```

To view the list of snapshots, run the following command:

```
mmlssnapshot fs1
```

The system displays the following output:

```
Snapshots in file system fs1:
Directory SnapId Status Created Fileset
snap1 1 Valid Thu May 12 04:27:18 2016 fset1
snap1 2 Valid Thu May 12 04:27:18 2016 fset2
snap1 3 Valid Thu May 12 04:27:18 2016 fset3
snapA 4 Valid Thu May 12 04:28:33 2016 fset1
snapB 5 Valid Thu May 12 04:28:33 2016 fset2
snapC 6 Valid Thu May 12 04:28:33 2016 fset3
```

**See also**

- “[mmdelsnapshot command](#)” on page 304
- “[mmlssnapshot command](#)” on page 439
- “[mmrestorefs command](#)” on page 525
- “[mmsnapdir command](#)” on page 565

**Location**

/usr/lpp/mmfs/bin

## **mmdefedquota**

---

### **mmdefedquota command**

Sets default quota limits.

#### **Synopsis**

```
mmdefedquota {-u | -g | -j} Device
```

or

```
mmdefedquota {-u | -g} Device:Fileset
```

#### **Availability**

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

#### **Description**

Use the **mmdefedquota** command to set or change default quota limits. Default quota limits can be set for new users, groups, and filesets for a specified file system. Default quota limits can also be applied at a more granular level for new users and groups in a specified fileset.

Default quota limits can be set or changed only if the **-Q yes** option is in effect for the file system, and quotas have been enabled with the **mmdefquotaon** command . To set default quotas at the fileset level, the **--perfileset-quota** option must also be in effect. If **--perfileset-quota** is in effect, all users and groups in the fileset **root** will not be impacted by default quota unless they are explicitly set. The **-Q yes** and **--perfileset-quota** options are specified when creating a file system with the **mmcrfs** command or changing file system attributes with the **mmchfs** command. Use the **mmfsinfo** command to display the current settings of these quota options.

The **mmdefedquota** command displays the current values for these limits, if any, and prompts you to enter new values using your default editor:

- current block usage (display only)
- current inode usage (display only)
- inode soft limit
- inode hard limit
- block soft limit

Displayed in **KB**, but may be specified using **g**, **G**, **k**, **K**, **m**, **M**, **p**, **P**, **t**, or **T**. If no suffix is provided, the number is assumed to be in **bytes**.

- block hard limit

Displayed in **KB**, but may be specified using **g**, **G**, **k**, **K**, **m**, **M**, **p**, **P**, **t**, or **T**. If no suffix is provided, the number is assumed to be in **bytes**.

**Note:** A block or inode limit of 0 indicates no limit.

The **mmdefedquota** command waits for the edit window to be closed before checking and applying new values. If an incorrect entry is made, reissue the command and enter the correct values.

When you set quota limits for a file system, consider replication within the file system. For more information, see the topic *Listing quotas* in the *IBM Spectrum Scale: Administration Guide*.

The **EDITOR** environment variable must contain a complete path name, for example:

```
export EDITOR=/bin/vi
```

## Parameters

### *Device*

The device name of the file system to have default quota values set for.

File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

### *Fileset*

The name of a fileset in the file system to have default quota values set for.

## Options

- g** Specifies that the default quota value is to be applied for new groups accessing the specified file system or fileset.
- j** Specifies that the default quota value is to be applied for new files in the specified file system.
- u** Specifies that the default quota value is to be applied for new users accessing the specified file system or fileset.

## Exit status

**0** Successful completion.

### **nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmdefedquota** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmdefedquota** command is issued.

## Examples

1. To set default quotas for new users of the file system **gpfs1**, issue this command:

```
mmdefedquota -u gpfs1
```

The system displays information in your default editor similar to:

```
*** Edit quota limits for USR DEFAULT entry
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs1: blocks in use: 0K, limits (soft = 0K, hard = 0K)
 inodes in use: 0, limits (soft = 0, hard = 0)
```

Change the soft block limit to 19 GB, the hard block limit to 20 GB, the inode soft limit to 1 KB, and the inode hard limit to 20 KB, as follows:

```
*** Edit quota limits for USR DEFAULT entry
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs1: blocks in use: 0K, limits (soft = 19G, hard = 20G)
 inodes in use: 0, limits (soft = 1K, hard = 20K)
```

After the edit window is closed, issue this command to confirm the change:

```
mmlsquota -d -u gpfs1
```

## **mmdefedquota**

The system displays information similar to:

| Default Block Limits(KB) |         |      | Default File Limits |          |            |
|--------------------------|---------|------|---------------------|----------|------------|
| Filesystem               | Fileset | type | quota               | limit    | quota      |
| gpfs1                    |         | USR  | 19922944            | 20971520 | 1024 20480 |

2. To set default quotas for new users of fileset **fset1** in file system **gpfs1**, issue this command:

```
mmdefedquota -u gpfs1:fset1
```

The system displays information in your default editor similar to:

```
*** Edit quota limits for USR DEFAULT entry for fileset fset1
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs1: blocks in use: 0K, limits (soft = 0K, hard = 31457280K)
 inodes in use: 0, limits (soft = 0, hard = 0)
```

Change the soft block limit to 3 GB and the hard block limit to 6 GB, as follows:

```
*** Edit quota limits for USR DEFAULT entry for fileset fset1
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs1: blocks in use: 0K, limits (soft = 3G, hard = 6G)
 inodes in use: 0, limits (soft = 0, hard = 0)
```

After the edit window is closed, issue this command to confirm the change:

```
mmlsquota -d gpfs1:fset1
```

The system displays information similar to:

| Default Block Limits(KB) |         |      |         | Default File Limits |       |       |
|--------------------------|---------|------|---------|---------------------|-------|-------|
| Filesystem               | Fileset | type | quota   | limit               | quota | limit |
| gpfs1                    | fset1   | USR  | 3145728 | 6291456             | 0     | 0     |
| gpfs1                    | fset1   | GRP  | 0       | 0                   | 0     | 0     |

## **See also**

- “mmchfs command” on page 179
- “mmcheckquota command” on page 169
- “mmcrfs command” on page 250
- “mmdefquotaoff command” on page 275
- “mmdefquotaoon command” on page 278
- “mmedquota command” on page 325
- “mmlsfs command” on page 411
- “mmlsquota command” on page 435
- “mmquotaoff command” on page 507
- “mmrepquota command” on page 517

## **Location**

/usr/lpp/mmfs/bin

---

## mmdefquotaoff command

Deactivates default quota limit usage.

### Synopsis

```
mmdefquotaoff [-u] [-g] [-j] [-v] [-d] {Device [Device...] | -a}
```

or

```
mmdefquotaoff [-u] [-g] [-v] [-d] {Device:Fileset ... | -a}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmdefquotaoff** command deactivates default quota limits for file systems and filesets. If default quota limits are deactivated, new users, groups, or filesets will then have a default quota limit of 0, indicating no limit.

If none of the following options are specified, the **mmdefquotaoff** command deactivates all default quotas:

- u
- j
- g

If the **-a** option is not used, *Device* must be the last parameter specified.

### Parameters

#### *Device*

The device name of the file system to have default quota values deactivated.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

#### *Fileset*

The name of a fileset in the file system to have default quota values deactivated.

### Options

- a** Deactivates default quotas for all GPFS file systems in the cluster. When used in combination with the **-g** option, only group quotas are deactivated. When used in combination with the **-u** or **-j** options, only user or fileset quotas, respectively, are deactivated.
- d** Resets quota limits to zero for users, groups, or filesets.

When **--profileset-quota** is not in effect for the file system, this option will reset quota limits to zero only for users, groups, or filesets that have default quotas established.

When **--profileset-quota** is in effect for the file system, this option will reset quota limits to zero for users, groups, or filesets that have default quotas established only if *both* the file system and fileset-level default quotas are zero. If either file system or fileset-level default quotas exist, the default quotas will be switched to the level that is non-zero.

If this option is not chosen, existing quota entries remain in effect.

- g** Specifies that default quotas for groups are to be deactivated.
- j** Specifies that default quotas for filesets are to be deactivated.

## **mmdefquotaoff**

- u Specifies that default quotas for users are to be deactivated.
- v Prints a message for each file system or fileset in which default quotas are deactivated.

### **Exit status**

0 Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the **mmdefquotaoff** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmdefquotaoff** command is issued.

### **Examples**

1. To deactivate default user quotas on file system **fs0**, issue this command:

```
mmdefquotaoff -u fs0
```

To confirm the change, issue this command:

```
mmlsquota -d -u fs0
```

The system displays information similar to:

| Default Block Limits(KB) |      |            | Default File Limits |       |       |
|--------------------------|------|------------|---------------------|-------|-------|
| Filesystem               | type | quota      | limit               | quota | limit |
| fs0                      | USR  | no default | limits              |       |       |

2. To deactivate default group quotas on all file systems, issue this command:

```
mmdefquotaoff -g -a
```

To confirm the change, issue this command:

```
mmlsquota -d -g
```

The system displays information similar to:

| Default Block Limits(KB) |      |            | Default File Limits |       |       |
|--------------------------|------|------------|---------------------|-------|-------|
| Filesystem               | type | quota      | limit               | quota | limit |
| fs0                      | GRP  | no default | limits              |       |       |

| Default Block Limits(KB) |      |            | Default File Limits |       |       |
|--------------------------|------|------------|---------------------|-------|-------|
| Filesystem               | type | quota      | limit               | quota | limit |
| fs1                      | GRP  | no default | limits              |       |       |

3. To deactivate both user and group default quotas for fileset **fset1** on file system **gpfs1**, issue this command:

```
mmdefquotaoff -d gpfs1:fset1
```

To confirm the change, issue this command:

```
mmlsquota -d gpfs1:fset1
```

The system displays information similar to:

| Default Block Limits(KB) |         |      |       |       | Default File Limits |       |             |
|--------------------------|---------|------|-------|-------|---------------------|-------|-------------|
| Filesystem               | Fileset | type | quota | limit | quota               | limit | entryType   |
| gpfs1                    | fset1   | USR  | 0     | 0     | 0                   | 0     | default off |
| gpfs1                    | fset1   | GRP  | 0     | 0     | 0                   | 0     | default off |

## See also

- “[mmcheckquota command](#)” on page 169
- “[mmdefedquota command](#)” on page 272
- “[mmdefquotaon command](#)” on page 278
- “[mmedquota command](#)” on page 325
- “[mmlsquota command](#)” on page 435
- “[mmquotaoff command](#)” on page 507
- “[mmrepquota command](#)” on page 517

## Location

/usr/lpp/mmfs/bin

## **mmdefquotaon**

---

### **mmdefquotaon command**

Activates default quota limit usage.

#### **Synopsis**

```
mmdefquotaon [-u] [-g] [-j] [-v] [-d] {Device [Device...] | -a}
```

or

```
mmdefquotaon [-u] [-g] [-v] [-d] {Device:Fileset ... | -a}
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmdefquotaon** command activates default quota limits for file systems and filesets. If default quota limits are not applied, new users, groups, or filesets will have a quota limit of 0, indicating no limit.

To use default quotas, the **-Q yes** option must be in effect for the file system. To use default quotas at the fileset level, the **--perfileset-quota** option must also be in effect. The **-Q yes** and **--perfileset-quota** options are specified when creating a file system with the **mmcrfs** command or changing file system attributes with the **mmchfs** command. Use the **mmlsfs** command to display the current settings of these quota options.

If none of the following options are specified, the **mmdefquotaon** command activates all default quota limits:

**-u**  
**-j**  
**-g**

If the **-a** option is not used, *Device* must be the last parameter specified.

Default quotas are established for new users, groups of users or filesets by issuing the **mmdefedquota** command. Under the **-d** option, all users without an explicitly set quota limit will have a default quota limit assigned.

#### **Parameters**

##### *Device*

The device name of the file system to have default quota values activated.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

##### *Fileset*

The name of a fileset in the file system to have default quota values activated.

#### **Options**

**-a** Activates default quotas for all GPFS file systems in the cluster. When used in combination with the **-g** option, only group quotas are activated. When used in combination with the **-u** or **-j** options, only user or fileset quotas, respectively, are activated.

**-d** Assigns default quota limits to existing users, groups, or filesets when the **mmdefedquota** command is issued.

When **--perfileset-quota** is not in effect for the file system, this option will only affect existing users, groups, or filesets with no established quota limits.

When **--perfileset-quota** is in effect for the file system, this option will affect existing users, groups, or filesets with no established quota limits, and it will also change existing users or groups that refer to default quotas at the file system level into users or groups that refer to fileset-level default quota. For more information about default quota priorities, see the topic *Default quotas* in the *IBM Spectrum Scale: Administration Guide*.

If this option is not chosen, existing quota entries remain in effect and are not governed by the default quota rules.

- g Specifies that default quotas for groups are to be activated.
- j Specifies that default quotas for filesets are to be activated.
- u Specifies that default quotas for users are to be activated.
- v Prints a message for each file system or fileset in which default quotas are activated.

## Exit status

0 Successful completion.

### nonzero

A failure has occurred.

## Security

You must have root authority to run the **mmdefquotaon** command.

The node on which you enter the command must be able to execute remote shell commands on any other administration node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmdefquotaon** command is issued.

## Examples

1. To activate default user quotas on file system **fs0**, issue this command:

```
mmdefquotaon -u fs0
```

To confirm the change, issue this command:

```
mmlsfs fs0 -Q
```

The system displays information similar to:

| flag  | value | description            |
|-------|-------|------------------------|
| <hr/> |       |                        |
| -Q    | user  | Quotas enforced        |
|       | user  | Default quotas enabled |

2. To activate default group quotas on all file systems in the cluster, issue this command:

```
mmdefquotaon -g -a
```

To confirm the change, individually for each file system, issue this command:

```
mmlsfs fs1 -Q
```

The system displays information similar to:

| flag  | value | description            |
|-------|-------|------------------------|
| <hr/> |       |                        |
| -Q    | group | Quotas enforced        |
|       | group | Default quotas enabled |

## **mmdefquotaon**

3. To activate user, group, and fileset default quotas on file system **fs2**, issue this command:

```
mmdefquotaon fs2
```

To confirm the change, issue this command:

```
mmlsfs fs2 -Q
```

The system displays information similar to:

| flag | value              | description            |
|------|--------------------|------------------------|
| -Q   | user;group;fileset | Quotas enforced        |
|      | user;group;fileset | Default quotas enabled |

4. To activate user default quota for fileset **fset1** on file system **gpfs1**, issue this command:

```
mmdefquotaon -d -u gpfs1:fset1
```

To confirm the change, issue this command:

```
mmlsquota -d gpfs1:fset1
```

The system displays information similar to:

| Filesystem | Fileset | type | quota | limit    | Default Block Limits(KB) | Default File Limits |
|------------|---------|------|-------|----------|--------------------------|---------------------|
| gpfs1      | fset1   | USR  | 0     | 31457280 | 0                        | 0 default on        |
| gpfs1      | fset1   | GRP  | 0     | 0        | 0                        | 0 default off       |

In this example, notice the entryType for user quota displays default on. To also activate group default quota for **fset1** on file system **gpfs1**, issue this command:

```
mmdefquotaon -d -g gpfs1:fset1
```

To confirm the change, issue this command:

```
mmlsquota -d gpfs1:fset1
```

The system displays information similar to :

| Filesystem | Fileset | type | quota | limit    | Default Block Limits(KB) | Default File Limits |
|------------|---------|------|-------|----------|--------------------------|---------------------|
| gpfs1      | fset1   | USR  | 0     | 31457280 | 0                        | 0 default on        |
| gpfs1      | fset1   | GRP  | 0     | 0        | 0                        | 0 default on        |

In this example, notice that the entryType for group quota also displays default on now.

## **See also**

- “mmcheckquota command” on page 169
- “mmchfs command” on page 179
- “mmcrfs command” on page 250
- “mmdefedquota command” on page 272
- “mmdefquotaoff command” on page 275
- “mmedquota command” on page 325
- “mmlsfs command” on page 411
- “mmlsquota command” on page 435
- “mmquotaoff command” on page 507
- “mmrepquota command” on page 517

## **Location**

/usr/lpp/mmfs/bin

## mmdefragfs command

Reduces disk fragmentation by increasing the number of full free blocks available to the file system.

### Synopsis

```
mmdefragfs Device [-i] [-u BlkUtilPct] [-P PoolName]
[-N {Node[,Node...] | NodeFile | NodeClass}] [--qos QoSClass]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmdefragfs** command to reduce fragmentation of a file system. The **mmdefragfs** command moves existing file system data within a disk to make more efficient use of disk blocks. The data is migrated to unused sub-blocks in partially allocated blocks, thereby increasing the number of free full blocks.

The **mmdefragfs** command can be run against a mounted or unmounted file system. However, best results are achieved when the file system is unmounted. When a file system is mounted, allocation status may change causing retries to find a suitable unused sub-block.

**Note:** On a file system that has a very low level of fragmentation, negative numbers can be seen in the output of **mmdefragfs** for free sub-blocks. This indicates that the block usage has in fact increased after running the **mmdefragfs** command. If negative numbers are seen, it does not indicate a problem and you do not need to rerun the **mmdefragfs** command.

### Parameters

#### *Device*

The device name of the file system to have fragmentation reduced. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

#### **-P** *PoolName*

Specifies the pool name to use.

#### **-N** {*Node[,Node...]* | *NodeFile* | *NodeClass*}

Specifies the nodes that can be used in this disk defragmentation. This parameter supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

#### **--qos** *QoSClass*

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos** command” on page 206. Specify one of the following QoS classes:

#### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

## **mmdefragfs**

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

## **Options**

**-i** Specifies to query the current disk fragmentation state of the file system. Does not perform the actual defragmentation of the disks in the file system.

**-u BlkUtilPct**

The average block utilization goal for the disks in the file system. The **mmdefragfs** command reduces the number of allocated blocks by increasing the percent utilization of the remaining blocks. The command automatically goes through multiple iterations until *BlkUtilPct* is achieved on all of the disks in the file system or until no progress is made in achieving *BlkUtilPct* from one iteration to the next, at which point it exits.

## **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmdefragfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To query the fragmentation state of file system **fs0**, issue this command:

```
mmdefragfs fs0 -i
```

The system displays information similar to:

| disk<br>name | disk<br>size<br>in<br>nSubblk | in<br>full<br>blocks | subblk<br>in<br>fragments | %<br>free | %<br>blk<br>util |
|--------------|-------------------------------|----------------------|---------------------------|-----------|------------------|
| nsd32        | 327680                        | 277504               | 12931                     | 84.688    | 96.054           |
| nsd33        | 327680                        | 315232               | 580                       | 96.201    | 99.823           |
| nsd21        | 327680                        | 301824               | 2481                      | 92.109    | 99.243           |
| nsd34        | 327680                        | 275904               | 13598                     | 84.199    | 95.850           |
| nsd30        | 327680                        | 275808               | 13380                     | 84.170    | 95.917           |
| nsd19        | 327680                        | 278496               | 12369                     | 84.990    | 96.225           |
| nsd31        | 327680                        | 276224               | 12012                     | 84.297    | 96.334           |
| (total)      | 2293760                       | 2000992              | 67351                     |           | 97.064           |

2. To reduce fragmentation of the file system **fs0** on all defined, accessible disks that are not stopped or suspended, issue this command:

```
mmdefragfs fs0
```

The system displays information similar to:

| disk<br>name | free<br>in full<br>blocks | subblk<br>in<br>blk<br>before freed | free<br>subblk<br>before after | %<br>fragments<br>free blk<br>after before | %<br>blk util<br>before after |
|--------------|---------------------------|-------------------------------------|--------------------------------|--------------------------------------------|-------------------------------|
|              |                           |                                     |                                |                                            |                               |

|           |       |       |    |      |     |       |       |       |       |
|-----------|-------|-------|----|------|-----|-------|-------|-------|-------|
| gpfs57nsd | 28896 | 29888 | 31 | 1462 | 463 | 50.39 | 52.12 | 94.86 | 98.31 |
| gpfs60nsd | 41728 | 43200 | 46 | 1834 | 362 | 59.49 | 61.59 | 93.55 | 98.66 |
| (total)   | 70624 | 73088 | 77 | 3296 | 825 |       |       | 93.63 | 98.84 |

3. To reduce fragmentation of all files in the **fs1** file system until the disks have 100% full block utilization, issue this command:

```
mmdefragfs fs1 -u 100
```

The system displays information similar to:

```
Defragmenting file system 'fs1'...
```

```
Defragmenting until full block utilization is 98.00%, currently 97.07%
27.35 % complete on Tue May 26 14:25:42 2009 (617882 inodes 4749 MB)
82.65 % complete on Tue May 26 14:26:02 2009 (1867101 inodes 10499 MB)
89.56 % complete on Tue May 26 14:26:23 2009 (2023206 inodes 14296 MB)
90.01 % complete on Tue May 26 14:26:43 2009 (2033337 inodes 17309 MB)
90.28 % complete on Tue May 26 14:27:03 2009 (2039551 inodes 19779 MB)
91.17 % complete on Tue May 26 14:27:23 2009 (2059629 inodes 23480 MB)
91.67 % complete on Tue May 26 14:27:43 2009 (2070865 inodes 26760 MB)
92.51 % complete on Tue May 26 14:28:03 2009 (2089804 inodes 29769 MB)
93.12 % complete on Tue May 26 14:28:23 2009 (2103697 inodes 32649 MB)
93.39 % complete on Tue May 26 14:28:43 2009 (2109629 inodes 34934 MB)
95.47 % complete on Tue May 26 14:29:04 2009 (2156805 inodes 36576 MB)
95.66 % complete on Tue May 26 14:29:24 2009 (2160915 inodes 38705 MB)
95.84 % complete on Tue May 26 14:29:44 2009 (2165146 inodes 40248 MB)
96.58 % complete on Tue May 26 14:30:04 2009 (2181719 inodes 41733 MB)
96.77 % complete on Tue May 26 14:30:24 2009 (2186053 inodes 43022 MB)
96.99 % complete on Tue May 26 14:30:44 2009 (2190955 inodes 43051 MB)
97.20 % complete on Tue May 26 14:31:04 2009 (2195726 inodes 43077 MB)
97.40 % complete on Tue May 26 14:31:24 2009 (2200378 inodes 43109 MB)
97.62 % complete on Tue May 26 14:31:44 2009 (2205201 inodes 43295 MB)
97.83 % complete on Tue May 26 14:32:05 2009 (2210003 inodes 43329 MB)
97.85 % complete on Tue May 26 14:32:25 2009 (2214741 inodes 43528 MB)
97.86 % complete on Tue May 26 14:32:55 2009 (2221888 inodes 43798 MB)
97.87 % complete on Tue May 26 14:33:35 2009 (2231453 inodes 44264 MB)
97.88 % complete on Tue May 26 14:34:26 2009 (2243181 inodes 45288 MB)
100.00 % complete on Tue May 26 14:35:10 2009
```

| disk<br>name | free subblk       |         |       | free                   |       |        | % util   |       |       |
|--------------|-------------------|---------|-------|------------------------|-------|--------|----------|-------|-------|
|              | in full<br>blocks |         | blk   | subblk in<br>fragments |       | %      | free blk | blk   | util  |
|              | before            | after   | freed | before                 | after | before | before   | after |       |
| nsd32        | 277504            | 287840  | 323   | 12931                  | 2183  | 84.69  | 87.84    | 96.05 | 99.33 |
| nsd33        | 315232            | 315456  | 7     | 580                    | 185   | 96.20  | 96.27    | 99.82 | 99.94 |
| nsd21        | 301824            | 303616  | 56    | 2481                   | 666   | 92.11  | 92.66    | 99.24 | 99.80 |
| nsd34        | 275904            | 285920  | 313   | 13598                  | 3159  | 84.20  | 87.26    | 95.85 | 99.04 |
| nsd30        | 275840            | 285856  | 313   | 13348                  | 2923  | 84.18  | 87.24    | 95.93 | 99.11 |
| nsd19        | 278592            | 288832  | 320   | 12273                  | 1874  | 85.02  | 88.14    | 96.25 | 99.43 |
| nsd31        | 276224            | 284608  | 262   | 12012                  | 3146  | 84.30  | 86.86    | 96.33 | 99.04 |
| (total)      | 2001120           | 2052128 | 1594  | 67223                  | 14136 |        |          | 97.07 | 99.38 |

Defragmentation complete, full block utilization is 99.04%.

## See also

- “mmdf command” on page 308

## Location

/usr/lpp/mmfs/bin

## **mmdelacl**

---

### **mmdelacl command**

Deletes a GPFS access control list.

#### **Synopsis**

**mmdelacl [-d] *Filename***

#### **Availability**

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

#### **Description**

Use the **mmdelacl** command to delete the extended entries of an access ACL of a file or directory, or to delete the default ACL of a directory.

#### **Parameters**

##### *Filename*

The path name of the file or directory for which the ACL is to be deleted. If the **-d** option is specified, *Filename* must contain the name of a directory.

#### **Options**

**-d** Specifies that the default ACL of a directory is to be deleted.

Since there can be only one NFS V4 ACL (no separate default), specifying the **-d** flag for a file with an NFS V4 ACL is an error. Deleting an NFS V4 ACL necessarily removes both the ACL and any inheritable entries contained in it.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

The **mmdelacl** command may be issued only by the file or directory owner, the root user, or by someone with control (c) authority in the ACL for the file.

You may issue the **mmdelacl** command only from a node in the GPFS cluster where the file system is mounted.

#### **Examples**

To delete the default ACL for a directory named **project2**, issue this command:

```
mmdelacl -d project2
```

To confirm the deletion, issue this command:

```
mmgetacl -d project2
```

The system displays information similar to:

```
#owner:uno
#group:system
```

**See also**

- “[mmeditacl command](#)” on page 322
- “[mmgetacl command](#)” on page 347
- “[mmputacl command](#)” on page 504

**Location**

/usr/lpp/mmfs/bin

## **mmdelcallback**

---

### **mmdelcallback command**

Deletes one or more user-defined callbacks from the GPFS system.

#### **Synopsis**

```
mmdelcallback CallbackIdentifier[,CallbackIdentifier...]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmdelcallback** command to delete one or more user-defined callbacks from the GPFS system.

#### **Parameters**

##### *CallbackIdentifier*

Specifies a user-defined unique name that identifies the callback to be deleted. Use the **mmlscallback** command to see the name of the callbacks that can be deleted.

- |     **Note:** Before you add or delete the **tiebreakerCheck** event, you must stop the GPFS daemon on all the nodes in the cluster.

#### **Exit status**

0     Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmdelcallback** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

#### **Examples**

To delete the **test1** callback from the GPFS system, issue this command:

```
mmdelcallback test1
```

The system displays information similar to:

```
mmdelcallback: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

#### **See also**

- “**mmaddcallback command**” on page 10
- “**mmlscallback command**” on page 395

#### **Location**

```
/usr/lpp/mmfs/bin
```

---

## mmdeldisk command

Deletes disks from a GPFS file system.

### Synopsis

```
mmdeldisk Device {"DiskName[;DiskName...]" | -F DescFile} [-a] [-c]
[-m | -r | -b] [-N {Node[,Node...] | NodeFile | NodeClass}]
[--inode-criteria CriteriaFile] [-o InodeResultFile]
[--qos QOSClass]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmdeldisk** command migrates all data that would otherwise be lost to the remaining disks in the file system. It then removes the disks from the file system descriptor, preserves replication at all times, and optionally rebalances the file system after removing the disks.

The **mmdeldisk** command has the following two functions:

- Copying unreplicated data off the disks and removing references to the disks (**deldisk** step).
- Rereplicating or rebalancing blocks across the remaining disks (**restripe** step).

These two functions can be done in one pass over the file system, or in two passes if the **-a** option is specified.

Run the **mmdeldisk** command when system demand is low.

If a replacement for a failing disk is available, use the **mmrpldisk** command in order to keep the file system balanced. Otherwise, use one of these procedures to delete a disk:

- If the file system is replicated, replica copies can be preserved at all times by using the default **-r** option or the **-b** option.
- Using the **-m** option will not preserve replication during the **deldisk** step because it will only copy the minimal amount of data off the disk being deleted so that every block has at least one copy. Also, using the **-a** option will not preserve replication during the **deldisk** step, but will then re-establish replication during the subsequent **restripe** step.
- If you want to move all data off the disk before running **mmdeldisk**, use **mmchdisk** to suspend all the disks that will be deleted and run **mmrestripefs** with the **-r** or **-b** option. This step is no longer necessary, now that **mmdeldisk** does the same function. If **mmdeldisk** fails (or is canceled), it leaves the disks in the suspended state, and **mmdeldisk** can be retried when the problem that caused **mmdeldisk** to stop is corrected.
- If the disk is permanently damaged and the file system is not replicated, or if the **mmdeldisk** command repeatedly fails, see the *IBM Spectrum Scale: Problem Determination Guide* and search for *Disk media failure*.

If the last disk in a storage pool is deleted, the storage pool is deleted. The **mmdeldisk** command is not permitted to delete the **system** storage pool. A storage pool must be empty in order for it to be deleted.

### Results

Upon successful completion of the **mmdeldisk** command, these tasks are completed:

- Data that has not been replicated from the target disks is migrated to other disks in the file system.
- Remaining disks are rebalanced, if specified.

## **mmdeldisk**

### **Parameters**

#### *Device*

The device name of the file system to delete the disks from. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**. This must be the first parameter.

#### **"DiskName[;DiskName...]"**

Specifies the names of the disks to be deleted from the file system. If there is more than one disk to be deleted, delimit each name with a semicolon (;) and enclose the list in quotation marks.

#### **-F DiskFile**

Specifies a file that contains the names of the disks (one name per line), to be deleted from the GPFS cluster.

#### **-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes that participate in the restripe of the file system after the specified disks have been removed. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

#### **--inode-criteria CriteriaFile**

Specifies the interesting inode criteria flag, where *CriteriaFile* is one of the following:

##### **BROKEN**

Indicates that a file has a data block with all of its replicas on disks that have been removed.

**Note:** **BROKEN** is always included in the list of flags even if it is not specified.

##### **dataUpdateMiss**

Indicates that at least one data block was not updated successfully on all replicas.

##### **exposed**

Indicates an inode with an exposed risk; that is, the file has data where all replicas are on suspended disks. This could cause data to be lost if the suspended disks have failed or been removed.

##### **i11Compressed**

Indicates an inode in which file compression or decompression is deferred, or in which a compressed file is partly decompressed to allow the file to be written into or memory-mapped.

##### **i11Placed**

Indicates an inode with some data blocks that might be stored in an incorrect storage pool.

##### **i11Replicated**

Indicates that the file has a data block that does not meet the setting for the replica.

##### **metaUpdateMiss**

Indicates that there is at least one metadata block that has not been successfully updated to all replicas.

##### **unbalanced**

Indicates that the file has a data block that is not well balanced across all the disks in all failure groups.

**Note:** If a file matches *any* of the specified interesting flags, all of its interesting flags (even those not specified) will be displayed.

#### **-o InodeResultFile**

Contains a list of the inodes that met the interesting inode flags that were specified on the **--inode-criteria** parameter. The output file contains the following:

**INODE\_NUMBER**

This is the inode number.

**DISKADDR**

Specifies a dummy address for later **tsfindinode** use.

**SNAPSHOT\_ID**

This is the snapshot ID.

**ISGLOBAL\_SNAPSHOT**

Indicates whether or not the inode is in a global snapshot. Files in the live file system are considered to be in a global snapshot.

**INDEPENDENT\_FSETID**

Indicates the independent files set to which the inode belongs.

**MEMO (INODE\_FLAGS FILE\_TYPE [ERROR])**

Indicates the inode flag and file type that will be printed:

**Inode flags:**

BROKEN  
exposed  
dataUpdateMiss  
i11Compressed  
i11Placed  
i11Replicated  
metaUpdateMiss  
unbalanced

**File types:**

BLK\_DEV  
CHAR\_DEV  
DIRECTORY  
FIFO  
LINK  
LOGFILE  
REGULAR\_FILE  
RESERVED  
SOCK  
\*UNLINKED\*  
\*DELETED\*

**Notes:**

1. An error message will be printed in the output file if an error is encountered when repairing the inode.
2. **DISKADDR**, **ISGLOBAL\_SNAPSHOT**, and **FSET\_ID** work with the **tsfindinode** tool (`/usr/lpp/mmfs/bin/tsfindinode`) to find the file name for each inode. **tsfindinode** uses the output file to retrieve the file name for each interesting inode.

**--qos QOSclass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

- other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

## **mmdeldisk**

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **Options**

- a** Specifies that the **mmdeldisk** command *not* wait for rereplicating or rebalancing to complete before returning. When this flag is specified, the **mmdeldisk** command runs asynchronously and returns after the file system descriptor is updated and the rebalancing scan is started, but it does not wait for rebalancing to finish. If no rebalancing is requested (-**r** option is not specified), this option has no effect.

If -**m** is specified, this option has no effect. If -**r** or -**b** is specified (no option defaulting to -**r**), then the **deldisk** step is done using -**m**, and the **restripe** step is done using the specified option.

- b** Rebalances the blocks onto the other disks while moving data off the disks being deleted. This might have to move much more data than the -**r** operation.

**Note:** Rebalancing of files is an I/O intensive and time consuming operation, and is important only for file systems with large files that are mostly invariant. In many cases, normal file update and creation will rebalance your file system over time, without the cost of the rebalancing.

- c** Specifies that processing continues even in the event that unreadable data exists on the disks being deleted. Data that has not been replicated is lost. Replicated data is not lost as long as the disks containing the replication are accessible.
- m** Does minimal data copying to preserve any data that is located only on the disks being removed. This is the fastest way to get a disk out of the system, but it could reduce replication of some blocks of the files and metadata.

**Note:** This might be I/O intensive if there is a lot of data to be copied or rereplicated off the disks that are being deleted.

- r** Preserves replication of all files and metadata during the **mmdeldisk** operation (except when the -**a** option is specified). This is the default.

**Note:** This might be I/O intensive if there is a lot of data to be copied or rereplicated off the disks that are being deleted.

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the **mmdeldisk** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Example**

To delete **gpfs1016nsd** from file system **fs1** and rebalance the files across the remaining disks, issue this command:

```
mmdeldisk fs1 gpfs1016nsd
```

The system displays information similar to:

```
Deleting disks ...
Scanning sp1 storage pool
Scanning user file metadata ...
100.00 % complete on Tue Mar 13 15:48:51 2012
Scan completed successfully.
Checking Allocation Map for storage pool 'sp1'
tsdeldisk64 completed.
mmdeldisk: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

## See also

- “[mmadddisk command](#)” on page 23
- “[mmchdisk command](#)” on page 161
- “[mmlsdisk command](#)” on page 402
- “[mmrpldisk command](#)” on page 539

## Location

/usr/lpp/mmfs/bin

## **mmdelfilesset**

---

### **mmdelfilesset command**

Deletes a GPFS fileset.

#### **Synopsis**

```
mmdelfilesset Device FilesetName [-f] [--qos QoSClass]
```

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

The **mmdelfilesset** command deletes a GPFS fileset. When deleting a fileset, consider these points:

- The root fileset cannot be deleted.
- A fileset that is not empty cannot be deleted unless the **-f** flag is specified.
- A fileset that is currently linked into the namespace cannot be deleted until it is unlinked with the **mmunlinkfilesset** command.
- A dependent fileset can be deleted at any time.
- An independent fileset cannot be deleted if it has any dependent filesets or fileset snapshots.
- Deleting a dependent fileset that is included in a fileset or global snapshot removes it from the active file system, but it remains part of the file system in a deleted state.
- Deleting an independent fileset that is included in any global snapshots removes it from the active file system, but it remains part of the file system in a deleted state.
- A fileset in the deleted state is displayed in the **mmlsfilesset** output with the fileset name in parenthesis. If the **-L** flag is specified, the latest including snapshot is also displayed. The **--deleted** option of the **mmlsfilesset** command can be used to display only deleted filesets.
- The contents of a deleted fileset are still available in the snapshot, through some path name containing a **.snapshots** component, because it was saved when the snapshot was created.
- When the last snapshot that includes the fileset has been deleted, the fileset is fully removed from the file system.

For information on GPFS filesets, see *Information Lifecycle Management for GPFS in IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system that contains the fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

##### *FilesetName*

Specifies the name of the fileset to be deleted.

**-f** Forces the deletion of the fileset. All fileset contents are deleted. Any child filesets are first unlinked.

##### **--qos** *QoSClass*

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos** command” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

- other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Security**

You must have root authority to run the **mmdefilesset** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

**Examples**

1. This sequence of commands illustrates what happens when attempting to delete a fileset that is linked.

- a. Command:

```
mmdefilesset gpfs1
```

The system displays output similar to:

| Name  | Status   | Path         |
|-------|----------|--------------|
| root  | Linked   | /gpfs1       |
| fset1 | Linked   | /gpfs1/fset1 |
| fset2 | Unlinked | --           |

- b. Command:

```
mmdefilesset gpfs1 fset1
```

The system displays output similar to:

Fileset fset1 must be unlinked to be deleted.

- c. Command:

```
mmdefilesset gpfs1 fset2
```

The system displays output similar to:

```
Checking fileset ...
Checking fileset complete.
Deleting fileset ...
Fileset 'fset2' deleted.
```

- d. To confirm the change, issue this command:

```
mmdefilesset gpfs1
```

The system displays output similar to:

| Name  | Status | Path         |
|-------|--------|--------------|
| root  | Linked | /gpfs1       |
| fset1 | Linked | /gpfs1/fset1 |

## **mmdelfilesset**

2. This sequence of commands illustrates what happens when attempting to delete a fileset that contains user files.

- a. Command:

```
mmlsfilesset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |          |              |
|----------------------------------|----------|--------------|
| Name                             | Status   | Path         |
| root                             | Linked   | /gpfs1       |
| fset1                            | Linked   | /gpfs1/fset1 |
| fset2                            | Unlinked | --           |

- b. Command:

```
mmdelfilesset gpfs1 fset2
```

The system displays output similar to:

Fileset 'fset2' contains user files,  
but can be deleted with the "-f" option.

- c. Command:

```
mmdelfilesset gpfs1 fset2 -f
```

The system displays output similar to:

```
Checking fileset ...
Checking fileset complete.
Deleting user files ...
100.00 % complete on Wed Feb 15 11:38:05 2012
Deleting fileset ...
Fileset 'fset2' deleted.
```

- d. To confirm the change, issue this command:

```
mmlsfilesset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |        |              |
|----------------------------------|--------|--------------|
| Name                             | Status | Path         |
| root                             | Linked | /gpfs1       |
| fset1                            | Linked | /gpfs1/fset1 |

## **See also**

- “mmchfileset command” on page 173
- “mmcrfileset command” on page 244
- “mmlinkfileset command” on page 390
- “mmlsfilesset command” on page 406
- “mmunlinkfileset command” on page 578

## **Location**

/usr/lpp/mmfs/bin

---

## mmdelfs command

Removes a GPFS file system.

### Synopsis

**mmdelfs** *Device* [-p]

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmdelfs** command removes all the structures for the specified file system from the nodes in the cluster.

Before you can delete a file system using the **mmdelfs** command, you must unmount it on all nodes.

### Results

Upon successful completion of the **mmdelfs** command, these tasks are completed on all nodes:

- Deletes the character device entry from **/dev**.
- Removes the mount point directory where the file system had been mounted.

### Parameters

#### *Device*

The device name of the file system to be removed. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

- p Indicates that the disks are permanently damaged and the file system information should be removed from the GPFS cluster data even if the disks cannot be marked as **available**.

### Exit status

0 Successful completion.

#### nonzero

A failure has occurred.

### Security

You must have root authority to run the **mmdelfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### Examples

To delete file system **fs0**, issue this command:

```
mmdelfs fs0
```

The system displays information similar to:

## **mmdelfs**

```
mmdelfs: 6027-1366 Marking the disks as available
GPFS: 6027-573 All data on the following disks of fs0 will be destroyed:
 gpfs9nsd
 gpfs10nsd
 gpfs15nsd
 gpfs17nsd
GPFS: 6027-574 Completed deletion of file system fs0.
mmdelfs: 6027-1371 Propagating the cluster configuration data to all
 affected nodes. This is an asynchronous process.
```

### **See also**

- “[mmcrfs command](#)” on page 250
- “[mmchfs command](#)” on page 179
- “[mmlsfs command](#)” on page 411

### **Location**

/usr/lpp/mmfs/bin

---

## mmdeinode command

Removes one or more nodes from a GPFS cluster.

### Synopsis

```
mmdeinode {-a | -N Node[,Node...] | NodeFile | NodeClass}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmdeinode** command to delete one or more nodes from the GPFS cluster. You may issue the **mmdeinode** command on any GPFS node.

A node cannot be deleted if any of the following are true:

1. It is a primary or secondary GPFS cluster configuration server.

The node being deleted cannot be the primary or secondary GPFS cluster configuration server unless you intend to delete the entire cluster.

You can determine whether a node is the primary or secondary configuration server by issuing the **mmlscluster** command. If the node is listed as of the servers and you still want to delete it without deleting the cluster, first use the **mmchcluster** command to assign another node as the server.

2. It is defined as an NSD server.

The node being deleted cannot be defined as an NSD server for any disk unless you intend to delete the entire cluster.

You can determine whether a node is an NSD server for one or more disks by issuing the **mmlsnsd** command. If the node is listed as an NSD server and you still want to delete it without deleting the cluster, first use the **mmchnsd** command to assign another node as an NSD server for the affected disks.

3. If the GPFS state is *unknown* and the node is reachable on the network.

You cannot delete a node if both of the following are true:

- The node responds to a TCP/IP ping command from another node.
- The status of the node shows *unknown* when you use the **mmgetstate** command from another node in the cluster.

**Note:** You will probably be able to delete such a node if you physically power it off.

4. If the node is defined as a Transparent cloud tiering node. You can determine whether a node is a Transparent cloud tiering node by issuing the **mmccloudgateway node list** command. If the node is listed as the Transparent cloud tiering node, and you still want to delete it without deleting the cluster, first use the **mmchnode** command to disable the Transparent cloud tiering node role.

- | • If the node is a Transparent cloud tiering node, disable Transparent cloud tiering from the node by using the **mmchnode --cloud-gateway-disable** command, and then uninstall the Transparent cloud tiering rpms. Doing so ensures that the **mmdeinode** command does not fail on a Transparent cloud tiering node.

You must follow these rules when deleting nodes:

1. Unless all nodes in the cluster are being deleted, run the **mmdeinode** command from a node that will remain in the cluster.
2. Before you can delete a node, unmount all of the GPFS file systems and stop GPFS on the node to be deleted.

## **mmdelnode**

3. Exercise caution when shutting down GPFS on quorum nodes. If the number of remaining quorum nodes falls below the requirement for a quorum, you will be unable to perform file system operations. For more information, see *Quorum in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

Each GPFS cluster is managed independently, so there is no automatic coordination and propagation of changes between clusters like there is between the nodes within a cluster. This means that if you permanently delete nodes that are being used as contact nodes by other GPFS clusters that can mount your file systems, you should notify the administrators of those GPFS clusters so that they can update their own environments.

## **Results**

Upon successful completion of the **mmdelnode** command, the specified nodes are deleted from the GPFS cluster.

## **Parameters**

**-a** Delete all nodes in the cluster.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the set of nodes to be deleted from the cluster.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmdelnode** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

You may issue the **mmdelnode** command from any node that will remain in the GPFS cluster.

## **Examples**

1. To delete all of the nodes in the cluster, issue this command:

```
mmdelnode -a
```

The system displays information similar to:

```
Verifying GPFS is stopped on all affected nodes ...
mmdelnode: Command successfully completed
mmdelnode: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

2. To delete nodes **k145n12**, **k145n13**, and **k145n14**, issue this command:

```
mmdelnode -N k145n12,k145n13,k145n14
```

The system displays information similar to:

```
Verifying GPFS is stopped on all affected nodes ...
mmdelnod: Command successfully completed
mmdelnod: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

## See also

- “[mmaddnode command](#)” on page 29
- “[mmcrcluster command](#)” on page 239
- “[mmchconfig command](#)” on page 132
- “[mmlsfs command](#)” on page 411
- “[mmlscluster command](#)” on page 397

## Location

/usr/lpp/mmfs/bin

## **mmdelnodaclass**

---

### **mmdelnodaclass command**

Deletes user-defined node classes.

#### **Synopsis**

```
mmdelnodaclass ClassName[,ClassName...]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmdelnodaclass** command to delete existing user-defined node classes.

#### **Parameters**

*ClassName*

Specifies an existing user-defined node class to delete.

If *ClassName* was used to change configuration attributes with **mmchconfig**, and the configuration attributes are still referencing *ClassName*, then *ClassName* cannot be deleted. Use the **mmchconfig** command to remove the references to *ClassName* before deleting this user-defined node class.

#### **Exit status**

0      Successful completion.

nonzero

A failure has occurred.

#### **Security**

You must have root authority to run the **mmdelnodaclass** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

#### **Examples**

To display the current user-defined node classes, issue this command:

```
mmlsnodaclass --user
```

The system displays information similar to:

| Node Class Name | Members             |
|-----------------|---------------------|
| <hr/>           |                     |
| siteA           | c6f1c3vp1,c6f1c3vp2 |
| siteB           | c6f1c3vp4,c6f1c3vp5 |

To delete the **siteA** node class, issue this command:

```
mmdelnodaclass siteA
```

The system displays information similar to:

```
mmdelnodaclass: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

To display the updated list of user-defined node classes, issue this command:

```
mmclsnodeclass --user
```

The system displays information similar to:

| Node Class Name | Members             |
|-----------------|---------------------|
| siteB           | c6f1c3vp4,c6f1c3vp5 |

## See also

- “[mmcrnodeclass command](#)” on page 260
- “[mmchnodeclass command](#)” on page 195
- “[mmclsnodeclass command](#)” on page 421

## Location

/usr/lpp/mmfs/bin

## **mmdelnsd**

---

### **mmdelnsd command**

Deletes Network Shared Disks (NSDs) from the GPFS cluster.

#### **Synopsis**

```
mmdelnsd {"DiskName[;DiskName...]" | -F DiskFile}
```

or

```
mmdelnsd -p NSDId [-N Node[,Node...]]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmdelnsd** command serves two purposes:

1. To delete NSDs from the GPFS cluster.
2. To remove the unique NSD volume ID left on a disk after the failure of a previous invocation of the **mmdelnsd** command. The NSD had been successfully deleted from the GPFS cluster but there was a failure to clear the unique volume ID from the disk.

NSDs being deleted cannot be part of any file system. To determine if an NSD belongs to a file system or not, issue the **mmlsnsd -d DiskName** command. If an NSD belongs to a file system, either the **mmdeldisk** or the **mmdelfs** command must be issued prior to deleting the NSDs from the GPFS cluster.

NSDs being deleted cannot be tiebreaker disks. To list the tiebreaker disks, issue the **mmlsconfig tiebreakerDisks** command. Use the **mmchconfig** command to assign new tiebreaker disks prior to deleting NSDs from the cluster. For information on tiebreaker disks, see *Quorum in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

#### **Results**

Upon successful completion of the **mmdelnsd** command, these tasks are completed:

- All references to the disks are removed from the GPFS cluster data.
- Each disk is cleared of its unique NSD volume ID.
- On Windows, the disk's GPT partition table is removed leaving the disk Unknown/Not Initialized.

#### **Parameters**

**DiskName[;DiskName...]**

Specifies the names of the NSDs to be deleted from the GPFS cluster. Specify the names generated when the NSDs were created. Use the **mmlsnsd -F** command to display disk names. If there is more than one disk to be deleted, delimit each name with a semicolon (;) and enclose the list of disk names in quotation marks.

**-F DiskFile**

Specifies a file containing the names of the NSDs, one per line, to be deleted from the GPFS cluster.

**-N Node[,Node...]**

Specifies the nodes to which the disk is attached. If no nodes are listed, the disk is assumed to be directly attached to the local node.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands in the IBM Spectrum Scale: Administration Guide*.

**-p NSDID**

Specifies the NSD volume ID of an NSD that needs to be cleared from the disk as indicated by the failure of a previous invocation of the **mmdelnsd** command.

**Exit status**

0      Successful completion.

**nonzero**

A failure has occurred.

**Security**

You must have root authority to run the **mmdelnsd** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

**Examples**

1. To delete **gpfs47nsd** from the GPFS cluster, issue this command:

```
mmdelnsd "gpfs47nsd"
```

The system displays output similar to:

```
mmdelnsd: Processing disk gpfs47nsd
mmdelnsd: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

2. If after running **mmdelnsd** to delete an NSD, you experience a failure, the disk was not found. Run **mmdelnsd -p NSD Volume ID**. For example:

```
mmdelnsd -p COA8910B626630E
```

This will remove the NSD definition from the GPFS configuration even if the NSD ID is not removed from the physical disk because it has been permanently lost.

**See also**

- “**mmcrnsd** command” on page 262
- “**mmlsnsd** command” on page 423

**Location**

/usr/lpp/mmfs/bin

## mmdelsnapshot

---

### mmdelsnapshot command

Deletes a GPFS snapshot.

#### Synopsis

```
mmdelsnapshot Device [[Fileset]:]Snapshot[[[Fileset]:]Snapshot...]
[-j FilesetName[,FilesetName...]] [--qos QOSClass]
[-N{all | mount | Node[,Node...]|NodeFile | NodeClass}
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

Use the **mmdelsnapshot** command to delete a GPFS snapshot.

Once the command is issued, the snapshot is marked for deletion and cannot be recovered.

If the node from which the **mmdelsnapshot** command is issued or the file system manager node fails, the snapshot might not be completely deleted. The **mmlsnapshot** command shows these snapshots with status DeleteRequired. Reissue the **mmdelsnapshot** from another node to complete the removal, or allow the snapshot to be cleaned up automatically by a later **mmdelsnapshot** command. A snapshot in this state cannot be accessed.

Any files open in the snapshot are forcibly closed. The user receives an **errno** of ESTALE on the next file access.

If a snapshot has file clones, you must delete the file clones or split them from their clone parents before you delete the snapshot. Use the **mmclone split** or **mmclone redirect** command to split file clones. Use a regular delete (**rm**) command to delete a file clone. If a snapshot is deleted that contains a clone parent, any attempts to read a block that refers to the missing snapshot returns an error. A policy file can be created to help determine whether a snapshot has file clones. See the *IBM Spectrum Scale: Administration Guide* for more information about file clones and policy files.

In IBM Spectrum Scale 4.2.1 and later, snapshot commands support the specification of multiple snapshots. Users can easily delete multiple snapshots for maintenance and cleanup operations. Also, system performance is increased by batching operations and reducing overhead.

In this release, the following new usages of the **mmdelsnapshot** command have been introduced:

```
mmdelsnapshot fs [[Fileset]:]Snapshot[[[Fileset]:]Snapshot...]
mmdelsnapshot fs [[Fileset]:]Snapshot[[[Fileset]:]Snapshot...]] -j Fileset
mmdelsnapshot fs Snapshot -j Fileset1,Fileset2,...
```

#### Parameters

##### *Device*

The device name of the file system for which the snapshot is to be deleted. File system names do not need to be fully qualified.

##### *Fileset*

Specifies the name of the fileset that contains the fileset snapshot to be deleted. If *Fileset* is not specified, the **mmdelsnapshot** command deletes a global snapshot named *Snapshot*.

**Note:** Ensure that multiple snapshots and multiple filesets are not used together.

**Snapshot**

Specifies the name of the snapshot to be deleted.

The snapshot names are separated by a comma.

The snapshot specifier describes global and fileset snapshots. For example, *Fileset1:Snapshot1* specifies a fileset snapshot named Snapshot1 for fileset Fileset1. If *Fileset1* is empty, *Snapshot1* is a global snapshot named Snapshot1.

**Note:** Ensure that the snapshot name does not include a colon (:), a comma (,), and whitespaces. If the snapshot name consists of a whitespace, the whitespace must be quoted and part of the snapshot name. For example, if the snap1 and snap A snapshots must be deleted, use the command  
**'mmdelsnapshot pk2 "snap A, snap1"'.**

**-j *FilesetName***

Specifies the name of the fileset that contains the fileset snapshot to be deleted (*SnapshotName*). If **-j** is not specified, the **mmdelsnapshot** command attempts to delete a global snapshot named *SnapshotName*.

**Note:** When a list of snapshots separated by a comma (,) is used with the **-j** option, the fileset is applicable to each snapshot that does not use the colon (:) syntax. The fileset name must not consist of a white space.

**-N{all | mount | Node[,Node...]|NodeFile | NodeClass}**

Specifies the nodes that participate in deleting the snapshot. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

**--qos *QoSClass***

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

**Exit status**

**0** Successful completion.

**Nonzero**

A failure occurred.

**Security**

You must have root authority to run the **mmdelsnapshot** command when you delete global snapshots.

## **mmdelsnapshot**

Independent fileset owners can run the **mmdelsnapshot** command to delete snapshots of filesets that they own.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. To delete the snapshot **snap1** for the file system **fs1**, run the following command:

```
mmdelsnapshot fs1 snap1
```

The output is similar to the following example:

```
Invalidating snapshot files...
Deleting snapshot files...
100.00 % complete on Tue Feb 28 10:40:59 2012
Delete snapshot snap1 complete, err = 0
```

Before you issue the command, the directory might have the following structure:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3
```

After you issue the command, the directory has the following structure:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3
/fs1/.snapshots
```

2. To delete the **snap1** snapshot from the **fset1**, **fset2**, and **fset3** filesets for the file system **fs1**, run the following command:

```
mmdelsnapshot fs1 snap1 -j fset1,fset2,fset3
```

The system displays the following output:

```
Invalidating snapshot files in fset1:snap1 fset2:snap1 fset3:snap1 (1..3) of 3...
Deleting files in snapshot fset1:snap1...
100.00 % complete on Thu May 12 04:43:07 2016 (100032 inodes with total 0 MB data processed)
Invalidate snapshot files in fset1:snap1/F/...
Deleting files in snapshot fset2:snap1...
100.00 % complete on Thu May 12 04:43:07 2016 (100032 inodes with total 0 MB data processed)
Invalidate snapshot files in fset2:snap1/F/...
Deleting files in snapshot fset3:snap1...
100.00 % complete on Thu May 12 04:43:07 2016 (100032 inodes with total 0 MB data processed)
Invalidate snapshot files in fset3:snap1/F/...
Delete snapshot fset1:snap1 successful.
Delete snapshot fset2:snap1 successful.
Delete snapshot fset3:snap1 successful.
```

3. To specify the snapshot names that must be deleted from each fileset in the file system **fs1**, run the following command:

```
mmdelsnapshot fs1 fset1:snapA,fset2:snapB,fset3:snapC
```

The system displays the following command:

```
Invalidate snapshot files in fset3:snapC fset1:snapA fset2:snapB (1..3) of 3...
Deleting files in snapshot fset1:snapA...
100.00 % complete on Thu May 12 04:44:46 2016 (100032 inodes with total 0 MB data processed)
Invalidate snapshot files in fset1:snapA/F/...
```

```
Deleting files in snapshot fset2:snapB...
 100.00 % complete on Thu May 12 04:44:46 2016 (100032 inodes with total 0 MB data processed)
Invalidateing snapshot files in fset2:snapB/F/...
Deleting files in snapshot fset3:snapC...
 100.00 % complete on Thu May 12 04:44:46 2016 (100032 inodes with total 0 MB data processed)
Invalidateing snapshot files in fset3:snapC/F/...
Delete snapshot fset1:snapA successful.
Delete snapshot fset2:snapB successful.
Delete snapshot fset3:snapC successful.
```

## See also

- “[mmclone command](#)” on page 215
- “[mmcrsnapshot command](#)” on page 267
- “[mmlssnapshot command](#)” on page 439
- “[mmrestorefs command](#)” on page 525
- “[mmsnapdir command](#)” on page 565

## Location

/usr/lpp/mmf5/bin

## mmdf

---

### mmdf command

Queries available file space on a GPFS file system.

#### Synopsis

```
| mmdf Device [-d] [-F] [-m] [-P PoolName] [-Y | --block-size {BlockSize | auto}]
| [--qos QoSClass]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

Use the **mmdf** command to display available file space on a GPFS file system. For each disk in the GPFS file system, the **mmdf** command displays this information, by failure group and storage pool:

- The size of the disk.
- The failure group of the disk.
- Whether the disk is used to hold data, metadata, or both.
- Available space in full blocks.
- Available space in fragments.

Displayed values are rounded down to a multiple of 1024 bytes. If the fragment size used by the file system is not a multiple of 1024 bytes, then the displayed values may be lower than the actual values. This can result in the display of a total value that exceeds the sum of the rounded values displayed for individual disks. The individual values are accurate if the fragment size is a multiple of 1024 bytes.

For the file system, the **mmdf** command displays the total number of inodes and the number available.

The **mmdf** command may be run against a mounted or unmounted file system.

#### Notes:

1. This command is I/O intensive and should be run when the system load is light.
2. An asterisk at the end of a line means that this disk is in a state where it is not available for new block allocation.

#### Parameters

##### *Device*

The device name of the file system to be queried for available file space. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

**-d** List only disks that can hold data.

**-F** List the number of inodes and how many of them are free.

**-m** List only disks that can hold metadata.

##### **-P** *PoolName*

Lists only disks that belong to the requested storage pool.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

#### --block-size {BlockSize | auto}

Specifies the unit in which the number of blocks is displayed. The value must be of the form [n]K, [n]M, [n]G or [n]T, where n is an optional integer in the range 1 to 1023. The default is 1K. If **auto** is specified, the number of blocks is automatically scaled to an easy-to-read value.

#### --qos QoSClass

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

##### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

## Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## Security

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

If you are a non-root user, you may specify only file systems that belong to the same cluster as the node on which the **mmdf** command was issued.

## Examples

1. To query all disks in the **fs2** file system that can hold data, issue this command:

```
mmdf fs2 -d
```

The system displays information similar to:

```
disk disk size failure holds holds free KB in free KB in
name in KB in group metadata data full blocks fragments
----- -----
Disks in storage pool: system (Maximum disk size allowed is 833 GB)
----- -----
(pool total) 0 0 (0%) 0 (0%)
----- -----
Disks in storage pool: sp1 (Maximum disk size allowed is 359 GB)
gpfs1002nsd 8897968 1 no yes 8342016 (94%) 928 (0%)
----- -----
(pool total) 8897968 8342016 (94%) 928 (0%)
```

## mmdf

|            |         |               |          |
|------------|---------|---------------|----------|
| (data)     | 8897968 | 8342016 (94%) | 928 (0%) |
| (metadata) | 0       | 0 (0%)        | 0 (0%)   |
| (total)    | 8897968 | 8342016 (94%) | 928 (0%) |

2. To query all disks in the **fs1** file system with the number of blocks automatically scaled to an easy-to-read value, issue this command:

```
mmdf fs1 --block-size auto
```

The system displays information similar to:

| disk name                                                           | disk size | failure group | holds metadata | holds data | free in full blocks | free in fragments |
|---------------------------------------------------------------------|-----------|---------------|----------------|------------|---------------------|-------------------|
| Disks in storage pool: system (Maximum disk size allowed is 437 GB) |           |               |                |            |                     |                   |
| gpfs1001nsd                                                         | 33.4G     | 1             | yes            | no         | 33.19G ( 99%)       | 2.5M ( 0%)        |
| gpfs1002nsd                                                         | 33.4G     | 1             | yes            | no         | 33.2G ( 99%)        | 2.5M ( 0%)        |
| gpfs1003nsd                                                         | 33.4G     | 2             | yes            | no         | 33.2G ( 99%)        | 2.5M ( 0%)        |
| gpfs1004nsd                                                         | 33.4G     | 2             | yes            | no         | 33.27G (100%)       | 2.5M ( 0%)        |
| gpfs1005nsd                                                         | 33.4G     | 2             | yes            | no         | 33.32G (100%)       | 1.562M ( 0%)      |
| (pool total)                                                        | 167G      |               |                |            | 166.2G (100%)       | 11.56M ( 0%)      |
| Disks in storage pool: sp1 (Maximum disk size allowed is 484 GB)    |           |               |                |            |                     |                   |
| gpfs1006nsd                                                         | 33.4G     | 4             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1007nsd                                                         | 33.4G     | 4             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1008nsd                                                         | 33.4G     | 4             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1010nsd                                                         | 33.4G     | 4             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1011nsd                                                         | 33.4G     | 4             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1012nsd                                                         | 33.4G     | 5             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1013nsd                                                         | 33.4G     | 5             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1014nsd                                                         | 33.4G     | 5             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1015nsd                                                         | 33.4G     | 5             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| gpfs1016nsd                                                         | 33.4G     | 5             | no             | yes        | 33.4G (100%)        | 1.562M ( 0%)      |
| (pool total)                                                        | 334G      |               |                |            | 334G (100%)         | 15.62M ( 0%)      |
| (data)                                                              | 334G      |               |                |            | 334G (100%)         | 15.62M ( 0%)      |
| (metadata)                                                          | 167G      |               |                |            | 166.2G (100%)       | 11.56M ( 0%)      |
| (total)                                                             | 501G      |               |                |            | 500.2G (100%)       | 27.19M ( 0%)      |

### Inode Information

|                             |        |
|-----------------------------|--------|
| Number of used inodes:      | 4043   |
| Number of free inodes:      | 497717 |
| Number of allocated inodes: | 501760 |
| Maximum number of inodes:   | 514048 |

3. To query **fs1** for inode information, issue this command:

```
mmdf fs1 -F
```

The system displays information similar to:

|                             |        |
|-----------------------------|--------|
| Number of used inodes:      | 4043   |
| Number of free inodes:      | 497717 |
| Number of allocated inodes: | 501760 |
| Maximum number of inodes:   | 514048 |

## See also

- “mmchfs command” on page 179
- “mmcrfs command” on page 250
- “mmdelfs command” on page 295

- “mmlsfs command” on page 411

**Location**

/usr/lpp/mmfs/bin

## mmdiag

### mmdiag command

Displays diagnostic information about the internal GPFS state on the current node.

#### Synopsis

```
mmdiag [--afm [fileset={all|device[:filesetName]}|gw][-Y]
| [--all] [--version] [--waiters] [--deadlock] [--threads]
| [--lroc] [--memory [-Y]] [--network [-Y]] [--config [-Y]] [--trace]
| [--assert] [--iohist [verbose]] [--tokenmgr] [--commands [-Y]]
| [--dapi [session|event|token|disposition|all]]
| [--rpc [node[=name]|size|message|all|nn{S|s|M|m|H|h|D|d}]]
| [--stats] [--nsd [all]]]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

Use the **mmdiag** command to query various aspects of the GPFS internal state for troubleshooting and tuning purposes. The **mmdiag** command displays information about the state of GPFS on the node where it is executed. The command obtains the required information by querying the GPFS daemon process (**mmfsd**), and thus functions only when the GPFS daemon is running.

#### Results

The **mmdiag** command displays the requested information and returns 0 if successful.

#### Parameters

##### --afm

Displays status and statistics of linked AFM and AFM DR filesets that are assigned to the gateway node. Accepts the following options:

**Note:** If you do not specify any option, status and statistics of all filesets are displayed.

##### fileset=all

Displays status and statistics of all active filesets.

##### fileset=device

Displays status and statistics of all active filesets on a specified device.

##### fileset=device:filename

Displays status and statistics of a specified fileset on the specified device.

##### gw

Displays gateway statistics like queue length and memory.

##### --all

Displays all available information. This option is the same as specifying all of the **mmdiag** parameters.

##### --assert

Displays current dynamic assert status and levels.

##### --commands

Displays all the commands currently running on the local node.

##### --config

Displays configuration parameters and their settings. The list of configuration parameters that are shown here consists of configuration parameters that are known to **mmfsd**. Note that some

configuration parameters (for example, trace settings) are scanned only by the layers of code above **mmfsd**, and those parameters are shown in **mmlsconfig** output but not here.

While **mmlsconfig** displays only a subset of configuration parameters (generally those parameters that have nondefault settings), the list here shows a larger parameter set. All of the documented **mmfsd** configuration parameters are shown, plus some of the undocumented parameters (generally those parameters that are likely to be helpful in tuning and troubleshooting).

Note that parameter values that are shown here are those parameters that are currently in effect (as opposed to the values shown in **mmlsconfig** output, which can show the settings that become effective on the next GPFS restart).

#### --deadlock

Displays the longest waiters that exceed the deadlock detection thresholds.

If a deadlock situation occurs, administrators can use this information from all nodes in a cluster to help decide how to break up the deadlock.

#### --dmapi

Displays various DMAPI information. If no other options are specified, summary information is displayed for sessions, pending events, cached tokens, stripe groups, and events that are waiting for reply. The **--dmapi** parameter accepts the following options:

##### **session**

Displays a list of sessions.

##### **event**

Displays a list of pending events.

##### **token**

Displays a list of cached tokens, stripe groups, and events that waiting for reply.

##### **disposition**

Displays the DAPI disposition for events.

##### **all**

Displays all of the **session**, **event**, **token**, and **disposition** information with more details.

#### | --iohist [verbose]

Displays recent I/O history. The information about I/O requests recently submitted by GPFS code is shown here. It can provide some insight into various aspects of GPFS IO, such as the type of data or metadata being read or written, the distribution of I/O sizes, and I/O completion times for individual I/Os. This information can be useful in performance tuning and troubleshooting.

##### | **verbose**

| Displays additional columns of information **info1**, **info2**, **context**, and **thread**. The contents of the columns are as follows:

##### | **info1**, **info2**

| The contents of columns **info1** and **info2** depend on the buffer type. The buffer type is displayed in the **Buf type** column of the command output:

| *Table 13. Contents of columns **input1** and **input2** depending on the value in column **Buf type***

| <b>Buf type (Buffer type)</b>                      | <b>info1</b>                 | <b>info2</b>                 |
|----------------------------------------------------|------------------------------|------------------------------|
| data                                               | The inode number of the file | The block number of the file |
| metadata                                           | The inode number of the file | (For internal use by IBM)    |
| LLIndBlock                                         | The inode number of the file | (For internal use by IBM)    |
| inode                                              | (For internal use by IBM)    | The inode number of the file |
| Other types, such as diskDesc, sgDesc, and others. | (For internal use by IBM)    | (For internal use by IBM)    |

## **mmdiag**

```
| context
| The I/O context that started this I/O.
| thread The name of the thread that started this I/O.
--lroc
 Displays status and statistics for local read-only cache (LROC) devices.
--memory
 Displays information about mmfsd memory usage. Several distinct memory regions are allocated and used by mmfsd, and it can be important to know the memory usage situation for each one.

Heap memory that is allocated by mmfsd
 This area is managed by the OS and is not associated with a preset limit that is enforced by GPFS.

Memory pools 1 and 2
 Both of these pools refer to a single memory area, also known as the shared segment. It is used to cache various kinds of internal GPFS metadata and for many other internal uses. This memory area is allocated by a special, platform-specific mechanism and is shared between user space and kernel code. The preset limit on the maximum shared segment size, current usage, and some prior usage information are shown here.

Memory pool 3
 This area is also known as the token manager pool. This memory area is used to store the token state on token manager servers. The preset limit on the maximum memory pool size, current usage, and some prior-usage information are shown here.

This information can be useful when you are troubleshooting ENOMEM errors that are returned by GPFS to a user application and memory allocation failures reported in a GPFS log file.

--network
 Displays information about mmfsd network connections and pending Remote Procedure Calls (RPCs). Basic information and statistics about all existing mmfsd network connections to other nodes is displayed, including information about broken connections. If any RPCs are pending (that is, sent but not yet replied to), the information about each one is shown, including the list of RPC destinations and the status of the request for each destination. This information can be helpful in following a multinode chain of dependencies during a deadlock or performance-problem troubleshooting.

--nsd [a11]
 Displays status and queue statistics for NSD queues that contain pending requests.

a11
 Displays status and queue statistics for all NSD queues.

--rpc
 Displays RPC performance statistics. The --rpc parameter accepts the following options:

node[=name]
 Displays all per node statistics (channel wait, send time TCP, send time verbs, receive time TCP, latency TCP, latency verbs, and latency mixed). If name is specified, all per node statistics for just the specified node are displayed.

size
 Displays per size range statistics.

message
 Displays per message type RPC execution time.

a11
 Displays everything.
```

*nn{S|s|M|m|H|h|D|d}*

Displays per node RPC latency statistics for the latest number of intervals, which are specified by *nn*, for the interval specified by one of the following characters:

**S|s**

Displays second intervals only.

**M|m**

Displays first the second intervals since the last-minute boundary followed by minute intervals.

**H|h**

Displays first the second and minute intervals since their last minute and hour boundary followed by hour intervals.

**D|d**

Displays first the second, minute, and hour intervals since their last minute, hour, and day boundary followed by day intervals.

Averages are displayed as a number of milliseconds with three decimal places (one-microsecond granularity).

#### --stats

Displays some general GPFS statistics.

GPFS uses a diverse array of objects to maintain the file system state and cache various types of metadata. The statistics about some of the more important object types are shown here.

#### OpenFile

This object is needed to access an inode. The target maximum number of cached OpenFile objects is governed by the **maxFilesToCache** configuration parameter. Note that more OpenFile objects can be cached, depending on the workload.

#### CompactOpenFile

These objects contain an abbreviated form of an OpenFile, and are collectively known as *stat cache*. The target maximum number of cached CompactOpenFile objects is governed by the **maxStatCache** parameter of the **mmchconfig** command.

#### OpenInstance

This object is created for each open file instance (file or directory opened by a distinct process).

#### BufferDesc

This object is used to manage buffers in the GPFS page pool.

#### IndBlockDesc

This object is used to cache indirect block data.

All of these objects use the shared segment memory. For each object type, a preset target exists, which is derived from configuration parameters and the memory available in the shared segment. The information about current object usage can be helpful in performance tuning.

#### --threads

Displays **mmfsd** thread statistics and the list of active threads. For each thread, its type and kernel thread ID are shown. All non-idle **mmfsd** threads are shown. For those threads that are currently waiting for an event, the wait reason and wait time in seconds are shown. This information provides more detail than the data displayed by **mmdiag --waiters**.

#### --tokenmgr

Displays information about token management. For each mounted GPFS file system, one or more token manager nodes is appointed. The first token manager is always colocated with the file system manager, while other token managers can be appointed from the pool of nodes with the *manager*

## **mmdiag**

designation. The information that is shown here includes the list of currently appointed token manager nodes and, if the current node is serving as a token manager, some statistics about prior token transactions.

### **--trace**

Displays current trace status and trace levels. During GPFS troubleshooting, it is often necessary to use the trace subsystem to obtain the debug data necessary to understand the problem. See *Trace facility* in *IBM Spectrum Scale: Problem Determination Guide*. It is important to have trace levels set correctly, per instructions provided by the IBM Support Center. The information that is shown here makes it possible to check the state of tracing and to see the trace levels currently in effect.

### **--version**

Displays information about the GPFS build currently running on this node. This information helps in troubleshooting installation problems. The information that is displayed here can be more comprehensive than the version information that is available from the OS package management infrastructure, in particular when an e-fix is installed.

### **--waiters**

Displays **mmfsd** threads that are waiting for events. This information can be helpful in troubleshooting deadlocks and performance problems. For each thread, the thread name, wait time in seconds, and wait reason are typically shown. Only non-idle threads that are currently waiting for some event to occur are displayed. Note that only **mmfsd** threads are shown; any application I/O threads that might be waiting in GPFS kernel code would not be present here.

- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

## **Exit status**

0      Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmdiag** command.

## **Examples**

1. To display a list of waiters, enter the following command:

```
mmdiag --waiters
```

The command displays output like the following example:

```
== mmdiag: waiters ==
0x11DA520 waiting 0.001147000 seconds, InodePrefetchWorker:
 for I/O completion
0x2AAAAB02830 waiting 0.002152000 seconds, InodePrefetchWorker:
 for I/O completion
0x2AAAAB103990 waiting 0.000593000 seconds, InodePrefetchWorker:
 for I/O completion
0x11F51E0 waiting 0.000612000 seconds, InodePrefetchWorker:
 for I/O completion
0x11EDE60 waiting 0.005736500 seconds, InodePrefetchWorker:
 on ThMutex 0x100073ABC8 (0xFFFFC2000073ABC8)
 (CacheReplacementListMutex)
```

In this example, all waiters have a short wait duration and represent a typical snapshot of normal GPFS operation.

2. To display information about memory use, enter the following command:

```
mmdiag --memory
```

The command displays output like the following example:

```
mmfsd heap size: 1503232 bytes
```

```
current mmfsd heap bytes in use: 1919624 total 1867672 payload
```

```
Statistics for MemoryPool id 1 ("Shared Segment (EPHEMERAL)")
```

```
 128 bytes in use
 557721725 hard limit on memory usage
 1048576 bytes committed to regions
 1 allocations
 1 frees
 0 allocation failures
```

```
Statistics for MemoryPool id 2 ("Shared Segment")
```

```
 8355904 bytes in use
 557721725 hard limit on memory usage
 8785920 bytes committed to regions
 1297534 allocations
 1296595 frees
 0 allocation failures
```

```
Statistics for MemoryPool id 3 ("Token Manager")
```

```
 496184 bytes in use
 510027355 hard limit on memory usage
 524288 bytes committed to regions
 1309 allocations
 130 frees
 0 allocation failures
```

In this example, a typical memory usage picture is shown. None of the memory pools are close to being full, and no prior allocation failures occurred.

3. To display information about the network, enter the following command:

```
mmdiag --network
```

The command displays information like the following example:

```
== mmdiag: network ==
```

```
Pending messages:
 (none)
```

```
Inter-node communication configuration:
```

|                |                               |
|----------------|-------------------------------|
| tscTcpPort     | 1191                          |
| my address     | 9.114.53.217/25 (eth2) <c0n2> |
| my addr list   | 9.114.53.217/25 (eth2)        |
| my node number | 4                             |

```
TCP Connections between nodes:
```

| Device null: | hostname                      | node   | destination | status | err | sock | sent(MB) | recv(MB) | ostype  |
|--------------|-------------------------------|--------|-------------|--------|-----|------|----------|----------|---------|
|              | c941f1n05.pok.stglabs.ibm.com | <c0n1> | 9.114.78.25 | broken | 233 | -1   | 0        | 0        | Linux/L |

```
Device eth2:
```

|  |                               |        |             |           |     |      |          |          |         |
|--|-------------------------------|--------|-------------|-----------|-----|------|----------|----------|---------|
|  | hostname                      | node   | destination | status    | err | sock | sent(MB) | recv(MB) | ostype  |
|  | c941f3n03.pok.stglabs.ibm.com | <c0n0> | 9.114.78.43 | connected | 0   | 61   | 0        | 0        | Linux/L |

```
c870f4ap06 <c0n3> 9.114.53.218
```

```
Connection details:
```

```
<c0n1> 9.114.78.25/0 (c941f1n05.pok.stglabs.ibm.com)
```

```
connection info:
```

```
 retry(success): 0(0)
```

```
<c0n0> 9.114.78.43/0 (c941f3n03.pok.stglabs.ibm.com)
```

```
connection info:
```

```
 retry(success): 0(0)
```

## mmdiag

```
tcp connection state: established tcp congestion state: open
packet statistics:
 lost: 0 unacknowledged: 0
 retrans: 0 unrecovered retrans: 0
network speed(μs):
 rtt(round trip time): 456 medium deviation of rtt: 127
pending data statistics(byte):
 read/write calls pending: 0
 GPFS Send-Queue: 0 GPFS Recv-Queue: 0
 Socket Send-Queue: 0 Socket Recv-Queue: 0
<c0n3> 9.114.53.218/0 (c870f4ap06)
connection info:
 retry(success): 0(0)
 tcp connection state: established tcp congestion state: open
packet statistics:
 lost: 0 unacknowledged: 0
 retrans: 0 unrecovered retrans: 0
network speed(μs):
 rtt(round trip time): 8813 medium deviation of rtt: 13754
pending data statistics(byte):
 read/write calls pending: 0
 GPFS Send-Queue: 0 GPFS Recv-Queue: 0
 Socket Send-Queue: 0 Socket Recv-Queue: 0
Device details:
 devicename speed mtu duplex rx_dropped rx_errors tx_dropped tx_errors
 eth2 1000 1500 full 0 0 0 0
diag verbs: VERBS RDMA class not initialized
```

4. To display information about status and statistics of all AFM and AFM DR relationships, enter the following command: **mmdiag --afm** The command displays output similar to the following example:

```
== mmdiag: afm ==
AFM Gateway: p7fbn10 Active

AFM-Cache: adrFset-4 (/gpfs/fs1/adrFset-4) in Device: fs1
 Mode: primary
 Home: p7fbn09 (nfs://p7fbn09/gpfs/fs1/adrFset-4)
 Fileset Status: Linked
 Handler-state: Mounted
 Cache-state: PrimInitInProg
 Q-state: Normal Q-length: 12126378 Q-executed: 40570
AFM-Cache: adrFset-5 (/gpfs/fs1/adrFset-5) in Device: fs1
 Mode: primary
 Home: p7fbn09 (nfs://p7fbn09/gpfs/fs1/adrFset-5)
 Fileset Status: Linked
 Handler-state: Mounted
 Cache-state: PrimInitInProg
 Q-state: Normal Q-length: 6164585 Q-executed: 7113648
AFM-Cache: adrFset-10 (/gpfs/fs1/adrFset-10) in Device: fs1
 Mode: primary
 Home: p7fbn09 (nfs://p7fbn09/gpfs/fs1/adrFset-10)
 Fileset Status: Linked
 Handler-state: Mounted
 Cache-state: PrimInitInProg
 Q-state: Normal Q-length: 16239687 Q-executed: 2415474
```

5. To display gateway statistics, enter the following command: **mmdiag --afm gw** The command displays output similar to the following example:

```
== mmdiag: afm ==
AFM Gateway: p7fbn10 Active

QLen: 33165776 QMem: 12560682162 SoftQMem: 12884901888 HardQMem 32212254720
Ping thread: Started
```

## Location

/usr/lpp/mmfs/bin

---

## mmdsh command

Runs commands on multiple nodes or network connected hosts at the same time.

### Synopsis

```
mmdsh -N {Node[,Node...] | NodeFile | NodeClass}
[-l LoginName] [--ignoreSignal {Signal[Signal...]}]
[-i] [-s] [-r RemoteShellPath] [-I {"File[;File]"}]
[-d]] [-v[-R ReportFile]] [-f FanOutValue] Command
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmdsh** command to remotely execute a command concurrently on each of the nodes that are specified in the **-N** option.

**Note:** For security reasons, the **mmdsh** command is limited to the list of allowable remote commands when sudo wrappers are implemented. For more information on how to configure sudo wrappers, see *Configuring sudo in IBM Spectrum Scale: Administration Guide*.

#### CAUTION:

The **mmdsh** command runs any authorized command that you specify concurrently against the list of nodes that you specify. To avoid accidentally damaging or corrupting your clusters or file systems, ensure that you have specified the correct command and the correct list of nodes before you run **mmdsh**.

### Parameters

**-N {Node[,Node...] | NodeFile | NodeClass}**

Runs the command on the nodes in the given node specification. The nodespecification argument can be a comma-separated list of nodes, a node file, or a node class. The nodes in the list or the file can be specified as long or short admin or daemon node names, node numbers, node number ranges, or IP addresses.

**-l *LoginName***

Allows the user to specify a log-in name for the nodes. The log-in names are entered in the command line.

**--ignoreSignal {*Signal[,Signal...]*}**

Rejects the listed signals.

**-i**

Displays the set of nodes before the command is run on those nodes.

**-s**

Suppresses the prepending of the hostname string to each line of output generated by running the command on the remote node.

**-r *RemoteShellPath***

Specifies the full path of the remote shell command that must be used.

**-I {"*File[;File]*"}**

## **mmdsh**

Name of the file to be copied on the remote node prior to executing the command. The file must be in /var/mmfs/tmp. Use this option to specify two files.

**-d**

Removes the files specified with the -I flag before exiting.

**-v**

- | Verifies that a node is reachable with an ICMP echo command (network ping) before adding it to the set of nodes on which the command must be run.

**-R ReportFile**

Reports the list of hosts removed from the working collective when host verification (host ping) fails. The report is written to the specified file with one host per line. The report is generated only when combined with the **-v** parameter.

**-f FanOutValue**

Specifies a fanout value used for concurrent execution. The default value is 64.

### **Command**

To be run on the remote hosts.

If the command is the reserved word `_SELECT_FROM_FILE_`, then the commands to be run on the different hosts are expected to be in the file pointed to by the environment variable `mmdshCommandsFile`. Each line of this file consists of hostname followed by a command string. If the command is the reserved word `_NOOP_`, then `mmdsh` is used to propagate files in parallel via the **-I** option.

**Note:** To determine the working collective, set the `WCOLL` environment variable that contains the host names (one name per line) where the commands must be executed, if it is not already obtained from the command line.

### **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the `mmdsh` command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### **Example**

1. To run the command on a list of nodes specified in NodeFile, run the following command:

```
mmdsh -N ./nodes uname -a
```

2. To run the command on a specified list of hosts, run the following command:

```
mmdsh -N host1,host2,host3 ls
```

3. To run the command on the quorum nodes specified by a node class, run the following command:

```
mmdsh -N quorumnodes date
```

4. To run the command on the hosts listed in the hostfile, run the following command:

```
mmdsh -N ./nodes -r /usr/bin/ssh uname -a
```

5. To run the command on a specified node along with a log-in name, run the following command:  
`mmdsh -N host -l admin ls`

## **mmeditacl**

---

### **mmeditacl command**

Creates or changes a GPFS access control list.

#### **Synopsis**

```
mmeditacl [-d] [-k {nfs4 | posix | native}] Filename
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmeditacl** command for interactive editing of the ACL of a file or directory. This command uses the default editor, specified in the EDITOR environment variable, to display the current access control information, and allows the file owner to change it. The command verifies the change request with the user before making permanent changes.

This command cannot be run from a Windows node.

The EDITOR environment variable must contain a complete path name, for example:

```
export EDITOR=/bin/vi
```

For information about NFS V4 ACLs, see the topics *Managing GPFS access control lists* and *NFS and GPFS* in the *IBM Spectrum Scale: Administration Guide*.

Users may need to see ACLs in their true form as well as how they are translated for access evaluations. There are four cases:

1. By default, **mmeditacl** returns the ACL in a format consistent with the file system setting, specified using the **-k** flag on the **mmcrfs** or **mmchfs** commands.
  - If the setting is **posix**, the ACL is shown as a traditional ACL.
  - If the setting is **nfs4**, the ACL is shown as an NFS V4 ACL.
  - If the setting is **all**, the ACL is returned in its true form.
2. The command **mmeditacl -k nfs4** always produces an NFS V4 ACL.
3. The command **mmeditacl -k posix** always produces a traditional ACL.
4. The command **mmeditacl -k native** always shows the ACL in its true form regardless of the file system setting.

The following describes how **mmeditacl** works for POSIX and NFS V4 ACLs:

| Command             | ACL   | mmcrfs -k | Display       | -d (default)   |
|---------------------|-------|-----------|---------------|----------------|
| mmeditacl           | posix | posix     | Access ACL    | Default ACL    |
| mmeditacl           | posix | nfs4      | NFS V4 ACL    | Error[1]       |
| mmeditacl           | posix | all       | Access ACL    | Default ACL    |
| mmeditacl           | nfs4  | posix     | Access ACL[2] | Default ACL[2] |
| mmeditacl           | nfs4  | nfs4      | NFS V4 ACL    | Error[1]       |
| mmeditacl           | nfs4  | all       | NFS V4 ACL    | Error[1]       |
| mmeditacl -k native | posix | any       | Access ACL    | Default ACL    |
| mmeditacl -k native | nfs4  | any       | NFS V4 ACL    | Error[1]       |
| mmeditacl -k posix  | posix | any       | Access ACL    | Default ACL    |
| mmeditacl -k posix  | nfs4  | any       | Access ACL[2] | Default ACL[2] |
| mmeditacl -k nfs4   | any   | any       | NFS V4 ACL    | Error[1]       |

[1] NFS V4 ACLs include inherited entries. Consequently, there cannot be a separate default ACL.

[2] Only the mode entries (owner, group, everyone) are translated.

The **rwx** values are derived from the NFS V4 file mode attribute. Since the NFS V4 ACL is more granular in nature, some information is lost in this translation.

---

In the case of NFS V4 ACLs, there is no concept of a default ACL. Instead, there is a single ACL and the individual access control entries can be flagged as being inherited (either by files, directories, both, or neither). Consequently, specifying the **-d** flag for an NFS V4 ACL is an error. By its nature, storing an NFS V4 ACL implies changing the inheritable entries (the GPFS default ACL) as well.

Depending on the file system's **-k** setting (**posix**, **nfs4**, or **all**), **mmeditacl** may be restricted. The **mmeditacl** command is not allowed to store an NFS V4 ACL if **-k posix** is in effect, and is not allowed to store a POSIX ACL if **-k nfs4** is in effect. For more information, see the description of the **-k** flag for the **mmchfs**, **mmcrfs**, and **mmlsfs** commands.

## Parameters

### *Filename*

The path name of the file or directory for which the ACL is to be edited. If the **-d** option is specified, *Filename* must contain the name of a directory.

## Options

**-d** Specifies that the default ACL of a directory is to be edited.

**-k {nfs4 | posix | native}**

#### **nfs4**

Always produces an NFS V4 ACL.

#### **posix**

Always produces a traditional ACL.

#### **native**

Always shows the ACL in its true form regardless of the file system setting.

This option should not be used for routine ACL manipulation. It is intended to provide a way to show the translations that are done. For example, if a **posix** ACL is translated by NFS V4. Beware that if the **-k nfs4** flag is used, but the file system does not allow NFS V4 ACLs, you will not be able to store the ACL that is returned. If the file system does support NFS V4 ACLs, the **-k nfs4** flag is an easy way to convert an existing **posix** ACL to **nfs4** format.

## Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## Security

You may issue the **mmeditacl** command only from a node in the GPFS cluster where the file system is mounted.

The **mmeditacl** command may be used to display an ACL. POSIX ACLs may be displayed by any user with access to the file or directory. NFS V4 ACLs have a **READ\_ACL** permission that is required for non-privileged users to be able to see an ACL. To change an existing ACL, the user must either be the owner, the root user, or someone with control permission (**WRITE\_ACL** is required where the existing ACL is of type NFS V4).

## **mmeditacl**

### **Examples**

To edit the ACL for a file named **project2.history**, issue this command:

```
mmeditacl project2.history
```

The current ACL entries are displayed using the default editor, provided that the EDITOR environment variable specifies a complete path name. When the file is saved, the system displays information similar to:

```
mmeditacl: 6027-967 Should the modified ACL be applied? (yes) or (no)
```

After responding **yes**, the ACLs are applied.

### **See also**

- “[mmdelacl command](#)” on page 284
- “[mmgetacl command](#)” on page 347
- “[mmputacl command](#)” on page 504

### **Location**

`/usr/lpp/mmfs/bin`

---

## mmedquota command

Sets quota limits.

### Synopsis

```
mmedquota {-u [-p [ProtoFileset:]ProtoUser] [Device:Fileset:]User ... |
-g [-p [ProtoFileset:]ProtoGroup] [Device:Fileset:]Group ... |
-j [-p ProtoFileset] Device:Fileset ... |
-d {-u User ... | -g Group ... | -j Device:Fileset ...} |
-t {{-u | -g | -j} [--reset]}}
```

### Availability

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

### Description

The **mmedquota** command serves two purposes:

1. Sets or changes quota limits or grace periods for users, groups, and filesets in the cluster from which the command is issued.
2. Reestablishes user, group, or fileset default quotas for all file systems with default quotas enabled in the cluster.

The **mmedquota** command displays the current values for these limits, if any, and prompts you to enter new values using your default editor:

- current block usage (the amount of disk space used by this user, group, or fileset, in 1KB units; display only)
- current inode usage (display only)
- node soft limit
- inode hard limit
- block soft limit (the amount of disk space that this user, group, or fileset is allowed to use during normal operation)

Displayed in **KB**, but may be specified using **g**, **G**, **k**, **K**, **m**, **M**, **p**, **P**, **t**, or **T**. If no suffix is provided, the number is assumed to be in **bytes**.

- block hard limit (the total amount of disk space that this user, group, or fileset is allowed to use during the grace period)

Displayed in **KB**, but may be specified using **g**, **G**, **k**, **K**, **m**, **M**, **p**, **P**, **t**, or **T**. If no suffix is provided, the number is assumed to be in **bytes**.

**Note:** A block or inode limit of 0 indicates no limit.

The **mmedquota** command waits for the edit window to be closed before checking and applying new values. If an incorrect entry is made, reissue the command and enter the correct values.

You can also use the **mmedquota** command to change the file system-specific grace periods for block and file usage if the default of one week is unsatisfactory. The grace period is the time during which users can exceed the soft limit. If the user, group, or fileset does not show reduced usage below the soft limit before the grace period expires, the soft limit becomes the new hard limit.

When you set quota limits for a file system, consider replication in the file system. See the topic *Listing quotas* in the *IBM Spectrum Scale: Administration Guide*.

The **EDITOR** environment variable must contain a complete path name, for example:  
`export EDITOR=/bin/vi`

## **mmedquota**

### **Parameters**

#### *Device*

Specifies the device name of the file system for which quota information is to be displayed. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

#### *Fileset*

Specifies the name of a fileset located on *Device* for which quota information is to be displayed.

#### *User*

Name or user ID of target user for quota editing.

#### *Group*

Name or group ID of target group for quota editing.

**-d** Reestablish default quota limits for a specific user, group, or fileset that has had an explicit quota limit set by a previous invocation of the **mmedquota** command.

**-g** Sets quota limits or grace times for groups.

**-j** Sets quota limits or grace times for filesets.

**-p** Applies already-established limits to a particular user, group or fileset.

When invoked with the **-u** option, [ProtoFileset:]ProtoUser limits are automatically applied to the specified *User* or space-delimited list of users.

When invoked with the **-g** option, [ProtoFileset:]ProtoGroup limits are automatically applied to the specified *Group* or space-delimited list of groups.

When invoked with the **-j** option, ProtoFileset limits are automatically applied to the specified fileset or space-delimited list of fileset names.

You can specify any user as a *ProtoUser* for another *User*, or any group as a *ProtoGroup* for another *Group*, or any fileset as a *ProtoFileset* for another *Fileset*.

**-p** cannot propagate a prototype quota from a user, group, or fileset on one file system to a user, group, or fileset on another file system.

**-t** Sets grace period during which quotas can exceed the soft limit before it is imposed as a hard limit. The default grace period is one week.

This flag is followed by one of the following flags: **-u**, **-g** or **-j**, to specify whether the changes apply to users, groups, or filesets respectively.

**-u** Sets quota limits or grace times for users.

#### **--reset**

With this option, when grace time is modified, all relative quota entries will be scanned and updated if necessary; without this option, when grace time is updated, quota entries will not be scanned and updated.

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the **mmedquota** command.

GPFS must be running on the node from which the **mmedquota** command is issued.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To set user quotas for userid **pfs001**, issue this command:

```
mmedquota -u pfs001
```

The system displays information similar to:

```
*** Edit quota limits for USR pfs001
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs3: (root): blocks in use: 8K, limits (soft = 0K, hard = 0K)
 inodes in use: 1, limits (soft = 0, hard = 0)
gpfs2: (fset4): blocks in use: 4104K, limits (soft = 102400K, hard = 153600K)
 inodes in use: 2, limits (soft = 100, hard = 150)
gpfs2: (fset3): blocks in use: 0K, limits (soft = 0K, hard = 0K)
 inodes in use: 0, limits (soft = 0, hard = 0)
gpfs2: (root): blocks in use: 0K, limits (soft = 0K, hard = 0K)
 inodes in use: 0, limits (soft = 0, hard = 0)
gpfs1: (fset1): blocks in use: 0K, limits (soft = 256K, hard = 256K)
 inodes in use: 0, limits (soft = 30, hard = 40)
gpfs1: (root): blocks in use: 0K, limits (soft = 256K, hard = 256K)
 inodes in use: 0, limits (soft = 40, hard = 45)
```

2. To reset default group quota values for the group **blueteam**, issue this command:

```
mmedquota -d -g blueteam
```

To verify the change, issue this command:

```
mmrepquota -q fs1
```

The system displays information similar to:

```
fs1: USR quota is on; default quota is on
fs1: GRP quota is on; default quota is on
fs1: FILESET quota is on; default quota is off
```

3. To change the grace periods for all users, issue this command:

```
mmedquota -t -u
```

The system displays information in your default editor similar to:

```
*** Edit grace times:
Time units may be : days, hours, minutes, or seconds
Grace period before enforcing soft limits for USRs:
gpfs0: block grace period: 7 days, file grace period: 7 days
```

4. To set user quotas for device **gpfs2**, fileset **fset3**, and userid **pfs001**, issue this command:

```
mmedquota -u gpfs2:fset3:pfs001
```

The system displays information similar to:

```
*** Edit quota limits for USR gpfs2:fset3:pfs001
NOTE: block limits will be rounded up to the next multiple of the block size.
 block units may be: K, M, G, T or P, inode units may be: K, M or G.
gpfs2: (fset3): blocks in use: 0K, limits (soft = 0K, hard = 0K)
 inodes in use: 0, limits (soft = 0, hard = 0)
```

5. To apply already-established limits of user **pfs002** to user **pfs001**, issue this command:

```
mmedquota -u -p pfs002 pfs001
```

The system displays information similar to:

## **mmedquota**

Already established limits of protouser pfs002 in root fileset are applied to all filesets (including root fileset) in all corresponding per-fileset quota enabled filesystems.

6. To apply already-established limits of user **pfs002** in fileset **fset2** to user **pfs001** in fileset **fset1** and file system **fs1**, issue this command:

```
mmedquota -u -p fset2:pfs002 fs1:fset1:pfs001
```

The system displays information similar to:

```
Limits of protouser pfs002 in fileset fset2 in filesystem fs1 are applied
to user pfs001 in fileset fset1
```

7. To apply an already-established fileset quota (from **fileset1** to **fileset2**) in two different file systems (**gpfstest1** and **gpfstest2**), issue this command:

```
mmedquota,-j -p fileset1 gpfstest1:fileset2 gpfstest2:fileset2
```

### **See also**

- “mmcheckquota command” on page 169
- “mmdefedquota command” on page 272
- “mmdefquotaoff command” on page 275
- “mmdefquotaoon command” on page 278
- “mmlsquota command” on page 435
- “mmquotaon command” on page 509
- “mmquotaoff command” on page 507
- “mmrepquota command” on page 517

### **Location**

/usr/lpp/mmfs/bin

---

## mmexportfs command

Retrieves the information needed to move a file system to a different cluster.

### Synopsis

```
mmexportfs {Device | all} -o ExportfsFile
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmexportfs** command, in conjunction with the **mmimportfs** command, can be used to move one or more GPFS file systems from one GPFS cluster to another GPFS cluster, or to temporarily remove file systems from the cluster and restore them at a later time. The **mmexportfs** command retrieves all relevant file system and disk information and stores it in the file specified with the **-o** parameter. This file must later be provided as input to the **mmimportfs** command. When running the **mmexportfs** command, the file system must be unmounted on all nodes.

When **all** is specified in place of a file system name, any disks that are not associated with a file system will be exported as well.

Exported file systems remain unusable until they are imported back with the **mmimportfs** command to the same or a different GPFS cluster.

### Results

Upon successful completion of the **mmexportfs** command, all configuration information pertaining to the exported file system and its disks is removed from the configuration data of the current GPFS cluster and is stored in the user specified file *ExportfsFile*.

### Parameters

**Device | all**

The device name of the file system to be exported. File system names need not be fully-qualified. *fs0* is as acceptable as */dev/fs0*. Specify **all** to export all GPFS file systems, as well as all disks that do not currently belong to a file system.

If the specified file system device is a IBM Spectrum Scale RAID-based file system, then all affected IBM Spectrum Scale RAID objects will be exported as well. This includes recovery groups, declustered arrays, vdks, and any other file systems that are based on these objects. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

This must be the first parameter.

**-o ExportfsFile**

The path name of a file to which the file system information is to be written. This file must be provided as input to the subsequent **mmimportfs** command.

### Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

## **mmexportfs**

### **Security**

You must have root authority to run the **mmexportfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

To export all file systems in the current cluster, issue this command:

```
mmexportfs all -o /u/admin/exportfile
```

The output is similar to this:

```
mmexportfs: Processing file system fs1 ...
```

```
mmexportfs: Processing file system fs2 ...
```

```
mmexportfs: Processing disks that do not belong to any file system ...
```

```
mmexportfs: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

### **See also**

- “[mmimportfs command](#)” on page 374

### **Location**

/usr/lpp/mmfs/bin

## mmfsck command

Checks and repairs a GPFS file system.

### Synopsis

```
mmfsck Device [-n | -y] [-s | -v | -V]
 [-c | -m | -o | --skip-inode-check | --skip-directory-check]
 [-t Directory]
 [-N {Node[,Node...] | NodeFile | NodeClass}]
 [--patch-file Path [--patch]] [--qos QosClass]
 [--threads ThreadLevel]
```

The file system must be unmounted before you can run the **mmfsck** command with any option other than **-o**.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmfsck** command in offline mode is intended to be used only in situations where there have been disk or communications failures that have caused **MMFS\_FSSTRUCT** error log entries to be issued, or where it is known that disks have been forcibly removed or otherwise permanently unavailable for use in the file system, and other unexpected symptoms are seen by users. In general it is unnecessary to run **mmfsck** in offline mode unless under the direction of the IBM Support Center.

**Note:** When an **fsstruct** error is show in **mmhealth**, you are asked to run a filesystem check. Once the problem is solved, you need to clear the **fsstruct** error from **mmhealth** manually by running the following command:

```
mmsysmonc event filesystem fsstruct_fixed <filesystem_name>
```

If neither the **-n** nor **-y** flag is specified, the **mmfsck** command runs interactively prompting you for permission to repair each consistency error as reported. It is suggested that in all but the most severely damaged file systems, you run the **mmfsck** command interactively (the default).

The occurrence of I/O errors, or the appearance of a message telling you to run the **mmfsck** command, may indicate file system inconsistencies. If either situation occurs, use the **mmfsck** command to check file system consistency and interactively repair the file system.

For information about file system maintenance and repair, see the topic *Checking and repairing a file system* in the *IBM Spectrum Scale: Administration Guide*. The **mmfsck** command checks for these inconsistencies:

- Blocks marked allocated that do not belong to any file. The corrective action is to mark the block free in the allocation map.
- Files for which an inode is allocated and no directory entry exists (orphaned files). The corrective action is to create directory entries for these files in a **lost+found** subdirectory of the fileset to which the orphaned file or directory belongs. The index number of the inode is assigned as the name. If you do not allow the **mmfsck** command to reattach an orphaned file, it asks for permission to delete the file.
- Directory entries pointing to an inode that is not allocated. The corrective action is to remove the directory entry.
- Incorrectly formed directory entries. A directory file contains the inode number and the generation number of the file to which it refers. When the generation number in the directory does not match the generation number stored in the file's inode, the corrective action is to remove the directory entry.

## **mmfsck**

- Incorrect link counts on files and directories. The corrective action is to update them with accurate counts.
- Policy files are not valid. The corrective action is to delete the file.
- Various problems related to filesets: missing or corrupted fileset metadata, inconsistencies in directory structure related to filesets, missing or corrupted fileset root directory, other problems in internal data structures. The repaired filesets will be renamed as **Fileset FilesetId** and put into unlinked state.

If you are repairing a file system due to node failure and the file system has quotas enabled, it is suggested that you run the **mmcheckquota** command to recreate the quota files.

Indications leading you to the conclusion that you should run the **mmfsck** command include:

- An **MMFS\_FSSTRUCT** along with an **MMFS\_SYSTEM\_UNMOUNT** error log entry on any node indicating some critical piece of the file system is inconsistent.
- Disk media failures
- Partial disk failure
- **EVALIDATE=214**, Invalid checksum or other consistency check failure on a disk data structure, reported in error logs or returned to an application.

For further information on recovery actions and how to contact the IBM Support Center, see the *IBM Spectrum Scale: Problem Determination Guide*.

If you are running the online **mmfsck** command to free allocated blocks that do not belong to any files, plan to make file system repairs when system demand is low. This is an I/O intensive activity and it can affect system performance.

## **Results**

If the file system is inconsistent, the **mmfsck** command displays information about the inconsistencies and (depending on the option entered) may prompt you for permission to repair them. The **mmfsck** command tries to avoid actions that may result in loss of data. In some cases, however, it may indicate the destruction of a damaged file.

All corrective actions, with the exception of recovering lost disk blocks (blocks that are marked as allocated but do not belong to any file), require that the file system be unmounted on all nodes. If the **mmfsck** command is run on a mounted file system, lost blocks are recovered but any other inconsistencies are only reported, not repaired.

If a bad disk is detected, the **mmfsck** command stops the disk and writes an entry to the error log. The operator must manually start and resume the disk when the problem is fixed.

The file system must be unmounted on all nodes before the **mmfsck** command can repair file system inconsistencies.

- | When the command is running in verbose or semi-verbose mode, the command provides a summary of the errors including the severity of each error:
  - | **CRITICAL**
    - | The command found a critical corruption in the file system. Using the file system without repairing the corruption might cause more corruptions or file system panics.
  - | **NONCRITICAL**
    - | The command found a non-critical corruption in the file system. You can use the file system without repairing the corruption, but some metadata might not be accessible.

## | HARMLESS

| The command found a harmless problem in the file system that indicates that some unused  
| metadata can be freed to reclaim space.

## Parameters

### *Device*

The device name of the file system to be checked and repaired. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

- n Specifies a **no** response to all file system error repair prompts from the **mmfsck** command. The option reports inconsistencies but it does not change the file system. To save this information, redirect it to an output file when you issue the **mmfsck** command.
- | **Note:** If the **mmfsck** command is run offline with the **-n** parameter and it detects errors, it panics the file system to force the file system to go through cleanup before any new command can be started.
- y Specifies a **yes** response to all file system error repair prompts from the **mmfsck** command. Use this option only on severely damaged file systems. It allows the **mmfsck** command to take any action necessary for repairs.
- s Specifies that the output is semi-verbose.
- v Specifies that the output is verbose.
- V Specifies that the output is verbose and contains information for debugging purposes.
- c When the file system log has been lost and the file system is replicated, this option specifies that the **mmfsck** command attempt corrective action by comparing the replicas of metadata and data. If this error condition occurs, it is indicated by an error log entry.
- m Has the same meaning as **-c**, except that **mmfsck** checks only the metadata replica blocks. It therefore runs faster than with **-c**.
- o Online mode does not perform a full file system consistency check, but blocks marked as allocated that do not belong to a file are recovered. Lost blocks do not constitute file system corruption.

### --skip-inode-check

Causes the command to run faster by skipping its inode-check phase. Include this option only if you know that the inodes are valid and that only directories need to be checked. In this mode, the product does not scan all parts of the file system and therefore might not detect all corruptions in the file system.

### --skip-directory-check

Causes the command to run faster by skipping its directory-check phase. Include this option if you want to check only the inodes. In this mode, the product does not scan all parts of the file system and therefore might not detect all corruptions in the file system.

### -t *Directory*

Specifies the directory that GPFS uses for temporary storage during **mmfsck** command processing. This directory must be available on all nodes that are participating in **mmfsck** and that are designated as either manager or quorum node. In addition to the location requirement, the storage directory has a minimum space requirement of 4GB. The default directory for **mmfsck** processing is **/tmp**.

### -N {*Node*[,*Node*...]} | *NodeFile* | *NodeClass*

Specify the nodes to participate in the check and repair of the file system. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

## **mmfsck**

For information on how to specify node names, see the topic *Specifying nodes as inputs to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### **--patch-file Path**

Specifies the name of a patch file. When the **--patch** parameter is not specified, information about file system inconsistencies (detected during an **mmfsck** run with the **-n** parameter) are stored in the patch file that is specified by *Path*. *Path* should be accessible from the file system manager node. The information stored in the patch file can be viewed as a report of the problems in the file system. For more information about patch files, see the topic *Checking and repairing a file system* in the *IBM Spectrum Scale: Administration Guide*.

When this option is specified with the **--patch** parameter, the information in the patch file is read and used to repair the file system.

### **--patch**

Specifies that the file system will be repaired using the information stored in the patch file that is specified with **--patch-file Path**.

### **--qos QoSClass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

#### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **--threads ThreadLevel**

The number of threads that are created to run **mmfsck**. The default is 16.

## **Exit status**

- 0** Successful completion.
- 2** The command was interrupted before it completed checks or repairs.
- 4** The command changed the file system and it must now be restarted.
- 8** The file system contains damage that has not been repaired.
- 16** The problem cannot be fixed.
- 64** Do a full offline file system check to verify the integrity of the file system.

The exit string is a combination of three different error indicators:

1. The first value is the Exit **errno** value.
2. The second value is an internal ancillary value that helps explain where the **errno** value came from.
3. The third value is the OR of several status bits.

## **Security**

You must have root authority to run the **mmfsck** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

## Examples

1. The following command checks file system fs2 and displays inconsistencies but does not try to make repairs:

```
mmfsck fs2 -vn
```

The command displays output similar to the following example:

```
Checking "fs2"
FsckFlags 0x2000009
Stripe group manager <c0n0>
NeedNewLogs 0

...
Error in inode 38662 snap 0: has empty EA overflow block
Delete EA overflow block? No

...
Lost blocks were found.
Correct the allocation map? No
Block 1:332288 has map status 0x9FFFFF01 should be 0xFFFFFFFF
Block 3:332288 has map status 0x9FFFFF01 should be 0xFFFFFFFF
Checking inode map file

Corrections are needed in the inode allocation map.
Correct the allocation map? No
Inode 131075 is not in use but marked as used (0) in map.
1 inodes are not in use but marked.
Root inode 131075 of fileset 'fset1' has been deleted.
Delete the inode reference from fileset metadata? No
Checking ACL file records
Starting Directory phase
Checking directories and files

...
Fileset 'fset1' with id 1 has corrupt record (reason 21)
Correct the fileset record? No
Fileset with id 2 is referenced in the file system
but is not recorded in the fileset metadata file.
Correct the metadata file? No

...
Error in inode 33536 snap 0: Directory block 0 has entry whose fileId does not match
fileId of directory
file entry inode 33537 "a"
Delete entry? No

Error in inode 33536 snap 0: Directory block 0 has entry whose fileId does not match
fileId of directory
file entry inode 33538 "b"
Delete entry? No

Error in inode 33536 snap 0: Directory block 0 has entry whose fileId does not match
fileId of directory
file entry inode 33539 "c"
Delete entry? No
```

## mmfsck

```
...
Error in inode 3 snap 0: Directory block 0 has entry referring to a deleted inode
subdir entry inode 131075 "fset1"
Delete entry? No

Error in inode 3 snap 0: Directory block 0 has entry with invalid fileId
subdir entry inode 33536 "dir1"
Delete entry? No

Error in inode 3 snap 0: has nlink field as 5
Correct to 3? No

...
Checking log files
Checking extended attributes file
Checking allocation summary file
Checking policy file
Validated policy /*no policy rules*/ :
Checking metadata of filesets
Inode space 1 does not have owning fileset.
Make fileset 'fset1' (inode -1) owner of the inode space? No
Root directory of fileset 'fset1' (inode -1) is invalid
Repair fileset root directory? No
Checking file reference counts

Error in inode 33536 snap 0: has corrupt fileId 2
Delete file? No

Error in inode 33537 snap 0: is unreferenced
Attach inode to lost+found of fileset root fileId 0? No

Error in inode 33538 snap 0: is unreferenced
Attach inode to lost+found of fileset root fileId 0? No

Error in inode 33539 snap 0: is unreferenced
Attach inode to lost+found of fileset root fileId 0? No

Error in inode 136448 snap 0: has corrupt fileId 1
Delete file? No

Error in inode 151296 snap 0: has corrupt fileId 1
Delete file? No
Checking file system replication status

No. SnapId InodeNum FileType Fix Error(s) Severity

1 0 38662 User N InodeMetadata Harmless
2 0 33536 Directory N Fatal + DirectoryEntry Noncritical
3 0 3 Root Directory N InodeMetadata + DirectoryEntry Noncritical
4 0 38 Fileset Metadata N ReservedFileRecord Critical
5 0 33537 User N OrphanInode Noncritical
6 0 33538 User N OrphanInode Noncritical
7 0 33539 User N OrphanInode Noncritical
8 0 136448 Directory N Fatal Noncritical
9 0 151296 Directory N Fatal Noncritical

197120 inodes
 59 allocated
 3 repairable
 0 repaired
 3 damaged
 0 deallocated
 3 orphaned
```

```

 0 attached
 0 corrupt ACL references

 4194304 subblocks
 217682 allocated
 14 unreferenced
 0 duplicates
 0 deletable
 0 deallocated

 1598 addresses
 0 suspended
 0 duplicates
 0 reserved file holes found
 0 reserved file holes repaired

```

Critical corruptions were found. Using file system without repairing the corruptions may cause more corruptions and/or file system panics.

File system contains unrepairs damage.

Exit status 0:0:8.

mmfsck: 6027-1639 Command failed. Examine previous error messages to determine cause.

2. The following command checks file system fs2, displays inconsistencies, and makes repairs:

```
mmfsck fs2 -vy
```

The command displays output similar to the following example:

```

Checking "fs2"
FsckFlags 0x400000A
Stripe group manager <c0n0>
NeedNewLogs 0

```

...

```

Error in inode 38662 snap 0: has empty EA overflow block
Delete EA overflow block? Yes

```

...

Lost blocks were found.

Correct the allocation map? Yes

Block 1:332288 has map status 0x9FFFFFF01 should be 0xFFFFFFFF

Block 3:332288 has map status 0x9FFFFFF01 should be 0xFFFFFFFF

Checking inode map file

Corrections are needed in the inode allocation map.

Correct the allocation map? Yes

Inode 131075 is not in use but marked as used (0) in map.

1 inodes are not in use but marked.

Root inode 131075 of filesset 'fset1' has been deleted.

Delete the inode reference from filesset metadata? Yes

Checking ACL file records

Starting Directory phase

Checking directories and files

...

Fileset 'fset1' with id 1 has corrupt record (reason 21)

Correct the filesset record? Yes

Fileset with id 2 is referenced in the file system

but is not recorded in the filesset metadata file.

Correct the metadata file? Yes

...

```

Error in inode 33536 snap 0: Directory block 0 has entry whose filessetId does not match
filessetId of directory
file entry inode 33537 "a"

```

## mmfsck

```
| Delete entry? Yes
| Error in inode 33536 snap 0: Directory block 0 has entry whose fileId does not match
| fileId of directory
| file entry inode 33538 "b"
| Delete entry? Yes
|
| Error in inode 33536 snap 0: Directory block 0 has entry whose fileId does not match
| fileId of directory
| file entry inode 33539 "c"
| Delete entry? Yes
|
| ...
|
| Error in inode 3 snap 0: Directory block 0 has entry referring to a deleted inode
| subdir entry inode 131075 "fset1"
| Delete entry? Yes
|
| Error in inode 3 snap 0: Directory block 0 has entry whose fileId does not match
| fileId of directory
| subdir entry inode 33536 "dir1"
| Delete entry? Yes
|
| Error in inode 3 snap 0: has nlink field as 5
| Correct to 3? Yes
|
| ...
|
| Checking log files
| Checking extended attributes file
| Checking allocation summary file
| Checking policy file
| Validated policy '/no policy rules/*' :
| Checking metadata of filesets
| Inode space 1 does not have owning fileset.
| Make fileset 'fset1' (inode -1) owner of the inode space? Yes
| Root directory of fileset 'Fileset2' (inode -1) is invalid
| Repair fileset root directory? Yes
| The number of filesets recorded in the fileset metadata file header (2)
| differs from the actual number (3).
| Correct the recorded number? Yes
| Checking file reference counts
|
| Error in inode 33536 snap 0: is unreferenced
| Attach inode to lost+found of fileset Fileset2 fileId 2? Yes
|
| Error in inode 33537 snap 0: is unreferenced
| Attach inode to lost+found of fileset root fileId 0? Yes
|
| Error in inode 33538 snap 0: is unreferenced
| Attach inode to lost+found of fileset root fileId 0? Yes
|
| Error in inode 33539 snap 0: is unreferenced
| Attach inode to lost+found of fileset root fileId 0? Yes
|
| Error in inode 136448 snap 0: is unreferenced
| Attach inode to lost+found of fileset fset1 fileId 1? Yes
|
| Error in inode 151296 snap 0: is unreferenced
| Attach inode to lost+found of fileset fset1 fileId 1? Yes
| Checking file system replication status
|
| No. SnapId InodeNum FileType Fix Error(s) Severity
| -----+-----+-----+-----+-----+-----+-----+
| 1 0 38662 User Y InodeMetadata
| 2 0 33536 Directory Y DirectoryEntry
| + OrphanInode
|
```

|   |   |        |                  |   |                                |             |
|---|---|--------|------------------|---|--------------------------------|-------------|
| 3 | 0 | 3      | Root Directory   | Y | InodeMetadata + DirectoryEntry | Noncritical |
| 4 | 0 | 38     | Fileset Metadata | Y | ReservedFileRecord             | Critical    |
| 5 | 0 | 33537  | User             | Y | OrphanInode                    | Noncritical |
| 6 | 0 | 33538  | User             | Y | OrphanInode                    | Noncritical |
| 7 | 0 | 33539  | User             | Y | OrphanInode                    | Noncritical |
| 8 | 0 | 136448 | Directory        | Y | OrphanInode                    | Noncritical |
| 9 | 0 | 151296 | Directory        | Y | OrphanInode                    | Noncritical |

```
197120 inodes
 61 allocated
 4 repairable
 4 repaired
 0 damaged
 0 deallocated
 6 orphaned
 6 attached
 0 corrupt ACL references
```

```
4194304 subblocks
 217682 allocated
 14 unreferenced
 0 duplicates
 0 deletable
 14 deallocated
```

```
 1598 addresses
 0 suspended
 0 duplicates
 0 reserved file holes found
 0 reserved file holes repaired
```

Critical corruptions were found in the file system which were repaired.

File system is clean.

3. The following command checks file system FSchk, displays inconsistencies, and stores a record of each inconsistency in the patch file path-towrite-patchfile:

```
mmfsck FSchk -nv --patch-file path-towrite-patchfile
```

The command displays output similar to the following example:

```
Creating patch file "path-towrite-patchfile" on node "node3".
Checking "FSchk"
FsckFlags 0xA000009
Stripe group manager <c0n0>
NeedNewLogs 0
```

...

```
Error in inode 50688 snap 0: Indirect block 0 level 1 has bad disk addr at offset 21
replica 1 addr 2:318976 is a duplicate address
Delete disk address? No
Cannot fix lost blocks if not deleting duplicate address.
```

```
Error in inode 50688 snap 0: Indirect block 0 level 1 has bad disk addr at offset 20
replica 0 addr 2:318976 is a duplicate address
Delete disk address? No
```

```
Error in inode 50688 snap 0: has illReplicated field as 0
Correct to 1? No
```

```
Error in inode 50688 snap 0: has nFullBlocks field as 800
Correct to 798? No
Node 192.168.56.163 (node3) ending inode scan 0 to 32895
```

```
Error in inode 50689 snap 0: Indirect block 0 level 1 has bad disk addr at offset 40
replica 0 addr 1:371200 is a duplicate address
Delete disk address? No
```

## mmfsck

```
Error in inode 50689 snap 0: has illReplicated field as 0
Correct to 1? No

Error in inode 50689 snap 0: has nFullBlocks field as 800
Correct to 799? No

Error in inode 50690 snap 0: Indirect block 0 level 1 has bad disk addr at offset 44
replica 0 addr 2:318976 is a duplicate address
Delete disk address? No

Error in inode 50690 snap 0: Indirect block 0 level 1 has bad disk addr at offset 792
replica 0 addr 1:371200 is a duplicate address
Delete disk address? No

Error in inode 50690 snap 0: has illReplicated field as 0
Correct to 1? No

Error in inode 50690 snap 0: has nFullBlocks field as 800
Correct to 798? No

...
Checking inode map file
Checking ACL file records
Starting Directory phase
Checking directories and files

...
Error in inode 50691 snap 0: Directory block 0 has entry with incorrect generation number
subdir entry inode 50692 "subdir7"
Delete entry? No

Error in inode 50691 snap 0: Directory block 0 has entry with incorrect generation number
subdir entry inode 13834 "subdir10"
Delete entry? No

Error in inode 50691 snap 0: has nlink field as 53
Correct to 51? No

...
Error in inode 13831 snap 0: is unreferenced
Attach inode to lost+found of fileset root fileId 0? No

Error in inode 13834 snap 0: is unreferenced
Attach inode to lost+found of fileset root fileId 0? No
Checking file system replication status
```

| No. | SnapId | InodeNum | FileType  | Fix Error(s)                           | Severity    |
|-----|--------|----------|-----------|----------------------------------------|-------------|
| 1   | 0      | 50688    | User      | N InodeMetadata<br>+ DuplicateDiskAddr | Critical    |
| 2   | 0      | 50689    | User      | N InodeMetadata<br>+ DuplicateDiskAddr | Critical    |
| 3   | 0      | 50690    | User      | N InodeMetadata<br>+ DuplicateDiskAddr | Critical    |
| 4   | 0      | 50691    | Directory | N InodeMetadata<br>+ DirectoryEntry    | Noncritical |
| 5   | 0      | 13831    | Directory | N OrphanInode                          | Noncritical |
| 6   | 0      | 13834    | Directory | N OrphanInode                          | Noncritical |

```
65792 inodes
 98 allocated
 4 repairable
 0 repaired
```

```

 0 damaged
 0 deallocated
 2 orphaned
 0 attached
 0 corrupt ACL references

 4194304 subblocks
 286532 allocated
 160 unreferenced
 64 duplicates
 0 deletable
 0 deallocated

 1598 addresses
 0 suspended
 5 duplicates
 0 reserved file holes found
 0 reserved file holes repaired

```

Critical corruptions were found. Using file system without repairing the corruptions may cause more corruptions and/or file system panics.

File system contains unrepaired damage.

Exit status 0:0:8.

Patch file written to "node3:path-towrite-patchfile" with 21 patch entries.

mmfsck: 6027-1639 Command failed. Examine previous error messages to determine cause.

4. The following command uses the information in patch file path-towrite-patchfile to repair the file system:

```
mmfsck FSchk -v --patch-file path-towrite-patchfile --patch
```

The command displays output similar to the following example:

```

Checking "FSchk"
FsckFlags 0x1C00000A
Stripe group manager <c0n0>
NeedNewLogs 0

...

```

```

Error in inode 50688 snap 0: Indirect block 0 level 1 has bad disk addr at offset 21
replica 1 addr 2:318976 is a duplicate address
Delete disk address? Yes

```

```

Error in inode 50688 snap 0: Indirect block 0 level 1 has bad disk addr at offset 20
replica 0 addr 2:318976 is a duplicate address
Delete disk address? Yes

```

```

Error in inode 50688 snap 0: has illReplicated field as 0
Correct to 1? Yes

```

```

Error in inode 50688 snap 0: has nFullBlocks field as 800
Correct to 798? Yes

```

```

Error in inode 50689 snap 0: Indirect block 0 level 1 has bad disk addr at offset 40
replica 0 addr 1:371200 is a duplicate address
Delete disk address? Yes

```

```

Error in inode 50689 snap 0: has illReplicated field as 0
Correct to 1? Yes

```

```

Error in inode 50689 snap 0: has nFullBlocks field as 800
Correct to 799? Yes

```

```

Error in inode 50690 snap 0: Indirect block 0 level 1 has bad disk addr at offset 44
replica 0 addr 2:318976 is a duplicate address
Delete disk address? Yes

```

```
Error in inode 50690 snap 0: Indirect block 0 level 1 has bad disk addr at offset 792
```

## **mmfsck**

```
| replica 0 addr 1:371200 is a duplicate address
| Delete disk address? Yes
|
| Error in inode 50690 snap 0: has illReplicated field as 0
| Correct to 1? Yes
|
| Error in inode 50690 snap 0: has nFullBlocks field as 800
| Correct to 798? Yes
|
| Lost blocks were found.
| Correct the allocation map? Yes
| Block 1:134144 has map status 0x00000000 should be 0xFFFFFFFF
| Block 3:134656 has map status 0x00000000 should be 0xFFFFFFFF
| Block 2:239104 has map status 0x00000000 should be 0xFFFFFFFF
| Block 2:318976 has map status 0x00000000 should be 0xFFFFFFFF
| Block 1:371200 has map status 0x00000000 should be 0xFFFFFFFF
|
| Error in inode 50691 snap 0: Directory block 0 has entry with incorrect generation number
| subdir entry inode 50692 "subdir7"
| Delete entry? Yes
|
| Error in inode 50691 snap 0: Directory block 0 has entry with incorrect generation number
| subdir entry inode 13834 "subdir10"
| Delete entry? Yes
|
| Error in inode 50691 snap 0: has nlink field as 53
| Correct to 51? Yes
| 52 % complete on Tue Feb 21 13:34:53 2017
| Scanning patch file for post-mount patches
|
| Error in inode 13831 snap 0: is unreferenced
| Attach inode to lost+found of fileset root fileId 0? Yes
|
| Error in inode 13834 snap 0: is unreferenced
| Attach inode to lost+found of fileset root fileId 0? Yes
| 100 % complete on Tue Feb 21 13:34:53 2017
|

No.	SnapId	InodeNum	FileType	Fix	Error(s)	Severity
1	0	50688	User	Y	InodeMetadata + DuplicateDiskAddr	Critical
2	0	50689	User	Y	InodeMetadata + DuplicateDiskAddr	Critical
3	0	50690	User	Y	InodeMetadata + DuplicateDiskAddr	Critical
4	0	50691	Directory	Y	InodeMetadata + DirectoryEntry	Noncritical
5	0	13831	Directory	Y	OrphanInode	Noncritical
6	0	13834	Directory	Y	OrphanInode	Noncritical

| Critical corruptions were found in the file system which were repaired.
| File system is patched.
```

## **See also**

- “mmcheckquota command” on page 169
- “mmcrfs command” on page 250
- “mmdelfs command” on page 295
- “mmdf command” on page 308
- “mmlsfs command” on page 411

## **Location**

/usr/lpp/mmfs/bin

---

## mmfsctl command

Issues a file system control request.

### Synopsis

```
mmfsctl Device {suspend | suspend-write | resume}
```

or

```
mmfsctl Device {exclude | include}
 {-d "DiskName[;DiskName...]" | -F DiskFile | -G FailureGroup}
```

or

```
mmfsctl Device syncFSconfig
 {-n RemoteNodesFile | -C RemoteClusterName} [-S SpecFile]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmfsctl** command to issue control requests to a particular GPFS file system. The command is used to temporarily suspend the processing of all application I/O requests, and later resume them, as well as to synchronize the file system's configuration state between peer clusters in disaster recovery environments.

See *Establishing disaster recovery for your GPFS cluster in IBM Spectrum Scale: Administration Guide*.

#### Using mmfsctl suspend and mmfsctl resume

Before creating a FlashCopy® image of the file system, the user must run **mmfsctl suspend** to temporarily quiesce all file system activity and flush the internal buffers on all nodes that mount this file system. The on-disk metadata will be brought to a consistent state, which provides for the integrity of the FlashCopy snapshot. If a request to the file system is issued by the application after the invocation of this command, GPFS suspends this request indefinitely, or until the user issues **mmfsctl resume**.

Once the FlashCopy image has been taken, the **mmfsctl resume** command can be issued to resume the normal operation and complete any pending I/O requests.

#### Using mmfsctl syncFSconfig

The **mmfsctl syncFSconfig** command extracts the file system's related information from the local GPFS configuration data, transfers this data to one of the nodes in the peer cluster, and attempts to import it there.

Once the GPFS file system has been defined in the primary cluster, users run this command to import the configuration of this file system into the peer recovery cluster. After producing a FlashCopy image of the file system and propagating it to the peer cluster using Peer-to-Peer Remote Copy (PPRC), users similarly run this command to propagate any relevant configuration changes made in the cluster after the previous snapshot.

The primary cluster configuration server of the peer cluster must be available and accessible using remote shell and remote copy at the time of the invocation of the **mmfsctl syncFSconfig** command, and remote nodes must be reachable by the ping utility. Also, the peer GPFS clusters should be defined to use the same remote shell and remote copy mechanism, and they must be set up to allow nodes in peer clusters to communicate without the use of a password.

**Note:** In a cluster that is CCR-enabled, you cannot run **mmfsctl syncFSconfig** on a file system that has tiebreaker disks.

## **mmfsctl**

Not all administrative actions performed on the file system necessitate this type of resynchronization. It is required only for those actions that modify the file system information maintained in the local GPFS configuration data. These actions include:

- Adding, removing, and replacing disks (commands **mmadddisk**, **mmdeldisk**, **mmrpldisk**)
- Modifying disk attributes (command **mmchdisk**)
- Changing the file system's mount point (command **mmchfs -T**)
- Changing the file system device name (command **mmchfs -W**)

The process of synchronizing the file system configuration data can be automated by utilizing the **syncfsconfig** user exit.

### **Using mmfsctl exclude**

The **mmfsctl exclude** command is to be used only in a disaster recovery environment, only after a disaster has occurred, and only after ensuring that the disks in question have been physically disconnected. Otherwise, unexpected results may occur.

The **mmfsctl exclude** command can be used to manually override the file system descriptor quorum after a site-wide disaster. See *Establishing disaster recovery for your GPFS cluster* in *IBM Spectrum Scale: Administration Guide*. This command enables users to restore normal access to the file system with less than a quorum of available file system descriptor replica disks, by effectively excluding the specified disks from all subsequent operations on the file system descriptor. After repairing the disks, the **mmfsctl include** command can be issued to restore the initial quorum configuration.

## **Parameters**

### *Device*

The device name of the file system. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**. If **all** is specified with the **syncFSconfig** option, this command is performed on all GPFS file systems defined in the cluster.

The following options can be specified after *Device*:

#### **suspend**

Instructs GPFS to flush the internal buffers on all nodes, bring the file system to a consistent state on disk, and suspend the processing of all subsequent application I/O requests.

#### **suspend-write**

Suspends the execution of all new write I/O requests coming from user applications, flushes all pending requests on all nodes, and brings the file system to a consistent state on disk.

#### **resume**

Instructs GPFS to resume the normal processing of I/O requests on all nodes.

#### **exclude**

Instructs GPFS to exclude the specified group of disks from all subsequent operations on the file system descriptor, and change their availability state to **down**, if the conditions in the following Note are met.

If necessary, this command assigns additional disks to serve as the disk descriptor replica holders, and migrate the disk descriptor to the new replica set. The excluded disks are not deleted from the file system, and still appear in the output of the **mmlsdisk** command.

**Note:** The **mmfsctl exclude** command is to be used only in a disaster recovery environment, only after a disaster has occurred, and only after ensuring that the disks in question have been physically disconnected. Otherwise, unexpected results may occur.

**include**

Informs GPFS that the previously excluded disks have become operational again. This command writes the up-to-date version of the disk descriptor to each of the specified disks, and clears the **excl** tag.

**-d "DiskName[;DiskName...]"**

Specifies the names of the NSDs to be included or excluded by the **mmfsctl** command. Separate the names with semicolons (;) and enclose the list of disk names in quotation marks.

**-F DiskFile**

Specifies a file containing the names of the NSDs, one per line, to be included or excluded by the **mmfsctl** command.

**-G FailureGroup**

A failure group identifier for the disks to be included or excluded by the **mmfsctl** command.

**syncFSconfig**

Synchronizes the configuration state of a GPFS file system between the local cluster and its peer in two-cluster disaster recovery configurations.

The following options can be specified after **syncFSconfig**:

**-n RemoteNodesFile**

Specifies a list of contact nodes in the peer recovery cluster that GPFS uses when importing the configuration data into that cluster. Although any node in the peer cluster can be specified here, users are advised to specify the identities of the peer cluster's primary and secondary cluster configuration servers, for efficiency reasons.

**-c RemoteClusterName**

Specifies the name of the GPFS cluster that owns the remote GPFS file system.

**-s SpecFile**

Specifies the description of changes to be made to the file system, in the peer cluster during the import step. The format of this file is identical to that of the *ChangeSpecFile* used as input to the **mmimportfs** command. This option can be used, for example, to define the assignment of the NSD servers for use in the peer cluster.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Results**

The **mmfsctl** command returns 0 if successful.

**Security**

You must have root authority to run the **mmfsctl** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **mmfsctl**

### **Examples**

This sequence of commands creates a FlashCopy image of the file system and propagates this image to the recovery cluster using the Peer-to-Peer Remote Copy technology. The following configuration is assumed:

| <b>Site</b>               | <b>LUNs</b>  |
|---------------------------|--------------|
| Primary cluster (site A)  | lunA1, lunA2 |
| Recovery cluster (site B) | lunB1        |

**lunA1** FlashCopy source

**lunA2** FlashCopy target, PPRC source

**lunB1** PPRC target

A single GPFS file system named **fs0** has been defined in the primary cluster over lunA1.

1. In the primary cluster, suspend all file system I/O activity and flush the GPFS buffers  
`mmfsctl fs0 suspend`

The output is similar to this:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
```

2. Establish a FlashCopy pair using lunA1 as the source and lunA2 as the target.
3. Resume the file system I/O activity:

```
mmfsctl fs0 resume
```

The output is similar to this:

```
Resuming operations.
```

4. Establish a Peer-to-Peer Remote Copy (PPRC) path and a synchronous PPRC volume pair lunA2-lunB1 (primary-secondary). Use the 'copy entire volume' option and leave the 'permit read from secondary' option disabled.
5. Wait for the completion of the FlashCopy background task. Wait for the PPRC pair to reach the duplex (fully synchronized) state.
6. Terminate the PPRC volume pair lunA2-lunB1.
7. If this is the first time the snapshot is taken, or if the configuration state of **fs0** changed since the previous FlashCopy snapshot, propagate the most recent configuration to site B:

```
mmfsctl fs0 syncFSconfig -n recovery_clust_nodelist
```

### **Location**

/usr/lpp/mmfs/bin

## mmgetacl command

Displays the GPFS access control list of a file or directory.

### Synopsis

```
mmgetacl [-d] [-o OutFilename] [-k {nfs4 | posix | native}] Filename
```

### Availability

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

### Description

Use the **mmgetacl** command to display the ACL of a file or directory.

For information about NFS V4 ACLs, see the topics *Managing GPFS access control lists* and *NFS and GPFS* in the *IBM Spectrum Scale: Administration Guide*.

Users may need to see ACLs in their true form as well as how they are translated for access evaluations. There are four cases:

1. By default, **mmgetacl** returns the ACL in a format consistent with the file system setting, specified using the **-k** flag on the **mmcrfs** or **mmchfs** commands.
  - If the setting is **posix**, the ACL is shown as a traditional ACL.
  - If the setting is **nfs4**, the ACL is shown as an NFS V4 ACL.
  - If the setting is **all**, the ACL is returned in its true form.
2. The command **mmgetacl -k nfs4** always produces an NFS V4 ACL.
3. The command **mmgetacl -k posix** always produces a traditional ACL.
4. The command **mmgetacl -k native** always shows the ACL in its true form regardless of the file system setting.

The following describes how **mmgetacl** works for POSIX and NFS V4 ACLs:

| Command            | ACL   | mmcrfs -k | Display       | -d (default)   |
|--------------------|-------|-----------|---------------|----------------|
| mmgetacl           | posix | posix     | Access ACL    | Default ACL    |
| mmgetacl           | posix | nfs4      | NFS V4 ACL    | Error[1]       |
| mmgetacl           | posix | all       | Access ACL    | Default ACL    |
| mmgetacl           | nfs4  | posix     | Access ACL[2] | Default ACL[2] |
| mmgetacl           | nfs4  | nfs4      | NFS V4 ACL    | Error[1]       |
| mmgetacl           | nfs4  | all       | NFS V4 ACL    | Error[1]       |
| mmgetacl -k native | posix | any       | Access ACL    | Default ACL    |
| mmgetacl -k native | nfs4  | any       | NFS V4 ACL    | Error[1]       |
| mmgetacl -k posix  | posix | any       | Access ACL    | Default ACL    |
| mmgetacl -k posix  | nfs4  | any       | Access ACL[2] | Default ACL[2] |
| mmgetacl -k nfs4   | any   | any       | NFS V4 ACL    | Error[1]       |

[1] NFS V4 ACLs include inherited entries. Consequently, there cannot be a separate default ACL.

[2] Only the mode entries (owner, group, everyone) are translated. The **rwx** values are derived from the NFS V4 file mode attribute. Since the NFS V4 ACL is more granular in nature, some information is lost in this translation.

### Parameters

#### *Filename*

The path name of the file or directory for which the ACL is to be displayed. If the **-d** option is specified, *Filename* must contain the name of a directory.

## **mmgetacl**

### **Options**

- d** Specifies that the default ACL of a directory is to be displayed.
- k {nfs4 | posix | native}**
  - nfs4**  
Always produces an NFS V4 ACL.
  - posix**  
Always produces a traditional ACL.
  - native**  
Always shows the ACL in its true form regardless of the file system setting.
- o OutFilename**  
The path name of a file to which the ACL is to be written.

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have read access to the directory where the file exists to run the **mmgetacl** command.

You may issue the **mmgetacl** command only from a node in the GPFS cluster where the file system is mounted.

### **Examples**

1. To display the ACL for a file named **project2.history**, issue this command:

```
mmgetacl project2.history
```

The system displays information similar to:

```
#owner:paul
#group:design
user::rwx
group::r-x-
other::r-x-
```

2. This is an example of an NFS V4 ACL displayed using **mmgetacl**. Each entry consists of three lines reflecting the greater number of permissions in a text format. An entry is either an **allow** entry or a **deny** entry. An **X** indicates that the particular permission is selected, a minus sign (-) indicates that it is not selected. The following access control entry explicitly allows **READ**, **EXECUTE** and **READ\_ATTR** to the **staff** group on a file:

```
group:staff:r-x-:allow
(X)READ/LIST (-)WRITE/CREATE (-)APPEND/MKDIR (-)SYNCHRONIZE (-)READ_ACL (X)READ_ATTR (-)READ_NAMED
(-)DELETE (-)DELETE_CHILD (-)CHOWN (X)EXEC/SEARCH (-)WRITE_ACL (-)WRITE_ATTR (-)WRITE_NAMED
```

3. This is an example of a directory ACLs, which may include *inherit* entries (the equivalent of a default ACL). These do not apply to the directory itself, but instead become the initial ACL for any objects created within the directory. The following access control entry explicitly denies **READ/LIST**, **READ\_ATTR**, and **EXEC/SEARCH** to the **sys** group.

```
group:sys:----:deny:DirInherit
(X)READ/LIST (-)WRITE/CREATE (-)APPEND/MKDIR (-)SYNCHRONIZE (-)READ_ACL (X)READ_ATTR (-)READ_NAMED
(-)DELETE (-)DELETE_CHILD (-)CHOWN (X)EXEC/SEARCH (-)WRITE_ACL (-)WRITE_ATTR (-)WRITE_NAMED
```

**See also**

- “mmeditacl command” on page 322
- “mmdelacl command” on page 284
- “mmputacl command” on page 504

**Location**

/usr/lpp/mmfs/bin

## **mmgetstate**

---

### **mmgetstate command**

Displays the state of the GPFS daemon on one or more nodes.

#### **Synopsis**

```
mmgetstate [-L] [-s] [-v] [-Y] [-a | -N {Node[,Node...] | NodeFile | NodeClass}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmgetstate** command displays the state of the GPFS daemon on the specified nodes.

#### **Parameters**

- | **-a** Displays the state of the GPFS daemon on all nodes in the cluster.
- | **-N {Node[,Node...] | NodeFile | NodeClass}**
  - Displays the state of the GPFS daemon information on the specified nodes.
  - For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.
  - This command does not support a *NodeClass* of **mount**.

#### **Options**

- | **-L** Displays the number of quorum nodes, the number of nodes that are up, the total number of nodes, the state of the GPFS daemon, and other information.
- | **-s** Displays summary information, such as the number of local and remote nodes that are joined in the cluster and the number of quorum nodes.
- | **-v** Displays intermediate error messages.
- | **-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

The command recognizes and displays the following GPFS states:

**active** The GPFS daemon is ready for operations.

**arbitrating**

A node is trying to form a quorum with the other available nodes.

**down** The GPFS daemon is not running on the node or is recovering from an internal error.

**unknown**

An unknown value. The command cannot connect with the node or some other error occurred.

- | **unresponsive**

The GPFS daemon is running but is not responding.

#### **Exit status**

- 0 Successful completion.

**nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmgetstate** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To display the quorum, the number of nodes up, and the total number of nodes for the GPFS cluster, issue the following command:

```
mmgetstate -a -L
```

The system displays output similar to the following example:

| Node number | Node name | Quorum | Nodes up | Total nodes | GPFS state | Remarks     |
|-------------|-----------|--------|----------|-------------|------------|-------------|
| 1           | c5n92     | 3      | 5        | 12          | active     |             |
| 2           | c5n94     | 3      | 5        | 12          | active     |             |
| 3           | c5n95     | 3      | 5        | 12          | active     | quorum node |
| 4           | c5n96     | 3      | 5        | 12          | active     |             |
| 5           | c5n97     | 3      | 5        | 12          | active     | quorum node |
| 6           | c5n98     | 3      | 5        | 12          | active     |             |
| 7           | c5n107    | 3      | 5        | 12          | active     | quorum node |
| 8           | c5n108    | 3      | 5        | 12          | active     |             |
| 9           | c5n109    | 3      | 5        | 12          | active     | quorum node |
| 10          | c5n110    | 3      | 5        | 12          | down       |             |
| 11          | c5n111    | 3      | 5        | 12          | active     | quorum node |
| 12          | c5n112    | 3      | 5        | 12          | active     |             |

| Node number | Node name | Quorum | Nodes up | Total nodes | GPFS state | Remarks     |
|-------------|-----------|--------|----------|-------------|------------|-------------|
| 1           | c5n92     | 3      | 5        | 12          | active     |             |
| 2           | c5n94     | 3      | 5        | 12          | active     |             |
| 3           | c5n95     | 3      | 5        | 12          | active     | quorum node |
| 4           | c5n96     | 3      | 5        | 12          | active     |             |
| 5           | c5n97     | 3      | 5        | 12          | active     | quorum node |
| 6           | c5n98     | 3      | 5        | 12          | active     |             |
| 7           | c5n107    | 3      | 5        | 12          | active     | quorum node |
| 8           | c5n108    | 3      | 5        | 12          | active     |             |
| 9           | c5n109    | 3      | 5        | 12          | active     | quorum node |
| 10          | c5n110    | 3      | 5        | 12          | down       |             |
| 11          | c5n111    | 3      | 5        | 12          | active     | quorum node |
| 12          | c5n112    | 3      | 5        | 12          | active     |             |

The 3 under the Quorum column means that you must have three quorum nodes up to achieve quorum.

2. In the following example, the cluster uses node quorum with tiebreaker disks. The asterisk (\*) in the Quorum field indicates that tiebreaker disks are being used:

```
mmgetstate -a -L
```

The system displays output similar to the following example:

| Node number | Node name | Quorum | Nodes up | Total nodes | GPFS state | Remarks     |
|-------------|-----------|--------|----------|-------------|------------|-------------|
| 1           | k5n91     | 5*     | 8        | 21          | active     |             |
| 2           | k5n92     | 5*     | 8        | 21          | active     | quorum node |
| 3           | k5n94     | 5*     | 8        | 21          | active     |             |
| 5           | k5n96     | 5*     | 8        | 21          | active     |             |
| 6           | k5n97     | 5*     | 8        | 21          | active     | quorum node |
| 7           | k5n98     | 5*     | 8        | 21          | active     |             |
| 8           | k5n99     | 5*     | 8        | 21          | active     | quorum node |

3. To display summary information, issue this command:

```
mmgetstate -s
```

The system displays output similar to the following example:

| Node number | Node name | GPFS state |
|-------------|-----------|------------|
| 5           | c5n97     | active     |

Summary information

| Number of nodes defined in the cluster: | 12 |
|-----------------------------------------|----|
|-----------------------------------------|----|

## **mmgetstate**

```
Number of local nodes active in the cluster: 12
Number of remote nodes joined in this cluster: 0
Number of quorum nodes defined in the cluster: 5
Number of quorum nodes active in the cluster: 5
Quorum = 3, Quorum achieved
```

### **See also**

- “mmchconfig command” on page 132
- “mmcrcluster command” on page 239
- “mmshutdown command” on page 552
- “mmstartup command” on page 569

### **Location**

/usr/lpp/mmfs/bin

---

## mmhadoopctl command

Installs and sets up the GPFS connector for a Hadoop distribution; starts or stops the GPFS connector daemon on a node.

### Synopsis

```
mmhadoopctl connector {start |stop |getstate}
```

or

```
mmhadoopctl connector syncconf [--nocheck] [path]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmhadoopctl** command to install and set up the GPFS connector for a Hadoop distribution, or to start or stop the GPFS connector daemon on a node.

### Parameters

#### **connector**

Controls the GPFS connector daemon with one of the following actions:

##### **start**

Starts the connector daemon.

##### **stop**

Stops the connector daemon.

##### **getstate**

Detects whether the connector daemon is running and shows its process ID.

#### **syncconf**

Synchronizes the connector configuration in the cluster.

#### **--nocheck**

If --nocheck is used, the sanity check is not performed. Therefore, use this option with caution.

#### **path**

Specifies the path to the config directory or file.

### Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### Security

You must have root authority to run the **mmhadoopctl** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **mmces**

### **Examples**

To start the GPFS connector daemon, issue this command:

```
mmhadoopctl connector start
```

The system displays output similar to this:

```
Hadoop connector 'gpfs-connector-daemon' started
```

### **Location**

```
/usr/lpp/mmfs/bin
```

---

## mmhealth command

Monitors health status of nodes.

### Synopsis

```
mmhealth node show [GPFS | NETWORK [UserDefinedSubComponent]
FILESYSTEM [UserDefinedSubComponent] | DISK [UserDefinedSubComponent]
CES | AUTH | AUTH_OBJ | BLOCK | CESNETWORK | NFS | OBJECT | SMB
HADOOP | CLOUDGATEWAY | GUI | PERFMON | THRESHOLD
AFM [UserDefinedSubComponent]]
[-N {Node[,Node..] | NodeFile | NodeClass}]
[-Y] [--verbose] [--unhealthy]
```

or

```
mmhealth node eventlog [[--hour | --day | --week | month] | [--clear] | [--verbose]]
[-N {Node[,Node..] | NodeFile | NodeClass}]
[-Y]
```

or

```
mmhealth event show [EventName | EventID] [-N {Node[,Node..] | NodeFile | NodeClass}]
```

or

```
mmhealth event hide [EventName [Entity_Name]]
```

or

```
mmhealth event unhide [EventName [Entity_Name]]
```

or

```
mmhealth event list HIDDEN
```

or

```
mmhealth cluster show [NODE | GPFS | NETWORK [UserDefinedSubComponent]
FILESYSTEM [UserDefinedSubComponent] | DISK [UserDefinedSubComponent]
CES | AUTH | AUTH_OBJ | BLOCK | CESNETWORK | NFS | OBJECT | SMB
HADOOP | CLOUDGATEWAY | GUI | PERFMON | THRESHOLD
AFM [UserDefinedSubComponent]]
[-Y] [--verbose]
```

or

```
mmhealth thresholds list [--verbose]
```

or

```
mmhealth thresholds add { metric [: sum | avg | min | max | rate]| measurement
[--errorlevel{threshold error limit}][--warnlevel{threshold warn limit }]
[--direction { high|low }]
[--sensitivity { bucketsize }] [--hysteresis { percentage }]
[--filterBy] [--groupBy] [--name { ruleName }]
[--errormsg { user defined action description }]
[--warnmsg { user defined action description }]
```

or

```
mmhealth thresholds delete { ruleName | all }
```

or

```
mmhealth config interval [OFF | LOW | MEDIUM | DEFAULT | HIGH]
```

## **mmhealth**

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

Use the **mmhealth** command to monitor the health of the node and services hosted on the node in IBM Spectrum Scale.

The IBM Spectrum Scale administrator can monitor the health of each node and the services hosted on that node using the **mmhealth** command. The **mmhealth** command also shows the events that are responsible for the unhealthy status of the services hosted on the node. This data can be used to monitor and analyze the reasons for the unhealthy status of the node. The **mmhealth** command acts as a problem determination tool to identify which services of the node are unhealthy, and find the events responsible for the unhealthy state of the service.

The **mmhealth** command also monitors the state of all the IBM Spectrum Scale RAID components such as array, pdisk, vdisk, and enclosure of the nodes that belong to the recovery group.

For more information about the system monitoring feature, see the *Monitoring system health by using the mmhealth command* section in the *IBM Spectrum Scale: Problem Determination Guide*.

The **mmhealth** command shows the details of threshold rules. This detail helps to avoid out-of-space errors for filesystems. The space availability of the filesystem component depends upon the occupancy level of fileset-inode spaces and the capacity usage in each data or metadata pool. The violation of any single rule triggers the parent filesystem's capacity-issue events. The capacity metrics are frequently compared with the rules boundaries by internal monitor process. If any of the metric values exceeds their threshold limit, then the system health (daemon/service) will receive an event notification from monitor process and generate a RAS event for the filesystem for space issues. For the predefined capacity utilization rules, the warn level is set to 80%, and the error level to 90%. For memory utilization rule, the warn level is set to 100 MB, and the error level to 50 MB. You can use the **mm1sfileset** and the **mm1spool** commands to track the inode and pool space usage.

### **Parameters**

#### **event**

Gives the details of various events:

##### **show**

Shows the detailed description of the specified event:

##### **EventName**

Displays the detailed description of the specified event name.

##### **EventID**

Displays the detailed description of the specified event ID.

##### **hide**

Hides the specified TIP events.

##### **unhide**

Reveals the TIP events that were previously hidden using the hide.

##### **list HIDDEN**

Shows all the TIP events that are added to the list of hidden events.

#### **node**

Displays the health status, specifically, at node level.

**show**

Displays the health status of the specified component with:

**GPFS | NETWORK | FILESYSTEM | DISK | CES | AUTH | AUTH\_OBJ | BLOCK | CESNETWORK | NFS | OBJECT | SMB | HADOOP | CLOUDGATEWAY | GUI | PERFMON | THRESHOLD | AFM**

Displays the detailed health status of the specified component.

**UserDefinedSubComponent**

Displays services that are named by the customer, categorized by one of the other hosted services. For example, a filesystem named gpfs0 is a subcomponent of filesystem.

- N Allows the system to make remote calls to the other nodes in the cluster for:

**Node[ ,Node....]**

Specifies the node or list of nodes that must be monitored for the health status.

**NodeFile**

Specifies a file, containing a list of node descriptors, one per line, to be monitored for health status.

**NodeClass**

Specifies a node class that must be monitored for the health status.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**--verbose**

Shows the detailed health status of a node, including its sub-components.

**--unhealthy**

Displays the unhealthy components only.

**eventlog**

Shows the event history for a specified period of time. If no time period is specified, it displays all the events by default:

**[--hour | --day | --week| --month]**

Displays the event history for the specified time period.

**[--clear]**

Clears the event log's database. This action cannot be reversed.

**CAUTION:**

The events database is used by the mmhealth node eventlog as well as the mmces events list. If you clear the database, it will also affect the mmces events list. Ensure that you use the --clear option with caution.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**[--verbose]**

Displays additional information about the event like component name and event ID in the eventlog.

## **mmhealth**

### **event**

Gives the details of various events:

#### **show**

Shows the detailed description of the specified event:

##### **EventName**

Displays the detailed description of the specified event name.

##### **EventID**

Displays the detailed description of the specified event ID.

#### **hide**

Hides the specified TIP events.

#### **unhide**

Reveals the TIP events that were previously hidden using the hide.

#### **list HIDDEN**

Shows all the TIP events that are added to the list of hidden events.

### **cluster**

Displays the health status of all nodes and monitored node components in the cluster.

#### **show**

Displays the health status of the specified component with:

**NODE | GPFS | NETWORK | FILESYSTEM | DISK | CES | AUTH | AUTH\_OBJ | BLOCK | CESNETWORK | NFS | OBJECT | SMB | HADOOP | CLOUDGATEWAY | GUI | PERFMON | THRESHOLD | AFM**

Displays the detailed health status of the specified component.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

#### **--verbose**

Shows the detailed health status of a node, including its sub-components.

### **thresholds list**

Displays the list of the threshold rules defined for the system.

### **thresholds add**

Creates a new thresholds rule for the specified metric or measurement, and activates monitoring process stores for this rule.

**Note:** A measurement is a value calculated using more than one metric in a pre-defined formula.

#### **metric [: SUM | AVG | MIN | MAX | RATE ]**

Creates a threshold for the specified metric. All metrics that are supported by the performance monitoring tool, and use raw values or are downsampled by aggregators (sum, avg, min, max, rate) can be used. For a list of metrics supported by the performance monitoring tool, see the *List of performance metrics* section in the *IBM Spectrum Scale: Problem Determination Guide*.

#### **measurement**

Creates a threshold for the specified measurement. The following measurements are supported:

#### **Fileset\_inode**

Fileset Inode Capacity Utilization. Calculated as:

$sum(gpfs\_fset\_allocInodes)-sum(gpfs\_fset\_freeInodes)/sum(gpfs\_fset\_maxInodes)$

**DataPool\_capUtil**

Data Pool Capacity Utilization. Calculated as:

$$\text{sum}(gpfs\_pool\_total\_dataKB)-\text{sum}(gpfs\_pool\_free\_dataKB)/\text{sum}(gpfs\_pool\_total\_dataKB)$$

**MetaDataPool\_capUtil**

MetaData Pool Capacity Utilization. Calculated as:

$$\text{sum}(gpfs\_pool\_total\_metaKB)-\text{sum}(gpfs\_pool\_free\_metaKB))/\text{sum}(gpfs\_pool\_total\_metaKB)$$

**--errorlevel**

Defines the threshold error limit. The threshold error limit can be a percentage or an integer, depending on the metric on which the threshold value is being set.

**--warnlevel**

Defines the threshold warn limit. The threshold warn limit can be a percentage or an integer, depending on the metric on which the threshold value is being set.

**--direction**

Defines the direction for the threshold limit. The allowed values are *high* or *low*.

**--groupby**

Groups the result based on the group key. The following values are allowed for the group key:

- *gpfs\_cluster\_name*
- *gpfs\_disk\_name*
- *gpfs\_diskpool\_name*
- *gpfs\_disk\_usage\_name*
- *gpfs\_fset\_name*
- *gpfs\_fs\_name*
- *mountPoint*
- *netdev\_name*
- *node*

**--filterby**

Filters the result based on the filter key. The following values are allowed for the filter key:

- *gpfs\_cluster\_name*
- *gpfs\_disk\_name*
- *gpfs\_diskpool\_name*
- *gpfs\_disk\_usage\_name*
- *gpfs\_fset\_name*
- *gpfs\_fs\_name*
- *mountPoint*
- *netdev\_name*
- *node*

**--sensitivity**

Defines the sample interval value in seconds. It is set to 300 by default. If a sensors is configured with interval period greater than 300 seconds, then the --sensitivity will be set to the same value as *sensors period*. The minimum value allowed is 120 seconds. If a sensors is configured with interval period less than 120 seconds, the --sensitivity will be set to 120 seconds.

**--hysteresis**

Defines the percentage that the observed value must be under (or over) the current threshold level to switch back to the previous state. The default value is 0.0, while the recommended value is 5.0

## **mmhealth**

### **--name**

Defines the name of the rule. It can be an alphanumeric string with up to 30 characters. If the rule name is not specified, default name will be set. The default name is set using the metric name followed by underscore and then a "custom" prefix.

### **--errormsg**

This is a user defined input. The message can be 256 bytes long. It must be added within double quotes (""), else the system will throw an error.

### **--warnmsg**

This is a user defined input. The message can be 256 bytes long. It must be added within double quotes (""), else the system will throw an error.

#### **Important:**

- The mathematical aggregations: AVG, SUM, MAX, MIN, RATE could be used to determine how to merge the metric values in the evaluation source. The aggregation operations are not supported for measurements.
- For each rule the user can configure up to two conditions, --error and --warn, triggering event state change. At least one level limit setting is required. For example, the threshold add command must have one of the following options:
  - mmhealth thresholds add { metric[:sum|avg|min|max|rate] | measurement  
--error{threshold error limit} ---direction {high|low}
  - mmhealth thresholds add { metric[:sum|avg|min|max|rate] | measurement  
--error{threshold error limit} --warn{threshold warn limit}
  - mmhealth thresholds add { metric[:sum|avg|min|max|rate] | measurement  
--error{threshold error limit} --warn{threshold warn limit} ---direction {high|low}}
- The customer can also influence the measuring quantity and precision by specifying sensitivity, groupby, filterby, histersesis, or rule name option setting.
- For each condition level the customer can leave an output message text using --errormsg or --warnmsg options, which will be integrated into the state change event notification message, triggered by the exceeding this condition.

## **thresholds delete**

Deletes the threshold rules from the system.

### **ruleName**

Deletes a specific threshold rule.

### **all**

Deletes all the threshold rules.

**Note:** Using the **mmhealth thresholds delete** command to delete a rule will accomplish the following tasks:

- The rule will be removed from the thresholds rules specification file and active monitoring process.
- All the current health information created by this particular rule will be removed as well.

## **config interval**

Sets the monitoring interval for the whole cluster.

### **off**

The monitoring will be off for the whole cluster.

### **low**

Monitoring is set for every (**default monitoring time \*10**) seconds.

### **medium**

Monitoring is set for every (**default monitoring time \*5**) seconds.

### **default**

Monitoring is set for every 15-30 seconds based on the service being monitored

**high**

Monitoring is set for every (**default monitoring time /2**) seconds.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Security**

You must have root authority to run the **mmhealth** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. See the information about the requirements for administering a GPFS system in the *IBM Spectrum Scale: Administration Guide*.

**Examples**

1. To show the health status of the current node, issue this command:

```
mmhealth node show
```

The system displays output similar to this:

```
Node name: test_node
Node status: HEALTHY
Status Change: 39 min. ago
```

| Component  | Status  | Status Change | Reasons |
|------------|---------|---------------|---------|
| GPFS       | HEALTHY | 39 min. ago   | -       |
| NETWORK    | HEALTHY | 40 min. ago   | -       |
| FILESYSTEM | HEALTHY | 39 min. ago   | -       |
| DISK       | HEALTHY | 39 min. ago   | -       |
| CES        | HEALTHY | 39 min. ago   | -       |
| PERFMON    | HEALTHY | 40 min. ago   | -       |
| THRESHOLD  | HEALTHY | 40 min. ago   | -       |

2. To view the health status of a specific node, issue this command:

```
mmhealth node show -N test_node2
```

The system displays output similar to this:

```
Node name: test_node2
Node status: CHECKING
Status Change: Now
```

| Component  | Status   | Status Change | Reasons |
|------------|----------|---------------|---------|
| GPFS       | CHECKING | Now           | -       |
| NETWORK    | HEALTHY  | Now           | -       |
| FILESYSTEM | CHECKING | Now           | -       |
| DISK       | CHECKING | Now           | -       |
| CES        | CHECKING | Now           | -       |
| PERFMON    | HEALTHY  | Now           | -       |

3. To view the health status of all the nodes, issue this command:

```
mmhealth node show -N all
```

The system displays output similar to this:

## mmhealth

```
Node name: test_node
Node status: DEGRADED
```

| Component  | Status  | Status Change | Reasons   |
|------------|---------|---------------|-----------|
| GPFS       | HEALTHY | Now           | -         |
| CES        | FAILED  | Now           | smbd_down |
| FileSystem | HEALTHY | Now           | -         |

```
Node name: test_node2
Node status: HEALTHY
```

| Component  | Status  | Status Change | Reasons |
|------------|---------|---------------|---------|
| GPFS       | HEALTHY | Now           | -       |
| CES        | HEALTHY | Now           | -       |
| FileSystem | HEALTHY | Now           | -       |

4. To view the detailed health status of the component and its sub-component, issue this command:

```
mmhealth node show ces
```

The system displays output similar to this:

```
Node name: test_node
```

| Component  | Status   | Status Change | Reasons |
|------------|----------|---------------|---------|
| CES        | HEALTHY  | 2 min. ago    | -       |
| AUTH       | DISABLED | 2 min. ago    | -       |
| AUTH_OBJ   | DISABLED | 2 min. ago    | -       |
| BLOCK      | DISABLED | 2 min. ago    | -       |
| CESNETWORK | HEALTHY  | 2 min. ago    | -       |
| NFS        | HEALTHY  | 2 min. ago    | -       |
| OBJECT     | DISABLED | 2 min. ago    | -       |
| SMB        | HEALTHY  | 2 min. ago    | -       |

5. To view the health status of only unhealthy components, issue this command:

```
mmhealth node show --unhealthy
```

The system displays output similar to this:

```
Node name: test_node
Node status: FAILED
Status Change: 1 min. ago
```

| Component  | Status | Status Change | Reasons                            |
|------------|--------|---------------|------------------------------------|
| GPFS       | FAILED | 1 min. ago    | gpfs_down, quorum_down             |
| FILESYSTEM | DEPEND | 1 min. ago    | unmounted_fs_check                 |
| CES        | DEPEND | 1 min. ago    | ces_network_ips_down, nfs_in_grace |

6. To view the health status of sub-components of a node's component, issue this command:

```
mmhealth node show --verbose
```

The system displays output similar to this:

```
Node name: gssio1-hs.gpfs.net
Node status: HEALTHY
```

| Component  | Status   | Reasons                               |
|------------|----------|---------------------------------------|
| GPFS       | DEGRADED | -                                     |
| NETWORK    | HEALTHY  | -                                     |
| bond0      | HEALTHY  | -                                     |
| ib0        | HEALTHY  | -                                     |
| ib1        | HEALTHY  | -                                     |
| FILESYSTEM | DEGRADED | stale_mount, stale_mount, stale_mount |
| Basic1     | FAILED   | stale_mount                           |
| Basic2     | FAILED   | stale_mount                           |
| Custom1    | HEALTHY  | -                                     |
| gpfs0      | FAILED   | stale_mount                           |
| gpfs1      | FAILED   | stale_mount                           |

| DISK                                | DEGRADED | disk_down                                                    |
|-------------------------------------|----------|--------------------------------------------------------------|
| rg_gssio1_hs_Basic1_data_0          | HEALTHY  | -                                                            |
| rg_gssio1_hs_Basic1_system_0        | HEALTHY  | -                                                            |
| rg_gssio1_hs_Basic2_data_0          | HEALTHY  | -                                                            |
| rg_gssio1_hs_Basic2_system_0        | HEALTHY  | -                                                            |
| rg_gssio1_hs_Custom1_data_0         | HEALTHY  | -                                                            |
| rg_gssio1_hs_Custom1_system_0       | DEGRADED | disk_down                                                    |
| rg_gssio1_hs_Data_8M_2p_1_gpfs0     | HEALTHY  | -                                                            |
| rg_gssio1_hs_Data_8M_3p_1_gpfs1     | HEALTHY  | -                                                            |
| rg_gssio1_hs_MetaData_1M_3W_1_gpfs0 | HEALTHY  | -                                                            |
| rg_gssio1_hs_MetaData_1M_4W_1_gpfs1 | HEALTHY  | -                                                            |
| rg_gssio2_hs_Basic1_data_0          | HEALTHY  | -                                                            |
| rg_gssio2_hs_Basic1_system_0        | HEALTHY  | -                                                            |
| rg_gssio2_hs_Basic2_data_0          | HEALTHY  | -                                                            |
| rg_gssio2_hs_Basic2_system_0        | HEALTHY  | -                                                            |
| rg_gssio2_hs_Custom1_data1_0        | HEALTHY  | -                                                            |
| rg_gssio2_hs_Custom1_system_0       | HEALTHY  | -                                                            |
| rg_gssio2_hs_Data_8M_2p_1_gpfs0     | HEALTHY  | -                                                            |
| rg_gssio2_hs_Data_8M_3p_1_gpfs1     | HEALTHY  | -                                                            |
| rg_gssio2_hs_MetaData_1M_3W_1_gpfs0 | HEALTHY  | -                                                            |
| rg_gssio2_hs_MetaData_1M_4W_1_gpfs1 | HEALTHY  | -                                                            |
| NATIVE RAID ARRAY                   | DEGRADED | gnr_pdisk_replaceable, gnr_rg_failed, enclosure_needsservice |
| rg_gssio2-hs/DA1                    | DEGRADED | -                                                            |
| rg_gssio2-hs/DA2                    | HEALTHY  | -                                                            |
| rg_gssio2-hs/NVR                    | HEALTHY  | -                                                            |
| rg_gssio2-hs/SSD                    | HEALTHY  | -                                                            |
| ENCLOSURE SV52122944                | DEGRADED | enclosure_needsservice                                       |
| SV53058375                          | DEGRADED | enclosure_needsservice                                       |
| PHYSICALDISK rg_gssio2-hs/e1d1s01   | FAILED   | -                                                            |
| rg_gssio2-hs/e1d1s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d1s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d1s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d1s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d1s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d1s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d2s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d3s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d4s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d5s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d5s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d5s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d5s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e1d5s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d1s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d2s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d3s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s10                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s11                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d4s12                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d5s07                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d5s08                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d5s09                | HEALTHY  | -                                                            |
| rg_gssio2-hs/e2d5s10                | HEALTHY  | -                                                            |

## mmhealth

|                                     |          |               |
|-------------------------------------|----------|---------------|
| rg_gssio2-hs/e2d5s10                | HEALTHY  | -             |
| rg_gssio2-hs/e2d5s11                | HEALTHY  | -             |
| rg_gssio2-hs/e2d5s12ssd             | HEALTHY  | -             |
| rg_gssio2-hs/n1s02                  | HEALTHY  | -             |
| rg_gssio2-hs/n2s02                  | HEALTHY  | -             |
| RECOVERYGROUP                       | DEGRADED | gnr_rg_failed |
| rg_gssio1-hs                        | FAILED   | gnr_rg_failed |
| rg_gssio2-hs                        | HEALTHY  | -             |
| VIRTUALDISK                         | DEGRADED | -             |
| rg_gssio2_hs_Basic1_data_0          | HEALTHY  | -             |
| rg_gssio2_hs_Basic1_system_0        | HEALTHY  | -             |
| rg_gssio2_hs_Basic2_data_0          | HEALTHY  | -             |
| rg_gssio2_hs_Basic2_system_0        | HEALTHY  | -             |
| rg_gssio2_hs_Custom1_data1_0        | HEALTHY  | -             |
| rg_gssio2_hs_Custom1_system_0       | HEALTHY  | -             |
| rg_gssio2_hs_Data_8M_2p_1_gpfs0     | HEALTHY  | -             |
| rg_gssio2_hs_Data_8M_3p_1_gpfs1     | HEALTHY  | -             |
| rg_gssio2_hs_MetaData_1M_3W_1_gpfs0 | HEALTHY  | -             |
| rg_gssio2_hs_MetaData_1M_4W_1_gpfs1 | HEALTHY  | -             |
| rg_gssio2_hs_Loghome                | HEALTHY  | -             |
| rg_gssio2_hs_logtip                 | HEALTHY  | -             |
| rg_gssio2_hs_logtipbackup           | HEALTHY  | -             |
| PERFMON                             | HEALTHY  | -             |

7. To view the eventlog history of the node for the last hour, issue this command:

```
mmhealth node eventlog --hour
```

The system displays output similar to this:

| Node name:                      | test-21.localnet.com | Event Name           | Severity | Details                                                    |
|---------------------------------|----------------------|----------------------|----------|------------------------------------------------------------|
|                                 |                      | monitor_started      | INFO     | The IBM Spectrum Scale monitoring service has been started |
| 2016-10-28 06:59:34.045980 CEST |                      | fs_remount_mount     | INFO     | The filesystem objfs was mounted internal                  |
| 2016-10-28 07:01:21.919943 CEST |                      | disk_found           | INFO     | The disk disk1 was found                                   |
| 2016-10-28 07:01:32.434703 CEST |                      | disk_found           | INFO     | The disk disk8 was found                                   |
| 2016-10-28 07:01:32.669125 CEST |                      | filesystem_found     | INFO     | Filesystem objfs was found                                 |
| 2016-10-28 07:01:36.975902 CEST |                      | unmounted_fs_check   | WARNING  | The filesystem objfs is probably needed, but not mounted   |
| 2016-10-28 07:01:37.226157 CEST |                      | mounted_fs_check     | INFO     | The filesystem objfs is mounted                            |
| 2016-10-28 07:01:52.1691 CEST   |                      | fs_remount_mount     | INFO     | The filesystem objfs was mounted normal                    |
| 2016-10-28 07:02:07.026093 CEST |                      | mounted_fs_check     | INFO     | The filesystem objfs is mounted                            |
| 2016-10-28 07:14:58.498854 CEST |                      | ces_network_ips_down | WARNING  | No CES relevant NICs detected                              |
| 2016-10-28 07:15:07.702351 CEST |                      | nodestatechange_info | INFO     | A CES node state change: Node 1 add startup flag           |
| 2016-10-28 07:15:37.322997 CEST |                      | nodestatechange_info | INFO     | A CES node state change: Node 1 remove startup flag        |
| 2016-10-28 07:15:43.741149 CEST |                      | ces_network_ips_up   | INFO     | CES-relevant IPs are served by found NICs                  |
| 2016-10-28 07:15:44.028031 CEST |                      | ces_network_vanished | INFO     | CES NIC eth0 has vanished                                  |

8. To view the eventlog history of the node for the last hour, issue this command:

```
mmhealth node eventlog --hour --verbose
```

The system displays output similar to this:

| Node name:                      | test-21.localnet.com | Component  | Event Name           | Event ID | Severity | Details                                                    |
|---------------------------------|----------------------|------------|----------------------|----------|----------|------------------------------------------------------------|
|                                 |                      | gpfs       | monitor_started      | 999726   | INFO     | The IBM Spectrum Scale monitoring service has been started |
| 2016-10-28 06:59:34.045980 CEST |                      | filesystem | fs_remount_mount     | 999306   | INFO     | The filesystem objfs was mounted internal                  |
| 2016-10-28 07:01:21.919943 CEST |                      | disk       | disk_found           | 999424   | INFO     | The disk disk1 was found                                   |
| 2016-10-28 07:01:32.434703 CEST |                      | disk       | disk_found           | 999424   | INFO     | The disk disk8 was found                                   |
| 2016-10-28 07:01:32.669125 CEST |                      | filesystem | filesystem_found     | 999299   | INFO     | Filesystem objfs was found                                 |
| 2016-10-28 07:01:36.975902 CEST |                      | filesystem | unmounted_fs_check   | 999298   | WARNING  | The filesystem objfs is probably needed, but not mounted   |
| 2016-10-28 07:01:37.226157 CEST |                      | filesystem | mounted_fs_check     | 999301   | INFO     | The filesystem objfs is mounted                            |
| 2016-10-28 07:01:52.283545 CEST |                      | filesystem | fs_remount_mount     | 999306   | INFO     | The filesystem objfs was mounted normal                    |
| 2016-10-28 07:02:07.026093 CEST |                      | filesystem | mounted_fs_check     | 999301   | INFO     | The filesystem objfs is mounted                            |
| 2016-10-28 07:14:58.498854 CEST |                      | cesnetwork | ces_network_ips_down | 999426   | WARNING  | No CES relevant NICs detected                              |
| 2016-10-28 07:15:07.702351 CEST |                      | gpfs       | nodestatechange_info | 999220   | INFO     | A CES node state change: Node 1 add startup flag           |
| 2016-10-28 07:15:37.322997 CEST |                      | gpfs       | nodestatechange_info | 999220   | INFO     | A CES node state change: Node 1 remove startup flag        |
| 2016-10-28 07:15:43.741149 CEST |                      | cesnetwork | ces_network_ips_up   | 999427   | INFO     | CES-relevant IPs are served by found NICs                  |
| 2016-10-28 07:15:44.028031 CEST |                      | cesnetwork | ces_network_vanished | 999434   | INFO     | CES NIC eth0 has vanished                                  |

9. To view the detailed description of an event, issue the **mmhealth event show** command. This is an example for *quorum\_down* event:

```
mmhealth event show quorum_down
```

The system displays output similar to this:

|              |                                                                                                                                                                       |
|--------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Event Name:  | quorum_down                                                                                                                                                           |
| Event ID:    | 999289                                                                                                                                                                |
| Description: | Reasons could be network or hardware issues, or a shutdown of the cluster service.<br>The event does not necessarily indicate an issue with the cluster quorum state. |
| Cause:       | The local node does not have quorum. The cluster service might not be running.                                                                                        |
| User Action: | Check if the cluster quorum nodes are running and can be reached over the network.<br>Check local firewall settings                                                   |
| Severity:    | ERROR                                                                                                                                                                 |
| State:       | DEGRADED                                                                                                                                                              |

10. To view the list of hidden events, issue the **mmhealth event list HIDDEN** command:

```
mmhealth event list HIDDEN
```

The system displays output similar to this:

| Event                | scope |
|----------------------|-------|
| gpfs_pagepool_small  | -     |
| nfsv4_acl_type_wrong | fs1   |
| nfsv4_acl_type_wrong | fs2   |

11. To view the detailed description of the cluster, issue the **mmhealth cluster show** command:

```
mmhealth cluster show
```

The system displays output similar to this:

| Component    | Total | Failed | Degraded | Healthy | Other |
|--------------|-------|--------|----------|---------|-------|
| NODE         | 50    | 1      | 1        | 48      | -     |
| GPFS         | 50    | 1      | -        | 49      | -     |
| NETWORK      | 50    | -      | -        | 50      | -     |
| FILESYSTEM   | 3     | -      | -        | 3       | -     |
| DISK         | 50    | -      | -        | 50      | -     |
| CES          | 5     | -      | 5        | -       | -     |
| CLOUDGATEWAY | 2     | -      | -        | 2       | -     |
| PERFMON      | 48    | -      | 5        | 43      | -     |
| THRESHOLD    | 4     | -      | -        | 4       | -     |

**Note:** The cluster must have the minimum release level as 4.2.2.0 or higher to use **mmhealth cluster show** command.

Also, this command does not support Windows operating system.

12. To view more information of the cluster health status, issue this command:

```
mmhealth cluster show --verbose
```

The system displays output similar to this:

| Component    | Total | Failed | Degraded | Healthy | Other |
|--------------|-------|--------|----------|---------|-------|
| NODE         | 50    | 1      | 1        | 48      | -     |
| GPFS         | 50    | 1      | -        | 49      | -     |
| NETWORK      | 50    | -      | -        | 50      | -     |
| FILESYSTEM   |       |        |          |         |       |
| FS1          | 15    | -      | -        | 15      | -     |
| FS2          | 5     | -      | -        | 5       | -     |
| FS3          | 20    | -      | -        | 20      | -     |
| DISK         | 50    | -      | -        | 50      | -     |
| CES          | 5     | -      | 5        | -       | -     |
| AUTH         | 5     | -      | -        | -       | 5     |
| AUTH_OBJ     | 5     | 5      | -        | -       | -     |
| BLOCK        | 5     | -      | -        | -       | 5     |
| CESNETWORK   | 5     | -      | -        | 5       | -     |
| NFS          | 5     | -      | -        | 5       | -     |
| OBJECT       | 5     | -      | -        | 5       | -     |
| SMB          | 5     | -      | -        | 5       | -     |
| CLOUDGATEWAY | 2     | -      | -        | 2       | -     |
| PERFMON      | 48    | -      | 5        | 43      | -     |
| THRESHOLD    | 4     | -      | -        | 4       | -     |

13. To create a new threshold rule, issue this command:

```
mmhealth thresholds add MetaDataPool_capUtil --errorlevel 90 --direction high
--groupby gpfs_fs_name,gpfs_diskpool_name
```

The system displays output similar to this:

```
New rule 'MetaDataPool_capUtil_custom' is created. The monitor process is activated
```

14. To view the list of threshold rules defined for the system, issue this command:

```
mmhealth thresholds list
```

## **mmhealth**

The system displays output similar to this:

| rule_name                   | metric               | error | warn   | direction | filterBy | groupBy                                             | sensitivity |
|-----------------------------|----------------------|-------|--------|-----------|----------|-----------------------------------------------------|-------------|
| InodeCapUtil_Rule           | Fileset_inode        | 90.0  | 80.0   | high      |          | gpfs_cluster_name, gpfs_fs_name, gpfs_fset_name     | 300         |
| MetaDataPool_capUtil_custom | MetaDataPool_capUtil | 90    | None   | high      |          | gpfs_fs_name, gpfs_diskpool_name                    | 300         |
| DataCapUtil_Rule            | DataPool_capUtil     | 90.0  | 80.0   | high      |          | gpfs_cluster_name, gpfs_fs_name, gpfs_diskpool_name | 300         |
| MemFree_Rule                | mem_memfree          | 50000 | 100000 | low       | node     |                                                     | 300         |
| MetadataCapUtil_Rule        | MetaDataPool_capUtil | 90.0  | 80.0   | high      |          | gpfs_cluster_name, gpfs_fs_name, gpfs_diskpool_name | 300         |

15. To view the detailed health status of filesystem component, issue this command:

```
mmhealth node show filesystem -v
```

The system displays output similar to this:

| Node name: gpfsgui-12.novalocal |           |                     |                      |                                                                                        |
|---------------------------------|-----------|---------------------|----------------------|----------------------------------------------------------------------------------------|
| Component                       | Status    | Status Change       | Reasons              |                                                                                        |
| FILESYSTEM                      | DEGRADED  | 2016-09-29 15:22:48 | pool-data_high_error |                                                                                        |
| fs1                             | FAILED    | 2016-09-29 15:22:48 | pool-data_high_error |                                                                                        |
| fs2                             | HEALTHY   | 2016-09-29 15:22:33 | -                    |                                                                                        |
| objfs                           | HEALTHY   | 2016-09-29 15:22:33 | -                    |                                                                                        |
| Event                           | Parameter | Severity            | Active Since         | Event Message                                                                          |
| pool-data_high_error            | fs1       | ERROR               | 2016-09-29 15:22:47  | The pool myPool of file system fs1 reached a nearly exhausted data level. 90.0         |
| inode_normal                    | fs1       | INFO                | 2016-09-29 15:22:47  | The inode usage of fileset root in file system fs1 reached a normal level.             |
| inode_normal                    | fs2       | INFO                | 2016-09-29 15:22:47  | The inode usage of fileset root in file system fs2 reached a normal level.             |
| inode_normal                    | objfs     | INFO                | 2016-09-29 15:22:47  | The inode usage of fileset root in file system objfs reached a normal level.           |
| inode_normal                    | objfs     | INFO                | 2016-09-29 15:22:47  | The inode usage of fileset Object_Fileset in file system objfs reached a normal level. |
| mounted_fs_check                | fs1       | INFO                | 2016-09-29 15:22:33  | The pool system of file system fs1 reached a normal data level.                        |
| mounted_fs_check                | fs2       | INFO                | 2016-09-29 15:22:33  | The pool system of file system fs2 reached a normal data level.                        |
| mounted_fs_check                | objfs     | INFO                | 2016-09-29 15:22:33  | The pool system of file system objfs reached a normal data level.                      |
| mounted_fs_check                | objfs     | INFO                | 2016-09-29 15:22:33  | The pool system of file system objfs reached a normal data level.                      |
| pool-data_normal                | fs1       | INFO                | 2016-09-29 15:22:47  | The pool system of file system fs1 reached a normal data level.                        |
| pool-data_normal                | fs2       | INFO                | 2016-09-29 15:22:47  | The pool system of file system fs2 reached a normal data level.                        |
| pool-data_normal                | objfs     | INFO                | 2016-09-29 15:22:47  | The pool system of file system objfs reached a normal data level.                      |
| pool-data_normal                | objfs     | INFO                | 2016-09-29 15:22:47  | The pool system of file system objfs reached a normal data level.                      |
| pool-metadata_normal            | fs1       | INFO                | 2016-09-29 15:22:47  | The pool system of file system fs1 reached a normal metadata level.                    |
| pool-metadata_normal            | fs1       | INFO                | 2016-09-29 15:22:47  | The pool myPool of file system fs1 reached a normal metadata level.                    |
| pool-metadata_normal            | fs2       | INFO                | 2016-09-29 15:22:47  | The pool system of file system fs2 reached a normal metadata level.                    |
| pool-metadata_normal            | objfs     | INFO                | 2016-09-29 15:22:47  | The pool system of file system objfs reached a normal metadata level.                  |
| pool-metadata_normal            | objfs     | INFO                | 2016-09-29 15:22:47  | The pool data of file system objfs reached a normal metadata level.                    |

16. To check the monitoring interval, issue the following command:

```
mmhealth config interval
```

The system displays output similar to this:

Monitor interval is DEFAULT.

17. To set the monitoring interval to low, issue the following command:

```
mmhealth config interval LOW
```

The system displays output similar to this:

Monitor interval changed to LOW.

## **See also**

- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmobj command” on page 466
- “mmsmb command” on page 554

## **Location**

/usr/lpp/mmfs/bin

---

## mmimgbackup command

Performs a backup of a single GPFS file system metadata image.

### Synopsis

```
mmimgbackup Device [-g GlobalWorkDirectory]
[-L n] [-N {Node[,Node...] | NodeFile | NodeClass}]
[-S SnapshotName] [--image ImageSetName] [--notsm | --tsm]
[--qos QOSClass] [--tsm-server ServerName] [POLICY-OPTIONS]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

### Description

The **mmimgbackup** command performs a backup of a single GPFS file system metadata image.

You must run the **mmbackupconfig** command before you run the **mmimgbackup** command. For more information, see the topic *Scale Out Backup and Restore (SOBAR)* in the *IBM Spectrum Scale: Administration Guide*.

### Parameters

#### Device

The device name of the file system whose metadata image is to be backed up. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

#### -g GlobalWorkDirectory

The directory to be used for temporary files that need to be shared between the **mmimgbackup** worker nodes and to hold backup images until sent to archive. The default is:

*mount\_point\_for\_Device/.mmimgbackup*

#### -L n

Controls the level of information displayed by the **mmimgbackup** command. The default for **mmimgbackup** is 1. Larger values indicate the display of more detailed information. *n* should be one of the following values:

- 0 Displays only serious errors.
- 1 Displays some information as the command executes, but not for each file. This is the default.
- 2 Displays each chosen file and the scheduled action.
- 3 Displays the same information as 2, plus each candidate file and the applicable rule.
- 4 Displays the same information as 3, plus each explicitly EXCLUDED file and the applicable rule.
- 5 Displays the same information as 4, plus the attributes of candidate and EXCLUDED files.
- 6 Displays the same information as 5, plus non-candidate files and their attributes.

#### -N {Node[,Node...] | NodeFile | NodeClass}

Specifies the nodes that will participate in the backup. This command supports all defined node classes. The default is to run only on the node where the **mmimgbackup** command is running or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

**Note:** If more than one node is specified, ensure all nodes have the same operating system or unexpected results may occur.

## **mmimgbackup**

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### **-S SnapshotName**

Archives the files in the specified global snapshot rather than in the active file system. The snapshot specified cannot be a fileset-level snapshot.

### **--image ImageSetName**

Saves the image files using the provided argument as the base set name. Image file names use the following format:

*ImageSetName*\_YYYYMMDD\_*hh.mm.ss.BBB.sbr*

or

*ImageSetName*\_YYYYMMDD\_*hh.mm.ss.idx*

where:

*ImageSetName*

The default *ImageSetName* is **ImageArchive**.

*YYYY*

A four-digit year.

*MM* A two-digit month.

*DD* A two-digit day.

*hh* A two-digit hour.

*mm* A two-digit minute.

*ss* A two-digit second.

*BBB*

A three-digit bucket number.

### **--notsm | --tsm**

Omits (enables) archiving an image fileset to IBM Spectrum Protect through the **dsmc** commands.

### **--qos QoSClass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos** command” on page 206. Specify one of the following QoS classes:

#### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **--tsm-server ServerName**

Specifies the server name to provide to the IBM Spectrum Protect **dsmc** command used to store the image data files in the IBM Spectrum Protect server.

### **POLICY-OPTIONS**

The following **mmapplypolicy** options may also be used with **mmimgbackup**:

**-A** *IscanBuckets*

Specifies the number of buckets of inode numbers (number of inode/filelists) to be created and processed by the parallel inode scan. The default is 17. A bucket will typically represent 1,000,000 in-use inodes.

**-a** *IscanThreads*

Specifies the number of threads and sort pipelines each node will run during parallel inode scan and policy evaluation. The default is 4.

**-D** *yyyy-mm-dd[@hh:mm[:ss]]*

Specifies the date and (UTC) time to be used by the **mmimgbackup** command when evaluating the policy rules. The default is the current date and time. If only a date is specified, the time will default to 00:00:00.

**-M** *name=value*

Indicates a user defined macro specification. There can be more than one **-M** argument. These macro specifications are passed on to the **m4** preprocessor as **-D** specifications.

**-n** *DirThreadLevel*

Specifies the number of threads that will be created and dispatched within each **mmimgbackup** process during the directory scan phase. The default is 24.

**-s** *LocalWorkDirectory*

Specifies the directory to be used for local temporary storage during command processing. The default directory is **/tmp**.

**--sort-buffer-size** *Size*

Specifies the size for the main memory buffer to be used by sort command.

## Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmimgbackup** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

To create a backup image with an *ImageSetName* of **sobar.cluster.fs9**, for the snapshot **snap1** of file system **fs9** where the image is stored in the file system **/backup\_images** on the IBM Spectrum Protect server, issue:

```
mmimgbackup fs9 -S snap1 -g /backup_images --image sober.cluster.fs9
```

To show that the images are stored on the IBM Spectrum Protect server, issue this command:

```
dsmls /backup_images
```

The system displays information similar to:

```
/backup_images/4063536/mmPolicy.4260220.51A8A6BF:
 1088 1088 0 r sober.cluster.fs9_20121129_16.19.55.idx
 4520 4520 0 r sober.cluster.fs9_20121129_16.19.55_000.sbr
```

## **mmimgbackup**

### **See also**

- “[mmapplypolicy command](#)” on page 56
- “[mmbbackupconfig command](#)” on page 81
- “[mmimimgrestore command](#)” on page 371
- “[mmrestoreconfig command](#)” on page 521

### **Location**

/usr/lpp/mmfs/bin

## mmimrestore command

Restores a single GPFS file system from a metadata image.

### Synopsis

```
mmimrestore Device ImagePath [-g GlobalWorkDirectory]
[-L n] [-N {Node[,Node...] | NodeFile | NodeClass}]
[--image ImageSetName] [--qos QoSClass] [POLICY-OPTIONS]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

### Description

The **mmimrestore** command restores a single GPFS file system from a metadata image.

The **mmrestoreconfig** command must be run prior to running the **mmimrestore** command. For more information, see *Scale Out Backup and Restore (SOBAR)* in *IBM Spectrum Scale: Administration Guide*.

### Parameters

#### *Device*

The device name of the file system whose metadata image is to be restored. The file system must be empty and mounted read-only. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

#### *ImagePath*

The fully-qualified path name to an image fileset containing GPFS backup images. The path must be accessible by every node participating in the restore.

#### -g *GlobalWorkDirectory*

The directory to be used for temporary files that need to be shared between the **mmimrestore** worker nodes. If not specified, the default working directory will be the *ImagePath* specified.

#### -L *n*

Controls the level of information displayed by the **mmimrestore** command. The default for **mmimrestore** is 1. Larger values indicate the display of more detailed information. *n* should be one of the following values:

- 0 Displays only serious errors.
- 1 Displays some information as the command executes, but not for each file. This is the default.
- 2 Displays each chosen file and the scheduled action.
- 3 Displays the same information as 2, plus each candidate file and the applicable rule.
- 4 Displays the same information as 3, plus each explicitly EXCLUDED file and the applicable rule.
- 5 Displays the same information as 4, plus the attributes of candidate and EXCLUDED files.
- 6 Displays the same information as 5, plus non-candidate files and their attributes.

#### -N {*Node[,Node...]* | *NodeFile* | *NodeClass*}

Specifies the nodes that will participate in the restore. This command supports all defined node classes. The default is to run only on the node where the **mmimrestore** command is running or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

**Note:** If more than one node is specified, ensure all nodes have the same operating system or unexpected results may occur.

## **mmimimgrestore**

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### **--image ImageSetName**

Specifies the image set name representing the metadata image to be restored. The *ImageSetName* must match the *ImageSetName\_YYYYMMDD\_hh.mm.ss* that was created during the backup with the **mmimimgbackup** command.

### **--qos QoSClass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos** command” on page 206. Specify one of the following QoS classes:

#### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **POLICY-OPTIONS**

The following **mmapplypolicy** options may also be used with **mmimimgrestore**:

#### **-m ThreadLevel**

The number of threads that will be created and dispatched within each image restore process during the policy execution phase of restore. The default is calculated to divide the work of processing all image files being restored evenly among all nodes specified with **-N**. The valid range is 1 to 20.

#### **-s LocalWorkDirectory**

Specifies the directory to be used for local temporary storage during command processing. The default directory is **/tmp**.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmimimgrestore** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

To restore file system **fs9** with data stored in the image with an *ImageSetName* of **sobar.cluster.fs9\_20121129\_16.19.55** and execute the restore only on AIX nodes, issue:

```
mmimimgrestore fs9 /backup_images/406*/mmP* --image sober.cluster.fs9_20121129_16.19.55 -N aixnodes
```

**See also**

- “[mmapplypolicy command](#)” on page 56
- “[mmbackupconfig command](#)” on page 81
- “[mmimgbackup command](#)” on page 367
- “[mmrestoreconfig command](#)” on page 521

**Location**

/usr/lpp/mmfs/bin

## **mmimportfs**

---

### **mmimportfs command**

Imports into the cluster one or more file systems that were created in another GPFS cluster.

#### **Synopsis**

```
mmimportfs {Device | all} -i ImportfsFile [-S ChangeSpecFile]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmimportfs** command, in conjunction with the **mmexportfs** command, can be used to move into the current GPFS cluster one or more file systems that were created in another GPFS cluster. The **mmimportfs** command extracts all relevant file system and disk information from the *ExportFilesysData* file specified with the **-i** parameter. This file must have been created by the **mmexportfs** command.

When **all** is specified in place of a file system name, any disks that are not associated with a file system will be imported as well.

If the file systems being imported were created on nodes that do not belong to the current GPFS cluster, the **mmimportfs** command assumes that all disks have been properly moved, and are online and available to the appropriate nodes in the current cluster.

If any node in the cluster, including the node on which you are running the **mmimportfs** command, does not have access to one or more disks, use the **-S** option to assign NSD servers to those disks.

The **mmimportfs** command attempts to preserve any NSD server assignments that were in effect when the file system was exported.

After the **mmimportfs** command completes, use **mmlsnsd** to display the NSD server names that are assigned to each of the disks in the imported file system. Use **mmchnsd** to change the current NSD server assignments as needed.

After the **mmimportfs** command completes, use **mmlsdisk** to display the failure groups to which each disk belongs. Use **mmchdisk** to make adjustments if necessary.

If you are importing file systems into a cluster that already contains GPFS file systems, it is possible to encounter name conflicts. You must resolve such conflicts before the **mmimportfs** command can succeed. You can use the **mmchfs** command to change the device name and mount point of an existing file system. If there are disk name conflicts, use the **mmcrnsd** command to define new disks and specify unique names (rather than let the command generate names). Then replace the conflicting disks using **mmrpldisk** and remove them from the cluster using **mmdelnsd**.

#### **Results**

Upon successful completion of the **mmimportfs** command, all configuration information pertaining to the file systems being imported is added to configuration data of the current GPFS cluster.

#### **Parameters**

**Device | all**

The device name of the file system to be imported. File system names need not be fully-qualified. *fs0* is as acceptable as */dev/fs0*. Specify **all** to import all GPFS file systems, as well as all disks that do not currently belong to a file system.

If the specified file system device is an IBM Spectrum Scale RAID-based file system, then all affected IBM Spectrum Scale RAID objects will be imported as well. This includes recovery groups, declustered arrays, vdisks, and any other file systems that are based on these objects. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

This must be the first parameter.

**-i ImportfsFile**

The path name of the file containing the file system information. This file must have previously been created with the **mmexportfs** command.

**-S ChangeSpecFile**

The path name of an optional file containing disk stanzas or recovery group stanzas, or both, specifying the changes that are to be made to the file systems during the import step.

Prior to GPFS 3.5, the disk information was specified in the form of disk descriptors defined as:

DiskName:ServerList:

For backward compatibility, the **mmimportfs** command will still accept the traditional disk descriptors, but their use is discouraged.

Disk stanzas have the following format:

**%nsd:**

```
nsd=NsibName
servers=ServerList
usage=DiskUsage
failureGroup=FailureGroup
pool=StoragePool
device=DiskName
```

where:

**nsd=DiskName**

Is the name of a disk from the file system being imported. This clause is mandatory for the **mmimportfs** command.

**servers=ServerList**

Is a comma-separated list of NSD server nodes. You can specify up to eight NSD servers in this list. The defined NSD will preferentially use the first server on the list. If the first server is not available, the NSD will use the next available server on the list.

When specifying server nodes for your NSDs, the output of the **mmlscluster** command lists the host name and IP address combinations recognized by GPFS. The utilization of aliased host names not listed in the **mmlscluster** command output may produce undesired results.

If you do not define a *ServerList*, GPFS assumes that the disk is SAN-attached to all nodes in the cluster. If all nodes in the cluster do not have access to the disk, or if the file system to which the disk belongs is to be accessed by other GPFS clusters, you must specify a *ServerList*.

To remove the NSD server list, do not specify a value for *ServerList* (remove or comment out the **servers=ServerList** clause of the NSD stanza).

**usage=DiskUsage**

Specifies the type of data to be stored on the disk. If this clause is specified, the value must match the type of usage already in effect for the disk; **mmimportfs** cannot be used to change this value.

**failureGroup=FailureGroup**

Identifies the failure group to which the disk belongs. If this clause is specified, the value must match the failure group already in effect for the disk; **mmimportfs** cannot be used to change this value.

## **mmimportfs**

**pool=StoragePool**

Specifies the storage pool to which the disk is to be assigned. If this clause is specified, the value must match the storage pool already in effect for the disk; **mmimportfs** cannot be used to change this value.

**device=DiskName**

The block device name of the underlying disk device. This clause is ignored by the **mmimportfs** command.

Recovery group stanzas have the following format:

**%rg: rgName=RecoveryGroupName  
servers=Primary[,Backup]**

where:

*RecoveryGroupName*

Specifies the name of the recovery group being imported.

**Primary[,Backup]**

Specifies the primary server and, optionally, a backup server to be associated with the recovery group.

### **Notes:**

1. You cannot change the name of a disk. You cannot change the disk usage or failure group assignment with the **mmimportfs** command. Use the **mmchdisk** command for this purpose.
2. All disks that do not have stanzas in *ChangeSpecFile* are assigned the NSD servers that they had at the time the file system was exported. All disks with NSD servers that are not valid are assumed to be SAN-attached to all nodes in the cluster. Use the **mmchnsd** command to assign new or change existing NSD server nodes.
3. Use the **mmchrecoverygroup** command to activate recovery groups that do not have stanzas in *ChangeSpecFile*. The **mmchrecoverygroup** command is documented in *IBM Spectrum Scale RAID: Administration*.

## **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmimportfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

To import all file systems in the current cluster, issue this command:

```
mmimportfs all -i /u/admin/exportfile
```

The output is similar to this:

```
mmimportfs: Processing file system fs1 ...
mmimportfs: Processing disk gpfs2nsd
mmimportfs: Processing disk gpfs3nsd
mmimportfs: Processing disk gpfs4nsd
```

```
mmimportfs: Processing file system fs2 ...
mmimportfs: Processing disk gpfs1nsd1
mmimportfs: Processing disk gpfs5nsd

mmimportfs: Processing disks that do not belong to any file system ...
mmimportfs: Processing disk gpfs6nsd
mmimportfs: Processing disk gpfs1001nsd

mmimportfs: Committing the changes ...

mmimportfs: The following file systems were successfully imported:
 fs1
 fs2
mmimportfs: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

**See also**

- “[mmexportfs command](#)” on page 329

**Location**

/usr/lpp/mmfs/bin

## **mmkeyserv**

---

### **mmkeyserv command**

Manages encryption key servers and clients.

#### **Synopsis**

```
mmkeyserv server {add | update}
 ServerName [--port RestPortNumber] [--user-id RestUserID]
 [--server-pwd PasswordFile] [--accept]
 [--backup ServerName[,ServerName...]] [--distribute | --nodistribute]
 [--timeout ConnectionTimeout] [--retry ConnectionAttempts]
 [--interval Microseconds]
```

or

```
mmkeyserv server delete ServerName
```

or

```
mmkeyserv server show [ServerName] [-Y]
```

or

```
mmkeyserv tenant add TenantName
 --server ServerName [--server-pwd PasswordFile]
```

or

```
mmkeyserv tenant delete TenantName
 --server ServerName [--server-pwd PasswordFile]
```

or

```
mmkeyserv tenant show [TenantName] [--server ServerName] [-Y]
```

or

```
mmkeyserv key create --server ServerName [--server-pwd PasswordFile]
 --tenant TenantName [--count NumberOfKeys]
```

or

```
mmkeyserv key delete --server ServerName [--server-pwd PasswordFile]
 {--all | --tenant TenantName | --file ListOfKeysFile}
```

or

```
mmkeyserv key show --server ServerName [--server-pwd PasswordFile]
 --tenant TenantName
```

or

```
mmkeyserv client create ClientName
 --server ServerName [--server-pwd PasswordFile]
 [--keystore-pwd PasswordFile]
```

or

```
mmkeyserv client delete ClientName
```

or

```
mmkeyserv client register ClientName
 --rkm-id RkmID --tenant TenantName [--server-pwd PasswordFile]
```

or

```
mmkeyserv client deregister ClientName
 --tenant TenantName [--server-pwd PasswordFile]
```

or

```
mmkeyserv client show [ClientName | --server ServerName] [-Y]
```

or

```
mmkeyserv rkm change RkmID {[--rkm-id NewRkmID]
 [--backup ServerName[,ServerName...]] [--distribute | --nodistribute]
 [--timeout ConnectionTimeout] [--retry ConnectionAttempts]
 [--interval Microseconds]}
```

or

```
mmkeyserv rkm show
```

## Availability

Available with IBM Spectrum Scale Advanced Edition or IBM Spectrum Scale Data Management Edition.

## Description

With the **mmkeyserv** command, you can configure a cluster and a remote key manager (RKM) server so that nodes in the cluster can retrieve master encryption keys when they need to. You must set up an RKM server before you run this command. The RKM server software must be IBM Security Key Lifecycle Manager (SKLM). Nodes in the cluster must have direct network access to the RKM server.

With this command you can connect to an RKM server, create GPFS tenants, create encryption keys, and create and register key clients. The command automatically generates and exchanges certificates and sets up a local keystore. You can also use this command to securely delete key clients, encryption keys, and tenants. You can run this command from any node in the cluster. Each node has a configuration file and a copy of the local keystore. Configuration changes affect all nodes in the cluster.

**Password files:** Several of the command options require a password file as a parameter. A password file is a text file that contains a password at the beginning. A password must be 1-20 characters in length. Because the password file is a security-sensitive file, it must have the following characteristics:

- It must be a regular file.
- It must be owned by the root user.
- Only the root user must have permission to read or write it.

The following terms are used:

### IBM Security Key Lifecycle Manager (SKLM)

Required key management server software.

### Master encryption key (MEK)

A key for encrypting file encryption keys.

### Remote key management (RKM) server

A server with software that authenticates clients and provides them with master encryption keys.

### RKM.conf file

A configuration file that the mmkeyserv command maintains. Each node in the cluster has a copy of this file. The full path is /var/mmfs/ss1/keyServ/RKM.conf.

### RKM stanza

A block of configuration information that describes a registered key client. RKM stanzas are stored in the RKM.conf file.

### RKM ID

An identifier for an RKM stanza.

**tenant** A device group in SKLM that contains MEKs for registered key clients.

## **mmkeyserv**

### **Parameters**

#### **server**

Manages a connection with an RKM server.

##### **add**

Adds an RKM server connection to the IBM Spectrum Scale cluster. You can adjust the values of some of these parameters later with the **mmkeyserv rkm change** command.

###### *ServerName*

Specifies the host name or IP address of the RKM server.

###### **--port RestPortNumber**

Specifies the port number for the Representational State Transfer (REST) interface. The default value is 9080. For SKLM v2.7 and later, the default port number is 443.

###### **--user-id RestUserID**

Specifies the user ID for the RKM server. The default value is SKLMAadmin.

###### **--server-pwd PasswordFile**

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

###### **--accept**

Configures the command to automatically accept certificates from the RKM server. The acceptance prompt is suppressed.

###### **--backup ServerName [,ServerName...]**

Specifies a comma-separated list of server names that you want to add to the list of backup RKM servers in the RKM.conf file. If an IBM Spectrum Scale node cannot retrieve a master encryption key from its main RKM server, it tries each backup server in the list until it either retrieves a key or exhausts the list.

**Note:** The **mmkeyserv** command itself does not attempt to contact backup servers or to replicate client information across background servers. The system administrator is responsible for maintaining replication across backup servers.

###### **--distribute | --nodistribute**

###### **--distribute**

Attempts to arrange the list of RKM server names (main RKM server and backup RKM servers) in the RKM.conf file in a different order on each node, so that each node connects with the servers in a different order. This option provides some performance advantage in retrieving MEKs. This option is the default.

###### **--nodistribute**

Does not attempt to arrange the list of backup RKM server names in the RKM.conf file.

###### **--timeout ConnectionTimeout**

Sets the connection timeout in seconds for retrieving an MEK from an RKM server. The valid range is 1 - 120 seconds. The default value is 60 seconds.

###### **--retry ConnectionAttempts**

Sets the number of attempts to retry a connection to an RKM server. The valid range is 1 - 10 retries. The default value is 3 retries.

###### **--interval Microseconds**

Specifies the number of microseconds to wait between connection retries. The valid range is 1 - 1000000000. The default value is 10000 (0.1 seconds).

**update**

Updates a connection between an IBM Spectrum Scale cluster and an RKM server. If you do not specify a parameter after **update**, the command does the following actions:

- It gets a fresh certificate from the RKM server.
- It updates the default port number for the Representational State Transfer (REST) interface to match the default on the key server. The default REST port number is 443 for SKLM v2.7. For earlier versions it is 9080.

*ServerName*

Specifies the host name or IP address of the RKM server.

**--port RestPortNumber**

Specifies the port number for the Representational State Transfer (REST) interface.

**--user-id RestUserID**

Specifies the user ID for the RKM server. The default value is SKLMAdmin.

**--server-pwd PasswordFile**

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password if one is required.

**--accept**

Configures the command to automatically accept certificates from the RKM server. The acceptance prompt is suppressed.

**--backup ServerName [,ServerName...]**

Specifies a comma-separated list of server names that you want to add to the backup RKM servers that are listed in the RKM.conf file. If an IBM Spectrum Scale node cannot retrieve a master encryption key from its main RKM server, it tries each backup server in the list until it either retrieves a key or exhausts the list.

**Important:** To remove a backup list, specify **delete** instead of a list of server names, as in the following example:

```
mmkeyserv server update ServerName --backup delete
```

The **backup** parameter does not change the RKM.conf file.

**Note:** The **mmkeyserv** command itself does not attempt to contact backup servers or to replicate client information across background servers. The system administrator is responsible for maintaining replication across backup servers.

**--distribute | --nodistribute**

You must specify one of these two options. There is no default value.

**--distribute**

Attempts to arrange the RKM server names (main RKM server and backup RKM servers) that are listed in the RKM.conf file in a different order on each node, so that each node connects with the servers in a different order. This option provides some performance advantage in retrieving MEKs.

**--nodistribute**

Does not attempt to arrange the backup RKM server names that are listed in the RKM.conf file.

Neither parameter changes the RKM.conf file.

**--timeout ConnectionTimeout**

Sets the connection timeout in seconds for retrieving an MEK from an RKM server. The valid range is 1 - 120 seconds. To restore the system default value (60 seconds), you can specify either 60 or the keyword **default** as the *ConnectionTimeout* value.

## **mmkeyserv**

| This parameter does not change the RKM.conf file.

| **--retry ConnectionAttempts**  
| Sets the number of attempts to retry a connection to an RKM server. The valid range is 1 - 10  
| retries. To restore the system default value (3 retries), you can specify either 3 or the keyword  
| **default** as the *ConnectionAttempts* value.

| This parameter does not change the RKM.conf file.

| **--interval Microseconds**  
| Specifies the number of microseconds to wait between connection retries. The valid range is 1  
| - 1000000000. To restore the system default value (10000, which is 0.1 seconds), you can  
| specify either 10000 or the keyword **default** as the *Microseconds* value.

| This parameter does not change the RKM.conf file.

### **delete**

Removes a connection with an RKM server from the cluster.

*ServerName*

Specifies the host name or IP address of the RKM server that you want to disconnect from.

### **show**

Displays information about RKM servers. The following table shows the display options:

*Table 14. mmkeyserv server show*

| Option                 | Displays information about |
|------------------------|----------------------------|
| <b>show</b>            | All servers.               |
| <b>show ServerName</b> | The specified server.      |

*ServerName*

Specifies the host name or IP address of an RKM server.

## **tenant**

Manages tenants on RKM servers. A tenant is an SKLM device group for holding encryption keys.

### **add**

Specifies the name of a tenant to add to the IBM Spectrum Scale cluster.

- If the tenant is already added to the cluster, the command returns with an error.
- If the tenant exists on the RKM server but is not added to the cluster, the command adds the tenant to the cluster.
- If the tenant does not exist on the RKM server, the command creates the tenant on the server and adds the tenant to the cluster.

*TenantName*

Specifies the name of the tenant that you want to create.

**--server ServerName**

Specifies the name of the RKM server to which the tenant belongs.

**--server-pwd PasswordFile**

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

### **delete**

Deletes a tenant from an RKM server.

**Note:**

- If you delete a tenant that has encryption keys on the key server, the command deletes the tenant from the cluster configuration but not from the key server.
- If you delete a tenant that has no encryption keys on the key server, the command deletes the tenant from both the cluster configuration and the key server.

**TenantName**

Specifies the name of the tenant that you want to delete.

**--server ServerName**

Specifies the name of the RKM server to which the tenant belongs.

**--server-pwd PasswordFile**

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

**show**

Displays information about tenants and RKM servers. The following table shows the results of various combinations of options:

*Table 15. mmkeyserv tenant show*

| Option                                     | Displays information about                     |
|--------------------------------------------|------------------------------------------------|
| <b>show</b>                                | All tenants from all RKM servers.              |
| <b>show TenantName</b>                     | The specified tenant.                          |
| <b>show --server ServerName</b>            | All tenants from the specified RKM server.     |
| <b>show TenantName --server ServerName</b> | The specified tenant and the specified server. |

**TenantName**

Specifies the name of a tenant.

**--server ServerName**

Specifies the name of an RKM server.

**key**

Manages encryption keys.

**create**

Creates encryption keys in a tenant and displays the key IDs on the console.

**Note:** Make a note of the key IDs. You must specify an encryption key ID and an RKM ID when you write an encryption policy rule.

**--server ServerName**

Specifies the host name or IP address of an RKM server.

**--server-pwd PasswordFile**

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

**--tenant TenantName**

Specifies the name of the tenant in which you want to create the encryption keys.

**--count NumberOfKeys**

Specifies the number of keys to create. The default value is 1.

**delete**

Deletes encryption keys from a tenant.

## **mmkeyserv**

### **CAUTION:**

**When you delete an encryption key, any data that was encrypted by that key becomes unrecoverable.**

#### **--server *ServerName***

Specifies the host name or IP address of an RKM server.

#### **--server-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

#### **--all --tenant *TenantName***

Deletes all the encryption keys in the specified tenant.

#### **--file *ListOfKeysFile***

Specifies a file that contains a list of the key IDs of encryption keys that you want to delete, one key per line.

#### **show**

Displays information about the encryption keys in a tenant.

#### **--server *ServerName***

Specifies the host name or IP address of an RKM server.

#### **--server-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

#### **--tenant *TenantName***

Specifies the name of the tenant that contains the keys that you want to display.

## **client**

Manages key clients. The following facts are important:

- You need only one key client per cluster per RKM server. However, you can create and use multiple key clients on the same RKM server.
- You can register only one key client per tenant per cluster. However, you can register one key client to more than one tenant in the same RKM server.

#### **create**

Creates a key client to communicate with the RKM server.

#### ***ClientName***

Specifies the name of the key client that you want to create. A key client name must be 1-16 characters in length and must be unique within an IBM Spectrum Scale cluster.

#### **--server *ServerName***

Specifies the name of the RKM server to which the key client belongs.

#### **--server-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

#### **--keystore-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

#### **delete**

Deletes a key client.

***ClientName***

Specifies the name of the key client that you want to delete.

**register**

Registers a key client to a tenant.

***ClientName***

Specifies the name of the key client that you want to register.

**--rkm-id *RkmID***

Specifies a new RKM ID. An RKM ID must be unique within the cluster, must be 1-21 characters in length, and can contain only alphanumeric characters or underscore (\_). It must begin with a letter or an underscore. An RKM ID identifies an RKM stanza in the RKM.conf file. The stanza contains the information that a node needs to retrieve a master encryption key (MEK) from an RKM.

**--tenant *TenantName***

Specifies the name of the tenant to which you want to register the client.

**--server-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

**deregister**

Unregisters a key client from a tenant.

***ClientName***

Specifies the name of the key client that you want to unregister.

**--tenant *TenantName***

Specifies the name of the tenant that you want to unregister the key client from.

**--server-pwd *PasswordFile***

Specifies a password file that contains a password for accessing the RKM server. See the requirements for password files in the Description section of this topic. If this parameter is omitted, the command prompts for a password.

**show**

Displays information about a key client.

***ClientName***

Specifies the name of the client whose information you want to display.

**--server *ServerName***

Specifies the name of the server to which the client belongs.

**rkm****change**

Changes the properties of an RKM stanza. For more information about an RKM stanza, see the Description section of this topic.

***RkmID***

Specifies the RKM ID of the stanza whose properties you want to change.

**--rkm-id *RkmID***

Specifies a new RKM ID. An RKM ID must be unique within the cluster, must be 1-21 characters in length, and can contain only alphanumeric characters or underscore (\_). It must begin with a letter or an underscore. An RKM ID identifies an RKM stanza in the RKM.conf file. The stanza contains the information that a node needs to retrieve a master encryption key (MEK) from an RKM.

## **mmkeyserv**

### **--backup ServerName[,ServerName...]**

Specifies a comma-separated list of server names that you want to add to the backup RKM servers that are listed in the RKM.conf file. If an IBM Spectrum Scale node cannot retrieve a master encryption key from its main RKM server, it tries each backup server in the list until it either retrieves a key or exhausts the list.

**Important:** To remove a backup list, specify **delete** instead of a list of server names, as in the following example:

```
mmkeyserv rkm change RkmID --backup delete
```

The **backup** parameter does not change the RKM.conf file.

**Note:** The **mmkeyserv** command itself does not attempt to contact backup servers or to replicate client information across background servers. The system administrator is responsible for maintaining replication across backup servers.

### **--distribute | --nodistribute**

#### **--distribute**

Attempts to arrange the list of RKM server names (main RKM server and backup RKM servers) in the RKM.conf file in a different order on each node so that each node connects with the servers in a different order. This option provides some performance advantage in retrieving MEKs.

#### **--nodistribute**

Does not attempt to arrange the list of backup RKM server names in the RKM.conf file.

### **--timeout ConnectionTimeout**

Sets the connection timeout in seconds for retrieving an MEK from an RKM server. The valid range is 1 - 120 seconds. To restore the system default value (60 seconds), you can specify either 60 or the keyword **default** as the *ConnectionTimeout* value.

This parameter does not change the RKM.conf file.

### **--retry ConnectionAttempts**

Sets the number of attempts to retry a connection to an RKM server. The valid range is 1 - 10 retries. To restore the system default value (3 retries), you can specify either 3 or the keyword **default** as the *ConnectionAttempts* value.

This parameter does not change the RKM.conf file.

### **--interval Microseconds**

Specifies the number of microseconds to wait between connection retries. The valid range is 1 - 1000000000. To restore the system default value (10000, which is 0.1 seconds), you can specify either 10000 or the keyword **default** as the *Microseconds* value.

This parameter does not change the RKM.conf file.

### **show**

Displays information about all the RKM stanzas in the RKM.conf file of the node.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

## **Exit status**

**0** Successful completion.

**Nonzero**

A failure occurred.

**Security**

You must have root authority to run the **mmkeyserv** command.

The node on which you enter the command must be able to execute remote shell commands on any other administration node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

**Examples**

Examples 1-5 illustrate the steps in configuring an RKM server and a key client and generating an encryption key:

1. The following command makes an RKM server known to an IBM Spectrum Scale cluster. The name `keyserver01` is the host name of the SKLM server:

```
mmkeyserv server add keyserver01
Enter password for the RKM server keyserver01:
The security certificate(s) from keyserver01.gpfs.net must be accepted to continue. View the
certificate(s) to determine whether you want to trust the certifying authority.
Do you want to view or trust the certificate(s)? (view/yes/no) view

Serial number: 01022a8adf20f3
SHA-256 digest: 2ca4a48a3038f37d430162be8827d91eb584e98f5b3809047ef4a1c72e15fc4c
Signature: 7f0312e7be18efd72c9d8f37dbb832724859ba4bb5827c230e2161473e0753b367ed49d9935
05bd23858541475de8e021e0930725abbd3d25b71edc8fc3de20b7c2db5cd4e865f41c7c410c1d710acf222e1c45189108e
40568ddcb21094264da60a1d96711015a7951eb2655363309d790ab44ee7b26adf8385e2c210b8268c5aede5f82f26855
4a6fc22ece6feeee2a6264706e71416a0dbe8c39ceacd86054d7cc34dda4fffea4605c037d32129055610821af85dd9819a
4d7e4baa70c51addcda720d33bc9f8bbde6d292c028b2f525a0275ebea968c26f8f0c4b604719ae3b04e71ed7a8188cd6ad
f68764374b29c91df3d101a941bf8b7189485ad7
Signature algorithm: SHA256WithRSASignature
Key size: 2048
Issuer: C=US, O=IBM, OU=SKLMNode, SKLMCell, Root Certificate, CN=c40bbc1xn3.gpfs.net
Subject: C=US, O=IBM, OU=SKLMNode, SKLMCell, CN=c40bbc1xn3.gpfs.net

Serial number: 01022a24475466
SHA-256 digest: 077c3b53c5046aa893b760c11cca3a993efbc729479771e03791f9ed4f716879
Signature: 227b5bef89f2e55ef628da6b50db1ab842095a54e1505655e3d95fee753a7f7554868aa79b
294c503dc34562cf69c2a20128796758838968565c0812c4aedbb0543d396646a269c02bf4c5ce5acba4409a10effbd47ca
38ce492698e2dcdc8390b9ae3f4a47c23ee3045ff0145218668f35a63edac68201789ed0db6e5c170f5c6db49769f0b4c9a
5f208746e4342294c447793ed087fa0ac762588faf420febeb3fc411e4e725bd46476e1f9f44759a696573af5dbbc95532
18c083c80440f2e542bf56cc5cc18156cce05efd6c2e5fea2b886c5c1e262c10af18b13ccf38c3533ba025b97bbe62f271
545b2ab5c1f50c1dca45ce504dfcfc257362e9b43
Signature algorithm: SHA256WithRSASignature
Key size: 2048
Issuer: C=US, O=IBM, OU=SKLMNode, SKLMCell, Root Certificate, CN=c40bbc1xn3.gpfs.net
Subject: C=US, O=IBM, OU=SKLMNode, SKLMCell, Root Certificate, CN=c40bbc1xn3.gpfs.net
```

Do you trust the certificate(s) above? (yes/no) yes

To display all RKM servers information, enter:

c34f2n03:~ # mmkeyserv server show

keyserver01

|                     |                      |
|---------------------|----------------------|
| Type:               | SKLM                 |
| Hostname:           | keyserver01.gpfs.net |
| User ID:            | SKLMAdmin            |
| REST port:          | 9080                 |
| Label:              | 1_keyserver01        |
| NIST:               | on                   |
| FIPS1402:           | off                  |
| Backup Key Servers: |                      |

## mmkeyserv

```
Distribute: yes
Retrieval Timeout: 120
Retrieval Retry: 3
Retrieval Interval: 10000
```

The following command displays information about all RKM servers that are known to the cluster. At the moment, the only one is keyserver01:

```
mmkeyserv server show
keyserver01
 Type: SKLM
 Hostname: keyserver01.gpfs.net
 User ID: SKLMAadmin
 REST port: 9080
 Label: 1_keyserver01
 NIST: on
 FIPS1402: off
 Backup Key Servers:
 Distribute: yes
 Retrieval Timeout: 120
 Retrieval Retry: 3
 Retrieval Interval: 10000
```

2. The following command creates a tenant in the server that you defined in Example 1, keyserver01. The name of the tenant is devG1:

```
mmkeyserv tenant add devG1 --server keyserver01
Enter password for the RKM server keyserver01:
```

The following command displays all the current tenants of keyserver01:

```
mmkeyserv tenant show
devG1
 Key Server: keyserver01.gpfs.net
 Registered Client: (none)
```

3. The following command adds a key client to the tenant that you created in Example 2. The command does not specify password files for the server and the new keystore, so the command prompts for the passwords. The name of the key client is c34f2n03Client1:

```
mmkeyserv client create c34f2n03Client1 --server keyserver01
Enter password for the RKM server keyserver01:
Create a pass phrase for keystore:
Confirm your pass phrase:
```

The following command displays information about all the key clients on the RKM server:

```
mmkeyserv client show
c34f2n03Client1
 Label: c34f2n03Client1
 Key Server: keyserver01.gpfs.net
 Tenants: (none)
```

4. The following command registers the key client from Example 3 to the tenant from Example 2. To ensure uniqueness in RKM IDs, it is a good practice to create the RKM ID name by combining the names of the RKM server and the tenant. However, the RKM ID cannot be longer than 21 characters. In this example the RKM ID is keyserver01\_devG1:

```
mmkeyserv client register c34f2n03Client1 --tenant devG1 --rkm-id keyserver01_devG1
Enter password for the RKM server :
```

```
mmkeyserv: [I] Client currently does not have access to the key. Continue the registration process ...
mmkeyserv: Successfully accepted client certificate
```

The following two commands now show that key client **c34f2n03Client1** is registered to tenant **devG1**:

```
mmkeyserv tenant show
devG1
 Key Server: keyserver01.gpfs.net
```

```
Registered Client: c34f2n03Client1

mmkeyserv client show
c34f2n03Client1
 Label: c34f2n03Client1
 Key Server: keyserver01.gpfs.net
 Tenants: devG1
```

The following command shows the contents of the new RKM stanza that was added to the RKM.conf file:

```
mmkeyserv rkm show
keyserver01_devG1 {
 type = SKLM
 kmipServerUri = tls://192.168.40.59:5696
 keyStore = /var/mmfs/ssl/keyServ/serverKmip.1_keyserver01.c34f2n03Client1.1.p12
 passphrase = pw4c34f2n03Client1
 clientCertLabel = c34f2n03Client1
 tenantName = devG1
}
```

You can also show the contents of the RKM.conf file by routing the contents of the file to the console:

```
cat /var/mmfs/ssl/keyServ/RKM.conf
keyserver01_devG1 {
 type = SKLM
 kmipServerUri = tls://192.168.40.59:5696
 keyStore = /var/mmfs/ssl/keyServ/serverKmip.1_keyserver01.c34f2n03Client1.1.p12
 passphrase = pw4c34f2n03Client1
 clientCertLabel = c34f2n03Client1
 tenantName = devG1
}
```

5. The following example creates an encryption key in the tenant from Example 4. In the third line, the command displays the new encryption key (KEY-d4e83148-e827-4f54-8e5b-5e1b5cc66de1):

```
mmkeyserv key create --server keyserver01.gpfs.net --tenant devG1
Enter password for the RKM server keyserver01.gpfs.net:
KEY-d4e83148-e827-4f54-8e5b-5e1b5cc66de1
```

## See also

### Location

/usr/lpp/mmfs/bin

## **mmlinkfileset**

---

### **mmlinkfileset command**

Creates a junction that references the root directory of a GPFS fileset.

#### **Synopsis**

```
mmlinkfileset Device FilesetName [-J JunctionPath]
```

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

The **mmlinkfileset** command creates a junction at *JunctionPath* that references the root directory of *FilesetName*. The junction is a special directory entry, much like a POSIX hard link, that connects a name in a directory of one fileset, the parent, to the root directory of a child fileset. From the user's viewpoint, a junction always appears as if it were a directory, but the user is not allowed to issue the **unlink** or **rmdir** commands on a junction. Instead, the **mmunlinkfileset** command must be used to remove a junction.

If *JunctionPath* is not specified, the junction is created in the current directory with the name *FilesetName*. The user may use the **mv** command on the directory to move to a new location in the parent fileset, but the **mv** command is not allowed to move the junction to a different fileset.

For information on GPFS filesets, see the *IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system that contains the fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

##### *FilesetName*

Specifies the name of the fileset to be linked. It must not already be linked into the namespace.

There are no restrictions on linking independent filesets, but a dependent fileset can only be linked inside its own inode space.

##### **-J** *JunctionPath*

Specifies the name of the junction. The name must not refer to an existing file system object.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmlinkfileset** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

This command links fileset **fset1** in file system **gpfs1** to junction path **/gpfs1/fset1**:

```
mmlinkfileset gpfs1 fset1 -J /gpfs1/fset1
```

The system displays output similar to:

```
Fileset 'fset1' linked at '/gpfs1/fset1'.
```

To confirm the change, issue this command:

```
mmlsfileset gpfs1
```

The system displays output similar to:

```
Filesets in file system 'gpfs1':
Name Status Path
root Linked /gpfs1
fset1 Linked /gpfs1/fset1
```

## See also

- “mmchfileset command” on page 173
- “mmcrlfileset command” on page 244
- “mmdelfileset command” on page 292
- “mmlsfileset command” on page 406
- “mmunlinkfileset command” on page 578

## Location

/usr/lpp/mmfss/bin

## **mmlsattr**

---

### **mmlsattr command**

Queries file attributes.

#### **Synopsis**

```
mmlsattr [-L] [-1]
[-d | --dump-attr]
[-n AttributeName | --get-attr AttributeName]
[-X | --hex-attr] [--hex-attr-name]
Filename [Filename...]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmlsattr** command to display attributes of a file.

#### **Results**

For the specified file, the **mmlsattr** command lists:

- The current number of copies of data for a file and the maximum value
- The number of copies of the metadata for a file and the maximum value
- Whether the Direct I/O caching policy is in effect for a file

#### **Parameters**

##### *Filename*

The name of the file to be queried. You must enter at least one file name; if you specify more than one, delimit each file name by a space. Wildcard characters are supported in file names; for example, **project\*.sched**.

- 1 Specifies that this command works only with regular files and directories and does not follow symlinks. The default is to follow symlinks.
- L Displays additional file attributes:
  - The assigned storage pool name of the file.
  - The name of the fileset that includes the file.
  - If a file is a snapshot file, the name of the snapshot that includes the file is shown. If the file is a regular file, an empty string is displayed.
  - Whether the file is exposed, ill replicated, ill placed, or unbalanced (displayed under the **flags** heading).
  - Whether the file is immutable.
  - Whether the file is in **appendOnly** mode.
  - The creation time of the file.
  - If the **compact** attribute is nonzero, then this parameter displays the number of directory slots that were set by the **mmchattr** command with the **--compact** option or by the **gpfs\_prealloc** subroutine. For more information, see the topics “**mmchattr command**” on page 122 and “**gpfs\_prealloc()** subroutine” on page 781.

-L can be combined with -d | --dump-attr to display all extended attribute names and values for each file.

- d | --dump-attr

Displays the names of all extended attributes for each file.

- n *AttributeName* | --get-attr *AttributeName***  
Displays the name and value of the specified extended attribute for each file.
- X | --hex-attr**  
Displays the attribute value in hex.
- hex-attr-name**  
Displays the attribute name in hex.

## Exit status

0 Successful completion.

### nonzero

A failure has occurred. The return code equals the number of files from which the command was not able to get attribute information.

## Security

You must have read access to run the **mmlsattr** command.

You may issue the **mmlsattr** command only from a node in the GPFS cluster where the file system is mounted.

## Examples

- To list the attributes of a file, issue this command:

```
mmlsattr -L newfile
```

The system displays information similar to:

```
file name: newfile
metadata replication: 1 max 2
data replication: 1 max 2
immutable: no
appendOnly: no
flags: directio
storage pool name: system
fileset name: root
snapshot name:
creation Time: Wed Feb 22 15:16:29 2012
Misc attributes: ARCHIVE
```

- To show the attributes for all files in the root directory of file system **fs0**, issue this command:

```
mmlsattr /fs0/*
```

The system displays information similar to:

```
replication factors
metadata(max) data(max) file [flags]

1 (1) 1 (1) /fs0/project4.sched
1 (1) 1 (1) /fs0/project4.hist
1 (1) 1 (1) /fs0/project5.plan
```

- To show all extended attribute names and values for the file **/ba1/newfile**, issue this command:

```
mmlsattr -d -L /ba1/newfile
```

The system displays information similar to:

```
file name: /ba1/newfile
metadata replication: 1 max 2
data replication: 1 max 2
immutable: no
appendOnly: no
flags: directio
```

## **mmIsattr**

```
| storage pool name: system
| fileset name: root
| snapshot name:
| creation time: Fri Jun 9 15:58:07 2017
| Misc attributes: ARCHIVE
| Encrypted: no
| user.attr1: "value1"
| user.attr: "val1"
| gpfs.DIRECTIO: "1"
| user.ea1: "value1"
```

### **See also**

- “[mmchattr command](#)” on page 122

### **Location**

/usr/lpp/mmfs/bin

---

## mmlscallback command

Lists callbacks that are currently registered in the GPFS system.

### Synopsis

| **mmlscallback [-Y] [CallbackIdentifier[,CallbackIdentifier...]] | user | system | all**

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlscallback** command to list some or all of the callbacks that are currently registered in the GPFS system.

### Parameters

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

*CallbackIdentifier*

Indicates the callback for which information is displayed.

**user**

Indicates all user-defined callbacks. This is the default.

**system**

Indicates all system-defined callbacks.

**all**

Indicates all callbacks currently registered with the system.

### Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

### Security

You must have root authority to run the **mmlscallback** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### Examples

To list all of the callbacks that are currently in the GPFS system, issue this command:

**mmlscallback**

The system displays information similar to:

## **mmiscallback**

```
test1 command = /tmp/myScript
 event = startup

test2 command = /tmp/myScript2
 event = shutdown
 parms = %upNodes
```

To list a specific callback (for example, **test2**) that is currently in the GPFS system, issue this command:

```
mmiscallback test2
```

The system displays information similar to:

```
test2 command = /tmp/myScript2
 event = shutdown
 parms = %upNodes
```

### **See also**

- “[mmaddcallback command](#)” on page 10
- “[mmdelcallback command](#)” on page 286

### **Location**

/usr/lpp/mmfs/bin

---

## **mmIscluster command**

Displays the current configuration information for a GPFS cluster.

### **Synopsis**

| **mmIscluster [-Y] [--ces] [--cnfs] [--cloud-gateway]**

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

Use the **mmIscluster** command to display the current configuration information for an IBM Spectrum Scale cluster.

For the IBM Spectrum Scale cluster, the **mmIscluster** command displays:

- The cluster name
- The cluster ID
- The UID domain
- The remote shell command being used
- The remote file copy command being used
- The repository type (CCR or server-based)
- The primary cluster configuration server (if server-based repository)
- The secondary cluster configuration server (if server-based repository)
- A list of nodes belonging to the IBM Spectrum Scale cluster

For each node, the command displays:

- The node number assigned to the node by IBM Spectrum Scale
- GPFS daemon node interface name
- Primary network IP address
- IBM Spectrum Scale administration node interface name
- Designation, such as whether the node is any of the following:
  - quorum node - A node in the cluster that is counted to determine if a quorum exists. Members of a cluster use the quorum node to determine if it is safe to continue I/O operations when a communications failure occurs.
  - manager node - The file system node that provides the file system manager services to all of the nodes using the file system, including: file system configuration, disk space allocation, token management, and quota management.
  - snmp\_collector node - The designated SNMP collector node for the cluster. The GPFS SNMP subagent runs on the designated SNMP collector node. For additional information see *GPFS SNMP support in IBM Spectrum Scale: Problem Determination Guide*.
  - gateway node - Ensures primary and Disaster Recovery cluster communication during failover.
  - perfmon node - Performance monitoring nodes collect metrics and performance information and sends the information to one or more performance collection nodes.

### **Parameters**

- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

## **mmlscluster**

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### **--ces**

Displays information about protocol nodes.

### **--cnfs**

Displays information about clustered NFS.

### **--cloud-gateway**

Displays information about Transparent cloud tiering nodes.

## **Exit status**

0      Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmlscluster** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To display the current configuration information for the GPFS cluster, issue this command:

```
mmlscluster
```

The system displays information similar to:

```
GPFS cluster information
=====
GPFS cluster name: cluster1.kgn.ibm.com
GPFS cluster id: 680681562214606028
GPFS UID domain: cluster1.kgn.ibm.com
Remote shell command: /usr/bin/ssh
Remote file copy command: /usr/bin/scp
Repository type: CCR

Node Daemon node name IP address Admin node name Designation

1 k164n04.kgn.ibm.com 89.116.68.68 k164n04.kgn.ibm.com quorum
2 k164n05.kgn.ibm.com 89.116.68.69 k164n05.kgn.ibm.com quorum
3 k164n06.kgn.ibm.com 89.116.68.70 k164n06.kgn.ibm.com quorum-manager
4 k164n07.kgn.ibm.com 89.116.68.71 k164n07.kgn.ibm.com quorum-manager-perfmon
5 k164n08.kgn.ibm.com 89.116.68.72 k164n08.kgn.ibm.com quorum-perfmon
6 k164n09.kgn.ibm.com 89.116.68.73 k164n09.kgn.ibm.com quorum-perfmon
7 k164n10.kgn.ibm.com 89.116.68.74 k164n10.kgn.ibm.com manager-perfmon
```

2. To display the configuration information about the Transparent cloud tiering nodes, issue this command:

```
mmlscluster --cloud-gateway
```

The system displays output similar to this:

```
GPFS cluster information
=====
GPFS cluster name: c350f1u1b11
GPFS cluster id: 9364209917238477017
```

| Node | Daemon node name          | Cloud node type |
|------|---------------------------|-----------------|
| 1    | c350f1u1b11               | Cloud-Gateway   |
| 2    | c350f8u17.pk.labs.ibm.com | Cloud-Gateway   |
| 3    | c350f8u18.pk.labs.ibm.com | Cloud-Gateway   |

## See also

- “mmaddnode command” on page 29
- “mmchcluster command” on page 128
- “mmcrccluster command” on page 239
- “mmdelnnode command” on page 297

## Location

/usr/lpp/mmf/bin

## **mmlsconfig**

---

### **mmlsconfig command**

Displays the current configuration data for a GPFS cluster.

#### **Synopsis**

**mmlsconfig** [*Attribute*[,*Attribute*...]] [-Y]

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmlsconfig** command to display the requested configuration attributes for a GPFS cluster. If no specific attributes are requested, the command displays all values that were set explicitly by the user. Depending on your configuration, additional information that is set by GPFS might be displayed. If a configuration attribute is not shown in the output of this command, the default value for that attribute, as documented in the **mmchconfig** command, is in effect.

#### **Parameters**

##### *Attribute*

Specifies the name of the attribute to be displayed. If an attribute has unique values that apply to only a subset of the nodes, the values are followed by the list of affected node names. You can specify more than one attribute in a comma-separated list.

**Note:** See the **mmchconfig** command for a list of supported attributes.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

#### **Exit status**

0      Successful completion.

##### **nonzero**

A failure occurred.

#### **Security**

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

#### **Examples**

1. To display the current configuration data for the GPFS cluster that you are running on, issue this command:

```
mmlsconfig
```

The system displays information similar to the following example:

```
Configuration data for cluster small.cluster:

myNodeConfigNumber 1
clusterName small.cluster
clusterId 6339012640885012929
autoload yes
minReleaseLevel 4.2.0.0
dmapiFileHandleSize 32
[c6f1c3vp3]
pagepool 512M
[common]
adminMode central
```

File systems in cluster small.cluster:

```

/dev/fs1
/dev/gpfs1
```

2. To display the current values for the **maxblocksize** and **pagepool** attributes, issue the following command:

```
mmlsconfig maxblocksize,pagepool
```

The system displays information similar to the following example:

```
|
| maxblocksize 1M
| pagepool 1G
| pagepool 512M [c6f1c3vp3]
```

3. To display the current value for the **cipherList** attribute, issue this command:

```
mmlsconfig cipherList
```

The system displays information similar to the following example:

```
cipherList AUTHONLY
```

## See also

- “mmchcluster command” on page 128
- “mmchconfig command” on page 132
- “mmcrcluster command” on page 239

## Location

/usr/lpp/mmfs/bin

## mmlsdisk

---

### mmlsdisk command

Displays the current configuration and state of the disks in a file system.

#### Synopsis

```
mmlsdisk Device [-d "DiskName[;DiskName...]" [-e | -Y] [-L]
```

or

```
| mmlsdisk Device [-d "DiskName[;DiskName...]" {-m | -M} [-Y]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

Use the **mmlsdisk** command to display the current state of the disks in the file system.

The **mmlsdisk** command may be run against a mounted or unmounted file system.

For each disk in the list, the **mmlsdisk** command displays the following:

- disk name
- driver type
- logical sector size (under the heading “sector size”)
- failure group
- whether it holds metadata
- whether it holds data
- status:

**ready** Normal status.

**suspended**

or

**to be emptied**

Indicates that data is to be migrated off this disk.

**being emptied**

Transitional status in effect while a disk deletion is pending.

**emptied**

Indicates that data is already migrated off this disk.

**replacing**

Transitional status in effect for old disk while replacement is pending.

**replacement**

Transitional status in effect for new disk while replacement is pending.

- availability:

**up** The disk is available to GPFS for normal **read** and **write** operations.

**down** No **read** and **write** operations can be performed on this disk.

**recovering**

An intermediate state for disks coming up, during which GPFS verifies and corrects data. **write** operations can be performed while a disk is in this state, but **read** operations cannot (because data on the disk being recovered might be stale until the **mmchdisk start** command completes).

**unrecovered**

The disk was not successfully brought up.

- disk ID
- storage pool to which the disk is assigned

**Parameters***Device*

The device name of the file system to which the disks belong. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

**-d "DiskName[;DiskName...]"**

The name of the disks for which you want to display current configuration and state information. When you enter multiple values for *DiskName*, separate them with semicolons and enclose the list in quotation marks.

`"gpfs3nsd;gpfs4nsd;gpfs5nsd"`

**Options**

- e** Displays all of the disks in the file system that do not have an availability of **up** and a status of **ready**. If all disks in the file system are **up** and **ready**, the message displayed is:  
6027-623 All disks up and ready
- L** Displays an extended list of the disk parameters, including the disk ID field and the **remarks** field. The **remarks** column shows the current file system descriptor quorum assignments, and displays the excluded disks. The **remarks** field contains **desc** for all disks assigned as the file system descriptor holders and **excl** for all excluded disks.
- M** Displays whether I/O requests to the disk are satisfied on the local node, or using an NSD server. If the I/O is done using an NSD server, shows the NSD server name and the underlying disk name on that server node.
- m** Displays whether I/O requests to the disk are satisfied on the local node, or using an NSD server. The scope of this option is the node on which the **mmlsdisk** command is issued.
- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Security**

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

As root, the command can also do an **mmlsdisk** on remote file systems.

## mmlsdisk

If you are a non-root user, you may specify only file systems that belong to the same cluster as the node on which the **mmlsdisk** command was issued.

The **mmlsdisk** command does not work if GPFS is down.

### Examples

1. To display the current state of **gpfs2nsd**, issue this command:

```
mmlsdisk /dev/fs0 -d gpfs2nsd
```

The system displays information similar to:

| disk name | driver type | sector size | failure group | holds metadata | holds data | status | storage availability pool |
|-----------|-------------|-------------|---------------|----------------|------------|--------|---------------------------|
| gpfs2nsd  | nsd         | 512         | 4002          | yes            | yes        | ready  | up system                 |

**Note:** In this output, “sector size” refers to logical sector size.

2. To display the current states of **gpfs2nsd**, **gpfs3nsd**, and **gpfs4nsd**, and display their respective disk ids and the descriptor quorum assignment, issue this command:

```
mmlsdisk /dev/fs0 -d "gpfs2nsd;gpfs3nsd;gpfs4nsd" -L
```

The system displays information similar to:

| disk name | driver type | sector size | failure group | holds metadata | holds data | status | storage availability | disk id | pool   | remarks |
|-----------|-------------|-------------|---------------|----------------|------------|--------|----------------------|---------|--------|---------|
| gpfs2nsd  | nsd         | 512         | 4002          | yes            | yes        | ready  | up                   | 2       | system | desc    |
| gpfs3nsd  | nsd         | 512         | 4002          | yes            | yes        | ready  | up                   | 3       | system |         |
| gpfs4nsd  | nsd         | 512         | 4002          | yes            | yes        | ready  | up                   | 4       | system |         |

Number of quorum disks: 3  
Read quorum value: 2  
Write quorum value: 2

**Note:** In this output, “sector size” refers to logical sector size.

3. After the **mmchdisk fs0 empty -d gpfs1nsd** command has been issued, you can view the current state of **gpfs1nsd** by issuing the following command:

```
mmlsdisk fs0 -L
```

In IBM Spectrum Scale V4.1.1 and later, the system displays information similar to the following example:

| disk name | driver type | sector size | failure group | holds metadata | holds data | status        | storage availability | disk id | pool   |
|-----------|-------------|-------------|---------------|----------------|------------|---------------|----------------------|---------|--------|
| gpfs1nsd  | nsd         | 512         | -1            | Yes            | Yes        | to be emptied | up                   | 1       | system |
| gpfs2nsd  | nsd         | 512         | -1            | Yes            | Yes        | to be emptied | up                   | 2       | system |
| gpfs3nsd  | nsd         | 512         | -1            | Yes            | Yes        | ready         | up                   | 3       | system |
| gpfs4nsd  | nsd         | 512         | -1            | Yes            | Yes        | ready         | up                   | 4       | system |

Number of quorum disks: 3  
Read quorum value: 2  
Write quorum value: 2

Attention: Due to an earlier configuration change the file system may contain data that is at risk of being lost.

4. To display whether the I/O is performed locally or using an NSD server, the NSD server name, and the underlying disk name for the file system named **test**, issue this command:

```
mmlsdisk test -M
```

The system displays information similar to:

| Disk name | IO performed on node | Device       | Availability |
|-----------|----------------------|--------------|--------------|
| gpfs7nsd  | localhost            | /dev/hdisk12 | up           |
| gpfs10nsd | k5n88.kgn.ibm.com    | /dev/hdisk13 | up           |
| gpfs4nsd  | localhost            | /dev/hdisk10 | up           |

5. To display the same information as in the previous example, but limited to the node on which the command is issued, issue this command:

```
mmlsdisk test -m
```

The system displays information similar to:

| Disk name | IO performed on node | Device       | Availability |
|-----------|----------------------|--------------|--------------|
| gpfs7nsd  | localhost            | /dev/hdisk12 | up           |
| gpfs10nsd | k5n88.kgn.ibm.com    | -            | up           |
| gpfs4nsd  | localhost            | /dev/hdisk10 | up           |

## See also

- “mmadddisk command” on page 23
- “mmchdisk command” on page 161
- “mmdeldisk command” on page 287
- “mmrpldisk command” on page 539

## Location

/usr/lpp/mmfs/bin

## **mmlsfileset**

---

### **mmlsfileset command**

Displays attributes and status for GPFS filesets.

#### **Synopsis**

```
| mmlsfileset Device
| [[Fileset[,Fileset...]] [-J Junction[,Junction...]] | -F FileName]
| [-d [--block-size {BlockSize | auto}]] [-i] [-L] [-X] [-Y]
| [--afm] [--deleted] [--iam-mode]
```

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

Use the **mmlsfileset** command to display information for the filesets that belong to a given GPFS file system. The default is to display information for all filesets in the file system. You may choose to display information for only a subset of the filesets.

The operation of the **-L** flag omits the attributes listed without it, namely status and junction path. In addition, if the fileset has status Deleted, then **-L** also displays the name of the latest snapshot that includes the fileset in place of the root inode number and parent fileset identifier.

The attributes displayed are:

- Name of the fileset
- Status of the fileset (when the **-L** flag is omitted)
- Junction path to the fileset (when the **-L** flag is omitted)
- Fileset identifier (when the **-L** flag is included)
- Root inode number, if not deleted (when the **-L** flag is included)
- Parent fileset identifier, if not deleted (when the **-L** flag is included)
- Latest including snapshot, if deleted (when the **-L** flag is included)
- Creation time (when the **-L** flag is included)
- Inode space (when the **-L** flag is included)
- Number of inodes in use (when the **-i** flag is included)
- Data size (when the **-d** flag is included)
- Comment (when the **-L** flag is included)
- Caching-related information (when the **--afm** flag is included)
- Value of the permission change flag (when the **-X** flag is used to generate stanza output)
- Integrated archive manager (IAM) mode information

For information on GPFS filesets, see the *IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system that contains the fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

##### *Fileset*

Specifies a comma-separated list of fileset names.

**-J Junction**

Specifies a comma-separated list of path names. They are not restricted to fileset junctions, but may name any file or directory within the filesets to be listed.

**-F FileName**

Specifies the name of a file containing either fileset names or path names. Each line must contain a single entry. All path names must be fully-qualified.

**-d** Displays the amount of storage in use for the fileset.

This operation requires an amount of time that is proportional to the size of the file system; therefore, it can take several minutes or even hours on a large and heavily-loaded file system.

This optional parameter can impact overall system performance. Avoid running the **mmlsfileset** command with this parameter frequently or during periods of high file system activity.

- | This option is not valid if the *Device* parameter is a remote file system.

**--block-size {BlockSize | auto}**

Specifies the unit in which the number of blocks is displayed. The value must be of the form [n]K, [n]M, [n]G or [n]T, where *n* is an optional integer in the range 1 to 1023. The default is 1K. If **auto** is specified, the number of blocks is automatically scaled to an easy-to-read value.

**-i** Displays the number of inodes in use for the fileset.

- | This option is not valid if the *Device* parameter is a remote file system.

This operation requires an amount of time that is proportional to the number of inodes in the file system; therefore, it can take several minutes or even hours on a large and heavily-loaded file system.

Information about the number of inodes in the fileset can be retrieved more efficiently with the following command, if quota management has been enabled for the file system:

**mmrepquota -j FileSystem**

**-L** Displays additional information for the fileset. This includes:

- Fileset identifier
- Root inode number
- Parent identifier
- Fileset creation time
- Inode space
- User defined comments, if any

If the fileset is a dependent fileset, dpnd will be displayed next to the inode space identifier.

**-X** Generates stanza output containing the following:

- The same information presented by the **-L** flag
- The value of the permission change flag
- The same information presented by the **--afm** flag

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**--afm**

Displays caching-related information for the fileset.

**--deleted**

Displays only the filesets with a status of Deleted.

## **mmlsfileset**

### **--iam-mode**

Displays integrated archive manager (IAM) mode information.

### **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

Fileset owners can run the **mmlsfileset** command with the **-L**, **-d**, and **-i** options.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. This command displays fileset information for all filesets in file system **gpfs1**:

```
mmlsfileset gpfs1
```

The system displays information similar to:

```
Filesets in file system 'gpfs1':
Name Status Path
root Linked /gpfs1
fset1 Linked /gpfs1/fset1
fset2 Linked /gpfs1/fset1/fset2
```

2. These commands display information for a file system with filesets and snapshots. Note that deleted filesets that are saved in snapshots are displayed with the name enclosed in parentheses.

- a. Command:

```
mmlsfileset fs1 -d -i
```

The system displays information similar to:

```
Filesets in file system 'fs1':
Name Status Path Inodes Data (in KB)
root Linked /gpfs 3 53528
(gone) Deleted /gpfs/.snapshots/Snap17/gone 0 0
TestF4 Linked /gpfs/test-f4 3 24
TestF3 Linked /gpfs/test-f4/dir1/f3 2 16
TestF2 Unlinked -- 98 784
TestF5 Linked <TestF2>/subdir/f5 1 8
```

- b. Command:

```
mmlsfileset fs1 --deleted
```

The system displays information similar to:

```
Filesets in file system 'fs1':
Name Status Path
(gone) Deleted /gpfs/.snapshots/Snap17/gone
```

- c. Command:

```
mmlsfileset fs1 --afm
```

The system displays information similar to:

```
Filesets in file system 'fs1':
Name Status Path afmTarget
root Linked /gpfs/fs1 --
ro1 Linked /gpfs/fs1/ro1 hs21n45:/gpfs/fs1/ro1
sw1 Linked /gpfs/fs1/sw1 hs21n45:/gpfs/fs1/sw1
lu1 Linked /gpfs/fs1/lu1 hs21n45:/gpfs/fs1/lu1
```

d. Command:

```
mmlsfileset fs1 -L
```

The system displays information similar to:

| Filesets in file system 'fs1': |    |           |          |                          |   |      | InodeSpace | MaxInodes | AllocInodes  | Comment       |
|--------------------------------|----|-----------|----------|--------------------------|---|------|------------|-----------|--------------|---------------|
| Name                           | Id | RootInode | ParentId | Created                  |   |      |            |           |              |               |
| root                           | 0  | 3         | --       | Mon Jan 23 18:59:36 2012 | 0 |      | 1000064    | 65792     | root fileset |               |
| TestF4                         | 2  | 59446     | 0        | Wed Feb 1 09:28:50 2012  | 0 |      | 0          | 0         | 0            | 4th in series |
| TestF3                         | 3  | 59435     | 2        | Wed Feb 1 09:28:52 2012  | 0 |      | 0          | 0         | 0            |               |
| (gone)                         | 1  | latest:   | Snap17   | Wed Feb 1 09:28:46 2012  | 0 |      | 0          | 0         | 0            | Not forgotten |
| TestF2                         | 4  | 59437     | --       | Wed Feb 1 09:28:52 2012  | 0 |      | 0          | 0         | 0            |               |
| TestF5                         | 5  | 7017      | 4        | Wed Feb 1 09:28:53 2012  | 0 |      | 0          | 0         | 0            | Number 5      |
| FsetF1-V2                      | 6  | 131075    | --       | Wed Feb 1 09:28:55 2012  | 1 |      | 100096     | 100096    | 100096       |               |
| FsetF2-V2                      | 7  | 262147    | --       | Wed Feb 1 09:28:56 2012  | 2 |      | 100096     | 100096    | 100096       |               |
| FsetF2-V2-lite                 | 9  | 263680    | --       | Wed Feb 1 09:28:59 2012  | 2 | dpnd | 0          | 0         | 0            |               |
| FsetF3-V2                      | 8  | 393219    | --       | Wed Feb 1 09:28:57 2012  | 3 |      | 100096     | 100096    | 100096       |               |

e. Command:

```
mmlsfileset fs1 sw1 --afm -L
```

The system displays information similar to:

Filesets in file system 'fs1':

| Attributes for fileset sw:   |                   |  |
|------------------------------|-------------------|--|
| =====                        |                   |  |
| Status                       | Linked            |  |
| Path                         | /gpfs/fs1/sw      |  |
| afm-associated               | Yes               |  |
| Target                       | c2m3n06:/gpfs/fs2 |  |
| Mode                         | single-writer     |  |
| File Lookup Refresh Interval | 30 (default)      |  |
| File Open Refresh Interval   | 30 (default)      |  |
| Dir Lookup Refresh Interval  | 60 (default)      |  |
| Dir Open Refresh Interval    | 60 (default)      |  |
| Async Delay                  | 15 (default)      |  |
| Expiration Timeout           | disable           |  |
| Recovery Point Objective     | disable (default) |  |
| Last pSnapId                 | 0                 |  |
| Display Home Snapshots       | no                |  |

f. Command:

```
mmlsfileset fs1 TestF2,TestF5 -J /gpfs/test-f4/dir1,/gpfs/test-f4/dir1/f3/dir2/
```

The system displays information similar to:

Filesets in file system 'fs1':

| Name   | Status   | Path                  |
|--------|----------|-----------------------|
| TestF2 | Unlinked | --                    |
| TestF5 | Linked   | <TestF2>/subdir/f5    |
| TestF4 | Linked   | /gpfs/test-f4         |
| TestF3 | Linked   | /gpfs/test-f4/dir1/f3 |

g. Command:

```
mmlsfileset gpfsha --deleted
```

The system displays information similar to:

Filesets in file system 'gpfsha':

| Name     | Status  | Path                             |
|----------|---------|----------------------------------|
| (fset17) | Deleted | /gpfsha/.snapshots/snap20/fset17 |

## See also

- “mmchfileset command” on page 173
- “mmcrlfileset command” on page 244
- “mmdelfileset command” on page 292
- “mmlinkfileset command” on page 390
- “mmunlinkfileset command” on page 578

## **mmIsfileset**

### **Location**

/usr/lpp/mmfs/bin

## mmlsfs command

Displays file system attributes.

### Synopsis

```
mmlsfs {Device | all | all_local | all_remote} [-A] [-B] [-d] [-D]
[-E] [-f] [-i] [-I] [-j] [-k] [-K] [-L] [-m] [-M] [-n] [-o]
[-P] [-Q] [-r] [-R] [-S] [-t] [-T] [-V] [-Y] [-z]
[--create-time] [--encryption] [--fastea] [--filesetdf]
[--inode-limit] [--is4KAligned] [--log-replicas] [--mount-priority]
[--perfileset-quota] [--rapid-repair] [--write-cache-threshold]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlsfs** command to list the attributes of a file system.

Depending on your configuration, additional information that is set by GPFS may be displayed to assist in problem determination when contacting the IBM Support Center.

### Results

If you do not specify any options, all attributes of the file system are displayed. When you specify options, only those attributes specified are listed, in the order issued in the command. Some parameters are preset for optimum performance and, although they display in the **mmlsfs** command output, you cannot change them.

### Parameters

The following parameter must be the first parameter:

*Device* | **all** | **all\_local** | **all\_remote**

*Device*

Indicates the device name of the file system for which information is displayed. File system names do not need to be fully qualified. *fs0* is as acceptable as */dev/fs0*.

**all**

Indicates all file systems that are known to this cluster.

**all\_local**

Indicates all file systems that are owned by this cluster.

**all\_remote**

Indicates all file systems that are owned by another cluster.

This must be the first parameter.

The following optional parameters, when used, must be provided after the *Device* | **all** | **all\_local** | **all\_remote** parameter:

- A** Displays if and when the file system is automatically mounted.
- | -**B** Displays the block size of the file system in bytes. For more information about block size, see the description of the **-B BlockSize** parameter in “mmcrfs command” on page 250.
- d** Displays the names of all of the disks in the file system.

## **mmIsfs**

- D** Displays the type of file locking semantics that are in effect (**nfs4** or **posix**).
- E** Displays the exact **mtime** values reported.
- f** Displays the minimum fragment (subblock) size of the file system in bytes.
- i** Displays the inode size, in bytes.
- I** Displays the indirect block size, in bytes.
- j** Displays the block allocation type.
- k** Displays the type of authorization supported by the file system.
- K** Displays the strict replication enforcement.
- L** Displays the internal log file size.
- m** Displays the default number of metadata replicas.
- M** Displays the maximum number of metadata replicas.
- n** Displays the estimated number of nodes for mounting the file system.
- o** Displays the additional mount options.
- P** Displays the storage pools defined within the file system.
- Q** Displays which quotas are currently enforced on the file system.
- r** Displays the default number of data replicas.
- R** Displays the maximum number of data replicas.
- S** Displays whether the updating of **atime** is suppressed for the **gpfs\_stat()**, **gpfs\_fstat()**, **stat()**, and **fstat()** calls.
- t** Displays the Windows drive letter.
- T** Displays the default mount point.
- V** Displays the current format version of the file system.
- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

- z** Displays whether DMAPI is enabled for this file system.

### **--create-time**

Displays the creation time of the file system.

### **--encryption**

Displays a **yes** or **no** value indicating whether encryption is enabled. This value cannot be changed with the **mmchfs** command. When the cluster is created this value is set to **no**. When an encryption policy is established for the file system, the value is set to **yes**.

### **--fastea**

Displays a **yes** or **no** value indicating whether fast external attributes is enabled. Displays a **migrating** value if migration was initiated with **mmigratefs --fastea** but is not yet complete.

### **--filesetdf**

Displays a **yes** or **no** value indicating whether **filesetdf** is enabled; if **yes**, the **mmdf** command reports numbers based on the quotas for the fileset and not for the total file system.

**--inode-limit**

Displays the maximum number of files in the file system.

**--is4KAligned**

Displays whether file systems are formatted to be 4K aligned.

**--log-replicas**

Displays the number of recovery log replicas. If a value of **0** is displayed, the number of recovery log replicas is the same as the number of metadata replicas currently in effect for the file system.

**--mount-priority**

Displays the assigned mount priority.

**--profileset-quota**

Displays the per-profileset quota.

**--rapid-repair**

Displays a **yes** or **no** value indicating whether the per-block replication tracking and repair feature is enabled.

**--write-cache-threshold**

Displays the threshold below which synchronous writes will be initially buffered in the highly-available write cache before being written back to primary storage.

## Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

**Note:** The command treats the following conditions as failures:

- The file system that you specified was not found.
- You specified **all**, **all\_local**, or **all\_remote** and no file systems were found.

## Security

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

As root, a user can also issue the **mmlsfs** on remote file systems.

If you are a non-root user, you may specify only file systems that belong to the same cluster as the node on which the **mmlsfs** command was issued.

## Examples

| The following command displays the attributes of the file system **gpfs1**:

| **mmlsfs gpfs1**

| The system displays information like the following:

| flag  | value | description                         |
|-------|-------|-------------------------------------|
| <hr/> |       |                                     |
| -f    | 8192  | Minimum fragment size in bytes      |
| -i    | 4096  | Inode size in bytes                 |
| -I    | 16384 | Indirect block size in bytes        |
| -m    | 2     | Default number of metadata replicas |
| -M    | 2     | Maximum number of metadata replicas |

## mmlsfs

|                         |                                                  |                                                                        |
|-------------------------|--------------------------------------------------|------------------------------------------------------------------------|
| -r                      | 2                                                | Default number of data replicas                                        |
| -R                      | 2                                                | Maximum number of data replicas                                        |
| -j                      | cluster                                          | Block allocation type                                                  |
| -D                      | nfs4                                             | File locking semantics in effect                                       |
| -k                      | all                                              | ACL semantics in effect                                                |
| -n                      | 32                                               | Estimated number of nodes that will mount file system                  |
| -B                      | 262144                                           | Block size                                                             |
| -Q                      | user;group;fileset<br>user;group;fileset<br>none | Quotas accounting enabled<br>Quotas enforced<br>Default quotas enabled |
| --perfileset-quota      | no                                               | Per-fileset quota enforcement                                          |
| --filesetdf             | no                                               | Fileset df enabled?                                                    |
| -V                      | 16.00 (4.2.2.0)                                  | File system version                                                    |
| --create-time           | Thu Oct 27 13:59:13 2016                         | File system creation time                                              |
| -z                      | no                                               | Is DMAPI enabled?                                                      |
| -L                      | 134217728                                        | Logfile size                                                           |
| -E                      | yes                                              | Exact mtime mount option                                               |
| -S                      | no                                               | SUPPRESS atime mount option                                            |
| -K                      | whenpossible                                     | Strict replica allocation option                                       |
| --fastea                | yes                                              | Fast external attributes enabled?                                      |
| --encryption            | no                                               | Encryption enabled?                                                    |
| --inode-limit           | 607488                                           | Maximum number of inodes in all inode spaces                           |
| --log-replicas          | 2                                                | Number of log replicas                                                 |
| --is4KAligned           | yes                                              | is4KAligned?                                                           |
| --rapid-repair          | yes                                              | rapidRepair enabled?                                                   |
| --write-cache-threshold | 65536                                            | HAWC Threshold (max 65536)                                             |
| -P                      | system                                           | Disk storage pools in file system                                      |
| -d                      | nsd20;nsd21;nsd3                                 | Disks in file system                                                   |
| -A                      | yes                                              | Automatic mount option                                                 |
| -o                      | none                                             | Additional mount options                                               |
| -T                      | /gpfs1                                           | Default mount point                                                    |
| --mount-priority        | 0                                                | Mount priority                                                         |

- If you issue the **mmlsfs** command with the **all** option:

```
mmlsfs all -A
```

The system displays information similar to:

File system attributes for /dev/fs1:

| flag | value | description            |
|------|-------|------------------------|
| -A   | yes   | Automatic mount option |

File system attributes for /dev/gpfs1:

| flag | value | description            |
|------|-------|------------------------|
| -A   | yes   | Automatic mount option |

## See also

- “mmcrfs command” on page 250
- “mmchfs command” on page 179
- “mmdelfs command” on page 295

## Location

/usr/lpp/mmfs/bin

---

## mmlslicense command

Displays information about the GPFS node licensing designation or about disk and cluster capacity.

### Synopsis

| **mmlslicense [-Y] [-L | --capacity [--formatted]]**

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlslicense** command to display the number of GPFS client, FPO, and server licenses assigned to the nodes in the cluster.

For information on IBM Spectrum Scale license designation, see *IBM Spectrum Scale license designation* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### Parameters

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.
 

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.
- L Displays information about the license type that is associated with each of the nodes in the cluster. An asterisk after the license type indicates insufficient license level for the roles that the node performs.
- capacity Displays disk and cluster size information.
- formatted When --capacity is specified, inserts commas to separate groups of three numerals for readability.

### Exit status

0 Successful completion.

#### nonzero

A failure occurred.

### Security

You must have root authority to run the **mmlslicense** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### Examples

1. To display the summary information about the type and number of GPFS licenses that are associated with the nodes in the cluster, enter the following command:

**mmlslicense**

## **mmlslicense**

The system displays information similar to the following example:

Summary information

```

Number of nodes defined in the cluster: 4
Number of nodes with server license designation: 1
Number of nodes with client license designation: 2
Number of nodes still requiring server license designation: 1
Number of nodes still requiring client license designation: 1
This node runs IBM Spectrum Scale Advanced Edition
```

2. To display information about the type of GPFS licenses that are associated with each of the nodes in the cluster, enter the following command:

```
mmlslicense -L
```

The system displays information similar to the following example:

| Node name           | Required license | Designated license |
|---------------------|------------------|--------------------|
| k145n05.kgn.ibm.com | server           | server             |
| k145n06.kgn.ibm.com | server           | client *           |
| k145n07.kgn.ibm.com | client           | client             |
| k145n08.kgn.ibm.com | client           | none *             |

Summary information

```

Number of nodes defined in the cluster: 4
Number of nodes with server license designation: 1
Number of nodes with client license designation: 2
Number of nodes still requiring server license designation: 1
Number of nodes still requiring client license designation: 1
This node runs IBM Spectrum Scale Advanced Edition
```

| 3. To display information for capacity-based licensing, enter the following command:

```
| mmlslicense --capacity --formatted
```

| The system displays information similar to the following example:

```
| NSD Summary:
| ======
| Total Number of NSDs: 2
| mynsd1: 10,737,418,240 Bytes
| mynsd2: 10,737,418,240 Bytes

| Cluster Summary:
| ======
| Cluster Total Capacity: 21,474,836,480 Bytes
```

### **See also**

- “mmchlicense command” on page 185

### **Location**

/usr/lpp/mmfs/bin

---

## mmlsmgr command

Displays which node is the file system manager for the specified file systems or which node is the cluster manager.

### Synopsis

| **mmlsmgr** [*Device* [*Device...*]] [-Y]

or

**mmlsmgr -C** *RemoteClusterName*

or

**mmlsmgr -c**

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlsmgr** command to display which node is the file system manager or cluster manager for the file system.

If you do not provide a *Device* operand, file system managers for all file systems within the current cluster for which a file system manager has been appointed are displayed.

### Parameters

#### *Device*

The device names of the file systems for which the file system manager information is displayed.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

If no file system is specified, information about all file systems is displayed.

#### **-C** *RemoteClusterName*

Displays the name of the nodes that are file system managers in cluster *RemoteClusterName*.

#### **-c**

Displays the current cluster manager node.

#### **-Y**

Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### Exit status

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## **mmlsmgr**

### **Security**

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

As root, a user can also issue the **mmlsmgr** on remote file systems.

If you are a non-root user, you may specify only file systems that belong to the same cluster as the node on which the **mmlsmgr** command was issued.

### **Examples**

1. To display the file system manager node information for all the file systems, issue this command:

```
mmlsmgr
```

The system displays information similar to:

| file system | manager node          |
|-------------|-----------------------|
| fs3         | 9.114.94.65 (c154n01) |
| fs2         | 9.114.94.73 (c154n09) |
| fs1         | 9.114.94.81 (c155n01) |

Cluster manager node: 9.114.94.65 (c154n01)

The output shows the device name of the file system and the file system manager's node number and name, in parenthesis, as they are recorded in the GPFS cluster data.

2. To display the file system manager information for file systems **gpfs2** and **gpfs3**, issue this command:

```
mmlsmgr gpfs2 gpfs3
```

The system displays information similar to:

| file system | manager node [from 199.116.68.69 (k156gn02)] |
|-------------|----------------------------------------------|
| gpfs2       | 199.116.68.70 (k154gn02)                     |
| gpfs3       | 199.116.68.72 (kolt2g_r1b42)                 |

### **See also**

- “**mcmgr** command” on page 188

### **Location**

/usr/lpp/mmfs/bin

## mmlsmount command

Lists the nodes that have a given GPFS file system mounted.

### Synopsis

```
mmlsmount {Device | all | all_local | all_remote | {-F DeviceFileName}} [-L]
 [-Y] [-C {all | all_remote | ClusterName[,ClusterName...]}]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmlsmount** command reports if a file system is in use at the time the command is issued. A file system is considered to be in use if it is explicitly mounted with the **mount** or **mmmount** command, or if it is mounted internally for the purposes of running some other GPFS command. For example, when you run the **mmrestripefs** command, the file system will be internally mounted for the duration of the command. If **mmlsmount** is issued in the interim, the file system will be reported as being in use by the **mmlsmount** command but, unless it is explicitly mounted, will not show up in the output of the **mount** or **df** commands.

### Parameters

*Device | all | all\_local | all\_remote | {-F DeviceFileName}*

Indicates the file system or file systems for which information is displayed.

*Device*

Indicates the device name of the file system for which information is displayed. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

**all**

Indicates all file systems known to this cluster.

**all\_local**

Indicates all file systems owned by this cluster.

**all\_remote**

Indicates all file systems owned by another cluster.

**-F DeviceFileName**

Specifies a file containing the device names, one per line, of the file systems for which information is displayed.

This must be the first parameter.

### Options

**-L** Specifies to list the nodes that have the file system mounted.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**-C {all | all\_remote | ClusterName[,ClusterName...]}**

Specifies the clusters for which mount information is requested. If one or more *ClusterName* is

## **mmlsmount**

specified, only the names of nodes that belong to these clusters and have the file system mounted are displayed. The dot character ('.') can be used in place of the cluster name to denote the local cluster.

Option **-C all\_remote** denotes all clusters other than the one from which the command was issued.

Option **-C all** refers to all clusters, local and remote, that can have the file system mounted. Option **-C all** is the default.

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## **Security**

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

If you are a non-root user, you may specify only file systems that belong to the same cluster as the node on which the **mmlsmount** command was issued.

## **Examples**

1. To see how many nodes have file system **fs2** mounted, issue this command:

```
mmlsmount fs2
```

The system displays output similar to:

```
File system fs2 is mounted on 3 nodes.
```

2. To display all mounted file systems:

```
mmlsmount all
```

The system displays output similar to:

```
File system fs1 is mounted on 17 nodes.
```

```
File system remotefs1 (remote.cluster:fs1) is mounted on 17 nodes.
```

3. To display all remotely mounted file systems:

```
mmlsmount all_remote
```

The system displays output similar to:

```
File system remotefs1 (remote.cluster:fs1) is mounted on 17 nodes.
```

4. To list the nodes having all file systems mounted:

```
mmlsmount all -L
```

The system displays output similar to:

```
File system fs1 is mounted on 3 nodes:
```

```
192.168.105.32 c6f1c3vp2
```

```
192.168.105.31 c6f1c3vp1
```

```
192.168.105.34 c6f1c3vp4
```

```
File system gpfs1 is not mounted.
```

## **See also**

- “**mmmount command**” on page 444
- “**mmumount command**” on page 575

## **Location**

/usr/lpp/mmfs/bin

---

## mmlsnodeclass command

Displays node classes defined in the system.

### Synopsis

| **mmlsnodeclass** [*ClassName*[,*ClassName*...]] | --user | --system | --all] [-Y]

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlsnodeclass** command to display node classes defined in the system.

### Parameters

#### *ClassName*

Displays the specified node class.

#### --user

Displays all user-defined node classes. This is the default.

#### --system

Displays all system-defined node classes.

#### --all

Displays both the system-defined and user-defined node classes.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### Exit status

**0** Successful completion.

#### nonzero

A failure has occurred.

### Security

You must have root authority to run the **mmlsnodeclass** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### Examples

1. To display the current user-defined node classes, issue this command:

```
mmlsnodeclass
```

The system displays information similar to:

## **mmlsnodeclass**

| Node Class Name | Members    |
|-----------------|------------|
| siteA           | linuxNodes |

2. To display all node classes defined in the system, issue this command:

```
mmlsnodeclass --all
```

The system displays information similar to:

| Node Class Name       | Members                                              |
|-----------------------|------------------------------------------------------|
| aixNodes              |                                                      |
| all                   | node03, node01, node02, node04, nodensd01, nodensd02 |
| cesNodes              | node01, node02, node04                               |
| clientLicense         |                                                      |
| clientNodes           | node03, node02, node04, nodensd01, nodensd02         |
| cnfsNodes             |                                                      |
| disabledCnfsNodes     |                                                      |
| enabledCnfsNodes      |                                                      |
| linuxNodes            | node03, node01, node02, node04, nodensd01, nodensd02 |
| managerNodes          | node01                                               |
| nonAixNodes           | node03, node01, node02, node04, nodensd01, nodensd02 |
| nonCesNodes           | node03, nodensd01, nodensd02                         |
| nonCnfsNodes          | node03, node01, node02, node04, nodensd01, nodensd02 |
| nonLinuxNodes         |                                                      |
| nonNsdNodes           | node003, node01, node02, node04                      |
| nonQuorumNodes        |                                                      |
| nonWindowsNodes       | node03, node01, node02, node04, nodensd01, nodensd02 |
| nsdNodes              | nodensd01, nodensd02                                 |
| quorumNodes           | node03, node01, node02, node04, nodensd01, nodensd02 |
| serverLicense         | node03, node01, node02, node04, nodensd01, nodensd02 |
| windowsNodes          |                                                      |
| siteA                 | linuxNodes                                           |
| object_database_node  | node01                                               |
| object_singleton_node | node01                                               |

3. To display only the nodes that are quorum nodes, issue this command:

```
mmlsnodeclass quorumNodes
```

The system displays information similar to:

| Node Class Name | Members                         |
|-----------------|---------------------------------|
| quorumNodes     | c8f2c1vp4, c8f2c4vp1, c8f2c4vp2 |

## **See also**

- “mmchnodeclass command” on page 195
- “mmcrnodeclass command” on page 260
- “mmdelnnodeclass command” on page 300

## **Location**

/usr/lpp/mmfs/bin

## mmIsnsd command

Displays Network Shared Disk (NSD) information for the GPFS cluster.

### Synopsis

```
mmIsnsd [-a | -F | -f Device | -d "DiskName[;DiskName...]"]
 [-L | -m | -M | -X] [-Y | -v]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmIsnsd** command to display the current information for the NSDs belonging to the GPFS cluster. The default is to display information for all NSDs defined to the cluster (-a). Otherwise, you may choose to display the information for a particular file system (-f) or for all disks that do not belong to any file system (-F).

### Parameters

- | **-a** Displays information for all of the NSDs belonging to the GPFS cluster. This is the default.
- | **-f *Device***
- |     Specifies the device name of the file system for which you want NSD information displayed. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.
- | **-F** Displays the NSDs that are not in use.
- | **-d *DiskName[;DiskName...]***
- |     Specifies the name of the NSDs for which you want information displayed. When you enter multiple *DiskNames*, separate them with semicolons and enclose the entire string of disk names in quotation marks:
- |     "gpfs3nsd;gpfs4nsd;gpfs5nsd"

### Options

- L** Displays the information in a long format that shows the NSD identifier.
- m** Maps the NSD name to its disk device name on the local node and, if applicable, on the NSD server nodes.
- M** Maps the NSD names to its disk device name on all nodes.  
This is a slow operation and its usage is suggested for problem determination only.
- v** Specifies that the output should contain error information, where available.
- X** Maps the NSD name to its disk device name on the local node and, if applicable, on the NSD server nodes. The **-X** option also displays extended information for the NSD volume ID and information such as NSD server status and Persistent Reserve (PR) enablement in the Remarks field. Using the **-X** option is a slow operation and is recommended only for problem determination.
- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

## **mmlsnsd**

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to issue the **mmlsnsd** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. To display the default information for all of the NSDs belonging to the cluster, issue this command:

```
mmlsnsd
```

The system displays information similar to:

| File system | Disk name | NSD servers            |
|-------------|-----------|------------------------|
| fs2         | hd3n97    | c5n97g, c5n98g, c5n99g |
| fs2         | hd4n97    | c5n97g, c5n98g, c5n99g |
| fs2         | hd5n98    | c5n98g, c5n97g, c5n99g |
| fs2         | hd6n98    | c5n98g, c5n97g, c5n99g |
| fs2         | hd7n97    | c5n97g, c5n98g, c5n99g |
| fs2         | hd8n97    | c5n97g, c5n98g, c5n99g |
| fs2         | hd9n97    | c5n97g, c5n98g, c5n99g |
| fs2         | hd10n98   | c5n98g, c5n97g, c5n99g |
| fs2         | hd11n98   | c5n98g, c5n97g         |
| fs2         | hd12n98   | c5n98g, c5n97g         |
| fs2         | sdbnsd    | c5n94g, c5n96g         |
| fs2         | sdcnsd    | c5n94g, c5n96g         |
| fs2         | sddnsd    | c5n94g, c5n96g         |
| fs2         | sdensd    | c5n94g, c5n96g         |
| fs2         | sdgnsd    | c5n94g, c5n96g         |
| fs2         | sdfnsd    | c5n94g, c5n96g         |
| fs2         | sdhnsd    | c5n94g, c5n96g         |
| (free disk) | hd2n97    | c5n97g, c5n98g         |

2. To display all of the NSDs attached to the node from which the command is issued, issue this command:

```
mmlsnsd -m
```

The system displays information similar to:

| Disk name | NSD volume ID    | Device       | Node name | Remarks     |
|-----------|------------------|--------------|-----------|-------------|
| hd10n98   | 0972846245C8E93C | /dev/hd10n98 | c5n97g    | server node |
| hd10n98   | 0972846245C8E93C | /dev/hd10n98 | c5n98g    | server node |
| hd11n98   | 0972846245C8E93F | /dev/hd11n98 | c5n97g    | server node |
| hd11n98   | 0972846245C8E93F | /dev/hd11n98 | c5n98g    | server node |
| hd12n98   | 0972846245C8E941 | /dev/hd12n98 | c5n97g    | server node |
| hd12n98   | 0972846245C8E941 | /dev/hd12n98 | c5n98g    | server node |
| hd2n97    | 0972846145C8E924 | /dev/hdisk2  | c5n97g    | server node |
| hd2n97    | 0972846145C8E924 | /dev/hdisk2  | c5n98g    | server node |
| hd3n97    | 0972846145C8E927 | /dev/hdisk3  | c5n97g    | server node |
| hd3n97    | 0972846145C8E927 | /dev/hdisk3  | c5n98g    | server node |
| hd4n97    | 0972846145C8E92A | /dev/hdisk4  | c5n97g    | server node |
| hd4n97    | 0972846145C8E92A | /dev/hdisk4  | c5n98g    | server node |
| hd5n98    | 0972846245EB501C | /dev/hdisk5  | c5n97g    | server node |

|        |                  |             |        |             |
|--------|------------------|-------------|--------|-------------|
| hd5n98 | 0972846245EB501C | /dev/hdisk5 | c5n98g | server node |
| hd6n98 | 0972846245DB3AD8 | /dev/hdisk6 | c5n97g | server node |
| hd6n98 | 0972846245DB3AD8 | /dev/hdisk6 | c5n98g | server node |
| hd7n97 | 0972846145C8E934 | /dev/hd7n97 | c5n97g | server node |

3. To display all of the NSDs in the GPFS cluster in extended format, issue this command:

```
mmIsnsd -L
```

The system displays information similar to:

| File system | Disk name | NSD volume ID    | NSD servers   |
|-------------|-----------|------------------|---------------|
| fs2         | hd3n97    | 0972846145C8E927 | c5n97g,c5n98g |
| fs2         | hd4n97    | 0972846145C8E92A | c5n97g,c5n98g |
| fs2         | hd5n98    | 0972846245EB501C | c5n98g,c5n97g |
| fs2         | hd6n98    | 0972846245DB3AD8 | c5n98g,c5n97g |
| fs2         | sdbnsd    | 0972845E45C8E8ED | c5n94g,c5n96g |
| fs2         | sdcnsd    | 0972845E45C8E8F6 | c5n94g,c5n96g |
| fs2         | sddnsd    | 0972845E45F83FDB | c5n94g,c5n96g |
| fs2         | sdensd    | 0972845E45C8E909 | c5n94g,c5n96g |
| fs2         | sdgnsd    | 0972845E45C8E912 | c5n94g,c5n96g |
| fs2         | sdfnsd    | 0972845E45F02E81 | c5n94g,c5n96g |
| fs2         | sdhnsd    | 0972845E45C8E91C | c5n94g,c5n96g |
| gpfs1       | hd2n97    | 0972846145C8E924 | c5n97g,c5n98g |

4. To display extended disk information about disks **hd3n97**, **sdfnsd**, and **hd5n98**, issue this command:

```
mmIsnsd -X -d "hd3n97;sdfnsd;hd5n98"
```

The system displays information similar to:

| Disk name | NSD volume ID    | Device      | Devtype | Node name | Remarks           |
|-----------|------------------|-------------|---------|-----------|-------------------|
| hd3n97    | 0972846145C8E927 | /dev/hdisk3 | hdisk   | c5n97g    | server node,pr=no |
| hd3n97    | 0972846145C8E927 | /dev/hdisk3 | hdisk   | c5n98g    | server node,pr=no |
| hd5n98    | 0972846245EB501C | /dev/hdisk5 | hdisk   | c5n97g    | server node,pr=no |
| hd5n98    | 0972846245EB501C | /dev/hdisk5 | hdisk   | c5n98g    | server node,pr=no |
| sdfnsd    | 0972845E45F02E81 | /dev/sdf    | generic | c5n94g    | server node       |
| sdfnsd    | 0972845E45F02E81 | /dev/sdm    | generic | c5n96g    | server node       |

5. The following shows the output of **mmIsnsd -X** with **mmchconfig usePersistentReserve=yes**.

| Disk name | NSD volume ID    | Device     | Devtype | Node name           | Remarks            |
|-----------|------------------|------------|---------|---------------------|--------------------|
| nsd0      | 947AC0A84F5FB55C | /dev/dm-0  | dmm     | c13clapv12.gpfs.net | server node,pr=yes |
| nsd0      | 947AC0A84F5FB55C | /dev/dm-19 | dmm     | c13clapv13.gpfs.net | server node,pr=yes |
| nsd0      | 947AC0A84F5FB55C | /dev/dm-6  | dmm     | c13clapv14.gpfs.net | server node,pr=yes |
| nsd0      | 947AC0A84F5FB55C | /dev/dm-15 | dmm     | c13clapv16.gpfs.net | server node,pr=yes |
| nsd1      | 947AC0A84F5FB564 | /dev/dm-1  | dmm     | c13clapv12.gpfs.net | server node,pr=yes |
| nsd1      | 947AC0A84F5FB564 | /dev/dm-4  | dmm     | c13clapv13.gpfs.net | server node,pr=yes |
| nsd1      | 947AC0A84F5FB564 | /dev/dm-9  | dmm     | c13clapv14.gpfs.net | server node,pr=yes |
| nsd1      | 947AC0A84F5FB564 | /dev/dm-13 | dmm     | c13clapv16.gpfs.net | server node,pr=yes |

## See also

- “**mmcrnsd** command” on page 262
- “**mmdelnsd** command” on page 302

## Location

/usr/lpp/mmfs/bin

## **mmlspolicy**

---

### **mmlspolicy command**

Displays policy information.

#### **Synopsis**

| **mmlspolicy** *Device* [-L] [-Y]

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

The **mmlspolicy** command displays policy information for a given file system. The information displayed includes:

- When the policy file was installed.
- The user who installed the policy file.
- The node on which the policy file was installed.
- The first line of the original policy file.

For information about GPFS policies and file placement, see *Information Lifecycle Management in IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system for which policy information is to be displayed. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

- L Displays the entire original policy file. If this flag is not specified, only the first line of the original policy file is displayed.
- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

#### **Exit status**

0      Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

#### **Examples**

1. This command displays basic information for the policy installed for file system **fs2**:

```
mmlspolicy fs2
```

The system displays output similar to:

```
Policy for file system '/dev/fs2':
 Installed by root@c103rp12.gpfs.net on Tue Mar 30 15:06:20 2010.
 First line of policy 'policy' is:
/* This is the policy for the fs2 GPFS file system. */
```

2. This command displays extended information for the policy installed for file system **fs2**:

```
mmfspolicy fs2 -L
```

The system displays output similar to:

```
/* This is the policy for the fs2 GPFS file system. */

/* File Placement Rules */
RULE SET POOL 'sp4' WHERE name like '%sp4%'
RULE SET POOL 'sp5' WHERE name like '%sp5%'
RULE 'default' SET POOL 'system'

/* Snapshot Placement Rules*/
RULE SET SNAP_POOL 'sp1' WHERE SNAP_NAME LIKE '%daily%'
RULE SET SNAP_POOL 'sp2'

/* Exclude Rule */
RULE 'Exclude root users files' EXCLUDE WHERE USER_ID = 0 AND
 name like '%org%'

/* Delete Rule */
RULE 'delete files' DELETE WHERE PATH_NAME like '%tmp%'

/* Migrate Rule */
RULE 'sp4.files' MIGRATE FROM POOL 'sp4' TO POOL 'sp5' WHERE
 name like '%sp4%'

/* End of Policy */
```

3. In this example, no policy file was installed for the specified file system:

```
mmfspolicy fs4 -L
```

The system displays output similar to:

```
No policy file was installed for file system 'fs4'.
Data will be stored in pool 'system'.
```

## See also

- “[mmapplypolicy command](#)” on page 56
- “[mmchpolicy command](#)” on page 201

## Location

/usr/lpp/mmfs/bin

## mmlspool

---

### mmlspool command

Displays information about the known storage pools.

#### Synopsis

```
mmlspool Device {StoragePool[,StoragePool...]} | all} [-L]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

The **mmlspool** command displays basic or detailed information about the storage pools in a file system.

#### Parameters

##### *Device*

Specifies the device name of the file system for which storage pool information is to be displayed. File system names do not need to be fully qualified; for example, **fs0** is as acceptable as **/dev/fs0**.

##### *StoragePool[,StoragePool...]*

Specifies one or more storage pools for which information is to be displayed.

##### **all**

Displays information about all the storage pools in specified file system.

##### **-L**

Displays detailed information about each storage pool.

#### Exit status

**0** Successful completion.

##### **nonzero**

A failure occurred.

#### Security

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

If you are a nonroot user, you may specify only file systems that belong to the same cluster as the node on which the **mmlspool** command was issued.

#### Examples

1. To show basic information about all storage pools in a file system, issue this command:

```
mmlspool /dev/fst all
```

The system displays information similar to:

| Name    | Id    |
|---------|-------|
| system  | 0     |
| sataXXX | 65537 |

2. To show more information, issue this command:

```
mmlspool fs1 p1 -L
```

The system displays information similar to this:

```
Pool:
 name = p1
 poolID = 65537
 blockSize = 256 KB
 usage = dataOnly
 maxDiskSize = 497 GB
 layoutMap = cluster
 allowWriteAffinity = no
 writeAffinityDepth = 0
 blockGroupFactor = 1
```

## See also

- “[mmlsattr command](#)” on page 392
- “[mmlscallback command](#)” on page 395
- “[mmlscluster command](#)” on page 397
- “[mmlsconfig command](#)” on page 400
- “[mmlsdisk command](#)” on page 402
- “[mmlspolicy command](#)” on page 426
- “[mmlsquota command](#)” on page 435
- “[mmlssnapshot command](#)” on page 439

See also the following *IBM Spectrum Scale RAID: Administration* topics:

- “[mmlsrecoverygroup command](#)”
- “[mmlsvdisk command](#)”

## Location

/usr/lpp/mmfs/bin

## mmls qos

### mmls qos command

Displays the I/O performance values of a file system, when you enable Quality of Service for I/O operations (QoS) with the mmchqos command.

#### Synopsis

```
mmls qos Device [--fine-stats DisplayRange]
[--pool {all | Pool}]
[--seconds Seconds]
[--sum-classes {yes | no}]
[--sum-nodes {yes | no}]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

With the **mmls qos** command, you can display the consumption of I/O operations by processes that access designated storage pools. With the **mmchqos** command, you can regulate I/O access to a specified storage pool by allocating shares of I/O operations to two QoS classes:

##### **maintenance**

The default QoS class for some I/O intensive, potentially long-running GPFS commands, such as **mmbackup**, **mmrestore**

##### **other**

The default QoS class for all other processes.

A third class, **misc**, is used to count the IOPS that some critical file system processes consume. You cannot assign IOPS to this class, but its count of IOPS is displayed in the output of the **mmls qos** command.

Remember the following points:

- Allocations persist across unmounting and remounting the file system.
- QoS stops applying allocations when you unmount the file system and resumes when you remount it.
- When you change allocations or mount the file system, a brief delay due to reconfiguration occurs before QoS starts applying allocations.

For more information about this command, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

When the file system is mounted, the command displays information about the QoS classes of both explicitly named pools and unnamed pools. *Unnamed pools* are storage pools that you have not specified by name in any **mmchqos** command. When the file system is unmounted, the command displays information about only the QoS classes of explicitly named pools.

#### Parameters

##### *Device*

The device name of the file system to which the QoS action applies.

##### **--fine-stats** *DisplayRange*

Displays the fine-grained statistics that are currently in memory. The *DisplayRange* values specify the indexes of the first and last blocks of statistics in the range of blocks that you want to display, such as 0-4. If you specify only one integer, such as 2, the command displays all the blocks in memory starting with that index and going to the end. The last line of the output displays the index of the next block in memory. You can avoid re-displaying statistics by having the next **mmls qos** command display statistics beginning at this block index.

Fine-grained statistics are taken at 1-second intervals and contain more information than regular statistics. They are intended to be used as input for programs that analyze and display data, such as the example plotting program at /usr/lpp/mmfs/samples/charts/qosplotfine.pl. For the content of the statistics, see the subtopic later in this topic.

**--pool**

| Displays the I/O performance values for all QoS pools if **all** is specified, or for the named pool if a pool name is specified. The default is **all**.

**--seconds**

| Displays the I/O performance values for the previous number of seconds. The valid range of seconds is 1-999. The default value is 60 seconds. The values are displayed for subperiods within the period that you specify. The subperiods might be every 5 seconds over the last 60 seconds, or every 60 seconds over the last 600 seconds. You cannot configure the number or length of subperiods.

**--sum-classes**

| Displays the I/O performance for each QoS class separately if **no** is specified, or summed across all the QoS classes if **yes** is specified. The default is **no**.

**--sum-nodes**

| If **yes** is specified, displays the I/O performance summed across all the nodes in the cluster. If **no** is specified, displays the I/O performance for each node separately. The default is **yes**.

**Exit status**

**0** Successful completion.

**Nonzero**

A failure occurred.

**Security**

You must have root authority to run the **mmls qos** command.

The node on which you enter the command must be able to execute remote shell commands on any other administration node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

**Analyzing regular output from mmls qos**

The **mmls qos** command always shows the first three lines of output:

**QOS config::**

Indicates whether QoS is actively regulating I/O consumption (enabled) or is quiescent (disabled).

**QOS values::**

Displays, for each storage pool that you configured, the name of the storage pool and the IOPS that you assigned to the **other** class and the **maintenance** class. In the following example fragment, the command shows that the system storage pool is configured with the value of **inf** for both QoS classes:

```
QOS values:: pool=system,other=inf,maintenance/all_local=inf
```

The qualifier **/all\_local** after **maintenance** indicates that the maintenance IOPS are applied to all the file systems owned by the cluster. This value is the default for the **maintenance** class.

**QOS status::**

Indicates whether QoS is regulating the consumption of IOPS ("throttling") and also whether QoS is recording ("monitoring") the consumption of IOPS of each storage pool.

## mmlsqos

The following sample output is complete:

```
mmlsqos fs --seconds 30
QOS config:: enabled
QOS values:: pool=system,other=inf,maintenance/all_local=inf:pool=fpodata,other=inf,maintenance/all_local=inf
QOS status:: throttling active, monitoring active
== for pool fpodata
01:31:45 misc iops=11 ioql=0.016539 qsdl=1.2e-06 et=5
== for pool system
01:31:45 misc iops=8.2 ioql=0.013774 qsdl=2e-06 et=5
```

The command `mmlsqos fs0 --seconds 30` requests a display of I/O performance values for all QoS pools over the previous 30 seconds. Because the parameters `--sum_classes` and `--sum_nodes` are missing, the command also requests I/O performance for each storage pool separately and summed across all the nodes of the cluster.

The information that is displayed for the two configured pools, `fpodata` and `system`, indicates that IOPS occurred only for processes in the `misc` class. The meaning of the categories in each line is as follows:

### First column

The time when the measurement period ends.

### Second column

The QoS class for which the measurement is made.

**iops**= The performance of the class in I/O operations per second.

**ioql**= The average number of I/O requests in the class that are pending for reasons other than being queued by QoS. This number includes, for example, I/O requests that are waiting for network or storage device servicing.

**qsdl**= The average number of I/O requests in the class that are queued by QoS. When the QoS system receives an I/O request from the file system, QoS first finds the class to which the I/O request belongs. It then finds whether the class has any I/O operations available for consumption. If not, then QoS queues the request until more I/O operations become available for the class. The `qsdl` value is the average number of I/O requests that are held in this queue.

**et**= The interval in seconds during which the measurement was made.

You can calculate the average service time for an I/O operation as  $(\text{ioql} + \text{qsdl})/\text{iops}$ . For a system that is running IO-intensive applications, you can interpret the value  $(\text{ioql} + \text{qsdl})$  as the number of threads in the I/O-intensive applications. This interpretation assumes that each thread spends most of its time in waiting for an I/O operation to complete.

## Analyzing fine-grained output from mmlsqos

The following code block shows an example of fine-grained output. Each line displays the sum of the data for all the QoS programs that are running on the specified node. If the `--pid-stats` parameter is specified in the `mmchqos` command, then each line displays the data for one QoS program that is running on the specified node.

```
Time, Class, Node, Iops, TotSctrs, Pool, Pid, RW, SctrI, AvgTm, SsTm, MinTm, MaxTm, AvgQd, SsQd
1480606227,misc,172.20.0.21,285,2360,system,0,R,16,0.000260,0.000197,0.000006,0.005894,0.000001,0.000000
1480606228,misc,172.20.0.21,631,5048,system,0,R,16,0.000038,0.000034,0.000006,0.002463,0.000000,0.000000
1480606239,other,172.20.0.21,13,112,system,26724,R,16,0.000012,0.000000,0.000008,0.000022,0.000001,0.000000
1480606239,other,172.20.0.21,1,512,system,26724,R,512,0.000375,0.000000,0.000375,0.000022,0.000002,0.000000
1480606239,other,172.20.0.21,30,15360,system,26724,W,512,0.000910,0.000004,0.000451,0.002042,0.000001,0.000000
1480606241,misc,172.20.0.21,5,48,system,14680072,W,16,0.000135,0.000000,0.000025,0.000258,0.000000,0.000000
1480606240,other,172.20.0.21,4,32,system,26724,R,16,0.000017,0.000000,0.000009,0.000022,0.000001,0.000000
1480606240,other,172.20.0.21,48,24576,system,26724,W,512,0.000916,0.000006,0.000490,0.002141,0.000001,0.000000
1480606241,other,172.20.0.21,10,80,system,26724,R,16,0.000017,0.000000,0.000007,0.000039,0.000001,0.000000
1480606241,other,172.20.0.21,34,17408,system,26724,W,512,0.000873,0.000007,0.000312,0.001957,0.000001,0.000000
```

```
1480606241,misc,172.20.0.21,12,104,system,15597572,W,16,0.000042,0.000000,0.000024,0.000096,0.000000,0.000000
1480606241,misc,172.20.0.21,1,512,system,15597572,W,512,0.000192,0.000000,0.000192,0.000192,0.000000,0.000000
1480606241,misc,172.20.0.21,9,72,system,14680071,W,16,0.000029,0.000000,0.000024,0.000040,0.000000,0.000000
conti=4
```

The following list describes the content in each column:

**Time** The time when the measurement was made, expressed in seconds since the epoch.

**Class** The QoS class of the process.

**Node** The IP address of the node on which the process was running.

**Iops** The number of IOPS that the process consumed during the sample period.

#### TotSctrs

The number of sectors for which data was read or written. By default, one sector is 512 bytes.

**Note:** The number of **TotSctrs** from one process or one node might not be equal to the real bytes read or written from applications. For cache I/O in IBM Spectrum Scale, if the modified data from one application I/O is less than **fgdbRangeSize** (by default, it is 4KB), IBM Spectrum Scale writes bytes defined by **fgdbRangeSize** from pagpool to backend disks. Therefore, if the application updates only a byte of the file, it shows eight sectors from the **mmls qos** output when taking the default **fgdbRangeSize**.

**Pool** The memory pool for which the IOPS were used.

**Pid** The process ID of the process for which the IOPS were used. When the **--pid-stats** parameter of the **mmchqos** command is not specified, this value is always 0.

**RW** Whether the process was doing a read operation or a write operation.

**SctrI** The number of sectors for which the IOPS were done. One of the following categories is displayed. The exact value of a category depends on the disk block size.

A single sector.

1/32 or less of a full block.

Less than a full block.

A full block.

For example, if the disk block size is 512, the command displays one of the following values: 1, 16, 511, or 512.

#### AvgTm

The mean time that was required for an I/O operation to be completed.

**SsTm** The sum of the squares of differences from the mean value that is displayed for **AvgTm**. You can use this value to calculate the variance of I/O service times.

#### MinTm

The minimum time that was required for an I/O operation to be completed.

#### MaxTm

The maximum time that was required for an I/O operation to be completed.

#### AvgQd

The mean time for which QoS imposed a delay of the read or write operation.

**SsQd** The sum of the squares of differences from the mean value that is displayed for **AvgQd**.

The last line in the example indicates that the index of the next block to be displayed is 4. You can avoid re-displaying statistics by having the next **mmls qos** command display statistics beginning at this block index.

## **mmlsqos**

### **Examples**

1. The following command displays the I/O performance values for all the pools in the file system over the previous 60 seconds. It does so for each QoS class separately and summed across all the nodes in the cluster.

```
mmlsqos fs0 --seconds 60
```

2. The following command displays the I/O performance values for the named pool over the previous 60 seconds. It does so for each QoS class separately and for each node separately.

```
mmlsqos fs0 --pool pname0 --sum-nodes no
```

### **See also**

- “[mmchqos command](#)” on page 206
- *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **Location**

/usr/lpp/mmfs/bin

## mmIsquota command

Displays quota information for a user, group, or fileset.

### Synopsis

```
| mmIsquota [-u User | -g Group] [-v | -q] [-e] [-C ClusterName]
| [-Y] [--block-size {BlockSize | auto}] [Device[:Fileset] ...]
|
| or
|
| mmIsquota -j Fileset [-v | -q] [-e] [-C ClusterName]
| [-Y] [--block-size {BlockSize | auto}] Device ...
|
| or
|
| mmIsquota -d {[-u] [-g] [-j]} [-C ClusterName]
| [-Y] [--block-size {BlockSize | auto}] [Device ...]
|
| or
|
| mmIsquota -d [-C ClusterName] [-Y] [--block-size {BlockSize | auto}] [Device[:Fileset] ...]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

For the specified *User*, *Group*, or *Fileset* the **mmIsquota** command displays information about quota limits and current usage on each file system in the cluster. This information is displayed only if quota limits have been established and the user has consumed some amount of storage. If you want quota information for a *User*, *Group*, or *Fileset* that has no file system storage allocated at the present time, you must specify **-v**.

If neither the **-g**, **-u**, or **-j** option is specified, the default is to display only user quotas for the user who issues the command.

For each file system in the cluster, the **mmIsquota** command displays:

1. Block limits:
  - quota type (USR or GRP or FILESET)
  - current usage
  - soft limit
  - hard limit
  - space in doubt
  - grace period
2. File limits:
  - current number of files
  - soft limit
  - hard limit
  - files in doubt
  - grace period

**Note:** In cases where small files do not have an additional block allocated for them, quota usage may show less space usage than expected.

3. Remarks

## **mmlsquota**

Because the sum of the *in-doubt* value and the current usage may not exceed the hard limit, the actual block space and number of files available to the user, group, or fileset may be constrained by the *in-doubt* value. If the *in-doubt* value approaches a significant percentage of the quota, run the **mmcheckquota** command to account for the lost space and files.

For more information, see *Listing quotas in IBM Spectrum Scale: Administration Guide*.

This command cannot be run from a Windows node.

### **Parameters**

#### **-C ClusterName**

Specifies the name of the cluster from which the quota information is obtained (from the file systems within that cluster). If **-C** is omitted, the local cluster is assumed. The cluster name specified by the **-C** flag must be part of the same multicluster group as the node issuing the **mmlsquota** command. A node that is part of a remote cluster can only see the file systems that it has been given authority to mount from the local cluster.

#### *Device*

Specifies the device name of the file system for which quota information is to be displayed. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

#### *Fileset*

Specifies the name of a fileset located on *Device* for which quota information is to be displayed.

**-d** Displays the default quota limits for user, group, or fileset quotas. When specified in combination with the **-u**, **-g**, or **-j** options, default file system quotas are displayed. When specified without any of the **-u**, **-g**, or **-j** options, default fileset-level quotas are displayed.

**-e** Specifies that **mmlsquota** is to collect updated quota usage data from all nodes before displaying results. If **-e** is not specified, there is the potential to display negative usage values as the quota server may process a combination of up-to-date and back-level information.

#### **-g Group**

Displays quota information for the user group or group ID specified in the *Group* parameter.

#### **-j Fileset**

Displays quota information for the named fileset.

**-q** Prints a terse message containing information only about file systems with usage over quota.

#### **-u User**

Displays quota information for the user name or user ID specified in the *User* parameter.

**-v** Displays quota information on file systems where the *User*, *Group* or *Fileset* limit has been set, but the storage has not been allocated.

#### **--block-size {BlockSize | auto}**

Specifies the unit in which the number of blocks is displayed. The value must be of the form **[n]K**, **[n]M**, **[n]G** or **[n]T**, where *n* is an optional integer in the range 1 to 1023. The default is 1K. If **auto** is specified, the number of blocks is automatically scaled to an easy-to-read value.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

#### **Note:**

- Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmc1idecode**. Use the **mmc1idecode** command to decode the field.
- **-Y** disregards the value of **--block-size** and returns the results in KB.

## Exit status

**0** Successful completion.

### nonzero

A failure has occurred.

## Security

If you are a root user:

- You may view quota information for all users, groups, and filesets.
- The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

If you are a non-root user, you may view only fileset quota information, your own quota information, and quota information for any groups to which you belong.

You must be a root user to use the **-d** option.

GPFS must be running on the node from which the **mmlsquota** command is issued.

## Examples

1. Userid **paul** issued this command:

```
mmlsquota
```

The system displays information similar to:

| Block Limits |      |          |        | File Limits |       |             |       |          |       |    |       |
|--------------|------|----------|--------|-------------|-------|-------------|-------|----------|-------|----|-------|
| Filesystem   | type | KB quota | limit  | in_doubt    | grace | files quota | limit | in_doubt | grace |    |       |
| fsn          | USR  | 728      | 100096 | 200192      | 4880  | none        | 35    | 30       | 50    | 10 | 6days |

This output shows the quotas for user **paul** in file system **fsn** set to a soft limit of 100096 KB, and a hard limit of 200192 KB. 728 KB is currently allocated to **paul**. 4880 KB is also in doubt, meaning that the quota system has not yet been updated as to whether this space has been used by the nodes, or whether it is still available. No grace period appears because the user has not exceeded his quota. If the user had exceeded the soft limit, the grace period would be set and the user would have that amount of time to bring his usage below the quota values. If the user failed to do so, the user would not be allocated any more space.

The soft limit for files (inodes) is set at 30 and the hard limit is 50. 35 files are currently allocated to this user, and the quota system does not yet know whether the 10 in doubt have been used or are still available. A grace period of six days appears because the user has exceeded his quota. The user would have this amount of time to bring his usage below the quota values. If the user fails to do so, the user is not allocated any more space.

2. To show the quotas for user **pfs001**, device **gpfs2**, and fileset **fset4**, issue this command:

```
mmlsquota -u pfs001 gpfs2:fset4
```

The system displays information similar to:

| Block Limits |         |      |          | File Limits |        |       |      |             |       |       |       |         |
|--------------|---------|------|----------|-------------|--------|-------|------|-------------|-------|-------|-------|---------|
| Filesystem   | Fileset | type | KB quota | limit       | doubt  | grace | in   | files quota | limit | doubt | grace | Remarks |
| gpfs2        | fset4   | USR  | 4104     | 10240       | 153600 | 0     | none | 1           | 1000  | 5000  | 0     | none    |

3. To show user and group default quotas for all filesets in the **gpfs1** file system, issue this command:

```
mmlsquota -d gpfs1
```

The system displays information similar to:

## mmlsquota

| Default Block Limits(KB) |         |      |       | Default File Limits |       |       |             |
|--------------------------|---------|------|-------|---------------------|-------|-------|-------------|
| Filesystem               | Fileset | type | quota | limit               | quota | limit | entryType   |
| gpfs1                    | root    | USR  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | root    | GRP  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | fset1   | USR  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | fset1   | GRP  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | fset2   | USR  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | fset2   | GRP  | 0     | 0                   | 0     | 0     | default off |

4. To show user and group default quotas for fileset **fset1** in the **gpfs1** file system, issue this command:

```
mmlsquota -d gpfs1:fset1
```

The system displays information similar to:

| Default Block Limits(KB) |         |      |       | Default File Limits |       |       |             |
|--------------------------|---------|------|-------|---------------------|-------|-------|-------------|
| Filesystem               | Fileset | type | quota | limit               | quota | limit | entryType   |
| gpfs1                    | fset1   | USR  | 0     | 0                   | 0     | 0     | default off |
| gpfs1                    | fset1   | GRP  | 0     | 0                   | 0     | 0     | default off |

5. To show the quotas for fileset **fset0** in file system **fs1**, issue this command:

```
mmlsquota -j fset0 fs1 --block-size auto
```

The system displays information similar to:

| Filesystem | type    | Block Limits |       |       |          | File Limits |       |       |          | Remarks  |
|------------|---------|--------------|-------|-------|----------|-------------|-------|-------|----------|----------|
|            |         | blocks       | quota | limit | in_doubt | files       | quota | limit | in_doubt |          |
| fs1        | FILESET | 89.25G       | 100G  | 200G  | 89.99M   | none        | 13729 | 4000  | 5000     | 0 7 days |

## See also

- “mmcheckquota command” on page 169
- “mmdefedquota command” on page 272
- “mmdefquotaoff command” on page 275
- “mmdefquotaoon command” on page 278
- “mmedquota command” on page 325
- “mmrepquota command” on page 517
- “mmquotaon command” on page 509
- “mmquotaoff command” on page 507

## Location

/usr/lpp/mmfs/bin

---

## mmlssnapshot command

Displays GPFS snapshot information.

### Synopsis

```
mmlssnapshot Device [-d [--block-size {BlockSize | auto}]]
[-s {all | global | [[Fileset]:]Snapshot[,[[Fileset]:]Snapshot...]} | -j Fileset[,Fileset...]]
[--qos QOSClass] [-Y]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmlssnapshot** command to display GPFS snapshot information for the specified file system or fileset. You can optionally display the amount of storage that is used by the snapshot.

### Parameters

#### *Device*

The device name of the file system for which snapshot information is to be displayed. File system names do not need to be fully qualified. *fs0* is as acceptable as */dev/fs0*.

**-d** Displays the amount of storage that is used by the snapshot.

This operation requires an amount of time that is proportional to the size of the file system; therefore, it can take several minutes or even hours on a large and heavily-loaded file system.

This optional parameter can impact overall system performance. Avoid running the **mmlssnapshot** command with this parameter frequently or during periods of high file system activity.

#### **--block-size {BlockSize | auto}**

Specifies the unit in which the number of blocks is displayed. The value must be of the form *[n]K*, *[n]M*, *[n]G* or *[n]T*, where *n* is an optional integer in the range 1 - 1023. The default is 1 K. If **auto** is specified, the number of blocks is automatically scaled to an easy-to-read value.

**-s** Displays the attributes for the specified snapshots.

#### **all**

Displays information for all snapshots. This option is the default.

#### **global**

Displays information for global snapshots.

#### **[[Fileset]:]**

:

A colon (:) followed by a snapshot name indicates a global snapshot. For example, :SS01 indicates a global snapshot with the name SS01. If a global snapshot with that name exists, the command displays information about it.

#### **Fileset:**

A fileset name followed by a colon (:) followed by a snapshot name indicates a fileset snapshot. For example, fset02:SS01 indicates a snapshot of fileset fset02 with the name SS01. If a snapshot of the fileset with that snapshot name exists, then the command displays information about it.

#### **Snapshot[,Snapshot...]**

Displays information for the specified snapshots.

#### **-j Fileset[,Fileset...]**

Displays only snapshots that contain the specified filesets; including all global snapshots.

## **mmlssnapshot**

### **--qos QoSClass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “[mmchqos command](#)” on page 206. Specify one of the following QoS classes:

#### **maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

**-Y** Displays the command output in a parseable format with a colon (:) as the field delimiter. Each column is described by a header.

**Note:** Fields having a colon are encoded to prevent confusion. If a field contains a % (percent sign) character, it is most likely encoded. Use the **mmc1idecode** command to decode the field.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure occurred.

## **Security**

You must be a root user or fileset owner to use the **-d** parameter.

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

If you are a non-root user, you can specify only file systems that belong to the same cluster as the node on which the **mmlssnapshot** command was issued.

## **Examples**

**Note:** Ensure that the snapshot name does not include a colon (:).

The following command displays information about all the existing snapshots in file system fs1:  
**mmlssnapshot fs1**

The following command displays information about snapshots named SS01. The snapshots can be global snapshots or fileset snapshots:

**mmlssnapshot fs1 -s SS01**

The following command displays information about a global snapshot named gSS01:

**mmlssnapshot fs1 -s :gSS01**

The following command displays information about a fileset snapshot with the name `fsSS01` that is a snapshot of fileset `fset02`:

```
mm1ssnapshot fs1 -s fset02:fsSS01
```

The following command displays information about global snapshots with the names `gSS01` and `gSS02` and a fileset snapshot of fileset `fset02` named `fsSS02`:

```
mm1ssnapshot fs1 -s :gSS01,:gSS02,fset02:fsSS02
```

The following command displays information about global snapshots and fileset snapshots that contain the filesets `fset02` and `fset03`:

```
mm1ssnapshot fs1 -j fset02,fset03
```

## See also

- “`mmcsnapshot` command” on page 267
- “`mmdelsnapshot` command” on page 304
- “`mmrestorefs` command” on page 525
- “`mmsnapdir` command” on page 565

## Location

`/usr/lpp/mmfs/bin`

## **mmmmigratefs**

---

### **mmmmigratefs command**

Performs needed conversions to support new file system features.

#### **Synopsis**

```
mmmmigratefs Device [--fastea] [--online | --offline]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmmmigratefs** command to enable features that require existing on-disk data structures to be converted to a new format.

Before issuing the **mmmmigratefs** command, see the topic about migration, coexistence, and compatibility in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*. You must ensure that all nodes in the cluster have been migrated to the latest level of GPFS code and that you have successfully run the **mmchconfig release=LATEST** command. You must also ensure that the new features have been enabled by running **mmchfs -V full**.

The **mmmmigratefs** command can be run with the file system mounted or unmounted. If **mmmmigratefs** is run without the **--online** or **--offline** parameters specified, the command will determine the mount status of the file system and run in the appropriate mode.

#### **Parameters**

##### **Device**

The device name of the file system to be migrated. File system names need not be fully qualified; for example, **fs0** is just as acceptable as **/dev/fs0**. This must be the first parameter.

##### **--fastea**

Convert the existing extended attributes to the new format required for storing the attributes in the file's inode and thereby allowing for faster extended-attribute access.

##### **--online**

Allows the **mmmmigratefs** command to run while the file system is mounted.

##### **--offline**

Allows the **mmmmigratefs** command to run while the file system is unmounted.

#### **Exit status**

**0**      Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmmmigratefs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

To enable fast extended attribute access for file system fs3, issue this command:

```
mmmmigratefs fs3 --fastea
```

The system displays information similar to the following:

```
Enabling fastea support
 11.19 % complete on Thu Nov 14 13:50:24 2013 (167936 inodes 328 MB)
 33.21 % complete on Thu Nov 14 13:51:56 2013 (498260 inodes 973 MB)
100.00 % complete on Thu Nov 14 13:52:04 2013
Finalizing upgrade
 11.19 % complete on Thu Nov 14 13:52:27 2013 (167936 inodes 328 MB)
 26.78 % complete on Thu Nov 14 13:53:07 2013 (401834 inodes 785 MB)
100.00 % complete on Thu Nov 14 13:53:27 2013
Feature 'fastea' is now enabled on "fs3".
```

## See also

- “mmchconfig command” on page 132
- “mmchfs command” on page 179

## Location

/usr/lpp/mmfs/bin

## **mmmount**

---

### **mmmount command**

Mounts GPFS file systems on one or more nodes in the cluster.

#### **Synopsis**

```
mmmount {Device | DefaultMountPoint | DefaultDriveLetter |
 a11 | a11_local | a11_remote | {-F DeviceFileName}}
 [-o MountOptions] [-a | -N {Node[,Node...] | NodeFile | NodeClass}]
```

or

```
mmmount Device {MountPoint | DriveLetter}
 [-o MountOptions] [-a | -N {Node[,Node...] | NodeFile | NodeClass}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmmount** command mounts the specified GPFS file system on one or more nodes in the cluster. If no nodes are specified, the file systems are mounted only on the node from which the command was issued. A file system can be specified using its device name or its default mount point, as established by the **mmcrfs**, **mmchfs** or **mmremoteefs** commands.

When **all** is specified in place of a file system name, all GPFS file systems will be mounted. This also includes remote GPFS file systems to which this cluster has access.

#### **Parameters**

*Device | DefaultMountPoint | DefaultDriveLetter | a11 | a11\_local | a11\_remote | {-F DeviceFileName}*

Indicates the file system or file systems to be mounted.

*Device*

The device name of the file system to be mounted. File system names need not be fully-qualified.  
**fs0** is just as acceptable as **/dev/fs0**.

*DefaultMountPoint*

The mount point associated with the file system as a result of the **mmcrfs**, **mmchfs**, or **mmremoteefs** commands.

*DefaultDriveLetter*

The Windows drive letter associated with the file system as a result of the **mmcrfs** or **mmchfs** command.

**a11**

Indicates all file systems known to this cluster.

**a11\_local**

Indicates all file systems owned by this cluster.

**a11\_remote**

Indicates all file systems owned by another cluster to which this cluster has access.

**-F DeviceFileName**

Specifies a file containing the device names, one per line, of the file systems to be mounted.

This must be the first parameter.

*DriveLetter*

The location where the file system is to be mounted. If not specified, the file system is mounted at its

default drive letter. This option can be used to mount a file system at a drive letter other than its default one or to mount a file system that does not have an established default drive letter.

#### **MountPoint**

The location where the file system is to be mounted. If not specified, the file system is mounted at its default mount point. This option can be used to mount a file system at a mount point other than its default mount point.

## **Options**

**-a** Mount the file system on all nodes in the GPFS cluster.

**-N {Node [,Node...] | NodeFile | NodeClass}**

Specifies the nodes on which the file system is to be mounted.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

**-o MountOptions**

Specifies the mount options to pass to the mount command when mounting the file system. For a detailed description of the available mount options, see *Mount options specific to IBM Spectrum Scale* in *IBM Spectrum Scale: Administration Guide*.

## **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmmount** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To mount all GPFS file systems on all of the nodes in the cluster, issue this command:

```
mmmount all -a
```

2. To mount file system **fs2** read-only on the local node, issue this command:

```
mmmount fs2 -o ro
```

3. To mount file system **fs1** on all NSD server nodes, issue this command:

```
mmmount fs1 -N nsdsvnodes
```

## **See also**

- “[mmumount command](#)” on page 575
- “[mmlsmount command](#)” on page 419

## **Location**

/usr/lpp/mmfs/bin

## **mmnetverify**

### **mmnetverify command**

Verifies network configuration and operation in a cluster.

#### **Synopsis**

```
mmnetverify [Operation [Operation...]] [-N {Node[,Node...] | all}]
[--target-nodes {Node[,Node...] | all}]
[--configuration-file File] [--log-file File]
[--verbose | -Y] [--min-bandwidth Number]
[--max-threads Number] [--ces-override]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

- | **Note:** The **mmnetverify** command is a diagnostic tool and is not intended to be issued continually during normal IBM Spectrum Scale cluster operations. Issue this command only to validate a cluster network before going into production or when you suspect network problems.

With the **mmnetverify** command, you can verify the network configuration and operation of a group of nodes before you organize them into an IBM Spectrum Scale cluster. You can also run the command to analyze network problems after you create a cluster.

If you have not created an IBM Spectrum Scale cluster yet, you must run the command with a configuration file. See the **--configuration-file** option in the **Parameters** section.

The command uses the concepts of local nodes and target nodes. A *local node* is a node from which a network test is run. The command can be started from one node and run from several separate local nodes. A *target node* is a node against which a test is run.

The command has the following requirements:

- IBM Spectrum Scale must be installed on all the nodes that are involved in the test, including both local nodes and target nodes.
- Each node must be able to issue remote shell commands to all nodes, including itself, without a password. (This requirement is tested in the **shell** check.)

The following table lists the types of output messages and where they are sent. Messages are not added to a log file unless you specify one on the command line:

*Table 16. Information and error messages*

| Message type                 | Printed to console                                                        | Added to the log file, if one is specified. |
|------------------------------|---------------------------------------------------------------------------|---------------------------------------------|
| Information messages         | Yes (stdout)                                                              | Yes                                         |
| Verbose information messages | Yes (stdout)<br><br>If <b>--verbose</b> is specified on the command line. | Yes                                         |
| Error messages               | Yes (stderr)                                                              | Yes                                         |

If sudo wrappers are enabled, the command uses a sudo wrapper when it communicates with remote nodes. Note the following restrictions:

- If you are working with an existing cluster, the cluster must be at IBM Spectrum Scale v4.2.3 or later.

- You must run the command on an administration node.
- The **-N** option is not supported.
- You must run the sudo command as the gpfadmin user.

## Parameters

### *Operation* [ **Operation**... ]

Specifies one or more operations, which are separated by blanks, that are to be verified against the target nodes. The operations are described in Table 18 on page 449. Shortcut terms are described in Table 17 on page 449.

- | If you do not specify any operations, the command does all the operations except **data-large**, **flood-node**, and **flood-cluster** against the target nodes.

### **-N** {*Node*[,*Node*...] | **a11**}

Specifies a list of nodes on which to run the command. If you specify more than one node, the command is run on all the specified nodes in parallel. Each node tests all the specified target nodes. If you do not include this parameter, the command runs the operations only from the node where you enter the command.

**Note:** The **--max-threads** parameter specifies how many nodes can run in parallel at the same time. If the limit is exceeded, the command still tests all the specified nodes, starting the surplus nodes as other nodes finish.

### *Node*[,*Node*...]

Specifies a list of nodes in the local cluster. This parameter accepts node classes. You can specify system node classes, such as **aixnodes**, and node classes that you define with the **mmcrnodeclass** command.

### **a11**

Specifies all the nodes in the local cluster.

### **--target-nodes** {*Node*[,*Node*...] | **a11**}

Specifies a list of nodes that are the targets of the testing. If you do not specify this parameter, the command runs the operations against all the nodes in the cluster.

### *Node*[,*Node*...]

Specifies a list of nodes in the local cluster. This parameter accepts node classes. You can specify system node classes, such as **aixnodes**, and node classes that you define with the **mmcrnodeclass** command.

### **a11**

Specifies all the nodes in the local cluster.

### **--configuration-file** *File*

Specifies the path of a configuration file. You must use a configuration file if you have not created an IBM Spectrum Scale cluster yet. You can also specify a configuration file if you have created a cluster but you do not want the command to run with the IBM Spectrum Scale cluster configuration values. Only the **node** parameter is required. The other parameters revert to their default values if they are not specified.

**Note:** When you specify this parameter, the **--ces-override** parameter is automatically set.

The following code block shows the format of the file:

```
node Node [AdminName]
rshPath Path
rcpPath Path
tscTcpPort Port
mmsdrservPort Port
tscCmdPortRange Min-Max
```

## **mmnetverify**

where:

### **node Node [AdminName]**

Specifies a node name, followed optionally by the node's admin name. If you do not specify an admin name, the command uses the node name as the admin name.

You can have multiple **node** parameters. Add a **node** parameter for each node that you want to be included in the testing, either as a local node or as a target node. You must include the node from which you are running the command.

### **rshPath Path**

Optional. Specifies the path of the remote shell command to be used. The default value is /usr/bin/ssh. Specify this parameter only if you want to use a different remote shell command.

### **rcpPath Path**

Optional. Specifies the path of the remote file copy command to be used. The default value is /usr/bin/scp. Specify this parameter only if you want to use a different remote copy command.

### **tscTcpPort Port**

Optional. Specifies the TCP port number to be used by the local GPFS daemon when it contacts a remote cluster. The default value is 1191. Specify this value only if you want to use a different port.

### **mmsdrservPort Port**

Optional. Specifies the TCP port number to be used by the **mmsdrserv** service to provide access to configuration data to the rest of the nodes in the cluster. The default value is the value that is stored in **mmfsdPort**.

### **tscCmdPortRange=Min-Max**

Specifies the range of port numbers to be used for extra TCP/IP ports that some administration commands need for their processing. Defining a port range makes it easier for you set firewall rules that allow incoming traffic on only those ports. For more information, see the topic *IBM Spectrum Scale port usage* in the *IBM Spectrum Scale: Administration Guide*.

If you used the **spectrumscale** installation toolkit to install a version of IBM Spectrum Scale that is earlier than version 5.0.0, then this attribute is initialized to 60000-61000. Otherwise, this attribute is initially undefined and the port numbers are dynamically assigned from the range of ephemeral ports that are provided by the operating system.

### **[--log-file File]**

Specifies the path of a file to contain the output messages from the network checks. If you do not specify this parameter, messages are displayed only on the console. See Table 16 on page 446.

### **--verbose**

Causes the command to generate verbose output messages. See Table 16 on page 446.

### **-Y**

Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

### **[--min-bandwidth Number]**

Specifies the minimum acceptable bandwidth for the data bandwidth check.

### **--max-threads Number**

Specifies how many nodes can run the command in parallel at the same time. The valid range is 1 - 64. The default value is 32. For more information, see the **-N** parameter.

### **--ces-override**

Causes the command to consider all the nodes in the configuration to be CES-enabled. This parameter overrides the requirement of the **protocol-ctdb** and the **protocol-object** network checks.

that the local node and the target nodes must be CES-enabled with the **mmchnode** command. This parameter is automatically set when you specify the **--configuration-file** parameter.

## The network checks

The following table lists the shortcut terms that you can specify for the network checks that are listed in Table 18:

**Note:** These network checks might cause entries to the mmfs log that say that connections are being ended. These entries are expected and do not indicate problems in your system.

*Table 17. Shortcut terms for network checks*

| Shortcut            | Checks that are performed                                        |
|---------------------|------------------------------------------------------------------|
| <b>local</b>        | <b>interface</b>                                                 |
| <b>connectivity</b> | <b>resolution</b> , <b>ping</b> , <b>shell</b> , and <b>copy</b> |
| <b>port</b>         | <b>daemon-port</b> , <b>sdserv-port</b> , and <b>tsccmd-port</b> |
| <b>data</b>         | <b>data-small</b> , <b>data-medium</b> , and <b>data-large</b>   |
| <b>bandwidth</b>    | <b>bandwidth-node</b> and <b>bandwidth-cluster</b>               |
| <b>protocol</b>     | <b>protocol-ctdb</b> and <b>protocol-object</b>                  |
| <b>flood</b>        | <b>flood-node</b> and <b>flood-cluster</b>                       |
| <b>all</b>          | All checks except <b>flood-node</b> and <b>flood-cluster</b>     |

The following table lists the parameters that you can specify for network checks. Separate these parameters with a blank on the command line. For example, the following command runs the **interface** and **copy** checks on the local node against all the nodes in the cluster:

```
mmnetverify interface copy
```

*Table 18. Network checks*

| Command-line option | Test description                | Test items                                                                                                                                                                                                                                                                                                           |
|---------------------|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>interface</b>    | Network interface configuration | The local node's daemon and admin network addresses are enabled.                                                                                                                                                                                                                                                     |
| <b>resolution</b>   | Host name resolution            | The command verifies the following conditions: <ul style="list-style-type: none"> <li>The target node's name and daemon node name can be resolved on the local node and they resolve to the same address.</li> <li>The target node's IP address resolves to either the node name or the daemon node name.</li> </ul> |
| <b>ping</b>         | Network connectivity via ping   | The local node can ping the target node with its name, daemon node name, rel_hostname, admin_shortname, and IP address entries in the mmsdrfs file.                                                                                                                                                                  |
| <b>shell</b>        | Remote shell command            | The command verifies the following conditions: <ul style="list-style-type: none"> <li>The local node can issue remote shell commands to the target node's admin interface without requiring a password.</li> <li>The target node's daemon and admin names refer to the same node.</li> </ul>                         |
| <b>copy</b>         | Remote copy                     | The local node can issue a remote copy command to the target node's admin interface without requiring a password.                                                                                                                                                                                                    |
| <b>time</b>         | Date and time                   | The time and date on the local node and target node do not differ by a wide margin.                                                                                                                                                                                                                                  |

## mmnetverify

*Table 18. Network checks (continued)*

| Command-line option | Test description              | Test items                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|---------------------|-------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>daemon-port</b>  | GPFS daemon connectivity      | The target node can establish a TCP connection to the local node on the <b>mmfsd</b> daemon port of the local node: <ul style="list-style-type: none"> <li>The target node uses the port that is specified in the cluster configuration property <b>tscTcpPort</b>. The default value is 1191.</li> <li>If <b>mmfsd</b> and <b>mmsdrserv</b> are not running on the local node, the command starts an echo server on the daemon port of the local node for this test.</li> </ul>                                     |
| <b>sdrserv-port</b> | GetObject daemon connectivity | The target node can establish a TCP connection directed to the local node on the port that is specified in the cluster configuration property <b>mmsdrservPort</b> : <ul style="list-style-type: none"> <li>The default value of this port is the value that is specified in the cluster configuration property <b>tscTcpPort</b>.</li> <li>If <b>mmfsd</b> and <b>mmsdrserv</b> are not running on the local node, the command starts an echo server on the daemon port of the local node for this test.</li> </ul> |
| <b>tsccmd-port</b>  | TS-command connectivity       | The target node can establish a TCP connection directed to the local node on a port in the range that is specified in the <b>tscCmdPortRange</b> property: <ul style="list-style-type: none"> <li>The command starts an echo server on the local node for this test.</li> <li>If the <b>tscCmdPortRange</b> property is set, then the echo server listens on a port in the specified range.</li> <li>If not, then the echo server listens on an ephemeral port that is provided by the operating system.</li> </ul>  |
| <b>data-small</b>   | Small data exchange           | The target node can establish a TCP connection to the local node and exchange a series of small-sized data messages without network errors: <ul style="list-style-type: none"> <li>The command starts an echo server on the local node for this test.</li> <li>If the <b>tscCmdPortRange</b> property is set, then the echo server listens on a port in the specified range.</li> <li>If not, then the echo server listens on an ephemeral port that is provided by the operating system.</li> </ul>                 |
| <b>data-medium</b>  | Medium data exchange          | The target node can establish a TCP connection to the local node and exchange a series of medium-sized data messages without network errors: <ul style="list-style-type: none"> <li>The command starts an echo server on the local node for this test.</li> <li>If the <b>tscCmdPortRange</b> property is set, then the echo server listens on a port in the specified range.</li> <li>If not, then the echo server listens on an ephemeral port that is provided by the operating system.</li> </ul>                |
| <b>data-large</b>   | Large data exchange           | The target node can establish a TCP connection to the local node and exchange a series of large-sized data messages without network errors: <ul style="list-style-type: none"> <li>The command starts an echo server on the local node for this test.</li> <li>If the <b>tscCmdPortRange</b> property is set, then the echo server listens on a port in the specified range.</li> <li>If not, then the echo server listens on an ephemeral port that is provided by the operating system.</li> </ul>                 |

Table 18. Network checks (continued)

| Command-line option      | Test description              | Test items                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|--------------------------|-------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>bandwidth-node</b>    | Network bandwidth one-to-one  | The target node can establish a TCP connection to the local node and send a large amount of data with adequate bandwidth: <ul style="list-style-type: none"><li>• Bandwidth is measured on the target node.</li><li>• If the <b>min-bandwidth</b> parameter was specified on the command line, the command verifies that the actual bandwidth exceeds the specified minimum bandwidth.</li></ul>                                                                                                                                                                                                                                                                                                                                                                          |
| <b>bandwidth-cluster</b> | Network bandwidth many-to-one | All target nodes can establish a TCP connection to the local node and send a large amount of data in parallel: <ul style="list-style-type: none"><li>• The bandwidth is measured on each target node.</li><li>• If the <b>min-bandwidth</b> parameter was specified on the command line, the command verifies that the actual bandwidth exceeds the specified minimum bandwidth.</li><li>• None of the target nodes has a significantly smaller bandwidth than the other target nodes.</li></ul>                                                                                                                                                                                                                                                                          |
| <b>gnr-bandwidth</b>     | Overall bandwidth             | All target nodes can establish a TCP connection to the local node and send data to it. <ul style="list-style-type: none"><li>• The bandwidth is measured on each target node.</li><li>• If the <b>min-bandwidth</b> parameter is specified on the command line, the command verifies that the actual bandwidth exceeds the specified minimum bandwidth.</li></ul> Note the following differences between this test and the <b>bandwidth-cluster</b> test: <ul style="list-style-type: none"><li>• The total bandwidth is measured rather than the bandwidth per node.</li><li>• All target nodes must be active and must participate in the test.</li><li>• The bandwidth value that is reported does not include ramp-up time for TCP to read full capability.</li></ul> |
| <b>flood-node</b>        | Flood one-to-one              | When the local node is flooded with datagrams, the target node can successfully send datagrams to the local node: <ul style="list-style-type: none"><li>• The target node tries to flood the local node with datagrams.</li><li>• The command records packet loss.</li><li>• The command verifies that some of the datagrams were received.</li></ul>                                                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>flood-cluster</b>     | Flood many-to-one             | When the local node is flooded with datagrams from all the target nodes in parallel, each target node can successfully send datagrams to the local node: <ul style="list-style-type: none"><li>• The command records packet loss for each target node.</li><li>• The command checks that each target node received some datagrams.</li><li>• The command checks that none of the target nodes has a packet loss significantly higher than the other nodes.</li></ul>                                                                                                                                                                                                                                                                                                      |

## **mmnetverify**

*Table 18. Network checks (continued)*

| Command-line option    | Test description             | Test items                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|------------------------|------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>protocol-ctdb</b>   | CTDB port connectivity       | The target node can establish a connection with the local node through the CTDB port. SMB uses CTDB for internode communications. The command does not run this test in the following situations: <ul style="list-style-type: none"><li>• The SMB protocol service is not enabled.</li><li>• The local node or a target node is not a CES-enabled node. A node is considered to be CES-enabled if it is in an existing cluster and if it was enabled with the <b>mmchnode --ces-enable</b> command. This requirement is not enforced in the following situations:<ul style="list-style-type: none"><li>– You override the requirement by specifying the <b>--ces-override</b> option.</li><li>– You specify the <b>--configuration-file</b> parameter.</li></ul></li></ul>                                                                                  |
| <b>protocol-object</b> | Object-protocol connectivity | The target node can establish a connection with the local node through the ports that are used by the object server, the container server, the account server, the object-server-sof, and the keystone daemons. The command does not run this test in the following situations: <ul style="list-style-type: none"><li>• The Object protocol service is not enabled.</li><li>• The local node or a target node is not a CES-enabled node. A node is considered to be CES-enabled if it is in an existing cluster and if it was enabled with the <b>mmchnode --ces-enable</b> command. This requirement is not enforced in the following situations:<ul style="list-style-type: none"><li>– You override the requirement by specifying the <b>--ces-override</b> option.</li><li>– You specify the <b>--configuration-file</b> parameter.</li></ul></li></ul> |

### **Exit status**

- 0**      The command completed successfully and all the tests were completed successfully.
- 1**      The command encountered problems with options or with running tests.
- 2**      The command completed successfully, but one or more tests was unsuccessful.

### **Security**

You must have root authority to run the **mmnetverify** command.

The node on which you enter the command must be able to execute remote shell commands on any other administration node in the cluster. It must be able to do so without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. The following command runs all checks, except **data-large**, **flood-node**, and **flood-cluster**, from the node where you enter the command against all the nodes in the cluster:  

```
mmnetverify
```
2. The following command runs connectivity checks from the node where you enter the command against all the nodes in the cluster:  

```
mmnetverify connectivity
```

3. The following command runs connectivity checks from nodes c49f04n11 and c49f04n12 against all the nodes in the cluster:

```
mmnetverify connectivity -N c49f04n11,c49f04n12
```

4. The following command runs all checks, except **flood-node** and **flood-cluster**, from the node where you enter the command against nodes c49f04n11 and c49f04n12:

```
mmnetverify all --target-nodes c49f04n11,c49f04n12
```

5. The following command runs network port checks on nodes c49f04n07 and c49f04n08, each checking against nodes c49f04n11 and c49f04n12:

```
mmnetverify port -N c49f04n07,c49f04n08 --target-nodes c49f04n11,c49f04n12
```

## See also

- “mmlsqos command” on page 430

## Location

/usr/lpp/mmfs/bin

## **mnnfs**

---

### **mnnfs command**

Manages NFS exports and configuration.

#### **Synopsis**

```
mnnfs export add Path [--client ClientOptions]
```

or

```
mnnfs export remove Path [--force]
```

or

```
mnnfs export change Path [--nfsadd ClientOptions]
 [--nfsremove {Client[,Client...]}]
 [--nfschange ClientOptions]
 [--nfsposition {PositionIndex | Client}]
```

or

```
mnnfs export list [-Y] [--nfsdefs Path]
 [--nfsdefs-match Path]
```

or

```
mnnfs export load ExportCFGFile
```

or

```
mnnfs config list [-Y] [--exportdefs]
```

or

```
| mnnfs configuration change "Option=Value:Option=Value1,Value2:Option=Value..."
```

### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

### **Description**

Use the **mnnfs export** commands to add, change, list, load, or remove NFS export declarations for IP addresses on nodes that are configured as CES types.

Use the **mnnfs config** commands to list and change NFS configuration.

The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms.

### **Parameters**

#### **export**

Manages the NFS export configuration for the cluster with one of the following actions:

##### **add**

Creates a new configuration file for the NFS server in case it does not yet exist. If there is already an export configuration file, then it is extended with the provided additional export parameters. This export configuration file is used by the NFS server to create an NFS export for the *Path* so that clients can connect to it. If there is already an existing export for the *Path* then an error is shown. Each export configuration set has internally its own unique identifier number. This number is automatically incremented for each added export. The **mnnfs export add** command

attempts to add the new export also to running NFS server instances, and may fail if one or more instances are not running. This is not a critical issue, since the configuration changes have been made in the repository and will be applied later when restarting the NFS server instance.

The authentication method must be established before an NFS export can be defined.

The export *Path* must be an existing path in the GPFS file system, which is intended to be exported or is already exported to external clients using the NFS protocol.

**Note:** The paths which are not within the GPFS file system cannot be exported using the commands.

Creating nested exports (such as `/path/to/folder` and `/path/to/folder/subfolder`) is strongly discouraged since this might lead to serious issues in data consistency. Be very cautious when creating and using nested exports.

If there is a need to have nested exports (such as `/path/to/folder` and `/path/to/folder/inside/subfolder`), NFSv4 client that mounts the parent (`/path/to/folder`) export will not be able to see the child export subtree (`/path/to/folder/inside/subfolder`) unless the same client is explicitly allowed to access the child export as well. This is okay as long as the client uses only NFSv4 mounts.

Some export configuration commands might allow multiple client declarations, and therefore they have separators to distinguish them.

The following separators can be used:

- **Colon** to separate multiple allowed values for a specified attribute. For example, the key/value pair "Protocols=3:4" allows the NFS protocols v3 and v4 to be declared for an export.
- **Comma** to separate key/value pairs within a client declaration list and a **Semicolon** to separate client declaration lists.

For example:

```
--client "192.0.2.0/20(Access_Type=RW,Protocols=3:4); \
198.51.100.0/20(Access_Type=R0,Protocols=3:4,Transports=TCP:UDP)"
```

**Note:** To take advantage of the GPFS independent fileset features such as quotas, snapshots, and data management, the export paths can be made from GPFS filesets (either dependent or independent).

#### **--client** *ClientOptions*

Declares the client specific settings. *ClientOptions* can be a list of one or more client definitions. To avoid incorrect parsing by the interpreter, quote the argument list. For a list of client definitions that can be specified with the **--client** option, see List of supported client options for the mmnfs export {add | change} command.

#### **remove**

Removes the requested export from the configuration file and also from running NFS server instances, and might fail if one or more instances are not running.

**Note:** This command does not remove data.

The **--force** option is not used.

#### **change**

Modifies an existing export configuration for the export that is specified by *Path*, if the export exists. If the export does not exist, then an error is shown.

**Note:** Only client-related attributes are modified by this command, but not the basic export settings such as path.

## **mmnfs**

The **--nfsposition** flag can be only used together with either **--nfsadd** or **--nfschange**. It cannot be used standalone or together with **--nfsremove**. When **--nfsadd** and **--nfsremove** are given on the same command line, then the remove procedure is executed first internally.

**Note:** The **mmnfs export change** will restart NFS services on all CES nodes on which the NFS server is currently running.

### **--nfsadd** *ClientOptions*

Adds a new client declaration for the specified *Path*. *ClientOptions* can be a list of one or more client definitions. To avoid incorrect parsing by the interpreter, quote the argument list. For a list of client definitions that can be specified with the **--client** option, see List of supported client options for the mmnfs export {add | change} command.

### **--nfsremove** {*Client*[,*Client*...]}

Removes the NFS client specified by *Client* from the export configuration for the *Path*. *Client* can be a single client specifier, or a comma-separated list of client specifiers. If a client is the only one within a client declaration section in the configuration file, then this client section is removed.

**Note:** When the last client within a client declaration is removed, then the export section is also removed.

### **--nfschange** *ClientOptions*

Modifies an existing client declaration for the specified *Path*. *ClientOptions* can be a list of one or more client definitions. To avoid incorrect parsing by the interpreter, quote the argument list. For a list of client definitions that can be specified with the **--client** option, see List of supported client options for the mmnfs export {add | change} command. If a client specifier declared in *ClientOptions* matches a client declaration section, which has further NFS clients that are assigned, then this section is split into a section containing that client specifier together with the modified attributes, and a section with the remaining client specifiers and their original attributes.

The **--nfsposition** option can be used to locate this modified entry.

### **--nfsposition** {*PositionIndex* | *Client*}

This option can be used only together with **--nfsadd** or **--nfschange**. It reorders the client declaration section within the export declaration file. The sequence of client entries can be important when there are different client sections affecting the same NFS client. For example, NFS position can be important when defining export attributes for a specific NFS client, and '\*' for all NFS clients, given the same export path.

#### *PositionIndex*

This must be an unsigned numeric value, and it indicates the absolute sequence number after reordering the client declarations. The indexing is zero-based. 0 means the top of the list, 1 means the second entry, and so on. Any number higher than the number of existing entries indicates the end-of-list position.

#### *Client*

The reference client specifier to indicate the new position of the current client entry. The modified client entry is placed at the location where the *Client* reference declaration is, and that section is ordered below next to the modified client section.

## **list**

Lists the declared exports based on the entries in the configuration file stored in the repository. The sequence of rows in the output for a given path reflects also the sequence of the internal client declaration list for each exported path. The sequence of client declarations within an export can be reordered using the **mmnfs export change** *Path* with the **--nfschange** and **--nfsposition** options. The output can be formatted human readable (default) or machine readable.

**--nfsdefs Path**

Lists the export configuration details for the specified *Path*. Without this option the **mmnfs export list** command shows a table with some basic configuration settings for all declared exports.

- Y Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmc1idecode**. Use the **mmc1idecode** command to decode the field.

**--nfsdefs-match Path**

Lists the export configuration details of all export paths, which contain the specified string.

**load**

Overwrites (deletes) all existing NFS export declarations in the repository, if any. The export declarations are fetched from a file provided to the load operation, which could contain a larger number of export declarations. Some basic format checks are done during export load. After loading export declarations from a file, the NFS service is restarted across all the nodes in the cluster.

*ExportCFGFile*

The file name for the new exports declarations. This file is loaded and stored in the repository to be published on all CES nodes running the NFS server. This load procedure can be used to load a set of export declarations and that removes any previous configuration. The NFS servers are restarted in order to apply the changes.

**List of supported client options for the mnmfs export {add | change} command:****ACCESS\_TYPE**

Allowed values are none, RW, RO, MDONLY, and MDONLY\_RO. The default value is none.

**PROTOCOLS**

Allowed values are 3, 4, NFS3, NFS4, V3, V4, NFSv3 , and NFSv4. The default value is 3,4.

**TRANSPORTS**

Allowed values are TCP and UDP. The default value is TCP.

**ANONYMOUS\_UID**

Allowed values are between -2147483648 and 4294967295. The default value is -2.

**ANONYMOUS\_GID**

Allowed values are between -2147483648 and 4294967295. The default value is -2.

**SECTYPE**

Allowed values are none, sys, krb5, krb5i, and krb5p. The default value is sys.

**PRIVILEGEDPORT**

Allowed values are true and false. The default value is false.

**MANAGE\_GIDS=true**

In this configuration, NFS server discards the list of GIDs passed via the RPC and fetch the group information. For getting the group list, the server makes use of the default authentication configured on the system (CES node). So, depending on the configuration, the list is obtained by using (/etc/passwd + /etc/groups) or from LDAP or from AD. To reach AD and get the list of GIDs, the client must use sec=krb5, else NFS server is not able to get the list from AD.

**MANAGE\_GIDS=false**

NFS server makes use of list of GIDs passed by RPC.

## **mmnfs**

### **SQUASH**

Allowed values are root, root\_squash, all, all\_squash, allsquash, no\_root\_squash, none , and noidsquash. The default value is root\_squash.

### **NFS\_COMMIT**

Allowed values are true and false. The default value is false.

**Important:** Use NFS\_COMMIT very carefully because it changes the behavior of how transmitted data is committed on the server side to a NFS v2 like sync-mode on every write action.

### **CLIENTS**

Allowed values are IP addresses in IPv4 or IPv6 notations, hostnames, netgroups, or \* for all. A netgroup name must not start with a numeric character and otherwise must only contain underscore and/or hyphens ("[a-zA-Z\_][0-9a-zA-Z\_-]\*"). The default value is \*.

### **config**

Manages NFS configuration for a CES cluster:

#### **list**

Displays the NFS configuration parameters and their values. This command also displays all the default export configurations. This is used as the defaults by the **mmnfs export add** command if no other client attributes are specified. The output can be formatted to be human readable or machine readable.

#### **--exportdefs**

If this option is specified, the command displays the default export configuration parameters.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

#### **change**

Modifies the NFS configuration parameters. NFS is restarted across all the nodes on which NFS is running, when this command is executed. Only some configuration options can be modified by this command.

The configuration options that can be modified and their allowed values are as follows:

#### **NFS\_PROTOCOLS**

Allowed values are 3, 4, NFS3, NFS4, V3, V4, NFSv3 , and NFSv4. The default value is 3,4.

#### **MNT\_PORT**

Specifies the port for the NFSv3 Mount protocol. Allowed values are between 0 and 65535. The default value is 0.

#### **NLM\_PORT**

Specifies the NLM port for NFSv3. Allowed values are between 0 and 65535. The default value is 0.

#### **RQUOTA\_PORT**

Specifies the RQUOTA port for NFSv3. Allowed values are between 0 and 65535. The default value is 0.

#### **STATD\_PORT**

Specifies the STATD port for NFSv3. Allowed values are between 0 and 65535. The default value is 0.

**LEASE\_LIFETIME**

Specifies the value in seconds during which an NFS server requires an NFS client to do some operation to renew its lease. This is specific to NFSv4. Allowed values are between 0 and 120. The default value is 60.

**GRACE\_PERIOD**

Specifies the time in seconds during which an NFS client is required to reclaim its locks, in the case of NFSv3, and additionally its state, in the case of NFSv4. Allowed values are between 0 and 180. The default is 60.

If the GRACE\_PERIOD is not displayed in the mmnfs config list output, it can be assumed to be the same as the LEASE\_LIFETIME. To display the GRACE\_PERIOD attribute, set it to default (or the desired value) with mmnfs config change GRACE\_PERIOD=DEFAULT (or mmnfs config change GRACE\_PERIOD=<value>).

**DOMAINNAME**

String. The default value is the "host's fully-qualified DNS domain name".

**IDMAPD\_DOMAIN**

String. Domain in the ID Mapd configuration. The default value is the host's fully-qualified DNS domain name.

**LOCAL\_REALMS**

String. Local-Realms in the ID Mapd configuration. The default value is the host's default realm name.

**LOG\_LEVEL**

Allowed values are NULL, FATAL, MAJ, CRIT, WARN, EVENT, INFO , DEBUG, MID\_DEBUG, and FULL\_DEBUG. The default value is EVENT.

**ENTRIES\_HWMARK**

The high water mark for NFS cache entries. Beyond this point, NFS tries to evict some objects from its cache. The default is 1500000.

**Note:** Specifying a port number in the NFS service configuration with the value of '0' means that the service is picking a port number dynamically. This port number might change across service restarts. If a firewall is to be established between the NFS server and the NFS clients, specific port number can be configured via the command to establish discrete firewall rules. Note that the NFS\_PORT 2049 is a well-known and established convention as NFS servers and clients typically expect this port number. Changing the NFS service port numbers impacts existing clients and a remount of the client is required.

The export defaults that can be set are:

**ACCESS\_TYPE**

Allowed values are none, RW, RO, MDONLY, and MDONLY\_RO. The default value is none.

**Note:** Changing this option to any value other than none exposes data to all NFS clients that can access your network, even if the export is created using **add -client ClientOptions** to limit that clients access. All clients, even if not declared with the **--client** in the mmnfs export add, has access to the data. The global value applies to an unseen \*, even if **showmount -eCESIP** does not display it. Use caution if you change it in this global definition.

**ANONYMOUS\_UID**

Allowed values are between -2147483648 and 4294967295. The default value is -2.

**ANONYMOUS\_GID**

Allowed values are between -2147483648 and 4294967295. The default value is -2.

**MANAGE\_GIDS**

Allowed values are true and false. The default value is false.

## **mmnfs**

### **NFS\_COMMIT**

Allowed values are true and false. The default value is false.

**Important:** Use NFS\_COMMIT very carefully because it changes the behavior of how transmitted data is committed on the server side to a NFS v2 like sync-mode on every write action.

### **PRIVILEGEDPORT**

Allowed values are true and false. The default value is false.

### **PROTOCOLS**

Allowed values are 3, 4, NFS3, NFS4, V3, V4, NFSv3 , and NFSv4. The default value is 3,4.

### **SECTYPE**

Allowed values are none, sys, krb5, krb5i, and krb5p. The default value is sys.

### **SQUASH**

Allowed values are root, root\_squash, all, all\_squash, allsquash, no\_root\_squash, none , and noidsquash. The default value is root\_squash.

**Note:** Changing this option to no\_root\_squash exposes data to root on all NFS clients, even if the export is created using **add --client ClientOptions** to limit the root access of that client.

### **TRANSPORTS**

Allowed values are TCP and UDP. The default value is TCP.

If export defaults are set, then new exports that are created pick up the export default values.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmnfs** command.

The node on which the command is issued must be able to execute remote shell commands on any other CES node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To create an NFS export (using a netgroup), issue this command:

```
mmnfs export add /mnt/gpfs0/netgroupopath \
--client "@netgroup(Access_Type=RO,Squash=allsquash)"
```

The system displays output similar to this:

The NFS export was created successfully.

**Note:** Instead of a netgroup, a client IP address can also be declared, like **--client "1.2.3.4"**.

2. To create an NFS export (using a client IP), issue this command:

```
mmnfs export add /mnt/gpfs0/netgroupopath --client "192.0.2.0/20(Access_Type=RW)"
```

The system displays output similar to this:

The NFS export was created successfully.

3. To add a client definition, issue these commands:

```
mmnfs export list --nfsdefs /gpfs/fs1/export_1
```

The system displays output similar to this:

| Path               | Delegations | Clients   | Access_Type | Protocols | Transports | Squash      | Anonymous_uid | Anonymous_gid | SecType | PrivilegedPort | DefaultDelegations | Manage_Gids | NFS_Commit |
|--------------------|-------------|-----------|-------------|-----------|------------|-------------|---------------|---------------|---------|----------------|--------------------|-------------|------------|
| /gpfs/fs1/export_1 | none        | 192.0.2.8 | RW          | 3,4       | TCP        | ROOT_SQUASH | -2            | SYS           | FALSE   | none           | FALSE              | FALSE       |            |

Now add a client definition that is more restrictive to a different client, 198.51.100.10, by issuing the command:

```
mmnfs export change /gpfs/fs1/export_1 \
--nfsadd "198.51.100.10(Access_Type=MDONLY,Squash=allsquash)""
mmnfs export list
```

The system displays output similar to this:

| Path                 | Delegations | Clients       |
|----------------------|-------------|---------------|
| /gpfs/fs1/controller | none        | *             |
| /gpfs/fs1/export_1   | none        | 192.0.2.8     |
| /gpfs/fs1/export_1   | none        | 198.51.100.10 |

Now add a client definition that is very permissive but we want it to be last on the list so that the more restrictive attributes for client 192.0.2.8 take precedence for that one client, by issuing the command:

```
mmnfs export change /gpfs/fs1/export_1 \
--nfsadd "192.0.2.0/20(Access_Type=RW,Squash=no_root_squash)" --nfsposition 4
mmnfs export list --nfsdefs /gpfs/fs1/export_1
```

The system displays output similar to this:

| Path               | Delegations | Clients       | Access_Type | Protocols | Transports | Squash         | Anonymous_uid | Anonymous_gid | SecType | PrivilegedPort | Default | Manage_Gids | NFS_Commit |
|--------------------|-------------|---------------|-------------|-----------|------------|----------------|---------------|---------------|---------|----------------|---------|-------------|------------|
| /gpfs/fs1/export_1 | none        | 192.0.2.8     | RW          | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | none    | FALSE       | FALSE      |
| /gpfs/fs1/export_1 | none        | 198.51.100.10 | MDONLY      | 3,4       | TCP        | allsquash      | -2            | -2            | SYS     | FALSE          | none    | FALSE       | FALSE      |
| /gpfs/fs1/export_1 | none        | 192.0.2.0/20  | RW          | 3,4       | TCP        | no_root_squash | -2            | -2            | SYS     | FALSE          | none    | FALSE       | FALSE      |

**Note:** The **mmnfs export change** restarts NFS services on all CES nodes on which the NFS server is currently running.

- To remove an NFS export, issue this command:

```
mmnfs export change /mnt/gpfs0/somepath --nfsremove "1.2.3.1"
```

The system displays output similar to this:

```
192.168.80.131: Redirecting to /bin/systemctl stop nfs-ganesha.service
192.168.80.135: Redirecting to /bin/systemctl stop nfs-ganesha.service
192.168.80.131: Redirecting to /bin/systemctl start nfs-ganesha.service
192.168.80.135: Redirecting to /bin/systemctl start nfs-ganesha.service
NFS Configuration successfully changed. NFS server restarted on all NFS
nodes on which NFS server is running.
```

**Note:** This command removes a single client definition, for the IP "1.2.3.1", from the NFS export. If this is the last client definition for the export, then the export is also removed.

- To modify an NFS export, issue this command:

```
mmnfs export change /mnt/gpfs0/p1 \
--nfschange "203.0.113.2(Access_Type=R0)" --nfsposition 0
```

The system displays output similar to this:

```
192.168.80.131: Redirecting to /bin/systemctl stop nfs-ganesha.service
192.168.80.135: Redirecting to /bin/systemctl stop nfs-ganesha.service
192.168.80.131: Redirecting to /bin/systemctl start nfs-ganesha.service
192.168.80.135: Redirecting to /bin/systemctl start nfs-ganesha.service
NFS Configuration successfully changed. NFS server restarted on all NFS
nodes on which NFS server is running.
```

## **mmnfs**

6. To list NFS exports, issue this command:

```
mmnfs export list
```

The system displays output similar to this:

| Path                 | Delegations | Clients |
|----------------------|-------------|---------|
| /gpfs/fs1/controller | none        | *       |
| /gpfs/fs1/export_1   | none        | *       |
| /gpfs/fs1/export_2   | none        | *       |

7. To list NFS exports, issue this command:

```
mmnfs export list --nfsdefs /mnt/gpfs0/p1
```

The system displays output similar to this:

| Path          | Delegations | Clients     | Access_Type | Protocols | Transports | Squash      | Anonymous_uid | Anonymous_gid | SecType | Privileged_Port | Default_Delegations | MANAGE_Gids | NFS_Commit |
|---------------|-------------|-------------|-------------|-----------|------------|-------------|---------------|---------------|---------|-----------------|---------------------|-------------|------------|
| /mnt/gpfs0/p1 | none        | 203.0.113.2 | RW          | 3,4       | TCP        | ROOT_SQUASH | -2            | -2            | SYS     | FALSE           | none                | FALSE       | FALSE      |
| /mnt/gpfs0/p1 | none        | *           | RW          | 3,4       | TCP        | ROOT_SQUASH | -2            | -2            | SYS     | FALSE           | none                | FALSE       | FALSE      |
| /mnt/gpfs0/p1 | none        | 203.0.113.1 | RW          | 3,4       | TCP        | ROOT_SQUASH | -2            | -2            | SYS     | FALSE           | none                | FALSE       | FALSE      |

8. To list the NFS configuration, issue this command:

```
mmnfs config list
```

The system displays output similar to this:

NFS Ganesh Configuration:

```
=====
```

NFS\_PROTOCOLS: 3,4

NFS\_PORT: 2049

MNT\_PORT: 0

NLM\_PORT: 0

RQUOTA\_PORT: 0

LEASE\_LIFETIME: 60

GRACE\_PERIOD: 60

DOMAINNAME: VIRTUAL1.COM

DELEGATIONS: Disabled

```
=====
```

STADT Configuration

```
=====
```

STADT\_PORT: 0

```
=====
```

CacheInode Configuration

```
=====
```

ENTRIES\_HWMARK: 1500000

```
=====
```

Export Defaults

```
=====
```

ACCESS\_TYPE: NONE

PROTOCOLS: 3,4

TRANSPORTS: TCP

ANONYMOUS\_UID: -2

ANONYMOUS\_GID: -2

SECTYPE: SYS

PRIVILEGEDPORT: FALSE

MANAGE\_GIDS: FALSE

SQUASH: ROOT\_SQUASH

NFS\_COMMIT: FALSE

```
=====
```

Log Configuration

```
=====
```

LOG\_LEVEL: EVENT

```
=====
```

```
Idmapd Configuration
=====
LOCAL-REALMS: LOCALREALM
DOMAIN: LOCALDOMAIN
=====
```

9. To change STATD\_PORT configuration, issue this command (When a port is assigned, STATD is started on the given port):

```
mmnfs config change STATD_PORT=32765
```

The system displays output similar to this:

```
NFS Configuration successfully changed. NFS server restarted on all NFS nodes
on which NFS server is running.
```

- | 10. To list export configuration details for the path /gpfs/gpfs1/export1, issue this command:

```
| mmnfs export list --nfsdefs /gpfs/gpfs1/export1
```

| The system displays output similar to this:

| Path                | Delegations | Clients   | Access_Type | Protocols | Transports | Squash         | Anonymous_uid | Anonymous_gid | SecType | PrivilegedPort | DefaultDelegations | Manage_Gids | NFS_Commit |
|---------------------|-------------|-----------|-------------|-----------|------------|----------------|---------------|---------------|---------|----------------|--------------------|-------------|------------|
| /gpfs/gpfs1/export1 | NONE        | c49f04n10 | RO          | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | *         | RW          | 3         | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | c49f04n12 | RO          | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | c49f04n11 | NONE        | 3,4       | TCP        | NO_ROOT_SQUASH | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |

- | 11. To list export configuration details for all paths that contain the string - /gpfs/gpfs1/export1, issue this command:

```
| mmnfs export list --nfsdefs-match /gpfs/gpfs1/export1
```

| The system displays output similar to this:

| Path                | Delegations | Clients   | Access_Type | Protocols | Transports | Squash         | Anonymous_uid | Anonymous_gid | SecType | PrivilegedPort | DefaultDelegations | Manage_Gids | NFS_Commit |
|---------------------|-------------|-----------|-------------|-----------|------------|----------------|---------------|---------------|---------|----------------|--------------------|-------------|------------|
| /gpfs/gpfs1/export1 | NONE        | c49f04n10 | RO          | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | *         | RW          | 3         | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | c49f04n12 | RO          | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | c49f04n11 | NONE        | 3,4       | TCP        | NO_ROOT_SQUASH | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |
| /gpfs/gpfs1/export1 | NONE        | *         | NONE        | 3,4       | TCP        | ROOT_SQUASH    | -2            | -2            | SYS     | FALSE          | NONE               | FALSE       | FALSE      |

## See also

- “mmces command” on page 103
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmobj command” on page 466
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

## Location

/usr/lpp/mmnfs/bin

## **mmnsddiscover**

---

### **mmnsddiscover command**

RedisCOVERS paths to the specified network shared disks.

#### **Synopsis**

```
mmnsddiscover [-a | -d "Disk[;Disk...]" | -F DiskFile] [-C ClusterName]
[-N {Node[,Node...] | NodeFile | NodeClass}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmnsddiscover** command is used to rediscover paths for GPFS NSDs on one or more nodes. If you do not specify a node, GPFS redisCOVERS NSD paths on the node from which you issued the command.

On server nodes, **mmnsddiscover** causes GPFS to rediscover access to disks, thus restoring paths which may have been broken at an earlier time. On client nodes, **mmnsddiscover** causes GPFS to refresh its choice of which NSD server to use when an I/O operation occurs.

In general, after the path to a disk is fixed, the **mmnsddiscover** command must be first run on the server that lost the path to the NSD. After that, run the command on all client nodes that need to access the NSD on that server. You can achieve the same effect with a single **mmnsddiscover** invocation if you utilize the **-N** option to specify a node list that contains all the NSD servers and clients that need to rediscover paths.

#### **Parameters**

**-a** RedisCOVERS paths for all NSDs. This is the default.

**-d "DiskName[;DiskName]"**

Specifies a list of NSDs whose paths are to be rediscovered.

**-F DiskFile**

Specifies a file that contains the names of the NSDs whose paths are to be rediscovered.

**-C ClusterName**

Specifies the name of the cluster to which the NSDs belong. This defaults to the local cluster if not specified.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes on which the rediscovery is to be done.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands in the IBM Spectrum Scale: Administration Guide*.

#### **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmnsddiscover** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For

more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To rediscover the paths for all of the NSDs in the local cluster on the local node, issue the command:

```
mmnsddiscover
```

The system displays output similar to:

```
mmnsddiscover: Attempting to rediscover the disks.
This may take a while ...
```

```
mmnsddiscover: Finished.
```

2. To rediscover the paths for all of the NSDs in the local cluster on all nodes in the local cluster, issue the command:

```
mmnsddiscover -a -N all
```

The system displays output similar to:

```
mmnsddiscover: Attempting to rediscover the disks.
This may take a while ...
```

```
mmnsddiscover: Finished.
```

3. To rediscover the paths for a given list of the NSDs on a node in the local cluster, issue the command:

```
mmnsddiscover -d "gpfs1nsd;gpfs2nsd" -N c6f2c2vp5
```

The system displays output similar to:

```
mmnsddiscover: Attempting to rediscover the disks. This may take a while ...
c6f2c2vp5.ppd.pok.ibm.com: GPFS: 6027-1805 [N] Rediscovered nsd server access to gpfs1nsd.
c6f2c2vp5.ppd.pok.ibm.com: GPFS: 6027-1805 [N] Rediscovered nsd server access to gpfs2nsd.
mmnsddiscover: Finished.
```

## See also

- “[mmchnsd command](#)” on page 198
- “[mmcrlnsd command](#)” on page 262
- “[mmdelnsd command](#)” on page 302
- “[mmlsnsd command](#)” on page 423

## Location

/usr/lpp/mmfs/bin

## **mmobj**

---

### **mmobj command**

Manages configuration of Object protocol service, and administers storage policies for object storage, unified file and object access, and multi-region object deployment.

#### **Synopsis**

```
mmobj swift base -g GPFSMountPoint --cluster-hostname CESHostName
 [-o ObjFileset] [-i MaxNumInodes] [--ces-group CESGroup]
 {{{{--local-keystone [--db-password Password] [--admin-token Token]}}
 | [--remote-keystone-url URL] [--configure-remote-keystone]}}
 --admin-password Password [--admin-user AdminUser]
 [--swift-user SwiftUser] [--swift-password SwiftPassword]
 [--enable-file-access] [--enable-s3] [--enable-multi-region]
 [--join-region-file RegionFile] [--region-number RegionNumber]
```

or

```
mmobj config list --ccrfile CCRFile [--section Section [--property propertyName]]
```

or

```
mmobj config change --ccrfile CCRFile --section Section --property propertyName --value Value
```

or

```
mmobj config change --ccrfile CCRFile --merge-file MergeFile
```

or

```
mmobj policy list [[--policy-name PolicyName] | [--policy-function PolicyFunction]] [--verbose]
```

or

```
mmobj policy create PolicyName [-f FilesetName]
 [--file-system FilesystemName] [-i MaxNumInodes]
 {{[--enable-compression --compression-schedule CompressionSchedule]}
 {{[--enable-encryption --encryption-keyfile EncryptionKeyFileName [--force-rule-append]]}}
 [--enable-file-access]
```

or

```
mmobj policy change PolicyName --default
```

or

```
mmobj policy change PolicyName --deprecate State
```

or

```
mmobj policy change --add-local-region
```

or

```
mmobj policy change --remove-region-number RegionNumber
```

| or

```
| mmobj policy change --compression-schedule CompressionSchedule
```

or

```
mmobj file-access enable
```

or

```
mmobj file-access disable [--objectizer]
```

or

```
mmobj file-access objectize
{ --object-path ObjectPath
| --storage-policy PolicyName
| --account-name AccountName
| --container-name ContainerName
| --object-name ObjectName]]}
[-N | --node NodeName]
```

or

```
mmobj file-access link-filesystem
--sourcefs-path FilesetPath
--account-name AccountName
--container-name ContainerName
--fileaccess-policy-name PolicyName
[--update-listing]
```

or

```
mmobj multiregion list
```

or

```
mmobj multiregion enable
```

or

```
mmobj multiregion export --region-file RegionFile
```

or

```
mmobj multiregion import --region-file RegionFile
```

or

```
mmobj multiregion remove --region-number RegionNumber --force
```

or

```
mmobj s3 enable | disable | list
```

## Availability

Available with IBM Spectrum Scale Standard Edition or higher.

## Description

Use the **mmobj** command to modify and change the Object protocol service configuration, and to administer storage policies for object storage, unified file and object access, and multi-region object deployment.

**Note:** The **mmobj config list** and the **mmobj config change** commands are used to list and change the configuration values for the underlying Swift service stored in the Cluster Configuration Repository (CCR).

At least one CES IP address is required in the node running the **mmobj swift base** command to set `object_singleton_node` and `object_database_node` attributes.

- The node with the `object_database_node` attribute runs the keystone database.
- The node with the `object_singleton_node` attribute runs unique object services across the CES cluster.

You can verify the address using the **mmces address list**. The IP address can be added using the **mmces address add** command.

## **mmobj**

If there is an existing object authentication configured, use the **mmuserauth service remove --data-access-method object** command to remove the object authentication. Then, use the **mmces service disable OBJ** command to perform the necessary cleanup before running the **mmobj swift base command** again.

Object authentication can be either local or remote. If the Object authentication is local, the Keystone identity service and the Keystone database will be running in and handled by the CES cluster. If the Object authentication is remote, the Keystone server must be fully configured and running before Object services are installed.

In the unified file and object access environment, the ibmobjectizer service can be used for making files created from the file interfaces, such as POSIX, NFS, and CIFS accessible through object interfaces, such as curl and SWIFT. The ibmobjectizer service runs periodically and makes files available for the object interface. The **file-access** option of the **mmobj** command can be used to enable and disable the file-access capability and the ibmobjectizer service. The ibmobjectizer service runs on the CES IP address with the **object\_singleton\_node** attribute.

The **file-access** option of the **mmobj** command can be used to enable object access on the files stored in an existing fileset under a container that is associated with a unified file and object access storage policy in an account.

In a data lake environment, you can access files that are already stored in an existing fileset via object interface through swift or curl by using the **mmobj file-access link-fileset** command. This command can be used to make the existing fileset data accessible under a unified file and object access storage policy container. By default, this command allows the existing fileset to have object access without updating the container and metadata. However, you can use the **--update-listing** option to update container listing with files that are existing in the linked fileset. If the **--update-listing** option of the **mmobj file-access link-fileset** command is used, then the **sourcefile-set-path** must be the exact fileset junction path and the fileset must be derived from the object file system.

**Note:** The **mmobj file-access link-fileset** command does not enable and disable the unified file and object access feature. It is only used to link the existing filesets to provide object access. Unlinking the object access-enabled filesets is not supported.

## **Parameters**

### **swift**

Configures the underlying Swift services:

#### **base**

Specifies the configuration of the Object protocol service.

#### **-g GPFSMountPoint**

Specifies the mount path of the GPFS file system used by Swift.

*GPFSMountPoint* is the mount path of the GPFS file system used by the Object store.

#### **--cluster-hostname CESHostName**

Specifies the hostname which is used to return one of the CES IP addresses. The returned CES IP address is used in the endpoint for the identity and Object-store values stored in Keystone.

*CESHostName* is the value for cluster host name used in the identity and object-store endpoint definitions in Keystone. Ideally, it should be a hostname which will return one the CES IP addresses, such as a round-robin DNS. It could also be a fixed IP of a load balancer that distributes requests to one of the CES nodes. It is not recommended to use an ordinary CES IP since all identity and object-store requests would be routed to the single node with that address and may cause performance issues.

**--local-keystone**

Specifies that a new Keystone server will be installed and configured locally in the cluster.

**--db-password *Password***

Specifies the password for the 'keystone' user in the postgres database. Defaults to the value for --admin-password.

**--admin-user *User***

Specifies the name of the admin user in Swift. The default value is admin.

**--admin-password *Password***

Specifies the password to be used when creating the administrator user in Keystone.

**--admin-token *Token***

Specifies the admin token to be used for the initial Keystone setup. If it is not specified, a random string will be used.

**--swift-user *User***

Specifies the user for the Swift services. The default value is swift.

**--swift-password *Password***

Specifies the password for the Swift user. The default value is the password value from --admin-password.

**--remote-keystone-url *URL***

Specifies the URL to an existing Keystone service.

**--configure-remote-keystone**

When a remote Keystone is used, this specifies that the remote Keystone should be configured as necessary. The required users, roles and endpoints needed by the Swift services will be added to the Keystone server. Keystone authentication information needs to be specified with the

**--admin-password** or the **--admin-token** flag to enable the configuration. If this flag is not specified, the remote Keystone is not modified and the administrator will need to add the appropriate entries for the Swift configuration after the install is complete.

**--admin-password *Password***

Specifies the password to be used when creating the administrator user in Keystone.

**-o *ObjFileset***

Specifies the name of the fileset to be created in GPFS for the object storage.

*ObjFileset* is the name of the independent fileset that will be created for the Object store. By default, `object_fileset` is created.

**-i**

Specifies the maximum number of inodes for the Object fileset.

***MaxNumInodes***

The maximum number of inodes for the Object fileset. By default, 8000000 is set.

**--ces-group**

Specifies the CES group that contains the IP addresses to be used for the swift ring files. This allows you to specify a subset of the overall collection of CES IP addresses to be used by the object protocol.

**--enable-s3**

Sets the s3 capability (Amazon S3 API support) to true. By default, S3 API is not enabled.

**--enable-file-access**

Sets the file-access capability initially to true. Further configuration is still necessary using the **mmobj file-access** command. By default, the file-access capability is not enabled.

## **mmobj**

### **--enable-multi-region**

Sets the multi-region capability initially to true. By default, multi-region capability is not enabled.

### **--join-region-file *RegionFile***

Specifies that this object installation will join an existing object multi-region Swift cluster. *RegionFile* is the region data file created by the **mmobj multiregion export** command from the existing multi-region cluster.

**Note:** The use of the **--configure-remote-keystone** flag is recommended so that the region-specific endpoints for this region are automatically created in Keystone.

### **--region-number *RegionNumber***

Specifies the Swift region number for this cluster. If it is not specified, the default value is to 1. In a multi-region configuration, this flag is required and must be a unique region number which is not used by another region in the multi-region environment.

## **config**

Administers the Object configuration:

### **list**

Lists configuration values of the underlying Swift/Keystone service stored in the CCR.

### **--section *Section***

Retrieves values for the specified section only.

The section is the heading enclosed in brackets ([] in the associated configuration file.

### **--property *PropertyName***

Retrieves values for the specified property only.

### **change**

Enables modifying Swift/Keystone configuration files. After you modify the configuration files, the CES monitoring framework downloads them from the CCR and distributes them to all the CES nodes in the cluster. The framework also automatically restarts the services which depend on the modified configuration files.

**Note:** It is recommended to not directly modify the configuration files in /etc/swift and /etc/keystone folders as they can be overwritten at any time by the files stored in the CCR.

### **--section *Section***

Specifies the section in the file that contains the parameter.

The section is the heading enclosed in brackets ([] in the associated configuration file.

### **--property *PropertyName***

Specifies the name of the property to be set.

### **--value *NewValue***

Specifies the value of the *PropertyName*.

### **--merge-file *MergeFile***

Specifies a file in the openstack-config .conf format that contains multiple values to be changed in a single operation. The properties in *MergeFile* can represent new properties or properties to be modified. If a section or property name in *MergeFile* begins with a '-' character, that section or property is deleted from the file. For example, a *MergeFile* with the following contents would delete the ldap section, set the connections property to 512, and delete the **noauth** property from the database section.

```
[-ldap]
[database]
connections = 512
-noauth =
```

**Parameter common for both mmobj config list and mmobj config change commands:**

**--ccrfile *CCRFile***

Indicates the name of the Swift, Keystone, or object configuration file stored in the CCR.

Some of the configuration files stored in the Cluster Configuration Repository (CCR) are:

- account-server.conf
- container-reconciler.conf
- container-server.conf
- object-expirer.conf
- object-server.conf
- proxy-server.conf
- swift.conf
- keystone.conf
- keystone-paste.ini
- spectrum-scale-object.conf
- object-server-sof.conf
- spectrum-scale-objectizer.conf

**policy**

Administers the storage policies for object storage:

**list**

Lists storage policies for object storage.

**--policy-name *PolicyName***

Lists details of the specified storage policy, if it exists.

**--policy-function *PolicyFunction***

Lists details of the storage policies with the specified function, if any exist.

**--verbose**

Lists the functions enabled for the storage policies.

**create**

Creates a storage policy for object storage. The associated configuration files are updated and the ring files are created for the storage policy. The CES monitoring framework distributes the changes to the protocol nodes and restarts the associated services.

***PolicyName***

Specifies the name of the storage policy.

The policy name must be unique (case insensitive), without spaces, and it must contain only letters, digits, or a dash.

**-f *FilesetName***

Specifies the name of an existing fileset to be used for this storage policy. An existing fileset can be used provided it is not being used for an existing storage policy.

If no fileset name is specified with the command, the policy name is reused for the fileset with the prefix obj\_.

**--file-system *FilesystemName***

Specifies the name of the file system on which the fileset is created.

**--i *MaxNumInodes***

Specifies the inode limit for the new inode space.

**--enable-compression**

Enables a compression policy. The Swift policy type is replication. If --enable-compression is used, --compression-schedule must be specified too and vice-versa.

## **mmobj**

Every object stored within a container that is linked to this storage policy is compressed on a scheduled basis. This occurs as a background process. For object retrieval, no decompression is needed because it occurs automatically in the background.

### **--compression-schedule: "MM:HH:dd:ww"**

Specifies the compression schedule if **--enable-compression** is used. Schedule needs to be given in the *MM:HH:dd:ww* format :

#### **MM = 0-59 minutes**

Minute after the hour the job should be executed. The range is 0 to 59.

#### **HH = 0-23**

Hour in which the job should be executed. Hours are represented as numbers from 0 to 23.

#### **dd = 1-31**

The day of a month on which the job should be executed. Days are represented as numbers from 1 to 31.

#### **ww = 0-7 (0=Sun, 7=Sun)**

The days of the week the job should be executed. One or more values can be specified (comma separated). Days are represented as numbers from 0 to 7. 0 and 7 represent Sunday. All days of a week are represented by \*. Optional. Default is 0.

- Use \* for specifying every instance of a unit. For example, dd = \* means that the job is scheduled to run every day.
- Comma separated lists are allowed. For example, dd = 1,3,5 means that the job is scheduled to run on every 1st, 3rd, 5th of a month.
- If ww and dd both are specified, the union is used.
- Specifying a range using - is not supported.
- Empty values are allowed for dd and ww. If empty, dd and or ww are not considered.
- Empty values for mm and hh are treated as \*.

### **--enable-encryption**

Enables an encryption policy.

### **--enable-file-access**

Enables a file-access policy. The file-access policies only exist in the region in which they were created. They do not support the multi-region capability. Objects stored within a container that is linked to this storage policy can be enabled for file protocol access.

### **--encryption-keyfile *EncryptionKeyFileName***

Specifies the encryption key file.

### **--force-rule-append**

Specifies whether to append and establish the rule if other rules are already existing

**Note:** The enabled functions are displayed in the functions column of the **mmobj policy list** command output as follows:

- **--enable-compression** compression
- **--enable-file-access** file-and-object-access

### **change**

Changes the state of the specified storage policy.

#### ***PolicyName***

Specifies the name of the storage policy that needs to be changed.

### **--default**

Sets the specified storage policy to be the default policy.

**Note:** You cannot set a deprecated storage policy as the default storage policy.

#### --deprecate *State*

Deprecates the specified storage policy. *State* can be either yes or no.

**yes**

Sets the state of the specified storage policy as deprecated.

**no** Sets the state of the specified storage policy as not deprecated.

**Note:** You cannot deprecate the default storage policy.

#### --add-local-region

In a multi-region environment, adds the region of the current cluster to the specified storage policy. The associated fileset previously defined for the storage policy must already exist or else it is created.

After the region is added, the multi-region configuration needs to be synced with the other regions by using the **mmobj multiregion** command.

**Note:** By default, a storage policy only stores objects in the region on which it was created. If the cluster is defined as multi-region, a storage policy can also be made multi-region by adding additional regions to its definition.

#### --remove-region-number *RegionNumber*

In a multi-region environment, removes a region from the specified storage policy. The associated fileset for the storage policy is not modified.

After the region is removed, the multi-region configuration needs to be synced with the other regions by using the **mmobj multiregion** command.

#### --compression-schedule: "MM:HH:dd:ww"

Specifies the compression schedule if **--enable-compression** is used. Schedule needs to be given in the *MM:HH:dd:ww* format :

**MM = 0-59 minutes**

Minute after the hour the job should be executed. The range is 0 to 59.

**HH = 0-23**

Hour in which the job should be executed. Hours are represented as numbers from 0 to 23.

**dd = 1-31**

The day of a month on which the job should be executed. Days are represented as numbers from 1 to 31.

**ww = 0-7 (0=Sun, 7=Sun)**

The days of the week the job should be executed. One or more values can be specified (comma separated). Days are represented as numbers from 0 to 7. 0 and 7 represent Sunday. All days of a week are represented by \*. Optional. Default is 0.

- Use \* for specifying every instance of a unit. For example, dd = \* means that the job is scheduled to run every day.
- Comma separated lists are allowed. For example, dd = 1,3,5 means that the job is scheduled to run on every 1st, 3rd, 5th of a month.
- If ww and dd both are specified, the union is used.
- Specifying a range using - is not supported.
- Empty values are allowed for dd and ww. If empty, dd and or ww are not considered.
- Empty values for mm and hh are treated as \*.

## **mmobj**

### **file-access**

Manages file-access capability, ibmobjectizer service and object access for files (objectizes) in a unified file and object access environment.

#### **enable**

Enables the file-access capability and the ibmobjectizer service.

If the file access capability is already enabled and the ibmobjectizer service is stopped, this option starts the ibmobjectizer service

#### **disable**

Disables the file-access capability and the ibmobjectizer service.

#### **--objectizer**

This option stops only the ibmobjectizer service, it does not disable the file-access capabilities.

#### **objectize**

Enables file(s) for object access (objectizes) in a unified file and object access environment.

#### **--object-path *ObjectPath***

The fully qualified path of a file or a directory that you want to enable for access through the object interface. If a fully-qualified path to a directory is specified, the command enables all the files from that directory for access through the object interface. This is a mandatory parameter for the **mmobj file-access** command if the **--storage-policy** parameter is not specified.

#### **--storage-policy *PolicyName***

The name of the storage policy for which you want to enable files for the object interface.

This is a mandatory parameter for the **mmobj file-access** command if the **--object-path** parameter is not specified. If only this parameter is specified, the command enables all files for object interface from the fileset associated with the specified storage policy.

#### **--account-name *AccountName***

The account name for which you want to enable files for access through the object interface.

The **--storage-policy** parameter is mandatory if you are using this parameter.

#### **--container-name *ContainerName***

The container name for which you want to enable files for access through the object interface.

You must specify the **--storage-policy** and the **--account-name** with this parameter.

#### **--object-name *ObjectName***

The object name for which you want to enable files for access through the object interface.

You must specify **--storage-policy**, **--account-name**, and **--container-name** parameters with this parameter.

#### **-N | --node *NodeName***

The node on which the command is to be executed. Optional. If this parameter is not specified, the command is executed on the current node if it is a protocol node. If the current node is not a protocol node, an available protocol node is selected.

### **link-fileset**

Enables object access to the given fileset path.

#### **--sourcefileset-path *FilesetPath***

The fully qualified non-object fileset path that has to be enabled for object access.

#### **--account-name *AccountName***

The name of the swift account or the keystone project. The contents of the fileset will be accessible from this account when it is accessed by using the object interface.

#### **--container-name *ContainerName***

The name of the container. The contents of the fileset belong to this container when it is accessed by using the object interface.

**--fileaccess-policy-name *PolicyName***

The name of the file-access enabled storage policy.

**--update-listing**

Updates the container database with the existing files in the provided fileset.

If the --update-listing option is used, then the **sourcefile-set-path** must be the exact fileset junction path and the fileset must be derived from the object file system.

**multiregion**

Administers multi-region object deployment. For more information on multi-region object deployment and capabilities, see *Overview of multi-region object deployment* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

**list**

Lists the information about the region.

**enable**

Enable the cluster for multi-region support.

Only the first cluster of the region can run the enable command. Subsequent regions join the multi-region cluster during installation with the use of the **--join-region-file** flag of the **mmobj swift base** command.

**export**

Exports the multi-region configuration environment so that other regions can be installed into the multi-region cluster or other regions can be synced to this region.

If successful, a region checksum is printed in the output. This checksum can be used to ensure different regions are in sync when the **mmobj multiregion import** command is run.

**Note:** When region-related information changes, such as CES IPs and storage policies, all regions must be updated with the changes.

**--region-file *RegionFile***

Specifies a path to store the multi-region data.

This file is created.

**import**

Imports the specified multi-region configuration environment into this region.

If successful, a region checksum for this region is printed in the output. If the local region configuration matches the imported configuration, the checksums match. If they differ, then it means that some configuration information in the local region needs to be synced to the other regions. This can happen when a configuration change in the local region, such as adding CES IPs or storage policies, has not yet been synced with the other regions. If this is the case, the multi-region configuration for the local region needs to be exported and synced to the other regions.

**--region-file *RegionFile***

Specifies the path to a multi-region data file created by using the **mmobj multiregion export** command.

**remove**

Completely removes a region from the multi-region environment. The removed region will no longer be accessible by other regions.

After the region is removed, the remaining regions need to have their multi-region information synced with this change by using the **mmobj multiregion export** and **mmobj multiregion import** commands.

## **mmobj**

### **--region-number *RegionNumber***

Specifies the region number that you need to remove from the multi-region configuration.

### **--force**

Indicates that all the configuration information for the specified region needs to be permanently deleted.

- s3** Enables and disables the S3 API without manually changing the configuration.

#### **enable**

Enables the S3 API.

#### **disable**

Disables the S3 API.

#### **list**

Verifies if the S3 API has been enabled or disabled.

## **Exit status**

- 0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmobj** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To specify configuration of Object protocol service with local keystone and S3 enabled, issue the following command:

```
mmobj swift base -g /gpfs/ObjectFS --cluster-hostname cluster-ces-ip.ibm --local-keystone
--enable-s3 --admin-password Passw0rd
```

The system displays output similar to this:

```
mmobj swift base: Creating fileset /dev/ObjectFS object_fileset
mmobj swift base: Configuring Keystone server in /gpfs/cesshared/ces/object/keystone
mmobj swift base: Configuring Swift services
Configuration complete
```

2. To list object configuration settings for proxy-server.conf, DEFAULT section, issue the following command:

```
mmobj config list --ccrfile proxy-server.conf --section DEFAULT
```

The system displays output similar to this:

```
[DEFAULT]
bind_port = 8080
workers = auto
user = swift
log_level = ERROR
```

3. To change the number of worker processes that each object server launches, update the object-server.conf file as shown here:

```
mmobj config change --ccrfile object-server.conf --section DEFAULT --property workers --value 16
```

4. To change the configuration value of `paste.filter_factory` which is in the `filter:s3_extension` section of the `keystone-paste.ini` configuration file, issue the following command:

```
mmobj config change --ccrfile keystone-paste.ini --section filter:s3_extension
--property paste.filter_factory --value keystone.contrib.s3:S3Extension.factory
```

5. To create a new storage policy `CompressionTest` with the compression function enabled and with the compression schedule specified, issue the following command:

```
| mmobj policy create CompressionTest --enable-compression
| --compression-schedule '20:5,11,17,23:*:*'
```

The system displays output similar to this:

```
[I] Getting latest configuration from ccr
[I] Creating fileset /dev/ObjectFS:obj_CompressionTest
[I] Creating new unique index and build the object rings
[I] Updating the configuration
[I] Uploading the changed configuration
```

This compression schedule indicates that the fileset associated with the policy will be compressed at the 20th minute of the 5th, 11th, 17th, and 23rd hour of the day, every day of every week.

6. To list storage policies for object storage with details of functions available with those storage policies, issue the following command:

```
mmobj policy list --verbose
```

The system displays output similar to this:

| Index       | Name         | Deprecated | Fileset         | Fileset Path                  | Functions              | Function Details | File System |
|-------------|--------------|------------|-----------------|-------------------------------|------------------------|------------------|-------------|
| 0           | SwiftDefault |            | object_fileset  | /ibm/ObjectFS/object_fileset  |                        |                  | ObjectFS    |
| 11751509160 | sof-policy   |            | obj_sof-policy  | /ibm/ObjectFS/obj_sof-policy  | file-and-object-access | regions="1"      | ObjectFS    |
| 11751509230 | mysofpolicy  |            | obj_mysofpolicy | /ibm/ObjectFS/obj_mysofpolicy | file-and-object-access | regions="1"      | ObjectFS    |
| 11751510260 | Test19       |            | obj_Test19      | /ibm/ObjectFS/obj_Test19      |                        | regions="1"      | ObjectFS    |

7. To change the default storage policy, issue the following command:

```
mmobj policy change sof-policy --default
```

The system displays `sof-policy` as the default storage policy.

8. To enable object access for an account, issue the following command:

```
mmobj file-access --storage-policy sof_policy --account-name admin
```

The system displays output similar to the following:

```
Loading objectization configuration from CCR
Fetching storage policy details
Creating container to database map
Performing objectization
Objectization complete
```

9. To enable object access for a container, issue the following command:

```
mmobj file-access --storage-policy sof_policy
--account-name admin --container-name container1
```

10. To enable object access on a file while specifying a storage policy, issue the following command:

```
mmobj file-access --storage-policy sof_policy
--account-name admin --container-name container1 --object-name file1.txt
```

11. To enable object access on a file, issue the following command:

```
mmobj file-access --object-path
/ibm/ObjectFS/obj_sofpolicy1/s69931509221z1device1/AUTH_12345/container1/file1.txt
```

12. To list the information about a region, issue the following command:

```
mmobj multiregion list
```

The system displays output similar to this:

| Region | Endpoint  | Cluster Name    | Cluster Id          |
|--------|-----------|-----------------|---------------------|
| 1      | RegionOne | europe.gpfs.net | 3694106483743716196 |
| 2      | Region2   | asia.gpfs.net   | 1860802811592373112 |

13. To set up the initial multi-region environment on the first region, issue the following command:

## **mmobj**

```
mmobj multiregion enable
```

The system displays output similar to this:

```
mmobj multiregion: Multi-region support is enabled in this cluster.
Region number: 1
```

14. To export multi-region data for use by other clusters to join multi-region, issue the following command:

```
mmobj multiregion export --region-file /tmp/region2.dat
```

The system displays output similar to this:

```
mmobj multiregion: The Object multi-region export file
was successfully created: /tmp/region2.dat
mmobj multiregion: Region checksum is: 34632-44791
```

15. To import the specified multi-region configuration environment created by the export command into a region, issue the following command:

```
mmobj multiregion import --region-file /tmp/region2.dat
```

The system displays output similar to this:

```
mmobj multiregion: The region config has been updated.
mmobj multiregion: Region checksum is: 34632-44791
```

16. To remove a region designated by region number 2 from a multi-region environment and to remove all configuration information of the specified region, issue the following command:

```
mmobj multiregion remove --remove-region-number 2 --force
```

The system displays output similar to this:

```
mmobj multiregion: Permanently removing region 2
(asia.gpfs.net 1860802811592373112) from multi-region configuration.
mmobj multiregion: Updating ring files.
mmobj multiregion: Successfully removed region 2.
Object services on region 2 will need to be unconfigured and its endpoint
removed from Keystone.
```

17. To create a policy with no force add where only the default GPFS policy rule is established, issue the following command:

```
mmobj policy create enc-policy-1 --enable-encryption
--encryption-keyfile /root/enc/enc.key
```

The system creates the policy, adds and establishes the rule within the GPFS policy rules, and displays the following output:

```
[I] Getting latest configuration from ccr
[I] Creating fileset /dev/ObjectFS:obj_enc-policy-1
[I] Creating GPFS Encryption Rule
[I] The following new encryption rule has been stored to
file:/var/mmfs/ces/policyencryption.rule
and is established within the GPFS policies.

/* begin of the encryption rule for fileset obj_enc-policy-1 */
rule 'enc-1_enc-policy-1' set encryption 'encryption_enc-policy-1'
for fileset('obj_enc-policy-1')
where name like '%'

rule 'enc-2_enc-policy-1' encryption 'encryption_enc-policy-1'
is ALGO 'DEFAULTNISTSP800131A'
KEYS('KEY-2712784a-d1ee-4c86-bf97-b88f918cbd12:sklm')
/* end of the encryption rule for fileset obj_enc-policy-1 */

[I] Creating new unique index and building the object rings
[I] Updating the configuration
[I] Uploading the changed configuration
```

18. To create a policy with force add, issue the following command:

```
mmobj policy create enc-policy-2 --enable-encryption
--encryption-keyfile /root/enc/enc.key
--force-rule-add
```

The system creates the policy, adds and establishes the rule within GPFS policy rules, and displays the following output:

```
[I] Getting latest configuration from ccr
[I] Creating fileset /dev/ObjectFS:obj_enc-policy-2
[I] Creating GPFS Encryption Rule
[I] The following new encryption rule has been stored to file:/var/mmfs/ces/policyencryption.rule
and is established within the GPFS policies.

/* begin of the encryption rule for fileset obj_enc-policy-2 */
rule 'enc-1_enc-policy-2' set encryption 'encryption_enc-policy-2'
for fileset('obj_enc-policy-2')
where name like '%'

rule 'enc-2_enc-policy-2' encryption 'encryption_enc-policy-2' is
ALGO 'DEFAULTNISTSP800131A'
KEYS('KEY-2712784a-d1ee-4c86-bf97-b88f918cbd12:sklm')
/* end of the encryption rule for fileset obj_enc-policy-2 */

[I] Creating new unique index and building the object rings
[I] Updating the configuration
[I] Uploading the changed configuration
```

19. To create a policy with no force add, issue the following command:

```
mmobj policy create enc-policy-3 --enable-encryption
--encryption-keyfile /root/enc/enc.key
```

The system creates the policy, adds the rule but does not establish the rule within the GPFS policy rules, and displays the following output:

```
[I] Getting latest configuration from ccr
[I] Creating fileset /dev/ObjectFS:obj_enc-policy-3
[I] Creating GPFS Encryption Rule
[I] The following new encryption rule has been stored to
file:/var/mmfs/ces/policyencryption.rule
but is not established within the GPFS policies.
```

```
/* begin of the encryption rule for fileset obj_enc-policy-3 */
rule 'enc-1_enc-policy-3' set encryption 'encryption_enc-policy-3'
for fileset('obj_enc-policy-3')
where name like '%'

rule 'enc-2_enc-policy-3' encryption 'encryption_enc-policy-3' is
ALGO 'DEFAULTNISTSP800131A'
KEYS('KEY-2712784a-d1ee-4c86-bf97-b88f918cbd12:sklm')
/* end of the encryption rule for fileset obj_enc-policy-3 */
```

```
[I] Creating new unique index and building the object rings
[I] Updating the configuration
[I] Uploading the changed configuration
```

20. To enable the ibmobjectizer service, issue the following command:

```
mmobj file-access enable
```

The system enables the file-access capability and starts the ibmobjectizer service.

21. To disable the ibmobjectizer service, issue the following command:

```
mmobj file-access disable
```

The system disables the file-access capability and stops the ibmobjectizer service.

22. To use the disable --objectizer-daemon parameter, issue the following command:

```
mmobj file-access disable --objectizer-daemon
```

The system disables the ibmobjectizer service.

23. To use the object-path parameter, issue the following command:

```
mmobj file-access objectize --object-path
/ibm/ObjectFS/fileset1/Auth_12345/container1/file1.txt
```

## **mmobj**

The system objectizes file1.txt at /ibm/ObjectFS/fileset1/Auth\_12345/container1/ and enables it for access through the object interface:

```
Loading objectization configuration from CCR
Fetching storage policy details
Performing objectization
Objectization complete
```

24. To use the storage-policy parameter for file objectization, issue the following command:

```
mmobj file-access objectize --storage-policy sof_policy --account-name admin
--container-name container1 --object-name file1.txt
```

The system objectizes file1.txt from the container named container1:

```
Loading objectization configuration from CCR
Fetching storage policy details
Performing objectization
Objectization complete
```

25. To use the node parameter for running the command on a remote node, issue the following command:

```
mmobj file-access objectize --node gpfs_node1
--storage-policy sof_policy --account-name admin
```

The command is run on the gpfs\_node1 node and all the files from all containers within the admin account are objectized. If **--node** or **-N** is not specified, the **mmobj file-access objecitze** command checks if the current node is CES node or not. If the current node is a CES node, the command is executed on the current node. If the current node is not a CES node, the command is executed on a random remote CES node:

```
Loading objectization configuration from CCR
Fetching storage policy details
Performing objectization
Objectization complete
```

## **See also**

- “mmces command” on page 103
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

## **Location**

/usr/lpp/mmfs/bin

---

## mmperfmon command

Configures the Performance Monitoring tool and lists the performance metrics.

### Synopsis

```
mmperfmon config generate --collectors CollectorNode[,CollectorNode...]
 [--config-file InputFile]
```

or

```
mmperfmon config add --sensors SensorFile
```

or

```
mmperfmon config update { [--collectors CollectorNode[,CollectorNode...]]
 [--config-file InputFile] [Attribute=value ...] }
```

or

```
mmperfmon config delete {--all |--sensors Sensor[,Sensor...]}
```

or

```
mmperfmon config show [--config-file OutputFile]
```

or

```
mmperfmon query Metric[,Metric...] | Key[,Key...] | NamedQuery
 [StartTime EndTime | Duration]
 [Options]
```

or

```
mmperfmon query compareNodes ComparisonMetric
 [StartTime EndTime | Duration]
 [Options]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher.

The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms.

### Description

**mmperfmon config** modifies the Performance Monitoring tool by updating the configuration stored in IBM Spectrum Scale. It can be used to generate an initial configuration, to update reporting periods of different sensors, or to restrict sensors to a given set of nodes.

**mmperfmon query** is used to query metrics in a cluster from the performance metrics collector. Output can be delivered either in a raw format, formatted table layout or as a CSV export.

In addition to metrics known by the performance collector, the **mmperfmon query** command can also run predefined named queries or use predefined computed metrics. You can specify a bucket size for each record to return in number of seconds and the number of buckets to retrieve. You can also specify the duration or a time range for which the query can run.

### Parameters

#### config

## **mmpfmon**

### **generate**

Generates the configuration of the Performance Monitoring tool.

**Note:** Once the configuration has been generated, do not forget to turn on monitoring through the **mmchnode** command.

--collectors *CollectorNode[,CollectorNode...]* specifies the set of collectors to which the sensors report their performance measurements. The number of collectors that each sensor shall report to may be specified through **colRedundancy** parameter in the template sensor configuration file (see --config-file).

--config-file *InputFile* specifies the template sensor configuration file to use. If this option is not provided, the /opt/IBM/zimon/defaults/ZIMonSensors.cfg file is used.

### **add**

Adds a new sensor to the Performance Monitoring tool.

--sensors *SensorFile* adds the sensors specified in *SensorFile* to the sensor configuration. Multiple sensors in the configuration file need to be separated by a comma. Following is a sample *SensorFile*:

```
sensors = {
 name = "MySensor"
 # sensor disabled by default
 period = 0
 type = "Generic"
}
```

The generic sensor and a sensor-specific configuration file need to be installed on all the nodes where the generic sensor is to be activated.

### **update**

Updates the existing configuration.

--collectors *CollectorNode[,CollectorNode...]* updates the collectors to be used by the sensors (see **config generate** for details).

--config-file *InputFile* specifies a template sensor configuration file to use. This overwrites the currently used configuration with the configuration specified in *InputFile*.

*Attribute=value ...* specifies a list of attribute value assignments. This sets the value of attribute *Attribute* to *value*.

### **delete**

Removes configuration of the Performance Monitoring tool or the specified sensors.

--sensors *Sensor[,Sensor...]* removes the sensors with the specified names from the performance monitoring configuration.

--all removes the entire performance monitoring configuration from IBM Spectrum Scale.

### **show**

Displays the currently active performance monitoring configuration. Specifies the following options:

--config-file *OutputFile* specifies that the output will be saved to the *OutputFile*. This option is optional.

### **query**

*Metric[,Metric...]* specifies a comma separated list of metrics for displaying in the output.

*Key[,Key...]* specifies a key that can consist of a node name, sensor group, or optional additional filters, and metrics that are separated by the pipe symbol (|). For example:

```
"cluster1.ibm.com|CTDBStats|locking|db_hop_count_bucket_00"
```

*NamedQuery* specifies the name of a predefined query.

*compareNodes* compares the specified metrics for all nodes in the system. The query creates one column per existing node and only one metric can be compared.

*ComparisonMetric* specifies the name of the metric to be compared when using the **compareNodes** query.

*StartTime* specifies the start timestamp for query in the YYYY-MM-DD-hh:mm:ss format.

*EndTime* specifies the end timestamp for query in the YYYY-MM-DD-hh:mm:ss format. If it is not specified, the query will return results until the present time.

*Duration* specifies the number of seconds into the past from present time or *EndTime*.

*Options* specifies the following options:

- -N or --Node *NODENAME* specifies the node from where the metrics should be retrieved.  
For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.
- --bucket-size *BUCKET\_SIZE* specifies the bucket size (number of seconds), default is 1.
- --number-buckets *NUMBER\_BUCKETS* specifies the number of buckets (records) to display, default is 10.
- --filter *FILTER* specifies the filter criteria for the query to run. To see the list of filters in the node use the **mmperfmon query --list filters** command.
- --format *FORMAT* specifies a common format for all columns.
- --csv provides the output in the CSV format.
- --raw provides the output in a raw format rather than a tabular format.
- --short displays the column names in a short form when there are too many to fit into a row.
- --nice displays the column headers in the output in a bold and underlined typeface.
- --resolve displays the resolved computed metrics and metrics that are used.
- --list {computed | metrics | keys | filters | queries | all} lists the following information:
  - computed displays the computed metrics.
  - metrics displays the metrics.
  - keys lists the keys.
  - filters lists the filters.
  - queries lists the available predefined queries.
  - all displays the computed metrics, metrics, keys, filters, and queries.

## Exit status

- 0 Successful completion.
- 1 Invalid arguments given
- 2 Invalid option
- 3 No node found with a running performance collector
- 4 Performance collector backend signaled bad query, for example, no data for this query.

## Security

You must have root authority to run the **mmperfmon** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For

## **mmpfmon**

more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. To generate configuration for the c89f8v03 collector node, issue the command:

```
mmpfmon config generate --collectors c89f8v03
```

The system displays output similar to this:

```
mmpfmon: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
Tue Oct 27 20:40:07 EDT 2015: mmcommon pushSdr_async: mmsdrfs propagation started
```

2. To add /tmp/SensorFile sensor to the Performance Monitoring tool issue the command:

```
mmpfmon config add --sensors /tmp/SensorFile
```

The system displays output similar to this:

```
mmpfmon: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
Tue Oct 27 20:44:33 EDT 2015: mmcommon pushSdr_async: mmsdrfs propagation started
mmpfmon config show | tail -12
{
 name = "NFSIO"
 period = 0
 proxyCmd = "/opt/IBM/zimon/GaneshaProxy"
 restrict = "cesNodes"
 type = "Generic"
},
{
 name = "TestAdd"
 period = 4
}
smbstat = ""
```

3. To update the NFSIO.period value to 5, issue the command:

```
mmpfmon config update NFSIO.period=5
```

The system displays output similar to this:

```
mmpfmon: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
Tue Oct 27 20:47:53 EDT 2015: mmcommon pushSdr_async: mmsdrfs propagation started
```

```
mmpfmon config show | tail -9
{
 name = "NFSIO"
 period = 5
 proxyCmd = "/opt/IBM/zimon/GaneshaProxy"
 restrict = "cesNodes"
 type = "Generic"
}
smbstat = ""
```

4. To remove the TestAdd sensor, issue the following command:

```
mmpfmon config delete --sensors TestAdd
```

The system displays output similar to this:

```
mmpfmon: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
Tue Oct 27 20:46:23 EDT 2015: mmcommon pushSdr_async: mmsdrfs propagation started
Tue Oct 27 20:46:28 EDT 2015: mmcommon pushSdr_async: mmsdrfs propagation completed; mmdsh rc=0
mmpfmon config show | tail -12
{
 name = "GPFSDiskCap"
 period = 0
```

```

},
{
 name = "NFSIO"
 period = 0
 proxyCmd = "/opt/IBM/zimon/GaneshaProxy"
 restrict = "cesNodes"
 type = "Generic"
}
smbstat = ""

```

5. To display the currently active performance monitoring configuration, issue the command:

```
mmperfmon config show
```

The system displays output similar to this:

```

cephMon = "/opt/IBM/zimon/CephMonProxy"
cephRados = "/opt/IBM/zimon/CephRadosProxy"
colCandidates = "c89f8v03"
colRedundancy = 1
collectors = {
 host = ""
 port = "4739"
}
config = "/opt/IBM/zimon/ZIMonSensors.cfg"
ctdbstat = ""
daemonize = T
hostname = ""
ipfixinterface = "0.0.0.0"
logfile = "/var/log/zimon/ZIMonSensors.log"
loglevel = "info"
mmcCmd = "/opt/IBM/zimon/MMCmProxy"
mmdfCmd = "/opt/IBM/zimon/MMDPProxy"
mmpmon = "/opt/IBM/zimon/MmpmonSockProxy"
piddir = "/var/run"
release = "4.2.0-0"
sensors = {
 name = "CPU"
 period = 1
},
{
 name = "Load"
 period = 1
},
{
 name = "Memory"
 period = 1
},
{
 name = "Network"
 period = 1
},
{
 name = "Netstat"
 period = 0
},
{
 name = "Diskstat"
 period = 0
},
{
 name = "DiskFree"
 period = 600
},
{
 name = "GPFSDisk"
 period = 0
},
{
 name = "GPFSFilesystem"

```

## mmpfmon

```
 period = 1
 },
 {
 name = "GPFSNSDDisk"
 period = 1
 restrict = "nsdNodes"
 },
 {
 name = "GPFSPoolIO"
 period = 0
 },
 {
 name = "GPFSVFS"
 period = 1
 },
 {
 name = "GPFSIOC"
 period = 0
 },
 {
 name = "GPFSVIO"
 period = 0
 },
 {
 name = "GPFSPDDisk"
 period = 1
 restrict = "nsdNodes"
 },
 {
 name = "GPFSvFLUSH"
 period = 0
 },
 {
 name = "GPFSNode"
 period = 1
 },
 {
 name = "GPFSNodeAPI"
 period = 1
 },
 {
 name = "GPFSFilesystemAPI"
 period = 1
 },
 {
 name = "GPFSLROC"
 period = 0
 },
 {
 name = "GPFSCHMS"
 period = 0
 },
 {
 name = "GPFSAFM"
 period = 0
 },
 {
 name = "GPFSAFMFS"
 period = 0
 },
 {
 name = "GPFSAFMFSET"
 period = 0
 },
 {
 name = "GPFSRPCS"
 period = 0
 }
```

```

},
{
 name = "GPFSFilesetQuota"
 period = 3600
},
{
 name = "GPFSDiskCap"
 period = 0
},
{
 name = "NFSIO"
 period = 0
 proxyCmd = "/opt/IBM/zimon/GaneshaProxy"
 restrict = "cesNodes"
 type = "Generic"
},
{
 name = "SwiftAccount"
 period = 1
 restrict = "cesNodes"
 type = "generic"
},
{
 name = "SwiftContainer"
 period = 1
 restrict = "cesNodes"
 type = "generic"
},
{
 name = "SwiftObject"
 period = 1
 restrict = "cesNodes"
 type = "generic"
},
{
 name = "SwiftProxy"
 period = 1
 restrict = "cesNodes"
 type = "generic"
}
smbstat = ""

```

6. To list metrics by key, for given node, sensor group and metrics, issue this command:

```
mmperfmon query "cluster1.ibm.com|CTDBDBStats|locking|db_hop_count_bucket_00"
```

The system displays output similar to this:

| Row | Timestamp           | db_hop_count_bucket_00 |
|-----|---------------------|------------------------|
| 1   | 2015-04-08-12:54:53 | 0                      |
| 2   | 2015-04-08-12:54:54 | 0                      |
| 3   | 2015-04-08-12:54:55 | 0                      |
| 4   | 2015-04-08-12:54:56 | 0                      |
| 5   | 2015-04-08-12:54:57 | 0                      |
| 6   | 2015-04-08-12:54:58 | 0                      |
| 7   | 2015-04-08-12:54:59 | 0                      |
| 8   | 2015-04-08-12:55:00 | 0                      |
| 9   | 2015-04-08-12:55:01 | 0                      |
| 10  | 2015-04-08-12:55:02 | 0                      |

7. To list the two metrics nfs\_read\_lat and nfs\_write\_lat for a specific time range, filtered by an export and NFS version with 60-seconds-buckets (one record represents 60 seconds), issue this command:

```
mmperfmon query nfs_read_lat,nfs_write_lat 2014-12-19-11:15:00 2014-12-19-11:20:00 --filter export=/ibm/gpfs/nfsexport,nfs_ver=NFSv3 -b 60
```

The system displays output similar to this:

## **mmpfmon**

| Row | Timestamp           | nfs_read_lat | nfs_write_lat |
|-----|---------------------|--------------|---------------|
| 1   | 2015-04-10-09:24:00 | 0            | 0             |
| 2   | 2015-04-10-09:24:10 | 0            | 0             |
| 3   | 2015-04-10-09:24:20 | 0            | 0             |
| 4   | 2015-04-10-09:24:30 | 0            | 0             |
| 5   | 2015-04-10-09:24:40 | 0            | 0             |
| 6   | 2015-04-10-09:24:50 | 0            | 0             |
| 7   | 2015-04-10-09:25:00 | 0            | 0             |
| 8   | 2015-04-10-09:25:10 | 0            | 0             |
| 9   | 2015-04-10-09:25:20 | 0            | 0             |
| 10  | 2015-04-10-09:25:30 | 0            | 0             |
| 11  | 2015-04-10-09:25:40 | 0            | 0             |
| 12  | 2015-04-10-09:25:50 | 45025738     | 1882453623    |
| 13  | 2015-04-10-09:26:00 | 0            | 0             |
| 14  | 2015-04-10-09:26:10 | 0            | 0             |
| 15  | 2015-04-10-09:26:20 | 0            | 0             |
| 16  | 2015-04-10-09:26:30 | 0            | 0             |
| 17  | 2015-04-10-09:26:40 | 0            | 0             |
| 18  | 2015-04-10-09:26:50 | 0            | 0             |

8. To list all available filters, issue this command:

```
mmpfmon query --list filters
```

The system displays output similar to this:

Available Filters:

```
node
 gpfss-21.localnet.com
 gpfss-22.localnet.com
protocol
 smb2
db_name
 account_policy
 autorid
 brlock
 ctdb
 dbwrap_watchers
 g_lock
 group_mapping
 leases
 locking
 netlogon_creds_cli
 notify_index
 passdb
 registry
 secrets
 serverid
 share_info
 smbXsrv_open_global
 smbXsrv_session_global
 smbXsrv_tcon_global
 smbXsrv_version_global
gpfss_fs_name
 fs0
 gpfss0
gpfss_cluster_name
 gpfss-cluster-2.localnet.com
mountPoint
 /
 /boot
 /dev
 /dev/shm
 /gpfss/fs0
 /mnt/gpfss0
 /run
 /sys/fs/cgroup
```

```

operation
 break
 cancel
 close
 create
 find
 flush
 getinfo
 ioctl
 keepalive
 lock
 logoff
 negprot
 notify
 read
 sesssetup
 setinfo
 tcon
 tdis
 write
sensor
 CPU
 CTDBDBStats
 CTDBStats
 DiskFree
 GPFSFilesystemAPI
 GPFSVFS
 Load
 Memory
 Network
 SMBGlobalStats
 SMBStats
netdev_name
 eth0
 lo

```

9. To run a named query for export /ibm/gpfs/nfsexport and **nfs\_ver NFSv3**, using default bucket size of 1 second, showing last 10 buckets , issue this command:

```
mmperfmon query nfsIOrate --filter export=/ibm/gpfs/nfsexport,nfs_ver=NFSv3,node=cluster1.ibm.com
```

The system displays output similar to this:

Legend:

|    |                                                                |
|----|----------------------------------------------------------------|
| 1: | cluster1.ibm.com NFSIO /ibm/gpfs/nfsexport NFSv3 nfs_read_ops  |
| 2: | cluster2.ibm.com NFSIO /ibm/gpfs/nfsexport NFSv3 nfs_write_ops |

| Row | Timestamp           | nfs_read_ops | nfs_write_ops |
|-----|---------------------|--------------|---------------|
| 1   | 2015-05-11-13:32:57 | 0            | 0             |
| 2   | 2015-05-11-13:32:58 | 90           | 90            |
| 3   | 2015-05-11-13:32:59 | 90           | 90            |
| 4   | 2015-05-11-13:33:00 | 90           | 91            |
| 5   | 2015-05-11-13:33:01 | 91           | 90            |
| 6   | 2015-05-11-13:33:02 | 91           | 92            |
| 7   | 2015-05-11-13:33:03 | 89           | 88            |
| 8   | 2015-05-11-13:33:04 | 91           | 92            |
| 9   | 2015-05-11-13:33:05 | 93           | 92            |
| 10  | 2015-05-11-13:33:06 | 89           | 89            |

10. To run a named query for export /ibm/gpfs/nfsexport and **nfs\_ver NFSv3**, using bucket size of 1 minute, showing last 20 buckets (= 20 minutes), issue this command:

```
mmperfmon query nfsIOrate --filter export=/ibm/gpfs/nfsexport,nfs_ver=NFSv3,node=cluster1.ibm.com -n 20 -b 60
```

The system displays output similar to this:

Legend:

|    |                                                                |
|----|----------------------------------------------------------------|
| 1: | cluster1.ibm.com NFSIO /ibm/gpfs/nfsexport NFSv3 nfs_read_ops  |
| 2: | cluster2.ibm.com NFSIO /ibm/gpfs/nfsexport NFSv3 nfs_write_ops |

## mmpfmon

| Row | Timestamp           | nfs_read_ops | nfs_write_ops |
|-----|---------------------|--------------|---------------|
| 1   | 2015-05-11-13:31:00 | 0            | 0             |
| 2   | 2015-05-11-13:32:00 | 280          | 280           |
| 3   | 2015-05-11-13:33:00 | 820          | 820           |
| 4   | 2015-05-11-13:34:00 | 0            | 0             |
| 5   | 2015-05-11-13:35:00 | 0            | 0             |
| 6   | 2015-05-11-13:36:00 | 0            | 0             |
| 7   | 2015-05-11-13:37:00 | 0            | 0             |
| 8   | 2015-05-11-13:38:00 | 0            | 0             |
| 9   | 2015-05-11-13:39:00 | 1000         | 1000          |
| 10  | 2015-05-11-13:40:00 | 1000         | 1000          |
| 11  | 2015-05-11-13:41:00 | 0            | 0             |
| 12  | 2015-05-11-13:42:00 | 0            | 0             |
| 13  | 2015-05-11-13:43:00 | 0            | 0             |
| 14  | 2015-05-11-13:44:00 | 2000         | 2000          |
| 15  | 2015-05-11-13:45:00 | 0            | 0             |
| 16  | 2015-05-11-13:46:00 | 0            | 0             |
| 17  | 2015-05-11-13:47:00 | 1000         | 1000          |
| 18  | 2015-05-11-13:48:00 | 1000         | 1000          |
| 19  | 2015-05-11-13:49:00 | 0            | 0             |
| 20  | 2015-05-11-13:50:00 | 0            | 0             |

11. To run a **compareNodes** query for the *cpu\_user* metric, issue this command:

```
mmpfmon query compareNodes cpu_user
```

The system displays output similar to this:

Legend:

```
1: cluster1.ibm.com|CPU|cpu_user
2: cluster2.ibm.com|CPU|cpu_user
```

| Row | Timestamp           | cluster1 | cluster2 |
|-----|---------------------|----------|----------|
| 1   | 2015-05-11-13:53:54 | 0.5      | 0.25     |
| 2   | 2015-05-11-13:53:55 | 0.5      | 0.25     |
| 3   | 2015-05-11-13:53:56 | 0.25     | 0.25     |
| 4   | 2015-05-11-13:53:57 | 0.5      | 0.25     |
| 5   | 2015-05-11-13:53:58 | 0.25     | 0.75     |
| 6   | 2015-05-11-13:53:59 | 0.5      | 0.25     |
| 7   | 2015-05-11-13:54:00 | 0.25     | 0.25     |
| 8   | 2015-05-11-13:54:01 | 0.5      | 0.25     |
| 9   | 2015-05-11-13:54:02 | 0.25     | 0.25     |
| 10  | 2015-05-11-13:54:03 | 0.5      | 0.25     |

12. To run an object query, issue the following command:

```
mmpfmon query objObj 2016-09-28-09:56:39 2016-09-28-09:56:43
```

The system displays output similar to this:

```
1: cluster1.ibm.com|SwiftObject|object_auditor_time
2: cluster1.ibm.com|SwiftObject|object_expirer_time
3: cluster1.ibm.com|SwiftObject|object_replication_partition_delete_time
4: cluster1.ibm.com|SwiftObject|object_replication_partition_update_time
5: cluster1.ibm.com|SwiftObject|object_DEL_time
6: cluster1.ibm.com|SwiftObject|object_DEL_err_time
7: cluster1.ibm.com|SwiftObject|object_GET_time
8: cluster1.ibm.com|SwiftObject|object_GET_err_time
9: cluster1.ibm.com|SwiftObject|object_HEAD_time
10: cluster1.ibm.com|SwiftObject|object_HEAD_err_time
11: cluster1.ibm.com|SwiftObject|object_POST_time
12: cluster1.ibm.com|SwiftObject|object_POST_err_time
13: cluster1.ibm.com|SwiftObject|object_PUT_time
14: cluster1.ibm.com|SwiftObject|object_PUT_err_time
15: cluster1.ibm.com|SwiftObject|object_REPLICATE_time
16: cluster1.ibm.com|SwiftObject|object_REPLICATE_err_time
17: cluster1.ibm.com|SwiftObject|object_updater_time

Row object_auditor_time object_expirer_time object_replication_partition_delete_time
object_replication_partition_update_time object_DEL_time object_DEL_err_time object_GET_time
object_GET_err_time object_HEAD_time object_HEAD_err_time object_POST_time object_POST_err_time
object_PUT_time object_PUT_err_time object_REPLICATE_time object_REPLICATE_err_time object_updater_time
1 2016-09-28 09:56:39 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.855923
0.000000 0.000000 0.000000 45.337915 0.000000 0.000000 0.000000 0.000000 0.000000
```

```

2 2016-09-28 09:56:40 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
3 2016-09-28 09:56:41 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.931925
0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
4 2016-09-28 09:56:42 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.855923
0.000000 0.000000 0.000000 516.280890 0.000000 0.000000 0.000000 0.000000

```

```

object_DEL_total_time = 0.0 object_PUT_total_time = 561.618805
object_GET_total_time = 0.0 object_POST_total_time = 0.0
object_HEAD_total_time = 1.786948 object_PUT_max_time = 516.28089
object_POST_max_time = 0.0 object_GET_max_time = 0.0
object_HEAD_max_time = 0.931025 object_DEL_max_time = 0.0
object_GET_avg_time = 0.0 object_DEL_avg_time = 0.0
object_PUT_avg_time = 280.809402 object_POST_avg_time = 0.0
object_HEAD_avg_time = 0.893474 object_DEL_time_count = 0.0
object_POST_time_count = 0 object_PUT_time_count = 2
object_HEAD_time_count = 2 object_GET_time_count = 0
object_DEL_min_time = 0.0 object_PUT_min_time = 45.337915
object_GET_min_time = 0.0 object_POST_min_time = 0.0
object_HEAD_min_time = 0.855923

```

## See also

- “mmdumpperfdata command”

## Location

/usr/lpp/mmfs/bin

## mmpmon

---

### mmpmon command

Manages performance monitoring and displays performance information.

#### Synopsis

```
mmpmon [-i CommandFile] [-d IntegerDelayValue] [-p]
 [-r IntegerRepeatValue] [-s] [-t IntegerTimeoutValue]
```

#### Availability

Available on all IBM Spectrum Scale editions.

#### Description

Before you attempt to use **mmpmon**, it is a good idea to review this command entry, then read the entire topic *Monitoring GPFS I/O performance with the mmpmon command* in the *IBM Spectrum Scale: Administration Guide*.

Use the **mmpmon** command to manage GPFS performance monitoring functions and display performance monitoring data. The **mmpmon** command reads requests from an input file or standard input (stdin), and writes responses to standard output (stdout). Error messages go to standard error (stderr). Prompts, if not suppressed, go to stderr.

You can run **mmpmon** so that it continually reads input from a pipe. That is, the driving script or application never sends an end of file. In this scenario, it is a good idea to set the **-r** option to 1, or to use the default value of 1. This setting prevents the command from caching input records. In doing so it avoids unnecessary memory consumption.

This command cannot be run from a Windows node.

#### Results

The performance monitoring request is sent to the GPFS daemon that is running on the same node that is running the **mmpmon** command.

All results from the request are written to stdout.

The command has two output formats:

- Human readable, intended for direct viewing.

In this format, the results are keywords that describe the value presented, followed by the value. Here is an example:

disks: 2

- Machine readable, an easily parsed format intended for further analysis by scripts or applications.

In this format, the results are strings with values presented as keyword/value pairs. The keywords are delimited by underscores (\_) and blanks to make them easier to locate.

For details on how to interpret the **mmpmon** command results, see the topic *Monitoring GPFS I/O performance with the mmpmon command* in the *IBM Spectrum Scale: Administration Guide*.

#### Parameters

##### **-i** *CommandFile*

The input file contains **mmpmon** command requests, one per line. Use of the **-i** flag implies use of

the **-s** flag. For interactive use, omit the **-i** flag. In this case, the input is then read from stdin, allowing **mmpmon** to take keyboard input or output piped from a user script or application program.

Leading blanks in the input file are ignored. A line beginning with a number sign (#) is treated as a comment. Leading blanks in a line whose first non-blank character is a number sign (#) are ignored.

The input requests to the **mmpmon** command are as follows:

**fs\_io\_s**

Displays I/O statistics per mounted file system.

**io\_s**

Displays I/O statistics for the entire node.

**nlist add name [name...]**

Adds node names to a list of nodes for **mmpmon** processing.

**nlist del**

Deletes a node list.

**nlist new name [name...]**

Creates a node list.

**nlist s**

Shows the contents of the current node list.

**nlist sub name [name...]**

Deletes node names from a list of nodes for **mmpmon** processing.

**once request**

Indicates that the request is to be performed only once.

**qosio**

Displays statistics for Quality of Service for I/O operations (QoS).

**reset**

Resets statistics to zero.

**rhist nr**

Changes the request histogram facility request size and latency ranges.

**rhist off**

Disables the request histogram facility. This value is the default.

**rhist on**

Enables the request histogram facility.

**rhist p**

Displays the request histogram facility pattern.

**rhist reset**

Resets the request histogram facility data to zero.

**rhist s**

Displays the request histogram facility statistics values.

**rpc\_s**

Displays the aggregation of execution time for remote procedure calls (RPCs).

**rpc\_s size**

Displays the RPC execution time according to the size of messages.

**ver**

Displays **mmpmon** version.

## **mmpmon**

**vio\_s [f rg RecoveryGroupName [da DeclusteredArray [v Vdisk]]] [reset]**

Displays IBM Spectrum Scale RAID vdisk I/O statistics. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

**vio\_s\_reset [f rg RecoveryGroupName [da DeclusteredArray [v Vdisk]]]**

Resets IBM Spectrum Scale RAID vdisk I/O statistics. For more information about IBM Spectrum Scale RAID, see *IBM Spectrum Scale RAID: Administration*.

## **Options**

**-d IntegerDelayValue**

Specifies a number of milliseconds to sleep after one invocation of all the requests in the input file. The default value is 1000. This value must be an integer greater than or equal to 500 and less than or equal to 8000000.

The command processes the input file in the following way:

1. The command processes each request in the input file sequentially. It reads a request, processes it, sends it to the GPFS daemon, and waits for the response. When it receives the response, the command processes it and displays the results of the request. The command then goes on to the next request in the input file.
2. When the command processes all the requests in the input file, it sleeps for the specified number of milliseconds.
3. When the command wakes, it begins another cycle of processing, beginning with Step 1. The number of repetitions depends on the value of the **-r** flag.

**-p** Indicates to generate output that can be parsed by a script or program. If this option is not specified, human-readable output is produced.

**-r IntegerRepeatValue**

Specifies the number of times to run all the requests in the input file.

The default value is one. Specify an integer between zero and 8000000. Zero means to run forever, in which case processing continues until it is interrupted. This feature is used, for example, by a driving script or application program that repeatedly reads the result from a pipe.

The **once** prefix directive can be used to override the **-r** flag. See the description of **once** in the topic *Monitoring GPFS I/O performance with the mmpmon command* in the *IBM Spectrum Scale: Administration Guide*.

**-s** Indicates to suppress the prompt on input.

Use of the **-i** flag implies use of the **-s** flag. For use in a pipe or with redirected input (<), the **-s** flag is preferred. If not suppressed, the prompts go to standard error (stderr).

**-t IntegerTimeoutValue**

Specifies a number of seconds that the command waits for responses from the GPFS daemon before it fails.

The default value is 60. This value must be an integer greater than or equal to 1 and less than or equal to 8000000.

## **Exit status**

- 0 Successful completion.
- 1 Various errors, including insufficient memory, input file not found, incorrect option, and others.
- 3 Either no commands were entered interactively, or the input file did not contain any **mmpmon** commands. The input file was empty, or consisted of all blanks or comments.
- 4 **mmpmon** terminated due to a request that was not valid.
- 5 An internal error occurred.

**111** An internal error occurred. A message follows.

## Restrictions

- Up to five instances of **mmpmon** can be run on a node concurrently. However, concurrent users might interfere with each other. See the topic *Monitoring GPFS I/O performance with the mmpmon command* in the *IBM Spectrum Scale: Administration Guide*.
- Do not alter the input file while **mmpmon** is running.
- The input file must contain valid input requests, one per line. When **mmpmon** finds an invalid request, it issues an error message and terminates. The command processes input requests that appear in the input file before the first invalid request.

## Security

The **mmpmon** command must be run by a user with root authority, on the node for which you want statistics.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

- Assume that **infile** contains these requests:

```
ver
io_s
fs_io_s
rhist off
```

The following command is issued:

```
mmpmon -i infile -r 10 -d 5000
```

The output (sent to stdout) is similar to the following example output:

```
mmpmon node 192.168.1.8 name node1 version 3.1.0
mmpmon node 192.168.1.8 name node1 io_s OK
timestamp: 1083350358/935524
bytes read: 0
bytes written: 0
opens: 0
closes: 0
reads: 0
writes: 0
readdir: 0
inode updates: 0
mmpmon node 192.168.1.8 name node1 fs_io_s status 1
no file systems mounted
mmpmon node 192.168.1.8 name node1 rhist off OK
```

The requests in the input file are run 10 times, with a delay of 5000 milliseconds (5 seconds) between invocations.

- This example uses the same parameters as the previous example, but with the **-p** flag:

```
mmpmon -i infile -p -r 10 -d 5000
```

The output (sent to stdout) is similar to the following example output:

```
_ver_n_192.168.1.8_nn_node1_v_2_lv_3_vt_0
-io_s_n_192.168.1.8_nn_node1_rc_0_t_1084195701_tu_350714_br_0_bw_0_oc_0
_cc_0_rdc_0_wc_0_dir_0_iu_0
-fs_io_s_n_192.168.1.8_nn_node1_rc_1_t_1084195701_tu_364489_cl_-fs_-rhist_
_n_192.168.1.8_nn_node1_req_off_rc_0_t_1084195701_tu_378217
```

## mmpmon

3. This example uses the **fs\_io\_s** option with a mounted file system:

```
mmpmon node 198.168.1.8 name node1 fs_io_s OK
cluster: node1.localdomain
filesystem: gpfs1
disks: 1
timestamp: 1093352136/799285
bytes read: 52428800
bytes written: 87031808
opens: 6
closes: 4
reads: 51
writes: 83
readdir: 0
inode updates: 11

mmpmon node 198.168.1.8 name node1 fs_io_s OK
cluster: node1.localdomain
filesystem: gpfs2
disks: 2
timestamp: 1093352136/799285
bytes read: 87031808
bytes written: 52428800
opens: 4
closes: 3
reads: 12834
writes: 50
readdir: 0
inode updates: 9
```

4. This example is the same as the previous one, but with the **-p** flag:

```
_fs_io_s_n_198.168.1.8_nn_node1_rc_0_t_1093352061_tu_93867_cl_node1.localdomain
_fs_gpfs1_d_1_br_52428800_bw_87031808_oc_6_cc_4_rdc_51_wc_83_dir_0_iu_10
_fs_io_s_n_198.168.1.8_nn_node1_rc_0_t_1093352061_tu_93867_cl_node1.localdomain
_fs_gpfs2_d_2_br_87031808_bw_52428800_oc_4_cc_3_rdc_12834_wc_50_dir_0_iu_8
```

This output consists of two strings.

5. This example uses the **io\_s** option with a mounted file system:

```
mmpmon node 198.168.1.8 name node1 io_s OK
timestamp: 1093351951/587570
bytes read: 139460608
bytes written: 139460608
opens: 10
closes: 7
reads: 12885
writes: 133
readdir: 0
inode updates: 14
```

6. This example is the same as the previous one, but with the **-p** flag:

```
_io_s_n_198.168.1.8_nn_node1_rc_0_t_1093351982_tu_356420_br_139460608
_bw_139460608_oc_10_cc_7_rdc_0_wc_133_dir_0_iu_14
```

This output consists of one string.

## Location

/usr/lpp/mmfs/bin

---

## mmprotocoltrace command

Starts, stops, and monitors tracing for the CES protocols.

### Synopsis

```
mmprotocoltrace start <identifier> [<identifier>...] [-c <clientIP>]
[-d <duration>] [-l <logFileDir>] [-n <nodes>] [-f]
```

or

```
mmprotocoltrace stop <identifier> [<identifier>...]
```

or

```
mmprotocoltrace status <identifier> [<identifier>...] [-v]
```

or

```
mmprotocoltrace clear <identifier> [<identifier>...] [-f]
```

or

```
mmprotocoltrace reset <identifier> [<identifier>...]
```

or

```
mmprotocoltrace {config| check}
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher.

**Notice:** This command has common function to other existing commands. As such, the function may, at any time in a future release, be rolled into other commands and immediately deprecated from use without prior notice. Information about the change and what commands replace it would be provided in some format at the time of that change. Users should avoid using this command in any type of automation or scripting or be advised a future change may break that automation without prior notice.

### Description

Use the **mmprotocoltrace** command to trace Winbind, NFS, SMB, Object, or network operations. You can start, stop, reset, check or, display the status of a trace with this command. It also controls the timeouts for the traces to avoid excessive logging.

**Note:** The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms.

For more information about this command, see the topic *CES tracing and debug data collection* in the *IBM Spectrum Scale: Problem Determination Guide*.

### Parameters

#### *options*

Specifies one of the following trace options:

##### **-d Duration**

Specifies the trace duration in minutes. The default is **10**.

## **mmprotocoltrace**

### **-l LogFileDir**

Specifies the name of the directory that contains the log and tar files that are created by the trace. The directory name cannot be a shared directory. The default is a directory in /tmp/mmfs that is named by the trace type and time.

### **-N Nodes**

Specifies a comma-separated list of names of the CES nodes where you want tracing to be done. The default is all the CES nodes. For more information, see the topic *Tips for using mmprotocoltrace* in the *IBM Spectrum Scale: Problem Determination Guide*.

### **-c ClientIPs**

Specifies a comma-separated list of client IP addresses to trace. The CES nodes that you specified in the -N parameter will trace their connections with these clients. This parameter applies only to SMB traces and Network traces. For more information, see the topic *Tips for using mmprotocoltrace* in the *IBM Spectrum Scale: Problem Determination Guide*.

### **-f**

Forces an action to occur. Affects the **clear** command. This parameter also disables prompt for smb and smbsyscalls.

### **-v**

Verbose output. Affects only the **status** command.

### *command*

Specifies one of the following trace commands:

#### **start**

Starts tracing for the specified component.

#### **stop**

Stops tracing for the specified component.

#### **status**

Displays the status of the specified component.

#### **check**

Checks and performs all planned tracing actions specified on the node.

#### **config**

Displays the protocol tracing configuration settings that are currently active for the node on which the command is run. For example, the limit for the maximal trace size. These settings are defined in the file /var/mmfs/ces/mmprotocoltrace.conf and are specific for each CES node.

#### **clear**

Clears the trace records from the trace list.

#### **reset**

Resets the nodes to the default state that is defined in the configuration file.

### *identifier*

Specifies one of the following components:

#### **nfs**

Traces the NFS service.

#### **smb**

Traces the SMB service. This enables a detailed tracing for incoming SMB connections from the specified SMB client IP addresses. When using this trace, it is recommended to only have few incoming connections from the specified IP addresses, as too many traced connections can impact the cluster performance.

#### **object**

Traces the Object service.

**network**

Traces the Network service.

**smbsyscalls**

Collects the strace-logs for SMB. This collects a trace for all system calls issued from SMB connections from the specified SMB client IP addresses. When using this trace it is recommended to only have few incoming connections from the specified IP addresses, as too many traced connections can impact the cluster performance.

**winbind**

Traces the winbind service that is used for user authentication.

**Exit status**

**0** Successful completion.

**Nonzero**

A failure occurred.

**Security**

You must have root authority to run the **mmprotocoltrace** command.

The node on which the command is run must be able to process remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. See the information about the requirements for administering a GPFS system in the *IBM Spectrum Scale: Administration Guide*.

**Examples**

1. To start an SMB trace, issue this command:

```
mmprotocoltrace start smb -c <clientIP>
```

The system displays output similar to this:

```
Trace 'fcb7cb07-c45e-43f8-8f1f-2de50cf15062' created successfully for 'smb'
```

2. To view the status of the SMB trace, issue this command:

```
mmprotocoltrace status smb
```

The system displays output similar to this:

```
Trace ID: fcb7cb07-c45e-43f8-8f1f-2de50cf15062
State: Active
User ID: root
Protocol: smb
Start Time: 10:57:43 04/03/2016
End Time: 11:07:43 04/03/2016
Client IPs: 10.0.100.42, 10.0.100.43
Origin Node: ch-42.localnet.com
Syscall: False
Syscall Only: False
Nodes:
 Node Name: ch-41.localnet.com
 State: ACTIVE
 Trace Location: /tmp/mmfs/smb.20160304_105742.trc

 Node Name: ch-42.localnet.com
 State: ACTIVE
 Trace Location: /tmp/mmfs/smb.20160304_105742.trc

 Node Name: ch-43.localnet.com
 State: ACTIVE
 Trace Location: /tmp/mmfs/smb.20160304_105742.trc
```

## **mmprotoctrace**

3. To stop the SMB trace, issue this command:

```
mmprotoctrace stop smb
```

The system displays output similar to this:

```
Stopping traces
Trace 'fcb7cb07-c45e-43f8-8f1f-2de50cf15062' stopped for smb
Waiting for traces to complete
Waiting for node 'node1'
Waiting for node 'node2'
Waiting for node 'node3'
Finishing trace 'fcb7cb07-c45e-43f8-8f1f-2de50cf15062'
Successfully copied file from 'node1:/tmp/mmfs/smb.20160304_105742.trc'
Successfully copied file from 'node2:/tmp/mmfs/smb.20160304_105742.trc'
Successfully copied file from 'node3:/tmp/mmfs/smb.20160304_105742.trc'
Trace tar file has been written to '/tmp/mmfs/smb.trace.20160304_105845.tar.gz'
```

4. To clear the SMB trace from the trace file, issue this command:

```
mmprotoctrace clear smb
```

The system displays output similar to this:

```
All traces cleared successfully
```

5. To trace the system calls for SMB, issue this command:

```
mmprotoctrace smbsyscalls -c <clientIP>
```

The system displays output similar to this:

```
Trace '9cd534c9-be3c-4478-ba45-2e00acd4b544' created successfully for 'smb'
```

## **See also**

- “mmaddcallback command” on page 10
- “mmces command” on page 103
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

## **Location**

/usr/lpp/mmfs/bin

---

## mmmsnap command

Creates or deletes identical snapshots on the cache and home clusters, or shows the status of snapshots that have been queued up on the gateway nodes.

### Synopsis

```
mmmsnap Device create -j FilesetName {[--comment Comment] [--uid ClusterUID] | --rpo] [--wait]
```

or

```
mmmsnap Device delete -s SnapshotName -j FilesetName
```

or

```
mmmsnap Device status -j FilesetName
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

### Description

The **mmmsnap** command creates or deletes identical snapshots on the cache and home clusters, or shows the status of snapshots that have been queued up on the gateway nodes. You can use this command only in a Single writer (SW) cache. You cannot use for Read only (RO), Independent writer (IW), or Local updates (LU) caches. Peer snapshots are not allowed on a Single writer (SW) cache that uses the NSD protocol for communicating with home.

### Parameters

#### *Device*

Specifies the device name of the file system.

#### **create**

Creates a fileset level snapshot in cache and a snapshot with the same name at home. The snapshot at home could be fileset level or file system level, depending on whether the exported path is an independent fileset or file system.

#### **-j FilesetName**

Specifies the name of the fileset.

#### **--comment Comment**

Optional; specifies user-defined text to be prepended to the snapshot name (thereby customizing the name of the snapshot).

#### **--uid ClusterUID**

Optional; specifies a unique identifier for the cache site. If not specified, this defaults to the GPFS cluster ID.

#### **--rpo**

Optional; specifies that user recovery point objective (RPO) snapshots are to be created for a primary fileset. This option cannot be specified with the **--comment** and **--uid** options.

#### **--wait**

Optional; makes the creation of cache and home snapshots a synchronous process. When specified, **mmmsnap** does not return until the snapshot is created on the home cluster. When not specified, **mmmsnap** is asynchronous and returns immediately rather than waiting for the snapshot to be created at home.

## **mmpsnap**

### **delete**

Deletes the local and remote copies of the specified snapshot; AFM automatically figures out the remote device and fileset.

#### **-s SnapshotName**

Specifies the name of the snapshot to be deleted. A snapshot name is constructed as follows:

{commentString}-psnap-{clusterId}-{fsUID}-{fsetID}-{timestamp}

Where *timestamp* has the form YY-MM-DD-HH-MM-SS.

In the following example, a comment string was not provided:

```
psnap-14133737607146558608-C0A8AA04:4EDD34DF-1-11-12-05-14-32-10
```

### **status**

Shows the status of snapshots that have been queued up on the gateway nodes. The status includes the following pieces of information:

- Last successful snapshot (obtained through **mmlsfileset --afm**)
- Status of the current snapshot process.

## **Exit status**

0      Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmpsnap** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To create a fileset level snapshot in cache of a single-writer fileset called **sw** in file system **fs1** issue this command:

```
mmpsnap fs1 create -j sw
```

The system displays output similar to the following:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
Snapshot psnap-13882361812785649740-C0A80E85:4F44B305-59-12-03-01-02-27-28 created with id 8.
Snapshot psnap-13882361812785649740-C0A80E85:4F44B305-59-12-03-01-02-27-28 created at the satellite.
Core snapshot has been queued.
```

2. To display the snapshot issue this command:

```
mmlssnapshot fs1 -j sw
```

The system displays output similar to the following:

```
Snapshots in file system fs1:
Directory SnapId Status Created Fileset
psnap-13882361812785649740-C0A80E85:4F44B305-59-12-03-01-02-27-28 8 Valid Thu Mar 1 02:27:29 2012 sw
```

3. To show that the snapshot is also created at home, issue this command:

```
mmlssnapshot fs1
```

The system displays output similar to the following:

```
Snapshots in file system fs1:
Directory SnapId Status Created Fileset
psnap-13882361812785649740-C0A80E85:4F44B305-59-12-03-01-02-27-28 8 Valid Thu Mar 1 02:23:16 2012
```

## See also

- “mmafmconfig command” on page 37
- “mmafmctl command” on page 40
- “mmafmlocal command” on page 54
- “mmchattr command” on page 122
- “mmchconfig command” on page 132
- “mmchfileset command” on page 173
- “mmchfs command” on page 179
- “mmcrfileset command” on page 244
- “mmcrfs command” on page 250
- “mmlsconfig command” on page 400
- “mmlsfileset command” on page 406
- “mmlsfs command” on page 411

## Location

/usr/lpp/mmfs/bin

## **mmpputacl**

---

### **mmpputacl command**

Sets the GPFS access control list for the specified file or directory.

#### **Synopsis**

**mmpputacl [-d] [-i Infilename] Filename**

#### **Availability**

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

#### **Description**

Use the **mmpputacl** command to set the ACL of a file or directory.

If the **-i** option is not used, the command expects the input to be supplied through standard input, and waits for your response to the prompt.

For information about NFS V4 ACLs, see the topic *Managing GPFS access control lists* in the *IBM Spectrum Scale: Administration Guide*.

Any output from the **mmpgetacl** command can be used as input to **mmpputacl**. The command is extended to support NFS V4 ACLs. In the case of NFS V4 ACLs, there is no concept of a default ACL. Instead, there is a single ACL and the individual access control entries can be flagged as being inherited (either by files, directories, both, or neither). Consequently, specifying the **-d** flag for an NFS V4 ACL is an error. By its nature, storing an NFS V4 ACL implies changing the inheritable entries (the GPFS default ACL) as well.

The following describes how **mmpputacl** works for POSIX and NFS V4 ACLs:

| Command             | POSIX ACL                                       | NFS V4 ACL                               |
|---------------------|-------------------------------------------------|------------------------------------------|
| <b>mmpputacl</b>    | Access ACL (Error if default ACL is NFS V4 [1]) | Stores the ACL (implies default as well) |
| <b>mmpputacl -d</b> | Default ACL (Error if access ACL is NFS V4 [1]) | Error: NFS V4 ACL (has no default ACL)   |

[1] The default and access ACLs are not permitted to be mixed types because NFS V4 ACLs include inherited entries, which are the equivalent of a default ACL. An **mmdelacl** of the NFS V4 ACL is required before an ACL is converted back to POSIX.

Depending on the file system's **-k** setting (**posix**, **nfs4**, or **all**), **mmpputacl** may be restricted. The **mmpputacl** command is not allowed to store an NFS V4 ACL if **-k posix** is in effect. The **mmpputacl** command is not allowed to store a POSIX ACL if **-k nfs4** is in effect. For more information, see the description of the **-k** flag for the **mmchfs**, **mmcrlfs**, and **mmlsfs** commands.

Note that the test to see if the given ACL is acceptable based on the file system's **-k** setting cannot be done until after the ACL is provided. For example, if **mmpputacl file1** is issued (no **-i** flag specified) the user then has to input the ACL before the command can verify that it is an appropriate ACL given the file system settings. Likewise, the command **mmpputacl -d dir1** (again the ACL was not given with the **-i** flag) requires that the ACL be entered before file system ACL settings can be tested. In this situation, the **-i** flag may be preferable to manually entering a long ACL, only to find out it is not allowed by the file system.

## Parameters

### *Filename*

The path name of the file or directory for which the ACL is to be set. If the **-d** option is specified, *Filename* must be the name of a directory.

## Options

**-d** Specifies that the default ACL of a directory is to be set. This flag cannot be used on an NFS V4 ACL.

### **-i** *InFilename*

The path name of a source file from which the ACL is to be read.

## Exit status

0      Successful completion.

### **nonzero**

A failure has occurred.

## Security

You may issue the **mmpputacl** command only from a node in the GPFS cluster where the file system is mounted.

You must be the file or directory owner, the root user, or someone with control permission in the ACL, to run the **mmpputacl** command.

## Examples

To use the entries in a file named **standard.acl** to set the ACL for a file named **project2.history**, issue this command:

```
mmpputacl -i standard.acl project2.history
```

where **standard.acl** contains:

```
user::rwxc
group::rwx-
other::--x-
mask::rw-c
user:alpha:rwxc
group:audit:rwx-
group:system:-w--
```

To confirm the change, issue this command:

```
mmgetacl project.history
```

The system displays information similar to:

```
#owner:paul
#group:design
user::rwxc
group::rwx-
other::--x-
mask::rw-c
user:alpha:rwxc
group:audit:rwx-
group:system:-w--
```

## See also

- “**mmeditacl** command” on page 322
- “**mmdelacl** command” on page 284

## **mmpputacl**

- “mmgetacl command” on page 347

## **Location**

/usr/lpp/mmfs/bin

---

## mmquotaoff command

Deactivates quota limit checking.

### Synopsis

```
mmquotaoff [-u] [-g] [-j] [-v] {Device [Device ...]} | -a}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmquotaoff** command disables quota limit checking by GPFS.

If none of: **-u**, **-j** or **-g** is specified, the **mmquotaoff** command deactivates quota limit checking for users, groups, and filesets.

If the **-a** option is not specified, *Device* must be the last parameter entered.

### Parameters

*Device* [*Device* ...]

The device name of the file system to have quota limit checking deactivated.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

### Options

- a** Deactivates quota limit checking for all GPFS file systems in the cluster. When used in combination with the **-g** option, only group quota limit checking is deactivated. When used in combination with the **-u** or **-j** options, only user or fileset quota limit checking, respectively, is deactivated.
- g** Specifies that only group quota limit checking is to be deactivated.
- j** Specifies that only quota checking for filesets is to be deactivated.
- u** Specifies that only user quota limit checking is to be deactivated.
- v** Prints a message for each file system in which quotas are deactivated.

### Exit status

0      Successful completion.

#### nonzero

A failure has occurred.

### Security

You must have root authority to run the **mmquotaoff** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmquotaoff** command is issued.

## **mmquotaoff**

### **Examples**

1. To deactivate user quota limit checking on file system **fs0**, issue this command:

```
mmquotaoff -u fs0
```

To confirm the change, issue this command:

```
mmlsfs fs0 -Q
```

The system displays information similar to:

| flag | value         | description     |
|------|---------------|-----------------|
| -Q   | group;fileset | Quotas enforced |

2. To deactivate group quota limit checking on all file systems, issue this command:

```
mmquotaoff -g -a
```

To confirm the change, individually for each file system, issue this command:

```
mmlsfs fs2 -Q
```

The system displays information similar to:

| flag | value        | description     |
|------|--------------|-----------------|
| -Q   | user;fileset | Quotas enforced |

3. To deactivate all quota limit checking on file system **fs0**, issue this command:

```
mmquotaoff fs0
```

To confirm the change, issue this command:

```
mmlsfs fs0 -Q
```

The system displays information similar to:

| flag | value | description     |
|------|-------|-----------------|
| -Q   | none  | Quotas enforced |

### **See also**

- “[mmcheckquota command](#)” on page 169
- “[mmdefedquota command](#)” on page 272
- “[mmdefquotaoff command](#)” on page 275
- “[mmdefquotaon command](#)” on page 278
- “[mmedquota command](#)” on page 325
- “[mmlsquota command](#)” on page 435
- “[mmquotaon command](#)” on page 509
- “[mmrepquota command](#)” on page 517

### **Location**

/usr/lpp/mmfs/bin

---

## mmquotaon command

Activates quota limit checking.

### Synopsis

**mmquotaon** [-u] [-g] [-j] [-v] {Device [Device...]} | -a

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmquotaon** command enables quota limit checking by GPFS.

If none of: **-u**, **-j** or **-g** is specified, the **mmquotaon** command activates quota limit checking for users, groups, and filesets.

If the **-a** option is not used, *Device* must be the last parameter specified.

After quota limit checking has been activated by issuing the **mmquotaon** command, issue the **mmcheckquota** command to count inode and space usage.

### Parameters

*Device* [Device...]

The device name of the file system to have quota limit checking activated.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

### Options

- a** Activates quota limit checking for all of the GPFS file systems in the cluster. When used in combination with the **-g** option, only group quota limit checking is activated. When used in combination with the **-u** or **-j** option, only user or fileset quota limit checking, respectively, is activated.
- g** Specifies that only group quota limit checking is to be activated.
- j** Specifies that only fileset quota checking is to be activated.
- u** Specifies that only user quota limit checking is to be activated.
- v** Prints a message for each file system in which quota limit checking is activated.

### Exit status

0      Successful completion.

#### nonzero

A failure has occurred.

### Security

You must have root authority to run the **mmquotaon** command.

## **mmquotaon**

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmquotaon** command is issued.

### **Examples**

1. To activate user quotas on file system **fs0**, issue this command:

```
mmquotaon -u fs0
```

To confirm the change, issue this command:

```
mmlsfs fs0 -Q
```

The system displays information similar to:

| flag | value | description     |
|------|-------|-----------------|
| -Q   | user  | Quotas enforced |

2. To activate group quota limit checking on all file systems, issue this command:

```
mmquotaon -g -a
```

To confirm the change, individually for each file system, issue this command:

```
mmlsfs fs1 -Q
```

The system displays information similar to:

| flag | value | description     |
|------|-------|-----------------|
| -Q   | group | Quotas enforced |

3. To activate user, group, and fileset quota limit checking on file system **fs2**, issue this command:

```
mmquotaon fs2
```

To confirm the change, issue this command:

```
mmlsfs fs2 -Q
```

The system displays information similar to:

| flag | value              | description     |
|------|--------------------|-----------------|
| -Q   | user;group;fileset | Quotas enforced |

### **See also**

- “[mmcheckquota command](#)” on page 169
- “[mmdefedquota command](#)” on page 272
- “[mmdefquotaoff command](#)” on page 275
- “[mmdefquotaon command](#)” on page 278
- “[mmedquota command](#)” on page 325
- “[mmlsquota command](#)” on page 435
- “[mmquotaoff command](#)” on page 507
- “[mmrepquota command](#)” on page 517

### **Location**

/usr/lpp/mmfs/bin

---

## **mmremotecluster command**

Manages information about remote GPFS clusters.

### **Synopsis**

```
mmremotecluster add RemoteClusterName [-n ContactNodes] [-k KeyFile]
```

or

```
mmremotecluster update RemoteClusterName [-c NewClusterName] [-n ContactNodes] [-k KeyFile]
```

or

```
mmremotecluster delete {RemoteClusterName | all}
```

or

```
mmremotecluster show [RemoteClusterName | all]
```

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

The **mmremotecluster** command is used to make remote GPFS clusters known to the local cluster, and to maintain the attributes associated with those remote clusters. The keyword appearing after **mmremotecluster** determines which action is performed:

#### **add**

Adds a remote GPFS cluster to the set of remote clusters known to the local cluster.

#### **delete**

Deletes the information for a remote GPFS cluster.

#### **show**

Displays information about a remote GPFS cluster.

#### **update**

Updates the attributes of a remote GPFS cluster.

To be able to mount file systems that belong to some other GPFS cluster, you must first make the nodes in this cluster aware of the GPFS cluster that owns those file systems. This is accomplished with the **mmremotecluster add** command. The information that the command requires must be provided to you by the administrator of the remote GPFS cluster. You will need this information:

- The name of the remote cluster.
- The names or IP addresses of a few nodes that belong to the remote GPFS cluster.
- The public key file generated by the administrator of the remote cluster by issuing the **mmauth genkey** command for the remote cluster.

Since each cluster is managed independently, there is no automatic coordination and propagation of changes between clusters like there is between the nodes within a cluster. This means that once a remote cluster is defined with the **mmremotecluster** command, the information about that cluster is automatically propagated across all nodes that belong to this cluster. But if the administrator of the remote cluster decides to rename it, or deletes some or all of the contact nodes, or change the public key file, the information in this cluster becomes obsolete. It is the responsibility of the administrator of the remote GPFS cluster to notify you of such changes so that you can update your information using the appropriate options of the **mmremotecluster update** command.

## **mmremotecluster**

### **Parameters**

#### *RemoteClusterName*

Specifies the cluster name associated with the remote cluster that owns the remote GPFS file system. The value **all** indicates all remote clusters defined to this cluster, when using the **mmremotecluster delete** or **mmremotecluster show** commands.

#### **-C NewClusterName**

Specifies the new cluster name to be associated with the remote cluster.

#### **-k KeyFile**

Specifies the name of the public key file provided to you by the administrator of the remote GPFS cluster.

#### **-n ContactNodes**

A comma separated list of nodes that belong to the remote GPFS cluster, in this format:

[tcpPort=NNNN,]node1[,node2 ...]

where:

tcpPort=NNNN

Specifies the TCP port number to be used by the local GPFS daemon when contacting the remote cluster. If not specified, GPFS will use the default TCP port number 1191.

node1[,node2...]

Specifies a list of nodes that belong to the remote cluster. The nodes can be identified through their host names or IP addresses.

### **Exit status**

- 0     Successful completion. After successful completion of the **mmremotecluster** command, the new configuration information is propagated to all nodes in the cluster.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the **mmremotecluster** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. This command adds remote cluster **k164.kgn.ibm.com** to the set of remote clusters known to the local cluster, specifying **k164n02** and **k164n03** as remote contact nodes. File **k164.id\_rsa.pub** is the name of the public key file provided to you by the administrator of the remote cluster.

```
mmremotecluster add k164.kgn.ibm.com -n k164n02,k164n03 -k k164.id_rsa.pub
```

The output is similar to this:

```
mmremotecluster: 6027-1371 Propagating the cluster configuration data to all affected nodes. This is an asynchronous process.
```

2. This command displays information for the remote cluster **k164.kgn.ibm.com**.

```
mmremotecluster show k164.kgn.ibm.com
```

The output is similar to this:

```
Cluster name: k164.kgn.ibm.com
Contact nodes: k164n02,k164n03
SHA digest: a3917c8282fca7a27d951566940768dc241902b
File systems: (none defined)
```

For more information on the SHA digest, see the *IBM Spectrum Scale: Problem Determination Guide* and search on *SHA digest*.

3. This command updates information for the remote cluster **k164.kgn.ibm.com**, changing the remote contact nodes to **k164n02** and **k164n01**. The TCP port to be used when contacting cluster **k164.kgn.ibm.com** is defined to be 6667.

```
mmremotecluster update k164.kgn.ibm.com -n tcpPort=6667,k164n02,k164n01
```

The output is similar to this:

```
mmremotecluster: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

The **mmremotecluster show** command can then be used to see the changes.

```
mmremotecluster show k164.kgn.ibm.com
```

The output is similar to this:

```
Cluster name: k164.kgn.ibm.com
Contact nodes: tcpPort=6667,k164n02,k164n01
SHA digest: a3917c8282fca7a27d951566940768dc241902b
File systems: (none defined)
```

For more information on the SHA digest, see the *IBM Spectrum Scale: Problem Determination Guide* and search on *SHA digest*.

4. This command deletes information for remote cluster **k164.kgn.ibm.com** from the local cluster.

```
mmremotecluster delete k164.kgn.ibm.com
```

The output is similar to this:

```
mmremotecluster: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

## See also

- “[mmauth command](#)” on page 67
- “[mmremotefs command](#)” on page 514

See also the topic about accessing GPFS file systems from other GPFS clusters in the *IBM Spectrum Scale: Administration Guide*.

## Location

/usr/lpp/mmfs/bin

## **mmremotefs**

---

### **mmremotefs command**

Manages information needed for mounting remote GPFS file systems.

#### **Synopsis**

```
mmremotefs add Device [-f RemoteDevice] [-C RemoteClusterName]
[-T MountPoint] [-t DriveLetter]
[-A {yes | no | automount}] [-o MountOptions] [--mount-priority Priority]
```

or

```
mmremotefs delete {Device | all | -C RemoteClusterName} [--force]
```

or

```
mmremotefs show [Device | all | -C RemoteClusterName]
```

or

```
mmremotefs update Device [-f RemoteDevice] [-C RemoteClusterName]
[-T MountPoint] [-t DriveLetter]
[-A {yes | no | automount}] [-o MountOptions] [--mount-priority Priority]
```

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

The **mmremotefs** command is used to make GPFS file systems that belong to other GPFS clusters known to the nodes in this cluster, and to maintain the attributes associated with these file systems. The keyword appearing after **mmremotefs** determines which action is performed:

#### **add**

Define a new remote GPFS file system.

#### **delete**

Delete the information for a remote GPFS file system.

#### **show**

Display the information associated with a remote GPFS file system.

#### **update**

Update the information associated with a remote GPFS file system.

Use the **mmremotefs** command to make the nodes in this cluster aware of file systems that belong to other GPFS clusters. The cluster that owns the given file system must have already been defined with the **mmremotecluster** command. The **mmremotefs** command is used to assign a local name under which the remote file system will be known in this cluster, the mount point where the file system is to be mounted in this cluster, and any local mount options that you may want.

Once a remote file system has been successfully defined and a local device name associated with it, you can issue normal commands using that local name, the same way you would issue them for file systems that are owned by this cluster.

When running the **mmremotefs** command delete and update options, the file system must be unmounted on the local cluster. However, it can be mounted elsewhere.

## Parameters

### *Device*

Specifies the name by which the remote GPFS file system will be known in the cluster.

### **-C** *RemoteClusterName*

Specifies the name of the GPFS cluster that owns the remote GPFS file system.

### **-f** *RemoteDevice*

Specifies the actual name of the remote GPFS file system. This is the device name of the file system as known to the remote cluster that owns the file system.

## Options

### **-A** {*yes* | *no* | *automount*}

Indicates when the file system is to be mounted:

#### **yes**

When the GPFS daemon starts.

#### **no**

Manual mount. This is the default.

#### **automount**

When the file system is first accessed.

### **-o** *MountOptions*

Specifies the mount options to pass to the mount command when mounting the file system. For a detailed description of the available mount options, see *GPFS-specific mount options* in *IBM Spectrum Scale: Administration Guide*.

### **-T** *MountPoint*

The local mount point directory of the remote GPFS file system. If it is not specified, the mount point will be set to *DefaultMountDir/Device*. The default value for *DefaultMountDir* is /gpfs, but it can be changed with the **mmchconfig** command.

### **-t** *DriveLetter*

Specifies the drive letter to use when the file system is mounted on Windows.

### **--mount-priority** *Priority*

Controls the order in which the individual file systems are mounted at daemon startup or when one of the **all** keywords is specified on the **mmmount** command.

File systems with higher *Priority* numbers are mounted after file systems with lower numbers. File systems that do not have mount priorities are mounted last. A value of zero indicates no priority.

### **--force**

The --force flag can only be used with the delete option. It will override an error that can occur when trying to delete a remote mount where the remote cluster was already removed. If the original delete attempt returns an error stating it cannot check to see if the mount is in use, then this is the condition to use. The --force flag overrides and allows the deletion to complete.

## Exit status

- 0 Successful completion. After successful completion of the **mmremotefs** command, the new configuration information is propagated to all nodes in the cluster.

### **nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmremotefs** command.

## **mmremotefs**

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

### **Examples**

This command adds remote file system **gpfsn**, owned by remote cluster **k164.kgn.ibm.com**, to the local cluster, assigning **rgpfsn** as the local name for the file system, and **/gpfs/rgpfsn** as the local mount point.

```
mmremotefs add rgpfsn -f gpfsn -C k164.kgn.ibm.com -T /gpfs/rgpfsn
```

The output is similar to this:

```
mmremotefs: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

The **mmremotefs show** command can be used to see the changes.

```
mmremotefs show rgpfsn
```

The output is similar to this:

| Local Name | Remote Name | Cluster name               | Mount Point | Mount Options | Automount | Drive |
|------------|-------------|----------------------------|-------------|---------------|-----------|-------|
| rgpfs1     | gpfs1       | gpfs-n60-win.fvtdomain.net | /rgpfs1     | rw            | no        | K     |

### **See also**

- “mmauth command” on page 67
- “mmremotecluster command” on page 511

See also the topic about accessing GPFS file systems from other GPFS clusters in the *IBM Spectrum Scale: Administration Guide*.

### **Location**

/usr/lpp/mmfs/bin

---

## mmrepquota command

Displays file system user, group, and fileset quotas.

### Synopsis

```
mmrepquota [-u] [-g] [-e] [-q] [-n] [-v] [-t]
[--block-size {BlockSize | auto} | -Y] {-a | Device:Fileset ...}
```

or

```
mmrepquota [-u] [-g] [-j] [-e] [-q] [-n] [-v] [-t]
[--block-size {BlockSize | auto} | -Y] {-a | Device...}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmrepquota** command reports file system usage and quota information for a user, group, or fileset.

This command cannot be run from a Windows node.

If **-g**, **-j**, or **-u** are not specified, then user, group, and fileset quotas are listed.

If **-a** is not specified, *Device* must be the last parameter entered.

For each file system in the cluster, the **mmrepquota** command displays:

1. Block limits (displayed in number of data blocks in 1KB units or a unit defined by the **--block-size** parameter):
  - quota type (USR, GRP or FILESET)
  - current usage (the amount of disk space used by this user, group, or fileset, in 1KB units or a unit defined by the **--block-size** parameter)
  - soft limit (the amount of disk space that this user, group, or fileset is allowed to use during normal operation, in 1KB units or a unit defined by the **--block-size** parameter)
  - hard limit (the total amount of disk space that this user, group, or fileset is allowed to use during the grace period, in 1KB units or a unit defined by the **--block-size** parameter)
  - space in doubt
  - grace period
2. File limits:
  - current number of files
  - soft limit
  - hard limit
  - files in doubt
  - grace period

**Note:** In cases where small files do not have an additional block allocated for them, quota usage may show less space usage than expected.

3. Entry Type

**default on**

Default quotas are enabled for this file system.

## **mmrepquota**

### **default off**

Default quotas are not enabled for this file system.

**e** Explicit quota limits have been set using the **mmedquota** command.

**d\_fsys** The quota limits are the default file system values set using the **mmdefedquota** command.

**d\_fset** The quota limits are the default fileset-level values set using the **mmdefedquota** command.

**i** Default quotas were not enabled when this initial entry was established. Initial quota limits have a value of zero indicating no limit.

Because the sum of the in-doubt value and the current usage may not exceed the hard limit, the actual block space and number of files available to the user, group, or fileset may be constrained by the *in-doubt* value. If the *in-doubt* values approach a significant percentage of the quota, run the **mmcheckquota** command to account for the lost space and files.

For more information, see *Listing quotas* in the *IBM Spectrum Scale: Administration Guide*.

## **Parameters**

### *Device*

The device name of the file system to be listed.

If more than one file system is listed, the names must be delimited by a space. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

### *Fileset*

Specifies an optional fileset to be listed.

- a** Lists quotas for all file systems in the cluster. A header line is printed automatically with this option.
- e** Specifies that the **mmrepquota** command is to collect updated quota usage data from all nodes before displaying results. If this option is not specified, there is the potential to display negative usage values as the quota server may process a combination of up-to-date and back-level information.
- g** Lists only group quotas.
- j** Lists only fileset quotas.
- n** Displays a numerical user ID.
- q** Shows whether quota enforcement is active.
- t** Lists global user, group, and fileset block and inode grace times.
- u** Lists only user quotas.
- v** Prints a header line for the file systems being queried and adds an entryType description for each quota entry.
- block-size {BlockSize | auto}**  
Specifies the unit in which the number of blocks is displayed. The value must be of the form **[n]K**, **[n]M**, **[n]G** or **[n]T**, where **n** is an optional integer in the range 1 to 1023. The default is **1K**. If **auto** is specified, the number of blocks is automatically scaled to an easy-to-read value.
- Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

## Exit status

**0** Successful completion.

### nonzero

A failure has occurred.

## Security

You must have root authority to run the **mmrepquota** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

GPFS must be running on the node from which the **mmrepquota** command is issued.

## Examples

1. To report on user quotas for file system **fs2** and display a header line, issue this command:

```
mmrepquota -u -v fs2
```

The system displays information similar to:

```
*** Report for USR quotas on fs2
```

| Name  | type | KB   | Block Limits |       |       | grace   | File Limits |       |       | Type             |
|-------|------|------|--------------|-------|-------|---------|-------------|-------|-------|------------------|
|       |      |      | quota        | limit | doubt |         | files       | quota | limit |                  |
| root  | USR  | 8    | 0            | 0     | 0     | none    | 1           | 0     | 0     | none default on  |
| user2 | USR  | 2016 | 256          | 512   | 0     | 6days   | 7           | 10    | 20    | 0 none d_fsys    |
| user3 | USR  | 104  | 256          | 512   | 0     | none    | 1           | 10    | 20    | 0 none d_fsys    |
| user4 | USR  | 0    | 256          | 512   | 0     | none    | 0           | 10    | 20    | 0 none d_fsys    |
| user5 | USR  | 368  | 256          | 512   | 0     | 23hours | 5           | 4     | 10    | 0 23hours d_fsys |
| user6 | USR  | 0    | 256          | 512   | 0     | none    | 0           | 10    | 20    | 0 none d_fsys    |
| user7 | USR  | 1024 | 1024         | 5120  | 4096  | none    | 1           | 0     | 0     | 19 none e        |

2. To report on quota enforcement for **fs2**, issue this command:

```
mmrepquota -q fs2
```

The system displays information similar to:

```
fs2: USR quota is on; default quota is on
fs2: GRP quota is on; default quota is on
fs2: FILESET quota is on; default quota is off
```

3. To report on user quotas for file system **gpfs2**, issue this command:

```
mmrepquota -u gpfs2
```

The system displays information similar to:

| Name   | fileset | type | Block Limits |       |        | grace  | File Limits |       |       | in     |
|--------|---------|------|--------------|-------|--------|--------|-------------|-------|-------|--------|
|        |         |      | KB           | quota | limit  |        | files       | quota | limit |        |
| root   | root    | USR  | 0            | 0     | 0      | 0 none | 1           | 0     | 0     | 0 none |
| root   | fset3   | USR  | 8            | 0     | 0      | 0 none | 1           | 0     | 0     | 0 none |
| root   | fset4   | USR  | 8192         | 0     | 0      | 0 none | 1           | 0     | 0     | 0 none |
| pfs001 | root    | USR  | 0            | 10240 | 0      | 0 none | 0           | 1000  | 5000  | 0 none |
| pfs001 | fset3   | USR  | 0            | 10240 | 0      | 0 none | 0           | 1000  | 5000  | 0 none |
| pfs001 | fset4   | USR  | 4104         | 10240 | 153600 | 0 none | 2           | 1000  | 5000  | 0 none |

4. To report on user quotas for file system **gpfs2** in fileset **fset4**, issue this command:

```
mmrepquota -u gpfs2:fset4
```

The system displays information similar to:

## **mmrepquota**

| Block Limits |         |      |      |       |        |       |       | File Limits |       |       |       |       |      |  |  |
|--------------|---------|------|------|-------|--------|-------|-------|-------------|-------|-------|-------|-------|------|--|--|
|              |         |      |      | in    |        |       |       |             |       |       |       | in    |      |  |  |
| Name         | fileset | type | KB   | quota | limit  | doubt | grace | files       | quota | limit | doubt | grace |      |  |  |
| root         | fset4   | USR  | 8192 | 0     | 0      | 0     | none  | 1           | 0     | 0     | 0     | 0     | none |  |  |
| pfs001       | fset4   | USR  | 4104 | 10240 | 153600 | 0     | none  | 2           | 1000  | 5000  | 0     | 0     | none |  |  |

5. To list global user, group, and fileset block and inode grace times, issue this command:

```
mmrepquota -u -t gpfs_s
```

The system displays information similar to:

| User:        | block default grace time 7days, inode default grace time 7days |    |       |       |          |       |       |       |       |          |       |
|--------------|----------------------------------------------------------------|----|-------|-------|----------|-------|-------|-------|-------|----------|-------|
| Group:       | block default grace time 7days, inode default grace time 7days |    |       |       |          |       |       |       |       |          |       |
| Fileset:     | block default grace time 7days, inode default grace time 7days |    |       |       |          |       |       |       |       |          |       |
| Block Limits |                                                                |    |       |       |          |       |       |       |       |          |       |
| Name         | type                                                           | KB | quota | limit | in_doubt | grace | files | quota | limit | in_doubt | grace |
| root         | USR                                                            | 0  | 0     | 0     | 0        | none  | 50    | 0     | 0     | 0        | none  |
| ftp          | USR                                                            | 0  | 0     | 0     | 0        | none  | 50    | 0     | 0     | 0        | none  |

6. To report on fileset quotas in file system fs1, issue this command:

```
mmrepquota -j fs1 --block-size auto
```

The system displays information similar to:

| Name  | fileset | type    | Block Limits |       |       |          | File Limits |       |       |       |          |        |
|-------|---------|---------|--------------|-------|-------|----------|-------------|-------|-------|-------|----------|--------|
|       |         |         | blocks       | quota | limit | in_doubt | grace       | files | quota | limit | in_doubt | grace  |
| root  | root    | FILESET | 256K         | 0     | 0     | 0        | none        | 1     | 0     | 0     | 0        | none   |
| fset0 | root    | FILESET | 89.25G       | 100G  | 200G  | 89.99M   | none        | 13729 | 4000  | 5000  | 0        | 7 days |
| fset1 | root    | FILESET | 0            | 100G  | 200G  | 0        | none        | 1     | 4000  | 5000  | 0        | none   |

**Note:** In any **mmrepquota** listing, when the 'type' is FILESET, the 'Name' column heading is meant to indicate the fileset name (root, fset0, fset1, in this example), and the value for the 'fileset' column heading (root, in this example) can be ignored.

## **See also**

- “**mmcheckquota** command” on page 169
- “**mmdefedquota** command” on page 272
- “**mmdefquotaoff** command” on page 275
- “**mmdefquotaoon** command” on page 278
- “**mmedquota** command” on page 325
- “**mmlsquota** command” on page 435
- “**mmquotaoff** command” on page 507
- “**mmquotaon** command” on page 509

## **Location**

/usr/lpp/mmfs/bin

---

## mmrestoreconfig command

Restores file system configuration information.

### Synopsis

```
mmrestoreconfig Device -i InputFile [-I {yes | test}]
[-Q {yes | no | only}] [-W NewDeviceName]
```

or

```
mmrestoreconfig Device -i InputFile --image-restore
[-I {yes | test}] [-W NewDeviceName]
```

or

```
mmrestoreconfig Device -i InputFile -F QueryResultFile
```

or

```
mmrestoreconfig Device -i InputFile -I continue
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

### Description

The **mmrestoreconfig** command allows you to either query or restore, or both query *and* restore, the output file of the **mmbbackupconfig** command.

In the **query phase**, the **mmrestoreconfig** command uses the output file generated by the **mmbbackupconfig** command as an input parameter, and then creates a configuration file. Users can then edit the configuration file to fit their current file system configuration. You can use the definitions in the configuration file to create the appropriate network shared disks (NSDs) and file systems required for the restore.

In the **image restore phase**, the **mmrestoreconfig** command uses the input file (output from the **mmbbackupconfig** command) to restore the backed up file system configuration in the newly created file system. The newly created file system must not be mounted prior to the **mmimimgrestore** command execution thus the quota settings are turned off for image restore. They can be reactivated after the **mmimimgrestore** command is completed using the **-Q only** flag of the **mmrestoreconfig** command.

This command cannot be run from a Windows node.

### Parameters

#### *Device*

Specifies the name of the file system to be restored.

#### **-i** *inputFile*

Specifies the file generated by the **mmbbackupconfig** command. The input file contains the file system configuration information.

#### **-I** {**yes** | **test**}

Specifies the action to be taken during the restore phase:

##### **yes**

Test and proceed on the restore process. This is the default action.

##### **test**

Test all the configuration settings before the actual restore is performed.

## **mmrestoreconfig**

Use **-I continue** to restart **mmrestoreconfig** from the last known successful configuration restore.

### **-F QueryResultFile**

Specifies the pathname of the configuration query result file generated by **mmrestoreconfig**. The configuration query result file is a report file that you can edit and use as a guide to **mmcrnsd** or **mmcdfs**.

### **--image-restore**

Restores the configuration data in the proper format for Scale Out Backup and Restore (SOBAR).

### **-Q {yes | no | only}**

Specifies whether quota settings are enforced during the file system restore. If set to **no**, the quota settings are ignored.

To restore quota settings after the **mmimgrstore** command has successfully run, the **-Q only** option must be specified.

### **-W newDeviceName**

Restores the backed up file system information to this new device name

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmrestoreconfig** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. Run **mmrestoreconfig -F QueryResultFile** to specify the pathname of the configuration query result file to be generated.

```
mmrestoreconfig gpfs1 -i inputFile -F reportfile
```

2. To test settings before running a restore:

```
mmrestoreconfig fs1 -i /tmp/fs1.mmbbackupconfig.out -I test
```

The system displays output similar to:

```

Configuration test restore of fs1 begins at Wed Mar 14 16:00:16 EDT 2012.

mmrestoreconfig: Checking disk settings for fs1:
mmrestoreconfig: Checking the number of storage pools defined for fs1.
The restored filesystem currently has 1 pools defined.
mmrestoreconfig: Checking storage pool names defined for fs1.
Storage pool 'system' defined.
mmrestoreconfig: Checking storage pool size for 'system'.
mmrestoreconfig: Storage pool size 127306752 was defined for 'system'.

mmrestoreconfig: Checking filesystem attribute configuration for fs1:
File system attribute to be restored: defaultDataReplicas
 Backup value: 2
 Current value: 1

mmrestoreconfig: Checking fileset configurations for fs1:
```

```

Fileset to restore: root.
Fileset status: Linked /fs1
Fileset mode: off
Fileset to restore: smallfileset.
Fileset status: Linked /fs1/smallfileset
Fileset mode: off

mmrestoreconfig: Checking policy rule configuration for fs1:
mmrestoreconfig: Testing policy configuration restore.
Validated policy `policyfile.backup': parsed 1 Placement Rules, 0 Restore Rules,
 0 Migrate/Delete/Exclude Rules,
 0 List Rules, 0 External Pool/List Rules

mmrestoreconfig: Checking quota settings for fs1:
mmrestoreconfig: Checking quota enablement for fs1.

mmrestoreconfig: Disabling the following settings:
mmrestoreconfig: Enabling the following settings: -u -g -j

mmrestoreconfig: Disabling the following default quota settings: -u -g -j

mmrestoreconfig: Enabling the following default quota settings: -u -g -j

mmrestoreconfig: Quota limits for fs1:

mmrestoreconfig: Default Quota limits for fs1:

mmrestoreconfig: Command successfully completed

```

3. Run **mmrestoreconfig** to restore the **gpfs1** file system:

```
mmrestoreconfig gpfs1 -i inputFile
```

4. To restore the **fs9** file system configuration data prior to the **mmimimgrestore** command, issue:

```
mmrestoreconfig fs9 -i fs9.backupconfig --image-restore
```

The system displays output similar to:

```

mmrestoreconfig: Quota and DMAPI are enabled.
mmrestoreconfig: Disabling quota and/or DMAPI ...

Configuration restore of fs9 begins at Thu Nov 29 17:09:55 EST 2012.

mmrestoreconfig: Checking disk settings for fs9:
mmrestoreconfig: Checking the number of storage pools defined for fs9.
mmrestoreconfig: Checking storage pool names defined for fs9.
mmrestoreconfig: Checking storage pool size for 'system'.

mmrestoreconfig: Checking filesystem attribute configuration for fs9:

mmrestoreconfig: Checking policy rule configuration for fs9:
mmrestoreconfig: No policy rules installed in backed up filesystem fs9.
mmrestoreconfig: Command successfully completed

```

5. To restore the quota settings for file system **fs9**, after the **mmimimgrestore** command, issue:

```
mmrestoreconfig fs9 -i fs9.backupconfig -Q only
```

The system displays output similar to:

```

Configuration restore of fs9 begins at Thu Nov 29 17:13:51 EST 2012.

mmrestoreconfig: Checking quota settings for fs9:
mmrestoreconfig: Checking quota enablement for fs9.

mmrestoreconfig: Restoring quota and defquota limits for fs9:
fs9: Start quota check
 11 % complete on Thu Nov 29 17:17:37 2012

```

## **mmrestoreconfig**

```
22 % complete on Thu Nov 29 17:17:37 2012
33 % complete on Thu Nov 29 17:17:37 2012
44 % complete on Thu Nov 29 17:17:37 2012
55 % complete on Thu Nov 29 17:17:38 2012
69 % complete on Thu Nov 29 17:17:38 2012
84 % complete on Thu Nov 29 17:17:38 2012
100 % complete on Thu Nov 29 17:17:39 2012
Finished scanning the inodes for fs9.
Merging results from scan.
mmrestoreconfig: Command successfully completed
```

### **See also**

- “[mmbbackupconfig command](#)” on page 81
- “[mmimbackup command](#)” on page 367
- “[mmimgrestore command](#)” on page 371

### **Location**

/usr/lpp/mmfs/bin

---

## mmrestorefs command

Restores a file system or an independent fileset from a snapshot.

### Synopsis

```
mmrestorefs Device SnapshotName [-j FilesetName]
[-N {Node[,Node...] | NodeFile | NodeClass}]
[--log-quiet] [--preserve-encryption-attributes]
[--suppress-external-attributes] [--threads MaxNumThreads]
[--work-unit FilesPerThread]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher. Available on AIX and Linux.

### Description

Use the **mmrestorefs** command to restore user data and attribute files to a file system or an independent fileset using those of the specified snapshot. Data will be restored by **mmrestorefs** without regard for file system or fileset quotas unless the **enforceFilesetQuotaOnRoot** configuration attribute of the **mmchconfig** command is set to **yes**. The **mmrestorefs** command does not restore the file system and fileset quota configuration information.

In versions before IBM Spectrum Scale 4.1.1, ensure that the file system is unmounted before you run the **mmrestorefs** command. For more information, see Table 19 on page 526 and Table 20 on page 526 below. When restoring from an independent fileset snapshot (using the **-j** option), link the fileset from nodes in the cluster that are to participate in the restore. It is preferable to run the **mmrestorefs** command when there are no user operations (either from commands, applications, or services) in progress on the file system or fileset. If there are user operations in progress on the file system or fileset while **mmrestorefs** is running, the restore might fail. For these failures, stop the user operations and run the **mmrestorefs** command again to complete the restore. For better performance, run the **mmrestorefs** command when the system is idle. While the restore is in progress, do not unlink the fileset, unmount the file system, or delete the fileset, fileset snapshot, or file system.

The **mmrestorefs** command cannot restore a fileset that was deleted after a global snapshot was created. In addition, the filesets in a global snapshot that are in deleted or unlinked state cannot be restored.

Snapshots are not affected by the **mmrestorefs** command. When a failure occurs during a restore, try repeating the **mmrestorefs** command except when there are ENOSPC or quota exceeded errors. In these cases, fix the errors then try the **mmrestorefs** command again.

For information on how GPFS policies and snapshots interact, see the *IBM Spectrum Scale: Administration Guide*.

Because snapshots are not copies of the entire file system, they should not be used as protection against media failures. For protection against media failures, see the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide* and search on “recoverability considerations”.

The **mmrestorefs** command can cause a compressed file in the active file system to become decompressed if it is overwritten by the restore process. To recompress the file, run the **mmstripefile** command with the **-z** option.

## **mmrestorefs**

### **CAUTION:**

- **Do not run file compression or decompression while an mmrestorefs command is running.** This caution applies to compression or decompression with the mmchattr command or with the mmapplypolicy command.
- **Do not run the mmrestripefs or mmrestripefile command while an mmrestorefs command is running.**

**Note:** The following table shows the requirements and the results when you restore a global snapshot:

*Table 19. Restoring a global snapshot*

| The product version level of the node that runs the mmrestorefs command | The file system must be in this state | Results                                                                   |
|-------------------------------------------------------------------------|---------------------------------------|---------------------------------------------------------------------------|
| Before V4.1.1                                                           | Unmounted                             | The file system manager performs the restore.                             |
| V4.1.1 or later                                                         | Mounted                               | By default the restore is performed on all nodes running V4.1.1 or later. |

The following table shows the requirements and the results when you restore a fileset snapshot:

*Table 20. Restoring a fileset snapshot*

| The product version level of the node that runs the mmrestorefs command | The file system must be in this state | Results                                                                   |
|-------------------------------------------------------------------------|---------------------------------------|---------------------------------------------------------------------------|
| V3.5                                                                    | Unmounted                             | The file system manager performs the restore.                             |
| V4.1.1 or later                                                         | Mounted                               | By default the restore is performed on all nodes running V4.1.1 or later. |

## **Parameters**

### *Device*

The device name of the file system that contains the snapshot to use for the restore. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

This must be the first parameter.

### *SnapshotName*

Specifies the name of the snapshot that will be used for the restore.

### **-j FilesetName**

Specifies the name of a fileset covered by this snapshot.

### **-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes that are to participate in the restore. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

Starting with IBM Spectrum Scale 4.1.1, **-N** can be used for both fileset and global snapshot restores. (In GPFS 4.1, **-N** can be used for fileset snapshot restore only. In GPFS 3.5 and earlier, there is no **-N** parameter.)

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### **--log-quiet**

Suppresses detailed thread log output.

### **--preserve-encryption-attributes**

Preserves the encryption extended attributes. Files that were removed after the snapshot was taken

are restored with the same encryption attributes (including FEK) of the file in the snapshot. If this option is not used, the file is recreated with the encryption policy in place at the time the file is restored.

#### --suppress-external-attributes

Specifies that external attributes will not be restored.

#### --threads *MaxNumThreads*

Specifies the maximum number of concurrent restore operations. The default is 24.

#### --work-unit *FilesPerThread*

Specifies the number of files each thread will process at a time. The default is 100.

## Exit status

0      Successful completion.

#### nonzero

A failure has occurred.

## Security

You must have root authority to run the **mmrestorefs** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

Suppose that you have the following directory structure:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3
```

The directory **userA** is then deleted, leaving the following structure:

```
/fs1/file1

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3
```

The directory **userB** is then created using the inode originally assigned to **userA**, and another snapshot is taken:

```
mmcrsnapshot fs1 snap2
```

The output is similar to this:

```
Writing dirty data to disk.
Quiescing all file system operations.
Writing dirty data to disk again.
Snapshot snap2 created with id 2.
```

The directory structure is similar to the following:

## **mmrestorefs**

```
/fs1/file1
/fs1/userB/file2b
/fs1/userB/file3b

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3

/fs1/.snapshots/snap2/file1
/fs1/.snapshots/snap2/userB/file2b
/fs1/.snapshots/snap2/userB/file3b
```

The file system is then restored from **snap1**:

```
mmrestorefs fs1 snap1
```

The resulting directory structure is similar to the following:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3

/fs1/.snapshots/snap1/file1
/fs1/.snapshots/snap1/userA/file2
/fs1/.snapshots/snap1/userA/file3

/fs1/.snapshots/snap2/file1
/fs1/.snapshots/snap2/userB/file2b
/fs1/.snapshots/snap2/userB/file3b
```

### **See also**

- “[mmcrsnapshot command](#)” on page 267
- “[mmdelsnapshot command](#)” on page 304
- “[mmlssnapshot command](#)” on page 439
- “[mmsnapdir command](#)” on page 565

### **Location**

```
/usr/lpp/mmfs/bin
```

## mmrestripefile command

Rebalances or restores the replication factor of the specified files, or performs any incomplete or deferred file compression or decompression.

### Synopsis

```
mmrestripefile {-m | -r | -p | -b | -l | -c | -z} {-F FilenameFile | Filename [Filename...]}
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **mmrestripefile** command attempts to repair the specified files, or performs any deferred or incomplete compression or decompression of the specified files. You can use **-F** option to specify a file that contains the list of file names to be processed, with one file name per line.

The repair options are rebalancing (**-b**), restoring replication factors (**-r**), migrating data (**-m**), and migrating file data to the proper pool (**-p**). The **-b** option not only rebalances files but also performs all the operations of the **-m** and **-r** options. For more information, see *Restriping a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

If you do not use replication, the **-r** and **-m** options are equivalent. Their behavior differs only on replicated files. After a successful rereplicate (**-r**), all suspended disks are empty. But a migrate operation (**-m**) leaves data on a suspended disk as long as at least one other replica of the data remains on a disk that is not suspended.

Use the **-l** option to relocate the block placement of the files.

Use the **-z** option to perform any deferred or incomplete compression or decompression of the files.

#### CAUTION:

Do not run the **mmrestripefs** or **mmrestripefile** command while an **mmrestorefs** command is running.

### Parameters

#### **-F** *FilenameFile*

Specifies a file that contains a list of names of files to be restriped, one name per line.

#### *Filename*

Specifies the names of one or more files to be restriped.

### Options

stop

- m Migrates critical data from any suspended disk for a list of specified files. Critical data is all data that would be lost if currently suspended disks were removed.
- r Migrates all data for a list of files from suspended disks. If a disk failure or removal makes some replicated data inaccessible, this command also restores replicated files to their designated level of replication. Use this option immediately after a disk failure to protect replicated data against a subsequent failure. You can also use this option before you take a disk offline for maintenance to protect replicated data against the failure of another disk during the maintenance process.
- p Moves the data of ill-placed files to the correct storage pool.

## **mmrestripefile**

| Some utilities, including the **mmchattr** command, can assign a file to a different storage pool without  
| moving the data of the file to the new pool. Such files are called *ill-placed*. The **-p** parameter causes  
| the command to move the data of ill-placed files to the correct storage pools.

- b Rebalances a list of files across all disks that are not suspended, even if they are stopped. Although blocks are allocated on a stopped disk, they are not written to a stopped disk, nor are reads allowed from a stopped disk, until that disk is started and replicated data is copied onto it.
- l Relocates the block placement of the file. The location of the blocks depends on the current write affinity depth, write affinity failure group setting, block group factor, and the node from which the command is run. For example, for an existing file, regardless of how its blocks are distributed on disks currently, if **mmrestripefile -l** is run from node A, the final block distribution looks as if the file was created from scratch on node A.

To specify the write affinity failure group where the replica is put, before you run **mmrestripefile -l**, enter a command like the following example:

```
mmchattr --write-affinity-failure-group "WadfgValueString" filename
```

where *WadfgValueString* is the failure group.

- c Compares the data of the specified files with their replicas and attempts to fix any inconsistencies. This option does not compare metadata blocks. It also does not compare files whose size is less than the inode size, because their data is stored in the inode. To find the inode size, enter the following command:

```
mm1sfs <file_system_name> -i
```

- z Performs any deferred or incomplete compression or decompression of files. For more information, see the topic *File compression* in the *IBM Spectrum Scale: Administration Guide*.

## **Exit status**

0 Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmrestripefile** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

This example illustrates restriping a file that is named `testfile0`. The following command confirms that the file is ill-placed:

```
mm1sattr -L testfile0
```

The system displays the following output:

```
file name: testfile0
metadata replication: 2 max 2
data replication: 2 max 2
immutable: no
appendOnly: no
```

```
flags: illplaced
storage pool name: system
fileset name: root
snapshot name:
```

To correct the problem, issue the following command:

```
mmrestripefile -p testfile0
```

To confirm the change, issue the following command:

```
mmlsattr -L testfile0
```

The system displays the following output:

```
file name: testfile0
metadata replication: 2 max 2
data replication: 2 max 2
immutable: no
appendOnly: no
flags:
storage pool name: system
fileset name: root
snapshot name:
```

The following command compresses or decompresses a file for which compression or decompression is deferred or incomplete:

```
mmrestripefile -z largefile.data
```

## See also

- “[mmadddisk command](#)” on page 23
- “[mmapplypolicy command](#)” on page 56
- “[mmchattr command](#)” on page 122
- “[mmchdisk command](#)” on page 161
- “[mmdeldisk command](#)” on page 287
- “[mmrpldisk command](#)” on page 539
- “[mmrestripefs command](#)” on page 532

## Location

/usr/lpp/mmf/bin

## **mmrestripefs**

---

### **mmrestripefs command**

Rebalances or restores the replication factor of all the files in a file system. Alternatively, this command performs any incomplete or deferred file compression or decompression of all the files in a file system.

#### **Synopsis**

```
mmrestripefs Device {-m | -r | -b | -R | -c [--read-only] | -p | -z}
[-N {Node[,Node...] | NodeFile | NodeClass}] [-o InodeResultFile]
[-P PoolName] [--inode-criteria CriteriaFile] [--qos QOSClass]
```

or

```
mmrestripefs Device {-r | -b | -R | -c [--read-only]} --metadata-only
[-N {Node[,Node...] | NodeFile | NodeClass}] [-o InodeResultFile]
[--inode-criteria CriteriaFile] [--qos QOSClass]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Issue the **mmrestripefs** command to rebalance or restore the replication of all files in a file system. The command moves existing file system data between different disks in the file system based on changes to the disk state made by the **mmchdisk**, **mmadddisk**, and **mmdeledisk** commands. It also attempts to restore the metadata or data replication of all the files in the file system.

Alternatively, you can issue the **mmrestripefs** command to perform any deferred or incomplete file compression or decompression in all the files of a file system.

You must specify one of the options (**-m**, **-r**, **-b**, **-R**, **-c**, **-p**, or **-z**) to indicate how much file system data to move or whether to perform file compression or decompression. You can issue this command against a mounted or unmounted file system.

If the file system uses replication, then restriping the file system also replicates it. Also, if the file system uses replication the **-r** option and the **-m** options treat suspended disks differently. The **-r** option removes all data from a suspended disk. But the **-m** option leaves data on a suspended disk if at least one replica of the data remains on a disk that is not suspended.

The **-b** option performs all the operations of the **-m** and **-r** options.

Use the **-z** option to perform any deferred or incomplete file compression or decompression.

#### **CAUTION:**

**Do not issue the mmrestripefs or mmrestripefile command while an mmrestorefs command is running.**

Consider the necessity of restriping and the current demands on the system. New data that is added to the file system is correctly striped. Restriping a large file system requires many insert and delete operations and might affect system performance. Plan to perform this task when system demand is low.

#### **Parameters**

##### *Device*

The device name of the file system to be restriped. File system names need not be fully qualified.

*Device* must be the first parameter. It can take the following parameters:

- m Migrates all critical data off of any suspended disk in this file system. Critical data is all data that would be lost if currently suspended disks were removed.
- r Migrates all data off suspended disks. It also restores all replicated files in the file system to their designated degree of replication when a previous disk failure or removal of a disk makes some replica data inaccessible. Use this parameter either immediately after a disk failure to protect replicated data against a subsequent failure, or before you take a disk offline for maintenance to protect replicated data against failure of another disk during the maintenance process.
- b Rebalances all files across all disks that are not suspended, even if they are stopped. Although blocks are allocated on a stopped disk, they are not written to a stopped disk, nor are reads allowed from a stopped disk, until that disk is started and replicated data is copied onto it. The **mmrestripefs** command rebalances and restripes the file system. Use this option to rebalance the file system after you add, change, or delete disks in a file system.

**Note:** Rebalancing of files is an I/O intensive and time-consuming operation, and is important only for file systems with large files that are mostly invariant. In many cases, normal file update and creation will rebalance your file system over time, without the cost of the rebalancing.

- R Changes the replication settings of each file, directory, and system metadata object so that they match the default file system settings (see the **mmchfs** command -m and -r options) as long as the maximum (-M and -R) settings for the object allow it. Next, it replicates or unreplicates the object as needed to match the new settings. This option can be used to replicate all of the existing files that were not previously replicated or to unreplicate the files if replication is no longer needed or wanted.
- c Scans the file system and compares replicas of metadata and data for conflicts. When conflicts are found, the -c option attempts to fix the replicas.

#### --read-only

Modifies the -c option so that it does not try to fix conflicting replicas. You can use this option only with the -c option.

#### --metadata-only

Limits the specified operation to metadata blocks. Data blocks are not affected. This option is valid only with the -r, -b, -R, or -c option.

**mmrestripefs** command with this option completes its operation quicker than a full stripe, replication, or replica compare of data and metadata.

Use this option when you want to prioritize the mmrestripefs operation on the metadata. This option ensures that the mmrestripefs operation has a reduced impact on the file system performance when compared to running the **mmrestripefs** command on the metadata and data.

After running the **mmrestripefs** command on the metadata with --metadata-only option, you can issue the **mmrestripefs** command without this option to stripe the data and any metadata that requires to be striped.

**Note:** This option does not run until all the nodes in the cluster are upgraded to IBM Spectrum Scale 4.2.1 release. If any of the nodes is not upgraded, the system displays the following error message:

```
mmrestripefs: The --metadata-only option support has not been enabled yet.
Run "mmchconfig release=LATEST" to activate the new function.
mmrestripefs: Command failed. Examine previous error messages to determine cause.
```

- p Directs **mmrestripefs** to repair the file placement within the storage pool.

Files that are assigned to one storage pool, but with data in a different pool, have their data migrated to the correct pool. Such files are referred to as ill-placed. Utilities, such as the **mmchattr** command, might change a file's storage pool assignment, but not move the data. The **mmrestripefs** command might then be invoked to migrate all of the data at once, rather than

## **mmrestripefs**

migrating each file individually. The placement option (**-p**) rebalances only the files that it moves. In contrast, the rebalance operation (**-b**) performs data placement on all files.

- z Performs any deferred or incomplete file compression or decompression of files in the file system. For more information, see the topic *File compression* in the *IBM Spectrum Scale: Administration Guide*.

### **-P PoolName**

Directs **mmrestripefs** to repair only files assigned to the specified storage pool. This option is convenient for migrating ill-placed data blocks between pools, for example after you change a file's storage pool assignment with **mmchattr** or **mmapplypolicy** with the **-I defer** flag.

Do not use for other tasks, in particular, for any tasks that require metadata processing, such as re-replication. By design, all GPFS metadata is kept in the system pool, even for files that have blocks in other storage pools. Therefore a command that must process all metadata must not be restricted to a specific storage pool.

### **-N {Node[,Node...] | NodeFile | NodeClass}**

Specify the nodes that participate in the stripe of the file system. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

### **-o InodeResultFile**

Contains a list of the inodes that met the interesting inode flags that were specified on the **--inode-criteria** parameter. The output file contains the following:

#### **INODE\_NUMBER**

This is the inode number.

#### **DISKADDR**

Specifies a dummy address for later **tsfindinode** use.

#### **SNAPSHOT\_ID**

This is the snapshot ID.

#### **ISGLOBAL\_SNAPSHOT**

Indicates whether or not the inode is in a global snapshot. Files in the live file system are considered to be in a global snapshot.

#### **INDEPENDENT\_FSETID**

Indicates the independent fileset to which the inode belongs.

#### **MEMO (INODE\_FLAGS FILE\_TYPE [ERROR])**

Indicates the inode flag and file type that will be printed:

##### **Inode flags:**

BROKEN  
exposed  
dataUpdateMiss  
illCompressed  
illPlaced  
illReplicated  
metaUpdateMiss  
unbalanced

##### **File types:**

BLK\_DEV  
CHAR\_DEV  
DIRECTORY  
FIFO  
LINK  
LOGFILE

```
REGULAR_FILE
RESERVED
SOCK
UNLINKED
DELETED
```

**Notes:**

1. An error message will be printed in the output file if an error is encountered when repairing the inode.
2. **DISKADDR**, **ISGLOBAL\_SNAPSHOT**, and **FSET\_ID** work with the **tsfindinode** tool (`/usr/lpp/mmfs/bin/tsfindinode`) to find the file name for each inode. **tsfindinode** uses the output file to retrieve the file name for each interesting inode.

**--inode-criteria CriteriaFile**

Specifies the interesting inode criteria flag, where *CriteriaFile* is one of the following:

**BROKEN**

Indicates that a file has a data block with all of its replicas on disks that have been removed.

**Note:** **BROKEN** is always included in the list of flags even if it is not specified.

**dataUpdateMiss**

Indicates that at least one data block was not updated successfully on all replicas.

**exposed**

Indicates an inode with an exposed risk; that is, the file has data where all replicas are on suspended disks. This could cause data to be lost if the suspended disks have failed or been removed.

**illCompressed**

Indicates an inode in which file compression or decompression is deferred, or in which a compressed file is partly decompressed to allow the file to be written into or memory-mapped.

**illPlaced**

Indicates an inode with some data blocks that might be stored in an incorrect storage pool.

**illReplicated**

Indicates that the file has a data block that does not meet the setting for the replica.

**metaUpdateMiss**

Indicates that there is at least one metadata block that has not been successfully updated to all replicas.

**unbalanced**

Indicates that the file has a data block that is not well balanced across all the disks in all failure groups.

**Note:** If a file matches *any* of the specified interesting flags, all of its interesting flags (even those not specified) will be displayed.

**--qos QOSclass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

## **mmrestripefs**

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

### **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

### **Security**

You must have root authority to issue the **mmrestripefs** command.

The node on which you issue the command must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### **Examples**

1. To move all critical data from any suspended disk in file system **fs1**, issue the following command:

```
mmrestripefs fs1 -m
```

The system displays information similar to the following output:

```
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
8.00 % complete on Tue Feb 24 16:56:55 2009 (708608 inodes 346 MB)
100.00 % complete on Tue Feb 24 16:56:56 2009
GPFS: 6027-552 Scan completed successfully.
```

2. To rebalance all files in file system **fs1** across all defined, accessible disks that are not stopped or suspended, issue the following command:

```
mmrestripefs fs1 -b
```

The system displays information similar to the following output:

```
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
3.00 % complete on Tue Feb 24 16:56:39 2009 (180224 inodes 161 MB)
100.00 % complete on Tue Feb 24 16:56:44 2009
GPFS: 6027-552 Scan completed successfully.
```

3. To compare and fix replica conflicts of metadata and data in file system **gpfs1**, issue the following command:

```
mmrestripefs gpfs1 -c
```

The system displays information similar to the following output:

```
Scanning file system metadata, phase 1 ...
Inode 0 in fileset 0 and snapshot 0 has mismatch in replicated disk address 2:104859136
Scan completed successfully.
Scanning file system metadata, phase 2 ...
Scan completed successfully.
Scanning file system metadata, phase 3 ...
Scan completed successfully.
Scanning file system metadata, phase 4 ...
Scan completed successfully.
Scanning user file metadata ...
100.00 % complete on Tue Jul 30 03:32:44 2013
Scan completed successfully.
```

4. To fix the pool placement of files in file system **fs1** and also determine which files are illReplicated (for example, as a result of a failed disk), issue the following command:

```
mmrestripefs fs1 -p --inode-criteria /tmp/crit -o /tmp/inodeResultFile
```

The system displays information similar to the following output:

```
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
Scanning file system metadata for data storage pool
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
100.00 % complete on Wed Apr 15 10:15:15 2015 (65792 inodes with total 400 MB data processed)
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-3902 Check file '/tmp/inodeResultFile' on vmip1 for inodes that were \
found matching the criteria.

#10:15:15# vmip1:/fs1 # cat /tmp/crit
illReplicated
#10:15:19# vmip1:/fs1 # cat /tmp/inodeResultFile
This inode list was generated in the Parallel Inode Traverse on Wed Apr 15 10:15:14 2015
INODE_NUMBER DISKADDR SNAPSHOT_ID ISGLOBAL_SNAPSHOT FSET_ID MEMO(INODE_FLAGS FILE_TYPE [ERROR])
24320 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24322 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24321 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24324 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24325 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24323 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24326 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24327 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24328 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
24329 0:0 0 1 0 illreplicated unbalanced REGULAR_FILE
```

## See also

- “mmadddisk command” on page 23
- “mmapplypolicy command” on page 56
- “mmchattr command” on page 122
- “mmchdisk command” on page 161
- “mmchfs command” on page 179
- “mmdeldisk command” on page 287
- “mmrpldisk command” on page 539
- “mmrestripefile command” on page 529

## **mmrestripefs**

### **Location**

/usr/lpp/mmfs/bin

---

## mmrpldisk command

Replaces the specified disk.

### Synopsis

```
mmrpldisk Device DiskName {DiskDesc | -F StanzaFile} [-v {yes | no}]
[-N {Node[,Node...]} | NodeFile | NodeClass}]
[--inode-criteria CriteriaFile] [-o InodeResultFile]
[--qos QOSClass]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmrpldisk** command to replace an existing disk in the GPFS file system with a new one. All data on the old disk is migrated to the new disk.

To replace a disk in a GPFS file system, you must first decide if you will:

1. Create a new disk using the **mmcrnsd** command.

In this case, use the rewritten disk stanza file produced by the **mmcrnsd** command or create a new disk stanza. When using the rewritten file, the disk usage and failure group specifications remain the same as specified on the **mmcrnsd** command.

2. Select a disk no longer in any file system. Issue the **mmlsnsd -F** command to display the available disks.

The disk may then be used to replace a disk in the file system using the **mmrpldisk** command.

#### Notes:

1. You cannot replace a disk when it is the only remaining disk in the file system.
2. Under no circumstances should you replace a stopped disk. You need to start a stopped disk before replacing it. If a disk cannot be started, delete it using the **mmdeldisk** command. See the *IBM Spectrum Scale: Problem Determination Guide* and search for "Disk media failure".
3. The file system may be mounted when running the **mmrpldisk** command.

### Results

Upon successful completion of the **mmrpldisk** command, the disk is replaced in the file system and data is copied to the new disk without restriping.

### Parameters

#### Device

The device name of the file system where the disk is to be replaced. File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

This must be the first parameter.

#### DiskName

The name of the disk to be replaced. To display the names of disks that belong to the file system, issue the **mmlsnsd -f**, **mmlsfs -d**, or **mmlsdisk** command. The **mmlsdisk** command will also show the current disk usage and failure group values for each of the disks.

#### DiskDesc

A descriptor for the replacement disk.

## **mmrpldisk**

Prior to GPFS 3.5, the disk information for the **mmrpldisk** command was specified in the form of a disk descriptor defined as follows (with the second, third, sixth, and seventh fields reserved):

DiskName:::DiskUsage:FailureGroup:::

For backward compatibility, the **mmrpldisk** command will still accept a traditional disk descriptor as input, but this use is discouraged.

### **-F StanzaFile**

Specifies a file containing the NSD stanzas for the replacement disk. NSD stanzas have this format:

```
%nsd:
 nsd=NsdName
 usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}
 failureGroup=FailureGroup
 pool=StoragePool
 servers=ServerList
 device=DiskName
```

where:

**nsd=NsdName**

The name of an NSD previously created by the **mmcrnsd** command. For a list of available disks, issue the **mmlsnsd -F** command. This clause is mandatory for the **mmrpldisk** command.

**usage={dataOnly | metadataOnly | dataAndMetadata | descOnly}**

Specifies the type of data to be stored on the disk:

**dataAndMetadata**

Indicates that the disk contains both data and metadata. This is the default for disks in the system pool.

**dataOnly**

Indicates that the disk contains data and does not contain metadata. This is the default for disks in storage pools other than the system pool.

**metadataOnly**

Indicates that the disk contains metadata and does not contain data.

**descOnly**

Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster-recovery configurations. For more information, see the *IBM Spectrum Scale: Administration Guide* and search for "Synchronous mirroring utilizing GPFS replication"

This clause is optional for the **mmrpldisk** command. If omitted, the new disk will inherit the usage type of the disk being replaced.

**failureGroup=FailureGroup**

Identifies the failure group to which the disk belongs. A failure group identifier can be a simple integer or a topology vector that consists of up to three comma-separated integers. The default is -1, which indicates that the disk has no point of failure in common with any other disk.

GPFS uses this information during data and metadata placement to ensure that no two replicas of the same block can become unavailable due to a single failure. All disks that are attached to the same NSD server or adapter must be placed in the same failure group.

If the file system is configured with data replication, all storage pools must have two failure groups to maintain proper protection of the data. Similarly, if metadata replication is in effect, the system storage pool must have two failure groups.

Disk that belong to storage pools in which write affinity is enabled can use topology vectors to identify failure domains in a shared-nothing cluster. Disks that belong to traditional storage pools must use simple integers to specify the failure group.

This clause is optional for the **mmrpldisk** command. If omitted, the new disk will inherit the failure group of the disk being replaced.

**pool=StoragePool**

Specifies the storage pool to which the disk is to be assigned. This clause is ignored by the **mmrpldisk** command.

**servers=ServerList**

A comma-separated list of NSD server nodes. This clause is ignored by the **mmrpldisk** command.

**device=DiskName**

The block device name of the underlying disk device. This clause is ignored by the **mmrpldisk** command.

**Note:** While it is not absolutely necessary to specify the same parameters for the new disk as the old disk, it is suggested that you do so. If the new disk is equivalent in size to the old disk, and if the disk usage and failure group parameters are the same, the data and metadata can be completely migrated from the old disk to the new disk. A disk replacement in this manner allows the file system to maintain its current data and metadata balance.

If the new disk has a different size, disk usage, parameter, or failure group parameter, the operation may leave the file system unbalanced and require a restripe. Additionally, a change in size or the disk usage parameter may cause the operation to fail since other disks in the file system may not have sufficient space to absorb more data or metadata. In this case, first use the **mmadddisk** command to add the new disk, the **mmdeledisk** command to delete the old disk, and finally the **mmrestripefs** command to rebalance the file system.

**-v {yes | no}**

Verify the new disk does not belong to an existing file system. The default is **-v yes**. Specify **-v no** only when you want to reuse a disk that is no longer needed for an existing file system. If the command is interrupted for any reason, use the **-v no** option on the next invocation of the command.

**Important:** Using **-v no** on a disk that already belongs to a file system will corrupt that file system. This will not be noticed until the next time that file system is mounted.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specify the nodes that participate in the migration of data from the old to the new disk. This command supports all defined node classes. The default is **all** or the current value of the **defaultHelperNodes** parameter of the **mmchconfig** command.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

**--inode-criteria CriteriaFile**

Specifies the interesting inode criteria flag, where *CriteriaFile* is one of the following:

**BROKEN**

Indicates that a file has a data block with all of its replicas on disks that have been removed.

**Note:** **BROKEN** is always included in the list of flags even if it is not specified.

**dataUpdateMiss**

Indicates that at least one data block was not updated successfully on all replicas.

**exposed**

Indicates an inode with an exposed risk; that is, the file has data where all replicas are on suspended disks. This could cause data to be lost if the suspended disks have failed or been removed.

## **mmrpldisk**

### **i11Compressed**

Indicates an inode in which file compression or decompression is deferred, or in which a compressed file is partly decompressed to allow the file to be written into or memory-mapped.

### **i11Placed**

Indicates an inode with some data blocks that might be stored in an incorrect storage pool.

### **i11Replicated**

Indicates that the file has a data block that does not meet the setting for the replica.

### **metaUpdateMiss**

Indicates that there is at least one metadata block that has not been successfully updated to all replicas.

### **unbalanced**

Indicates that the file has a data block that is not well balanced across all the disks in all failure groups.

**Note:** If a file matches *any* of the specified interesting flags, all of its interesting flags (even those not specified) will be displayed.

### **-o InodeResultFile**

Contains a list of the inodes that met the interesting inode flags that were specified on the **--inode-criteria** parameter. The output file contains the following:

#### **INODE\_NUMBER**

This is the inode number.

#### **DISKADDR**

Specifies a dummy address for later **tsfindinode** use.

#### **SNAPSHOT\_ID**

This is the snapshot ID.

#### **ISGLOBAL\_SNAPSHOT**

Indicates whether or not the inode is in a global snapshot. Files in the live file system are considered to be in a global snapshot.

#### **INDEPENDENT\_FSETID**

Indicates the independent fileset to which the inode belongs.

#### **MEMO (INODE\_FLAGS FILE\_TYPE [ERROR])**

Indicates the inode flag and file type that will be printed:

##### **Inode flags:**

BROKEN  
exposed  
dataUpdateMiss  
i11Compressed  
i11Placed  
i11Replicated  
metaUpdateMiss  
unbalanced

##### **File types:**

BLK\_DEV  
CHAR\_DEV  
DIRECTORY  
FIFO  
LINK  
LOGFILE  
REGULAR\_FILE

```
RESERVED
SOCK
UNLINKED
DELETED
```

**Notes:**

1. An error message will be printed in the output file if an error is encountered when repairing the inode.
2. **DISKADDR**, **ISGLOBAL\_SNAPSHOT**, and **FSET\_ID** work with the **tsfindinode** tool (**/usr/lpp/mmfs/bin/tsfindinode**) to find the file name for each inode. **tsfindinode** uses the output file to retrieve the file name for each interesting inode.

**--qos QOSClass**

Specifies the Quality of Service for I/O operations (QoS) class to which the instance of the command is assigned. If you do not specify this parameter, the instance of the command is assigned by default to the **maintenance** QoS class. This parameter has no effect unless the QoS service is enabled. For more information, see the topic “**mmchqos command**” on page 206. Specify one of the following QoS classes:

**maintenance**

This QoS class is typically configured to have a smaller share of file system IOPS. Use this class for I/O-intensive, potentially long-running GPFS commands, so that they contribute less to reducing overall file system performance.

**other** This QoS class is typically configured to have a larger share of file system IOPS. Use this class for administration commands that are not I/O-intensive.

For more information, see the topic *Setting the Quality of Service for I/O operations (QoS)* in the *IBM Spectrum Scale: Administration Guide*.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Security**

You must have root authority to run the **mmrpldisk** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

**Examples**

1. To replace disk **hd27n01** in **fs1** with a new disk, **hd16vsdn10** allowing the disk usage and failure group parameters to default to the corresponding values of **hd27n01**, and have only nodes **c154n01**, **c154n02**, and **c154n09** participate in the migration of the data, issue this command:

```
mmrpldisk fs1 hd27n01 hd16vsdn10 -N c154n01,c154n02,c154n09
```

The system displays information similar to:

```
Replacing hd27n01 ...
```

```
The following disks of fs1 will be formatted on node c155n01.ppd.pok.ibm.com:
hd16vsdn10: size 17793024 KB
Extending Allocation Map
Checking Allocation Map for storage pool 'system'
```

## **mmrpldisk**

```
 7 % complete on Wed May 16 16:36:30 2007
 18 % complete on Wed May 16 16:36:35 2007
 34 % complete on Wed May 16 16:36:40 2007
 49 % complete on Wed May 16 16:36:45 2007
 65 % complete on Wed May 16 16:36:50 2007
 82 % complete on Wed May 16 16:36:55 2007
 98 % complete on Wed May 16 16:37:00 2007
 100 % complete on Wed May 16 16:37:01 2007
Completed adding disks to file system fs1.
Scanning system storage pool
Scanning file system metadata, phase 1 ...
 2 % complete on Wed May 16 16:37:04 2007
 7 % complete on Wed May 16 16:37:11 2007
 14 % complete on Wed May 16 16:37:18 2007
 20 % complete on Wed May 16 16:37:24 2007
 27 % complete on Wed May 16 16:37:31 2007
 34 % complete on Wed May 16 16:37:37 2007
 50 % complete on Wed May 16 16:37:50 2007
 61 % complete on Wed May 16 16:38:00 2007
 68 % complete on Wed May 16 16:38:06 2007
 79 % complete on Wed May 16 16:38:16 2007
 90 % complete on Wed May 16 16:38:26 2007
 100 % complete on Wed May 16 16:38:32 2007
Scan completed successfully.
Scanning file system metadata, phase 2 ...
Scanning file system metadata for fs1spl storage pool
Scan completed successfully.
Scanning file system metadata, phase 3 ...
Scan completed successfully.
Scanning file system metadata, phase 4 ...
Scan completed successfully.
Scanning user file metadata ...
 3 % complete on Wed May 16 16:38:38 2007
 25 % complete on Wed May 16 16:38:47 2007
 53 % complete on Wed May 16 16:38:53 2007
 87 % complete on Wed May 16 16:38:59 2007
 97 % complete on Wed May 16 16:39:06 2007
 100 % complete on Wed May 16 16:39:07 2007
Scan completed successfully.
Done
mmrpldisk: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

2. To replace disk **vmip3\_nsd1** from storage pool GOLD on file system **fs2** and to search for any interesting files handled during the **mmrpldisk** at the same time, issue this command:

```
mmrpldisk fs2 vmip3_nsd1 -F f /tmp/crit --inode-criteria
```

The system displays information similar to:

```
Replacing vmip3_nsd1 ...
GPFS: 6027-531 The following disks of fs2 will be formatted on node vmip1:
 vmip2_nsd3: size 5120 MB
Extending Allocation Map
Checking Allocation Map for storage pool GOLD
 59 % complete on Wed Apr 15 10:52:44 2015
 100 % complete on Wed Apr 15 10:52:49 2015
GPFS: 6027-1503 Completed adding disks to file system fs2.
GPFS: 6027-589 Scanning file system metadata, phase 1 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 2 ...
Scanning file system metadata for GOLD storage pool
Scanning file system metadata for BRONZE storage pool
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 3 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-589 Scanning file system metadata, phase 4 ...
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-565 Scanning user file metadata ...
```

```
6.47 % complete on Wed Apr 15 10:53:11 2015 (65792 inodes with total 448 MB data processed)
6.49 % complete on Wed Apr 15 10:55:01 2015 (65792 inodes with total 448 MB data processed)
100.00 % complete on Wed Apr 15 10:55:03 2015 (65792 inodes with total 448 MB data processed)
GPFS: 6027-552 Scan completed successfully.
GPFS: 6027-3902 Check file '/var/mmfs/tmp/fs2.pit.interestingInodes.12884901928' on vmip1 for inodes \
that were found matching the criteria.
Checking Allocation Map for storage pool GOLD
56 % complete on Wed Apr 15 10:55:08 2015
100 % complete on Wed Apr 15 10:55:12 2015
Done
mmrpldisk: 6027-1371 Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
#11:57:08# vmip1:/fs2 # cat /tmp/crit
illReplicated
illPlaced
dataUpdateMiss
metaUpdateMiss
exposed
BROKEN
#11:09:24# vmip1:/fs2 # cat /var/mmfs/tmp/fs2.pit.interestingInodes.12884901928
This inode list was generated in the Parallel Inode Traverse on Wed Apr 15 10:55:02 2015
INODE_NUMBER DISKADDR SNAPSHOT_ID ISGLOBAL_SNAPSHOT FSET_ID MEMO(INODE_FLAGS FILE_TYPE [ERROR])
50177 0:0 0 1 0 illplaced REGULAR_FILE
```

**Note:** The **mmrpldisk** command will report any interesting inodes that it finds during routine processing, but the list might not be 100% accurate or complete.

## See also

- “[mmadddisk command](#)” on page 23
- “[mmchdisk command](#)” on page 161
- “[mmcrnsd command](#)” on page 262
- “[mmlsdisk command](#)” on page 402
- “[mmlsnsd command](#)” on page 423
- “[mmrestripefs command](#)” on page 532

## Location

/usr/lpp/mmfs/bin

## **mmsdrrestore**

---

### **mmsdrrestore command**

Restores the latest GPFS system files on the specified nodes.

#### **Synopsis**

```
mmsdrrestore [-p NodeName] [-F mmsdrfsFile] [-R remoteFileCopyCommand]
[-a | -N {Node[,Node...] | NodeFile | NodeClass}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

The **mmsdrrestore** command is intended for use by experienced system administrators.

Use the **mmsdrrestore** command to restore the latest GPFS system files on the specified nodes. If no nodes are specified, the command restores the configuration information only on the node on which it is run. If the local GPFS configuration file is missing, the file that is specified with the **-F** option from the node that is specified with the **-p** option is used instead. This command works best when used with the **mmsdrbackup** user exit. See “**mmsdrbackup** user exit” on page 826.

#### **Parameters**

##### **-p *NodeName***

Specifies the node from which to obtain a valid GPFS configuration file. The node must be either the primary configuration server or a node that has a valid backup copy of the **mmsdrfs** file. If not specified, the local node is used.

##### **-F *mmsdrfsFile***

Specifies the path name of the GPFS configuration file for the **mmsdrrestore** command to use. This configuration file might be the current one on the primary server, or it might be a configuration file that is obtained from the **mmsdrbackup** user exit. If not specified, */var/mmfss/gen/mmsdrfs* is used.

If the configuration file is a Cluster Configuration Repository (CCR) backup file, then you must also specify the **-a** option. All the nodes in the cluster are restored. However, if a configuration of the cluster is still available, you can restore the configuration of an individual node by running **mmsdrrestore -p**.

##### **-R *remoteFileCopyCommand***

Specifies the fully qualified path name for the remote file copy program to be used for obtaining the GPFS configuration file. The default is */usr/bin/rcp*.

##### **-a** Restores the GPFS configuration files on all nodes in the cluster.

##### **-N {*Node*[,*Node...*] | *NodeFile* | *NodeClass*}**

Restores the GPFS configuration files on a set of nodes.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*. This command does not support a *NodeClass* of mount.

**Note:** The **-N** option cannot be used if CCR is in effect for the cluster.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure occurred.

## Security

You must have root authority to run the **mmsdrrestore** command.

The node on which the command is issued must be able to run remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To restore the latest GPFS system files on the local node using the GPFS configuration file `/var/mmfs/gen/mmsdrfs` from the node that is named **primaryServer**, issue the following command:

```
mmsdrrestore -p primaryServer
```

The system displays output similar to the following example:

```
Tue Jul 3 18:19:53 CDT 2012: mmsdrrestore: Processing node k164n04.kgn.ibm.com
mmsdrrestore: Node k164n04.kgn.ibm.com successfully restored.
```

2. To restore the GPFS system files on all nodes in the cluster using GPFS configuration file `/GPFSconfigFiles/mmsdrfs.120605` on the node that is named **GPFSarchive**, issue the following command from the node named **localNode**:

```
mmsdrrestore -p GPFSarchive -F /GPFSconfigFiles/mmsdrfs.120605 -a
```

The system displays output similar to the following example:

```
Tue Jul 3 18:29:28 CDT 2012: mmsdrrestore: Processing node k164n04.kgn.ibm.com
Tue Jul 3 18:29:30 CDT 2012: mmsdrrestore: Processing node k164n05.kgn.ibm.com
Tue Jul 3 18:29:31 CDT 2012: mmsdrrestore: Processing node k164n06.kgn.ibm.com
mmsdrrestore: Command successfully completed
```

3. The following command restores the GPFS system files from the information in a Cluster Configuration Repository (CCR) backup file. When you restore from a CCR backup file, you must specify the **-a** option. All the nodes in the cluster are restored:

```
mmsdrrestore -F /GPFSbackupFiles/CCRBackup.2015.10.14.10.01.25.tar.gz -a
```

The command displays output similar to the following example:

```
Restoring CCR backup
CCR backup has been restored
```

## See also

- “[mmsdrbackup user exit](#)” on page 826

## Location

`/usr/lpp/mmfs/bin`

## **mmsetquota**

---

### **mmsetquota command**

Sets quota limits.

#### **Synopsis**

```
mmsetquota Device{[:FilesetName]
 [--user IdOrName[,IdOrName]] [--group IdOrName[,IdOrName]]}
 {[--block SoftLimit[:HardLimit]] [--files SoftLimit[:HardLimit]]}
```

or

```
mmsetquota Device[:FilesetName] --default {user | group}
 {[--block SoftLimit[:HardLimit]] [--files SoftLimit[:HardLimit]]}
```

or

```
mmsetquota Device --default fileset
 {[--block SoftLimit[:HardLimit]] [--files SoftLimit[:HardLimit]]}
```

or

```
mmsetquota Device --grace {user | group | fileset}
 {[--block GracePeriod] [--files GracePeriod]}
```

or

```
mmsetquota -F StanzaFile
```

#### **Availability**

Available on all IBM Spectrum Scale editions. Available on AIX and Linux.

#### **Description**

The **mmsetquota** command sets quota limits, default quota limits, or grace periods for users, groups, and filesets in the specified file system.

When setting quota limits for a file system, replication within the file system should be considered. For explanation, see *Listing quotas in IBM Spectrum Scale: Administration Guide*

#### **Parameters**

##### *Device*

Specifies the device name of the file system.

##### *FilesetName*

Specifies the name of a fileset located on *Device* for which quota information is to be set.

##### *IdOrName*

Specifies a numeric ID, user name, or group name.

##### *SoftLimit*

Specifies the amount of data or the number of files the user, group, or fileset will be allowed to use.

##### *HardLimit*

Specifies the amount of data or the number of files the user, group, or fileset will be allowed to use during a grace period. If omitted, the default is no limit. See note.

##### *GracePeriod*

Specifies the file-system grace period during which quotas can exceed the soft limit before it is imposed as a hard limit. See note.

***StanzaFile***

Specifies a file containing quota stanzas.

**--block**

Specifies the quota limits or grace period for data.

**--files**

Specifies the quota limits or grace period for files.

**--default**

Sets the default quota for the user, group, or fileset.

**--grace**

Sets the grace period for the user, group, or fileset.

**-F *StanzaFile***

Specifies a file containing the quota stanzas for set quota, set default quota, or set grace period.

Quota stanzas have this format:

```
%quota:
 device=Device
 command={setQuota|setDefaultQuota|setGracePeriod}
 type={USR|GRP|FILESET}
 id=IdList
 fileset=FilesetName
 blockQuota=Number
 blockLimit=Number
 blockGrace=Period
 filesQuota=Number
 filesLimit=Number
 filesGrace=Period
```

where:

**device=Device**

The device name of file system.

**command={setQuota|setDefaultQuota|setGracePeriod}**

Specifies the command to be executed for this stanza.

**setQuota**

Sets the quota limits. This command ignores **blockGrace** and **filesGrace** attributes.

**setDefaultQuota**

Sets the default quota limits. This command ignores **id**, **blockGrace** and **filesGrace** attributes.

**setGracePeriod**

Sets the grace periods. The command ignores **id**, **fileset**, and quota limit attributes. Grace periods can be set for each quota type in the file system.

**type={USR|GRP|FILESET}**

Specifies whether the command applies to user, group, or fileset.

**id=IdList**

Specifies a list of numeric IDs or user, group, or fileset names.

**fileset=FilesetName**

Specifies the fileset name for the **perfileset** quota setting. This attribute is ignored for

**type=FILESET****blockQuota=Number**

Specifies the block soft limit. The number can be specified using the suffix K, M, G, or T. See note.

## **mmsetquota**

### **blockLimit=Number**

Specifies the block hard limit. The number can be specified using the suffix K, M, G, or T. See note.

### **filesQuota=Number**

Specifies the inode soft limit. The number can be specified using the suffix K, M, or G. See note.

### **filesLimit=Number**

Specifies the inode hard limit. The number can be specified using the suffix K, M, or G. See note.

### **blockGrace=Period**

Specifies the file-system grace period during which the block quotas can exceed the soft limit before it is imposed as a hard limit. The period can be specified in days, hours, minutes, or seconds.

### **filesGrace=Period**

Specifies the file-system grace period during which the files quota can exceed the soft limit before it is imposed as a hard limit. The period can be specified in days, hours, minutes, or seconds.

### **Note:**

- | • The maximum files limit is 2147483647.
- | • The maximum block limit is 999999999999999K. For values greater than 976031318016K (909T) and up to the maximum limit of 999999999999999K (about 931322T), you must specify the equivalent value with the suffix K, M, or G.

## **Exit status**

**0**      Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmsetquota** command.

GPFS must be running on the node from which the **mmsetquota** command is issued.

You may issue the **mmsetquota** command only from a node in the GPFS cluster where the file system is mounted.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. The following command sets the block soft and hard limit to 25G and 30G and files soft and hard limit to 10K and 11K, respectively for user user234:

```
mmsetquota fs1 --user user234 --block 25G:30G --files 10K:11K
```

To verify the change, issue the following command:

```
mmquota -u user234 fs1
 Block Limits File Limits
Filesystem Fileset type KB quota limit in_doubt grace | files quota limit in_doubt grace Remarks
 fs1 root USR 143688 26214400 31457280 0 none | 1 10240 11264 0 none
 fs1 mkfiles2 USR no limits
 fs1 ifset1 USR no limits
 fs1 1111 USR no limits
 fs1 ifset2 USR no limits
```

2. If perfileset quota is enabled, the following command sets block soft and hard limit to 5G and 7G, respectively, for group fvt090 and for fileset ifset2:

```
mmsetquota fs1:ifset2 --group fvt090 --block 5G:7G
```

To verify the change, issue the following command:

```
mmlsquota -g fvt090 fs1
```

| Disk quotas for group fvt090 (gid 2590): |          |      |              |         |         |             |       |       |       |       |          |       |
|------------------------------------------|----------|------|--------------|---------|---------|-------------|-------|-------|-------|-------|----------|-------|
| Filesystem                               | Fileset  | type | Block Limits |         |         | File Limits |       |       |       |       |          |       |
|                                          |          |      | KB           | quota   | limit   | in_doubt    | grace | files | quota | limit | in_doubt | grace |
| fs1                                      | root     | GRP  | none         |         |         |             |       |       |       |       |          |       |
| fs1                                      | mkfiles2 | GRP  | none         |         |         |             |       |       |       |       |          |       |
| fs1                                      | ifset1   | GRP  | none         |         |         |             |       |       |       |       |          |       |
| fs1                                      | 1111     | GRP  | none         |         |         |             |       |       |       |       |          |       |
| fs1                                      | ifset2   | GRP  | 143704       | 5242880 | 7340032 | 0           | none  | 1     | 0     | 0     | 0        | none  |

3. To change the user grace period for block data to 10 days, issue the following command:

```
mmsetquota fs1 --grace user --block 10days
```

4. All of the previous examples can be done in one invocation of **mmsetquota** by using quota stanza file. The stanza file /tmp/quotaExample may look like this:

```
%quota:
device=fs1
command=setquota
type=USR
id=user234
blockQuota=25G
blockLimit=30G
filesQuota=10K
filesLimit=11K

%quota:
device=fs1
command=setquota
type=GRP
id=fvt090
fileset=ifset2
blockQuota=5G
blockLimit=7G

%quota:
device=fs1
command=setgraceperiod
type=user
blockGrace=10days
```

Then issue the command:

```
mmsetquota -F /tmp/quotaExample
```

## See also

- “mmcheckquota command” on page 169
- “mmdefedquota command” on page 272
- “mmdefquotaoff command” on page 275
- “mmdefquotao command” on page 278
- “mmedquota command” on page 325
- “mmlsquota command” on page 435
- “mmquotaon command” on page 509
- “mmquotaooff command” on page 507
- “mmrepquota command” on page 517

## Location

/usr/lpp/mmfs/bin

## **mmshutdown**

---

### **mmshutdown command**

Unmounts all GPFS file systems and stops GPFS on one or more nodes.

#### **Synopsis**

```
mmshutdown [-t UnmountTimeout] [-a | -N {Node[,Node...] | NodeFile | NodeClass}]
```

#### **Availability**

Available on all IBM Spectrum Scale editions.

#### **Description**

Use the **mmshutdown** command to stop the GPFS daemons on one or more nodes. If no operand is specified, GPFS is stopped only on the node from which the command was issued.

The **mmshutdown** command first attempts to unmount all GPFS file systems. If the unmount does not complete within the specified *timeout* period, the GPFS daemons shut down anyway.

#### **Results**

Upon successful completion of the **mmshutdown** command, these tasks are completed:

- GPFS file systems are unmounted.
- GPFS daemons are stopped.

#### **Parameters**

**-a** Stop GPFS on all nodes in a GPFS cluster.

**-N {*Node[ ,Node... ]* | *NodeFile* | *NodeClass*}**

Directs the **mmshutdown** command to process a set of nodes.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

#### **Options**

**-t *UnmountTimeout***

The maximum amount of time, in seconds, that the unmount command is given to complete. The default timeout period is equal to:

$60 + 3 \times \text{number of nodes}$

If the unmount does not complete within the specified amount of time, the command times out and the GPFS daemons shut down.

#### **Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmshutdown** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To stop GPFS on all nodes in the GPFS cluster, issue this command:

```
mmshutdown -a
```

The system displays information similar to:

```
Thu Mar 15 14:02:50 EDT 2012: mmshutdown: Starting force unmount of GPFS file systems
c6f1c3vp2(gpfs.net): forced unmount of /gpfs/fs1
c6f1c3vp1(gpfs.net): forced unmount of /gpfs/fs1
Thu Mar 15 14:03:00 EDT 2012: mmshutdown: Shutting down GPFS daemons
c6f1c3vp3(gpfs.net): Shutting down!
c6f1c3vp4(gpfs.net): Shutting down!
c6f1c3vp4(gpfs.net): 'shutdown' command about to kill process 23649
c6f1c3vp4(gpfs.net): Unloading modules from /lib/modules/2.6.27.19-5-ppc64/extra
c6f1c3vp4(gpfs.net): Unloading module mmfs26
c6f1c3vp1(gpfs.net): Shutting down!
c6f1c3vp4(gpfs.net): Unloading module mmfslinux
c6f1c3vp1(gpfs.net): 'shutdown' command about to kill process 7667944
c6f1c3vp2(gpfs.net): Shutting down!
c6f1c3vp2(gpfs.net): 'shutdown' command about to kill process 5701764
c6f1c3vp2(gpfs.net): Master did not clean up; attempting cleanup now
c6f1c3vp2(gpfs.net): Thu Mar 15 14:04:05.114 2012: mmfsd is shutting down.
c6f1c3vp2(gpfs.net): Thu Mar 15 14:04:05.115 2012: Reason for shutdown: mmfsadm shutdown command timed out
c6f1c3vp2(gpfs.net): Thu Mar 15 14:04:06 EDT 2012: mmcommon mmfsdown invoked. Subsystem: mmfs Status: down
c6f1c3vp2(gpfs.net): Thu Mar 15 14:04:06 EDT 2012: mmcommon: Unmounting file systems ...
Thu Mar 15 14:04:10 EDT 2012: mmshutdown: Finished
```

2. To stop GPFS on only node **k164n04**, issue this command:

```
mmshutdown -N k164n04
```

The system displays information similar to:

```
Thu Mar 15 14:00:12 EDT 2012: mmshutdown: Starting force unmount of GPFS file systems
k164n04: forced unmount of /gpfs/fs1
Thu Mar 15 14:00:22 EDT 2012: mmshutdown: Shutting down GPFS daemons
k164n04: Shutting down!
k164n04: 'shutdown' command about to kill process 7274548
Thu Mar 15 14:00:45 EDT 2012: mmshutdown: Finished
```

## See also

- “mmgetstate command” on page 350
- “mmlscluster command” on page 397
- “mmstartup command” on page 569

## Location

/usr/lpp/mmfs/bin

## mmsmb

---

### mmsmb command

Administers SMB shares, export ACLs, and global configuration.

#### Synopsis

```
mmsmb export list [ListofSMBExports][-Y][--option Arg][--export-regex Arg]
[--header N][--all][--key-info Arg]
```

or

```
mmsmb export add SMBExport Path [--option SMBOption=Value | --key-info SMBOption]
```

or

```
mmsmb export change SMBExport [--option SMBOption=Value | --remove SMBOption |
--key-info SMBOption]
```

or

```
mmsmb export remove SMBExport [--force]
```

or

```
mmsmb config list [ListofSMBOptions | -Y | --supported | --header N
| --key-info SMBOption]
```

or

```
mmsmb config change [--option SMBOption=Value | --remove SMBOption |
--key-info SMBOption]
```

or

```
mmsmb exportacl getid { Name | --user UserName | --group GroupName
| --system SystemName }
```

or

```
mmsmb exportacl list { [ExportName] | [ExportName --viewsddl] }[--viewsids]
```

or

```
mmsmb exportacl add ExportName { Name | --user UserName | --group GroupName
| --system SystemName | --SID SID } [--viewsids] --access Access --permissions
Permissions [--force]
```

or

```
mmsmb exportacl change ExportName { Name | --user UserName | --group GroupName
| --system SystemName | --SID SID } [--viewsids] --access Access --permissions Permissions
```

or

```
mmsmb exportacl remove ExportName { Name | --user UserName | --group GroupName
| --system SystemName | --SID SID } [--viewsids] [--access Access]
[--permissions Permissions]
```

or

```
mmsmb exportacl replace ExportName { Name | --user UserName | --group GroupName
| --system SystemName | --SID SID } --access Access--permissions Permissions [--viewsids]
[--force]
```

or

```
mmsmb exportacl delete ExportName [--viewsids] [--force]
```

**Note:** For **mmsmb export**, you can specify --option --remove multiple times but you cannot specify both of these options simultaneously.

## Availability

Available with IBM Spectrum Scale Standard Edition or higher.

The protocol functions provided in this command, or any similar command, are generally referred to as CES (Cluster Export Services). For example, protocol node and CES node are functionally equivalent terms.

## Description

Use the **mmsmb** command to administer SMB shares and global configuration.

Using the **mmsmb export** command, you can do the following tasks:

- Create the specified SMB share. The **mmsmb export add** command creates the specified export for the specified path. Any supported SMB option can be specified by repeating --option. Also the substitution values %D for the domain, %U for session user name and %G for the primary group of %U are supported as part of the specified path. The % character is not allowed in any other context. If the export exists but the path does not exist or if it is not inside the GPFS file system, the command returns with an error. When one or more substitution variables are used, only the sub-path to the first substitution variable is checked. If authentication on the cluster is not enabled, this command will terminate with an error.
- Change the specified SMB share using the **mmsmb export change** command.
- Delete an SMB share using the **mmsmb export remove** command. Existing connections to the deleted SMB share will be disconnected. This can result in data loss for files being currently open on the affected connections.
- List the SMB shares by using the **mmsmb export list** command. The command displays the configuration of options for each SMB share. If no specific options are specified, the command displays all SMB shares with the SMB options browseable; guest ok; smb encrypt as a table. Each row represents an SMB share and each column represents an SMB option.

Using the **mmsmb config** command, you can do the following tasks:

- Change, add or remove the specified SMB option for the SMB configuration. Use the **mmsmb config change** command to change the global configuration.
- List the global configuration of SMB shares. Use the **mmsmb config list** command to display the global configuration options of SMB shares. If no specific options are specified, the command displays all SMB option-value pairs.

Using the **mmsmb exportacl** command, you can do the following tasks:

- Retrieve the ID of the specified user/group/system.
- List, change, add, remove, replace and delete the ACL associated with an export.
- The add option has two mandatory arguments: --access and --permissions.

## Parameters

### **mmsmb export**

#### **list**

Lists the SMB shares.

#### **ListofSMBExports**

Specifies the list of SMB shares that needs to be listed as blank separated strings.

#### **--option {key | all | unsupported}**

## **mmsmb**

**key** Specifies only the supported SMB option to be listed.

**all** Displays all used SMB options.

### **unsupported**

Detects and displays all SMB shares with unsupported SMB options. The path of all unsupported options are listed for each SMB share.

### **--export-regex arg**

*arg* is a regular expression against which the exportnames are matched. Only matching exports are shown. If this option is not specified and if *ListofSMBExports* are also not specified, all existing exports are displayed.

### **--all**

Displays all defined SMB options. Similar to **--option all**.

### **add**

Creates the specified SMB share on a GPFS file system with NFSv4 ACLs enforced. You can verify whether your GPFS files system has been configured correctly by using the **mm1sfs** command. For example, **mm1sfs gpfs0 -k**.

### **path**

Specifies the path of the SMB share that needs to be added.

### **--option SMBOption=value**

Specifies the SMB option for the SMB protocol. If it is not a supported SMB option or the value is not allowable for this SMB option, the command terminates with an error. If this option is not specified, the default options: guest ok = no and smb encrypt = auto are set.

### **change**

Modifies the specified SMB share.

### **--option SMBOption=value**

Specifies the SMB option for the SMB protocol. If the SMB option is not configured for the specified export, it will be added with the specified value. If the SMB option is not supported or the value is not allowable for this SMB option the command terminates with an error. If no value is specified, the specified SMB option is set to default by removing the current setting from the configuration.

### **--remove SMBOption**

Specifies the SMB option that is to be removed. If the SMB option is supported it will be removed from the specified export. The default value becomes active. If the SMB option is not supported, the command terminates with an error.

### **remove**

Deletes the specified SMB share.

### **--force**

Suppresses confirmation questions.

### **SMBExport**

Specifies the SMB share that needs to be listed.

## **List of supported SMB options for the mmsmb export {list | add | change | remove} command:**

### **admin users**

Using this option, administrative users can be defined in the format of **admin users=user1,user2,...,usern**.

### **browsable**

If the value is set as yes, the export is shown in the Windows Explorer browser when browsing the file server. By default, this option is enabled.

**comment**

Description of the export.

**csc policy**

csc policy stands for client-side caching policy, and specifies how clients that are capable of offline caching cache the files in the share. The valid values are: manual and disable. Setting csc profile = disable disables offline caching. For example, this can be used for shares containing roaming profiles. By default, this option is set to the value manual.

**fileid:algorithm**

This option allows to control the level of enforced data integrity. If the data integrity is ensured on the application level, it can be beneficial in cluster environments to reduce the level of enforced integrity for performance reasons.

fsname is the default value that ensures data integrity in the entire cluster by managing concurrent access to files and directories cluster-wide.

**Note:** Data integrity is ensured if an application does not use multiple processes to access the data at the same time, for example, reading of file content does not happen while another process is still writing to the file. Without locking, the consistency of files is no longer guaranteed on protocol level. If data integrity is not ensured on application level this can lead to data corruption. For example, if two processes modify the same file in parallel, assuming that they have exclusive access.

**gpfs:leases**

gpfs:leases are cross protocol oplocks (opportunistic locks), that means an SMB client can lock a file that provides the user improved performance while reading or writing to the file because no other user read or write to this file. If the value is set as yes, clients accessing the file over the other protocols can break the oplock of an SMB client and the user gets informed when another user is accessing the same file at the same time.

**gpfs:recalls**

If the value is set as yes files that have been migrated from disk will be recalled on access. By default, this is enabled. If recalls = no files will not be recalled on access and the client will receive ACCESS\_DENIED message.

**gpfs:sharemodes**

An application can set share modes. If you set gpfs:sharemodes = yes, using the mmsmb export change *SMBexport --option "gpfs:sharemodes = yes"* the sharemodes specified by the application will be respected by all protocols and not only by the SMB protocol. If you set gpfs:sharemodes = no the sharemodes specified by the application will only be respected by the SMB protocol. For example, the NFS protocol will ignore the sharemode set by the application.

The application can set the following sharemodes: SHARE\_READ or SHARE\_WRITE or SHARE\_READ and SHARE\_WRITE or no sharemodes.

**gpfs:syncio**

If the value is set as yes, it specifies the files in an export, for which the setting is enabled, are opened with the O\_SYNC flag. Accessing a file is faster if gpfs:syncio is set to yes.

Performance for certain workloads can be improved when SMB accesses the file with the O\_SYNC flag set. For example, updating only small blocks in a large file as observed with database applications. The underlying GPFS behavior is then changed to not read a complete block if there is only a small update to it. By default, this option is disabled.

**hide unreadable**

If the value is set as yes, all files and directories that the user has no permission to read are hidden from directory listings in the export. The hideunreadable=yes option is also known as access-based enumeration because when a user is listing (enumerating) the directories and files within the export, they only see the files and directories that they have read access to. By default, this option is disabled.

## **mmsmb**

**Note:** Enabling this option has a negative impact on directory listing performance, especially for large directories as the ACLs for all directory entries have to be read and evaluated. This option is disabled by default.

### **oplocks**

If the value is set as yes, a client may request an opportunistic lock (**oplock**) from an SMB server when it opens a file. If the server grants the request, the client can cache large chunks of the file without informing the server what it is doing with the cached chunks until the task is completed. Caching large chunks of a file saves a lot of network I/O round-trip time and enhances performance. By default, this option is enabled.

**Note:** While **oplocks** can enhance performance, they can also contribute to data loss in case of SMB connection breaks/timeouts. To avoid the loss of data in case of an interface node failure or storage timeout, you might want to disable **oplocks**.

Opportunistic locking allows a client to notify the SMB server that it will be the exclusive writer of the file. It also notifies the SMB server that it will cache its changes to that file on its own system and not on the SMB server to speed up file access for that client. When the SMB server is notified about a file being opportunistically locked by a client, it marks its version of the file as having an opportunistic lock and waits for the client to complete work on the file. The client has to send the final changes back to the SMB server for synchronization. If a second client requests access to that file before the first client has finished working on it, the SMB server can send an oplock break request to the first client. This request informs the client to stop caching its changes and return the current state of the file to the server so that the interrupting client can use it. An opportunistic lock, however, is not a replacement for a standard deny-mode lock. There are many use cases when the interrupting process to be granted an oplock break only to discover that the original process also has a deny-mode lock on the file.

### **posix locking**

If the value is set as yes, it will be tested if a byte-range (fcnt1) lock is already present on the requested portion of the file before granting a byte-range lock to an SMB client. For improved performance on SMB-only shares this option can be disabled, which is the default behavior. Disabling locking on cross-protocol shares can result in data integrity issues when clients concurrently set locks on a file via multiple protocols, for example, SMB and NFS.

### **read only**

If the value is set as yes, files cannot be modified or created on this export independent of the ACLs. By default, the value is no.

### **smb encrypt**

This option controls whether the remote client is allowed or required to use SMB encryption. Possible values are This is set when the export is created with default value, which is auto. Clients may chose to encrypt the entire session, not just traffic to a specific export. The server would return access denied message to all non-encrypted requests on such an export. Selecting encrypted traffic reduces throughput as smaller packet sizes must be used as well as the overhead of encrypting and signing all the data. If SMB encryption is selected, the message integrity is guaranteed so that signing is implicit. When set to auto, SMB encryption is offered, but not enforced. When set to mandatory, SMB encryption is required and if set to disabled, SMB encryption cannot be negotiated.

### **syncops:onclose**

This option ensures that the file system synchronizes data to the disk each time a file is closed after writing. The written data is flushed to the disk. By default, this option is enabled.

**Note:** Disabling this option increases the risk of data loss in case of a node failure.

## **mmsmb config**

### **list**

Lists the global configuration options of SMB shares.

***List of SMB Options***

Specifies the list of SMB options that needs to be listed.

**--supported**

Displays all changeable SMB options and their values.

**change**

Modifies the global configuration options of SMB shares.

**--option *SMBOption=value***

Sets the value of the specified SMB option. If no value is given, the SMB option is removed.

Specifies the SMB option for the SMB protocol. If the SMB option is not configured for global configuration, it will be added with the specified value. If no value is specified, the specified SMB option is set to default and removed from the global configuration. If the SMB option is not supported or the value is not allowable for the SMB option, the command terminates with an error.

**--remove *SMBOption***

Specifies the SMB option that is to be removed. If the SMB option is supported it will be removed from the global configuration. The default value becomes active. If the SMB option is not supported, the command terminates with an error.

**List of supported SMB options by the mmsmb config {list | change} command:**

**gpfs:dfreequota**

**gpfs:dfreequota** stands for disk Free Quota. If the value is set to yes the free space and size reported to a SMB client for a share will be adjusted according to the applicable quotas. The applicable quotas are the quota of the user requesting this information, the quota of the user's primary group and the quota of the fileset containing the export.

**restrict anonymous**

The setting of this parameter determines whether access to information is allowed or restricted for anonymous users. The options are:

restrict anonymous = 2: anonymous users are restricted from accessing information.

restrict anonymous = 0: anonymous users are allowed to access information. This is the default setting.

restrict anonymous = 1: is not supported

**server string**

**server string** stands for Server Description. It specifies the server description for SMB protocol. Server description with special characters must be provided in single quotes.

**mmsmb exportacl****getid**

Retrieve the ID of the specified user/group/system.

```
mmsmb exportacl getid myUser
mmsmb exportacl getid --user myUser
mmsmb exportacl getid --group myGroup
mmsmb exportacl getid --system mySystem
```

**list**

List can only take viewing options.

|                                                 |                                                         |
|-------------------------------------------------|---------------------------------------------------------|
| <b>mmsmb exportacl list myExport</b>            | Show the export ACL for this exportname                 |
| <b>mmsmb exportacl list</b>                     | Show all the export ACLs                                |
| <b>mmsmb exportacl list myExport --viewsddl</b> | Show the export ACL for this exportname in sddl format. |

**add**

Add will add a new permission to the export ACL. It will include adding a user, group or system. The options are:

## **mmsmb**

- {User/group/system name} If you do not specify the type of name, the system will prioritize in this order:
  - --user
  - --group
  - --system, or
  - --SID (this can be the SID for a user, group or system).

Mandatory arguments are:

- --access: ALLOWED or DENIED
- --permissions: One of FULL, CHANGE, or READ or any combination of RWXDPO.

### **Examples:**

```
mmsmb exportacl add myExport0 myUser --access ALLOWED --permissions FULL
mmsmb exportacl add myExport1 --user user01 --access ALLOWED --permissions RWO
mmsmb exportacl add myExport2 --group group01 --access DENIED -- permissions RWXDP
```

## **change**

Change will update the specified ACE in an export ACL. The options are:

- {User/group/system name} If you do not specify the type of name, the system will prioritize in this order:
  - --user
  - --group
  - --system
  - --SID (This can be the SID for a user, group or system).

Mandatory arguments are:

- --access: ALLOWED or DENIED
- --permissions: One of FULL, CHANGE, or READ or any combination of RWXDPO.

### **Examples:**

```
mmsmb exportacl change myExport --user myUser --access ALLOWED --permissions RWX
mmsmb exportacl change myExport --group allUsers --access ALLOWED --permissions R
```

## **remove**

Remove will remove the ACE for the specified user/group/system from the ACL.

The user, group or system will be removed automatically for a specified name.

In the event that the system is unable to locate an ACE within the export ACL that it can remove the permissions for as instructed, an error will be issued to the user informing them of this.

The options are:

- {User/group/system name} If you do not specify the type of name, the system will prioritize in this order:
  - --user
  - --group
  - --system
  - --SID (This can be the SID for a user, group or system).

Optional arguments are:

- --access: ALLOWED or DENIED
- --permissions: One of FULL, CHANGE, or READ or any combination of RWXDPO.

### **Examples:**

```
mmsmb exportacl remove myExport01 UserName --access ALLOWED --permissions CHANGE
mmsmb exportacl remove myExport02 GroupName --access DENIED --permissions READ
mmsmb exportacl remove myExport03 UserName
```

**replace**

The replace command replaces all the permissions in a export ACL with those indicated in its ACE specification. It is therefore a potentially destructive command and will include a confirmation. This confirmation can be overridden with the --force command. The options are:

- {User/group/system name} If you do not specify the type of name, the system will prioritize in this order:
  - --user
  - --group
  - --system
- --SID (This can be the SID for a user, group or system)
- --force.

Mandatory arguments are:

- --access: ALLOWED or DENIED
- --permissions: One of FULL, CHANGE, or READ or any combination of RWXDP0.

**Examples:**

```
mmsmb exportacl replace myExport01 --user user01 --access ALLOWED --permissions FULL
mmsmb exportacl replace myExport02 --group group01 --access ALLOWED --permissions READ --force
mmsmb exportacl replace myUser --access ALLOWED --permissions FULL
```

**delete**

The Delete command will remove an entire export ACL. It therefore does not require a system, id or user identified as they are not appropriate in this case. All it needs is the name of an export for which the export ACL will be deleted.

Delete will include a confirmation. This can be overridden using the --force parameter.

**Examples:**

```
mmsmb exportacl delete myExport01
mmsmb exportacl delete myExport02 --force
```

**Parameters common for both mmsmb export and mmsmb config commands**

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**--header *n***

Repeats the output table header every *n* lines for a table that is spread over multiple pages. The value *n* can be of any integer value.

**--key-info *arg***

Displays the supported SMB options and their possible values.

*arg* *SMBoption* | supported

***SMBoption***

Specifies the SMB option.

**supported**

Displays descriptions for all the supported SMB options.

## **mmsmb**

### **Exit status**

**0** Successful completion.

#### **nonzero**

A failure has occurred.

### **Security**

You must have root authority to run the **mmsmb** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

### **Examples**

#### **mmsmb config list**

1. Show descriptions of all the supported SMB configuration options.

```
mmsmb config list --key-info supported
```

The system displays output similar to this:

```
Supported smb options with allowed values:
gpfs:dfreequota = yes, no
restrict anonymous = 0, 2
server string = any
```

2. List the SMB option that specifies whether anonymous access is allowed or not.

```
mmsmb config list "gpfs:dfreequota"
```

The system displays output similar to this:

```
SMB option value
gpfs:dfreequota yes
```

3. Display the SMB configuration options in a machine-readable format.

```
mmsmb config list -Y
```

The system displays output similar to this:

```
add share command:aio read size:aio write size:aio_pthread%3A aio open:async smb echo handler:auth methods:change
notify:change share command:client NTLMv2 auth:ctdb locktime warn threshold:debug hires timestamp:delete share
command:dfree cache time:disable netbios:disable spoolss:dapi support:ea support:fileid%3A mapping:force unknown
acl user:gencache%3A stabilize_count:gpfs%3Adfreequota:gpfs%3Ahsm:gpfs%3A leases:gpfs%3A prealloc:gpfs%3A sharemodes:
```

#### **mmsmb config change**

1. Show descriptions of all the supported SMB configuration options.

```
mmsmb config change --key-info supported
```

The system displays output similar to this:

```
Supported smb options with allowed values:
gpfs:dfreequota = yes, no
restrict anonymous = 0, 2
server string = any
```

**Note:** The output of this command depends on the configuration options supported by the system.

2. Change an SMB configuration option.

```
mmsmb config change --option "restrict anonymous=2"
```

You can confirm the change by using this **mmsmb config list "restrict anonymous"** command.

3. Remove an SMB configuration option.

```
mmsmb config change --remove "server string"
```

The system displays output similar to this:

**Warning:**  
Unused options suppressed in display:  
server string

#### **mmsmb export list**

1. To list all SMB options for export `myExport`, issue this command:

```
mmsmb export list myExport --option all
```

The system displays output similar to this:

|          |                   |      |         |          |            |
|----------|-------------------|------|---------|----------|------------|
| export   | path              | smb  | encrypt | guest ok | browseable |
| myExport | /gpfs/fs0/combo-1 | auto | no      | yes      |            |

2. To list SMB option `csc policy` for SMB share `myexport` and all SMB shares starting with `foo` followed by any quantity of 1 and ending on 2 or 3, issue this command:

```
mmsmb export list --option "csc policy" myexport --export-regex "foo1*[23]$"'
```

The system displays output similar to this:

|          |                                      |            |
|----------|--------------------------------------|------------|
| export   | path                                 | csc policy |
| myExport | /mnt/gpfs0/shared                    | - - -      |
| foo11b23 | /mnt/gpfs0/testExport/testSmbEncrypt | disable    |
| foo11b23 | /mnt/gpfs0/testExport/testSmbEncrypt | documents  |

3. To list all exports where unsupported SMB options are set, issue this command:

```
mmsmb export list --option unsupported
```

The system displays output similar to this:

|                     |                   |               |
|---------------------|-------------------|---------------|
| export              | path              | mangled names |
| hasForbiddenOptions | /mnt/gpfs0/shared | yes           |

#### **mmsmb export add**

1. To create an export `myExport` as default SMB share for path `/ibm/gpfs0/myFolder` with default options, issue this command:

```
mmsmb export add myExport "/ibm/gpfs0/myFolder"
```

The system displays output similar to this:

```
mmsmb export add: The SMB export was created successfully.
```

#### **mmsmb export change**

1. To change the SMB option `oplocks` to value `yes` for SMB share `myExport`, issue this command:

```
mmsmb export change myExport --option "oplocks"="yes"
```

You can confirm the change by using this `mmsmb export list --option "oplocks" myExport` command. The system displays output similar to this:

|          |                   |         |
|----------|-------------------|---------|
| export   | path              | oplocks |
| myExport | /mnt/gpfs0/shared | yes     |

2. To remove the SMB option `oplocks` from SMB share `myExport`, issue the following commands:

```
mmsmb export change myExport --remove "oplocks"
```

You can confirm the change by using this `mmsmb export list --option all myExport` command.

#### **mmsmb export remove**

1. To delete the SMB share `myExport` with confirmation, issue this command:

## **mmsmb**

```
mmsmb export remove myExport
```

The system asks for confirmation, similar to this:

```
Do you really want to perform the operation (yes/no - default no):
```

2. To delete the SMB share `myExport` without confirmation:

```
mmsmb export remove myExport --force
```

The system removes the SMB share without any confirmation.

### **See also**

- “`mmnfs` command” on page 454
- “`mnces` command” on page 103
- “`mmlsfs` command” on page 411

### **Location**

`/usr/lpp/mmfs/bin`

---

## mmsnapdir command

Controls how the special directories that connect to snapshots appear.

### Synopsis

```
mmsnapdir Device [-r | -a]
 [--show-global-snapshots {rootfileset | allfilesets}]
 [{[--fileset-snapdir FilesetSnapDirName] [--global-snapdir GlobalSnapDirName]}]
 [{--snapdir | -s SnapDirName}]
```

or

```
mmsnapdir Device [-q]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmsnapdir** command to control how the special directories that connect to snapshots appear. Both the name of the directories and where they are accessible can be changed.

Global snapshots appear in a subdirectory in the root directory of the file system, whose default name is **.snapshots**. Fileset snapshots appear in a similar **.snapshots** subdirectory located in the root directory of each independent fileset. These special subdirectories are collectively referred to as *snapdirs*. Note that the root directory of the file system and the root directory of the root fileset are the same, so global snapshots and fileset snapshots of the root fileset will both appear in the same snapdir.

If you prefer to access the snapshots from each directory rather than traversing through the root directory, you can use an invisible directory to make the connection by issuing the **mmsnapdir** command with the **-a** option (see “Examples”). The **-a** option enables an invisible directory in each directory in the active file system (they do not appear in directories in snapshots) that contains a subdirectory for each existing snapshot of the file system (in the root fileset) or fileset (in other independent filesets). These subdirectories correspond to the copy of the active directory in the snapshot with the same name. For example, if you enter **ls -a /fs1/userA**, the (invisible) **.snapshots** directory is not listed. However, you can use **ls /fs1/userA/.snapshots**, for example, to confirm that **.snapshots** is present and contains the snapshots holding copies of **userA**. When the **-a** option is enabled, the paths **/fs1/.snapshots/Snap17/userA** and **/fs1/userA/.snapshots/Snap17** refer to the same directory, namely **userA** at the time when **Snap17** was created. The **-r** option (root-directories-only), which is the default, reverses the effect of the **-a** option (all-directories), and disables access to snapshots via snapdirs in non-root directories.

If you prefer to access global snapshots from the root directory of all independent filesets, use the **mmsnapdir** command with the **--show-global-snapshots allfilesets** option. With this option, global snapshots will also appear in the snapdir in the fileset root directory. The global snapshots will also appear in the snapdirs in each non-root directory if all-directories (the **-a** option) is enabled. To return to the default setting use **--show-global-snapshots rootfileset**, and global snapshots will only be available in root of the file system, or the root fileset, if all-directories is enabled.

The name of the snapdir directories can be changed using the **--snapdir** (or **-s**) option. This name is used for both global and fileset snapshots in both fileset root directories and, if all-directories is enabled, non-root directories also. The snapdir name for global and fileset snapshots can be specified separately using the **--global-snapdir** and **--fileset-snapdir** options. If these names are different, two snapdirs will appear in the file system root directory, with the global and fileset snapshots listed separately. When **--show-global-snapshots** is set to **allfilesets**, two snapdirs will appear in fileset root directories also, and when all-directories (the **-a** option) is specified, the two snapdirs will be available in non-root directories as well. If **--global-snapdir** is specified by itself, the fileset snapdir name is left unchanged, and vice versa.

## mmsnapdir

if **--fileset-snapdir** option is used. Setting both snapdirs to the same name is equivalent to using the **--snapdir** option. The snapdir name enabled in non-root directories by all-directories is always the same as the name used in root directories.

For more information on global snapshots, see *Creating and maintaining snapshots of file systems* in the *IBM Spectrum Scale: Administration Guide*.

For more information on fileset snapshots, see *Fileset-level snapshots* in the *IBM Spectrum Scale: Administration Guide*.

## Parameters

### *Device*

The device name of the file system. File system names need not be fully-qualified. **fs0** is just as acceptable as **/dev/fs0**.

This must be the first parameter.

- a** Adds a snapshots subdirectory to all subdirectories in the file system.
- r** Reverses the effect of the **-a** option. All invisible snapshot directories are no longer accessible. The snapshot directory under the file system root directory is not affected.

### **--show-global-snapshots {rootfileset | allfilesets}**

This option controls whether global snapshots are accessible through a subdirectory under the root directory of all independent filesets (**allfilesets**) or only in the file system root (**rootfileset**). For example, issuing the following command:

```
mmsnapdir fs1 --show-global-snapshots allfilesets
```

specifies that the root directory of each independent fileset will contain a **.gsnaps** subdirectory listing all global snapshots, such as **/fs1/junctions/FsetA/.gsnaps** and **/fs1/junctions/FsetA/.fsnaps**. This can be used to make global snapshots accessible to clients, for example NFS users, that do not have access to the file system root directory.

Specifying **rootfileset** reverses this feature, restoring the default condition, so that global snapshots are only visible in the root fileset.

### **--fileset-snapdir FilesetSnapDirName**

### **--global-snapdir GlobalSnapDirName**

The **--global-snapdir** option specifies the name for the directory where global snapshots are listed.

The **--fileset-snapdir** option specifies the name for the directory where fileset snapshots are listed.

These options can be specified together or separately, in which case only the corresponding snapdir is changed. Neither option may be specified with **--snapdir**, which sets both to the same name.

For example, after issuing the command:

```
mmsnapdir fs1 --fileset-snapdir .fsnaps --global-snapdir .gsnaps
```

the directory **/fs1/.gsnaps** will list all global snapshots and **/fs1/.fsnaps** will only list fileset snapshots of the root fileset. Fileset snapshots of other independent filesets will be listed in **.fsnaps** under the root directory of each independent fileset, such as **/fs1/junctions/FsetA/.fsnaps**.

### **--snapdir | -s SnapDirName**

Changes the name of the directory for both global and fileset snapshots to *SnapDirName*. This affects both the directory in the file system root as well as the invisible directory in the other file system directories if the **-a** option has been enabled. The root and non-root snapdirs cannot be given different names.

- q** Displays current snapshot settings. The **-q** option cannot be specified with any other options. This is the default if no other options are specified.

## Exit status

**0** Successful completion.

### nonzero

A failure has occurred.

## Security

If you are a root user, the node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see the topic *Requirements for administering a GPFS file system* in the *IBM Spectrum Scale: Administration Guide*.

You must be a root user to use all of the **mmsnapdir** options. Non-root users can only use the **-q** option.

If you are a non-root user, you may only specify file systems that belong to the same cluster as the node on which the **mmsnapdir** command was issued.

## Examples

1. To rename the **.snapshots** directory (the default snapshots directory name) to **.link** for file system **fs1**, issue the command:

```
mmsnapdir fs1 -s .link
```

After the command has been issued, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3
```

```
/fs1/.link/snap1/file1
/fs1/.link/snap1/userA/file2
/fs1/.link/snap1/userA/file3
```

2. To add the **.link** subdirectory to all subdirectories in the file system, issue:

```
mmsnapdir fs1 -a
```

After the command has been issued, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3
/fs1/userA/.link/snap1/file2
/fs1/userA/.link/snap1/file3
```

```
/fs1/.link/snap1/file1
/fs1/.link/snap1/userA/file2
/fs1/.link/snap1/userA/file3
```

The **.link** subdirectory under the root directory and under each subdirectory of the tree provides two different paths to each snapshot copy of a file. For example, **/fs1/userA/.link/snap1/file2** and **/fs1/.link/snap1/userA/file2** are two different paths that access the same snapshot copy of **/fs1/userA/file2**.

3. To reverse the effect of the previous command, issue:

```
mmsnapdir fs1 -r
```

After the command has been issued, the directory structure would appear similar to:

```
/fs1/file1
/fs1/userA/file2
/fs1/userA/file3
```

## **mmsnapdir**

```
/fs1/.link/snap1/file1
/fs1/.link/snap1/userA/file2
/fs1/.link/snap1/userA/file3
```

4. To display the current snapshot settings, issue:

```
mmsnapdir fs1 -q
```

The system displays output similar to:

```
Snapshot directory for "fs1" is ".link" (root directory only)
```

If there are independent filesets, fileset snapshots, or the global and fileset snapshot directory names are different in the file system, the system displays output similar to:

```
Fileset snapshot directory for "fs1" is ".link" (root directory only)
```

```
Global snapshot directory for "fs1" is ".link" in root directory only
```

### **See also**

- “mmcrsnapshot command” on page 267
- “mmdelsnapshot command” on page 304
- “mmlssnapshot command” on page 439
- “mmrestorefs command” on page 525

### **Location**

```
/usr/lpp/mmfs/bin
```

---

## mmstartup command

Starts the GPFS subsystem on one or more nodes.

### Synopsis

```
mmstartup [-a | -N {Node[,Node...] | NodeFile | NodeClass}] [-E EnvVar=value ...]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Use the **mmstartup** command to start the GPFS daemons on one or more nodes. If no operand is specified, GPFS is started only on the node from which the command was issued.

### Parameters

**-a** Start GPFS on all nodes in a GPFS cluster.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Directs the **mmstartup** command to process a set of nodes.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

**-E EnvVar=value**

Specifies the name and value of an environment variable to be passed to the GPFS daemon. You can specify multiple **-E** options.

### Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

### Security

You must have root authority to run the **mmstartup** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

### Examples

To start GPFS on all nodes in the GPFS cluster, issue this command:

```
mmstartup -a
```

The system displays information similar to:

```
Thu Aug 12 13:22:40 EDT 2004: 6027-1642 mmstartup: Starting GPFS ...
```

## **mmstartup**

### **See also**

- “[mmgetstate command](#)” on page 350
- “[mmlscluster command](#)” on page 397
- “[mmshutdown command](#)” on page 552

### **Location**

/usr/lpp/mmfs/bin

## mmtracectl command

Sets up and enables GPFS tracing.

### Synopsis

```
mmtracectl { --start | --stop | --off | --set }
[--trace={io | all | def | "Class Level [Class Level ...]"}]
[--trace-recycle={off | local | global | globalOnShutdown}]
[--aix-trace-buffer-size=BufferSize]
[--tracedev-buffer-size=BufferSize]
[--trace-file-size=FileSize] [--trace-dispatch={yes | no}]
[--tracedev-compression-level=Level]
[--tracedev-write-mode={blocking | overwrite}]
[--tracedev-timeformat={relative | absolute | calendar}]
[--tracedev-overwrite-buffer-size=Size]
[--format | --noformat]
[-N {Node [,Node...] | NodeFile | NodeClass}]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

**Attention:** Use this command only under the direction of the IBM Support Center.

Use the **mmtracectl** command to perform the following functions:

- Start or stop tracing.
- Turn tracing on (start or set trace recycle) or off on the next session. This is a persistent setting to automatically start trace each time GPFS starts.
- Allow for predefined trace levels: **io**, **all**, and **def**, as well as user-specified trace levels.
- Allow for changing the size of trace buffer sizes for AIX and all others using the **tracedev** option.
- Trace recycle functions, which allow for never cycling traces (**off** option), cycling traces on all nodes when GPFS ends abnormally (**global** option), and cycling traces any time GPFS goes down on all nodes (**globalOnShutdown** option).
- For Linux nodes only, this command allows you to change:
  - The trace writing mode
  - The raw data compression level

**Note:** Tracing on Windows requires support programs provided by Microsoft. For details about this prerequisite, see the section about configuring Windows and installing tracing support programs in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### Results

GPFS tracing can be started, stopped, or related configuration options can be set.

### Parameters

**--start | --stop | --off | --set**

Specifies the actions that the **mmtracectl** command performs, where:

**--start**

Starts the trace.

**--stop**

Stops the trace.

## **mmtracectl**

### **--off**

Clears all of the setting variables and stops the trace.

### **--set**

Sets the trace variables.

### **--trace={io | all | def | "Class Level [Class Level ...]"}**

Allows for predefined and user-specified trace levels, where:

**io** Indicates trace-level settings tailored for input and output (I/O).

#### **all**

Sets trace levels to their highest setting (9).

#### **def**

Indicates that the default trace settings will be used.

#### **"Class Level [Class Level ...]"**

Specifies a trace class and level.

### **--trace-recycle={off | local | global | globalOnShutdown}**

Controls trace recycling during daemon termination. The following values are recognized:

#### **off**

Does not recycle traces. This is the default value until **mmtracectl --start** is run. If no trace-recycle value has been explicitly set when **mmtracectl --start** is run, see the following description for **local**. The **mmtracectl --off** command will remove any explicit value for **trace-recycle**, thus effectively setting **trace-recycle** back to the **off** value.

#### **local**

Recycles traces on the local node when **mmfsd** goes down abnormally. This setting also starts traces automatically any time that **mmstartup** is run to start the GPFS daemon, which could include an autoload on a reboot. If there is no **trace-recycle** value explicitly set at the time that **mmtracectl --start** is run, **mmchconfig** will be run to explicitly set the **trace-recycle** value to **local**. The starting of the actual tracing will be delayed while that configuration change is being made.

#### **global**

Recycles traces on all nodes in the cluster when an abnormal daemon shutdown occurs.

#### **globalOnShutdown**

Recycles traces on all nodes in the cluster for normal and abnormal daemon shutdowns.

### **--aix-trace-buffer-size=BufferSize**

Controls the size of the trace buffer in memory for AIX.

### **--tracedev-buffer-size=BufferSize**

Specifies the trace buffer size for Linux trace in blocking mode. If **--tracedev-write-mode** is set to blocking, this parameter will be used. It should be no less than 4K and no more than 64M. The default is 4M.

**Note:** This option applies only to Linux nodes.

### **--trace-file-size=FileSize**

Controls the size of the trace file. The default is 128M on Linux and 64M on other platforms.

### **--trace-dispatch={yes | no}**

Enables AIX thread dispatching trace hooks.

### **--tracedev-compression-level=Level**

Specifies the trace raw data compression level. Valid values are 0 to 9. A value of zero indicates no compression. A value of 9 provides the highest compression ratio, but at a lower speed. The default is 6.

**Note:** This option applies only to Linux nodes.

**--tracedev-write-mode={blocking | overwrite}**

Specifies when to overwrite the old data, where:

**blocking**

Specifies that if the trace buffer is full, wait until the trace data is written to the local disk and the buffer becomes available again to overwrite the old data.

**overwrite**

Specifies that if the trace buffer is full, overwrite the old data. This is the default.

**Note:** This option applies only to Linux nodes.

**--tracedev-timeformat={relative | absolute | calendar}**

Controls time formatting in the trace records. The following values are accepted:

**relative**

Displays the trace time stamp in relative format, showing the number of seconds from the beginning time stamp. This is the default.

**absolute**

Displays the trace time stamp in absolute format, showing the number of seconds since 1/1/1970.

**calendar**

Displays the trace time stamp in local calendar format, showing day of the week, month, day, hours, minutes, seconds, and year.

**--tracedev-overwrite-buffer-size=Size**

Specifies the trace buffer size for Linux trace in overwrite mode. If **--tracedev-write-mode** is set to overwrite, this parameter will be used. It should be no less than 16M. The default is 64M.

**Note:** This option applies only to Linux nodes.

**--format | --noformat**

Enables or disables formatting.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes that will participate in the tracing of the file system. This option supports all defined node classes (with the exception of **mount**). The default value is **all**.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

## Exit status

0      Successful completion.

**nonzero**

A failure has occurred.

## Security

You must have root authority to run the **mmtracectl** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

To set trace levels to the defined group **def** and start the traces on all nodes when GPFS starts, issue the following command:

## **mmtracectl**

```
mmtracectl --set --trace=def --trace-recycle=global
```

The system displays the following output:

```
mmchconfig: Command successfully completed
mmchconfig: Propagating the cluster configuration data to all
affected nodes. This is an asynchronous process.
```

To confirm the change, issue the following command:

```
| mmIsconfig trace,traceRecycle
```

The system displays output like the following example:

```
trace all 4 tm 2 thread 1 mutex 1 vnode 2 ksvfs 3 klockl 2 io 3 palloc 1 mb 1 lock 2 fsck 3
traceRecycle global
```

To manually start traces on all nodes, issue the following command:

```
mmtracectl --start
```

### **See also**

- “mmchconfig command” on page 132

See the **mmtrace** shell script.

### **Location**

```
/usr/lpp/mmfs/bin
```

---

## mmumount command

Unmounts GPFS file systems on one or more nodes in the cluster.

### Synopsis

```
mmumount {Device | MountPoint | DriveLetter |
 all | all_local | all_remote | {-F DeviceFileName} |
 [-f] [-a] [-N {Node[,Node...]} | NodeFile | NodeClass]}
```

or

```
mmumount Device -f -C {all_remote | ClusterName} [-N Node[,Node...]]
```

### Availability

Available on all IBM Spectrum Scale editions.

### Description

Another name for the **mmumount** command is the **mmunmount** command. Either name can be used.

The **mmumount** command unmounts a previously mounted GPFS file system on one or more nodes in the cluster. If no nodes are specified, the file systems are unmounted only on the node from which the command was issued. The file system can be specified using its device name or the mount point where it is currently mounted.

Use the first form of the command to unmount file systems on nodes that belong to the local cluster.

Use the second form of the command with the **-C** option when it is necessary to force an unmount of file systems that are owned by the local cluster, but are mounted on nodes that belong to another cluster.

When a file system is unmounted by force with the second form of the **mmumount** command, the affected nodes may still show the file system as mounted, but the data will not be accessible. It is the responsibility of the system administrator to clear the mount state by issuing the **umount** command.

When multiple nodes are affected and the unmount target is identified via a mount point or a Windows drive letter, the mount point is resolved on each of the target nodes. Depending on how the file systems were mounted, this may result in different file systems being unmounted on different nodes. When in doubt, always identify the target file system with its device name.

### Parameters

*Device | MountPoint | DriveLetter | all | all\_local | all\_remote | {-F DeviceFileName}*  
Indicates the file system or file systems to be unmounted.

*Device*

Is the device name of the file system to be unmounted. File system names do not need to be fully qualified. **fs0** is as acceptable as **/dev/fs0**.

*MountPoint*

Is the location where the GPFS file system to be unmounted is currently mounted.

*DriveLetter*

Identifies a file system by its Windows drive letter.

**all**

Indicates all file systems that are known to this cluster.

**all\_local**

Indicates all file systems that are owned by this cluster.

## **mmumount**

### **all\_remote**

Indicates all file systems that are owned by another cluster to which this cluster has access.

### **-F DeviceFileName**

Specifies a file containing the device names, one per line, of the file systems to be unmounted.

This must be the first parameter.

## **Options**

**-a** Unmounts the file system on all nodes in the GPFS cluster.

**-f** Forces the unmount to take place even though the file system may be still in use.

Use this flag with *extreme caution*. Using this flag may cause outstanding write operations to be lost. Because of this, forcing an unmount can cause data integrity failures and should be used with caution.

The **mmumount** command relies on the native **umount** command to carry out the unmount operation. The semantics of forced unmount are platform-specific. On some platforms (such as Linux), even when forced unmount is requested, a file system cannot be unmounted if it is still referenced by the system kernel. Examples of such cases are:

- Open files are present in the file system
- A process uses a subdirectory in the file system as the current working directory
- The file system is NFS-exported

To unmount a file system successfully in such a case, it may be necessary to identify and stop the processes that are referencing the file system. System utilities like **lsof** and **fuser** could be used for this purpose.

### **-C {all\_remote | ClusterName}**

Specifies the cluster on which the file system is to be unmounted by force. **all\_remote** denotes all clusters other than the one from which the command was issued.

### **-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the nodes on which the file system is to be unmounted.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands in the IBM Spectrum Scale: Administration Guide*.

This command does not support a *NodeClass* of **mount**.

When the **-N** option is specified in conjunction with **-C ClusterName**, the specified node names are assumed to refer to nodes that belong to the specified remote cluster (as identified by the **mmlsmount** command). The **mmumount** command cannot verify the accuracy of this information. *NodeClass* and *NodeFile* are not supported in conjunction with the **-C** option.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmumount** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system in IBM Spectrum Scale: Administration Guide*.

## Examples

1. To unmount file system **fs1** on all nodes in the cluster, issue this command:

```
mmumount fs1 -a
```

The system displays output similar to:

```
Fri Feb 10 15:51:25 EST 2006: mmumount: Unmounting file systems ...
```

2. To force unmount file system **fs2** on the local node, issue this command:

```
mmumount fs2 -f
```

The system displays output similar to:

```
Fri Feb 10 15:52:20 EST 2006: mmumount: Unmounting file systems ...
forced unmount of /fs2
```

## See also

- “mmmount command” on page 444
- “mmlsmount command” on page 419

## Location

/usr/lpp/mmfs/bin

## **mmunlinkfileset**

---

### **mmunlinkfileset command**

Removes the junction to a GPFS fileset.

#### **Synopsis**

```
mmunlinkfileset Device {FilesetName | -J JunctionPath} [-f]
```

#### **Availability**

Available with IBM Spectrum Scale Standard Edition or higher.

#### **Description**

The **mmunlinkfileset** command removes the junction to the fileset. The junction can be specified by path or by naming the fileset that is its target. The unlink fails if there are files open in the fileset, unless the **-f** flag is specified. The root fileset may not be unlinked.

**Attention:** If you are using the IBM Spectrum Protect Backup-Archive client, use caution when you unlink filesets that contain data backed up by IBM Spectrum Protect. IBM Spectrum Protect tracks files by pathname and does not track filesets. As a result, when you unlink a fileset, it appears to IBM Spectrum Protect that you deleted the contents of the fileset. Therefore, the IBM Spectrum Protect Backup-Archive client inactivates the data on the IBM Spectrum Protect server which may result in the loss of backup data during the expiration process.

For information on GPFS filesets, see the *IBM Spectrum Scale: Administration Guide*.

#### **Parameters**

##### *Device*

The device name of the file system that contains the fileset.

File system names need not be fully-qualified. **fs0** is as acceptable as **/dev/fs0**.

##### *FilesetName*

Specifies the name of the fileset to be removed.

##### **-J JunctionPath**

Specifies the name of the junction to be removed.

A junction is a special directory entry that connects a name in a directory of one fileset to the root directory of another fileset.

**-f** Forces the unlink to take place even though there may be open files. This option forcibly closes any open files, causing an **errno** of **ESTALE** on their next use of the file.

#### **Exit status**

**0** Successful completion.

##### **nonzero**

A failure has occurred.

#### **Security**

You must have root authority to run the **mmunlinkfileset** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## Examples

1. This command indicates the current configuration of filesets for file system **gpfs1**:

```
mmlsfileset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |        |              |
|----------------------------------|--------|--------------|
| Name                             | Status | Path         |
| root                             | Linked | /gpfs1       |
| fset1                            | Linked | /gpfs1/fset1 |

This command unlinks fileset **fset1** from file system **gpfs1**:

```
mmunlinkfileset gpfs1 fset1
```

The system displays output similar to:

Fileset 'fset1' unlinked.

To confirm the change, issue this command:

```
mmlsfileset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |          |        |
|----------------------------------|----------|--------|
| Name                             | Status   | Path   |
| root                             | Linked   | /gpfs1 |
| fset1                            | Unlinked | --     |

2. This command indicates the current configuration of filesets for file system **gpfs1**:

```
mmlsfileset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |        |              |
|----------------------------------|--------|--------------|
| Name                             | Status | Path         |
| root                             | Linked | /gpfs1       |
| fset1                            | Linked | /gpfs1/fset1 |

This command unlinks junction path **/gpfs1/fset1** from file system **gpfs1**:

```
mmunlinkfileset gpfs1 -J /gpfs1/fset1
```

The system displays output similar to:

Fileset 'fset1' unlinked.

To confirm the change, issue this command:

```
mmlsfileset gpfs1
```

The system displays output similar to:

| Filesets in file system 'gpfs1': |          |        |
|----------------------------------|----------|--------|
| Name                             | Status   | Path   |
| root                             | Linked   | /gpfs1 |
| fset1                            | Unlinked | --     |

## See also

- “mmchfileset command” on page 173
- “mmcrfileset command” on page 244
- “mmdelfileset command” on page 292
- “mmlinkfileset command” on page 390
- “mmlsfileset command” on page 406

**mmunlinkfileset**

**Location**

/usr/lpp/mmfs/bin

## mmuserauth command

Manages the authentication configuration of file and object access protocols. The configuration allows protocol access methods to authenticate users who need to access data that is stored on the system over these protocols.

### Synopsis

```
mmuserauth service create --data-access-method{file|object}
 --type {ldap|local|ad|nis|userdefined}
 --servers[IP address/hostname]
 [--base-dn]
 {[--enable-anonymous-bind] | [--user-name] [--password]}
 [--enable-server-tls] [--enable-ks-ssl]
 [--enable-kerberos] [--enable-nfs-kerberos] [--enable-ks-casigning]
 [--user-dn] [--group-dn] [--netgroup-dn]
 [--netbios-name] [--domain]
 [--idmap-role{master|subordinate}] [--idmap-range] [--idmap-range-size]
 [--user-objectclass] [--group-objectclass] [--user-name-attrib]
 [--user-id-attrib] [--user-mail-attrib] [--user-filter]
 [--ks-dns-name] [--ks-admin-user] [--ks-admin-pwd]
 [--ks-swift-user] [--ks-swift-pwd] [--ks-ext-endpoint]
 [--kerberos-server] [--kerberos-realm]
 [--unixmap-domains] [ldapmap-domains]
```

Or

```
mmuserauth service list [-Y] [--data-access-method {file|object|all}]
```

Or

```
mmuserauth service check [--data-access-method {file|object|all}] [-r|--rectify]
 [-N|--nodes {node-list|cesNodes}] [--server-reachability]
```

Or

```
mmuserauth service remove --data-access-method {file|object|all} [--idmapdelete]
```

### Availability

Available with IBM Spectrum Scale Standard Edition or higher.

### Description

Use the **mmuserauth** commands to create and manage IBM Spectrum Scale protocol authentication and ID mappings.

### Parameters

#### service

Manages the authentication configuration of file and object access protocols with one of the following actions:

##### create

Configures authentication for file and object access protocols. The authentication method for file and object access protocols cannot be configured together. The **mmuserauth service create** command needs to be submitted separately for configuring authentication for the file and object access protocols each.

##### list

Displays the details of the authentication method that is configured for both file and object access protocols.

## **mmuserauth**

### **check**

Verifies the authentication method configuration details for file and object access protocols. Validates the connectivity to the configured authentication servers. It also supports corrections to the configuration details on the erroneously configured protocol nodes.

### **remove**

Removes the authentication method configuration of file and object access protocols and ID maps if any.

If you plan to remove both, authentication method configuration and ID maps, remove authentication method configuration followed by the removal of ID maps. That is, at first you need to submit the **mmuserauth service remove** command without the **--idmapdelete** option to remove the authentication method configuration and then submit the same command with the **--idmapdelete** option to remove ID maps.

**Note:** CAUTION: Deleting the authentication method configuration with the ID maps can lead to irrecoverable loss of access to data. Use this option with proper planning.

### **--data-access-method {file|object}**

Specifies the access protocols for which the authentication method needs to be configured. The IBM Spectrum Scale system supports file access protocols such as SMB and NFS along with Object access protocols to access data that is stored on the system.

The file data access method is meant for authorizing the users who access data over SMB and NFS protocols.

### **--type {ldap|local|ad|nis|userdefined}**

Specifies the authentication method to be configured for accessing data over file and object access protocols.

ldap - Defines an external LDAP server as the authentication server. This authentication type is valid for both file and object access protocols.

ad - Defines an external Microsoft Active Directory server as the authentication server. This authentication type is valid for both file and object access protocols.

local - Defines an internal database stored on IBM Spectrum Scale protocol nodes for authenticating user accessing data over object access protocol. This authentication type is valid for Object access protocol only.

nis - Defines an external NIS server as the authentication server. This authentication type is only valid for NFS file access protocol only.

userdefined - Defines user-defined ( system administrator defined ) authentication method for data access. This authentication type is valid for both file and object access protocols.

### **--servers [AuthServer1[:Port],AuthServer2[:Port],AuthServer3[:Port] ...]**

Specifies the host name or IP address of the authentication server that is used for file and object access protocols.

This option is only valid with **--type {ldap|ad|nis}**.

With **--type ldap**, the input value format is "serverName/serverIP:[port]".

Specifying the port value is optional. Default port is 389.

For example,

**--servers ldapserver.mydomain.com:1389.**

For file access protocol, multiple LDAP servers can be specified by using a comma as a separator.

For the object access protocol, only one authentication server must be specified. If multiple servers are specified by using a comma, only the first server in the list is considered as the authentication server for configuration.

With --type ad, the input value format is "serverName/serverIP".

For example,

```
--servers ldapserver.mydomain.com.
```

For the file access protocol, only one authentication server must be specified. Specifying multiple servers is invalid. The AD server accepted while configuration is used to fetch details required for validation and configuration of the authentication method. Post successful configuration, each CES node will query DNS to lookout available Domain Controllers serving the AD domain it is joined to. Amongst the returned list, the node binds with the best available domain controller.

For the object access protocol, only one authentication server must be specified. If multiple servers are specified by using a comma, only the first server in the list is considered as the authentication server.

With --type nis, the input value format is "serverName/serverIP".

For example,

```
--servers nisserver.mydomain.com.
```

Multiple NIS servers can be specified by using a comma separator. At least one of the specified servers must be available and reachable while configuring the authentication method. This is important for the verification of the specified NIS domain, against which the availability of either passwdbyname or netgroup map is validated.

#### **--base-dn *ldapBase***

Specifies the LDAP base DN of the authentication server. This option is only valid with --type {ldap|ad} for --data-access-method object and --type ldap for --data-access-method file.

#### **--enable-anonymous-bind**

Specifies whether to enable anonymous binding with authentication server for various validation operations.

This option is only valid with --type {ldap|ad} and --data-access-method {object}. This option is mutually exclusive with --user-name and --password.

#### **--user-name *userName***

Specifies the user name to be used to perform operations against the authentication server. This option is only valid with --type {ldap|ad} and --data-access-method {file|object}.

This option along with --password is mutually exclusive with --enable-anonymous-bind. The specified user name must have sufficient permissions to read user and group attributes from the authentication server.

In case of --type {ad|ldap} with --data-access-method object, the user name must be specified in complete DN format. In case of --type ad with --data-access-method file, the specified username is used to join the cluster to AD domain. It results in creating a machine account for the cluster based on the --netbios-name specified in the command. After successful configuration, the cluster connects with its machine account, and not the user used during the domain join. So the specified username after domain join has no role to play in communication with the AD domain controller and can be even deleted from the AD server. The cluster can still keep using AD for authentication via the machine account created.

#### **--password *userPassword***

Specifies the password of the user name that is specified with the --user-name option. This option is only valid with --type {ldap|ad} and --data-access-method {file|object}.

With --type ad for file authentication, the specified password is only required during the domain joining period. After joining the domain, the password of the machine account of the cluster is used for accessing Active Directory.

The password must be in clear text. To hide the password, submit the command without this option and then the system prompts you to enter the password.

## **mmuserauth**

### **--enable-server-tls**

Specifies whether to enable TLS communication with the authentication server. With --data-access-method object, this option is only valid with --type {ldap|ad}. With --data-access-method file, this option is only valid with --type {ldap}.

This option is disabled by default.

For file access protocol configuration, ensure that the CA certificate is placed in the /var/mmfs/tmp/ directory with the name ldap\_cacert.pem on the node that the command is to be run.

For object access protocol configuration, ensure that the CA certificate is placed in the /var/mmfs/tmp/ directory with the name object\_ldap\_cacert.pem on the node that the command is to be run.

### **--enable-nfs-kerberos**

Specifies whether to enable Kerberized logins for users gaining access by using the NFSv3 and NFSv4 file access protocols.

This option is only valid with --type {ad} and --data-access-method {file}.

This option is disabled by default.

**Note:** Kerberized NFSv3 and NFSv4 access is only supported for users from AD domains that are configured for fetching the UID/GID information from Active Directory (RFC2307 schema attributes). Such AD domain definition is specified by using the --unixmap-domains option.

### **--user-dn ldapUserDN**

Specifies the LDAP group DN. Restricts search of groups within the specified sub-tree. For CIFS access, the value of this parameter is ignored and a search is performed on the baseDN.

This option is only valid with --type {ldap} and --data-access-method {file}. If this parameter is not set, the system uses the value that is set for baseDN as the default value.

### **--group-dn ldapGroupDN**

Specifies the LDAP group suffix. Restricts search of groups within a specified sub-tree.

This option is only valid with --type {ldap} and --data-access-method {file}. If this parameter is not set, the system uses the value that is set for baseDN as the default value.

### **--netgroup-dn ldapGroupDN**

Specifies the LDAP netgroup suffix. The system searches the netgroups based on this suffix. The value must be specified in complete DN format.

This option is only valid with --type {ldap} and --data-access-method {file}. Default value is baseDN.

### **--user-objectclass userObjectClass**

Specifies the object class of user on the authentication server. Only users with specified object class along with other filter are treated as valid users.

If the --data-access-method is object, this option is only valid with --type {ldap|ad}.

If the --data-access-method is file, this option is only valid with --type {ldap}. With --type ldap, the default value is posixAccount and with --type ad the default value is organizationalPerson.

### **--group-objectclass groupObjectClass**

Specifies the object class of group on the authentication server. This option is only valid with --type {ldap} and --data-access-method {file}.

### **--netbios-name netBiosName**

Specifies the unique identifier of the resources on a network that are running NetBIOS. This option is only valid with --type {ad|ldap} and --data-access-method {file}.

The NetBIOS name is limited to 15 ASCII characters and must not contain any white space or one of the following characters: / : \* ? . " ; |

If AD is selected as the authentication method, the NetBIOS name must be selected carefully. If there are name collisions across multiple IBM Spectrum Scale clusters, or between the AD Domain and the NetBIOS name, the configuration does not work properly. Consider the following points while planning for a naming strategy:

- There must not be NetBIOS name collision between two IBM Spectrum Scale clusters that are configured against the same Active Directory server.
- The domain join of the latter machines revokes the join of the former one.
- The NetBIOS name and the domain name must not collide.
- The NetBIOS name and the short name of the Domain Controllers hosting the domain must not collide.

#### **--domain *domainName***

Specifies the name of the NIS domain. This option is only valid with --type {nis} and --data-access-method {file}.

The NIS domain that is specified must be served by one of the servers specified with --server. This option is mandatory when NIS-based authentication is configured for file access.

#### **--idmap-role {master|subordinate}**

Specifies the ID map role of the IBM Spectrum Scale system. ID map role of a stand-alone or singular system deployment must be selected "master". The value of the ID map role is important in AFM-based deployments.

This option is only valid with --type {ad} and --data-access-method {file}.

You can use AD with automatic ID mapping to set up two or more storage subsystems in AFM relationship. The two or more systems configured in a master-subordinate relationship provides a means to synchronize the UIDs and GIDs generated for NAS clients on one system with UIDs and GIDs on the other systems. In the AFM relationship, only one system can be configured as master and other systems must be configured as subordinates. The ID map role of master and subordinate systems are the following:

- **Master:** System creates ID maps on its own.
- **Subordinate:** System does not create ID maps on its own. ID maps must be exported from the master to the subordinate.

While using automatic ID mapping, in order to have same ID maps on systems sharing AFM relationship, you need to export the ID mappings from master to subordinate. The NAS file services are inactive on the subordinate system. If you need to export and import ID maps from one system to another, contact the IBM Support Center.

#### **--idmap-range *lowerValue-higherValue***

Specifies the range of values from which the IBM Spectrum Scale UIDs and GIDs are assigned by the system to the Active Directory users and groups. This option is only valid with --type {ad} and --data-access-method {file}. The default value is 10000000-299999999. The lower value of the range must be at least 1000. After configuring the IBM Spectrum Scale system with AD authentication, only the higher value can be increased (this essentially increases the number of ranges).

#### **--idmap-range-size *rangeSize***

Specifies the total number of UIDs and GIDs that are assignable per domain. For example, if --idmap-range is defined as 10000000-299999999, and range size is defined as 1000000, 290 domains can be mapped, each consisting of 1000000 IDs.

Choose a value for range size that allows for the highest anticipated RID value among all of the anticipated AD users and AD groups in all of the anticipated AD domains. Choose the range size value carefully because range size cannot be changed after the first AD domain is defined on the IBM Spectrum Scale system.

This option is only valid with --type {ad} and --data-access-method {file}. Default value is 1000000.

## **mmuserauth**

### **--unixmap-domains unixDomainMap**

Specifies the AD domains for which user ID and group ID should be fetched from the AD server. This option is only valid with --type {ad} and --data-access-method {file}. The unixDomainMap takes value in this format: DOMAIN1(L1-H1) [;DOMAIN2(L2-H2) [;DOMAIN3(L3-H3) ....]]

#### **DOMAIN**

Use DOMAIN to specify an AD domain for which ID mapping services are to be configured. The name of the domain to be specified must be the NetBIOS domain name. The UIDs and GIDs of the users and groups for the specified DOMAIN are read from the UNIX attributes that are populated in the RFC2307 schema extension of AD server. Any users or groups, from this domain, with missing UID/GID attributes are denied access. Use the L-H format to specify the ID range. All the users or groups from DOMAIN that need access to exports need to have their UIDs or GIDs in the specified range.

The specified range should not intersect with:

- The range specified by using the --idmap-range option of the command .
- The range specified for other AD DOMAIN for which ID mapping needs to be done from Active Directory (RFC2307 schema attributes) specified in --unixmap-domains option.
- The range specified for other AD DOMAIN for which ID mapping needs to be done from LDAP server specified in the --ldapmap-domains option.

The command reports a failure if you attempt to run the command with such configurations. This is intended to avoid ID collisions among users and groups from different domains.

For example,

```
--unixmap-domains "MYDOMAIN1(20000-50000);MYDOMAIN2(100000-200000)"
```

### **--ldapmap-domains ldapDomainMap**

Specifies the AD domains for which user ID and group ID should be fetched from a separate standalone LDAP server. This option is only valid with --type {ad} and --data-access-method {file}. ldapDomainMap takes value of the format as follows,

```
DOMAIN1 (type=stand-alone:ldap_srv=ldapServer:range=Range:usr_dn=userDN:grp_dn=groupDN:[bind_dn=bindDN]
:[bind_dn_pwd=bindDNpassword]) [;DOMAIN2(type=stand-alone:ldap_srv=ldapServer:range=Range:usr_dn=userDN
:grp_dn=groupDN:[bind_dn=bindDN] :[bind_dn_pwd=bindDNpassword]) [;DOMAIN3(type=stand-alone:ldap_srv=ldapServer
:range=Range:usr_dn=userDN:grp_dn=groupDN:[bind_dn=bindDN] :[bind_dn_pwd=bindDNpassword])...]]
```

#### **DOMAIN**

Use DOMAIN to specify an AD domain for which ID mapping services are to be configured. The name of the domain to be specified must be the Pre-Win2K domain name. The UID and GID of the users and groups for the specified DOMAIN are read from the objects stored on LDAP server in RFC2307 schema attributes. Any users or groups, from this domain, with missing UID/GID attributes are denied access.

#### **type**

Defines the type of LDAP server to use.

Supported value: stand-alone.

#### **range**

Attribute takes value in the L-H format. Defines the user or group from DOMAIN that needs access to exports need to have their UIDs or GIDs in the specified range. The specified range should not intersect with,

- The range specified using --idmap-range option of the command
- The range specified for other AD DOMAIN for which ID mapping needs to be done from Active Directory (RFC2307 schema attributes) specified in --unixmap-domains option
- The range specified for other AD DOMAIN for which ID mapping needs to be done from LDAP server specified in --ldapmap-domains option

This is intended to avoid ID collisions among users and groups from different domains.

**ldap\_srv**

Defines the name or IP address of the LDAP server to fetch the UID or GID for of a user or group records in RFC2307 schema format. The user and group objects should be in RFC2307 schema format. Specifying only single LDAP server is supported.

**user\_dn**

Defines the bind tree on the LDAP server where user objects shall be found.

**grp\_dn**

Defines the bind tree on the LDAP server where the group objects shall be found.

**bind\_dn**

Optional attribute.

Defines the user DN that should be used for authentication against the LDAP server. If not specified anonymous, bind shall be performed against the LDAP server.

**bind\_dn\_pwd**

Optional attribute.

Defines the password of the user DN specified in bind\_dn to be used for authentication against the LDAP server. Must be specified when bind\_dn attribute is specified for binding with the LDAP server in the DOMAIN definition.

Password cannot contain these special characters such as semicolon (;) or colon (:).

For example,

```
--ldapmap-domains "MYDOMAIN1(type=stand-alone:range=10000-50000
:ldap_srv=myldapserver.mydomain.com :usr_dn=ou=People,dc=mydomain,dc=com
:grp_dn=ou=Groups,dc=mydomain,dc=com :bind_dn=cn=manager,dc=mydomain,dc=com
:bind_dn_pwd=MYPASSWORD);MYDOMAIN2(type=stand-alone :range=70000-100000
:ldap_srv=myldapserver.example.com:usr_dn=ou=People,dc=example,dc=com
:grp_dn=ou=Groups,dc=example,dc=com)"
```

**--enable-kerberos**

Specifies whether to enable Kerberized logins for users who are gaining access by using file access protocols.

This option is only valid with --type {ldap} and **--data-access-method {file}**.

This option is disabled by default.

**Note:** Ensure that the legitimate keytab file is placed in the /var/mmfs/tmp directory and is named as krb5.keytab on the node that the authentication method configuration command is to be run.

**--kerberos-server kerberosServer**

Specifies the Kerberos server. This option is only valid with --type {ldap} and **--data-access-method {file}**.

**--kerberos-realm kerberosRealm**

Indicates the Kerberos server authentication administrative domain. The realm name is usually the all-uppercase version of the domain name. This option is case sensitive.

**--user-name-attrb UserNameAttribute**

Specifies the attribute to be used to search for user name on authentication server.

If the **--data-access-method** is object, this option is only valid with --type {ldap|ad}.

If the **--data-access-method** is file, this option is only valid with --type {ldap}. With --type ldap, default value is cn and with --type ad, the default value is sAMAccountName.

**--user-id-attrb UserIDAttribute**

Specifies the attribute to be used to search for user ID on the authentication server.

If **--data-access-method** is object, this option is only valid with --type {ldap|ad}.

## **mmuserauth**

If --data-access-method is file, this option is only valid with --type {ldap}. For --type ldap, default value is uid and for --type ad the default value is CN.

### **--user-mail-attrib UserMailAttribute**

Specifies the attribute to be used to search for email on authentication server. If the --data-access-method is object, this option is only valid with --type {ldap|ad}. For --data-access-method file, this option is only valid with --type {ldap}. Default value is mail.

### **--user-filter userFilter**

Specifies the additional filter to be used to search for user in the authentication server. The filter must be specified in LDAP filter format. This option is only valid with --type {ldap|ad} and --data-access-method {object}. By default, no filter is used.

### **--ks-dns-name keystoneDnsName**

Specifies the DNS name for keystone service. The specified name must be resolved on all protocol nodes for proper functioning. This is optional with --data-access-method {object}. If the value is not specified for this parameter, the **mmuserauth service create** command uses the value that is used during the IBM Spectrum Scale system installation.

### **--ks-admin-user keystoneAdminName**

Specifies the Keystone server administrative user. This user must be a valid user on authentication server if --type {ldap|ad} is specified. In case of --type local, new user along with the password specified in --ks-admin-pwd is created, and admin role is assigned in Keystone. This option is mandatory with --data-access-method {object}.

For --type {ldap|ad}, do not specify user name in DN format for --ks-admin-user. The name must be the base or short name that is written against the specified user-id-attrib or user-name-attrib of user on the LDAP server.

### **--ks-admin-pwd keystoneAdminPwd**

Specifies the password of the Keystone administrative user. This option is mandatory and valid with --type {local} and --data-access-method {object}. To hide the password due to security reasons, call the command without this option and the command prompts to enter the password when the **mmuserauth service create** command is issued.

### **--enable-ks-ssl**

Specifies whether the SSL communication must be enabled with the Keystone service. It enables a secured way to access the Keystone service over the HTTPS protocol. The default communication option with the Keystone service is over HTTP protocol, which has security risks.

This option is only valid with --data-access-method {object}.

By default, this option is disabled.

With --type local | ad | ldap, ensure that the valid certificate files are placed in the /var/mmfs/tmp directory on the node that the command has to be run:

The SSL certificate at: /var/mmfs/tmp/ssl\_cert.pem

The private key at: /var/mmfs/tmp/ssl\_key.pem

The CA certificate at: /var/mmfs/tmp/ssl\_cacert.pem

With --type userdefined, ensure that the valid certificate files are placed in the /var/mmfs/tmp directory on the node that the command has to be run:

The CA certificate at: /var/mmfs/tmp/ssl\_cacert.pem

### **--ks-swift-user keystoneSwiftName**

Specifies the username to be used as swift user in proxy-server.conf. If AD or LDAP-based authentication is used, this user must be available in the AD or LDAP authentication server. If local authentication method is used, new user with this name is created in the local database. This option is only valid with --data-access-method {object}.

For --type {ldap|ad}, do not specify user name in DN format for --ks-swift-user. The name must be the base or short name that is written against the specified user-id-attrib or user-name-attrib of user on the LDAP server.

**--ks-swift-pwd** *keystoneSwiftPwd*

Specifies the password of the ks-swift-user. If AD or LDAP-based authentication is used, this must be set for ks-swift user in AD or LDAP server. If local authentication method is used, the ks-swift-user with this password is created in the local database. This option is only valid with --data-access-method {object}.

**--enable-ks-casigning**

Specifies whether to use a CA-signed certificate for PKI (signing).

This option is only valid with --data-access-method {object} and --type {ad|ldap|local}.

Valid certificate files must be placed in the /var/mmfs/tmp directory on the node that the command has to be run.

The signing certificate is at: /var/mmfs/tmp/signing\_cert.pem

The private key is at: /var/mmfs/tmp/signing\_key.pem

The CA certificate is at: /var/mmfs/tmp/signing\_cacert.pem

**--ks-ext-endpoint** *externaleEndpoint*

Specifies the endpoint URL of external keystone. Only API v3 and HTTP are supported. This option is only valid with --data-access-method {object} and --type {userdefined}

**--idmapdelete**

Specifies whether to delete the current ID maps (SID to UID/GID mappings) from the ID mapping databases for the file access method and user-role-project-domain mappings stored in local keystone database for object access method.

This option is only valid with the **mmuserauth service remove** command. Unless the --data-access-method parameter is specified on the command line, ID maps for file and object access protocols are erased by default. To delete the ID maps of a particular access protocol, explicitly specify the --data-access-method parameter on command line along with the valid access protocol name.

The authentication method configuration and ID maps cannot be deleted together. The authentication method configuration must be deleted before the ID maps.

**Note:** CAUTION: Deleting ID maps can lead to irrecoverable loss of access to data. Use this option with proper planning.

**-N|--nodes** {node-list|cesNodes}

Verifies the authentication configuration on each node. If the specified node is not protocol node, it is ignored. If protocol node is specified, then the system checks configuration on all protocol nodes. If you do not specify a node, the system checks the configuration of only the current node.

**-Y** Displays the command output in a parseable format with a colon (:) as a field delimiter. Each column is described by a header.

**Note:** Fields that have a colon (:) are encoded to prevent confusion. For the set of characters that might be encoded, see the command documentation of **mmclidecode**. Use the **mmclidecode** command to decode the field.

**-r|--rectify**

Rectifies the authentication configurations and missing SSL and TLS certificates.

**--server-reachability**

Without this flag, the **mmuserauth service check** command only validates whether the authentication

## **mmuserauth**

configuration files are consistent across the protocol nodes. Use this flag to ensure if the external authentication server is reachable by each protocol node.

Exit status

**0** Successful completion.

**nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **mmuserauth** command.

The node on which the command is issued must be able to run remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS file system* in *IBM Spectrum Scale: Administration Guide*.

## **Examples**

1. To configure Microsoft Active Directory (AD) based authentication with automatic ID mapping for file access, issue this command:

```
mmuserauth service create --type ad --data-access-method file --netbios-name
ess --user-name administrator --idmap-role master --servers myAdserver
--password Passw0rd --idmap-range-size 10000000 --idmap-range 10000000-299999999
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : AD
PARAMETERS VALUES

ENABLE_NFS_KERBEROS false
SERVERS myAdserver
USER_NAME administrator
NETBIOS_NAME ess
IDMAP_ROLE master
IDMAP_RANGE 10000000-299999999
IDMAP_RANGE_SIZE 1000000
UNIXMAP_DOMAINS none

OBJECT access not configured
PARAMETERS VALUES

```

2. To configure Microsoft Active Directory (AD) based authentication with RFC2307 ID mapping for file access, issue this command:

```
mmuserauth service create --type ad --data-access-method file
--netbios-name ess --user-name administrator --idmap-role master
--servers myAdserver --password Passw0rd --idmap-range-size 1000000
--idmap-range 10000000-299999999 --unixmap-domains 'DOMAIN(5000-20000)'
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : AD
PARAMETERS VALUES

ENABLE_NFS_KERBEROS false
SERVERS myAdserver
USER_NAME administrator
NETBIOS_NAME ess
IDMAP_ROLE master
IDMAP_RANGE 10000000-299999999
IDMAP_RANGE_SIZE 1000000
UNIXMAP_DOMAINS DOMAIN(5000-20000)

OBJECT access not configured
PARAMETERS VALUES
```

3. To configure Microsoft Active Directory (AD) based authentication with LDAP ID mapping for file access, issue this command:

```
mmuserauth service create --data-access-method file --type ad --servers myADserver
--user-name administrator --password Passw0rd --netbios-name ess --idmap-role master
--ldapmap-domains "SONAS(type=stand-alone: range=1000-10000:ldap_srv=9.118.46.17:
usr_dn=ou=People,dc=example,dc=com:grp_dn=ou=Groups,dc=example,
dc=com:bind_dn=cn=manager,dc=example,dc=com:bind_dn_pwd=password)"
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : AD
PARAMETERS VALUES

ENABLE_NFS_KERBEROS false
SERVERS myADserver
USER_NAME administrator
NETBIOS_NAME ess
IDMAP_ROLE master
IDMAP_RANGE 10000000-299999999
IDMAP_RANGE_SIZE 1000000
UNIXMAP_DOMAINS none
LDAPMAP_DOMAINS DOMAIN(type=stand-alone: range=1000-10000:
ldap_srv=myLDAPserver:usr_dn=ou=People,dc=example,dc=com:grp_dn=
ou=Groups,dc=example,dc=com:bind_dn=cn=manager,dc=example,dc=com)
```

```
OBJECT access not configured
PARAMETERS VALUES
```

4. To configure Microsoft Active Directory (AD) based authentication with LDAP ID mapping for file access (anonymous binding with LDAP), issue this command:

```
mmuserauth service create --data-access-method file --type ad
--servers myADserver --user-name administrator --password Passw0rd
--netbios-name ess --idmap-role master --ldapmap-domains
"SONAS(type=stand-alone: range=1000-10000:ldap_srv=9.118.46.17:
usr_dn=ou=People,dc=example,dc=com:grp_dn=ou=Groups,dc=example,dc=com)"
```

## **mmuserauth**

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : AD
PARAMETERS VALUES

ENABLE_NFS_KERBEROS false
SERVERS myADserver
USER_NAME administrator
NETBIOS_NAME ess
IDMAP_ROLE master
IDMAP_RANGE 10000000-299999999
IDMAP_RANGE_SIZE 1000000
UNIXMAP_DOMAINS none
LDAPMAP_DOMAINS DOMAIN(type=stand-alone: range=1000-10000:ldap_srv=myLDAPserver:
 usr_dn=ou=People,dc=example,dc=com:grp_dn=ou=Groups,dc=example,dc=com)

OBJECT access not configured
PARAMETERS VALUES

```

5. To configure LDAP-based authentication with TLS encryption for file access, issue this command:

```
mmuserauth service create --type ldap --data-access-method file
--servers es-pune-host-01 --base-dn dc=example,dc=com
--user-name cn=manager,dc=example,dc=com --password secret
--netbios-name ess --enable-server-tls
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

**Note:** Before issuing the **mmuserauth service create** command to configure LDAP with TLS, ensure that the CA certificate for LDAP server is placed under /var/mmfs/tmp directory with the name "ldap\_cacert.pem" specifically on the protocol node where the command is issued.

To verify the authentication configuration, use the **mmuserauth service list** as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : LDAP
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS true
ENABLE_KERBEROS false
USER_NAME cn=manager,dc=example,dc=com
SERVERS myLDAPserver
NETBIOS_NAME ess
BASE_DN dc=example,dc=com
USER_DN none
GROUP_DN none
NETGROUP_DN none
USER_OBJECTCLASS posixAccount
GROUP_OBJECTCLASS posixGroup
USER_NAME_ATTRIB cn
USER_ID_ATTRIB uid
KERBEROS_SERVER none
KERBEROS_REALM none
```

```
OBJECT access not configured
PARAMETERS VALUES

```

- To configure LDAP-based authentication with Kerberos for file access, issue this command:

```
mmuserauth service create --type ldap --data-access-method file
--servers myLDAPserver --base-dn dc=example,dc=com
--user-name cn=manager,dc=example,dc=com --password secret
--netbios-name ess --enable-kerberos
--kerberos-server myKerberosServer --kerberos-realm example.com
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

**Note:** Before issuing the **mmuserauth service create** command to configure LDAP with Kerberos, ensure that the keytab file is also placed under /var/mmfs/tmp directory name as "krb5.keytab" specifically on the node where the command is run.

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : LDAP
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS false
ENABLE_KERBEROS true
USER_NAME cn=manager,dc=example,dc=com
SERVERS myLDAPserver
NETBIOS_NAME ess
BASE_DN dc=example,dc=com
USER_DN none
GROUP_DN none
NETGROUP_DN none
USER_OBJECTCLASS posixAccount
GROUP_OBJECTCLASS posixGroup
USER_NAME_ATTRIB cn
USER_ID_ATTRIB uid
KERBEROS_SERVER myKerberosServer
KERBEROS_REALM example.com
```

```
OBJECT access not configured
PARAMETERS VALUES

```

- To configure LDAP with TLS and Kerberos for file access, issue this command:

```
mmuserauth service create --type ldap --data-access-method file
--servers myLDAPserver --base-dn dc=example,dc=com
--user-name cn=manager,dc=example,dc=com --password secret
--netbios-name ess --enable-server-tls --enable-kerberos
--kerberos-server myKerberosServer --kerberos-realm example.com
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

## **mmuserauth**

```
FILE access configuration : LDAP
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS true
ENABLE_KERBEROS true
USER_NAME cn=manager,dc=example,dc=com
SERVERS es-pune-host-01
NETBIOS_NAME ess
BASE_DN dc=example,dc=com
USER_DN none
GROUP_DN none
NETGROUP_DN none
USER_OBJECTCLASS posixAccount
GROUP_OBJECTCLASS posixGroup
USER_NAME_ATTRIB cn
USER_ID_ATTRIB uid
KERBEROS_SERVER myKerberosServer
KERBEROS_REALM example.com

OBJECT access not configured
PARAMETERS VALUES

```

8. To configure LDAP without TLS and Kerberos for file access, issue this command:

```
mmuserauth service create --type ldap --data-access-method file
--servers myLDAPserver --base-dn dc=example,dc=com --user-name
cn=manager,dc=example,dc=com --password secret --netbios-name ess
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : LDAP
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS false
ENABLE_KERBEROS false
USER_NAME cn=manager,dc=example,dc=com
SERVERS myLDAPserver
NETBIOS_NAME ess
BASE_DN dc=example,dc=com
USER_DN none
GROUP_DN none
NETGROUP_DN none
USER_OBJECTCLASS posixAccount
GROUP_OBJECTCLASS posixGroup
USER_NAME_ATTRIB cn
USER_ID_ATTRIB uid
KERBEROS_SERVER none
KERBEROS_REALM none

OBJECT access not configured
PARAMETERS VALUES

```

9. To configure NIS-based authentication for file access, issue this command:

```
mmuserauth service create --type nis --data-access-method file
--servers myNISserver --domain nisdomain
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : NIS
PARAMETERS VALUES

SERVERS myNISserver
DOMAIN nisdomain

OBJECT access not configured
PARAMETERS VALUES

```

10. To configure user-defined authentication for file access, issue this command:

```
mmuserauth service create --data-access-method file --type userdefined
```

The system displays output similar to this:

```
File Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access configuration : USERDEFINED
PARAMETERS VALUES

OBJECT access not configured
PARAMETERS VALUES

```

11. To configure local authentication for object access, issue this command:

```
mmuserauth service create --data-access-method object --type local
--ks-dns-name ksDNSname --ks-admin-user admin
--ks-admin-pwd Passw0rd
```

The system displays output similar to this:

```
Object configuration with local (Database) as identity backend is completed
successfully.
Object Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access not configured
PARAMETERS VALUES

OBJECT access configuration : LOCAL
PARAMETERS VALUES

```

## **mmuserauth**

```

ENABLE_KS_SSL false
ENABLE_KS_CASIGNING false
KS_ADMIN_USER admin
```

12. To configure AD without TLS authentication for object access, issue this command:

```
mmuserauth service create --type ad --data-access-method object
--user-name "cn=Administrator,cn=Users,dc=example,dc=com" --password Passw0rd --base-dn "dc=example,DC=com"
--ks-dns-name ksDNSname --ks-admin-user admin --servers myADserver --user-id-attrib cn
--user-name-attrib sAMAccountName --user-objectclass organizationalPerson --user-dn "cn=Users,dc=example,dc=com"
--ks-swift-user swift --ks-swift-pwd Passw0rd
```

The system displays output similar to this:

```
Object configuration with LDAP (Active Directory) as identity backend is completed
successfully.
```

```
Object Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```
FILE access not configured
PARAMETERS VALUES

OBJECT access configuration: AD
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS false
ENABLE_KS_SSL false
USER_NAME cn=Administrator,cn=Users,dc=example,dc=com
SERVERS myADserver
BASE_DN dc=IBM,DC=local
USER_DN cn=users,dc=example,dc=com
USER_OBJECTCLASS organizationalPerson
USER_NAME_ATTRIB sAMAccountName
USER_ID_ATTRIB cn
USER_MAIL_ATTRIB mail
USER_FILTER none
ENABLE_KS_CASIGNING false
KS_ADMIN_USER admin
```

13. To configure AD with TLS authentication for object access, issue this command:

```
mmuserauth service create --type ad --data-access-method object
--user-name "cn=Administrator,cn=Users,dc=example,dc=com" --password Passw0rd --base-dn
"dc=example,DC=com" --enable-server-tls --ks-dns-name ksDNSname --ks-admin-user admin --servers
myADserver --user-id-attrib cn --user-name-attrib sAMAccountName --user-objectclass organizationalPerson
--user-dn "cn=Users,dc=example,dc=com" --ks-swift-user swift --ks-swift-pwd Passw0rd
```

The system displays output similar to this:

```
Object configuration with LDAP (Active Directory) as identity backend is completed
successfully.
```

```
Object Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```

FILE access not configured
PARAMETERS VALUES

OBJECT access configuration: AD
PARAMETERS VALUES

ENABLE_ANONYMOUS_BIND false
ENABLE_SERVER_TLS true
ENABLE_KS_SSL false
USER_NAME cn=Administrator,cn=Users,dc=example,dc=com
SERVERS myADserver
BASE_DN dc=IBM,DC=local
USER_DN cn=users,dc=example,dc=com
USER_OBJECTCLASS organizationalPerson
USER_NAME_ATTRIB SAMAccountName
USER_ID_ATTRIB cn
USER_MAIL_ATTRIB mail
USER_FILTER none
ENABLE_KS_CASIGNING false
KS_ADMIN_USER admin

```

14. To configure LDAP-based authentication for object access, issue this command:

```
mmuserauth service create --type ldap --data-access-method object
--user-name "cn=manager,dc=example,dc=com" --password "Passw0rd"
--base-dn dc=example,dc=com --ks-dns-name ksDNSname --ks-admin-user admin
--servers myLDAPserver --user-dn "ou=People,dc=example,dc=com"
--ks-swift-user swift --ks-swift-pwd Passw0rd
```

The system displays output similar to this:

```
Object configuration with LDAP as identity backend is completed successfully.
Object Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

```

FILE access not configured
PARAMETERS VALUES

OBJECT access configuration : LDAP
PARAMETERS VALUES

ENABLE_SERVER_TLS false
ENABLE_KS_SSL false
USER_NAME cn=manager,dc=example,dc=com
SERVERS myLDAPserver
BASE_DN dc=example,dc=com
USER_DN ou=people,dc=example,dc=com
USER_OBJECTCLASS posixAccount
USER_NAME_ATTRIB cn
USER_ID_ATTRIB uid
USER_MAIL_ATTRIB mail
USER_FILTER none
ENABLE_KS_CASIGNING false
KS_ADMIN_USER admin

```

15. To configure LDAP with TLS-based authentication for object access, issue this command:

```
mmuserauth service create --type ldap --data-access-method object
--user-name "cn=manager,dc=example,dc=com" --password "Passw0rd"
--base-dn dc=example,dc=com --enable-server-tls
--ks-dns-name ksDNSname --ks-admin-user admin --servers myLDAPserver
--user-dn "ou=People,dc=example,dc=com" --ks-swift-user swift
--ks-swift-pwd Passw0rd
```

## **mmuserauth**

The system displays output similar to this:

```
Object configuration with LDAP as identity backend is completed successfully.
Object Authentication configuration completed successfully.
```

To verify the authentication configuration, use the **mmuserauth service list** command as shown in the following example:

```
mmuserauth service list
```

The system displays output similar to this:

| FILE access not configured         |                              |
|------------------------------------|------------------------------|
| PARAMETERS                         | VALUES                       |
| OBJECT access configuration : LDAP |                              |
| PARAMETERS                         | VALUES                       |
| ENABLE_SERVER_TLS                  | true                         |
| ENABLE_KS_SSL                      | false                        |
| USER_NAME                          | cn=manager,dc=example,dc=com |
| SERVERS                            | myLDAPserver                 |
| BASE_DN                            | dc=example,dc=com            |
| USER_DN                            | ou=people,dc=example,dc=com  |
| USER_OBJECTCLASS                   | posixAccount                 |
| USER_NAME_ATTRIB                   | cn                           |
| USER_ID_ATTRIB                     | uid                          |
| USER_MAIL_ATTRIB                   | mail                         |
| USER_FILTER                        | none                         |
| ENABLE_KS_CASIGNING                | false                        |
| KS_ADMIN_USER                      | admin                        |

16. To remove the authentication method that is configured for file access, issue this command:

```
mmuserauth service remove --data-access-method file
```

The system displays output similar to this:

```
mmuserauth service remove: Command successfully completed
```

**Note:** Authentication configuration and ID maps cannot be deleted together. To remove ID maps, remove the authentication configuration first and then remove ID maps. Also, you cannot delete ID maps that are used for file and object access together. That is, when you delete the ID maps, the value that is specified for **--data-access-method** must be either **file** or **object**.

17. To remove the authentication method that is configured for object access, issue this command:

```
mmuserauth service remove --data-access-method object
```

The system displays output similar to this:

```
mmuserauth service remove: Command successfully completed
```

**Note:** Authentication configuration and ID maps cannot be deleted together. To remove ID maps, remove the authentication configuration first and then remove the ID maps. Also, you cannot delete ID maps that are used for file and object access together. That is, when you delete the ID maps, the value that is specified for **--data-access-method** must be either **file** or **object**.

18. To check whether the authentication configuration is consistent across the cluster and the required services are enabled and running, issue this command:

```
mmuserauth service check --data-access-method file --nodes cesNodes --rectify
```

The system displays output similar to this:

```
Userauth file check on node: dgnode3
Checking SSSD_CONF: OK
LDAP server status
```

```
LDAP server 192.0.2.18 : OK
Service 'sssd' status: OK
Userauth file check on node: dgnode2
dgnode2: not CES node. Ignoring...
```

19. To check whether the file authentication configuration is consistent across the cluster and the required services are enabled and running, and if you do not want to correct the situation, issue this command:
 

```
mmuserauth service check --data-access-method file --nodes cesNodes --rectify
```
20. To check that all object configuration files (including certificates) are present, and if not, rectify the situation by issuing the following command:
 

```
mmuserauth service check --data-access-method object --rectify
```

The system displays output similar to this:

```
Userauth object check on node: node1
Checking keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
Checking /etc/keystone/ssl/certs/object_ldap_cacert.pem: OK
Service 'openstack-keystone' status: OK
```

21. To check if the external authentication server is reachable by each protocol node, issue the following command:

```
mmuserauth service check --server-reachability
```

- a. If file is not configured, object is configured, and there are no errors, the system displays output similar to this:

```
Userauth object check on node: vmnode2
Checking keystone.conf: OK
Checking wsgi-keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
```

```
LDAP servers status
LDAP server 9.118.37.234 : OK
Service 'httpd' status: OK
```

- b. If file is not configured, object is configured, and there is a single error, the system displays output similar to this:

```
Userauth object check on node: vmnode2
Checking keystone.conf: OK
Checking wsgi-keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
```

```
LDAP servers status
LDAP server 9.118.37.234 : ERROR
Service 'httpd' status: OK
```

- c. If file and object are configured and there are no errors, the system displays output similar to this:

```
Userauth file check on node: vmnode2
Checking nsswitch file: OK
```

```
AD servers status
NETLOGON connection: OK
Domain join status: OK
Machine password status: OK
Service 'gpfs-winbind' status: OK
```

```
Userauth object check on node: vmnode2
Checking keystone.conf: OK
```

## **mmuserauth**

```
Checking wsgi-keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
```

```
LDAP servers status
LDAP server 9.118.37.234 : OK
Service 'httpd' status: OK
```

- d. If file and object are configured and there is a single error, the system displays output similar to this:

```
Userauth file check on node: vmnode2
Checking nsswitch file: OK
```

```
AD servers status
NETLOGON connection: OK
Domain join status: OK
Machine password status: ERROR
Service 'gpfs-winbind' status: OK
```

```
Userauth object check on node: vmnode2
```

```
Checking keystone.conf: OK
Checking wsgi-keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
```

```
LDAP servers status
LDAP server 9.118.37.234 : OK
```

- e. If file and object are configured and there are multiple errors, the system displays output similar to this:

```
Userauth file check on node: vmnode2
Checking nsswitch file: OK
```

```
AD servers status
NETLOGON connection: ERROR
Domain join status: ERROR
Machine password status: ERROR
Service 'gpfs-winbind' status: OK
```

```
Userauth object check on node: vmnode2
```

```
Checking keystone.conf: OK
Checking wsgi-keystone.conf: OK
Checking /etc/keystone/ssl/certs/signing_cert.pem: OK
Checking /etc/keystone/ssl/private/signing_key.pem: OK
Checking /etc/keystone/ssl/certs/signing_cacert.pem: OK
```

```
LDAP servers status
LDAP server 9.118.37.234 : ERROR
Service 'httpd' status: OK
```

**Note:** The --rectify or -r option cannot fix server reachability errors. Specifying that option with --server-reachability may fix the erroneous config files and service-related errors only.

## **See also**

- “mmces command” on page 103
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454

- “[mmobj command](#)” on page 466
- “[mmsmb command](#)” on page 554

**Location**

/usr/lpp/mmfs/bin

## **mmwinserctl**

---

### **mmwinserctl command**

Manages the **mmwinser** Windows service.

#### **Synopsis**

```
mmwinserctl set [--account AccountName [--password Password]] [--remote-shell {yes | no}]
[-N {Node[,Node...] | NodeFile | NodeClass}] [-v]
```

or

```
mmwinserctl {enable | disable | query} [-N {Node[,Node...] | NodeFile | NodeClass}] [-v]
```

#### **Availability**

Available on all IBM Spectrum Scale editions. Available on Windows.

#### **Description**

**mmwinser** is a GPFS for Windows service that is needed for the proper functioning of the GPFS daemon on nodes running Windows. Optionally, the service can be configured to provide a remote execution facility for GPFS administration commands.

Use the **mmwinserctl** command to manage the **mmwinser** service. You can set the log on account and password for the service, enable or disable the service, enable or disable the service's remote execution facility, or query its current state.

The **mmwinserctl** command must be run on a Windows node and it has no effect on nodes running other operating systems.

If the remote execution facility of **mmwinser** is enabled, a Windows GPFS cluster can be configured to use **mmwinrsh** and **mmwinrcp** as the remote shell and remote file copy commands:

- **mmwinrsh** (`/usr/lpp/mmfs/bin/mmwinrsh`) uses Windows Named Pipes to pass the command to the target node.
- **mmwinrcp** (`/usr/lpp/mmfs/bin/mmwinrcp`) is a wrapper module that invokes the Cygwin **cp** command to copy the files that are needed by the **mm** commands. The path names on remote hosts are translated into path names based on the standard Windows ADMIN\$ share.

An account must be given the right to log on as a service before it can be used to run **mmwinser**. The right to log on as a service is controlled by the Local Security Policy of each Windows node. You can use the Domain Group Policy to set the Local Security Policy on all Windows nodes in a GPFS cluster.

For more information on the **mmwinser** service, see *Configuring the GPFS Administration service* in the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

#### **Parameters**

##### **set**

Sets the service configuration options and restarts the service if it is running.

##### **enable**

Sets the service to automatic startup and starts the service.

##### **disable**

Sets the service to disabled and stops the service.

##### **query**

Returns information about the service's configuration and current state.

**--account AccountName**

Specifies the log on account name for the **mmwinser** service. By default, **mmwinser** is configured to run using the **LocalSystem** account.

**--password Password**

Specifies the log on password for the **mmwinser** service.

**--remote-shell {yes | no}**

Specifies whether or not remote connections are allowed.

**-N {Node[,Node...] | NodeFile | NodeClass}**

Specifies the list of nodes on which to perform the action. The default is the node on which the **mmwinserctl** command is issued.

If the node on which the **mmwinserctl** command is issued belongs to a GPFS cluster, the nodes specified with the **-N** parameter must belong to the cluster.

If the node on which the **mmwinserctl** command is issued does not belong to a GPFS cluster, the nodes specified with the **-N** parameter must be identified by their host names or IP addresses. Node classes and node numbers cannot be used.

For general information on how to specify node names, see *Specifying nodes as input to GPFS commands* in the *IBM Spectrum Scale: Administration Guide*.

**-v** Displays progress and intermediate error messages.

**Exit status**

**0** Successful completion.

**nonzero**

A failure has occurred.

**Security**

You must be a member of the Domain Admins group to run the **mmwinserctl** command.

**Examples**

1. To specify 'gpfs\root' as the log on account name for the **mmwinser** service and enable the remote command execution facility on nodes **ls21n19** and **ls21n20**, issue:

```
mmwinserctl set -N ls21n19,ls21n20 --account gpfs\root -password abcdefg -remote-shell yes
```

The system displays information similar to:

| Node name | Service state | Remote shell | Account name |
|-----------|---------------|--------------|--------------|
| ls21n19   | START_PENDING | yes          | gpfs\root    |
| ls21n20   | START_PENDING | yes          | gpfs\root    |

2. To display the current state of the **mmwinser** service on all nodes in the cluster, issue:

```
mmwinserctl query -N all
```

The system displays information similar to:

| Node name | Service state | Remote shell | Account name |
|-----------|---------------|--------------|--------------|
| ls21n19   | RUNNING       | yes          | gpfs\root    |
| ls21n20   | RUNNING       | yes          | gpfs\root    |
| ls21n14   | RUNNING       | yes          | LocalSystem  |

**Location**

/usr/lpp/mmfs/bin

## spectrumscale

---

### spectrumscale command

Installs and configures GPFS; adds nodes to a cluster; deploys and configures protocols, performance monitoring tools, and authentication services; and upgrades GPFS and protocols.

#### Synopsis

```
spectrumscale setup [-i SSHIdentity] [-s ServerIP] [--storesecret]
```

or

```
spectrumscale node add [-g] [-q] [-m] [-a] [-n] [-p [ExportIP]] Node
```

or

```
spectrumscale node load [-q] [-m] [-a] [-n] NodeFile
```

or

```
spectrumscale node delete [-f] Node
```

or

```
spectrumscale node clear [-f]
```

or

```
spectrumscale node list
```

or

```
spectrumscale config gpfs [-l] [-c ClusterName] [-p {default | randomio}]
[-r RemoteShell] [-rc RemoteFileCopy]
[-e EphemeralPortRange]
```

or

```
spectrumscale config protocols [-l] [-f FileSystem] [-m MountPoint] [-e ExportIPPool]
```

or

```
spectrumscale config object [-f FileSystem] [-m MountPoint] [-e EndPoint] [-o ObjectBase]
[-i InodeAllocation] [-t AdminToken]
[-au AdminUser] [-ap AdminPassword]
[-su SwiftUser] [-sp SwiftPassword]
[-dp DatabasePassword]
[-mr MultiRegion] [-rn RegionNumber]
[-s3 {on | off}]
```

or

```
spectrumscale config perfmon [-r {on | off}] [-d {on | off}] [-l]
```

or

```
spectrumscale config ntp [-e {on | off} [-l List] [-s Upstream_Servers]]
```

or

```
spectrumscale config clear {gpfs | protocols | object}
```

or

```
spectrumscale config update
```

or

```
| spectrumscale config populate --node Node
```

or

```
spectrumscale nsd add -p Primary [-s Secondary] [-fs FileSystem]
[-po Pool]
[-u {dataOnly | dataAndMetadata | metaDataOnly | descOnly | localCache}]
[-fg FailureGroup] [--no-check]
PrimaryDevice [PrimaryDevice ...]
```

or

```
spectrumscale nsd balance [--node Node | --all]
```

or

```
spectrumscale nsd delete NSD
```

or

```
spectrumscale nsd modify [-n Name]
[-u {dataOnly | dataAndMetadata | metaDataOnly | descOnly}]
[-po Pool] [-fs FileSystem] [-fg FailureGroup]
NSD
```

or

```
spectrumscale nsd servers
```

or

```
spectrumscale nsd clear [-f]
```

or

```
spectrumscale nsd list
```

or

```
spectrumscale filesystem modify [-b {64K | 128K | 256K | 512K | 1M | 2M | 4M | 8M | 16M}]
[-m MountPoint]
FileSystem
```

or

```
spectrumscale filesystem list
```

or

```
spectrumscale auth file {ldap | ad | nis | none}
```

or

```
spectrumscale auth object [--https] [--pki] {local | external | ldap | ad}
```

or

```
spectrumscale auth commitsettings
```

or

```
spectrumscale auth clear
```

or

```
spectrumscale enable {object | nfs | smb}
```

or

```
spectrumscale disable {object | nfs | smb}
```

## **spectrumscale**

### **CAUTION:**

Disabling object service discards OpenStack Swift configuration and ring files from the CES cluster. If OpenStack Keystone configuration is configured locally, disabling object storage also discards the Keystone configuration and database files from the CES cluster. However, the data is not removed. For subsequent object service enablement with a clean configuration and new data, remove object store fileset and set up object environment. See the `mmobj swift base` command. For more information, contact the IBM Support Center.

or

```
spectrumscale install [-pr] [-po] [-s] [-f]
```

or

```
spectrumscale deploy [-pr] [-po] [-s] [-f]
```

or

```
spectrumscale upgrade [-pr | -po | -ve] [-f]
```

or

```
spectrumscale installgui {start | stop | status}
```

## **Availability**

The **spectrumscale** is available as follows:

- Available with IBM Spectrum Scale Standard Edition or higher.

## **Description**

Use the **spectrumscale** command (also called the installation toolkit) to do the following:

- Install and configure GPFS.
- Add GPFS nodes to an existing cluster.
- Deploy and configure SMB, NFS, OpenStack Swift, and performance monitoring tools on top of GPFS.
- Configure authentication services for protocols.
- Upgrade GPFS and protocols.

**Note:** The following prerequisites and assumptions apply:

- The installation toolkit requires the following package:
  - python-2.7
- TCP traffic from the nodes should be allowed through the firewall to communicate with the install toolkit on port 8889 for communication with the chef zero server and port 10080 for package distribution.
- The nodes themselves have external Internet access or local repository replicas that can be reached by the nodes to install necessary packages (dependency installation). For more information, see the *Repository setup* section of the *Installation prerequisites* topic in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.
- To install protocols, there must be a GPFS cluster running a minimum version of 4.1.1.0 with CCR enabled.
- The node that you plan to run the install toolkit from must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages.

## Parameters

### **setup**

Installs Chef and its components, as well as configure the install node in the cluster definition file. The IP address passed in should be the node from which the installation toolkit will be run. The SSH key passed in should be the key the installer should use to have passwordless SSH onto all other nodes. This is the first command you will run to set up IBM Spectrum Scale. This option accepts the following arguments:

**-i *SSHIdentity***

Adds the path to the SSH identity file into the configuration.

**-s *ServerIP***

Adds the control node IP into the configuration.

**--storesecret**

Disables the prompts for the encryption secret.

**CAUTION:**

If you use this option, passwords will not be securely stored.

This is the first command to run to set up IBM Spectrum Scale.

### **node**

Used to add, remove, or list nodes in the cluster definition file. This command only interacts with this configuration file and does not directly configure nodes in the cluster itself. The nodes that have an entry in the cluster definition file will be used during install, deploy, or upgrade. This option accepts the following arguments:

**add *Node***

Adds the specified node and configures it according to the following arguments:

**-g** Adds GPFS Graphical User Interface servers to the cluster definition file.

**-q** Configures the node as a quorum node.

**-m** Configures the node as a manager node.

**-a** Configures the node as an admin node.

**-n** Specifies the node as NSD.

**-p *[ExportIP]***

Configures the node as a protocol node and optionally assigns it an IP.

*Node*

Specifies the node name.

**load *NodeFile***

Loads the specified file containing a list of nodes, separated per line; adds the nodes in the file and configures them according to the following:

**-q** Configures the node as a quorum node.

**-m** Configures the node as a manager node.

**-a** Configures the node as an admin node.

**-n** Specifies the node to be NSD.

**delete *Node***

Removes the specified node from the configuration. The following option is accepted.

**-f** Forces the action without manual confirmation.

## spectrumscale

### clear

Clears the current node configuration. The following option is accepted:

- f Forces the action without manual confirmation.

### list

Lists the nodes configured in your environment.

### config

Used to set properties in the cluster definition file that will be used during install, deploy, or upgrade. This command only interacts with this configuration file and does not directly configure these properties on the GPFS cluster. This option accepts the following arguments:

#### gpfs

Sets any of the following GPFS-specific properties to be used during GPFS installation and configuration:

- l Lists the current settings in the configuration.

**-c ClusterName**

**-p**

Specifies the profile to be set on cluster creation. The following values are accepted:

**default**

Specifies that the **GpfsProtocolDefaults** profile is to be used.

**randomio**

Specifies that the **GpfsProtocolRandomIO** profile is to be used.

**-r RemoteShell**

Specifies the remote shell binary to be used by GPFS. If no remote shell is specified in the cluster definition file, /usr/bin/ssh will be used as the default.

**-rc RemoteFileCopy**

Specifies the remote file copy binary to be used by GPFS. If no remote file copy binary is specified in the cluster definition file, /usr/bin/scp will be used as the default.

**-e EphemeralPortRange**

Specifies an ephemeral port range to be set on all GPFS nodes. If no port range is specified in the cluster definition, 60000-61000 will be used as default.

For information about ephemeral port range, see the topic about GPFS port usage in the *Miscellaneous advanced administration tasks in IBM Spectrum Scale: Administration Guide*.

### protocols

Provides details of the GPFS environment that will be used during protocol deployment, according to the following options:

- l Lists the current settings in the configuration.

**-f FileSystem**

Specifies the file system.

**-m MountPoint**

Specifies the mount point.

**-e ExportIPPool**

Specifies a comma-separated list of additional CES export IPs to configure on the cluster.

### object

Sets any of the following Object-specific properties to be used during Object deployment and configuration:

- l Lists the current settings in the configuration.

- f *FileSystem*  
Specifies the file system.
- m *MountPoint*  
Specifies the mount point.
- e *EndPoint*  
Specifies the host name that will be used for access to the object store. This should be a round-robin DNS entry that maps to all CES IP addresses or the address of a load balancer front end; this will distribute the load of all keystone and object traffic that is routed to this host name. Therefore the endpoint is an IP address in a DNS or in a load balancer that maps to a group of export IPs (that is, CES IPs that were assigned on the protocol nodes).
- o *ObjectBase*  
Specifies the object base.
- i *InodeAllocation*  
Specifies the inode allocation.
- t *AdminToken*  
Specifies the admin token.
- au *AdminUser*  
Specifies the user name for the admin.
- ap *AdminPassword*  
Specifies the admin user password. This credential is for the Keystone administrator. This user can be local or on remote authentication server based on the authentication type used.  
  
**Note:** You will be prompted to enter a Secret Encryption Key which will be used to securely store the password. Choose a memorable pass phrase which you will be prompted for each time you enter the password.
- su *SwiftUser*  
Specifies the Swift user name. This credential is for the Swift services administrator. All Swift services are run in this user's context. This user can be local or on remote authentication server based on the authentication type used.
- sp *SwiftPassword*  
Specifies the Swift user password.  
  
**Note:** You will be prompted to enter a Secret Encryption Key which will be used to securely store the password. Choose a memorable pass phrase which you will be prompted for each time you enter the password.
- dp *DataBasePassword*  
Specifies the object database user password.  
  
**Note:** You will be prompted to enter a Secret Encryption Key which will be used to securely store the password. Choose a memorable pass phrase which you will be prompted for each time you enter the password.
- mr *MultiRegion*  
Enables the **multi-region** option.
- rn *RegionNumber*  
Specifies the region number.
- s3 **on | off**  
Specifies whether s3 is to be turned on or off.

## | **perfmon**

| Sets Performance Monitoring specific properties to be used during installation and configuration:

## spectrumscale

**-r on | off**

Specifies if the install toolkit can reconfigure Performance Monitoring.

**Note:** When set to on, reconfiguration might move the collector to different nodes and it might reset sensor data. Custom sensors and data might be erased.

**-d on | off**

Specifies if Performance Monitoring should be disabled (not installed).

**Note:** When set to on, pmcollector and pmsensor packages are not installed or upgraded. Existing sensor or collector state remains as is.

**-l** Lists the current settings in the configuration.

### ntp

Used to add, list, or remove NTP nodes to the configuration. NTP nodes will be configured on the cluster as follows: the admin node will point to the upstream NTP servers that you provide to determine the correct time. The rest of the nodes in the cluster will point to the admin node to obtain the time.

**-sUpstream\_Server**

Specifies the host name that will be used. You can use an upstream server that you have already configured, but it cannot be part of your Spectrum Scale cluster.

**Note:** NTP works best with at least four upstream servers. If you provide fewer than four, you will receive a warning during installation advising that you add more.

**-lList**

Lists the current settings of your NTP setup.

**-e on | off**

Specifies whether NTP is enabled or not. If this option is turned to off, you will receive a warning during installation.

### clear

Removes specified properties from the cluster definition file:

#### gpfs

Removes GPFS related properties from the cluster definition file:

**-c** Clears the GPFS cluster name.

**-p** Clears the GPFS profile to be applied on cluster creation. The following values are accepted:

#### default

Specifies that the **GpfProtocolDefaults** profile is to be cleared.

#### randomio

Specifies that the **GpfProtocolRandomIO** profile is to be cleared.

**-r RemoteShell**

Clears the absolute path name of the remote shell command GPFS uses for node communication. For example, /usr/bin/ssh.

**-rc RemoteFileCopy**

Clears the absolute path name of the remote copy command GPFS uses when transferring files between nodes. For example, /usr/bin/scp.

**-e EphemeralPortRange**

Clears the GPFS daemon communication port range.

**--all**

Clears all settings in the cluster definition file.

**protocols**

Removes protocols related properties from the cluster definition file:

**-f** Clears the shared file system name.

**-m** Clears the shared file system mount point.

**-e** Clears a comma-separated list of additional CES export IPs to configure on the cluster.

**--all**

Clears all settings in the cluster definition file.

**object**

Removes object related properties from the cluster definition file:

**-f** Clears the object file system name.

**-m** Clears the absolute path to your file system on which the objects reside.

**-e** Clears the host name which maps to all CES IP addresses in a round-robin manner.

**-o** Clears the GPFS fileset to be created or used as the object base.

**-i** Clears the GPFS fileset inode allocation to be used by the object base.

**-t** Clears the admin token to be used by Keystone.

**-au**

Clears the user name for the admin user.

**-ap**

Clears the password for the admin user.

**-su**

Clears the user name for the Swift user.

**-sp**

Clears the password for the Swift user.

**-dp**

Clears the password for the object database.

**-s3**

Clears the S3 API setting, if it is enabled.

**-mr**

Clears the multi-region data file path.

**-rn**

Clears the region number for the multi-region configuration.

**--all**

Clears all settings in the cluster definition file.

**update**

Updates operating system and CPU architecture fields in the cluster definition file. This update is automatically done if you run the upgrade precheck with the **spectrumscale upgrade --precheck** command while upgrading to IBM Spectrum Scale release 4.2.2 or later.

**populate**

Populates the cluster definition file with the current cluster state. In the following upgrade scenarios, you might need to update the cluster definition file with the current cluster state:

- A manually created cluster in which you want to use the installation toolkit to perform administration tasks on the cluster such as adding protocols, adding nodes, and upgrading.

## **spectrumscale**

- A cluster created using the installation toolkit in which manual changes were done without using the toolkit wherein you want to synchronize the installation toolkit with the updated cluster configuration that was performed manually.

### **--node Node**

Specifies an existing node in the cluster that is used to query the cluster information.

## **nsd**

Used to add, remove, list or balance NSDs, as well as add file systems in the cluster definition file. This command only interacts with this configuration file and does not directly configure NSDs on the cluster itself. The NSDs that have an entry in the cluster definition file will be used during install. This option accepts the following arguments:

### **add**

Adds an NSD to the configuration, according to the following specifications:

#### **-p Primary**

Specifies the primary NSD server name.

#### **-s Secondary**

Specifies the secondary NSD server names. You can use a comma-separated list to specify multiple secondary NSD servers.

#### **-fs FileSystem**

Specifies the file system to which the NSD is assigned.

#### **-po Pool**

Specifies the file system pool.

#### **-u**

Specifies NSD usage. The following values are accepted:

**dataOnly**

**dataAndMetadata**

**metaDataOnly**

**descOnly**

**localCache**

#### **-fg FailureGroup**

Specifies the failure group to which the NSD belongs.

#### **--no-check**

Specifies not to check for the device on the server.

#### **PrimaryDevice**

Specifies the device name on the primary NSD server.

## **balance**

Balances the NSD preferred node between the primary and secondary nodes. The following options are accepted:

### **--node Node**

Specifies the node to move NSDs from when balancing.

### **--all**

Specifies that all NSDs are to be balanced.

## **delete NSD**

Removes the specified NSD from the configuration.

## **modify NSD**

Modifies the NSD parameters on the specified NSD, according to the following options:

- n *Name***  
Specifies the name.
- u** The following values are accepted:
  - dataOnly**
  - dataAndMetadata**
  - metadataOnly**
  - descOnly**
- po *Pool***  
Specifies the pool
- fs *FileSystem***  
Specifies the file system.
- fg *FailureGroup***  
Specifies the failure group.

**servers**

Adds and removes servers, and sets the primary server for NSDs.

**clear**

Clears the current NSD configuration. The following option is accepted:

**-f** Forces the action without manual confirmation.

**list**

Lists the NSDs configured in your environment.

**filesystem**

Used to list or modify file systems in the cluster definition file. This command only interacts with this configuration file and does not directly modify file systems on the cluster itself. To modify the properties of a file system in the cluster definition file, the file system must first be added with **spectrumscale nsd**. This option accepts the following arguments:

**modify**

Modifies the file system attributes. This option accepts the following arguments:

- b** Specifies the file system block size. This argument accepts the following values: 64K, 128K, 256K, 512K, 1M, 2M, 4M, 8M, 16M.
- m *MountPoint***  
Specifies the mount point.

*FileSystem*

Specifies the file system to be modified.

**list**

Lists the file systems configured in your environment.

**auth**

Used to configure either Object or File authentication on protocols in the cluster definition file. This command only interacts with this configuration file and does not directly configure authentication on the protocols. To configure authentication on the GPFS cluster during a deploy, authentication settings must be provided through the use of a template file. This option accepts the following arguments:

**file**

Specifies file authentication.

One of the following must be specified:

**ldap**

## **spectrumscale**

**ad**

**nis**

**none**

### **object**

Specifies object authentication.

Either of the following options are accepted:

**--https**

**--pki**

One of the following must be specified:

**local**

**external**

**ldap**

**ad**

Both file and object authentication can be set up with the authentication backend server specified.

Running this command will open a template settings file to be filled out before installation.

### **commitsettings**

Merges authentication settings into the main cluster definition file.

### **clear**

Clears your current authentication configuration.

### **enable**

Used to enable Object, SMB or NFS in the cluster definition file. This command only interacts with this configuration file and does not directly enable any protocols on the GPFS cluster itself. The default configuration is that all protocols are disabled. If a protocol is enabled in the cluster definition file, this protocol will be enabled on the GPFS cluster during deploy. This option accepts the following arguments:

**obj**

Object

**nfs**

NFS

**smb**

SMB

### **disable**

Used to disable Object, SMB or NFS in the cluster definition file. This command only interacts with this configuration file and does not directly disable any protocols on the GPFS cluster itself. The default configuration is that all protocols are disabled, so this command is only necessary if a protocol has previously been enabled in the cluster definition file, but is no longer required.

**Note:** Disabling a protocol in the cluster definition will not disable this protocol on the GPFS cluster during a deploy, it merely means that this protocol will not be enabled during a deploy.

This option accepts the following arguments:

**obj**

Object

**CAUTION:**

Disabling object service discards OpenStack Swift configuration and ring files from the CES cluster. If OpenStack Keystone configuration is configured locally, disabling object storage also discards the Keystone configuration and database files from the CES cluster. However, the data is not removed. For subsequent object service enablement with a clean configuration and new data, remove object store fileset and set up object environment. See the `mmobj swift base` command. For more information, contact the IBM Support Center.

**nfs**

NFS

**smb**

SMB

**install**

Installs, creates a GPFS cluster, creates NSDs and adds nodes to an existing GPFS cluster. The installation toolkit will use the environment details in the cluster definition file to perform these tasks. If all configuration steps have been completed, this option can be run with no arguments (and pre-install and post-install checks will be performed automatically).

For a “dry-run,” the following arguments are accepted:

**-pr**

Performs a pre-install environment check.

**-po**

Performs a post-install environment check.

**-s SecretKey**

Specifies the secret key on the command line required to decrypt sensitive data in the cluster definition file and suppresses the prompt for the secret key.

**-f**

Forces action without manual confirmation.

**deploy**

Creates file systems, deploys protocols, and configures protocol authentication on an existing GPFS cluster. The installation toolkit will use the environment details in the cluster definition file to perform these tasks. If all configuration steps have been completed, this option can be run with no arguments (and pre-deploy and post-deploy checks will be performed automatically). However, the secret key will be prompted for unless it is passed in as an argument using the **-s** flag.

For a “dry-run,” the following arguments are accepted:

**-pr**

Performs a pre-deploy environment check.

**-po**

Performs a post-deploy environment check.

**-s SecretKey**

Specifies the secret key on the command line required to decrypt sensitive data in the cluster definition file and suppresses the prompt for the secret key.

**-f**

Forces action without manual confirmation.

**upgrade**

Upgrades all components of an existing GPFS cluster. This command can still be used even if all protocols are not enabled. If a protocol is not enabled, then the packages will still be upgraded, but the service won't be started.

The installation toolkit will use environment details in the cluster definition file to perform these tasks. To perform environment health checks prior to and after the upgrade run the **spectrumscale upgrade** command using the **-pr** and **-po** arguments. This is not required, however, because **upgrade** with no arguments will also run this. The following arguments are accepted:

## **spectrumscale**

### **-ve**

Shows the current versions of installed packages and the available version to upgrade to

### **-pr**

Performs health checks on the cluster prior to the upgrade

### **-po**

Performs health checks on the cluster after the upgrade has been completed

### **-f** Forces action without manual confirmation.

## **installgui**

Invokes the installation GUI that can be used to install the IBM Spectrum Scale software on cluster nodes, create an IBM Spectrum Scale cluster, and configure NTP. The installation GUI is used only for installing the system and a separate management GUI needs to be used for configuring and managing the system. The installation GUI cannot be used to upgrade the software in an existing IBM Spectrum Scale system. For more information, see *Installing IBM Spectrum Scale by using the graphical user interface (GUI) in IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

### **start**

Starts the installation GUI

### **status**

Displays the status of the processes that are running on the installation GUI

### **stop**

Stops the installation GUI through the CLI. The installation process through the GUI automatically stops when you exit the installation GUI.

## **Exit status**

**0** Successful completion.

### **nonzero**

A failure has occurred.

## **Security**

You must have root authority to run the **spectrumscale** command.

The node on which the command is issued must be able to execute remote shell commands on any other node in the cluster without the use of a password and without producing any extraneous messages. For more information, see *Requirements for administering a GPFS(tm) file system in IBM Spectrum Scale: Administration Guide*.

## **Examples**

### **Creating a new IBM Spectrum Scale cluster**

1. To instantiate your chef zero server, issue a command similar to the following:  
`spectrumscale setup -s 192.168.0.1`
2. To designate NSD server nodes in your environment to use for the installation, issue this command:  
`./spectrumscale node add FQDN -n`
3. To add four non-shared NSDs seen by a primary NSD server only, issue this command:  
`./spectrumscale nsd add -p FQDN_of_Primary_NSD_Server /dev/dm-1 /dev/dm-2 /dev/dm-3 /dev/dm-4`
4. To add four non-shared NSDs seen by both a primary NSD server and a secondary NSD server, issue this command:  
`./spectrumscale nsd add -p FQDN_of_Primary_NSD_Server -s FQDN_of_Secondary_NSD_Server\ /dev/dm-1 /dev/dm-2 /dev/dm-3 /dev/dm-4`

5. To define a shared root file system using two NSDs and a file system `fs1` using two NSDs, issue these commands:

```
./spectrumscale nsd list
./spectrumscale file system list
./spectrumscale nsd modify nsd1 -fs cesSharedRoot
./spectrumscale nsd modify nsd2 -fs cesSharedRoot
./spectrumscale nsd modify nsd3 -fs fs1
./spectrumscale nsd modify nsd4 -fs fs1
```

6. To designate GUI nodes in your environment to use for the installation, issue this command:

```
./spectrumscale node add FQDN -g -a
```

7. To designate additional client nodes in your environment to use for the installation, issue this command:

```
./spectrumscale node add FQDN
```

8. To allow the installation toolkit to reconfigure Performance Monitoring if it detects any existing configurations, issue this command:

```
./spectrumscale config perfmon -r on
```

9. To name your cluster, issue this command:

```
./spectrumscale config config gpfs -c Cluster_Name
```

10. To review the configuration prior to installation, issue these commands:

```
./spectrumscale node list
./spectrumscale nsd list
./spectrumscale filesystem list
./spectrumscale config gpfs --list
```

11. To start the installation on your defined environment, issue these commands:

```
./spectrumscale install --precheck
./spectrumscale install
```

12. To deploy file systems after a successful installation, do one of the following depending on your requirement:

- If you want to use only the file systems, issue these commands:

```
./spectrumscale deploy --precheck
./spectrumscale deploy
```

- If you want to deploy protocols also, see the examples in the *Deploying protocols on an existing cluster* section.

## Deploying protocols on an existing cluster

**Note:** If your cluster contains ESS, see the *Adding protocols to a cluster containing ESS* section.

1. To instantiate your chef zero server, issue a command similar to the following:

```
spectrumscale setup -s 192.168.0.1
```

2. To designate protocol nodes in your environment to use for the deployment, issue this command:

```
./spectrumscale node add FQDN -p
```

3. To designate GUI nodes in your environment to use for the deployment, issue this command:

```
./spectrumscale node add FQDN -g -a
```

4. To configure protocols to point to a file system that will be used as a shared root, issue this command:

```
./spectrumscale config protocols -f FS_Name -m FS_Mountpoint
```

5. To configure a pool of export IPs, issue this command:

```
./spectrumscale config protocols -e Comma_Separated_List_of_Exportpool_IPs
```

6. To enable NFS on all protocol nodes, issue this command:

```
./spectrumscale enable nfs
```

7. To enable SMB on all protocol nodes, issue this command:

## **spectrumscale**

```
./spectrumscale enable smb
```

8. To enable Object on all protocol nodes, issue these commands:

```
./spectrumscale enable object
./spectrumscale config object -au Admin_User -ap Admin_Password -dp Database_Password
./spectrumscale config object -e FQDN
./spectrumscale config object -f FS_Name -m FS_Mountpoint
./spectrumscale config object -o Object_Filesset
```

9. To review the configuration prior to deployment, issue these commands:

```
./spectrumscale config protocols
./spectrumscale config object
./spectrumscale node list
```

10. To deploy protocols on your defined environment, issue these commands:

```
./spectrumscale deploy --precheck
./spectrumscale deploy
```

## **Deploying protocol authentication**

**Note:** For the following example commands, it is assumed that the protocols cluster was deployed successfully using the **spectrumscale** command options.

1. To enable file authentication with AD server on all protocol nodes, issue this command:

```
./spectrumscale auth file ad
```

Fill out the template and save the information, and then issue the following commands:

```
./spectrumscale deploy --precheck
./spectrumscale deploy
```

2. To enable Object authentication with AD server on all protocol nodes, issue this command:

```
./spectrumscale auth object ad
```

Fill out the template and save the information, and then issue the following commands:

```
./spectrumscale deploy --precheck
./spectrumscale deploy
```

## **Upgrading an IBM Spectrum Scale cluster**

1. Extract the IBM Spectrum Scale package for the required code level by issuing a command similar to the following depending on the package name:

```
./Spectrum_Scale_Protocols_Standard-4.2.x.x-xxxxx
```

2. Copy the cluster definition file from the prior installation to the latest installer location by issuing this command:

```
| cp -p /usr/lpp/mmfs/4.2.2.0/installer/configuration/clusterdefinition.txt \
| /usr/lpp/mmfs/4.2.3.0/installer/configuration/
```

**Note:** This is a command example of when you are upgrading from 4.2.2.0 to 4.2.3.0.

You can populate the cluster definition file with the current cluster state by issuing the **spectrumscale config populate** command.

3. Run the upgrade precheck from the installer directory of the latest code level extraction by issuing commands similar to the following:

```
cd /usr/lpp/mmfs/Latest_Code_Level_Directory/installer
./spectrumscale upgrade --precheck
```

**Note:** If you are upgrading to IBM Spectrum Scale version 4.2.2, the upgrade precheck updates the operating system and CPU architecture fields in the cluster definition file. You can also update the operating system and CPU architecture fields in the cluster definition file by issuing the **spectrumscale config update** command.

4. Run the upgrade by issuing this command:

```
cd /usr/lpp/mmfs/Latest_Code_Level_Directory/installer
./spectrumscale upgrade
```

### **Adding to an installation process**

1. To add nodes to an installation, do the following:
  - a. Add one or more node types using the following commands:
    - Client nodes:  
./spectrumscale node add *FQDN*
    - NSD nodes:  
./spectrumscale node add *FQDN* -n
    - Protocol nodes:  
./spectrumscale node add *FQDN* -p
    - GUI nodes:  
./spectrumscale node add *FQDN* -g -a
  - b. Install GPFS on the new nodes using the following commands:  
./spectrumscale install --precheck  
./spectrumscale install
  - c. If protocol nodes are being added, deploy protocols using the following commands:  
./spectrumscale deploy --precheck  
./spectrumscale deploy
2. To add NSDs to an installation, do the following:
  - a. Verify that the NSD server connecting this new disk runs an OS compatible with the installation toolkit and that the NSD server exists within the cluster.
  - b. Add NSDs to the installation using the following command:  
./spectrumscale nsd add -p *FQDN\_of\_Primary\_NSD\_Server Path\_to\_Disk\_Device\_File*
  - c. Run the installation using the following commands:  
./spectrumscale install --precheck  
./spectrumscale install
3. To add file systems to an installation, do the following:
  - a. Verify that free NSDs exist and that they can be listed by the installation toolkit using the following commands:  
mm1snsd  
./spectrumscale nsd list
  - b. Define the file system using the following command:  
./spectrumscale nsd modify *NSD* -fs *File\_System\_Name*
  - c. Deploy the file system using the following commands:  
./spectrumscale deploy --precheck  
./spectrumscale deploy
4. To enable another protocol on an existing cluster that has protocols enabled, do the following steps depending on your configuration:
  - a. Enable NFS on all protocol nodes using the following command:  
./spectrumscale enable nfs
  - b. Enable SMB on all protocol nodes using the following command:  
./spectrumscale enable smb
  - c. Enable Object on all protocol nodes using the following commands:  
./spectrumscale enable object  
./spectrumscale config object -au *Admin\_User* -ap *Admin\_Password* -dp *Database\_Password*  
./spectrumscale config object -e *FQDN*  
./spectrumscale config object -f *FS\_Name* -m *FS\_Mountpoint*  
./spectrumscale config object -o *Object\_Filesset*

## **spectrumscale**

- d. Enable the new protocol using the following commands:

```
./spectrumscale deploy --precheck
./spectrumscale deploy
```

### **Adding protocols to a cluster containing ESS**

For information on preparing a cluster that contains ESS for adding protocols, see *Preparing a cluster that contains ESS for adding protocols* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

After you have prepared your cluster that contains ESS for adding protocols, you can use commands similar to the ones listed in the *Deploying protocols on an existing cluster* section.

### **Diagnosing an error during install, deploy, or upgrade**

1. Note the screen output indicating the error. This helps in narrowing down the general failure.  
When a failure occurs, the screen output also shows the log file containing the failure.
2. Open the log file in an editor such as vi.
3. Go to the end of the log file and search upwards for the text FATAL.
4. Find the topmost occurrence of FATAL (or first FATAL error that occurred) and look above and below this error for further indications of the failure.

For more information, see *Finding deployment related error messages more easily and using them for failure analysis* in *IBM Spectrum Scale: Problem Determination Guide*.

### **See also**

- *Installing IBM Spectrum Scale on Linux nodes and deploying protocols* in *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.
- *Configuring with the spectrumscale installation toolkit* in the *IBM Spectrum Scale: Administration Guide*.
- “mmchconfig command” on page 132
- “mmlscluster command” on page 397
- “mmlsconfig command” on page 400
- “mmnfs command” on page 454
- “mmobj command” on page 466
- “mmsmb command” on page 554
- “mmuserauth command” on page 581

### **Location**

| /usr/lpp/mmfs/4.2.3.0/installer

---

## Chapter 2. IBM Spectrum Scale Data Management API for GPFS information

The Data Management Application Programming Interface (DMAPI) for (GPFS) is based on The Open Group's System Management: Data Storage Management (XDSM) API Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X specification. The implementation is compliant with the standard. Some optional features are not implemented.

The XDSM DMAPI model is intended mainly for a single-node environment. Some of the key concepts, such as sessions, event delivery, and recovery, required enhancements for a multiple-node environment such as GPFS.

---

### Overview of IBM Spectrum Scale Data Management API for GPFS

The Data Management Application Programming Interface (DMAPI) for GPFS allows you to monitor events associated with a GPFS file system or with an individual file. You can also manage and maintain file system data.

See the IBM Spectrum Scale FAQ in IBM Knowledge Center([www.ibm.com/support/knowledgecenter/STXKQY/gpfsclustersfaq.html](http://www.ibm.com/support/knowledgecenter/STXKQY/gpfsclustersfaq.html)) for the current limitations of DMAPI-managed file systems.

**Note:** IBM Spectrum Protect for Space Management for GPFS file systems is not available for Windows.

DMAPI for GPFS is compliant with the Open Group's XDSM Standard and includes these features:

- “GPFS-specific DMAPI events”
- “DMAPI functions” on page 622
- “DMAPI configuration attributes” on page 626
- “DMAPI restrictions for GPFS” on page 627

### GPFS-specific DMAPI events

There are three GPFS-specific DMAPI events: events implemented in DMAPI for GPFS, optional events that are not implemented in DMAPI for GPFS, and GPFS-specific attribute events that are not part of the DMAPI standard.

For more information, see:

- “Events implemented in DMAPI for GPFS”
- “Optional events that are not implemented in DMAPI for GPFS” on page 622
- “GPFS-specific attribute events that are not part of the DMAPI standard” on page 622

### Events implemented in DMAPI for GPFS

These are the events, as defined in the *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X, implemented in DMAPI for GPFS:

#### File system administration events

- mount
- preunmount
- unmount
- nospace

### **Namespace events**

- create, postcreate
- remove, postremove
- rename, postrename
- symlink, postsymlink
- link, postlink

### **Data events**

- read
- write
- truncate

### **Metadata events**

- attribute
- destroy
- close

### **Pseudo event**

- user event

GPFS guarantees that asynchronous events are delivered, except when the GPFS daemon fails. Events are enqueued to the session before the corresponding file operation completes. For further information on failures, see “Failure and recovery of IBM Spectrum Scale Data Management API for GPFS” on page 657.

## **Optional events that are not implemented in DMAPI for GPFS**

The following optional events, as defined in the *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X, are **not** implemented in DMAPI for GPFS:

### **File system administration event**

- debut

### **Metadata event**

- cancel

## **GPFS-specific attribute events that are not part of the DAPI standard**

GPFS generates the following attribute events for DAPI that are specific to GPFS and not part of the DAPI standard:

- Pre-permission change
- Post-permission change

For additional information, refer to “GPFS-specific DAPI events” on page 655.

## **DMAPI functions**

All mandatory DAPI functions and most optional functions that are defined in the *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X, are implemented in DAPI for GPFS.

For C declarations of all the functions implemented in DAPI for GPFS, refer to the **dmapi.h** file located in the **/usr/lpp/mmfs/include** directory.

For changes and restrictions on functions in DMAPI for GPFS, see “Usage restrictions on DMAPI functions” on page 639, and “Semantic changes to DMAPI functions” on page 654.

## Mandatory functions implemented in DMAPI for GPFS

These mandatory functions, as defined in the *System Management: Data Storage Management (XDSM) API Common Applications Environment (CAE) Specification C429*, The Open Group, ISBN 1-85912-190-X, are implemented in DMAPI for GPFS.

For C declarations of all the mandatory functions implemented in DMAPI for GPFS, refer to the **dmapi.h** file located in the **/usr/lpp/mmfs/include** directory. However, for a quick description of the mandatory functions and their applications, refer to the following set of functions:

### **dm\_create\_session**

Create a new session.

### **dm\_create\_userevent**

Create a pseudo-event message for a user.

### **dm\_destroy\_session**

Destroy an existing session.

### **dm\_fd\_to\_handle**

Create a file handle using a file descriptor.

### **dm\_find\_eventmsg**

Return the message for an event.

### **dm\_get\_allocinfo**

Get a file's current allocation information.

### **dm\_get\_bulkattr**

Get bulk attributes of a file system.

### **dm\_get\_config**

Get specific data on DMAPI implementation.

### **dm\_get\_config\_events**

List all events supported by the DMAPI implementation.

### **dm\_get\_dirattrs**

Return a directory's bulk attributes.

### **dm\_get\_eventlist**

Return a list of an object's enabled events.

### **dm\_get\_events**

Return the next available event messages.

### **dm\_get\_fileattr**

Get file attributes.

### **dm\_get\_mountinfo**

Return details from a mount event.

### **dm\_get\_region**

Get a file's managed regions.

### **dm\_getall\_disp**

For a given session, return the disposition of all file system's events.

### **dm\_getall\_sessions**

Return all extant sessions.

### **dm\_getall\_tokens**

Return a session's outstanding tokens.

- dm\_handle\_cmp**  
Compare file handles.
- dm\_handle\_free**  
Free a handle's storage.
- dm\_handle\_hash**  
Hash the contents of a handle.
- dm\_handle\_is\_valid**  
Check a handle's validity.
- dm\_handle\_to\_fhandle**  
Return the file system handle associated with an object handle.
- dm\_handle\_to\_path**  
Return a path name from a file system handle.
- dm\_init\_attrloc**  
Initialize a bulk attribute location offset.
- dm\_init\_service**  
Initialization processing that is implementation-specific.
- dm\_move\_event**  
Move an event from one session to another.
- dm\_path\_to\_fhandle**  
Create a file system handle using a path name.
- dm\_path\_to\_handle**  
Create a file handle using a path name.
- dm\_query\_right**  
Determine an object's access rights.
- dm\_query\_session**  
Query a session.
- dm\_read\_invis**  
Read a file without using DMAPI events.
- dm\_release\_right**  
Release an object's access rights.
- dm\_request\_right**  
Request an object's access rights.
- dm\_respond\_event**  
Issue a response to an event.
- dm\_send\_msg**  
Send a message to a session.
- dm\_set\_disp**  
For a given session, set the disposition of all file system's events.
- dm\_set\_eventlist**  
For a given object, set the list of events to be enabled.
- dm\_set\_fileattr**  
Set a file's time stamps, ownership and mode.
- dm\_set\_region**  
Set a file's managed regions.

**dm\_write\_invis**

Write to a file without using DMAPI events.

**Optional functions implemented in DMAPI for GPFS**

These optional functions, as defined in the *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X, are implemented in DMAPI for GPFS.

For C declarations of these optional functions implemented in DMAPI for GPFS, refer to the **dmapi.h** file located in the **/usr/lpp/mmfs/include** directory. However, for a quick description of the optional functions and their applications, refer to the following set of functions:

**dm\_downgrade\_right**

Change an exclusive access right to a shared access right.

**dm\_get\_bulkall**

Return a file system's bulk data management attributes.

**dm\_get\_dmattr**

Return a data management attribute.

**dm\_getall\_dmattr**

Return all data management attributes of a file.

**dm\_handle\_to\_fsid**

Get a file system ID using its handle.

**dm\_handle\_to\_igen**

Get inode generation count using a handle.

**dm\_handle\_to\_ino**

Get inode from a handle.

**dm\_make\_handle**

Create a DMAPI object handle.

**dm\_make\_fshandle**

Create a DMAPI file system handle.

**dm\_punch\_hole**

Make a hole in a file.

**dm\_probe\_hole**

Calculate the rounded result of the area where it is assumed that a hole is to be punched.

**dm\_remove\_dmattr**

Delete a data management attribute.

**dm\_set\_dmattr**

Define or update a data management attribute.

**dm\_set\_return\_on\_destroy**

Indicate a DM attribute to return with destroy events.

**dm\_sync\_by\_handle**

Synchronize the in-memory state of a file with the physical medium.

**dm\_upgrade\_right**

Change a currently held access right to be exclusive.

**Optional functions that are not implemented in DMAPI for GPFS**

There are optional functions that are not implemented in DMAPI for GPFS.

The following optional functions, as defined in the *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X, are **not** implemented in DMAPI for GPFS:

**dm\_clear\_inherit**

Reset the inherit-on-create status of an attribute.

**dm\_create\_by\_handle**

Define a file system object using a DM handle.

**dm\_getall\_inherit**

Return a file system's inheritable attributes.

**dm\_mkdir\_by\_handle**

Define a directory object using a handle.

**dm\_obj\_ref\_hold**

Put a hold on a file system object.

**dm\_obj\_ref\_query**

Determine if there is a hold on a file system object.

**dm\_obj\_ref\_rele**

Release the hold on a file system object.

**dm\_pending**

Notify FS of slow DM application processing.

**dm\_set\_inherit**

Indicate that an attribute is inheritable.

**dm\_symlink\_by\_handle**

Define a symbolic link using a DM handle.

## GPFS-specific DAPI functions

There are several GPFS-specific DAPI functions that are not part of the DAPI open standard.

The GPFS-specific functions are listed and described in “Definitions for GPFS-specific DAPI functions” on page 641.

## DAPI configuration attributes

The *System Management: Data Storage Management (XDSM) API* Common Applications Environment (CAE) Specification C429, The Open Group, ISBN 1-85912-190-X defines a set of configuration attributes to be exported by each DAPI implementation. These attributes specify which optional features are supported and give bounds on various resources.

The Data Management (DM) application can query the attribute values using the function **dm\_get\_config**. It can also query which events are supported, using the function **dm\_get\_config\_events**.

The functions **dm\_get\_config** and **dm\_get\_config\_events** receive a file handle from input arguments *happ* and *hlen*. In GPFS, both functions ignore the handle, as the configuration is not dependent on the specific file or file system. This enables the DM application to query the configuration during initialization, when file handles may not yet be available.

**Note:** To ensure that the most current values are being used, the DM application should always query the configuration at runtime by using **dm\_get\_config**.

Table 21 on page 627 shows the attribute values that are used in DAPI for GPFS:

Table 21. DMAPI configuration attributes

| Name                            | Value |
|---------------------------------|-------|
| DM_CONFIG_BULKALL               | 1     |
| DM_CONFIG_CREATE_BY_HANDLE      | 0     |
| DM_CONFIG_DTIME_OVERLOAD        | 1     |
| DM_CONFIG_LEGACY                | 1     |
| DM_CONFIG_LOCK_UPGRADE          | 1     |
| DM_CONFIG_MAX_ATTR_ON_DESTROY   | 1022  |
| DM_CONFIG_MAX_ATTRIBUTE_SIZE    | 1022  |
| DM_CONFIG_MAX_HANDLE_SIZE       | 32    |
| DM_CONFIG_MAX_MANAGED_REGIONS   | 32    |
| DM_CONFIG_MAX_MESSAGE_DATA      | 4096  |
| DM_CONFIG_OBJ_REF               | 0     |
| DM_CONFIG_PENDING               | 0     |
| DM_CONFIG_PERS_ATTRIBUTE        | 1     |
| DM_CONFIG_PERS_EVENTS           | 1     |
| DM_CONFIG_PERS_INHERIT_ATTRIBS  | 0     |
| DM_CONFIG_PERS_MANAGED_REGIONS  | 1     |
| DM_CONFIG_PUNCH_HOLE            | 1     |
| DM_CONFIG_TOTAL_ATTRIBUTE_SPACE | 7168  |
| DM_CONFIG_WILL_RETRY            | 0     |

Attribute value **DM\_CONFIG\_TOTAL\_ATTRIBUTE\_SPACE** is per file. The entire space is available for opaque attributes. Non-opaque attributes (event list and managed regions) use separate space.

## DMAPI restrictions for GPFS

All DMAPI APIs must be called from nodes that are in the cluster where the file system is created. DMAPI APIs may **not** be invoked from a remote cluster.

Furthermore, GPFS also places the following DMAPI API restrictions:

- Running **dm\_get\_events** with the **DM\_EV\_WAIT** flag set causes the calling process to wait uninterruptibly.
- Interacting with a handle after calling **dm\_handle\_free** will result in undefined behavior.

In addition to the DMAPI API restrictions, GPFS places the following restrictions on the use of file system snapshots when you have DMAPI enabled:

- Snapshots cannot coexist with file systems using GPFS 3.1 or earlier.
- GPFS 3.2 and later permits snapshots with DMAPI-enabled file systems. However, GPFS places the following restrictions on DMAPI access to the snapshot files:
  - The DM server may read files in a snapshot using **dm\_read\_invis**.
  - The DM server is not allowed to modify or delete the file using **dm\_write\_invis** or **dm\_punch\_hole**.
  - The DM server is not allowed to establish a managed region on the file.
  - Snapshot creation or deletion does not generate DMAPI namespace events.
  - Snapshots of a file are not managed regardless of the state of the original file and they do not inherit the DMAPI attributes of the original file.

---

## Concepts of IBM Spectrum Scale Data Management API for GPFS

The XDSM standard is intended mainly for a single-node environment. Some of the key concepts in the standard such as sessions, event delivery, mount and unmount, and failure and recovery, are not well defined for a multiple-node environment such as GPFS.

For a list of restrictions and coexistence considerations, see “Usage restrictions on DMAPI functions” on page 639.

All DAPI APIs must be called from nodes that are in the cluster where the file system is created.

Key concepts of DAPI for GPFS include these areas:

- “Sessions”
- “Data management events”
- “Mount and unmount” on page 630
- “Tokens and access rights” on page 631
- “Parallelism in Data Management applications” on page 632
- “Data Management attributes” on page 633
- “Support for NFS” on page 633
- “Quota” on page 633
- “Memory mapped files” on page 633

### Sessions

In GPFS, a session is associated only with the node on which the session was created. This node is known as the *session node*.

Events are generated at any node where the file system is mounted. The node on which a given event is generated is called the *source node* of that event. The event is delivered to a session queue on the session node.

There are restrictions as to which DAPI functions can and cannot be called from a node other than the session node. In general, functions that change the state of a session or event can only be called on the session node. For example, the maximum number of DAPI sessions that can be created on a node is 4000. See “Usage restrictions on DAPI functions” on page 639 for details.

Session ids are unique over time within a GPFS cluster. When an existing session is assumed, using `dm_create_session`, the new session id returned is the same as the old session id.

A session fails when the GPFS daemon fails on the session node. Unless this is a total failure of GPFS on all nodes, the session is recoverable. The DM application is expected to assume the old session, possibly on another node. This will trigger the reconstruction of the session queue. All pending synchronous events from surviving nodes are resubmitted to the recovered session queue. Such events will have the same token id as before the failure, except mount events. Asynchronous events, on the other hand, are lost when the session fails. See “Failure and recovery of IBM Spectrum Scale Data Management API for GPFS” on page 657 for information on failure and recovery.

### Data management events

Data management events arrive on a session queue from any of the nodes in the GPFS cluster.

The source node of the event is identified by the `ev_nodeid` field in the header of each event message in the structure `dm_eventmsg`. The identification is the GPFS cluster data node number, which is attribute `node_number` in the `mmsdrfs2` file for a PSSP node or `mmsdrfs` file for any other type of node.

Data Management events are generated only if the following two conditions are true:

1. The event is enabled.
2. It has a disposition.

A file operation will fail with the **EIO** error if there is no disposition for an event that is enabled and would otherwise be generated.

A list of enabled events can be associated individually with a file and globally with an entire file system. The XDSM standard leaves undefined the situation where the individual and the global event lists are in conflict. In GPFS, such conflicts are resolved by always using the individual event list, if it exists.

**Note:** The XDSM standard does not provide the means to remove the individual event list of a file. Thus, there is no way to enable or disable an event for an entire file system without explicitly changing each conflicting individual event list.

In GPFS, event lists are persistent.

Event dispositions are specified per file system and are not persistent. They must be set explicitly after the session is created.

Event generation mechanisms have limited capacity. In case resources are exceeded, new file operations will wait indefinitely for free resources.

File operations wait indefinitely for a response from synchronous events. The **dmapEventTimeout** configuration attribute on the **mmchconfig** command, can be used to set a timeout on events that originate from NFS file operations. This is necessary because NFS servers have a limited number of threads that cannot be blocked for long periods of time. Refer to “GPFS configuration attributes for DMAPI” on page 635 and “Support for NFS” on page 633.

The XDSM standard permits asynchronous events to be discarded at any time. In GPFS, asynchronous events are guaranteed when the system runs normally, but may be lost during abnormal conditions, such as failure of GPFS on the session node. Asynchronous events are delivered in a timely manner. That is, an asynchronous event is enqueued to the session before the corresponding file operation completes.

Figure 1 on page 630, shows the flow of a typical synchronous event in a multiple-node GPFS environment. The numbered arrows in the figure correspond to the following steps:

1. The user application on the source node performs a file operation on a GPFS file. The file operation thread generates a synchronous event and blocks, waiting for a response.
2. GPFS on the source node sends the event to GPFS on the session node, according to the disposition for that event. The event is enqueued to the session queue on the session node.
3. The Data Management application on the session node receives the event (using **dm\_get\_events**) and handles it.
4. The Data Management application on the session node responds to the event (using **dm\_respond\_event**).
5. GPFS on the session node sends the response to GPFS on the source node.
6. GPFS on the source node passes the response to the file operation thread and unblocks it. The file operation continues.

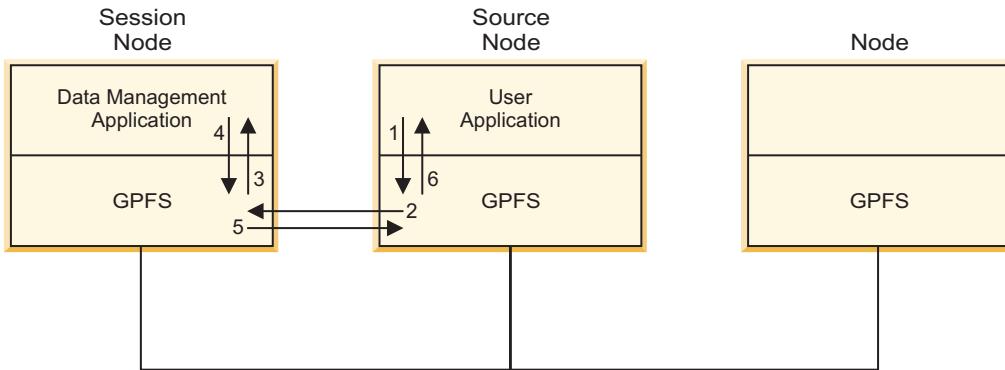


Figure 1. Flow of a typical synchronous event in a multiple-node GPFS environment

## Reliable DMAPI destroy events

A metadata destroy event is generated when the operating system has destroyed an object. This type of event is different from a remove event, which is a namespace event and is not related to the destruction of an object. A reliable destroy event supports synchronous destroy events in the same way that other synchronous events do. When a synchronous event is generated, a user process is suspended in the kernel; it will be suspended until a DM application issues an explicit response to the event. The DM application at the session that supports the reliable destroy event must be capable of handling the synchronous destroy event. In other words, it must respond to the **DM\_EVENT\_DESTROY** event with **DM\_RESPOND\_EVENT**. Otherwise, the event will wait forever at the session node for a response. Based on this, it is recommended that the cluster not be made up of nodes that are running back-level code and new code, because the destroy event is not reliable in a mixed environment.

## Mount and unmount

The XDSM standard implicitly assumes that there is a single mount, pre-unmount and unmount event per file system. In GPFS, a separate mount event is generated by each mount operation on each node. Similarly, if the pre-unmount and unmount events are enabled, they are generated by each unmount operation on each node. Thus, there may be multiple such events for the same file system.

To provide additional information to the DM application, the mode field in the respective event message structures (**me\_mode** for mount, and **ne\_mode** for pre-unmount and unmount) has a new flag, **DM\_LOCAL\_MOUNT**, which is not defined in the standard. When the flag is set, the mount or unmount operation is local to the session node. In addition, the new field **ev\_nodeid** in the header of the event message can be used to identify the source node where the mount or unmount operation was invoked. The identification is the GPFS cluster data node number, which is attribute **node\_number** in the **mmsdrfs2** file for a PSSP node or **mmsdrfs** file for any other type of node.

The mount event is sent to multiple sessions that have a disposition for it. If there is no disposition for the mount event, the mount operation fails with an **EIO** error.

There is no practical way to designate the *last* unmount, since there is no serialization of all mount and unmount operations of each file system. Receiving an unmount event with the value 0 in the **ne\_retcodes** field is no indication that there will be no further events from the file system.

An unmount initiated internally by the GPFS daemon, due to file system forced unmount or daemon shutdown, will not generate any events. Consequently, there need not be a match between the number of mount events and the number of pre-unmount or unmount events for a given file system.

The **dmapiMountTimeout** attribute on the **mmchconfig** command enables blocking the mount operation for a limited time until some session has set the mount disposition. This helps GPFS and the DM application synchronize during initialization. See “GPFS configuration attributes for DMAPI” on page 635 and “Initializing the Data Management application” on page 637.

Mount events are enqueued on the session queue ahead of any other events. This gives mount events a higher priority, which improves the response time for mount events when the queue is very busy.

If the **DM\_UNMOUNT\_FORCE** flag is set in the pre-unmount event message, the response of the DM application to the pre-unmount event is ignored, and the forced unmount proceeds. If the **DM\_LOCAL\_MOUNT** flag is also set, the forced unmount will result in the loss of all access rights of the given file system that are associated with any local session.

If the unmount is not forced (the **DM\_UNMOUNT\_FORCE** flag is not set), and the **DM\_LOCAL\_MOUNT** flag is set, the DM application is expected to release all access rights on files of the given file system associated with any local session. If any access rights remain held after the **DM\_RESP\_CONTINUE** response is given, the unmount will fail with **EBUSY**. This is because access rights render the file system busy, similar to other locks on files.

The function **dm\_get\_mountinfo** can be called from any node, even if the file system is not mounted on that node. The **dm\_mount\_event** structure returned by the **dm\_get\_mountinfo** function provides the following enhanced information. The **me\_mode** field contains two new flags, **DM\_LOCAL\_MOUNT** and **DM\_REMOTE\_MOUNT**. At least one of the two flags is always set. When both flags are set simultaneously, it is an indication that the file system is mounted on the local node, as well as one or more other (remote) nodes. When only **DM\_LOCAL\_MOUNT** is set, it is an indication that the file system is mounted on the local node but not on any other node. When only **DM\_REMOTE\_MOUNT** is set, it is an indication that the file system is mounted on some remote node, but not on the local node.

In the latter case (only **DM\_REMOTE\_MOUNT** is set), the fields **me\_roothandle** and **me\_handle2** (the mount point handle) in the **dm\_mount\_event** structure are set to **DM\_INVALID\_HANDLE**. Also in this case, the **me\_name1** field (the mount point path) is taken from the stanza in the file **/etc/filesystems** on one of the remote nodes (with the use of GPFS cluster data, the stanzas on all nodes are identical).

The enhanced information provided by the **dm\_get\_mountinfo** function can be useful during the processing of mount and pre-unmount events. For example, before responding to a mount event from a remote (non-session) node, **dm\_get\_mountinfo** could be invoked to find out whether the file system is already mounted locally at the session node, and if not, initiate a local mount. On receiving a pre-unmount event from the local session node, it is possible to find out whether the file system is still mounted elsewhere, and if so, fail the local unmount or delay the response until after all remote nodes have unmounted the file system.

**Note:** The **DM\_REMOTE\_MOUNT** flag is redundant in the **dm\_mount\_event** structure obtained from the mount event (as opposed to the **dm\_get\_mountinfo** function).

## Tokens and access rights

A DMAPI token is an identifier of an outstanding event (a synchronous event that the DM application has received and is currently handling). The token is unique over time in the cluster. The token becomes invalid when the event receives a response.

The main purpose of tokens is to convey access rights in DMAPI functions. Access rights are associated with a specific event token. A function requiring access rights to some file may present an event token that has the proper access rights.

DMAPI functions can also be invoked using **DM\_NO\_TOKEN**, in which case sufficient access protection is provided for the duration of the operation. This is semantically equivalent to holding an access right, but no access right on the file is actually acquired.

In GPFS, when an event is received, its token has no associated access rights.

DM access rights are implemented in GPFS using an internal lock on the file. Access rights can be acquired, changed, queried, and released only at the session node. This is an implementation restriction caused by the GPFS locking mechanisms.

In GPFS, it is not possible to set an access right on an entire file system from the file system handle. Thus, DMAPI function calls that reference a file system, using a file system handle, are not allowed to present a token and must specify **DM\_NO\_TOKEN**. For the same reason, functions that acquire or change access rights are not allowed to present a file system handle.

Holding access rights renders the corresponding file system busy at the session node, preventing normal (non-forced) unmount. This behavior is similar to that of other locks on files. When receiving a pre-unmount event, the DM application is expected to release all access rights before responding. Otherwise, the unmount operation will fail with an **EBUSY** error.

All access rights associated with an event token are released when the response is given. There is no transfer of access rights from DMAPI to the file operation thread. The file operation will acquire any necessary locks after receiving the response of the event.

## Parallelism in Data Management applications

Given the multiple-node environment of GPFS, it is desirable to exploit parallelism in the Data Management application as well.

This can be accomplished in several ways:

- On a given session node, multiple DM application threads can access the same file in parallel, using the same session. There is no limit on the number of threads that can invoke DMAPI functions simultaneously on each node.
- Multiple sessions, each with event dispositions for a different file system, can be created on separate nodes. Thus, files in different file systems can be accessed independently and simultaneously, from different session nodes.
- Dispositions for events of the same file system can be partitioned among multiple sessions, each on a different node. This distributes the management of one file system among several session nodes.
- Although GPFS routes all events to a single session node, data movement may occur on multiple nodes. The function calls **dm\_read\_invis**, **dm\_write\_invis**, **dm\_probe\_hole**, and **dm\_punch\_hole** are honored from a root process on another node, provided it presents a session ID for an established session on the session node.

A DM application may create a *worker process*, which exists on any node within the GPFS cluster. This worker process can move data to or from GPFS using the **dm\_read\_invis** and **dm\_write\_invis** functions. The worker processes must adhere to these guidelines:

1. They must run as root.
2. They must present a valid session ID that was obtained on the session node.
3. All writes to the same file which are done in parallel must be done in multiples of the file system block size, to allow correct management of disk blocks on the writes.
4. No DMAPI calls other than **dm\_read\_invis**, **dm\_write\_invis**, **dm\_probe\_hole**, and **dm\_punch\_hole** may be issued on nodes other than the session node. This means that any rights required on a file must be obtained within the session on the session node, prior to the data movement.

5. There is no persistent state on the nodes hosting the worker process. It is the responsibility of the DM application to recover any failure which results from the failure of GPFS or the data movement process.

## Data Management attributes

Data Management attributes can be associated with any individual file. There are opaque and non-opaque attributes.

An opaque attribute has a unique name, and a byte string value which is not interpreted by the DMAP API implementation. Non-opaque attributes, such as managed regions and event lists, are used internally by the DMAP API implementation.

DM attributes are persistent. They are kept in a hidden file in the file system.

GPFS provides two *quick access* single-bit opaque DM attributes for each file, stored directly in the inode. These attributes are accessible through regular DMAP API functions, by specifying the reserved attribute names `_GPFSSQA1` and `_GPFSSQA2` (where `_GPF` is a reserved prefix). The attribute data must be a single byte with contents 0 or 1.

## Support for NFS

A DM application could be slow in handling events. NFS servers have a limited number of threads which must not all be blocked simultaneously for extended periods of time. GPFS provides a mechanism to guarantee progress of NFS file operations that generate data events without blocking the server threads indefinitely.

The mechanism uses a timeout on synchronous events. Initially the NFS server thread is blocked on the event. When the timeout expires, the thread unblocks and the file operation fails with an `ENOTREADY` error code. The event itself continues to exist and will eventually be handled. When a response for the event arrives at the source node it is saved. NFS is expected to periodically retry the operation. The retry will either find the response which has arrived between retries, or cause the operation to fail again with `ENOTREADY`. After repeated retries, the operation is eventually expected to succeed.

The interval is configurable using the `dmapEventTimeout` configuration attribute on the `mmchconfig` command. See “GPFS configuration attributes for DMAP API” on page 635. The default is no timeout.

The timeout mechanism is activated only for data events (read, write, truncate), and only when the file operation comes from NFS.

- | For the parameter change to take effect, restart the GPFS daemon on all nodes.

## Quota

GPFS supports user quota. When `dm_punch_hole` is invoked, the file owner's quota is adjusted by the disk space that is freed. The quota is also adjusted when `dm_write_invis` is invoked and additional disk space is consumed.

Since `dm_write_invis` runs with root credentials, it will never fail due to insufficient quota. However, it is possible that the quota of the file owner will be exceeded as a result of the invisible write. In that case the owner will not be able to perform further file operations that consume quota.

## Memory mapped files

In GPFS, a read event or a write event will be generated (if enabled) at the time the memory mapping of a file is established.

No events will be generated during actual mapped access, regardless of the setting of the event list or the managed regions. Access to the file with regular file operations, while the file is memory mapped, will generate events, if such events are enabled.

To protect the integrity of memory mapped access, the DM application is not permitted to punch a hole in a file while the file is memory mapped. If the DM application calls `dm_punch_hole` while the file is memory mapped, the error code `EBUSY` will be returned.

## Administration of IBM Spectrum Scale Data Management API for GPFS

To set up the DMAPI for GPFS, install the DAPI files that are included in the GPFS installation package, and then choose the configuration attributes for DAPI with the `mmchconfig` command. For each file system that you want DAPI access, enable DAPI with the `-z` flag of the `mmcrfs` or `mmchfs` command.

All DAPI APIs must be called from nodes that are in the cluster where the file system is created. DAPI APIs may **not** be invoked from a remote cluster. The GPFS daemon and each DAPI application must be synchronized to prevent failures.

Administration of DAPI for GPFS includes:

- “Required files for implementation of Data Management applications”
- “GPFS configuration attributes for DAPI” on page 635
- “Enabling DAPI for a file system” on page 636
- “Initializing the Data Management application” on page 637

## Required files for implementation of Data Management applications

The installation image for GPFS contains the required files for implementation of Data Management applications.

For more information about installation, see the *IBM Spectrum Scale: Concepts, Planning, and Installation Guide*.

The required files are:

### **dmapi.h**

The header file that contains the C declarations of the DAPI functions.

This header file must be included in the source files of the DM application.

The file is installed in directory: `/usr/lpp/mmfs/include`.

### **dmapi\_types.h**

The header file that contains the C declarations of the data types for the DAPI functions and event messages.

The header file `dmapi.h` includes this header file.

The file is installed in directory: `/usr/lpp/mmfs/include`.

### **libdmapi.a**

The library that contains the DAPI functions.

The library `libdmapi.a` consists of a single shared object, which is built with auto-import of the system calls that are listed in the export file `dmapi.exp`.

The file is installed in directory: `/usr/lpp/mmfs/lib`.

### **dmapi.exp**

The export file that contains the DAPI system call names.

The file **dapi.exp** needs to be explicitly used only if the DM application is to be explicitly built with static binding, using the binder options **-bnso -bI:dapi.exp**.

The file is installed in directory: **/usr/lpp/mmfs/lib**.

#### **dmapicalls, dmapicalls64**

Module loaded during processing of the DMAPI functions.

The module is installed in directory: **/usr/lpp/mmfs/bin**.

#### **Notes:**

- On Linux nodes running DAPI, the required files **libdapi.a**, **dapi.exp**, **dmapicalls**, and **dmapicalls64** are replaced by **libdapi.so**.
- If you are compiling with a non-IBM compiler on AIX nodes, you must compile DAPI applications with **-D\_AIX**.

## **GPFS configuration attributes for DAPI**

GPFS uses several attributes for DAPI that define various timeout intervals. These attributes can be changed with the **mmchconfig** command.

The DAPI configuration attributes are:

#### **dapiEventDataEventRetry**

Controls how GPFS handles the data event when it is enabled again right after this event is handled by the DAPI application. Valid values are:

- 1** Specifies that GPFS will always regenerate the event as long as it is enabled. This value should only be used when the DAPI application recalls and migrates the same file in parallel by many processes at the same time.
- 0** Specifies to never regenerate the event. This value should not be used if a file could be migrated and recalled at the same time.

#### **Positive Number**

Specifies how many times the data event should be retried. The default is 2, which should be enough to cover most DAPI applications. Unless a special situation occurs, you can increase this to a larger number or even set this to **-1** to always regenerate the events. Unless you perform careful testing, IBM recommends that you never change the default setting.

#### **dapiEventTimeout**

Controls the blocking of file operation threads of NFS, while in the kernel waiting for the handling of a DAPI synchronous event. The parameter value is the maximum time, in milliseconds, the thread will block. When this time expires, the file operation returns **ENOTREADY**, and the event continues asynchronously. The NFS server is expected to repeatedly retry the operation, which eventually will find the response of the original event and continue. This mechanism applies only to read, write, and truncate events, and only when such events come from NFS server threads.

The timeout value is given in milliseconds. The value 0 indicates immediate timeout (fully asynchronous event). A value greater than or equal to 86400000 (which is 24 hours) is considered 'infinity' (no timeout, fully synchronous event). The default value is 86400000. See also "Support for NFS" on page 633.

| For the parameter change to take effect, restart the GPFS daemon on all nodes.

#### **dapiFileHandleSize**

Controls the size of file handles generated by GPFS. The default DAPI file handle size is 32 bytes. For clusters created prior to GPFS 3.2, the default DAPI file handle size is 16 bytes.

**Note:** To change the DMAP API file handle size, GPFS must be stopped on all nodes in the cluster.

#### dmapiMountEvent

Controls the generation of the **mount**, **preunmount**, and **unmount** events. Valid values are:

- all**      Specifies that **mount**, **preunmount**, and **unmount** events are generated on each node.  
This is the default behavior.

#### LocalNode

Specifies that **mount**, **preunmount**, and **unmount** events are generated only if the node is a session node.

#### SessionNode

Specifies that **mount**, **preunmount**, and **unmount** events are generated on each node and are delivered to the session node, but the session node will respond with **DM\_RESP\_CONTINUE** to the event node without delivering the event to the DMAP API application, unless the event is originated from the **SessionNode** itself.

#### dmapiMountTimeout

Controls the blocking of mount operations, waiting for a disposition for the mount event to be set. This timeout is activated at most once on each node, by the first mount of a file system which has DMAP API enabled, and only if there has never before been a mount disposition. Any mount operation on this node that starts while the timeout period is active will wait for the mount disposition. The parameter value is the maximum time, in seconds, that the mount operation will wait for a disposition. When this time expires and there still is no disposition for the mount event, the mount operation fails, returning the **EIO** error.

The timeout value is given in full seconds. The value 0 indicates immediate timeout (immediate failure of the mount operation). A value greater than or equal to 86400 (which is 24 hours) is considered 'infinity' (no timeout, indefinite blocking until there is a disposition). The default value is 60. See also "Mount and unmount" on page 630 and "Initializing the Data Management application" on page 637.

#### dmapiSessionFailureTimeout

Controls the blocking of file operation threads, while in the kernel, waiting for the handling of a DMAP API synchronous event that is enqueued on a session that has suffered a failure. The parameter value is the maximum time, in seconds, the thread will wait for the recovery of the failed session. When this time expires and the session has not yet recovered, the event is aborted and the file operation fails, returning the **EIO** error.

The timeout value is given in full seconds. The value 0 indicates immediate timeout (immediate failure of the file operation). A value greater than or equal to 86400 (which is 24 hours) is considered 'infinity' (no timeout, indefinite blocking until the session recovers). The default value is 0. See also "Failure and recovery of IBM Spectrum Scale Data Management API for GPFS" on page 657 for details on session failure and recovery.

## Enabling DMAP API for a file system

DMAP API must be enabled individually for each file system.

DMAP API can be enabled for a file system when the file system is created, using the **-z yes** option on the **mmcrfs** command. The default is **-z no**. The setting can be changed when the file system is not mounted anywhere, using the **-z yes | no** option on the **mmchfs** command. The setting is persistent.

The current setting can be queried using the **-z** option on the **mmlsfs** command.

While DMAP API is disabled for a given file system, no events are generated by file operations of that file system. Any DMAP API function calls referencing that file system fail with an **EPERM** error.

When **mmchfs -z no** is used to disable DMAPI, existing event lists, extended attributes, and managed regions in the given file system remain defined, but will be ignored until DAPI is re-enabled. The command **mmchfs -z no** should be used with caution, since punched holes, if any, are no longer protected by managed regions.

For more information about GPFS commands, see Chapter 1, “Command reference,” on page 1.

## Initializing the Data Management application

All DAPI APIs must be called from nodes that are in the cluster where the file system is created. DAPI APIs may **not** be invoked from a remote cluster.

During initialization of GPFS, it is necessary to synchronize the GPFS daemon and the DM application to prevent mount operations from failing. There are two mechanisms to accomplish this:

1. The shell script **gpfsready** invoked by the GPFS daemon during initialization.
2. A timeout interval, allowing mount operations to wait for a disposition to be set for the mount event.

During GPFS initialization, the daemon invokes the shell script **gpfsready**, located in directory **/var/mmfs/etc**. This occurs as the file systems are starting to be mounted. The shell script can be modified to start or restart the DM application. Upon return from this script, a session should have been created and a disposition set for the mount event. Otherwise, mount operations may fail due to a lack of disposition.

In a multiple-node environment such as GPFS, usually only a small subset of the nodes are session nodes, having DM applications running locally. On a node that is not a session node, the **gpfsready** script can be modified to synchronize between the local GPFS daemon and a remote DM application. This will prevent mount from failing on any node.

A sample shell script **gpfsready.sample** is installed in directory **/usr/lpp/mmfs/samples**.

If no mount disposition has ever been set in the cluster, the first external mount of a DAPI-enabled file system on each node will activate a timeout interval on that node. Any mount operation on that node that starts during the timeout interval will wait for the mount disposition until the timeout expires. The timeout interval is configurable using the **dmapimounttimeout** configuration attribute on the **mmchconfig** command (the interval can even be made infinite). A message is displayed at the beginning of the wait. If there is still no disposition for the mount event when the timeout expires, the mount operation will fail with an **EIO** error code. See “GPFS configuration attributes for DAPI” on page 635 for more information on **dmapimounttimeout**.

---

## Specifications of enhancements for IBM Spectrum Scale Data Management API for GPFS

DAPI for GPFS provides numerous enhancements in data structures and functions.

These enhancements are provided mainly by the multiple-node environment. Some data structures have additional fields. Many functions have usage restrictions, changes in semantics, and additional error codes. The enhancements are in these areas:

- “Enhancements to data structures” on page 638
- “Usage restrictions on DAPI functions” on page 639
- “Definitions for GPFS-specific DAPI functions” on page 641
- “Semantic changes to DAPI functions” on page 654
- “GPFS-specific DAPI events” on page 655
- “Additional error codes returned by DAPI functions” on page 656

## Enhancements to data structures

This is a description of GPFS enhancements to data structures defined in the XDSM standard.

For complete C declarations of all the data structures that are used in DMAPI for GPFS, refer to the **dapi\_types.h** file located in the **/usr/lpp/mmfs/include** directory.

- All file offsets and sizes in DAPI data structures are 64 bits long.
- Names or path names that are passed in event messages are character strings, terminated by a null character. The length of the name buffer, as specified in the **dm\_vardata\_t** structure, includes the null character.
- The **dm\_region\_t** structure has a new 4-byte field, **rg\_opaque**. The DAPI implementation does not interpret **rg\_opaque**. The DM application can use this field to store additional information within the managed region.
- The **dt\_change** field in the **dm\_stat** structure is not implemented in the inode. The value will change each time it is returned by the **dm\_get\_fileattr** function.
- The **dt\_dtime** field in the **dm\_stat** structure is overloaded on the **dt\_ctime** field.
- The **dm\_eventmsg** structure has a 4 byte field, **ev\_nodeid** that uniquely identifies the node that generated the event. The id is the GPFS cluster data node number, which is attribute **node\_number** in the **mmsdrfs2** file for a PSSP node or **mmsdrfs** file for any other type of node.
- The **ne\_mode** field in the **dm\_namesp\_event** structure has an additional flag, **DM\_LOCAL\_MOUNT**. For the events preunmount and unmount when this flag is set, the unmount operation is local to the session node. See “Mount and unmount” on page 630. The **me\_mode** field in the **dm\_mount\_event** structure has two additional flags; **DM\_LOCAL\_MOUNT**, and **DM\_REMOTE\_MOUNT**. See “Mount and unmount” on page 630.
- There are two ‘quick access’ single-bit opaque DM attributes for each file, stored directly in the inode. See “Data Management attributes” on page 633.
- The data type **dm\_eventset\_t** is implemented as a bit map, containing one bit for each event that is defined in DAPI. The bit is set if, and only if, the event is present.

Variables of type **dm\_eventset\_t** should be manipulated only using special macros. The XDSM standard provides a basic set of such macros. GPFS provides a number of additional macros. The names of all such macros begin with the prefix **DMEV\_**.

This is the list of additional macros that are provided in DAPI for GPFS:

**DMEV\_ALL(eset)**

Add all events to **eset**

**DMEV\_ISZERO(eset)**

Check if **eset** is empty

**DMEV\_ISALL(eset)**

Check if **eset** contains all events

**DMEV\_ADD(eset1, eset2)**

Add to **eset2** all events in **eset1**

**DMEV\_Rem(eset1, eset2)**

Remove from **eset2** all events in **eset1**

**DMEV\_RES(eset1, eset2)**

Restrict **eset2** by **eset1**

**DMEV\_ISEQ(eset1, eset2)**

Check if **eset1** and **eset2** are equal

**DMEV\_ISDISJ(eset1, eset2)**

Check if **eset1** and **eset2** are disjoint

**DMEV\_ISSUB(eset2)**

Check if **eset1** is a subset of **eset2**

**DMEV\_NORM(eset)**

Normalize the internal format of **eset**, clearing all unused bits

- DMAPI for GPFS provides a set of macros for comparison of token ids (value of type **dm\_token\_t**).

**DM\_TOKEN\_EQ (x,y)**

Check if **x** and **y** are the same

**DM\_TOKEN\_NE (x,y)**

Check if **x** and **y** are different

**DM\_TOKEN\_LT (x,y)**

Check if **x** is less than **y**

**DM\_TOKEN\_GT (x,y)**

Check if **x** is greater than **y**

**DM\_TOKEN\_LE (x,y)**

Check if **x** is less than or equal to **y**

**DM\_TOKEN\_GE (x,y)**

Check if **x** is greater than or equal to **y**

## Usage restrictions on DAPI functions

There are usage restrictions on the DAPI for GPFS functions.

- The maximum number of DAPI sessions that can be created on a node is 4000.
- Root credentials are a prerequisite for invoking any DAPI function, otherwise the function fails with an **EPERM** error code.
- DAPI functions are unable to run if the GPFS kernel extension is not loaded, or if the runtime module **dmapicalls** is not installed. An **ENOSYS** error code is returned in this case.
- Invoking a DAPI function that is not implemented in GPFS results in returning the **ENOSYS** error code.
- DAPI functions will fail, with the **ENOTREADY** error code, if the local GPFS daemon is not running.
- DAPI functions will fail, with the **EPERM** error code, if DAPI is disabled for the file system that is referenced by the file handle argument.
- DAPI functions cannot access GPFS reserved files, such as quota files, inode allocation maps, and so forth. The **EBADF** error code is returned in this case.
- GPFS does not support access rights on entire file systems (as opposed to individual files). Hence, DAPI function calls that reference a file system (with a file system handle) cannot present a token, and must use **DM\_NO\_TOKEN**. Functions affected by this restriction are:

- **dm\_set\_eventlist**
- **dm\_get\_eventlist**
- **dm\_set\_disp**
- **dm\_get\_mountinfo**
- **dm\_set\_return\_on\_destroy**
- **dm\_get\_bulkattr**
- **dm\_get\_bulkall**

If a token is presented, these functions fail with the **EINVAL** error code.

- DAPI functions that acquire, change, query, or release access rights, must not present a file system handle. These functions are:
  - **dm\_request\_right**
  - **dm\_upgrade\_right**

- **dm\_downgrade\_right**
- **dm\_release\_right**
- **dm\_query\_right**

If a file system handle is presented, these functions fail with the **EINVAL** error code.

- The function **dm\_request\_right**, when invoked without wait (the *flags* argument has a value of 0), will almost always fail with the **EAGAIN** error. A GPFS implementation constraint prevents this function from completing successfully without wait, even if it is known that the requested access right is available. The **DM\_RR\_WAIT** flag must always be used. If the access right is available, there will be no noticeable delay.
- DMAPI function calls that reference a specific token, either as input or as output, can be made only on the session node. Otherwise, the call fails with the **EINVAL** error code.
- DAPI function calls that reference an individual file by handle must be made on the session node. The corresponding file system must be mounted on the session node. The call fails with **EINVAL** if it is not on the session node, and with **EBADF** if the file system is not mounted.
- DAPI function calls that reference a file system by handle (as opposed to an individual file) can be made on any node, not just the session node. The relevant functions are:
  - **dm\_set\_eventlist**
  - **dm\_get\_eventlist**
  - **dm\_set\_disp**
  - **dm\_get\_mountinfo**
  - **dm\_set\_return\_on\_destroy**
  - **dm\_get\_bulkattr**
  - **dm\_get\_bulkall**

For **dm\_get\_bulkattr** and **dm\_get\_bulkall**, the system file must be mounted on the node that is making the call. For the other functions, the file system must be mounted on some node, but not necessarily on the node that is making the call. As specified previously, all such function calls must use **DM\_NO\_TOKEN**. The function fails with the **EBADF** error code if the file system is not mounted as required.

- The function **dm\_punch\_hole** will fail with the **EBUSY** error code if the file to be punched is currently memory-mapped.
- The function **dm\_move\_event** can only be used when the source session and the target session are on the same node. The function must be called on the session node. Otherwise, the function fails with the **EINVAL** error code.
- The function **dm\_create\_session**, when providing an existing session id in the argument *oldsid*, can only be called on the session node, except after session node failure. Otherwise, the call will return the **EINVAL** error code.
- The function **dm\_destroy\_session** can only be called on the session node, otherwise the call will fail with the **EINVAL** error code.
- The function **dm\_set\_fileattr** cannot change the file size. If the **dm\_at\_size** bit in the attribute mask is set, the call fails with the **EINVAL** error code.
- DAPI functions that reference an event with a token fail with the **ESRCH** error code, if the event is not in an outstanding state. This is related to session recovery. See “Failure and recovery of IBM Spectrum Scale Data Management API for GPFS” on page 657 for details on session failure and recovery.

For additional information about:

- Semantic changes to the DAPI for GPFS functions, see “Semantic changes to DAPI functions” on page 654.
- C declarations of all functions in DAPI for GPFS, refer to the **dmapi.h** file located in the **/usr/lpp/mmfs/include** directory.

## Definitions for GPFS-specific DMAPI functions

The GPFS-specific DMAPI functions are not part of the DAPI open standard.

You can use the following GPFS-specific DAPI functions to work with file system snapshots:

- “[dm\\_handle\\_to\\_snap](#)” on page 642
- “[dm\\_make\\_xhandle](#)” on page 643

You can use the following GPFS-specific DAPI functions to make asynchronous updates to attributes, managed regions, and event lists on files:

- “[dm\\_remove\\_dmatr\\_nosync](#)” on page 645
- “[dm\\_set\\_dmatr\\_nosync](#)” on page 647
- “[dm\\_set\\_eventlist\\_nosync](#)” on page 649
- “[dm\\_set\\_region\\_nosync](#)” on page 651

You can use the following GPFS-specific DAPI function to make the previously listed asynchronous updates persistent by flushing them to disk:

- “[dm\\_sync\\_dmatr\\_by\\_handle](#)” on page 653

## **dm\_handle\_to\_snap**

Extracts a snapshot ID from a handle.

### **Synopsis**

```
int dm_handle_to_snap(
 void *hanp, /* IN */
 size_t hlen, /* IN */
 dm_snap_t *isnapp /* OUT */
);
```

### **Description**

Use the **dm\_handle\_to\_snap** function to extract a snapshot ID from a handle. **dm\_handle\_to\_snap()** is a GPFS-specific DMAPI function. It is not part of the open standard.

### **Parameters**

#### **void \*hanp (IN)**

A pointer to an opaque DM handle previously returned by DMAPI.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_snap\_t \*isnapp (OUT)**

A pointer to the snapshot ID.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

#### **[EBADF]**

The file handle does not refer to an existing or accessible object.

#### **[EFAULT]**

The system detected an invalid address in attempting to use an argument.

#### **[EINVAL]**

The argument *token* is not a valid token.

#### **[ENOMEM]**

DMAPI could not obtain the required resources to complete the call.

#### **[ENOSYS]**

Function is not supported by the DM implementation.

#### **[EPERM]**

The caller does not hold the appropriate privilege.

### **See also**

“dm\_make\_xhandle” on page 643

## **dm\_make\_xhandle**

Converts a file system ID, inode number, inode generation count, and snapshot ID into a handle.

### **Synopsis**

```
int dm_make_xhandle(
 dm_fsid_t *fsidp, /* IN */
 dm_ino_t *inop, /* IN */
 dm_igen_t *igenp, /* IN */
 dm_snap_t *isnapp, /* IN */
 void **hanpp, /* OUT */
 size_t *hlenp) /* OUT */
```

### **Description**

Use the **dm\_make\_xhandle()** function to convert a file system ID, inode number, inode generation count, and snapshot ID into a handle. **dm\_make\_xhandle()** is a GPFS-specific DMAP API function. It is not part of the open standard.

### **Parameters**

**dm\_fsid\_t \*fsidp (IN)**

The file system ID.

**dm\_ino\_t \*inop (IN)**

The inode number.

**dm\_igen\_t \*igenp (IN)**

The inode generation count.

**dm\_snap\_t \*isnapp (IN)**

The snapshot ID.

**void \*\*hanpp (OUT)**

A DMAP API initialized pointer that identifies a region of memory containing an opaque DM handle.

The caller is responsible for freeing the allocated memory.

**size\_t \*hlenp (OUT)**

The length of the handle in bytes.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

**[EBADF]**

The file handle does not refer to an existing or accessible object.

**[EFAULT]**

The system detected an invalid address in attempting to use an argument.

**[EINVAL]**

The argument *token* is not a valid token.

**[ENOMEM]**

DM API could not obtain the required resources to complete the call.

**[ENOSYS]**

Function is not supported by the DM implementation.

**[EPERM]**

The caller does not hold the appropriate privilege.

## **See also**

“dm\_handle\_to\_snap” on page 642

## **dm\_remove\_dmatr\_nosync**

Asynchronously removes the specified attribute.

### **Synopsis**

```
int dm_remove_dmatr_nosync(
 dm_sessid_t sid,
 void *hanp,
 size_t hlen,
 dm_token_t token,
 int setdtime,
 dm_attrname_t *attrnamep
);
```

### **Description**

Use the **dm\_remove\_dmatr\_nosync** function to asynchronously remove the attribute specified by *attrname*.

**dm\_remove\_dmatr\_nosync** is a GPFS-specific DMAPI function; it is not part of the open standard. It has the same purpose, parameters, and return values as the standard DMAPI **dm\_remove\_dmatr** function, except that the update that it performs is not persistent until some other activity on that file (or on other files in the file system) happens to flush it to disk. To be certain that your update is made persistent, use one of the following functions:

- Standard DMAPI **dm\_sync\_by\_handle** function, which flushes the file data and attributes
- GPFS-specific **dm\_sync\_dmatr\_by\_handle** function, which flushes only the attributes.

### **Parameters**

#### **dm\_sessid\_t sid (IN)**

The identifier for the session of interest.

#### **void \*hanp (IN)**

The handle for the file for which the attributes should be removed.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_token\_t \*token (IN)**

The token referencing the access right for the handle. The access right must be **DM\_RIGHT\_EXCL**, or the token **DM\_NO\_TOKEN** may be used and the interface acquires the appropriate rights.

#### **int setdtime (IN)**

If *setdtime* is non-zero, updates the file's attribute time stamp.

#### **dm\_attrname\_t \*attrnamep (IN)**

The attribute to be removed.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

#### **[EACCES]**

The access right referenced by the token for the handle is not **DM\_RIGHT\_EXCL**.

#### **[EBADF]**

The file handle does not refer to an existing or accessible object.

#### **[EFAULT]**

The system detected an invalid address in attempting to use an argument.

**[EINVAL]**

The argument *token* is not a valid token.

**[EINVAL]**

The session is not valid.

**[EIO]** I/O error resulted in failure of operation.**[ENOSYS]**

The DMAPI implementation does not support this optional function.

**[EPERM]**

The caller does not hold the appropriate privilege.

**[EROFS]**

The operation is not allowed on a read-only file system.

**See also**

[“dm\\_set\\_dmatr\\_nosync” on page 647](#), [“dm\\_sync\\_dmatr\\_by\\_handle” on page 653](#)

## **dm\_set\_dmatr\_nosync**

Asynchronously creates or replaces the value of the named attribute with the specified data.

### **Synopsis**

```
int dm_set_dmatr_nosync(
 dm_sessid_t sid,
 void *hanp,
 size_t hlen,
 dm_token_t token,
 dm_attrname_t *attrnamep,
 int setftime,
 size_t buflen,
 void *bufp
);
```

### **Description**

Use the **dm\_set\_dmatr\_nosync** function to asynchronously create or replace the value of the named attribute with the specified data.

**dm\_set\_dmatr\_nosync** is a GPFS-specific DMAPI function; it is not part of the open standard. It has the same purpose, parameters, and return values as the standard DMAPI **dm\_set\_dmatr** function, except that the update that it performs is not persistent until some other activity on that file (or on other files in the file system) happens to flush it to disk. To be certain that your update is made persistent, use one of the following functions:

- Standard DMAPI **dm\_sync\_by\_handle** function, which flushes the file data and attributes
- GPFS-specific **dm\_sync\_dmatr\_by\_handle** function, which flushes only the attributes.

### **Parameters**

#### **dm\_sessid\_t sid (IN)**

The identifier for the session of interest.

#### **void \*hanp (IN)**

The handle for the file for which the attributes should be created or replaced.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_token\_t \*token (IN)**

The token referencing the access right for the handle. The access right must be **DM\_RIGHT\_EXCL**, or the token **DM\_NO\_TOKEN** may be used and the interface acquires the appropriate rights.

#### **dm\_attrname\_t \*attrnamep (IN)**

The attribute to be created or replaced.

#### **int setftime (IN)**

If *setftime* is non-zero, updates the file's attribute time stamp.

#### **size\_t buflen (IN)**

The size of the buffer in bytes.

#### **void \*bufp (IN)**

The buffer containing the attribute data.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

**[E2BIG]**

The attribute value exceeds one of the implementation defined storage limits.

**[E2BIG]**

*buflen* is larger than the implementation defined limit. The limit can be determined by calling the `dm_get_config()` function.

**[EACCES]**

The access right referenced by the token for the handle is not `DM_RIGHT_EXCL`.

**[EBADF]**

The file handle does not refer to an existing or accessible object.

**[EFAULT]**

The system detected an invalid address in attempting to use an argument.

**[EIO]** An attempt to write the new or updated attribute resulted in an I/O error.**[EINVAL]**

The argument *token* is not a valid token.

**[EINVAL]**

The session is not valid.

**[ENOMEM]**

The DMAPI could not acquire the required resources to complete the call.

**[ENOSPC]**

An attempt to write the new or updated attribute resulted in an error due to no free space being available on the device.

**[ENOSYS]**

The DMAPI implementation does not support this optional function.

**[EPERM]**

The caller does not hold the appropriate privilege.

**[EROFS]**

The operation is not allowed on a read-only file system.

**See also**

“`dm_remove_dmattrib_nosync`” on page 645, “`dm_sync_dmattrib_by_handle`” on page 653

## **dm\_set\_eventlist\_nosync**

Asynchronously sets the list of events to be enabled for an object.

### **Synopsis**

```
int dm_set_eventlist_nosync(
 dm_sessid_t sid,
 void *hanp,
 size_t hlen,
 dm_token_t token,
 dm_eventset_t *eventsetp,
 u_int maxevent
);
```

### **Description**

Use the **dm\_set\_eventlist\_nosync** function to asynchronously set the list of events to be enabled for an object.

**dm\_set\_eventlist\_nosync** is a GPFS-specific DMAPI function; it is not part of the open standard. It has the same purpose, parameters, and return values as the standard DMAPI **dm\_set\_eventlist** function, except that the update that it performs is not persistent until some other activity on that file (or on other files in the file system) happens to flush it to disk. To be certain that your update is made persistent, use one of the following functions:

- Standard DMAPI **dm\_sync\_by\_handle** function, which flushes the file data and attributes
- GPFS-specific **dm\_sync\_dmatr\_by\_handle** function, which flushes only the attributes.

### **Parameters**

#### **dm\_sessid\_t sid (IN)**

The identifier for the session of interest.

#### **void \*hanp (IN)**

The handle for the object. The handle can be either the system handle or a file handle.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_token\_t \*token (IN)**

The token referencing the access right for the handle. The access right must be **DM\_RIGHT\_EXCL**, or the token **DM\_NO\_TOKEN** may be used and the interface acquires the appropriate rights.

#### **dm\_eventset\_t \*eventsetp (IN)**

The list of events to be enabled for the object.

#### **u\_int maxevent (IN)**

The number of events to be checked for dispositions in the event set. The events from 0 to *maxevent-1* are examined.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

#### **[EACCES]**

The access right referenced by the token for the handle is not **DM\_RIGHT\_EXCL**.

#### **[EBADF]**

The file handle does not refer to an existing or accessible object.

#### **[EFAULT]**

The system detected an invalid address in attempting to use an argument.

**[EINVAL]**

The argument *token* is not a valid token.

**[EINVAL]**

The session is not valid.

**[EINVAL]**

Tried to set event on a global handle.

**[ENOMEM]**

The DMAPI could not acquire the required resources to complete the call.

**[ENXIO]**

The implementation of the DAPI does not support enabling event delivery on the specified handle.

**[EPERM]**

The caller does not hold the appropriate privilege.

**[EROFS]**

The operation is not allowed on a read-only file system.

**See also**

“dm\_sync\_dmatr\_by\_handle” on page 653

## **dm\_set\_region\_nosync**

Asynchronously replaces the set of managed regions for a file.

### **Synopsis**

```
int dm_set_region_nosync(
 dm_sessid_t sid,
 void *hanp,
 size_t hlen,
 dm_token_t token,
 u_int nelem,
 dm_region_t *regbufp,
 dm_boolean_t *exactflagp
);
```

### **Description**

Use the **dm\_set\_region\_nosync** function to asynchronously replace the set of managed regions for a file.

**dm\_set\_region\_nosync** is a GPFS-specific DMAPI function; it is not part of the open standard. It has the same purpose, parameters, and return values as the standard DMAP API **dm\_set\_region** function, except that the update that it performs is not persistent until some other activity on that file (or on other files in the file system) happens to flush it to disk. To be certain that your update is made persistent, use one of the following functions:

- Standard DMAP API **dm\_sync\_by\_handle** function, which flushes the file data and attributes
- GPFS-specific **dm\_sync\_dmatr\_by\_handle** function, which flushes only the attributes.

### **Parameters**

#### **dm\_sessid\_t sid (IN)**

The identifier for the session of interest.

#### **void \*hanp (IN)**

The handle for the regular file to be affected.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_token\_t \*token (IN)**

The token referencing the access right for the handle. The access right must be **DM\_RIGHT\_EXCL**, or the token **DM\_NO\_TOKEN** may be used and the interface acquires the appropriate rights.

#### **u\_int nelem (IN)**

The number of input regions in *regbufp*. If *nelem* is 0, then all existing managed regions are cleared.

#### **dm\_region\_t \*regbufp (IN)**

A pointer to the structure defining the regions to be set. May be NULL if *nelem* is zero.

#### **dm\_boolean\_t \*exactflagp (OUT)**

If **DM\_TRUE**, the file system did not alter the requested managed region set.

Valid values for the *rg\_flags* field of the region structure are created by OR'ing together one or more of the following values:

##### **DM\_REGION\_READ**

Enable synchronous event for read operations that overlap this managed region.

##### **DM\_REGION\_WRITE**

Enable synchronous event for write operations that overlap this managed region.

##### **DM\_REGION\_TRUNCATE**

Enable synchronous event for truncate operations that overlap this managed region.

**DM\_REGION\_NOEVENT**

Do not generate any events for this managed region.

**Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

**[E2BIG]**

The number of regions specified by nelem exceeded the implementation capacity.

**[EACCES]**

The access right referenced by the token for the handle is not **DM\_RIGHT\_EXCL**.

**[EBADF]**

The file handle does not refer to an existing or accessible object.

**[EFAULT]**

The system detected an invalid address in attempting to use an argument.

**[EINVAL]**

The argument *token* is not a valid token.

**[EINVAL]**

The file handle does not refer to a regular file.

**[EINVAL]**

The regions passed in are not valid because they overlap or some other problem.

**[EINVAL]**

The session is not valid.

**[EIO]** An I/O error resulted in failure of operation.**[ENOMEM]**

The DMAPI could not acquire the required resources to complete the call.

**[EPERM]**

The caller does not hold the appropriate privilege.

**[EROFS]**

The operation is not allowed on a read-only file system.

**See also**

“dm\_sync\_dmatr\_by\_handle” on page 653

## **dm\_sync\_dmattr\_by\_handle**

Synchronizes one or more files' in-memory attributes with those on the physical medium.

### **Synopsis**

```
int m_sync_dmattr_by_handle(
 dm_sessid_t sid,
 void *hamp,
 size_t hlen,
 dm_token_t token
);
```

### **Description**

Use the **dm\_sync\_dmattr\_by\_handle** function to synchronize one or more files' in-memory attributes with those on the physical medium.

**dm\_sync\_dmattr\_by\_handle** is a GPFS-specific DMAPI function; it is not part of the open standard. It has the same purpose, parameters, and return values as the standard DAPI **dm\_sync\_by\_handle** function, except that it flushes only the attributes, not the file data.

Like **dm\_sync\_by\_handle**, **dm\_sync\_dmattr\_by\_handle** commits all previously unsynchronized updates for that node, not just the updates for one file. Therefore, if you update a list of files and call **dm\_sync\_dmattr\_by\_handle** on the last file, the attribute updates to all of the files in the list are made persistent.

### **Parameters**

#### **dm\_sessid\_t sid (IN)**

The identifier for the session of interest.

#### **void \*hamp (IN)**

The handle for the file whose attributes are to be synchronized.

#### **size\_t hlen (IN)**

The length of the handle in bytes.

#### **dm\_token\_t \*token (IN)**

The token referencing the access right for the handle. The access right must be **DM\_RIGHT\_EXCL**, or the token **DM\_NO\_TOKEN** may be used and the interface acquires the appropriate rights.

### **Return values**

Zero is returned on success. On error, -1 is returned, and the global *errno* is set to one of the following values:

#### **[EACCES]**

The access right referenced by the token for the handle is not **DM\_RIGHT\_EXCL**.

#### **[EBADF]**

The file handle does not refer to an existing or accessible object.

#### **[EFAULT]**

The system detected an invalid address in attempting to use an argument.

#### **[EINVAL]**

The argument *token* is not a valid token.

#### **[ENOMEM]**

The DAPI could not acquire the required resources to complete the call.

#### [ENOSYS]

The DMAPI implementation does not support this optional function.

#### [EPERM]

The caller does not hold the appropriate privilege.

#### See also

"`dm_remove_dmatr_nosync`" on page 645, "`dm_set_dmatr_nosync`" on page 647, "`dm_set_eventlist_nosync`" on page 649, and "`dm_set_region_nosync`" on page 651

## Semantic changes to DMAPI functions

There are semantic changes to functions in DAPI for GPFS. These changes are entailed mostly by the multiple-node environment.

For a list of additional error codes that are used in DAPI for GPFS, see "Additional error codes returned by DAPI functions" on page 656. For C declarations of all the DAPI for GPFS functions, refer to the `dmapi.h` file located in the `/usr/lpp/mmfs/include` directory.

- The following DAPI functions can be invoked on any node, not just the session node, as long as the session exists on some node in the GPFS cluster.
  - `dm_getall_disp`
  - `dm_query_session`
  - `dm_send_msg`
- DAPI functions that reference a file system, as opposed to an individual file, can be made on any node, not just the session node. Being able to call certain functions on any node has advantages. The DM application can establish event monitoring when receiving a mount event from any node. Also, a distributed DM application can change event lists and dispositions of any file system from any node.
  - `dm_set_eventlist`
  - `dm_get_eventlist`
  - `dm_set_disp`
  - `dm_get_mount_info`
  - `dm_set_return_on_destroy`
  - `dm_get_bulkattr`
  - `dm_get_bulkall`
- The following functions, that construct a handle from its components, do not check if the resulting handle references a valid file. Validity is checked when the handle is presented in function calls that actually reference the file.
  - `dm_make_handle`
  - `dm_make_fshandle`
  - `dm_make_xhandle`
- The following data movement functions may be invoked on any node within the GPFS cluster, provided they are run as root and present a session ID for an established session on the session node. For guidelines on how to perform data movement from multiple nodes, see "Parallelism in Data Management applications" on page 632.
  - `dm_read_invis`
  - `dm_write_invis`
  - `dm_probe_hole`
  - `dm_punch_hole`

- The following functions that extract components of the handle, do not check whether the specified handle references a valid file. Validity is checked when the handle is presented in function calls that actually reference the file.
  - `dm_handle_to_fsid`
  - `dm_handle_to_igen`
  - `dm_handle_to_ino`
  - `dm_handle_to_snap`
- `dm_handle_to_fshandle` converts a file handle to a file system handle without checking the validity of either handle.
- `dm_handle_is_valid` does not check if the handle references a valid file. It verifies only that the internal format of the handle is correct.
- `dm_init_attrloc` ignores all of its arguments, except the output argument *locp*. In DMAP API for GPFS, the location pointer is initialized to a constant. Validation of the session, token, and handle arguments is done by the bulk access functions.
- When `dm_query_session` is called on a node other than the session node, it returns only the first eight bytes of the session information string.
- `dm_create_session` can be used to move an existing session to another node, if the current session node has failed. The call must be made on the new session node. See “Failure and recovery of IBM Spectrum Scale Data Management API for GPFS” on page 657 for details on session node failure and recovery.
- Assuming an existing session, using `dm_create_session` does not change the session id. If the argument *sessinfop* is `NULL`, the session information string is not changed.
- The argument *maxevent* in the functions `dm_set_disp` and `dm_set_eventlist` is ignored. In GPFS the set of events is implemented as a bitmap, containing a bit for each possible event.
- The value pointed to by the argument *nelemp*, on return from the functions `dm_get_eventlist` and `dm_get_config_events`, is always `DM_EVENT_MAX-1`. The argument *nelem* in these functions is ignored.
- The *dt\_nevents* field in the `dm_stat_t` structure, which is returned by the `dm_get_fileattr` and `dm_get_bulkall` functions, has a value of `DM_EVENT_MAX-1` when the file has a file-system-wide event enabled by calling the `dm_set_eventlist` function. The value will always be 3 when there is no file-system-wide event enabled. A value of 3 indicates that there could be a managed region enabled for the specific file, which might have enabled a maximum of three events: READ, WRITE, and TRUNCATE.
- The functions `dm_get_config` and `dm_get_config_events` ignore the arguments *hanp* and *hlen*. This is because the configuration is not dependent on the specific file or file system.
- The function `dm_set_disp`, when called with the global handle, ignores any events in the event set being presented, except the mount event. When `dm_set_disp` is called with a file system handle, it ignores the mount event.
- The function `dm_handle_hash`, when called with an individual file handle, returns the inode number of the file. When `dm_handle_hash` is called with a file system handle, it returns the value 0.
- The function `dm_get_mountinfo` returns two additional flags in the `me_mode` field in the `dm_mount_event` structure. The flags are `DM_MOUNT_LOCAL` and `DM_MOUNT_REMOTE`. See “Mount and unmount” on page 630 for details.

## GPFS-specific DMAPI events

The GPFS-specific events are not part of the DMAPI open standard. You can use these GPFS events to filter out events that are not critical to file management and to prevent system overloads from trivial information.

The DMAPI standard specifies that the system must generate ATTRIBUTE events each time the "changed time" (`ctime`) attribute for a file changes. For systems that write files in parallel, like GPFS, this generates ATTRIBUTE events from every node writing to the file. Consequently, it is easy for ATTRIBUTE events to

overwhelm a data management server. However, the only **ctime** changes that are critical to GPFS are changes to either the permissions or ACLs of a file. In most cases, GPFS can ignore other **ctime** changes.

To distinguish file permission and ACL changes from other **ctime** updates, the following DMAPI metadata attribute events allow GPFS to filter **ctime** updates. Using these events, DM servers are able to track file permission changes without overwhelming the system with irrelevant ATTRIBUTE events. However, these events are not part of the CAE Specification C429 open standard and they were implemented specifically for GPFS systems.

## Metadata Events

### DM\_EVENT\_PREPERMCHANGE

Pre-permission change event. Event is triggered before file permission change.

### DM\_EVENT\_POSTPERMCHANGE

Post-permission change event. Event is triggered after file permission change.

**Note:** If you only want to track permission and ACL changes, turn off the **DM\_EVENT\_ATTRIBUTE** and turn on both the **DM\_EVENT\_PREPERMCHANGE** and **DM\_EVENT\_POSTPERMCHANGE** events.

## Additional error codes returned by DAPI functions

DMAPI for GPFS uses additional error codes, not specified in the XDSM standard, for most DAPI functions.

For C declarations of all the DAPI for GPFS functions, refer to the **dmapi.h** file located in the **/usr/lpp/mmfs/include** directory.

For all DAPI functions, these error codes are used:

### ENOSYS

The GPFS kernel extension is not loaded, or the runtime module **dmapicalls** is not installed.

### ENOSYS

An attempt has been made to invoke a DAPI function that is not implemented in GPFS.

### ENOTREADY

The local GPFS daemon is not running or is initializing.

### ENOMEM

DAPI could not acquire the required resources to complete the call. **ENOMEM** is defined in the XDSM standard for some DAPI functions, but not for all.

### ESTALE

An error has occurred which does not fit any other error code specified for this function.

For DAPI functions that provide a file handle as an input argument, these error codes are used:

### EINVAL

The format of the file handle is not valid.

This error is returned without attempting to locate any object that is referenced by the handle.

The **EINVAL** error code is to be distinguished from the **EBADF** error code, which, as specified in the XDSM standard, indicates that the object does not exist or is inaccessible. Thus, GPFS provides a refinement, distinguishing between format and access errors related to handles.

### EPERM

DAPI is disabled for the file system that is referenced by the file handle.

For DAPI functions that provide a token as an input argument, these error codes are used:

## ESRCH

The event referenced by the token is not in outstanding state.

This is to be distinguished from the **EINVAL** error code, which is returned when the token itself is not valid. **ESRCH** is defined in the XDSM standard for some DMAPI functions, but not for all relevant functions. In GPFS, the **ESRCH** error code occurs mostly after recovery from session failure. See “Event recovery” on page 659 for details.

For these **specific DAPI functions**, the error code listed is used:

Table 22. Specific DAPI functions and associated error codes.

| Name of function                                                                                                                                                                                                                                                                                               | Error codes and descriptions                                                                                                                                                                            |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>dm_downgrade_right()</code><br><code>dm_upgrade_right()</code>                                                                                                                                                                                                                                           | <b>EINVAL</b> - The session or token is not valid.                                                                                                                                                      |
| <code>dm_get_region()</code>                                                                                                                                                                                                                                                                                   | <b>EPERM</b> - The caller does not hold the appropriate privilege.                                                                                                                                      |
| <code>dm_init_service()</code>                                                                                                                                                                                                                                                                                 | <b>EFAULT</b> - The system detected an invalid address in attempting to use an argument.                                                                                                                |
| <code>dm_move_event()</code><br><code>dm_respond_event()</code>                                                                                                                                                                                                                                                | <b>EINVAL</b> - The token is not valid.                                                                                                                                                                 |
| <code>dm_punch_hole()</code>                                                                                                                                                                                                                                                                                   | <b>EBUSY</b> - The file is currently memory mapped.                                                                                                                                                     |
| <code>dm_probe_hole()</code><br><code>dm_punch_hole()</code>                                                                                                                                                                                                                                                   | <b>EINVAL</b> - The argument <i>len</i> is too large, and will overflow if cast into <b>offset_t</b> .<br><b>EINVAL</b> - The argument <i>off</i> is negative.                                          |
| <code>dm_write_invis()</code>                                                                                                                                                                                                                                                                                  | <b>EINVAL</b> - The argument <i>flags</i> is not valid.                                                                                                                                                 |
| <code>dm_read_invis()</code><br><code>dm_write_invis()</code>                                                                                                                                                                                                                                                  | <b>EINVAL</b> - The argument <i>len</i> is too large, and will overflow if placed into the <b>uio_resid</b> field in the structure <b>uio</b> .<br><b>EINVAL</b> - The argument <i>off</i> is negative. |
| <code>dm_sync_by_handle()</code>                                                                                                                                                                                                                                                                               | <b>EROFS</b> - The operation is not allowed on a read-only file system.                                                                                                                                 |
| <code>dm_find_eventmsg()</code><br><code>dm_get_bulkall()</code><br><code>dm_get_bulkattr()</code><br><code>dm_get_dirattrs()</code><br><code>dm_get_events()</code><br><code>dm_get_mountinfo()</code><br><code>dm_getall_disp()</code><br><code>dm_getall_dmatr()</code><br><code>dm_handle_to_path()</code> | <b>EINVAL</b> - The argument <i>buflen</i> is too large; it must be smaller than <b>INT_MAX</b> .                                                                                                       |
| <code>dm_get_alloc_info()</code><br><code>dm_getall_sessions()</code><br><code>dm_getall_tokens()</code>                                                                                                                                                                                                       | <b>EINVAL</b> - The argument <i>nelem</i> is too large; DAPI cannot acquire sufficient resources.                                                                                                       |

## Failure and recovery of IBM Spectrum Scale Data Management API for GPFS

Failure and recovery of DAPI applications in the multiple-node GPFS environment is different than in a single-node environment.

The failure model in XDSM is intended for a single-node environment. In this model, there are two types of failures:

### DM application failure

The DM application has failed, but the file system works normally. Recovery entails restarting the DM application, which then continues handling events. Unless the DM application recovers, events may remain pending indefinitely.

## Total system failure

The file system has failed. All non-persistent DMAPI resources are lost. The DM application itself may or may not have failed. Sessions are not persistent, so recovery of events is not necessary. The file system cleans its state when it is restarted. There is no involvement of the DM application in such cleanup.

The simplistic XDSM failure model is inadequate for GPFS. In a multiple-node environment, GPFS can fail on one node, but survive on other nodes. This type of failure is called *single-node failure (or partial system failure)*. GPFS is built to survive and recover from single-node failures, without meaningfully affecting file access on surviving nodes.

Designers of Data Management applications for GPFS must comply with the enhanced DMAPI failure model, in order to support recoverability of GPFS. These areas are addressed:

- “Single-node failure”
- “Session failure and recovery” on page 659
- “Event recovery” on page 659
- “Loss of access rights” on page 660
- “DODeferred deletions” on page 660
- “DM application failure” on page 660

## Single-node failure

In DMAPI for GPFS, single-node failure means that DMAPI resources are lost on the failing node, but not on any other node.

The most common single-node failure is when the local GPFS daemon fails. This renders any GPFS file system at that node inaccessible. Another possible single-node failure is file system forced unmount. When just an individual file system is forced unmounted on some node, its resources are lost, but the sessions on that node, if any, survive.

Single-node failure has a different effect when it occurs on a session node or on a source node:

### session node failure

When the GPFS daemon fails, all session queues are lost, as well as all nonpersistent local file system resources, particularly DM access rights. The DM application may or may not have failed. The missing resources may in turn cause DMAPI function calls to fail with errors such as ENOTREADY or ESRCH.

Events generated at other source nodes remain pending despite any failure at the session node. Moreover, client threads remain blocked on such events.

### source node failure

Events generated by that node are obsolete. If such events have already been enqueued at the session node, the DM application will process them, even though this may be redundant since no client is waiting for the response.

According to the XDSM standard, sessions are not persistent. This is inadequate for GPFS. Sessions must be persistent to the extent of enabling recovery from single-node failures. This is in compliance with a basic GPFS premise that single-node failures do not affect file access on surviving nodes. Consequently, after session node failure, the session queue and the events on it must be reconstructed, possibly on another node.

Session recovery is triggered by the actions of the DM application. The scenario depends on whether or not the DM application itself has failed.

If the DM application has failed, it must be restarted, possibly on another node, and assume the old session by id. This will trigger reconstruction of the session queue and the events on it, using backup information replicated on surviving nodes. The DM application may then continue handling events. The session id is never changed when a session is assumed.

If the DM application itself survives, it will notice that the session has failed by getting certain error codes from DMAPI function calls (**ENOTREADY**, **ESRCH**). The application could then be moved to another node and recover the session queue and events on it. Alternatively, the application could wait for the GPFS daemon to recover. There is also a possibility that the daemon will recover before the DM application even notices the failure. In these cases, session reconstruction is triggered when the DM application invokes the first DMAPI function after daemon recovery.

## Session failure and recovery

A session fails when the GPFS daemon of the session node fails.

Session failure results in the loss of all DM access rights associated with events on the queue, and all the tokens become invalid. After the session has recovered, any previously outstanding synchronous events return to the initial (non-outstanding) state, and must be received again.

Session failure may also result in partial loss of the session information string. In such case, GPFS will be able to restore only the first eight characters of the session string. It is suggested to not have the DM application be dependent on more than eight characters of the session string.

In extreme situations, failure may also result in the loss of event dispositions for some file system. This happens only if the GPFS daemon fails simultaneously on all nodes where the file system was mounted. When the file system is remounted, a mount event will be generated, at which point the dispositions could be reestablished by the DM application.

During session failure, events originating from surviving nodes remain pending, and client threads remain blocked on such events. It is therefore essential that the DM application assume the old session and continue processing the pending events. To prevent indefinite blocking of clients, a mechanism has been implemented whereby pending events will be aborted and corresponding file operations failed with the **EIO** error if the failed session is not recovered within a specified time-out interval. The interval is configurable using the **dmapiSessionFailureTimeout** attribute on the **mmchconfig** command. See “GPFS configuration attributes for DMAPI” on page 635. The default is immediate timeout.

GPFS keeps the state of a failed session for 24 hours, during which the session should be assumed. When this time has elapsed, and the session has not been assumed, the session is discarded. An attempt to assume a session after it has been discarded will fail.

## Event recovery

Synchronous events are recoverable after session failure.

The state of synchronous events is maintained both at the source node and at the session node. When the old session is assumed, pending synchronous events are resubmitted by surviving source nodes.

All the events originating from the session node itself are lost during session failure, including user events generated by the DM application. All file operations on the session node fail with the **ESTALE** error code.

When a session fails, all of its tokens become obsolete. After recovery, the **dm\_getall\_tokens** function returns an empty list of tokens, and it is therefore impossible to identify events that were outstanding when the failure occurred. All recovered events return to the initial non-received state, and must be explicitly received again. The token id of a recovered event is the same as prior to the failure (except for the mount event).

If the token of a recovered event is presented in any DMAPI function before the event is explicitly received again, the call will fail with the **ESRCH** error code. The **ESRCH** error indicates that the event exists, but is not in the outstanding state. This is to be distinguished from the **EINVAL** error code, which indicates that the token id itself is not valid (there is no event).

The semantics of the **ESRCH** error code in GPFS are different from the XDSM standard. This is entailed by the enhanced failure model. The DM application may not notice that the GPFS daemon has failed and recovered, and may attempt to use a token it has received prior to the failure. For example, it may try to respond to the event. The **ESRCH** error code tells the DM application that it must receive the event again, before it can continue using the token. Any access rights associated with the token prior to the failure are lost. See “Loss of access rights.”

When a mount event is resubmitted to a session during session recovery, it will have a different token id than before the failure. This is an exception to the normal behavior, since all other recovered events have the same token id as before. The DM application thus cannot distinguish between recovered and new mount events. This should not be a problem, since the DM application must in any case be able to handle multiple mount events for the same file system.

Unmount events will not be resubmitted after session recovery. All such events are lost. This should not be a problem, since the event cannot affect the unmount operation, which has already been completed by the time the event was generated. In other words, despite being synchronous, semantically the unmount event resembles an asynchronous post event.

## **Loss of access rights**

When the GPFS daemon fails on the session node, all file systems on the node are forced unmounted. As a result, all DM access rights associated with any local session are lost.

After daemon recovery, when the old sessions are assumed and the events are resubmitted, there is no way of identifying events that were already being handled prior to the failure (outstanding events), nor is there a guarantee that objects have not been accessed or modified after the access rights were lost. The DM application must be able to recover consistently without depending on persistent access rights. For example, it could keep its own state of events in progress, or process events idempotently.

Similarly, when a specific file system is forced unmounted at the session node, all DM access rights associated with the file system are lost, although the events themselves prevail on the session queue. After the file system is remounted, DAPI calls using existing tokens may fail due to insufficient access rights. Also, there is no guarantee that objects have not been accessed or modified after the access rights were lost.

## **DODeferred deletions**

The asynchronous recovery code supports deferred deletions if there are no external mounts at the time of recovery.

Once a node successfully generates a mount event for an external mount, the **sgmgr** node will start deferred deletions if it is needed. Any internal mounts would bypass deferred deletions if the file system is DAPI enabled.

## **DM application failure**

If only the DM application fails, the session itself remains active, events remain pending, and client threads remain blocked waiting for a response. New events will continue to arrive at the session queue.

**Note:** GPFS is unable to detect that the DM application has failed.

The failed DM application must be recovered on the same node, and continue handling the events. Since no DMAPI resources are lost in this case, there is little purpose in moving the DM application to another node. Assuming an existing session on another node is not permitted in GPFS, except after session node failure.

If the DM application fails simultaneously with the session node, the **gpfsready** shell script can be used to restart the DM application on the failed node. See “[Initializing the Data Management application](#)” on page 637. In the case of simultaneous failures, the DM application can also be moved to another node and assume the failed session there. See “[Single-node failure](#)” on page 658.



---

## Chapter 3. GPFS programming interfaces

A list of all the GPFS programming interfaces and a short description of each is presented in this topic.

The GPFS APIs are not supported on Windows.

Table 23 summarizes the GPFS programming interfaces.

*Table 23. GPFS programming interfaces*

| Interface                                                                        | Purpose                                                                                                     |
|----------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| " <a href="#">gpfs_acl_t structure</a> " on page 666                             | Contains buffer mapping for the <code>gpfs_getacl()</code> and <code>gpfs_putacl()</code> subroutines.      |
| " <a href="#">gpfs_clone_copy() subroutine</a> " on page 667                     | Creates a file clone of a read-only clone parent file.                                                      |
| " <a href="#">gpfs_clone_snap() subroutine</a> " on page 669                     | Creates a read-only clone parent from a source file.                                                        |
| " <a href="#">gpfs_clone_split() subroutine</a> " on page 671                    | Splits a file clone from its clone parent.                                                                  |
| " <a href="#">gpfs_clone_unsnap() subroutine</a> " on page 673                   | Changes a clone parent with no file clones back to a regular file.                                          |
| " <a href="#">gpfs_close_inodescan() subroutine</a> " on page 675                | Closes an inode scan.                                                                                       |
| " <a href="#">gpfs_cmp_fssnapid() subroutine</a> " on page 676                   | Compares two file system snapshot IDs.                                                                      |
| " <a href="#">gpfs_declone() subroutine</a> " on page 678                        | Removes file clone references to clone parent blocks.                                                       |
| " <a href="#">gpfs_direntx_t structure</a> " on page 680                         | Contains attributes of a GPFS directory entry.                                                              |
| " <a href="#">gpfs_direntx64_t structure</a> " on page 682                       | Contains attributes of a GPFS directory entry.                                                              |
| " <a href="#">gpfs_fcntl() subroutine</a> " on page 684                          | Performs operations on an open file.                                                                        |
| " <a href="#">gpfs_fgetattr() subroutine</a> " on page 687                       | Retrieves all extended file attributes in opaque format.                                                    |
| " <a href="#">gpfs_fputattr() subroutine</a> " on page 689                       | Sets all the extended file attributes for a file.                                                           |
| " <a href="#">gpfs_fputattrswithpathname() subroutine</a> " on page 691          | Sets all of the extended file attributes for a file and invokes the policy engine for <b>RESTORE</b> rules. |
| " <a href="#">gpfs_free_fssnaphandle() subroutine</a> " on page 693              | Frees a GPFS file system snapshot handle.                                                                   |
| " <a href="#">gpfs_fssnap_handle_t structure</a> " on page 694                   | Contains a handle for a GPFS file system or snapshot.                                                       |
| " <a href="#">gpfs_fssnap_id_t structure</a> " on page 695                       | Contains a permanent identifier for a GPFS file system or snapshot.                                         |
| " <a href="#">gpfs_fstat() subroutine</a> " on page 696                          | Returns exact file status for a GPFS file.                                                                  |
| " <a href="#">gpfs_fstat_x() subroutine</a> " on page 698                        | Returns extended status information for a GPFS file with specified accuracy.                                |
| " <a href="#">gpfs_get_fsnname_from_fssnaphandle() subroutine</a> " on page 700  | Obtains a file system name from its snapshot handle.                                                        |
| " <a href="#">gpfs_get_fssnaphandle_by_fssnapid() subroutine</a> " on page 701   | Obtains a file system snapshot handle using its snapshot ID.                                                |
| " <a href="#">gpfs_get_fssnaphandle_by_name() subroutine</a> " on page 703       | Obtains a file system snapshot handle using its name.                                                       |
| " <a href="#">gpfs_get_fssnaphandle_by_path() subroutine</a> " on page 705       | Obtains a file system snapshot handle using its path name.                                                  |
| " <a href="#">gpfs_get_fssnapid_from_fssnaphandle() subroutine</a> " on page 707 | Obtains a file system snapshot ID using its handle.                                                         |

*Table 23. GPFS programming interfaces (continued)*

| Interface                                                                    | Purpose                                                                                                |
|------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| " <code>gpfs_get_pathname_from_fssnaphandle()</code> subroutine" on page 709 | Obtains a file system path name using its snapshot handle.                                             |
| " <code>gpfs_get_snapdirname()</code> subroutine" on page 711                | Obtains the name of the directory containing global snapshots.                                         |
| " <code>gpfs_get_snapname_from_fssnaphandle()</code> subroutine" on page 713 | Obtains a snapshot name using its file system snapshot handle.                                         |
| " <code>gpfs_getacl()</code> subroutine" on page 715                         | Retrieves the access control information for a GPFS file.                                              |
| " <code>gpfs_iattr_t</code> structure" on page 717                           | Contains attributes of a GPFS inode.                                                                   |
| " <code>gpfs_iattr64_t</code> structure" on page 720                         | Contains attributes of a GPFS inode.                                                                   |
| " <code>gpfs_iclose()</code> subroutine" on page 724                         | Closes a file given its inode file handle.                                                             |
| " <code>gpfs_ifile_t</code> structure" on page 726                           | Contains a handle for a GPFS inode.                                                                    |
| " <code>gpfs_igetattrs()</code> subroutine" on page 727                      | Retrieves extended file attributes in opaque format.                                                   |
| " <code>gpfs_igetattrsx()</code> subroutine" on page 729                     | Retrieves extended file attributes; provides an option to include DMAPI attributes.                    |
| " <code>gpfs_igetfilesetname()</code> subroutine" on page 731                | Returns the name of the fileset defined by a fileset ID.                                               |
| " <code>gpfs_igetstoragepool()</code> subroutine" on page 733                | Returns the name of the storage pool for the given storage pool ID.                                    |
| " <code>gpfs_iopen()</code> subroutine" on page 735                          | Opens a file or directory by inode number.                                                             |
| " <code>gpfs_iopen64()</code> subroutine" on page 737                        | Opens a file or directory by inode number.                                                             |
| " <code>gpfs_iputattrsx()</code> subroutine" on page 739                     | Sets the extended file attributes for a file.                                                          |
| " <code>gpfs_iread()</code> subroutine" on page 742                          | Reads a file opened by <code>gpfs_iopen()</code> .                                                     |
| " <code>gpfs_ireaddir()</code> subroutine" on page 744                       | Reads the next directory entry.                                                                        |
| " <code>gpfs_ireaddir64()</code> subroutine" on page 746                     | Reads the next directory entry.                                                                        |
| " <code>gpfs_ireadlink()</code> subroutine" on page 748                      | Reads a symbolic link by inode number.                                                                 |
| " <code>gpfs_ireadlink64()</code> subroutine" on page 750                    | Reads a symbolic link by inode number.                                                                 |
| " <code>gpfs_ireadx()</code> subroutine" on page 752                         | Performs block level incremental read of a file within an incremental inode scan.                      |
| " <code>gpfs_iscan_t</code> structure" on page 755                           | Contains a handle for an inode scan of a GPFS file system or snapshot.                                 |
| " <code>gpfs_lib_init()</code> subroutine" on page 756                       | Sets up a GPFS interface for additional calls.                                                         |
| " <code>gpfs_lib_term()</code> subroutine" on page 757                       | Cleans up after GPFS interface calls have been completed.                                              |
| " <code>gpfs_next_inode()</code> subroutine" on page 758                     | Retrieves the next inode from the inode scan.                                                          |
| " <code>gpfs_next_inode64()</code> subroutine" on page 760                   | Retrieves the next inode from the inode scan.                                                          |
| " <code>gpfs_next_inode_with_xatrrs()</code> subroutine" on page 762         | Retrieves the next inode and its extended attributes from the inode scan.                              |
| " <code>gpfs_next_inode_with_xatrrs64()</code> subroutine" on page 764       | Retrieves the next inode and its extended attributes from the inode scan.                              |
| " <code>gpfs_next_xattr()</code> subroutine" on page 766                     | Returns individual attributes and their values.                                                        |
| " <code>gpfs_opaque_acl_t</code> structure" on page 768                      | Contains buffer mapping for the <code>gpfs_getacl()</code> and <code>gpfs_putacl()</code> subroutines. |
| " <code>gpfs_open_inodescan()</code> subroutine" on page 769                 | Opens an inode scan of a file system or snapshot.                                                      |
| " <code>gpfs_open_inodescan64()</code> subroutine" on page 772               | Opens an inode scan of a file system or snapshot.                                                      |
| " <code>gpfs_open_inodescan_with_xatrrs()</code> on page 775                 | Opens an inode file and extended attributes for an inode scan.                                         |

*Table 23. GPFS programming interfaces (continued)*

| Interface                                                                 | Purpose                                                                                             |
|---------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|
| " <code>gpfs_open_inodescan_with_xattr64()</code> subroutine" on page 778 | Opens an inode file and extended attributes for an inode scan.                                      |
| " <code>gpfs_prealloc()</code> subroutine" on page 781                    | Preallocates storage for a regular file or preallocates directory entries for a directory.          |
| " <code>gpfs_putacl()</code> subroutine" on page 784                      | Restores the access control information for a GPFS file.                                            |
| " <code>gpfs_quotactl()</code> subroutine" on page 786                    | Manipulates disk quotas on file systems.                                                            |
| " <code>gpfs_quotaInfo_t</code> structure" on page 789                    | Contains buffer mapping for the <code>gpfs_quotactl()</code> subroutine.                            |
| " <code>gpfs_seek_inode()</code> subroutine" on page 791                  | Advances an inode scan to the specified inode number.                                               |
| " <code>gpfs_seek_inode64()</code> subroutine" on page 793                | Advances an inode scan to the specified inode number.                                               |
| " <code>gpfs_stat()</code> subroutine" on page 795                        | Returns exact file status for a GPFS file.                                                          |
| " <code>gpfs_stat_inode()</code> subroutine" on page 797                  | Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan. |
| " <code>gpfs_stat_inode64()</code> subroutine" on page 799                | Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan. |
| " <code>gpfs_stat_inode_with_xattr()</code> subroutine" on page 801       | Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan. |
| " <code>gpfs_stat_inode_with_xattr64()</code> subroutine" on page 803     | Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan. |
| " <code>gpfs_stat_x()</code> subroutine" on page 805                      | Returns extended status information for a GPFS file with specified accuracy.                        |
| " <code>gpfsFcctlHeader_t</code> structure" on page 807                   | Contains declaration information for the <code>gpfs_fcctl()</code> subroutine.                      |
| " <code>gpfsGetDataBlkDiskIdx_t</code> structure" on page 808             | Obtains the FPO data block location of a file.                                                      |
| " <code>gpfsGetFilesetName_t</code> structure" on page 811                | Obtains the fileset name of a file.                                                                 |
| " <code>gpfsGetReplication_t</code> structure" on page 812                | Obtains the replication factors of a file.                                                          |
| " <code>gpfsGetSetXAttr_t</code> structure" on page 814                   | Obtains or sets extended attribute values.                                                          |
| " <code>gpfsGetSnapshotName_t</code> structure" on page 816               | Obtains the snapshot name of a file.                                                                |
| " <code>gpfsGetStoragePool_t</code> structure" on page 817                | Obtains the storage pool name of a file.                                                            |
| " <code>gpfsListXAttr_t</code> structure" on page 818                     | Lists extended attributes.                                                                          |
| " <code>gpfsRestripeData_t</code> structure" on page 819                  | Restripes the data blocks of a file.                                                                |
| " <code>gpfsSetReplication_t</code> structure" on page 821                | Sets the replication factors of a file.                                                             |
| " <code>gpfsSetStoragePool_t</code> structure" on page 823                | Sets the assigned storage pool of a file.                                                           |

## **gpfs\_acl\_t**

---

### **gpfs\_acl\_t structure**

Contains buffer mapping for the **gpfs\_getacl()** and **gpfs\_putacl()** subroutines.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct gpfs_acl
{
 gpfs_aclLen_t acl_len; /* Total length of this ACL in bytes */
 gpfs_aclLevel_t acl_level; /* Reserved (must be zero) */
 gpfs_aclVersion_t acl_version; /* POSIX or NFS4 ACL */
 gpfs_aclType_t acl_type; /* Access, Default, or NFS4 */
 gpfs_aclCount_t acl_nace; /* Number of Entries that follow */

 union
 {
 gpfs_ace_v1_t ace_v1[1]; /* when GPFS_ACL_VERSION_POSIX */
 gpfs_ace_v4_t ace_v4[1]; /* when GPFS_ACL_VERSION_NFS4 */
 v4Level1_t v4Level1; /* when GPFS_ACL_LEVEL_V4FLAGS */
 };
} gpfs_acl_t;
```

#### **Description**

The **gpfs\_acl\_t** structure contains size, version, and ACL type information for the **gpfs\_getacl()** and **gpfs\_putacl()** subroutines.

#### **Members**

##### **acl\_len**

The total length (in bytes) of this **gpfs\_acl\_t** structure.

##### **acl\_level**

Reserved for future use. Currently must be zero.

##### **acl\_version**

This field contains the version of the GPFS ACL. GPFS supports the following ACL versions: **GPFS\_ACL\_VERSION\_POSIX** and **GPFS\_ACL\_VERSION\_NFS4**. On input to the **gpfs\_getacl()** subroutine, set this field to zero.

##### **acl\_type**

On input to the **gpfs\_getacl()** subroutine, set this field to:

- Either **GPFS\_ACL\_TYPE\_ACCESS** or **GPFS\_ACL\_TYPE\_DEFAULT** for POSIX ACLs
- **GPFS\_ACL\_TYPE\_NFS4** for NFS ACLs.

These constants are defined in the **gpfs.h** header file.

##### **acl\_nace**

The number of ACL entries that are in the array (**ace\_v1** or **ace\_v4**).

#### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_clone\_copy() subroutine**

Creates a file clone of a read-only clone parent file.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_clone_copy(const char *sourcePathP, const char *destPathP);
```

### **Description**

The **gpfs\_clone\_copy()** subroutine creates a writeable file clone from a read-only clone parent file.

### **Parameters**

#### **sourcePathP**

The path of a read-only source file to clone. The source file can be a file in a snapshot or a clone parent file created with the **gpfs\_clone\_snap()** subroutine.

#### **destPathP**

The path of the destination file to create. The destination file will become the file clone.

### **Exit status**

If the **gpfs\_clone\_copy()** subroutine is successful, it returns a value of 0 and creates a file clone from the clone parent.

If the **gpfs\_clone\_copy()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EACCESS**

Permission denied when writing to the destination path or reading from the source path.

#### **EEXIST**

The destination file already exists.

#### **EFAULT**

The input argument points outside the accessible address space.

#### **EINVAL**

The source or destination does not refer to a regular file or a GPFS file system.

#### **EISDIR**

The specified destination file is a directory.

#### **ENAMETOOLONG**

The source or destination path name is too long.

## **gpfs\_clone\_copy()**

### **ENOENT**

The source file does not exist.

### **ENOSPC**

The file system has run out of disk space.

### **ENOSYS**

The `gpfs_clone_copy()` subroutine is not available.

### **EPERM**

The source file is a directory or is not a regular file.

### **EXDEV**

The source file and destination file are not in the same file system.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

### **Related reference:**

“`gpfs_clone_snap()` subroutine” on page 669

Creates a read-only clone parent from a source file.

“`gpfs_clone_split()` subroutine” on page 671

Splits a file clone from its clone parent.

“`gpfs_clone_unsnap()` subroutine” on page 673

Changes a clone parent with no file clones back to a regular file.

---

## **gpfs\_clone\_snap() subroutine**

Creates a read-only clone parent from a source file.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_clone_snap(const char *sourcePathP, const char *destPathP);
```

### **Description**

The **gpfs\_clone\_snap()** subroutine creates a read-only clone parent from a source file.

### **Parameters**

#### **sourcePathP**

The path of the source file to clone.

#### **destPathP**

The path of the destination file to create. The destination file will become a read-only clone parent file.

If **destPathP** is NULL, then the source file will be changed in place into a read-only clone parent. When using this method to create a clone parent, the specified file cannot be open for writing or have hard links.

### **Exit status**

If the **gpfs\_clone\_snap()** subroutine is successful, it returns a value of 0 and creates a read-only clone parent from the source file.

If the **gpfs\_clone\_snap()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EACCESS**

Permission denied when writing to the destination path or reading from the source path.

#### **EEXIST**

The destination file already exists.

#### **EFAULT**

The input argument points outside accessible address space.

#### **EINVAL**

The source or destination does not refer to a regular file or a GPFS file system.

#### **EISDIR**

The specified destination file is a directory.

## **gpfs\_clone\_snap()**

### **ENAMETOOLONG**

The source or destination path name is too long.

### **ENOENT**

The source file does not exist.

### **ENOSPC**

The file system has run out of disk space.

### **ENOSYS**

The `gpfs_clone_snap()` subroutine is not available.

### **EPERM**

The source file is a directory or is not a regular file, or you tried to create a clone file with depth greater than 1000.

### **EXDEV**

The source file and destination file are not in the same file system.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

### **Related reference:**

“`gpfs_clone_copy()` subroutine” on page 667

Creates a file clone of a read-only clone parent file.

“`gpfs_clone_split()` subroutine” on page 671

Splits a file clone from its clone parent.

“`gpfs_clone_unsnap()` subroutine” on page 673

Changes a clone parent with no file clones back to a regular file.

---

## **gpfs\_clone\_split() subroutine**

Splits a file clone from its clone parent.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_clone_split(gpfs_file_t fileDesc, int ancLimit);
```

### **Description**

The **gpfs\_clone\_split()** subroutine splits a file clone from its clone parent. The **gpfs\_declone()** subroutine must be called first to remove all references to the clone parent.

### **Parameters**

#### **fileDesc**

File descriptor for the file clone to split from its clone parent.

#### **ancLimit**

The ancestor limit specified with one of these values:

##### **GPFS\_CLONE\_ALL**

Remove references to all clone parents.

##### **GPFS\_CLONE\_PARENT\_ONLY**

Remove references from the immediate clone parent only.

### **Exit status**

If the **gpfs\_clone\_split()** subroutine is successful, it returns a value of 0.

If the **gpfs\_clone\_split()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EACCESS**

Permission denied when writing to the target file.

#### **EBADF**

The file descriptor is not valid or is not a GPFS file.

#### **EINVAL**

An argument to the function was not valid.

#### **ENOSYS**

The **gpfs\_clone\_split()** subroutine is not available.

#### **EPERM**

The file descriptor does not refer to a regular file or a file clone.

## **gpfs\_clone\_split()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

#### **Related reference:**

“[gpfs\\_clone\\_copy\(\) subroutine](#)” on page 667

Creates a file clone of a read-only clone parent file.

“[gpfs\\_clone\\_snap\(\) subroutine](#)” on page 669

Creates a read-only clone parent from a source file.

“[gpfs\\_clone\\_unsnap\(\) subroutine](#)” on page 673

Changes a clone parent with no file clones back to a regular file.

---

## **gpfs\_clone\_unsnap() subroutine**

Changes a clone parent with no file clones back to a regular file.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_clone_unsnap(gpfs_file_t fileDesc);
```

### **Description**

The **gpfs\_clone\_unsnap()** subroutine changes a clone parent with no file clones back to a regular file.

### **Parameters**

#### **fileDesc**

File descriptor for the clone parent to convert back to a regular file.

### **Exit status**

If the **gpfs\_clone\_unsnap()** subroutine is successful, it returns a value of 0.

If the **gpfs\_clone\_unsnap()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EACCESS**

Permission denied when writing to the target file.

#### **EBADF**

The file descriptor is not valid or is not a GPFS file.

#### **EINVAL**

An argument to the function was not valid.

#### **ENOSYS**

The **gpfs\_clone\_unsnap()** subroutine is not available.

#### **EPERM**

The file descriptor does not refer to a regular file or a clone parent.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

#### **Related reference:**

## **gpfs\_clone\_unsnap()**

“[gpfs\\_clone\\_copy\(\) subroutine](#)” on page 667  
Creates a file clone of a read-only clone parent file.  
“[gpfs\\_clone\\_snap\(\) subroutine](#)” on page 669  
Creates a read-only clone parent from a source file.  
“[gpfs\\_clone\\_split\(\) subroutine](#)” on page 671  
Splits a file clone from its clone parent.

---

## **gpfs\_close\_inodescan() subroutine**

Closes an inode scan.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
void gpfs_close_inodescan(gpfs_iscan_t *iscan);
```

### **Description**

The **gpfs\_close\_inodescan()** subroutine closes the scan of the inodes in a file system or snapshot that was opened with the **gpfs\_open\_inodescan()** subroutine. The **gpfs\_close\_inodescan()** subroutine frees all storage used for the inode scan and invalidates the **iscan** handle.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **iscan**

Pointer to the inode scan handle.

### **Exit status**

The **gpfs\_close\_inodescan()** subroutine returns void.

### **Exceptions**

None.

### **Error status**

None.

### **Examples**

For an example using **gpfs\_close\_inodescan()**, see /usr/lpp/mmfs/samples/util/tsgetusage.c.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_cmp\_fssnapid()**

---

### **gpfs\_cmp\_fssnapid() subroutine**

Compares two file system snapshot IDs.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_cmp_fssnapid(const gpfs_fssnap_id_t *fssnapId1,
 const gpfs_fssnap_id_t *fssnapId2,
 int *result);
```

#### **Description**

The **gpfs\_cmp\_fssnapid()** subroutine compares two snapshot IDs for the same file system to determine the order in which the two snapshots were taken. The **result** parameter is set as follows:

- **result** less than zero indicates that snapshot 1 was taken before snapshot 2.
- **result** equal to zero indicates that snapshot 1 and 2 are the same.
- **result** greater than zero indicates that snapshot 1 was taken after snapshot 2.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **fssnapId1**

File system snapshot ID of the first snapshot.

##### **fssnapId2**

File system snapshot ID of the second snapshot.

##### **result**

Pointer to an integer indicating the outcome of the comparison.

#### **Exit status**

If the **gpfs\_cmp\_fssnapid()** subroutine is successful, it returns a value of 0 and the **result** parameter is set.

If the **gpfs\_cmp\_fssnapid()** subroutine is unsuccessful, it returns a value of -1 and the global error variable **errno** is set to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **EDOM**

The two snapshots cannot be compared because they were taken from two different file systems.

**ENOSYS**

The `gpfs_cmp_fssnapid()` subroutine is not available.

**GPFS\_EINVAL\_FSSNAPID**

`fssnapId1` or `fssnapId2` is not a valid snapshot ID.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_declone()**

---

### **gpfs\_declone() subroutine**

Removes file clone references to clone parent blocks.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_declone(gpfs_file_t fileDesc, int ancLimit, gpfs_off64_t nBlocks,
 gpfs_off64_t *offsetP);
```

#### **Description**

The **gpfs\_declone()** subroutine removes all file clone references to a clone parent by copying the clone parent blocks to the file clone.

#### **Parameters**

##### **fileDesc**

The file descriptor for the file clone.

##### **ancLimit**

The ancestor limit specified with one of these values:

##### **GPFS\_CLONE\_ALL**

Remove references to all clone parents.

##### **GPFS\_CLONE\_PARENT\_ONLY**

Remove references from the immediate clone parent only.

##### **nBlocks**

The maximum number of GPFS blocks to copy.

##### **offsetP**

A pointer to the starting offset within the file clone. This pointer will be updated to the offset of the next block to process or -1 if there no more blocks.

#### **Exit status**

If the **gpfs\_declone()** subroutine is successful, it returns a value of 0.

If the **gpfs\_declone()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **EACCESS**

Permission denied when writing to the target file.

##### **EBADF**

The file descriptor is not valid or is not a GPFS file.

**EFAULT**

The input argument points outside the accessible address space.

**EINVAL**

An argument to the function was not valid.

**ENOSPC**

The file system has run out of disk space.

**ENOSYS**

The `gpfs_declone()` subroutine is not available.

**EPERM**

The file descriptor does not refer to a regular file.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_direntx\_t**

---

### **gpfs\_direntx\_t structure**

Contains attributes of a GPFS directory entry.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct gpfs_direntx
{
 int d_version; /* this struct's version */
 unsigned short d_reclen; /* actual size of this struct including
 null terminated variable length d_name */
 unsigned short d_type; /* Types are defined below */
 gpfs_ino_t d_ino; /* File inode number */
 gpfs_gen_t d_gen; /* Generation number for the inode */
 char d_name[256]; /* null terminated variable length name */
} gpfs_direntx_t;

/* File types for d_type field in gpfs_direntx_t */
#define GPFS_DE_OTHER 0
#define GPFS_DE_DIR 4
#define GPFS_DE_REG 8
#define GPFS_DE_LNK 10
#define GPFS_DE_DEL 16
```

#### **Description**

The **gpfs\_direntx\_t** structure contains the attributes of a GPFS directory entry.

#### **Members**

##### **d\_version**

The version number of this structure.

##### **d\_reclen**

The actual size of this structure including the null-terminated variable-length **d\_name** field.

To allow some degree of forward compatibility, careful callers should use the **d\_reclen** field for the size of the structure rather than the **sizeof()** function.

##### **d\_type**

The type of directory.

##### **d\_ino**

The directory inode number.

##### **d\_gen**

The directory generation number.

##### **d\_name**

Null-terminated variable-length name of the directory.

#### **Examples**

For an example using **gpfs\_direntx\_t**, see `/usr/lpp/mmfs/samples/util/tsfindinode.c`.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX  
/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_direntx64\_t**

---

### **gpfs\_direntx64\_t structure**

Contains attributes of a GPFS directory entry.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct gpfs_direntx64
{
 int d_version; /* this struct's version */
 unsigned short d_reclen; /* actual size of this struct including
 null terminated variable length d_name */
 unsigned short d_type; /* Types are defined below */
 gpfs_ino64_t d_ino; /* File inode number */
 gpfs_gen64_t d_gen; /* Generation number for the inode */
 unsigned int d_flags; /* Flags are defined below */
 char d_name[1028]; /* null terminated variable length name
 /* (1020+null+7 byte pad to double word) */
 /* to handle up to 255 UTF-8 chars */
} gpfs_direntx64_t;

/* File types for d_type field in gpfs_direntx64_t */
#define GPFS_DE_OTHER 0
#define GPFS_DE_DIR 4
#define GPFS_DE_REG 8
#define GPFS_DE_LNK 10
#define GPFS_DE_DEL 16

/* Define flags for gpfs_direntx64_t */
#define GPFS_DEFFLAG_NONE 0x0000 /* Default value, no flags set */
#define GPFS_DEFFLAG_JUNCTION 0x0001 /* DirEnt is a fileset junction */
#define GPFS_DEFFLAG_IJUNCTION 0x0002 /* DirEnt is a inode space junction */
#define GPFS_DEFFLAG_ORPHAN 0x0004 /* DirEnt is an orphan (pcache) */
```

#### **Description**

The **gpfs\_direntx64\_t** structure contains the attributes of a GPFS directory entry.

#### **Members**

##### **d\_version**

The version number of this structure.

##### **d\_reclen**

The actual size of this structure including the null-terminated variable-length **d\_name** field.

To allow some degree of forward compatibility, careful callers should use the **d\_reclen** field for the size of the structure rather than the **sizeof()** function.

##### **d\_type**

The type of directory.

##### **d\_ino**

The directory inode number.

##### **d\_gen**

The directory generation number.

##### **d\_flags**

The directory flags.

**d\_name**

Null-terminated variable-length name of the directory.

**Examples**

See the **gpfs\_direntx\_t** example in `/usr/lpp/mmfs/samples/util/tsfindinode.c`.

**Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_fcntl()**

---

### **gpfs\_fcntl() subroutine**

Performs operations on an open file.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_fcntl(gpfs_file_t fileDesc, void* fcntlArgP);
```

#### **Description**

The **gpfs\_fcntl()** subroutine is used to pass file access pattern information and to control certain file attributes on behalf of an open file. More than one operation can be requested with a single invocation of **gpfs\_fcntl()**. The type and number of operations is determined by the second operand, **fcntlArgP**, which is a pointer to a data structure built in memory by the application. This data structure consists of:

- A fixed length header, mapped by **gpfsFcntlHeader\_t**.
- A variable list of individual file access hints, directives or other control structures:
  - File access hints:
    - **gpfsAccessRange\_t**
    - **gpfsFreeRange\_t**
    - **gpfsMultipleAccessRange\_t**
    - **gpfsClearFileCache\_t**
  - File access directives:
    - **gpfsCancelHints\_t**
    - **gpfsDataShipMap\_t**
    - **gpfsDataShipStart\_t**
    - **gpfsDataShipStop\_t**
  - Platform-independent extended attribute operations:
    - **gpfsGetSetXAttr\_t**
    - **gpfsListXAttr\_t**
  - Other file attribute operations:
    - **gpfsGetDataBlkDiskIdx\_t**
    - **gpfsGetFilesetName\_t**
    - **gpfsGetReplication\_t**
    - **gpfsGetSnapshotName\_t**
    - **gpfsGetStoragePool\_t**
    - **gpfsRestripeData\_t**
    - **gpfsSetReplication\_t**
    - **gpfsSetStoragePool\_t**

These hints, directives and other operations may be mixed within a single **gpfs\_fcntl()** subroutine, and are performed in the order that they appear. A subsequent hint or directive may cancel out a preceding one.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

## Parameters

### **fileDesc**

The file descriptor identifying the file to which GPFS applies the hints and directives.

### **fcntlArgP**

A pointer to the list of operations to be passed to GPFS.

## Exit status

If the **gpfs\_fcntl()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fcntl()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### **EBADF**

The file descriptor is not valid.

### **EINVAL**

The file descriptor does not refer to a GPFS file or a regular file.

The system call is not valid.

### **ENOSYS**

The **gpfs\_fcntl()** subroutine is not supported under the current file system format.

## Examples

1. This programming segment releases all cache data held by the file *handle* and tell GPFS that the subroutine will write the portion of the file with file offsets between 2 GB and 3 GB minus one:

```
struct
{
 gpfsFcntlHeader_t hdr;
 gpfsClearFileCache_t rel;
 gpfsAccessRange_t acc;
} arg;

arg.hdr.totalLength = sizeof(arg);
arg.hdr.fcntlVersion = GPFS_FCNTRL_CURRENT_VERSION;
arg.hdr.fcntlReserved = 0;
arg.rel.structLen = sizeof(arg.rel);
arg.rel.structType = GPFS_CLEAR_FILE_CACHE;
arg.acc.structLen = sizeof(arg.acc);
arg.acc.structType = GPFS_ACCESS_RANGE;
arg.acc.start = 2LL * 1024LL * 1024LL * 1024LL;
arg.acc.length = 1024 * 1024 * 1024;
arg.acc.isWrite = 1;
rc = gpfs_fcntl(handle, &arg);
```

2. This programming segment gets the storage pool name and fileset name of a file from GPFS.

```
struct {
 gpfsFcntlHeader_t hdr;
 gpfsGetStoragePool_t pool;
 gpfsGetFilesetName_t fileset;
} fcntlArg;
fcntlArg.hdr.totalLength = sizeof(fcntlArg.hdr) + sizeof(fcntlArg.pool) + sizeof(fcntlArg.fileset);
```

## **gpfs\_fcntl()**

```
fcntlArg.hdr.fcntlVersion = GPFS_FCNTRL_CURRENT_VERSION;
fcntlArg.hdr.fcntlReserved = 0;

fcntlArg.pool.structLen = sizeof(fcntlArg.pool);
fcntlArg.pool.structType = GPFS_FCNTRL_GET_STORAGEPOOL;

fcntlArg.fileset.structLen = sizeof(fcntlArg.fileset);
fcntlArg.fileset.structType = GPFS_FCNTRL_GET_FILESETNAME;

rc = gpfs_fcntl(fd, &fcntlArg);
```

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_fgetattrs() subroutine

Retrieves all extended file attributes in opaque format.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_fgetattrs(gpfs_file_t fileDesc,
 int flags,
 void *bufferP,
 int bufferSize,
 int *attrSizeP);
```

### Description

The **gpfs\_fgetattrs()** subroutine, together with **gpfs\_fputattrs()**, is intended for use by a backup program to save (**gpfs\_fgetattrs()**) and restore (**gpfs\_fputattrs()**) extended file attributes such as ACLs, DMAPI attributes, and other information for the file. If the file has no extended attributes, the **gpfs\_fgetattrs()** subroutine returns a value of 0, but sets **attrSizeP** to 0 and leaves the contents of the buffer unchanged.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### fileDesc

The file descriptor identifying the file whose extended attributes are being retrieved.

#### flags

Must have one of the following values:

##### **GPFS\_ATTRFLAG\_DEFAULT**

Saves the attributes for file placement and the currently assigned storage pool.

##### **GPFS\_ATTRFLAG\_NO\_PLACEMENT**

Does not save attributes for file placement or the currently assigned storage pool.

##### **GPFS\_ATTRFLAG\_IGNORE\_POOL**

Saves attributes for file placement but does not save the currently assigned storage pool.

##### **GPFS\_ATTRFLAG\_USE\_POLICY**

Uses the restore policy rules to determine the pool ID.

##### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

Includes the DMAPI attributes.

##### **GPFS\_ATTRFLAG\_FINALIZE\_ATTRS**

Finalizes immutability attributes.

##### **GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**

Skips immutable attributes.

##### **GPFS\_ATTRFLAG\_INCL\_ENCR**

Includes encryption attributes.

##### **GPFS\_ATTRFLAG\_SKIP\_CLONE**

Skips clone attributes.

## **gpfs\_fgetattrs()**

### **GPFS\_ATTRFLAG MODIFY\_CLONEPARENT**

Allows modification on the clone parent.

### **bufferP**

Pointer to a buffer to store the extended attribute information.

### **bufferSize**

The size of the buffer that was passed in.

### **attrSizeP**

If successful, returns the actual size of the attribute information that was stored in the buffer. If the **bufferSize** was too small, returns the minimum buffer size.

## **Exit status**

If the **gpfs\_fgetattrs()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fgetattrs()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EBADF**

The file descriptor is not valid.

### **EFAULT**

The address is not valid.

### **EINVAL**

The file descriptor does not refer to a GPFS file.

### **ENOSPC**

**bufferSize** is too small to return all of the attributes. On return, **attrSizeP** is set to the required size.

### **ENOSYS**

The **gpfs\_fgetattrs()** subroutine is not supported under the current file system format.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_fputattrs() subroutine

Sets all the extended file attributes for a file.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_fputattrs(gpfs_file_t fileDesc,
 int flags,
 void *bufferP);
```

### Description

The **gpfs\_fputattrs()** subroutine, together with **gpfs\_fgetattr()**, is intended for use by a backup program to save (**gpfs\_fgetattr()**) and restore (**gpfs\_fputattrs()**) all of the extended attributes of a file. This subroutine also sets the storage pool for the file and sets data replication to the values that are saved in the extended attributes.

If the saved storage pool is not valid or if the **GPFS\_ATTRFLAG\_IGNORE\_POOL** flag is set, GPFS will select the storage pool by matching a **PLACEMENT** rule using the saved file attributes. If GPFS fails to match a placement rule or if there are no placement rules installed, GPFS assigns the file to the system storage pool.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### fileDesc

The file descriptor identifying the file whose extended attributes are being set.

#### flags

Must have one of the following values:

##### **GPFS\_ATTRFLAG\_DEFAULT**

Restores the previously assigned storage pool and previously assigned data replication.

##### **GPFS\_ATTRFLAG\_NO\_PLACEMENT**

Does not change storage pool and data replication.

##### **GPFS\_ATTRFLAG\_IGNORE\_POOL**

Selects storage pool and data replication by matching the saved attributes to a placement rule instead of restoring the saved storage pool.

##### **GPFS\_ATTRFLAG\_USE\_POLICY**

Uses the restore policy rules to determine the pool ID.

##### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

Includes the DAPI attributes.

##### **GPFS\_ATTRFLAG\_FINALIZE\_ATTRS**

Finalizes immutability attributes.

##### **GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**

Skips immutable attributes.

## **gpfs\_fputattrs()**

### **GPFS\_ATTRFLAG\_INCL\_ENCR**

Includes encryption attributes.

### **GPFS\_ATTRFLAG\_SKIP\_CLONE**

Skips clone attributes.

### **GPFS\_ATTRFLAG MODIFY\_CLONEPARENT**

Allows modification on the clone parent.

Non-placement attributes such as ACLs are always restored, regardless of value of the flag.

### **bufferP**

A pointer to the buffer containing the extended attributes for the file.

If you specify a value of NULL, all extended ACLs for the file are deleted.

## **Exit status**

If the **gpfs\_fputattrs()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fputattrs()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EBADF**

The file descriptor is not valid.

### **EINVAL**

The buffer pointed to by **bufferP** does not contain valid attribute data, or the file descriptor does not refer to a GPFS file.

### **ENOSYS**

The **gpfs\_fputattrs()** subroutine is not supported under the current file system format.

## **Examples**

To copy extended file attributes from file **f1** to file **f2**:

```
char buf[4096];
int f1, f2, attrSize, rc;

rc = gpfs_fgetattrs(f1, GPFS_ATTRFLAG_DEFAULT, buf, sizeof(buf), &attrSize);
if (rc != 0)
 ...
if (attrSize != 0)
 rc = gpfs_fputattrs(f2, 0, buf); // copy attributes from f1 to f2
else
 rc = gpfs_fputattrs(f2, 0, NULL); // f1 has no attributes
 // delete attributes on f2
```

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_fputattrswithpathname() subroutine**

Sets all of the extended file attributes for a file and invokes the policy engine for **RESTORE** rules.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_fputattrswithpathname(gpfs_file_t fileDesc,
 int flags,
 void *bufferP,
 const char *pathName);
```

### **Description**

The **gpfs\_fputattrswithpathname()** subroutine sets all of the extended attributes of a file. In addition, **gpfs\_fputattrswithpathname()** invokes the policy engine using the saved attributes to match a **RESTORE** rule to set the storage pool and the data replication for the file. The caller should include the full path to the file (including the file name) to allow rule selection based on file name or path. If the file fails to match a **RESTORE** rule or if there are no **RESTORE** rules installed, GPFS selects the storage pool and data replication as it does when calling **gpfs\_fputattrs()**.

**Note:** Compile any program that uses this subroutine with the `-lgpfs` flag from one the following libraries:

- `libgpfs.a` for AIX
- `libgpfs.so` for Linux

### **Parameters**

#### **fileDesc**

Is the file descriptor that identifies the file whose extended attributes are to be set.

#### **flags**

Must have one of the following values:

##### **GPFS\_ATTRFLAG\_DEFAULT**

Uses the saved attributes to match a **RESTORE** rule to set the storage pool and the data replication for the file.

##### **GPFS\_ATTRFLAG\_NO\_PLACEMENT**

Does not change storage pool and data replication.

##### **GPFS\_ATTRFLAG\_IGNORE\_POOL**

Checks the file to see if it matches a **RESTORE** rule. If the file fails to match a **RESTORE** rule, GPFS ignores the saved storage pool and selects a pool by matching the saved attributes to a **PLACEMENT** rule.

##### **GPFS\_ATTRFLAG\_USE\_POLICY**

Uses the restore policy rules to determine the pool ID.

##### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

Includes the DAPI attributes.

##### **GPFS\_ATTRFLAG\_FINALIZE\_ATTRS**

Finalizes immutability attributes.

##### **GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**

Skips immutable attributes.

## **gpfs\_fputattrswithpathname()**

### **GPFS\_ATTRFLAG\_INCL\_ENCR**

Includes encryption attributes.

### **GPFS\_ATTRFLAG\_SKIP\_CLONE**

Skips clone attributes.

### **GPFS\_ATTRFLAG MODIFY\_CLONEPARENT**

Allows modification on the clone parent.

Non-placement attributes such as ACLs are always restored, regardless of value of the flag.

### **bufferP**

A pointer to the buffer containing the extended attributes for the file.

If you specify a value of NULL, all extended ACLs for the file are deleted.

### **pathName**

A pointer to the path name to a file or directory.

## **Exit status**

If the **gpfs\_fputattrswithpathname()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fputattrswithpathname()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EBADF**

The file descriptor is not valid.

### **EINVAL**

The buffer to which **bufferP** points does not contain valid attribute data.

### **ENOENT**

No such file or directory.

### **ENOSYS**

The **gpfs\_fputattrswithpathname()** subroutine is not supported under the current file system format.

## **Examples**

Refer to “**gpfs\_fputattr()** subroutine” on page 689 for examples.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_free\_fssnaphandle() subroutine**

Frees a GPFS file system snapshot handle.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
void gpfs_free_fssnaphandle(gpfs_fssnap_handle_t *fssnapHandle);
```

### **Description**

The **gpfs\_free\_fssnaphandle()** subroutine frees the snapshot handle that is passed. The return value is always void.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

### **Exit status**

The **gpfs\_free\_fssnaphandle()** subroutine always returns void.

### **Exceptions**

None.

### **Error status**

None.

### **Examples**

For an example using **gpfs\_free\_fssnaphandle()**, see */usr/lpp/mmfs/samples/util/tstimes.c*.

### **Location**

*/usr/lpp/mmfs/lib/libgpfs.a* for AIX

*/usr/lpp/mmfs/lib/libgpfs.so* for Linux

## **gpfs\_fssnap\_handle\_t**

---

### **gpfs\_fssnap\_handle\_t structure**

Contains a handle for a GPFS file system or snapshot.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct gpfs_fssnap_handle gpfs_fssnap_handle_t;
```

#### **Description**

A file system or snapshot is uniquely identified by an **fssnapId** of type **gpfs\_fssnap\_id\_t**. While the **fssnapId** is permanent and global, a shorter **fssnapHandle** is used by the backup application programming interface to identify the file system and snapshot being accessed. The **fssnapHandle**, like a POSIX file descriptor, is volatile and may be used only by the program that created it.

There are three ways to create a file system snapshot handle:

1. By using the name of the file system and snapshot
2. By specifying the path through the mount point
3. By providing an existing file system snapshot ID

Additional subroutines are provided to obtain the permanent, global **fssnapId** from the **fssnapHandle**, or to obtain the path or the names for the file system and snapshot, if they are still available in the file system.

The file system must be mounted in order to use the backup programming application interface. If the **fssnapHandle** is created by the path name, the path may be relative and may specify any file or directory in the file system. Operations on a particular snapshot are indicated with a path to a file or directory within that snapshot. If the **fssnapHandle** is created by name, the file system's unique name may be specified (for example, **fs1**) or its device name may be provided (for example, **/dev/fs1**). To specify an operation on the active file system, the pointer to the snapshot's name should be set to NULL or a zero-length string provided.

The name of the directory under which all snapshots appear may be obtained by the **gpfs\_get\_snapdirname()** subroutine. By default this is **.snapshots**, but it can be changed using the **mmsnapdir** command. The **gpfs\_get\_snapdirname()** subroutine returns the currently set value, which is the one that was last set by the **mmsnapdir** command, or the default, if it was never changed.

#### **Members**

##### **gpfs\_fssnap\_handle**

File system snapshot handle

#### **Examples**

For an example using **gpfs\_fssnap\_handle\_t**, see **/usr/lpp/mmfs/samples/util/tsgetusage.c**.

#### **Location**

**/usr/lpp/mmfs/lib/libgpfs.a** for AIX

**/usr/lpp/mmfs/lib/libgpfs.so** for Linux

---

## **gpfs\_fssnap\_id\_t structure**

Contains a permanent identifier for a GPFS file system or snapshot.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Structure**

```
typedef struct gpfs_fssnap_id
{
 char opaque[48];
} gpfs_fssnap_id_t;
```

### **Description**

A file system or snapshot is uniquely identified by an **fssnapId** of type **gpfs\_fssnap\_id\_t**. The **fssnapId** is a permanent and global identifier that uniquely identifies an active file system or a read-only snapshot of a file system. Every snapshot of a file system has a unique identifier that is also different from the identifier of the active file system itself.

The **fssnapId** is obtained from an open **fssnapHandle**. Once obtained, the **fssnapId** should be stored along with the file system's data for each backup. The **fssnapId** is required to generate an incremental backup. The **fssnapId** identifies the previously backed up file system or snapshot and allows the inode scan to return only the files and data that have changed since that previous scan.

### **Members**

#### **opaque**

A 48 byte area for containing the snapshot identifier.

### **Examples**

For an example using **gpfs\_fssnap\_id\_t**, see `/usr/lpp/mmfssamples/util/tsinode.c`.

### **Location**

`/usr/lpp/mmfss/lib/libgpfs.a` for AIX

`/usr/lpp/mmfss/lib/libgpfs.so` for Linux

## gpfs\_fstat()

---

### gpfs\_fstat() subroutine

Returns exact file status for a GPFS file.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Synopsis

```
#include <gpfs.h>
int gpfs_fstat(gpfs_file_t fileDesc,
 gpfs_stat64_t *buffer);
```

#### Description

The **gpfs\_fstat()** subroutine is used to obtain exact information about the file associated with the **fileDesc** parameter. This subroutine is provided as an alternative to the **stat()** subroutine, which may not provide exact **mtime** and **atime** values. For more information, see the topic *Exceptions to Open Group technical standards* in the *IBM Spectrum Scale: Administration Guide*.

**read**, **write**, or **execute** permission for the named file is not required, but all directories listed in the path leading to the file must be searchable. The file information is written to the area specified by the **buffer** parameter.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### Parameters

##### fileDesc

The file descriptor identifying the file for which exact status information is requested.

##### buffer

A pointer to the **gpfs\_stat64\_t** structure in which the information is returned. The **gpfs\_stat64\_t** structure is described in the *sys/stat.h* file.

#### Exit status

If the **gpfs\_fstat()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fstat()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### Exceptions

None.

#### Error status

Error codes include but are not limited to the following:

##### EBADF

The file descriptor is not valid.

##### EINVAL

The file descriptor does not refer to a GPFS file or a regular file.

**ENOSYS**

The **gpfs\_fstat()** subroutine is not supported under the current file system format.

**ESTALE**

The cached file system information was not valid.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_fstat\_x()**

---

### **gpfs\_fstat\_x() subroutine**

Returns extended status information for a GPFS file with specified accuracy.

#### **Library**

GPFS library. Runs on AIX, Linux, and Windows.

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_fstat_x(gpfs_file_t fileDesc,
 unsigned int *st_litemask,
 gpfs_iattr64_t *iattr,
 size_t iattrBufLen);
```

#### **Description**

The **gpfs\_fstat\_x()** subroutine is similar to the **gpfs\_fstat()** subroutine but returns more information in a **gpfs\_iattr64** structure that is defined in **gpfs.h**. This subroutine is supported only on the Linux operating system.

Your program must verify that the version of the **gpfs\_iattr64** structure that is returned in the field **ia\_version** is the same as the version that you are using. Versions are defined in **gpfs.h** with the pattern **GPFS\_IA64\_VERSION\***.

File permissions **read**, **write**, and **execute** are not required for the specified file, but all the directories in the specified path must be searchable.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- **libgpfs.so** for Linux

#### **Parameters**

##### **gpfs\_file\_t fileDesc**

The file descriptor of a file for which information is requested.

##### **\*st\_litemask**

A pointer to a bitmask specification of the items that you want to be returned exactly. Bitmasks are defined in **gpfs.h** with the pattern **GPFS\_SLITE\_\*BIT**. The subroutine returns exact values for the items that you specify in the bitmask. The subroutine also sets bits in the bitmask to indicate any other items that are exact.

##### **\*iattr**

A pointer to a **gpfs\_iattr64\_t** structure in which the information is returned. The structure is described in the **gpfs.h** file.

##### **iattrBufLen**

The length of your **gpfs\_iattr64\_t** structure, as given by **sizeof(myStructure)**. The subroutine does not write data past this limit. The field **ia\_recflen** in the returned structure is the length of the **gpfs\_iattr64\_t** structure that the subroutine is using.

#### **Exit status**

If the **gpfs\_fstat\_x()** subroutine is successful, it returns a value of 0.

If the **gpfs\_fstat\_x()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following errors:

### **EBADF**

The file handle does not refer to an existing or accessible object.

### **ENOSYS**

The `gpfs_fstat_x()` subroutine is not supported under the current file system format.

### **ESTALE**

The cached file system information was invalid.

## Location

/usr/lpp/mmfs/lib

## **gpfs\_get\_fsname\_from\_fssnaphandle()**

---

### **gpfs\_get\_fsname\_from\_fssnaphandle() subroutine**

Obtains a file system name from its snapshot handle.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
const char *gpfs_get_fsname_from_fssnaphandle(gpfs_fssnap_handle_t *fssnapHandle);
```

#### **Description**

The **gpfs\_get\_fsname\_from\_fssnaphandle()** subroutine returns a pointer to the name of file system that is uniquely identified by the file system snapshot handle.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **fssnapHandle**

File system snapshot handle.

#### **Exit status**

If the **gpfs\_get\_fsname\_from\_fssnaphandle()** subroutine is successful, it returns a pointer to the name of the file system identified by the file system snapshot handle.

If the **gpfs\_get\_fsname\_from\_fssnaphandle()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **ENOSYS**

The **gpfs\_get\_fsname\_from\_fssnaphandle()** subroutine is not available.

##### **EPERM**

The caller does not have superuser privileges.

##### **GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

#### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_get\_fssnaphandle\_by\_fssnapid() subroutine**

Obtains a file system snapshot handle using its snapshot ID.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
gpfs_fssnap_handle_t *gpfs_get_fssnaphandle_by_fssnapid(const gpfs_fssnap_id_t *fssnapId);
```

### **Description**

The **gpfs\_get\_fssnaphandle\_by\_fssnapid()** subroutine creates a handle for the file system or snapshot that is uniquely identified by the permanent, unique snapshot ID.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapId**

File system snapshot ID

### **Exit status**

If the **gpfs\_get\_fssnaphandle\_by\_fssnapid()** subroutine is successful, it returns a pointer to the file system snapshot handle.

If the **gpfs\_get\_fssnaphandle\_by\_fssnapid()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOMEM**

Space could not be allocated for the file system snapshot handle.

#### **ENOSYS**

The **gpfs\_get\_fssnaphandle\_by\_fssnapid()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

#### **GPFS\_EINVAL\_FSSNAPID**

The file system snapshot ID is not valid.

## **gpfs\_get\_fssnaphandle\_by\_fssnapid()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_get\_fssnaphandle\_by\_name() subroutine**

Obtains a file system snapshot handle using its name.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
gpfs_fssnap_handle_t *gpfs_get_fssnaphandle_by_name(const char *fsName,
 const char *snapName);
```

### **Description**

The **gpfs\_get\_fssnaphandle\_by\_name()** subroutine creates a handle for the file system or snapshot that is uniquely identified by the file system's name and the name of the snapshot.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fsName**

A pointer to the name of the file system whose snapshot handle is desired.

#### **snapName**

A pointer to the name of the snapshot whose snapshot handle is desired, or NULL to access the active file system rather than a snapshot within the file system.

### **Exit status**

If the **gpfs\_get\_fssnaphandle\_by\_name()** subroutine is successful, it returns a pointer to the file system snapshot handle.

If the **gpfs\_get\_fssnaphandle\_by\_name()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOENT**

The file system name is not valid.

#### **ENOMEM**

Space could not be allocated for the file system snapshot handle.

#### **ENOSYS**

The **gpfs\_get\_fssnaphandle\_by\_name()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

## **gpfs\_get\_fssnaphandle\_by\_name()**

### **GPFS\_EINVAL\_SNAPNAME**

The snapshot name is not valid.

## **Examples**

For an example using `gpfs_get_fssnaphandle_by_name()`, see `/usr/lpp/mmfs/samples/util/tsinode.c`.

## **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## **gpfs\_get\_fssnaphandle\_by\_path() subroutine**

Obtains a file system snapshot handle using its path name.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
gpfs_fssnap_handle_t *gpfs_get_fssnaphandle_by_path(const char *pathName);
```

### **Description**

The **gpfs\_get\_fssnaphandle\_by\_path()** subroutine creates a handle for the file system or snapshot that is uniquely identified by a path through the file system's mount point to a file or directory within the file system or snapshot.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **pathName**

A pointer to the path name to a file or directory within the desired file system or snapshot.

### **Exit status**

If the **gpfs\_get\_fssnaphandle\_by\_path()** subroutine is successful, it returns a pointer to the file system snapshot handle.

If the **gpfs\_get\_fssnaphandle\_by\_path()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOENT**

The path name is not valid.

#### **ENOMEM**

Space could not be allocated for the file system snapshot handle.

#### **ENOSYS**

The **gpfs\_get\_fssnaphandle\_by\_path()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

### **Examples**

For an example using **gpfs\_get\_fssnaphandle\_by\_path()**, see */usr/lpp/mmfs/samples/util/tsgetusage.c*.

## **gpfs\_get\_fssnaphandle\_by\_path()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_get\_fssnapid\_from\_fssnaphandle() subroutine**

Obtains a file system snapshot ID using its handle.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_get_fssnapid_from_fssnaphandle(gpfs_fssnap_handle_t *fssnapHandle,
 gpfs_fssnap_id_t *fssnapId);
```

### **Description**

The **gpfs\_get\_fssnapid\_from\_fssnaphandle()** subroutine obtains the permanent, globally unique file system snapshot ID of the file system or snapshot identified by the open file system snapshot handle.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

#### **fssnapId**

File system snapshot ID.

### **Exit status**

If the **gpfs\_get\_fssnapid\_from\_fssnaphandle()** subroutine is successful, it returns a pointer to the file system snapshot ID.

If the **gpfs\_get\_fssnapid\_from\_fssnaphandle()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EFAULT**

Size mismatch for **fssnapId**.

#### **EINVAL**

NULL pointer given for returned **fssnapId**.

#### **ENOSYS**

The **gpfs\_get\_fssnapid\_from\_fssnaphandle()** subroutine is not available.

#### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

## **gpfs\_get\_fssnapid\_from\_fssnaphandle()**

### **Examples**

For an example using `gpfs_get_fssnapid_from_fssnaphandle()`, see `/usr/lpp/mmfs/samples/util/tsinode.c`.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_get.pathname\_from\_fssnaphandle() subroutine**

Obtains a file system path name using its snapshot handle.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
const char *gpfs_get.pathname_from_fssnaphandle(gpfs_fssnap_handle_t *fssnapHandle);
```

### **Description**

The **gpfs\_get.pathname\_from\_fssnaphandle()** subroutine obtains the path name of the file system or snapshot identified by the open file system snapshot handle.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

### **Exit status**

If the **gpfs\_get.pathname\_from\_fssnaphandle()** subroutine is successful, it returns a pointer to the path name of the file system or snapshot.

If the **gpfs\_get.pathname\_from\_fssnaphandle()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOSYS**

The **gpfs\_get.pathname\_from\_fssnaphandle()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

#### **GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **Examples**

For an example using **gpfs\_get.pathname\_from\_fssnaphandle()**, see /usr/lpp/mmfs/samples/util/tsfindinode.c.

## **gpfs\_get.pathname\_from\_fssnaphandle()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_get\_snapdirname() subroutine**

Obtains the name of the directory containing global snapshots.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_get_snapdirname(gpfs_fssnap_handle_t *fssnapHandle,
 char *snapdirName,
 int bufLen);
```

### **Description**

The **gpfs\_get\_snapdirname()** subroutine obtains the name of the directory that contains global snapshots.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

#### **snapdirName**

Buffer into which the name of the snapshot directory will be copied.

#### **bufLen**

The size of the provided buffer.

### **Exit status**

If the **gpfs\_get\_snapdirname()** subroutine is successful, it returns a value of 0 and the **snapdirName** and **bufLen** parameters are set.

If the **gpfs\_get\_snapdirname()** subroutine is unsuccessful, it returns a value of -1 and the global error variable **errno** is set to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOMEM**

Unable to allocate memory for the request.

#### **ENOSYS**

The **gpfs\_get\_snapdirname()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

## **gpfs\_get\_snapdirname()**

### **ERANGE**

The buffer is too small to return the snapshot directory name.

### **ESTALE**

The cached file system information was not valid.

### **GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **E2BIG** The buffer is too small to return the snapshot directory name.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_get\_snapname\_from\_fssnaphandle() subroutine**

Obtains a snapshot name using its file system snapshot handle.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
const char *gpfs_get_snapname_from_fssnaphandle(gpfs_fssnap_handle_t *fssnapHandle);
```

### **Description**

The **gpfs\_get\_snapname\_from\_fssnaphandle()** subroutine obtains a pointer to the name of a GPFS snapshot given its file system snapshot handle. If the **fssnapHandle** identifies an active file system, as opposed to a snapshot of a file system, **gpfs\_get\_snapname\_from\_fssnaphandle()** returns a pointer to a zero-length snapshot name and a successful return code.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

### **Exit status**

If the **gpfs\_get\_snapname\_from\_fssnaphandle()** subroutine is successful, it returns a pointer to the name of the snapshot.

If the **gpfs\_get\_snapname\_from\_fssnaphandle()** subroutine is unsuccessful, it returns NULL and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOSYS**

The **gpfs\_get\_snapname\_from\_fssnaphandle()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

#### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

#### **GPFS\_E\_INVAL\_SNAPNAME**

The snapshot has been deleted.

## **gpfs\_get\_snapname\_from\_fssnaphandle()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_getacl() subroutine

Retrieves the access control information for a GPFS file.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_getacl(const char *pathname,
 int flags,
 void *acl);
```

### Description

The **gpfs\_getacl()** subroutine, together with the **gpfs\_putacl()** subroutine, is intended for use by a backup program to save (**gpfs\_getacl()**) and restore (**gpfs\_putacl()**) the ACL information for the file.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### pathname

The path identifying the file for which the ACLs are being obtained.

#### flags

Consists of one of these values:

0 Indicates that the **acl** parameter is to be mapped with the **gpfs\_opaque\_acl\_t** structure.

The **gpfs\_opaque\_acl\_t** structure should be used by backup and restore programs.

#### GPFS\_GETACL\_STRUCT

Indicates that the **acl** parameter is to be mapped with the **gpfs\_acl\_t** structure.

The **gpfs\_acl\_t** structure is provided for applications that need to interpret the ACL.

#### acl

Pointer to a buffer mapped by the structure **gpfs\_opaque\_acl\_t** or **gpfs\_acl\_t**, depending on the value of **flags**.

The first four bytes of the buffer must contain its total size.

### Exit status

If the **gpfs\_getacl()** subroutine is successful, it returns a value of 0.

If the **gpfs\_getacl()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### Exceptions

None.

### Error status

Error codes include but are not limited to the following:

## **gpfs\_getacl()**

### **EINVAL**

The path name does not refer to a GPFS file or a regular file.

### **ENOMEM**

Unable to allocate memory for the request.

### **ENOTDIR**

File is not a directory.

### **ENOSPC**

The buffer is too small to return the entire ACL. The required buffer size is returned in the first four bytes of the buffer pointed to by **acl**.

### **ENOSYS**

The **gpfs\_getacl()** subroutine is not supported under the current file system format.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_iattr\_t structure

Contains attributes of a GPFS inode.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Structure

```
typedef struct gpfs_iattr
{
 int ia_version; /* this struct version */
 int ia_reclen; /* sizeof this structure */
 int ia_checksum; /* validity check on iattr struct */
 gpfs_mode_t ia_mode; /* access mode */
 gpfs_uid_t ia_uid; /* owner uid */
 gpfs_gid_t ia_gid; /* owner gid */
 gpfs_ino_t ia_inode; /* file inode number */
 gpfs_gen_t ia_gen; /* inode generation number */
 gpfs_nlink_t ia_nlink; /* number of links */
 short ia_flags; /* Flags (defined below) */
 int ia_blocksize; /* preferred block size for io */
 gpfs_mask_t ia_mask; /* Initial attribute mask (not used) */
 unsigned int ia_pad1; /* reserved space */
 gpfs_off64_t ia_size; /* file size in bytes */
 gpfs_off64_t ia_blocks; /* 512 byte blocks of disk held by file */
 gpfs_timestruc_t ia_atime; /* time of last access */
 gpfs_timestruc_t ia_mtime; /* time of last data modification */
 gpfs_timestruc_t ia_ctime; /* time of last status change */
 gpfs_dev_t ia_rdev; /* id of device */
 unsigned int ia_xperm; /* extended attributes (defined below) */
 unsigned int ia_modsnapid; /* snapshot id of last modification */
 unsigned int ia_filesetid; /* fileset ID */
 unsigned int ia_datapoolid; /* storage pool ID for data */
 unsigned int ia_pad2; /* reserved space */
} gpfs_iattr_t;

/* Define flags for inode attributes */
#define GPFS_IAFLAG_SNAPDIR 0x0001 /* (obsolete) */
#define GPFS_IAFLAG_USRQUOTA 0x0002 /* inode is a user quota file */
#define GPFS_IAFLAG_GRPQUOTA 0x0004 /* inode is a group quota file */
#define GPFS_IAFLAG_ERROR 0x0008 /* error reading inode */

/* Define flags for inode replication attributes */
#define GPFS_IAFLAG_FILESET_ROOT 0x0010 /* root dir of a fileset */
#define GPFS_IAFLAG_NO_SNAP_RESTORE 0x0020 /* don't restore from snapshots */
#define GPFS_IAFLAG_FILESETQUOTA 0x0040 /* inode is a fileset quota file */
#define GPFS_IAFLAG_COMANAGED 0x0080 /* file data is co-managed */
#define GPFS_IAFLAG_ILLPLACED 0x0100 /* may not be properly placed */
#define GPFS_IAFLAG_REPLMETA 0x0200 /* metadata replication set */
#define GPFS_IAFLAG_REPLDATA 0x0400 /* data replication set */
#define GPFS_IAFLAG_EXPOSED 0x0800 /* may have data on suspended disks */
#define GPFS_IAFLAG_ILLREPLICATED 0x1000 /* may not be properly replicated */
#define GPFS_IAFLAG_UNBALANCED 0x2000 /* may not be properly balanced */
#define GPFS_IAFLAG_DATAUPDATEMISS 0x4000 /* has stale data blocks on
 unavailable disk */
#define GPFS_IAFLAG_METAUPDATEMISS 0x8000 /* has stale metadata on
 unavailable disk */

#define GPFS_IAFLAG_IMMUTABLE 0x00010000 /* Immutability */
#define GPFS_IAFLAG_INDEFRETENT 0x00020000 /* Indefinite retention */
#define GPFS_IAFLAG_SECUREDELETE 0x00040000 /* Secure deletion */

#define GPFS_IAFLAG_TRUNCMANAGED 0x00080000 /* dmapi truncate event enabled */
#define GPFS_IAFLAG_READMANAGED 0x00100000 /* dmapi read event enabled */
#define GPFS_IAFLAG_WRITEMANAGED 0x00200000 /* dmapi write event enabled */
```

## **gpfs\_iattr\_t**

```
#define GPFS_IAFLAG_APPENDONLY 0x00400000 /* AppendOnly only */
#define GPFS_IAFLAG_DELETED 0x00800000 /* inode has been deleted */

/* Define flags for extended attributes */
#define GPFS_IAXPERM_ACL 0x0001 /* file has acls */
#define GPFS_IAXPERM_XATTR 0x0002 /* file has extended attributes */
#define GPFS_IAXPERM_DMATTR 0x0004 /* file has dm attributes */
#define GPFS_IAXPERM_DOSATTR 0x0008 /* file has non-default dos attrs */
#define GPFS_IAXPERM_RPATTR 0x0010 /* file has restore policy attrs */
```

## **Description**

The **gpfs\_iattr\_t** structure contains the various attributes of a GPFS inode.

## **Members**

### **ia\_version**

The version number of this structure.

### **ia\_reclen**

The size of this structure.

### **ia\_checksum**

The checksum for this **gpfs\_iattr** structure.

### **ia\_mode**

The access mode for this inode.

### **ia\_uid**

The owner user ID for this inode.

### **ia\_gid**

The owner group ID for this inode.

### **ia\_inode**

The file inode number.

### **ia\_gen**

The inode generation number.

### **ia\_nlink**

The number of links for this inode.

### **ia\_flags**

The flags defined for inode attributes.

### **ia\_blocksize**

The preferred block size for I/O.

### **ia\_mask**

The initial attribute mask (not used).

### **ia\_pad1**

Reserved space.

### **ia\_size**

The file size in bytes.

### **ia\_blocks**

The number of 512 byte blocks of disk held by the file.

### **ia\_atime**

The time of last access.

**ia\_mtime**

The time of last data modification.

**ia\_ctime**

The time of last status change.

**ia\_rdev**

The ID of the device.

**ia\_xperm**

Indicator - nonzero if file has extended ACL.

**ia\_modsnapid**

Internal snapshot ID indicating the last time that the file was modified. Internal snapshot IDs for the current snapshots are displayed by the **mmlssnapshot** command.

**ia\_filesetid**

The fileset ID for the inode.

**ia\_datapoolid**

The storage pool ID for data for the inode.

**ia\_pad2**

Reserved space.

## Examples

For an example using **gpfs\_iattr\_t**, see `/usr/lpp/mmfs/samples/util/tsinode.c`.

## Location

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## gpfs\_iattr64\_t

---

### gpfs\_iattr64\_t structure

Contains attributes of a GPFS inode.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Structure

```
typedef struct gpfs_iattr64
{
 int ia_version; /* this struct version */
 int ia_reclen; /* sizeof this structure */
 int ia_checksum; /* validity check on iattr struct */
 gpfs_mode_t ia_mode; /* access mode */
 gpfs_uid64_t ia_uid; /* owner uid */
 gpfs_gid64_t ia_gid; /* owner gid */
 gpfs_ino64_t ia_inode; /* file inode number */
 gpfs_gen64_t ia_gen; /* inode generation number */
 gpfs_nlink64_t ia_nlink; /* number of links */
 gpfs_off64_t ia_size; /* file size in bytes */
 gpfs_off64_t ia_blocks; /* 512 byte blocks of disk held by file */
 gpfs_timestruc64_t ia_atime; /* time of last access */
 unsigned int ia_winflags; /* windows flags (defined below) */
 unsigned int ia_pad1; /* reserved space */
 gpfs_timestruc64_t ia_mtime; /* time of last data modification */
 unsigned int ia_flags; /* flags (defined below) */

/* next four bytes were ia_pad2 */
 unsigned char ia_repl_data; /* data replication factor */
 unsigned char ia_repl_data_max; /* data replication max factor */
 unsigned char ia_repl_meta; /* meta data replication factor */
 unsigned char ia_repl_meta_max; /* meta data replication max factor */
 gpfs_timestruc64_t ia_ctime; /* time of last status change */
 int ia_blocksize; /* preferred block size for io */
 unsigned int ia_pad3; /* reserved space */
 gpfs_timestruc64_t ia_CREATETIME; /* creation time */
 gpfs_mask_t ia_mask; /* initial attribute mask (not used) */
 int ia_pad4; /* reserved space */
 unsigned int ia_reserved[GPFS_IA64_RESERVED]; /* reserved space */
 unsigned int ia_xperm; /* extended attributes (defined below) */
 gpfs_dev_t ia_dev; /* id of device containing file */
 gpfs_dev_t ia_rdev; /* device id (if special file) */
 unsigned int ia_pcacheflags; /* pcache inode bits */
 gpfs_snapid64_t ia_modsnapid; /* snapshot id of last modification */
 unsigned int ia_filesetid; /* filesset ID */
 unsigned int ia_datapoolid; /* storage pool ID for data */
 gpfs_ino64_t ia_inode_space_mask; /* inode space mask of this file system */
 /* This value is saved in the iattr structure
 during backup and used during restore */
 gpfs_off64_t ia_dirminsize; /* dir pre-allocation size in bytes */
 unsigned int ia_unused[GPFS_IA64_UNUSED]; /* reserved space */
} gpfs_iattr64_t;

#ifndef GPFS_64BIT_INODES
#define GPFS_IA_VERSION
#define GPFS_IA_VERSION GPFS_IA_VERSION64
#define gpfs_iattr_t gpfs_iattr64_t
#endif

/* Define flags for inode attributes */
#define GPFS_IAFLAG_SNAPDIR 0x0001 /* (obsolete) */
#define GPFS_IAFLAG_USRQUOTA 0x0002 /* inode is a user quota file */
#define GPFS_IAFLAG_GRPQUOTA 0x0004 /* inode is a group quota file */
#define GPFS_IAFLAG_ERROR 0x0008 /* error reading inode */
```

```

/* Define flags for inode replication attributes */
#define GPFS_IAFLAG_FILESET_ROOT 0x0010 /* root dir of a fileset */
#define GPFS_IAFLAG_NO_SNAP_RESTORE 0x0020 /* don't restore from snapshots */
#define GPFS_IAFLAG_FILESETQUOTA 0x0040 /* inode is a fileset quota file */
#define GPFS_IAFLAG_COMANAGED 0x0080 /* file data is co-managed */
#define GPFS_IAFLAG_ILLPLACED 0x0100 /* may not be properly placed */
#define GPFS_IAFLAG_REPLMETA 0x0200 /* metadata replication set */
#define GPFS_IAFLAG_REPLDATA 0x0400 /* data replication set */
#define GPFS_IAFLAG_EXPOSED 0x0800 /* may have data on suspended disks */
#define GPFS_IAFLAG_ILLREPLICATED 0x1000 /* may not be properly replicated */
#define GPFS_IAFLAG_UNBALANCED 0x2000 /* may not be properly balanced */
#define GPFS_IAFLAG_DATAUPDATEMISS 0x4000 /* has stale data blocks on
 unavailable disk */
#define GPFS_IAFLAG_METAUPDATEMISS 0x8000 /* has stale metadata on
 unavailable disk */

#define GPFS_IAFLAG_IMMUTABLE 0x00010000 /* Immutability */
#define GPFS_IAFLAG_INDEFRETENT 0x00020000 /* Indefinite retention */
#define GPFS_IAFLAG_SECUREDELETE 0x00040000 /* Secure deletion */

#define GPFS_IAFLAG_TRUNCMANAGED 0x00080000 /* dmapi truncate event enabled */
#define GPFS_IAFLAG_READMANAGED 0x00100000 /* dmapi read event enabled */
#define GPFS_IAFLAG_WRITEMANAGED 0x00200000 /* dmapi write event enabled */

#define GPFS_IAFLAG_APPENDONLY 0x00400000 /* AppendOnly only */
#define GPFS_IAFLAG_DELETED 0x00800000 /* inode has been deleted */
#ifndef ZIP
#define GPFS_IAFLAG_ILLCOMPRESSED 0x01000000 /* may not be properly compressed */
#endif
#define GPFS_IAFLAG_FPOILLPLACED 0x02000000 /* may not be properly placed per
 FPO attributes (bgf, wad, wadfg) */

/* Define flags for window's attributes */
#define GPFS_IWINFLAG_ARCHIVE 0x0001 /* Archive */
#define GPFS_IWINFLAG_HIDDEN 0x0002 /* Hidden */
#define GPFS_IWINFLAG_NOTINDEXED 0x0004 /* Not content indexed */
#define GPFS_IWINFLAG_OFFLINE 0x0008 /* Off-line */
#define GPFS_IWINFLAG_READONLY 0x0010 /* Read-only */
#define GPFS_IWINFLAG_REPARSE 0x0020 /* Reparse point */
#define GPFS_IWINFLAG_SYSTEM 0x0040 /* System */
#define GPFS_IWINFLAG_TEMPORARY 0x0080 /* Temporary */
#define GPFS_IWINFLAG_COMPRESSED 0x0100 /* Compressed */
#define GPFS_IWINFLAG_ENCRYPTED 0x0200 /* Encrypted */
#define GPFS_IWINFLAG_SPARSE 0x0400 /* Sparse file */
#define GPFS_IWINFLAG_HASSTREAMS 0x0800 /* Has streams */

/* Define flags for extended attributes */
#define GPFS_IAXPERM_ACL 0x0001 /* file has acls */
#define GPFS_IAXPERM_XATTR 0x0002 /* file has extended attributes */
#define GPFS_IAXPERM_DMATTR 0x0004 /* file has dm attributes */
#define GPFS_IAXPERM_DOSATTR 0x0008 /* file has non-default dos attrs */
#define GPFS_IAXPERM_RPATTR 0x0010 /* file has restore policy attrs */

/* Define flags for pcache bits defined in the inode */
#define GPFS_ICAFLAG_CACHED 0x0001 /* "cached complete" */
#define GPFS_ICAFLAG_CREATE 0x0002 /* "created" */
#define GPFS_ICAFLAG_DIRTY 0x0004 /* "data dirty" */
#define GPFS_ICAFLAG_LINK 0x0008 /* "hard linked" */
#define GPFS_ICAFLAG_SETATTR 0x0010 /* "attr changed" */
#define GPFS_ICAFLAG_LOCAL 0x0020 /* "local" */
#define GPFS_ICAFLAG_APPEND 0x0040 /* "append" */
#define GPFS_ICAFLAG_STATE 0x0080 /* "has remote state" */

```

## Description

The `gpfs_iattr64_t` structure contains the various attributes of a GPFS inode.

## **gpfs\_iattr64\_t**

### **Members**

#### **ia\_version**

The version number of this structure.

#### **ia\_reclen**

The size of this structure.

#### **ia\_checksum**

The checksum for this **gpfs\_iattr64** structure.

#### **ia\_mode**

The access mode for this inode.

#### **ia\_uid**

The owner user ID for this inode.

#### **ia\_gid**

The owner group ID for this inode.

#### **ia\_inode**

The file inode number.

#### **ia\_gen**

The inode generation number.

#### **ia\_nlink**

The number of links for this inode.

#### **ia\_size**

The file size in bytes.

#### **ia\_blocks**

The number of 512 byte blocks of disk held by the file.

#### **ia\_atime**

The time of last access.

#### **ia\_winflags**

The Windows flags.

#### **ia\_pad1**

Reserved space.

#### **ia\_mtime**

The time of last data modification.

#### **ia\_flags**

The flags defined for inode attributes.

#### **ia\_repl\_data**

The data replication factor.

#### **ia\_repl\_data\_max**

The maximum data replication factor.

#### **ia\_repl\_meta**

The metadata replication factor.

#### **ia\_repl\_meta\_max**

The maximum metadata replication factor.

#### **ia\_ctime**

The time of last status change.

#### **ia\_blocksize**

The preferred block size for I/O.

**ia\_pad3**

Reserved space.

**ia\_createtime**

The creation time.

**ia\_mask**

The initial attribute mask (not used).

**ia\_pad4**

Reserved space.

**ia\_reserved**

Reserved space.

**ia\_xperm**

Indicator - nonzero if file has extended ACL.

**ia\_dev**

The ID of the device containing the file.

**ia\_rdev**

The ID of the device.

**ia\_pcacheflags**

The pcache inode bits.

**ia\_modsnapid**

Internal snapshot ID indicating the last time that the file was modified. Internal snapshot IDs for the current snapshots are displayed by the **mmlssnapshot** command.

**ia\_filesetid**

The fileset ID for the inode.

**ia\_datapoolid**

The storage pool ID for data for the inode.

**ia\_dirminsize**

Directory preallocation size in bytes.

**ia\_inode\_space\_mask**

The inode space mask of this file system. This value is saved in the **iattr** structure during backup and used during restore.

**ia\_unused**

Reserved space.

## Examples

See the **gpfs\_iattr\_t** example in /usr/lpp/mmfs/samples/util/tsinode.c.

## Location

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_iclose()**

---

### **gpfs\_iclose() subroutine**

Closes a file given its inode file handle.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
void gpfs_iclose(gpfs_ifile_t *ifile);
```

#### **Description**

The **gpfs\_iclose()** subroutine closes an open file descriptor created by **gpfs\_iopen()**.

For an overview of using **gpfs\_iclose()** in a backup application, see the topic *Using APIs to develop backup applications* in the *IBM Spectrum Scale: Administration Guide*

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **ifile**

Pointer to **gpfs\_ifile\_t** from **gpfs\_iopen()**.

#### **Exit status**

The **gpfs\_iclose()** subroutine returns void.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **ENOSYS**

The **gpfs\_iclose()** subroutine is not available.

##### **EPERM**

The caller does not have superuser privileges.

##### **ESTALE**

Cached file system information was not valid.

#### **Examples**

For an example using **gpfs\_iclose()**, see `/usr/lpp/mmfs/samples/util/tsreaddir.c`.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX  
/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_ifile\_t**

---

### **gpfs\_ifile\_t structure**

Contains a handle for a GPFS inode.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct gpfs_ifile gpfs_ifile_t;
```

#### **Description**

The **gpfs\_ifile\_t** structure contains a handle for the file of a GPFS inode.

#### **Members**

##### **gpfs\_ifile**

The handle for the file of a GPFS inode.

#### **Examples**

For an example using **gpfs\_ifile\_t**, see `/usr/lpp/mmfs/samples/util/tsfindinode.c`.

#### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_igetattrs() subroutine

Retrieves extended file attributes in opaque format.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_igetattrs(gpfs_ifile_t *ifile,
 void *buffer,
 int bufferSize,
 int *attrSize);
```

### Description

The **gpfs\_igetattrs()** subroutine retrieves all extended file attributes in opaque format. This subroutine is intended for use by a backup program to save all extended file attributes (ACLs, attributes, and so forth). If the file does not have any extended attributes, the subroutine sets **attrSize** to zero.

#### Notes:

1. This call does not return extended attributes used for the Data Storage Management (XDSM) API (also known as DMAP).
2. Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:
  - libgpfs.a for AIX
  - libgpfs.so for Linux

### Parameters

#### **ifile**

Pointer to **gpfs\_ifile\_t** from **gpfs\_iopen()**.

#### **buffer**

Pointer to buffer for returned attributes.

#### **bufferSize**

Size of the buffer.

#### **attrSize**

Pointer to returned size of attributes.

### Exit status

If the **gpfs\_igetattrs()** subroutine is successful, it returns a value of 0.

If the **gpfs\_igetattrs()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### Exceptions

None.

### Error status

Error codes include but are not limited to the following:

## **gpfs\_igetattrs()**

### **ENOSPC**

The buffer is too small to return all attributes. Field **attrSize** will be set to the size necessary.

### **ENOSYS**

The **gpfs\_igetattrs()** subroutine is not available.

### **EPERM**

The caller does not have superuser privileges.

### **ESTALE**

Cached file system information was not valid.

### **GPFS\_EINVAL\_IFILE**

Incorrect **ifile** parameters.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_igetattrsx() subroutine

Retrieves extended file attributes; provides an option to include DMAPI attributes.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_igetattrsx(gpfs_ifile_t *ifile,
 int flags,
 void *buffer,
 int bufferSize,
 int *attrSize);
```

### Description

The **gpfs\_igetattrsx()** subroutine retrieves all extended file attributes in opaque format. It provides the same function as **gpfs\_igetattrs()** but includes a **flags** parameter that allows the caller to back up and restore DMAPI attributes.

This function is intended for use by a backup program to save (and restore, using the related subroutine **gpfs\_iputattrsx()**) all extended file attributes (ACLs, user attributes, and so forth) in one call. If the file does not have any extended attributes, the subroutine sets **attrSize** to zero.

#### Notes:

1. This call can optionally return extended attributes used for the Data Storage Management (XDSM) API (also known as DMAPI).
2. Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:
  - libgpfs.a for AIX
  - libgpfs.so for Linux

### Parameters

#### **ifile**

Pointer to **gpfs\_ifile\_t** from **gpfs\_iopen()**.

#### **flags**

Flags must have one of the following values:

##### **GPFS\_ATTRFLAG\_NO\_PLACEMENT**

File attributes for placement are not saved, and neither is the current storage pool.

##### **GPFS\_ATTRFLAG\_IGNORE\_PLACEMENT**

File attributes for placement are saved, but the current storage pool is not.

##### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

File attributes for DMAPI are included in the returned buffer.

##### **GPFS\_ATTRFLAG\_USE\_POLICY**

Uses the restore policy rules to determine the pool ID.

##### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

Includes the DMAPI attributes.

##### **GPFS\_ATTRFLAG\_FINALIZE\_ATTRS**

Finalizes immutability attributes.

## **gpfs\_igetattrsx()**

### **GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**

Skips immutable attributes.

### **GPFS\_ATTRFLAG\_INCL\_ENCR**

Includes encryption attributes.

### **GPFS\_ATTRFLAG\_SKIP\_CLONE**

Skips clone attributes.

### **GPFS\_ATTRFLAG MODIFY\_CLONEPARENT**

Allows modification on the clone parent.

### **buffer**

A pointer to the buffer for returned attributes.

### **bufferSize**

Size of the buffer.

### **attrSize**

Pointer to returned size of attributes.

## **Exit status**

If the **gpfs\_igetattrsx()** subroutine is successful, it returns a value of 0.

If the **gpfs\_igetattrsx()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EINVAL**

Not a GPFS file, or the flags provided are not valid.

### **ENOSPC**

The buffer is too small to return all attributes. Field **attrSize** will be set to the size necessary.

### **ENOSYS**

The **gpfs\_igetattrsx()** subroutine is not available.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_igetfilesetname() subroutine**

Returns the name of the fileset defined by a fileset ID.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_igetfilesetname(gpfs_iscan_t *iscan,
 unsigned int fileId,
 void *buffer,
 int bufferSize);
```

### **Description**

The **gpfs\_igetfilesetname()** subroutine is part of the backup by inode interface. The caller provides a pointer to the scan descriptor used to obtain the fileset ID. This library routine will return the name of the fileset defined by the fileset ID. The name is the null-terminated string provided by the administrator when the fileset was defined. The maximum string length is **GPFS\_MAXNAMLEN**, which is defined in /usr/lpp/mmfs/include/gpfs.h.

#### **Notes:**

1. This routine is not thread safe. Only one thread at a time is allowed to invoke this routine for the given scan descriptor.
2. Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:
  - libgpfs.a for AIX
  - libgpfs.so for Linux

### **Parameters**

#### **iscan**

Pointer to **gpfs\_iscan\_t** used to obtain the fileset ID.

#### **filesetId**

The fileset ID.

#### **buffer**

Pointer to buffer for returned attributes.

#### **bufferSize**

Size of the buffer.

### **Exit status**

If the **gpfs\_igetfilesetname()** subroutine is successful, it returns a value of 0.

If the **gpfs\_igetfilesetname()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

## **gpfs\_igetfilesetname()**

### **Error status**

Error codes include but are not limited to the following:

**E2BIG** The buffer is too small to return the fileset name.

**EINTR**

The call was interrupted. This routine is not thread safe.

**EINVAL**

The fileset ID is not valid.

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The `gpfs_igetfilesetname()` subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ESTALE**

The cached file system information was not valid.

**GPFS\_E\_INVAL\_ISCAN**

The `iscan` parameters were not valid.

### **Examples**

This programming segment gets the fileset name based on the given fileset ID. The returned fileset name is stored in `FileSetNameBuffer`, which has a length of `FileSetNameSize`.

```
gpfs_iscan_t *fsInodeScanP;
gpfs_igetfilesetname(fsInodeScanP, FileSetId, &FileSetNameBuffer, FileSetNameSize);
```

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_igetstoragepool() subroutine**

Returns the name of the storage pool for the given storage pool ID.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_igetstoragepool(gpfs_iscan_t *iscan,
 unsigned int dataPoolId,
 void *buffer,
 int bufferSize);
```

### **Description**

The **gpfs\_igetstoragepool()** subroutine is part of the backup by inode interface. The caller provides a pointer to the scan descriptor used to obtain the storage pool ID. This routine returns the name of the storage pool for the given storage pool ID. The name is the null-terminated string provided by the administrator when the storage pool was defined. The maximum string length is **GPFS\_MAXNAMLEN**, which is defined in /usr/lpp/mmfs/include/gpfs.h.

#### **Notes:**

1. This routine is not thread safe. Only one thread at a time is allowed to invoke this routine for the given scan descriptor.
2. Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:
  - libgpfs.a for AIX
  - libgpfs.so for Linux

### **Parameters**

#### **iscan**

Pointer to **gpfs\_iscan\_t** used to obtain the storage pool ID.

#### **dataPoolId**

The storage pool ID.

#### **buffer**

Pointer to buffer for returned attributes.

#### **bufferSize**

Size of the buffer.

### **Exit status**

If the **gpfs\_igetstoragepool()** subroutine is successful, it returns a value of 0.

If the **gpfs\_igetstoragepool()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

## **gpfs\_igetstoragepool()**

### **Error status**

Error codes include but are not limited to the following:

**E2BIG** The buffer is too small to return the storage pool name.

**EINTR** The call was interrupted. This routine is not thread safe.

**EINVAL**

The storage pool ID is not valid.

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The `gpfs_igetstoragepool()` subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ESTALE**

The cached storage pool information was not valid.

**GPFS\_E\_INVAL\_ISCAN**

The `iscan` parameters were not valid.

### **Examples**

This programming segment gets the storage pool name based on the given storage pool ID. The returned storage pool name is stored in **StoragePoolNameBuffer** which has the length of **StoragePoolNameSize**.

```
gpfs_iscan_t *fsInodeScanP;
gpfs_igetstoragepool(fsInodeScanP,StgpoolIdBuffer, &StgpoolNameBuffer,StgpoolNameSize);
```

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_iopen() subroutine

Opens a file or directory by inode number.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
gpfs_ifile_t *gpfs_iopen(gpfs_fssnap_handle_t *fssnapHandle,
 gpfs_ino_t ino,
 int open_flags,
 const gpfs_iattr_t *statxbuf,
 const char *symLink);
```

### Description

The **gpfs\_iopen()** subroutine opens a user file or directory for backup. The file is identified by its inode number **ino** within the file system or snapshot identified by the **fssnapHandle**. The **fssnapHandle** parameter must be the same one that was used to create the inode scan that returned the inode number **ino**.

To read the file or directory, the **open\_flags** must be set to **GPFS\_O\_BACKUP**. The **statxbuf** and **symLink** parameters are reserved for future use and must be set to NULL.

For an overview of using **gpfs\_iopen()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### fssnapHandle

File system snapshot handle.

#### ino

The inode number.

#### open\_flags

##### GPFS\_O\_BACKUP

Read files for backup.

##### O\_RDONLY

For **gpfs\_iread()**.

#### statxbuf

This parameter is reserved for future use and should always be set to NULL.

#### symLink

This parameter is reserved for future use and should always be set to NULL.

### Exit status

If the **gpfs\_iopen()** subroutine is successful, it returns a pointer to the inode's file handle.

## **gpfs\_iopen()**

If the **gpfs\_iopen()** subroutine is unsuccessful, it returns NULL and the global error variable **errno** is set to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EINVAL**

Missing or incorrect parameter.

### **ENOENT**

The file does not exist in the file system.

### **ENOMEM**

Unable to allocate memory for the request.

### **ENOSYS**

The **gpfs\_iopen()** subroutine is not available.

### **EPERM**

The caller does not have superuser privileges.

### **ESTALE**

Cached file system information was not valid.

### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **GPFS\_E\_INVAL\_INUM**

Users are not authorized to open the reserved inodes.

## **Examples**

For an example using **gpfs\_iopen()**, see `/usr/lpp/mmfs/samples/util/tsreaddir.c`.

## **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_iopen64() subroutine

Opens a file or directory by inode number.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
gpfs_ifile_t *gpfs_iopen64(gpfs_fssnap_handle_t *fssnapHandle,
 gpfs_ino64_t ino,
 int open_flags,
 const gpfs_iattr64_t *statxbuf,
 const char *symLink);
```

### Description

The **gpfs\_iopen64()** subroutine opens a user file or directory for backup. The file is identified by its inode number **ino** within the file system or snapshot identified by the **fssnapHandle**. The **fssnapHandle** parameter must be the same one that was used to create the inode scan that returned the inode number **ino**.

To read the file or directory, the **open\_flags** must be set to **GPFS\_O\_BACKUP**. The **statxbuf** and **symLink** parameters are reserved for future use and must be set to NULL.

For an overview of using **gpfs\_iopen64()** in a backup application, see *Using APIs to develop backup applications in IBM Spectrum Scale: Administration Guide*.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### fssnapHandle

The file system snapshot handle.

#### ino

The inode number.

#### open\_flags

##### GPFS\_O\_BACKUP

Read files for backup.

##### O\_RDONLY

For **gpfs\_iread()**.

#### statxbuf

This parameter is reserved for future use and should always be set to NULL.

#### symLink

This parameter is reserved for future use and should always be set to NULL.

### Exit status

If the **gpfs\_iopen64()** subroutine is successful, it returns a pointer to the inode's file handle.

## **gpfs\_iopen64()**

If the **gpfs\_iopen64()** subroutine is unsuccessful, it returns NULL and the global error variable **errno** is set to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EFORMAT**

The file system version number is not valid.

### **EINVAL**

Missing or incorrect parameter.

### **ENOENT**

The file does not exist in the file system.

### **ENOMEM**

Unable to allocate memory for the request.

### **ENOSYS**

The **gpfs\_iopen64()** subroutine is not available.

### **EPERM**

The caller does not have superuser privileges.

### **ESTALE**

Cached file system information was not valid.

### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **GPFS\_E\_INVAL\_IATTR**

The **iattr** structure was corrupted.

### **GPFS\_E\_INVAL\_INUM**

Users are not authorized to open the reserved inodes.

**Note:** **gpfs\_iopen64()** calls the standard library subroutines **dup()**, **open()**, and **malloc()**; if one of these called subroutines returns an error, **gpfs\_iopen64()** also returns that error.

## **Examples**

See the **gpfs\_iopen()** example in `/usr/lpp/mmfs/samples/util/tsreaddir.c`.

## **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_iputattrsx() subroutine

Sets the extended file attributes for a file.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_iputattrsx(gpfs_ifile_t *ifile,
 int flags,
 void *buffer,
 const char *pathName);
```

### Description

The **gpfs\_iputattrsx()** subroutine, together with **gpfs\_igetattrsx()**, is intended for use by a backup program to save (**gpfs\_igetattrsx()**) and restore (**gpfs\_iputattrsx()**) all of the extended attributes of a file. This subroutine also sets the storage pool for the file and sets data replication to the values that are saved in the extended attributes.

This subroutine can optionally invoke the policy engine to match a **RESTORE** rule using the file's attributes saved in the extended attributes to set the file's storage pool and data replication as when calling **gpfs\_fputattrswithpathname()**. When used with the policy engine, the caller should include the full path to the file, including the file name, to allow rule selection based on file name or path.

By default, the routine will not use **RESTORE** policy rules for data placement. The **pathName** parameter will be ignored and may be set to **NULL**.

If the call does not use **RESTORE** policy rules, or if the file fails to match a **RESTORE** rule, or if there are not **RESTORE** rules installed, then the storage pool and data replication are selected as when calling **gpfs\_fputattrs()**.

The buffer passed in should contain extended attribute data that was obtained by a previous call to **gpfs\_fgetattrs()**.

**Note:** This call will restore extended attributes used for the Data Storage Management (XDSM) API (also known as DMAP) if they are present in the buffer.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### ifile

A pointer to **gpfs\_ifile\_t** from **gpfs\_iopen()**.

#### flags

Flags must have one of the following values:

##### **GPFS\_ATTRFLAG\_NO\_PLACEMENT**

File attributes are restored, but the storage pool and data replication are unchanged.

##### **GPFS\_ATTRFLAG\_IGNORE\_POOL**

File attributes are restored, but the storage pool and data replication are selected by matching the saved attributes to a placement rule instead of restoring the saved storage pool.

## **gpfs\_iputattrsx()**

### **GPFS\_ATTRFLAG\_USE\_POLICY**

File attributes are restored, but the storage pool and data replication are selected by matching the saved attributes to a RESTORE rule instead of restoring the saved storage pool.

### **GPFS\_ATTRFLAG\_USE\_POLICY**

Uses the restore policy rules to determine the pool ID.

### **GPFS\_ATTRFLAG\_INCL\_DMAPI**

Includes the DMAP API attributes.

### **GPFS\_ATTRFLAG\_FINALIZE\_ATTRS**

Finalizes immutability attributes.

### **GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**

Skips immutable attributes.

### **GPFS\_ATTRFLAG\_INCL\_ENCR**

Includes encryption attributes.

### **GPFS\_ATTRFLAG\_SKIP\_CLONE**

Skips clone attributes.

### **GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT**

Allows modification on the clone parent.

### **buffer**

A pointer to the buffer containing the extended attributes for the file.

### **pathName**

A pointer to a file path and file name. NULL is a valid value for **pathName**.

**Note:** **pathName** is a UTF-8 encoded string. On Windows, applications can convert UTF-16 (Unicode) to UTF-8 using the platform's **WideCharToMultiByte** function.

## **Exit status**

If the **gpfs\_iputattrsx()** subroutine is successful, it returns a value of 0.

If the **gpfs\_iputattrsx()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **Exceptions**

None.

## **Error status**

Error codes include but are not limited to the following:

### **EINVAL**

The buffer pointed to by **buffer** does not contain valid attribute data, or invalid flags were provided.

### **ENOSYS**

The **gpfs\_iputattrsx()** subroutine is not supported under the current file system format.

### **EPERM**

The caller of the subroutine must have superuser privilege.

### **ESTALE**

The cached *fs* information was not valid.

**GPFS\_EINVAL\_IFILE**

The **ifile** parameters provided were not valid.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## gpfs\_iread()

---

### gpfs\_iread() subroutine

Reads a file opened by `gpfs_iopen()`.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Synopsis

```
#include <gpfs.h>
int gpfs_iread(gpfs_ifile_t *ifile,
 void *buffer,
 int bufferSize,
 gpfs_off64_t *offset);
```

#### Description

The `gpfs_iread()` subroutine reads data from the file indicated by the `ifile` parameter returned from `gpfs_iopen()`. This subroutine reads data beginning at parameter `offset` and continuing for `bufferSize` bytes into the buffer specified by `buffer`. If successful, the subroutine returns a value that is the length of the data read, and sets parameter `offset` to the offset of the next byte to be read. A return value of 0 indicates end-of-file.

For an overview of using `gpfs_iread()` in a backup application, see *Using APIs to develop backup applications in IBM Spectrum Scale: Administration Guide*.

**Note:** Compile any program that uses this subroutine with the `-lgpfs` flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### Parameters

##### `ifile`

Pointer to `gpfs_ifile_t` from `gpfs_iopen()`.

##### `buffer`

Buffer for the data to be read.

##### `bufferSize`

Size of the buffer (that is, the amount of data to be read).

##### `offset`

Offset of where within the file to read. If `gpfs_iread()` is successful, `offset` is updated to the next byte after the last one that was read.

#### Exit status

If the `gpfs_iread()` subroutine is successful, it returns the number of bytes read.

If the `gpfs_iread()` subroutine is unsuccessful, it returns a value of -1 and sets the global error variable `errno` to indicate the nature of the error.

#### Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EISDIR

The specified file is a directory.

### EINVAL

Missing or incorrect parameter.

### ENOSYS

The `gpfs_iread()` subroutine is not available.

### EPERM

The caller does not have superuser privileges.

### ESTALE

Cached file system information was not valid.

### GPFS\_E\_INVAL\_IFILE

Incorrect `ifile` parameter.

## Location

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_ireaddir()**

---

### **gpfs\_ireaddir() subroutine**

Reads the next directory entry.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_ireaddir(gpfs_ifile_t *idir,
 const gpfs_direntx_t **dirent);
```

#### **Description**

The **gpfs\_ireaddir()** subroutine returns the next directory entry in a file system. For an overview of using **gpfs\_ireaddir()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **idir**

Pointer to **gpfs\_ifile\_t** from **gpfs\_iopen()**.

##### **dirent**

Pointer to returned pointer to directory entry.

#### **Exit status**

If the **gpfs\_ireaddir()** subroutine is successful, it returns a value of 0 and sets the **dirent** parameter to point to the returned directory entry. If there are no more GPFS directory entries, **gpfs\_ireaddir()** returns a value of 0 and sets the **dirent** parameter to NULL.

If the **gpfs\_ireaddir()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **ENOMEM**

Unable to allocate memory for the request.

##### **ENOSYS**

The **gpfs\_ireaddir()** subroutine is not available.

##### **ENOTDIR**

File is not a directory.

##### **EPERM**

The caller does not have superuser privileges.

**ESTALE**

The cached file system information was not valid.

**GPFS\_EINVAL\_IFILE**

Incorrect **ifile** parameter.

**Examples**

For an example using **gpfs\_ireaddir()**, see */usr/lpp/mmfs/samples/util/tsreaddir.c*.

**Location**

*/usr/lpp/mmfs/lib/libgpfs.a* for AIX

*/usr/lpp/mmfs/lib/libgpfs.so* for Linux

## **gpfs\_ireaddir64()**

---

### **gpfs\_ireaddir64() subroutine**

Reads the next directory entry.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_ireaddir64(gpfs_ifile_t *idir,
 const gpfs_direntx64_t **dirent);
```

#### **Description**

The **gpfs\_ireaddir64()** subroutine returns the next directory entry in a file system. For an overview of using **gpfs\_ireaddir64()** in a backup application, see *Using APIs to develop backup applications in IBM Spectrum Scale: Administration Guide*.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **idir**

A pointer to **gpfs\_ifile\_t** from **gpfs\_iopen64()**.

##### **dirent**

A pointer to the returned pointer to the directory entry.

#### **Exit status**

If the **gpfs\_ireaddir64()** subroutine is successful, it returns a value of 0 and sets the **dirent** parameter to point to the returned directory entry. If there are no more GPFS directory entries, **gpfs\_ireaddir64()** returns a value of 0 and sets the **dirent** parameter to NULL.

If the **gpfs\_ireaddir64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **ENOMEM**

Unable to allocate memory for the request.

##### **ENOSYS**

The **gpfs\_ireaddir64()** subroutine is not available.

##### **ENOTDIR**

File is not a directory.

##### **EPERM**

The caller does not have superuser privileges.

**ESTALE**

The cached file system information was not valid.

**GPFS\_EINVAL\_IFILE**

Incorrect **ifile** parameter.

**Examples**

See the **gpfs\_ireaddir()** example in `/usr/lpp/mmfs/samples/util/tsreaddir.c`.

**Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_ireadlink()**

---

### **gpfs\_ireadlink() subroutine**

Reads a symbolic link by inode number.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_ireadlink(gpfs_fssnap_handle_t *fssnapHandle,
 gpfs_ino_t ino,
 char *buffer,
 int bufferSize);
```

#### **Description**

The **gpfs\_ireadlink()** subroutine reads a symbolic link by inode number. Like **gpfs\_iopen()**, it uses the same **fssnapHandle** parameter that was used by the inode scan.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **fssnapHandle**

File system snapshot handle.

##### **ino**

inode number of the link file to read.

##### **buffer**

Pointer to buffer for the returned link data.

##### **bufferSize**

Size of the buffer.

#### **Exit status**

If the **gpfs\_ireadlink()** subroutine is successful, it returns the number of bytes read.

If the **gpfs\_ireadlink()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **EINVAL**

Missing or incorrect parameter.

##### **ENOENT**

No such file or directory.

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The **gpfs\_ireadlink()** subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ERANGE**

On AIX, the buffer is too small to return the symbolic link.

**ESTALE**

Cached file system information was not valid.

**GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_ireadlink64()**

---

### **gpfs\_ireadlink64() subroutine**

Reads a symbolic link by inode number.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_ireadlink64(gpfs_fssnap_handle_t *fssnapHandle,
 gpfs_ino64_t ino,
 char *buffer,
 int bufferSize);
```

#### **Description**

The **gpfs\_ireadlink64()** subroutine reads a symbolic link by inode number. Like **gpfs\_iopen64()**, it uses the same **fssnapHandle** parameter that was used by the inode scan.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **fssnapHandle**

The file system snapshot handle.

##### **ino**

The inode number of the link file to read.

##### **buffer**

A pointer to buffer for the returned link data.

##### **bufferSize**

The size of the buffer.

#### **Exit status**

If the **gpfs\_ireadlink64()** subroutine is successful, it returns the number of bytes read.

If the **gpfs\_ireadlink64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **EINVAL**

Missing or incorrect parameter.

##### **ENOENT**

No such file or directory.

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The **gpfs\_ireadlink64()** subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ERANGE**

On AIX, the buffer is too small to return the symbolic link.

**ESTALE**

The cached file system information was not valid.

**GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

**Note:** **gpfs\_ireadlink64()** calls the standard library subroutine **readlink()**; if this called subroutine returns an error, **gpfs\_ireadlink64()** also returns that error.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## gpfs\_ireadx()

---

### gpfs\_ireadx() subroutine

Performs block level incremental read of a file within an incremental inode scan.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Synopsis

```
#include <gpfs.h>
gpfs_off64_t gpfs_ireadx(gpfs_ifile_t *ifile,
 gpfs_iscan_t *iscan,
 void *buffer,
 int bufferSize,
 gpfs_off64_t *offset,
 gpfs_off64_t termOffset,
 int *hole);
```

#### Description

The **gpfs\_ireadx()** subroutine performs a block level incremental read on a file opened by **gpfs\_iopen()** within a given incremental scan opened using **gpfs\_open\_inodescan()**.

For an overview of using **gpfs\_ireadx()** in a backup application, see *Using APIs to develop backup applications in IBM Spectrum Scale: Administration Guide*.

The **gpfs\_ireadx()** subroutine returns the data that has changed since the **prev\_fssnapId** specified for the inode scan. The file is scanned starting at **offset** and terminating at **termOffset**, looking for changed data. Once changed data is located, the **offset** parameter is set to its location, the new data is returned in the **buffer** provided, and the amount of data returned is the subroutine's value.

If the change to the data is that it has been deleted (that is, the file has been truncated), no data is returned, but the **hole** parameter is returned with a value of 1, and the size of the **hole** is returned as the subroutine's value. The returned size of the hole may exceed the **bufferSize** provided. If no changed data was found before reaching the **termOffset** or the end-of-file, then the **gpfs\_ireadx()** subroutine return value is 0.

Block level incremental backups are not available on small files (a file size smaller than the file system block size), directories, or if the file has been deleted. The **gpfs\_ireadx()** subroutine can still be used, but it returns all of the file's data, operating like the standard **gpfs\_iread()** subroutine. However, the **gpfs\_ireadx()** subroutine will still identify sparse files and explicitly return information on holes in the files, rather than returning the NULL data.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### Parameters

##### **ifile**

Pointer to **gpfs\_ifile\_t** returned from **gpfs\_iopen()**.

##### **iscan**

Pointer to **gpfs\_iscan\_t** from **gpfs\_open\_inodescan()**.

##### **buffer**

Pointer to buffer for returned data, or NULL to query the next increment to be read.

**bufferSize**

Size of buffer for returned data.

**offset**

On input, the offset to start the scan for changes. On output, the offset of the changed data, if any was detected.

**termOffset**

Read terminates before reading this offset. The caller may specify **ia\_size** from the file's **gpfs\_iattr\_t** or 0 to scan the entire file.

**hole**

Pointer to a flag returned to indicate a hole in the file. A value of 0 indicates that the **gpfs\_ireadx()** subroutine returned data in the **buffer**. A value of 1 indicates that **gpfs\_ireadx()** encountered a hole at the returned **offset**.

## Exit status

If the **gpfs\_ireadx()** subroutine is successful, it returns the number of bytes read and returned in **bufP**, or the size of the hole encountered in the file.

If the **gpfs\_ireadx()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

**EDOM**

The file system stripe ID from the **iscanId** does not match the **ifile**'s.

**EINVAL**

Missing or incorrect parameter.

**EISDIR**

The specified file is a directory.

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The **gpfs\_ireadx()** subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ERANGE**

The file system snapshot ID from the **iscanId** is more recent than the **ifile**'s.

**ESTALE**

Cached file system information was not valid.

**GPFS\_E\_INVAL\_IFILE**

Incorrect **ifile** parameter.

**GPFS\_E\_INVAL\_ISCAN**

Incorrect **iscan** parameter.

## **gpfs\_ireadx()**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_iscan\_t structure**

Contains a handle for an inode scan of a GPFS file system or snapshot.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Structure**

```
typedef struct gpfs_iscan gpfs_iscan_t;
```

### **Description**

The `gpfs_iscan_t` structure contains a handle for an inode scan of a GPFS file system or snapshot.

### **Members**

#### **gpfs\_iscan**

The handle for an inode scan for a GPFS file system or snapshot.

### **Examples**

For an example using `gpfs_iscan_t`, see `/usr/lpp/mmfs/samples/util/tstimes.c`.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_lib\_init()**

---

### **gpfs\_lib\_init() subroutine**

Sets up a GPFS interface for additional calls.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_lib_init(int flags);
```

#### **Description**

The **gpfs\_lib\_init()** subroutine, together with the **gpfs\_lib\_term()** subroutine, is intended for use by a program that makes repeated calls to a GPFS programming interface. This subroutine sets up the internal structure to speed up additional interface calls.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **flags**

Reserved for future use. Must be zero.

#### **Exit status**

If the **gpfs\_lib\_init()** subroutine is successful, it returns a value of 0.

If the **gpfs\_lib\_init()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

#### **Exceptions**

None.

#### **Error status**

Error codes include but are not limited to the following:

##### **EINVAL**

A nonzero value was passed as the **flags** parameter.

##### **ENOSYS**

The **gpfs\_lib\_init()** subroutine is not supported under the current file system format.

#### **Examples**

For an example using **gpfs\_lib\_init()**, see [/usr/lpp/mmfs/samples/util/tsfindinode.c](#).

#### **Location**

[/usr/lpp/mmfs/lib/libgpfs.a](#) for AIX

[/usr/lpp/mmfs/lib/libgpfs.so](#) for Linux

---

## **gpfs\_lib\_term() subroutine**

Cleans up after GPFS interface calls have been completed.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_lib_term(int flags);
```

### **Description**

The **gpfs\_lib\_term()** subroutine, together with the **gpfs\_lib\_init()** subroutine, is intended for use by a program that makes repeated calls to a GPFS programming interface. This subroutine cleans up the internal structure previously set up by **gpfs\_lib\_init()**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **flags**

Reserved for future use. Must be zero.

### **Exit status**

If the **gpfs\_lib\_term()** subroutine is successful, it returns a value of 0.

If the **gpfs\_lib\_term()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EINTR**

The **gpfs\_lib\_term()** subroutine was interrupted by a signal that was caught. Cleanup was done.

#### **EINVAL**

A nonzero value was passed as the **flags** parameter.

### **Examples**

For an example using **gpfs\_lib\_term()**, see */usr/lpp/mmfs/samples/util/tsfindinode.c*.

### **Location**

*/usr/lpp/mmfs/lib/libgpfs.a* for AIX

*/usr/lpp/mmfs/lib/libgpfs.so* for Linux

## gpfs\_next\_inode()

---

### gpfs\_next\_inode() subroutine

Retrieves the next inode from the inode scan.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Synopsis

```
#include <gpfs.h>
int gpfs_next_inode(gpfs_iscan_t *iscan,
 gpfs_ino_t termIno,
 const gpfs_iattr_t **iattr);
```

#### Description

The **gpfs\_next\_inode()** subroutine obtains the next inode from the specified inode scan and sets the **iattr** pointer to the inode's attributes. The **termIno** parameter can be used to terminate the inode scan before the last inode in the file system or snapshot being scanned. A value of 0 may be provided to indicate the last inode in the file system or snapshot. If there are no more inodes to be returned before the termination inode, the **gpfs\_next\_inode()** subroutine returns a value of 0 and the inode's attribute pointer is set to NULL.

For an overview of using **gpfs\_next\_inode()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_next\_inode()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For an incremental backup, only inodes of files that have changed since the specified previous snapshot will be returned. Any operation that changes the file's **mtime** or **ctime** is considered a change and will cause the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

Incremental backups return deleted files, but full backups do not. A deleted file is indicated by the field **ia\_nlinks** having a value of 0.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan()** with **fssnapHandle** of the snapshot and pass **fssnapid** of the **fssnapHandle** as **prev\_fssnapId**. Repeated invocations of **gpfs\_next\_inode()** then return the inodes copied to the snapshot, skipping holes.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### Parameters

##### iscan

Pointer to the inode scan handle.

**termIno**

The inode scan terminates before this inode number. The caller may specify **maxIno** from **gpfs\_open\_inodescan()** or zero to scan the entire inode file.

**iattr**

Pointer to the returned pointer to the inode's **iattr**.

**Exit status**

If the **gpfs\_next\_inode()** subroutine is successful, it returns a value of 0 and a pointer. The pointer points to NULL if there are no more inodes. Otherwise, the pointer points to the returned inode's attributes.

If the **gpfs\_next\_inode()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

**Exceptions**

None.

**Error status**

Error codes include but are not limited to the following:

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The **gpfs\_next\_inode()** subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ESTALE**

Cached file system information was not valid.

**GPFS\_E\_INVAL\_FSSNAPID**

The file system snapshot ID is not valid.

**GPFS\_E\_INVAL\_ISCAN**

Incorrect parameters.

**Examples**

For an example using **gpfs\_next\_inode()**, see */usr/lpp/mmfs/samples/util/tstimes.c*.

**Location**

*/usr/lpp/mmfs/lib/libgpfs.a* for AIX

*/usr/lpp/mmfs/lib/libgpfs.so* for Linux

## gpfs\_next\_inode64()

---

### gpfs\_next\_inode64() subroutine

Retrieves the next inode from the inode scan.

#### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### Synopsis

```
#include <gpfs.h>
int gpfs_next_inode64(gpfs_iscan_t *iscan,
 gpfs_ino64_t termIno,
 const gpfs_iattr64_t **iattr);
```

#### Description

The **gpfs\_next\_inode64()** subroutine obtains the next inode from the specified inode scan and sets the **iattr** pointer to the inode's attributes. The **termIno** parameter can be used to stop the inode scan before the last inode in the file system or snapshot being scanned. A value of 0 can be provided to indicate the last inode in the file system or snapshot. If there are no more inodes to be returned before the termination inode, the **gpfs\_next\_inode64()** subroutine returns a value of 0 and the inode's attribute pointer is set to NULL.

For an overview of using **gpfs\_next\_inode64()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan64()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode64()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_next\_inode64()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For an incremental backup, only inodes of files that have changed since the specified previous snapshot will be returned. Any operation that changes the file's **mtime** or **ctime** is considered a change and will cause the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

Incremental backups return deleted files, but full backups do not. A deleted file is indicated by the field **ia\_nlinks** having a value of 0.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan64()** with **fssnapHandle** of the snapshot and pass **fssnapid** of the **fssnapHandle** as **prev\_fssnapId**. Repeated invocations of **gpfs\_next\_inode64()** then return the inodes copied to the snapshot, skipping holes.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### Parameters

##### iscan

A pointer to the inode scan handle.

**termIno**

The inode scan terminates before this inode number. The caller may specify **maxIno** from **gpfs\_open\_inodescan64()** or zero to scan the entire inode file.

**iattr**

A pointer to the returned pointer to the inode's **iattr**.

**Exit status**

If the **gpfs\_next\_inode64()** subroutine is successful, it returns a value of 0 and a pointer. The pointer points to NULL if there are no more inodes. Otherwise, the pointer points to the returned inode's attributes.

If the **gpfs\_next\_inode64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

**Exceptions**

None.

**Error status**

Error codes include but are not limited to the following:

**ENOMEM**

Unable to allocate memory for the request.

**ENOSYS**

The **gpfs\_next\_inode64()** subroutine is not available.

**EPERM**

The caller does not have superuser privileges.

**ESTALE**

The cached file system information was not valid.

**GPFS\_E\_INVAL\_FSSNAPID**

The file system snapshot ID is not valid.

**GPFS\_E\_INVAL\_ISCAN**

Incorrect parameters.

**Examples**

See the **gpfs\_next\_inode()** example in /usr/lpp/mmfs/samples/util/tstimes.c.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_next\_inode\_with\_xattrs()**

---

### **gpfs\_next\_inode\_with\_xattrs() subroutine**

Retrieves the next inode and its extended attributes from the inode scan.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_next_inode_with_xattrs(gpfs_iscan_t *iscan,
 gpfs_ino_t termIno,
 const gpfs_iattr_t **iattr,
 const char **xattrBuf,
 unsigned int *xattrBufLen);
```

#### **Description**

The **gpfs\_next\_inode\_with\_xattrs()** subroutine retrieves the next inode and its extended attributes from the inode scan. The set of extended attributes returned are defined when the inode scan was opened. The scan stops before the last inode that was specified or the last inode in the inode file being scanned.

The data returned by **gpfs\_next\_inode()** is overwritten by subsequent calls to **gpfs\_next\_inode()**, **gpfs\_seek\_inode()**, or **gpfs\_stat\_inode()**.

The **termIno** parameter provides a way to partition an inode scan so it can be run on more than one node.

The returned values for **xattrBuf** and **xattrBufLen** must be provided to **gpfs\_next\_xattr()** to obtain the extended attribute names and values. The buffer used for the extended attributes is overwritten by subsequent calls to **gpfs\_next\_inode()**, **gpfs\_seek\_inode()**, or **gpfs\_stat\_inode()**.

The returned pointers to the extended attribute name and value will be aligned to a double-word boundary.

#### **Parameters**

##### **iscan**

A pointer to the inode scan descriptor.

##### **termIno**

The inode scan stops before this inode number. The caller can specify **maxIno** from **gpfs\_open\_inodescan()** or zero to scan the entire inode file.

##### **iattr**

A pointer to the returned pointer to the file's **iattr**.

##### **xattrBuf**

A pointer to the returned pointer to the **xiaattr** buffer.

##### **xattrBufLen**

The returned length of the **xiaattr** buffer.

#### **Exit status**

If the **gpfs\_next\_inode\_with\_xattrs()** subroutine is successful, it returns a value of 0 and **iattr** is set to point to **gpfs\_iattr\_t**. The pointer points to NULL if there are no more inodes, otherwise, the pointer points to **gpfs\_iattr\_t**.

If the **gpfs\_next\_inode\_with\_xattrs()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to NULL to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EFAULT

The buffer data was overwritten.

### ENOMEM

The buffer is too small, unable to allocate memory for the request.

### ENOSYS

The **gpfs\_next\_inode\_with\_xattrs()** subroutine is not available.

### EPERM

The caller does not have superuser privileges.

### ESTALE

The cached file system information was not valid.

### GPFS\_EINVAL\_ISCAN

Incorrect parameters.

### GPFS\_EINVAL\_XATTR

Incorrect parameters.

## Examples

For an example using **gpfs\_next\_inode\_with\_xattrs()**, see */usr/lpp/mmfssamples/util/tsinode.c*.

## Location

*/usr/lpp/mmfsslib/libgpfs.a* for AIX

*/usr/lpp/mmfsslib/libgpfs.so* for Linux

## **gpfs\_next\_inode\_with\_xattrs64()**

---

### **gpfs\_next\_inode\_with\_xattrs64() subroutine**

Retrieves the next inode and its extended attributes from the inode scan.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_next_inode_with_xattrs64(gpfs_iscan_t *iscan,
 gpfs_ino64_t termIno,
 const gpfs_iattr64_t **iattr,
 const char **xattrBuf,
 unsigned int *xattrBufLen);
```

#### **Description**

The **gpfs\_next\_inode\_with\_xattrs64()** subroutine retrieves the next inode and its extended attributes from the inode scan. The set of extended attributes returned are defined when the inode scan was opened. The scan stops before the last inode that was specified or the last inode in the inode file being scanned.

The data returned by **gpfs\_next\_inode64()** is overwritten by subsequent calls to **gpfs\_next\_inode64()**, **gpfs\_seek\_inode64()**, or **gpfs\_stat\_inode64()**.

The **termIno** parameter provides a way to partition an inode scan so it can be run on more than one node.

The returned values for **xattrBuf** and **xattrBufLen** must be provided to **gpfs\_next\_xattr()** to obtain the extended attribute names and values. The buffer used for the extended attributes is overwritten by subsequent calls to **gpfs\_next\_inode64()**, **gpfs\_seek\_inode64()**, or **gpfs\_stat\_inode64()**.

The returned pointers to the extended attribute name and value will be aligned to a double-word boundary.

#### **Parameters**

##### **iscan**

A pointer to the inode scan descriptor.

##### **termIno**

The inode scan stops before this inode number. The caller can specify **maxIno** from **gpfs\_open\_inodescan64()** or zero to scan the entire inode file.

##### **iattr**

A pointer to the returned pointer to the file's **iattr**.

##### **xattrBuf**

A pointer to the returned pointer to the **xiaattr** buffer. Initialize this parameter to a valid value or NULL before calling **gpfs\_next\_inode\_with\_xattrs64()**.

##### **xattrBufLen**

The returned length of the **xiaattr** buffer. Initialize this parameter to a valid value or NULL before calling **gpfs\_next\_inode\_with\_xattrs64()**.

## Exit status

If the **gpfs\_next\_inode\_with\_xattrs64()** subroutine is successful, it returns a value of 0 and **iatr** is set to point to **gpfs\_iattr\_t**. The pointer points to NULL if there are no more inodes, otherwise, the pointer points to **gpfs\_iattr\_t**.

If the **gpfs\_next\_inode\_with\_xattrs64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to NULL to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EFAULT

The buffer data was overwritten.

### ENOMEM

Unable to allocate memory for the request.

### ENOSYS

The **gpfs\_next\_inode\_with\_xattrs64()** subroutine is not available.

### EPERM

The caller does not have superuser privileges.

### ESTALE

The cached file system information was not valid.

### GPFS\_E\_INVAL\_ISCAN

Incorrect parameters.

### GPFS\_E\_INVAL\_XATTR

Incorrect parameters.

## Examples

See the **gpfs\_next\_inode\_with\_xattrs()** example in `/usr/lpp/mmfs/samples/util/tsinode.c`.

## Location

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_next\_xattr()**

---

### **gpfs\_next\_xattr() subroutine**

Returns individual attributes and their values.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_next_xattr(gpfs_iscan_t *iscan,
 const char **xattrBuf,
 unsigned int *xattrBufLen,
 const char **name,
 unsigned int *valueLen,
 const char **value);
```

#### **Description**

The **gpfs\_next\_xattr()** subroutine iterates over the extended attributes buffer returned by the **gpfs\_next\_inode\_with\_xattrs()** or **gpfs\_next\_inode\_with\_xattrs64()** subroutine to return the individual attributes and their values. The attribute names are null-terminated strings, whereas the attribute value contains binary data.

**Note:** The caller is not allowed to modify the returned attribute names or values. The data returned by **gpfs\_next\_xattr()** might be overwritten by subsequent calls to **gpfs\_next\_xattr()** or other GPFS library calls.

#### **Parameters**

##### **iscan**

A pointer to the inode descriptor.

##### **xattrBuf**

A pointer to the pointer to the attribute buffer.

##### **xattrBufLen**

A pointer to the attribute buffer length.

##### **name**

A pointer to the attribute name.

##### **valueLen**

A pointer to the length of the attribute value.

##### **value**

A pointer to the attribute value.

#### **Exit status**

If the **gpfs\_next\_xattr()** subroutine is successful, it returns a value of 0 and a pointer to the attribute name. It also sets:

- The **valueLen** parameter to the length of the attribute value
- The **value** parameter to point to the attribute value
- The **xattrBufLen** parameter to the remaining length of buffer
- The **xattrBuf** parameter to index the next attribute in buffer

If the **gpfs\_next\_xattr()** subroutine is successful, but there are no more attributes in the buffer, it returns a value of 0 and the attribute name is set to NULL. It also sets:

- The **valueLen** parameter to 0
- The **value** parameter to NULL
- The **xattrBufLen** parameter to 0
- The **xattrBuf** parameter to NULL

If the **gpfs\_next\_xattr()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EINVAL

Incorrect parameters.

### ENOSYS

The **gpfs\_next\_xattr()** subroutine is not available.

## Examples

For an example using **gpfs\_next\_xattr()**, see */usr/lpp/mmfs/samples/util/tsinode.c*.

## Location

*/usr/lpp/mmfs/lib/libgpfs.a* for AIX

*/usr/lpp/mmfs/lib/libgpfs.so* for Linux

## **gpfs\_opaque\_acl\_t**

---

### **gpfs\_opaque\_acl\_t structure**

Contains buffer mapping for the **gpfs\_getacl()** and **gpfs\_putacl()** subroutines.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct
{
 int acl_buffer_len;
 unsigned short acl_version;
 unsigned char acl_type;
 char acl_var_data[1];
} gpfs_opaque_acl_t;
```

#### **Description**

The **gpfs\_opaque\_acl\_t** structure contains size, version, and ACL type information for the **gpfs\_getacl()** and **gpfs\_putacl()** subroutines.

#### **Members**

##### **acl\_buffer\_len**

On input, this field must be set to the total length, in bytes, of the data structure being passed to GPFS. On output, this field contains the actual size of the requested information. If the initial size of the buffer is not large enough to contain all of the information, the **gpfs\_getacl()** invocation must be repeated with a larger buffer.

##### **acl\_version**

This field contains the current version of the GPFS internal representation of the ACL. On input to the **gpfs\_getacl()** subroutine, set this field to zero.

##### **acl\_type**

On input to the **gpfs\_getacl()** subroutine, set this field to either **GPFS\_ACL\_TYPE\_ACCESS** or **GPFS\_ACL\_TYPE\_DEFAULT**, depending on which ACL is requested. These constants are defined in the **gpfs.h** header file.

##### **acl\_var\_data**

This field signifies the beginning of the remainder of the ACL information.

#### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_open\_inodescan() subroutine

Opens an inode scan of a file system or snapshot.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
gpfs_iscan_t *gpfs_open_inodescan(gpfs_fssnap_handle_t *fssnapHandle,
 const gpfs_fssnap_id_t *prev_fssnapId,
 gpfs_ino_t *maxIno);
```

### Description

The **gpfs\_open\_inodescan()** subroutine opens a scan of the inodes in the file system or snapshot identified by the **fssnapHandle** parameter. The scan traverses all user files, directories and links in the file system or snapshot. The scan begins with the user file with the lowest inode number and returns the files in increasing order. The **gpfs\_seek\_inode()** subroutine may be used to set the scan position to an arbitrary inode. System files, such as the block allocation maps, are omitted from the scan. The file system must be mounted to open an inode scan.

For an overview of using **gpfs\_open\_inodescan()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_open\_inodescan()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For the incremental backup, any operation that changes the file's **mtime** or **ctime** causes the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

A full inode scan (**prev\_fssnapId** set to NULL) does not return any inodes of nonexistent or deleted files, but an incremental inode scan (**prev\_fssnapId** not NULL) does return inodes for files that have been deleted since the previous snapshot. The inodes of deleted files have a link count of zero.

If the snapshot indicated by **prev\_fssnapId** is available, the caller may benefit from the extended read subroutine, **gpfs\_ireadx()**, which returns only the changed blocks within the files. Without the previous snapshot, all blocks within the changed files are returned.

Once a full or incremental backup completes, the **new\_fssnapId** must be saved in order to reuse it on a subsequent incremental backup. This **fssnapId** must be provided to the **gpfs\_open\_inodescan()** subroutine, as the **prev\_fssnapId** input parameter.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan()** with **fssnapHandle** of the snapshot and pass **fssnapId** of the **fssnapHandle** as the **prev\_fssnapId**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

## **gpfs\_open\_inodescan()**

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

File system snapshot handle.

#### **prev\_fssnapId**

Pointer to file system snapshot ID or NULL. If **prev\_fssnapId** is provided, the inode scan returns only the files that have changed since the previous backup. If the pointer is NULL, the inode scan returns all user files.

If it is the same as the fssnapid of the fssnapHandle parameter, the scan only returns the inodes copied into the corresponding snapshot.

#### **maxIno**

Pointer to inode number or NULL. If provided, **gpfs\_open\_inodescan()** returns the maximum inode number in the file system or snapshot being scanned.

### **Exit status**

If the **gpfs\_open\_inodescan()** subroutine is successful, it returns a pointer to an inode scan handle.

If the **gpfs\_open\_inodescan()** subroutine is unsuccessful, it returns a NULL pointer and the global error variable **errno** is set to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EDOM**

The file system snapshot ID passed for **prev\_fssnapId** is from a different file system.

#### **EINVAL**

Incorrect parameters.

#### **ENOMEM**

Unable to allocate memory for the request.

#### **ENOSYS**

The **gpfs\_open\_inodescan()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

#### **ERANGE**

The **prev\_fssnapId** parameter is the same as or more recent than **snapId** being scanned.

#### **ESTALE**

Cached file system information was not valid.

#### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

#### **GPFS\_E\_INVAL\_FSSNAPID**

The file system snapshot ID passed for **prev\_fssnapId** is not valid.

## **Examples**

For an example using `gpfs_open_inodescan()`, see `/usr/lpp/mmfs/samples/util/tstimes.c`.

## **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_open\_inodescan64()**

---

### **gpfs\_open\_inodescan64() subroutine**

Opens an inode scan of a file system or snapshot.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
gpfs_iscan_t *gpfs_open_inodescan64(gpfs_fssnap_handle_t *fssnapHandle,
 const gpfs_fssnap_id_t *prev_fssnapId,
 gpfs_ino64_t *maxIno);
```

#### **Description**

The **gpfs\_open\_inodescan64()** subroutine opens a scan of the inodes in the file system or snapshot identified by the **fssnapHandle** parameter. The scan traverses all user files, directories and links in the file system or snapshot. The scan begins with the user file with the lowest inode number and returns the files in increasing order. The **gpfs\_seek\_inode64()** subroutine may be used to set the scan position to an arbitrary inode. System files, such as the block allocation maps, are omitted from the scan. The file system must be mounted to open an inode scan.

For an overview of using **gpfs\_open\_inodescan64()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan64()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode64()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_open\_inodescan64()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For the incremental backup, any operation that changes the file's **mtime** or **ctime** causes the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

A full inode scan (**prev\_fssnapId** set to NULL) does not return any inodes of nonexistent or deleted files, but an incremental inode scan (**prev\_fssnapId** not NULL) does return inodes for files that have been deleted since the previous snapshot. The inodes of deleted files have a link count of zero.

If the snapshot indicated by **prev\_fssnapId** is available, the caller may benefit from the extended read subroutine, **gpfs\_ireadx()**, which returns only the changed blocks within the files. Without the previous snapshot, all blocks within the changed files are returned.

Once a full or incremental backup completes, the **new\_fssnapId** must be saved in order to reuse it on a subsequent incremental backup. This **fssnapId** must be provided to the **gpfs\_open\_inodescan64()** subroutine, as the **prev\_fssnapId** input parameter.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan()** with **fssnapHandle** of the snapshot and pass **fssnapid** of the **fssnapHandle** as the **prev\_fssnapId**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

## Parameters

### **fssnapHandle**

The file system snapshot handle.

### **prev\_fssnapId**

A pointer to file system snapshot ID or NULL. If **prev\_fssnapId** is provided, the inode scan returns only the files that have changed since the previous backup. If the pointer is NULL, the inode scan returns all user files.

If it is same as the fssnapid of the fssnapHandle parameter, the scan only returns the inodes copied into the corresponding snapshot.

### **maxIno**

A pointer to inode number or NULL. If provided, **gpfs\_open\_inodescan64()** returns the maximum inode number in the file system or snapshot being scanned.

## Exit status

If the **gpfs\_open\_inodescan64()** subroutine is successful, it returns a pointer to an inode scan handle.

If the **gpfs\_open\_inodescan64()** subroutine is unsuccessful, it returns a NULL pointer and the global error variable **errno** is set to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### **EDOM**

The file system snapshot ID passed for **prev\_fssnapId** is from a different file system.

### **EINVAL**

Incorrect parameters.

### **ENOMEM**

Unable to allocate memory for the request.

### **ENOSYS**

The **gpfs\_open\_inodescan64()** subroutine is not available.

### **EPERM**

The caller does not have superuser privileges.

### **ERANGE**

The **prev\_fssnapId** parameter is the same as or more recent than **snapId** being scanned.

### **ESTALE**

The cached file system information was not valid.

### **GPFS\_E\_INVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **GPFS\_E\_INVAL\_FSSNAPID**

The file system snapshot ID passed for **prev\_fssnapId** is not valid.

## **gpfs\_open\_inodescan64()**

**Note:** `gpfs_open_inodescan64()` calls the standard library subroutines `dup()` and `malloc()`; if one of these called subroutines returns an error, `gpfs_open_inodescan64()` also returns that error.

### **Examples**

See the `gpfs_open_inodescan()` example in `/usr/lpp/mmfs/samples/util/tstimes.c`.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_open\_inodescan\_with\_xattrs() subroutine

Opens an inode file and extended attributes for an inode scan.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
gpfs_iscan_t *gpfs_open_inodescan_with_xattrs(gpfs_fssnap_handle_t *fssnapHandle,
 const gpfs_fssnap_id_t *prev_fssnapId,
 int nxAttrs,
 const char *xattrsList[],
 gpfs_ino_t *maxIno);
```

### Description

The **gpfs\_open\_inodescan\_with\_xattrs()** subroutine opens an inode file and extended attributes for an inode scan identified by the **fssnapHandle** parameter. The scan traverses all user files, directories and links in the file system or snapshot. The scan begins with the user file with the lowest inode number and returns the files in increasing order. The **gpfs\_seek\_inode()** subroutine can be used to set the scan position to an arbitrary inode. System files, such as the block allocation maps, are omitted from the scan. The file system must be mounted to open an inode scan.

For an overview of using **gpfs\_open\_inodescan\_with\_xattrs()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan\_with\_xattrs()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_open\_inodescan\_with\_xattrs()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For the incremental backup, any operation that changes the file's **mtime** or **ctime** causes the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

A full inode scan (**prev\_fssnapId** set to NULL) returns all inodes of existing files. An incremental inode scan (**prev\_fssnapId** not NULL) returns inodes for files that have changed since the previous snapshot. The inodes of deleted files have a link count of zero.

If the snapshot indicated by **prev\_fssnapId** is available, the caller may benefit from the extended read subroutine, **gpfs\_ireadx()**, which returns only the changed blocks within the files. Without the previous snapshot, all blocks within the changed files are returned.

Once a full or incremental backup completes, the **new\_fssnapId** must be saved in order to reuse it on a subsequent incremental backup. This **fssnapId** must be provided to the **gpfs\_open\_inodescan\_with\_xattrs()** subroutine, as the **prev\_fssnapId** input parameter.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan()** with **fssnapHandle** of the snapshot and pass **fssnapid** of the **fssnapHandle** as the **prev\_fssnapId**.

## **gpfs\_open\_inodescan\_with\_xattrs()**

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **fssnapHandle**

The file system snapshot handle.

#### **prev\_fssnapId**

A pointer to file system snapshot ID or NULL. If **prev\_fssnapId** is provided, the inode scan returns only the files that have changed since the previous backup. If the pointer is NULL, the inode scan returns all user files.

If it is the same as the fssnapid of the fssnapHandle parameter, the scan only returns the inodes copied into the corresponding snapshot.

#### **nxAttrs**

The count of extended attributes to be returned. If **nxAttrs** is set to 0, call returns no extended attributes, like **gpfs\_open\_inodescan()**. If **nxAttrs** is set to -1, call returns all extended attributes.

#### **xattrsList**

A pointer to an array of pointers to names of extended attributes to be returned. **nxAttrsList** may be null if **nxAttrs** is set to 0 or -1.

#### **maxIno**

A pointer to inode number or NULL. If provided, **gpfs\_open\_inodescan\_with\_xattrs()** returns the maximum inode number in the file system or snapshot being scanned.

### **Exit status**

If the **gpfs\_open\_inodescan\_with\_xattrs()** subroutine is successful, it returns a pointer to **gpfs\_iscan\_t**.

If the **gpfs\_open\_inodescan\_with\_xattrs()** subroutine is unsuccessful, it returns a NULL pointer and the global error variable **errno** is set to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EDOM**

The file system snapshot ID passed for **prev\_fssnapId** is from a different file system.

#### **EINVAL**

Incorrect parameters.

#### **ENOMEM**

Unable to allocate memory for the request.

#### **ENOSYS**

The **gpfs\_open\_inodescan\_with\_xattrs()** subroutine is not available.

#### **EPERM**

The caller does not have superuser privileges.

#### **ERANGE**

The **prev\_fssnapId** parameter is the same as or more recent than **snapId** being scanned.

**ESTALE**

The cached file system information was not valid.

**GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

**GPFS\_EINVAL\_FSSNAPID**

The file system snapshot ID passed for **prev\_fssnapId** is not valid.

**Note:** `gpfs_open_inodescan_with_xattrs()` calls the standard library subroutines `dup()` and `malloc()`; if one of these called subroutines returns an error, `gpfs_open_inodescan_with_xattrs()` also returns that error.

## Examples

For an example using `gpfs_open_inodescan_with_xattrs()`, see `/usr/lpp/mmfs/samples/util/tsinode.c`.

## Location

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfs\_open\_inodescan\_with\_xattrs64()**

---

### **gpfs\_open\_inodescan\_with\_xattrs64() subroutine**

Opens an inode file and extended attributes for an inode scan.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
gpfs_iscan_t *gpfs_open_inodescan_with_xattrs64(gpfs_fssnap_handle_t *fssnapHandle,
 const gpfs_fssnap_id_t *prev_fssnapId,
 int nxAttrs,
 const char *xattrList[],
 gpfs_ino64_t *maxIno);
```

#### **Description**

The **gpfs\_open\_inodescan\_with\_xattrs64()** subroutine opens an inode file and extended attributes for an inode scan identified by the **fssnapHandle** parameter. The scan traverses all user files, directories and links in the file system or snapshot. The scan begins with the user file with the lowest inode number and returns the files in increasing order. The **gpfs\_seek\_inode64()** subroutine may be used to set the scan position to an arbitrary inode. System files, such as the block allocation maps, are omitted from the scan. The file system must be mounted to open an inode scan.

For an overview of using **gpfs\_open\_inodescan\_with\_xattrs64()** in a backup application, see *Using APIs to develop backup applications* in *IBM Spectrum Scale: Administration Guide*.

To generate a full backup, invoke **gpfs\_open\_inodescan\_with\_xattrs64()** with NULL for the **prev\_fssnapId** parameter. Repeated invocations of **gpfs\_next\_inode64()** then return inode information about all existing user files, directories, and links in inode number order.

To generate an incremental backup, invoke **gpfs\_open\_inodescan\_with\_xattrs64()** with the **fssnapId** that was obtained from a **fssnapHandle** at the time the previous backup was created. The snapshot that was used for the previous backup does not need to exist at the time the incremental backup is generated. That is, the backup application needs to remember only the **fssnapId** of the previous backup; the snapshot itself can be deleted as soon as the backup is completed.

For the incremental backup, any operation that changes the file's **mtime** or **ctime** causes the file to be included. Files with no changes to the file's data or file attributes, other than a change to **atime**, are omitted from the scan.

A full inode scan (**prev\_fssnapId** set to NULL) returns all inodes of existing files. An incremental inode scan (**prev\_fssnapId** not NULL) returns inodes for files that have changed since the previous snapshot. The inodes of deleted files have a link count of zero.

If the snapshot indicated by **prev\_fssnapId** is available, the caller may benefit from the extended read subroutine, **gpfs\_ireadx()**, which returns only the changed blocks within the files. Without the previous snapshot, all blocks within the changed files are returned.

Once a full or incremental backup completes, the **new\_fssnapId** must be saved in order to reuse it on a subsequent incremental backup. This **fssnapId** must be provided to the **gpfs\_open\_inodescan\_with\_xattrs64()** subroutine, as the **prev\_fssnapId** input parameter.

To read only the inodes that have been copied to a snapshot, use **gpfs\_open\_inodescan()** with **fssnapHandle** of the snapshot and pass **fssnapid** of the **fssnapHandle** as the **prev\_fssnapId**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

## Parameters

### **fssnapHandle**

The file system snapshot handle.

### **prev\_fssnapId**

A pointer to file system snapshot ID or NULL. If **prev\_fssnapId** is provided, the inode scan returns only the files that have changed since the previous backup. If the pointer is NULL, the inode scan returns all user files.

If it is same as the fssnapid of the fssnapHandle parameter, the scan only returns the inodes copied into the corresponding snapshot.

### **nxAttrs**

The count of extended attributes to be returned. If **nxAttrs** is set to 0, call returns no extended attributes, like **gpfs\_open\_inodescan64()**. If **nxAttrs** is set to -1, call returns all extended attributes

### **xattrsList**

A pointer to an array of pointers to names of extended attributes to be returned. **nxAttrsList** may be null if **nxAttrs** is set to 0 or -1.

### **maxIno**

A pointer to inode number or NULL. If provided, **gpfs\_open\_inodescan\_with\_xattrs64()** returns the maximum inode number in the file system or snapshot being scanned.

## Exit status

If the **gpfs\_open\_inodescan\_with\_xattrs64()** subroutine is successful, it returns a pointer to **gpfs\_iscan\_t**.

If the **gpfs\_open\_inodescan\_with\_xattrs64()** subroutine is unsuccessful, it returns a NULL pointer and the global error variable **errno** is set to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### **EDOM**

The file system snapshot ID passed for **prev\_fssnapId** is from a different file system.

### **EINVAL**

Incorrect parameters.

### **ENOMEM**

Unable to allocate memory for the request.

### **ENOSYS**

The **gpfs\_open\_inodescan\_with\_xattrs64()** subroutine is not available.

### **EPERM**

The caller does not have superuser privileges.

### **ERANGE**

The **prev\_fssnapId** parameter is the same as or more recent than **snapId** being scanned.

## **gpfs\_open\_inodescan\_with\_xattrs64()**

### **ESTALE**

The cached file system information was not valid.

### **GPFS\_EINVAL\_FSSNAPHANDLE**

The file system snapshot handle is not valid.

### **GPFS\_EINVAL\_FSSNAPID**

The file system snapshot ID passed for **prev\_fssnapId** is not valid.

**Note:** `gpfs_open_inodescan_with_xattrs64()` calls the standard library subroutines `dup()` and `malloc()`; if one of these called subroutines returns an error, `gpfs_open_inodescan_with_xattrs64()` also returns that error.

## **Examples**

See the `gpfs_open_inodescan_with_xattrs()` example in `/usr/lpp/mmfs/samples/util/tsinode.c`.

## **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_malloc() subroutine

Preallocates storage for a regular file or preallocates directory entries for a directory.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_malloc(gpfs_file_t fileDesc,
 gpfs_off64_t startOffset,
 gpfs_off64_t bytesToPrealloc);
```

### Description

The **gpfs\_malloc()** subroutine preallocates disk storage for a file or a directory.

In the case of a regular file, preallocation can improve I/O performance by creating a block of available disk storage in the file immediately instead of increasing the file size incrementally as data is written. The preallocated disk storage begins at the specified offset in the file and extends for the specified number of bytes, rounded up to a GPFS block boundary. Existing data is not modified. Reading any of the preallocated blocks returns zeroes. To determine how much space was actually preallocated, call the **stat()** subroutine and compare the reported file size and number of blocks used with their values before the preallocation.

In the case of a directory, preallocation can improve metadata performance by setting the minimum compaction size of the directory. Preallocation is most effective in systems in which many files are added to and removed from a directory in a short time. The minimum compaction size is the number of directory slots, including both full and empty slots, that a directory is allowed to retain when it is compacted. In IBM Spectrum Scale v4.1 or later, by default, a directory is automatically compacted as far as possible.

The **gpfs\_malloc()** subroutine sets the minimum compaction size of a directory to the specified number of slots and adds directory slots if needed to reach the minimum size. For example, if a directory contains 5,000 files and you set the minimum compaction size to 50,000, then the file system adds 45,000 directory slots. The directory can grow beyond 50,000 entries, but the file system does not allow the directory to be compacted below 50,000 slots.

You must specify the minimum compaction size as a number of bytes. Determine the number of bytes to allocate with the following formula:

`bytesToPrealloc = n * ceiling((namelen + 13) / 32) * 32`

where:

**n**      Specifies the number of directory entries that you want.

#### **ceiling()**

Is a function that rounds a fractional number up to the next highest integer. For example, `ceiling(1.03125)` returns 2.

#### **namelen**

Specifies the expected average length of file names.

For example, if you want 20,000 entries with an average file name length of 48, then `bytesToPrealloc = 20000 * 2 * 32 = 1,280,000`.

## **gpfs\_prealloc()**

To restore the default behavior of the file system for a directory, in which a directory is compacted as far as possible, call the subroutine with bytesToPrealloc set to 0.

The number of bytes that are allocated for a directory is stored in the ia\_dirminsize field in the `gpfs_iattr64_t` structure that is returned by the `gpfs_fstat_x()` subroutine. To convert the number of bytes to the number of directory slots, apply the following rule:

```
numSlots = numBytesReturned / 32
```

To convert the number of bytes to the number of files in the directory, apply a version of the formula that is described above:

```
numFiles = numBytesReturned / (ceiling(namelen + 13) * 32)
```

If the file is not a directory or is a directory with no preallocation, the `ia_dirminsize` field is set to 0. This attribute is reported for information only and is meaningful only in the active file system. It is not backed up and restored and it is ignored in snapshots. It usually differs from the file size that is recorded in `ia_size` or returned by `stat()`.

To set the minimum compaction size of a directory from the command line, set the `compact` parameter in the `mmchattr` command.

**Note:** Compile any program that uses this subroutine with the `-lgpfs` flag from the following library:

- `libgpfs.a` for AIX
- `libgpfs.so` for Linux

## **Parameters**

### **fileDesc**

The file descriptor returned by `open()`. Note the following points:

- A file must be opened for writing.
- A directory can be opened for reading, but the caller must have write permission to the directory.

### **startOffset**

For a file, the byte offset into the file at which to begin preallocation. For a directory, set this parameter to 0.

### **bytesToPrealloc**

The number of bytes to be preallocated. For a file, this value is rounded up to a GPFS block boundary. For a directory, calculate the number of bytes with the formula that is described above.

## **Exit status**

If the `gpfs_prealloc()` subroutine is successful, it returns a value of 0.

If the `gpfs_prealloc()` subroutine is unsuccessful, it returns a value of -1 and sets the global error variable `errno` to indicate the nature of the error. If `errno` is set to one of the following, some storage may have been preallocated:

- `EDQUOT`
- `ENOSPC`

The pre-allocation or minimum compaction setting of a directory can be obtained using the “`gpfs_fstat_x()` subroutine” on page 698. The `gpfs_iattr64` structure it returns, defined in `gpfs.h`, is extended to include `ia_dirminsize` giving the pre-allocation size of a directory. For non-directories or directories without a pre-allocation set the value is zero. The size is expressed in bytes, which is interpreted in the same way as for the `gpfs_prealloc()` function. This will be 32 times the value reported by `mmlsattr` and will generally differ from the file size reported in `ia_size` and by `stat()`.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EACCES

The file or directory is not opened for writing.

### EBADF

The file descriptor is not valid.

### EDQUOT

A disk quota has been exceeded

### EINVAL

The file descriptor does not refer to a file or directory; a negative value was specified for **startOffset** or **bytesToPrealloc**.

### ENOSPC

The file system has run out of disk space.

### ENOSYS

The **gpfs\_prealloc()** subroutine is not supported under the current file system format.

## Examples

```
#include <stdio.h>
#include <fcntl.h>
#include <errno.h>
#include <gpfs.h>

int rc;
int fileHandle = -1;
char* fileNameP = "datafile";
offset_t startOffset = 0;
offset_t bytesToAllocate = 20*1024*1024; /* 20 MB */

fileHandle = open(fileNameP, O_RDWR|O_CREAT, 0644);
if (fileHandle < 0)
{
 perror(fileNameP);
 exit(1);
}

rc = gpfs_prealloc(fileHandle, startOffset, bytesToAllocate);
if (rc < 0)
{
 fprintf(stderr, "Error %d preallocation at %lld for %lld in %s\n",
 errno, startOffset, bytesToAllocate, fileNameP);
 exit(1);
}
```

## Location

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_putacl()**

---

### **gpfs\_putacl() subroutine**

Restores the access control information for a GPFS file.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_putacl(const char *pathname,
 int flags,
 void *acl);
```

#### **Description**

The **gpfs\_putacl()** subroutine together with the **gpfs\_getacl()** subroutine is intended for use by a backup program to save (**gpfs\_getacl()**) and restore (**gpfs\_putacl()**) the ACL information for the file.

#### **Notes:**

1. The use of **gpfs\_fgetattrs()** and **gpfs\_fputattrs()** is preferred.
2. You must have **write** access to the file.
3. Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:
  - libgpfs.a for AIX
  - libgpfs.so for Linux

#### **Parameters**

##### **pathname**

Path name of the file for which the ACLs is to be set.

##### **flags**

Consists of one of these values:

- 0 Indicates that the **acl** parameter is to be mapped with the **gpfs\_opaque\_acl\_t** structure.  
The **gpfs\_opaque\_acl\_t** structure should be used by backup and restore programs.

##### **GPFS\_PUTACL\_STRUCT**

Indicates that the **acl** parameter is to be mapped with the **gpfs\_acl\_t** structure.

The **gpfs\_acl\_t** structure is provided for applications that need to change the ACL.

##### **acl**

Pointer to a buffer mapped by the structure **gpfs\_opaque\_acl\_t** or **gpfs\_acl\_t**, depending on the value of **flags**.

This is where the ACL data is stored, and should be the result of a previous invocation of **gpfs\_getacl()**.

#### **Exit status**

If the **gpfs\_putacl()** subroutine is successful, it returns a value of 0.

If the **gpfs\_putacl()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## Exceptions

None.

## Error status

Error codes include but are not limited to the following:

### EINVAL

The path name does not refer to a GPFS file or a regular file.

### ENOMEM

Unable to allocate memory for the request.

### ENOSYS

The `gpfs_putacl()` subroutine is not supported under the current file system format.

### ENOTDIR

File is not a directory.

## Location

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfs\_quotactl()**

---

### **gpfs\_quotactl() subroutine**

Manipulates disk quotas on file systems.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Synopsis**

```
#include <gpfs.h>
int gpfs_quotactl(const char *pathname,
 int cmd,
 int id,
 void *bufferP);
```

#### **Description**

The **gpfs\_quotactl()** subroutine manipulates disk quotas. It enables, disables, and manipulates disk quotas for file systems on which quotas have been enabled.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

#### **Parameters**

##### **pathname**

Specifies the path name of any file within the mounted file system to which the quota control command is to applied.

##### **cmd**

Specifies the quota control command to be applied and whether it is applied to a user, group, or fileset quota.

The **cmd** parameter can be constructed using **GPFS\_QCMD(qcmd, Type)** contained in **gpfs.h**. The **qcmd** parameter specifies the quota control command. The **Type** parameter specifies one of the following quota types:

- user (**GPFS\_USRQUOTA**)
- group (**GPFS\_GRPQUOTA**)
- fileset (**GPFS\_FILESETQUOTA**)

The valid values for the **qcmd** parameter specified in **gpfs.h** are:

##### **Q\_QUOTAON**

Enables quotas.

Enables disk quotas for the file system specified by the **pathname** parameter and type specified in **Type**. The **id** and **bufferP** parameters are unused. Root user authority is required to enable quotas.

##### **Q\_QUOTAOFF**

Disables quotas.

Disables disk quotas for the file system specified by the **pathname** parameter and type specified in **Type**. The **id** and **bufferP** parameters are unused. Root user authority is required to disable quotas.

##### **Q\_GETQUOTA**

Gets quota limits and usage information.

Retrieves quota limits and current usage for a user, group, or fileset specified by the **id** parameter. The **bufferP** parameter points to a **gpfs\_quotaInfo\_t** structure to hold the returned information. The **gpfs\_quotaInfo\_t** structure is defined in `gpfs.h`.

Root authority is required if the **id** value is not the current id (user id for **GPFS\_USRQUOTA**, group id for **GPFS\_GRPQUOTA**) of the caller.

#### **Q\_SETQUOTA**

Sets quota limits

Sets disk quota limits for a user, group, or fileset specified by the **id** parameter. The **bufferP** parameter points to a **gpfs\_quotaInfo\_t** structure containing the new quota limits. The **gpfs\_quotaInfo\_t** structure is defined in `gpfs.h`. Root user authority is required to set quota limits.

#### **Q\_SETUSE**

Sets quota usage

Sets disk quota usage for a user, group, or fileset specified by the **id** parameter. The **bufferP** parameter points to a **gpfs\_quotaInfo\_t** structure containing the new quota usage. The **gpfs\_quotaInfo\_t** structure is defined in `gpfs.h`. Root user authority is required to set quota usage.

#### **Q\_SYNC**

Synchronizes the disk copy of a file system quota

Updates the on disk copy of quota usage information for a file system. The **id** and **bufferP** parameters are unused. Root user authority is required to synchronize a file system quota.

**id** Specifies the user, group, or fileset ID to which the quota control command applies. The **id** parameters is interpreted by the specified quota type.

#### **bufferP**

Points to the address of an optional, command-specific data structure that is copied in or out of the system.

### **Exit status**

If the **gpfs\_quotactl()** subroutine is successful, it returns a value of 0.

If the **gpfs\_quotactl()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EACCES**

Search permission is denied for a component of a path prefix.

#### **EFAULT**

An invalid **bufferP** parameter is supplied. The associated structure could not be copied in or out of the kernel.

#### **EINVAL**

One of the following errors:

- The file system is not mounted.

## **gpfs\_quotactl()**

- Invalid command or quota type.
- Invalid input limits: negative limits or soft limits are greater than hard limits.
- UID is not defined.

### **ENOENT**

No such file or directory.

### **EPERM**

The quota control command is privileged and the caller did not have root user authority.

### **GPFS\_E\_NO\_QUOTA\_INST**

The file system does not support quotas. This is the actual **errno** generated by GPFS.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_quotaInfo\_t structure

Contains buffer mapping for the **gpfs\_quotactl()** subroutine.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Structure

```
typedef struct gpfs_quotaInfo
{
 gpfs_off64_t blockUsage; /* current block count */
 gpfs_off64_t blockHardLimit; /* absolute limit on disk blks alloc */
 gpfs_off64_t blockSoftLimit; /* preferred limit on disk blks */
 gpfs_off64_t blockInDoubt; /* distributed shares + "lost" usage for blks */
 int inodeUsage; /* current # allocated inodes */
 int inodeHardLimit; /* absolute limit on allocated inodes */
 int inodeSoftLimit; /* preferred inode limit */
 int inodeInDoubt; /* distributed shares + "lost" usage for inodes */
 gpfs_uid_t quoId; /* uid, gid or fileset id */
 int entryType; /* entry type, not used */
 unsigned int blockGraceTime; /* time limit for excessive disk use */
 unsigned int inodeGraceTime; /* time limit for excessive inode use */
} gpfs_quotaInfo_t;
```

### Description

The **gpfs\_quotaInfo\_t** structure contains detailed information for the **gpfs\_quotactl()** subroutine.

### Members

#### **blockUsage**

The current block count in 1 KB units.

#### **blockHardLimit**

The absolute limit on disk block allocation.

#### **blockSoftLimit**

The preferred limit on disk block allocation.

#### **blockInDoubt**

The distributed shares and block usage that have not been accounted for.

#### **inodeUsage**

The current number of allocated inodes.

#### **inodeHardLimit**

The absolute limit on allocated inodes.

#### **inodeSoftLimit**

The preferred inode limit.

#### **inodeInDoubt**

The distributed inode share and inode usage that have not been accounted for.

#### **quoId**

The user ID, group ID, or fileset ID.

#### **entryType**

Not used

#### **blockGraceTime**

The time limit (in seconds since the Epoch) for excessive disk use.

## **gpfs\_quotaInfo\_t**

### **inodeGraceTime**

The time limit (in seconds since the Epoch) for excessive inode use.

Epoch is midnight on January 1, 1970 UTC (Coordinated Universal Time).

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_seek\_inode() subroutine**

Advances an inode scan to the specified inode number.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_seek_inode(gpfs_iscan_t *iscan,
 gpfs_ino_t ino);
```

### **Description**

The **gpfs\_seek\_inode()** subroutine advances an inode scan to the specified inode number.

The **gpfs\_seek\_inode()** subroutine is used to start an inode scan at some place other than the beginning of the inode file. This is useful to restart a partially completed backup or an interrupted dump transfer to a mirror. It could also be used to do an inode scan in parallel from multiple nodes, by partitioning the inode number space into separate ranges for each participating node. The maximum inode number is returned when the scan was opened and each invocation to obtain the next inode specifies a termination inode number to avoid returning the same inode more than once.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **iscan**

Pointer to the inode scan handle.

#### **ino**

The next inode number to be scanned.

### **Exit status**

If the **gpfs\_seek\_inode()** subroutine is successful, it returns a value of 0.

If the **gpfs\_seek\_inode()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **ENOSYS**

The **gpfs\_seek\_inode()** subroutine is not available.

#### **GPFS\_EINVAL\_ISCAN**

Incorrect parameters.

## **gpfs\_seek\_inode()**

### **Examples**

For an example using `gpfs_seek_inode()`, see `/usr/lpp/mmfs/samples/util/tsinode.c`.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_seek\_inode64() subroutine

Advances an inode scan to the specified inode number.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_seek_inode64(gpfs_iscan_t *iscan,
 gpfs_ino64_t ino);
```

### Description

The **gpfs\_seek\_inode64()** subroutine advances an inode scan to the specified inode number.

The **gpfs\_seek\_inode64()** subroutine is used to start an inode scan at some place other than the beginning of the inode file. This is useful to restart a partially completed backup or an interrupted dump transfer to a mirror. It could also be used to do an inode scan in parallel from multiple nodes, by partitioning the inode number space into separate ranges for each participating node. The maximum inode number is returned when the scan was opened and each invocation to obtain the next inode specifies a termination inode number to avoid returning the same inode more than once.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### iscan

A pointer to the inode scan handle.

#### ino

The next inode number to be scanned.

### Exit status

If the **gpfs\_seek\_inode64()** subroutine is successful, it returns a value of 0.

If the **gpfs\_seek\_inode64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### Exceptions

None.

### Error status

Error codes include but are not limited to the following:

#### ENOSYS

The **gpfs\_seek\_inode64()** subroutine is not available.

#### GPFS\_EINVAL\_ISCAN

Incorrect parameters.

## **gpfs\_seek\_inode64()**

### **Examples**

See the `gpfs_seek_inode()` example in `/usr/lpp/mmfs/samples/util/tsinode.c`.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

---

## gpfs\_stat() subroutine

Returns exact file status for a GPFS file.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_stat(const char *pathname,
 gpfs_stat64_t *buffer);
```

### Description

The **gpfs\_stat()** subroutine is used to obtain exact information about the file named by the **pathname** parameter. This subroutine is provided as an alternative to the **stat()** subroutine, which may not provide exact **mtime** and **atime** values. For more information, see *Exceptions to Open Group technical standards* in *IBM Spectrum Scale: Administration Guide*.

**read**, **write**, or **execute** permission for the named file is not required, but all directories listed in the path leading to the file must be searchable. The file information is written to the area specified by the **buffer** parameter.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### pathname

The path identifying the file for which exact status information is requested.

#### buffer

A pointer to the **gpfs\_stat64\_t** structure in which the information is returned. The **gpfs\_stat64\_t** structure is described in the **sys/stat.h** file.

### Exit status

If the **gpfs\_stat()** subroutine is successful, it returns a value of 0.

If the **gpfs\_stat()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### Exceptions

None.

### Error status

Error codes include but are not limited to the following:

#### EINVAL

The path name does not refer to a GPFS file or a regular file.

#### ENOENT

The file does not exist.

## **gpfs\_stat()**

### **ENOSYS**

The **gpfs\_stat()** subroutine is not supported under the current file system format.

### **ESTALE**

The cached file system information was not valid.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_stat\_inode() subroutine

Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_stat_inode(gpfs_iscan_t *iscan,
 gpfs_ino_t ino,
 gpfs_ino_t termIno,
 const gpfs_iattr_t **iattr);
```

### Description

The **gpfs\_stat\_inode()** subroutine is used to seek the specified inode and to retrieve that inode and its extended attributes from the inode scan. This subroutine combines **gpfs\_seek\_inode()** and **get\_next\_inode()**, but will only return the specified inode.

The data returned by **gpfs\_next\_inode()** is overwritten by subsequent calls to **gpfs\_next\_inode()**, **gpfs\_seek\_inode()**, or **gpfs\_stat\_inode()**.

The **termIno** parameter provides a way to partition an inode scan so it can be run on more than one node. It is only used by this call to control prefetching.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### iscan

A pointer to an inode scan descriptor.

#### ino

The inode number to be returned.

#### termIno

Prefetches inodes up to this inode. The caller might specify **maxIno** from **gpfs\_open\_inodescan()** or 0 to allow prefetching over the entire inode file.

#### iattr

A pointer to the returned pointer to the file's **iattr**.

### Exit status

If the **gpfs\_stat\_inode()** subroutine is successful, it returns a value of 0 and the **iattr** parameter is set to point to **gpfs\_iattr\_t**. If the **gpfs\_stat\_inode()** subroutine is successful, but there are no more inodes before the **termIno** parameter, or if the requested inode does not exist, it returns a value of 0 and the **iattr** parameter is set to NULL.

If the **gpfs\_stat\_inode()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **gpfs\_stat\_inode()**

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EPERM**

The caller must have superuser privilege.

#### **ENOSYS**

The **gpfs\_stat\_inode()** subroutine is not supported under the current file system format.

#### **ESTALE**

The cached file system information was not valid.

#### **ENOMEM**

The buffer is too small.

#### **GPFS\_E\_INVAL\_ISCAN**

Incorrect parameters.

#### **GPFS\_E\_HOLE\_IN\_IFILE**

The inode scan is reading only the inodes that have been copied to a snapshot and this inode has not yet been copied.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_stat\_inode64() subroutine

Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_stat_inode64(gpfs_iscan_t *iscan,
 gpfs_ino64_t ino,
 gpfs_ino64_t termIno,
 const gpfs_iattr64_t **iattr);
```

### Description

The **gpfs\_stat\_inode64()** subroutine is used to seek the specified inode and to retrieve that inode and its extended attributes from the inode scan. This subroutine combines **gpfs\_seek\_inode64()** and **get\_next\_inode64()**, but will only return the specified inode.

The data returned by **gpfs\_next\_inode64()** is overwritten by subsequent calls to **gpfs\_next\_inode64()**, **gpfs\_seek\_inode64()**, or **gpfs\_stat\_inode64()**.

The **termIno** parameter provides a way to partition an inode scan so it can be run on more than one node. It is only used by this call to control prefetching.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### iscan

A pointer to an inode scan descriptor.

#### ino

The inode number to be returned.

#### termIno

Prefetches inodes up to this inode. The caller might specify **maxIno** from **gpfs\_open\_inodescan()** or 0 to allow prefetching over the entire inode file.

#### iattr

A pointer to the returned pointer to the file's **iattr**.

### Exit status

If the **gpfs\_stat\_inode64()** subroutine is successful, it returns a value of 0 and the **iattr** parameter is set to point to **gpfs\_iattr\_t**.

If the **gpfs\_stat\_inode64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### Exceptions

None.

## **gpfs\_stat\_inode64()**

### **Error status**

Error codes include but are not limited to the following:

#### **EPERM**

The caller must have superuser privilege.

#### **ENOSYS**

The **gpfs\_stat\_inode()** subroutine is not supported under the current file system format.

#### **ESTALE**

The cached file system information was not valid.

#### **ENOMEM**

The buffer is too small.

#### **GPFS\_EINVAL\_ISCAN**

Incorrect parameters.

#### **GPFS\_EHOLE\_IN\_IFILE**

The inode scan is reading only the inodes that have been copied to a snapshot and this inode has not yet been copied.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_stat\_inode\_with\_xattrs() subroutine

Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Synopsis

```
#include <gpfs.h>
int gpfs_stat_inode_with_xattrs(gpfs_iscan_t *iscan,
 gpfs_ino_t ino,
 gpfs_ino_t termIno,
 const gpfs_iattr_t **iattr,
 const char **xattrBuf,
 unsigned int *xattrBufLen);
```

### Description

The **gpfs\_stat\_inode\_with\_xattrs()** subroutine is used to seek the specified inode and to retrieve that inode and its extended attributes from the inode scan. This subroutine combines **gpfs\_seek\_inode()** and **get\_next\_inode()**, but will only return the specified inode.

The data returned by **gpfs\_next\_inode()** is overwritten by subsequent calls to **gpfs\_next\_inode()**, **gpfs\_seek\_inode()**, or **gpfs\_stat\_inode\_with\_xattrs()**.

The **termIno** parameter provides a way to partition an inode scan such that it can be run on more than one node. It is only used by this call to control prefetching.

The returned values for **xattrBuf** and **xattrBufLen** must be provided to **gpfs\_next\_xattr()** to obtain the extended attribute names and values. The buffer used for the extended attributes is overwritten by subsequent calls to **gpfs\_next\_inode()**, **gpfs\_seek\_inode()**, or **gpfs\_stat\_inode\_with\_xattrs()**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### Parameters

#### iscan

A pointer to an inode scan descriptor.

#### ino

The inode number to be returned.

#### termIno

Prefetches inodes up to this inode. The caller might specify **maxIno** from **gpfs\_open\_inodescan()** or 0 to allow prefetching over the entire inode file.

#### iattr

A pointer to the returned pointer to the file's **iattr**.

#### xattrBuf

A pointer to the returned pointer to the **xattr** buffer.

#### xattrBufLen

The returned length of the **xattr** buffer.

## **gpfs\_stat\_inode\_with\_xattrs()**

### **Exit status**

If the **gpfs\_stat\_inode\_with\_xattrs()** subroutine is successful, it returns a value of 0 and the **iatr** parameter is set to point to **gpfs\_iattr\_t**. If the **gpfs\_stat\_inode\_with\_xattrs()** subroutine is successful, but there are no more inodes before the **termIno** parameter, or if the requested inode does not exist, it returns a value of 0 and the **iatr** parameter is set to NULL.

If the **gpfs\_stat\_inode\_with\_xattrs()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EPERM**

The caller must have superuser privilege.

#### **ENOSYS**

The **gpfs\_stat\_inode\_with\_xattrs()** subroutine is not supported under the current file system format.

#### **ESTALE**

The cached file system information was not valid.

#### **ENOMEM**

The buffer is too small.

#### **GPFS\_EINVAL\_ISCAN**

Incorrect parameters.

#### **GPFS\_EHOLE\_IN\_IFILE**

The inode scan is reading only the inodes that have been copied to a snapshot and this inode has not yet been copied.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfs\_stat\_inode\_with\_xattrs64() subroutine**

Seeks the specified inode and retrieves that inode and its extended attributes from the inode scan.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Synopsis**

```
#include <gpfs.h>
int gpfs_stat_inode_with_xattrs64(gpfs_iscan_t *iscan,
 gpfs_ino64_t ino,
 gpfs_ino64_t termIno,
 const gpfs_iattr64_t **iattr,
 const char **xattrBuf,
 unsigned int *xattrBufLen);
```

### **Description**

The **gpfs\_stat\_inode\_with\_xattrs64()** subroutine is used to seek the specified inode and to retrieve that inode and its extended attributes from the inode scan. This subroutine combines **gpfs\_seek\_inode64()** and **get\_next\_inode64()**, but will only return the specified inode.

The data returned by **get\_next\_inode64()** is overwritten by subsequent calls to **gpfs\_next\_inode64()**, **gpfs\_seek\_inode64()**, or **gpfs\_stat\_inode\_with\_xattrs64()**.

The **termIno** parameter provides a way to partition an inode scan so it can be run on more than one node. It is only used by this call to control prefetching.

The returned values for **xattrBuf** and **xattrBufLen** must be provided to **gpfs\_next\_xattr()** to obtain the extended attribute names and values. The buffer used for the extended attributes is overwritten by subsequent calls to **gpfs\_next\_inode64()**, **gpfs\_seek\_inode64()**, or **gpfs\_stat\_inode\_with\_xattrs64()**.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- libgpfs.a for AIX
- libgpfs.so for Linux

### **Parameters**

#### **iscan**

A pointer to an inode scan descriptor.

#### **ino**

The inode number to be returned.

#### **termIno**

Prefetches inodes up to this inode. The caller might specify **maxIno** from **gpfs\_open\_inodescan64()** or 0 to allow prefetching over the entire inode file.

#### **iattr**

A pointer to the returned pointer to the file's **iattr**.

#### **xattrBuf**

A pointer to the returned pointer to the **xattr** buffer.

#### **xattrBufLen**

The returned length of the **xattr** buffer.

## **gpfs\_stat\_inode\_with\_xattrs64()**

### **Exit status**

If the **gpfs\_stat\_inode\_with\_xattrs64()** subroutine is successful, it returns a value of 0 and the **iattr** parameter is set to point to **gpfs\_iattr\_t**. If the **gpfs\_stat\_inode\_with\_xattrs64()** subroutine is successful, but there are no more inodes before the **termIno** parameter, or if the requested inode does not exist, it returns a value of 0 and the **iattr** parameter is set to NULL.

If the **gpfs\_stat\_inode\_with\_xattrs64()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following:

#### **EPERM**

The caller must have superuser privilege.

#### **ENOSYS**

The **gpfs\_stat\_inode\_with\_xattrs64()** subroutine is not supported under the current file system format.

#### **ESTALE**

The cached file system information was not valid.

#### **ENOMEM**

The buffer is too small.

#### **GPFS\_EINVAL\_ISCAN**

Incorrect parameters.

#### **GPFS\_EHOLE\_IN\_IFILE**

The inode scan is reading only the inodes that have been copied to a snapshot and this inode has not yet been copied.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfs\_stat\_x() subroutine

Returns extended status information for a GPFS file with specified accuracy.

### Library

GPFS library. Runs on Linux, AIX, and Windows.

### Synopsis

```
#include <gpfs.h>
int gpfs_stat_x(const char *pathname,
 unsigned int *st_litemask,
 gpfs_iattr64_t *iattr,
 size_t iattrBufLen);
```

### Description

The **gpfs\_stat\_x()** subroutine is similar to the **gpfs\_stat()** subroutine but returns more information in a **gpfs\_iattr64** structure that is defined in **gpfs.h**. This subroutine is supported only on the Linux operating system.

Your program must verify that the version of the **gpfs\_iattr64** structure that is returned in the field **ia\_version** is the same as the version that you are using. Versions are defined in **gpfs.h** with the pattern **GPFS\_IA64\_VERSION\***.

File permissions **read**, **write**, and **execute** are not required for the specified file, but all the directories in the specified path must be searchable.

**Note:** Compile any program that uses this subroutine with the **-lgpfs** flag from the following library:

- **libgpfs.so** for Linux

### Parameters

#### \***pathname**

A pointer to the path of the file for which information is requested.

#### \***st\_litemask**

A pointer to a bitmask specification of the items that you want to be returned exactly. Bitmasks are defined in **gpfs.h** with the pattern **GPFS\_SLITE\_\*BIT**. The subroutine returns exact values for the items that you specify in the bitmask. The subroutine also sets bits in the bitmask to indicate any other items that are exact.

#### \***iattr**

A pointer to a **gpfs\_iattr64\_t** structure in which the information is returned. The structure is described in the **gpfs.h** file.

#### **iattrBufLen**

The length of your **gpfs\_iattr64\_t** structure, as given by **sizeof(myStructure)**. The subroutine does not write data past this limit. The field **ia\_reclen** in the returned structure is the length of the **gpfs\_iattr64\_t** structure that the system is using.

### Exit status

If the **gpfs\_stat\_x()** subroutine is successful, it returns a value of 0.

If the **gpfs\_stat\_x()** subroutine is unsuccessful, it returns a value of -1 and sets the global error variable **errno** to indicate the nature of the error.

## **gpfs\_stat\_x()**

### **Exceptions**

None.

### **Error status**

Error codes include but are not limited to the following errors:

#### **EINVAL**

The path name does not refer to a GPFS file or a regular file.

#### **ENOENT**

The file does not exist.

#### **ENOSYS**

The `gpfs_stat_x()` subroutine is not supported under the current file system format.

#### **ESTALE**

The cached file system information was invalid.

### **Location**

`/usr/lpp/mmfs/lib`

---

## **gpfsFcntlHeader\_t structure**

Contains declaration information for the **gpfs\_fcntl()** subroutine.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Structure**

```
typedef struct
{
 int totalLength;
 int fcntlVersion;
 int errorOffset;
 int fcntlReserved;
} gpfsFcntlHeader_t;
```

### **Description**

The **gpfsFcntlHeader\_t** structure contains size, version, and error information for the **gpfs\_fcntl()** subroutine.

### **Members**

#### **totalLength**

This field must be set to the total length, in bytes, of the data structure being passed in this subroutine. This includes the length of the header and all hints and directives that follow the header.

The total size of the data structure *cannot* exceed the value of **GPFS\_MAX\_FCNTL\_LENGTH**, as defined in the header file `gpfs_fcntl.h`. The current value of **GPFS\_MAX\_FCNTL\_LENGTH** is 64 KB.

#### **fcntlVersion**

This field must be set to the current version number of the **gpfs\_fcntl()** subroutine, as defined by **GPFS\_FCNTL\_CURRENT\_VERSION** in the header file `gpfs_fcntl.h`. The current version number is one.

#### **errorOffset**

If an error occurs processing a system call, GPFS sets this field to the offset within the parameter area where the error was detected.

For example,

1. An incorrect version number in the header would cause **errorOffset** to be set to zero.
2. An error in the first hint following the header would set **errorOffset** to **sizeof(header)**.

If no errors are found, GPFS does not alter this field.

#### **fcntlReserved**

This field is currently unused.

For compatibility with future versions of GPFS, set this field to zero.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfsGetDataBlkDiskIdx\_t**

---

### **gpfsGetDataBlkDiskIdx\_t structure**

Obtains the FPO data block location of a file.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 FilemapIn filemapIn;
 FilemapOut filemapOut;
} gpfsGetDataBlkDiskIdx_t;
```

#### **Description**

The **gpfsGetDataBlkDiskIdx\_t** structure is used to obtain a file's blocks on-disk or in-memory location information.

#### **Members**

##### **structLen**

Length of the **gpfsGetDataBlkDiskIdx\_t** structure.

##### **structType**

The structure identifiers are **GPFS\_FCNTL\_GET\_DATABLKDISKIDX** and **GPFS\_FCNTL\_GET\_DATABLKLOC**. **GPFS\_FCNTL\_GET\_DATABLKDISKIDX** can be used to retrieve the disk ID on which the data blocks are located. **GPFS\_FCNTL\_GET\_DATABLKLOC** can be used to retrieve the disk ID on which the data blocks are located and the node ID on which data block buffers (in memory) are located.

##### **filemapIn**

Input parameters:

```
struct FilemapIn {
 long long startOffset;
 long long skipfactor;
 long long length;
 int mreplicas;
 int reserved;
} FilemapIn;
```

##### **startOffset**

Start offset in bytes.

##### **skipfactor**

Number of bytes to skip before the next offset read. Could be 0 in which case the meta block size will be used as the **skipfactor**.

##### **length**

Number of bytes to read; **startOffset + length / skipfactor = numBlksReturned**

##### **mreplicas**

Number of replicas that the user wants.

0 - all; 1 - primary; 2 - primary and 1st replica; 3 - primary, 1st and 2nd replica

##### **reserved**

Reserved field that should not be used.

**filemapOut**

Output data:

```
struct FilemapOut {
 int numReplicasReturned;
 int numBlksReturned;
 int blockSize;
 int blockSizeHigh;
 char buffer [GPFS_MAX_FCNTL_LENGTH-1024];
} FilemapOut;
```

**numReplicasReturned**

Number of replicas returned.

**numBlksReturned**

Number of data blocks returned.

**blockSize**

Low 32bits of meta block size. Only valid if **skipfactor** in the input is 0. Otherwise, **blockSize** will be 0.

**blockSizeHigh**

High 32bits of meta block size. Only valid if **structType** is **GPFS\_FCNTL\_GET\_DATABLKLOC** and **skipfactor** in the input is 0. Otherwise, **blockSizeHigh** will be 0.

**buffer [GPFS\_MAX\_FCNTL\_LENGTH-1024]**

Buffer in which the disk ID and node ID are stored. If the buffer cannot store all of the requested information, you will have to calculate a new **startOffset** and **length** in the input data according to the output data and then repeat the call.

If **structType** is **GPFS\_FCNTL\_GET\_DATABLKDISKIDX**, the format is:

```
If number of replicas returned is 1,
 offset(64bits) disid1(32bits)
 offset(64bits) diskid1(32bits)
 ...
 If number of replicas returned is 2,
 offset(64bits) diskid1(32bits) diskid2(32bits)
 offset(64bits) diskid1(32bits) diskid2(32bits)
 ...
 If number of replicas returned is 3,
 offset(64bits) diskid1(32bits) diskid2(32bits) diskid3(32bits)
 offset(64bits) diskid1(32bits) diskid2(32bits) diskid3(32bits)
 ...
```

If **structType** is **GPFS\_FCNTL\_GET\_DATABLKLOC**, the format is:

```
If number of replicas returned is 1,
 disid1(32bits) nodeid1(32bits)
 disid1(32bits) nodeid1(32bits)
 ...
 If number of replicas returned is 2,
 disid1(32bits) diskid2(32bits) nodeid1(32bits) nodeid2(32bits)
 disid1(32bits) diskid2(32bits) nodeid1(32bits) nodeid2(32bits)
 ...
 If number of replicas returned is 3,
 disid1(32bits) diskid2(32bits) diskid3(32bits) nodeid1(32bits) nodeid2(32bits) nodeid3(32bits)
 disid1(32bits) diskid2(32bits) diskid3(32bits) nodeid1(32bits) nodeid2(32bits) nodeid3(32bits)
 ...
```

## **gpfsGetDataBlkDiskIdx\_t**

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfsGetFilesetName\_t structure**

Obtains the fileset name of a file.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 char buffer[GPFS_FCNTRL_MAX_NAME_BUFFER];
} gpfsGetFilesetName_t;
```

### **Description**

The `gpfsGetFilesetName_t` structure is used to obtain a file's fileset name.

### **Members**

#### **structLen**

Length of the `gpfsGetFilesetName_t` structure.

#### **structType**

Structure identifier `GPFS_FCNTRL_GET_FILESETNAME`.

#### **buffer**

The size of the buffer may vary, but must be a multiple of eight. Upon successful completion of the call, the buffer contains a null-terminated character string for the name of the requested object.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfsGetReplication\_t**

---

### **gpfsGetReplication\_t structure**

Obtains the replication factors of a file.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 int metadataReplicas;
 int maxMetadataReplicas;
 int dataReplicas;
 int maxDataReplicas;
 int status;
 int reserved;
} gpfsGetReplication_t;
```

#### **Description**

The **gpfsGetReplication\_t** structure is used to obtain a file's replication factors.

#### **Members**

##### **structLen**

Length of the **gpfsGetReplication\_t** structure.

##### **structType**

Structure identifier **GPFS\_FCNTRL\_GET\_REPLICATION**.

##### **metadataReplicas**

Returns the current number of copies of indirect blocks for the file.

##### **maxMetadataReplicas**

Returns the maximum number of copies of indirect blocks for a file.

##### **dataReplicas**

Returns the current number of copies of the data blocks for a file.

##### **maxDataReplicas**

Returns the maximum number of copies of data blocks for a file.

##### **status**

Returns the status of the file.

##### **reserved**

Unused, but should be set to 0.

#### **Error status**

These values are returned in the **status** field:

##### **GPFS\_FCNTRL\_STATUS\_EXPOSED**

This file may have some data where the only replicas are on suspended disks; implies some data may be lost if suspended disks are removed.

##### **GPFS\_FCNTRL\_STATUS\_ILLREPLICATE**

This file may not be properly replicated; that is, some data may have fewer or more than the desired number of replicas, or some replicas may be on suspended disks.

**GPFS\_FCNTL\_STATUS\_UNBALANCED**

This file may not be properly balanced.

**GPFS\_FCNTL\_STATUS\_DATAUPDATEMISS**

This file has stale data blocks on at least one of the disks that are marked as unavailable or recovering.

**GPFS\_FCNTL\_STATUS\_METAUPDATEMISS**

This file has stale indirect blocks on at least one unavailable or recovering disk.

**GPFS\_FCNTL\_STATUS\_ILLPLACED**

This file may not be properly placed; that is, some data may be stored in an incorrect storage pool.

**Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfsGetSetXAttr\_t**

---

### **gpfsGetSetXAttr\_t structure**

Obtains or sets extended attribute values.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 int nameLen;
 int bufferLen;
 unsigned int flags;
 int errReasonCode;
 char buffer[0];
} gpfsGetSetXAttr_t;
```

#### **Description**

The **gpfsGetSetXAttr\_t** structure is used to obtain extended attributes.

#### **Members**

##### **structLen**

Length of the **gpfsGetSetXAttr\_t** structure.

##### **structType**

Structure identifier **GPFS\_FCNTL\_GET\_XATTR** or **GPFS\_FCNTL\_SET\_XATTR**.

##### **nameLen**

Length of the attribute name. May include a trailing '\0' character.

##### **bufferLen**

For **GPFS\_FCNTL\_GET\_XATTR**: Input, length of the buffer; output, length of the attribute value.

For **GPFS\_FCNTL\_SET\_XATTR**: Input, length of the attribute value. Specify **-1** to delete an attribute.

##### **errReasonCode**

Reason code.

##### **flags**

The following flags are recognized:

- **GPFS\_FCNTL\_XATTRFLAG\_NONE**
- **GPFS\_FCNTL\_XATTRFLAG\_SYNC**
- **GPFS\_FCNTL\_XATTRFLAG\_CREATE**
- **GPFS\_FCNTL\_XATTRFLAG\_REPLACE**
- **GPFS\_FCNTL\_XATTRFLAG\_DELETE**
- **GPFS\_FCNTL\_XATTRFLAG\_NO\_CTIME**
- **GPFS\_FCNTL\_XATTRFLAG\_RESERVED**

##### **buffer**

Buffer for the attribute name and value.

For **GPFS\_FCNTL\_GET\_XATTR**:

Input: The name begins at offset 0 and must be null terminated.

Output: The name is returned unchanged; the value begins at **nameLen** rounded up to a multiple of 8.

For **GPFS\_FCNTL\_SET\_XATTR**:

Input: The name begins at offset 0 and must be null terminated. The value begins at **nameLen** rounded up to a multiple of 8.

The actual length of the buffer should be **nameLen** rounded up to a multiple of 8, plus the length of the attribute value rounded up to a multiple of 8.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

## **gpfsGetSnapshotName\_t**

---

### **gpfsGetSnapshotName\_t structure**

Obtains the snapshot name of a file.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 char buffer[GPFS_FCNTRL_MAX_NAME_BUFFER];
} gpfsGetSnapshotName_t;
```

#### **Description**

The **gpfsGetSnapshotName\_t** structure is used to obtain a file's snapshot name. If the file is not part of a snapshot, a zero length snapshot name will be returned.

#### **Members**

##### **structLen**

Length of the **gpfsGetSnapshotName\_t** structure.

##### **structType**

Structure identifier **GPFS\_FCNTRL\_GET\_SNAPSHOTNAME**.

##### **buffer**

The size of the buffer may vary, but must be a multiple of eight. Upon successful completion of the call, the buffer contains a null-terminated character string for the name of the requested object.

#### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfsGetStoragePool\_t structure**

Obtains the storage pool name of a file.

### **Library**

GPFS Library (`libgpfs.a` for AIX, `libgpfs.so` for Linux)

### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 char buffer[GPFS_FCNTL_MAX_NAME_BUFFER];
} gpfsGetStoragePool_t;
```

### **Description**

The `gpfsGetStoragePool_t` structure is used to obtain a file's storage pool name.

### **Members**

#### **structLen**

Length of the `gpfsGetStoragePool_t` structure.

#### **structType**

Structure identifier `GPFS_FCNTL_GET_STORAGEPOOL`.

#### **buffer**

The size of the buffer may vary, but must be a multiple of eight. Upon successful completion of the call, the buffer contains a null-terminated character string for the name of the requested object.

### **Location**

`/usr/lpp/mmfs/lib/libgpfs.a` for AIX

`/usr/lpp/mmfs/lib/libgpfs.so` for Linux

## **gpfsListXAttr\_t**

---

### **gpfsListXAttr\_t structure**

Lists extended attributes.

#### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

#### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 int bufferLen;
 int errReasonCode;
 char buffer[0];
} gpfsListXAttr_t;
```

#### **Description**

The **gpfsListXAttr\_t** structure is used to list extended attributes.

#### **Members**

##### **structLen**

Length of the **gpfsListXAttr\_t** structure.

##### **structType**

Structure identifier **GPFS\_FCNTL\_LIST\_XATTR**.

##### **bufferLen**

Input: Length of the buffer. Output: Length of the returned list of names.

The actual length of the buffer required depends on the number of attributes set and the length of each attribute name. If the buffer provided is too small for all of the returned names, the **errReasonCode** will be set to **GPFS\_FCNTL\_ERR\_BUFFER\_TOO\_SMALL**, and **bufferLen** will be set to the minimum size buffer required to list all attributes. An initial buffer length of 0 may be used to query the attributes and determine the correct buffer size for this file.

##### **errReasonCode**

Reason code.

##### **buffer**

Buffer for the returned list of names. Each attribute name is prefixed with a one-byte name length. The attribute name may contain embedded null bytes. The attribute name may be null-terminated in which case the name length includes this null byte. The next attribute name follows immediately in the buffer (and is prefixed with its own length). Following the last name, a '\0' is appended to terminate the list. The returned **bufferLen** includes the final '\0'.

#### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfsRestripeData\_t structure**

Restripes the data blocks of a file.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 int options;
 int errReason;
 int errValue1;
 int errValue2;
 int reserved1;
 int reserved2;
} gpfsRestripeData_t;
```

### **Description**

The **gpfsRestripeData\_t** structure is used to restripe a file's data blocks to updates its replication and migrate its data. The data movement is always done immediately.

### **Members**

#### **structLen**

Length of the **gpfsRestripeData\_t** structure.

#### **structType**

Structure identifier **GPFS\_FCNTRL\_RESTRIPE\_DATA**.

#### **options**

Options for restripe command. See the **mmrestripefile** command for complete definitions.

##### **GPFS\_FCNTRL\_RESTRIPE\_M**

Migrate critical data off of suspended disks.

##### **GPFS\_FCNTRL\_RESTRIPE\_R**

Replicate data against subsequent failure.

##### **GPFS\_FCNTRL\_RESTRIPE\_P**

Place file data in assigned storage pool.

##### **GPFS\_FCNTRL\_RESTRIPE\_B**

Rebalance file data.

#### **errReason**

Reason code describing the failure. Possible codes are defined in “Error status” on page 820.

#### **errValue1**

Returned value depending upon **errReason**.

#### **errValue2**

Returned value depending upon **errReason**.

#### **reserved1**

Unused, but should be set to 0.

#### **reserved2**

Unused, but should be set to 0.

## **gpfsRestripeData\_t**

### **Error status**

These values are returned in the **errReason** field:

#### **GPFS\_FCNTL\_ERR\_NO\_REPLICA\_GROUP**

Not enough replicas could be created because the desired degree of replication is larger than the number of failure groups.

#### **GPFS\_FCNTL\_ERR\_NO\_REPLICA\_SPACE**

Not enough replicas could be created because there was not enough space left in one of the failure groups.

#### **GPFS\_FCNTL\_ERR\_NO\_BALANCE\_SPACE**

There was not enough space left on one of the disks to properly balance the file according to the current stripe method.

#### **GPFS\_FCNTL\_ERR\_NO\_BALANCE\_AVAILABLE**

The file could not be properly balanced because one or more disks are unavailable.

#### **GPFS\_FCNTL\_ERR\_ADDR\_BROKEN**

All replicas were on disks that have since been deleted from the stripe group.

#### **GPFS\_FCNTL\_ERR\_NO\_IMMUTABLE\_DIR**

No immutable attribute can be set on directories.

#### **GPFS\_FCNTL\_ERR\_NO\_IMMUTABLE\_SYSFILE**

No immutable attribute can be set on system files.

#### **GPFS\_FCNTL\_ERR\_IMMUTABLE\_FLAG**

Immutable and indefinite retention flag is wrong.

#### **GPFS\_FCNTL\_ERR\_IMMUTABLE\_PERM**

Immutable and indefinite retention flag is wrong.

#### **GPFS\_FCNTL\_ERR\_APPENDONLY\_CONFLICT**

The **appendOnly** flag should be set separately.

#### **GPFS\_FCNTL\_ERR\_NOIMMUTABLE\_ONSNAP**

Cannot set immutable or **appendOnly** on snapshots.

#### **GPFS\_FCNTL\_ERR\_FILE\_HAS\_XATTRS**

An attempt to change **maxDataReplicas** or **maxMetadataReplicas** was made on a file that has extended attributes.

#### **GPFS\_FCNTL\_ERR\_NOT\_GPFS\_FILE**

This file is not part of a GPFS file system.

### **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## gpfsSetReplication\_t structure

Sets the replication factors of a file.

### Library

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### Structure

```
typedef struct {
 int structLen;
 int structType;
 int metadataReplicas;
 int maxMetadataReplicas;
 int dataReplicas;
 int maxDataReplicas;
 int errReason;
 int errValue1;
 int errValue2;
 int reserved;
} gpfsSetReplication_t;
```

### Description

The **gpfsGetReplication\_t** structure is used to set a file's replication factors. However, the directive does not cause the file to be restriped immediately. Instead, the caller must append a **gpfsRestripeData\_t** directive or invoke an explicit restripe using the **mmrestripefs** or **mmrestripefile** command.

### Members

#### **structLen**

Length of the **gpfsSetReplication\_t** structure.

#### **structType**

Structure identifier **GPFS\_FCNTRL\_SET\_REPLICATION**.

#### **metadataReplicas**

Specifies how many copies of the file system's metadata to create. Enter a value of 1 or 2, but not greater than the value of the **maxMetadataReplicas** attribute of the file. A value of 0 indicates not to change the current value.

#### **maxMetadataReplicas**

The maximum number of copies of indirect blocks for a file. Space is reserved in the inode for all possible copies of pointers to indirect blocks. Valid values are 1 and 2, but cannot be less than **DefaultMetadataReplicas**. The default is 1. A value of 0 indicates not to change the current value.

#### **dataReplicas**

Specifies how many copies of the file data to create. Enter a value of 1 or 2, but not greater than the value of the **maxDataReplicas** attribute of the file. A value of 0 indicates not to change the current value.

#### **maxDataReplicas**

The maximum number of copies of data blocks for a file. Space is reserved in the inode and indirect blocks for all possible copies of pointers to data blocks. Valid values are 1 and 2, but cannot be less than **DefaultDataReplicas**. The default is 1. A value of 0 indicates not to change the current value.

#### **errReason**

Reason code describing the failure. Possible codes are defined in "Error status" on page 822.

#### **errValue1**

Returned value depending upon **errReason**.

## **gpfsSetReplication\_t**

### **errValue2**

Returned value depending upon **errReason**.

### **reserved**

Unused, but should be set to 0.

## **Error status**

These values are returned in the **errReason** field:

### **GPFS\_FCNTL\_ERR\_NONE**

Command was successful or no reason information was returned.

### **GPFS\_FCNTL\_ERR\_METADATA\_REPLICAS\_RANGE**

Field **metadataReplicas** is out of range. Fields **errValue1** and **errValue2** contain the valid lower and upper range boundaries.

### **GPFS\_FCNTL\_ERR\_MAXMETADATA\_REPLICAS\_RANGE**

Field **maxMetadataReplicas** is out of range. Fields **errValue1** and **errValue2** contain the valid lower and upper range boundaries.

### **GPFS\_FCNTL\_ERR\_DATA\_REPLICAS\_RANGE**

Field **dataReplicas** is out of range. Fields **errValue1** and **errValue2** contain the valid lower and upper range boundaries.

### **GPFS\_FCNTL\_ERR\_MAXDATA\_REPLICAS\_RANGE**

Field **maxDataReplicas** is out of range. Fields **errValue1** and **errValue2** contain the valid lower and upper range boundaries.

### **GPFS\_FCNTL\_ERR\_FILE\_NOT\_EMPTY**

An attempt to change **maxMetadataReplicas** or **maxDataReplicas** or both was made on a file that is not empty.

### **GPFS\_FCNTL\_ERR\_REPLICAS\_EXCEED\_FGMAX**

Field **metadataReplicas**, or **dataReplicas**, or both exceed the number of failure groups. Field **errValue1** contains the maximum number of metadata failure groups. Field **errValue2** contains the maximum number of data failure groups.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## **gpfsSetStoragePool\_t structure**

Sets the assigned storage pool of a file.

### **Library**

GPFS Library (libgpfs.a for AIX, libgpfs.so for Linux)

### **Structure**

```
typedef struct {
 int structLen;
 int structType;
 int errReason;
 int errValue1;
 int errValue2;
 int reserved;
 char buffer[GPFS_FCNTRL_MAX_NAME_BUFFER];
} gpfsSetStoragePool_t;
```

### **Description**

The **gpfsSetStoragePool\_t** structure is used to set a file's assigned storage pool. However, the directive does not cause the file data to be migrated immediately. Instead, the caller must append a **gpfsRestripeData\_t** directive or invoke an explicit restripe with the **mmrestripefs** or **mmrestripefile** command. The caller must have su or root privileges to change a storage pool assignment.

### **Members**

#### **structLen**

Length of the **gpfsSetStoragePool\_t** structure.

#### **structType**

Structure identifier **GPFS\_FCNTRL\_SET\_STORAGEPOOL**.

#### **errReason**

Reason code describing the failure. Possible codes are defined in "Error status."

#### **errValue1**

Returned value depending upon **errReason**.

#### **errValue2**

Returned value depending upon **errReason**.

#### **reserved**

Unused, but should be set to 0.

#### **buffer**

The name of the storage pool for the file's data. Only user files may be reassigned to different storage pool. System files, including all directories, must reside in the system pool and may not be moved.

The size of the buffer may vary, but must be a multiple of eight.

### **Error status**

These values are returned in the **errReason** field:

#### **GPFS\_FCNTRL\_ERR\_NONE**

Command was successful or no reason information was returned.

#### **GPFS\_FCNTRL\_ERR\_NOPERM**

User does not have permission to perform the requested operation.

## **gpfsSetStoragePool\_t**

### **GPFS\_FCNTL\_ERR\_INVALID\_STORAGE\_POOL**

Invalid storage pool name was given.

### **GPFS\_FCNTL\_ERR\_INVALID\_STORAGE\_POOL\_TYPE**

Invalid storage pool. File cannot be assigned to given pool.

### **GPFS\_FCNTL\_ERR\_INVALID\_STORAGE\_POOL\_ISDIR**

Invalid storage pool. Directories cannot be assigned to given pool.

### **GPFS\_FCNTL\_ERR\_INVALID\_STORAGE\_POOL\_ISLNK**

Invalid storage pool. System files cannot be assigned to given pool.

### **GPFS\_FCNTL\_ERR\_INVALID\_STORAGE\_POOL\_ISSYS**

Invalid storage pool. System files cannot be assigned to given pool.

### **GPFS\_FCNTL\_ERR\_STORAGE\_POOL\_NOTENABLED**

File system has not been upgraded to support storage pools.

## **Location**

/usr/lpp/mmfs/lib/libgpfs.a for AIX

/usr/lpp/mmfs/lib/libgpfs.so for Linux

---

## Chapter 4. GPFS user exits

Apart from the user exits define by using the **mmaddcallback** command, GPFS provides three more user exits: **mmsdrbackup**, **nsddevices**, and **syncfsconfig**.

Table 24 summarizes the GPFS-specific user exits.

*Table 24. GPFS user exits*

| User exit                            | Purpose                                                                              |
|--------------------------------------|--------------------------------------------------------------------------------------|
| "mmsdrbackup user exit" on page 826  | Performs a backup of the GPFS configuration data.                                    |
| "nsddevices user exit" on page 827   | Identifies local physical devices that are used as GPFS Network Shared Disks (NSDs). |
| "syncfsconfig user exit" on page 828 | Keeps file system configuration data in replicated clusters synchronized.            |

## **mmsdrbackup**

---

### **mmsdrbackup user exit**

Performs a backup of the GPFS configuration data.

#### **Description**

The **/var/mmfs/etc/mmsdrbackup** user exit, when properly installed on the primary GPFS configuration server, is called asynchronously every time there is a change to the GPFS master configuration file. You can use this user exit to create a backup of the GPFS configuration data.

Read the sample file **/usr/lpp/mmfs/samples/mmsdrbackup.sample** for a detailed description on how to code and install this user exit.

The type of backup that is created depends on the configuration of the cluster:

- If the Cluster Configuration Repository (CCR) is enabled, then a CCR backup is created. This type of backup applies to IBM Spectrum Scale V4.2.0 or later.
- Otherwise, a mmsdrfs backup is created.

For more information about the CCR, see “[mmcrcluster command](#)” on page 239.

**Note:** The **mmsdrbackup** user exit is supported from IBM Spectrum Scale 4.2.0 or later. It must not be used on clusters which have nodes running earlier versions of IBM Spectrum Scale.

#### **Parameters**

The generation number of the most recent version of the GPFS configuration data.

#### **Exit status**

The **mmsdrbackup** user exit returns a value of zero.

#### **Location**

`/var/mmfs/etc`

---

## nsddevices user exit

Identifies local physical devices that are used as GPFS Network Shared Disks (NSDs).

### Description

The **/var/mmfs/etc/nsddevices** user exit, when properly installed, is invoked synchronously by the GPFS daemon during its disk discovery processing. The purpose of this procedure is to discover and verify the physical devices on each node that correspond to the disks previously defined to GPFS with the **mmcrnsd** command. The **nsddevices** user exit can be used to either replace or to supplement the disk discovery procedure of the GPFS daemon.

Read the sample file **/usr/lpp/mmfs/samples/nsddevices.sample** for a detailed description on how to code and install this user exit.

### Parameters

None.

### Exit status

The **nsddevices** user exit should return either zero or one.

When the **nsddevices** user exit returns a value of zero, the GPFS disk discovery procedure is bypassed.

When the **nsddevices** user exit returns a value of one, the GPFS disk discovery procedure is performed and the results are concatenated with the results from the **nsddevices** user exit.

### Location

`/var/mmfs/etc`

## **syncfsconfig**

---

### **syncfsconfig user exit**

Keeps file system configuration data in replicated clusters synchronized.

#### **Description**

The **/var/mmfs/etc-syncfsconfig** user exit, when properly installed, will be synchronously invoked after each command that may change the configuration of a file system. Examples of such commands are: **mmadddisk**, **mmdeldisk**, **mmchfs**, and so forth. The **syncfsconfig** user exit can be used to keep the file system configuration data in replicated GPFS clusters automatically synchronized.

Read the sample file **/usr/lpp/mmfs/samples-syncfsconfig.sample** for a detailed description on how to code and install this user exit.

#### **Parameters**

None.

#### **Exit status**

The **syncfsconfig** user exit should always return a value of zero.

#### **Location**

**/var/mmfs/etc**

---

## **Chapter 5. IBM Spectrum Scale management API commands**

The following topics describe the IBM Spectrum Scale management API commands in detail.

For more information about using the REST API commands, see *IBM Spectrum Scale management API* in *IBM Spectrum Scale: Administration Guide*.

---

### **Version 1 - deprecated**

Lists the IBM Spectrum Scale management API version 1 commands.

## CES addresses: GET

Gets information about CES addresses.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET cesaddresses** request gets information about the CES (Cluster Export Services) addresses in the cluster. For more information about the fields in the returned data structure, see the topic “mmces command” on page 103.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/cesaddresses`

where

#### **cesaddresses**

Is the cesaddresses resource.

### Request headers

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "cesaddresses": [
 {
 "attributes": "Attributes",
 "cesAddress": "IP",
 "cesGroup": "Group",
 "cesNode": "Node",
 "links": {
 "self": "URL"
 }
 }
],
 "status":
 {
 "code": ReturnCode,
 "message": "ReturnMessage",
 }
}
```

#### **"cesaddresses":**

An array of information about existing CES addresses. Each array element describes one CES address. For more information about the fields in this structure, see the link at the end of this topic.

##### **"Attributes": "Attributes"**

Protocol attributes that are associated with the CES address.

##### **"cesAddress": "IP"**

The IP address.

##### **"cesGroup": "Group"**

The group to which the CES address is assigned.

```

"cesNode": "Node"
 The CES node to which the address is assigned.

"links": ""
 Links.

"self": "URL"
 The URL of the resource.

"status":
 Return status.

"code": ReturnCode,
 The return code for the operation.

"message": "ReturnMessage"
 The return message.

```

## Examples

The following example gets information about the CES addresses in the cluster:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/cesaddresses'
```

In the JSON data that is returned, the return code is 0, indicating that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **cesaddresses** array contains two elements. Each element describes one of the two CES addresses in the cluster:

```
{
 "cesaddresses": [
 {
 "attributes": "",
 "cesAddress": "198.51.100.8",
 "cesGroup": "",
 "cesNode": 0,
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/cesaddresses/198.51.100.8"
 }
 },
 {
 "attributes": "",
 "cesAddress": "198.51.100.10",
 "cesGroup": "",
 "cesNode": 0,
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/cesaddresses/198.51.100.10"
 }
 }
],
 "status": {
 "code": 0,
 "message": ""
 }
}
```

## See also

- “mmces command” on page 103

## CES addresses/{cesAddress}: GET

Gets information about a CES address.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The GET `cesaddresses/{cesAddress}` request gets information about the specified CES (Cluster Export Services) address. For more information about the fields in the data structures that are returned, see the topic “mmces command” on page 103.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/cesaddresses/CesAddress`

where

#### **cesaddresses**

Is the cesaddresses resource. Required.

#### **CesAddress**

Specifies the CES address about which you want to get information. Required.

### Request headers

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "cesaddresses": [
 {
 "attributes": "Attributes",
 "cesAddress": "IP",
 "cesGroup": "Group",
 "cesNode": "Node",
 "links": {
 "self": "URL"
 }
 }
],
 "status":
 {
 "code": ReturnCode,
 "message": "ReturnMessage",
 }
}
```

#### **"cesaddresses":**

An array of information about CES addresses. The array contains one element, which describes the CES address that is specified in the request URL. For more information about the fields in this structure, see the links at the end of this topic.

#### **"Attributes": "Attributes"**

Protocol attributes that are associated with the CES address.

**"cesAddress": "IP"**  
The IP address.

**"cesGroup": "Group"**  
The group to which the CES address is assigned.

**"cesNode": "Node"**  
The CES node to which the address is assigned.

**"links": ""**  
Links.

**"self": "URL"**  
The URL of the resource element.

**"status":**  
Return status.

**"code": *ReturnCode*,**  
The return code for the operation.

**"message": "*ReturnMessage*"**  
The return message.

## Examples

The following example gets information about the CES address 198.51.100.8:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/cesaddresses/198.51.100.8'
```

In the JSON data that is returned, the return code is 0, indicating that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **cesaddresses** array contains one element, which describes the CES address that is specified in the request URL :

```
{
 "cesaddresses": [
 {
 "attributes": "",
 "cesAddress": "198.51.100.8",
 "cesGroup": "",
 "cesNode": 0,
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/cesaddresses/198.51.100.16"
 }
 }
],
 "status": {
 "code": 0,
 "message": ""
 }
}
```

## See also

- “mmces command” on page 103

## CES services: GET

Gets information about CES services.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET cesservices** request gets information about the CES (Cluster Export Services) services in the cluster. For more information about the fields in the returned data structure, see the topic “mmces command” on page 103.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/cesservices`

where

#### cesservices

Is the cesservices resource.

### Request headers

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "cesservices": {
 "protocolNodes": [
 {
 "nodeName": "Node"
 "serviceStates": [
 {
 "running": "{yes | no}",
 "service": "Service"
 }
]
 }
],
 "protocolStates": [
 {
 "enabled": "{yes | no}",
 "links": {
 "self": "URL"
 },
 "service": "Service"
 }
]
 },
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

#### cesservices:

Information about CES services. The information consists of two arrays, **protocolNodes** and

**protocolStates**. Each array contains multiple elements, with one element for each CES service. For more information about the fields in these structures, see the link at the end of this topic.

#### **protocolNodes**

An array of information about protocol nodes. Each array element describes one protocol node.

##### **"nodeName": "Node"**

The name of the node.

##### **serviceStates:**

An array of information about the services for which the node is a protocol node.

##### **"running": "{yes | no}"**

Indicates whether the service is running.

##### **"service": "Service"**

Identifies the service, such as **CES**, **NETWORK**, **NFS**, or **SMB**.

#### **protocolStates:**

An array of information about the services. Each element describes one service:

##### **"enabled": "{yes | no}"**

Indicates that the services is enabled or disabled.

##### **"links":**

Links.

##### **"self": "URL"**

The URL of the service.

##### **"service": "Service"**

Identifies the service, such as **CES**, **NETWORK**, **NFS**, or **SMB**.

#### **"status":**

Return status.

##### **"message": "ReturnMessage"**

The return message.

##### **"code": "ReturnCode",**

The return message.

## Examples

The following example gets information about the current CES addresses:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/sclemgmt/v1/cesservices'
```

In the JSON data that is returned, the return code is 0, which indicates that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **protocolNodes** array is empty, indicating that no nodes are designated as CES service nodes. The elements in the **protocolStates** array indicate that the BLOCK, NFS, OBJ, and SMB services are not enabled:

```
{
 "cesservices": {
 "protocolNodes": [],
 "protocolStates": [
 {
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/sclemgmt/v1/cesservices/block"
 },
 "service": "BLOCK"
 },
 {
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/sclemgmt/v1/cesservices/nfs"
 },
 "service": "NFS"
 },
 {
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/sclemgmt/v1/cesservices/obj"
 },
 "service": "OBJ"
 },
 {
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/sclemgmt/v1/cesservices/smb"
 },
 "service": "SMB"
 }
]
 }
}
```

```

{
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/scalemgmt/v1/cesservices/nfs"
 },
 "service": "NFS"
},
{
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/scalemgmt/v1/cesservices/obj"
 },
 "service": "OBJ"
},
{
 "enabled": "no",
 "links": {
 "self": "https://198.51.100.04:8191/scalemgmt/v1/cesservices/smb"
 },
 "service": "SMB"
}
],
},
"status": {
 "code": 0,
 "message": ""
}
}
}

```

## See also

- “mmces command” on page 103

## CES services/{service}: GET

Gets information about a CES service.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET cesservices/{service}** request gets information about the specified CES (Cluster Export Services) service. For more information about the fields in the data structure that is returned, see the topic “mmces command” on page 103.

### Request URL

`https://REST_API_host:port/sclemgmt/v1/cesservices/Service`

where

#### cesservices

Is the cesservices resource. Required.

#### Service

Specifies the service about which you want to get information. Required.

### Request headers

Accept: application/json

### Request data

No request data.

### Response data

```
{
 "cesservices": {
 "protocolNodes": [
 {
 "nodeName": "Node",
 "serviceStates": [
 {
 "running": "{yes | no}",
 "service": "Service"
 }
]
 }
],
 "protocolStates": [
 {
 "enabled": "{yes | no}",
 "links": {
 "self": "URL"
 },
 "service": "Service"
 }
],
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 }
 }
}
```

## **"cesservices":**

Information about CES services. It consists of two arrays, **protocolNodes** and **protocolStates**. Each array contains one element, which describes the CES service that is specified in the request URL. For more information about the fields in these structures, see the link at the end of this topic.

### **"protocolNodes":**

An array of information about the CES nodes.

#### **"nodeName": "Node"**

The name of the CES node.

#### **"serviceStates":**

An array of information about the CES service that is associated with the CES node.

##### **"running": "{yes | no}"**

Indicates whether the service is running.

##### **"service": "Service"**

Identifies the service, such as **CES**, **NETWORK**, **NFS**, or **SMB**.

### **"protocolStates":**

An array of information about the supported protocols.

#### **"enabled": "{yes | no}"**

Indicates that the services is enabled or disabled.

### **"links":**

An array of links:

#### **"self": "URL"**

The URL of the resource.

#### **"service": "Service"**

Identifies the service, such as **CES**, **NETWORK**, **NFS**, or **SMB**.

## **"status":**

Return status.

### **"code": *ReturnCode*,**

The return message.

### **"message": "*ReturnMessage*"**

The return message.

## **Examples**

The following example gets information about the NFS service that is specified in the request URL:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scaladmgt/v1/cesservices/NFS'
```

In the JSON data that is returned, the return code is 0, indicating that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **protocolNodes** array contains information about the single NFS node. The **ProtocolStates** array indicates that the NFS service is enabled:

```
{
 "cesservices": {
 "protocolNodes": [
 {
 "nodeName": "host5.division.company.com",
 "serviceStates": [
 {
 "running": "no",
 "service": "NFS"
 }
]
 }
]
 }
}
```

```
]
 }
],
"protocolStates": [
{
 "enabled": "yes",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/cesservices/nfs"
 },
 "service": "NFS"
}
],
"status": {
 "code": 0,
 "message": ""
}
}
```

## See also

- “mmces command” on page 103

## **Config: GET**

Gets information about the product configuration.

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

The **GET config** request gets information about the local cluster. For more information about the fields in the data structures that are returned, see the topics “**mmchconfig command**” on page 132 and “**mmlsconfig command**” on page 400.

### **Request URL**

`https://REST_API_host:port/scalemgmt/v1/config`

where

#### **config**

Is the configuration resource.

### **Request headers**

`Content-Type: application/json`

`Accept: application/json`

### **Request data**

No request data.

### **Response data**

The return information and the information that the command retrieves are returned in the same way as they are for the other requests. The parameters that are returned are the same as the configuration attributes that are displayed by the **mmlsconfig** command. For more information, see the “**mmlsconfig command**” on page 400.

### **See also**

- “**mmchconfig command**” on page 132
- “**mmlsconfig command**” on page 400

## Cluster: GET

Gets information about the cluster.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET cluster** request gets information about the local cluster. For more information about the fields in the data structures that are returned, see the topics “[mmlscluster command](#)” on page 397, “[mmchfs command](#)” on page 179, and “[mmlsfs command](#)” on page 411.

### Request URL

`https://REST_API_host:port/sclemgmt/v1/cluster`

where

#### cluster

Is the cluster resource.

### Request headers

`Content-Type: application/json`  
`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "cluster": {
 "cesSummary": {
 "enabledServices": "ServiceList",
 "addressPolicy": "{none | balanced-load | node-affinity | even-coverage}",
 "cesSharedRoot": "Directory",
 "logLevel": Level
 }
 "clusterSummary": {
 "clusterId": ID,
 "clusterName": "ClusterName",
 "primaryServer": "Server",
 "rcpPath": "RemoteFileCopyCommand",
 "rcpSudoWrapper": "Path",
 "repositoryType": "CCR | non-CCR",
 "rshPath": "RemoteShellCommand",
 "rshSudoWrapper": "Path",
 "secondaryServer": "Server",
 "uidDomain": "Domain"
 }
 "cnfsSummary": {
 "cnfsGanesha": "ServerAddress",
 "cnfsMonitorEnabled": "enabled | disabled",
 "cnfsMountPort": "Port",
 "cnfsNFSdprocs": "Number",
 "cnfsReboot": "{true | false}",
 "cnfsSharedRoot": "Directory"
 }
 "links": {
 "self": "URL"
 },
 },
}
```

```

"nodes": [
 {
 "adminLoginName": "ID",
 "adminNodeName": "NodeName",
 "cesNode": {
 "cesGroup": "Group",
 "cesIpList": "IPList",
 "cesState": "{S | X | R | F | P | E | e}",
 "ipAddress": "{IPAddress}"
 },
 "cloudGatewayNode": {
 },
 "cnfsNode": {
 "cnfsGroupId": "GroupID",
 "cnfsIpList": "IPList",
 "cnfsState": "{enabled | disabled}"
 },
 "daemonNodeName": "NodeName",
 "designation": "RoleList",
 "ipAddress": "IPAddress",
 "links": {
 "node": "URL"
 },
 "nodeNumber": Number,
 "otherNodeRoles": RoleList
 }
],
"status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
}
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

**Note:** The structures **cesNode**, **cesSummary**, **cnfsNode**, **cnfsSummary**, and **gatewayNode** appear in the output only when the corresponding role is active on the node.

#### "cluster":

A data structure that describes the local cluster. It contains the following data structures: **cesSummary**, **clusterSummary**, **cnfsSummary**, **links**, and **nodes**.

##### "cesSummary":

CES cluster information.

##### "enabledServices": "ServiceList"

A comma-separated list of the services that are enabled. The list can include **NFS**, **SMB**, **OBJ**, **BLOCK**, or **None**.

##### "addressPolicy": "{none | balanced-load | node-affinity | even-coverage}"

The address policy for distributing CES addresses.

##### "cesSharedRoot": "Directory"

The CES shared root directory, which is used for storing CES shared configuration data.

##### "logLevel": *Level*

The CES log level.

##### "clusterSummary":

Cluster information.

##### "clusterID": "ID"

The ID of the cluster.

```

"clusterName": "ClusterName"
 The name of the cluster.

"primaryServer": "Server"
 The primary server node for GPFS cluster data.

"rcpPath": "RemoteFileCopyCommand"
 The remote file copy command.

"rcpSudoWrapper": ""
 The fully qualified path of the remote file copy program for sudo wrappers.

"repositoryType": "CCR | non-CCR"
 The type of repository that the cluster uses for storing configuration data.

"rshPath": "RemoteShellCommand"
 The remote shell command.

"rshSudoWrapper": "Path"
 The fully qualified path of the remote shell program for sudo wrappers.

"secondaryServer": "Server"
 The secondary server node for GPFS cluster data.

"uidDomain": "Domain"
 The UID domain name of the cluster.

"cnfsSummary":
 CNFS cluster information.

 "cnfsGanesha" : "IPAddress"
 The server address of the Ganesha CNFS server.

 "cnfsMonitorEnabled" : "enabled" | "disabled"
 The state of CNFS monitoring.

 "cnfsMountPort" : "Port"
 The port number of the rpc.mountd daemon.

 "cnfsNFSDprocs" : "Number"
 The number of nfsd server threads.

 "cnfsReboot" : "true" | "false"
 Whether the node reboots when CNFS monitoring detects an unrecoverable problem.

 "cnfsSharedRoot" : "Path"
 A directory that the CNFS subsystem uses.

"links":
 Links.

 "self": "URL"
 The URL of the cluster resource.

"nodes":
 An array of elements that describe the nodes in the cluster. Each element describes one node.

 "adminLoginName": "ID"
 The login ID for GPFS administration commands.

 "adminNodeName": "NodeName"
 The node name that administration commands use to communicate between nodes.

"cesNode":
 CES (Cluster Export Services) information.

```

**"cesGroup": "Group"**  
     The CES group that contains the IP addresses that are used for the swift ring files.

**"cesIpList": "IPList"**  
     A list of the CES IP addresses.

**"cesState": "{S | X | R | F | P | E | e}"**  
     The state of the CES node, where:

|          |                                                      |
|----------|------------------------------------------------------|
| <b>S</b> | Indicates that the node is suspended                 |
| <b>X</b> | Indicates that the network is down.                  |
| <b>R</b> | Indicates that the node does not have a shared root. |
| <b>F</b> | Indicates that the node has failed.                  |
| <b>P</b> | Indicates that the node is starting up.              |
| <b>E</b> | Indicates that the node is using Big Endian.         |
| <b>e</b> | Indicates that the node is using Little Endian.      |

**"ipAddress": "IPAddress"**  
     The IP address of the CES node.

**"cloudGatewayNode":**  
     Cloud gateway information.

**"cnfsNode":**  
     CNFS (Clustered NFS) node information.

**"cnfsGroupID": "GroupID"**  
     The failover recovery group for the node.

**"cnfsIpList": "IPList"**  
     A list of the CNFS IP addresses.

**"cnfsState": "{enabled | disabled}"**  
     The state of the CNFS node.

**"daemonNodeName": "NodeName"**  
     The node name that GPFS daemons use to communicate between nodes.

**"designation": "RoleList"**  
     A list of node roles. The following roles are possible:
 

- **manager**: The node is a manager node.
- **quorum**: The node is a quorum node.
- **quorumManager**: The node is a quorum node and a manager node.
- (blank): The node is not a manager node and not a quorum node.

**"ipAddress": "IPAddress"**  
     The IP address of the node.

**"links":**

**"node" : "URL"**  
     The URL of the node.

**"nodeNumber": "Number"**  
     The node number.

**"otherNodeRoles": "RoleList"**  
     Node roles other than manager-client or quorum-nonquorum, including **cesNode**, **cloudNodeMarker**, **cnfsNode**, **ctdbNode**, **gatewayNode**, **ioNode**, **perfmonNode**, **snmpAgent**, and **tealAgent**.

```

"status":
 Return status.

"message": "ReturnMessage",
 The return message.

"code": ReturnCode
 The return code.

```

## Examples

The following example gets information about the local cluster:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/cluster'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. In the nodes array, only the third node includes a **cesNode** structure, because it is a CES node:

```

:
{
 "cluster": {
 "cesSummary": {
 "EnabledServices": "None",
 "addressPolicy": "even-coverage",
 "cesSharedRoot": "/fs1/CES",
 "logLevel": 0
 },
 "clusterSummary": {
 "clusterId": 739880187759692000,
 "clusterName": "solo.division.company.com",
 "primaryServer": "host3.division.company.com",
 "rcpPath": "/usr/bin/scp",
 "rcpSudoWrapper": "no",
 "repositoryType": "CCR",
 "rshPath": "/usr/bin/ssh",
 "rshSudoWrapper": "no",
 "secondaryServer": "",
 "uidDomain": "solo.division.company.com"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/cluster"
 },
 "nodes": [
 {
 "adminLoginName": "",
 "adminNodeName": "host1",
 "daemonNodeName": "host1",
 "designation": "quorumManager",
 "ipAddress": "198.51.100.1",
 "links": {
 "node": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host1"
 },
 "nodeNumber": 4,
 "otherNodeRoles": ""
 },
 {
 "adminLoginName": "",
 "adminNodeName": "host2",
 "daemonNodeName": "host2",
 "designation": "quorumManager",
 "ipAddress": "198.51.100.10",
 "links": {
 "node": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host2"
 }
 }
]
 }
}

```

```

},
"nodeNumber": 5,
"otherNodeRoles": ""
},
{
"adminLoginName": "",
"adminNodeName": "host3.division.company.com",
"cesNode": {
 "cesGroup": "",
 "cesIpList": "",
 "cesState": "e",
 "ipAddress": "198.51.100.14"
},
"daemonNodeName": "host3.division.company.com",
"designation": "quorumManager",
"ipAddress": "198.51.100.14",
"links": {
 "node": [
 {
 "nodeNumber": 1,
 "otherNodeRoles": "cesNode"
 }
]
},
"status": {
 "code": 0,
 "message": ""
}
}
}

```

## See also

- “mmlscluster command” on page 397
- “mmchfs command” on page 179
- “mmlsfs command” on page 411

## Filesets: GET

Gets information about filesets.

## Availability

Available on all IBM Spectrum Scale editions.

## Description

The **GET filesets** request gets information about the filesets in the specified file system. For more information about the fields in the data structures that are returned, see the topics “[mmcrfileset command](#)” on page 244, “[mmchfileset command](#)” on page 173, and “[mmlsfileset command](#)” on page 406.

## Request URL

`https://REST_API_host:port/sclemgmt/v1/filesets  
?filesystemName=FileSystemName`

where:

## filesets

Is the filesets resource. Required.

**filesystemName**=*FileSystemName*

Specifies the name of the file system to which the filesets belong. Required.

## Request headers

Accept: application/json

## Request data

No request data.

## Response data

```
{ "filesets": [
 {
 "afm": {
 "afmAsyncDelay": Delay,
 "afmDirLookupRefreshInterval": Interval,
 "afmDirOpenRefreshInterval": Interval,
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": Interval,
 "afmFileLookupRefreshInterval": Interval,
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
 independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": Threads,
 "afmParallelReadChunkSize": Size,
 "afmParallelReadThreshold": Threshold,
 "afmParallelWriteChunkSize": Size,
 "afmParallelWriteThreshold": Threshold,
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": Interval,
 "afmShowHomeSnapshots": "{yes | no}",
 "afmTarget": "Protocol://{Host | Map}/{Path}"
 }
 },
 "config": {
 "comment": "Comment",
 "filesetName": "Fileset",
 "filesystemName": "Device",
 "iamMode": "Mode",
 "inodeSpace": Inodes,
 "maxFilesets": 1000,
 "maxFilesystems": 1000
 }
}
```

```

 "maxNumInodes": Inodes,
 "allocInodes": Inodes,
 "owner": "Owner",
 "path": "Path",
 "permissionChangeMode": "Mode",
 "permissions": "Permissions"
}
"links": {
 "self": "URI"
}
"state": {
 "afmState": "State",
 "created": "DateTime",
 "dataInKB": Data,
 "freeInodes": Inodes,
 "id": ID,
 "inodeSpaceMask": inodeSpace,
 "inodes": Inodes,
 "isInodeSpaceOwner": "isOwner",
 "parentId": "ID",
 "rootInode": Inode,
 "snapID": ID,
 "status": "Status"
}
},
"status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
}
}
}

"filesets":

```

An array of information about the filesets in the specified file system. Each array element describes one fileset and can contain the data structures **afm**, **config**, **links**, and **state**. For more information about the fields in these data structures, see the links at the end of this topic:

#### "afm":

Information about Active File Management (AFM).

##### "afmAsyncDelay": *Delay*

The time in seconds by which to delay write operations because of the lag in updating remote clusters.

##### "afmDirLookupRefreshInterval": *Interval*

The interval in seconds between data revalidations of directories caused by lookup operations such as **ls** and **stat**.

##### "afmDirOpenRefreshInterval": *Interval*

The interval in seconds between data revalidations of directories caused by lookup operations.

##### "afmEnableAutoEviction": "{yes | no}"

Enables or disables eviction on the fileset.

##### "afmExpirationTimeout": *Interval*

The timeout in seconds after which cached data is considered expired. Used with **afmDisconnectTimeout**, which is set with **mmchconfig**.

##### "afmFileLookupRefreshInterval": *Interval*

The interval in seconds between data revalidations of files caused by lookup operations such as **ls** and **stat**.

##### "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}"

The mode in which the cache operates.

**"afmNumFlushThreads": *numThreads***  
The number of threads used on each gateway to synchronize updates to the home cluster.

**"afmParallelReadChunkSize": *Size***  
The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes during parallel reads.

**"afmParallelReadThreshold": *Threshold***  
The minimum file size, in megabytes, at which to do parallel reads.

**"afmParallelWriteChunkSize": *Size***  
The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes during parallel writes.

**"afmParallelWriteThreshold": *Threshold***  
The minimum file size, in megabytes, at which to do parallel writes.

**"afmPrefetchThreshold": {1 | 1-99 | 100}**  
The percentage of file size that must be cached before the entire file is prefetched.

- 0** Enables full file prefetching.
- 1-99** The percentage of file size that must be cached before the entire file is prefetched.
- 100** Disables full file prefetching.

**"afmPrimaryID": "*ID***  
The unique primary ID of the primary filesset for asynchronous data replication.

**"afmRPO": *Interval***  
The recovery point objective (RPO) interval, in minutes, for a primary filesset.

**"afmShowHomeSnapshots": "{yes | no}"**  
Shows or hides the home snapshot directory in cache. Specify **yes** to show, **no** to hide.

**"afmTarget": *Interval*"Protocol://{**Host** | **Map**}/*Path*"**  
The home that is associated with the cache.

**"config":**  
Information about the filesset configuration.

**"comment": "Comment",**  
A comment that appears in the output of the **mmlsfilesset** command.

**"filessetName": "Fileset",**  
The name of the filesset.

**"filesystemName": "Device"**  
Required. The file system in which the filesset is located.

**"iamMode": "Mode"**  
The integrated archive manager (IAM) mode for the filesset.

- ad | advisory**
- nc | noncompliant**
- co | compliant**

**"inodeSpace": "Inodes"**  
The number of inodes that are allocated for use by the filesset.

**"maxNumInodes": "Inodes"**  
The inode limit for the inode space owned by the specified filesset.

**"allocInodes": "Inodes"**  
The number of inodes that are allocated for use by the filesset.

**"owner": "Owner"**  
     The owner of the fileset.

**"path": "Path"**  
     The absolute path of the fileset.

**"permissionChangeMode": "Mode"**  
     The permission change mode. Controls how **chmod** and ACL commands affect objects in the fileset.

- chmodOnly**  
         Only the **chmod** command can change access permissions.
- setAclOnly**  
         Only the ACL commands and API can change access permissions.
- chmodAndSetAcl**  
         Both the **chmod** command and ACL commands can change access permissions.
- chmodAndUpdateAcl**  
         Both the **chmod** command and ACL commands can change access permissions.

**"permissions": "Permissions"**  
     The file permissions of the fileset.

**"links":**  
     Links.

**"self": "URL"**  
     The URL of the resource.

**state:**  
     Information about the fileset state.

**"afmState": "State"**  
     The AFM status.

**"created": "DateTime",**  
     The date and time when the fileset was created.

**"dataInKB": "Data",**  
     The data size of the fileset.

**"freeInodes": *Inodes***  
     The number of allocated but unused inodes in the fileset.

**"id": *ID***  
     The fileset identifier.

**"inodeSpaceMask": *inodeSpace*,**  
     The inode space mask of the fileset.

**"inodes": *Inodes*,**

**"isInodeSpaceOwner": "isOwner",**  
     Indicates whether the fileset has its own inode space.

**"parentId": *ParentID***  
     The parent identifier of the fileset.

**"rootInode": "Inode"**  
     The number of the root inode of the fileset.

**"snapID": *ID***  
     The snapshot ID.

```

"status": "Status"
 Specifies whether the file is linked or unlinked.

"status":
 Return status.

"code": ReturnCode,
 The return message.

"message": "ReturnMessage"
 The return message.

```

## Examples

The following example gets information about the filesets in file system fs1:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/filesets?filesystemName=fs1'
```

In the JSON data that is returned, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The following **filesets** array contains three elements, each of which describes a fileset. The **afm** data structures are empty because AFM (Active File Management) is not enabled:

```
{
 "filesets": [
 {
 "afm": {},
 "config": {
 "comment": "",
 "filesetName": "fileset1",
 "filesystemName": "fs1",
 "inodeSpace": 1,
 "path": "--"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset1?filesystemName=fs1"
 },
 "state": {
 "created": "Tue Sep 6 08:58:09 2016",
 "freeInodes": 100031,
 "id": 1,
 "inodeSpaceMask": 24576,
 "isInodeSpaceOwner": 1,
 "root_inode": 524291,
 "snapId": 0,
 "status": "Unlinked"
 }
 },
 {
 "afm": {},
 "config": {
 "comment": "",
 "filesetName": "fileset02",
 "filesystemName": "fs1",
 "inodeSpace": 3,
 "path": "--"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset02?filesystemName=fs1"
 },
 "state": {
 "created": "Thu Sep 22 02:51:23 2016",
 "freeInodes": 100031,
 "id": 3,
 "inodeSpaceMask": 24576
 }
 }
]
}
```

```

 "inodeSpaceMask": 24576,
 "isInodeSpaceOwner": 1,
 "rootInode": 1572867,
 "snapId": 0,
 "status": "Unlinked"
 }
},
{
 "afm": {},
 "config": {
 "comment": "root fileset",
 "filesetName": "root",
 "filesystemName": "fs1",
 "inodeSpace": 0,
 "path": "/fs1"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/filesets/root?filesystemName=fs1"
 },
 "state": {
 "created": "Mon Aug 29 07:54:39 2016",
 "freeInodes": 454068,
 "id": 0,
 "inodeSpaceMask": 24576,
 "isInodeSpaceOwner": 1,
 "rootInode": 3,
 "snapId": 0,
 "status": "Linked"
 }
},
],
"status": {
 "code": 0,
 "message": ""
}
}

```

## See also

- “mmchfileset command” on page 173
- “mmcrfileset command” on page 244
- “mmlsfileset command” on page 406

## Filesets/{filesetName}: GET

Gets information about a fileset.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET filesets/{filesetName}** request gets information about the specified fileset. For more information about the fields in the data structures that are returned, see the topics “[mmcrfileset command](#)” on page 244, “[mmchfileset command](#)” on page 173, and “[mmlsfileset command](#)” on page 406.

### Request URL

```
https://REST_API_host:port/sclemgmt/v1/filesets/Fileset
?filesystemname=filesystemName
```

where:

#### filesets

Is the filesets resource. Required.

#### Fileset

Specifies the fileset that you want to get information about. Required.

#### filesystemName

Specifies the name of the file system to which the fileset belongs. Required.

### Request headers

Accept: application/json

### Request data

No request data.

### Response data

```
{ "filesets": [{ "afm": { "afmAsyncDelay": Delay, "afmDirLookupRefreshInterval": Interval, "afmDirOpenRefreshInterval": Interval, "afmEnableAutoEviction": "{yes | no}", "afmExpirationTimeout": Interval, "afmFileLookupRefreshInterval": Interval, "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}", "afmNumFlushThreads": Threads, "afmParallelReadChunkSize": Size, "afmParallelReadThreshold": Threshold, "afmParallelWriteChunkSize": Size, "afmParallelWriteThreshold": Threshold, "afmPrefetchThreshold": {1 | 1-99 | 100}, "afmPrimaryID": "ID", "afmRPO": Interval, "afmShowHomeSnapshots": "{yes | no}", "afmTarget": "Protocol://{Host | Map}/Path" }, "config": { "comment": "Comment", "filesetName": "Fileset", } }]}
```

```

 "filesystemName": "Device",
 "iamMode": "Mode",
 "inodeSpace": Inodes,
 "maxNumInodes": Inodes,
 "allocInodes": Inodes,
 "owner": "Owner",
 "path": "Path",
 "permissionChangeMode": "Mode",
 "permissions": "Permissions"
}
"links": {
 "self": "URI"
}
"state": {
 "afmState": "State",
 "created": "DateTime",
 "dataInKB": Data,
 "freeInodes": Inodes,
 "id": ID,
 "inodeSpaceMask": inodeSpace,
 "inodes": Inodes,
 "isInodeSpaceOwner": "isOwner",
 "parentId": "ID",
 "rootInode": Inode,
 "snapID": ID,
 "status": "Status"
}
},
"status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
}
}

```

## "filesets":

An array that contains elements that describe filesets. In this type of request, the array contains only one element, which describes the fileset that is specified in the HTTP request. Each array element contains the data structures **afm**, **config**, **links**, and **state**. For more information about the fields in these data structures, see the links at the end of this topic.

### "afm":

Information about Active File Management (AFM).

#### "afmAsyncDelay": *Delay*

The time in seconds by which to delay write operations because of the lag in updating remote clusters.

#### "afmDirLookupRefreshInterval": *Interval*

The interval in seconds between data revalidations of directories caused by lookup operations such as **ls** and **stat**.

#### "afmDirOpenRefreshInterval": *Interval*

The interval in seconds between data revalidations of directories caused by lookup operations.

#### "afmEnableAutoEviction": "{yes | no}"

Enables or disables eviction on the fileset.

#### "afmExpirationTimeout": *Interval*

The timeout in seconds after which cached data is considered expired. Used with **afmDisconnectTimeout**, which is set with **mmchconfig**.

#### "afmFileLookupRefreshInterval": *Interval*

The interval in seconds between data revalidations of files caused by lookup operations such as **ls** and **stat**.

**"afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}"**  
 The mode in which the cache operates.

**"afmNumFlushThreads": *numThreads***  
 The number of threads used on each gateway to synchronize updates to the home cluster.

**"afmParallelReadChunkSize": *Size***  
 The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes during parallel reads.

**"afmParallelReadThreshold": *Threshold***  
 The minimum file size, in megabytes, at which to do parallel reads.

**"afmParallelWriteChunkSize": *Size***  
 The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes during parallel writes.

**"afmParallelWriteThreshold": *Threshold***  
 The minimum file size, in megabytes, at which to do parallel writes.

**"afmPrefetchThreshold": {1 | 1-99 | 100}**  
 The percentage of file size that must be cached before the entire file is prefetched.
 

- 0** Enables full file prefetching.
- 1-99** The percentage of file size that must be cached before the entire file is prefetched.
- 100** Disables full file prefetching.

**"afmPrimaryID": "*ID***  
 The unique primary ID of the primary fileset for asynchronous data replication.

**"afmRPO": *Interval***  
 The recovery point objective (RPO) interval, in minutes, for a primary fileset.

**"afmShowHomeSnapshots": "{yes | no}"**  
 Shows or hides the home snapshot directory in cache. Specify **yes** to show, **no** to hide.

**"afmTarget": *Interval*"Protocol://{**Host** | **Map**}/**Path**"**  
 The home that is associated with the cache.

**"config":**  
 Information about the fileset configuration.

**"comment": "*Comment*",**  
 A comment that appears in the output of the **mmlsfileset** command.

**"filesetName": "*Fileset*",**  
 The name of the fileset.

**"filesystemName": "*Device*"**  
 Required. The file system in which the fileset is located.

**"iamMode": "*Mode*"**  
 The integrated archive manager (IAM) mode for the fileset.
 

- ad | advisory**
- nc | noncompliant**
- co | compliant**

**"inodeSpace": "*Inodes*"**  
 The number of inodes that are allocated for use by the fileset.

**"maxNumInodes": "Inodes"**  
     The inode limit for the inode space owned by the specified fileset.

**"allocInodes": "Inodes"**  
     The number of inodes that are allocated for use by the fileset.

**"owner": "Owner,"**  
     The owner of the fileset.

**"path": "Path"**  
     The absolute path of the fileset.

**"permissionChangeMode": "Mode"**  
     The permission change mode. Controls how **chmod** and ACL commands affect objects in the fileset.

- chmodOnly**  
         Only the **chmod** command can change access permissions.
- setAclOnly**  
         Only the ACL commands and API can change access permissions.
- chmodAndSetAcl**  
         Both the **chmod** command and ACL commands can change access permissions.
- chmodAndUpdateAcl**  
         Both the **chmod** command and ACL commands can change access permissions.

**"permissions": "Permissions"**  
     The file permissions of the fileset.

**"links":**  
     Links.

**"self": "URL"**  
     The URL of the resource.

**state:**  
     Information about the status of the fileset.

**"afmState": "State"**  
     The AFM status.

**"created": "DateTime",**  
     The date and time when the fileset was created.

**"dataInKB": "Data",**  
     The data size of the fileset.

**"freeInodes": "Inodes"**  
     The number of allocated but unused inodes in the fileset.

**"id": ID**  
     The fileset identifier.

**"inodeSpaceMask": *inodeSpace*,**  
     The inode space mask of the fileset.

**"inodes": *Inodes*,**

**"isInodeSpaceOwner": *isOwner*,**  
     Indicates whether the fileset has its own inode space.

**"parentId": *ParentID***  
     The parent identifier of the fileset.

**"rootInode": "Inode",**  
 The number of the root inode of the fileset.  
**"snapID": ID,**  
 The snapshot ID.  
**"status": "Status",**  
 Specifies whether the file is linked or unlinked.  
**"status":**  
 Return information.  
**"code": *ReturnCode*,**  
 The return message.  
**"message": "ReturnMessage"**  
 The return message.

## Examples

The following example gets information about the fileset `fileset1` in file system `fs1`:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset1?
filesystemName=fs1'
```

In the JSON data that is returned, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The `afm` data structure is empty because AFM (Active File Management) is not enabled for this fileset:

```
{
 "filesets": {
 "afm": {},
 "config": {
 "comment": "",
 "filesetName": "fileset1",
 "filesystemName": "fs1",
 "inodeSpace": 3,
 "path": "--"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset1?
 filesystemName=fs1"
 },
 "state": {
 "created": "Thu Sep 22 02:51:23 2016",
 "freeInodes": 100031,
 "id": 3,
 "inodeSpaceMask": 24576,
 "isInodeSpaceOwner": 1,
 "rootInode": 1572867,
 "snapId": 0,
 "status": "Unlinked"
 }
 },
 "status": {
 "code": 0,
 "message": ""
 }
}
```

## See also

- “`mmpchfileset` command” on page 173
- “`mmcrfileset` command” on page 244

- “`mmlsfileset` command” on page 406

## Filesets: POST

Creates a fileset.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **POST filesets** command creates the specified fileset. For more information about the fields in the request data structures, see the topics “[mmcrfileset command](#)” on page 244, “[mmchfileset command](#)” on page 173, and “[mmcrfs command](#)” on page 250.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/filesets`

where

#### filesets

Is the filesets resource.

### Request headers

`Content-Type: application/json`  
`Accept: text/html`

### Request data

Specify only the fields that you want to change. Required fields are marked with a superscript R. Corequisite fields are marked with a superscript 1. See the footnotes following the data:

```
{
 "afm":
 {
 "afmAsyncDelay": SecondsDelay,
 "afmDirLookupRefreshInterval": SecondsInterval,
 "afmDirOpenRefreshInterval": SecondsInterval,
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": SecondsInterval,
 "afmFileLookupRefreshInterval": SecondsInterval,
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
 independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": NumThreads,
 "afmParallelReadChunkSize": NumBytes,
 "afmParallelReadThreshold": NumMegabytes,
 "afmParallelWriteChunkSize": NumBytes,
 "afmParallelWriteThreshold": NumMegabytes,
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": MinutesInterval,
 "afmShowHomeSnapshots": "{yes | no}",
 "afmTarget": "Protocol://{Host | Map}/Path"
 },
 "config":R
 {
 "comment": "Comment",
 "filesetName": "Fileset",R
 "filesystemName": "Device",
 "iamMode": "{advisory | noncompliant | compliant}",
 "inodeSpace": "{new | ExistingFileset}",
 "maxNumInodes": "Inodes",
 "allocInodes": "Inodes",
 "owner": "UserID[:GroupID]",1
 }
}
```

```

 "path": "Path",1
 "permissionChangeMode": "chmodOnly | setAclOnly | chmodAndSetAcl",
 "permissions": "nnn"1
}
}

```

<sup>R</sup>Required structure or field.

<sup>1</sup>Optional, but if one of these three is specified, all must be specified: **owner**, **path**, and **permissions**. Use these parameters when you want to link the new filesset immediately.

## Request parameters

For more information about the fields in the request data structures, see the links at the end of this topic.

### "afm":

This structure is optional.

#### "afmAsyncDelay": *SecondsDelay*

Optional. The time in seconds by which to delay write operations because of the lag in updating remote clusters.

#### "afmDirLookupRefreshInterval": *SecondsInterval*

Optional. The interval in seconds between data revalidations of directories caused by lookup operations such as **ls** and **stat**.

#### "afmDirOpenRefreshInterval": *SecondsInterval*

Optional. The interval in seconds between data revalidations of directories caused by lookup operations.

#### "afmEnableAutoEviction": "{yes | no}"

Optional. Enables or disables eviction on the filesset.

#### "afmExpirationTimeout": *SecondsInterval*

Optional. The timeout in seconds after which cached data is considered expired. Used with **afmDisconnectTimeout**, which is set with **mmchconfig**.

#### "afmFileLookupRefreshInterval": *SecondsInterval*

Optional. The interval in seconds between data revalidations of files caused by lookup operations such as **ls** and **stat**.

#### "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}"

Optional. The mode in which the cache operates.

#### "afmNumFlushThreads": *NumThreads*

Optional. The number of threads used on each gateway to synchronize updates to the home cluster.

#### "afmParallelReadChunkSize": *NumBytes*

Optional. The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes during parallel reads.

#### "afmParallelReadThreshold": *NumMegabytes*

Optional. The minimum file size, in megabytes, at which to do parallel reads.

#### "afmParallelWriteChunkSize": *NumBytes*

Optional. The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes during parallel writes.

#### "afmParallelWriteThreshold": *NumMegabytes*

Optional. The minimum file size, in megabytes, at which to do parallel writes.

#### "afmPrefetchThreshold": {1 | 1-99 | 100}

Optional. The percentage of file size that must be cached before the entire file is prefetched.

0       Enables full file prefetching.  
 1-99    The percentage of file size that must be cached before the entire file is prefetched.  
 100     Disables full file prefetching.

**"afmPrimaryID": "ID"**  
 Optional. The unique primary ID of the primary fileset for asynchronous data replication.

**"afmRPO": "MinutesInterval"**  
 Optional. The recovery point objective (RPO) interval, in minutes, for a primary fileset.

**"afmShowHomeSnapshots": "{yes | no}"**  
 Optional. Shows or hides the home snapshot directory in cache. Specify **yes** to show, **no** to hide.

**"afmTarget": "Protocol://Host | Map}/Path"**  
 Optional. The home that is associated with the cache.

**"config":**  
 This structure is required. For more information about the fields in this structure, see the topic "mmcrfileset command" on page 244.

**"comment": "Comment"**  
 Optional. A comment that appears in the output of the **mmlsfileset** command. It must contain fewer than 256 characters.

**"filesetName": "Fileset"**  
 Required. The name of the newly created fileset.

**"filesystemName": "Device"**  
 Required. The file system that contains the new fileset.

**"iamMode": "{advisory | noncompliant | compliant}"**  
 Optional. The integrated archive manager (IAM) mode for the fileset.

**"inodeSpace": "{new | ExistingFileset}"**  
 Optional. Specifies the type of fileset to create, which controls how nodes are allocated.

**"maxNumInodes": "NumInodes"**  
 Optional. The maximum number of inodes for the inode space that the fileset owns.

**"allocInodes": "NumInodes"**  
 Optional. The number of additional inodes to preallocate for the inode space

**"owner": "UserID[:GroupID]"**  
 Optional, but if specified, you must also specify **path** and **permissions**. The user ID of the owner of the new fileset, followed optionally by the group ID, as in "root:root".

**"path": "Path"**  
 Optional, but if specified, you must also specify **owner** and **permissions**. The fully qualified path of the fileset, including the fileset name.

**"permissionChangeMode": "chmodOnly | setAc1Only | chmodAndSetAc1"**  
 Optional. Specifies that **chmod** and ACL operations are permitted.

**"permissions": "nnn"**  
 Optional, but if specified, you must also specify **owner** and **path**. File permissions for the fileset, such as "755".

## Response data

```
{
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

```

"status":
 Return information:
"code": ReturnCode,
 The return code for the request.
"message": ReturnMessage
 The return message.

```

## Examples

1. The following example creates and links a fileset `fileset1` in file system `fs1`:

```

curl -X POST --header 'Content-Type: application/json'
--header 'Accept: application/json' -d '{
 "config": {
 "filesetName": "fileset1",
 "filesystemName": "fs1",
 "owner": "root:root",
 "path": "/gpfs/fs1/fileset1",
 "permissions": "755"
 }
}' 'https://198.51.100.1:8191/scalemgmt/v1/filesets'

```

In the returned data, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 201, which indicates that the fileset was successfully created.

```
{
 "status": {
 "message": "Fileset fileset1 created with id 8 root inode 71053.
Fileset fileset1 linked at /gpfs/fs1/fileset1",
 "code": 0
 }
}
```

2. The following example illustrates how to specify AFM attributes when you create a fileset. The example creates and links a fileset `single-writer1` in file system `fs1`:

```

curl -X POST --header 'Content-Type: application/json'
--header 'Accept: application/json' -d '{
 "config": {
 "filesetName": "single-writer1",
 "inodeSpace": "new",
 "filesystemName": "fs1",
 "owner": "root:root",
 "path": "/fs1/single-writer1",
 "permissions": "755"
 }
 "afm": {
 "afmTarget": "c870mgrs2:/install",
 "afmMode": "single-writer",
 "afmAsyncDelay": 2147483647,
 "afmDirLookupRefreshInterval": 30,
 "afmFileLookupRefreshInterval": 180
 }
}' 'https://198.51.100.1:8191/scalemgmt/v1/filesets'

```

In the returned data, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 201, which indicates that the fileset was successfully created.

```
{
 "status": {
 "message": "Fileset single-writer1 created with id 9 root inode 71253.
Fileset single-writer1 linked at /gpfs/fs1/single-writer1",
 "code": 0
 }
}
```

## **See also**

- “[mmcrfileset command](#)” on page 244
- “[mmchfileset command](#)” on page 173
- “[mmcrfs command](#)” on page 250

## Filesets/{filesetName}: PUT

Changes the properties of a fileset.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The PUT `filesets/{filesetName}` command changes the properties of the specified fileset. For more information about the fields in the request data structures, see the topics “`mmchfileset` command” on page 173 and “`mmcrfileset` command” on page 244.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/filesets/FileName`

where

#### filesets

Is the filesets resource. Required.

#### FileName

Specifies the fileset whose properties you want to modify. Required.

### Request headers

`Content-Type: application/json`

`Accept: text/html`

### Request data

Specify only the fields that you want to change. To link a fileset, specify the `path`, `owner`, and `permissions`. If the fileset is already linked, it is unlinked and then relinked with the new options. To unlink a fileset, set the `path` to an empty string "".

Required structures and fields are marked with a superscript R. See the footnote following the data:

```
{
 "afm":
 {
 "afmAsyncDelay": SecondsDelay,
 "afmDirLookupRefreshInterval": SecondsInterval,
 "afmDirOpenRefreshInterval": SecondsInterval,
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": SecondsInterval,
 "afmFileLookupRefreshInterval": SecondsInterval,
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
 independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": NumThreads,
 "afmParallelReadChunkSize": NumBytes,
 "afmParallelReadThreshold": NumMegabytes,
 "afmParallelWriteChunkSize": NumBytes,
 "afmParallelWriteThreshold": NumMegabytes,
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": MinutesInterval,
 "afmShowHomeSnapshots": "{yes | no}",
 "afmTarget": "Protocol://{{Host | Map}}/Path"
 }
 "config":R
 {
 "comment": "Comment",
 }
```

```

 "filesetName": "Fileset",
 "filesystemName": "Device",R,
 "iamMode": "{advisory | noncompliant | compliant}",
 "inodeSpace": "{new | ExistingFileset}",
 "maxNumInodes": "Inodes",
 "allocInodes": "Inodes",
 "owner": "UserID[:GroupID]",
 "path": "Path",
 "permissionChangeMode": "chmodOnly | setAclOnly | chmodAndSetAcl",
 "permissions": "nnn"
}
}
RRequired structure or field.

```

## Request parameters

For more information about the fields in this structure, see the links at the end of this topic.

**"afm":**

Optional. Information about automatic file management.

**"afmAsyncDelay": *secondsDelay***

Optional. The time in seconds by which to delay write operations because of the lag in updating remote clusters.

**"afmDirLookupRefreshInterval": *secondsInterval***

Optional. The interval in seconds between data revalidations of directories caused by lookup operations such as **ls** and **stat**.

**"afmDirOpenRefreshInterval": *secondsInterval***

Optional. The interval in seconds between data revalidations of directories caused by lookup operations.

**"afmEnableAutoEviction": "{yes | no}"**

Optional. Enables or disables eviction on the fileset.

**"afmExpirationTimeout": *secondsInterval***

Optional. The timeout in seconds after which cached data is considered expired. Used with **afmDisconnectTimeout**, which is set with **mmchconfig**.

**"afmFileLookupRefreshInterval": *secondsInterval***

Optional. The interval in seconds between data revalidations of files caused by lookup operations such as **ls** and **stat**.

**"afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}"**

Optional. The mode in which the cache operates.

**"afmNumFlushThreads": *numThreads***

Optional. The number of threads used on each gateway to synchronize updates to the home cluster.

**"afmParallelReadChunkSize": *numBytes***

Optional. The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes during parallel reads.

**"afmParallelReadThreshold": *numMegabytes***

Optional. The minimum file size, in megabytes, at which to do parallel reads.

**"afmParallelWriteChunkSize": *numBytes***

Optional. The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes during parallel writes.

**"afmParallelWriteThreshold": *numMegabytes***

Optional. The minimum file size, in megabytes, at which to do parallel writes.

```

"afmPrefetchThreshold": {1 | 1-99 | 100}
 Optional. The percentage of file size that must be cached before the entire file is prefetched.
 0 Enables full file prefetching.
 1-99 The percentage of file size that must be cached before the entire file is prefetched.
 100 Disables full file prefetching.

"afmPrimaryID": "ID"
 Optional. The unique primary ID of the primary filesset for asynchronous data replication.

"afmRPO": minutesInterval
 Optional. The recovery point objective (RPO) interval, in minutes, for a primary filesset.

"afmShowHomeSnapshots": "{yes | no}"
 Optional. Shows or hides the home snapshot directory in cache. Specify yes to show, no to hide.

"afmTarget": minutesInterval"Protocol://{Host | Map}/Path"
 Optional. The home that is associated with the cache.

"config":
 Required. Configuration information.

"comment": "Comment"
 Optional. A comment that appears in the output of the mmlsfileset command. It must contain fewer than 256 characters.

"filesetName": "Fileset"
 Optional. Use this field to rename a filesset.

"filesystemName": "Device"
 Required. The file system that contains the new filesset.

"iamMode": "{advisory | noncompliant | compliant}"
 Optional. The integrated archive manager (IAM) mode for the filesset.

"inodeSpace": "{new | ExistingFileset}"
 Optional. Specifies the type of filesset.

"maxNumInodes": "numInodes"
 Optional. The maximum number of inodes for the inode space that the filesset owns.

"allocInodes": "numInodes"
 Optional. The number of additional inodes to preallocate for the inode space

"owner": "ownerID[:groupID]"
 Required. The user ID of the owner of the new filesset, followed optionally by the group ID, as in "root:root".

"path": "Path"
 Required. The fully qualified path of the filesset, including the filesset name.

"permissionChangeMode": "chmodOnly | setAc1Only | chmodAndSetAc1"
 Optional. Specifies that chmod and ACL operations are permitted.

"permissions": "nnn"
 Required. File permissions for the filesset, such as "755".

```

## Response data

```
{
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

```

"status":
 Return information:
"code": "ReturnCode,
 The return code for the request.
"message": "ReturnMessage"
 The return message.

```

## Examples

1. The following example changes a comment in fileset `fileset1` in file system `fs1`:

```

PUT --header 'Content-Type: application/json' --header 'Accept: application/json' -d
'{
 "config": {
 "comment": "This is a comment.",
 "filesetName": "fileset1",
 "filesystemName": "fs1"
 }
}' 'https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset1'

```

The return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully updated the properties of the fileset.

```
{
 "status": {
 "message": "Fileset fileset1 changed. ",
 "code": 0
 }
}
```

2. The following example illustrates how to change an AFM attribute in a fileset. The example changes AFM attributes in fileset `single-writer1` in file system `fs1`:

```

PUT --header 'Content-Type: application/json' --header 'Accept: application/json' -d
'{
 "config": {
 "filesystemName": "fs1"
 },
 "afm": {
 "afmAsyncDelay": 2147483,
 "afmDirLookupRefreshInterval": 60,
 "afmFileLookupRefreshInterval": 80
 }
}' 'https://198.51.100.1:8191/scalemgmt/v1/filesets/single-writer1'

```

The return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully updated the properties of the fileset.

```
{
 "status": {
 "message": "Fileset single-writer1 changed. ",
 "code": 0
 }
}
```

## See also

- “[mmchfileset command](#)” on page 173
- “[mmcrfileset command](#)” on page 244

## Filesets/{filesetName}: DELETE

Deletes a fileset.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **DELETE filesets/filesetName** command deletes the specified fileset. All files and folders under the fileset are deleted. For more information about deleting a fileset, see the topic “[mmdelfileset command](#)” on page 292.

### Request URL

```
https://REST_API_host:port/scalemgmt/v1/filesets/FileName?
filesystemname=FileSystemName&qos=other
```

where:

#### filesets

Is the filesets resource. Required.

#### filesystemname=*filesystemName*

Specifies the name of the file system that contains the fileset. Required.

#### qos=other

Specifies the QoS class of the maintenance command that deletes the fileset. Optional.

### Request headers

Accept: text/html

### Response data

```
{
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

#### status:

Return status.

##### "code": *ReturnCode*,

The return code for the request.

##### "message": "ReturnMessage"

The return message.

### Examples

The following example deletes a fileset fileset1 in file system fs1. The QoS class is set to other:

```
curl -X DELETE --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/filesets/fileset1?
filesystemName=fs1&qosClass=other'
```

In the returned JSON data, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully deleted the fileset.

```
"status": {
 "code": 0,
 "message": "Checking fileset ... Checking fileset complete. Deleting fileset
... Fileset fileset1 deleted"
}
```

## See also

- “mmdelfiles set command” on page 292

## Filesystems: GET

Gets information about file systems.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET filesystems** request gets information about file systems in the cluster. For more information about the fields in the data structures that are returned, see the topics “mmcrfs command” on page 250, “mmchfs command” on page 179, and “mmlsfs command” on page 411.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/filesystems`

where

#### filesystems

Is the filesystems resource.

### Request headers

Content-Type: application/json

Accept: application/json

### Request data

No request data.

### Response data

```
{
 "filesystems": [
 {
 "ACLSemantics": "{posix | afs4 | all}",
 "DMAPIEnabled": "{yes | no}",
 "UID": "ID",
 "additionalMountOptions": "Options",
 "automaticMountOption": "{yes | no | automount}",
 "blockAllocationType": "{scatter | cluster}",
 "blockSize": "Size",
 "createTime": "DateTime",
 "defaultDataReplicas": "Number",
 "defaultMetadataReplicas": "Number",
 "defaultMountPoint": "Path",
 "defaultQuotasEnabled": "[user] [;group] [;fileset]",
 "disks": "DiskList",
 "encryption": "{yes | no}",
 "exactMTIME": "{yes | no}",
 "fastEAEEnabled": "{yes | no}",
 "fileLockingSemantics": "{nfs4 | posix}",
 "filesetdfEnabled": "{yes | no}",
 "filesystemHighestSupported": "Version",
 "filesystemName": "FileSystem",
 "filesystemVersion": "Version",
 "filesystemVersionLocal": "Version",
 "filesystemVersionManager": "Version",
 "filesystemVersionOriginal": "Version",
 "indirectBlockSize": "Size",
 "inodeSize": "Size",
 "is4KAligned": "{yes | no}",
 "links": {
 "self": "URI",
 "children": "URI"
 }
 }
]
}
```

```

 "self": "URL"
 },
 "logReplicas": Number,
 "logfileSize": Size,
 "maxDataReplicas": Number,
 "maxMetadataReplicas": Number,
 "maxNumberOfInodes": Number,
 "maxSnapshotId": Number,
 "minFragmentSize": Size,
 "mountPriority": Priority,
 "numNodes": Number,
 "otherPools": [
 "blockSize": "Size",
 "minFragmentSize": "Size",
]
 "perFilesetQuotas": "{yes | no}",
 "quotasAccountingEnabled": "[user] [;group] [;fileset]"
 "quotasEnforced": "[user] [;group] [;fileset]"
 "rapidRepairEnabled": "{yes | no}",
 "storagePools": "Version",
 "strictReplication": "{no | whenpossible | always }",
 "suppressATime": "{yes | no}",
 "writeCacheThreshold": Number
}
],
"status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
},
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

#### **"filesystems":**

An array of elements that describe file systems. Each element describes one file system.

##### **"ACLSemantics": "{ posix | afs4 | all}"**

The type of authorization that the file system supports. ACL stands for "access control list".

##### **"DMAPISupported": "yes | no"**

Whether DMAP (Data Management API) is enabled on the file system.

##### **"UID": "ID"**

The UID of the file system.

##### **"additionalMountOptions": "Options"**

The mount options to pass to the mount command when the file system is being mounted.

##### **"automaticMountOption": "{yes | no | automount }"**

When the file system is to be mounted.

##### **"blockAllocationType": "{scatter | cluster}"**

The block allocation map type.

##### **"blockSize": Size**

The block size of the disks in the storage pool.

##### **"createTime": "DateTime"**

The date and time when the file system was created.

##### **"defaultDataReplicas": Number**

The default number of copies of each data block for a file.

##### **"defaultMetadataReplicas": Number**

The default number of copies of inodes, directories, and indirect blocks for a file.

**"defaultMountPoint": "Path"**  
     The default parent directory for IBM Spectrum Scale file systems.

**"defaultQuotasEnabled": "[user] [;group] [;fileset]"**  
     Which sets of default quotas are enabled.

**"disks": "DiskList"**  
     A semicolon-separated list of the disks that are included in the file system.

**"encryption": "{yes | no}"**  
     Whether encryption is enabled for the file system.

**"exactMTIME": "{yes | no}"**  
     Whether to report exact **mtime** values or to periodically update the **mtime** value for the file system.

**"fastEAEEnabled": "{yes | no}"**  
     Whether fast external attributes are enabled.

**"fileLockingSemantics": "{nfs4 | posix}"**  
     The type of file-locking semantics that are in effect.

**"filesetdfEnabled": "{yes | no}"**  
     Whether **filesetdf** is enabled.

**"filesystemHighestSupported": "Version"**  
     The highest file system version supported.

**"filesystemName": "FileSystem"**  
     The name of the file system.

**"filesystemVersion": "Version"**  
     The version of the file system.

**"filesystemVersionLocal": "Version"**  
     The version of the file system on the local cluster.

**"filesystemVersionManager": "Version"**  
     The version of the file system.

**"filesystemVersionOriginal": "Version"**  
     The version of the file system originally installed on the cluster.

**"indirectBlockSize": "Size"**  
     The indirect block size, in bytes.

**"inodeSize": "Size"**  
     The inode size, in bytes.

**"is4KAligned": "{yes | no}"**  
     Whether the file system is formatted to be 4K aligned.

**"links":**  
     Links.

**"self": "URL"**  
     The URL of the resource element.

**"logReplicas": "Number"**  
     The number of recovery log replicas.

**"logFileSize": "Size"**  
     The internal log file size.

**"maxDataReplicas": "Number"**  
     The maximum number of data replicas.

**"maxMetadataReplicas":*Number***  
The maximum number of metadata replicas.

**"maxNumberOfInodes":*Number***  
The maximum number of files in the file system.

**"maxSnapshotId":*Number***  
The maximum snapshot ID.

**"minFragmentSize":*Size***  
The minimum fragment size, in bytes.

**"mountPriority":*Priority***  
The mount priority of the file system.

**"numNodes":*Number***  
The number of nodes in the file system.

**"otherPools":**  
Lists the **blockSize** and **minFragmentSize** of pools that have a different **blockSize** and **minFragmentSize** than the system pool.

**"blockSize":*Size***  
The block size of the disks in the storage pool.

**"minFragmentSize":*Size***  
The minimum fragment size, in bytes.

**"perFilesetQuotas": "{yes | no}"**  
Whether the scope of user and group quota limit checks is the individual fileset level (**yes**) or the entire file system (**no**).

**"quotasAccountingEnabled": "[user] [;group] [;fileset]"**  
The types of quotas for which accounting is enabled.

**"quotasEnforced": "[user] [;group] [;fileset]"**  
The types of quotas that are enforced.

**"rapidRepairEnabled": "{yes | no}"**  
Whether the file system keeps track of incomplete replication on an individual file block basis (**yes**) or a whole file basis (**no**).

**"storagePools": "PoolList"**  
A list of the names of the storage pools that file system uses.

**"strictReplication": "{no | whenpossible | always }"**  
Whether strict replication is enforced.

**"suppressATime": "{yes | no}"**  
Whether the periodic updating of the value of **atime** is suppressed.

**"writeCacheThreshold":*Number***  
The maximum length in bytes of write requests that are initially buffered in the highly-available write cache before being written back to primary storage.

**"status":**  
Return status.

**"message": "ReturnMessage",**  
The return message.

**"code": *ReturnCode***  
The return code.

## Examples

The following example gets information about file systems in the cluster:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/filesystems'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The `filesystems` array contains one element, which describes the single file system in the cluster:

```
{
 "filesystems": [
 {
 "ACLSemantics": "all",
 "DMAPIEnabled": "no",
 "UID": "09724E2B:57C42283",
 "additionalMountOptions": "none",
 "automaticMountOption": "no",
 "blockAllocationType": "cluster",
 "blockSize": 262144,
 "create-time": "Mon Aug 29 07:54:43 2016",
 "defaultDataReplicas": 1,
 "defaultMetadataReplicas": 1,
 "defaultMountPoint": "/fs1",
 "defaultQuotasEnabled": "user;group;fileset",
 "disks": "DATA_host3_ib0_sdb;DATA_host3_ib0_sdd;DATA_host3_ib0_sde;
META_host3_ib0_sdc;DATA_host3_ib0_sdf",
 "encryption": "no",
 "exactMTIME": "yes",
 "fastEAEnabled": "yes",
 "fileLockingSemantics": "nfs4",
 "filesetdfEnabled": "no",
 "filesystemHighestSupported": "16.00 (4.2.2.0)",
 "filesystemName": "fs1",
 "filesystemVersion": "16.00 (4.2.2.0)",
 "filesystemVersionLocal": "16.00 (4.2.2.0)",
 "filesystemVersionManager": "16.00 (4.2.2.0)",
 "filesystemVersionOriginal": "16.00 (4.2.2.0)",
 "indirectBlockSize": 16384,
 "inodeSize": 4096,
 "is4KAligned": "yes",
 "links": {
 "self": "https://198.51.100.1/scalemgmt/v1/filesystems/fs1"
 },
 "logReplicas": 0,
 "logfileSize": 4194304,
 "maxDataReplicas": 3,
 "maxMetadataReplicas": 3,
 "maxNumberOfInodes": 4963712,
 "maxSnapshotId": 13,
 "minFragmentSize": 8192,
 "mountPriority": 0,
 "numNodes": 32,
 "perfilesetQuotas": "yes",
 "quotasAccountingEnabled": "user;group;fileset",
 "quotasEnforced": "user;group;fileset",
 "rapidRepairEnabled": "yes",
 "storagePools": "system",
 "strictReplication": "whenpossible",
 "suppressATime": "no",
 "write-cache-threshold": 0
 }
 "status": {
```

```
 "code": 0,
 "message": ""
 }
}
```

## See also

- “mmcrfs command” on page 250
- “mmchfs command” on page 179
- “mmlsfs command” on page 411

## Filesystems/{filesystemName}: GET

Gets information about a file system.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET filesystems** request gets information about the specified file system. For more information about the fields in the data structures that are returned, see the topics “mmcrfs command” on page 250, “mmchfs command” on page 179, and “mmlsfs command” on page 411.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/filesystems/FileSystem`

where

#### **filesystems**

Is the filesystems resource. Required.

#### **FileSystem**

The file system about which you want to get information. Required.

### Request headers

Accept: application/json

### Request data

No request data.

### Response data

```
{
 "filesystems": [
 {
 "ACLSemantics": "{posix | afs4 | all}",
 "DMAPIEnabled": "{yes | no}",
 "UID": "ID",
 "additionalMountOptions": "Options",
 "automaticMountOption": "{yes | no | automount}",
 "blockAllocationType": "{scatter | cluster}",
 "blockSize": "Size",
 "createTime": "DateTime",
 "defaultDataReplicas": "Number",
 "defaultMetadataReplicas": "Number",
 "defaultMountPoint": "Path",
 "defaultQuotasEnabled": "[user] [,group] [,fileset]",
 "disks": "DiskList",
 "encryption": "{yes | no}",
 "exactMTIME": "{yes | no}",
 "fastEAEnabled": "{yes | no}",
 "fileLockingSemantics": "{nfs4 | posix}",
 "filesetdfEnabled": "{yes | no}",
 "filesystemHighestSupported": "Version",
 "filesystemName": "FileSystem",
 "filesystemVersion": "Version",
 "filesystemVersionLocal": "Version",
 "filesystemVersionManager": "Version",
 "filesystemVersionOriginal": "Version",
 "indirectBlockSize": "Size",
 "inodeSize": "Size",
 "label": "Label",
 "mountPoint": "Path",
 "name": "Name",
 "node": "Node",
 "path": "Path",
 "quota": "Quota",
 "replica": "Replica",
 "size": "Size",
 "snapshot": "Snapshot",
 "storagePool": "StoragePool",
 "type": "Type",
 "version": "Version",
 "volume": "Volume"
 }
]
}
```

```

 "is4KAligned": "{yes | no}",
 "links": {
 "self": "URL"
 },
 "logReplicas": Number,
 "logfileSize": Size,
 "maxDataReplicas": Number,
 "maxMetadataReplicas": Number,
 "maxNumberofInodes": Number,
 "maxSnapshotId": Number,
 "minFragmentSize": Size,
 "mountPriority": Priority,
 "numNodes": Number,
 "otherPools": [
 "blockSize": "Size",
 "minFragmentSize": "Size",
],
 "perFilesetQuotas": "{yes | no}",
 "quotasAccountingEnabled": "[user] [;group] [;fileset]"
 "quotasEnforced": "[user] [;group] [;fileset]"
 "rapidRepairEnabled": "{yes | no}",
 "storagePools": "Version",
 "strictReplication": "{no | whenpossible | always }",
 "suppressATime": "{yes | no}",
 "writeCacheThreshold": Number
},
],
"status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
},
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"filesystems":**

An array of elements that describe file systems. Each element describes one file system.

**"ACLSemantics": "{ posix | afs4 | all }"**

The type of authorization that the file system supports. ACL stands for "access control list".

**"DMAPISupported": "yes | no"**

Whether DMAPI (Data Management API) is enabled on the file system.

**"UID": "ID"**

The UID of the file system.

**"additionalMountOptions": "Options"**

The mount options to pass to the mount command when the file system is being mounted.

**"automaticMountOption": "{yes | no | automount }"**

When the file system is to be mounted.

**"blockAllocationType": "{scatter | cluster}"**

The block allocation map type.

**"blockSize": Size**

The block size of the disks in the storage pool.

**"createTime": "DateTime"**

The date and time when the file system was created.

**"defaultDataReplicas": Number**

The default number of copies of each data block for a file.

**"defaultMetadataReplicas":*Number***  
The default number of copies of inodes, directories, and indirect blocks for a file.

**"defaultMountPoint":*Path***  
The default parent directory for IBM Spectrum Scale file systems.

**"defaultQuotasEnabled": "[*user*] [*group*] [*;fileset*]"**  
Which sets of default quotas are enabled.

**"disks": "*DiskList*"**  
A semicolon-separated list of the disks that are included in the file system.

**"encryption": "{*yes* | *no*}"**  
Whether encryption is enabled for the file system.

**"exactMTime": "{*yes* | *no*}"**  
Whether to report exact **mtime** values or to periodically update the **mtime** value for the file system.

**"fastEAEEnabled": "{*yes* | *no*}"**  
Whether fast external attributes are enabled.

**"fileLockingSemantics": "{*nfs4* | *posix*}"**  
The type of file-locking semantics that are in effect.

**"filesetdfEnabled": "{*yes* | *no*}"**  
Whether **filesetdf** is enabled.

**"filesystemHighestSupported": "*Version*"**  
The highest file system version supported.

**"filesystemName": "*FileSystem*"**  
The name of the file system.

**"filesystemVersion": "*Version*"**  
The version of the file system.

**"filesystemVersionLocal": "*Version*"**  
The version of the file system on the local cluster.

**"filesystemVersionManager": "*Version*"**  
The version of the file system.

**"filesystemVersionOriginal": "*Version*"**  
The version of the file system originally installed on the cluster.

**"indirectBlockSize":*Size***  
The indirect block size, in bytes.

**"inodeSize":*Size***  
The inode size, in bytes.

**"is4KAligned": "{*yes* | *no*}"**  
Whether the file system is formatted to be 4K aligned.

**"links":**  
Links.

**"self": "*URL*"**  
The URL of the resource element.

**"logReplicas":*Number***  
The number of recovery log replicas.

**"logfileSize":*Size***  
The internal log file size.

**"maxDataReplicas":*Number***  
The maximum number of data replicas.

**"maxMetadataReplicas":*Number***  
The maximum number of metadata replicas.

**"maxNumberOfInodes":*Number***  
The maximum number of files in the file system.

**"maxSnapshotId":*Number***  
The maximum snapshot ID.

**"minFragmentSize":*Size***  
The minimum fragment size, in bytes.

**"mountPriority":*Priority***  
The mount priority of the file system.

**"numNodes":*Number***  
The number of nodes in the file system.

**"otherPools":**  
Lists the **blockSize** and **minFragmentSize** of pools that have a different **blockSize** and **minFragmentSize** than the system pool.

**"blockSize":*Size***  
The block size of the disks in the storage pool.

**"minFragmentSize":*Size***  
The minimum fragment size, in bytes.

**"perFilesetQuotas": "{yes | no}"**  
Whether the scope of user and group quota limit checks is the individual fileset level (**yes**) or the entire file system (**no**).

**"quotasAccountingEnabled": "[user] [;group] [;fileset]"**  
The types of quotas for which accounting is enabled.

**"quotasEnforced": "[user] [;group] [;fileset]"**  
The types of quotas that are enforced.

**"rapidRepairEnabled": "{yes | no}"**  
Whether the file system keeps track of incomplete replication on an individual file block basis (**yes**) or a whole file basis (**no**).

**"storagePools": "PoolList"**  
A list of the names of the storage pools that file system uses.

**"strictReplication": "{no | whenpossible | always }"**  
Whether strict replication is enforced.

**"suppressATime": "{yes | no}"**  
Whether the periodic updating of the value of **atime** is suppressed.

**"writeCacheThreshold":*Number***  
The maximum length in bytes of write requests that are initially buffered in the highly-available write cache before being written back to primary storage.

**"status":**  
Return status.

**"message": "ReturnMessage",**  
The return message.

**"code": *ReturnCode***  
The return code.

## Examples

The following example gets information about file system `fs1`:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/filesystems/fs1'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The `filesystems` array contains one element, which describes file system `fs1`, the only file system in the cluster:

```
{
 "filesystems": [
 {
 "ACLSemantics": "all",
 "DMAPIEnabled": "no",
 "UID": "09724E2B:57C42283",
 "additionalMountOptions": "none",
 "automaticMountOption": "no",
 "blockAllocationType": "cluster",
 "blockSize": 262144,
 "createTime": "Mon Aug 29 07:54:43 2016",
 "defaultDataReplicas": 1,
 "defaultMetadataReplicas": 1,
 "defaultMountPoint": "/fs1",
 "defaultQuotasEnabled": "user;group;fileset",
 "disks": "DATA_host3_ib0_sdb;DATA_host3_ib0_sdd;DATA_host3_ib0_sde;
META_host3_ib0_sdc;DATA_host3_ib0_sdf",
 "encryption": "no",
 "exactMTIME": "yes",
 "fastEAEnabled": "yes",
 "fileLockingSemantics": "nfs4",
 "filesetdfEnabled": "no",
 "filesystemHighestSupported": "16.00 (4.2.2.0)",
 "filesystemName": "fs1",
 "filesystemVersion": "16.00 (4.2.2.0)",
 "filesystemVersionLocal": "16.00 (4.2.2.0)",
 "filesystemVersionManager": "16.00 (4.2.2.0)",
 "filesystemVersionOriginal": "16.00 (4.2.2.0)",
 "indirectBlockSize": 16384,
 "inodeSize": 4096,
 "is4KAligned": "yes",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/filesystems/fs1"
 },
 "logReplicas": 0,
 "logfileSize": 4194304,
 "maxDataReplicas": 3,
 "maxMetadataReplicas": 3,
 "maxNumberOfInodes": 4963712,
 "maxSnapshotId": 13,
 "minFragmentSize": 8192,
 "mountPriority": 0,
 "numNodes": 32,
 "perfilesetQuotas": "yes",
 "quotasAccountingEnabled": "user;group;fileset",
 "quotasEnforced": "user;group;fileset",
 "rapidRepairEnabled": "yes",
 "storagePools": "system",
 "strictReplication": "whenpossible",
 "suppressATime": "no",
 "writeCacheThreshold": 0
 }
 "status": {
```

```
 "code": 0,
 "message": ""
 }
}
```

## See also

- “mmcrfs command” on page 250
- “mmchfs command” on page 179
- “mmlsfs command” on page 411

## Info: GET

Gets information about the REST API server.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET info** request gets information about the REST API server. The information includes the version number, the resources that are supported, and the HTTP methods that are supported for each resource.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/info`

where

#### info

Is the information resource.

### Request headers

`Content-Type: application/json`

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "info": {
 "links": {
 "self": "URL"
 },
 "name": "APIName",
 "paths": {
 "/resourcePath": "MethodList",
 ["/resourcePath": "MethodList"]
 ...
 },
 "restVersion": "Version"
 }
 "status": {
 "message": "ReturnMessage",
 "code": "ReturnCode"
 }
}

"info":
 Information about the REST API server.

"links"
 Links.

 "self": "URL"
 The URL of the info resource.

 "name": "APIName"
 The name of the API, such as "Spectrum Scale REST Management".
```

```

"paths":
 Supported resources and elements.

"resourcePath": Methodlist

 resourcePath
 The URI of a REST API resource or element, such as /filesets/{filesetName}.

 MethodList
 A comma-separated list of HTTP functions that the resource or element supports, such as "GET, DELETE, POST".

"restVersion": "Version"
 The version of the REST API server.

"status":
 Return status.

"message": "ReturnMessage",
 The return message.

"code": "ReturnCode"
 The return code.

```

## Examples

The following example gets information about the REST API server:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/sclemgmt/v1/info'
```

In the JSON data that is returned, the return code is 0, indicating that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **paths** structure contains a list of the resources and elements that REST API supports and the HTTP methods that each resource or element supports:

```
{
 "info": {
 "links": {
 "self": "https://198.51.100.1:8191/sclemgmt/v1/info"
 },
 "name": "Spectrum Scale REST Management",
 "paths": {
 "/cesaddresses": [
 "GET"
],
 "/cesaddresses/{cesAddress)": [
 "GET"
],
 "/cesservices": [
 "GET"
],
 "/cesservices/{service)": [
 "GET"
],
 "/cluster": [
 "GET"
],
 "/config": [
 "GET"
],
 "/filesets": [
 "GET",
 "POST"
],
 "/filesets/{fileName)": [

```

```

 "DELETE",
 "GET",
 "PUT"
],
"/filesystems": [
 "GET"
],
"/filesystems/{filesystemName)": [
 "GET"
],
"/info": [
 "GET"
],
"/nodes": [
 "GET"
],
"/nodes/{name)": [
 "GET"
],
"/quotas/": [
 "GET",
 "POST"
],
"/snapshots": [
 "GET",
 "POST"
],
"/snapshots/{snapshotName)": [
 "DELETE",
 "GET"
]
},
"restVersion": "1.0.26"
},
"status": {
 "code": 0,
 "message": ""
}
}

```

## Nodes: GET

Gets information about nodes.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET nodes** request gets information about nodes in the cluster. For more information about the fields in the data structures that are returned, see the topics “[mmchnode command](#)” on page 190 and “[mmgetstate command](#)” on page 350.

### Request URL

`https://REST_API_host:port/sclemgmt/v1/nodes`

where

#### nodes

Is the nodes resource.

### Request headers

`Content-Type: application/json`

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "nodes": [
 {
 "adminLogin": "ID",
 "cesNode": "{true | false}",
 "clientNode": "{true | false}",
 "cnfsInterface": "IPAddressList",
 "cnfsState": "{active | inactive}",
 "daemonIPAddress": "IPAddress",
 "daemonInterface": "{HostName | IPAddress}",
 "designatedLicense": "LicenseType",
 "gatewayNode": "{true | false}",
 "gpfsState": "{active | arbitrating | down | unknown}",
 "links": {
 "self": "URL"
 },
 "managerNode": "{true | false}",
 "nodeName": "Node",
 "nodeNumber": Number,
 "osName": "OS",
 "productVersion": "Version",
 "quorumNode": "{true | false}",
 "snmpNode": "{true | false}"
 }
],
 "status": {
 "code": ReturnCode,
 "message": ReturnMessage
 },
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"nodes":**

An array of elements that describe nodes. Each element describes one node.

**"adminLogin": "ID"**

Specifies the login ID for GPFS administration commands.

**"cesNode": "{true | false}"**

Specifies whether the node is a CES node.

**"clientNode": "{true | false}"**

Specifies whether the node is a client node.

**"cnfsInterface": "IPAddressList"**

A comma-separated list of host names or IP addresses to be used for CNFS (clustered NFS) servers.

**"cnfsState": "{active | inactive}"**

Specifies whether CNFS is enabled on the node.

**"daemonIPAddress": "{IPAddress}"**

Specifies the IP address to be used by GPFS daemons for node-to-node communication.

**"daemonInterface": "{HostName | IPAddress}"**

Specifies the DNS name of the node.

**"designatedLicense": "LicenseType"**

Specifies the license type of the node.

**"gatewayNode": "{true | false}"**

Specifies whether the node is a gateway node.

**"gpfsState": "{active | arbitrating | down | unknown | }"**

Specifies the state of the GPFS daemon.

**"links":**

An array of links.

**"self": "URL"**

The URL of the node.

**"managerNode": "{true | false}"**

Specifies whether the node is a manager node.

**"nodeName": "Node"**

Specifies the name of the node.

**"nodeNumber": Number**

Specifies the number of the node.

**"osName": "OS"**

Specifies the name of the operating system that runs on the node.

**"productVersion": "Version"**

Specifies the product version.

**"quorumNode": "{true | false}"**

Specifies whether the node is a quorum node.

**"snmpNode": "{true | false}"**

Specifies whether the node is an SNMP (Simple Network Management Protocol) collector node.

**"status":**

Return status.

```
"message": "ReturnMessage",
```

The return message.

```
"code": "ReturnCode
```

The return code.

## Examples

The following example gets information about the nodes in the cluster:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/nodes'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The array contains three elements, each of which describes one of the nodes in the cluster:

```
{
 "nodes": [
 {
 "adminLogin": "root",
 "cesNode": "false",
 "clientNode": "false",
 "cnfsInterface": "",
 "cnfsState": "disabled",
 "daemonIPAddress": "198.51.100.10",
 "daemonInterface": "host3.division.company.com",
 "designatedLicense": "server",
 "gatewayNode": "false",
 "gpfsState": "active",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host3.division.company.com"
 },
 "managerNode": "true",
 "nodeName": "host3.division.company.com",
 "nodeNumber": 1,
 "osName": "Linux",
 "productVersion": "4.2.2.0",
 "quorumNode": "true",
 "snmpNode": "false"
 },
 {
 "adminLogin": "root",
 "cesNode": "false",
 "clientNode": "true",
 "cnfsInterface": "",
 "cnfsState": "disabled",
 "daemonIPAddress": "198.51.100.12",
 "daemonInterface": "host5",
 "designatedLicense": "server",
 "gatewayNode": "false",
 "gpfsState": "active",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host5"
 },
 "managerNode": "false",
 "nodeName": "host5",
 "nodeNumber": 4,
 "osName": "Linux",
 "productVersion": "4.2.2.0",
 "quorumNode": "false",
 "snmpNode": "false"
 },
 {
 "adminLogin": "root",
 "cesNode": "false",
 "clientNode": "true",
 "cnfsInterface": "",
 "cnfsState": "disabled",
 "daemonIPAddress": "198.51.100.13",
 "daemonInterface": "host6",
 "designatedLicense": "server",
 "gatewayNode": "false",
 "gpfsState": "active",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host6"
 },
 "managerNode": "false",
 "nodeName": "host6",
 "nodeNumber": 5,
 "osName": "Linux",
 "productVersion": "4.2.2.0",
 "quorumNode": "false",
 "snmpNode": "false"
 }
]
}
```

```

"clientNode": "true",
"cnfsInterface": "",
"cnfsState": "disabled",
"daemonIPAddress": "198.51.100.14",
"daemonInterface": "host6",
"designatedLicense": "server",
"gatewayNode": "false",
"gpfssState": "active",
"links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/nodes/host6"
},
"managerNode": "false",
"nodeName": "host6",
"nodeNumber": 5,
"osName": "Linux",
"productVersion": "4.2.2.0",
"quorumNode": "false",
"snmpNode": "false"
}
],
"status": {
 "code": 0,
 "message": ". "
}
}

```

## See also

- “mmchnode command” on page 190
- “mmgetstate command” on page 350

## **Nodes/{name}: GET**

Gets information about a node.

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

The **GET nodes/{nodeName}** request gets information about the specified node. For more information about the fields in the data structures that are returned, see the topics “mmchnode command” on page 190 and “mmgetstate command” on page 350.

### **Request URL**

`https://REST_API_host:port/sclemgmt/v1/nodes/NodeName`

where

#### **nodes**

Is the nodes resource. Required.

#### **NodeName**

Specifies the node about which you want to get information. Required.

### **Request headers**

`Content-Type: application/json`

`Accept: application/json`

### **Request data**

No request data.

### **Response data**

```
{
 "nodes": [
 {
 "adminLogin": "ID",
 "cesNode": "{true | false}",
 "clientNode": "{true | false}",
 "cnfsInterface": "IPAddressList",
 "cnfsState": "{active | inactive}",
 "daemonIPAddress": "IPAddress",
 "daemonInterface": "[HostName | IPAddress]",
 "designatedLicense": "LicenseType",
 "gatewayNode": "{true | false}",
 "gpfsState": "{active | arbitrating | down | unknown}",
 "links": {
 "self": "URL"
 },
 "managerNode": "{true | false}",
 "nodeName": "Node",
 "nodeNumber": "Number",
 "osName": "OS",
 "productVersion": "Version",
 "quorumNode": "{true | false}",
 "snmpNode": "{true | false}"
 }
],
 "status": {
```

```

 "code":ReturnCode,
 "message":ReturnMessage
},
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"nodes":**

An array of elements that describe nodes. Each element describes one node.

**"adminLogin": "ID"**

Specifies the login ID for GPFS administration commands.

**"cesNode": "{true | false}"**

Specifies whether the node is a CES node.

**"clientNode": "{true | false}"**

Specifies whether the node is a client node.

**"cnfsInterface": "IPAddressList"**

A comma-separated list of host names or IP addresses to be used for CNFS (clustered NFS) servers.

**"cnfsState": "{active | inactive}"**

Specifies whether CNFS is enabled on the node.

**"daemonIPAddress": "{IPAddress}"**

Specifies the IP address to be used by GPFS daemons when they are communicating between nodes.

**"daemonInterface": "{HostName | IPAddress}"**

Specifies the DNS name of the node.

**"designatedLicense": "LicenseType"**

Specifies the license type of the node.

**"gatewayNode": "{true | false}"**

Specifies whether the node is a gateway node.

**"gpfsState": "{active | arbitrating | down | unknown | }"**

Specifies the state of the GPFS daemon.

**"links":**

An array of links.

**"self": "URL"**

The URL of the node that is specified in the request.

**"managerNode": "{true | false}"**

Specifies whether the node is a manager node.

**"nodeName": "Node"**

Specifies the name of the node.

**"nodeNumber": *Number***

Specifies the number of the node.

**"osName": "OS"**

Specifies the name of the operating system that runs on the node.

**"productVersion": "Version"**

Specifies the product version.

**"quorumNode": "{true | false}"**

Specifies whether the node is a quorum node.

```

"snmpNode": "{true | false}"
 Specifies whether the node is an SNMP (Simple Network Management Protocol) collector node.

"status":
 Return status.

"message": "ReturnMessage",
 The return message.

"code": ReturnCode
 The return code.

```

## Examples

The following example gets information about the node node1:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/nodes/node1'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The **nodes** array contains one element, which describes the specified node:

```
{
 "nodes": [
 {
 "adminLogin": "root",
 "cesNode": "false",
 "clientNode": "false",
 "cnfsInterface": "",
 "cnfsState": "disabled",
 "daemonIPAddress": "198.51.100.11",
 "daemonInterface": "host3.division.company.com",
 "designatedLicense": "server",
 "gatewayNode": "false",
 "gpfsState": "active",
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/nodes/node1"
 },
 "managerNode": "true",
 "nodeName": "node1",
 "nodeNumber": 1,
 "osName": "Linux",
 "productVersion": "4.2.2.0",
 "quorumNode": "true",
 "snmpNode": "false"
 }
],
 "status": {
 "code": 0,
 "message": ". "
 }
}
```

## See also

- “mmchnode command” on page 190
- “mmgetstate command” on page 350

## Quotas: GET

Gets information about a quota.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET quotas** request gets information about the quotas in the current cluster. For more information about the fields in the data structures that are returned, see the topics “[mmlsquota command](#)” on page 435, “[mmrepquota command](#)” on page 517, and “[gpfs\\_quotaInfo\\_t structure](#)” on page 789.

### Request URL

```
https://REST_API_host:port/scalemgmt/v1/quotas
?queryType=Type"aType=Qtype&filesystemName=fileSystem
&filesetName=Fileset
```

where:

**Note:** If you omit all parameters, the request gets information about all the quotas and non-default quota limits in the cluster, including both global quotas and fileset quotas.

#### quotas

Is the name of the quotas resource.

#### queryType=Type

Specifies whether the request is for information about quota defaults or for information about non-default quota limits:

- To get information about quota defaults, set **queryType=defaults**.
- To get information about non-default quota limits, omit the **queryType** parameter or set **queryType** to **limits**.

Optional.

#### quotaType=Qtype

Specifies the type of quota to get information about:

- If the request is for information about non-default quota limits, set **quotaType** to **user**, **group**, **fileset**, or **all**.
- If the request is for quota defaults, choose one of the following actions:
  - If the request is for the cluster or for a file system, set **quotaType** to **user**, **group**, **fileset**, or **all**.
  - If the request is for a specified fileset, set **quotaType=all**.

Optional.

#### filesystemName=fileSystem

Specifies the file system for which the quotas are set. Optional.

#### filesetName=Fileset

Specifies the fileset for which quotas are set. Optional.

### Request headers

Accept: application/json

### Request data

No request data.

## Response data

```
{
 "links": {
 "self": "URI"
 }
 "quotas": [
 {
 "blockGrace": "Period",
 "blockInDoubt": Number,
 "blockLimit": Number,
 "blockQuota": Number,
 "blockUsage": Number,
 "defQuota": "{on | off}",
 "filesGrace": "Period",
 "filesInDoubt": Number,
 "filesLimit": Number,
 "filesQuota": Number,
 "filesUsage": Number,
 "filesetId": ID,
 "filesetName": "FilesetName",
 "filesystemName": "Device",
 "objectId": ID,
 "objectName": "Name",
 "quotaType": "{USR | GRP | FILESET}",
 }
],
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

## Response parameters

### "links":

Links:

### "self": "URL"

The URL of the resource.

### "quotas":

An array of information about the quotas in the specified node or fileset. Each array element describes one quota limit or quota default. For more information about the fields in this data structure, see the links at the end of this topic.

#### "blockGrace": "Period"

The time limit for excessive disk use. The limit is expressed in seconds since the Epoch.

#### "blockInDoubt": Number

The distributed shares and block usage that are not accounted for, in kilobytes.

#### "blockLimit": Number

The absolute limit (hard limit) on disk blocks allocated, in kilobytes.

#### "blockQuota": Number

The preferred limit (soft limit) on disk blocks allocated, in kilobytes.

#### "blockUsage": Number

The current block count, in kilobytes.

#### "defQuota": "{on | off}"

Specifies whether default quota limits are activated ("on") or deactivated ("off").

#### "filesGrace": "Period"

The time limit for excessive file use. The limit is expressed in seconds since the Epoch.

**"filesInDoubt": *Number***  
The distributed file share and file usage that are not accounted for.

**"filesLimit": *Number***  
The absolute limit (hard limit) on allocated files.

**"filesQuota": *Number***  
The preferred limit (soft limit) on allocated files.

**"filesUsage": *Number***  
The current number of allocated files.

**"filesetId": *ID***  
The ID of the fileset for which quotas are set.

**"filesetName": "*FilesetName*"**  
The fileset for which the quotas are set.

**"filesystemName": "*Device*"**  
The file system for which quotas are set.

**"objectId": *ID***  
The ID of the user, group, or fileset for which quotas are set.

**"objectName": "*Name*"**  
The name of the user, group, or fileset for which quotas are set.

**"quotaType": "{**USR** | **GRP** | **FILESET**}"**  
The type of quota about which information is returned

**"status":**  
Return status.

**"code": *ReturnCode*,**  
The return message.

**"message": "*ReturnMessage*"**  
The return message.

## Examples

The following example changes gets information about the quota limits in file system fs1:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/quotas
?queryType=limits"aType=all&filesystemName=fs1'
```

In the JSON data that is returned, the return code is 0, which indicates that the command is successful. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The following example shows only part of the output:

```
{
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/quotas?queryType=limits"aType=all
&filesystemName=fs1"
 },
 "quotas": [
 {
 "blockGrace": "none",
 "blockInDoubt": 0,
 "blockLimit": 31457280,
 "blockQuota": 26214400,
 "blockUsage": 0,
 "defQuota": "on",
 "filesGrace": "none",
 "filesInDoubt": 0,
```

```

 "filesLimit": 10485760,
 "filesQuota": 1048576,
 "filesUsage": 0,
 "filesetId": 1,
 "filesetName": "ifset1",
 "filesystemName": "fs1",
 "objectId": 6004,
 "objectName": 6004,
 "quota": "on",
 "quotaType": "USR"
 },
 {
 "blockGrace": "none",
 "blockInDoubt": 0,
 "blockLimit": 41943040,
 "blockQuota": 26214400,
 "blockUsage": 0,
 "defQuota": "on",
 "filesGrace": "none",
 "filesInDoubt": 0,
 "filesLimit": 10485760,
 "filesQuota": 1048576,
 "filesUsage": 1,
 "filesetId": "",
 "filesetName": "",
 "filesystemName": "fs1",
 "objectId": 1,
 "objectName": "ifset1",
 "quota": "on",
 "quotaType": "FILESET"
 },
 {
 "blockGrace": "none",
 "blockInDoubt": 0,
 "blockLimit": 62914560,
 "blockQuota": 26214400,
 "blockUsage": 0,
 "defQuota": "on",
 "filesGrace": "none",
 "filesInDoubt": 0,
 "filesLimit": 10485760,
 "filesQuota": 1048576,
 "filesUsage": 0,
 "filesetId": 0,
 "filesetName": "root",
 "filesystemName": "fs1",
 "objectId": 1001,
 "objectName": "suser1",
 "quota": "on",
 "quotaType": "GRP"
 }
],
"status": {
 "code": 0,
 "message": ""
}
}

```

## See also

- “mmlsquota command” on page 435
- “gpfs\_quotaInfo\_t structure” on page 789

## Quotas: POST

Sets quota limits, quota defaults, or grace periods.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **POST quotas** request sets quota limits, quota defaults, or quota grace periods in the current cluster. For more information about the fields in the data structure that is sent with the request, see the topics “mmsetquota command” on page 548, “mmrepquota command” on page 517, and “mmlsquota command” on page 435.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/quotas`

where

#### quotas

Is the name of the quotas resource.

### Request headers

Content-Type: application/json

Accept: application/json

### Request data

```
{
 "blockGracePeriod": "Period",
 "blockHardLimit": Number,
 "blockSoftLimit": Number,
 "filesGracePeriod": "Period",
 "filesHardLimit": Number,
 "filesSoftLimit": Number,
 "filesetName": "FilesetName",
 "filesystemName": "Device",
 "objectName": "Name",
 "operationType": "{setQuota | setDefaultQuota | setGracePeriod}",
 "quotaType": "{user | group | fileset}"
}
```

#### "blockGracePeriod": "Period"

The time limit for excessive disk use. This parameter specifies the grace period during which block use can exceed the soft limit before it is imposed as a hard limit. Specify in days, hours, minutes, or seconds.

**Note:** The grace period cannot be set together with quota limits in the same request.

#### "blockHardLimit": Number

The absolute limit on disk blocks allocated. This parameter specifies the number of disk blocks that the system can use during a grace period. If omitted, defaults to no limit. Specify K (default), M, G, or T.

#### "blockSoft": Number

The preferred limit on disk blocks allocated. Specify K (default), M, G, or T.

#### "filesGracePeriod": "Period"

The time limit for excessive file use. This parameter specifies the grace period during which the number of files can exceed the soft limit before it is imposed as a hard limit. Specify in days, hours, minutes, or seconds.

**"filesHardLimit": Number**

The absolute limit on allocated files. This parameter specifies the number of files that the system can use during a grace period. If omitted, defaults to no limit. Specify units (default), K, M, or G.

**"filesSoftLimit": Number**

The preferred limit on allocated files. Specify units (default), K, M, or G.

**"filesetName": "FilesetName"**

The fileset that is the target of the operation.

**"filesystemName": "Device"**

The file system that is the target of the operation.

**"objectName": "Name"**

The name of the user, group, or fileset that is the target of the operation.

**"operationType": "{setQuota | setDefaultQuota | setGracePeriod}"**

The operation to be performed.

**"quotaType": "{user | group | fileset}"**

The type of quota that is the target of the operation.

## Response data

```
{
 "status":
 {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

**"status":**

Return status.

**"code": *ReturnCode*,**

The return code for the operation.

**"returnMessage": "MessageSnapshot"**

The return message.

## Examples

1. The following example sets a quota for the fileset fileset1:

```
curl -X POST --header 'Content-Type: application/json' --header 'Accept: application/json' -d '{
 "blockHardLimit": "30G",
 "blockSoftLimit": "25G",
 "filesetName": "fileset1",
 "filesystemName": "fs1",
 "operationType": "setQuota",
 "quotaType": "fileset"
}' 'https://198.51.100.1:8191/scalemgmt/v1/quotas/'
```

In the JSON data that is returned, the return code is 0, which indicates that the command is successful. The response code, not shown, is 201, which indicates that the command successfully created the quota:

```
{
 "status": {
 "code": 0,
 "message": "Command successful"
 }
}
```

2. The following example sets the block grace period and file grace period for users on file system fs1:

```
curl -k -X POST --header 'Content-Type: application/json' --header 'Accept: application/json' -d '{
 "filesGracePeriod": "30days",
 "blockGracePeriod": "25days",
 "filesystemName": "fs1",
 "operationType": "setGracePeriod",
 "quotaType": "user"
' 'https://198.51.100.4:8191/scalemgmt/v1/quotas/'
```

In the JSON data that is returned, the return code is 0, which indicates that the command is successful. The response code, not shown, is 201, which indicates that the command successfully set the grace periods:

```
{
 "status": {
 "code": 0,
 "message": "Command successful"
 }
}
```

## See also

- “mmlsquota command” on page 435
- “gpfs\_quotaInfo\_t structure” on page 789

## Snapshots: GET

Gets information about snapshots.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET snapshots** request gets information about snapshots in the specified file system. For more information about the fields in the data structures that are returned, see the topics “[mmcrsnapshot command](#)” on page 267 and “[mmlsnapshot command](#)” on page 439.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/snapshots/?filesystemName=Device&filesetsFilter=Filter`

where

#### snapshots

Is the snapshots resource. Required.

#### filesystemName=Device

Specifies the name of the file system to which the snapshot belongs. Required.

#### filesetsFilter=Filter

Specifies a comma-separated list of filesets, which can include global filesets. The request gets information about all snapshots that contain at least one of the filesets. If this parameter is omitted, the request gets information about all the snapshots in the file system. Optional.

### Request headers

Content-Type: application/json  
Accept: application/json

### Request data

No request data.

### Response data

```
{
 "snapshots": [
 {
 "config": {
 "filesetName": "Fileset",
 "filesystemName": "Device",
 "snapshotName": "Snapshot"
 },
 "links": {
 "self": "URL"
 },
 "status": {
 "created": "DateTime",
 "quotas": "Quotas",
 "snapID": "ID",
 "snapType": "Type",
 "status": "Status"
 }
 }
],
 "status": {
```

```
 "code":ReturnCode,
 "message":"ReturnMessage"
 },
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"snapshots":**

An array of elements that describe snapshots, with one array element for each snapshot. The array element contains three data structures: **config**, **links**, and **status**.

**"config":**

The configuration of the snapshot.

**"filesetName":*"Fileset"***

For a fileset snapshot, the fileset that is a target of the snapshot.

**"filesystemName":*"Device"***

The file system that is the target of the snapshot.

**"snapshotName":*"Snapshot"***

The snapshot name.

**"links":**

Links.

**"self":*"URL"***

The URL of the resource.

**"status":**

The status of the snapshot.

**"created":*"DateTime"***

The date and time when the snapshot was created.

**"quotas":*"Quotas"***

Any quotas that are applied to the fileset.

**"snapID":*"ID"***

The snapshot ID.

**"snapType":*"Type"***

The AFM type of the snapshot, including *"afm\_snap"*, *"afm\_recovery"*, *"afm\_failover"*, *"afm\_rpo"*, *"afm\_baserpo"*, and *"Invalid"*.

**"status":*"Status"***

The snapshot status.

**"status":**

Return status.

**"message": "ReturnMessage",**

The return message.

**"code": *ReturnCode***

The return code.

## Examples

The following example gets information about the snapshots in filesystem fs1:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/snapshots?filesystemName=fs1'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. In this example, the output array contains one element, which describes the single snapshot in the file system:

```
{
 "snapshots": [
 {
 "config": {
 "filesetName": "",
 "filesystemName": "fs1",
 "snapshotName": "snapshot1"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/snapshots/mysnap1?filesystemName=fs1"
 },
 "status": {
 "created": "Tue Sep 20 21:46:35 2016",
 "quotas": "",
 "snapID": 3,
 "snapType": "",
 "status": "Valid"
 }
 }
],
 "status": {
 "code": 0,
 "message": ""
 }
}
```

## See also

- “mmcrsnapshot command” on page 267
- “mmlssnapshot command” on page 439

## Snapshots/{snapshotName}: GET

Gets information about a snapshot.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET snapshots/{snapshotName}** request gets information about the specified snapshot. For more information about the fields in the data structures that are returned, see the topics “mmcrsnapshot command” on page 267 and “mmlssnapshot command” on page 439.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/snapshots/Snapshot?  
filesystemName=Device&filesetName=Fileset`

where

#### snapshots

Is the snapshots resource. Required.

#### Snapshot

Specifies the snapshot about which you want to get information. Required.

#### filesystemName=Device

Specifies the file system to which the snapshot belongs. Required.

#### filesetName=Fileset

For a fileset snapshot, specifies the fileset name. Optional

### Request headers

Content-Type: application/json

Accept: application/json

### Request data

No request data.

### Response data

```
{
 "snapshots": [
 {
 "config": {
 "filesetName": "Fileset",
 "filesystemName": "Device",
 "snapshotName": "Snapshot"
 }
 "links": {
 "self": "URL"
 }
 "status": {
 "created": "DateTime",
 "quotas": "Quotas",
 "snapID": "ID",
 "snapType": "Type",
 "status": "Status"
 }
 }
]
 "status": {
```

```
 "code":ReturnCode,
 "message":ReturnMessage
 },
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"snapshots":**

An array that contains a single element that describes the specified snapshot. The array element contains three data structures: **config**, **links**, and **status**.

**"config":**

A data structure that describes the configuration of the snapshot.

**"filesetName": "Fileset"**

For a fileset snapshot, the fileset that is a target of the snapshot.

**"filesystemName": "Device"**

The file system that contains the snapshot.

**"snapshotName": "Snapshot"**

The snapshot name.

**"links":**

An array of links.

**"self": "URL"**

The URL of the resource element.

**"status":**

**"created": "DateTime"**

The date and time when the snapshot was created.

**"quotas": "Quotas"**

Any quotas that are applied to the fileset.

**"snapID": "ID"**

The snapshot ID.

**"snapType": "Type"**

The AFM type of the snapshot, including "afm\_snap", "afm\_recovery", "afm\_failover", "afm\_rpo", "afm\_baserpo", and "Invalid".

**"status": "Status"**

The snapshot status.

**"status":**

Return status.

**"message": "ReturnMessage",**

The return message.

**"code": *ReturnCode***

The return code.

## Examples

The following example gets information about snapshot snapshot1 in file system fs1:

```
curl -X GET --header 'Accept: application/json'
'https://198.51.100.1:8191/scalemgmt/v1/snapshots/snapshot1?filesystemName=fs1'
```

In the JSON data that is returned, the return code is 0, which indicates success. The response code, not shown, is 200, which indicates that the command successfully retrieved the information. The array contains one element, which describes the snapshot:

```
{
 "snapshots": [
 {
 "config": {
 "filesetName": "",
 "filesystemName": "fs1",
 "snapshotName": "snapshot1"
 },
 "links": {
 "self": "https://198.51.100.1:8191/scalemgmt/v1/snapshots/snapshot1?filesystemName=fs1"
 },
 "status": {
 "created": "Tue Sep 20 21:46:35 2016",
 "quotas": "",
 "snapID": 3,
 "snapType": "",
 "status": "Valid"
 }
 }
,
 "status": {
 "code": 0,
 "message": ""
 }
 }
}
```

## See also

- “mmcrsnapshot command” on page 267
- “mmlssnapshot command” on page 439

## Snapshots: POST

Creates a snapshot.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The POST snapshots command creates a snapshot of the specified fileset or file system.

### Request URL

`https://REST_API_host:port/scalemgmt/v1/snapshots`

where

#### snapshots

Is the snapshots resource.

### Request headers

Content-Type: application/json  
Accept: text/html

### Request data

```
{
 "config": {
 "filesetName": "Fileset",
 "filesystemName": "FileSystem",
 "snapshotName": "Snapshot"
 }
}

"config":
 A data structure that contains input parameters for creating the snapshot.

 "filesetName": "Fileset"
 The name of the fileset that is the target of the snapshot. Required for a fileset snapshot. If this parameter is omitted then the request creates a global snapshot.

 "filesystemName": "FileSystem"
 The name of the file system for which the snapshot is created. Required.

 "snapshotName": "Snapshot"
 The name of the new snapshot. Required.
```

### Response data

```
{
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}

status:
 Return status.

 "code": ReturnCode,
 The return code for the request.

 "message": "ReturnMessage"
 The return message.
```

## Examples

The following example creates a snapshot `snapshot1` in file system `fs1`. Because the `filesetName` field is included, the command takes a snapshot of the specified filesset.

```
curl -X POST --header 'Content-Type: application/json' --header 'Accept: application/json' -d '{
 "config": {
 "filesetName": "filesset1",
 "filesystemName": "fs1",
 "snapshotName": "snapshot1"
 }
' 'https://198.51.100.1:8191/scalemgmt/v1/snapshots'
```

In the returned JSON data, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 201, which indicates that the command successfully created the snapshot:

```
{
 "status": {
 "code": 0,
 "message": "Flushing dirty data for snapshot filesset1:snapshot1...
Quiescing all file system operations.
Snapshot filesset1:snapshot1 created with id 18."
 }
}
```

## See also

- “`mmcrsnapshot` command” on page 267

## Snapshots/{snapshotName}: DELETE

Deletes a snapshot.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **DELETE snapshots/snapshotName** command deletes the specified snapshot. For more information on deleting snapshots, see the topic “[mmdelsnapshot command](#)” on page 304.

### Request URL

Use this URL to delete a global snapshot:

```
https://REST_API_host:port/sclemgmt/v1/snapshots/SnapshotName
?filesystemName=FileSystemName
```

Use this URL to delete a fileset snapshot:

```
https://REST_API_host:port/sclemgmt/v1/snapshots/SnapshotName
?filesystemName=FileSystemName&filesetName=FilesetName
```

where:

#### snapshots

Is the snapshots resource. Required.

#### filesystemName=FileSystemName

Specifies the name of the file system for which the snapshot was created. Required.

#### filesetName=FilesetName

Specifies the name of the fileset that is the target of the snapshot. Required only for a fileset snapshot.

### Response data

```
{
 "status": {
 "code": ReturnCode,
 "message": "ReturnMessage"
 }
}
```

#### status:

Return status.

##### "code": *ReturnCode*,

The return code for the request.

##### "message": "ReturnMessage"

The return message.

### Examples

The following example deletes a fileset snapshot fileset1:snapshot1 in file system fs1:

```
curl -X DELETE --header 'Accept: application/json'
'https://198.51.100.1:8191/sclemgmt/v1/snapshots/snapshot1?
filesystemName=fs1&filesetName=fileset1'
```

In the returned JSON data, the return code is 0, which indicates that IBM Spectrum Scale processed the request successfully. The response code, not shown, is 200, which indicates that the command successfully deleted the snapshot:

```
"status": {
 "code": 0,
 "message": "Invalidating snapshot files in fileset1:snapshot1...
Deleting files in snapshot fileset1:snapshot1...
100.00 % complete on Mon Sep 19 18:49:03 2016 (
100032 inodes with total 0 MB data processed)
Invalidating snapshot files in fileset1:snapshot1/F/...
Delete snapshot fileset1:snapshot1 successful. "
}
```

## See also

- “mmdelsnapshot command” on page 304

---

## Version 2

Lists the IBM Spectrum Scale management API version 2 commands.

## | CES/addresses: GET

| Gets information about CES addresses.

## | Availability

| Available on all IBM Spectrum Scale editions.

## | Description

| The **GET ces/addresses** request gets information about CES (Cluster Export Services) addresses. For more information about the fields in the data structures that are returned, see “mmces command” on page 103.

## | Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/ces/addresses`

| where

### | ces/addresses

|     Specifies CES address as the resource. Required.

## | Request headers

| `Accept: application/json`

## | Parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 25. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

## | Request data

| No request data.

## | Response data

| The following list of attributes are available in the response data:

```
| {
| "status":
| {
| "code": ReturnCode
| "message": "ReturnMessage",
| }
| "paging":
| {
| "next": "URL"
| },
| "cesaddresses": [
| }
```

```

 {
 "oid": "Integer",
 "nodeNumber": "Integer",
 "nodeName": "Node name",
 "attributes": "Attributes",
 "cesAddress": "IP",
 "cesGroup": "Group",
 }
],
}

```

The details of the parameters are provided in the following list:

- "status":**  
Return status.
- "code": *ReturnCode*,**  
The HTTP status code that was returned by the request.
- "message": "ReturnMessage"**  
The return message.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "cesaddresses":**  
An array of information about CES addresses. For more information about the fields in this structure, see the links at the end of this topic.
  - "oid": "Integer"**  
Internal identifier that is used for paging.
  - "nodeNumber": "Integer"**  
The number of the cesNode this address belongs to.
  - "nodeName": "Node"**  
The CES node to which the address is assigned.
  - "Attributes": "Attributes"**  
Protocol attributes that are associated with the CES address.
  - "cesAddress": "IP"**  
The CES address that is assigned to the node.
  - "cesGroup": "Group"**  
The group to which the CES address is assigned.

## Examples

The following example gets information about the CES addresses available in the system:

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/ces/addresses'
```

The request URL with no field or filter parameter returns only the details that uniquely identify the object.

Response data:

| Note: In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

| The **cesaddresses** array returns four objects in the following example. Each object contains details about  
| one CES address.

```
| {
| "cesaddresses" : [{
| "cesAddress" : "198.51.100.35",
| "nodeName" : "mari-13.localnet.com"
| }, {
| "cesAddress" : "198.51.100.37",
| "nodeName" : "mari-12.localnet.com"
| }, {
| "cesAddress" : "198.51.100.31",
| "nodeName" : "mari-12.localnet.com"
| }, {
| "cesAddress" : "198.51.100.33",
| "nodeName" : "mari-14.localnet.com"
| }, {
| "cesAddress" : "198.51.100.27",
| "nodeName" : "mari-14.localnet.com"
| }
| },
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }
```

| Using the field parameter ":all:" returns entire details of the CES addresses available in the system:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/ces/addresses?fields=:all:'
```

| Response data:

```
| {
| "cesaddresses" : [{
| "attributes" : "",
| "cesAddress" : "198.51.100.35",
| "cesGroup" : "",
| "nodeName" : "mari-13.localnet.com",
| "nodeNumber" : 3,
| "oid" : 1
| }, {
| "attributes" : "",
| "cesAddress" : "198.51.100.37",
| "cesGroup" : "",
| "nodeName" : "mari-12.localnet.com",
| "nodeNumber" : 2,
| "oid" : 2
| }, {
| "attributes" : "",
| "cesAddress" : "198.51.100.31",
| "cesGroup" : "",
| "nodeName" : "mari-12.localnet.com",
| "nodeNumber" : 2,
| "oid" : 3
| }, {
| "attributes" : "",
| "cesAddress" : "198.51.100.33",
| "cesGroup" : "",
| "nodeName" : "mari-14.localnet.com",
| "nodeNumber" : 4,
| "oid" : 4
| }
```

```
| }, {
| "attributes" : "",
| "cesAddress" : "198.51.100.27",
| "cesGroup" : "",
| "nodeName" : "mari-14.localnet.com",
| "nodeNumber" : 4,
| "oid" : 5
| }
|],
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }
```

| **Related reference:**

- | “mmces command” on page 103
- | Manage CES (Cluster Export Services) configuration.

## | CES/addresses/{cesAddress}: GET

| Gets information about a CES address.

### | Availability

| Available on all IBM Spectrum Scale editions.

### | Description

| The **GET ces/addresses/{cesAddress}** request gets information about a specific CES (Cluster Export Services) address. For more information about the fields in the data structures that are returned, see “mmces command” on page 103.

### | Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/ces/addresses/CesAddress`

| where

#### | **ces/addresses**

| Specifies CES address as the resource. Required.

#### | **CesAddress**

| Specifies the CES address about which you want to get information. Required.

### | Request headers

| `Accept: application/json`

### | Parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 26. List of parameters*

| Parameter name | Description and applicable keywords | Required/optional |
|----------------|-------------------------------------|-------------------|
| cesAddress     | The IP address to query.            | Required.         |

### | Request data

| No request data.

### | Response data

| The following list of attributes are available in the response data:

```
| {
| "status":
| {
| "code": ReturnCode
| "message": "ReturnMessage",
| }
| "paging":
| {
| "next": "URL"
| },
| "cesaddresses": [
| {
| "oid": "Integer"
| "nodeNumber": "Integer"
| }
|]
| }
```

```

 "nodeName": Node name,
 "attributes": "Attributes",
 "cesAddress": "IP",
 "cesGroup": "Group",
 }
],
}

```

The details of the parameters are provided in the following list:

- "status":**  
Return status.
- "code": *ReturnCode*,**  
The HTTP status code that was returned by the request.
- "message": "ReturnMessage"**  
The return message.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "cesaddresses":**  
An array of information about CES addresses. For more information about the fields in this structure, see the links at the end of this topic.
  - "oid": "Integer"**  
Internal identifier that is used for paging.
  - "nodeNumber": "Integer"**  
The number of the cesNode this address belongs to.
  - "nodeName": "Node"**  
The CES node to which the address is assigned.
  - "Attributes": "Attributes"**  
Protocol attributes that are associated with the CES address.
  - "cesAddress": "IP"**  
The CES address that is assigned to the node.
  - "cesGroup": "Group"**  
The group to which the CES address is assigned.

## Examples

The following example gets information about the CES address 198.51.100.8.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/ces/addresses/198.51.100.8'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
}
```

```
| "cesaddresses": [
| {
| "attributes" : "",
| "cesAddress" : "198.51.100.8",
| "cesGroup" : "",
| "nodeName" : "mari-13.localnet.com",
| "nodeNumber" : 3,
| "oid" : 1
| }
|]
| }
```

| **Related reference:**

| “mmces command” on page 103

| Manage CES (Cluster Export Services) configuration.

## | CES/services: GET

| Gets information about CES services.

## | Availability

| Available on all IBM Spectrum Scale editions.

## | Description

| The **GET ces/services** request gets information about the CES (Cluster Export Services) services in the cluster. For more information about the fields in the returned data structure, see “mmces command” on page 103.

## | Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/ces/services`

| where

### | ces/services

| Specifies CES services as the resource. Required.

## | Request headers

| `Accept: application/json`

## | Parameters

| No parameters.

## | Request data

| No request data.

## | Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "cesservices": {
 "protocolStates": [
 {
 "service": "Service",
 "enabled": "{yes | no}",
 }
]
 "protocolNodes": [
 {
 "nodeName": "Node",
 "serviceStates": [
 {
 "service": "Service",
 "running": "{yes | no}",
 }
]
 },
],
 },
}
```

| The details of the parameters are provided in the following list:

| **status:**

|     Return status.

|     **"code":** *ReturnCode*,

|         The return message.

|     **"message":** "ReturnMessage"

|         The return message.

| **cesservices:**

|     Information about CES services. The information consists of two arrays, **protocolNodes** and **protocolStates**. Each array contains multiple elements, with one element for each CES service. For more information about the fields in these structures, see the link at the end of this topic.

| **protocolStates:**

|     An array of information about the services. Each element describes one service:

|     **"service":**"Service"

|         Identifies the service such as CES, NETWORK, NFS, or SMB.

|     **"enabled":**"{yes | no}"

|         Indicates whether the service is enabled.

| **protocolNodes**

|     An array of information about protocol nodes. Each array element describes one protocol node.

|     **"nodeName":**"Node"

|         The name of the node.

| **serviceStates:**

|     An array of information about the services for which the node is a protocol node.

|     **"service":**"Service"

|         Name of the service.

|     **"running":**"{yes | no}"

|         Indicates whether the service is running.

## Examples

| The following example gets information about the CES services.

| Request data:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalermgmt/v2/ces/services'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

| The **protocolNodes** array returns the details of the protocols configured in each protocol node. The **protocolStates** array provides the status of the protocols.

```
{
 "cesservices" : {
 "protocolNodes" : [{
 "nodeName" : "mari-12.localnet.com",
 "serviceStates" : [{
 "running" : "yes",
 "service" : "OBJ"
 }]
 }]
 }
}
```

```

 }, {
 "running" : "yes",
 "service" : "SMB"
 }, {
 "running" : "yes",
 "service" : "NFS"
 }]
 }, {
 "nodeName" : "mari-13.localnet.com",
 "serviceStates" : [{
 "running" : "yes",
 "service" : "OBJ"
 }, {
 "running" : "yes",
 "service" : "SMB"
 }, {
 "running" : "yes",
 "service" : "NFS"
 }]
 }, {
 "nodeName" : "mari-14.localnet.com",
 "serviceStates" : [{
 "running" : "yes",
 "service" : "OBJ"
 }, {
 "running" : "yes",
 "service" : "SMB"
 }, {
 "running" : "yes",
 "service" : "NFS"
 }]
 }, {
 "nodeName" : "mari-15.localnet.com",
 "serviceStates" : [{
 "running" : "yes",
 "service" : "OBJ"
 }, {
 "running" : "yes",
 "service" : "SMB"
 }, {
 "running" : "yes",
 "service" : "NFS"
 }]
 },
 "protocolStates" : [{
 "enabled" : "yes",
 "service" : "OBJ"
 }, {
 "enabled" : "no",
 "service" : "BLOCK"
 }, {
 "enabled" : "yes",
 "service" : "SMB"
 }, {
 "enabled" : "yes",
 "service" : "NFS"
 }]
 },
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
 }
}

```

**Related reference:**

- | “[mmces command](#)” on page 103
- | Manage CES (Cluster Export Services) configuration.

## | CES/services/{service}: GET

| Gets information about a CES service.

### | Availability

| Available on all IBM Spectrum Scale editions.

### | Description

| The **GET ces/services/service** request gets information about a CES (Cluster Export Services) service in the cluster. For more information about the fields in the returned data structure, see “mmces command” on page 103.

### | Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/ces/services/service`

| where

| **ces/services**

|     Specifies ces/services as the resource.

| **service**

|     Specifies the service about which you want to get information. Required.

### | Request headers

| `Accept: application/json`

### | Parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 27. List of parameters*

| Parameter name | Description and applicable keywords                   | Required/optional |
|----------------|-------------------------------------------------------|-------------------|
| service        | Name of the service about which you need the details. | Required.         |

### | Request data

| No request data.

### | Response data

```
| {
| "status": {
| "code": ReturnCode,
| "message": "ReturnMessage"
| }
| "cesservices": {
| "protocolStates": [
| {
| "service": "Service"
| "enabled": "{yes | no}",
| }
|]
| "protocolNodes": [
| {
| "nodeName": "Node"
| }
|]
| }
| }
```

```

 "serviceStates": [
 {
 "service": "Service"
 "running": "{yes | no}",
 }
],
 },
},
}

status:
Return status.

"code": ReturnCode,
The return message.

"message": "ReturnMessage"
The return message.

cesservices:
Information about CES services. The information consists of two arrays, protocolNodes and protocolStates. Each array contains multiple elements, with one element for each CES service. For more information about the fields in these structures, see the link at the end of this topic.

protocolStates:
An array of information about the services. Each element describes one service:

"service": "Service"
Identifies the service such as CES, NETWORK, NFS, or SMB.

"enabled": "{yes | no}"
Indicates whether the service is enabled.

protocolNodes
An array of information about protocol nodes. Each array element describes one protocol node.

"nodeName": "Node"
The name of the node.

serviceStates:
An array of information about the services for which the node is a protocol node.

"service": "Service"
Name of the service.

"running": "{yes | no}"
Indicates whether the service is running.

```

## Examples

The following example gets information about the SMB service.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/ces/services/SMB'
```

Response data:

- Note:** In the JSON data that is returned, the return code indicates whether the command is successful.
- The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

The **protocolNodes** array returns the details of the node on which the SMB protocol is configured. The **protocolStates** array provides the status of the SMB protocol.

```
{
 "status": {
 "code": "200",
 "message": "The request finished successfully"
 },
 "cesservices": {
 "protocolStates": [
 {
 "service": "SMB",
 "enabled": "yes"
 }
],
 "protocolNodes": [
 {
 "nodeName": "testnode-1.localnet.com",
 "serviceStates": [
 {
 "service": "SMB",
 "running": "yes"
 }
]
 }
]
 }
}
```

**Related reference:**

“mmces command” on page 103

Manage CES (Cluster Export Services) configuration.

## Cluster: GET

Gets details of the cluster.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET cluster** request gets information about the local cluster. For more information about the fields in the data structures that are returned, see the following topics: “[mmlscluster command](#)” on page 397, “[mmchfs command](#)” on page 179, and “[mmlsfs command](#)” on page 411.

### Request URL

`https://<IP address or host name of API server>:<port>/scalemgmt/v2/cluster`

where

#### cluster

Specifies the IBM Spectrum Scale cluster as the resource of the GET call.

### Request headers

`Accept: application/json`

### Request data

No request data.

### Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "cluster": {
 "clusterSummary": {
 "clusterName": "ClusterName",
 "clusterId": "ID",
 "uidDomain": "Domain",
 "rshPath": "RemoteShellCommand",
 "rshSudoWrapper": "Path",
 "rcpPath": "RemoteFileCopyCommand",
 "rcpSudoWrapper": "Path",
 "repositoryType": "CCR | non-CCR",
 "primaryServer": "Server",
 "secondaryServer": "Server",
 },
 "cnfsSummary": {
 "cnfsSharedRoot": "Directory",
 "cnfsMountPort": "Port",
 "cnfsNFSDprocs": "Number",
 "cnfsReboot": "{true | false}",
 "cnfsMonitorEnabled": "enabled | disabled",
 "cnfsGanesha": "ServerAddress",
 },
 "cesSummary": {
 "cesSharedRoot": "Directory",
 "enabledServices": "ServiceList",
 "logLevel": "Level",
 }
 }
}
```

```

 "addressPolicy": "{none | balanced-load | node-affinity | even-coverage}",
 }
},
}

| For more information about the fields in the following data structures, see the links at the end of this
| topic.

| Note: The structures cesNode, cesSummary, cnfsNode, cnfsSummary, and gatewayNode appear in the
| output only when the corresponding role is active on the node.

| "cluster":
| A data structure that describes the local cluster. It contains the following data structures:
| cesSummary, clusterSummary, cnfsSummary, links, and nodes.

| "clusterSummary":
| Cluster information.

| "clusterId": "ID"
| The ID of the cluster.

| "clusterName": "ClusterName"
| The name of the cluster.

| "primaryServer": "Server"
| The primary server node for GPFS cluster data.

| "rcpPath": "RemoteFileCopyCommand"
| The remote file copy command.

| "rcpSudoWrapper": ""
| The fully qualified path of the remote file copy program for sudo wrappers.

| "repositoryType": "CCR | non-CCR"
| The type of repository that the cluster uses for storing configuration data.

| "rshPath": "RemoteShellCommand"
| The remote shell command.

| "rshSudoWrapper": "Path"
| The fully qualified path of the remote shell program for sudo wrappers.

| "secondaryServer": "Server"
| The secondary server node for GPFS cluster data.

| "uidDomain": "Domain"
| The UID domain name of the cluster.

| "cnfsSummary":
| CNFS cluster information.

| "cnfsGanesha" : "IPAddress"
| The server address of the Ganesha CNFS server.

| "cnfsMonitorEnabled" : "enabled" | "disabled"
| The state of CNFS monitoring.

| "cnfsMountPort" : "Port"
| The port number of the rpc.mountd daemon.

| "cnfsNFSDprocs" : "Number"
| The number of nfsd server threads.

| "cnfsReboot" : "true" | "false"
| Whether the node restarts when CNFS monitoring detects an unrecoverable problem.

```

```

| "cnfsSharedRoot" : "Path"
| A directory that the CNFS subsystem uses.

| "cesSummary":
| CES cluster information.

| "enabledServices": "ServiceList"
| A comma-separated list of the services that are enabled. The list can include NFS, SMB, OBJ,
| BLOCK, or None.

| "addressPolicy": "{none | balanced-load | node-affinity | even-coverage}"
| The address policy for distributing CES addresses.

| "cesSharedRoot": "Directory"
| The CES shared root directory, which is used for storing CES shared configuration data.

| "logLevel": Level
| The CES log level.

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": ReturnCode
| The return code.

```

## Examples

The following example gets information about the local cluster.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/cluster'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
The response code 200 indicates that the command successfully retrieved the information. Error code 400  
represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "cluster": {
 "clusterSummary": {
 "clusterName": "gpfs-cluster-2.localnet.com",
 "clusterId": "13445038716632536363",
 "uidDomain": "localnet.com",
 "rshPath": "/usr/bin/ssh",
 "rshSudoWrapper": "no",
 "rcpPath": "/usr/bin/scp",
 "rcpSudoWrapper": "no",
 "repositoryType": "CCR",
 "primaryServer": "testnode-1.localnet.com",
 "secondaryServer": "testnode-2.localnet.com"
 },
 "cnfsSummary": {
 "cnfsSharedRoot": "/mnt/objfs",
 "cnfsMountdPort": "111",
 "cnfsNFSdprocs": "32",
 }
 }
}
```

```
| "cnfsReboot": "yes",
| "cnfsMonitorEnabled": "yes",
| "cnfsGanesha": "198.51.100.8"
|,
| "cesSummary": {
| "cesSharedRoot": "/mnt/objfs/ces",
| "enabledServices": "SMB,NFS",
| "logLevel": "0",
| "addressPolicy": "even-coverage"
| }
| }
|
| "mmlscluster command" on page 397
| Displays the current configuration information for a GPFS cluster.
|
| "mmchfs command" on page 179
| Changes the attributes of a GPFS file system.
|
| "mmlsfs command" on page 411
| Displays file system attributes.
```

## | **Config: GET**

| Gets information about the configuration of the IBM Spectrum Scale cluster.

## | **Availability**

| Available on all IBM Spectrum Scale editions.

## | **Description**

| The **GET config** request gets information about the local cluster. For more information about the fields in  
| the data structures that are returned, see the following topics: “**mmchconfig command**” on page 132 and  
| “**mmlsconfig command**” on page 400.

| This API call gets its data from the **mmlsconfig** command. The **mmlsconfig** command displays the  
| PERSISTED configuration values; not the ACTIVE values. That is, the configuration changes made  
| through the **mmchconfig** command come into effect based on how you use **-I** and **-i** parameters of  
| the command. The following three scenarios are applicable:  
| • **mmchconfig -I**: A value is changed immediately in the daemon but not persistent so that it is reset after  
| a restart. Those settings do not show in **mmlsconfig** at all. In this case, this API call shows an old  
| setting that is temporarily changed to something else.  
| • **mmchconfig -i**: The setting takes effect immediately and is also persisted so that **mmlsconfig** shows it.  
| • **mmchconfig**: The setting takes effect on the next restart and is also persisted so that **mmlsconfig** shows  
| it. In this case, the API call shows a future setting that will get applied after the restart.

## | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/config`

| where

### | **config**

| Specifies that the GET request fetches the configuration details of the cluster.

## | **Request headers**

| Content-Type: application/json  
| Accept: application/json

## | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 28. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |

## | **Request data**

| No request data.

## | **Response data**

| The following list of attributes are available in the response data:

```

{
 "status": [
 {
 "code": ReturnCode,
 "message": "ReturnMessage",
 }
],
 "config": [
 {
 "clusterConfig": "Configuration details"
 }
]
}

"status":
 Return status.

"code": ReturnCode,
 The HTTP status code that was returned by the request.

"message": "ReturnMessage"
 The return message.

"config":
 An array of information about the cluster configuration.

"clusterConfig": "Configuration details"
 The cluster configuration as shown by mmlsconfig -Y command.

The return information and the information that the command retrieves are returned in the same way as they are for the other requests. The parameters that are returned are the same as the configuration attributes that are displayed by the mmlsconfig command. For more information, see the “mmlsconfig command” on page 400.

```

## Examples

- The following example gets information about the cluster configuration.
- Request data:  

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/config'
```
- Response data:  

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.
- The **config** array provides information about the cluster configuration.

```
{
 "config" : {
 "clusterConfig" : {
 "FIPS1402mode" : "no",
 "adminMode" : "central",
 "afmAsyncDelay" : "15",
 "afmAsyncOpWaitTimeout" : "120",
 "afmAtimeXattr" : "no",
 "afmDIO" : "0",
 "afmDirLookupRefreshInterval" : "60",
 "afmDirOpenRefreshInterval" : "60",
 "afmDisconnectTimeout" : "60",
 "afmEnableADR" : "no",
 "afmExpirationTimeout" : "disable",
 "afmFileLookupRefreshInterval" : "30",
 }
 }
}
```

```
"afmFileOpenRefreshInterval" : "30",
"afmFlushThreadDelay" : "5",
"afmGatewayQueueTransfer" : "yes",
"afmHardMemThreshold" : "5368709120",
"afmMaxThrottle" : "0",
"afmMaxWriteMergeLen" : "1073741824",
"afmMode" : "read-only",
"afmMountRetryInterval" : "300",
"afmNumIOFlushThreads" : "32",
"afmNumReadThreads" : "1",
"afmNumWriteThreads" : "1",
"afmParallelReadChunkSize" : "134217728",
"afmParallelReadThreshold" : "1024",
"afmParallelWriteChunkSize" : "134217728",
"afmParallelWriteThreshold" : "1024",
"afmRPO" : "disable",
"afmReadBufferSize" : "33554432",
"afmReadFlags" : "cached",
"afmReadSparseThreshold" : "134217728",
"afmShowHomeSnapshot" : "no",
"afmSkipWrites" : "(undefined)",
"afmSyncOpWaitTimeout" : "60",
"afmWriteBufferSize" : "2097152",
"atimeDeferredSeconds" : "86400",
"autoload" : "yes",
"automountDir" : "/gpfs/automountdir",
"ccrEnabled" : "yes",
"cesSharedRoot" : "/mnt/objfs/ces",
"cifsBypassShareLocksOnRename" : "yes",
"cifsBypassTraversalChecking" : "yes",
"cipherList" : "AUTHONLY",
"clusterId" : "317908494281353658",
"clusterName" : "gpfsgui-cluster-1.localnet.com",
"cnfsMountdPort" : "(undefined)",
"cnfsNFSDprocs" : "32",
"cnfsReboot" : "yes",
"cnfsSharedRoot" : "(undefined)",
"commandAudit" : "syslogOnly",
"dataCollectionPendingDelay" : "10",
"dataDiskCacheProtectionMethod" : "0",
"dataDiskWaitTimeForRecovery" : "3600",
"dataStructureDump" : "/tmp/mmfs",
"deadlockBreakupDelay" : "0",
"deadlockDataCollectionDailyLimit" : "3",
"deadlockDataCollectionMinInterval" : "3600",
"deadlockDetectionThreshold" : "300",
"deadlockDetectionThresholdForShortWaiters" : "60",
"deadlockDetectionThresholdIfOverloaded" : "1800",
"deadlockOverloadThreshold" : "1",
"debugDataControl" : "light",
"debugDataControlOnSignal" : "heavy",
"defaultHelperNodes" : "(undefined)",
"defaultMountDir" : "/gpfs",
"disableInodeUpdateOnFdatasync" : "no",
"dmapiDataEventRetry" : "2",
"dmapiEventTimeout" : "86400000",
"dmapiFileHandleSize" : "32",
"dmapiMountEvent" : "all",
"dmapiMountTimeout" : "60",
"dmapiSessionFailureTimeout" : "0",
"enableIPv6" : "no",
"enforceFilesetQuotaOnRoot" : "no",
"expelDataCollectionDailyLimit" : "3",
"expelDataCollectionMinInterval" : "3600",
"failureDetectionTime" : "35",
"fastestPolicyCmpThreshold" : "50",
"fastestPolicyMaxValidPeriod" : "600",
```

```

 "fastestPolicyMinDiffPercent" : "50",
 "fastestPolicyNumReadSamples" : "5",
 "fileHeatLossPercent" : "10",
 "fileHeatPeriodMinutes" : "0",
 "forceLogWriteOnFdatasync" : "no",
 "indefiniteRetentionProtection" : "no",
 "indirectBlocksPerLockRange" : "10",
 "maxDownDisksForRecovery" : "16",
 "maxFailedNodesForRecovery" : "3",
 "maxFcntlRangesPerFile" : "200",
 "maxFilesToCache" : "4000",
 "maxMBpS" : "2048",
 "maxStatCache" : "1000",
 "maxblocksize" : "16384K",
 "metadataDiskWaitTimeForRecovery" : "2400",
 "minDiskWaitTimeForRecovery" : "1800",
 "minReleaseLevel" : "4.2.3.0",
 "minimumTargetTimeUnderRangeLock" : "30",
 "mmapRangeLock" : "yes",
 "mmfsLogLevel" : "detail",
 "nistCompliance" : "SP800-131A",
 "noSpaceEventInterval" : "120",
 "nsdRAIDBufferSizePct" : "50",
 "nsdRAIDTracks" : "0",
 "nsdServerWaitTimeForMount" : "300",
 "nsdServerWaitTimeWindowOnMount" : "600",
 "nsdbufspace" : "30",
 "pagepool" : "1G",
 "pagepoolMaxPhysMemPct" : "75",
 "pitWorkerThreadsPerNode" : "0",
 "prefetchThreads" : "72",
 "readReplicaPolicy" : "default",
 "restripeOnDiskFailure" : "no",
 "rpcPerfNumberDayIntervals" : "30",
 "rpcPerfNumberHourIntervals" : "24",
 "rpcPerfNumberMinuteIntervals" : "60",
 "rpcPerfNumberSecondIntervals" : "60",
 "rpcPerfRawExecBufferSize" : "2",
 "rpcPerfRawStatBufferSize" : "6",
 "sidAutoMapRangeLength" : "15000000",
 "sidAutoMapRangeStart" : "15000000",
 "subnets" : "(undefined)",
 "syncSambaMetadataOps" : "yes",
 "systemLogLevel" : "notice",
 "tctEnable" : "yes",
 "tiebreakerDisks" : "(undefined)",
 "tmMaxPhysMemPct" : "75",
 "tscCmdPortRange" : "(undefined)",
 "uidDomain" : "localnet.com",
 "unmountOnDiskFail" : "no",
 "usePersistentReserve" : "no",
 "verbsPorts" : "(undefined)",
 "verbsRdma" : "disable",
 "verbsRdmaCm" : "disable",
 "verbsRdmaPkey" : "32767",
 "verbsRdmaRoCEToS" : "-1",
 "verbsRdmaSend" : "no",
 "verbsRdmasesPerConnection" : "16",
 "verbsRdmasesPerNode" : "1000",
 "verbsSendBufferMemoryMB" : "0",
 "worker1Threads" : "48"
 }
},
"status" : {

```

```
| "code" : 200,
| "message" : "The request finished successfully"
| }
|
| "mmlsconfig command" on page 400
| Displays the current configuration data for a GPFS cluster.
| "mmchconfig command" on page 132
| Changes GPFS configuration parameters.
```

## | **Filesystems: GET**

| Gets information about file systems.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems** request gets information about file systems in the cluster. For more information about the fields in the data structures that are returned, see the topics “**mmcrfs command**” on page 250, “**mmchfs command**” on page 179, and “**mmlsfs command**” on page 411.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems`

| where

#### | **filesystems**

| Specifies file systems as the resource of the GET call.

### | **Request headers**

| `Accept: application/json`

### | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 29. List of parameters*

| Parameter name | Description and applicable keywords                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. 'all' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                 | Optional.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
| "paging":
| {
| "next": "URL"
| },
| "filesystems": [
| {
| "oid":Internal ID,
| }
|]
| }
```

```

"uuid": "ID",
"name": "Name",
"version": "Version",
"type": "Type",
"createTime": "DateTime",
"block"
{
 "pools": "Pools",
 "disks": "List of disks",
 "blockSize": "Block size",
 "metaDataBlockSize": "Metadata block size",
 "indirectBlockSize": "Indirect block size",
 "minFragmentSize": "Minimum fragment size",
 "inodeSize": "Inode size",
 "logFileSize": "Log file size",
 "writeCacheThreshold": "Threshold"
},
"mount"
{
 "mountPoint": "Path",
 "automaticMountOption": "{yes | no | automount}",
 "additionalMountOptions": "Options",
 "mountPriority": "Priority",
 "driveLetter": "Drive letter",
 "remoteDeviceName": "Device name",
 "readOnly": "Read-only",
}
"replication"
{
 "defaultMetadataReplicas": "Number",
 "maxMetadataReplicas": "Number",
 "defaultDataReplicas": "Number",
 "maxDataReplicas": "Number",
 "strictReplication": "{no | whenpossible | always}",
 "logReplicas": "Number",
}
"quota"
(
 "quotasAccountingEnabled": "[user] [;group] [;fileset]"
 "quotasEnforced": "[user] [;group] [;fileset]"
 "defaultQuotasEnabled": "[user] [;group] [;fileset]"
 "perfilesetQuotas": "{yes | no}",
 "filesetdfEnabled": "{yes | no}",
)
"settings"
(
 "blockAllocationType": "{scatter | cluster}",
 "fileLockingSemantics": "{nfs4 | posix}",
 "numNodes": "Number",
 "exactMTime": "{true | false}",
 "suppressATime": "{yes | no}",
 "fastEAEEnabled": "{yes | no}",
 "encryption": "{yes | no}",
 "maxNumberOfInodes": "Number",
 "is4KAligned": "{yes | no}",
 "rapidRepairEnabled": "{yes | no}",
 "stripeMethod": "{yes | no}",
 "stripedLogs": "{yes | no}",
 "ACLSemantics": "{posix | afs4 | all}",
 "DMAPIEnabled": "{yes | no}",
)
}
],
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

```

| "status":
| Return status.
|
| "message": "ReturnMessage",
| The return message.
|
| "code": "ReturnCode"
| The return code.
|
| "filesystems":
| An array of elements that describe file systems. Each element describes one file system.
|
| "oid": "Internal ID",
| Internal identifier that is used for paging.
|
| "uuid": "ID"
| The UUID of the file system.
|
| "name": "Name"
| Name of the file system.
|
| "version": "Version"
| File system version.
|
| "type": "local | remote"
| File system type.
|
| "createTime": "DateTime"
| Creation time.
|
| "block"
|
| "pools": "Pools"
| List of pools of this file system.
|
| "disks": "List of disks"
| A semicolon-separated list of the disks that are included in the file system.
|
| "blockSize": "Block size"
| The block size of the disks in the storage pool.
|
| "metaDataBlockSize": "Metadata block size"
| Block size of metadata pool.
|
| "indirectBlockSize": "Indirect block size"
| Indirect block size in bytes.
|
| "minFragmentSize": "Minimum fragment size"
| Minimum fragment size in bytes.
|
| "inodeSize": "Inode size"
| Inode size in bytes.
|
| "logfileSize": "Log file size"
| The size of the internal log files in bytes.
|
| "writeCacheThreshold": "Threshold"
| Specifies the maximum length (in bytes) of write requests that are initially buffered in the
| highly available write cache before being written back to primary storage.
|
| "mount"
|
| "mountPoint": "Path"
| Default mount point.

```

```

| "automaticMountOption":"{yes | no | automount }"
| Indicates when the file system is to be mounted: yes when GPFS daemon starts, no to manual
| mount, automount when the file system is first accessed.
|
| "additionalMountOptions":"Options"
| The mount options to pass to the mount command while mounting the file system.
|
| "mountPriority":"Priority"
| Controls the order in which the individual file systems are mounted at daemon startup or
| when one of the all keywords is specified on the mmmount command.
|
| "driveLetter":"Drive letter"
| Drive letter when mounted on windows nodes.
|
| "remoteDeviceName":"Device name"
| Device name on the owning cluster of this cross-cluster mounted file system.
|
| "readOnly":"Read-only"
| Read-only file system.
|
| "replication"
|
| "defaultMetadataReplicas":"Number"
| Default number of metadata replicas.
|
| "maxMetadataReplicas":"Number"
| Maximum number of metadata replicas.
|
| "defaultDataReplicas":"Number"
| Default number of data replicas.
|
| "maxDataReplicas":"Number"
| Maximum number of data replicas.
|
| "strictReplication":"{no | whenpossible | always }"
| Whether strict replication is to be enforced. 'no': not enforced, 'whenpossible': enforced if at
| least two failure groups present, 'always': always enforced.
|
| "logReplicas":"Number"
| Number of log replicas.
|
| "quotas"
|
| "quotasAccountingEnabled":"[user] [;group] [;fileset]"
| The types of quotas for which accounting is enabled. Accounting means that usage is tracked
| against quotas. If action needs to be taken when quotas are exceeded, quotas must be also
| enforced.
|
| "quotasEnforced":"[user] [;group] [;fileset]"
| Quotas that are enforced for any of 'user', 'group' and 'fileset'.
|
| "defaultQuotasEnabled":"[user] [;group] [;fileset]"
| Default quotas that are enabled.
|
| "perfilesetQuotas":"{yes | no}"
| Specifies whether the scope of user and group quota limit checks on the individual fileset
| level rather than the entire file system.
|
| "filesetdfEnabled":"{yes | no}"
| Specifies whether the quotas are enforced for a fileset, the numbers reported by the df
| command are based on the quotas for the fileset rather than the entire file system.
|
| "settings"
|
| "blockAllocationType":"{scatter | cluster}"
| Block allocation type.

```

```

| "fileLockingSemantics": "{nfs4 | posix}"
| File locking semantics in effect.
|
| "numNodes": Number,
| Estimated number of nodes that will mount file system.
|
| "exactMTime": "{true | false}"
| true reports exact mtime values and false mtime is only periodically updated.
|
| "suppressATime": "{yes | no | realtime}"
| Suppress the periodic updating of the value of atime as reported by stat calls.
|
| "fastEAEEnabled": "{yes | no}"
| Specifies whether fast extended attributes are enabled.
|
| "encryption": "{yes | no}"
| Specifies whether encryption is enabled.
|
| "maxNumberOfInodes": "Number"
| Specifies the maximum number of files that can be created inside the inode space of the root
| fileset of this file system.
|
| "is4KAligned": "{yes | no}"
| Whether the file system is formatted to be 4K aligned.
|
| "rapidRepairEnabled": "{yes | no}"
| Whether to keep track of incomplete replication on an individual file block basis as opposed
| to the entire file.
|
| "stripeMethod": "{yes | no}"
| Algorithm that is used for striping data.
|
| "stripedLogs": "{yes | no}"
| Metadata logs are striped across available metadata disks.
|
| "ACLSemantics": "{posix | afs4 | all }"
| ACL semantics supported by this file system.
|
| "DMAPIEnabled": "{yes | no}"
| Whether Data Management API is enabled on this file system.

```

## Examples

The following example gets information about file systems in the cluster.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems'
```

The request URL with no field or filter parameter returns only the details that uniquely identify the object.

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "filesystems" : [{
 "name" : "gpfs0"
 }, {
 "name" : "objfs"
 }, {
```

```

 "name" : "gpfs1"
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}

```

Using the field parameter ":all:" returns entire details of the file systems that are configured in the system. The **filesystems** array returns file system objects in the following example. Each object contains details about one file system:

```

{
 "filesystems" : [{
 "oid" : 1,
 "uuid" : "0A00640B:58B82AF4",
 "name" : "gpfs0",
 "version" : "17.00",
 "type" : "local",
 "createTime" : "2017-03-02 15:23:48,000",
 "block" : {
 "blockSize" : 262144,
 "disks" : "disk1;disk8",
 "indirectBlockSize" : 16384,
 "inodeSize" : 4096,
 "logfileSize" : 4194304,
 "metaDataBlockSize" : 262144,
 "minFragmentSize" : 8192,
 "pools" : "system;data",
 "writeCacheThreshold" : 0
 },
 "mount" : {
 "additionalMountOptions" : "none",
 "automaticMountOption" : "yes",
 "mountPoint" : "/mnt/gpfs0",
 "mountPriority" : 0,
 "readOnly" : false,
 "remoteDeviceName" : "gpfsgui-cluster-1.localnet.com:gpfs0"
 },
 "replication" : {
 "defaultDataReplicas" : 1,
 "defaultMetadataReplicas" : 1,
 "logReplicas" : 0,
 "maxDataReplicas" : 2,
 "maxMetadataReplicas" : 2,
 "strictReplication" : "whenpossible"
 },
 "quota" : {
 "defaultQuotasEnabled" : "none",
 "filesetdfEnabled" : false,
 "perfilesetQuotas" : true,
 "quotasAccountingEnabled" : "user;group;fileset",
 "quotasEnforced" : "user;group;fileset"
 },
 "settings" : {
 "aclSemantics" : "nfs4",
 "blockAllocationType" : "cluster",
 "dmapiEnabled" : false,
 "encryption" : false,
 "exactMTIME" : true,
 "fastEAEEnabled" : true,
 "fileLockingSemantics" : "nfs4",
 "is4KAligned" : true,
 "maxNumberOfInodes" : 1747648,
 "numNodes" : 32,
 "rapidRepairEnabled" : true,
 "stripeMethod" : "roundRobin",
 }
 }
}

```

```

 "stripedLogs" : true,
 "suppressATime" : "no"
 }
},
{
 "oid" : 2,
 "uuid" : "0A00640C:58B82C5A",
 "name" : "objfs",
 "version" : "17.00",
 "type" : "local",
 "createTime" : "2017-03-02 15:29:46,000",
 "block" : {
 "blockSize" : 262144,
 "disks" : "disk2;disk7",
 "indirectBlockSize" : 16384,
 "inodeSize" : 4096,
 "logfileSize" : 4194304,
 "metaDataBlockSize" : 262144,
 "minFragmentSize" : 8192,
 "pools" : "system;data",
 "writeCacheThreshold" : 0
 },
 "mount" : {
 "additionalMountOptions" : "none",
 "automaticMountOption" : "yes",
 "mountPoint" : "/mnt/objfs",
 "mountPriority" : 0,
 "readOnly" : false,
 "remoteDeviceName" : "gpfsgui-cluster-1.localnet.com:objfs"
 },
 "replication" : {
 "defaultDataReplicas" : 1,
 "defaultMetadataReplicas" : 1,
 "logReplicas" : 0,
 "maxDataReplicas" : 2,
 "maxMetadataReplicas" : 2,
 "strictReplication" : "whenpossible"
 },
 "quota" : {
 "defaultQuotasEnabled" : "none",
 "filesetdfEnabled" : false,
 "perfilesetQuotas" : false,
 "quotasAccountingEnabled" : "none",
 "quotasEnforced" : "none"
 },
 "settings" : {
 "aclSemantics" : "all",
 "blockAllocationType" : "cluster",
 "dmapiEnabled" : false,
 "encryption" : false,
 "exactMTIME" : true,
 "fastEAEnabled" : true,
 "fileLockingSemantics" : "nfs4",
 "is4KAligned" : true,
 "maxNumberOfInodes" : 1747648,
 "numNodes" : 32,
 "rapidRepairEnabled" : true,
 "stripeMethod" : "roundRobin",
 "stripedLogs" : true,
 "suppressATime" : "no"
 }
},
{
 "oid" : 3,
 "uuid" : "0A00640D:58B95D19",
 "name" : "gpfs1",
 "version" : "17.00",
 "type" : "local",
 "createTime" : "2017-03-03 13:10:01,000",

```

```

"block" : {
 "blockSize" : 262144,
 "disks" : "disk3;disk4;disk5",
 "indirectBlockSize" : 16384,
 "inodeSize" : 4096,
 "logfileSize" : 4194304,
 "metaDataBlockSize" : 262144,
 "minFragmentSize" : 8192,
 "pools" : "system;data",
 "writeCacheThreshold" : 0
},
"mount" : {
 "additionalMountOptions" : "nfssync,nodev",
 "automaticMountOption" : "yes",
 "mountPoint" : "/mnt/gpfs1",
 "mountPriority" : 0,
 "readOnly" : false,
 "remoteDeviceName" : "gpfsgui-cluster-1.localnet.com:gpfs1"
},
"replication" : {
 "defaultDataReplicas" : 1,
 "defaultMetadataReplicas" : 1,
 "logReplicas" : 0,
 "maxDataReplicas" : 2,
 "maxMetadataReplicas" : 2,
 "strictReplication" : "whenpossible"
},
"quota" : {
 "defaultQuotasEnabled" : "none",
 "filesetdfEnabled" : false,
 "perfilesetQuotas" : false,
 "quotasAccountingEnabled" : "user;group;fileset",
 "quotasEnforced" : "user;group;fileset"
},
"settings" : {
 "aclSemantics" : "nfs4",
 "blockAllocationType" : "scatter",
 "dmapiEnabled" : false,
 "encryption" : false,
 "exactMTIME" : true,
 "fastEAEnabled" : true,
 "fileLockingSemantics" : "nfs4",
 "is4KAligned" : true,
 "maxNumberOfInodes" : 100032,
 "numNodes" : 32,
 "rapidRepairEnabled" : true,
 "stripeMethod" : "roundRobin",
 "stripedLogs" : true,
 "suppressATime" : "no"
}
}],
"status" : {
 "code" : 200,
 "message" : "The request finished successfully"
}
}

```

“mmcrfs command” on page 250

Creates a GPFS file system.

“mmchfs command” on page 179

Changes the attributes of a GPFS file system.

“mmlsfs command” on page 411

Displays file system attributes.

## | **Filesystems/{filesystemName}: GET**

| Gets information about a file system.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/filesystemName** request gets information about a particular file system. For more information about the fields in the data structures that are returned, see the topics “mmcrfs command” on page 250, “mmchfs command” on page 179, and “mmlfs command” on page 411.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/fileSystemName`

| where

#### | **filesystems**

| Specifies file system as the resource of the GET call. Required.

#### | **FileSystemName**

| The file system about which you want to get information. Required.

### | **Request headers**

| `Accept: application/json`

### | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 30. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. |                   |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                |                   |

### | **Request data**

| No request data.

### | **Response data**

```
{
 "status": {
 "code":ReturnCode,
 "message":ReturnMessage
 },
 "paging":
```

```

{
 "next": "URL"
},
filesystems: [
 {
 "oid": "Internal ID",
 "uuid": "ID",
 "name": "Name",
 "version": "Version",
 "type": "Type",
 "createTime": "DateTime",
 "block": {
 "pools": "Pools",
 "disks": "List of disks",
 "blockSize": "Block size",
 "metaDatBlockSize": "Metadata block size",
 "indirectBlockSize": "Indirect block size",
 "minFragmentSize": "Minimum fragment size",
 "inodeSize": "Inode size",
 "logFileSize": "Log file size",
 "writeCacheThreshold": "Threshold"
 },
 "mount": {
 "mountPoint": "Path",
 "automaticMountOption": "{yes | no | automount}",
 "additionalMountOptions": "Options",
 "mountPriority": "Priority",
 "driveLetter": "Drive letter",
 "remoteDeviceName": "Device name",
 "readOnly": "Read-only"
 }
 },
 "replication": {
 "defaultMetadataReplicas": "Number",
 "maxMetadataReplicas": "Number",
 "defaultDataReplicas": "Number",
 "maxDataReplicas": "Number",
 "strictReplication": "{no | whenpossible | always}",
 "logReplicas": "Number"
 },
 "quota": {
 "quotasAccountingEnabled": "[user] [;group] [;fileset]",
 "quotasEnforced": "[user] [;group] [;fileset]",
 "defaultQuotasEnabled": "[user] [;group] [;fileset]",
 "perfilesetQuotas": "{yes | no}",
 "filesetdfEnabled": "{yes | no}"
 }
},
"settings": {
 "blockAllocationType": "{scatter | cluster}",
 "fileLockingSemantics": "{nfs4 | posix}",
 "numNodes": "Number",
 "exactMTIME": "{true | false}",
 "suppressATime": "{yes | no}",
 "fastEAEEnabled": "{yes | no}",
 "encryption": "{yes | no}",
 "maxNumberOfInodes": "Number",
 "is4KAligned": "{yes | no}",
 "rapidRepairEnabled": "{yes | no}",
 "stripeMethod": "{yes | no}",
 "stripedLogs": "{yes | no}",
 "ACLSemantics": "{posix | afs4 | all }",
 "DMAPIEnabled": "{yes | no}"
}

```

```

| }
|],
| }
| For more information about the fields in the following data structures, see the links at the end of this
| topic.

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": "ReturnCode"
| The return code.

| "filesystems":
| An array of elements that describe file systems. Each element describes one file system.

| "oid": "Internal ID",
| Internal identifier that is used for paging.

| "uuid": "ID"
| The UUID of the file system.

| "name": "Name"
| Name of the file system.

| "version": "Version"
| File system version.

| "type": "local | remote"
| File system type.

| "createTime": "DateTime"
| Creation time.

| "block"

| "pools": "Pools"
| List of pools of this file system.

| "disks": "List of disks"
| A semicolon-separated list of the disks that are included in the file system.

| "blockSize": "Block size"
| The block size of the disks in the storage pool.

| "metaDataBlockSize": "Metadata block size"
| Block size of metadata pool.

| "indirectBlockSize": "Indirect block size"
| Indirect block size in bytes.

| "minFragmentSize": "Minimum fragment size"
| Minimum fragment size in bytes.

| "inodeSize": "Inode size"
| Inode size in bytes.

| "logfileSize": "Log file size"
| The size of the internal log files in bytes.

| "writeCacheThreshold": "Threshold"
| Specifies the maximum length (in bytes) of write requests that are initially buffered in the
| highly available write cache before being written back to primary storage.

```

```

| "mount"
|
| "mountPoint": "Path"
| Default mount point.
|
| "automaticMountOption": "{yes | no | automount }"
| Indicates when the file system is to be mounted: yes when GPFS daemon starts, no to manual
| mount, automount when the file system is first accessed.
|
| "additionalMountOptions": "Options"
| The mount options to pass to the mount command while mounting the file system.
|
| "mountPriority": "Priority"
| Controls the order in which the individual file systems are mounted at daemon startup or
| when one of the all keywords is specified on the mmmount command.
|
| "driveLetter": "Drive letter"
| Drive letter when mounted on windows nodes.
|
| "remoteDeviceName": "Device name"
| Device name on the owning cluster of this cross-cluster mounted file system.
|
| "readOnly": "Read-only"
| Read-only file system.

| "replication"
|
| "defaultMetadataReplicas": "Number"
| Default number of metadata replicas.
|
| "maxMetadataReplicas": "Number"
| Maximum number of metadata replicas.
|
| "defaultDataReplicas": "Number"
| Default number of data replicas.
|
| "maxDataReplicas": "Number"
| Maximum number of data replicas.
|
| "strictReplication": "{no | whenpossible | always }"
| Whether strict replication is to be enforced. 'nowhenpossible': enforced if at
| least two failure groups present, 'always': always enforced.
|
| "logReplicas": "Number"
| Number of log replicas.

| "quotas"
|
| "quotasAccountingEnabled": "[user] [;group] [;fileset]"
| The types of quotas for which accounting is enabled. Accounting means that usage is tracked
| against quotas. If action needs to be taken when quotas are exceeded, quotas must be also
| enforced.
|
| "quotasEnforced": "[user] [;group] [;fileset]"
| Quotas that are enforced for any of 'user', 'group' and 'fileset'.
|
| "defaultQuotasEnabled": "[user] [;group] [;fileset]"
| Default quotas that are enabled.
|
| "perfilesetQuotas": "{yes | no}"
| Specifies whether the scope of user and group quota limit checks on the individual fileset
| level rather than the entire file system.
|
| "filesetdfEnabled": "{yes | no}"
| Specifies whether the quotas are enforced for a fileset, the numbers reported by the df
| command are based on the quotas for the fileset rather than the entire file system.

```

```

| "settings"
| "blockAllocationType": "{scatter | cluster}"
| Block allocation type.
| "fileLockingSemantics": "{nfs4 | posix}"
| File locking semantics in effect.
| "numNodes": Number,
| Estimated number of nodes that will mount file system.
| "exactMTime": "{true | false}"
| true reports exact mtime values and false mtime is only periodically updated.
| "suppressATime": "{yes | no | realtime}"
| Suppress the periodic updating of the value of atime as reported by stat calls.
| "fastEAEEnabled": "{yes | no}"
| Specifies whether fast extended attributes are enabled.
| "encryption": "{yes | no}"
| Specifies whether encryption is enabled.
| "maxNumberOfInodes": Number
| Specifies the maximum number of files that can be created inside the inode space of the root
| fileset of this file system.
| "is4KAligned": "{yes | no}"
| Whether the file system is formatted to be 4K aligned.
| "rapidRepairEnabled": "{yes | no}"
| Whether to keep track of incomplete replication on an individual file block basis as opposed
| to the entire file.
| "stripeMethod": "{yes | no}"
| Algorithm that is used for striping data.
| "stripedLogs": "{yes | no}"
| Metadata logs are striped across available metadata disks.
| "ACLSemantics": "{posix | afs4 | a11 }"
| ACL semantics supported by this file system.
| "DMAPISupported": "{yes | no}"
| Whether Data Management API is enabled on this file system.

```

## Examples

- | The following example gets information about file system gpfs0.
- | Request data:
 

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0?fields=:all:'
```
- | Using the field parameter ":all:" returns entire details of the file system.
- | Response data:
  - | **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.
  - | The response code 200 indicates that the command successfully retrieved the information. Error code 400
 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0?lastId=1001"
 },
 "filesystems": [
 {
 "oid": "2",
 "uuid": "0A00641F:58982753",
 "name": "gpfs0",
 "version": "17.00",
 "type": "local",
 "createTime": "Mon Feb 06 08:35:47 2017",
 "block": {
 "pools": "system;data",
 "disks": "disk1;disk2",
 "blockSize": "262144",
 "metaDataBlockSize": "262144",
 "indirectBlockSize": "16384",
 "minFragmentSize": "8192",
 "inodeSize": "4096",
 "logFileSize": "4194304",
 "writeCacheThreshold": "0"
 },
 "mount": {
 "mountPoint": "/mnt/gpfs0",
 "automaticMountOption": "yes",
 "additionalMountOptions": "none",
 "mountPriority": "0",
 "driveLetter": "F",
 "remoteDeviceName": "gpfs0",
 "readOnly": "false"
 },
 "replication": {
 "defaultMetadataReplicas": "1",
 "maxMetadataReplicas": "2",
 "defaultDataReplicas": "1",
 "maxDataReplicas": "2",
 "strictReplication": "whenpossible",
 "logReplicas": "0"
 },
 "quota": {
 "quotasAccountingEnabled": "user;group;fileset",
 "quotasEnforced": "user;group;fileset",
 "defaultQuotasEnabled": "none",
 "perfilesetQuotas": "true",
 "filesetdfEnabled": "false"
 },
 "settings": {
 "blockAllocationType": "cluster",
 "fileLockingSemantics": "nfs4",
 "numNodes": "32",
 "exactMTime": "true",
 "suppressATime": "no",
 "fastEAEnabled": "true",
 "encryption": "false",
 "maxNumberOfInodes": "3211520",
 "is4KAligned": "true",
 "rapidRepairEnabled": "true",
 "stripeMethod": "round-robin",
 "stripedLogs": "true",
 "aclsemantics": "nfs4",
 "dmapienabled": "true"
 }
 }
]
}

```

```
| }
|]
|
| "mmcrfs command" on page 250
| Creates a GPFS file system.
| "mmchfs command" on page 179
| Changes the attributes of a GPFS file system.
| "mmlsfs command" on page 411
| Displays file system attributes.
```

## | **Filesystems/{filesystemName}/acl/{path}: GET**

| Gets information about the access control list (ACL) set for a file or directory.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/filesystemName/acl/path** request gets information about ACLs set for files or  
| directories within a particular file system. For more information about the fields in the data structures  
| that are returned, see the topics “mmgetacl command” on page 347 and “mmpputacl command” on page  
| 504.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/FileSystemName/acl/path`

| where

#### | **filesystems/filesystemName**

| The file system to which the file or directory belongs. Required.

#### | **acl/path**

| The path of the file or directory about which you want to get the ACL information. Required.

### | **Request headers**

| `Accept: application/json`

### | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 31. List of parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                   | Required/optional |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | Name of the file system.                                                                                                                                                                                                              | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields.                                                                                                                                      | Optional.         |
| path           | The file path relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to "%2F" in the request URL. | Required.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
```

```

 "message": "ReturnMessage"
},
acl:
{
 "type": "{NFSv4} ",
 "entries": [
 "type": "{allow | deny | alarm | audit }",
 "who": "User or group",
 "permissions": "Access permissions",
 "flags": "fFags",
],
}
}

| For more information about the fields in the following data structures, see the links at the end of this
| topic.

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": "ReturnCode
| The return code.

| "acl":
| An array of elements that describe ACL.

| "type": "NFSv4"
| Type of the ACL.

| "entries": "Access control entries"

| "type": "allow | deny | alarm | audit"
| Type of the entry.

| "who": " special:owner@ | special:group@ | special:everyone@ | user:{name} | group:{name}"
| The name of the user or group of users for which the ACL is applicable.

| "permission": "(r) read | (w) write | (m) mkdir, | (x) execute | (d) delete | (D) delete child | (a) read attr | (A) write attr (n) read named | (N) write Named | (c) read acl | (C) write acl | (o) change owner| (s) synchronize "
| The access permissions.

| "flags": "(f) file inherit | (d) dir inherit | (i) inherit only | (I) inherited | (S) successful access | (F) failed access"
| Special flags and inheritance definition.

```

## Examples

- | The following example gets ACL information for the file system gpfs0 and path mnt/gpfs0.
- | Request data:
 

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/acl/mnt%2Fgpfs0'
```
- | Response data:
  - | **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.
  - | The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "acl": {
| "type": "NFSv4",
| "entries": [
| {
| "type": "allow",
| "who": "user:testuser",
| "permissions": "rxancs",
| "flags": "fd"
| }
|]
| }
| }
| "mmgetacl command" on page 347
| Displays the GPFS access control list of a file or directory.
| "mmputacl command" on page 504
| Sets the GPFS access control list for the specified file or directory.
```

## | **Filesystems/{filesystemName}/acl/{path}: PUT**

| Sets access control list (ACL) for a file or directory. You can set only NFSv4 ACLs. The POSIX ACLs are not supported.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **PUT filesystems/filesystemName/acl/path** request sets ACL for files or directories within a particular file system. For more information about the fields in the data structures that are returned, see the topics “mmgetacl command” on page 347 and “mmputacl command” on page 504.

| **Note:** Only the users with *dataaccess* role can set ACL for a file or directory.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/filesystems/FileSystemName/acl/path`

| where

#### | **filesystems/filesystemName**

| The file system in which the file or directory is located. Required.

#### | **acl/path**

| The path of the file or directory for which you want to set the ACL. Required.

### | **Request headers**

| Content-Type: application/json

| Accept: application/json

### | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 32. List of parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                   | Required/optional |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                                                                                                           | Required.         |
| path           | The file path relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to "%2F" in the request URL. | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation.                                                                                        | Required.         |

## Request data

```
{
 "type": "{NFSv4} ",
 "entries": [
 {"type": "{allow | deny | alarm | audit }",
 "who": "User or group",
 "permissions": "Access permissions",
 "flags": "Flags",
],
}
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"type": "NFSv4"**  
| Type of the ACL.

| **"entries": "Access control entries"**

| **"type": "allow | deny | alarm | audit"**  
| Type of the entry.

| **"who": " special:owner@ | special:group@ | special:everyone@ | user:{name} | group:{name}"**  
| The name of the user or group of users for which the ACL is applicable.

| **"permission": "(r) read | (w) write | (m) mkdir, | (x) execute | (d) delete | (D) delete child | (a) read attr | (A) write attr (n) read named | (N) write Named | (c) read acl | (C) write acl | (o) change owner| (s) synchronize "**  
| The access permissions.

| **"flags": "(f) file inherit | (d) dir inherit | (i) inherit only | (I) inherited | (S) successful access | (F) failed access"**  
| Special flags and inheritance definition.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": "",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

- For more information about the fields in the following data structures, see the links at the end of this topic.
- "jobs":**  
An array of elements that describe jobs. Each element describes one job.
- "status":**  
Return status.
- "message": "ReturnMessage",**  
The return message.
- "code": *ReturnCode***  
The return code.
- "result"**
- "commands": "String"**  
Array of commands that are run in this job.
- "progress": "String"**  
Progress information for the request.
- "exitCode": "Exit code"**  
Exit code of command. Zero is success, nonzero denotes failure.
- "stderr": "Error"**  
CLI messages from *stderr*.
- "stdout": "String"**  
CLI messages from *stdout*.
- "request"**
- "type": "{GET | POST | PUT | DELETE}"**  
HTTP request type.
- "url": "URL"**  
The URL through which the job is submitted.
- "data": "**  
Optional.
- "jobId": "ID",**  
The unique ID of the job.
- "submitted": "Time"**  
The time at which the job was submitted.
- "completed": "Time"**  
The time at which the job was completed.
- "status": "RUNNING | COMPLETED | FAILED"**  
Status of the job.

## Examples

- The following example sets ACL information for the file system *gpfs0* and path *mnt/gpfs0*.

- Request data:

```
{
 "type": "NFSv4",
 "entries": [
 {
 "type": "allow",
 "who": "user:testuser",
 }
]
}
```

```

 "permissions": "rxancs",
 "flags": "fd"
 }
]
}

```

Corresponding request URL:

```

curl -k -u admin:admin001 -X PUT --header 'content-type:application/json' --header 'accept:application/json' -d '{
 "type": "NFSv4",
 "entries": [
 {
 "type" : "allow",
 "who" : "special:owner@",
 "permissions" : "rwxmDaAnNcCos",
 "flags" : ""
 },
 {
 "type" : "allow",
 "who" : "special:group@",
 "permissions" : "rxancs",
 "flags" : ""
 },
 {
 "type" : "allow",
 "who" : "special:everyone@",
 "permissions" : "rxancs",
 "flags" : ""
 },
 {
 "type" : "allow",
 "who" : "user:scalemgmt",
 "permissions" : "rxancs",
 "flags" : "fd"
 }
]
}' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/acl/mnt%2Fgpfs0'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "jobs" : [{
 "jobId" : 1000000000002,
 "status" : "RUNNING",
 "submitted" : "2017-03-14 15:50:00,493",
 "completed" : "N/A",
 "request" : {
 "data" : {
 "entries" : [
 {
 "type" : "allow",
 "who" : "special:owner@",
 "permissions" : "rwxmDaAnNcCos",
 "flags" : ""
 },
 {
 "type" : "allow",
 "who" : "special:group@",
 "permissions" : "rxancs",
 "flags" : ""
 },
 {
 "type" : "allow",
 "who" :

```

```
| "who" : "special:everyone@",
| "permissions" : "rxancs",
| "flags" : ""
| },
| {
| "type" : "allow",
| "who" : "user:scalemgmt",
| "permissions" : "rxancs",
| "flags" : "fd"
| }],
| "type" : "NFSv4"
| },
| "type" : "PUT",
| "url" : "/scalemgmt/v2/filesystems/gpfs0/acl/mnt%2Fgpfs0"
|,
| "result" : { }
|],
| "status" : {
| "code" : 202,
| "message" : "The request was accepted for processing"
| }
```

“mmgetacl command” on page 347

Displays the GPFS access control list of a file or directory.

“mmpputacl command” on page 504

Sets the GPFS access control list for the specified file or directory.

## **Filesystems/{filesystemName}/disks: GET**

| Gets details of the disks that are part of a file system.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **GET filesystems/{filesystemName}/disks** request gets details of the disks that belong to the specified file system. For more information about the fields in the data structures that are returned, see the topics “mmlsdisk command” on page 402, “mmlsnsd command” on page 423 and “mmlsfs command” on page 411.

### **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/disks`

| where

#### **filesystems/filesystemName**

| Specifies the file system to which the disks belong. Required.

#### **disks**

| The disks that belong to the specified file system. Required.

### **Request headers**

| `Accept: application/json`

### **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 33. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |

### **Request data**

| No request data.

### **Response data**

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| }
```

```

},
"paging":
{
 "next": "URL"
},
disks: [
{
 "name": "DiskName",
 "filesystem": "FileSystemName",
 "failureGroup": "FailureGroupID",
 "type": "{dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}",
 "storagePool": "StoragePool",
 "status": "Status",
 "availability": "Availability",
 "quorumdisk": "QuorumDisk",
 "remarks": "Remarks",
 "size": "Size",
 "availableBlocks": "AvailableBlocks",
 "availableFragments": "AvailableFragments",
 "nsdServers": "NSDServers",
 "nsdVolumeId": "NSDVolumeID",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**  
|     Return status.

| **"message": "ReturnMessage",**  
|     The return message.

| **"code": *ReturnCode***  
|     The return code.

| **"paging"**  
|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by  
|     the query.

| **"disks":**  
|     An array of elements that describe one disk.

|     **"name": "DiskName"**  
|         Name of the disk.

|     **"filesystem": "FileSystemName"**  
|         The file system to which the disk belongs.

|     **"failureGroup": "FailureGroupID"**  
|         Failure group of the disk.

|     **"type": "{dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}"**  
|         Specifies the type of data to be stored on the disk:

|         **dataOnly**  
|             Indicates that the disk contains data and does not contain metadata. This is the default  
|             for disks in storage pools other than the system pool.

|         **metadataOnly**  
|             Indicates that the disk contains metadata and does not contain data.

**dataAndMetadata**  
 Indicates that the disk contains both data and metadata. This is the default for disks in the system pool.

**descOnly**  
 Indicates that the disk contains no data and no file metadata. Such a disk is used solely to keep a copy of the file system descriptor, and can be used as a third failure group in certain disaster recovery configurations. For more information, see the help topic *Synchronous mirroring utilizing GPFS replication* in *IBM Spectrum Scale: Administration Guide*.

**localCache**  
 Indicates that the disk is to be used as a local read-only cache device.

**"storagePool": "StoragePool"**  
 Storage pool to which the disk belongs.

**"status": "Status"**  
 The status of the disk.

**"availability": "Availability"**  
 The availability of the disk.

**"quorumdisk": "Quorum disk"**  
 The quorum status of the disk.

**"remarks": "Remarks"**  
 User-defined remarks about the disk.

**"size": "Size"**  
 The size of the disk.

**"availableBlocks": "Available blocks"**  
 The available blocks of the disk.

**"availableFragments": "Available fragments"**  
 The available fragments of the disk.

**"nsdServers": "NSD servers"**  
 The name of the NSD servers.

**"nsdVolumeID": "NSD volume ID"**  
 The volume ID of the NSD.

## Examples

- The following example gets information about disks in the file system gpfso.
  - Request data:
 

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfso/disks'
```
  - The request URL with no field or filter parameter returns only the details that uniquely identify the object.
  - Response data:
    - Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.
    - The **disks** array returns details of the disks. Each **config** object contains details about one disk.

```

{
 "disks" : [{
 "config" : {
 "fileSystem" : "gpfs0",
 "name" : "disk1"
 }
 }, {
 "config" : {
 "fileSystem" : "gpfs0",
 "name" : "disk8"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}

```

| Using the field parameter ":all:" returns entire details of the disks that are available in the system.

```

curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/disks?fields=:all:'
{
 "disks" : [{
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.15 GB",
 "availableFragments" : "552 KB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs0",
 "name" : "disk1",
 "nsdServers" : "mari-11.localnet.com,mari-15.localnet.com,mari-14.localnet.com,
 mari-13.localnet.com,mari-12.localnet.com",
 "nsdVolumeId" : "0A00640B58B82A8C",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "system",
 "type" : "nsd"
 }
 }, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.45 GB",
 "availableFragments" : "664 KB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs0",
 "name" : "disk8",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AD9",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}

```

- | “[mmlsdisk command](#)” on page 402
- | Displays the current configuration and state of the disks in a file system.
- | “[mmcrnsd command](#)” on page 262
- | Creates Network Shared Disks (NSDs) used by GPFS.
- | “[mmcrfs command](#)” on page 250
- | Creates a GPFS file system.
- | “[mmadddisk command](#)” on page 23
- | Adds disks to a GPFS file system.

## | **Filesystems/{filesystemName}/disks/{diskName}: GET**

| Gets details of a specific disk that is part of a file system.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/{filesystemName}/disks/{diskName}** request gets details of a specific disk. For more information about the fields in the data structures that are returned, see the topics “**mmlsdisk command**” on page 402, “**mmlsnrd command**” on page 423 and “**mmlsfs command**” on page 411.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/fileSystemName/disks/diskName`

| where

| **filesystems/filesystemName**

| The file system to which the disks belong. Required.

| **disks/diskName**

| The disk about which you want to get information. Required.

### | **Request headers**

| `Accept: application/json`

### | **Parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 34. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| diskName       | The name of the disk.                                                                            | Required.         |

### | **Request data**

| No request data.

### | **Response data**

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "paging":
 {
 "next": "URL"
 },
}
```

```
 "disks": [
 {
 "name": "DiskName",
 "filesystem": "FileSystemName",
 "failureGroup": "FailureGroupID",
 "type": "{dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}",
 "storagePool": "StoragePool",
 "status": "Status",
 "availability": "Availability",
 "quorumdisk": "QuorumDisk",
 "remarks": "Remarks",
 "size": "Size",
 "availableBlocks": "AvailableBlocks",
 "availableFragments": "AvailableFragments"
 },
 {
 "nsdServers": "NSDServers",
 "nsdVolumeId": "NSDVolumeID",
 }
],
 }
}
```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

```
"status":
 Return status.

"message": "ReturnMessage",
 The return message.

"code": ReturnCode
 The return code.

"paging"
 The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by
the query.

"disks":
 An array of elements that describe one disk.

 "name": "DiskName"
 Name of the disk.

 "filesystem": "FileSystemName"
 The file system to which the disk belongs.

 "failureGroup": "FailureGroupID"
 Failure group of the disk.

 "type": "{dataOnly | metadataOnly | dataAndMetadata | descOnly | localCache}"
 Specifies the type of data to be stored on the disk:

 dataOnly
 Indicates that the disk contains data and does not contain metadata. This is the default
for disks in storage pools other than the system pool.

 metadataOnly
 Indicates that the disk contains metadata and does not contain data.

 dataAndMetadata
 Indicates that the disk contains both data and metadata. This is the default for disks in
the system pool.

 descOnly
 Indicates that the disk contains no data and no file metadata. Such a disk is used solely
to keep a copy of the file system descriptor, and can be used as a third failure group in
```

certain disaster recovery configurations. For more information, see the help topic *Synchronous mirroring utilizing GPFS replication* in *IBM Spectrum Scale: Administration Guide*.

**localCache**  
 Indicates that the disk is to be used as a local read-only cache device.

**"storagePool": "StoragePool"**  
 Storage pool to which the disk belongs.

**"status": "Status"**  
 The status of the disk.

**"availability": "Availability"**  
 The availability of the disk.

**"quorundisk": "Quorum disk"**  
 The quorum status of the disk.

**"remarks": "Remarks"**  
 User-defined remarks about the disk.

**"size": "Size"**  
 The size of the disk.

**"availableBlocks": "Available blocks"**  
 The available blocks of the disk.

**"availableFragments": "Available fragments"**  
 The available fragments of the disk.

**"nsdServers": "NSD servers"**  
 The name of the NSD servers.

**"nsdVolumeID": "NSD volume ID"**  
 The volume ID of the NSD.

## Examples

The following example gets information about the disk `disk8` in the file system `gpfs0`.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'

'https://198.51.100.1:443/scaladmgt/v2/filesystems/gpfs0/disks/disk8'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://198.51.100.1:443/scaladmgt/v2/filesystems/gpfs0/disks/disk8?lastId=1001"
 },
 "disks": [
 {
 "config": {
 "name": "disk8",
 "fileSystem": "gpfs0",
 "failureGroup": "1",
 "type": "dataOnly",
 "localCache": true
 }
 }
]
}
```

```
 "storagePool": "data",
 "status": "ready",
 "availability": "up",
 "quorumDisk": "no",
 "remarks": "This is a comment",
 "size": "10.00 GB",
 "availableBlocks": "730.50 MB",
 "availableFragments": "1.50 MB",
 "nsdServers": "gpfsgui-21.localnet.com",
 "nsdVolumeId": "0A0064155874F5AA"
 }
}
]
}

"mmlsdisk command" on page 402
Displays the current configuration and state of the disks in a file system.

"mmcrnsd command" on page 262
Creates Network Shared Disks (NSDs) used by GPFS.

"mmcrfs command" on page 250
Creates a GPFS file system.

"mmadddisk command" on page 23
Adds disks to a GPFS file system.
```

## | **Filesystems/{filesystemName}/filesets: GET**

| Gets information of all filesets that are part of a specific file system.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/{filesystemName}/filesets** request gets information about all filesets in the specified file system. For more information about the fields in the data structures that are returned, see the topics “mmcrfileset command” on page 244, “mmchfileset command” on page 173, and “mmlsfileset command” on page 406.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/filesystems/{filesystemName}/filesets`

| where:

#### | **filesystem/{filesystemName}**

| Specifies the name of the file system to which the filesets belong. Required.

#### | **filesets**

| Specifies that you need to get details of all filesets that belong to the specified file system. Required.

### | **Request headers**

| `Accept: application/json`

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 35. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |

### | **Request data**

| No request data.

### | **Response data**

```
| { "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
```

```

| "paging":
| {
| "next": "URL"
| },
| "filesets": [
| {
| "config":
| {
| "filesetName": "Fileset",
| "filesystemName": "Device",
| "path": "Path",
| "inodeSpace": "Inodes",
| "maxNumInodes": "Inodes",
| "permissionChangeMode": "Mode",
| "comment": "Comment",
| "iamMode": "Mode",
| "oid": "ID",
| "id": "ID",
| "status": "Status"
| "parentId": "ID",
| "created": "DateTime",
| "isInodeSpaceOwner": "isOwner",
| "inodeSpaceMask": "inodeSpace",
| "snapID": "ID",
| "rootInode": "Inode",
| }
| }
| "afm": {
| "afmTarget": "Protocol://{Host | Map}/Path",
| "afmAsyncDelay": "Delay",
| "afmDirLookupRefreshInterval": "Interval",
| "afmDirOpenRefreshInterval": "Interval",
| "afmEnableAutoEviction": "{yes | no}",
| "afmExpirationTimeout": "Interval",
| "afmFileLookupRefreshInterval": "Interval",
| "afmMode": "{single-writer | sw | read-only | ro | local-updates | 1u | independent-writer | Primary | Secondary}",
| "afmNumFlushThreads": "Threads",
| "afmParallelReadChunkSize": "Size",
| "afmParallelReadThreshold": "Threshold",
| "afmParallelWriteChunkSize": "Size",
| "afmParallelWriteThreshold": "Threshold",
| "afmPrefetchThreshold": "{1 | 1-99 | 100}",
| "afmPrimaryID": "ID",
| "afmRPO": "Interval",
| "afmShowHomeSnapshots": "{yes | no}",
| }
| }
|],
| }
| }

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": "ReturnCode"
| The return code.

| "paging"
| The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

| "filesets":
| An array of information about the filesets in the specified file system. Each array element describes

```

one fileset and can contain the data structures **afm**, **config**, **links**, and **state**. For more information about the fields in these data structures, see the links at the end of this topic.

**"filesets":**

Information about the fileset configuration.

**"filesetName": "Fileset",**

The name of the fileset.

**"filesystemName": "Device"**

Required. The file system in which the fileset is located.

**"path": "Path"**

The absolute path of the fileset.

**"inodeSpace": "Inodes"**

The number of inodes that are allocated for use by the fileset.

**"maxNumInodes": "Inodes"**

The inode limit for the inode space that is owned by the specified fileset.

**"permissionChangeMode": "Mode"**

The permission change mode. Controls how **chmod** and ACL commands affect objects in the fileset.

**chmodOnly**

Only the **chmod** command can change access permissions.

**setAclOnly**

Only the ACL commands and API can change access permissions.

**chmodAndSetAcl**

Both the **chmod** command and ACL commands can change access permissions.

**chmodAndUpdateAcl**

Both the **chmod** command and ACL commands can change access permissions.

**"comment": "Comment",**

A comment that appears in the output of the **mmlsfileset** command.

**"iamMode": "Mode"**

The integrated archive manager (IAM) mode for the fileset.

**ad | advisory**

**nc | noncompliant**

**co | compliant**

**"oid": "ID"**

Internal identifier of the fileset that is used for paging.

**"id": ID**

The fileset identifier.

**"status": "Status"**

Specifies whether the file is linked or unlinked.

**"parentId": ParentID**

The parent identifier of the fileset.

**"created": "DateTime",**

The date and time when the fileset was created.

**"isInodeSpaceOwner": "isOwner",**

Indicates whether the fileset has its own inode space.

```

| "isInodeSpaceMask": "isMask",
| The inode space mask.
|
| "snapID": ID
| The snapshot ID.
|
| "afm":
| Information about Active File Management (AFM).
|
| "afmTarget": "Interval"Protocol://{Host | Map}/Path
| The home that is associated with the cache.
|
| "afmAsyncDelay": "Delay"
| The time in seconds by which to delay write operations because of the lag in updating
| remote clusters.
|
| "afmDirLookupRefreshInterval": "Interval"
| The interval in seconds between data revalidations of directories that are caused by lookup
| operations such as ls and stat.
|
| "afmDirOpenRefreshInterval": "Interval"
| The interval in seconds between data revalidations of directories that are caused by lookup
| operations.
|
| "afmEnableAutoEviction": "{yes | no}"
| Enables or disables eviction on the fileset.
|
| "afmExpirationTimeout": "Interval"
| The timeout in seconds after which cached data is considered expired. Used with
| afmDisconnectTimeout, which is set with mmchconfig.
|
| "afmFileLookupRefreshInterval": "Interval"
| The interval in seconds between data revalidations of files that are caused by lookup
| operations such as ls and stat.
|
| "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
| independent-writer | Primary | Secondary}"
| The mode in which the cache operates.
|
| "afmNumFlushThreads": numThreads
| The number of threads that are used on each gateway to synchronize updates to the home
| cluster.
|
| "afmParallelReadChunkSize": Size
| The minimum chunk size of read data, in bytes, that must be distributed among the gateway
| nodes during parallel reads.
|
| "afmParallelReadThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel reads.
|
| "afmParallelWriteChunkSize": Size
| The minimum chunk size of write data, in bytes, that must be distributed among the gateway
| nodes during parallel writes.
|
| "afmParallelWriteThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel writes.
|
| "afmPrefetchThreshold": {1 | 1-99 | 100}
| The percentage of file size that must be cached before the entire file is prefetched.
|
| 0 Enables full file prefetching.
|
| 1-99 The percentage of file size that must be cached before the entire file is prefetched.
|
| 100 Disables full file prefetching.

```

```

| "afmPrimaryID": "ID"
| The unique primary ID of the primary fileset for asynchronous data replication.
| "afmRPO": Interval
| The recovery point objective (RPO) interval, in minutes, for a primary fileset.
| "afmShowHomeSnapshots": "{yes | no}"
| Shows or hides the home snapshot directory in cache. Specify yes to show, no to hide.

```

## Examples

- | The following example gets information about the filesets in file system gpfs0.
- | Request data:
 

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets'
```
- | The request URL with no field or filter parameter returns only the details that uniquely identify the object.
- | Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "filesets": [{
 "config": {
 "filesetName": "root",
 "filesystemName": "gpfs0"
 }
 }, {
 "config": {
 "filesetName": "fset1",
 "filesystemName": "gpfs0"
 }
 }, {
 "config": {
 "filesetName": "fset2",
 "filesystemName": "gpfs0"
 }
 }],
 "status": {
 "code": 200,
 "message": "The request finished successfully"
 }
}
```

- | Using the field parameter ":all:" returns entire details of the filesets that are part of the specified file system. For example:

```

| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets?fields=:all:'
{
 "filesets": [{
 "afm": {
 "afmAsyncDelay": 0,
 "afmDirLookupRefreshInterval": 0,
 "afmDirOpenRefreshInterval": 0,
 "afmEnableAutoEviction": false,
 "afmExpirationTimeout": 0,
 "afmFileLookupRefreshInterval": 0,
 "afmMode": ""
 }
 }]
}
```

```

 "afmNumFlushThreads" : 0,
 "afmParallelReadChunkSize" : 0,
 "afmParallelReadThreshold" : 0,
 "afmParallelWriteChunkSize" : 0,
 "afmParallelWriteThreshold" : 0,
 "afmPrefetchThreshold" : 0,
 "afmPrimaryID" : "",
 "afmRPO" : 0,
 "afmShowHomeSnapshots" : false
 },
 "config" : {
 "comment" : "root fileset",
 "created" : "2017-03-02 15:23:53,000",
 "filesetName" : "root",
 "filesystemName" : "gpfs0",
 "iamMode" : "off",
 "id" : 0,
 "inodeSpace" : 0,
 "inodeSpaceMask" : 0,
 "isInodeSpaceOwner" : true,
 "maxNumInodes" : 65792,
 "oid" : 1,
 "path" : "/mnt/gpfs0",
 "permissionChangeMode" : "chmodAndSetacl",
 "rootInode" : 3,
 "snapId" : 0,
 "status" : "Linked"
 }
}, {
 "afm" : {
 "afmAsyncDelay" : 0,
 "afmDirLookupRefreshInterval" : 0,
 "afmDirOpenRefreshInterval" : 0,
 "afmEnableAutoEviction" : false,
 "afmExpirationTimeout" : 0,
 "afmFileLookupRefreshInterval" : 0,
 "afmMode" : "",
 "afmNumFlushThreads" : 0,
 "afmParallelReadChunkSize" : 0,
 "afmParallelReadThreshold" : 0,
 "afmParallelWriteChunkSize" : 0,
 "afmParallelWriteThreshold" : 0,
 "afmPrefetchThreshold" : 0,
 "afmPrimaryID" : "",
 "afmRPO" : 0,
 "afmShowHomeSnapshots" : false
 },
 "config" : {
 "comment" : "",
 "created" : "2017-03-13 10:49:38,000",
 "filesetName" : "fset1",
 "filesystemName" : "gpfs0",
 "iamMode" : "off",
 "id" : 1,
 "inodeSpace" : 1,
 "inodeSpaceMask" : 0,
 "isInodeSpaceOwner" : true,
 "maxNumInodes" : 1048576,
 "oid" : 8,
 "parentId" : 0,
 "path" : "/mnt/gpfs0/fset1",
 "permissionChangeMode" : "chmodAndSetacl",
 "rootInode" : 131075,
 "snapId" : 0,
 "status" : "Linked"
 }
}, {
}

```

```

 "afm" : {
 "afmAsyncDelay" : 0,
 "afmDirLookupRefreshInterval" : 0,
 "afmDirOpenRefreshInterval" : 0,
 "afmEnableAutoEviction" : false,
 "afmExpirationTimeout" : 0,
 "afmFileLookupRefreshInterval" : 0,
 "afmMode" : "",
 "afmNumFlushThreads" : 0,
 "afmParallelReadChunkSize" : 0,
 "afmParallelReadThreshold" : 0,
 "afmParallelWriteChunkSize" : 0,
 "afmParallelWriteThreshold" : 0,
 "afmPrefetchThreshold" : 0,
 "afmPrimaryID" : "",
 "afmRPO" : 0,
 "afmShowHomeSnapshots" : false
 },
 "config" : {
 "comment" : "My second fileset!",
 "created" : "2017-03-13 11:02:02,000",
 "filesetName" : "fset2",
 "filesystemName" : "gpfs0",
 "iamMode" : "off",
 "id" : 2,
 "inodeSpace" : 2,
 "inodeSpaceMask" : 0,
 "isInodeSpaceOwner" : true,
 "maxNumInodes" : 633280,
 "oid" : 10,
 "parentId" : 0,
 "path" : "/mnt/gpfs0/fset2",
 "permissionChangeMode" : "chmodAndSetacl",
 "rootInode" : 262147,
 "snapId" : 0,
 "status" : "Linked"
 }
},
"status" : {
 "code" : 200,
 "message" : "The request finished successfully"
}
}

```

[“mmcrfileset command” on page 244](#)

Creates a GPFS fileset.

[“mmchfileset command” on page 173](#)

Changes the attributes of a GPFS fileset.

[“mmlsfileset command” on page 406](#)

Displays attributes and status for GPFS filesets.

## | **Filesystems/{filesystemName}/filesets: POST**

| Creates a fileset.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The POST **filesystems/{filesystemName}/filesets** command creates a new fileset in the specified file system. For more information about the fields in the request data structures, see the topics “mmcrfileset command” on page 244, “mmchfileset command” on page 173, and “mmcrfs command” on page 250.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/{filesystemName}/filesets`

| where

#### | **filesystem/filesetName**

| Specifies the file system in which the fileset is to be created. Required.

#### | **filesets**

| Specifies fileset as the resource. Required.

### | **Request headers**

| `Accept: application/json`

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 36. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                           | Required/optional |
|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                   | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to theIBM Spectrum Scale system to perform the requested operation. | Required.         |

### | **Request data**

```
{
 "filesetName": "Fileset name",
 "path": "Path",
 "owner": "UserID[:GroupID]",
 "permissions": "Permissions"
 "inodeSpace": "{new | ExistingFileset}",
 "maxNumInodes": "Inodes",
 "allocInodes": "Inodes",
 "permissionChangeMode": "chmodOnly | setAclOnly | chmodAndSetAcl",
 "comment": "Comment",
 "iamMode": "{advisory | noncompliant | compliant}",
 "afmTarget": "Protocol://[{Host | Map}]/Path"
 "afmAsyncDelay": "SecondsDelay",
 "afmDirLookupRefreshInterval": "SecondsInterval",
```

```

 "afmDirOpenRefreshInterval": "SecondsInterval",
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": "AFM expiration timeout",
 "afmFileLookupRefreshInterval": Seconds Interval,
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
 independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": NumThreads,
 "afmParallelReadChunkSize": NumBytes,
 "afmParallelReadThreshold": NumMegabytes,
 "afmParallelWriteChunkSize": NumBytes,
 "afmParallelWriteThreshold": NumMegabytes,
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": MinutesInterval,
 "afmShowHomeSnapshots": "{yes | no}",
}

```

| Details of the parameters are given in the following list:

- | **"filesetName": "Fileset"**  
| The name of the fileset.
- | **"path": "Path"**  
| The path the fileset will be linked to. If not specified the fileset will be linked to the default path, for example, /mnt/gpfs0/fset1.
- | **"owner": "UserID[:GroupID]"**  
| The user ID of the owner of the new fileset, followed optionally by the group ID, as in "root".
- | **"permissions": "Integer"**  
| The permissions that are set using **chmod**. If not specified, a default ACL will be applied.
- | **"inodeSpace": "Inodes"**  
| The number of inodes that are allocated for use by the fileset.
- | **"maxNumInodes": "Inodes"**  
| The inode limit for the inode space that is owned by the specified fileset.
- | **"permissionChangeMode": "Mode"**  
| The permission change mode. Controls how **chmod** and ACL commands affect objects in the fileset.
  - | **chmodOnly**  
Only the **chmod** command can change access permissions.
  - | **setAclOnly**  
Only the ACL commands and API can change access permissions.
  - | **chmodAndSetAcl**  
Both the **chmod** command and ACL commands can change access permissions.
  - | **chmodAndUpdateAcl**  
Both the **chmod** command and ACL commands can change access permissions.
- | **"comment": "Comment"**,  
| A comment that appears in the output of the **mmlsfileset** command.
- | **"iamMode": "Mode"**  
| The integrated archive manager (IAM) mode for the fileset.
  - | **advisory**
  - | **noncompliant**
  - | **compliant**
- | **"afmTarget": "Home fileset"**  
| The home that is associated with the cache.

- | "afmAsyncDelay": *Delay*  
|     The time in seconds by which to delay write operations because of the lag in updating remote clusters.
- | "afmDirLookupRefreshInterval": *Interval*  
|     The interval in seconds between data revalidations of directories that are caused by lookup operations such as `ls` and `stat`.
- | "afmDirOpenRefreshInterval": *Interval*  
|     The interval in seconds between data revalidations of directories that are caused by lookup operations.
- | "afmEnableAutoEviction": "{**yes** | **no**}"  
|     Enables or disables eviction on the fileset.
- | "afmExpirationTimeout": *Interval*  
|     The timeout in seconds after which cached data is considered expired. Used with `afmDisconnectTimeout`, which is set with `mmchconfig`.
- | "afmFileLookupRefreshInterval": *Interval*  
|     The interval in seconds between data revalidations of files that are caused by lookup operations such as `ls` and `stat`.
- | "afmMode": "{**single-writer** | **sw** | **read-only** | **ro** | **local-updates** | **lu** | **independent-writer** | **Primary** | **Secondary**}"  
|     The mode in which the cache operates.
- | "afmNumFlushThreads": *numThreads*  
|     The number of threads that are used on each gateway to synchronize updates to the home cluster.
- | "afmParallelReadChunkSize": *Size*  
|     The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes during parallel reads.
- | "afmParallelReadThreshold": *Threshold*  
|     The minimum file size, in megabytes, at which to do parallel reads.
- | "afmParallelWriteChunkSize": *Size*  
|     The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes during parallel writes.
- | "afmParallelWriteThreshold": *Threshold*  
|     The minimum file size, in megabytes, at which to do parallel writes.
- | "afmPrefetchThreshold": {**1** | **1-99** | **100**}  
|     The percentage of file size that must be cached before the entire file is prefetched.
  - | **0**     Enables full file prefetching.
  - | **1-99**     The percentage of file size that must be cached before the entire file is prefetched.
  - | **100**     Disables full file prefetching.
- | "afmPrimaryID": "*ID*"  
|     The unique primary ID of the primary fileset for asynchronous data replication.
- | "afmRPO": *Interval*  
|     The recovery point objective (RPO) interval, in minutes, for a primary fileset.
- | "afmShowHomeSnapshots": "{**yes** | **no**}"  
|     Shows or hides the home snapshot directory in cache. Specify **yes** to show, **no** to hide.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

### "jobs":

An array of elements that describe jobs. Each element describes one job.

#### "status":

Return status.

#### "message": "ReturnMessage",

The return message.

#### "code": "ReturnCode"

The return code.

#### "result"

##### "commands": "String"

Array of commands that are run in this job.

##### "progress": "String"

Progress information for the request.

##### "exitCode": "Exit code"

Exit code of command. Zero is success, nonzero denotes failure.

##### "stderr": "Error"

CLI messages from stderr.

##### "stdout": "String"

CLI messages from stdout.

#### "request"

##### "type": "{GET | POST | PUT | DELETE}"

HTTP request type.

```

| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example creates and links a fileset fset1 in file system gpfs0.

Request data:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "filesetName": "fset1",
 "path": "/mnt/gpfs0/fset1",
 "owner": "scalemgmt",
 "permissions": "700"
}' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets'
```

Response data:

```

{
 "jobs" : [{
 "jobId" : 1000000000003,
 "status" : "COMPLETED",
 "submitted" : "2017-03-14 15:54:23,042",
 "completed" : "N/A",
 "request" : {
 "data" : {
 "filesetName" : "fset1",
 "owner" : "scalemgmt",
 "path" : "/mnt/gpfs0/fset1",
 "permissions" : 700
 },
 "type" : "POST",
 "url" : "/scalemgmt/v2/filesystems/gpfs0/filesets"
 },
 "result" : { }
 }],
 "status" : {
 "code" : 202,
 "message" : "The request was accepted for processing"
 }
}
```

[“mmcrfileset command” on page 244](#)

Creates a GPFS fileset.

[“mmchfileset command” on page 173](#)

Changes the attributes of a GPFS fileset.

[“mmlsfileset command” on page 406](#)

Displays attributes and status for GPFS filesets.

## | **Filesystems/{filesystemName}/filesets/{filesetName}: DELETE**

| Deletes a specific fileset.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **DELETE filesystems/{filesystemName}/filesets/{filesetName}** command deletes the specified fileset.  
| For more information on deleting snapshots, see “[mmdelfileset command](#)” on page 292.

### | **Request URL**

| Use this URL to delete a global snapshot:

| `https://management API host:port/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName`

| where:

| **filesystems/filesystemName**

|     Specifies the file system to which the fileset belongs. Required.

| **filesets/filesetName**

|     Specifies the fileset to be deleted. Required.

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 37. List of request parameters*

| Parameter name | Description and applicable keywords                                                         | Required/optional |
|----------------|---------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote: | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                          | Required.         |
| qosClass       | The QoS class in which the operation is performed.                                          | Optional.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
| {
| "result": "",
| {
| "commands": "String",
| "progress": "String",
| "exitCode": "Exit code",
| }
| }
|]
| }
```

```

 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**  
|     Return status.

| **"message": "ReturnMessage",**  
|     The return message.

| **"code": *ReturnCode***  
|     The return code.

| **"jobs":**  
|     An array of elements that describe jobs. Each element describes one job.

| **"result"**

|     **"commands": "String'**  
|         Array of commands that are run in this job.

|     **"progress": "String'**  
|         Progress information for the request.

|     **"exitCode": "Exit code"**  
|         Exit code of command. Zero is success, nonzero denotes failure.

|     **"stderr": "Error"**  
|         CLI messages from *stderr*.

|     **"stdout": "String"**  
|         CLI messages from *stdout*.

| **"request"**

|     **"type": "{GET | POST | PUT | DELETE}"**  
|         HTTP request type.

|     **"url": "URL"**  
|         The URL through which the job is submitted.

|     **"data": " "**  
|         Optional.

|     **"jobId": "ID",**  
|         The unique ID of the job.

|     **"submitted": "Time"**  
|         The time at which the job was submitted.

```

| "completed":Time"
| The time at which the job was completed.
| "status":"RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example deletes the fileset *myFset1* from the file system *gpfs0*.

Request data:

```

| curl -k -u admin:admin001 -X DELETE --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {
 "commands": "[['mmcrfileset gpfs0 restfs1001', ...],",
 "progress": "[('2/3) Linking fileset')",
 "exitCode": "0",
 "stderr": "[['EFSSG0740C There are not enough resources available to create a
 new independent file set.', ...]]",
 "stdout": "[['EFSSG4172I The file set {0} must be independent.', ...]]"
 },
 "request": {
 "type": "DELETE",
 "url": "/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1",
 "data": "",
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}

```

## See also

- “[mmdelfileset command](#)” on page 292

## **Filesystems/{filesystemName}/filesets/{filesetName}: GET**

| Gets information about a fileset.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **GET filesystems/{filesystemName}/filesets/{filesetName}** request gets information about the specified fileset in a file system. For more information about the fields in the data structures that are returned, see the topics “[mmcrfileset command](#)” on page 244, “[mmchfileset command](#)” on page 173, and “[mmlsfileset command](#)” on page 406.

### **Request URL**

| `https://<IP or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName`

| where:

#### **filesystem/filesystemName**

| Specifies the name of the file system to which the filesets belong. Required.

#### **filesets/filesetName**

| Specifies the fileset about which you need information. Required.

### **Request headers**

| `Accept: application/json`

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 38. List of parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                               | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |

### **Request data**

| No request data.

### **Response data**

```
| { "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
| "paging":
| {
| "next": "URL"
| }
```

```

 },
 "filesets": [
 {
 "config": {
 "filesetName": "Fileset",
 "filesystemName": "Device",
 "path": "Path",
 "inodeSpace": "Inodes",
 "maxNumInodes": "Inodes",
 "permissionChangeMode": "Mode",
 "comment": "Comment",
 "iamMode": "Mode",
 "oid": "ID",
 "id": "ID",
 "status": "Status"
 "parentId": "ID",
 "created": "DateTime",
 "isInodeSpaceOwner": "isOwner",
 "inodeSpaceMask": "inodeSpace",
 "snapID": "ID",
 "rootInode": "Inode",
 }
 },
 "afm": {
 "afmTarget": "Protocol://{Host | Map}/Path",
 "afmAsyncDelay": "Delay",
 "afmDirLookupRefreshInterval": "Interval",
 "afmDirOpenRefreshInterval": "Interval",
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": "Interval",
 "afmFileLookupRefreshInterval": "Interval",
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": "Threads",
 "afmParallelReadChunkSize": "Size",
 "afmParallelReadThreshold": "Threshold",
 "afmParallelWriteChunkSize": "Size",
 "afmParallelWriteThreshold": "Threshold",
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": "Interval",
 "afmShowHomeSnapshots": "{yes | no}",
 }
 }
],
}
}

"status":
Return status.

"message": "ReturnMessage",
The return message.

"code": "ReturnCode"
The return code.

"paging":
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

"filesets":
An array of information about the filesets in the specified file system. Each array element describes one fileset and can contain the data structures afm, config, links, and state. For more information about the fields in these data structures, see the links at the end of this topic:
```

```

| "filesets":
| Information about the fileset configuration.
|
| "filesetName": "Fileset",
| The name of the fileset.
|
| "filesystemName": "Device"
| Required. The file system in which the fileset is located.
|
| "path": "Path"
| The absolute path of the fileset.
|
| "inodeSpace": "Inodes"
| The number of inodes that are allocated for use by the fileset.
|
| "maxNumInodes": "Inodes"
| The inode limit for the inode space that is owned by the specified fileset.
|
| "permissionChangeMode": "Mode"
| The permission change mode. Controls how chmod and ACL commands affect objects in the
| fileset.
|
| chmodOnly
| Only the chmod command can change access permissions.
|
| setAclOnly
| Only the ACL commands and API can change access permissions.
|
| chmodAndSetAcl
| Both the chmod command and ACL commands can change access permissions.
|
| chmodAndUpdateAcl
| Both the chmod command and ACL commands can change access permissions.
|
| "comment": "Comment",
| A comment that appears in the output of the mmlsfileset command.
|
| "iamMode": "Mode"
| The integrated archive manager (IAM) mode for the fileset.
|
| advisory
|
| noncompliant
|
| compliant
|
| "oid": "ID"
| Internal identifier of the fileset that is used for paging.
|
| "id": "ID"
| The fileset identifier.
|
| "status": "Status"
| Specifies whether the file is linked or unlinked.
|
| "parentId": "ParentID"
| The parent identifier of the fileset.
|
| "created": "DateTime",
| The date and time when the fileset was created.
|
| "isInodeSpaceOwner": "isOwner",
| Indicates whether the fileset has its own inode space.
|
| "isInodeSpaceMask": "isMask",
| The inode space mask.

```

```

| "snapID": ID
| The snapshot ID.

| "afm":
| Information about Active File Management (AFM).

| "afmTarget": Interval"Protocol://{Host | Map}/Path"
| The home that is associated with the cache.

| "afmAsyncDelay": Delay
| The time in seconds by which to delay write operations because of the lag in updating
| remote clusters.

| "afmDirLookupRefreshInterval": Interval
| The interval in seconds between data revalidations of directories that are caused by lookup
| operations such as ls and stat.

| "afmDirOpenRefreshInterval": Interval
| The interval in seconds between data revalidations of directories that are caused by lookup
| operations.

| "afmEnableAutoEviction": "{yes | no}"
| Enables or disables eviction on the fileset.

| "afmExpirationTimeout": Interval
| The timeout in seconds after which cached data is considered expired. Used with
| afmDisconnectTimeout, which is set with mmchconfig.

| "afmFileLookupRefreshInterval": Interval
| The interval in seconds between data revalidations of files that are caused by lookup
| operations such as ls and stat.

| "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
| independent-writer | Primary | Secondary}"
| The mode in which the cache operates.

| "afmNumFlushThreads": numThreads
| The number of threads that are used on each gateway to synchronize updates to the home
| cluster.

| "afmParallelReadChunkSize": Size
| The minimum chunk size of read data, in bytes, that must be distributed among the gateway
| nodes during parallel reads.

| "afmParallelReadThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel reads.

| "afmParallelWriteChunkSize": Size
| The minimum chunk size of write data, in bytes, that must be distributed among the gateway
| nodes during parallel writes.

| "afmParallelWriteThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel writes.

| "afmPrefetchThreshold": {1 | 1-99 | 100}
| The percentage of file size that must be cached before the entire file is prefetched.

| 0 Enables full file prefetching.

| 1-99 The percentage of file size that must be cached before the entire file is prefetched.

| 100 Disables full file prefetching.

| "afmPrimaryID": "ID"
| The unique primary ID of the primary fileset for asynchronous data replication.

```

```

| "afmRPO": Interval
| The recovery point objective (RPO) interval, in minutes, for a primary fileset.
| "afmShowHomeSnapshots": "{yes | no}"
| Shows or hides the home snapshot directory in cache. Specify yes to show, no to hide.

```

## Examples

The following example gets information about the fileset myFset1 in file system gpfs0.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myfset1'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/filesets/myfset1?lastId=1001"
 },
 "filesets": [
 {
 "config": {
 "filesetName": "myFset1",
 "filesystemName": "gpfs0",
 "path": "/mnt/gpfs0/myFset1",
 "inodeSpace": "0",
 "maxNumInodes": "4096",
 "permissionChangeMode": "chmodAndSetAcl",
 "comment": "Comment1",
 "iamMode": "off",
 "oid": "158",
 "id": "5",
 "status": "Linked",
 "parentId": "1",
 "created": "2016-12-13 13.59.15",
 "isInodeSpaceOwner": "1",
 "inodeSpaceMask": "0",
 "snapId": "0",
 "rootInode": "131075"
 },
 "afm": {
 "afmTarget": "nfs://198.51.100.11/gpfs/afmHomeFs/afmHomeFileset2",
 "afmAsyncDelay": "0",
 "afmDirLookupRefreshInterval": "60",
 "afmDirOpenRefreshInterval": "60",
 "afmEnableAutoEviction": "false",
 "afmExpirationTimeout": "100",
 "afmFileLookupRefreshInterval": "30",
 "afmMode": "read-only",
 "afmNumFlushThreads": "4",
 "afmParallelReadChunkSize": "134217728",
 "afmParallelReadThreshold": "1024",
 "afmParallelWriteChunkSize": "0",
 "afmParallelWriteThreshold": "0",
 "afmPrefetchThreshold": "0",
 "afmPrimaryID": "string",
 }
 }
]
}
```

```
| "afmRPO": "0",
| "afmShowHomeSnapshots": "false"
| }
|]
|
| "mmcrfileset command" on page 244
| Creates a GPFS fileset.
|
| "mmchfileset command" on page 173
| Changes the attributes of a GPFS fileset.
|
| "mmlsfileset command" on page 406
| Displays attributes and status for GPFS filesets.
```

## **Filesystems/{filesystemName}/filesets/{filesetName}: PUT**

| Modifies a fileset.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The PUT `filesystems/{filesystemName}/filesets/{filesetName}` request modifies a specific fileset, which is part of the specified file system. For more information about the fields in the data structures that are returned, see the topics “`mmcrfileset` command” on page 244, “`mmchfileset` command” on page 173, and “`mmlsfileset` command” on page 406.

### **Request URL**

| `https://<IP or host name of API server>:port/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName`

| where:

| **filesystems/filesystemName**

| Specifies the name of the file system to which the fileset belongs. Required.

| **filesets/filesetName**

| Specifies fileset to be modified. Required.

### **Request headers**

| `Accept: application/json`

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 39. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                    | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                                                                             | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation. | Required.         |

### **Request data**

```
{
 "newFilesetName": "Fileset",
 "maxNumInodes": "Inodes",
 "allocInodes": "Inodes",
 "permissionChangeMode": "chmodOnly | setAclOnly | chmodAndSetAcl",
 "comment": "Comment",
 "iamMode": "{advisory | noncompliant | compliant}",
 "afmTarget": "Protocol://{Host | Map}/Path"
 "afmAsyncDelay": SecondsDelay,
 "afmDirLookupRefreshInterval": SecondsInterval,
```

```

 "afmDirOpenRefreshInterval": SecondsInterval,
 "afmEnableAutoEviction": "{yes | no}",
 "afmExpirationTimeout": SecondsInterval,
 "afmFileLookupRefreshInterval": SecondsInterval,
 "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu |
 independent-writer | Primary | Secondary}",
 "afmNumFlushThreads": NumThreads,
 "afmParallelReadChunkSize": NumBytes,
 "afmParallelReadThreshold": NumMegabytes,
 "afmParallelWriteChunkSize": NumBytes,
 "afmParallelWriteThreshold": NumMegabytes,
 "afmPrefetchThreshold": {1 | 1-99 | 100},
 "afmPrimaryID": "ID",
 "afmRPO": MinutesInterval,
 "afmShowHomeSnapshots": "{yes | no}",
 }
}

"newFilesetName": "Fileset",
 The name of the fileset.

"maxNumInodes": "Inodes"
 The inode limit for the inode space that is owned by the specified fileset.

"permissionChangeMode": "Mode"
 The permission change mode. Controls how chmod and ACL commands affect objects in the fileset.

 chmodOnly
 Only the chmod command can change access permissions.

 setAclOnly
 Only the ACL commands and API can change access permissions.

 chmodAndSetAcl
 Both the chmod command and ACL commands can change access permissions.

 chmodAndUpdateAcl
 Both the chmod command and ACL commands can change access permissions.

"comment": "Comment",
 A comment that appears in the output of the mmlsfileset command.

"iamMode": "Mode"
 The integrated archive manager (IAM) mode for the fileset.

 advisory
 noncompliant
 compliant

"afmTarget": Protocol://{Host | Map}/Path
 The home that is associated with the cache.

"afmAsyncDelay": Delay
 The time in seconds by which to delay write operations because of the lag in updating remote clusters.

"afmDirLookupRefreshInterval": Interval
 The interval in seconds between data revalidations of directories that are caused by lookup operations such as ls and stat.

"afmDirOpenRefreshInterval": Interval
 The interval in seconds between data revalidations of directories that are caused by lookup operations.

```

```

| "afmEnableAutoEviction": "{yes | no}"
| Enables or disables eviction on the fileset.

| "afmExpirationTimeout": Interval
| The timeout in seconds after which cached data is considered expired. Used with
| afmDisconnectTimeout, which is set with mmchconfig.

| "afmFileLookupRefreshInterval": Interval
| The interval in seconds between data revalidations of files that are caused by lookup operations such
| as ls and stat.

| "afmMode": "{single-writer | sw | read-only | ro | local-updates | lu | independent-writer |
| Primary | Secondary}"
| The mode in which the cache operates.

| "afmNumFlushThreads": numThreads
| The number of threads that are used on each gateway to synchronize updates to the home cluster.

| "afmParallelReadChunkSize": Size
| The minimum chunk size of read data, in bytes, that must be distributed among the gateway nodes
| during parallel reads.

| "afmParallelReadThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel reads.

| "afmParallelWriteChunkSize": Size
| The minimum chunk size of write data, in bytes, that must be distributed among the gateway nodes
| during parallel writes.

| "afmParallelWriteThreshold": Threshold
| The minimum file size, in megabytes, at which to do parallel writes.

| "afmPrefetchThreshold": {1 | 1-99 | 100}
| The percentage of file size that must be cached before the entire file is prefetched.

| 0 Enables full file prefetching.
| 1-99 The percentage of file size that must be cached before the entire file is prefetched.
| 100 Disables full file prefetching.

| "afmPrimaryID": "ID"
| The unique primary ID of the primary fileset for asynchronous data replication.

| "afmRPO": Interval
| The recovery point objective (RPO) interval, in minutes, for a primary fileset.

| "afmShowHomeSnapshots": "{yes | no}"
| Shows or hides the home snapshot directory in cache. Specify yes to show, no to hide.

```

## Response data

```

{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 },
],
}

```

```

 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
}
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

**"status":**  
| Return status.

**"message": "ReturnMessage",**  
| The return message.

**"code": "ReturnCode**  
| The return code.

**"jobs":**  
| An array of elements that describe jobs. Each element describes one job.

**"result"**

**"commands": "String"**  
| Array of commands that are run in this job.

**"progress": "String"**  
| Progress information for the request.

**"exitCode": "Exit code"**  
| Exit code of command. Zero is success, nonzero denotes failure.

**"stderr": "Error"**  
| CLI messages from *stderr*.

**"stdout": "String"**  
| CLI messages from *stdout*.

**"request"**

**"type": "{GET | POST | PUT | DELETE}"**  
| HTTP request type.

**"url": "URL"**  
| The URL through which the job is submitted.

**"data": " "**  
| Optional.

**"jobId": "ID",**  
| The unique ID of the job.

**"submitted": "Time"**  
| The time at which the job was submitted.

**"completed": "Time"**  
| The time at which the job was completed.

```
| "status":"RUNNING | COMPLETED | FAILED"
| Status of the job.
```

## Examples

| The following example gets information about the fileset myFset1 in file system gpfs0.

| Request data:

```
| curl -k -u admin:admin001 -X PUT --header 'content-type:application/json' --header 'accept:application/json'
| -d '{
| "newFilesetName": "newFset1",
| "maxNumInodes": "100M",
| "allocInodes": "5M",
| "permissionChangeMode": "chmodAndSetAcl",
| "comment": "New comment",
| "iamMode": "compliant",
| "afmTarget": "nfs://10.0.100.11/gpfs/afmHomeFs/afmHomeFileset2",
| "afmAsyncDelay": "0",
| "afmDirLookupRefreshInterval": "60",
| "afmDirOpenRefreshInterval": "60",
| "afmEnableAutoEviction": "false",
| "afmExpirationTimeout": "100",
| "afmFileLookupRefreshInterval": "30",
| "afmMode": "read-only",
| "afmNumFlushThreads": "4",
| "afmParallelReadChunkSize": "134217728",
| "afmParallelReadThreshold": "1024",
| "afmParallelWriteChunkSize": "0",
| "afmParallelWriteThreshold": "0",
| "afmPrefetchThreshold": "0",
| "afmPrimaryID": "string",
| "afmRPO": "0",
| "afmShowHomeSnapshots": "false"
| }' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "job": [
| {
| "result": {
| "commands": "[\"mmcrfileset gpfs0 restfs1001\", ...]",
| "progress": "[\"(2/3) Linking fileset\"]",
| "exitCode": "0",
| "stderr": "[\"EFSSG0740C There are not enough resources available to create
| a new independent file set.\", ...]",
| "stdout": "[\"EFSSG4172I The file set {0} must be independent.\", ...]"
| },
| "request": {
| "type": "POST",
| "url": "/scalemgmt/v2/filesystems/gpfs0/filesets",
| "data": "{\"config\":{\"filesetName\":\"restfs1001\",\"owner\":\"root\",\"path\":
| \"/mnt/gpfs0/rest1001\",\"permissions\":\"555\"}}"
| },
| "jobId": "12345",
| "submitted": "2016-11-14 10.35.56",
| "completed": "2016-11-14 10.35.56",
| }
|]
| }
```

```
| "status": "COMPLETED"
| }
| }
```

### See also

- “[mmchfileset command](#)” on page 173
- “[mmcrfileset command](#)” on page 244
- “[mmlsfileset command](#)” on page 406

## **Filesystems/{filesystemName}/filesets/{filesetName}/link: DELETE**

| Unlinks a fileset.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **DELETE filesystems/{filesystemName}/filesets/{filesetName}/link** request unlinks an existing fileset, which is part of the specified file system. For more information about the fields in the data structures that are returned, see “[mmunlinkfileset command](#)” on page 578.

### **Request URL**

| `https://management API host:port/sclemgmt/v2/filesystems/filesystemName/filesets/filesetName/link`

| where:

#### **filesystem/filesystemName**

| Specifies the name of the file system to which the fileset belongs. Required.

#### **filesets/filesetName**

| Specifies fileset to be unlinked. Required.

### **Request headers**

| `Accept: application/json`

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 40. List of request parameters*

| Parameter name | Description and applicable keywords                                                         | Required/optional |
|----------------|---------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote: | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                          | Required.         |
| force          | Specifies whether to force the unlink operation. Default value is false.                    | Optional.         |

### **Request data**

```
| { "force": "True | False"
| }
| "force": "True | False"
| Forces the unlink to take place even though there are open files.
```

### **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
```

```

{
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"jobs":**

| An array of elements that describe jobs. Each element describes one job.

| **"status":**

| Return status.

| **"message": "ReturnMessage",**  
|       The return message.

| **"code": "ReturnCode"**  
|       The return code.

| **"result"**

| **"commands": "String"**  
|       Array of commands that are run in this job.

| **"progress": "String"**  
|       Progress information for the request.

| **"exitCode": "Exit code"**  
|       Exit code of command. Zero is success, nonzero denotes failure.

| **"stderr": "Error"**  
|       CLI messages from *stderr*.

| **"stdout": "String"**  
|       CLI messages from *stdout*.

| **"request"**

| **"type": "{GET | POST | PUT | DELETE}"**  
|       HTTP request type.

| **"url": "URL"**  
|       The URL through which the job is submitted.

| **"data": " "**  
|       Optional.

```

| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example unlinks the fileset myFset1.

Request data:

```

| curl -k -u admin:admin001 -X DELETE --header 'accept:application/json' -d '{
| "force": "true"
| }' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/link'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "jobs" : [{
 "jobId" : 10000000000006,
 "status" : "RUNNING",
 "submitted" : "2017-03-14 16:05:30,960",
 "completed" : "N/A",
 "request" : {
 "data" : {
 "force" : false
 },
 "type" : "DELETE",
 "url" : "/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/link"
 },
 "result" : { }
 }],
 "status" : {
 "code" : 202,
 "message" : "The request was accepted for processing"
 }
}

```

## See also

- “[mmunlinkfileset command](#)” on page 578

## | **Filesystems/{filesystemName}/filesets/{filesetName}/link: POST**

| Links an existing fileset.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **POST filesystems/{filesystemName}/filesets/{filesetName}/link** request links an existing fileset, which is part of the specified file system. For more information about the fields in the data structures that are returned, see “[mmlinkfileset command](#)” on page 390.

### | **Request URL**

| <https://management API host:port/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName/link>

| where:

| **filesystem/filesystemName**

|     Specifies the name of the file system to which the fileset belong. Required.

| **filesets/filesetName**

|     Specifies fileset to be linked. Required.

| **link**

|     Specifies action to be performed in the POST call. Required.

### | **Request headers**

| Accept: application/json

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 41. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                           | Required/optional |
|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                   | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                                                                            | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to theIBM Spectrum Scale system to perform the requested operation. | Required.         |

### | **Request data**

| { "path": "Path"  
| }

| **"path": "Path"**

|     The absolute path of the fileset.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

### "jobs":

An array of elements that describe jobs. Each element describes one job.

#### "status":

Return status.

"message": "ReturnMessage",  
The return message.

"code": *ReturnCode*  
The return code.

#### "result"

"commands": "String"  
Array of commands that are run in this job.

"progress": "String"  
Progress information for the request.

"exitCode": "Exit code"  
Exit code of command. Zero is success, nonzero denotes failure.

"stderr": "Error"  
CLI messages from *stderr*.

"stdout": "String"  
CLI messages from *stdout*.

#### "request"

"type": "{GET | POST | PUT | DELETE}"  
HTTP request type.

```

| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example links the fileset myFset1.

Request data:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "path": "/mnt/gpfs0/fset1"
}' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/link'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "jobs" : [{
 "jobId" : 1000000000005,
 "status" : "RUNNING",
 "submitted" : "2017-03-14 15:59:34,180",
 "completed" : "N/A",
 "request" : {
 "data" : {
 "path" : "/mnt/gpfs0/fset1"
 },
 "type" : "POST",
 "url" : "/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/link"
 },
 "result" : { }
 }],
 "status" : {
 "code" : 202,
 "message" : "The request was accepted for processing"
 }
}
```

## See also

- “mmlinkfileset command” on page 390

## **Filesystems/{filesystemName}/filesets/{filesetName}/quotas: GET**

| Gets information about quota set at the fileset level. With no parameters, all quota limits are returned.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **GET filesystems/filesystemName/filesets/filesetName/quotas** request gets information about quotas  
| set for filesets within a particular file system. For more information about the fields in the data structures  
| that are returned, see the topics “[mmsetquota command](#)” on page 548 and “[mmrepquota command](#)” on  
| page 517.

| The quota definition at the file system level must be *profileset* to retrieve the quota information using this  
| API call.

### **Request URL**

| `https://<API server>:<port>/scalemgmt/v2/filesystems/FileSystemName/filesets/filesetName/quotas`

| where

#### **filesystems/filesystemName**

| The file system that contains the fileset. Required.

#### **filesets/filesetName**

| The fileset about which you need to get the quota information. Required.

#### **quotas**

| Specifies that you need to get the quota details. Required.

### **Request headers**

| Accept: application/json

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 42. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                               | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |

## Request data

| No request data.

## Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging": "NFSv4",
| "quotas": [
| "quotaID": "ID",
| "filesystemName": "File system name",
| "filesetName": "Fileset name",
| "quotaType": "Type",
| "objectName": "Name",
| "objectId": "ID",
| "blockUsage": "Usage",
| "blockQuota": "Soft limit",
| "blockLimit": "Hard limit",
| "blockInDoubt": "Space in doubt",
| "blockGrace": "Grace period",
| "filesUsage": "Number of files in usage",
| "filesQuota": "Soft limit",
| "filesLimit": "Hard limit",
| "filesInDoubt": "Files in doubt",
| "filesGrace": "Grace period",
| "defaultQuota": "Default",
|],
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| "status":  
| Return status.

| "message": "ReturnMessage",  
| The return message.

| "code": "ReturnCode"  
| The return code.

| "paging": "NFSv4"

| "quotas": "

| "quotaID": "ID"  
| Internal ID used for paging.

| "filesystemName": "File system name"  
| The file system for which the quota is applicable.

| "filesetName": "Fileset name"  
| The fileset for which the quota is applicable.

| "quotaType": "USR | GRP | FILESET"  
| The quota type.

| "objectName": "Name"  
| Name of the fileset, user, or user group for which the quota is applicable.

| "objectId": "ID"  
| Unique identifier of the fileset, user, or user group.

```

| "blockUsage": "Usage"
| Current capacity quota usage.
|
| "blockQuota": "Soft limit"
| The soft limit set for the fileset, user, or user group.
|
| "blockLimit": "Hard limit"
| The hard limit set for the capacity quota usage. A grace period starts when the hard limit is
| reached.
|
| "blockInDoubt": "Space in doubt"
| The capacity that is in doubt.
|
| "blockGrace": "Grace period"
| The grace period set for the capacity quota.
|
| "filesUsage": "Number of files in usage"
| Number of inodes.
|
| "filesQuota": "Soft limit"
| The soft limit set for the inode quota.
|
| "filesLimit": "Hard limit"
| The hard limit set for the inode quota.
|
| "filesInDoubt": "Files in doubt"
| The files that are in doubt.
|
| "filesGrace": "Grace period"
| The grace period set for the inode usage.
|
| "defaultQuota": "Default"
| Default quota.

```

## Examples

The following example gets quota information for the fileset myFset1, which is part of the file system gpfs0.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/quotas'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/quotas?lastId=1001"
 },
 "quotas": [
 {
 "quotaId": "4711",
 "filesystemName": "gpfs0",
 "filesetName": "myFset1",
 "quotaType": "USR",
 "objectName": "myFset1",
 }
]
}
```

```
| "objectId": "128",
| "blockUsage": "0",
| "blockQuota": "2048",
| "blockLimit": "4096",
| "blockInDoubt": "1024",
| "blockGrace": "none",
| "filesUsage": "32",
| "filesQuota": "50",
| "filesLimit": "100",
| "filesInDoubt": "3",
| "filesGrace": "none",
| "defaultQuota": false
| }
|]
|
| "mmsetquota command" on page 548
| Sets quota limits.
|
| "mmrepquota command" on page 517
| Displays file system user, group, and fileset quotas.
```

## **Filesystems/{filesystemName}/filesets/{filesetName}/quotas: POST**

| Sets quota limits or default quota limit for a fileset.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The POST **filesystems/filesystemName/filesets/filesetName/quotas** request define quota limits for a fileset, which is part of the specified file system. For more information about the fields in the data structures that are returned, see the topics “mmsetquota command” on page 548 and “mmrepquota command” on page 517.

| The quota definition at the file system level must be *perfileset* to successfully run this API command.

### **Request URL**

| `https://<API server>:<port>/scalemgmt/v2/filesystems/FileSystemName/filesets/filesetName/quotas`

| where

| **filesystems/filesystemName**

| The file system that contains the fileset. Required.

| **filesets/filesetName**

| The fileset for which you need to set the quota information. Required.

| **quotas**

| Specifies that you need to set the quota details. Required.

### **Request headers**

| Accept: application/json

### **Request data**

```
| {
| "operationType": "Type",
| "quotaType": "Type",
| "blockSoftLimit": "Soft limit",
| "blockHardLimit": "Hard limit",
| "filesSoftLimit": "Soft limit",
| "filesHardLimit": "Hard limit",
| "filesGracePeriod": "Grace period",
| "blockGracePeriod": "Default",
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"operationType": "Operation type"**

| In this case, set quota.

| **"quotaType": "USR | GRP | FILESET"**

| The quota type.

| **"objectName": "Name"**

| Name of the fileset, user, or user group for which the quota is applicable.

| **"blockSoftLimit": "Soft limit"**

| The soft limit set for the fileset, user, or user group.

```

| "blockHardLimit": "Hard limit"
| The hard limit set for the capacity quota usage. A grace period starts when the hard limit is reached.
| "filesSoftLimit": "Soft limit"
| The soft limit set for the inode quota.
| "filesHardLimit": "Hard limit"
| The hard limit set for the inode quota.
| "filesGracePeriod": "Grace period"
| The grace period set for the inode usage.
| "blockGracePeriod": "Grace period"
| The grace period set for the capacity quota.

```

## Response data

```

{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

```

| "status":
| Return status.
| "message": "ReturnMessage",
| The return message.
| "code": ReturnCode
| The return code.
| "jobs":
| An array of elements that describe jobs. Each element describes one job.
| "result"
| "commands": "String"
| Array of commands that are run in this job.

```

```

| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example sets quota for the fileset myFset1, which is part of the file system gpfs0:

Request data:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "operationType": "setQuota",
 "quotaType": "user",
 "objectName": "adam",
 "blockSoftLimit": "1M",
 "blockHardLimit": "2M",
 "filesSoftLimit": "1K",
 "filesHardLimit": "2K",
 "filesGracePeriod": "null",
 "blockGracePeriod": "null"
}' "https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/quotas"

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "jobs" : [{
 "jobId" : 1000000000008,
 "status" : "RUNNING",
 "submitted" : "2017-03-14 16:07:43,474",
 "completed" : "N/A",
 }
]

```

```
| "request" : {
| "data" : {
| "operationType": "setQuota",
| "quotaType": "user",
| "objectName": "adam",
| "blockSoftLimit": "1M",
| "blockHardLimit": "2M",
| "filesSoftLimit": "1K",
| "filesHardLimit": "2K",
| "filesGracePeriod": "null",
| "blockGracePeriod": "null"
| },
| "type" : "POST",
| "url" : "/scalemgmt/v2/filesystems/gpfs0/quotas"
| },
| "result" : { }
|],
| "status" : {
| "code" : 202,
| "message" : "The request was accepted for processing"
| }
```

“mmsetquota command” on page 548

Sets quota limits.

“mmrepquota command” on page 517

Displays file system user, group, and fileset quotas.

## **Filesystems/{filesystemName}/filesets/{filesetName}/snapshots: GET**

| Gets information about snapshots of a specific fileset.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **GET filesystems/{filesystemName}/filesets/{filesetName}/snapshots** request gets information about snapshots in the specified fileset. For more information about the fields in the data structures that are returned, see the topics “[mmcrsnapshot command](#)” on page 267 and “[mmlssnapshot command](#)” on page 439.

### **Request URL**

| `https://IP of API server:<port>/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName/snapshots`

| where

| **filesystems/filesystemName**

|     Specifies the file system to which the fileset belongs. Required.

| **filesets/filesetName**

|     Specifies the fileset of which the snapshot is taken. Required.

| **snapshots**

|     Specifies snapshot as the resource of this GET call. Required.

### **Request headers**

| Content-Type: application/json

| Accept: application/json

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 43. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                               | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |

### **Request data**

| No request data.

## Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "snapshots": [
| {
| "config": {
| "snapshotName": "Snapshot",
| "filesystemName": "Device",
| "filesetName": "Fileset",
| "oid": "ID",
| "snapID": "ID",
| "status": "Status",
| "created": "DateTime",
| "quotas": "Quotas",
| "snapType": "Type",
| }
| }
|]
| }
```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**  
|     Return status.  
  
| **"message": "ReturnMessage",**  
|     The return message.  
  
| **"code": "ReturnCode"**  
|     The return code.  
  
| **"paging"**  
|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by  
|     the query.  
  
| **"snapshotName": "Snapshot"**  
|     The snapshot name.  
  
| **"filesystemName": "Device"**  
|     The file system that is the target of the snapshot.  
  
| **"filesetName": "Fileset"**  
|     For a fileset snapshot, the fileset that is a target of the snapshot.  
  
| **"oid": "ID"**  
|     Internal identifier that is used for paging.  
  
| **"snapID": "ID"**  
|     The snapshot ID.  
  
| **"status": "Status"**  
|     The snapshot status.  
  
| **"created": "DateTime"**  
|     The date and time when the snapshot was created.  
  
| **"quotas": "Quotas"**  
|     Any quotas that are applied to the fileset.

| "snapType": "Type"  
|     The AFM type of the snapshot, including "afm\_snap", "afm\_recovery", "afm\_failover", "afm\_rpo",  
|     "afm\_baserpo", and "Invalid".

## | Examples

| The following example gets information about the snapshots of the fileset *myFset1*, which belongs to the  
| file system *gpfs0*.

| Request data:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots'
```

| The request URL with no field or filter parameter returns only the details that uniquely identify the  
| object.

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "snapshots" : [{
 "config" : {
 "filesetName" : "myFset1",
 "filesystemName" : "gpfs0",
 "snapshotName" : "mySnap1"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}
```

| Using the field parameter ":all:" returns entire details of the snapshot as shown in the following example:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots?fields=:all:'
| {
 "snapshots" : [{
 "config" : {
 "created" : "2017-03-21 15:52:14,000",
 "filesetName" : "myFset1",
 "filesystemName" : "gpfs0",
 "oid" : 2,
 "quotas" : "",
 "snapID" : 2,
 "snapType" : "",
 "snapshotName" : "mySnap1",
 "status" : "Valid"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}
```

| "mmcrsnapshot command" on page 267

| Creates a snapshot of a file system or fileset at a single point in time.

- | “mmlssnapshot command” on page 439
- | Displays GPFS snapshot information.

## **Filesystems/{filesystemName}/filesets/{filesetName}/snapshots: POST**

| Creates a fileset snapshot.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The POST `filesystems/{filesystemName}/filesets/{filesetName}/snapshots` command creates a snapshot of  
| the specified fileset.

### **Request URL**

| `https://IP of API server:<port>/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName/snapshots`

| where

#### **filesystems/filesystemName**

| Specifies the file system to which the fileset belongs. Required.

#### **filesets/filesetName**

| Specifies the fileset of which the snapshot is taken. Required.

#### **snapshots**

| Specifies snapshot as the resource of this POST call. Required.

### **Request headers**

| Content-Type: application/json

| Accept: application/json

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 44. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local;, or :all_remote:                                                    | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                                                                             | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation. | Required.         |

### **Request data**

```
| {
| "snapshotName": ""
| }
```

#### **"snapshotName":**

| Name of the snapshot to be created.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"status":**  
Return status.

**"message": "ReturnMessage",**  
The return message.

**"code": "ReturnCode"**  
The return code.

**"jobs":**  
An array of elements that describe jobs. Each element describes one job.

**"result"**

**"commands": "String"**  
Array of commands that are run in this job.

**"progress": "String"**  
Progress information for the request.

**"exitCode": "Exit code"**  
Exit code of command. Zero is success, nonzero denotes failure.

**"stderr": "Error"**  
CLI messages from *stderr*.

**"stdout": "String"**  
CLI messages from *stdout*.

**"request"**

**"type": "{GET | POST | PUT | DELETE}"**  
HTTP request type.

```

| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example creates a snapshot *snap2* of the fileset *myFset1* in file system *gpfs0*.

Request data:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{ \
 "snapshotName": "snap2" \
}' 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "POST",
 "url": "/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots",
 "data": "{\"snapshotName\": \"snap2\"}"
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}

```

“[mmcrsnapshot command](#)” on page 267  
Creates a snapshot of a file system or fileset at a single point in time.  
“[mmlssnapshot command](#)” on page 439  
Displays GPFS snapshot information.

## | **Filesystems/{filesystemName}/filesets/{filesetName}/snapshots/{snapshotName}: DELETE**

| Deletes a fileset snapshot.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **DELETE filesystems/{filesystemName}/filesets/{filesetName}/snapshots/{snapshotName}** command  
| deletes the specific fileset snapshot. For more information on deleting snapshots, see the topic  
| “[mmdelsnapshot command](#)” on page 304.

### | **Request URL**

| Use this URL to delete a snapshot:

| `https://IP address of API server:<port>/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName/snapshots/snapshotName`

| where:

#### | **filesystems/filesystemName**

| Specifies the file system to which the fileset belongs. Required.

#### | **filesets/filesetName**

| Specifies the fileset of which the snapshot is deleted. Required.

#### | **snapshots/snapshotName**

| Specifies the snapshot to be deleted. Required.

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 45. List of request parameters*

| Parameter name | Description and applicable keywords                                                         | Required/optional |
|----------------|---------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote: | Required.         |
| filesetName    | The fileset name. This is the path of the fileset.                                          | Required.         |
| snapshotName   | The snapshot name. This is the path of the snapshot.                                        | Required.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
|]
| }
```

```

{
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

**"status":**  
| Return status.

**"message": "ReturnMessage",**  
| The return message.

**"code": *ReturnCode***  
| The return code.

**"jobs":**  
| An array of elements that describe jobs. Each element describes one job.

**"result"**

**"commands": "String"**  
|   Array of commands that are run in this job.

**"progress": "String"**  
|   Progress information for the request.

**"exitCode": "Exit code"**  
|   Exit code of command. Zero is success, nonzero denotes failure.

**"stderr": "Error"**  
|   CLI messages from *stderr*.

**"stdout": "String"**  
|   CLI messages from *stdout*.

**"request"**

**"type": "{GET | POST | PUT | DELETE}"**  
|   HTTP request type.

**"url": "URL"**  
|   The URL through which the job is submitted.

**"data": " "**  
|   Optional.

```

| " jobId": "ID",
| The unique ID of the job.
| " submitted": "Time"
| The time at which the job was submitted.
| " completed": "Time"
| The time at which the job was completed.
| " status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## | Examples

| The following example deletes the fileset snapshot *snap1* from the file system *gpfs0*.

| Request URL:

```
| curl -k -u admin:admin001 -X DELETE --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```

| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "job": [
| {
| "result": {},
| "request": {
| "type": "DELETE",
| "url": "scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1",
| "data": "{}"
| },
| "jobId": "12345",
| "submitted": "2016-11-14 10.35.56",
| "completed": "2016-11-14 10.35.56",
| "status": "COMPLETED"
| }
|]
| }
|
```

| “[mmcrsnapshot command](#)” on page 267  
| Creates a snapshot of a file system or fileset at a single point in time.  
| “[mmlssnapshot command](#)” on page 439  
| Displays GPFS snapshot information.

## **Filesystems/{filesystemName}/filesets/{filesetName}/snapshots/{snapshotName}: GET**

Gets information about a specific fileset snapshot.

### **Availability**

Available on all IBM Spectrum Scale editions.

### **Description**

The GET `filesystems/filesystemName/filesets/filesetName/snapshots/snapshotName` request gets information about a specific fileset snapshot in a specific file system. For more information about the fields in the data structures that are returned, see the topics “`mmcrsnapshot` command” on page 267 and “`mmlssnapshot` command” on page 439.

### **Request URL**

`https://IP address of API server:<port>/scalemgmt/v2/filesystems/filesystemName/filesets/filesetName/snapshots/snapshotName`

where

#### **filesystems/filesystemName**

Specifies the file system to which the fileset belongs. Required.

#### **filesets/{filesetName}**

Specifies the fileset of which the snapshot is taken. Required.

#### **snapshots/snapshotName**

Specifies a particular snapshot as the resource of this GET call. Required.

### **Request headers**

`Content-Type: application/json`

`Accept: application/json`

### **Request parameters**

The following parameters can be used in the request URL to customize the request:

*Table 46. List of request parameters*

| Parameter name              | Description and applicable keywords                                                                                                  | Required/optional |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| <code>filesystemName</code> | The file system name. You can also use keywords such as <code>:all:</code> , <code>:all_local:</code> , or <code>:all_remote:</code> | Required.         |
| <code>filesetName</code>    | The fileset name. This is the path of the fileset.                                                                                   | Required.         |
| <code>snapshotName</code>   | The snapshot name. This is the path of the snapshot.                                                                                 | Required.         |
| <code>fields</code>         | Comma separated list of fields to be included in response. <code>'all'</code> selects all available fields.                          | Optional.         |
| <code>filter</code>         | Filter objects by expression. For example,<br><code>'status=HEALTHY,entityType=FILESET'</code>                                       | Optional.         |

## Request data

| No request data.

## Response data

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
| "paging":
| {
| "next": "URL"
| },
| "snapshots": [
| {
| "config": {
| "snapshotName": "Snapshot",
| "filesystemName": "Device",
| "filesetName": "Fileset",
| "oid": "ID",
| "snapID": "ID",
| "status": "Status",
| "created": "DateTime",
| "quotas": "Quotas",
| "snapType": "Type",
| }
| }
|]
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**  
|   Return status.  
  
| **"message": "ReturnMessage",**  
|   The return message.  
  
| **"code": *ReturnCode***  
|   The return code.  
  
| **"paging"**  
|   The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.  
  
| **"snapshotName": "Snapshot"**  
|   The snapshot name.  
  
| **"filesystemName": "Device"**  
|   The file system that is the target of the snapshot.  
  
| **"filesetName": "Fileset"**  
|   For a fileset snapshot, the fileset that is a target of the snapshot.  
  
| **"oid": "ID"**  
|   Internal identifier that is used for paging.  
  
| **"snapID": "ID"**  
|   The snapshot ID.  
  
| **"status": "Status"**  
|   The snapshot status.

- | **"created":** "DateTime"
  - | The date and time when the snapshot was created.
- | **"quotas":** "Quotas"
  - | Any quotas that are applied to the fileset.
- | **"snapType":** "Type"
  - | The AFM type of the snapshot, including "afm\_snap", "afm\_recovery", "afm\_failover", "afm\_rpo", "afm\_baserpo", and "Invalid".

## | Examples

- | The following example gets information about the snapshots *snap1* of the fileset *myFset1*, which belongs to the file system *gpfs0*.

- | Request data:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1'
```

- | Response data:

- | **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.
- | The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "paging": {
| "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1?lastId=1001"
| },
| "snapshots": [
| {
| "config": {
| "snapshotName": "snap1",
| "filesystemName": "gpfs0",
| "filesetName": "myFset1",
| "oid": "123",
| "snapID": "5",
| "status": "Valid",
| "created": "2017-01-09 14.55.37",
| "quotas": "string",
| "snapType": "string"
| }
| }
|]
| }
```

| "mmcrsnapshot command" on page 267

| Creates a snapshot of a file system or fileset at a single point in time.

| "mmlssnapshot command" on page 439

| Displays GPFS snapshot information.

## | **Filesystems/{filesystemName}/owner/{path}: GET**

| Gets information about owner of a file or directory.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/filesystemName/owner/path** request gets information about owner of files or directories within a particular file system.

### | **Request URL**

| `https://<IP address of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/owner/path`

| where

#### | **filesystems/filesystemName**

| The file system about which you want to get the information. Required.

#### | **owner/path**

| The path of the file or directory about which you want to get the owner information. Required.

### | **Request headers**

| Accept: application/json

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 47. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                   | Required/optional |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                                                                                                           | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields.                                                                                                                                      | Optional.         |
| path           | The file path relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to "%2F" in the request URL. | Required.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage"
| }
```

```

 },
 "owner": [
 "user": "User name",
 "uid": "User ID",
 "group": "Group name",
 "gid": "Group ID",
],
 },
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

```

| "status":
| Return status.
|
| "message": "ReturnMessage",
| The return message.
|
| "code": "ReturnCode
| The return code.
|
| "Owner": "Owner details"
|
| "user": "User name"
| Name of the owner.
|
| "uid": "User ID"
| Unique identifier of the owner.
|
| "group": "Group name"
| Name of the user group that owns the file or directory.
|
| "gid": "Group ID"
| Unique identifier of the user group that owns the file or directory.

```

## Examples

| The following example gets owner information for the files and directories of the file system gpfso.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfso/owner/mnt%2Fgpfso%2Fxaz'
```

| Response URL:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```

| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "owner": {
| "path": "/mnt/gpfso/xaz",
| "user": "testuser55",
| "uid": "1234",
| "group": "mygroup",
| "gid": "4711"
| }
| }

```

| "mmcrfs command" on page 250  
| Creates a GPFS file system.

- | “mmchfs command” on page 179
- | Changes the attributes of a GPFS file system.
- | “mmlsfs command” on page 411
- | Displays file system attributes.

## **Filesystems/{filesystemName}/owner/{path}: PUT**

| Sets owner for a file or directory.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The PUT **filesystems/filesystemName/owner/path** request sets owner for files or directories within a particular file system.

| The file and directory ownership can be controlled by Linux commands such as **stat** and **chown**. You can use either ID or name of the user or group to define the ownership. If you use both ID and name, the ID takes precedence. Only the user with *DataAccess* role can change the owner of a file or a non-empty directory. Empty directories can be changed by other roles such as Administrator,ProtocolAdmin, StorageAdmin, and SecurityAdmin.

### **Request URL**

| `https://<IP address of API server>:<port>/scalemgmt/v2/filesystems/FileSystemName/owner/path`

| where

| **filesystems/filesystemName**

|     Specifies the file system to which the file or directory belongs. Required.

| **owner/path**

|     The path of the file or directory for which you want to set the owner. Required.

### **Request headers**

| `Accept: application/json`

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 48. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                                 | Required/optional |
|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local;, or :all_remote:                                                                                                                                                         | Required.         |
| path           | The file path relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: <code>mnt/gpfs0/rest01</code> . The forward slashes in the path are encoded to "%2F" in the request URL. | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation.                                                                                                      | Required.         |

## Request data

```
| {
| "path": "Path",
| "user": "User name",
| "uid": "User ID",
| "group": "Group name",
| "gid": "Group ID",
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"path": "Path"**  
| The path of the file or directory.  
  
| **"user": "User name"**  
| Name of the owner.  
  
| **"uid": "User ID"**  
| Unique identifier of the owner.  
  
| **"group": "Group name"**  
| Name of the user group that owns the file or directory.  
  
| **"gid": "Group ID"**  
| Unique identifier of the user group that owns the file or directory.

## Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "owner": [
| "path": "Path",
| "user": "User name",
| "uid": "User ID",
| "group": "Group name",
| "gid": "Group ID",
|],
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**  
| Return status.  
  
| **"message": "ReturnMessage"**,  
| The return message.  
  
| **"code": "ReturnCode"**  
| The return code.  
  
| **"Owner": "Owner details"**  
  
| **"path": "Path"**  
| The path of the file or directory.  
  
| **"user": "User name"**  
| Name of the owner.  
  
| **"uid": "User ID"**  
| Unique identifier of the owner.

```

| "group": "Group name"
| Name of the user group that owns the file or directory.
| "gid": "Group ID"
| Unique identifier of the user group that owns the file or directory.

```

## Examples

| The following example sets owner for the files and directories of the file system gpfs0.

| Request data:

```

| curl -k -u admin:admin001 -X PUT --header 'content-type:application/json' --header 'accept:application/json'
| -d '{
| "path": "/mnt/gpfs0/xaz",
| "user": "testuser55",
| "uid": "1234",
| "group": "mygroup",
| "gid": "4711"
| }' "https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/owner/%2Fmnt%2Fgpfs0%2Fxaz"

```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "POST",
 "url": "/scalemgmt/v2/filesystems/gpfs0/owner",
 "data": "{\n \"path\": \"/mnt/gpfs0/xaz\",\n \"user\": \"testuser55\",\n \"uid\": \"1234\",\n \"group\": \"mygroup\",\n \"gid\": \"4711\"\n }",
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
 }
]
}

```

| “mmcrfs command” on page 250

| Creates a GPFS file system.

| “mmchfs command” on page 179

| Changes the attributes of a GPFS file system.

| “mmlsfs command” on page 411

| Displays file system attributes.

## | **Filesystems/{filesystemName}/quotas: GET**

- | Gets information about quota set at the file system level. With no parameters, all quota limits are returned.

### | **Availability**

- | Available on all IBM Spectrum Scale editions.

### | **Description**

- | The **GET filesystems/filesystemName/quotas** request gets information about quotas set at file system level. “mmsetquota command” on page 548 and “mmrepquota command” on page 517
- | The *perfileset* quota must be disabled to display the user and group quotas. The fileset quota is also available in the response data.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/filesystems/filesystemName/quotas`

| where

#### | **filesystems/filesystemName**

| The file system about which you need the information. Required.

#### | **quotas**

| Specifies that you need to get the quota details. Required.

### | **Request headers**

| Accept: application/json

### | **Request parameters**

- | The following parameters can be used in the request URL to customize the request:

| *Table 49. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>  'status=HEALTHY,entityType=FILESET'              | Optional.         |

### | **Request data**

- | No request data.

## Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging": "{NFSv4} ",
| "quotas": [
| "quotaID": "ID",
| "filesystemName": "File system name",
| "filesetName": "Fileset name",
| "quotaType": "Type",
| "objectName": "Name",
| "objectId": "ID",
| "blockUsage": "Usage",
| "blockQuota": "Soft limit",
| "blockLimit": "Hard limit",
| "blockInDoubt": "Space in doubt",
| "blockGrace": "Grace period",
| "filesUsage": "Number of files in usage",
| "filesQuota": "Soft limit",
| "filesLimit": "Hard limit",
| "filesInDoubt": "Files in doubt",
| "filesGrace": "Grace period",
| "defaultQuota": "Default",
|],
| }
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

```
| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": "ReturnCode"
| The return code.

| "paging": "NFSv4"

| "quotas": ""

| "quotaID": "ID"
| Internal ID used for paging.

| "filesystemName": "File system name"
| The file system for which the quota is applicable.

| "filesetName": "Fileset name"
| The fileset for which the quota is applicable.

| "quotaType": "USR | GRP | FILESET"
| The quota type.

| "objectName": "Name"
| Name of the fileset, user, or user group for which the quota is applicable.

| "objectId": "ID"
| Unique identifier of the fileset, user, or user group.

| "blockUsage": "Usage"
| Current capacity quota usage.
```

```

| "blockQuota": "Soft limit"
| The soft limit set for the fileset, user, or user group.
|
| "blockLimit": "Hard limit"
| The hard limit set for the capacity quota usage. A grace period starts when the hard limit is
| reached.
|
| "blockInDoubt": "Space in doubt"
| The capacity that is in doubt.
|
| "blockGrace": "Grace period"
| The grace period set for the capacity quota.
|
| "filesUsage": "Number of files in usage"
| Number of inodes.
|
| "filesQuota": "Soft limit"
| The soft limit set for the inode quota.
|
| "filesLimit": "Hard limit"
| The hard limit set for the inode quota.
|
| "filesInDoubt": "Files in doubt"
| The files that are in doubt.
|
| "filesGrace": "Grace period"
| The grace period set for the inode usage.
|
| "defaultQuota": "Default"
| Default quota.

```

## Examples

| The following example gets quota information for the the file system gpfs0.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/quotas'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/quotas?lastId=1001"
 },
 "quotas": [
 {
 "quotaId": "4711",
 "filesystemName": "gpfs0",
 "filesetName": "myFset1",
 "quotaType": "USR",
 "objectName": "myFset1",
 "objectId": "128",
 "blockUsage": "0",
 "blockQuota": "2048",
 "blockLimit": "4096",
 }
]
}
```

```
| "blockInDoubt": "1024",
| "blockGrace": "none",
| "filesUsage": "32",
| "filesQuota": "50",
| "filesLimit": "100",
| "filesInDoubt": "3",
| "filesGrace": "none",
| "defaultQuota": false
| }
|]
| }
```

“mmsetquota command” on page 548  
Sets quota limits.

“mmrepquota command” on page 517  
Displays file system user, group, and fileset quotas.

## | **Filesystems/{filesystemName}/quotas: POST**

| Sets quota limits or default quota limit at file system level.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **POST filesystems/filesystemName/quotas** request define quota limits for a file system. For more information about the fields in the data structures that are returned, see the topics “[mmsetquota command](#)” on page 548 and “[mmrepquota command](#)” on page 517.

| The *perfileset* quota must be disabled to successfully complete this API command.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/filesystems/filesystemName/quotas`

| where

#### | **filesystems/filesystemName**

| Specifies that you need to set quota for the particular file system. Required.

#### | **quotas**

| Specifies that you need to set the quota details. Required.

### | **Request headers**

| Accept: application/json

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 50. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                    | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation. | Required.         |

### | **Request data**

```
{
 "operationType": "Type",
 "quotaType": "Type",
 "blockSoftLimit": "Soft limit",
 "blockHardLimit": "Hard limit",
 "filesSoftLimit": "Soft limit",
 "filesHardLimit": "Hard limit",
 "filesGracePeriod": "Grace period",
 "blockGracePeriod": "Default",
}
```

- | For more information about the fields in the following data structures, see the links at the end of this topic.
- | **"operationType":***Operation type*
  - | In this case, set quota.
- | **"quotaType":***USR | GRP | FILESET*
  - | The quota type.
- | **"objectName":***Name*
  - | Name of the fileset, user, or user group for which the quota is applicable.
- | **"blockSoftLimit":***Soft limit*
  - | The soft limit set for the fileset, user, or user group.
- | **"blockHardLimit":***Hard limit*
  - | The hard limit set for the capacity quota usage. A grace period starts when the hard limit is reached.
- | **"filesSoftLimit":***Soft limit*
  - | The soft limit set for the inode quota.
- | **"filesHardLimit":***Hard limit*
  - | The hard limit set for the inode quota.
- | **"filesGracePeriod":***Grace period*
  - | The grace period set for the inode usage.
- | **"blockGracePeriod":***Grace period*
  - | The grace period set for the capacity quota.

## | Response data

```
{
 "status": {
 "code":ReturnCode,
 "message":ReturnMessage
 },
 "jobs": [
 {
 "result": "",
 {
 "commands":String,
 "progress":String,
 "exitCode":Exit code,
 "stderr":Error,
 "stdout":String,
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 },
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
]
```

- | For more information about the fields in the following data structures, see the links at the end of this topic.
- | **"status":**
  - | Return status.

```

| "message": "ReturnMessage",
| The return message.
|
| "code": ReturnCode
| The return code.
|
| "jobs":
| An array of elements that describe jobs. Each element describes one job.
|
| "result"
|
| "commands": "String"
| Array of commands that are run in this job.
|
| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example sets quota for the file system gpfs0.

Request data:

```

| curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
| -d '{
| "operationType": "setQuota",
| "quotaType": "user",
| "objectName": "adam",
| "blockSoftLimit": "1M",
| "blockHardLimit": "2M",
| "filesSoftLimit": "1K",
| "filesHardLimit": "2K",
| "filesGracePeriod": "null",
| "blockGracePeriod": "null"
| }' "https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/quotas"

```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "job": [
| {
| "result": {},
| "request": {
| "type": "POST",
| "url": "/scalemgmt/v2/filesystems/gpfs0/quotas",
| "data": "{
| \"operationType\": \"setQuota\",
| \"quotaType\": \"user\",
| \"objectName\": \"adam\",
| \"blockSoftLimit\": \"1M\",
| \"blockHardLimit\": \"2M\",
| \"filesSoftLimit\": \"1K\",
| \"filesHardLimit\": \"2K\",
| \"filesGracePeriod\": \"null\",
| \"blockGracePeriod\": \"null\"\"}
| },
| "jobId": "12345",
| "submitted": "2016-11-14 10.35.56",
| "completed": "2016-11-14 10.35.56",
| "status": "COMPLETED"
| }
|]
| }
|
| "mmsetquota command" on page 548
| Sets quota limits.
|
| "mmrepquota command" on page 517
| Displays file system user, group, and fileset quotas.
```

## | **Filesystems/{filesystemName}/snapshots: GET**

| Gets information about snapshots of a specific file system.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/{filesystemName}/snapshots** request gets information about snapshots in the specified file system. For more information about the fields in the data structures that are returned, see the topics “**mmcrsnapshot command**” on page 267 and “**mmlssnapshot command**” on page 439.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/snapshots`

| where

| **filesystems/filesystemName**

|     Specifies the file system of which the snapshot is taken. Required.

| **snapshots**

|     Specifies snapshot as the resource of this GET call. Required.

### | **Request headers**

| `Content-Type: application/json`

| `Accept: application/json`

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 51. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
```

```

| "paging":
| {
| "next": "URL"
| },
| snapshots: [
| {
| "config: {
| "snapshotName": "Snapshot"
| "filesystemName": "Device",
| "filesetName": "Fileset",
| "oid": "ID",
| "snapID": "ID",
| "status": "Status"
| "created": "DateTime",
| "quotas": "Quotas",
| "snapType": "Type",
| }
| }
|]
| }

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**  
|     Return status.

| **"message": "ReturnMessage",**  
|     The return message.

| **"code": *ReturnCode***  
|     The return code.

| **"paging"**  
|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

| **"snapshotName": "Snapshot"**  
|     The snapshot name.

| **"filesystemName": "Device"**  
|     The file system that is the target of the snapshot.

| **"filesetName": "Fileset"**  
|     For a fileset snapshot, the fileset that is a target of the snapshot.

| **"oid": "ID"**  
|     Internal identifier that is used for paging.

| **"snapID": "ID"**  
|     The snapshot ID.

| **"status": "Status"**  
|     The snapshot status.

| **"created": "DateTime"**  
|     The date and time when the snapshot was created.

| **"quotas": "Quotas"**  
|     Any quotas that are applied to the fileset.

| **"snapType": "Type"**  
|     The AFM type of the snapshot, including "afm\_snap", "afm\_recovery", "afm\_failover", "afm\_rpo", "afm\_baserpo", and "Invalid".

## Examples

The following example gets information about the snapshots of the file system gpfs0.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/snapshots'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/snapshots?lastId=1001"
 },
 "snapshots": [
 {
 "config": {
 "snapshotName": "snap1",
 "filesystemName": "gpfs0",
 "filesetName": "myFset1",
 "oid": "123",
 "snapID": "5",
 "status": "Valid",
 "created": "2017-01-09 14.55.37",
 "quotas": "string",
 "snapType": "string"
 }
 }
]
}

“mmcrsnapshot command” on page 267
Creates a snapshot of a file system or fileset at a single point in time.
“mmlssnapshot command” on page 439
Displays GPFS snapshot information.
```

## **Filesystems/{filesystemName}/snapshots: POST**

| Creates a file system snapshot.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The POST `filesystems/{filesystemName}/snapshots` command creates a snapshot of the specified file system.

### **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/snapshots`

| where

#### **filesystems/filesystemName**

| Specifies the file system of which the snapshot needs to be taken. Required.

#### **snapshots**

| Specifies snapshot as the resource of this POST call. Required.

### **Request headers**

| Content-Type: application/json

| Accept: application/json

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 52. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                           | Required/optional |
|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:                                                   | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to theIBM Spectrum Scale system to perform the requested operation. | Required.         |

### **Request data**

```
| {
| "snapshotName": ""
| }

| "snapshotName":
| Name of the snapshot to be created.
```

### **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
```

```

{
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**

| Return status.

| **"message": "ReturnMessage",**

| The return message.

| **"code": *ReturnCode***

| The return code.

| **"jobs":**

| An array of elements that describe jobs. Each element describes one job.

| **"result"**

| **"commands": "String"**

| Array of commands that are run in this job.

| **"progress": "String"**

| Progress information for the request.

| **"exitCode": "Exit code"**

| Exit code of command. Zero is success, nonzero denotes failure.

| **"stderr": "Error"**

| CLI messages from *stderr*.

| **"stdout": "String"**

| CLI messages from *stdout*.

| **"request"**

| **"type": "{GET | POST | PUT | DELETE}"**

| HTTP request type.

| **"url": "URL"**

| The URL through which the job is submitted.

| **"data": " "**

| Optional.

```

| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example creates a snapshot *snap2* of the file system *gpfs0*.

Request data:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{ \
 "snapshotName": "snap2"
}' 'https://198.51.100.1:443/scalermgmt/v2/filesystems/gpfs0/snapshots'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "POST",
 "url": "/scalermgmt/v2/filesystems/gpfs0/snapshots",
 "data": "{\"snapshotName\": \"snap2\"}"
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}

```

“[mmcrsnapshot command](#)” on page 267  
Creates a snapshot of a file system or fileset at a single point in time.

“[mmlssnapshot command](#)” on page 439  
Displays GPFS snapshot information.

## **Filesystems/{filesystemName}/snapshots/{snapshotName}: DELETE**

| Deletes a file system snapshot.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **DELETE filesystems/{filesystemName}/snapshots/{snapshotName}** command deletes the specified  
| file system snapshot. For more information on deleting snapshots, see “[mmdelsnapshot command](#)” on  
| page 304.

### **Request URL**

| Use this URL to delete a global snapshot:

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/snapshots/snapshotName`

| where:

#### **filesystems/filesystemName**

| Specifies the file system of which the snapshot is deleted. Required.

#### **snapshots/snapshotName**

| Specifies snapshot to be deleted. Required.

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 53. List of request parameters*

| Parameter name | Description and applicable keywords                                                         | Required/optional |
|----------------|---------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote: | Required.         |
| snapshotName   | The snapshot name. This is the path of the snapshot.                                        | Required.         |

### **Request data**

| No request data.

### **Response data**

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 }
 }
]
}
```

```

 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**

|     Return status.

| **"message": "ReturnMessage",**

|     The return message.

| **"code": *ReturnCode***

|     The return code.

| **"jobs":**

|     An array of elements that describe jobs. Each element describes one job.

| **"result"**

|     **"commands": "String"**

|         Array of commands that are run in this job.

|     **"progress": "String"**

|         Progress information for the request.

|     **"exitCode": "Exit code"**

|         Exit code of command. Zero is success, nonzero denotes failure.

|     **"stderr": "Error"**

|         CLI messages from *stderr*.

|     **"stdout": "String"**

|         CLI messages from *stdout*.

| **"request"**

|     **"type": "{GET | POST | PUT | DELETE}"**

|         HTTP request type.

|     **"url": "URL"**

|         The URL through which the job is submitted.

|     **"data": " "**

|         Optional.

|     **"jobId": "ID"**,

|         The unique ID of the job.

|     **"submitted": "Time"**

|         The time at which the job was submitted.

|     **"completed": "Time"**

|         The time at which the job was completed.

```
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.
```

## | Examples

| The following example deletes the fileset snapshot *snap1* from the file system *gpfs0*.

| Request URL:

```
| curl -k -u admin:admin001 -X DELETE --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/snapshots/snap1'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "job": [
| {
| "result": {},
| "request": {
| "type": "DELETE",
| "url": "/scalemgmt/v2/filesystems/gpfs0/snapshots/snap1",
| },
| "jobId": "12345",
| "submitted": "2016-11-14 10.35.56",
| "completed": "2016-11-14 10.35.56",
| "status": "COMPLETED"
| }
|]
| }
```

| “mmcrsnapshot command” on page 267  
| Creates a snapshot of a file system or fileset at a single point in time.  
| “mmlssnapshot command” on page 439  
| Displays GPFS snapshot information.

## | **Filesystems/{filesystemName}/snapshots/{snapshotName}: GET**

| Gets information about a specific file system snapshot.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET filesystems/filesystemName/snapshots/snapshotName** request gets information about a specific file system snapshot. For more information about the fields in the data structures that are returned, see the topics “mmcrsnapshot command” on page 267 and “mmlsnapshot command” on page 439.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/filesystems/filesystemName/snapshots/snapshotName`

| where

#### | **filesystems/filesystemName**

| Specifies the file system of which the snapshot is taken. Required.

#### | **snapshots/snapshotName**

| Specifies a particular snapshot as the resource of this GET call. Required.

### | **Request headers**

| Content-Type: application/json

| Accept: application/json

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 54. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| filesystemName | The file system name. You can also use keywords such as :all:, :all_local:, or :all_remote:      | Required.         |
| snapshotName   | The snapshot name. This is the path of the snapshot.                                             | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |

### | **Request data**

| No request data.

### | **Response data**

```
| {
| "status": {
| "code":ReturnCode,
| "message":ReturnMessage
| },
| "paging":
| {
```

```

 "next": "URL"
 },
 snapshots: [
 {
 "config: {
 "snapshotName": "Snapshot",
 "filesystemName": "Device",
 "filesetName": "Fileset",
 "oid": "ID",
 "snapID": "ID",
 "status": "Status",
 "created": "DateTime",
 "quotas": "Quotas",
 "snapType": "Type",
 }
 }
]
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**

|     Return status.

| **"message": "ReturnMessage",**  
|     The return message.

| **"code": *ReturnCode***  
|     The return code.

| **"paging"**

|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by  
|     the query.

| **"snapshotName": "Snapshot"**  
|     The snapshot name.

| **"filesystemName": "Device"**  
|     The file system that is the target of the snapshot.

| **"filesetName": "Fileset"**  
|     For a fileset snapshot, the fileset that is a target of the snapshot.

| **"oid": "ID"**  
|     Internal identifier that is used for paging.

| **"snapID": "ID"**  
|     The snapshot ID.

| **"status": "Status"**  
|     The snapshot status.

| **"created": "DateTime"**  
|     The date and time when the snapshot was created.

| **"quotas": "Quotas"**  
|     Any quotas that are applied to the fileset.

| **"snapType": "Type"**  
|     The AFM type of the snapshot, including "afm\_snap", "afm\_recovery", "afm\_failover", "afm\_rpo",  
|     "afm\_baserpo", and "Invalid".

## Examples

| The following example gets information about the snapshots *snap1* of the file system *gpfs0*.

| Request data:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/filesets/myFset1/snapshots/snap1?lastId=1001"
 },
 "snapshots": [
 {
 "config": {
 "snapshotName": "snap1",
 "filesystemName": "gpfs0",
 "filesetName": "myFset1",
 "oid": "123",
 "snapID": "5",
 "status": "Valid",
 "created": "2017-01-09 14.55.37",
 "quotas": "string",
 "snapType": "string"
 }
 }
]
}
```

| “mmcrsnapshot command” on page 267

| Creates a snapshot of a file system or fileset at a single point in time.

| “mmlssnapshot command” on page 439

| Displays GPFS snapshot information.

## | Info: GET

| Gets information about IBM Spectrum Scale management API version and supported commands.

## | Availability

| Available on all IBM Spectrum Scale editions.

## | Description

| The **GET info** request gets information about IBM Spectrum Scale management API. The information includes the version number, the resources that are supported, and the HTTP methods that are supported for each resource.

## | Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/info`

| where

### | **info**

| Is the information resource.

## | Request headers

| Content-Type: application/json

| Accept: application/json

## | Request data

| No request data.

## | Response data

```
| {
| "status": {
| "message": "ReturnMessage",
| "code": "ReturnCode"
| }
| "info": {
| "restAPIVersion": "Version",
| "name": "API name",
| "path": "Resource path"
| }
| }

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": ReturnCode
| The return code.

| "info":
| Information about the REST API server.

| "restAPIVersion": "Version"
| The IBM Spectrum Scale management API version.

| "name": "APIName"
| The name of the API, such as "Spectrum Scale REST Management".
```

```
| "path": "Resource path
| Supported resources and elements.
```

## Examples

| The following example gets information about the REST API server.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/info'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "info" : {
 "name" : "Spectrum Scale REST Management",
 "paths" : {
 "/ces/addresses" : ["GET"],
 "/ces/addresses/{cesAddress}" : ["GET"],
 "/ces/services" : ["GET"],
 "/ces/services/{service}" : ["GET"],
 "/cluster" : ["GET"],
 "/config" : ["GET"],
 "/filesystems" : ["GET"],
 "/filesystems/{filesystemName}" : ["GET"],
 "/filesystems/{filesystemName}/acl/{path}" : ["GET", "PUT"],
 "/filesystems/{filesystemName}/disks" : ["GET"],
 "/filesystems/{filesystemName}/disks/{diskName}" : ["GET"],
 "/filesystems/{filesystemName}/filesets" : ["GET", "POST"],
 "/filesystems/{filesystemName}/filesets/{filesetName}" : ["GET", "PUT", "DELETE"],
 "/filesystems/{filesystemName}/filesets/{filesetName}/link" : ["POST", "DELETE"],
 "/filesystems/{filesystemName}/filesets/{filesetName}/quotas" : ["GET", "POST"],
 "/filesystems/{filesystemName}/filesets/{filesetName}/snapshots" : ["GET", "POST"],
 "/filesystems/{filesystemName}/filesets/{filesetName}/snapshots/{snapshotName}" : ["GET", "DELETE"],
 "/filesystems/{filesystemName}/owner/{path}" : ["GET", "PUT"],
 "/filesystems/{filesystemName}/quotas" : ["GET", "POST"],
 "/filesystems/{filesystemName}/snapshots" : ["GET", "POST"],
 "/filesystems/{filesystemName}/snapshots/{snapshotName}" : ["GET", "DELETE"],
 "/info" : ["GET"],
 "/jobs" : ["GET"],
 "/jobs/{jobId}" : ["GET"],
 "/nfs(exports" : ["GET", "POST"],
 "/nfs(exports/{exportPath}" : ["GET", "PUT", "DELETE"],
 "/nodes" : ["GET"],
 "/nodes/{name}" : ["GET"],
 "/nodes/{name}/health/events" : ["GET"],
 "/nodes/{name}/health/states" : ["GET"],
 "/nsds" : ["GET"],
 "/nsds/{nsdsName}" : ["GET"],
 "/smb/shares" : ["GET", "POST"],
 "/smb/shares/{shareName}" : ["GET", "PUT", "DELETE"]
 },
 "restApiVersion" : "2.0.0"
 },
 "status" : {
 "code" : 200,

```

## Jobs: GET

| Gets details about the asynchronous jobs.

### Availability

| Available on all IBM Spectrum Scale editions.

### Description

| The GET **jobs** request gets information about the asynchronous jobs.

### Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/jobs`

| where

#### jobs

| Specifies running jobs as the resource of this GET call. Required.

### Request headers

| Content-Type: application/json

| Accept: application/json

### Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 55. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

### Request data

| No request data.

### Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "jobs": [
| {
| "result": "",
| {
| "commands": "String",
| }
| }
|]
| }
```

```

 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "''",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**

| Return status.

| **"message": "ReturnMessage",**

| The return message.

| **"code": *ReturnCode***

| The return code.

| **"paging"**

| The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

| **"jobs":**

| An array of elements that describe jobs. Each element describes one job.

| **"result"**

| **"commands": "String"**

| Array of commands that are run in this job.

| **"progress": "String"**

| Progress information for the request.

| **"exitCode": "Exit code"**

| Exit code of command. Zero is success, nonzero denotes failure.

| **"stderr": "Error"**

| CLI messages from *stderr*.

| **"stdout": "String"**

| CLI messages from *stdout*.

| **"request"**

| **"type": "{GET | POST | PUT | DELETE}"**

| HTTP request type.

| **"url": "URL"**

| The URL through which the job is submitted.

| **"data": " "**

| Optional.

```

| " jobId": "ID",
| The unique ID of the job.
| " submitted": "Time"
| The time at which the job was submitted.
| " completed": "Time"
| The time at which the job was completed.
| " status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

| The following example gets information about the jobs.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/jobs'
```

| The request URL with no field or filter parameter returns only the details that uniquely identify the object.

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```

| {
| "jobs" : [{
| "jobId" : 1000000000001,
| "status" : "COMPLETED",
| "submitted" : "2017-03-21 15:49:20,100",
| "completed" : "2017-03-21 15:49:24,251"
| }, {
| "jobId" : 1000000000002,
| "status" : "COMPLETED",
| "submitted" : "2017-03-21 15:52:13,776",
| "completed" : "2017-03-21 15:52:16,030"
| }],
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }
```

| Using the field parameter ":all:" returns entire details of the jobs. For example:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/jobs?fields=:all:'
| {
| "jobs" : [{
| "jobId" : 1000000000001,
| "status" : "COMPLETED",
| "submitted" : "2017-03-21 15:49:20,100",
| "completed" : "2017-03-21 15:49:24,251",
| "request" : {
| "data" : {
| "snapshotName" : "mySnap1"
| },
| "type" : "POST",
| "url" : "/scalemgmt/v2/filesystems/gpfs0/snapshots"
| },
| "requestId" : "12345678901234567890123456789012"
| }]
| }
```

```

 "result" : {
 "progress" : [],
 "commands" : ["mmcrlsnapshot 'gpfs0' 'mySnap1' "],
 "stdout" : ["EFSSG0019I The snapshot mySnap1 has been successfully created."],
 "stderr" : [],
 "exitCode" : 0
 }
 }, {
 "jobId" : 1000000000002,
 "status" : "COMPLETED",
 "submitted" : "2017-03-21 15:52:13,776",
 "completed" : "2017-03-21 15:52:16,030",
 "request" : {
 "data" : {
 "snapshotName" : "mySnap1"
 },
 "type" : "POST",
 "url" : "/scalemgmt/v2/filesystems/gpfs0/filesets/fset1/snapshots"
 },
 "result" : {
 "progress" : [],
 "commands" : ["mmcrlsnapshot 'gpfs0' 'mySnap1' -j 'fset1' "],
 "stdout" : ["EFSSG0019I The snapshot mySnap1 has been successfully created."],
 "stderr" : [],
 "exitCode" : 0
 }
 }
],
"status" : {
 "code" : 200,
 "message" : "The request finished successfully"
}
}

```

## Jobs/{jobID}: GET

| Gets details about an asynchronous job.

### Availability

| Available on all IBM Spectrum Scale editions.

### Description

| The GET jobs/jobID request gets information about the asynchronous job that is specified in the request.

### Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/jobs/jobID`

| where

#### jobs

| Specifies running jobs as the resource of this GET call. Required.

#### jobs/jobID

| Specifies the job about which you need to get information. Required.

### Request headers

| Content-Type: application/json  
| Accept: application/json

### Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 56. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| jobId          | Unique identifier of the asynchronous job.                                                       | Required.         |
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |

### Request data

| No request data.

### Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "jobs": [
| {
| "result": "",
| {
| "key": "value"
| }
| }
|]
| }
```

```

 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": """",
 }
}
"jobId": "ID",
"submitted": "Time",
"completed": "Time",
"status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**

|     Return status.

| **"message": "ReturnMessage",**

|     The return message.

| **"code": *ReturnCode***

|     The return code.

| **"paging"**

|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by  
|     the query.

| **"jobs":**

|     An array of elements that describe jobs. Each element describes one job.

| **"result"**

|     **"commands": "String"**

|         Array of commands that are run in this job.

|     **"progress": "String"**

|         Progress information for the request.

|     **"exitCode": "Exit code"**

|         Exit code of command. Zero is success, nonzero denotes failure.

|     **"stderr": "Error"**

|         CLI messages from *stderr*.

|     **"stdout": "String"**

|         CLI messages from *stdout*.

| **"request"**

|     **"type": "{GET | POST | PUT | DELETE}"**

|         HTTP request type.

|     **"url": "URL"**

|         The URL through which the job is submitted.

|     **"data": " "**

|         Optional.

```

| " jobId": "ID",
| The unique ID of the job.
| " submitted": "Time"
| The time at which the job was submitted.
| " completed": "Time"
| The time at which the job was completed.
| " status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## | Known limitation

| If there are more than one GUI nodes in a cluster, you need to use the filter parameter to get the details  
| of the particular job. For example, if a job is submitted on GUI2 and it returns a jobId=20000000000003. If  
| you query GUI1 for that job ID using `/jobs/20000000000003`, it returns "Invalid request" because it cannot  
| find the job since it was submitted to the other GUI. A workaround is to use  
| `"/jobs?filter=jobId=20000000000003"` instead of `"/jobs/20000000000003"` in the request URL to retrieve the job  
| correctly from the other GUI.

## | Examples

| The following example gets information about the job 12345.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/jobs/12345'
```

| **Note:** If there are more than one GUI nodes in the cluster, use the following request URL to retrieve the  
| details of the job:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/jobs?filter=jobId=12345'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/filesystems/gpfs0/filesets?lastId=1001"
 },
 "jobs": [
 {
 "result": {
 "commands": "[\"mmcrfileset gpfs0 restfs1001\", ...]",
 "progress": "[\"(2/3) Linking fileset\"]",
 "exitCode": "0",
 "stderr": "[\"EFSSG0740C There are not enough resources available to create
a new independent file set.\", ...]",
 "stdout": "[\"EFSSG4172I The file set {0} must be independent.\", ...]"
 },
 "request": {
 "type": "GET",
 "url": "/scalemgmt/v2/filesystems/gpfs0/filesets",
 "data": "{\"config\":{\"filesetName\":\"restfs1001\",\"owner\":\"root\",\"path\":"
 }
 }
]
}
```

```
| \"/mnt/gpfs0/rest1001\", \"permissions\": \"555\"}"
| },
| \"jobId\": \"12345\",
| \"submitted\": \"2016-11-14 10.35.56\",
| \"completed\": \"2016-11-14 10.35.56\",
| \"status\": \"COMPLETED"
| }
| }
```

## | NFS/exports: GET

| Gets information about NFS exports configured in the system.

## | Availability

| Available on all IBM Spectrum Scale editions.

## | Description

| The **GET nfs/exports** request gets information about NFS exports. For more information about the fields  
| in the data structures that are returned, see “mmnfs command” on page 454.

## | Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/nfs/exports`

| where

### | nfs/exports

|     Specifies NFS export as the resource. Required.

## | Request headers

| `Accept: application/json`

## | Request data

| No request data.

## | Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 57. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

## | Response data

| The following list of attributes are available in the response data:

```
| {
| "status":
| {
| "code": ReturnCode
| "message": "ReturnMessage",
| }
| "exports":
| {
| "config": [
| {
| "filesystemName": "FileSystemName"
| }
|]
| }
| }
```

```

 "path": "Path",
 "delegations": "Delegations",
],
 nfsClients:[
 {
 "clientName": "String",
 "exportID": "Path",
 "access_type": "none | RW | RO | MDONLY | MDONLY_RO",
 "squash": "root | root_squash | all | all_squash | allsquash | no_root_squash | none"
 "anonUid": "ID",
 "anonGid": "ID",
 "privilegedPort": "True | False",
 "sectype": "sys | krb5 | krb5i | krb5p | None",
 "protocols": "NFS3 | NFS4",
 "transportProtocol": "TCP | UDP",
 "delegations": "none | read | write | readwrite",
 "manageGids": "True | False",
 "nfsCommit": "True | False",
 }
 }
}

"status":
 Return status.

"code": ReturnCode,
 The HTTP status code that was returned by the request.

"message": "ReturnMessage"
 The return message.

"exports":
 An array of information about NFS exports. The array contains elements that describe the NFS exports. For more information about the fields in this structure, see the links at the end of this topic.

 config:
 Manages NFS configuration in a cluster.

 "filesystemName": "FileSystemName"
 The name of the file system to which the export belongs.

 "path": "Path",
 Specifies the path for the export.

 "delegations": "Delegations"
 Specifies what delegate file operations are permitted.

 nfsClients

 "clientName": "String"
 The host name or IP address of the NFS server to be accessed.

 "exportID": "Path"
 The internal unique ID of the export.

 "access_type": "none | RW | RO | MDONLY | MDONLY_RO"
 Specifies the type of the access for the client.

 "squash": "root | root_squash | all | all_squash | allsquash | no_root_squash | none"
 Specifies whether the squashing mechanism is applied to the connecting client.

 "anonUid": "ID"
 This option explicitly sets the UID of the anonymous account.

 "anonGid": "ID"
 This option explicitly sets the GID of the anonymous account.

```

```

| "privilegedPort": "True | False"
| This option to specify whether a client that uses an ephemeral port must be rejected.
| "sectype": "sys | krb5 | krb5i | krb5p | None"
| The supported authentication method for the client.
| "protocols": "NFS3 | NFS4 "
| Supported NFS protocol version.
| "transportProtocol": "TCP | UDP"
| Specifies the transport layer.
| "delegations": "none | read | write | readwrite"
| Specifies what delegate file operations are permitted.
| "manageGids": "True | False"
| Whether this NFS client is allowed to manage GIDs.
| "nfsCommit": "True | False"
| Whether to commit the transmitted data on the server side.

```

## Examples

The following example gets information about the NFS exports that are configured in the system:

Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nfs/exports'
```

The request URL with no field or filter parameter returns only the details that uniquely identify the object.

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "nfsexports" : [{
 "config" : {
 "path" : "/mnt/gpfs0/nfs01"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}
```

Using the field parameter ":all:" returns entire details of the NFS exports as shown in the following example:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nfs/exports?fields=:all:'
{
 "nfsexports" : [{
 "config" : {
 "delegations" : "NONE",
 "filesetName" : "",
 "filesystemName" : "gpfs0",
 "path" : "/mnt/gpfs0/nfs01"
 }
 }]
}
```

```
| },
| "nfsClients" : [{
| "access_type" : "RW",
| "anonGid" : "-2",
| "anonUid" : "-2",
| "clientName" : "198.51.100.123",
| "delegations" : "NONE",
| "exportid" : "1",
| "manageGids" : false,
| "nfsCommit" : false,
| "privilegedPort" : false,
| "protocols" : "3,4",
| "sectype" : "SYS",
| "squash" : "ROOT_SQUASH",
| "transportProtocol" : "TCP"
| }
|],
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }
```

| **Related reference:**

- | “[mmnfs command](#)” on page 454
- | Manages NFS exports and configuration.

## | NFS/exports: POST

| Creates an NFS export.

## | Availability

| Available on all IBM Spectrum Scale editions.

## | Description

| The **POST nfs/export** request creates a new NFS export. For more information about the fields in the data structures that are returned, see “mmnfs command” on page 454.

## | Request URL

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/nfs/exports`

| where

### | nfs/exports

|     Specifies the NFS export as the resource. Required.

## | Request headers

| `Accept: application/json`

## | Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 58. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| body           | Body of the request that contains the required parameters to be passed on to the IBM Spectrum Scale system to perform the requested operation. | Required.         |

## | Request data

```
| {
| "path": "Path",
| "nfsclients": "Client details"
| }
```

| **"path": "Path "**  
|     Specifies the path of the NFS export.

### | nfsclients

|     Details of the NFS clients.

## | Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
| {
| "result": "",
| {
```

```

 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": """
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
}
],
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

| **"status":**

|     Return status.

| **"message": "ReturnMessage",**

|     The return message.

| **"code": *ReturnCode***

|     The return code.

| **"jobs":**

|     An array of elements that describe jobs. Each element describes one job.

| **"result"**

|     **"commands": "String"**

|         Array of commands that are run in this job.

|     **"progress": "String"**

|         Progress information for the request.

|     **"exitCode": "Exit code"**

|         Exit code of command. Zero is success, nonzero denotes failure.

|     **"stderr": "Error"**

|         CLI messages from *stderr*.

|     **"stdout": "String"**

|         CLI messages from *stdout*.

| **"request"**

|     **"type": "{GET | POST | PUT | DELETE}"**

|         HTTP request type.

|     **"url": "URL"**

|         The URL through which the job is submitted.

|     **"data": " "**

|         Optional.

| **"jobId": "ID",**

|         The unique ID of the job.

```

| "submitted": "Time"
| The time at which the job was submitted.
| "completed": "Time"
| The time at which the job was completed.
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example creates an NFS export.

Request URL:

```

curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "path": "/mnt/gpfs0/fset1",
 "nfsClients": "['198.51.100.111:443(access_type=ro,squash=no_root_squash)',",
 "'198.51.100.110:443(access_type=rw,squash=no_root_squash)']"
}'
'https://198.51.100.1:443/scalemgmt/v2/nfs/exports'

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {
 "commands": "['mmcrfileset gpfs0 restfs1001', ...]",
 "progress": "['(2/3) Linking fileset']",
 "exitCode": "0",
 "stderr": "['EFSSG0740C There are not enough resources available to create
 a new independent file set.', ...]",
 "stdout": "['EFSSG4172I The file set {0} must be independent.', ...]"
 },
 "request": {
 "type": "POST",
 "url": "/scalemgmt/v2/filesystems/gpfs0/filesets",
 "data": {
 "path": "/mnt/gpfs0/fset1",
 "nfsClients": "['198.51.100.111:443(access_type=ro,squash=no_root_squash)',",
 "'198.51.100.110:443(access_type=rw,squash=no_root_squash)']"
 }
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}

```

**Related reference:**

[“mmnfs command” on page 454](#)

Manages NFS exports and configuration.

## | **NFS/exports/{exportPath}: GET**

| Gets information about an NFS export that is configured in the system.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET nfs/exports/exportPath** request gets information about an NFS export. For more information  
| about the fields in the data structures that are returned, see “mmnfs command” on page 454.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nfs/exports`

| where

#### | **nfs/exports**

|     Specifies NFS export as the resource. Required.

#### | **exportPath**

|     Specifies the NFS export about which you need the details. Required.

### | **Request headers**

| `Accept: application/json`

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 59. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                                                  | Required/optional |
|----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields.                                                                                                                                                                     | Optional.         |
| exportPath     | The path of the NFS export. The export path is relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to %2F in the request URL. | Required.         |

### | **Request data**

| No request data.

### | **Response data**

| The following list of attributes are available in the response data:

```
| {
| "status":
| {
| "code": ReturnCode
```

```

 "message": "ReturnMessage",
}
"exports":
{
 "config": [
 {
 "filesystemName": "FileSystemName",
 "path": "Path",
 "delegations": "Delegations",
 }],
 nfsClients:[
 {
 "clientName": "String",
 "exportID": "Path",
 "access_type": "none | RW | RO | MDONLY | MDONLY_RO",
 "squash": "root | root_squash | all | all_squash | allsquash | no_root_squash | none",
 "anonUid": "ID",
 "anonGid": "ID",
 "privilegedPort": "True | False",
 "sectype": "sys | krb5 | krb5i | krb5p | None",
 "protocols": "NFS3 | NFS4",
 "transportProtocol": "TCP | UDP",
 "delegations": "none | read | write | readwrite",
 "manageGids": "True | False",
 "nfsCommit": "True | False",
 }
 }
}

"status":
 Return status.

"code": ReturnCode,
 The HTTP status code that was returned by the request.

"message": "ReturnMessage"
 The return message.

"exports":
 An array of information about NFS exports. The array contains elements that describe the NFS exports. For more information about the fields in this structure, see the links at the end of this topic.

 config:
 Manages NFS configuration in a cluster.

 "filesystemName": "FileSystemName"
 The name of the file system to which the export belongs.

 "path": "Path",
 Specifies the path for the export.

 "delegations": "Delegations"
 Specifies what delegate file operations are permitted.

 nfsClients
 "clientName": "String"
 The host name or IP address of the NFS server to be accessed.

 "exportID": "Path"
 The internal unique ID of the export.

 "access_type": "none | RW | RO | MDONLY | MDONLY_RO"
 Specifies the type of the access for the client.

 "squash": "root | root_squash | all | all_squash | allsquash | no_root_squash | none"
 Specifies whether the squashing mechanism is applied to the connecting client.

```

```

| "anonUid": "ID"
| This option explicitly sets the UID of the anonymous account.
|
| "anonGid": "ID"
| This option explicitly sets the GID of the anonymous account.
|
| "privilegedPort": "True | False"
| This option to specify whether a client that uses an ephemeral port must be rejected.
|
| "sectype": "sys | krb5 | krb5i | krb5p | None"
| The supported authentication method for the client.
|
| "protocols": "NFS3 | NFS4 "
|
| Supported NFS protocol version.
|
| "transportProtocol": "TCP | UDP"
| Specifies the transport layer.
|
| "delegations": "none | read | write | readwrite"
| Specifies what delegate file operations are permitted.
|
| "manageGids": "True | False"
| Whether this NFS client is allowed to manage GIDs.
|
| "nfsCommit": "True | False"
|
| Whether to commit the transmitted data on the server side.

```

## Examples

The following example gets information about the NFS export /mnt/gpfs0/fset1 that is configured in the system.

Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffset1'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffset1?lastId=1001"
 },
 "exports": [
 {
 "config": {
 "filesystemName": "gpfs0",
 "path": "/mnt/gpfs0/fset1",
 "delegations": "NONE"
 },
 "nfsClients": [
 {
 "clientName": "198.51.100.8",
 "exportid": "1",
 "access_type": "RO",
 "squash": "root_squash",
 }
]
 }
]
}
```

```
| "anonUid": "1",
| "anonGid": "1",
| "privilegedPort": "false",
| "sectype": "sys",
| "protocols": "NFS4",
| "transportProtocol": "TCP",
| "delegations": "none",
| "manageGids": "yes",
| "nfsCommit": "This is a comment"
| }
|]
| }
}
```

| **Related reference:**

- | “mmnfs command” on page 454
- | Manages NFS exports and configuration.

## | NFS/exports/{exportPath}: PUT

| Modifies an existing NFS export.

### | Availability

| Available on all IBM Spectrum Scale editions.

### | Description

| The PUT nfs/export/exportPath request modifies an existing NFS export. For more information about the fields in the data structures that are returned, see “mmnfs command” on page 454.

### | Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nfs/exports/exportPath`

| where

#### | nfs/exports

|     Specifies the NFS export as the resource. Required.

#### | exportPath

|     Specifies the NFS export, which you need to modify. Required.

### | Request headers

| Accept: application/json

### | Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 60. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                                                  | Required/optional |
|----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| exportPath     | The path of the NFS export. The export path is relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to %2F in the request URL. | Required.         |
| body           | Body of the request that contains the required parameters to be passed on to theIBM Spectrum Scale system to perform the requested operation.                                                                                                                        | Required.         |

### | Request data

```
| {
| "nfsadd": "New client",
| "nfschange": "Change export",
| "nfsremove": "Remove NFS client",
| "nfsposition": "NFS position"
| }
```

| Detailed explanation of the fields are given in the following list:

```

| "nfsadd": "New client"
| Adds a new client declaration for the specified path. Client options can be a list of one or more client
| definitions.

| "nfschange": "Change export"
| Modifies an existing client declaration for the specified path. Client options can be a list of one or
| more client definitions.

| "nfsremove": "Remove NFS client"
| Removes the NFS client specified by client from the export configuration for the path.

| "nfsposition": "NFS position"
| It reorders the client declaration section within the export declaration file. This option can be used
| only together with --nfsadd or --nfschange.

```

## Response data

```

{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 jobs: [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 },
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

```

"status":
 Return status.

"message": "ReturnMessage",
 The return message.

"code": "ReturnCode"
 The return code.

"jobs":
 An array of elements that describe jobs. Each element describes one job.

"result"

"commands": "String"
 Array of commands that are run in this job.

```

```

| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example modifies the NFS export /mnt/gpfs0/fset1 that is configured in the system.

Request URL:

```

curl -k -u admin:admin001 -X PUT --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "nfsadd": "198.51.100.10:443(sectype=sys,krb5)",
 "nfschange": "198.51.100.11:443(sectype=sys,krb5)",
 "nfsremove": "198.51.100.21:443",
 "nfsposition": "1"
}' "https://198.51.100.1:443/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffs1"

```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "PUT",
 "url": "/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffs1",
 }
 }
]
}

```

```
| "data": "{\n| \"nfsadd\": \"198.51.100.10:443(sectype=sys,krb5)\",\n| \"nfschange\": \"198.51.100.11:443(sectype=sys,krb5)\",\n| \"nfsremove\": \"198.51.100.21:443\",\n| \"nfsposition\": \"1\"\"}\n| },\n| \"jobId\": \"12345\",\n| \"submitted\": \"2016-11-14 10.35.56\",\n| \"completed\": \"2016-11-14 10.35.56\",\n| \"status\": \"COMPLETED\"\n| }\n|]\n| }
```

| **Related reference:**

- | “mmnfs command” on page 454
- | Manages NFS exports and configuration.

## **NFS/exports/{exportPath}: DELETE**

| Deletes an NFS export.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **DELETE nfs/exports/exportPath** command deletes the specified NFS export. For more information, | see “mmnfs command” on page 454.

### **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nfs/exports/exportPath`

| where:

#### **nfs/exports**

| Specifies NFS export as the resource. Required.

#### **exportPath**

| Specifies the NFS export to be deleted. Required.

### **Request headers**

| `Accept: application/json`

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 61. List of request parameters*

| Parameter name | Description and applicable keywords                                                                                                                                                                                                   | Required/optional |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| path           | The file path relative to file system's mount point. The path of the file or directory is specified with forward slashes (/). For example: mnt/gpfs0/rest01. The forward slashes in the path are encoded to "%2F" in the request URL. | Required.         |

### **Response data**

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 },
],
}
```

```

 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
}
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**  
|     Return status.

| **"message": "ReturnMessage",**  
|     The return message.

| **"code": *ReturnCode***  
|     The return code.

| **"jobs":**  
|     An array of elements that describe jobs. Each element describes one job.

| **"result"**

|     **"commands": "String"**  
|         Array of commands that are run in this job.

|     **"progress": "String"**  
|         Progress information for the request.

|     **"exitCode": "Exit code"**  
|         Exit code of command. Zero is success, nonzero denotes failure.

|     **"stderr": "Error"**  
|         CLI messages from *stderr*.

|     **"stdout": "String"**  
|         CLI messages from *stdout*.

| **"request"**

|     **"type": "{GET | POST | PUT | DELETE}"**  
|         HTTP request type.

|     **"url": "URL"**  
|         The URL through which the job is submitted.

|     **"data": " "**  
|         Optional.

| **"jobId": "ID",**  
|     The unique ID of the job.

| **"submitted": "Time"**  
|     The time at which the job was submitted.

| **"completed": "Time"**  
|     The time at which the job was completed.

| **"status": "RUNNING | COMPLETED | FAILED"**  
|     Status of the job.

## Examples

The following example deletes the NFS export /mnt/gpfs0/fset1.

```
curl -k -u admin:admin001 -X DELETE --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffset1'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "DELETE",
 "url": "/scalemgmt/v2/nfs/exports/mnt%2Fgpfs0%2Ffset1",
 "data": "{}"
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}
```

**Related reference:**

“mmnfs command” on page 454  
Manages NFS exports and configuration.

## Nodes: GET

| Gets information about nodes in the cluster.

## Availability

| Available on all IBM Spectrum Scale editions.

## Description

| The **GET nodes** request gets information about all nodes in the cluster. For more information about the fields in the data structures that are returned, see the topics “[mmchnode command](#)” on page 190 and “[mmgetstate command](#)” on page 350.

## Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nodes`

| where

### nodes

| Specifies nodes as the resource of this GET call. Required.

## Request headers

| Content-Type: application/json

| Accept: application/json

## Request data

| No request data.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "paging": {
 "next": "URL"
 },
 "nodes": [
 {
 "adminNodeName": "IP or host name",
 "nodeNumber": "Node ID",
 "config": {
 "adminLoginName": "Admin login name",
 "designatedLicense": "Designated license",
 "requiredLicense": "Required license",
 },
 "status": {
 "osName": "Operating system",
 "nodeState": "Health status",
 "gpfsState": "GPFS status",
 "productVersion": "Version",
 },
 "network": {
 "adminIPAddress": "IP address or host name",
 "daemonNodeName": "GPFS daemon node name",
 "daemonIPAddress": "IP address",
 }
 }
]
}
```

```

 "getcnfsNodeName": "Host name used by cNFS",
 },
 "roles":
 {
 "snmpNode": "yes | no",
 "managerNode": "yes | no",
 "gatewayNode": "yes | no",
 "cnfsNode": "yes | no",
 "clientNode": "yes | no",
 "cesNode": "yes | no",
 "quorumNode": "yes | no",
 "cloudGatewayNode": "yes | no",
 "otherNodeRoles": "Roles",
 "designation": "quorum | quorumManager | manager",
 },
 "cnfsInfo":
 {
 "cnfsState": "cNFS state",
 "cnfsGroupId": "cNFS group ID",
 "cnfsIplist": "cNFS IP addresses",
 },
 "cesInfo":
 {
 "cesState": "CES state",
 "cesGroup": "CES group",
 "cesIpList": "CES IPs",
 "ipAddress": "IP address of this node",
 }
}
]
}
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

| **"status":**

|     Return status.

| **"message": "ReturnMessage",**

|     The return message.

| **"code": "ReturnCode**

|     The return code.

| **"paging"**

|     The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

| **"nodes":**

| **"adminNodeName": "IP or host name"**

|     The host name that is used by the GPFS admin network.

| **"nodeNumber": "Node ID"**

|     The GPFS node ID.

| **"config":**

| **"adminLoginName": "Admin login name"**

|     The name of the admin login.

| **"designatedLicense": "Designated license"**

|     The license this node is running on.

| **"requiredLicense": "client | FPO | server"**

|     Controls the type of GPFS required license that is associated with the nodes in the cluster.

```

| "status":
| "osName": "Operating system"
| The name of Operating System running on this node.
| "nodeState": "Health status"
| The state of the node as reported by the mmhealth node show command.
| "gpfsState": "GPFS status"
| The state of GPFS on this node as reported the mmhealth node show command.
| "productVersion": "Version"
| The IBM Spectrum Scale version that is installed on this node.

| "network":
| "adminIPAddress": "IP address or host name"
| The IP address that is used by the GPFS admin network.
| "daemonNodeName": "GPFS daemon node name"
| The host name that is used by the GPFS daemon network.
| "daemonIPAddress": "IP address"
| The IP address that is used by the GPFS daemon network.
| "getcnfsNodeName": "Host name used by cNFS"
| The hostname that is used by cNFS.

| "roles":
| "snmpNode": "yes | no"
| Specifies whether this node is an SNMP node.
| "managerNode": "yes | no"
| Specifies whether this node is a manager node.
| "gatewayNode": "yes | no"
| Specifies whether this node is a gateway node.
| "cnfsNode": "yes | no"
| Specifies whether this node is a cNFS node.
| "clientNode": "yes | no"
| Specifies whether this node is a client node.
| "cesNode": "yes | no"
| Specifies whether this node is a CES node.
| "quorumNode": "yes | no"
| Specifies whether this node is a quorum node.
| "cloudGatewayNode": "yes | no"
| Specifies whether this node is a cloud gateway node.
| "otherNodeRoles": "Roles"
| Other roles of this node.
| "designation": "quorum | quorumManager |manager"
| The designated role.

| "cnfsInfo":
| "cnfsState": "cNFS state"
| "cnfsGroupId": "cNFS group ID"
| "cnfsIplist": "cNFS IP addresses"
|

```

```

| "cesInfo":
| """cesState""" : "CES state"
| The state of CNFS on this node.
| "cesGroup" : "CES group"
| The CNFS group ID.
| "cesIpList" : "CES IPs"
| A list of CNFS IP addresses.
| "ipAddress" : "IP address of this node"
| The actual IP address of this node.

```

## Examples

| The following example gets information about the nodes.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/nodes'
```

| The request URL with no field or filter parameter returns only the details that uniquely identify the object.

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "nodes" : [
 {
 "adminNodeName" : "mari-11.localnet.com"
 },
 {
 "adminNodeName" : "mari-12.localnet.com"
 },
 {
 "adminNodeName" : "mari-13.localnet.com"
 },
 {
 "adminNodeName" : "mari-14.localnet.com"
 },
 {
 "adminNodeName" : "mari-15.localnet.com"
 }
],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}
```

| Using the field parameter ":all:" returns entire details of the nodes that are available in the cluster. For example:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/nodes?fields=:all:'
{
 "nodes" : [
 {
 "nodeNumber" : 1,
 "adminNodeName" : "mari-11.localnet.com",
 "config" : {
 "adminLoginName" : "root",
 "designatedLicense" : "server",
 "requiredLicense" : "server"
 },
 "status" : {

```

```

 "gpfsState" : "HEALTHY",
 "nodeState" : "HEALTHY",
 "osName" : "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "productVersion" : "4.2.3.0"
},
"network" : {
 "adminIPAddress" : "198.52.100.11",
 "daemonIPAddress" : "198.52.100.11",
 "daemonNodeName" : "mari-11.localnet.com"
},
"roles" : {
 "cesNode" : false,
 "clientNode" : true,
 "cloudGatewayNode" : false,
 "cnfsNode" : false,
 "designation" : "quorum",
 "gatewayNode" : false,
 "managerNode" : false,
 "otherNodeRoles" : "perfmonNode",
 "quorumNode" : true,
 "snmpNode" : false
}
}, {
 "nodeNumber" : 2,
 "adminNodeName" : "mari-12.localnet.com",
 "config" : {
 "adminLoginName" : "root",
 "designatedLicense" : "server",
 "requiredLicense" : "server"
 },
 "status" : {
 "gpfsState" : "HEALTHY",
 "nodeState" : "HEALTHY",
 "osName" : "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "productVersion" : "4.2.3.0"
 },
 "network" : {
 "adminIPAddress" : "198.52.100.12",
 "daemonIPAddress" : "198.52.100.12",
 "daemonNodeName" : "mari-12.localnet.com"
 },
 "roles" : {
 "cesNode" : true,
 "clientNode" : true,
 "cloudGatewayNode" : false,
 "cnfsNode" : false,
 "designation" : "quorum",
 "gatewayNode" : false,
 "managerNode" : false,
 "otherNodeRoles" : "perfmonNode,cesNode,cloudNodeMarker",
 "quorumNode" : true,
 "snmpNode" : false
 },
 "cesInfo" : {
 "cesGroup" : "",
 "cesIpList" : "198.51.100.43,198.51.100.22,198.51.100.28,198.51.100.34,198.51.100.40,
 198.51.100.46,198.51.100.19,198.51.100.25,198.51.100.31,198.51.100.37",
 "cesState" : "e",
 "ipAddress" : "198.52.100.12"
 }
}, {
 "nodeNumber" : 3,
 "adminNodeName" : "mari-13.localnet.com",
 "config" : {
 "adminLoginName" : "root",
 "designatedLicense" : "server",
 "requiredLicense" : "server"
 }
}

```

```

},
"status" : {
 "gpfsState" : "HEALTHY",
 "nodeState" : "HEALTHY",
 "osName" : "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "productVersion" : "4.2.3.0"
},
"network" : {
 "adminIPAddress" : "198.52.100.13",
 "daemonIPAddress" : "198.52.100.13",
 "daemonNodeName" : "mari-13.localnet.com"
},
"roles" : {
 "cesNode" : true,
 "clientNode" : true,
 "cloudGatewayNode" : false,
 "cnfsNode" : false,
 "designation" : "quorum",
 "gatewayNode" : false,
 "managerNode" : false,
 "otherNodeRoles" : "perfmonNode,cesNode,cloudNodeMarker",
 "quorumNode" : true,
 "snmpNode" : false
},
"cesInfo" : {
 "cesGroup" : "",
 "cesIpList" : "198.51.100.41,198.51.100.47,198.51.100.20,198.51.100.26,198.51.100.32,
 198.51.100.38,198.51.100.44,198.51.100.23,198.51.100.29,198.51.100.35",
 "cesState" : "e",
 "ipAddress" : "198.52.100.13"
}
}, {
 "nodeNumber" : 4,
 "adminNodeName" : "mari-14.localnet.com",
 "config" : {
 "adminLoginName" : "root",
 "designatedLicense" : "server",
 "requiredLicense" : "server"
 },
 "status" : {
 "gpfsState" : "HEALTHY",
 "nodeState" : "HEALTHY",
 "osName" : "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "productVersion" : "4.2.3.0"
 },
 "network" : {
 "adminIPAddress" : "198.52.100.14",
 "daemonIPAddress" : "198.52.100.14",
 "daemonNodeName" : "mari-14.localnet.com"
 },
 "roles" : {
 "cesNode" : true,
 "clientNode" : true,
 "cloudGatewayNode" : false,
 "cnfsNode" : false,
 "designation" : "client",
 "gatewayNode" : false,
 "managerNode" : false,
 "otherNodeRoles" : "perfmonNode,cesNode,cloudNodeMarker",
 "quorumNode" : false,
 "snmpNode" : false
 },
 "cesInfo" : {
 "cesGroup" : "",
 "cesIpList" : "198.51.100.45,198.51.100.24,198.51.100.30,198.51.100.36,198.51.100.42,
 198.51.100.48,198.51.100.21,198.51.100.27,198.51.100.33,198.51.100.39",
 "cesState" : "e"
 }
}

```

```

 "ipAddress" : "198.52.100.14"
 }
}, {
 "nodeNumber" : 5,
 "adminNodeName" : "mari-15.localnet.com",
 "config" : {
 "adminLoginName" : "root",
 "designatedLicense" : "server",
 "requiredLicense" : "server"
 },
 "status" : {
 "gpfsState" : "HEALTHY",
 "nodeState" : "HEALTHY",
 "osName" : "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "productVersion" : "4.2.3.0"
 },
 "network" : {
 "adminIPAddress" : "198.52.100.15",
 "daemonIPAddress" : "198.52.100.15",
 "daemonNodeName" : "mari-15.localnet.com"
 },
 "roles" : {
 "cesNode" : true,
 "clientNode" : true,
 "cloudGatewayNode" : false,
 "cnfsNode" : false,
 "designation" : "client",
 "gatewayNode" : false,
 "managerNode" : false,
 "otherNodeRoles" : "perfmonNode,cesNode,cloudNodeMarker",
 "quorumNode" : false,
 "snmpNode" : false
 },
 "cesInfo" : {
 "cesGroup" : "",
 "cesIpList" : "198.51.100.10,198.51.100.14,198.51.100.12,198.51.100.18,198.51.100.16,
 198.51.100.49,198.51.100.11,198.51.100.15,198.51.100.13,198.51.100.17",
 "cesState" : "e",
 "ipAddress" : "198.52.100.15"
 }
},
"status" : {
 "code" : 200,
 "message" : "The request finished successfully"
}
}

```

“mmchnode command” on page 190

Changes node attributes.

“mmgetstate command” on page 350

Displays the state of the GPFS daemon on one or more nodes.

## Nodes/{name}: GET

Gets information about a node.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET nodes/nodeName** request gets information about the specified node. For more information about the fields in the data structures that are returned, see the topics “mmchnode command” on page 190 and “mmgetstate command” on page 350.

### Request URL

`https://<IP address or host name of API server>:<port>/scalemgmt/v2/nodes/nodeName`

where

#### nodes

Is the nodes resource. Required.

#### nodeName

Specifies the node about which you want to get information. Required.

### Request headers

Content-Type: application/json  
Accept: application/json

### Request data

No request data.

### Request parameters

The following parameters can be used in the request URL to customize the request:

*Table 62. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |
| name           | Name of the node.                                                                                | Required.         |

### Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "paging": {
 }
```

```

 "next": "URL"
 },
 nodes: [
 {
 "adminNodename": "IP or host name",
 "nodeNumber": "Node ID",
 "config": {
 {
 "adminLoginName": "Admin login name",
 "designatedLicense": "Designated license",
 "requiredLicense": "Required license",
 },
 "status": {
 {
 "osName": "Operating system",
 "nodeState": "Health status",
 "gpfsState": "GPFS status",
 "productVersion": "Version",
 },
 },
 "network": {
 {
 "adminIPAddress": "IP address or host name",
 "daemonNodeName": "GPFS daemon node name",
 "daemonIPAddress": "IP address",
 "getcnfsNodeName": "Host name used by cNFS",
 },
 },
 "roles": {
 {
 "snmpNode": "yes | no",
 "managerNode": "yes | no",
 "gatewayNode": "yes | no",
 "cnfsNode": "yes | no",
 "clientNode": "yes | no",
 "cesNode": "yes | no",
 "quorumNode": "yes | no",
 "cloudGatewayNode": "yes | no",
 "otherNodeRoles": "Roles",
 "designation": "quorum | quorumManager | manager",
 },
 },
 "cnfsInfo": {
 {
 "cnfsState": "cNFS state",
 "cnfsGroupId": "cNFS group ID",
 "cnfsIplist": "cNFS IP addresses",
 },
 },
 "cesInfo": {
 {
 "cesState": "CES state",
 "cesGroup": "CES group",
 "cesIpList": "CES IPs",
 "ipAddress": "IP address of this node",
 }
 }
 }
 }
]
}

```

| For more information about the fields in the following data structures, see the links at the end of this  
| topic.

```

| "status":
| Return status.

| "message": "ReturnMessage",
| The return message.

| "code": ReturnCode
| The return code.

```

```

| "paging":
| The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by
| the query.

| "nodes":
| "adminNodeName": "IP or host name"
| The host name that is used by the GPFS admin network.

| "nodeNumber": "Node ID"
| The GPFS node ID.

| "config":
| "adminLoginName": "Admin login name"
| The name of the admin login.

| "designatedLicense": "Designated license"
| The license this node is running on.

| "requiredLicense": "client | FPO | server"
| Controls the type of GPFS required license that is associated with the nodes in the cluster.

| "status":
| "osName": "Operating system"
| The name of Operating System running on this node.

| "nodeState": "Health status"
| The state of the node as reported by the mmhealth node show command.

| "gpfsState": "GPFS status"
| The state of GPFS on this node as reported the mmhealth node show command.

| "productVersion": "Version"
| The IBM Spectrum Scale version that is installed on this node.

| "network":
| "adminIPAddress": "IP address or host name"
| The IP address that is used by the GPFS admin network.

| "daemonNodeName": "GPFS daemon node name"
| The host name that is used by the GPFS daemon network.

| "daemonIPAddress": "IP address"
| The IP address that is used by the GPFS daemon network.

| "getcnfsNodeName": "Host name used by cNFS"
| The hostname that is used by cNFS.

| "roles":
| "snmpNode": "yes | no"
| Specifies whether this node is an SNMP node.

| "managerNode": "yes | no"
| Specifies whether this node is a manager node.

| "gatewayNode": "yes | no"
| Specifies whether this node is a gateway node.

| "cnfsNode": "yes | no"
| Specifies whether this node is a cNFS node.

| "clientNode": "yes | no"
| Specifies whether this node is a client node.

```

```

| "cesNode":"yes | no"
| Specifies whether this node is a CES node.
|
| "quorumNode":"yes | no"
| Specifies whether this node is a quorum node.
|
| "cloudGatewayNode":"yes | no"
| Specifies whether this node is a cloud gateway node.
|
| "otherNodeRoles":"Roles"
| Other roles of this node.
|
| "designation":"quorum | quorumManager |manager"
| The designated role.
|
| "cnfsInfo":
|
| "cnfsState":"cNFS state"
|
| "cnfsGroupId":"cNFS group ID"
|
| "cnfsIplist":"cNFS IP addresses"
|
| "cesInfo":
|
| "cesState":"CES state"
| The state of CNFS on this node.
|
| "cesGroup":"CES group"
| The CNFS group ID.
|
| "cesIpList":"CES IPs"
| A list of CNFS IP addresses.
|
| "ipAddress":"IP address of this node"
| The actual IP address of this node.

```

## Examples

The following example gets information about the node testnode1-d.localnet.com.

Request data:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/node/testnode1-d.localnet.com'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/node/testnode1-d.localnet.com?lastId=1001"
 },
 "nodes": [
 {
 "adminNodeName": "testnode1-a.localnet.com",
 "nodeNumber": "1",
 "config": {
 "adminLoginName": "root",
 "designatedLicense": "server",

```

```

 "requiredLicense": "true"
 },
 "status": {
 "osName": "Red Hat Enterprise Linux Server 7.2 (Maipo)",
 "nodeState": "HEALTHY",
 "gpfsState": "HEALTHY",
 "productVersion": "4.2.3.0"
 },
 "network": {
 "adminIPAddress": "10.0.200.21",
 "daemonNodeName": "testnode1-d.localnet.com",
 "daemonIPAddress": "10.0.100.21",
 "getcnfsNodeName": "string"
 },
 "roles": {
 "snmpNode": "true",
 "managerNode": "false",
 "gatewayNode": "false",
 "cnfsNode": "true",
 "clientNode": "false",
 "cesNode": "false",
 "quorumNode": "true",
 "cloudGatewayNode": "true",
 "otherNodeRoles": "perfmonNode,cesNode",
 "designation": "quorum"
 },
 "cnfsInfo": {
 "cnfsState": "enabled",
 "cnfsGroupId": "5",
 "cnfsIplist": "10.0.100.1,10.0.100.2"
 },
 "cesInfo": {
 "cesState": "e",
 "cesGroup": "LAN_4",
 "cesIpList": "192.168.2.11,192.168.2.13,192.168.2.15",
 "ipAddress": "10.0.100.23"
 }
}
]
}
}

```

“mmchnode command” on page 190

Changes node attributes.

“mmgetstate command” on page 350

Displays the state of the GPFS daemon on one or more nodes.

## **Nodes/{name}/health/events: GET**

| Gets details about system health events that are reported in a node.

### **Availability**

| Available on all IBM Spectrum Scale editions.

### **Description**

| The **GET nodes/nodeName/health/events** request gets information about the system health events that are reported in the specified node. For more information about the fields in the data structures that are returned, see the topics “mmhealth command” on page 355 and “mmces command” on page 103.

### **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nodes/nodeName/health/events`

| where

#### **nodes/nodeName**

| Specifies the node about which you want to get information. Required.

#### **health/events**

| Specifies that you need to get the details of the system health events reported on the node. Required.

### **Request headers**

| Content-Type: application/json

| Accept: application/json

### **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 63. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |
| name           | Name of the node.                                                                                | Required.         |

### **Request data**

| No request data.

### **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
```

```

 "next": "URL"
 },
 events: [
 {
 "oid": "ID",
 "component": "Component",
 "reportingnode": "Node Name",
 "type": "{TIP | STATE_CHANGE | NOTICE}",
 "activeSince": "Time",
 "name": "Event Name",
 "message": "Message",
 "entityType": "Entity Type",
 "entityName": "Name",
 "userAction": "Action",
 "description": "Description",
 "state": "Event State",
 "severity": "Severity",
 "parentName": "Name",
 }
],
}

```

| For more information about the fields in the following data structures, see the links at the end of this topic.

```

 "status":
 Return status.

 "message": "ReturnMessage",
 The return message.

 "code": "ReturnCode"
 The return code.

 "paging"
 The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

 "events":
 An array of elements that describe the system health event. Each element describes one node.

 "oid": "ID"
 The internal unique ID of the event.

 "component": "Component"
 The component to which the event belongs.

 "reportingnode": "Node Name"
 The node in which the event is reported.

 "type": "{TIP | STATE_CHANGE | NOTICE}"
 Specifies the event type.

 "activeSince": "Time"
 The time at which the event occurred.

 "name": "Event Name"
 The name of the event.

 "message": "Message"
 The event message.

 "entityType": "Entity Type"
 The type of the entity for which the event occurred.

```

```

| "entityName": Name
| The name of the entity for which the event occurred.
| "userAction": Action
| The user action that is required to resolve the event.
| "description": Description
| Description of the event.
| "state": Event State
| The state of the event.
| "severity": Severity
| Severity of the event.
| "parentName": Name
| Name of the parent event to which this event belongs.

```

## Examples

The following example gets information about the events reported in the node testnode1-d.localnet.com

Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/node/testnode1-d.localnet.com/health/events'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "events" : [{
 "activeSince" : "2017-03-13 10:55:15,809",
 "component" : "FILESYSTEM",
 "description" : "The inode usage in the fileset reached a normal level.",
 "entityName" : "fset1",
 "entityType" : "FILESET",
 "message" : "The inode usage of fileset fset1 in file system gpfs0 reached a normal level.",
 "name" : "inode_normal",
 "oid" : 690,
 "parentName" : "gpfs0",
 "reportingNode" : "mari-11.localnet.com",
 "severity" : "INFO",
 "state" : "HEALTHY",
 "type" : "STATE_CHANGE",
 "userAction" : "N/A"
 }, {
 "activeSince" : "2017-03-13 11:05:15,912",
 "component" : "FILESYSTEM",
 "description" : "The inode usage in the fileset reached a normal level.",
 "entityName" : "fset2",
 "entityType" : "FILESET",
 "message" : "The inode usage of fileset fset2 in file system gpfs0 reached a normal level.",
 "name" : "inode_normal",
 "oid" : 691,
 "parentName" : "gpfs0",
 "reportingNode" : "mari-11.localnet.com",
 "severity" : "INFO",
 "state" : "HEALTHY",
 "type" : "STATE_CHANGE",
 "userAction" : "N/A"
 }, {
 "activeSince" : "2017-03-15 18:44:42,568",
 }
]
```

```

| "component" : "GUI",
| "description" : "The GUI service is running",
| "entityName" : "mari-11.localnet.com",
| "entityType" : "NODE",
| "message" : "GUI service as expected, state is started",
| "name" : "gui_up",
| "oid" : 882,
| "reportingNode" : "mari-11.localnet.com",
| "severity" : "INFO",
| "state" : "HEALTHY",
| "type" : "STATE_CHANGE",
| "userAction" : "N/A"
| },
| {
| "activeSince" : "2017-03-16 07:07:06,418",
| "component" : "GUI",
| "description" : "The GUI checks the reachability of all nodes.",
| "entityName" : "mari-11.localnet.com",
| "entityType" : "NODE",
| "message" : "The GUI can reach the node mari-12.localnet.com",
| "name" : "gui_reachable_node",
| "oid" : 884,
| "reportingNode" : "mari-11.localnet.com",
| "severity" : "INFO",
| "state" : "HEALTHY",
| "type" : "STATE_CHANGE",
| "userAction" : "None needed."
| }
|],
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }

"mmhealth command" on page 355
Monitors health status of nodes.

"mmc ces command" on page 103
Manage CES (Cluster Export Services) configuration.

```

## Nodes/{name}/health/states: GET

Gets details about system health states for the node or nodes of node class.

### Availability

Available on all IBM Spectrum Scale editions.

### Description

The **GET nodes/nodeName/health/states** request gets information about the system health states of the specified node or node class. For more information about the fields in the data structures that are returned, see the topics “mmhealth command” on page 355 and “mmces command” on page 103.

### Request URL

`https://<IP address or host name of API server>:<port>/scalemgmt/v2/nodes/nodeName/health/states`

where

#### nodes/nodeName

Specifies the node about which you want to get information. Required.

#### health/states

Specifies that you need to get the details of the system health state of the node. Required.

### Request headers

Content-Type: application/json

Accept: application/json

### Request parameters

The following parameters can be used in the request URL to customize the request:

Table 64. List of request parameters

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example, 'status=HEALTHY,entityType=FILESET'                   | Optional.         |
| name           | Name of the node.                                                                                | Required.         |

### Request data

No request data.

### Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "paging":
 {
```

```

 "next": "URL"
 },
 events: [
 {
 "oid": "ID",
 "component": "Component",
 "reportingnode": "Node Name",
 "type": "{TIP | STATE_CHANGE | NOTICE}",
 "activeSince": "Time",
 "name": "Event Name",
 "message": "Message",
 "entityType": "Entity Type",
 "entityName": "Name",
 "userAction": "Action",
 "description": "Description",
 "state": "Event State",
 "severity": "Severity",
 "parentName": "Name",
 }
],
}

```

For more information about the fields in the following data structures, see the links at the end of this topic.

```

 "status":
 Return status.

 "message": "ReturnMessage",
 The return message.

 "code": "ReturnCode"
 The return code.

 "paging"
 The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.

 "events":
 An array of elements that describe the system health event. Each element describes one node.

 "oid": "ID"
 The internal unique ID of the event.

 "component": "Component"
 The component to which the event belongs.

 "reportingnode": "Node Name"
 The node in which the event is reported.

 "type": "{TIP | STATE_CHANGE | NOTICE}"
 Specifies the event type.

 "activeSince": "Time"
 The time at which the event occurred.

 "name": "Event Name"
 The name of the event.

 "message": "Message"
 The event message.

 "entityType": "Entity Type"
 The type of the entity for which the event occurred.

```

```

| "entityName": "Name
| The name of the entity for which the event occurred.
| "userAction": "Action
| The user action that is required to resolve the event.
| "description": "Description
| Description of the event.
| "state": "Event State
| The state of the event.
| "severity": "Severity
| Severity of the event.
| "parentName": "Name
| Name of the parent event to which this event belongs.

```

## Examples

| The following example gets information about the events reported in the node testnode1-d.localnet.com.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/node/testnode1-d.localnet.com/health/states'
```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "states" : [{
 "activeSince" : "2017-03-02 15:21:57,431",
 "component" : "THRESHOLD",
 "entityName" : "mari-11.localnet.com",
 "entityType" : "NODE",
 "oid" : 4,
 "reportingNode" : "mari-11.localnet.com",
 "state" : "HEALTHY"
 }, {
 "activeSince" : "2017-03-02 15:21:58,284",
 "component" : "NETWORK",
 "entityName" : "mari-11.localnet.com",
 "entityType" : "NODE",
 "oid" : 5,
 "reportingNode" : "mari-11.localnet.com",
 "state" : "HEALTHY"
 }, {
 "activeSince" : "2017-03-02 15:24:57,143",
 "component" : "PERFMON",
 "entityName" : "mari-11.localnet.com",
 "entityType" : "NODE",
 "oid" : 6,
 "reportingNode" : "mari-11.localnet.com",
 "state" : "HEALTHY"
 }, {
 "activeSince" : "2017-03-02 15:21:58,284",
 "component" : "NETWORK",
 "entityName" : "eth0",
 "entityType" : "NIC",
 "oid" : 9
 }
]
```

```
| "parentName" : "mari-11.localnet.com",
| "reportingNode" : "mari-11.localnet.com",
| "state" : "HEALTHY"
| }
|],
| "status" : {
| "code" : 200,
| "message" : "The request finished successfully"
| }
| }
```

“mmhealth command” on page 355  
Monitors health status of nodes.

“mmces command” on page 103  
Manage CES (Cluster Export Services) configuration.

## | NSDs: GET

- | Gets information about NSDs that are part of the system.

## I Availability

- | Available on all IBM Spectrum Scale editions.

## I Description

- | The **GET nsds** request gets information about NSDs that are configured in the system. For more  
| information about the fields in the data structures that are returned, see the topics “**mmcrnsd command**”  
| on page 262, “**mmlsdisk command**” on page 402, and “**mmlsnsd command**” on page 423.

## | Request URL

- | https://<IP address or host name of API server>:<port>/scalemgmt/v2/nsds

| where

| nsds

- | Specifies NSDs as the resource of the GET call. Required.

## | Request headers

- | Accept: application/json

## I Request data

- | No request data.

## Request parameters

- The following parameters can be used in the request URL to customize the request:

| *Table 65. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

## | Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "nsds":
| {
| "config": [
| {
| "id": "1",
| "label": "Label 1",
| "value": "Value 1",
| "type": "Type 1",
| "order": 1
| },
| {
| "id": "2",
| "label": "Label 2",
| "value": "Value 2",
| "type": "Type 2",
| "order": 2
| }
|]
| }
| }
```

```

 "name": "Name",
 "fileSystem": FilesystemName,
 "failureGroup": "Failure Group",
 "type": "Type",
 "storagePool": "Storage Pool",
 "status": "Status",
 "availability": Availability,
 "quorumDisk": "Quorum Disk",
 "size": "Size",
 "availableBlocks": "Available Blocks",
 "availableFragments": "Available Fragments",
 "nsdServers": "NSD Servers",
 "nsdVolumeId": NSD Volume ID
 "oid": "OID",
],
}

```

The details of these parameters are given in the following list:

- "status":**  
Return status.
- "message": "ReturnMessage",**  
The return message.
- "code": *ReturnCode***  
The return code.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "nsds":**  
An array of information about NSDs. The array contains elements that describe the NSDs. For more information about the fields in this structure, see the links at the end of this topic.
  - "name": "Name"**  
The name of the disk.
  - "fileSystem": *FilesystemName***  
The file system to which the NSD belongs.
  - "failureGroup": "Failure Group"**  
ID of the failure group to which the NSD belongs.
  - "type": "Type"**  
The type of data to be stored on the disk.
  - "storagePool": "Storage Pool"**  
The pool to which the disk belongs.
  - "status": "Status"**  
Status of the disk.
  - "quorumDisk": "Quorum Disk"**  
The quorum status of the disk.
  - "size": "Size"**  
Size of the disk.
  - "availableBlocks": "Available Blocks"**  
Available blocks in the disk.
  - "availableFragments": "Available Fragments"**  
Available fragments of the disk.

```

| "nsdServers": "NSD Servers"
| The server that manages the I/O.
|
| "nsdVolumeId": NSD Volume ID
| ID of the NSD volume.
|
| "oid": "OID"
| Internal identifier that is used for paging.

```

## Examples

| The following example gets information about the NSDs that are configured in the system.

| Request URL:

```
| curl -k -u admin:admin001 -X GET --header 'accept:application/json'
| 'https://198.51.100.1:443/scalemgmt/v2/nsds'
```

| The request URL with no field or filter parameter returns only the details that uniquely identify the object.

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```
{
 "nsds" : [{
 "config" : {
 "fileSystem" : "gpfs0",
 "name" : "disk1"
 }
 }, {
 "config" : {
 "fileSystem" : "gpfs0",
 "name" : "disk8"
 }
 }, {
 "config" : {
 "fileSystem" : "objfs",
 "name" : "disk2"
 }
 }, {
 "config" : {
 "fileSystem" : "gpfs1",
 "name" : "disk3"
 }
 }, {
 "config" : {
 "fileSystem" : "gpfs1",
 "name" : "disk4"
 }
 }, {
 "config" : {
 "fileSystem" : "gpfs1",
 "name" : "disk5"
 }
 }, {
 "config" : {
 "name" : "disk6"
 }
 }, {
 "config" : {
 "fileSystem" : "objfs",

```

```

 "name" : "disk7"
 },
}],
"status" : {
 "code" : 200,
 "message" : "The request finished successfully"
}
}

Using the field parameter ":all:" returns entire details of the NSDs that are available in the cluster. For example:
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nsds?fields=:all:'
{
 "nsds" : [{
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.15 GB",
 "availableFragments" : "552 kB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs0",
 "name" : "disk1",
 "nsdServers" : "mari-11.localnet.com,mari-15.localnet.com,mari-14.localnet.com,
 mari-13.localnet.com,mari-12.localnet.com",
 "nsdVolumeId" : "0A00640B58B82A8C",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "system",
 "type" : "nsd"
 }
 }, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.45 GB",
 "availableFragments" : "664 kB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs0",
 "name" : "disk8",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AD9",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
 }, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.30 GB",
 "availableFragments" : "504 kB",
 "failureGroup" : "1",
 "fileSystem" : "objfs",
 "name" : "disk2",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82A96",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "system",
 }
 }
}

```

```

 "type" : "nsd"
 }
}, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.42 GB",
 "availableFragments" : "712 kB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs1",
 "name" : "disk3",
 "nsdServers" : "mari-11.localnet.com,mari-15.localnet.com,mari-14.localnet.com,
 mari-13.localnet.com,mari-12.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AA1",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "system",
 "type" : "nsd"
 }
}, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.94 GB",
 "availableFragments" : "504 kB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs1",
 "name" : "disk4",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AAC",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
}, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "9.94 GB",
 "availableFragments" : "536 kB",
 "failureGroup" : "1",
 "fileSystem" : "gpfs1",
 "name" : "disk5",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AB7",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : "10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
}, {
 "config" : {
 "availability" : "unknown",
 "availableBlocks" : "Not Available",
 "availableFragments" : "Not Available",
 "failureGroup" : "1",
 "name" : "disk6",
 "nsdServers" : "mari-11.localnet.com,mari-15.localnet.com,mari-12.localnet.com,
 mari-13.localnet.com,mari-14.localnet.com",
 "nsdVolumeId" : "0A00640B58B82AC2",
 "quorumDisk" : "no",
 }
}

```

```

 "size" : " 10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
}, {
 "config" : {
 "availability" : "up",
 "availableBlocks" : "7.22 GB",
 "availableFragments" : "4.62 MB",
 "failureGroup" : "1",
 "fileSystem" : "objfs",
 "name" : "disk7",
 "nsdServers" : "mari-11.localnet.com,mari-12.localnet.com,mari-13.localnet.com,
 mari-14.localnet.com,mari-15.localnet.com",
 "nsdVolumeId" : "0A00640B58B82ACD",
 "quorumDisk" : "no",
 "remarks" : "desc",
 "size" : " 10.00GiB",
 "status" : "ready",
 "storagePool" : "data",
 "type" : "nsd"
 }
}
], {
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}

```

**Related reference:**

- | “[mmcrnsd command](#)” on page 262
- | Creates Network Shared Disks (NSDs) used by GPFS.
- | “[mmlsnrd command](#)” on page 423
- | Displays Network Shared Disk (NSD) information for the GPFS cluster.
- | “[mmlsdisk command](#)” on page 402
- | Displays the current configuration and state of the disks in a file system.

## | NSDs/{nsdName}: GET

| Gets information about a specific NSD that is part of the system.

### | Availability

| Available on all IBM Spectrum Scale editions.

### | Description

| The **GET nsds/nsdName** request gets information about an NSD that is configured in the system. For more information about the fields in the data structures that are returned, see the topics “mmcrnsd command” on page 262, “mmlsdisk command” on page 402, and “mmlsnsd command” on page 423.

### | Request URL

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/nsds/nsdName`

| where

#### | nsds/nsdName

| Specifies the NSD about which you need to get the details. Required.

### | Request headers

| Accept: application/json

### | Request parameters

| The following parameters can be used in the request URL to customize the request:

| *Table 66. List of request parameters*

| Parameter name | Description and applicable keywords                                                            | Required/optional |
|----------------|------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. 'all' selects all available fields. | Optional.         |
| nsdName        | Name of the NSD.                                                                               | Required.         |

### | Request data

| No request data.

### | Response data

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "nsds":
| {
| "config": [
| {
| "name": "Name"
| "fileSystem": "FilesystemName"
| "failureGroup": "Failure Group",
| "type": "Type",
| }
|]
| }
| }
```

```

 "storagePool": "Storage Pool",
 "status": "Status",
 "availability": Availability,
 "quorumDisk": "Quorum Disk",
 "size": "Size",
 "availableBlocks": "Available Blocks",
 "availableFragments": "Available Fragments",
 "nsdServers": "NSD Servers",
 "nsdVolumeId": NSD Volume ID
 "oid": "OID",
],
}

```

The details of these parameters are given in the following list:

- "status":**  
Return status.
- "message": "ReturnMessage",**  
The return message.
- "code": *ReturnCode***  
The return code.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "nsds":**  
An array of information about NSDs. The array contains elements that describe the NSDs. For more information about the fields in this structure, see the links at the end of this topic.
  - "name": "Name"**  
The name of the disk.
  - "fileSystem": *FilesystemName***  
The file system to which the NSD belongs.
  - "failureGroup": "Failure Group"**  
ID of the failure group to which the NSD belongs.
  - "type": "Type"**  
The type of data to be stored on the disk.
  - "storagePool": "Storage Pool"**  
The pool to which the disk belongs.
  - "status": "Status"**  
Status of the disk.
  - "quorumDisk": "Quorum Disk"**  
The quorum status of the disk.
  - "size": "Size"**  
Size of the disk.
  - "availableBlocks": "Available Blocks"**  
Available blocks in the disk.
  - "availableFragments": "Available Fragments"**  
Available fragments of the disk.
  - "nsdServers": "NSD Servers"**  
The server that manages the I/O.

- "nsdVolumeId": NSD Volume ID**  
ID of the NSD volume.
- "oid": "OID"**  
Internal identifier that is used for paging.

## Examples

- The following example gets information about the NSD nsd1 that are configured in the system.

- Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/nsds/nsd1'
```

- Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/nsds/nsd1?lastId=1001"
 },
 "nsds": [
 {
 "config": {
 "name": "nsd1",
 "fileSystem": "gpfs0",
 "failureGroup": "1",
 "type": "dataOnly",
 "storagePool": "data",
 "status": "ready",
 "availability": "up",
 "quorumDisk": "no",
 "remarks": "This is a comment",
 "size": "10.00 GB",
 "availableBlocks": "730.50 MB",
 "availableFragments": "1.50 MB",
 "nsdServers": "gpfsgui-21.localnet.com",
 "nsdVolumeId": "0A0064155874F5AA"
 }
 }
]
}
```

### Related reference:

- “mmcrnsd command” on page 262  
Creates Network Shared Disks (NSDs) used by GPFS.
- “mmlsnsd command” on page 423  
Displays Network Shared Disk (NSD) information for the GPFS cluster.
- “mmlsdisk command” on page 402  
Displays the current configuration and state of the disks in a file system.

## | **SMB/shares: GET**

| Gets information about SMB shares.

## | **Availability**

| Available on all IBM Spectrum Scale editions.

## | **Description**

| The **GET smb/shares** request gets information about SMB shares. For more information about the fields  
| in the data structures that are returned, see “[mmsmb command](#)” on page 554.

## | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/smb/shares`

| where

### | **smb/shares**

|     Specifies SMB shares as the resource. Required.

## | **Request headers**

| `Accept: application/json`

## | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 67. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| filter         | Filter objects by expression. For example,<br>'status=HEALTHY,entityType=FILESET'                | Optional.         |

## | **Request data**

| No request data.

## | **Response data**

| The following list of attributes are available in the response data:

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "shares":
| {
```

```

"config": [
 {
 "shareName": "Share name"
 "filesystemName": File system name
 "filesetName": Fileset Name
 "path": "Path",
 },
 "smbOptions": "Options"
 (
 "browseable": "Browseable"
 "guestok": "yes | no"
 "smbencrypt": "auto | default | mandatory | disabled"
 "adminusers": "Users"
 "comment": "Comments"
 "cscpolicy": "Client-side caching policy"
 "fileidalgorithm": "fsname | hostname | fsnamenodirs | fsnamenorootdir"
 "gpfseleases": "yes | no"
 "gpfssrecalls": "yes | no"
 "gpfssharemodes": "yes | no"
 "gpfsyncio": "yes | no"
 "hideunreadable": "yes | no"
 "oplocks": "yes | no"
 "posixlocking": "yes | no"
 "readonly": "yes | no"
 "syncopsonclose": "yes | no"
)
}

```

The details of the parameters are given in the following list:

- "status":**  
Return status.
- "message": "ReturnMessage",**  
The return message.
- "code": *ReturnCode***  
The return code.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "shares":**  
An array of information about SMB shares addresses. The array contains elements that describe the SMB shares. For more information about the fields in this structure, see the links at the end of this topic.
- config":**  
Configuration details of SMB share.
- "sharename": "String"**  
Name of the SMB share.
- "filesystemName": "String"**  
The file system to which the SMB share belongs.
- "filesetName": "String"**  
The fileset to which the SMB share belongs.
- "path": "String"**  
The path for which SMB share is created.
- "smbOptions":**  
An array of information about SMB shares addresses. The array contains elements that describe the SMB shares. For more information about the fields in this structure, see the links at the end of this topic.

**"browseable": "*"Browseable"***  
If the value is set as yes, the export is shown in the Windows Explorer browser while browsing the file server.

**"guestok": "*{yes | no}*"**  
The guest ok parameter means that the access is permitted as the default guest user. By default, this parameter is set to yes and it cannot be changed.

**"smbencrypt": "*{auto | default | mandatory | disabled}*"**  
This option controls whether the remote client is allowed or required to use SMB encryption.  
This option controls whether the remote client is allowed or required to use SMB encryption.

**"adminusers": "*Users*"**  
Using this option, administrative users can be defined in the format of admin users=user1,user2,..usern. The users must be domain users. ,

**"comment": "*Comments*"**  
User-defined text.

**"cscpolicy": "*{manual | disable | documents | programs}*"**  
The client-side caching policy specifies how the clients that are capable of offline caching caches the files in the share.

**"fileidalgorithm": "*{fsname | hostname | fsnamenodirs | fsnamenorootdir}*"**  
This option allows to control the level of enforced data integrity. If the data integrity is ensured on the application level, it can be beneficial in cluster environments to reduce the level of enforced integrity for performance reasons.

**"gpfsleases": "*yes | no*"**  
These are cross protocol oplocks (opportunistic locks). That is, an SMB client can lock a file that provides the user improved performance while reading or writing to the file because no other user read or write to this file. If the value is set as yes, clients accessing the file over the other protocols can break the lock of an SMB client and the user gets informed when another user is accessing the same file at the same time.

**"gpfsrecalls": "*{yes | no}*"**  
If the value is set as yes, files that are migrated from disk recalled on access. If recalls = no, files are not recalled on access and the client receives.

**"gpfssharemodes": "*{yes | no}*"**  
An application can set share modes. If you set gpfs:sharemodes = yes, the share modes that are specified by the application is respected by all protocols and not only by the SMB protocol. If you set gpfs:sharemodes = no, then the share modes that are specified by the application is only respected by the SMB protocol.

**"gpfssyncio": "*{yes | no}*"**  
If the value is set as yes, the files in an export for which the setting is enabled, are opened with the OSYNC flag. Accessing a file is faster if gpfs:syncio is set to yes. Performance for certain workloads can be improved when SMB accesses the file with the OSYNC flag set. For example, updating only small blocks in a large file as observed with database applications. The underlying GPFS behavior is then changed to not read a complete block, if there is only a small update to it.

**"hideunreadable": "*{yes | no}*"**  
If the value is set as yes, all files and directories that the user has no permission to read is hidden from directory listings in the export. The hideunreadable=yes option is also known as access-based enumeration. When users are listing (enumerating) the directories and files within the export, they see only the files and directories that they have read access to.

**"oplocks": "*{yes | no}*"**  
If the value is set as yes, a client might request an opportunistic lock (oplock) from an SMB server when it opens a file. If the server grants the request, the client can cache large chunks

of the file without informing the server what it is doing with the cached chunks until the task is completed. Caching large chunks of a file saves a lot of network I/O round-trip time and enhances performance.

**"posixlocking": "{yes | no}"**

If the value is set as yes, it is tested if a byte-range (fctl) lock is already present on the requested portion of the file before granting a byte-range lock to an SMB client. For improved performance on SMB-only shares this option can be disabled. Disabling locking on cross-protocol shares can result in data integrity issues when clients concurrently set locks on a file through multiple protocols, for example, SMB and NFS.

**"readonly": "{yes | no}"**

If the value is set as yes, files cannot be modified or created on this export independent of the ACLs.

**"syncopsonclose": "{yes | no}"**

This option ensures that the file system synchronizes data to the disk each time a file is closed after writing. The written data is flushed to the disk.

## Examples

The following example gets information about the SMB shares that are configured in the system.

Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/smb/shares'
```

The request URL with no field or filter parameter returns only the details that uniquely identify the object.

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "smbShares" : [{
 "config" : {
 "shareName" : "share01"
 }
 }],
 "status" : {
 "code" : 200,
 "message" : "The request finished successfully"
 }
}
```

Using the field parameter ":all:" returns entire details of the SMB shares. For example:

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/smb/shares?lastId=1001"
 },
 "shares": [
 {
 "config": {
 "shareName": "smbShare",
 "filesystemName": "gpfs0",
 "path": "/gpfs0"
 }
 }
]
}
```

```
 "filesetName": "fset1",
 "path": "/mnt/gpfs0/fset1"
 },
 "smbOptions": {
 "browseable": "yes",
 "guestok": "yes",
 "smbencrypt": "auto",
 "adminusers": "admin",
 "comment": "This is a comment",
 "cscpolicy": "manual",
 "fileidalgorithm": "fsname",
 "gpfsleases": "yes",
 "gpfsrecalls": "yes",
 "gpfssharemodes": "yes",
 "gpfssyncio": "yes",
 "hideunreadable": "yes",
 "oplocks": "yes",
 "posixlocking": "yes",
 "readonly": "yes",
 "syncponclose": "yes"
 }
}
]
```

**Related reference:**

- | “[mmsmb command](#)” on page 554
- | Administers SMB shares, export ACLs, and global configuration.

## | **SMB/shares/{shareName}: GET**

| Gets information about a specific SMB share.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **GET smb/shares/shareName** request gets information about the specified SMB share. For more information about the fields in the data structures that are returned, see “mmsmb command” on page 554.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/smb/shares/shareName`

| where

#### | **smb/shares**

| Specifies the SMB share as the resource. Required.

#### | **shareName**

| Specifies the SMB share about which you need to get the details. Required.

### | **Request headers**

| `Accept: application/json`

### | **Request parameters**

| The following parameters can be used in the request URL to customize the request:

| *Table 68. List of request parameters*

| Parameter name | Description and applicable keywords                                                              | Required/optional |
|----------------|--------------------------------------------------------------------------------------------------|-------------------|
| fields         | Comma separated list of fields to be included in response. ':all:' selects all available fields. | Optional.         |
| shareName      | Name of the SMB share.                                                                           | Required.         |

### | **Request data**

| No request data.

### | **Response data**

| The following list of attributes are available in the response data:

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "paging":
| {
| "next": "URL"
| },
| "shares":
```

```

{
 "config": [
 {
 "shareName": "Share name",
 "filesystemName": File system name,
 "filesetName": Fileset Name,
 "path": "Path",
 },
 "smbOptions": "Options"
 (
 "browseable": "Browseable"
 "guestok": "yes | no"
 "smbencrypt": "auto | default | mandatory | disabled"
 "adminusers": "Users"
 "comment": "Comments"
 "cscpolicy": "Client-side caching policy"
 "fileidalgorithm": "fsname | hostname | fsnamenodirs | fsnamenorootdir"
 "gpfsleases": "yes | no"
 "gpfsrecalls": "yes | no"
 "gpfssharemodes": "yes | no"
 "gpfssyncio": "yes | no"
 "hideunreadable": "yes | no"
 "oplocks": "yes | no"
 "posixlocking": "yes | no"
 "readonly": "yes | no"
 "syncopsonclose": "yes | no"
 }
}

```

The details of the parameters are given in the following list:

- "status":**  
Return status.
- "message": "ReturnMessage",**  
The return message.
- "code": *ReturnCode***  
The return code.
- "paging"**  
The URL to retrieve the next page. Paging is enabled when more than 1000 objects are returned by the query.
- "shares":**  
An array of information about SMB shares addresses. The array contains elements that describe the SMB shares. For more information about the fields in this structure, see the links at the end of this topic.
- config":**  
Configuration details of SMB share.
- "sharename": "String"**  
Name of the SMB share.
- "filesystemName": "String"**  
The file system to which the SMB share belongs.
- "filesetName": "String"**  
The fileset to which the SMB share belongs.
- "path": "String"**  
The path for which SMB share is created.

```

"smbOptions":
 An array of information about SMB shares addresses. The array contains elements that describe
 the SMB shares. For more information about the fields in this structure, see the links at the end of
 this topic.

"browseable": "Browseable"
 If the value is set as yes, the export is shown in the Windows Explorer browser while
 browsing the file server.

"guestok": "{yes | no}"
 The guest ok parameter means that the access is permitted as the default guest user. By
 default, this parameter is set to yes and it cannot be changed.

"smbencrypt": "{auto | default | mandatory | disabled}"
 This option controls whether the remote client is allowed or required to use SMB encryption.
 This option controls whether the remote client is allowed or required to use SMB encryption.

"adminusers": "Users"
 Using this option, administrative users can be defined in the format of admin
 users=user1;user2,..usern. The users must be domain users. ,

"comment": "Comments"
 User-defined text.

"cscpolicy": "{manual | disable | documents | programs}"
 The client-side caching policy specifies how the clients that are capable of offline caching
 caches the files in the share.

"fileidalgorithm": "{fsname | hostname | fsnamenodirs | fsnamenorootdir}"
 This option allows to control the level of enforced data integrity. If the data integrity is
 ensured on the application level, it can be beneficial in cluster environments to reduce the
 level of enforced integrity for performance reasons.

"gpfsleases": "yes | no"
 These are cross protocol oplocks (opportunistic locks). That is, an SMB client can lock a file
 that provides the user improved performance while reading or writing to the file because no
 other user read or write to this file. If the value is set as yes, clients accessing the file over the
 other protocols can break the lock of an SMB client and the user gets informed when another
 user is accessing the same file at the same time.

"gpfsrecalls": "{yes | no}"
 If the value is set as yes, files that are migrated from disk recalled on access. If recalls = no,
 files are not recalled on access and the client receives.

"gpfssharemodes": "{yes | no}"
 An application can set share modes. If you set gpfs:sharemodes = yes, the share modes that
 are specified by the application is respected by all protocols and not only by the SMB
 protocol. If you set gpfs:sharemodes = no, then the share modes that are specified by the
 application is only respected by the SMB protocol.

"gpfssyncio": "{yes | no}"
 If the value is set as yes, the files in an export for which the setting is enabled, are opened
 with the OSYNC flag. Accessing a file is faster if gpfs:syncio is set to yes. Performance for
 certain workloads can be improved when SMB accesses the file with the OSYNC flag set. For
 example, updating only small blocks in a large file as observed with database applications.
 The underlying GPFS behavior is then changed to not read a complete block, if there is only
 a small update to it.

"hideunreadable": "{yes | no}"
 If the value is set as yes, all files and directories that the user has no permission to read is
 hidden from directory listings in the export. The hideunreadable=yes option is also known as

```

access-based enumeration. When users are listing (enumerating) the directories and files within the export, they see only the files and directories that they have read access to.

**"oplocks": "{yes | no}"**

If the value is set as yes, a client might request an opportunistic lock (oplock) from an SMB server when it opens a file. If the server grants the request, the client can cache large chunks of the file without informing the server what it is doing with the cached chunks until the task is completed. Caching large chunks of a file saves a lot of network I/O round-trip time and enhances performance.

**"posixlocking": "{yes | no}"**

If the value is set as yes, it is tested if a byte-range (fctl) lock is already present on the requested portion of the file before granting a byte-range lock to an SMB client. For improved performance on SMB-only shares this option can be disabled. Disabling locking on cross-protocol shares can result in data integrity issues when clients concurrently set locks on a file through multiple protocols, for example, SMB and NFS.

**"readonly": "{yes | no}"**

If the value is set as yes, files cannot be modified or created on this export independent of the ACLs.

**"syncopsonclose": "{yes | no}"**

This option ensures that the file system synchronizes data to the disk each time a file is closed after writing. The written data is flushed to the disk.

## Examples

The following example gets information about the SMB share share1.

Request URL:

```
curl -k -u admin:admin001 -X GET --header 'accept:application/json'
'https://198.51.100.1:443/scalemgmt/v2/smb/shares/smbShare1'
```

Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
{
 "status": {
 "code": "200",
 "message": "..."
 },
 "paging": {
 "next": "https://localhost:443/scalemgmt/v2/smb/shares/smbShare1?lastId=1001"
 },
 "shares": [
 {
 "config": {
 "shareName": "smbShare1",
 "filesystemName": "gpfs0",
 "filesetName": "fset1",
 "path": "/mnt/gpfs0/fset1"
 },
 "smbOptions": {
 "browseable": "yes",
 "guestok": "yes",
 "smbencrypt": "auto",
 "adminusers": "admin",
 "comment": "This is a comment",
 "cscpolicy": "manual",
 "fileidalgorithm": "fsname",
 "shareName": "smbShare1",
 "filesystemName": "gpfs0",
 "filesetName": "fset1",
 "path": "/mnt/gpfs0/fset1"
 }
 }
]
}
```

```
| "gpfsleases": "yes",
| "gpfsrecalls": "yes",
| "gpfssharemodes": "yes",
| "gpfssyncio": "yes",
| "hideunreadable": "yes",
| "oplocks": "yes",
| "posixlocking": "yes",
| "readonly": "yes",
| "syncopsonclose": "yes"
| }
| }
|]
| }
```

| **Related reference:**

| “[mmsmb command](#)” on page 554

| Administers SMB shares, export ACLs, and global configuration.

## | **SMB/shares: POST**

| Creates a new SMB share.

## | **Availability**

| Available on all IBM Spectrum Scale editions.

## | **Description**

| The **POST smb/shares** request creates a new SMB share. For more information about the fields in the data structures that are returned, see “[mmsmb command](#)” on page 554.

## | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/smb/shares/shareName/scalemgmt/v2/smb/shares`

| where

### | **smb/shares**

|     Specifies SMB share as the target. Required.

## | **Request headers**

| `Accept: application/json`

## | **Request data**

| The following list of attributes are available in the response data:

```
{ "shareName": "ShareName",
 "path": "Path",
 "smbOptions": {
 "browseable": "Browseable",
 "guestok": "{yes | no}",
 "smbencrypt": "{auto | default | mandatory | disabled}",
 "adminusers": "Users",
 "comment": "Comments",
 "cscpolicy": "Client-side caching policy",
 "fileidalgorithm": "fsname | hostname | fsnamenodirs | fsnamenorootdir",
 "gpfsleases": "{yes | no}",
 "gpfsrecalls": "{yes | no}",
 "gpfssharemodes": "{yes | no}",
 "gpfssyncio": "{yes | no}",
 "hideunreadable": "{yes | no}",
 "oplocks": "{yes | no}",
 "posixlocking": "{yes | no}",
 "readonly": "{yes | no}",
 "syncoponclose": "{yes | no}"
 }
}
```

| The details of the parameters are given in the following list:

### | **config:**

|     Configuration details of SMB share.

### | **"sharename": "String"**

|     Name of the SMB share.

### | **"filesystemName": "String"**

|     The file system to which the SMB share belongs.

```

| "filesetName": "String"
| The fileset to which the SMB share belongs.
|
| "path": "String"
| The path for which SMB share is created.
|
| "smbOptions":
| An array of information about SMB shares addresses. The array contains elements that describe the
| SMB shares. For more information about the fields in this structure, see the links at the end of this
| topic.
|
| "browseable": "{yes | no}"
| If the value is set as yes, the export is shown in the Windows Explorer browser while browsing
| the file server.
|
| "guestok": "{yes | no}"
| The guest ok parameter means that the access is permitted as the default guest user. By default,
| this parameter is set to yes and it cannot be changed.
|
| "smbencrypt": "{auto | default | mandatory | disabled}"
| This option controls whether the remote client is allowed or required to use SMB encryption. This
| option controls whether the remote client is allowed or required to use SMB encryption. Possible
| values are auto, mandatory, and disabled.
|
| "adminusers": "Users"
| Using this option, administrative users can be defined in the format of admin users=user1;user2,
| ..usern. The users must be domain users. ,
|
| "comment": "Comments"
| User-defined text.
|
| "cscpolicy": "{manual | disable | documents | programs}"
| The client-side caching policy specifies how the clients that are capable of offline caching caches
| the files in the share.
|
| "fileidalgorithm": "{fsname | hostname | fsnamenodirs | fsnamenorootdir}"
| This option allows to control the level of enforced data integrity. If the data integrity is ensured
| on the application level, it can be beneficial in cluster environments to reduce the level of
| enforced integrity for performance reasons.
|
| "gpfsleases": "{yes | no}"
| These are cross protocol oplocks (opportunistic locks). That is, an SMB client can lock a file that
| provides the user improved performance while reading or writing to the file because no other
| user read or write to this file. If the value is set as yes, clients accessing the file over the other
| protocols can break the lock of an SMB client and the user gets informed when another user is
| accessing the same file at the same time.
|
| "gpfsrecalls": "{yes | no}"
| If the value is set as yes, files that are migrated from disk recalled on access. If recalls = no, files
| are not recalled on access and the client receives.
|
| "gpfssharemodes": "{yes | no}"
| An application can set share modes. If you set gpfs:sharemodes = yes, the share modes that are
| specified by the application is respected by all protocols and not only by the SMB protocol. If you
| set gpfs:sharemodes = no, then the share modes that are specified by the application is only
| respected by the SMB protocol.
|
| "gpfssyncio": "{yes | no}"
| If the value is set as yes, the files in an export for which the setting is enabled, are opened with
| the OSYNC flag. Accessing a file is faster if gpfs:syncio is set to yes. Performance for certain
| workloads can be improved when SMB accesses the file with the OSYNC flag set. For example,

```

updating only small blocks in a large file as observed with database applications. The underlying GPFS behavior is then changed to not read a complete block, if there is only a small update to it.

**"hideunreadable": "{yes | no}"**

If the value is set as yes, all files and directories that the user has no permission to read is hidden from directory listings in the export. The hideunreadable=yes option is also known as access-based enumeration. When a user is listing (enumerating) the directories and files within the export, they see only the files and directories that they have read access to.

**"oplocks": "{yes | no}"**

If the value is set as yes, a client might request an opportunistic lock (oplock) from an SMB server when it opens a file. If the server grants the request, the client can cache large chunks of the file without informing the server what it is doing with the cached chunks until the task is completed. Caching large chunks of a file saves a lot of network I/O round-trip time and enhances performance.

**"posixlocking": "{yes | no}"**

If the value is set as yes, it is tested if a byte-range (fctl) lock is already present on the requested portion of the file before granting a byte-range lock to an SMB client. For improved performance on SMB-only shares this option can be disabled. Disabling locking on cross-protocol shares can result in data integrity issues when clients concurrently set locks on a file through multiple protocols, for example, SMB and NFS.

**"readonly": "{yes | no}"**

If the value is set as yes, files cannot be modified or created on this export independent of the ACLs.

**"syncponclose": "{yes | no}"**

This option ensures that the file system synchronizes data to the disk each time a file is closed after writing. The written data is flushed to the disk.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": "",
 }
 },
 {
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

```

| "status":
| Return status.
|
| "message": "ReturnMessage",
| The return message.
|
| "code": ReturnCode
| The return code.
|
| "jobs":
| An array of elements that describe jobs. Each element describes one job.
|
| "result"
|
| "commands": "String"
| Array of commands that are run in this job.
|
| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

| The following example creates a new SMB share.

| Request data:

```

| curl -k -u admin:admin001 -X POST --header 'content-type:application/json' --header 'accept:application/json'
| -d '{
| "shareName": "smbShare",
| "path": "/mnt/gpfs0/fset1",
| "smbOptions": {
| "browseable": "yes",
| "guestok": "yes",
| "smbencrypt": "auto",
| "adminusers": "admin",
| }
| }
```

```

 "comment": "This is a comment",
 "cscpolicy": "manual",
 "fileidalgorithm": "fsname",
 "gpfsleases": "yes",
 "gpfsrecalls": "yes",
 "gpfssharemodes": "yes",
 "gpfssyncio": "yes",
 "hideunreadable": "yes",
 "oplocks": "yes",
 "posixlocking": "yes",
 "readonly": "yes",
 "syncoponclose": "yes"
}
}' "https://198.51.100.1:443/scalemgmt/v2/smb/shares"
```

#### Response data:

**Note:** In the JSON data that is returned, the return code indicates whether the command is successful. The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "POST",
 "url": "/scalemgmt/v2/smb/shares",
 "data": "{
 \"shareName\": \"smbShare\",
 \"path\": \"/mnt/gpfs0/fset1\",
 \"smbOptions\": {
 \"browseable\": \"yes\",
 \"guestok\": \"yes\",
 \"smbencrypt\": \"auto\",
 \"adminusers\": \"admin\",
 \"comment\": \"This is a comment\",
 \"cscpolicy\": \"manual\",
 \"fileidalgorithm\": \"fsname\",
 \"gpfsleases\": \"yes\",
 \"gpfsrecalls\": \"yes\",
 \"gpfssharemodes\": \"yes\",
 \"gpfssyncio\": \"yes\",
 \"hideunreadable\": \"yes\",
 \"oplocks\": \"yes\",
 \"posixlocking\": \"yes\",
 \"readonly\": \"yes\",
 \"syncoponclose\": \"yes\"}
 },
 " jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
 }
}
```

#### Related reference:

- “[mmsmb command](#)” on page 554
- Administers SMB shares, export ACLs, and global configuration.

## | **SMB/shares/{shareName}: PUT**

| Modifies an existing SMB share.

## | **Availability**

| Available on all IBM Spectrum Scale editions.

## | **Description**

| The **PUT smb/shares/shareName** request modifying an existing SMB share. For more information about  
| the fields in the data structures that are returned, see “[mmsmb command](#)” on page 554.

## | **Request URL**

| `https://<IP address or host name of API server>:<port>/sclemgmt/v2/smb/shares/shareName`

| where

### | **smb/shares**

|     Specifies SMB share as the resource. Required.

### | **shareName**

|     Specifies the SMB share to be modified. Required.

## | **Request headers**

| `Accept: application/json`

## | **Request data**

### | **Request data**

| The following list of attributes are available in the response data:

```
| {
| "smbOptions": [
| "browseable": "{yes | no}"
| "guestok": "{yes | no}"
| "smbencrypt": "{auto | default | mandatory | disabled}"
| "adminusers": "Users"
| "comment": "Comments"
| "cscpolicy": "Client-side caching policy"
| "fileidalgorithm": "{fsname | hostname | fsnamenodirs | fsnamenorootdir}"
| "gpfsleases": "{yes | no}"
| "gpfsrecalls": "{yes | no}"
| "gpfssharemodes": "{yes | no}"
| "gpfssyncio": "{yes | no}"
| "hideunreadable": "{yes | no}"
| "oplocks": "{yes | no}"
| "posixlocking": "{yes | no}"
| "readonly": "{yes | no}"
| "syncopsongclose": "{yes | no}"
| }
| "removeSmboptions": ["List of SMB options to be removed"]
| }
```

| The details of the parameters are given in the following list:

#### | **"smbOptions":**

|     An array of information about SMB shares addresses. The array contains elements that describe the  
|     SMB shares. For more information about the fields in this structure, see the links at the end of this  
|     topic.

```

"browseable": "Browseable"
If the value is set as yes, the export is shown in the Windows Explorer browser while browsing the file server.

"guestok": "{yes | no}"
The guest ok parameter means that the access is permitted as the default guest user. By default, this parameter is set to yes and it cannot be changed.

"smbencrypt": "{auto | default | mandatory | disabled}"
This option controls whether the remote client is allowed or required to use SMB encryption. This option controls whether the remote client is allowed or required to use SMB encryption. Possible values are auto, mandatory, and disabled.

"adminusers": "Users"
Using this option, administrative users can be defined in the format of admin users=user1;user2, ..usern. The users must be domain users. ,

"comment": "Comments"
User-defined text.

"cscpolicy": "{manual | disable | documents | programs}"
The client-side caching policy specifies how the clients that are capable of offline caching caches the files in the share.

"fileidalgorithm": "{fsname | hostname | fsnamenodirs | fsnamenorootdir}"
This option allows to control the level of enforced data integrity. If the data integrity is ensured on the application level, it can be beneficial in cluster environments to reduce the level of enforced integrity for performance reasons.

"gpfsleases": "{yes | no}"
These are cross protocol oplocks (opportunistic locks). That is, an SMB client can lock a file that provides the user improved performance while reading or writing to the file because no other user read or write to this file. If the value is set as yes, clients accessing the file over the other protocols can break the lock of an SMB client and the user gets informed when another user is accessing the same file at the same time.

"gpfsrecalls": "{yes | no}"
If the value is set as yes, files that are migrated from disk recalled on access. If recalls = no, files are not recalled on access and the client receives.

"gpfssharemodes": "{yes | no}"
An application can set share modes. If you set gpfs:sharemodes = yes, the share modes that are specified by the application is respected by all protocols and not only by the SMB protocol. If you set gpfs:sharemodes = no, then the share modes that are specified by the application is only respected by the SMB protocol.

"gpfsyncio": "{yes | no}"
If the value is set as yes, the files in an export for which the setting is enabled, are opened with the OSYNC flag. Accessing a file is faster if gpfs:syncio is set to yes. Performance for certain workloads can be improved when SMB accesses the file with the OSYNC flag set. For example, updating only small blocks in a large file as observed with database applications. The underlying GPFS behavior is then changed to not read a complete block, if there is only a small update to it.

"hideunreadable": "{yes | no}"
If the value is set as yes, all files and directories that the user has no permission to read is hidden from directory listings in the export. The hideunreadable=yes option is also known as access-based enumeration. When a user is listing (enumerating) the directories and files within the export, they see only the files and directories that they have read access to.

"oplocks": "{yes | no}"
If the value is set as yes, a client might request an opportunistic lock (oplock) from an SMB server when it opens a file. If the server grants the request, the client can cache large chunks of

```

the file without informing the server what it is doing with the cached chunks until the task is completed. Caching large chunks of a file saves a lot of network I/O round-trip time and enhances performance.

**"posixlocking": "{yes | no}"**  
If the value is set as yes, it is tested if a byte-range (fcntl) lock is already present on the requested portion of the file before granting a byte-range lock to an SMB client. For improved performance on SMB-only shares this option can be disabled. Disabling locking on cross-protocol shares can result in data integrity issues when clients concurrently set locks on a file through multiple protocols, for example, SMB and NFS.

**"readonly": "{yes | no}"**  
If the value is set as yes, files cannot be modified or created on this export independent of the ACLs.

**"syncopsonclose": "{yes | no}"**  
This option ensures that the file system synchronizes data to the disk each time a file is closed after writing. The written data is flushed to the disk.

**"removeSmboptions": ["List of SMB options to be removed"]**  
Specifies the SMB options to be removed.

## Response data

```
{
 "status": {
 "code": "ReturnCode",
 "message": "ReturnMessage"
 },
 "jobs": [
 {
 "result": "",
 {
 "commands": "String",
 "progress": "String",
 "exitCode": "Exit code",
 "stderr": "Error",
 "stdout": "String",
 },
 "request": " ",
 {
 "type": "{GET | POST | PUT | DELETE}",
 "url": "URL",
 "data": """
 }
 "jobId": "ID",
 "submitted": "Time",
 "completed": "Time",
 "status": "Job status",
 }
],
}
```

For more information about the fields in the following data structures, see the links at the end of this topic.

**"status":**  
Return status.

**"message": "ReturnMessage"**,  
The return message.

**"code": ReturnCode**  
The return code.

```

| "jobs":
| An array of elements that describe jobs. Each element describes one job.
|
| "result"
|
| "commands": "String"
| Array of commands that are run in this job.
|
| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

The following example modifies the SMB share `smbShare1`.

Request data:

```

curl -k -u admin:admin001 -X PUT --header 'content-type:application/json' --header 'accept:application/json'
-d '{
 "smbOptions": {
 "browseable": "yes",
 "guestok": "yes",
 "smbencrypt": "auto",
 "adminusers": "admin",
 "comment": "This is a comment",
 "cscpolicy": "manual",
 "fileidalgorithm": "fsname",
 "gpfsleases": "yes",
 "gpfsrecalls": "yes",
 "gpfssharemodes": "yes",
 "gpfssyncio": "yes",
 "hideunreadable": "yes",
 "oplocks": "yes",
 "posixlocking": "yes",
 "readonly": "yes",
 }
}'

```

```

 "syncopsonclose": "yes"
 },
 "removeSmbOptions": [
 "string"
]
}
}' "https://198.51.100.1:443/scalemgmt/v2/smb/shares/smbShare1"

```

| Response data:

| **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.  
| The response code 200 indicates that the command successfully retrieved the information. Error code 400  
| represents an invalid request and 500 represents internal server error.

```

{
 "status": {
 "code": "200",
 "message": "..."
 },
 "job": [
 {
 "result": {},
 "request": {
 "type": "PUT",
 "url": "/scalemgmt/v2/smb/shares/smbShare1",
 "data": {
 "smbOptions": {
 "browseable": "yes",
 "guestok": "yes",
 "smbencrypt": "auto",
 "adminusers": "admin",
 "comment": "This is a comment",
 "cscpolicy": "manual",
 "fileidalgorithm": "fsname",
 "gpfsleases": "yes",
 "gpfsrecalls": "yes",
 "gpfssharemodes": "yes",
 "gpfssyncio": "yes",
 "hideunreadable": "yes",
 "oplocks": "yes",
 "posixlocking": "yes",
 "readonly": "yes",
 "syncopsonclose": "yes"
 },
 "removeSmbOptions": [
 "string"
]
 }
 },
 "jobId": "12345",
 "submitted": "2016-11-14 10.35.56",
 "completed": "2016-11-14 10.35.56",
 "status": "COMPLETED"
 }
]
}

```

| **Related reference:**

| “[mmsmb command](#)” on page 554

| Administers SMB shares, export ACLs, and global configuration.

## | **SMB/shares/{shareName}: DELETE**

| Deletes an SMB share.

### | **Availability**

| Available on all IBM Spectrum Scale editions.

### | **Description**

| The **DELETE smb/shares/shareName** command deletes the specified SMB share. For more information, see “[mmsmb command](#)” on page 554.

### | **Request URL**

| `https://<IP address or host name of API server>:<port>/scalemgmt/v2/smb/shares/shareName`

| where:

#### | **smb/shares**

|     Specifies the SMB share as the resource. Required.

#### | **shareName**

|     Specifies the SMB share to be deleted. Required.

### | **Request headers**

| `Accept: application/json`

### | **Response data**

```
| {
| "status": {
| "code": "ReturnCode",
| "message": "ReturnMessage"
| },
| "jobs": [
| {
| "result": "",
| {
| "commands": "String",
| "progress": "String",
| "exitCode": "Exit code",
| "stderr": "Error",
| "stdout": "String",
| },
| "request": " ",
| {
| "type": "{GET | POST | PUT | DELETE}",
| "url": "URL",
| "data": "",
| }
| },
| {"jobId": "ID",
| "submitted": "Time",
| "completed": "Time",
| "status": "Job status",
| }
|],
| }
```

| For more information about the fields in the following data structures, see the links at the end of this topic.

#### | **"status":**

|     Return status.

```

| "message": "ReturnMessage",
| The return message.
|
| "code": ReturnCode
| The return code.
|
| "jobs":
| An array of elements that describe jobs. Each element describes one job.
|
| "result"
|
| "commands": "String"
| Array of commands that are run in this job.
|
| "progress": "String"
| Progress information for the request.
|
| "exitCode": "Exit code"
| Exit code of command. Zero is success, nonzero denotes failure.
|
| "stderr": "Error"
| CLI messages from stderr.
|
| "stdout": "String"
| CLI messages from stdout.
|
| "request"
|
| "type": "{GET | POST | PUT | DELETE}"
| HTTP request type.
|
| "url": "URL"
| The URL through which the job is submitted.
|
| "data": " "
| Optional.
|
| "jobId": "ID",
| The unique ID of the job.
|
| "submitted": "Time"
| The time at which the job was submitted.
|
| "completed": "Time"
| The time at which the job was completed.
|
| "status": "RUNNING | COMPLETED | FAILED"
| Status of the job.

```

## Examples

- | The following example deletes the SMB share share1.
- | Request data:
 

```
curl -k -u admin:admin001 -X DELETE --header 'accept:application/json'
'https://198.51.100.1:443/sclemgmt/v2/smb/shares/share1'
```
- | Response data:
  - | **Note:** In the JSON data that is returned, the return code indicates whether the command is successful.
  - | The response code 200 indicates that the command successfully retrieved the information. Error code 400 represents an invalid request and 500 represents internal server error.

```
| {
| "status": {
| "code": "200",
| "message": "..."
| },
| "job": [
| {
| "result": {},
| "request": {
| "type": "DELETE",
| "url": "/scalemgmt/v2/filesystems/gpfs0/filesets",
| "data": "{}"
| },
| "jobId": "12345",
| "submitted": "2016-11-14 10.35.56",
| "completed": "2016-11-14 10.35.56",
| "status": "COMPLETED"
| }
|]
| }
```

| **Related reference:**

- | “[mmsmb command](#)” on page 554
- | Administers SMB shares, export ACLs, and global configuration.

---

## **Accessibility features for IBM Spectrum Scale**

Accessibility features help users who have a disability, such as restricted mobility or limited vision, to use information technology products successfully.

---

### **Accessibility features**

The following list includes the major accessibility features in IBM Spectrum Scale:

- Keyboard-only operation
- Interfaces that are commonly used by screen readers
- Keys that are discernible by touch but do not activate just by touching them
- Industry-standard devices for ports and connectors
- The attachment of alternative input and output devices

IBM Knowledge Center, and its related publications, are accessibility-enabled. The accessibility features are described in IBM Knowledge Center ([www.ibm.com/support/knowledgecenter](http://www.ibm.com/support/knowledgecenter)).

---

### **Keyboard navigation**

This product uses standard Microsoft Windows navigation keys.

---

### **IBM and accessibility**

See the IBM Human Ability and Accessibility Center ([www.ibm.com/able](http://www.ibm.com/able)) for more information about the commitment that IBM has to accessibility.



---

## Notices

This information was developed for products and services offered in the US. This material might be available from IBM in other languages. However, you may be required to own a copy of the product or product version in that language in order to access it.

IBM may not offer the products, services, or features discussed in this document in other countries. Consult your local IBM representative for information on the products and services currently available in your area. Any reference to an IBM product, program, or service is not intended to state or imply that only that IBM product, program, or service may be used. Any functionally equivalent product, program, or service that does not infringe any IBM intellectual property right may be used instead. However, it is the user's responsibility to evaluate and verify the operation of any non-IBM product, program, or service.

IBM may have patents or pending patent applications covering subject matter described in this document. The furnishing of this document does not grant you any license to these patents. You can send license inquiries, in writing, to:

*IBM Director of Licensing IBM Corporation North Castle Drive, MD-NC119 Armonk, NY 10504-1785 US*

For license inquiries regarding double-byte character set (DBCS) information, contact the IBM Intellectual Property Department in your country or send inquiries, in writing, to:

*Intellectual Property Licensing Legal and Intellectual Property Law IBM Japan Ltd. 19-21, Nihonbashi-Hakozakicho, Chuo-ku Tokyo 103-8510, Japan*

INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. Some jurisdictions do not allow disclaimer of express or implied warranties in certain transactions, therefore, this statement may not apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes will be incorporated in new editions of the publication. IBM may make improvements and/or changes in the product(s) and/or the program(s) described in this publication at any time without notice.

Any references in this information to non-IBM websites are provided for convenience only and do not in any manner serve as an endorsement of those websites. The materials at those websites are not part of the materials for this IBM product and use of those websites is at your own risk.

IBM may use or distribute any of the information you provide in any way it believes appropriate without incurring any obligation to you.

Licensees of this program who wish to have information about it for the purpose of enabling: (i) the exchange of information between independently created programs and other programs (including this one) and (ii) the mutual use of the information which has been exchanged, should contact:

*IBM Director of Licensing IBM Corporation North Castle Drive, MD-NC119 Armonk, NY 10504-1785 US*

Such information may be available, subject to appropriate terms and conditions, including in some cases, payment of a fee.

The licensed program described in this document and all licensed material available for it are provided by IBM under terms of the IBM Customer Agreement, IBM International Program License Agreement or any equivalent agreement between us.

The performance data discussed herein is presented as derived under specific operating conditions. Actual results may vary.

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements or other publicly available sources. IBM has not tested those products and cannot confirm the accuracy of performance, compatibility or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

Statements regarding IBM's future direction or intent are subject to change or withdrawal without notice, and represent goals and objectives only.

All IBM prices shown are IBM's suggested retail prices, are current and are subject to change without notice. Dealer prices may vary.

This information is for planning purposes only. The information herein is subject to change before the products described become available.

This information contains examples of data and reports used in daily business operations. To illustrate them as completely as possible, the examples include the names of individuals, companies, brands, and products. All of these names are fictitious and any similarity to actual people or business enterprises is entirely coincidental.

#### COPYRIGHT LICENSE:

This information contains sample application programs in source language, which illustrate programming techniques on various operating platforms. You may copy, modify, and distribute these sample programs in any form without payment to IBM, for the purposes of developing, using, marketing or distributing application programs conforming to the application programming interface for the operating platform for which the sample programs are written. These examples have not been thoroughly tested under all conditions. IBM, therefore, cannot guarantee or imply reliability, serviceability, or function of these programs. The sample programs are provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of the sample programs.

Each copy or any portion of these sample programs or any derivative work must include a copyright notice as follows:

© (your company name) (year).

Portions of this code are derived from IBM Corp.

Sample Programs. © Copyright IBM Corp. \_enter the year or years\_.

If you are viewing this information softcopy, the photographs and color illustrations may not appear.

---

## Trademarks

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at Copyright and trademark information at [www.ibm.com/legal/copytrade.shtml](http://www.ibm.com/legal/copytrade.shtml).

Intel is a trademark of Intel Corporation or its subsidiaries in the United States and other countries.

Java and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle and/or its affiliates.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Microsoft and Windows are trademarks of Microsoft Corporation in the United States, other countries, or both.

UNIX is a registered trademark of the Open Group in the United States and other countries.

---

## **Terms and conditions for product documentation**

Permissions for the use of these publications are granted subject to the following terms and conditions.

### **Applicability**

These terms and conditions are in addition to any terms of use for the IBM website.

### **Personal use**

You may reproduce these publications for your personal, noncommercial use provided that all proprietary notices are preserved. You may not distribute, display or make derivative work of these publications, or any portion thereof, without the express consent of IBM.

### **Commercial use**

You may reproduce, distribute and display these publications solely within your enterprise provided that all proprietary notices are preserved. You may not make derivative works of these publications, or reproduce, distribute or display these publications or any portion thereof outside your enterprise, without the express consent of IBM.

### **Rights**

Except as expressly granted in this permission, no other permissions, licenses or rights are granted, either express or implied, to the publications or any information, data, software or other intellectual property contained therein.

IBM reserves the right to withdraw the permissions granted herein whenever, in its discretion, the use of the publications is detrimental to its interest or, as determined by IBM, the above instructions are not being properly followed.

You may not download, export or re-export this information except in full compliance with all applicable laws and regulations, including all United States export laws and regulations.

IBM MAKES NO GUARANTEE ABOUT THE CONTENT OF THESE PUBLICATIONS. THE PUBLICATIONS ARE PROVIDED "AS-IS" AND WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE.

---

## **IBM Online Privacy Statement**

IBM Software products, including software as a service solutions, ("Software Offerings") may use cookies or other technologies to collect product usage information, to help improve the end user experience, to tailor interactions with the end user or for other purposes. In many cases no personally identifiable information is collected by the Software Offerings. Some of our Software Offerings can help enable you to

collect personally identifiable information. If this Software Offering uses cookies to collect personally identifiable information, specific information about this offering's use of cookies is set forth below.

This Software Offering does not use cookies or other technologies to collect personally identifiable information.

If the configurations deployed for this Software Offering provide you as customer the ability to collect personally identifiable information from end users via cookies and other technologies, you should seek your own legal advice about any laws applicable to such data collection, including any requirements for notice and consent.

For more information about the use of various technologies, including cookies, for these purposes, See IBM's Privacy Policy at <http://www.ibm.com/privacy> and IBM's Online Privacy Statement at <http://www.ibm.com/privacy/details> the section entitled "Cookies, Web Beacons and Other Technologies" and the "IBM Software Products and Software-as-a-Service Privacy Statement" at <http://www.ibm.com/software/info/product-privacy>.

---

# Glossary

This glossary provides terms and definitions for IBM Spectrum Scale.

The following cross-references are used in this glossary:

- *See* refers you from a nonpreferred term to the preferred term or from an abbreviation to the spelled-out form.
- *See also* refers you to a related or contrasting term.

For other terms and definitions, see the IBM Terminology website ([www.ibm.com/software/globalization/terminology](http://www.ibm.com/software/globalization/terminology)) (opens in new window).

## B

### block utilization

The measurement of the percentage of used subblocks per allocated blocks.

## C

### cluster

A loosely-coupled collection of independent systems (nodes) organized into a network for the purpose of sharing resources and communicating with each other. See also *GPFS cluster*.

### cluster configuration data

The configuration data that is stored on the cluster configuration servers.

### Cluster Export Services (CES) nodes

A subset of nodes configured within a cluster to provide a solution for exporting GPFS file systems by using the Network File System (NFS), Server Message Block (SMB), and Object protocols.

### cluster manager

The node that monitors node status using disk leases, detects failures, drives recovery, and selects file system managers. The cluster manager must be a quorum node. The selection of the cluster manager node favors the quorum-manager node with the lowest node number among the nodes that are operating at that particular time.

**Note:** The cluster manager role is not moved to another node when a node with a lower node number becomes active.

### control data structures

Data structures needed to manage file data and metadata cached in memory. Control data structures include hash tables and link pointers for finding cached data; lock states and tokens to implement distributed locking; and various flags and sequence numbers to keep track of updates to the cached data.

## D

### Data Management Application Program Interface (DMAPI)

The interface defined by the Open Group's XDSM standard as described in the publication *System Management: Data Storage Management (XDSM) API Common Application Environment (CAE) Specification C429*, The Open Group ISBN 1-85912-190-X.

### deadman switch timer

A kernel timer that works on a node that has lost its disk lease and has outstanding I/O requests. This timer ensures that the node cannot complete the outstanding I/O requests (which would risk causing file system corruption), by causing a panic in the kernel.

### dependent fileset

A fileset that shares the inode space of an existing independent fileset.

### disk descriptor

A definition of the type of data that the disk contains and the failure group to which this disk belongs. See also *failure group*.

### disk leasing

A method for controlling access to storage devices from multiple host systems. Any host that wants to access a storage device configured to use disk leasing registers for a lease; in the event of a perceived failure, a host system can deny access,

preventing I/O operations with the storage device until the preempted system has reregistered.

## disposition

The session to which a data management event is delivered. An individual disposition is set for each type of event from each file system.

## domain

A logical grouping of resources in a network for the purpose of common management and administration.

## E

### ECKD™

See *extended count key data (ECKD)*.

### ECKD device

See *extended count key data device (ECKD device)*.

### encryption key

A mathematical value that allows components to verify that they are in communication with the expected server. Encryption keys are based on a public or private key pair that is created during the installation process. See also *file encryption key*, *master encryption key*.

### extended count key data (ECKD)

An extension of the count-key-data (CKD) architecture. It includes additional commands that can be used to improve performance.

### extended count key data device (ECKD device)

A disk storage device that has a data transfer rate faster than some processors can utilize and that is connected to the processor through use of a speed matching buffer. A specialized channel program is needed to communicate with such a device. See also *fixed-block architecture disk device*.

## F

### fallback

Cluster recovery from failover following repair. See also *failover*.

### failover

(1) The assumption of file system duties by another node when a node fails. (2) The process of transferring all control of the ESS to a single cluster in the ESS

when the other clusters in the ESS fails. See also *cluster*. (3) The routing of all transactions to a second controller when the first controller fails. See also *cluster*.

### failure group

A collection of disks that share common access paths or adapter connection, and could all become unavailable through a single hardware failure.

### FEK

See *file encryption key*.

### fileset

A hierarchical grouping of files managed as a unit for balancing workload across a cluster. See also *dependent fileset*, *independent fileset*.

### fileset snapshot

A snapshot of an independent fileset plus all dependent filesets.

### file clone

A writable snapshot of an individual file.

### file encryption key (FEK)

A key used to encrypt sectors of an individual file. See also *encryption key*.

### file-management policy

A set of rules defined in a policy file that GPFS uses to manage file migration and file deletion. See also *policy*.

### file-placement policy

A set of rules defined in a policy file that GPFS uses to manage the initial placement of a newly created file. See also *policy*.

### file system descriptor

A data structure containing key information about a file system. This information includes the disks assigned to the file system (*stripe group*), the current state of the file system, and pointers to key files such as quota files and log files.

### file system descriptor quorum

The number of disks needed in order to write the file system descriptor correctly.

### file system manager

The provider of services for all the nodes using a single file system. A file system manager processes changes to the state or description of the file system, controls the regions of disks that are allocated to each node, and controls token management and quota management.

## **fixed-block architecture disk device (FBA disk device)**

A disk device that stores data in blocks of fixed size. These blocks are addressed by block number relative to the beginning of the file. See also *extended count key data device*.

## **fragment**

The space allocated for an amount of data too small to require a full block. A fragment consists of one or more subblocks.

## **G**

### **global snapshot**

A snapshot of an entire GPFS file system.

### **GPFS cluster**

A cluster of nodes defined as being available for use by GPFS file systems.

### **GPFS portability layer**

The interface module that each installation must build for its specific hardware platform and Linux distribution.

### **GPFS recovery log**

A file that contains a record of metadata activity, and exists for each node of a cluster. In the event of a node failure, the recovery log for the failed node is replayed, restoring the file system to a consistent state and allowing other nodes to continue working.

## **I**

### **ill-placed file**

A file assigned to one storage pool, but having some or all of its data in a different storage pool.

### **ill-replicated file**

A file with contents that are not correctly replicated according to the desired setting for that file. This situation occurs in the interval between a change in the file's replication settings or suspending one of its disks, and the restripe of the file.

### **independent fileset**

A fileset that has its own inode space.

### **indirect block**

A block containing pointers to other blocks.

### **inode**

The internal structure that describes the

individual files in the file system. There is one inode for each file.

### **inode space**

A collection of inode number ranges reserved for an independent fileset, which enables more efficient per-fileset functions.

## **ISKLM**

IBM Security Key Lifecycle Manager. For GPFS encryption, the ISKLM is used as an RKM server to store MEKs.

## **J**

### **journalized file system (JFS)**

A technology designed for high-throughput server environments, which are important for running intranet and other high-performance e-business file servers.

### **junction**

A special directory entry that connects a name in a directory of one fileset to the root directory of another fileset.

## **K**

### **kernel**

The part of an operating system that contains programs for such tasks as input/output, management and control of hardware, and the scheduling of user tasks.

## **M**

### **master encryption key (MEK)**

A key used to encrypt other keys. See also *encryption key*.

### **MEK**

See *master encryption key*.

### **metadata**

Data structures that contain information that is needed to access file data. Metadata includes inodes, indirect blocks, and directories. Metadata is not accessible to user applications.

### **metanode**

The one node per open file that is responsible for maintaining file metadata integrity. In most cases, the node that has had the file open for the longest period of continuous time is the metanode.

### **mirroring**

The process of writing the same data to multiple disks at the same time. The

mirroring of data protects it against data loss within the database or within the recovery log.

### **Microsoft Management Console (MMC)**

A Windows tool that can be used to do basic configuration tasks on an SMB server. These tasks include administrative tasks such as listing or closing the connected users and open files, and creating and manipulating SMB shares.

### **multi-tailed**

A disk connected to multiple nodes.

## **N**

### **namespace**

Space reserved by a file system to contain the names of its objects.

### **Network File System (NFS)**

A protocol, developed by Sun Microsystems, Incorporated, that allows any host in a network to gain access to another host or netgroup and their file directories.

### **Network Shared Disk (NSD)**

A component for cluster-wide disk naming and access.

### **NSD volume ID**

A unique 16 digit hex number that is used to identify and access all NSDs.

**node** An individual operating-system image within a cluster. Depending on the way in which the computer system is partitioned, it may contain one or more nodes.

### **node descriptor**

A definition that indicates how GPFS uses a node. Possible functions include: manager node, client node, quorum node, and nonquorum node.

### **node number**

A number that is generated and maintained by GPFS as the cluster is created, and as nodes are added to or deleted from the cluster.

### **node quorum**

The minimum number of nodes that must be running in order for the daemon to start.

### **node quorum with tiebreaker disks**

A form of quorum that allows GPFS to run with as little as one quorum node

available, as long as there is access to a majority of the quorum disks.

### **non-quorum node**

A node in a cluster that is not counted for the purposes of quorum determination.

## **P**

**policy** A list of file-placement, service-class, and encryption rules that define characteristics and placement of files. Several policies can be defined within the configuration, but only one policy set is active at one time.

### **policy rule**

A programming statement within a policy that defines a specific action to be performed.

**pool** A group of resources with similar characteristics and attributes.

### **portability**

The ability of a programming language to compile successfully on different operating systems without requiring changes to the source code.

### **primary GPFS cluster configuration server**

In a GPFS cluster, the node chosen to maintain the GPFS cluster configuration data.

### **private IP address**

A IP address used to communicate on a private network.

### **public IP address**

A IP address used to communicate on a public network.

## **Q**

### **quorum node**

A node in the cluster that is counted to determine whether a quorum exists.

**quota** The amount of disk space and number of inodes assigned as upper limits for a specified user, group of users, or fileset.

### **quota management**

The allocation of disk blocks to the other nodes writing to the file system, and comparison of the allocated space to quota limits at regular intervals.

## R

### **Redundant Array of Independent Disks (RAID)**

A collection of two or more disk physical drives that present to the host an image of one or more logical disk drives. In the event of a single physical device failure, the data can be read or regenerated from the other disk drives in the array due to data redundancy.

### **recovery**

The process of restoring access to file system data when a failure has occurred. Recovery can involve reconstructing data or providing alternative routing through a different server.

### **remote key management server (RKM server)**

A server that is used to store master encryption keys.

### **replication**

The process of maintaining a defined set of data in more than one location. Replication involves copying designated changes for one location (a source) to another (a target), and synchronizing the data in both locations.

### **RKM server**

See *remote key management server*.

### **rule**

A list of conditions and actions that are triggered when certain conditions are met. Conditions include attributes about an object (file name, type or extension, dates, owner, and groups), the requesting client, and the container name associated with the object.

## S

### **SAN-attached**

Disk that are physically attached to all nodes in the cluster using Serial Storage Architecture (SSA) connections or using Fibre Channel switches.

### **Scale Out Backup and Restore (SOBAR)**

A specialized mechanism for data protection against disaster only for GPFS file systems that are managed by IBM Spectrum Protect Hierarchical Storage Management (HSM).

### **secondary GPFS cluster configuration server**

In a GPFS cluster, the node chosen to maintain the GPFS cluster configuration

data in the event that the primary GPFS cluster configuration server fails or becomes unavailable.

### **Secure Hash Algorithm digest (SHA digest)**

A character string used to identify a GPFS security key.

### **session failure**

The loss of all resources of a data management session due to the failure of the daemon on the session node.

### **session node**

The node on which a data management session was created.

### **Small Computer System Interface (SCSI)**

An ANSI-standard electronic interface that allows personal computers to communicate with peripheral hardware, such as disk drives, tape drives, CD-ROM drives, printers, and scanners faster and more flexibly than previous interfaces.

### **snapshot**

An exact copy of changed data in the active files and directories of a file system or fileset at a single point in time. See also *fileset snapshot*, *global snapshot*.

### **source node**

The node on which a data management event is generated.

### **stand-alone client**

The node in a one-node cluster.

### **storage area network (SAN)**

A dedicated storage network tailored to a specific environment, combining servers, storage products, networking products, software, and services.

### **storage pool**

A grouping of storage space consisting of volumes, logical unit numbers (LUNs), or addresses that share a common set of administrative characteristics.

### **stripe group**

The set of disks comprising the storage assigned to a file system.

### **striping**

A storage process in which information is split into blocks (a fixed amount of data) and the blocks are written to (or read from) a series of disks in parallel.

## **subblock**

The smallest unit of data accessible in an I/O operation, equal to one thirty-second of a data block.

## **system storage pool**

A storage pool containing file system control structures, reserved files, directories, symbolic links, special devices, as well as the metadata associated with regular files, including indirect blocks and extended attributes. The **system storage pool** can also contain user data.

## **T**

### **token management**

A system for controlling file access in which each application performing a read or write operation is granted some form of access to a specific block of file data. Token management provides data consistency and controls conflicts. Token management has two components: the token management server, and the token management function.

### **token management function**

A component of token management that requests tokens from the token management server. The token management function is located on each cluster node.

### **token management server**

A component of token management that controls tokens relating to the operation of the file system. The token management server is located at the file system manager node.

### **transparent cloud tiering (TCT)**

A separately installable add-on feature of IBM Spectrum Scale that provides a native cloud storage tier. It allows data center administrators to free up on-premise storage capacity, by moving out cooler data to the cloud storage, thereby reducing capital and operational expenditures. .

## **twin-tailed**

A disk connected to two nodes.

## **U**

### **user storage pool**

A storage pool containing the blocks of data that make up user files.

## **V**

### **VFS** See *virtual file system*.

### **virtual file system (VFS)**

A remote file system that has been mounted so that it is accessible to the local user.

### **virtual node (vnode)**

The structure that contains information about a file system object in a virtual file system (VFS).

---

# Index

## Special characters

--aix-trace-buffer-size  
  changing 572  
--trace-dispatch  
  changing 572  
--tracedev-buffer-size  
  changing 572  
--tracedev-compression-level  
  changing 572  
--tracedev-overwrite-buffer-size  
  changing 573  
--tracedev-write-mode  
  changing 573  
-traceFileSize  
  changing 572

## A

access control information 715  
  restoring 784  
access control lists  
  creating 504  
  deleting 284  
  displaying 347  
  editing 322  
  getting 350  
access rights  
  locking 631  
  loss of 660  
  restrictions 631  
accessibility features for IBM Spectrum Scale 1123  
acl  
  API  
    GET 946  
    PUT 949  
ACL information 666, 768  
  restoring 784  
  retrieving 715  
AD server  
  querying 32  
adding  
  disks 23  
adding nodes to a GPFS cluster 29  
additional calls, setup of interface for 756  
adminMode attribute 135  
AFM 37, 40, 54, 501  
API  
  DELETE fileset 975  
  DELETE fileset/snapshots 1011  
  DELETE filesystem/snapshots 1037  
  DELETE nfs/exports 1068  
  DELETE smb/shares/shareName 1120  
  GET acl path 946  
  GET ces/addresses 909  
  GET ces/addresses/{cesAddress} 913  
  GET ces/services 916  
  GET ces/services/{service} 919  
  GET cluster 922  
  GET config 926  
  GET disks 954, 959

API (*continued*)  
  GET fileset 978  
  GET fileset/snapshots 1004, 1014  
  GET filesets 963  
  GET filesystem/snapshots 1031, 1040  
  GET filesystems 931, 939  
  GET jobs 1045, 1049  
  GET nfs/exports 1053, 1060  
  GET nodes 1071, 1078  
  GET nodes/nodeName/health/events 1083  
  GET nodes/nodeName/health/states 1087  
  GET nsds 1091, 1097  
  GET owner path 1017  
  GET quotas 996, 1000, 1023  
  GET smb/shares 1100, 1105  
  POST fileset/snapshots 1008  
  POST filesets 970  
  POST filesystem/snapshots 1034  
  POST link fileset 993  
  POST nfs/exports 1057  
  POST quotas 1027  
  POST smb/shares 1110  
  POST unlink fileset 990  
  PUT acl path 949  
  PUT filesets 984  
  PUT nfs/exports 1064  
  PUT owner path 1020  
  PUT smb/shares/shareName 1115  
appendOnly file attribute 122, 392  
application failure 660  
argument  
  buflen 657  
  flags 657  
  hanp 655  
  hlen 655  
  len 657  
  nelem 655, 657  
  nelemp 655  
  off 657  
  sessinfop 655  
asynchronous jobs  
  API  
    GET 1045, 1049  
atime 181, 769, 772, 775, 778, 795, 797, 799, 805  
atimeDeferredSeconds attribute 137  
attribute bit  
  dm\_at\_size 640  
attributes  
  adminMode 135  
  atimeDeferredSeconds 137  
  autoload 137  
  automountDir 137  
  cesSharedRoot 138  
  cipherList 138  
  cnfsGrace 138  
  cnfsMountdPort 139  
  cnfsNFSDprocs 139  
  cnfsReboot 139  
  cnfsSharedRoot 139  
  cnfsVersions 139  
  commandAudit 139

attributes (*continued*)

- configuration 626
- dataDiskWaitTimeForRecovery 139
- dataStructureDump 140
- deadlockBreakupDelay 140
- deadlockDataCollectionDailyLimit 140
- deadlockDataCollectionMinInterval 140
- deadlockDetectionThreshold 140
- deadlockDetectionThresholdForShortWaiters 140
- deadlockOverloadThreshold 140
- debugDataControl 140
- defaultHelperNodes 141
- defaultMountDir 141
- description 633
- disableInodeUpdateOnFdatasync 141
- dmapiEventDataRetry 141
- dmapiEventTimeout 142
- dmapiMountEvent 142
- dmapiMountTimeout 142
- dmapiSessionFailureTimeout 143
- enableIPv6 143
- enforceFilesetQuotaOnRoot 143
- expelDataCollectionDailyLimit 143
- expelDataCollectionMinInterval 144
- extended 633
- failureDetectionTime 144
- fastestPolicyCmpThreshold 144
- fastestPolicyMaxValidPeriod 144
- fastestPolicyMinDiffPercent 144
- fastestPolicyNumReadSamples 144
- fileHeatLossPercent 144
- fileHeatPeriodMinutes 144
- FIPS1402mode 144
- forceLogWriteOnFdatasync 144
- GPFS-specific 638
- lrocData 145
- lrocDataMaxFileSize 145
- lrocDataStubFileSize 145
- lrocDirectories 145
- lrocInodes 145
- maxblocksize 145
- maxDownDisksForRecovery 146
- maxFailedNodesForRecovery 146
- maxFcntlRangesPerFile 146
- maxFilesToCache 146
- maxMBpS 146
- maxStatCache 147
- metadataDiskWaitTimeForRecovery 147
- minDiskWaitTimeForRecovery 147
- mmapRangeLock 148
- nistCompliance 148
- non-opaque 627, 633
- noSpaceEventInterval 148
- nsdBufSpace 149
- nsdRAIDBufferPoolSizePct 149
- nsdRAIDTracks 149
- nsdServerWaitTimeForMount 150
- nsdServerWaitTimeWindowOnMount 150
- numaMemoryInterleave 150
- opaque 627, 633
- pagepool 150
- pagepoolMaxPhysMemPct 150
- prefetchThreads 151
- readReplicaPolicy 151
- release 152
- stripeOnDiskFailure 152
- rpcPerfNumberDayIntervals 152

attributes (*continued*)

- rpcPerfNumberHourIntervals 152
- rpcPerfNumberMinuteIntervals 153
- rpcPerfNumberSecondIntervals 153
- rpcPerfRawExecBufferSize 153
- rpcPerfRawStatBufferSize 153
- sidAutoMapRangeLength 154
- sidAutoMapRangeStart 154
- subnets 154
- systemLogLevel 155
- tiebreakerDisks 155
- uidDomain 156
- unmountOnDiskFail 156
- usePersistentReserve 156
- verbsPorts 157
- verbsRdma 157
- verbsRdmaCm 157
- verbsRdmaRoCEToS 158
- verbsRdmaSend 158
- verbsRdmasPerConnection 158
- verbsRdmasPerNode 158
- verbsSendBufferMemoryMB 158
- worker1Threads 159
- autoload attribute 137
- automated installation toolkit 604
- automatic mount, indicating 254
- automountDir attribute 137

## B

- backing up a file system
  - configuration information 81
- backup server 72
- BigInsights Hadoop distribution
  - mmhadoopctl 353
- block level incremental backups 752
- block size
  - choosing 254
  - effect on maximum mounted file system size 145, 254

## C

- callbacks 10
  - daRebuildFailed 17
  - nsdCksumMismatch 19
  - pdFailed 19
  - pdPathDown 19
  - pdRecovered 19
  - pdReplacePdisk 19
  - postRGRelinquish 20
  - postRGTakeover 20
  - preRGRelinquish 19
  - preRGTakeover 20
  - rgOpenFailed 20
  - rgPanic 21
- ces
  - config 554
  - mmsmb 554
- CES
  - configuration 103, 454
  - mmces 103
  - mmcesdr 113
  - mmnfs 454
  - mmobj 466
  - mmprotocoltrace 497
  - mmuserauth 581

CES (*continued*)  
     protocol tracing 497  
     topic 581  
 CES BLOCK service  
     iSCSI 83  
 cesaddress: GET  
     REST API 832  
 cesaddresses: GET  
     REST API 830  
 cesservice: GET  
     REST API 837  
 ccesservices: GET  
     REST API 834  
 cesSharedRoot attribute 138  
 changing  
     an administration or daemon interface for a node 190  
 attributes  
     adminMode 135  
     atimeDeferredSeconds 137  
     autoload 137  
     automountDir 137  
     cesSharedRoot 138  
     cipherList 138  
     cluster configuration 132  
     cnfsGrace 138  
     cnfsMountdPort 139  
     cnfsNFSDprocs 139  
     cnfsReboot 139  
     cnfsSharedRoot 139  
     cnfsVersions 139  
     dataDiskWaitTimeForRecovery 139  
     dataStructureDump 140  
     deadlockBreakupDelay 140  
     deadlockDataCollectionDailyLimit 140  
     deadlockDataCollectionMinInterval 140  
     deadlockDetectionThreshold 140  
     deadlockDetectionThresholdForShortWaiters 140  
     deadlockOverloadThreshold 140  
     debugDataControl 140  
     defaultHelperNodes 141  
     defaultMountDir 141  
     disableInodeUpdateOnFdatasync 141  
     dmapiDataEventRetry 141  
     dmapiEventTimeout 142  
     dmapiMountEvent 142  
     dmapiMountTimeout 142  
     dmapiSessionFailureTimeout 143  
     enableIPv6 143  
     enforceFilesetQuotaOnRoot 143  
     expelDataCollectionDailyLimit 143  
     expelDataCollectionMinInterval 144  
     failureDetectionTime 144  
     fastestPolicyCmpThreshold 144  
     fastestPolicyMaxValidPeriod 144  
     fastestPolicyMinDiffPercent 144  
     fastestPolicyNumReadSamples 144  
     fileHeatLossPercent 144  
     fileHeatPeriodMinutes 144  
     FIPS1402mode 144  
     forceLogWriteOnFdatasync 144  
     lrocData 145  
     lrocDataMaxFileSize 145  
     lrocDataStubFileSize 145  
     lrocDirectories 145  
     lrocInodes 145  
     maxblocksize 145  
     maxDownDisksForRecovery 146

changing (*continued*)  
     attributes (*continued*)  
         maxFailedNodesForRecovery 146  
         maxFcntlRangesPerFile 146  
         maxFilesToCache 146  
         maxMBPs 146  
         maxStatCache 147  
         metadataDiskWaitTimeForRecovery 147  
         minDiskWaitTimeForRecovery 147  
         mmapRangeLock 148  
         nistCompliance 148  
         noSpaceEventInterval 148  
         nsdBufSpace 149  
         nsdRAIDBufferPoolSizePct 149  
         nsdRAIDTracks 149  
         nsdServerWaitTimeForMount 150  
         nsdServerWaitTimeWindowOnMount 150  
         numaMemoryInterleave 150  
         pagepool 150  
         pagepoolMaxPhysMemPct 150  
         prefetchThreads 151  
         readReplicaPolicy 151  
         release 152  
         stripeOnDiskFailure 152  
         rpcPerfNumberDayIntervals 152  
         rpcPerfNumberHourIntervals 152  
         rpcPerfNumberMinuteIntervals 153  
         rpcPerfNumberSecondIntervals 153  
         rpcPerfRawExecBufferSize 153  
         rpcPerfRawStatBufferSize 153  
         sidAutoMapRangeLength 154  
         sidAutoMapRangeStart 154  
         subnets 154  
         systemLogLevel 155  
         tiebreakerDisks 155  
         uidDomain 156  
         unmountOnDiskFail 156  
         usePersistentReserve 156  
         verbsPorts 157  
         verbsRdma 157  
         verbsRdmaCm 157  
         verbsRdmaRoCEToS 158  
         verbsRdmaSend 158  
         verbsRdmaPerConnection 158  
         verbsRdmaPerNode 158  
         verbsSendBufferMemoryMB 158  
         workerThreads 159  
     disk parameters 161  
     disk states 161  
     fileset attributes 173  
     tracing attributes 571  
     user-defined node classes 195  
 changing Quality of Service for I/O operations (QoS)  
     level 206  
 changing storage pool properties 204  
 cipherList attribute 68, 138  
 cleanup after GPFS interface calls 757  
 Client license 185  
 client node  
     refresh NSD server 464  
 clone, file  
     copy 215, 667  
     decloning 678  
     redirect 215  
     show 215  
     snap 215, 669  
     split 215, 671

clone, file (*continued*)  
  unsnap 673

cluster  
  changing configuration attributes 132  
  changing tracing attributes 571  
  configuration data 329

cluster configuration attributes  
  displaying 400

cluster configuration data 344

cluster configuration server 239, 343

Cluster Export Services  
  config 554  
  configuration 103, 454  
  mmces 103  
  mmcesdr 113  
  mmnfs 454  
  mmobj 466  
  mmprotocoltrace 497  
  mmsmb 554  
  mmuserauth 581  
  protocol tracing 497  
  topic 581

cluster): GET  
  REST API 841

cnfsGrace attribute 138

cnfsMountdPort attribute 139

cnfsNFSDprocs attribute 139

cnfsReboot attribute 139

cnfsSharedRoot attribute 139

cnfsVersions attribute 139

commandAudit attribute 139

commands 1  
  GPFS REST API 829  
  gpfs.snap 6  
  mmaddcallback 10  
  mmaddir 23, 828  
  mmaddnode 29  
  mmadquery 32  
  mmafmconfig 37  
  mmafmctl 40  
  mmafmlocal 54  
  mmapplypolicy 56  
  mmauth 67  
  mmbackup 72  
  mmbackupconfig 81  
  mmblock 83  
  mmbuildgpl 88  
  mmcallhome 90  
  mmces 103  
  mmcesdr 113  
  mmchattr 122  
  mmchcluster 128  
  mmchconfig 132, 635  
  mmchdisk 161  
  mmcheckquota 169  
  mmchfileset 173  
  mmchfs 179, 636  
  mmchlicense 185  
  mmchmgr 188  
  mmchnode 190  
  mmchnodeclass 195  
  mmchnsd 198  
  mmchpolicy 201  
  mmchpool 204  
  mmchqos 206  
  mmclidecode 213  
  mmclone 215

commands (*continued*)  
  mmcloudgateway 218  
  mmccluster 239  
  mmcrfileset 244  
  mmcrfs 250, 636  
  mmcrnodeclass 260  
  mmcrnsd 23, 262, 539, 827, 828  
  mmcrsnapshot 267  
  mmdefquota 272  
  mmdefquotaoff 275  
  mmdefquotao 278  
  mmdefragfs 281  
  mmdelacl 284  
  mmdelcallback 286  
  mmdeldisk 287, 828  
  mmdelfileset 292  
  mmdeflfs 295  
  mmdelnode 297  
  mmdelnodeclass 300  
  mmdelnsd 302  
  mmdelesnapshot 304  
  mmdf 308  
  mmdiag 312  
  mmdsh 319  
  mmeditacl 322  
  mmdefquota 325  
  mmexportfs 329  
  mmfsck 331  
  mmfsctl 343  
  mmgetacl 347  
  mmgetstate 350  
  mmhadoopctl 353  
  mmhealth 355, 378  
  mmimgbackup 367  
  mmimgrestore 371  
  mmimportfs 374  
  mmlinkfileset 390  
  mmllsattr 392  
  mmllscallback 395  
  mmllscluster 397  
  mmllsconfig 400  
  mmllsdisk 402  
  mmllsfileset 406  
  mmllsfs 411  
  mmllslicense 415  
  mmllsmgr 417  
  mmllsmount 419  
  mmllsnodeclass 421  
  mmllsnsd 423  
  mmllspolicy 426  
  mmllspool 428  
  mmllsqos 430  
  mmllsquota 435  
  mmllssnapshot 439, 719, 723  
  mmigratefs 442  
    fastea 442  
  mmmount 444  
  mmnetverify 446  
  mmnfs 454  
  mmnsdiscover 464  
  mmobj 466  
  mmperfmon query 481  
  mmppmon 492  
  mmprotocoltrace 497  
  mmpsnap 501  
  mmputacl 504  
  mmquotaoff 507

commands (*continued*)

- mmquotaon 509
- mmremotecluster 511
- mmremotefs 514
- mmrepquota 517
- mmrestoreconfig 521
- mmrestorefs 525
- mmrestripefile 529
- mmrestripefs 532
- mmrpldisk 539
- mmsdrrestore 546
- mmsetquota 548
- mmshutdown 552
- mmsmb 554
- mmsnapdir 565, 694
- mmstartup 569
- mmtracectl 571
- mmumount 575
- mmunlinkfileset 578
- mmuserauth 581
- mmwinstervctl 602
- spectrumscale 604

config: GET

- REST API 840

configuration attributes

- DMAPI 635
- dmapiEnable 636
- dmapiEventTimeout 633
- NFS (Network File System) 635
- dmapiMountTimeout 631, 636
- dmapiSessionFailureTimeout 636, 659

connector for Hadoop distributions, GPFS

- mmhadoopctl 353

consistency checks 32

contact node 298

creating

- access control lists 504
- file systems 250
- filesets 244

ctime 769, 772, 775, 778

## D

daRebuildFailed callback 17

Data Management API

- failure 660
- restarting 660

data replica 181

data structures

- defined 637
- specific to GPFS implementation 638

data type

- `dm_eventset_t` 638

dataDiskWaitTimeForRecovery attribute 139

dataStructureDump attribute 140

deadlockBreakupDelay attribute 140

deadlockDataCollectionDailyLimit attribute 140

deadlockDataCollectionMinInterval attribute 140

deadlockDetectionThreshold attribute 140

deadlockDetectionThresholdForShortWaiters attribute 140

deadlockOverloadThreshold attribute 140

debugDataControl attribute 140

declone 678

default quotas

- activating 278
- deactivating 275
- editing 272

defaultHelperNodes attribute 141

defaultMountDir attribute 141

definitions

- GPFS-specific DAPI functions 641, 642, 643, 645, 647, 649, 651, 653

DELETE

- fileset/snapshots 1011
- filessets 975
- filesystem/snapshots 1037
- nfs/exports 1068
- smb/shares/shareName 1120

DELETE fileset

- GPFS REST API 868

DELETE snapshot

- GPFS REST API 907

deleting

- disks 287
- file systems 295
- filesets 292
- nodes from a cluster 297
- snapshots 304

deleting links

- snapshots 565

deleting, Network Shared Disks (NSDs) 302

deny-write open lock 179

description

- `dmapiDataEventRetry` 635
- `dmapiFileHandleSize` 635
- `dmapiMountEvent` 636

directives

- subroutine for passing 684

directory

- `/usr/lpp/mmfs/bin` 635
- `/usr/lpp/mmfs/include` 634
- `/usr/lpp/mmfs/lib` 634
- `/usr/lpp/mmfs/samples` 637
- `/var/mmfs/etc` 637

directory entry

- reading 744, 746

disableInodeUpdateOnFdatasync attribute 141

disaster recovery 343

disk access

- path discovery 464

disk descriptor 162

disk parameter

- changing 161

disk state

- changing 161
- suspended 161

disk storage

- pre-allocating 781

disk usage 161, 539

disks

- adding 23, 539
- API
  - GET 954, 959
  - configuration 402
  - deleting 287
  - displaying state 402
  - reducing fragmentation 281
  - replacing 539
- displaying
  - access control lists 347
  - cluster configuration attributes 400
  - disk state 402
  - filesets 406
  - GPFS cluster configuration information 397

displaying (*continued*)  
     NSD belonging to a GPFS cluster 423  
     quotas 435  
     snapshots 439  
 displaying Quality of Service for I/O operations (QoS)  
     settings 430  
 DM application threads 632  
 DM application, role in session failure 628  
 DM\_EVENT\_POSTPERMCHANGE 656  
 DM\_EVENT\_PREPERMCHANGE 656  
 dm\_handle\_to\_snap  
     definitions 642  
 dm\_make\_xhandle  
     definitions 643  
 DM\_NO\_TOKEN 631  
 dm\_remove\_dmatr\_nosync  
     definitions 645  
 dm\_set\_dmatr\_nosync  
     definitions 647  
 dm\_set\_eventlist\_nosync  
     definitions 649  
 dm\_set\_region\_nosync  
     definitions 651  
 dm\_sync\_dmatr\_by\_handle  
     definitions 653  
 DMAPI 182  
     administration 634  
     applications 634  
     compiling on AIX nodes 635  
     configuration attributes 626, 635  
     failure 657, 660  
     features 621  
     files on Linux nodes 635  
     functions 622  
     initializing 637  
     overview 621  
     recovery 657  
     restarting 660  
     restrictions 627  
 DAPI events  
     GPFS-specific 621  
     GPFS-specific attribute events that are not part of the DAPI standard 622  
     implemented in DAPI for GPFS 621  
     optional events not implemented in DAPI for GPFS 622  
 DAPI events, GPFS-specific 655  
 DAPI functions  
     error code  
         EIO 656  
         ENOMEM 656  
         ENOSYS 656  
         ENOTREADY 656  
         EPERM 656  
         ESTALE 656  
 DAPI functions, GPFS-specific 626  
     definitions 641  
     dm\_handle\_to\_snap 642  
     dm\_make\_xhandle 643  
     dm\_remove\_dmatr\_nosync 645  
     dm\_set\_dmatr\_nosync 647  
     dm\_set\_eventlist\_nosync 649  
     dm\_set\_region\_nosync 651  
     dm\_sync\_dmatr\_by\_handle 653  
 DAPI token, description 631  
 dmapiDataEventRetry  
     description 635  
 dmapiDataEventRetry attribute 141  
 dmapiEventTimeout attribute 142  
 dmapiFileSize  
     description 635  
 dmapiMountEvent attribute 142  
     description 636  
 dmapiMountTimeout attribute 142  
 dmapiSessionFailureTimeout attribute 143  
 DODeferred deletions 660  
 dumps, storage of information 140

## E

editing  
     default quotas 272  
 enableIPv6 attribute 143  
 enabling DAPI  
     migrating a file system 636  
     mmchfs command 636  
     mmcrfs command 636  
 enforceFilesetQuotaOnRoot attribute 143  
 environment  
     multiple-node 628, 657  
     single-node 628, 657  
 error code  
     EAGAIN 640  
     EBADF 639, 640, 656  
     EBUSY 631, 634  
     EINVAL 640, 656, 660  
     EIO 630, 637  
     ENOSYS 639  
     ENOTREADY 633, 639, 659  
     EPERM 639, 656  
     ESRCH 640, 656, 659, 660  
 error code, definitions 656  
 events  
     as defined in XDSM standard 621  
     asynchronous 622, 629  
     description 628  
     disposition 628  
     enabled 629  
     GPFS-specific attribute events that are not part of the DAPI standard 622  
     GPFS-specific DAPI events 621, 655  
     implemented  
         data events 622  
         file system administration 621  
         metadata events 622  
         namespace events 622  
         pseudo events 622  
     implemented in DAPI for GPFS 621  
     mount 630  
     not implemented  
         file system administration 622  
         metadata 622  
     optional events not implemented in DAPI for GPFS 622  
     pre-unmount 630  
     preunmount 638  
     reliable DAPI destroy 630  
     source node 657  
     synchronous 629  
     unmount 630, 638  
 events, metadata  
     DM\_EVENT\_POSTPERMCHANGE 656  
     DM\_EVENT\_PREPERMCHANGE 656  
 expelDataCollectionDailyLimit attribute 143  
 expelDataCollectionMinInterval attribute 144

extended ACLs  
   retrieve 687  
   set 689, 691, 739

extended attributes 392

extended file attributes 689, 691, 739  
   retrieve 687  
   set 689, 691, 739

## F

failure  
   dm application 657  
   GPFS daemon 622, 628  
   partial system 658  
   session 628, 629  
   session node 658  
   single-node 658  
   source node 658  
   total system 658

failure group 161, 539

failureDetectionTime attribute 144

fastestPolicyCmpThreshold attribute 144

fastestPolicyMaxValidPeriod attribute 144

fastestPolicyMinDiffPercent attribute 144

fastestPolicyNumReadSamples attribute 144

field  
   dt\_change 638  
   dt\_ctime 638  
   dt\_dtime 638  
   dt\_nevents 655  
   ev\_nodeid 638  
   me\_handle2 631  
   me\_mode 631, 638, 655  
   me\_name1 631  
   me\_roothandle 631  
   ne\_mode 638  
   rg\_opaque 638  
   uio\_resid 657

file  
   /etc/filesystems 631  
   access control information 711, 715, 784  
   ACL information 711, 715, 784  
   block level incremental read 752  
   dmapi\_types.h 634  
   dmapi.exp export 634  
   dmapi.h 634  
   dmapicalls 635, 639  
   extended attributes 687, 689, 691, 739  
   reading 742

file access pattern information 684

file attribute  
   extended 727, 729  
   querying 392

file attributes  
   appendOnly 122, 392

file clone  
   copy 215, 667  
   decloning 678  
   redirect 215  
   show 215  
   snap 215, 669  
   split 215, 671  
   unsnap 673

file descriptor  
   closing 724  
   opening 735, 737

file handle  
   error code 656

file status information 696, 698, 795, 797, 799, 805  
   gpfs\_stat\_inode\_with\_xattrs() 801  
   gpfs\_stat\_inode\_with\_xattrs64() 803

file system descriptor quorum 344

file system handle 631  
   usage of 654

file system manager  
   changing nodes 188  
   displaying current 417

file system name 700, 703

file system snapshot  
   API  
     DELETE 1037  
     GET 1031, 1040  
     POST 1034

file system snapshot handle 694

file system space  
   querying 308

file system]: GET  
   REST API 876

file systems  
   adding disks 23  
   API  
     GET 931, 939  
   backing up 72  
   block size 250  
   change manager node 188  
   changing attributes 179  
   changing attributes for files 122  
   checking 331, 374  
   control request 343  
   creating 250  
   creating snapshot 267  
   deleting 295  
   deleting disks 287  
   displaying attributes 411  
   displaying format version 411  
   exporting 329

file system manager  
   displaying 417

format version 179

formatting 251

handle 694

importing 374

inconsistencies 331

links to snapshots 565

listing mounted 419

migrating 179, 442

mounted file system sizes 145, 254

mounting 444

moving to another cluster 179, 329

mtime value 180

querying space 308

quotas 278

rebalancing 532

reducing fragmentation 281

remote 67, 511, 514

repairing 331

restoring configuration information 521

restoring with snapshots 525

stripe 26

restriping 532

unmounting 552, 575

file systems]: GET  
   REST API 870

fileHeatLossPercent attribute 144  
 fileHeatPeriodMinutes attribute 144  
 files  
     orphaned 331  
     rebalancing 529  
     restriping 529  
 files, memory mapped 634  
 files, required 634  
 fileset  
     API  
         GET 963, 978  
         POST 970, 990, 993  
 fileset quota 273, 517  
 fileset snapshot  
     API  
         DELETE 1011  
         GET 1004, 1014  
         POST 1008  
 filesets  
     API  
         DELETE 975  
         PUT 984  
     changing attributes 173  
     creating 244  
     deleting 292  
     displaying 406  
     ID 731  
     linking 390  
     name 731  
     restoring with snapshots 525  
     unlinking 578  
 FIPS1402mode attribute 144  
 flag  
     DM\_LOCAL\_MOUNT 630, 631, 638  
     DM\_MOUNT\_LOCAL 655  
     DM\_MOUNT\_REMOTE 655  
     DM\_REMOTE\_MOUNT 631, 638  
     DM\_RR\_WAIT 640  
     DM\_UNMOUNT\_FORCE 631  
 FlashCopy image 343  
 forceLogWriteOnFdatasync attribute 144  
 FPO license 185  
 full backup 758, 760, 769, 772, 775, 778  
 function  
     dm\_create\_session 655  
     dm\_downgrade\_right 639, 657  
     dm\_find\_eventmsg 657  
     dm\_get\_alloc\_info 657  
     dm\_get\_bulkall 639, 640, 654, 655, 657  
     dm\_get\_bulkattr 639, 640, 654, 657  
     dm\_get\_config 626  
     dm\_get\_config\_events 626, 655  
     dm\_get\_dirattrs 657  
     dm\_get\_eventlist 639, 640, 654, 655  
     dm\_get\_events 657  
     dm\_get\_fileattr 638, 655  
     dm\_get\_mount\_info 640  
     dm\_get\_mountinfo 631, 638, 639, 654, 655, 657  
     dm\_get\_region 657  
     dm\_getall\_disp 654, 657  
     dm\_getall\_dmatr 657  
     dm\_getall\_sessions 657  
     dm\_getall\_tokens 657, 659  
     dm\_handle\_hash 655  
     dm\_handle\_is\_valid 655  
     dm\_handle\_to\_fhandle 655  
     dm\_handle\_to\_fsid 655

function (*continued*)  
     dm\_handle\_to\_igen 655  
     dm\_handle\_to\_ino 655  
     dm\_handle\_to\_path 657  
     dm\_handle\_to\_snap 655  
     dm\_init\_attrloc 655  
     dm\_init\_service 657  
     dm\_make\_fhandle 654  
     dm\_make\_handle 654  
     dm\_make\_xhandle 654  
     dm\_mount\_event 631  
     dm\_move\_event 640, 657  
     dm\_probe\_hole 654, 657  
     dm\_punch\_hole 633, 634, 640, 654, 657  
     dm\_query\_right 639  
     dm\_query\_session 654, 655  
     dm\_read\_invis 654, 657  
     dm\_release\_right 639  
     dm\_request\_right 639, 640  
     dm\_respond\_event 657  
     dm\_send\_msg 654  
     dm\_set\_disp 639, 640, 654, 655  
     dm\_set\_eventlist 639, 640, 654, 655  
     dm\_set\_file\_attr 640  
     dm\_set\_return\_on\_destroy 639, 640, 654  
     dm\_sync\_by\_handle 657  
     dm\_upgrade\_right 639, 657  
     dm\_write\_invis 633, 654, 657  
 functions  
     implemented 623, 625  
     mandatory 623  
     not implemented 626  
     optional 625, 626  
     restrictions 639  
 functions, GPFS-specific DAPI 626  
     definitions 641  
     dm\_handle\_to\_snap 642  
     dm\_make\_xhandle 643  
     dm\_remove\_dmatr\_nosync 645  
     dm\_set\_dmatr\_nosync 647  
     dm\_set\_eventlist\_nosync 649  
     dm\_set\_region\_nosync 651  
     dm\_sync\_dmatr\_by\_handle 653

## G

gathering data to solve GPFS problems 6  
 genkey 68  
 GET  
     acl path 946  
     ces/addresses 909  
     ces/addresses/{cesAddress} 913  
     ces/services 916  
     ces/services/{service} 919  
     cluster 922  
     config 926  
     disks 954, 959  
     fileset 978  
     fileset/snapshots 1004, 1014  
     filesets 963  
     filesystem/snapshots 1031, 1040  
     filesystems 931, 939  
     jobs 1045, 1049  
     nfs/exports 1053, 1060  
     nodes 1071, 1078  
     nodes/nodeName/health/events 1083  
     nodes/nodeName/health/states 1087

GET (*continued*)  
 nsds 1091, 1097  
 owner path 1017  
 quotas 996, 1000, 1023  
 smb/shares 1100, 1105

GET fileset  
 GPFS REST API 853

GET filesets  
 GPFS REST API 847

GPFS  
 access rights  
     loss of 660  
 Data Management API 621  
 DM application failure 660  
 DMAPI 621  
     failure 657  
     recovery 657  
 enhancements 637  
 failure  
     single-node 658  
 file system 621  
 implementation 621, 637  
 installation toolkit 604  
 license designation 185, 415  
 licensing 185, 415  
 mmhealth command 355  
 mmprotocoltrace command 497  
 programming interfaces 675, 676, 678, 680, 682, 684, 687, 689, 691, 693, 694, 695, 696, 698, 700, 701, 703, 705, 707, 709, 711, 713, 715, 717, 720, 724, 726, 727, 729, 731, 733, 735, 737, 739, 742, 744, 746, 748, 750, 752, 755, 756, 757, 758, 760, 762, 764, 766, 768, 769, 772, 775, 778, 781, 784, 786, 789, 791, 793, 795, 797, 799, 801, 803, 805, 807, 808, 811, 812, 814, 816, 817, 818, 819, 821, 823  
 session  
     failure 659  
     recovery 659  
 stopping 552

GPFS cluster  
 creating 239

GPFS cluster configuration data 329

GPFS cluster configuration information  
 displaying 397

GPFS cluster configuration server  
 changing 128  
 primary 129  
 secondary 129

GPFS cluster configuration servers  
 choosing 239

GPFS cluster data 262

GPFS commands 1

GPFS configuration data 826, 828

GPFS connector for Hadoop distributions  
 mmhadoopctl 353

GPFS daemon  
 starting 569  
 stopping 552

GPFS daemon failure 628

GPFS daemon status 350

GPFS directory entry 680, 682

GPFS enhancements  
 implementation of 638

GPFS file system snapshot handle 700, 701, 703, 705, 707, 709, 713  
 free 693

GPFS portability layer 88

GPFS programming interfaces 663

GPFS REST API  
 commands 829  
 DELETE fileset 868  
 DELETE snapshot 907  
 GET fileset 853  
 GET filesets 847  
 POST filesets 859  
 programming 829  
 PUT filesets 864  
 quotas 892

GPFS subroutines 663

GPFS user exits 825

gpfs\_acl\_t 666

GPFS\_ATTRFLAG\_DEFAULT  
 gpfs\_fgetattr() 687  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691

GPFS\_ATTRFLAG\_FINALIZE\_ATTRS  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_IGNORE\_PLACEMENT  
 gpfs\_igetattrsx() 729

GPFS\_ATTRFLAG\_IGNORE\_POOL  
 GPFS\_ATTRFLAG\_FINALIZE\_ATTRS  
     gpfs\_fgetattr() 687  
 GPFS\_ATTRFLAG\_INCL\_DMAPI  
     gpfs\_fgetattr() 687  
 GPFS\_ATTRFLAG\_INCL\_ENCR  
     gpfs\_fgetattr() 687

GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT  
 gpfs\_fgetattr() 687

GPFS\_ATTRFLAG\_SKIP\_CLONE  
 gpfs\_fgetattr() 687

GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE  
 gpfs\_fgetattr() 687

GPFS\_ATTRFLAG\_USE\_POLICY  
 gpfs\_fgetattr() 687

gpfs\_fgetattrs() 687  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_INCL\_DMAPI  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_INCL\_ENCR  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_NO\_PLACEMENT  
 gpfs\_fgetattr() 687  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739

GPFS\_ATTRFLAG\_SKIP\_CLONE  
 gpfs\_fputattr() 689  
 gpfs\_fputattrswithpathname() 691

**GPFS\_ATTRFLAG\_SKIP\_CLONE** (*continued*)  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739  
**GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE**  
 gpfs\_fputattrs() 689  
 gpfs\_fputattrswith pathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739  
**GPFS\_ATTRFLAG\_USE\_POLICY**  
 gpfs\_fputattrs() 689  
 gpfs\_fputattrswith pathname() 691  
 gpfs\_igetattrsx() 729  
 gpfs\_iputattrsx() 739  
 gpfs\_clone\_copy() 667  
 gpfs\_clone\_snap() 669  
 gpfs\_clone\_split() 671  
 gpfs\_clone\_unsnap() 673  
 gpfs\_close\_inodescan() 675  
 gpfs\_cmp\_fssnapid() 676  
 gpfs\_declone() 678  
 gpfs\_direntx\_t 680  
 gpfs\_direntx64\_t 682  
 gpfs\_fcntl() 684  
 gpfs\_fgetattrs() 687  
     GPFS\_ATTRFLAG\_DEFAULT 687  
     GPFS\_ATTRFLAG\_FINALIZE\_ATTRS 687  
     GPFS\_ATTRFLAG\_IGNORE\_POOL 687  
     GPFS\_ATTRFLAG\_INCL\_DMAPI 687  
     GPFS\_ATTRFLAG\_INCL\_ENCR 687  
     GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT 687  
     GPFS\_ATTRFLAG\_NO\_PLACEMENT 687  
     GPFS\_ATTRFLAG\_SKIP\_CLONE 687  
     GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE 687  
     GPFS\_ATTRFLAG\_USE\_POLICY 687  
 gpfs\_fputattrs() 689  
     GPFS\_ATTRFLAG\_DEFAULT 689  
     GPFS\_ATTRFLAG\_FINALIZE\_ATTRS 689  
     GPFS\_ATTRFLAG\_IGNORE\_POOL 689  
     GPFS\_ATTRFLAG\_INCL\_DMAPI 689  
     GPFS\_ATTRFLAG\_INCL\_ENCR 689  
     GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT 689  
     GPFS\_ATTRFLAG\_NO\_PLACEMENT 689  
     GPFS\_ATTRFLAG\_SKIP\_CLONE 689  
     GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE 689  
     GPFS\_ATTRFLAG\_USE\_POLICY 689  
 gpfs\_fputattrswith pathname() 691  
     GPFS\_ATTRFLAG\_DEFAULT 691  
     GPFS\_ATTRFLAG\_FINALIZE\_ATTRS 691  
     GPFS\_ATTRFLAG\_IGNORE\_POOL 691  
     GPFS\_ATTRFLAG\_INCL\_DMAPI 691  
     GPFS\_ATTRFLAG\_INCL\_ENCR 691  
     GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT 691  
     GPFS\_ATTRFLAG\_NO\_PLACEMENT 691  
     GPFS\_ATTRFLAG\_SKIP\_CLONE 691  
     GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE 691  
     GPFS\_ATTRFLAG\_USE\_POLICY 691  
 gpfs\_free\_fssnaphandle() 693  
 gpfs\_fssnap\_handle\_t 694  
 gpfs\_fssnap\_id\_t 695  
 gpfs\_fstat\_x() 698  
 gpfs\_fstat() 696  
 gpfs\_get\_fsname\_from\_fssnaphandle() 700  
 gpfs\_get\_fssnaphandle\_by\_fssnapid() 701  
 gpfs\_get\_fssnaphandle\_by\_name() 703  
 gpfs\_get\_fssnaphandle\_by\_path() 705  
 gpfs\_get\_fssnapid\_from\_fssnaphandle() 707  
 gpfs\_get.pathname\_from\_fssnaphandle() 709  
 gpfs\_get\_snapdirname() 711  
 gpfs\_get\_snapname\_from\_fssnaphandle() 713  
 gpfs\_getacl() 715  
 gpfs\_iattr\_t 717  
 gpfs\_iattr64\_t 720  
 gpfs\_iclose() 724  
 gpfs\_ifile\_t 726  
 gpfs\_igetattrs() 727  
 gpfs\_igetattrsx() 729  
     GPFS\_ATTRFLAG\_FINALIZE\_ATTRS 729  
     GPFS\_ATTRFLAG\_IGNORE\_PLACEMENT 729  
     GPFS\_ATTRFLAG\_INCL\_DMAPI 729  
     GPFS\_ATTRFLAG\_INCL\_ENCR 729  
     GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT 729  
     GPFS\_ATTRFLAG\_NO\_PLACEMENT 729  
     GPFS\_ATTRFLAG\_SKIP\_CLONE 729  
     GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE 729  
     GPFS\_ATTRFLAG\_USE\_POLICY 729  
 gpfs\_getfilesetname() 731  
 gpfs\_igetstoragepool() 733  
 gpfs\_iopen() 735  
 gpfs\_iopen64() 737  
 gpfs\_iputattrsx() 739  
     GPFS\_ATTRFLAG\_FINALIZE\_ATTRS 739  
     GPFS\_ATTRFLAG\_IGNORE\_POOL 739  
     GPFS\_ATTRFLAG\_INCL\_DMAPI 739  
     GPFS\_ATTRFLAG\_INCL\_ENCR 739  
     GPFS\_ATTRFLAG\_MODIFY\_CLONEPARENT 739  
     GPFS\_ATTRFLAG\_NO\_PLACEMENT 739  
     GPFS\_ATTRFLAG\_SKIP\_CLONE 739  
     GPFS\_ATTRFLAG\_SKIP\_IMMUTABLE 739  
     GPFS\_ATTRFLAG\_USE\_POLICY 739  
 gpfs\_read() 742  
 gpfs\_readdir() 744  
 gpfs\_readdir64() 746  
 gpfs\_ireadlink() 748  
 gpfs\_ireadlink64() 750  
 gpfs\_ireadx() 752  
 gpfs\_iscan\_t 755  
 gpfs\_lib\_init() 756  
 gpfs\_lib\_term() 757  
 gpfs\_next\_inode\_with\_xatrrs() 762  
 gpfs\_next\_inode\_with\_xatrrs64() 764  
 gpfs\_next\_inode() 758  
 gpfs\_next\_inode64() 760  
 gpfs\_next\_xattr() 766  
 gpfs\_opaque\_acl\_t 768  
 gpfs\_open\_inodescan\_with\_xatrrs() 775  
 gpfs\_open\_inodescan\_with\_xatrrs64() 778  
 gpfs\_open\_inodescan() 769  
 gpfs\_open\_inodescan64() 772  
 gpfs\_prealloc() 781  
 gpfs\_putacl() 784  
 gpfs\_quotactl() 786  
 gpfs\_quotaInfo\_t 789  
 gpfs\_seek\_inode() 791  
 gpfs\_seek\_inode64() 793  
 gpfs\_stat\_inode\_with\_xatrrs() 801  
 gpfs\_stat\_inode\_with\_xatrrs64() 803  
 gpfs\_stat\_inode() 797  
 gpfs\_stat\_inode64() 799  
 gpfs\_stat\_x() 805  
 gpfs\_stat() 795  
 GPFS-specific DMAPI events 621, 655  
 GPFS-specific DMAPI functions 626  
     definitions 641  
     dm\_handle\_to\_snap 642

GPFS-specific DMAPI functions (*continued*)  
  dm\_make\_xhandle 643  
  dm\_remove\_dmatr\_nosync 645  
  dm\_set\_dmatr\_nosync 647  
  dm\_set\_eventlist\_nosync 649  
  dm\_set\_region\_nosync 651  
    dm\_sync\_dmatr\_by\_handle 653  
gpfs.snap command 6  
gpfsFcntlHeader\_t 807  
gpfsGetDataBlkDiskIdx\_t 808  
gpfsGetFilesetName\_t 811  
gpfsGetReplication\_t 812  
gpfsGetXAttr\_t 814  
gpfsGetSnapshotName\_t 816  
gpfsGetStoragePool\_t 817  
gpfsListXAttr\_t 818  
gpfsRestripeData\_t 819  
gpfsSetReplication\_t 821  
gpfsSetStoragePool\_t 823  
grace period  
  changing 325, 548  
  setting 325, 548  
group quota 273, 275, 278, 326, 435, 509, 517

## H

Hadoop distributions, GPFS connector for  
  mmhadoopctl 353

hints

  subroutine for passing 684

hole 752

## I

I/O caching policy  
  changing 122

IBM Spectrum Protect

  using the mmbackup command 72

IBM Spectrum Scale 367, 371, 621, 622, 623, 625, 626, 627, 628, 630, 631, 632, 633, 634, 635, 636, 637, 638, 639, 641, 642, 643, 645, 647, 649, 651, 653, 655, 656

  access rights

    loss of 660

  commands

    mmcloudgateway 218

  Data Management API 621

  DM application failure 660

  DMAPI 621, 628

    failure 657

    recovery 657

  DMAPI functions 654

  DODeferred deletions 660

  failure

    single-node 658

  installation toolkit 604

  license designation 185, 415

  licensing 185, 415

  programming interfaces 675, 676, 678, 680, 682, 684, 687,

    689, 691, 693, 694, 695, 696, 698, 700, 701, 703, 705, 707,

    709, 711, 713, 715, 717, 720, 724, 726, 727, 729, 731, 733,

    735, 737, 739, 742, 744, 746, 748, 750, 752, 755, 756, 757,

    758, 760, 762, 764, 766, 768, 769, 772, 775, 778, 781, 784,

    786, 789, 791, 793, 795, 797, 799, 801, 803, 805, 807, 808,

    811, 812, 814, 816, 817, 818, 819, 821, 823

  recovery

    synchronous event 659

IBM Spectrum Scale (*continued*)

  session

    failure 659

    recovery 659

IBM Spectrum Scale information units ix

IBM Spectrum Scale management API

  DELETE fileset 975

  DELETE fileset/snapshots 1011

  DELETE filesystem/snapshots 1037

  DELETE nfs/exports 1068

  DELETE smb/shares/shareName 1120

  GET acl path 946

  GET ces/addresses 909

  GET ces/addresses/{cesAddress} 913

  GET ces/services 916

  GET ces/services/{service} 919

  GET cluster 922

  GET config 926

  GET disks 954, 959

  GET fileset 978

  GET fileset/snapshots 1004, 1014

  GET filesets 963

  GET filesystem/snapshots 1031, 1040

  GET filesystems 931, 939

  GET jobs 1045, 1049

  GET nfs/exports 1053, 1060

  GET nodes 1071, 1078

  GET nodes/nodeName/health/events 1083

  GET nodes/nodeName/health/states 1087

  GET nsds 1091, 1097

  GET owner path 1017

  GET quotas 996, 1000, 1023

  GET smb/shares 1100, 1105

  POST fileset/snapshots 1008

  POST filesets 970

  POST filesystem/snapshots 1034

  POST link fileset 993

  POST nfs/exports 1057

  POST quotas 1027

  POST smb/shares 1110

  POST unlink fileset 990

  PUT acl path 949

  PUT filesets 984

  PUT nfs/exports 1064

  PUT owner path 1020

  PUT smb/shares/shareName 1115

IBM Spectrum Scale REST API 829, 892, 896, 905

IBM Spectrum Scale user exits 825

in-doubt value 169, 436, 518

incremental backup 758, 760, 769, 772, 775, 778

info: GET

  REST API 882, 1043

inode

  attributes 717, 720

inode file handle 724, 726

inode number 735, 737, 748, 750, 758, 760, 762, 764, 766, 769, 772, 775, 778, 791, 793

inode scan 752, 758, 760, 762, 764, 766, 791, 793

  closing 675

  opening 769, 772, 775, 778

inode scan handle 675, 755

installation 604

  installation requirements 634

  interface calls, cleanup after 757

  interface for additional calls, setup of 756

iscan handle 675

iSCSI

  mmblock 83

## K

kernel memory 122

## L

license 185

linking

  filesets 390

links to snapshots

  creating 565

  deleting 565

listing

  snapshots 439

  user-defined callbacks 395

listing Quality of Service for I/O operations (QoS)

  settings 430

lrocData attribute 145

lrocDataMaxFileSize attribute 145

lrocDataStubFileSize attribute 145

lrocDirectories attribute 145

lrocInodes attribute 145

## M

macro

  DM\_TOKEN\_EQ (x,y) 639

  DM\_TOKEN\_GE (x,y) 639

  DM\_TOKEN\_GT (x,y) 639

  DM\_TOKEN\_LE (x,y) 639

  DM\_TOKEN\_LT (x,y) 639

  DM\_TOKEN\_NE (x,y) 639

  DMEV\_ADD(eset1, eset2) 638

  DMEV\_ALL(eset) 638

  DMEV\_ISALL(eset) 638

  DMEV\_ISDISJ(eset1, eset2) 638

  DMEV\_ISEQ(eset1, eset2) 638

  DMEV\_ISSUB(eset2) 639

  DMEV\_ISZERO(eset) 638

  DMEV\_NORM(eset) 639

  DMEV\_Rem(eset1, eset2) 638

  DMEV\_RES(eset1, eset2) 638

macros, GPFS 638

macros, XDSM standard 638

management API

  DELETE fileset 975

  DELETE fileset/snapshots 1011

  DELETE filesystem/snapshots 1037

  DELETE nfs/exports 1068

  DELETE smb/shares/shareName 1120

  GET acl path 946

  GET ces/addresses 909

  GET ces/addresses/{cesAddress} 913

  GET ces/services 916

  GET ces/services/{service} 919

  GET cluster 922

  GET config 926

  GET disks 954, 959

  GET fileset 978

  GET fileset/snapshots 1004, 1014

  GET filesets 963

  GET filesystem/snapshots 1031, 1040

  GET filesystems 931, 939

management API (*continued*)

  GET jobs 1045, 1049

  GET nfs/exports 1053, 1060

  GET nodes 1071, 1078

  GET nsds 1091, 1097

  GET owner path 1017

  GET quotas 996, 1000, 1023

  GET smb/shares 1100, 1105

  POST fileset/snapshots 1008

  POST filesets 970

  POST filesystem/snapshots 1034

  POST link fileset 993

  POST nfs/exports 1057

  POST quotas 1027

  POST smb/shares 1110

  POST unlink fileset 990

  PUT acl path 949

  PUT filesets 984

  PUT nfs/exports 1064

  PUT owner path 1020

  PUT smb/shares/shareName 1115

maxblocksize attribute 145

maxDownDisksForRecovery attribute 146

maxFailedNodesForRecovery attribute 146

maxFcntlRangesPerFile attribute 146

maxFilesToCache attribute 146

maximum number of files

  changing 179

  displaying 411

maxMBps attribute 146

maxStatCache attribute 147

memory mapped files 634

metadata 123

metadata events

  DM\_EVENT\_POSTPERMCHANGE 656

  DM\_EVENT\_PREPERMCHANGE 656

metadata replica 180

metadataDiskWaitTimeForRecovery attribute 147

minDiskWaitTimeForRecovery attribute 147

mmaddcallback 10

mmadddisk 23, 828

mmaddnode 29

mmadquery 32

mmafmconfig 37

mmafmctl 40

mmafmlocal 54

mmaapplypolicy 56

mmapRangeLock attribute 148

mmauth 67

mmbackup 72

mmbackupconfig 81

mmblock 83

mmbuildgpl 88

mmcallehome 90

mmces 103

mmcesdr 113

mmchattr 122

mmchcluster 128

mmchconfig 132

mmchdisk 161

mmcheckquota 169

mmchfileset 173

mmchfs 179

mmchlicense 185

mmchmgr 188

mmchnode 190  
 mmchnodeclass 195  
 mmchnsd 198  
 mmchpolicy 201  
 mmchpool 204  
 mmchqos 206  
 mmclidecode 213  
 mmclone 215  
 mmccluster 239  
 mmcrfleset 244  
 mmcrfs 250  
 mmcrnodeclass 260  
 mmcrnsd 262, 539, 827, 828  
 mmcrsnapshot 267  
 mmdefedquota 272  
 mmdefquotaoff 275  
 mmdefquotaon 278  
 mmdefragfs 281  
 mmdelacl 284  
 mmdelcallback 286  
 mmdeldisk 287, 828  
 mmdelfleset 292  
 mmdeps 295  
 mmdeinode 297  
 mmdeinodeclass 300  
 mmdeinsn 302  
 mmdeleset 304  
 mmdf 308  
 mmdiag 312  
 mmsh 319  
 mmeditacl 322  
 mmmedquota 325  
 mmexportfs 329  
 MMFS\_FSSTRUCT 331  
 MMFS\_SYSTEM\_UNMOUNT 332  
 mmfsck 331  
 mmfsctl 343  
 mmgetacl 347  
 mmgetstate 350  
 mmhadoopctl 353  
 mmhealth 355, 378  
 mmimgbackup 367  
 mmimgrestore 371  
 mmimportfs 374  
 mmunlinkfleset 390  
 mmllsattr 392  
 mmllscallback 395  
 mmllscluster 397  
 mmllsconfig 400  
 mmllsdisk 402  
 mmllsfileset 406  
 mmllsfs 411  
 mmllslicense 415  
 mmllsmgr 417  
 mmllsmount 419  
 mmllsnodaclass 421  
 mmllsnsd 423  
 mmllspolicy 426  
 mmllspool 428  
 mmllsqos 430  
 mmllsquota 435  
 mmllssnapshot 439, 719, 723  
 mmigratefs 442  
 mmmount 444  
 mmnetverify 446  
 mmnfs 454  
 mmnsddiscover 464  
 mmobj 466  
 mmperfmon query 481  
 mmpmon 492  
 mmprotocoltrace 497  
 mmppsnap 501  
 mmputacl 504  
 mmquotaoff 507  
 mmquotaon 509  
 mmremotecluster 511  
 mmremotesfs 514  
 mmrepquota 517  
 mmrestoreconfig 521  
 mmrestores 525  
 mmrestripefile 529  
 mmrestripefs 532  
 mmrpldisk 539  
 mmsdrrestore 546  
 mmsetquota 548  
 mmshutdown 552  
 mmsmb 554  
 mmsnapdir 565, 694  
 mmstartup 569  
 mmtracectl 571  
 mmumount 575  
 mmunlinkfleset 578  
 mmuserauth 581  
 mmwinserv service  
     managing 602  
 mmwinservcctl 602  
 monitoring  
     performance 492  
 mount point directory 251  
 mounting a file system 179  
 mtime 254, 769, 772, 775, 778, 795, 797, 799, 805  
 multi-region object deployment  
     mmobj command 466  
 multiple sessions 632  
 multiple-node environment 628, 657  
     model for DAPI 657

## N

Network Shared Disks (NSDs) 827  
     changing configuration attributes 198  
     creating 262  
     displaying 423  
 Network Shared Disks (NSDs), deleting 302  
 network verification tool 446  
 NFS (Network File System) 633  
 nfs export  
     API  
         DELETE 1068  
         GET 1053, 1060  
         POST 1057  
         PUT 1064  
 NFS V4 179, 254  
 NFS V4 ACL 180, 284, 322, 323, 347, 348, 504  
 nistCompliance attribute 148  
 node classes, user-defined  
     changing 195  
     creating 260  
     deleting 300  
     listing 421  
 node descriptor 29, 239  
 node designation 29, 239  
 node failure detection 144  
 node id 638

node}: GET  
     REST API 889  
 nodes  
     adding to a cluster 29  
     API  
         GET 1071, 1078  
         GET nodes/nodeName/health/events 1083  
         GET nodes/nodeName/health/states 1087  
     deleting from a cluster 297  
 nodes}: GET  
     REST API 885  
 noSpaceEventInterval attribute 148  
 nsd  
     API  
         GET 954, 959, 1091, 1097  
 NSD path 464  
 NSD server 374  
 NSD server list  
     changing 198  
 NSD server nodes  
     changing 198  
     choosing 262  
 NSD volume ID 262, 302  
 nsdBufSpace attribute 149  
 nsdCksumMismatch callback 19  
 nsdRAIDBufferSizePct attribute 149  
 nsdRAIDTracks attribute 149  
 nsdServerWaitTimeForMount attribute 150  
 nsdServerWaitTimeWindowOnMount attribute 150  
 numaMemoryInterleave attribute 150

## O

owner  
     API  
         GET 1017  
         PUT 1020

## P

pagepool attribute 150  
 pagepoolMaxPhysMemPct attribute 150  
 parallel environment, DM applications 632  
 pdFailed callback 19  
 pdPathDown callback 19  
 pdRecovered callback 19  
 pdReplacePdisk callback 19  
 peer recovery cluster 343  
 Peer-to-Peer Remote Copy (PPRC) 343  
 performance 629  
 performance, monitoring 492  
 policy  
     applying 56  
 pool  
     displaying 428  
 POST  
     fileset/snapshots 1008  
     filesets 970  
     filesystem/snapshots 1034  
     link fileset 993  
     nfs/exports 1057  
     quotas 896, 1027  
     smb/shares 1110  
     unlink fileset 990  
 POST filesets  
     GPFS REST API 859

postRGRelinquish callback 20  
 postRGTakeover callback 20  
 prefetchThreads attribute 151  
 preRGRelinquish callback 19  
 preRGTakeover callback 20  
 primary GPFS cluster configuration server 241  
 problem determination information, placement of 140  
 Programming  
     GPFS REST API 829  
 public/private key pair 67  
 PUT  
     acl path 949  
     filesets 984  
     nfs/exports 1064  
     owner path 1020  
     smb/shares/shareName 1115  
 PUT filesets  
     GPFS REST API 864

## Q

Quality of Service for I/O operations (QoS) level  
     changing 206  
 Quality of Service for I/O operations (QoS) settings  
     listing 430  
 quorum 344  
 quorum node 239, 298, 350  
 quota 633  
 quota files  
     replacing 169  
 quota information 789  
 quotas  
     activating 509  
     API  
         GET 996, 1000, 1023  
         POST 1027  
     changing 325, 548, 786  
     checking 169  
     creating reports 517  
     deactivating 507  
     displaying 435  
     GPFS REST API 892  
     POST 896  
     REST API 896  
     setting 325, 548

## R

RAID stripe size 250  
 readReplicaPolicy attribute 151  
 rebalancing a file 529  
 rebalancing a file system 532  
 recovery  
     DODeferred deletions 660  
     mount event 660  
     synchronous event 659  
     unmount event 660  
 recovery groups  
     stanza files 376  
 refresh NSD server  
     mmnsddiscover 464  
 registering user event commands 10  
 release attribute 152  
 reliable DMAPI destroy events 630  
 remote copy command  
     changing 128

remote copy command (*continued*)
   
     choosing 239

remote file systems 67

remote shell command
   
     changing 128

replacing disks 539

replicated cluster 828

replication 180
   
     querying 392

replication attributes
   
     changing 122

replication factor 122

replication, strict 255

REST API
   
     cesaddress: GET 832

    cesaddresses: GET 830

    cesservice: GET 837

    cesservices: GET 834

    cluster: GET 841

    config: GET 840

    file system: GET 876

    file systems: GET 870

    info: GET 882, 1043

    node: GET 889

    nodes: GET 885

    quotas 896

    snapshots: GET 899

    snapshots: POST 905

    snapshots/{snapshot}: GET 902

restoring configuration information 521

restoring NSD path
   
     mmnsddiscover 464

restrictions
   
     functions 639

stripeOnDiskFailure attribute 152

restriping a file 529

restriping a file system 532

rgOpenFailed callback 20

rgPanic callback 21

root credentials 639

rpcPerfNumberDayIntervals attribute 152

rpcPerfNumberHourIntervals attribute 152

rpcPerfNumberMinuteIntervals attribute 153

rpcPerfNumberSecondIntervals attribute 153

rpcPerfRawExecBufferSize attribute 153

rpcPerfRawStatBufferSize attribute 153

**S**

secondary GPFS cluster configuration server 241

semantic changes
   
     for the GPFS implementation 654

Server license 185

server node
   
     restoring NSD path 464

server node, NSD
   
     choosing 262

session
   
     failure 629, 655, 659

    recovery 659

session node 628, 654, 658

session, assuming a 628, 655

sessions
   
     description 628

    failure 628

    information string, changing 655

sessions (*continued*)
   
     maximum per node 628, 639

    state of 628

setup of interface for additional calls 756

shell script
   
     gpfsready 637

sidAutoMapRangeLength attribute 154

sidAutoMapRangeStart attribute 154

single-node 658

single-node environment 628, 657

SMB
   
     export 481

smb shares
   
     API

        DELETE 1120

        GET 1100, 1105

        POST 1110

        PUT 1115

    snapshot directory 711

    snapshot handle 694, 700, 701, 703, 705, 707, 709, 713

    free 693

    snapshot ID 694, 695, 701, 707

        comparing 676

        internal 719, 723

    snapshot name 703, 713

    snapshots

        coexistence 627

        creating 267

        deleting 267, 304

        directory 267

        displaying 439

        fileset 267

        global 267

        listing 439

        restoring a file system 525

        restoring a fileset 525

    snapshots: POST

        REST API 905

    snapshots/{snapshot}: GET

        REST API 902

    snapshots): GET

        REST API 899

    sort-command parameter of mmapplypolicy command 64

source node 628, 658

sparse file 752

spectrumscale 604

spectrumscale installation toolkit 604

stanza files
   
     recovery group 376

starting GPFS 569

storage
   
     pre-allocating 781

storage pool 428

storage pool properties
   
     changing 204

storage pools
   
     ID 733

    name 733

strict replication 180, 255

structure
   
     dm\_eventmsg 638

    dm\_mount\_event 631, 638, 655

    dm\_namesp\_event 638

    dm\_region\_t 638

    dm\_stat 638

    dm\_stat\_t 655

    dm\_vardata\_t 638

structure (*continued*)  
  `ui0` 657  
structures  
  `gpfs_acl_t` 666  
  `gpfs_direntx_t` 680  
  `gpfs_direntx64_t` 682  
  `gpfs_fssnap_handle_t` 694  
  `gpfs_fssnap_id_t` 695  
  `gpfs_iattr_t` 717  
  `gpfs_iattr64_t` 720  
  `gpfs_ifile_t` 726, 727, 729  
  `gpfs_iscan_t` 755  
  `gpfs_opaque_acl_t` 768  
  `gpfs_quotaInfo_t` 789  
  `gpfsFcntlHeader_t` 807  
  `gpfsGetDataBlkDiskIdx_t` 808  
  `gpfsGetFilesetName_t` 811  
  `gpfsGetReplication_t` 812  
  `gpfsGetSetXAttr_t` 814  
  `gpfsGetSnapshotName_t` 816  
  `gpfsGetStoragePool_t` 817  
  `gpfsListXAttr_t` 818  
  `gpfsRestripeData_t` 819  
  `gpfsSetReplication_t` 821  
  `gpfsSetStoragePool_t` 823  
subnets attribute 154  
subroutine  
  `gpfs_close_inodescan()` 675  
  `gpfs_cmp_fssnapid()` 676  
subroutines  
  `gpfs_clone_copy()` 667  
  `gpfs_clone_snap()` 669  
  `gpfs_clone_split()` 671  
  `gpfs_clone_unsnap()` 673  
  `gpfs_declone()` 678  
  `gpfs_fcntl()` 684  
  `gpfs_fgetatrrs()` 687  
  `gpfs_fputatrrs()` 689  
  `gpfs_fputattrswithpathname()` 691  
  `gpfs_free_fssnaphandle()` 693  
  `gpfs_fstat_x()` 698  
  `gpfs_fstat()` 696  
  `gpfs_get_fnname_from_fssnaphandle()` 700  
  `gpfs_get_fssnaphandle_by_fssnapid()` 701  
  `gpfs_get_fssnaphandle_by_name()` 703  
  `gpfs_get_fssnaphandle_by_path()` 705  
  `gpfs_get_fssnapid_from_fssnaphandle()` 707  
  `gpfs_get.pathname_from_fssnaphandle()` 709  
  `gpfs_get_snapdirname()` 711  
  `gpfs_get_snapname_from_fssnaphandle()` 713  
  `gpfs_getacl()` 715  
  `gpfs_iclose()` 724  
  `gpfs_igetattrs()` 727  
  `gpfs_igetatrrsx()` 729  
  `gpfs_igetfilesetname()` 731  
  `gpfs_igetstoragepool()` 733  
  `gpfs_iopen()` 735  
  `gpfs_iopen64()` 737  
  `gpfs_iputatrrsx()` 739  
  `gpfs_iread()` 742  
  `gpfs_ireaddir()` 744  
  `gpfs_ireaddir64()` 746  
  `gpfs_ireadlink()` 748  
  `gpfs_ireadlink64()` 750  
  `gpfs_ireadx()` 752  
  `gpfs_lib_init()` 756  
  `gpfs_lib_term()` 757  
subroutines (*continued*)  
  `gpfs_next_inode_with_xattrs()` 762  
  `gpfs_next_inode_with_xattrs64()` 764  
  `gpfs_next_inode()` 758  
  `gpfs_next_inode64()` 760  
  `gpfs_next_xattr()` 766  
  `gpfs_open_inodescan_with_xatrrs()` 775  
  `gpfs_open_inodescan_with_xatrrs64()` 778  
  `gpfs_open_inodescan()` 769  
  `gpfs_open_inodescan64()` 772  
  `gpfs_malloc()` 781  
  `gpfs_putacl()` 784  
  `gpfs_quotactl()` 786  
  `gpfs_seek_inode()` 791  
  `gpfs_seek_inode64()` 793  
  `gpfs_stat_inode_with_xatrrs()` 801  
  `gpfs_stat_inode_with_xatrrs64()` 803  
  `gpfs_stat_inode()` 797  
  `gpfs_stat_inode64()` 799  
  `gpfs_stat_x()` 805  
  `gpfs_stat()` 795  
symbolic link  
  reading 748, 750  
syncFSconfig 343  
system snapshots 6  
systemLogLevel attribute 155

## T

`tiebreakerDisks` attribute 155  
`timeout` period 552  
`token, usage` 631  
tokens  
  input arguments 656  
trace-recycle  
  changing 572  
tracing attributes, changing 571  
traditional ACL 180, 322, 323, 347, 348  
traditional ACLs  
  NFS V4 ACL 254  
  Windows 254  
transparent cloud tiering  
  commands  
    `mmcloudgateway` 218

## U

`UID domain` 241  
`uidDomain` attribute 156  
unified file and object access  
  `mmobj` command 466  
unlinking  
  filesets 578  
`unmountOnDiskFail` attribute 156  
usage restrictions 639  
`usePersistentReserve` attribute 156  
user event commands, registering 10  
user exit  
  GPFS 828  
  IBM Spectrum Scale 828  
user exits 825  
  GPFS 827  
  IBM Spectrum Scale 827  
`mmsdrbackup` 826  
`nsddevices` 827  
`syncfsconfig` 828

user quota 273, 275, 278, 326, 435, 509, 517  
user space buffer 122  
user-defined callbacks  
  deleting 286  
  listing 395  
user-defined node classes  
  changing 195  
  creating 260  
  deleting 300  
  listing 421  
using the gpfs.snap command  
  gathering data 6

## V

verbsPorts attribute 157  
verbsRdma attribute 157  
verbsRdmaCm attribute 157  
verbsRdmaRoCEToS attribute 158  
verbsRdmaSend attribute 158  
verbsRdmasesPerConnection attribute 158  
verbsRdmasesPerNode attribute 158  
verbsSendBufferMemoryMB attribute 158  
verification tool, network 446

## W

worker1Threads attribute 159

## X

XDSM standard 626, 628, 657, 658







Product Number: 5725-Q01  
5641-GPF  
5725-S28  
5765-ESS

Printed in USA

SA23-1456-16

